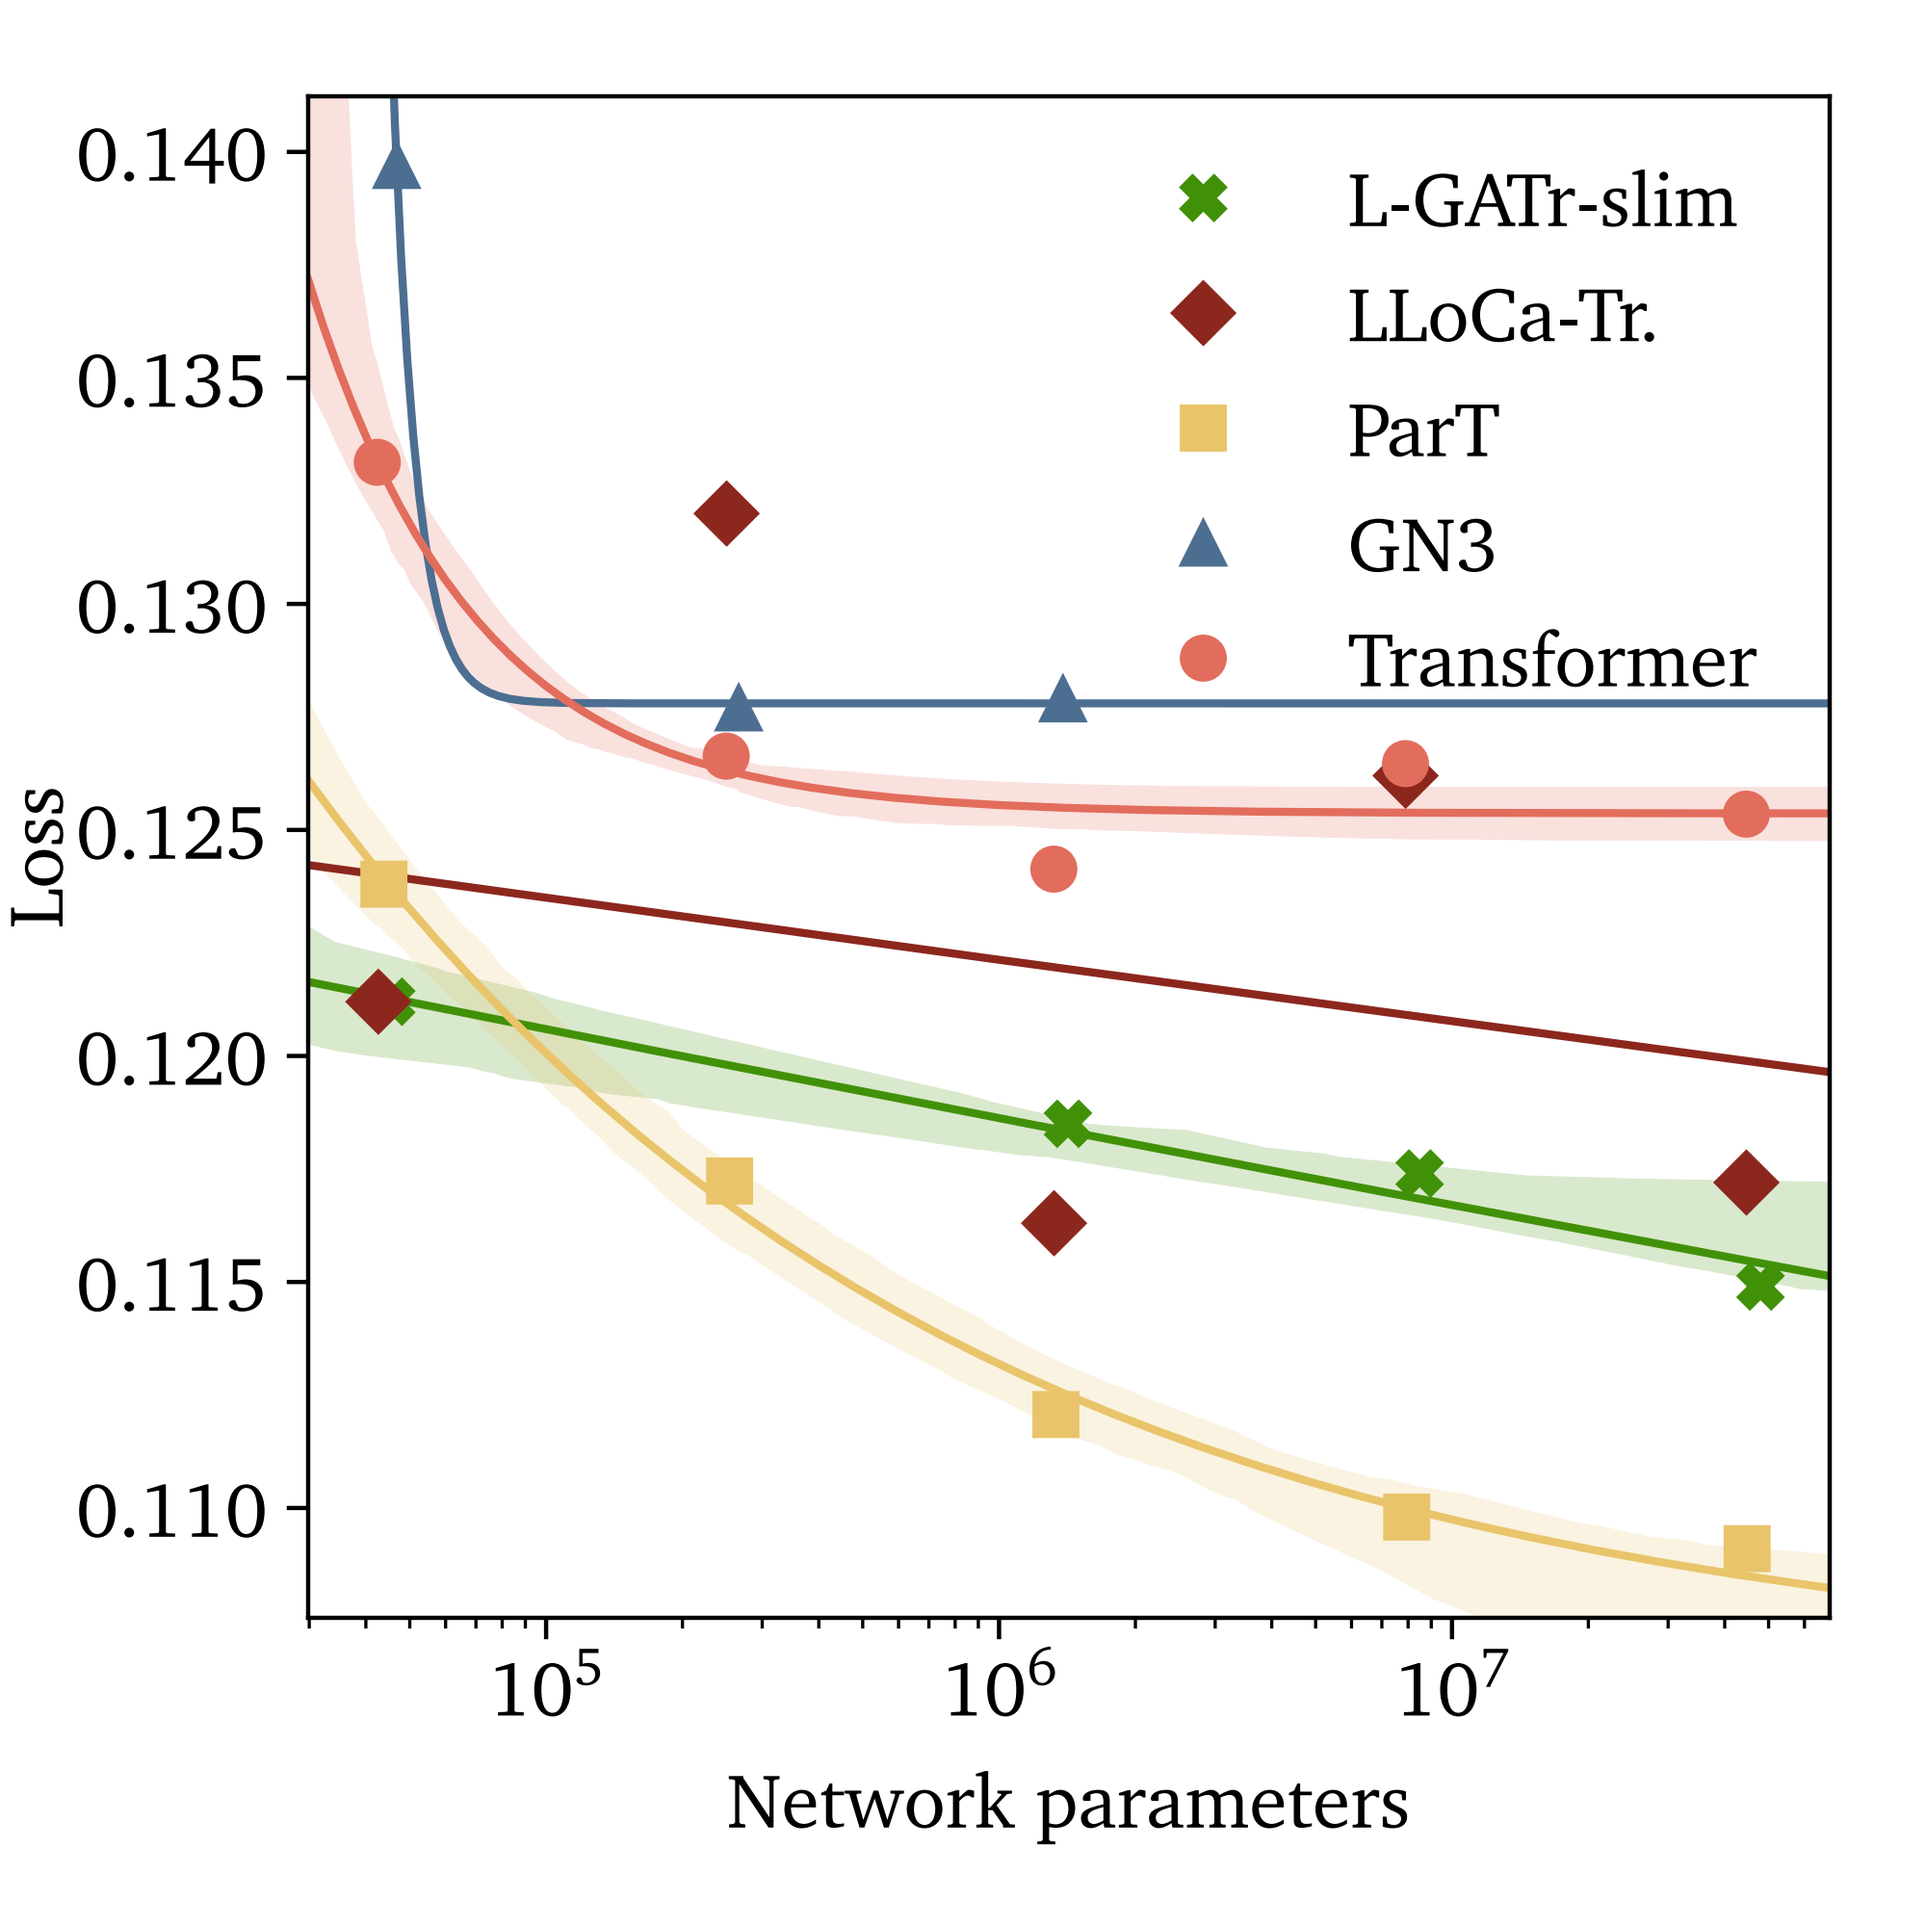
0.140
0.135
0.130
0.125
0.120
0.115
0.110
Loss
10^5
10^6
10^7
Network parameters
L-GATr-slim
LLoCa-Tr.
ParT
GN3
Transformer

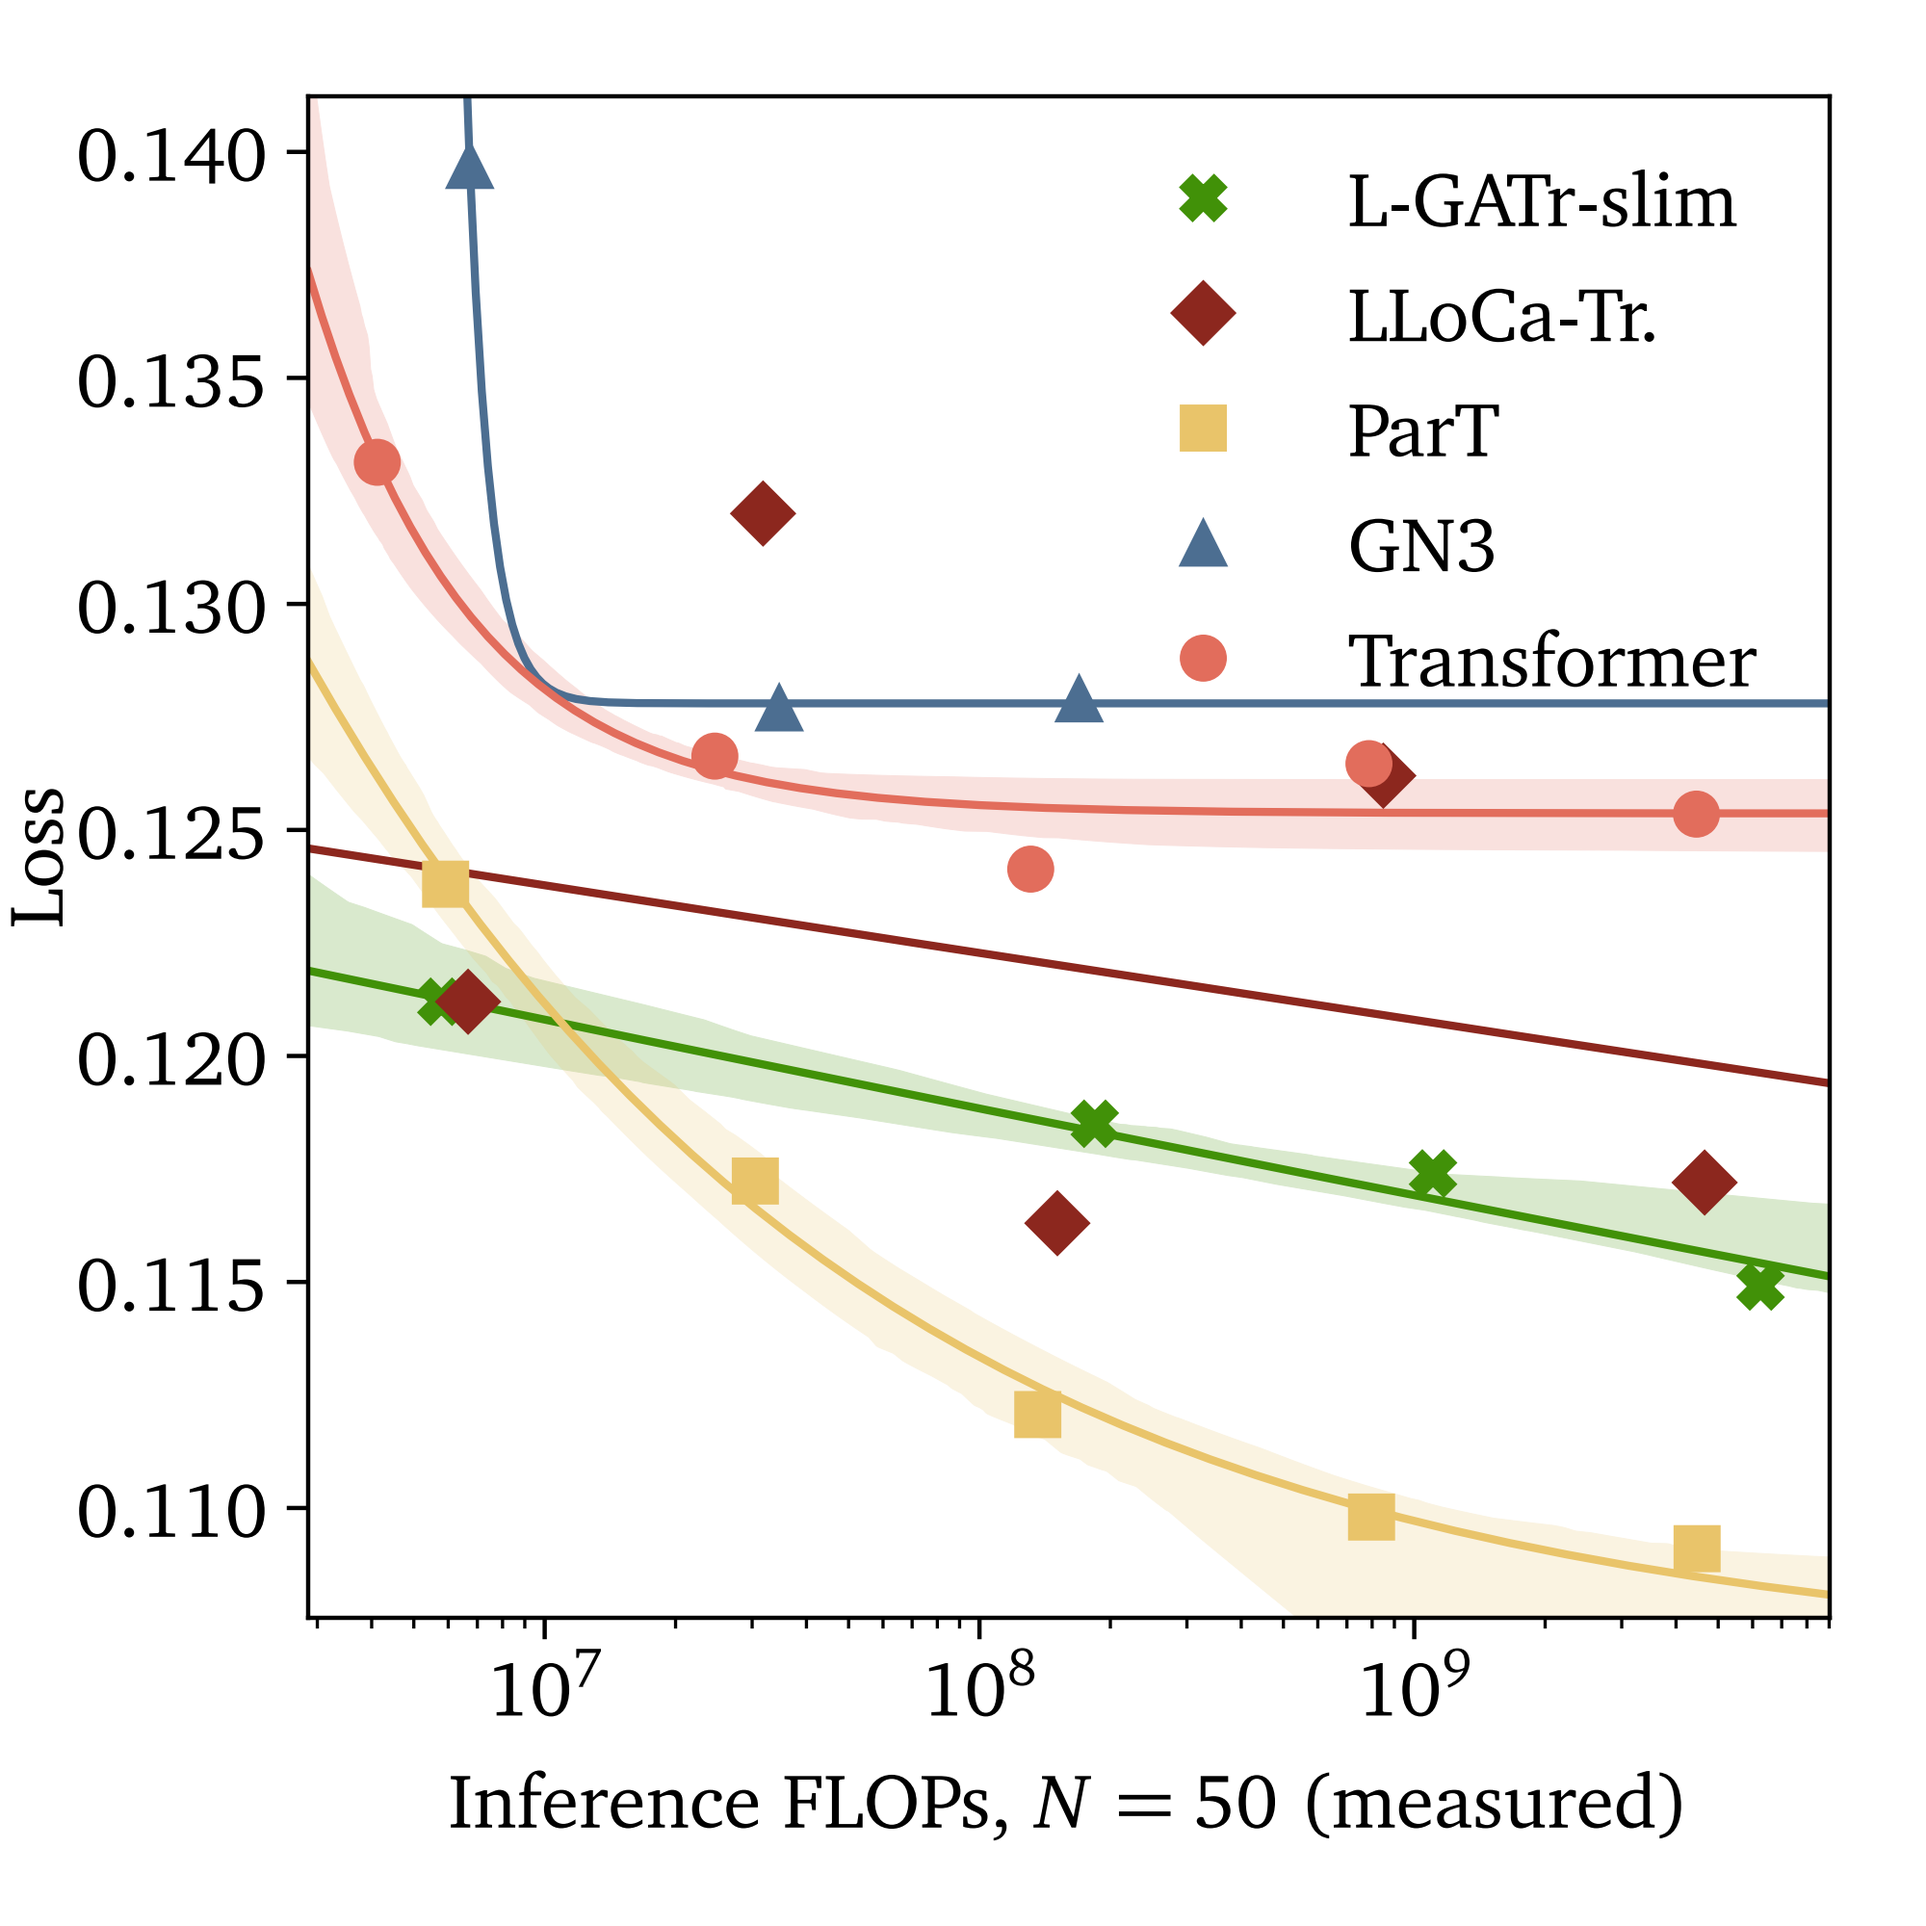
0.140
0.135
0.130
0.125
0.120
0.115
0.110
Loss
L-GATr-slim
LLoCa-Tr.
ParT
GN3
Transformer
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (measured)

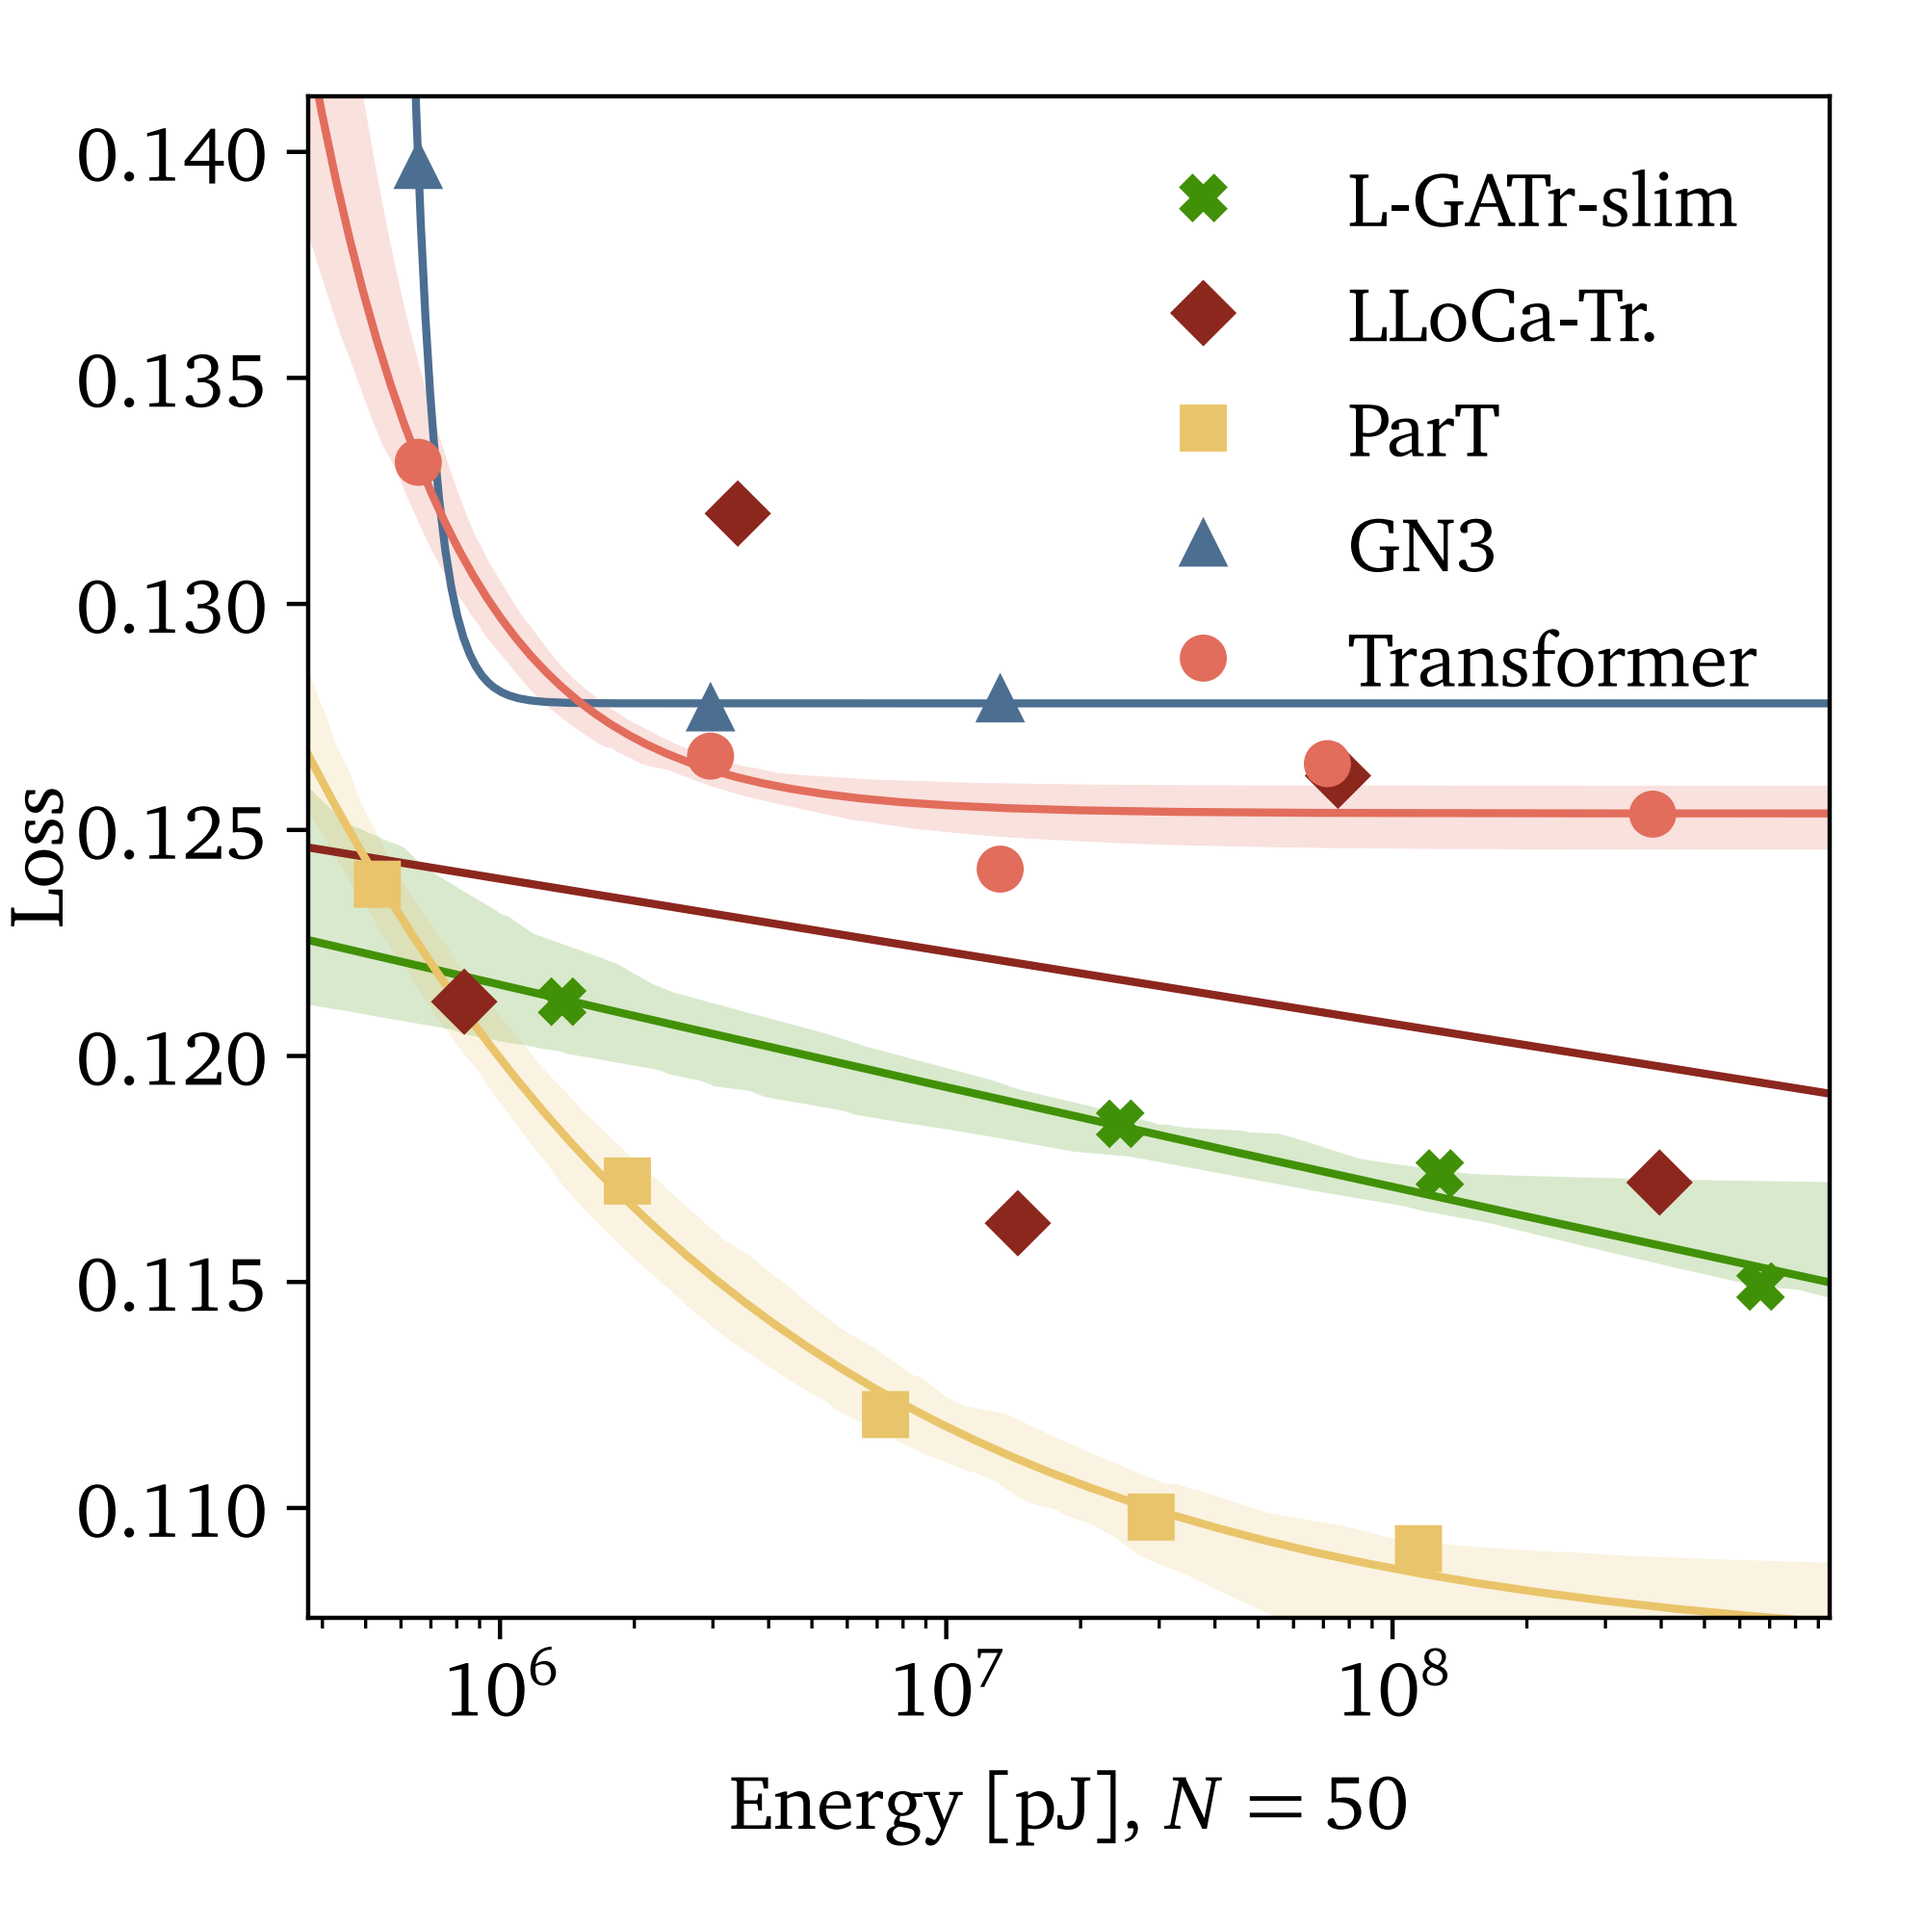
0.140
0.135
0.130
0.125
0.120
0.115
0.110
Loss
L-GATr-slim
LLoCa-Tr.
ParT
GN3
Transformer
10^6
10^7
10^8
Energy [pJ], $N = 50$

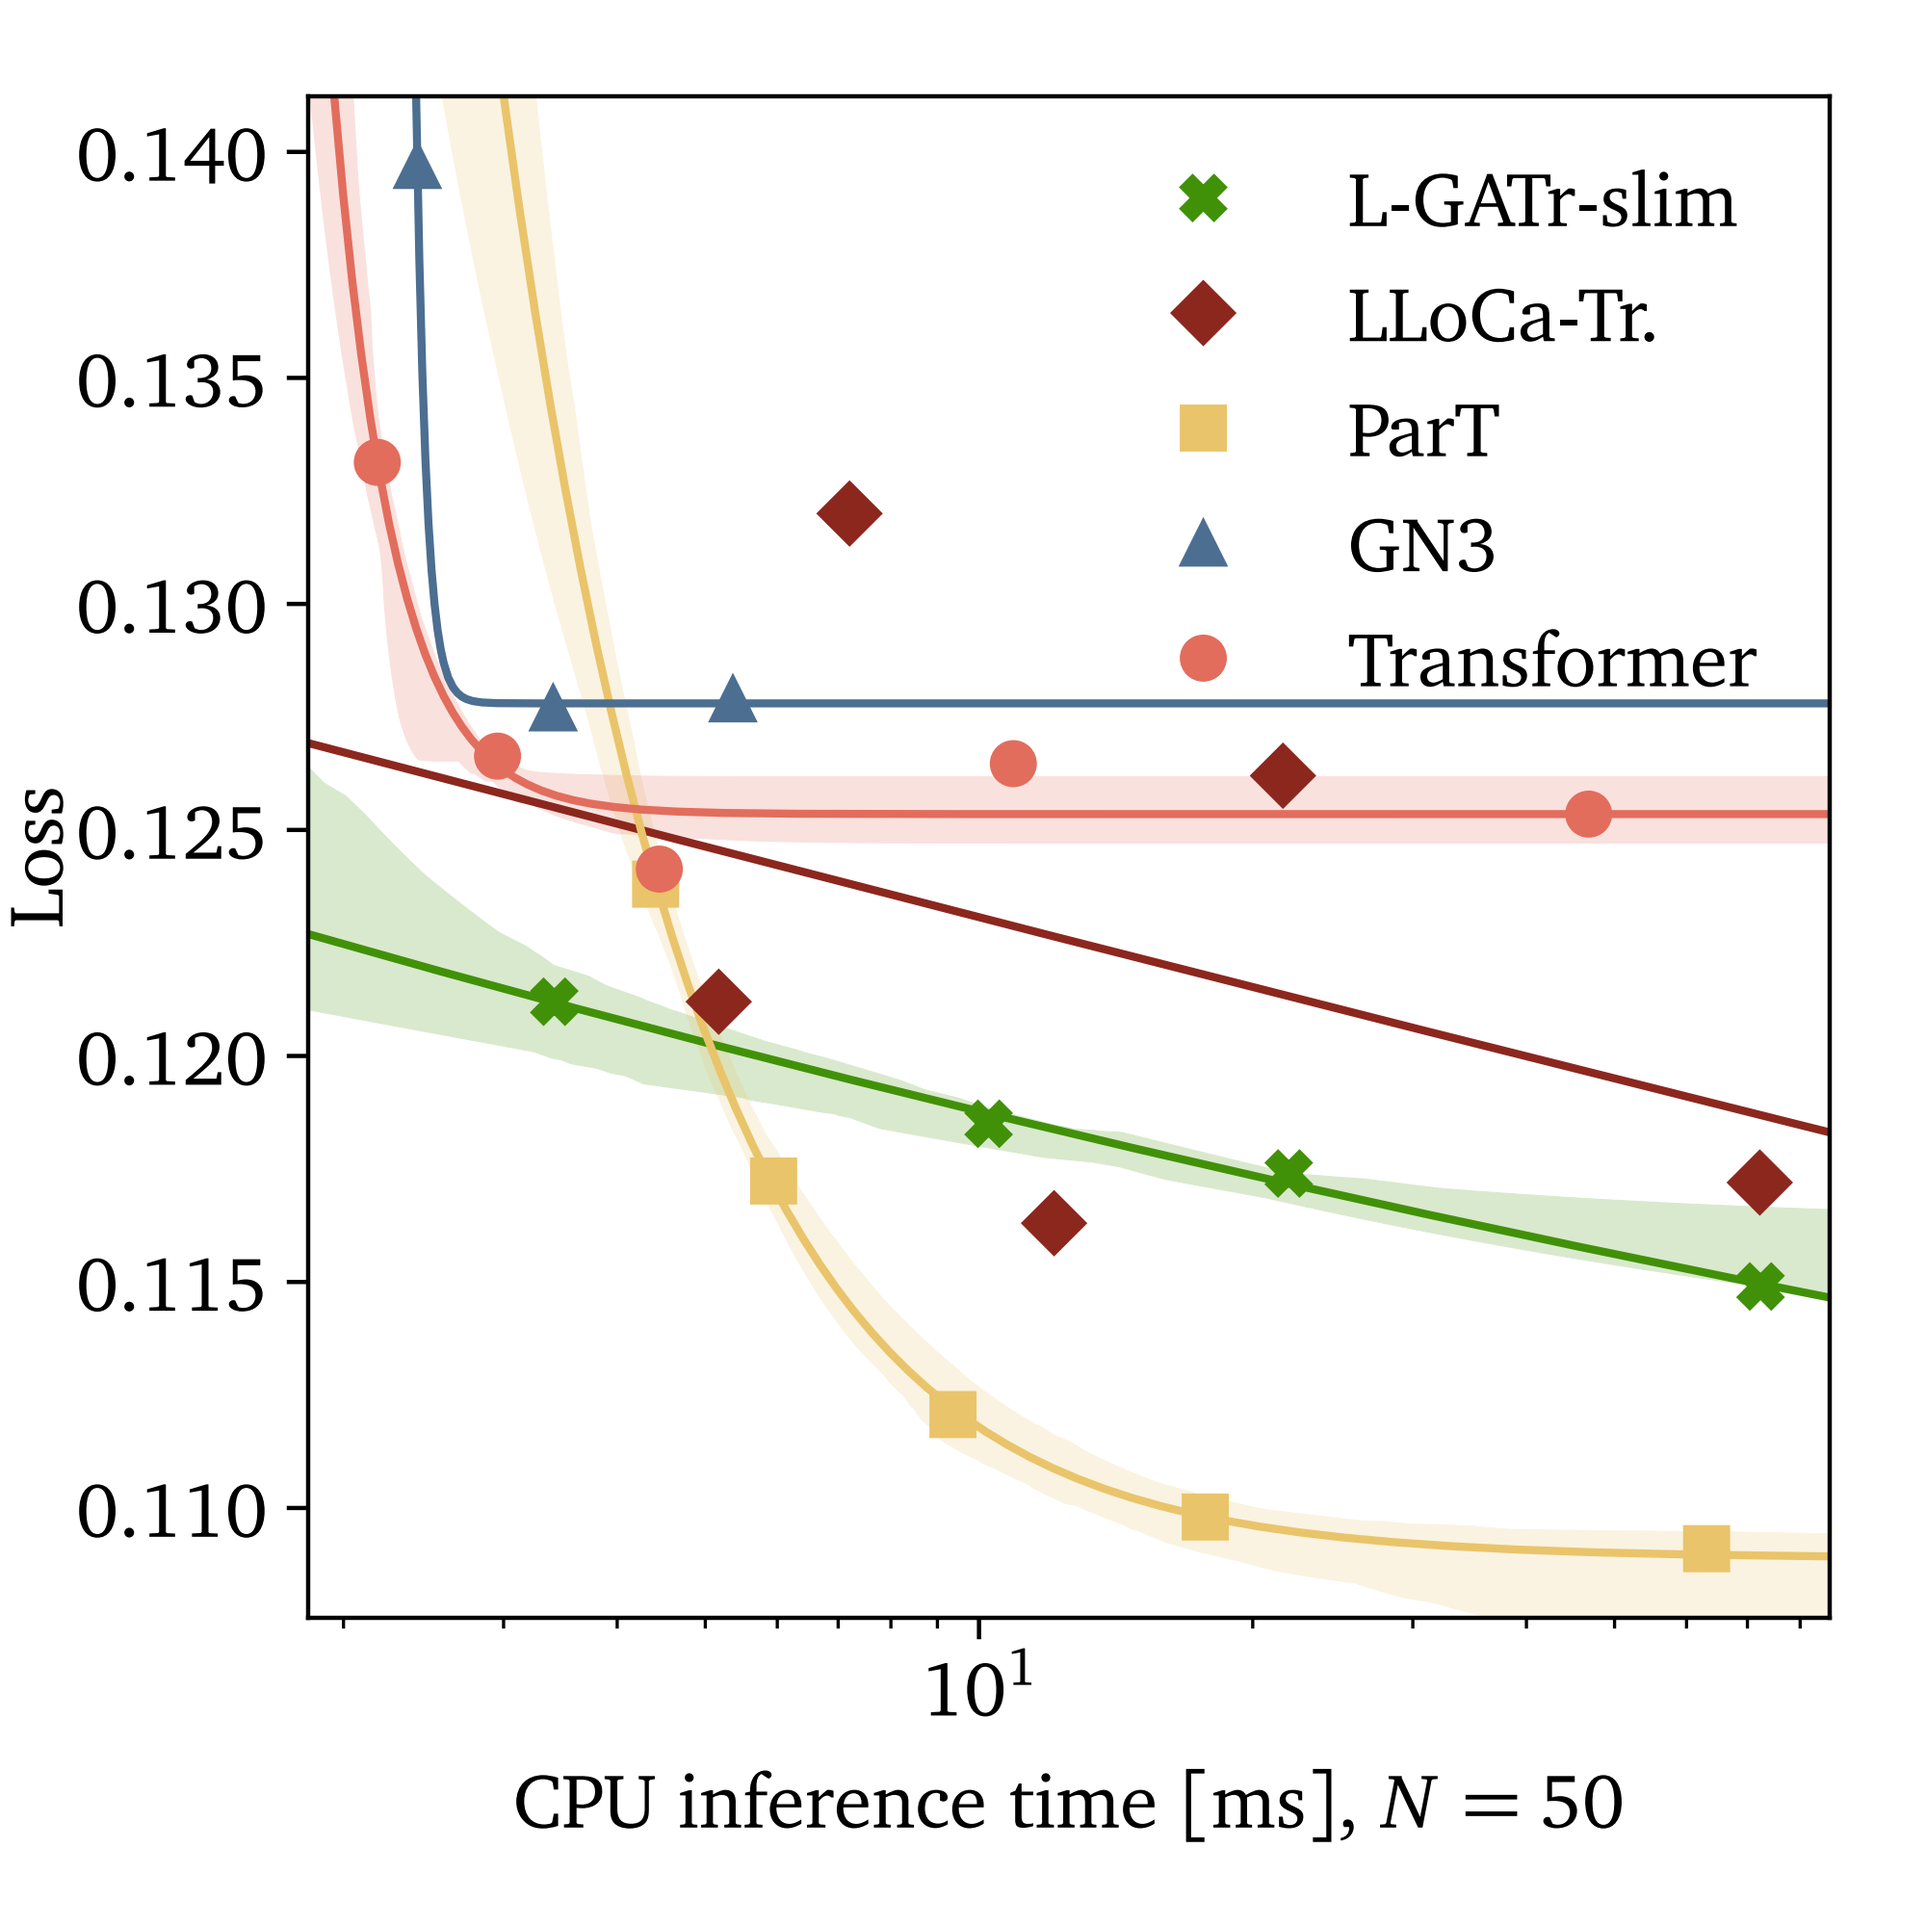
0.140
0.135
0.130
0.125
0.120
0.115
0.110
Loss
10^1
CPU inference time [ms], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
GN3
Transformer

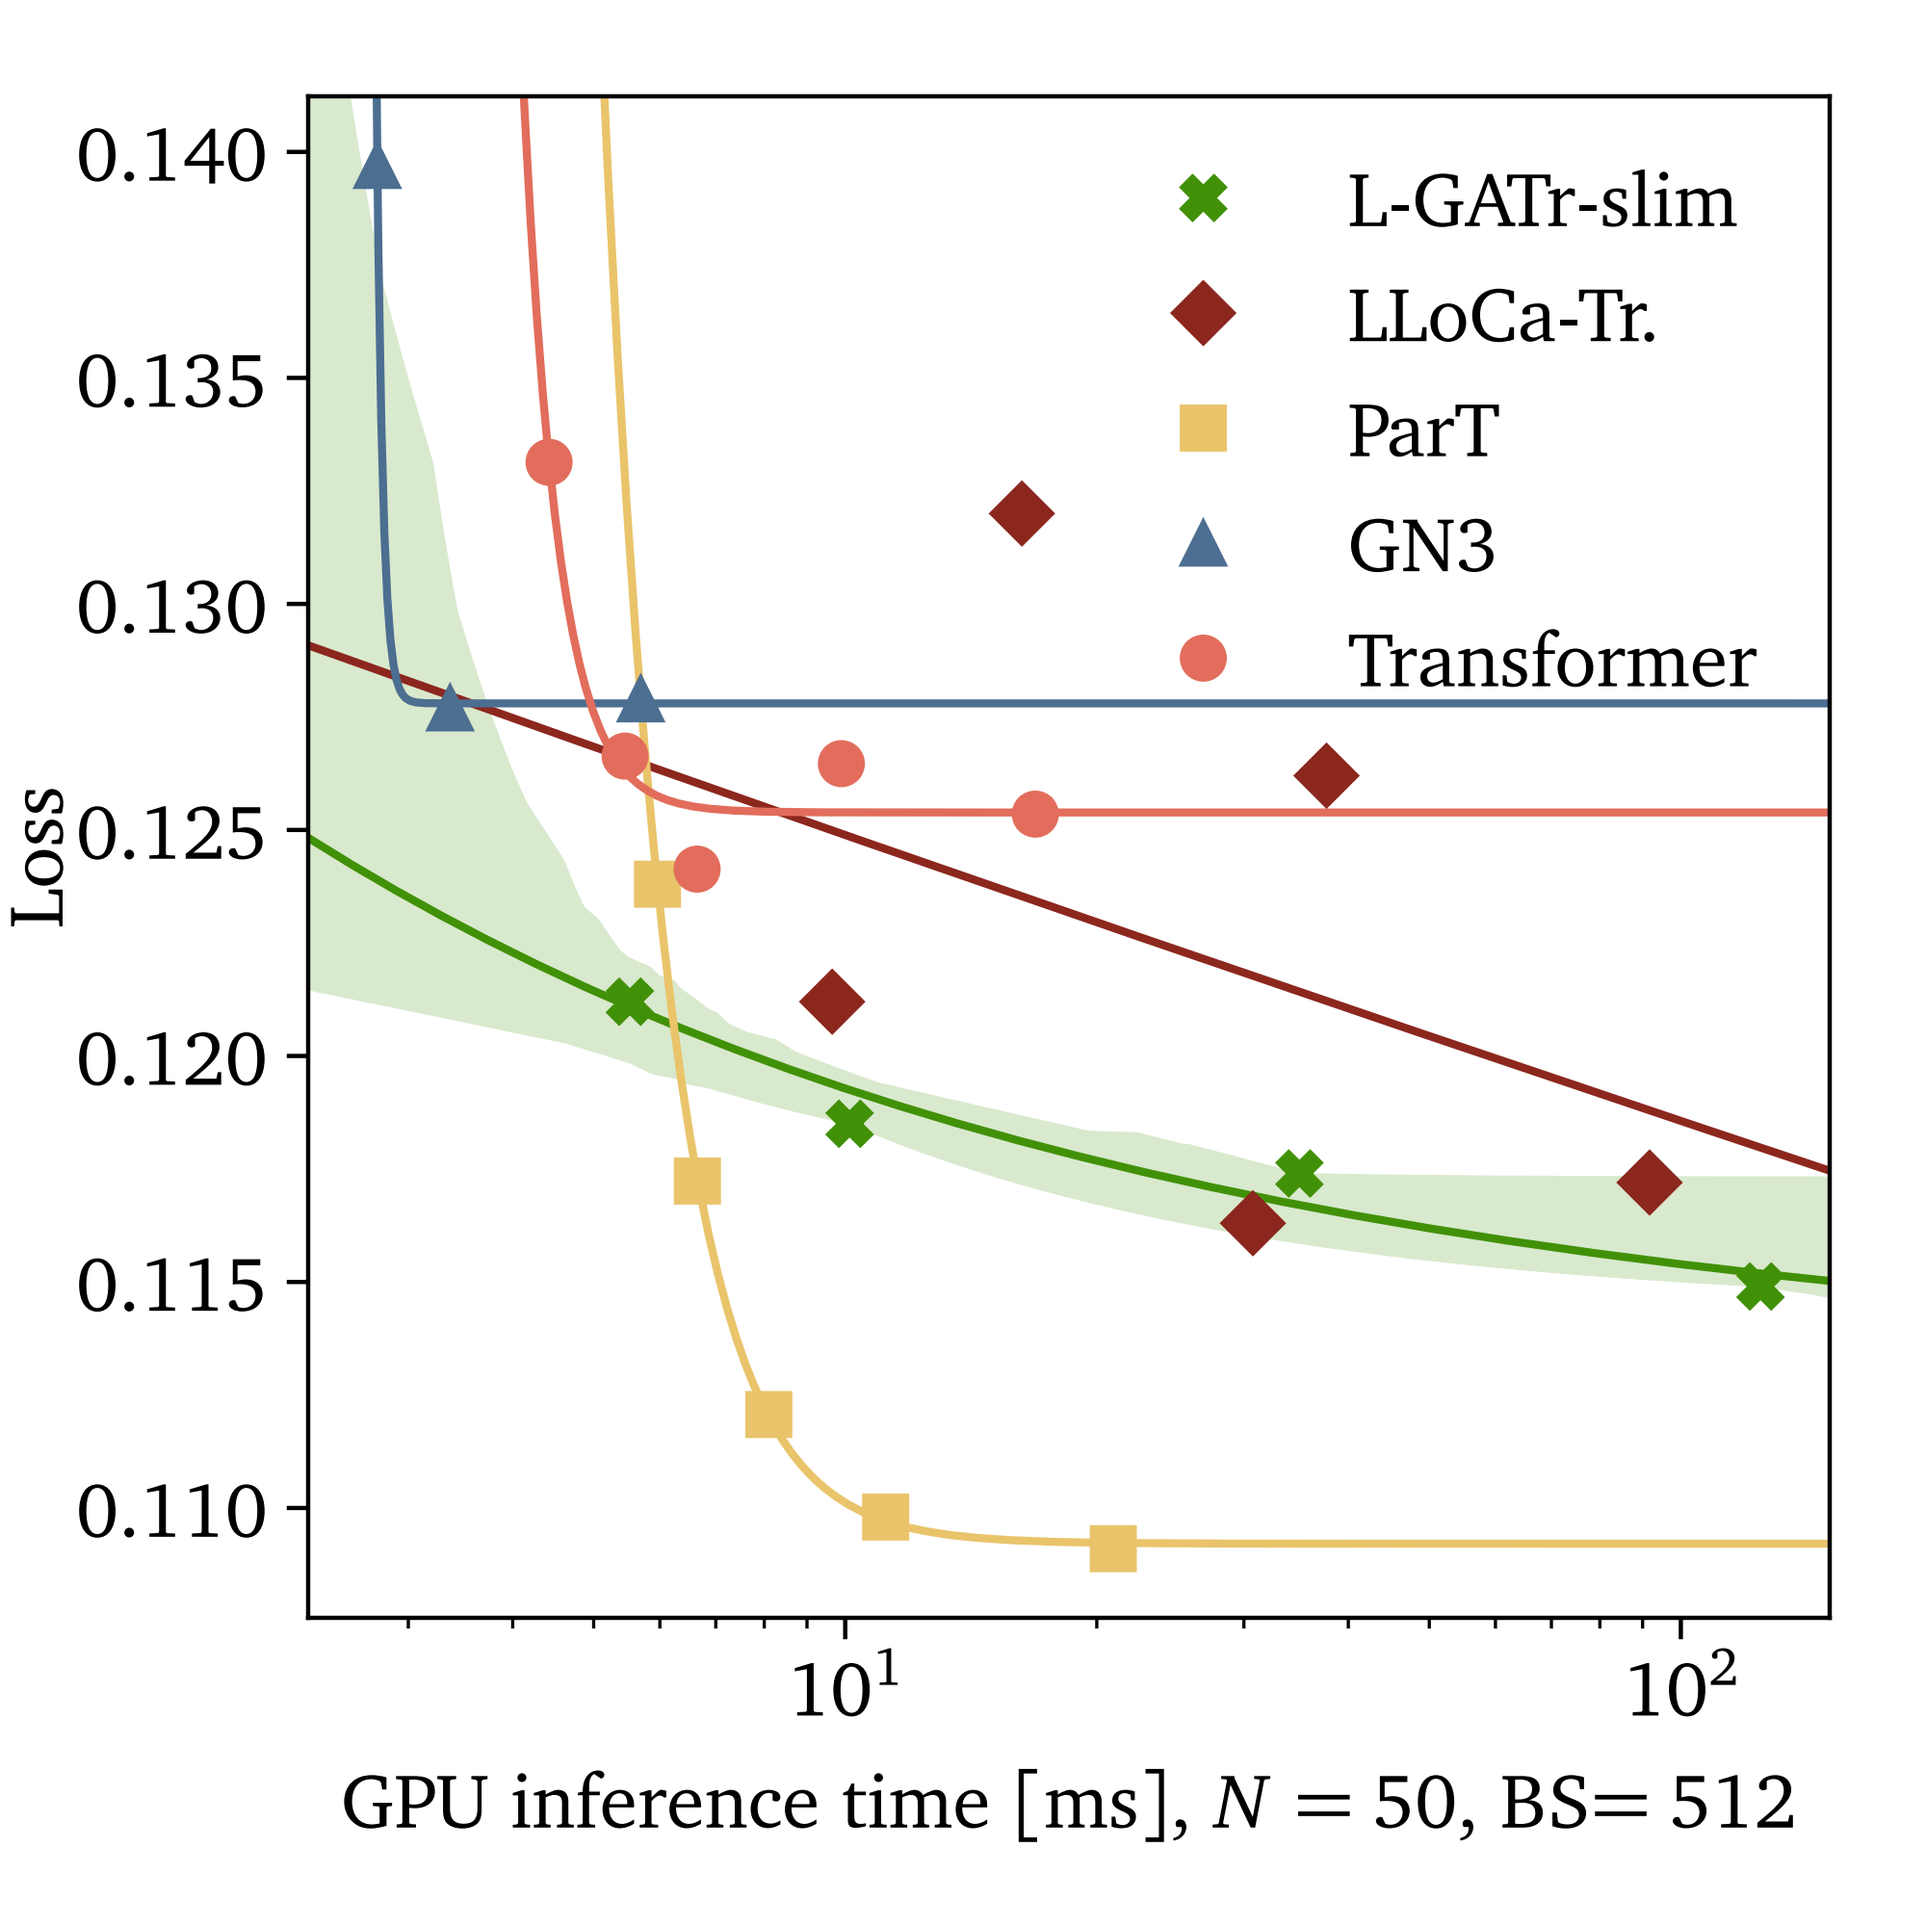
L-GATr-slim
LLoCa-Tr.
ParT
GN3
Transformer
Loss
0.140
0.135
0.130
0.125
0.120
0.115
0.110
10^1
10^2
GPU inference time [ms], $N = 50$, BS$= 512$

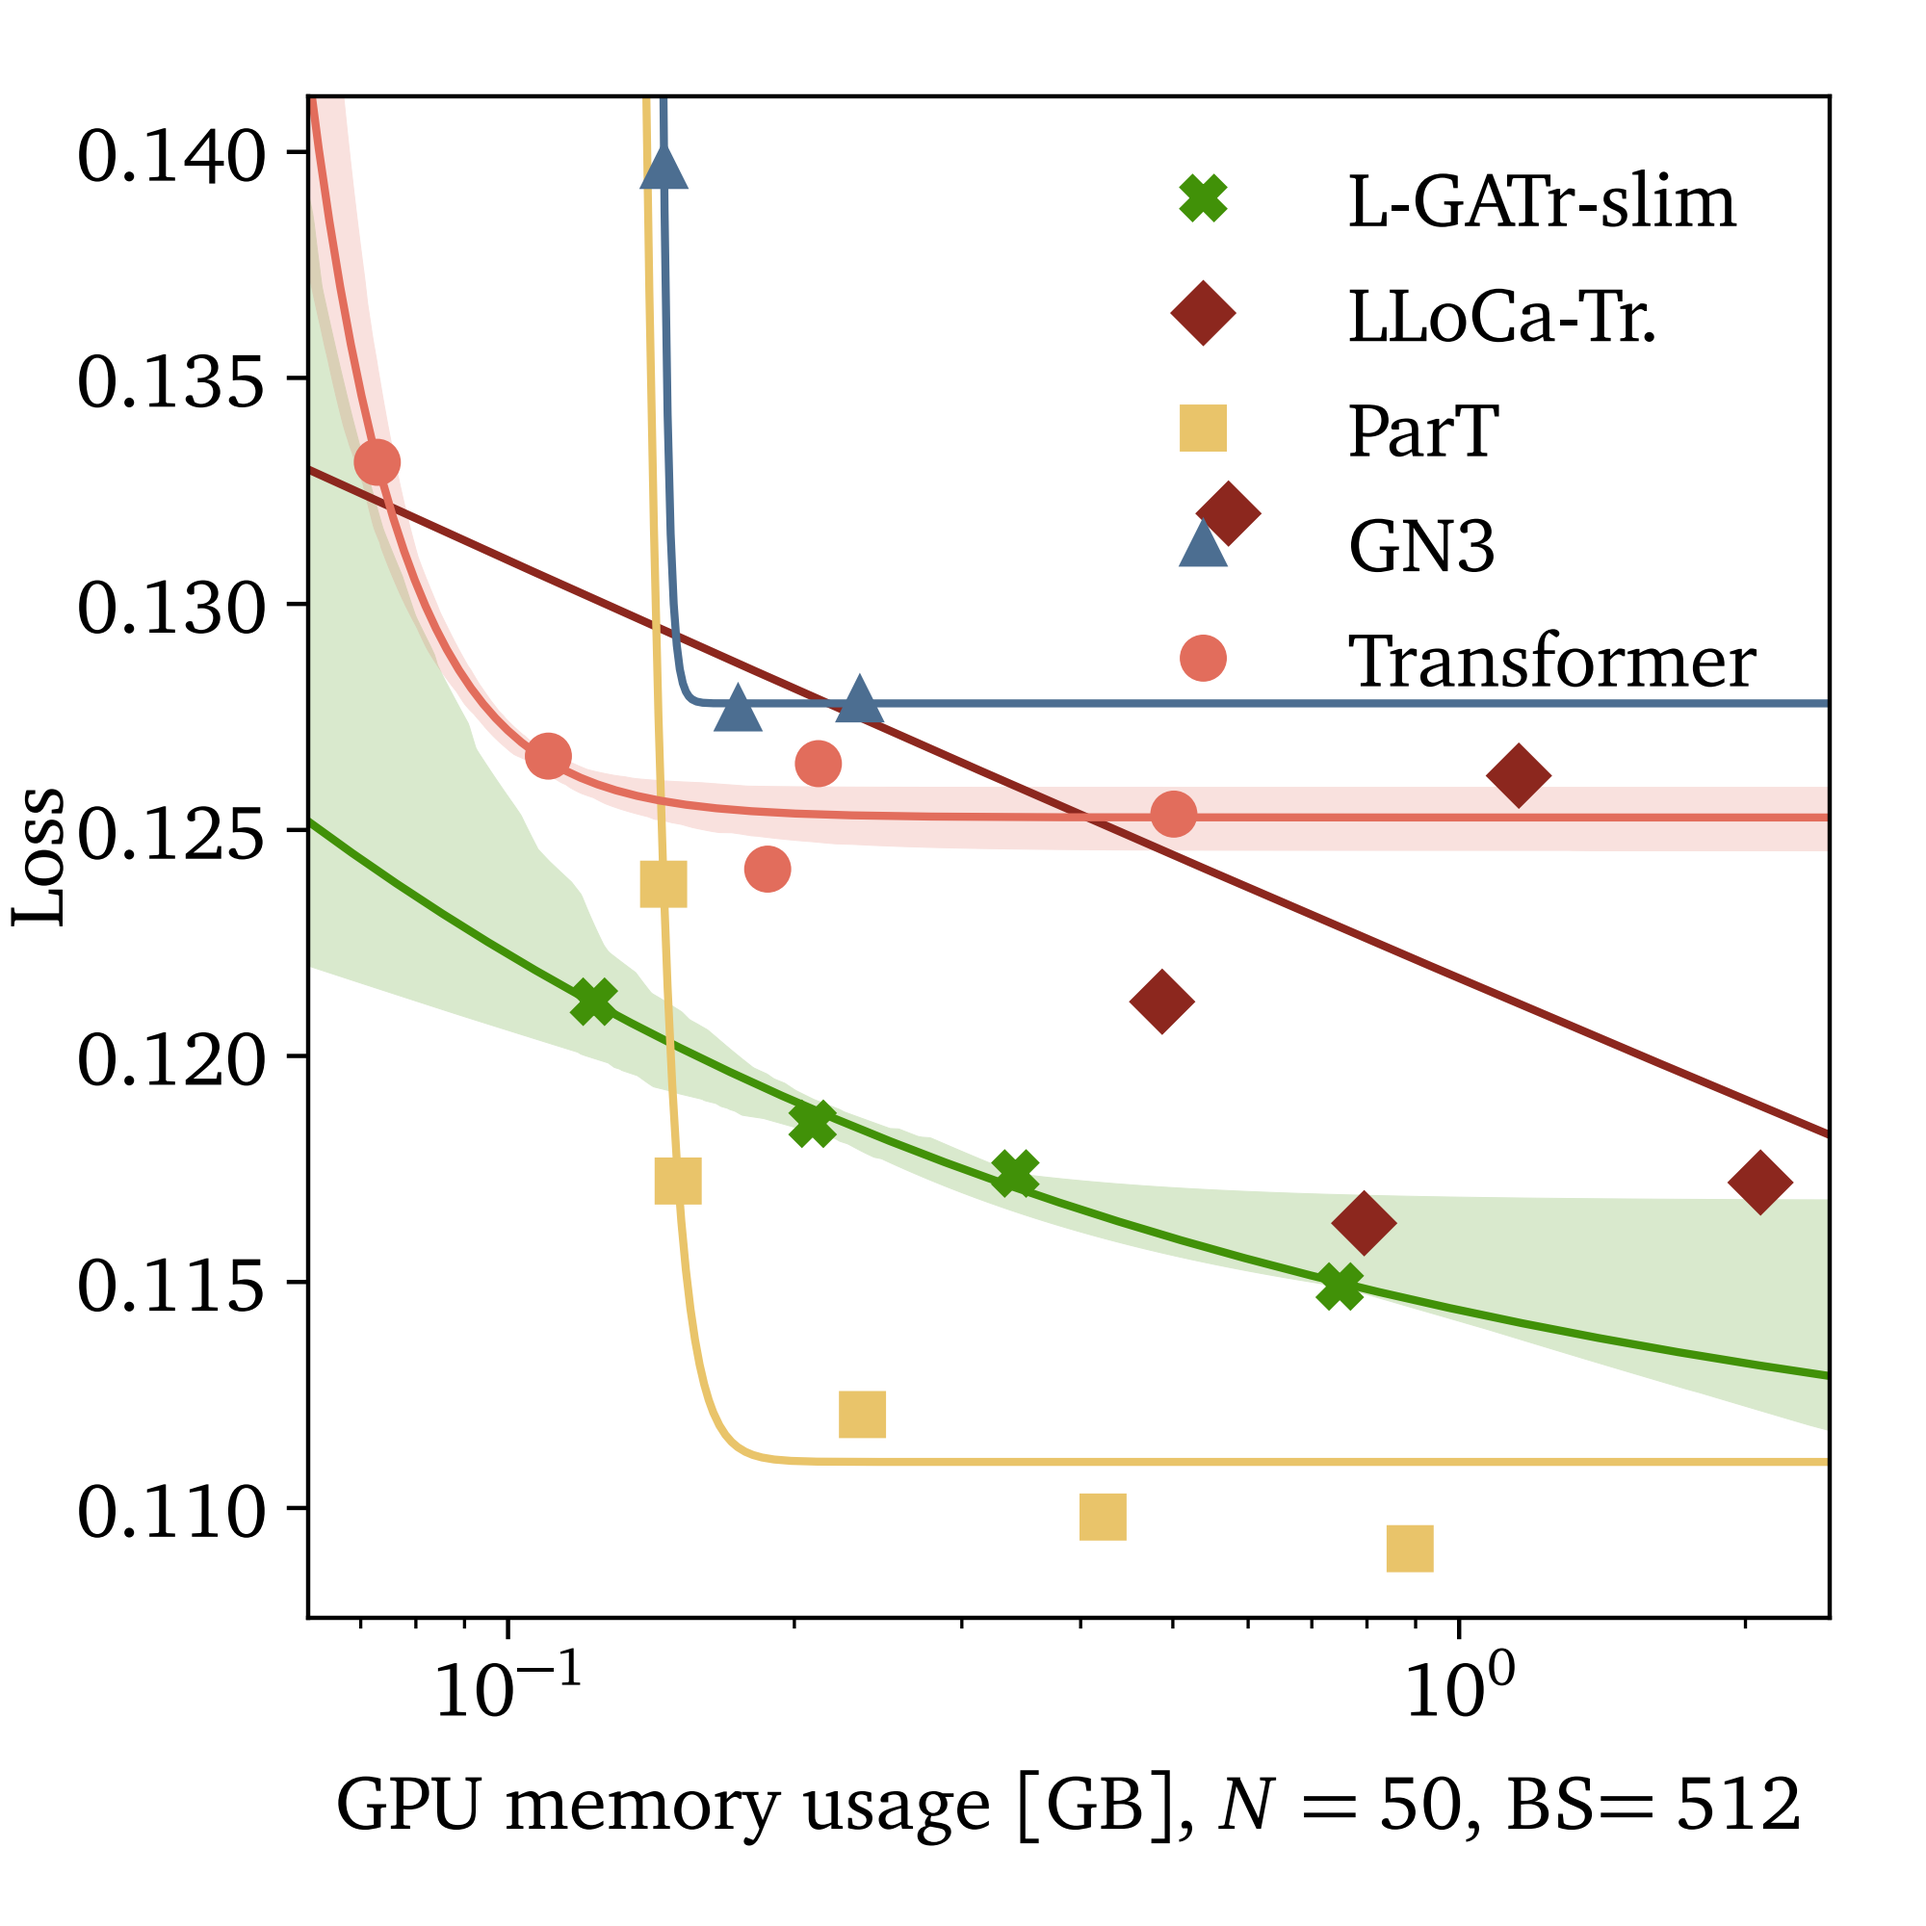
0.140
0.135
0.130
0.125
0.120
0.115
0.110
Loss
10^{-1}
10^{0}
GPU memory usage [GB], $N = 50$, BS$= 512$
L-GATr-slim
LLoCa-Tr.
ParT
GN3
Transformer

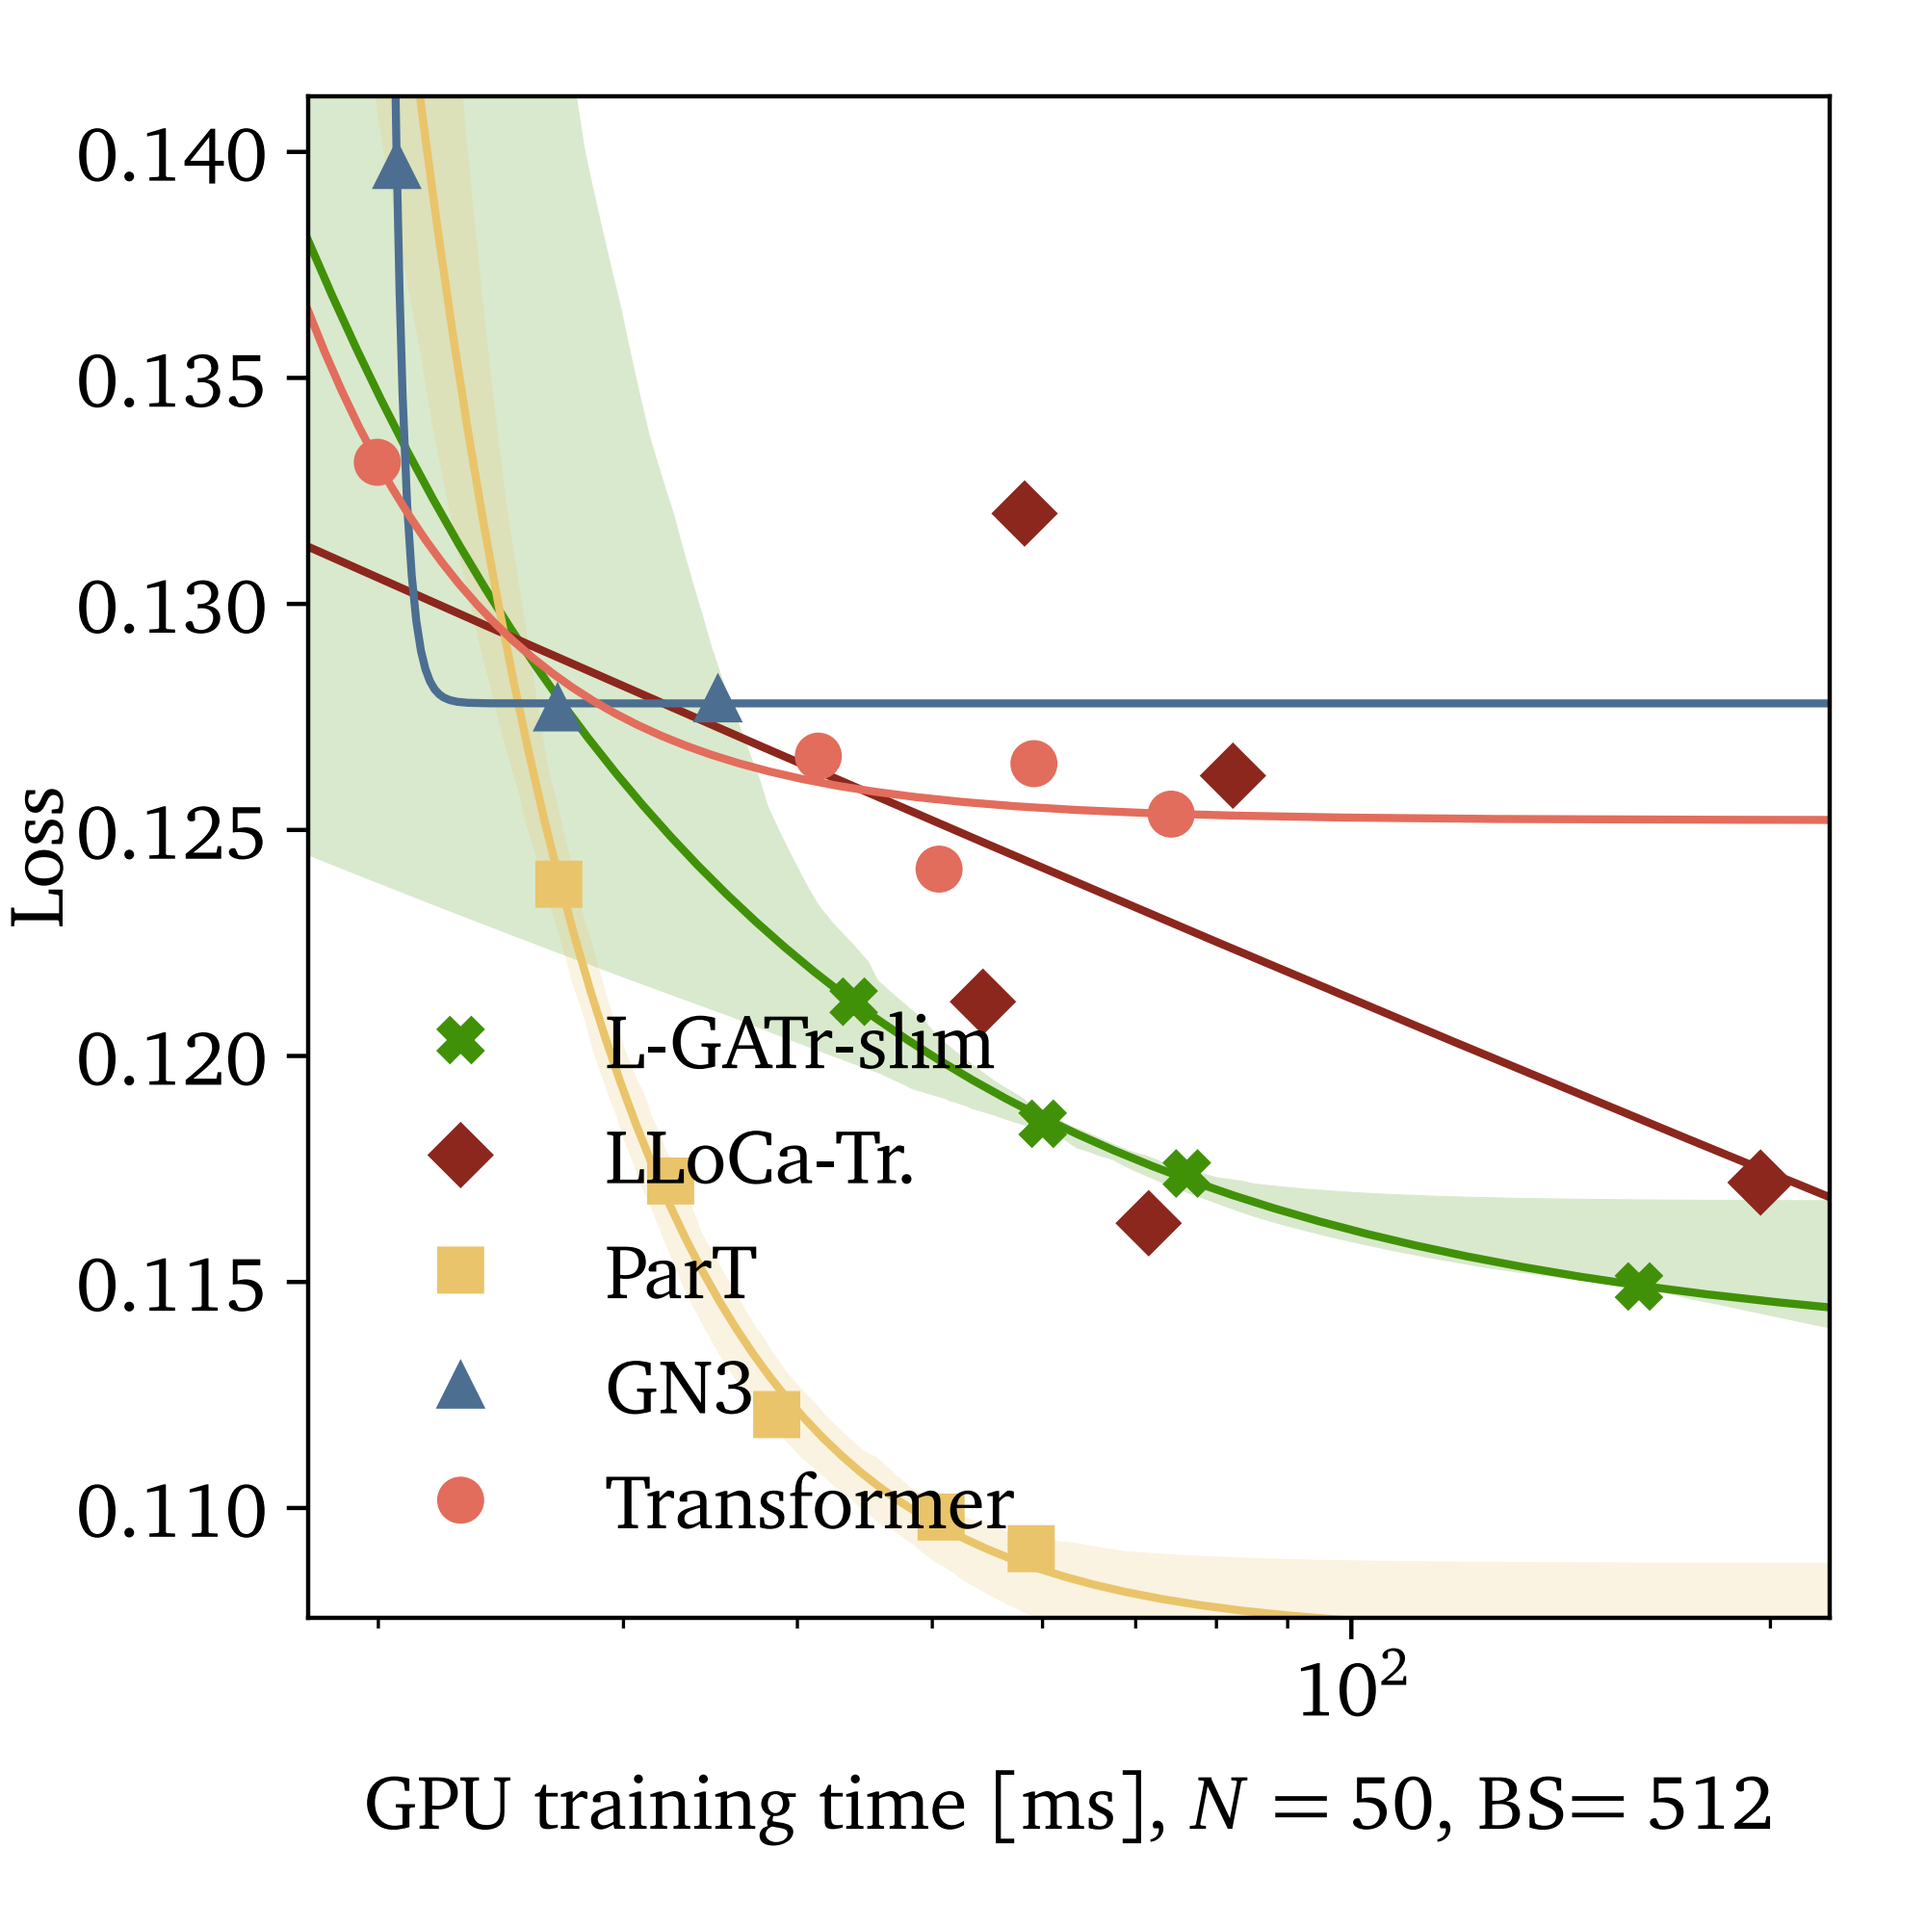
0.140
0.135
0.130
0.125
0.120
0.115
0.110
Loss
L-GATr-slim
LLoCa-Tr.
ParT
GN3
Transformer
10^2
GPU training time [ms], $N = 50$, BS$= 512$

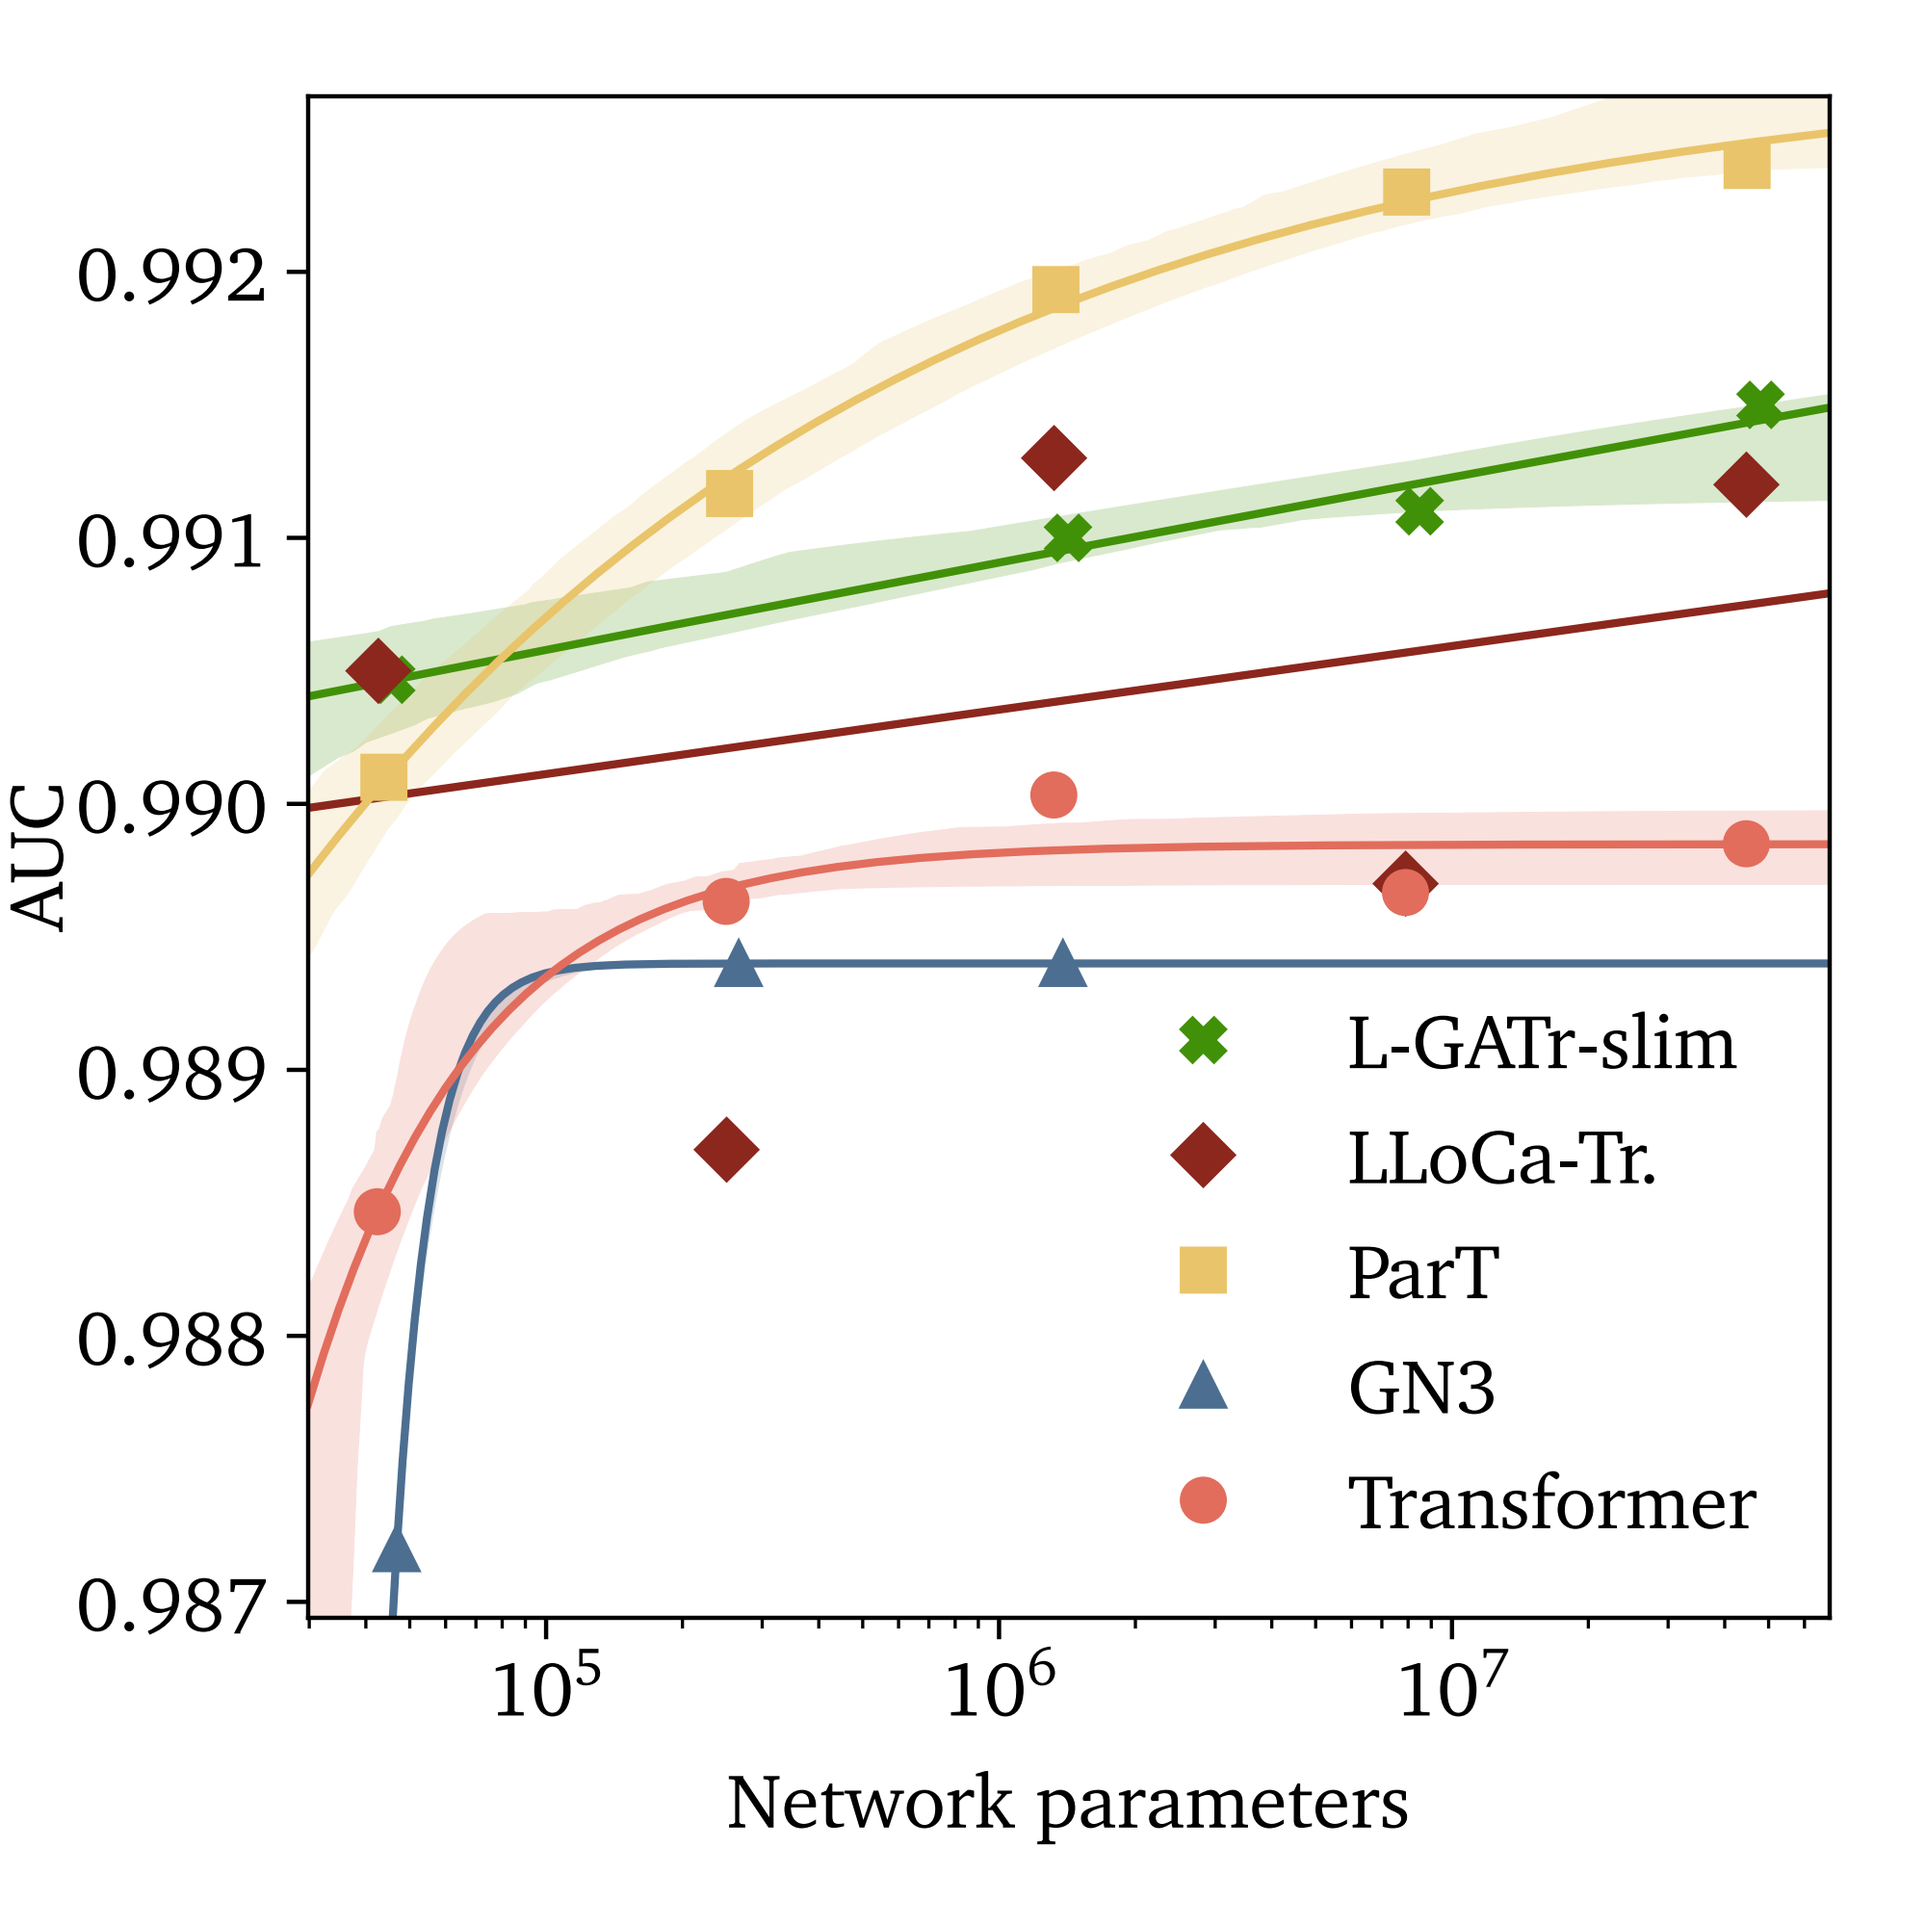
0.992
0.991
0.990
0.989
0.988
0.987
AUC
10^5
10^6
10^7
Network parameters
L-GATr-slim
LLoCa-Tr.
ParT
GN3
Transformer

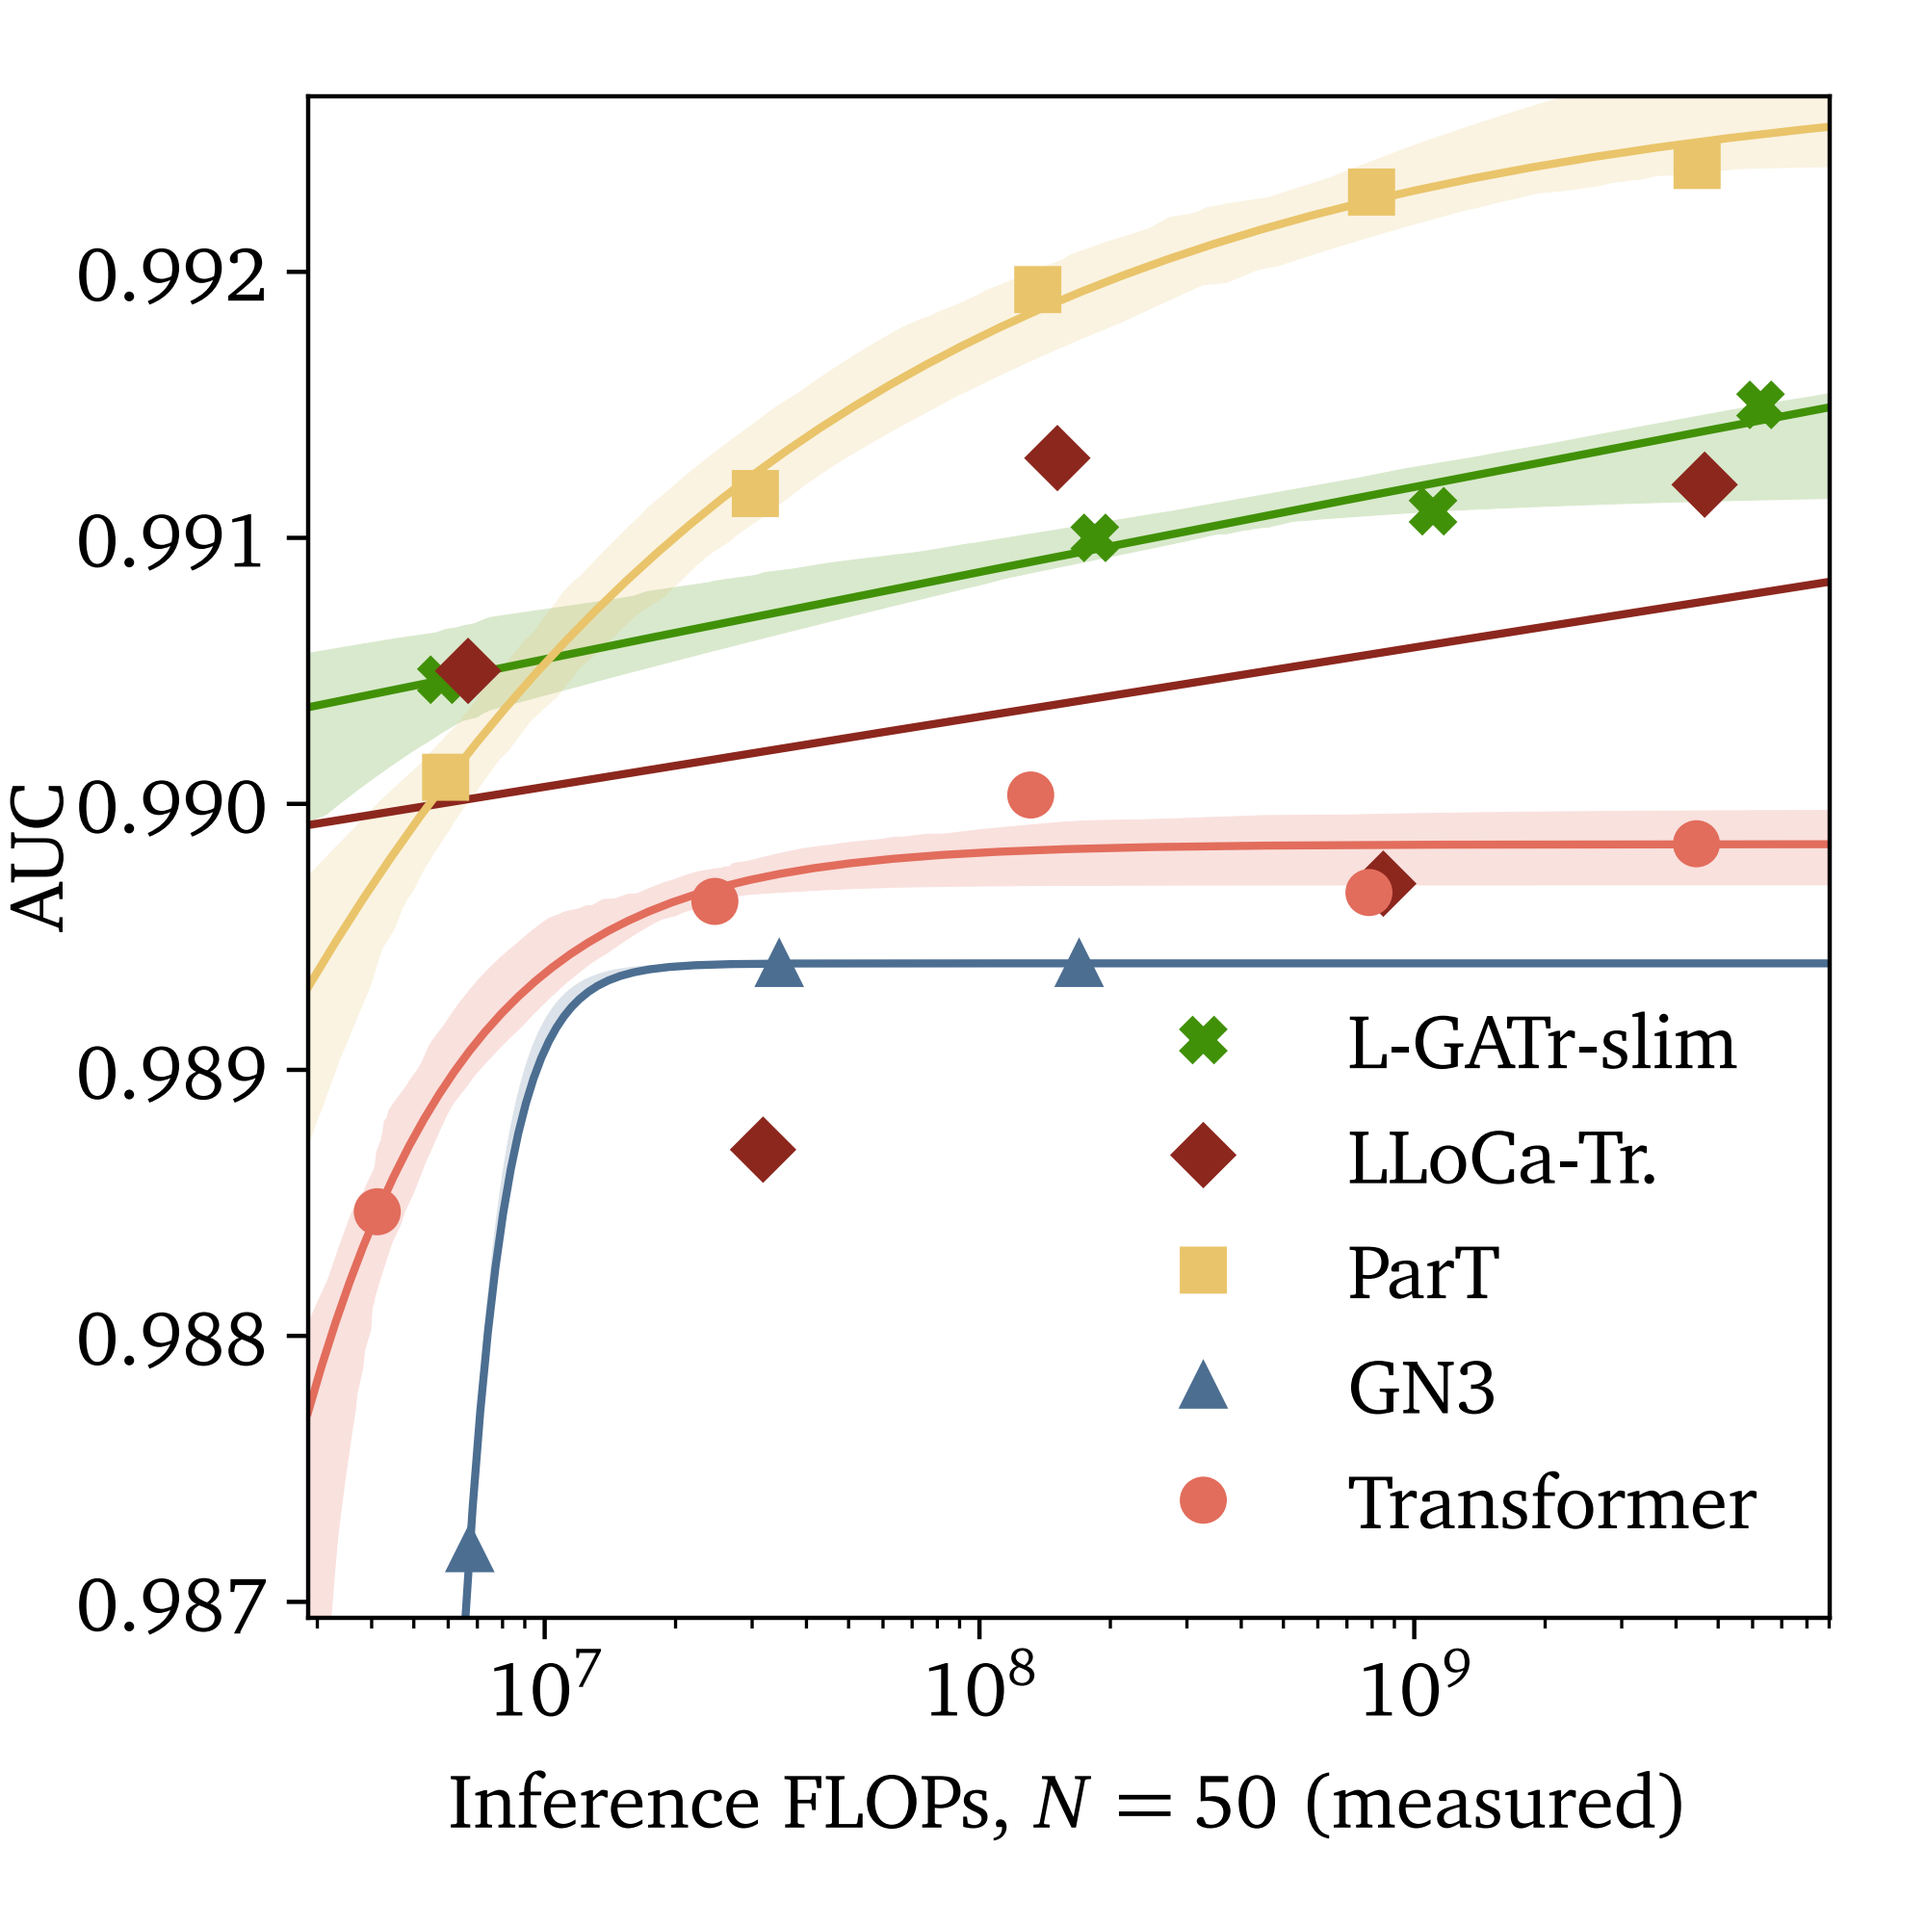
0.992
0.991
0.990
0.989
0.988
0.987
AUC
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (measured)
L-GATr-slim
LLoCa-Tr.
ParT
GN3
Transformer

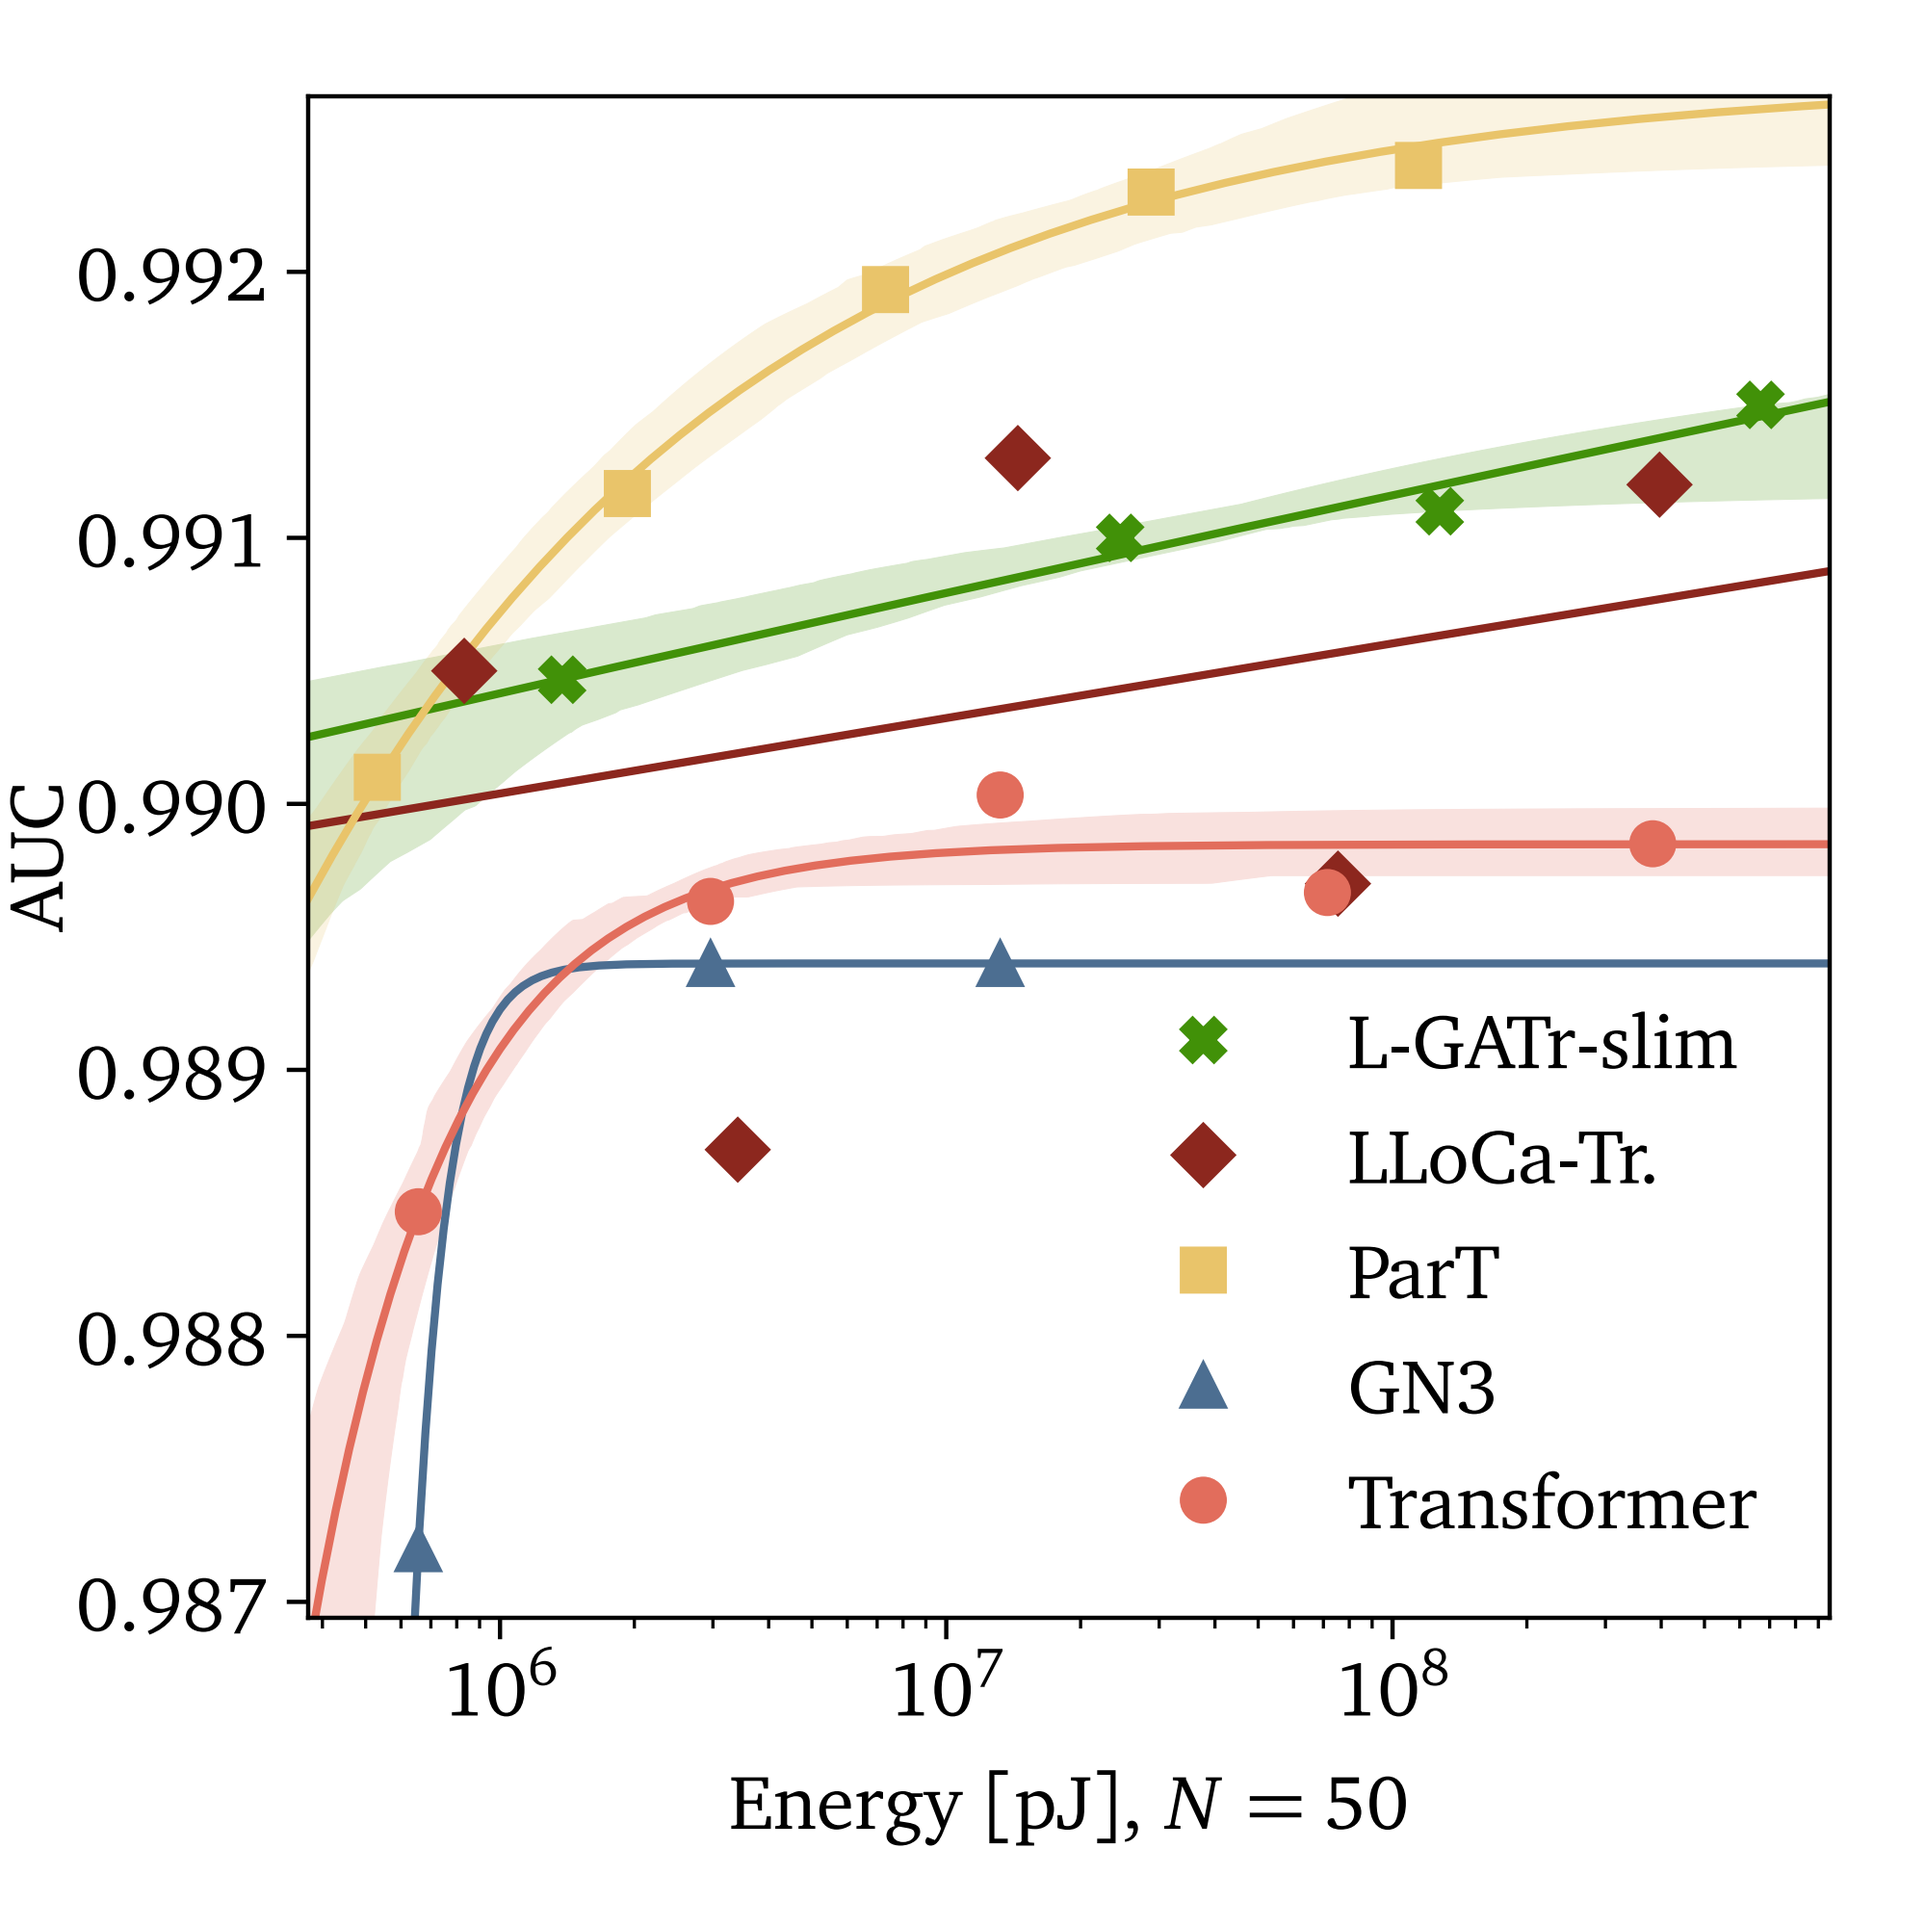
0.992
0.991
0.990
0.989
0.988
0.987
AUC
10^6
10^7
10^8
Energy [pJ], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
GN3
Transformer

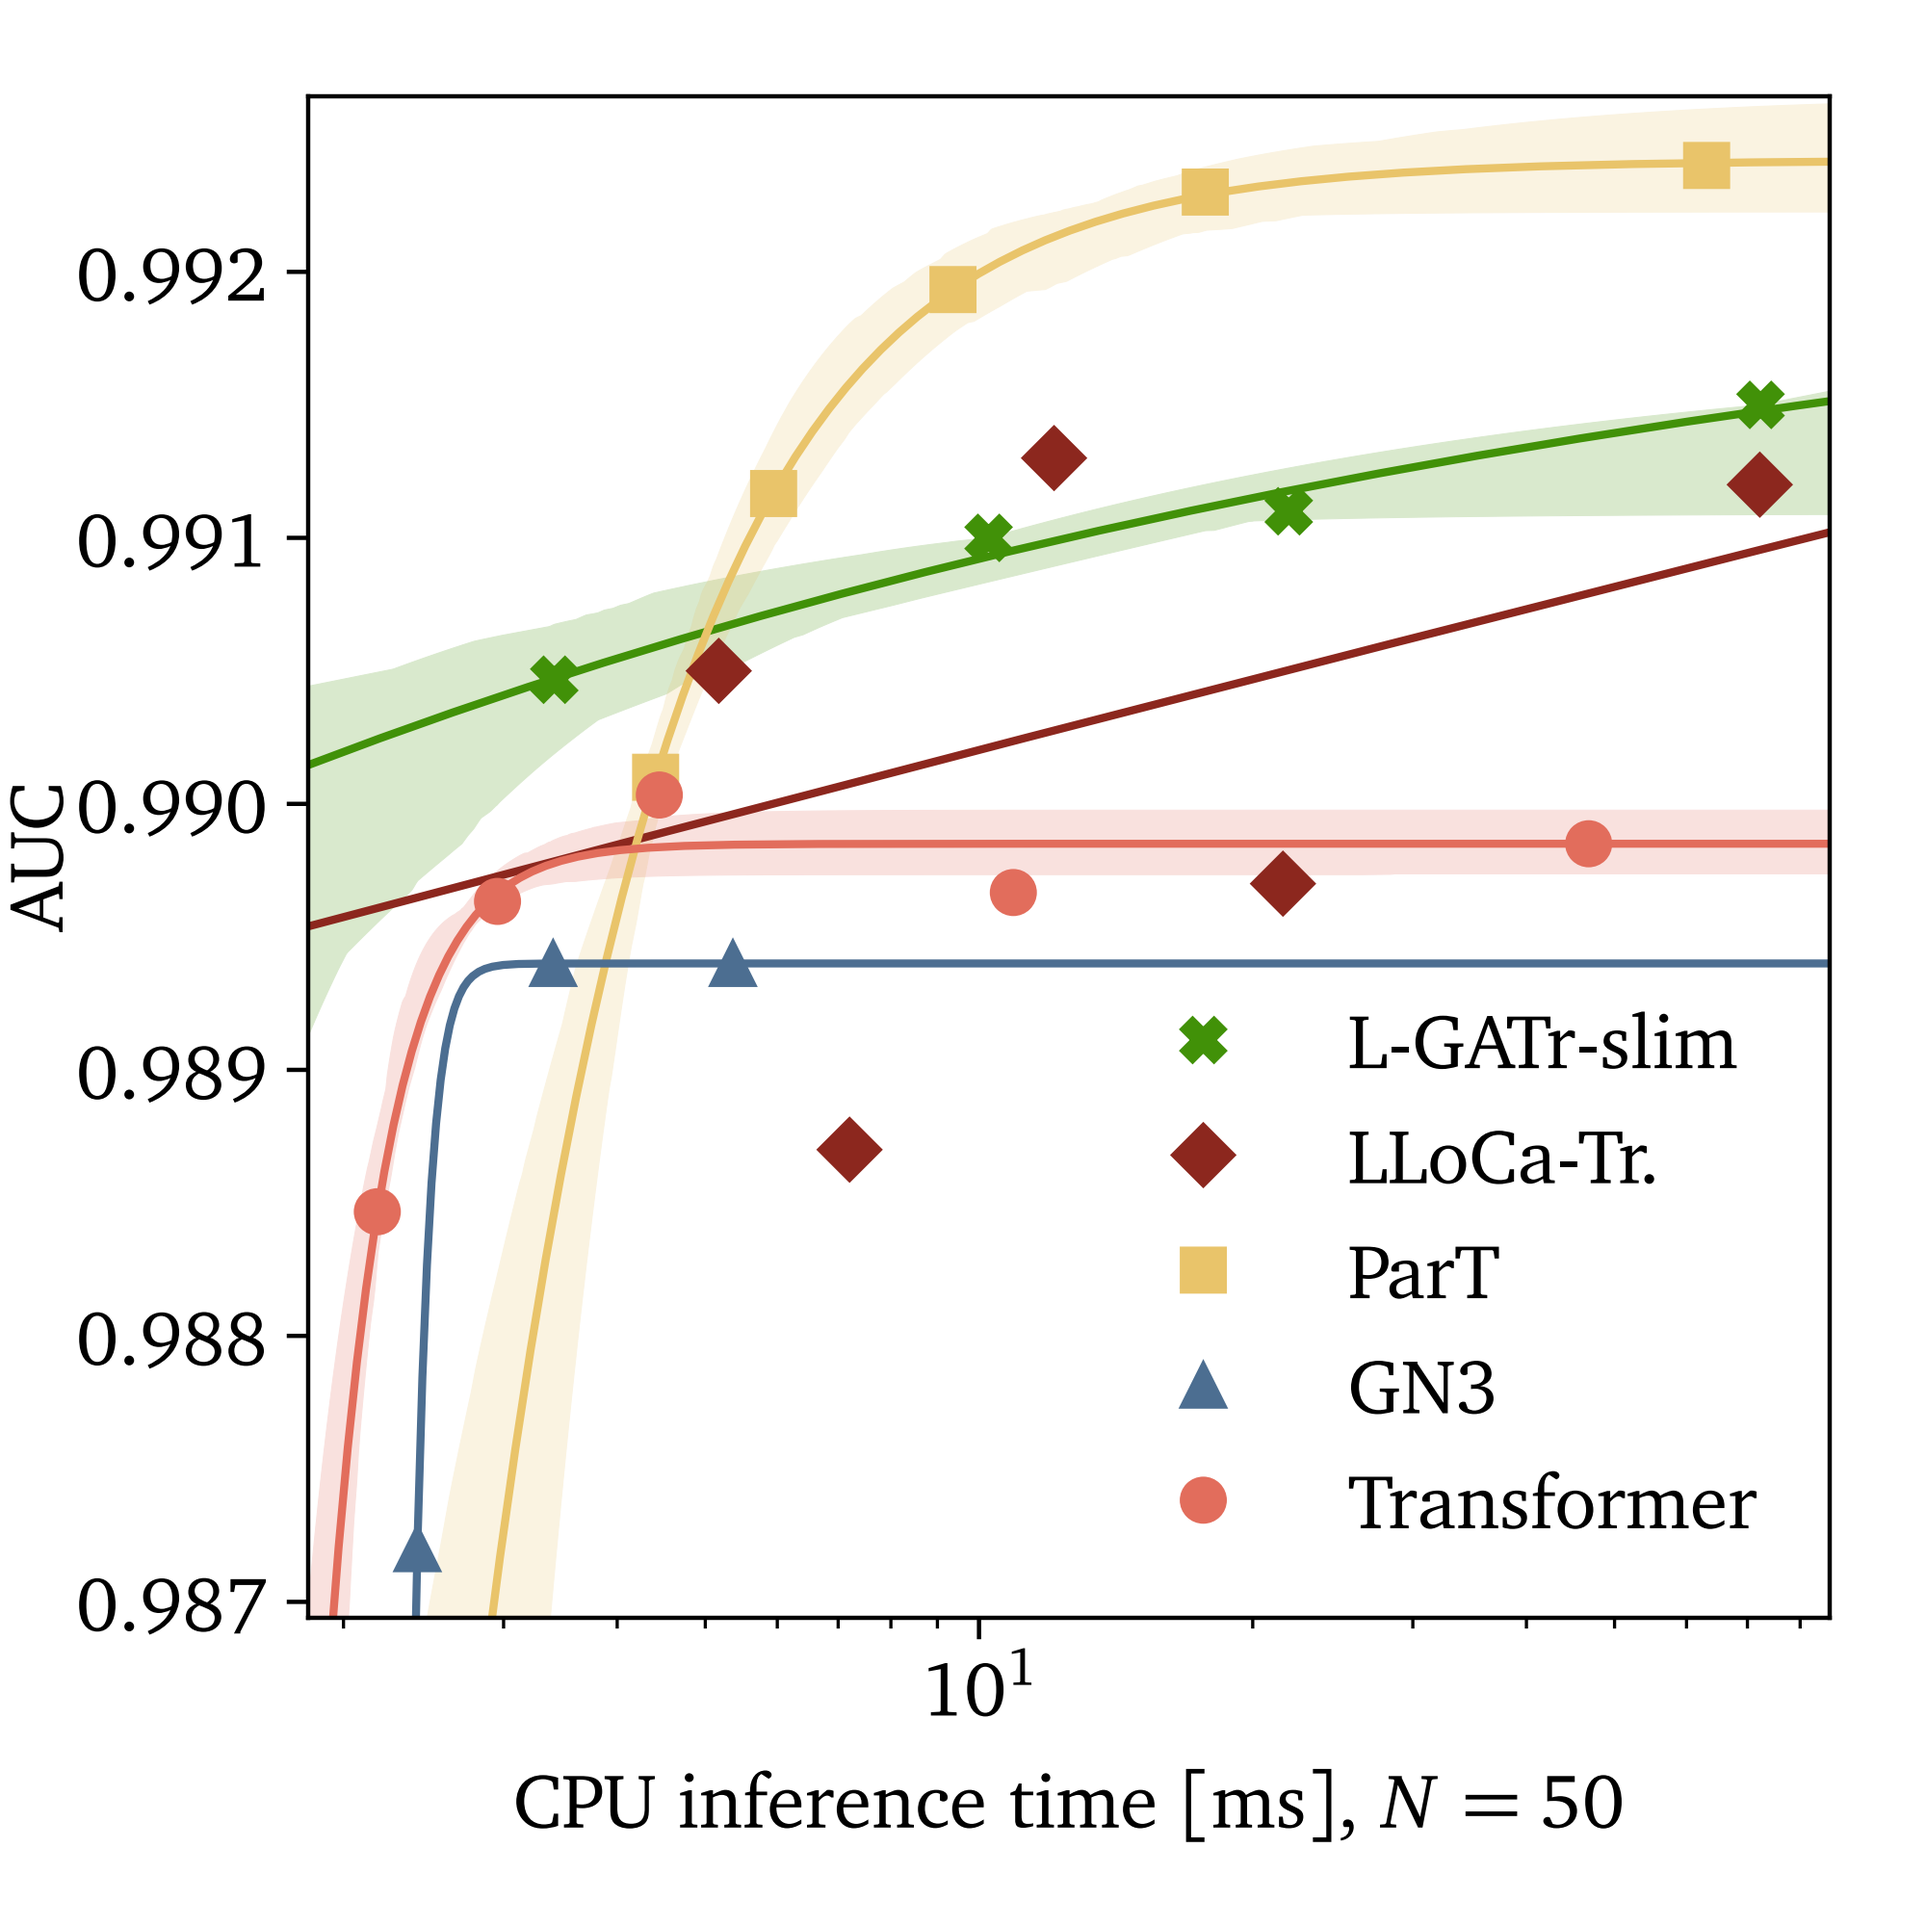
0.992
0.991
0.990
0.989
0.988
0.987
AUC
10^1
CPU inference time [ms], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
GN3
Transformer

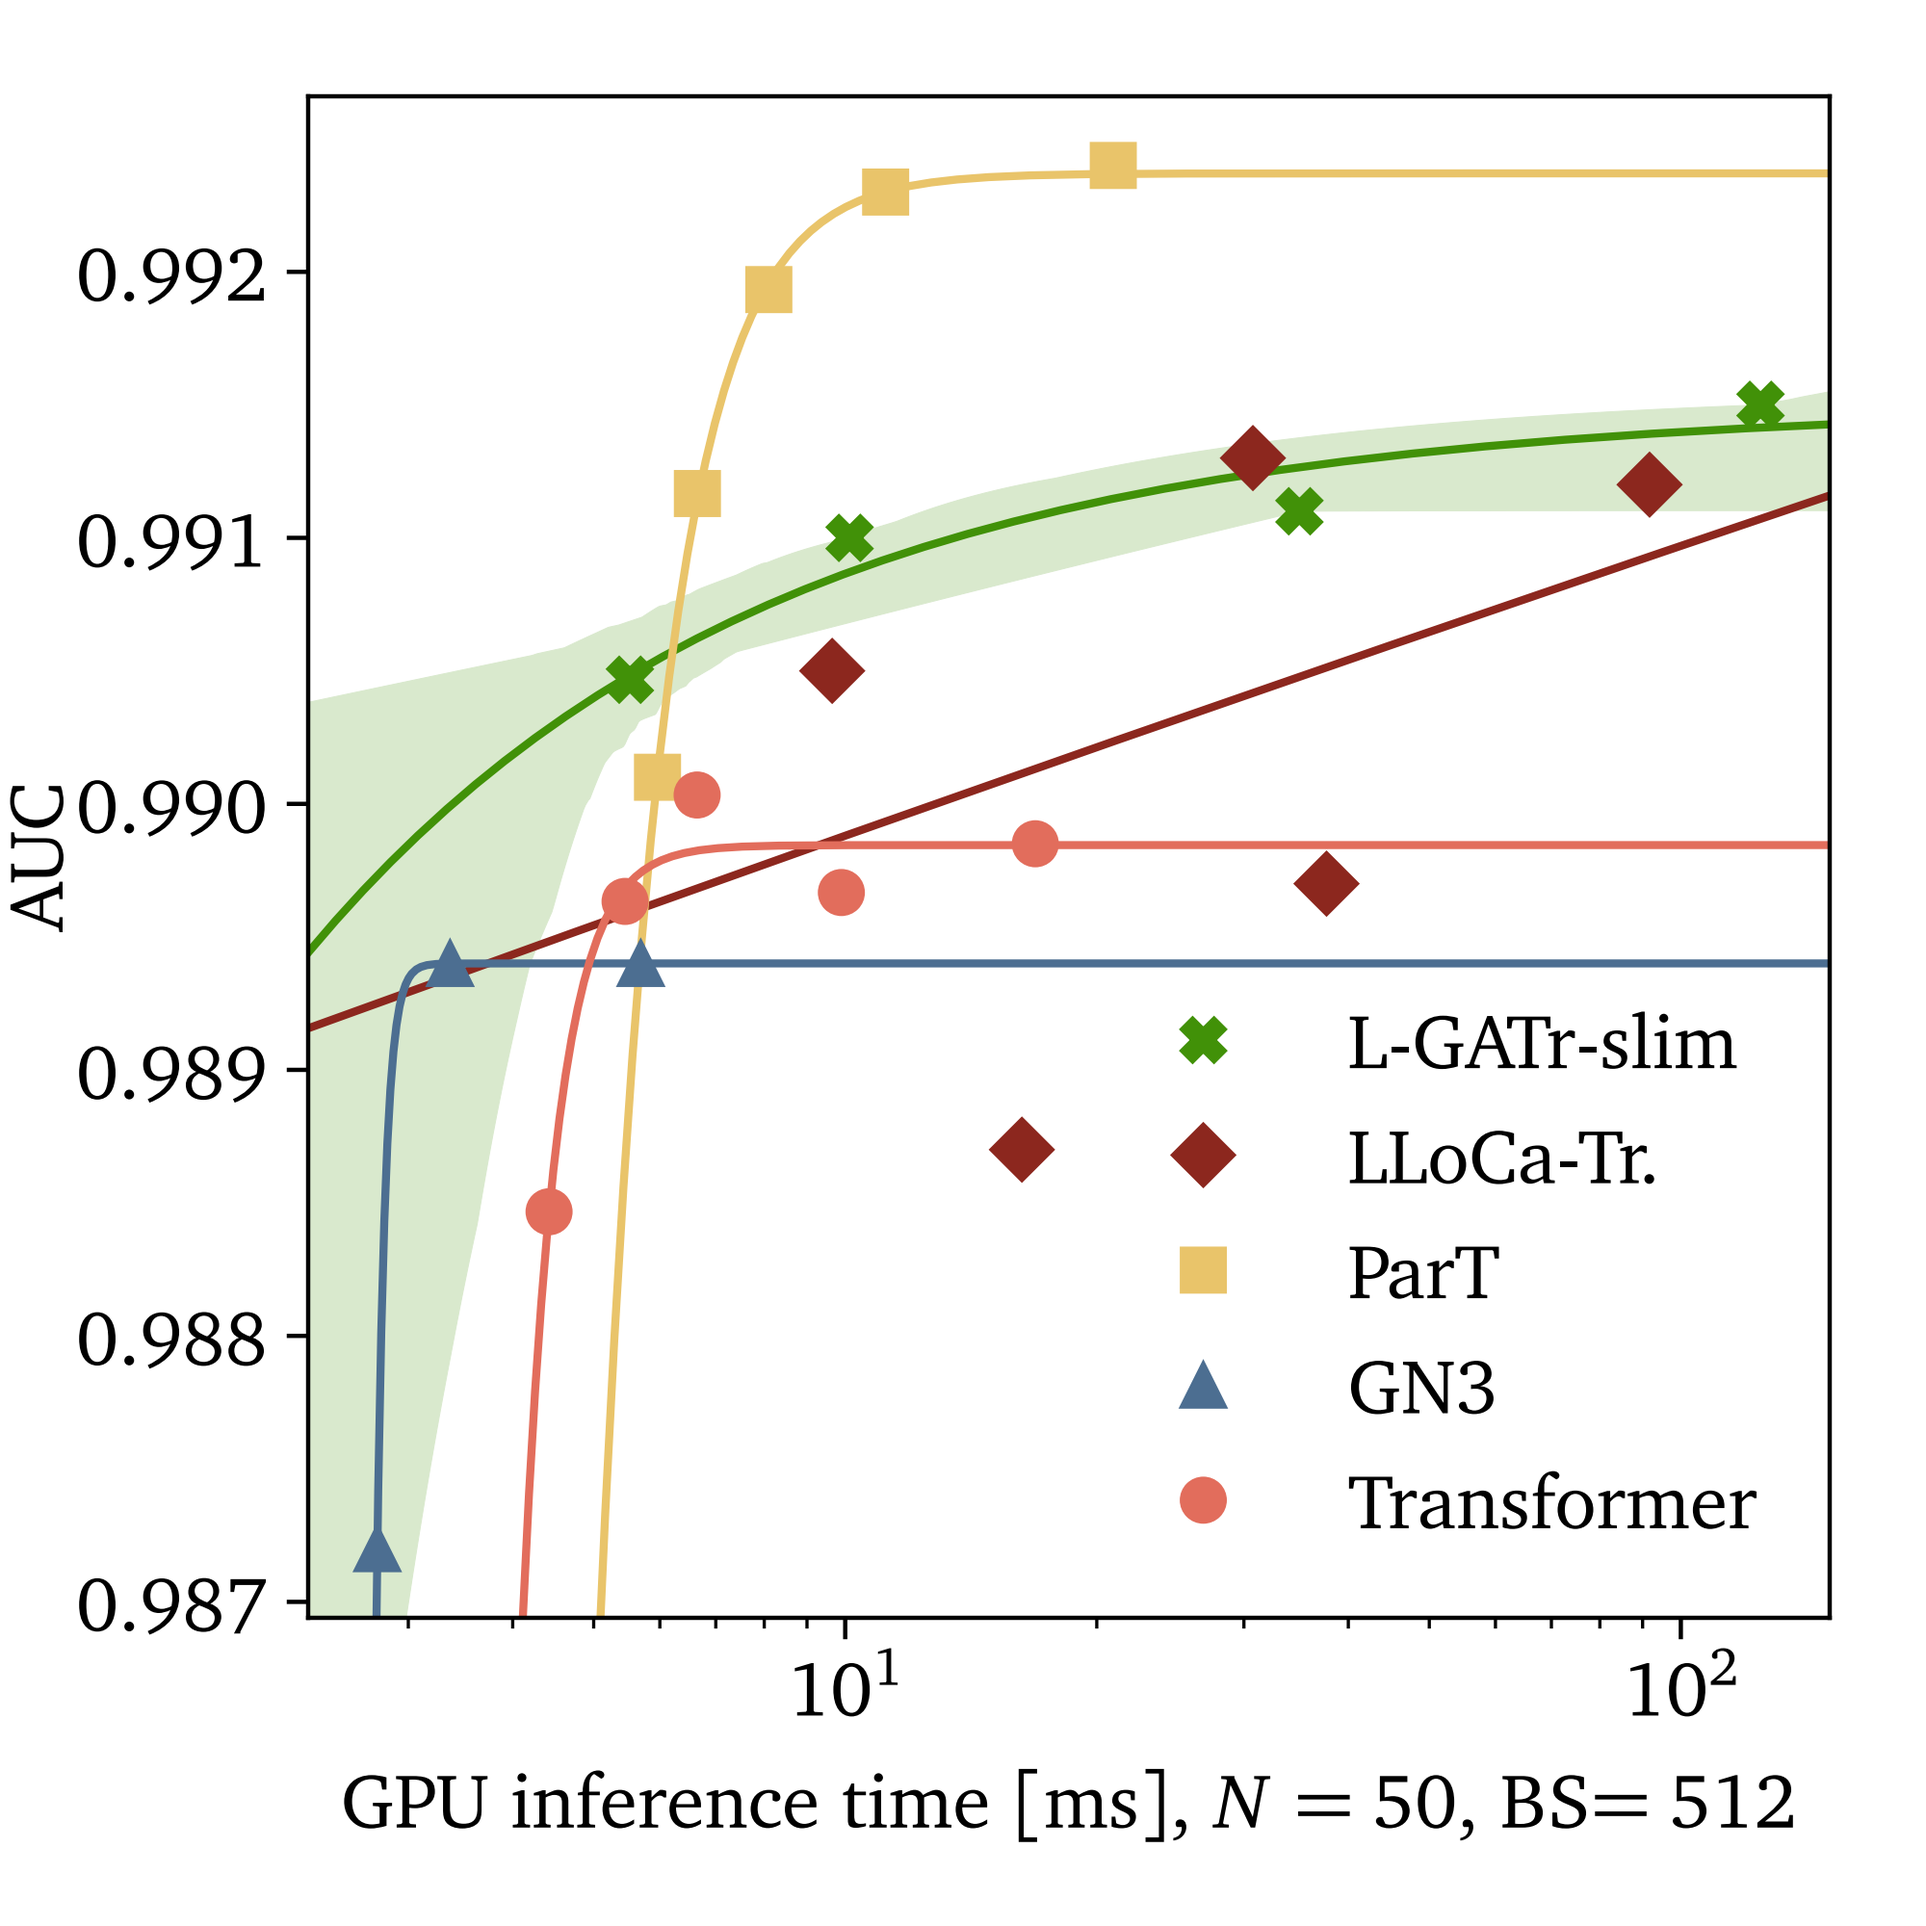
0.992
0.991
0.990
0.989
0.988
0.987
AUC
10^1
10^2
GPU inference time [ms], $N = 50$, BS= 512
L-GATr-slim
LLoCa-Tr.
ParT
GN3
Transformer

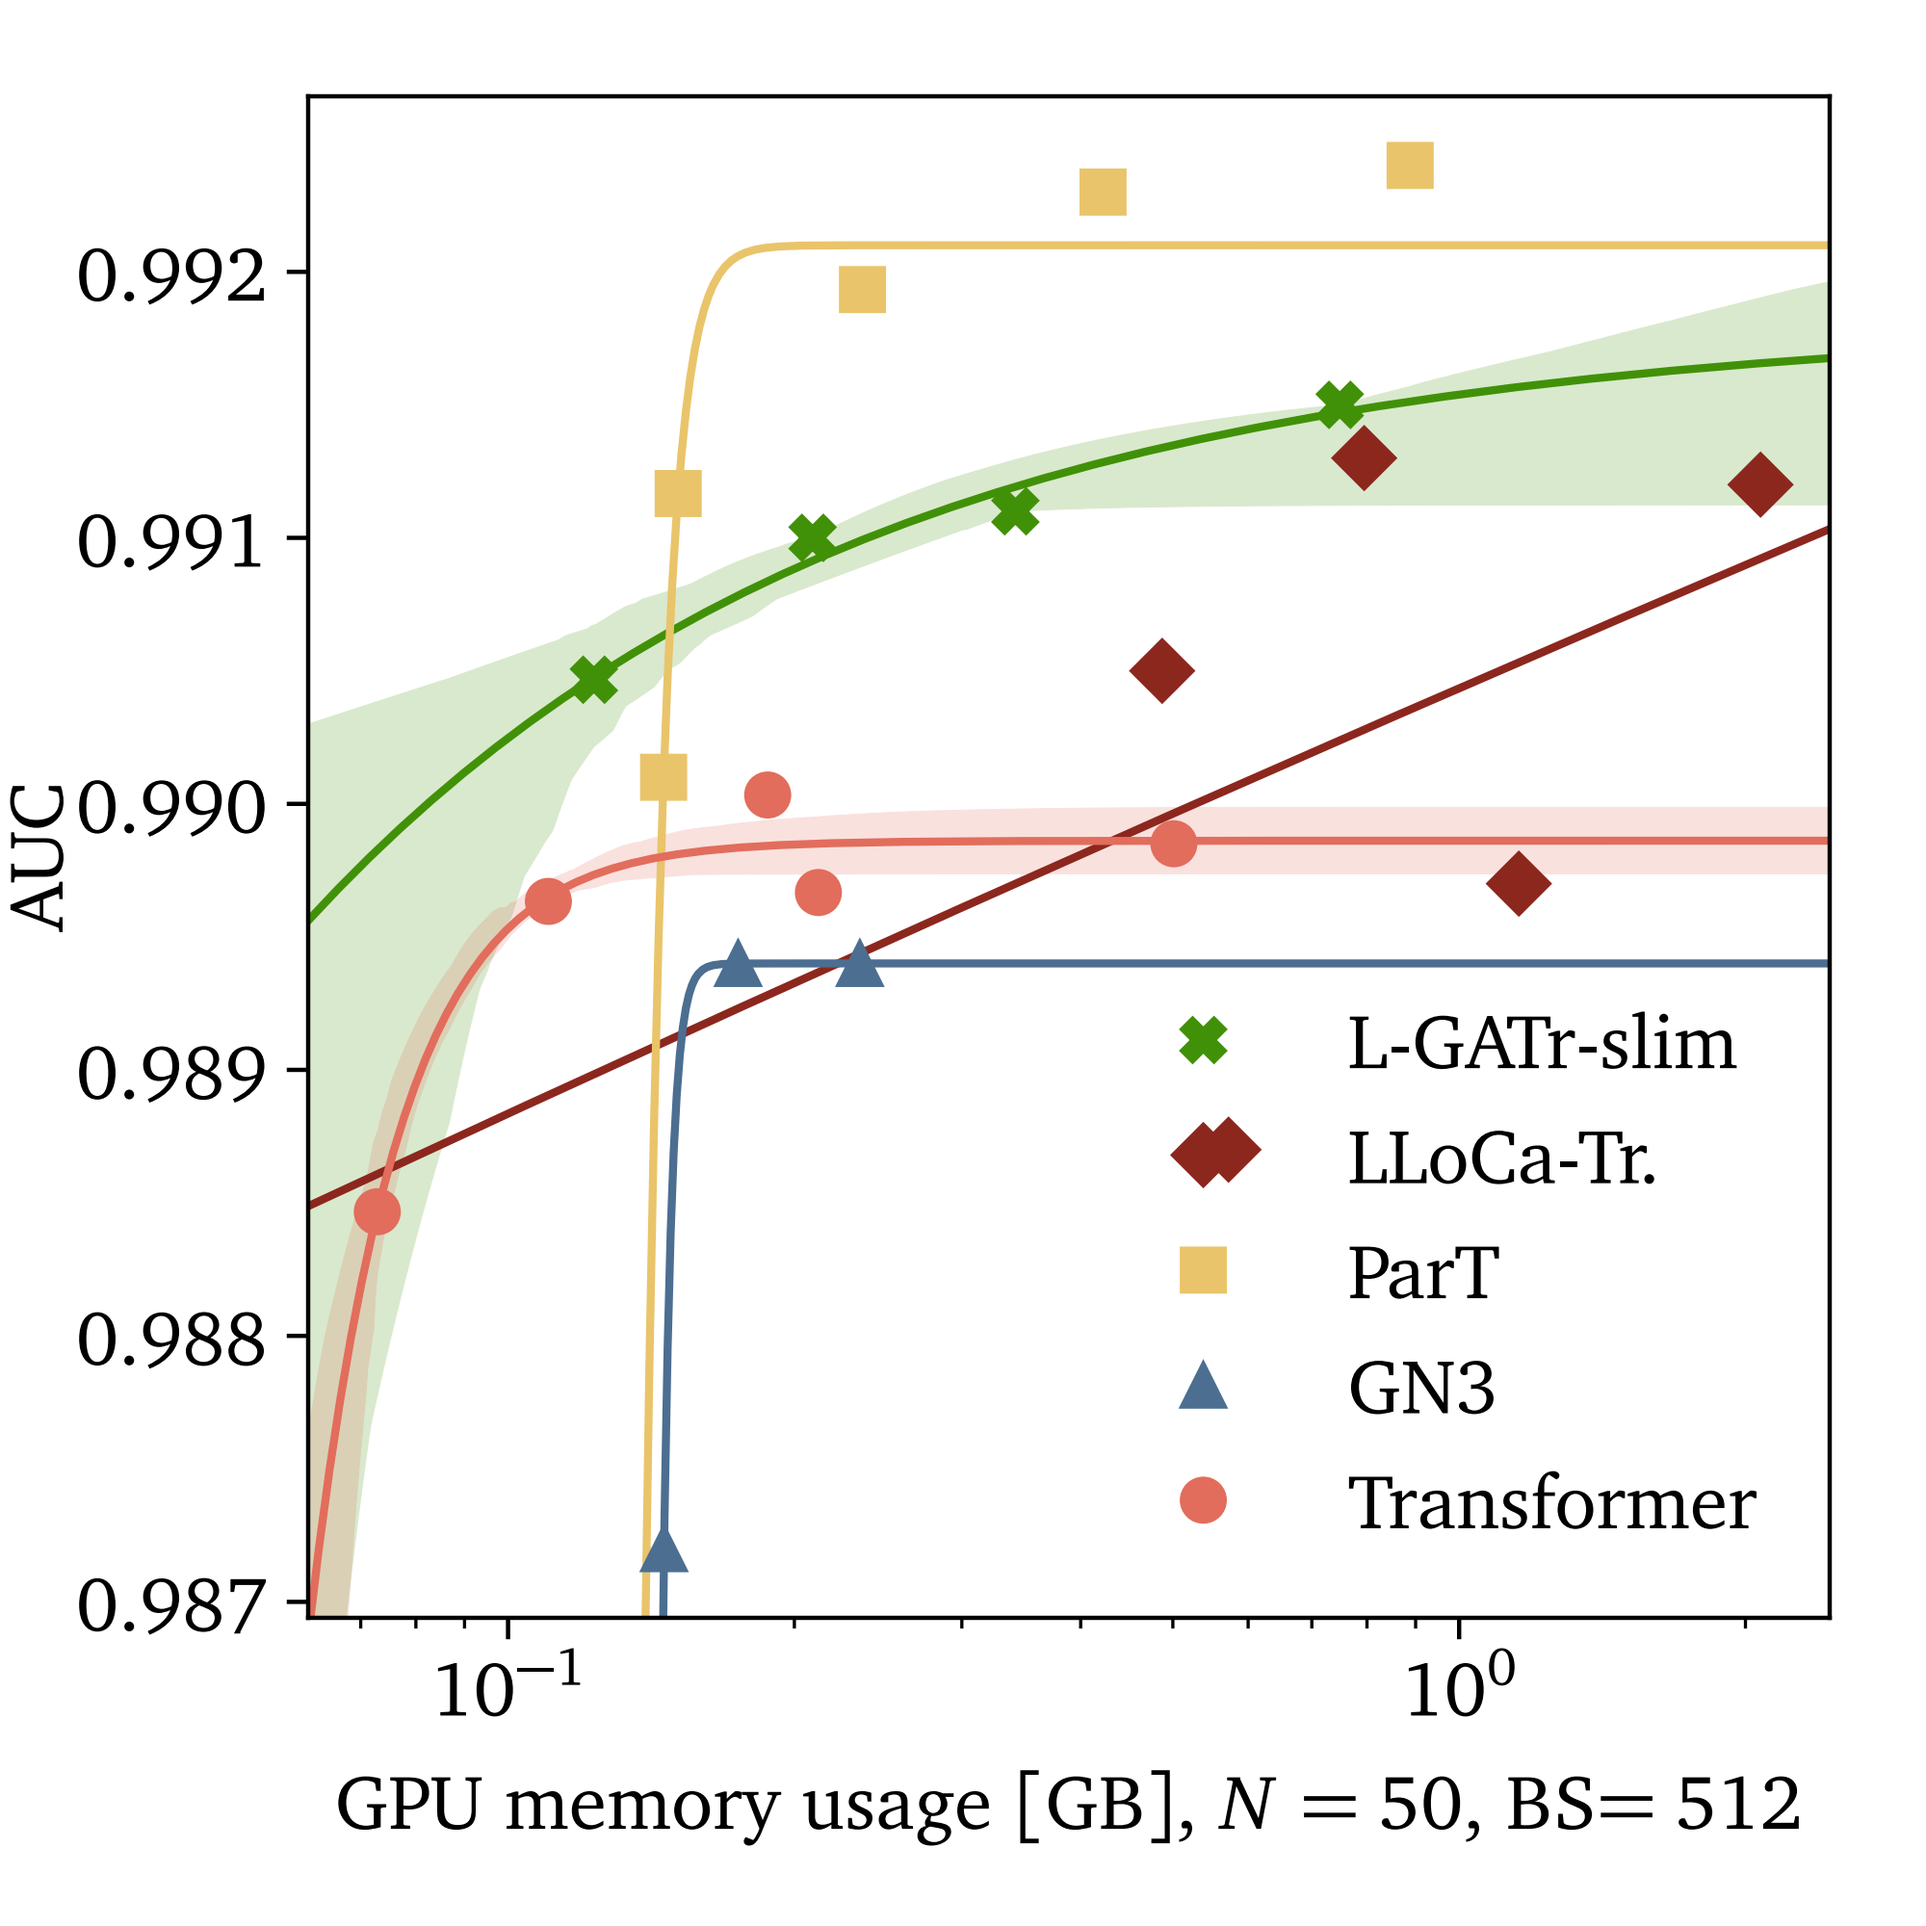
0.992
0.991
0.990
0.989
0.988
0.987
AUC
10^{-1}
10^{0}
GPU memory usage [GB], $N = 50$, BS$= 512$
L-GATr-slim
LLoCa-Tr.
ParT
GN3
Transformer

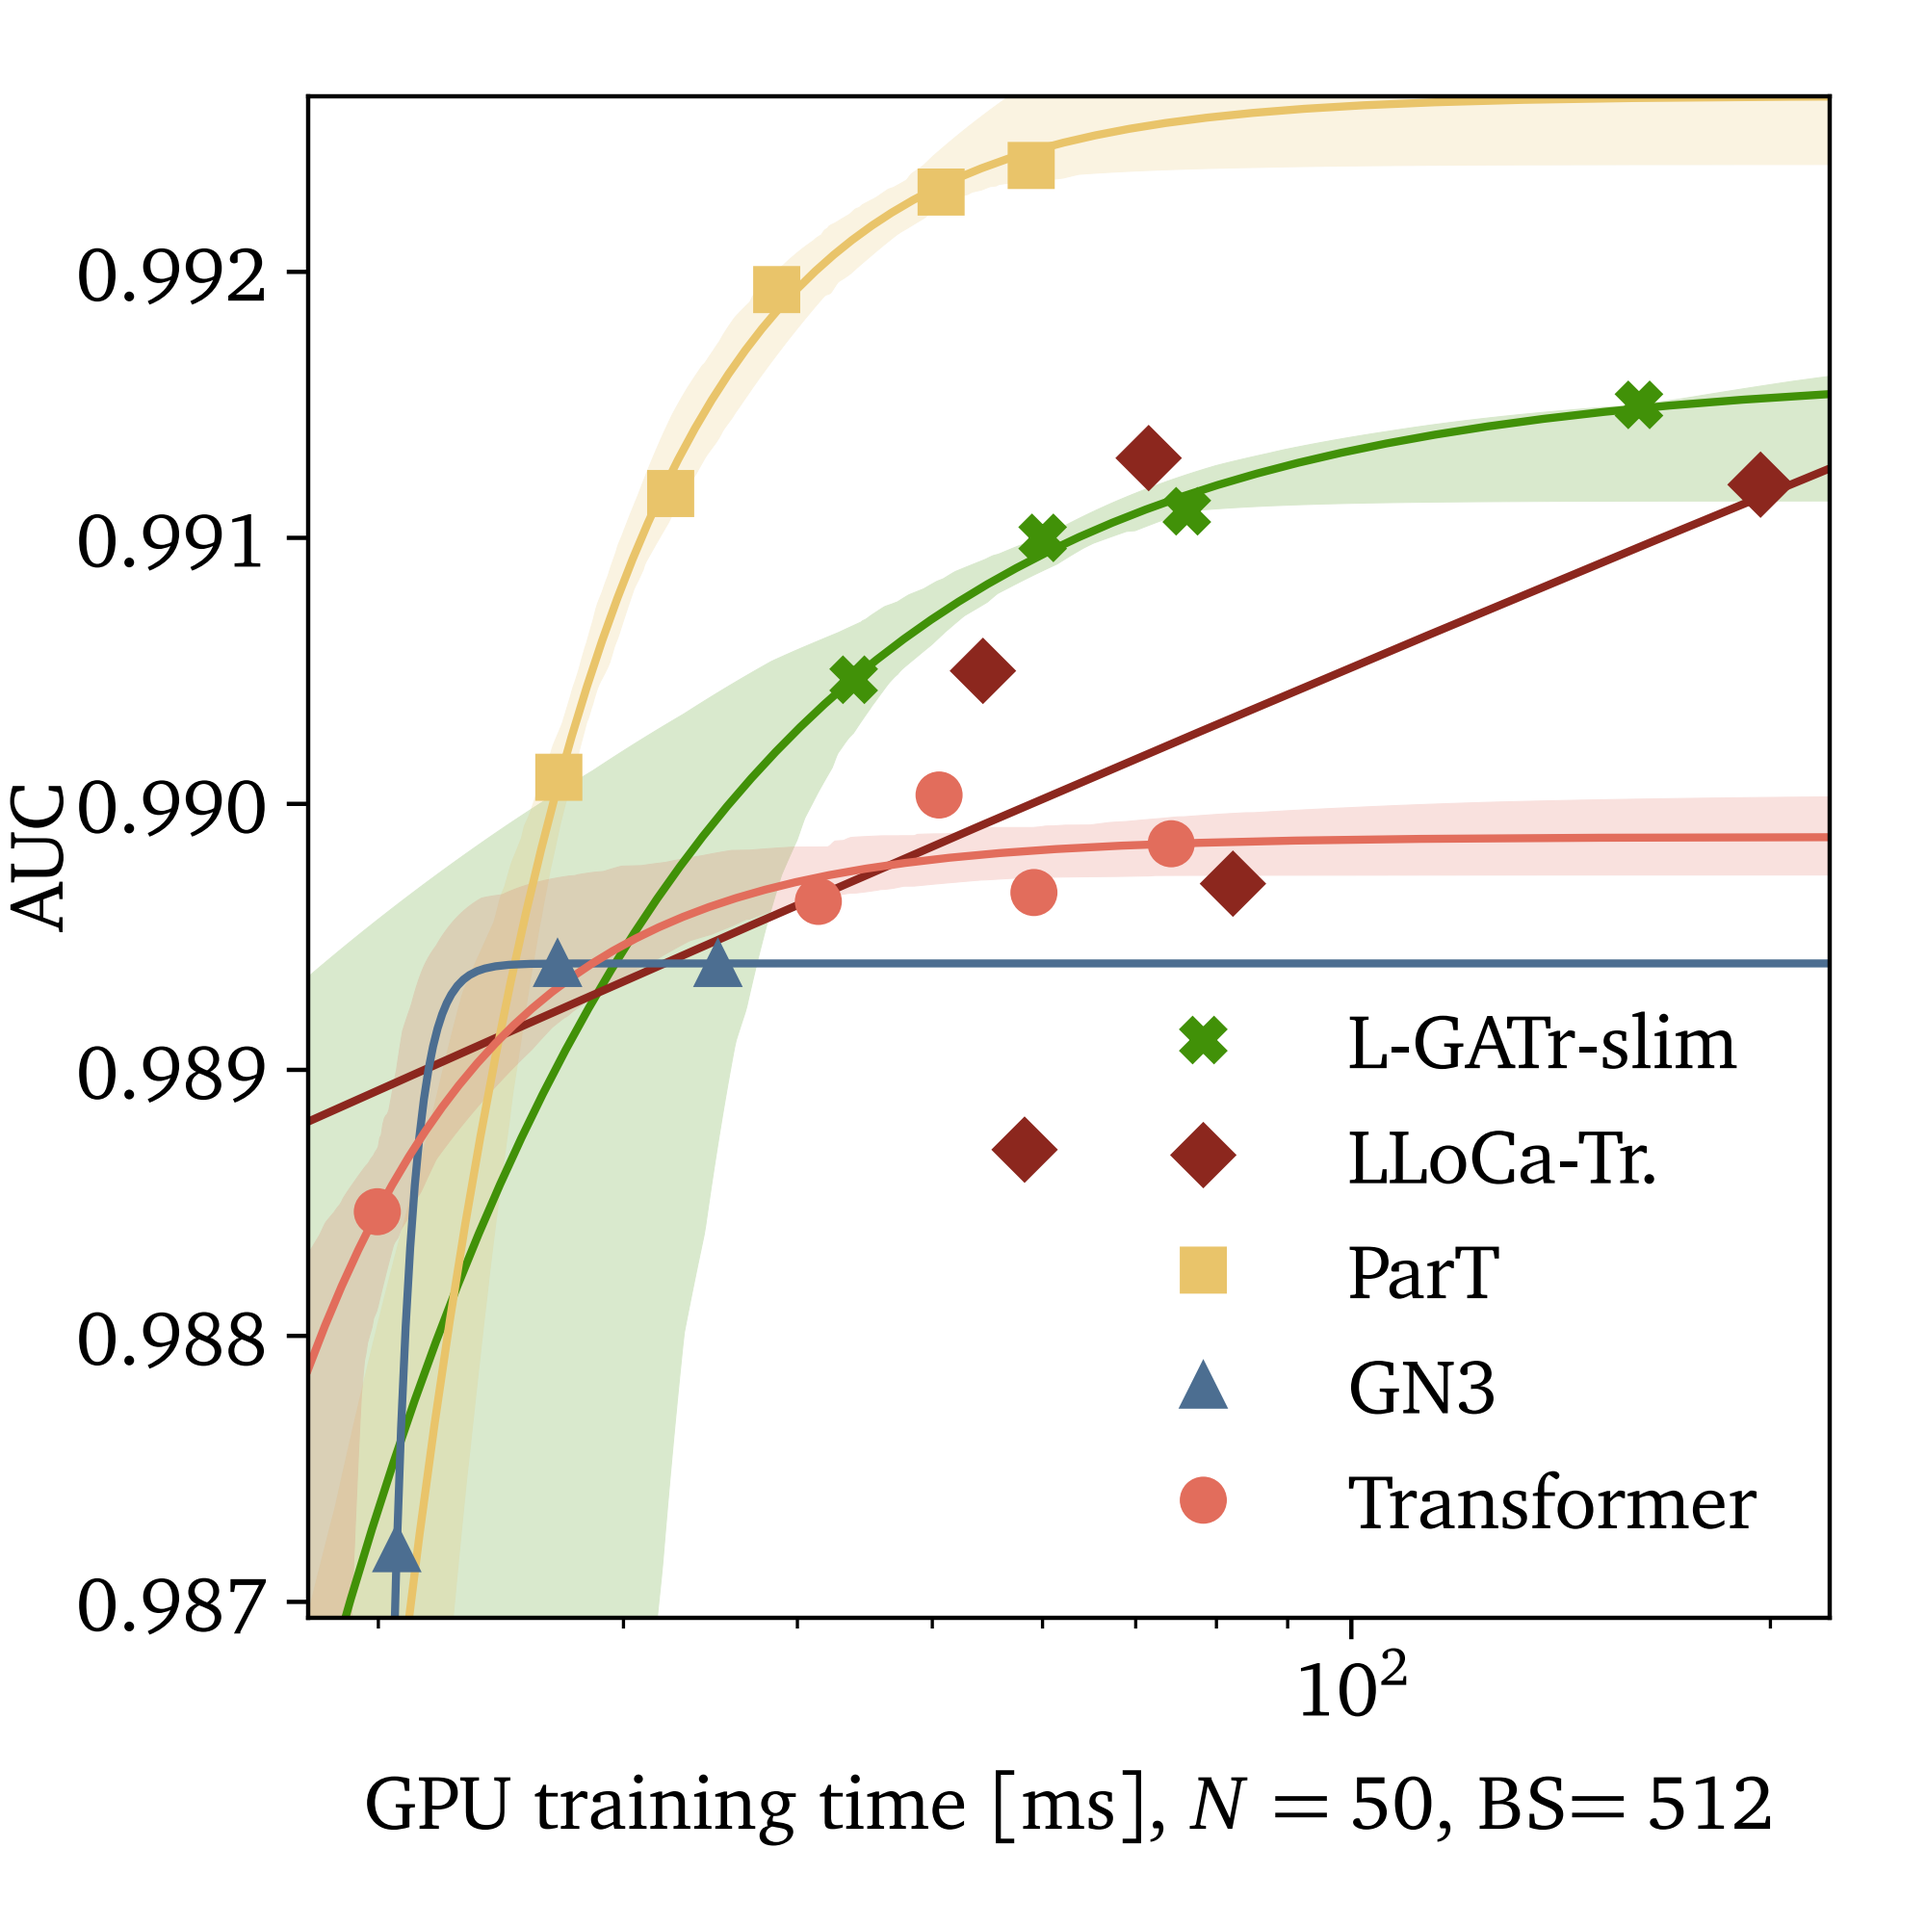
0.992
0.991
0.990
0.989
0.988
0.987
AUC
10^2
GPU training time [ms], $N = 50$, BS$= 512$
L-GATr-slim
LLoCa-Tr.
ParT
GN3
Transformer

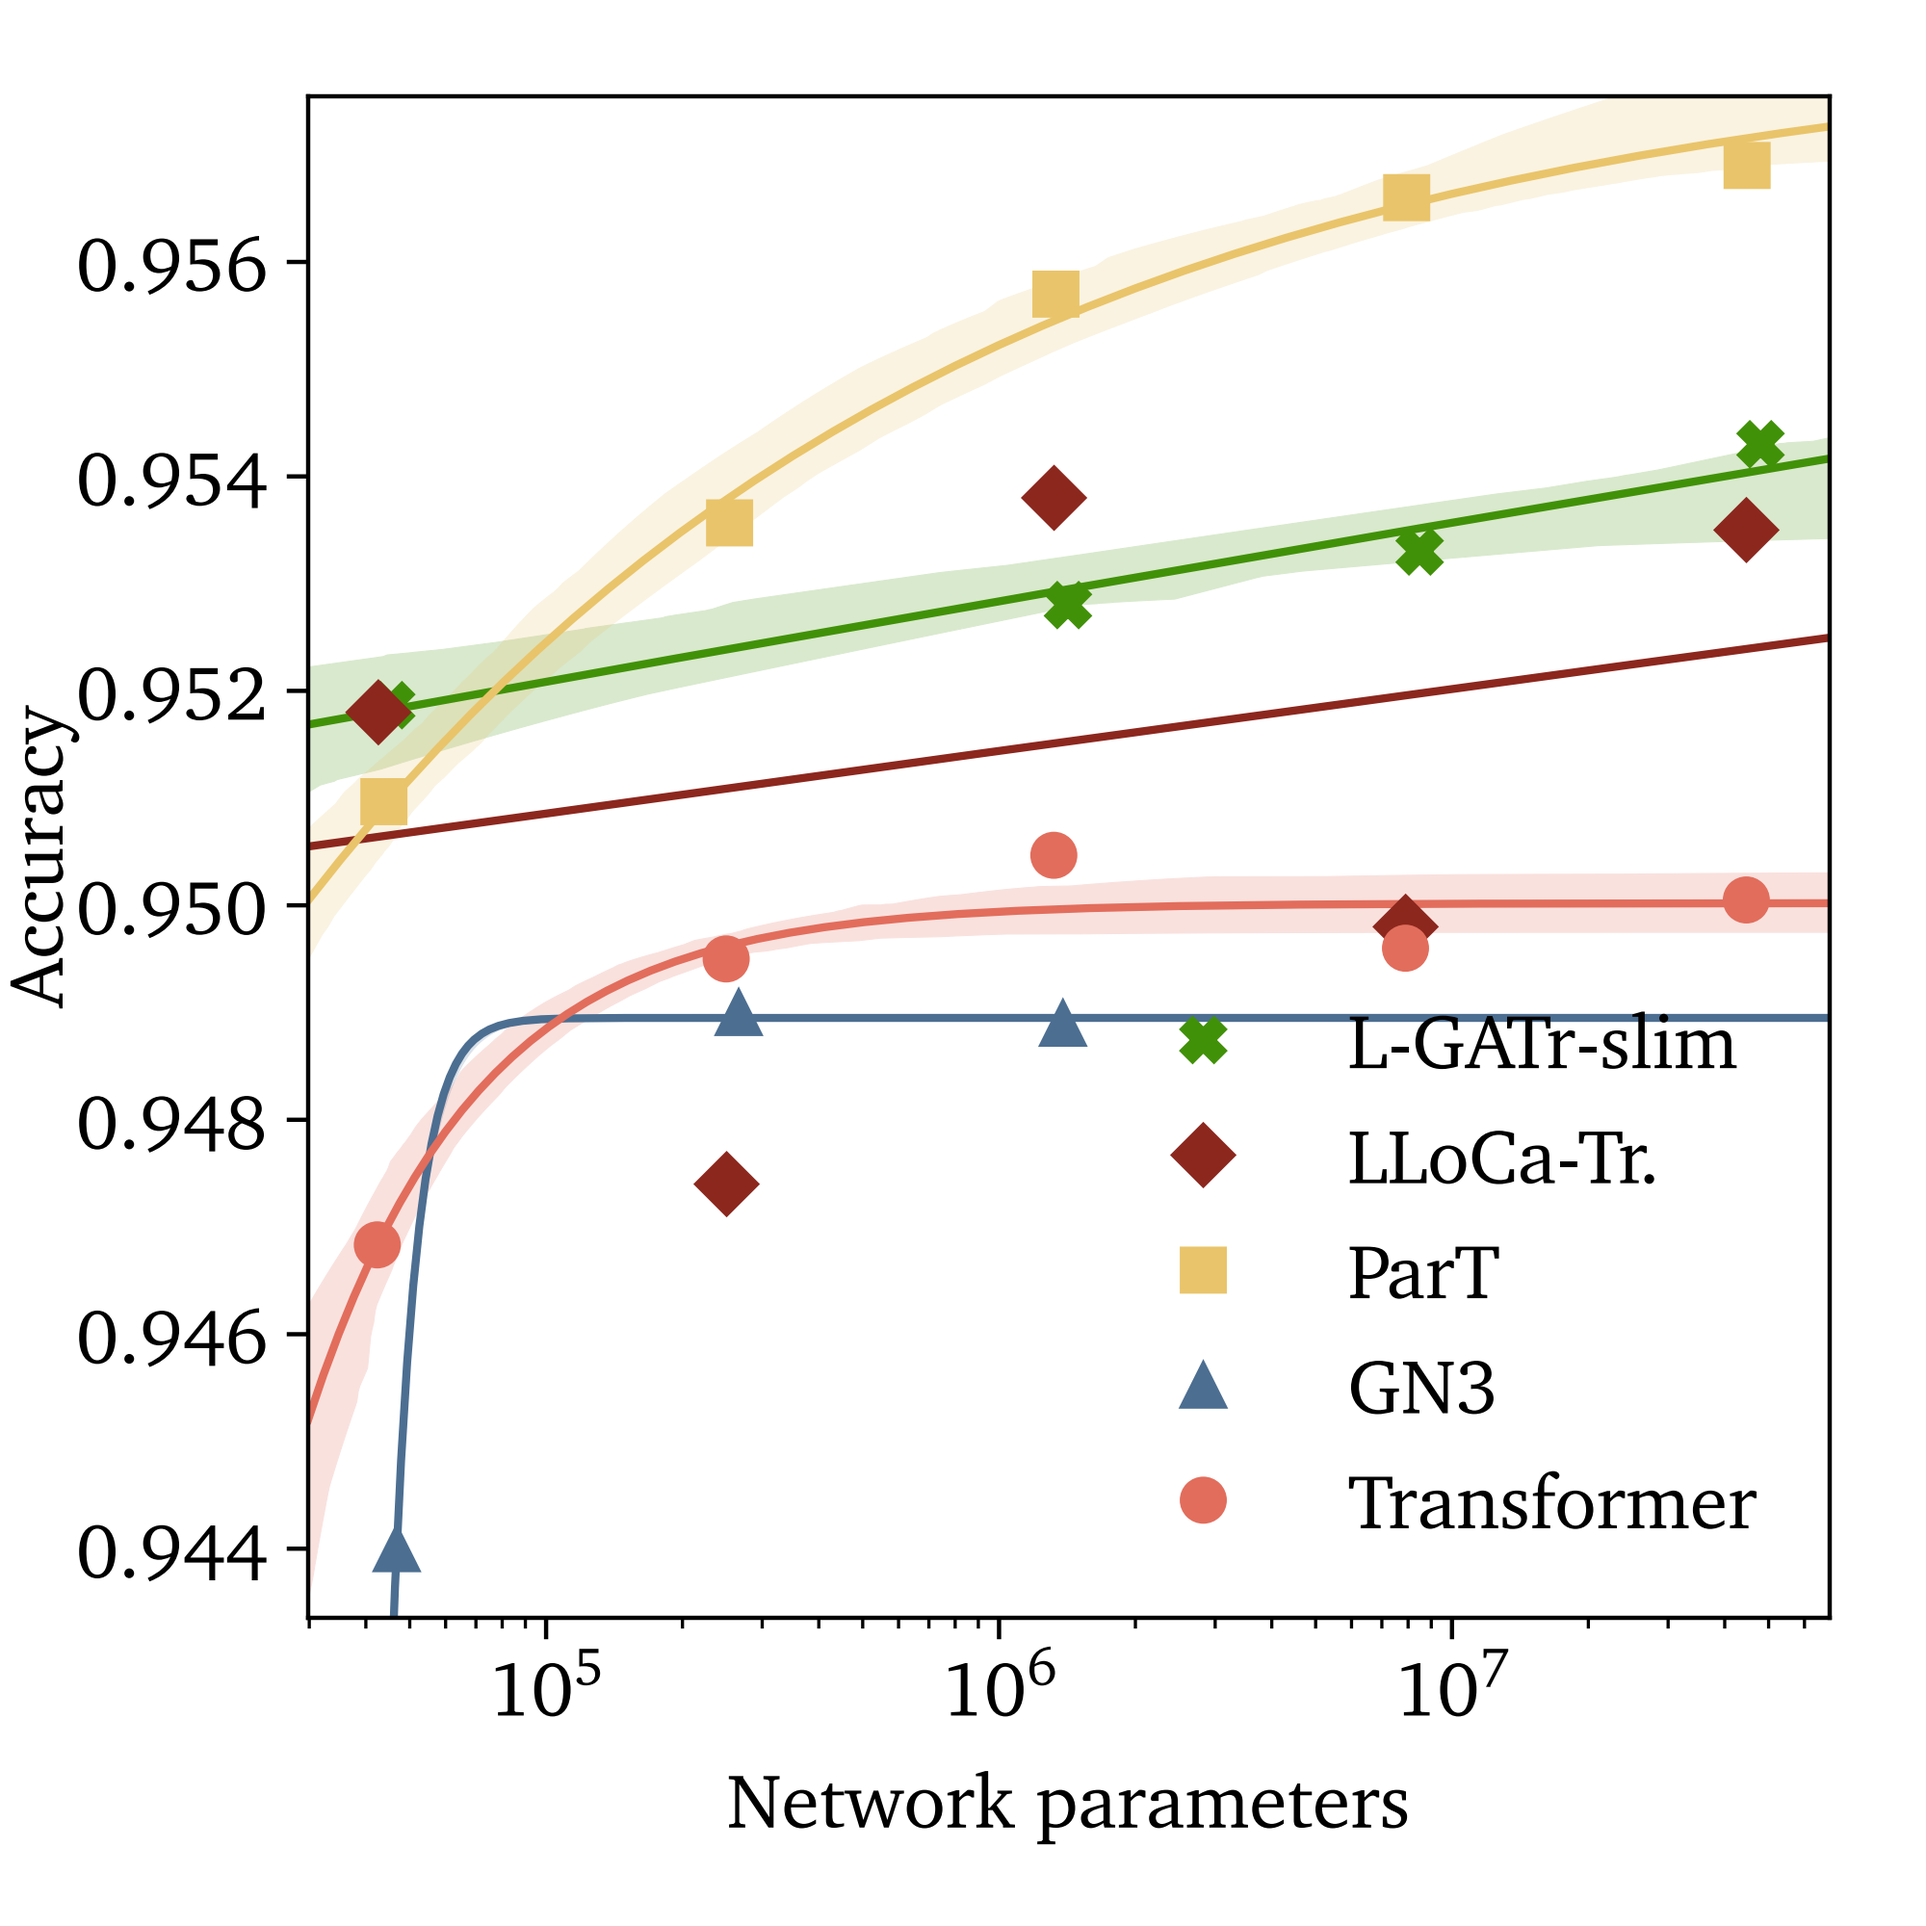
0.956
0.954
0.952
0.950
0.948
0.946
0.944
Accuracy
10^5
10^6
10^7
Network parameters
L-GATr-slim
LLoCa-Tr.
ParT
GN3
Transformer

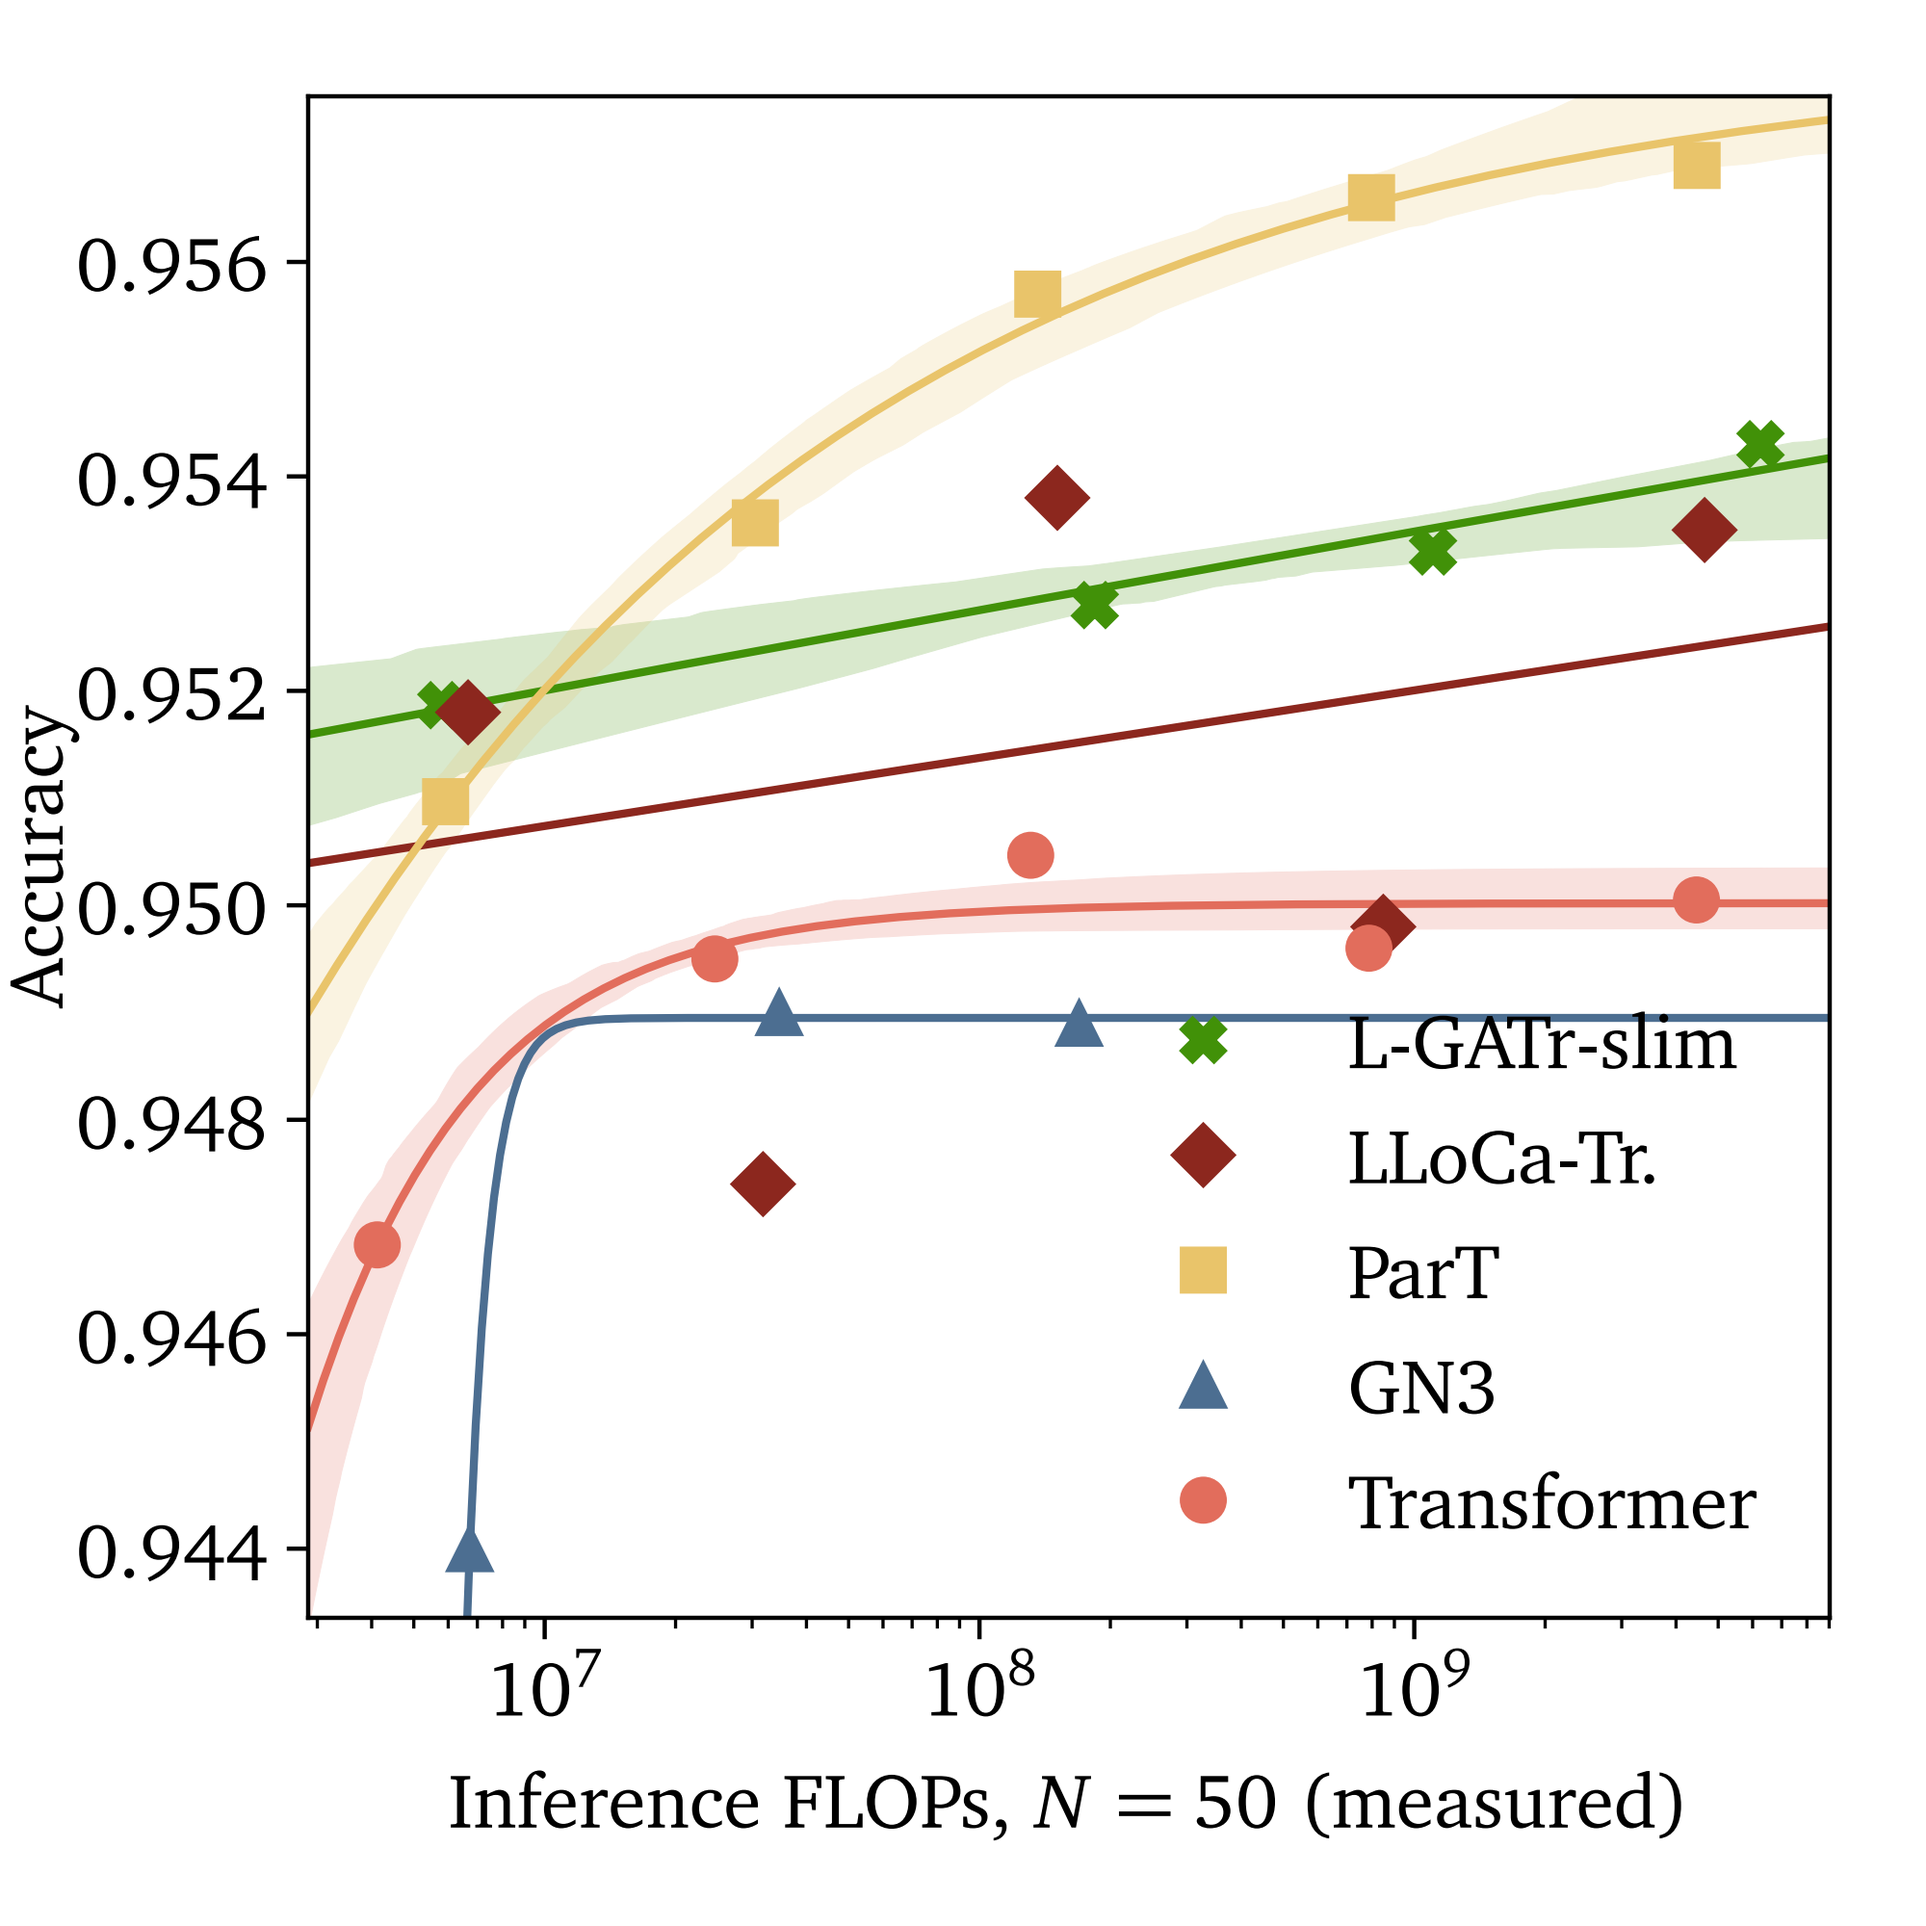
0.956
0.954
0.952
0.950
0.948
0.946
0.944
Accuracy
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (measured)
L-GATr-slim
LLoCa-Tr.
ParT
GN3
Transformer

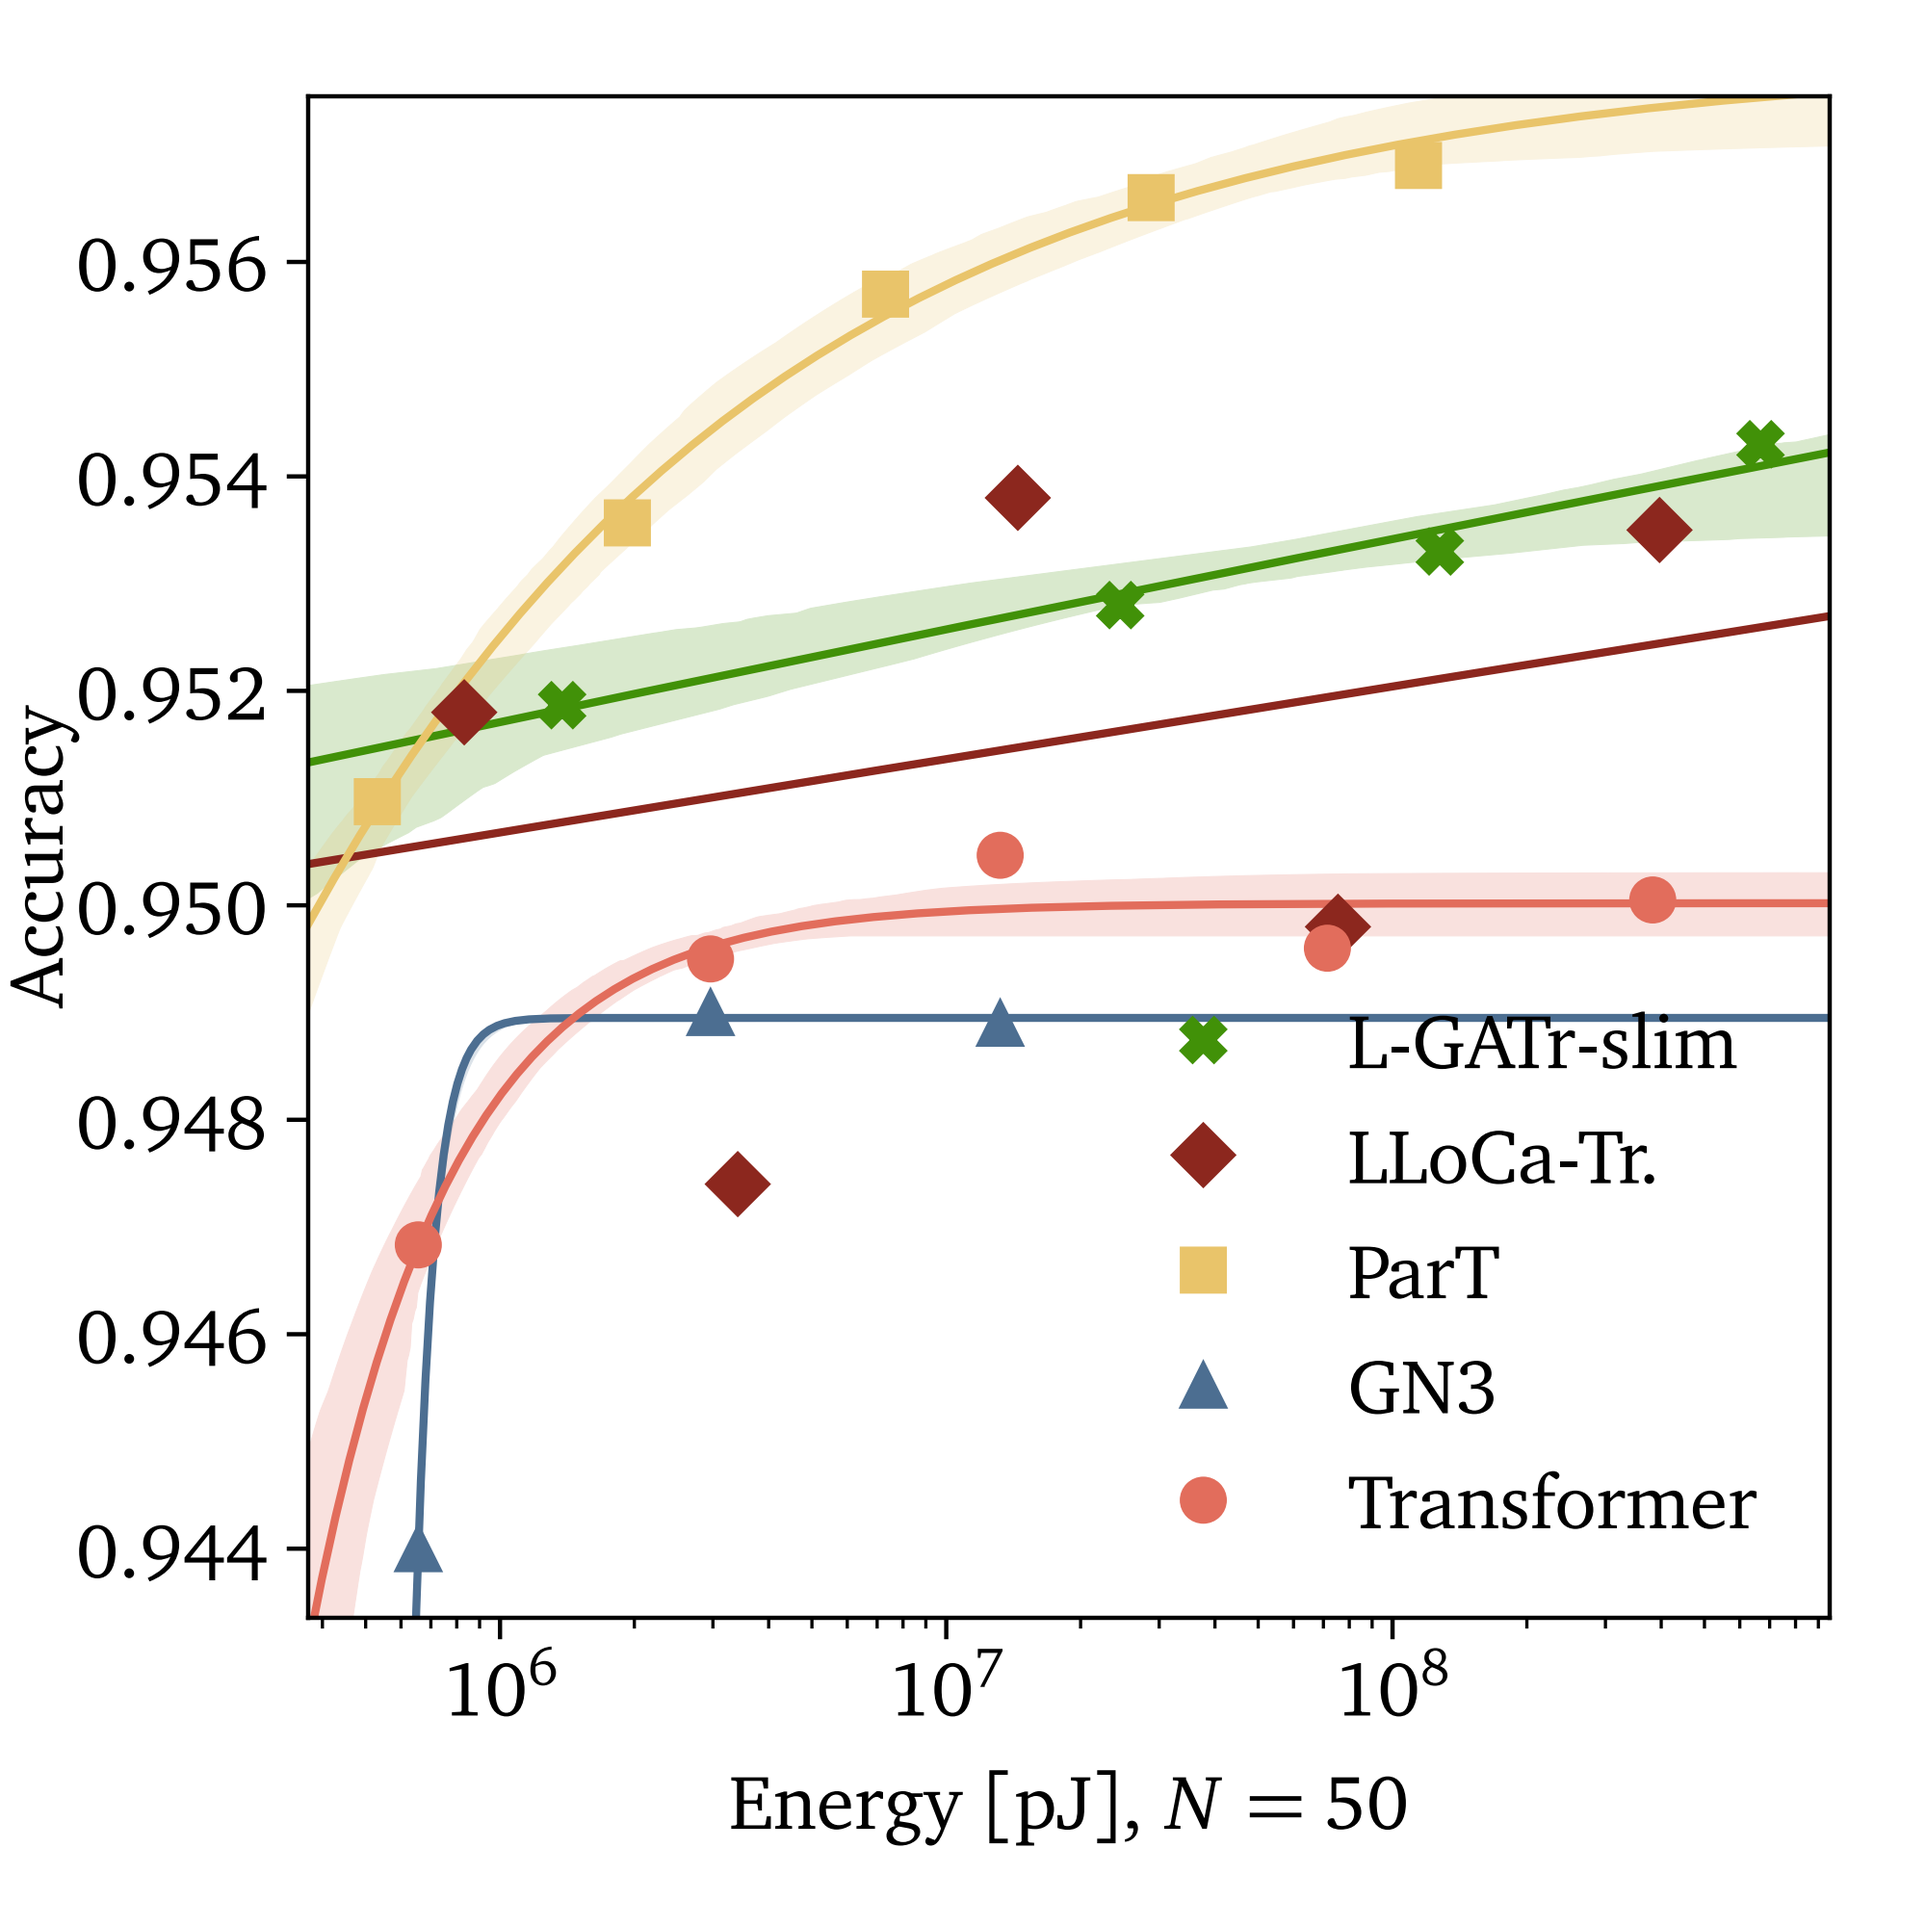
0.956
0.954
0.952
0.950
0.948
0.946
0.944
Accuracy
10^6
10^7
10^8
Energy [pJ], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
GN3
Transformer

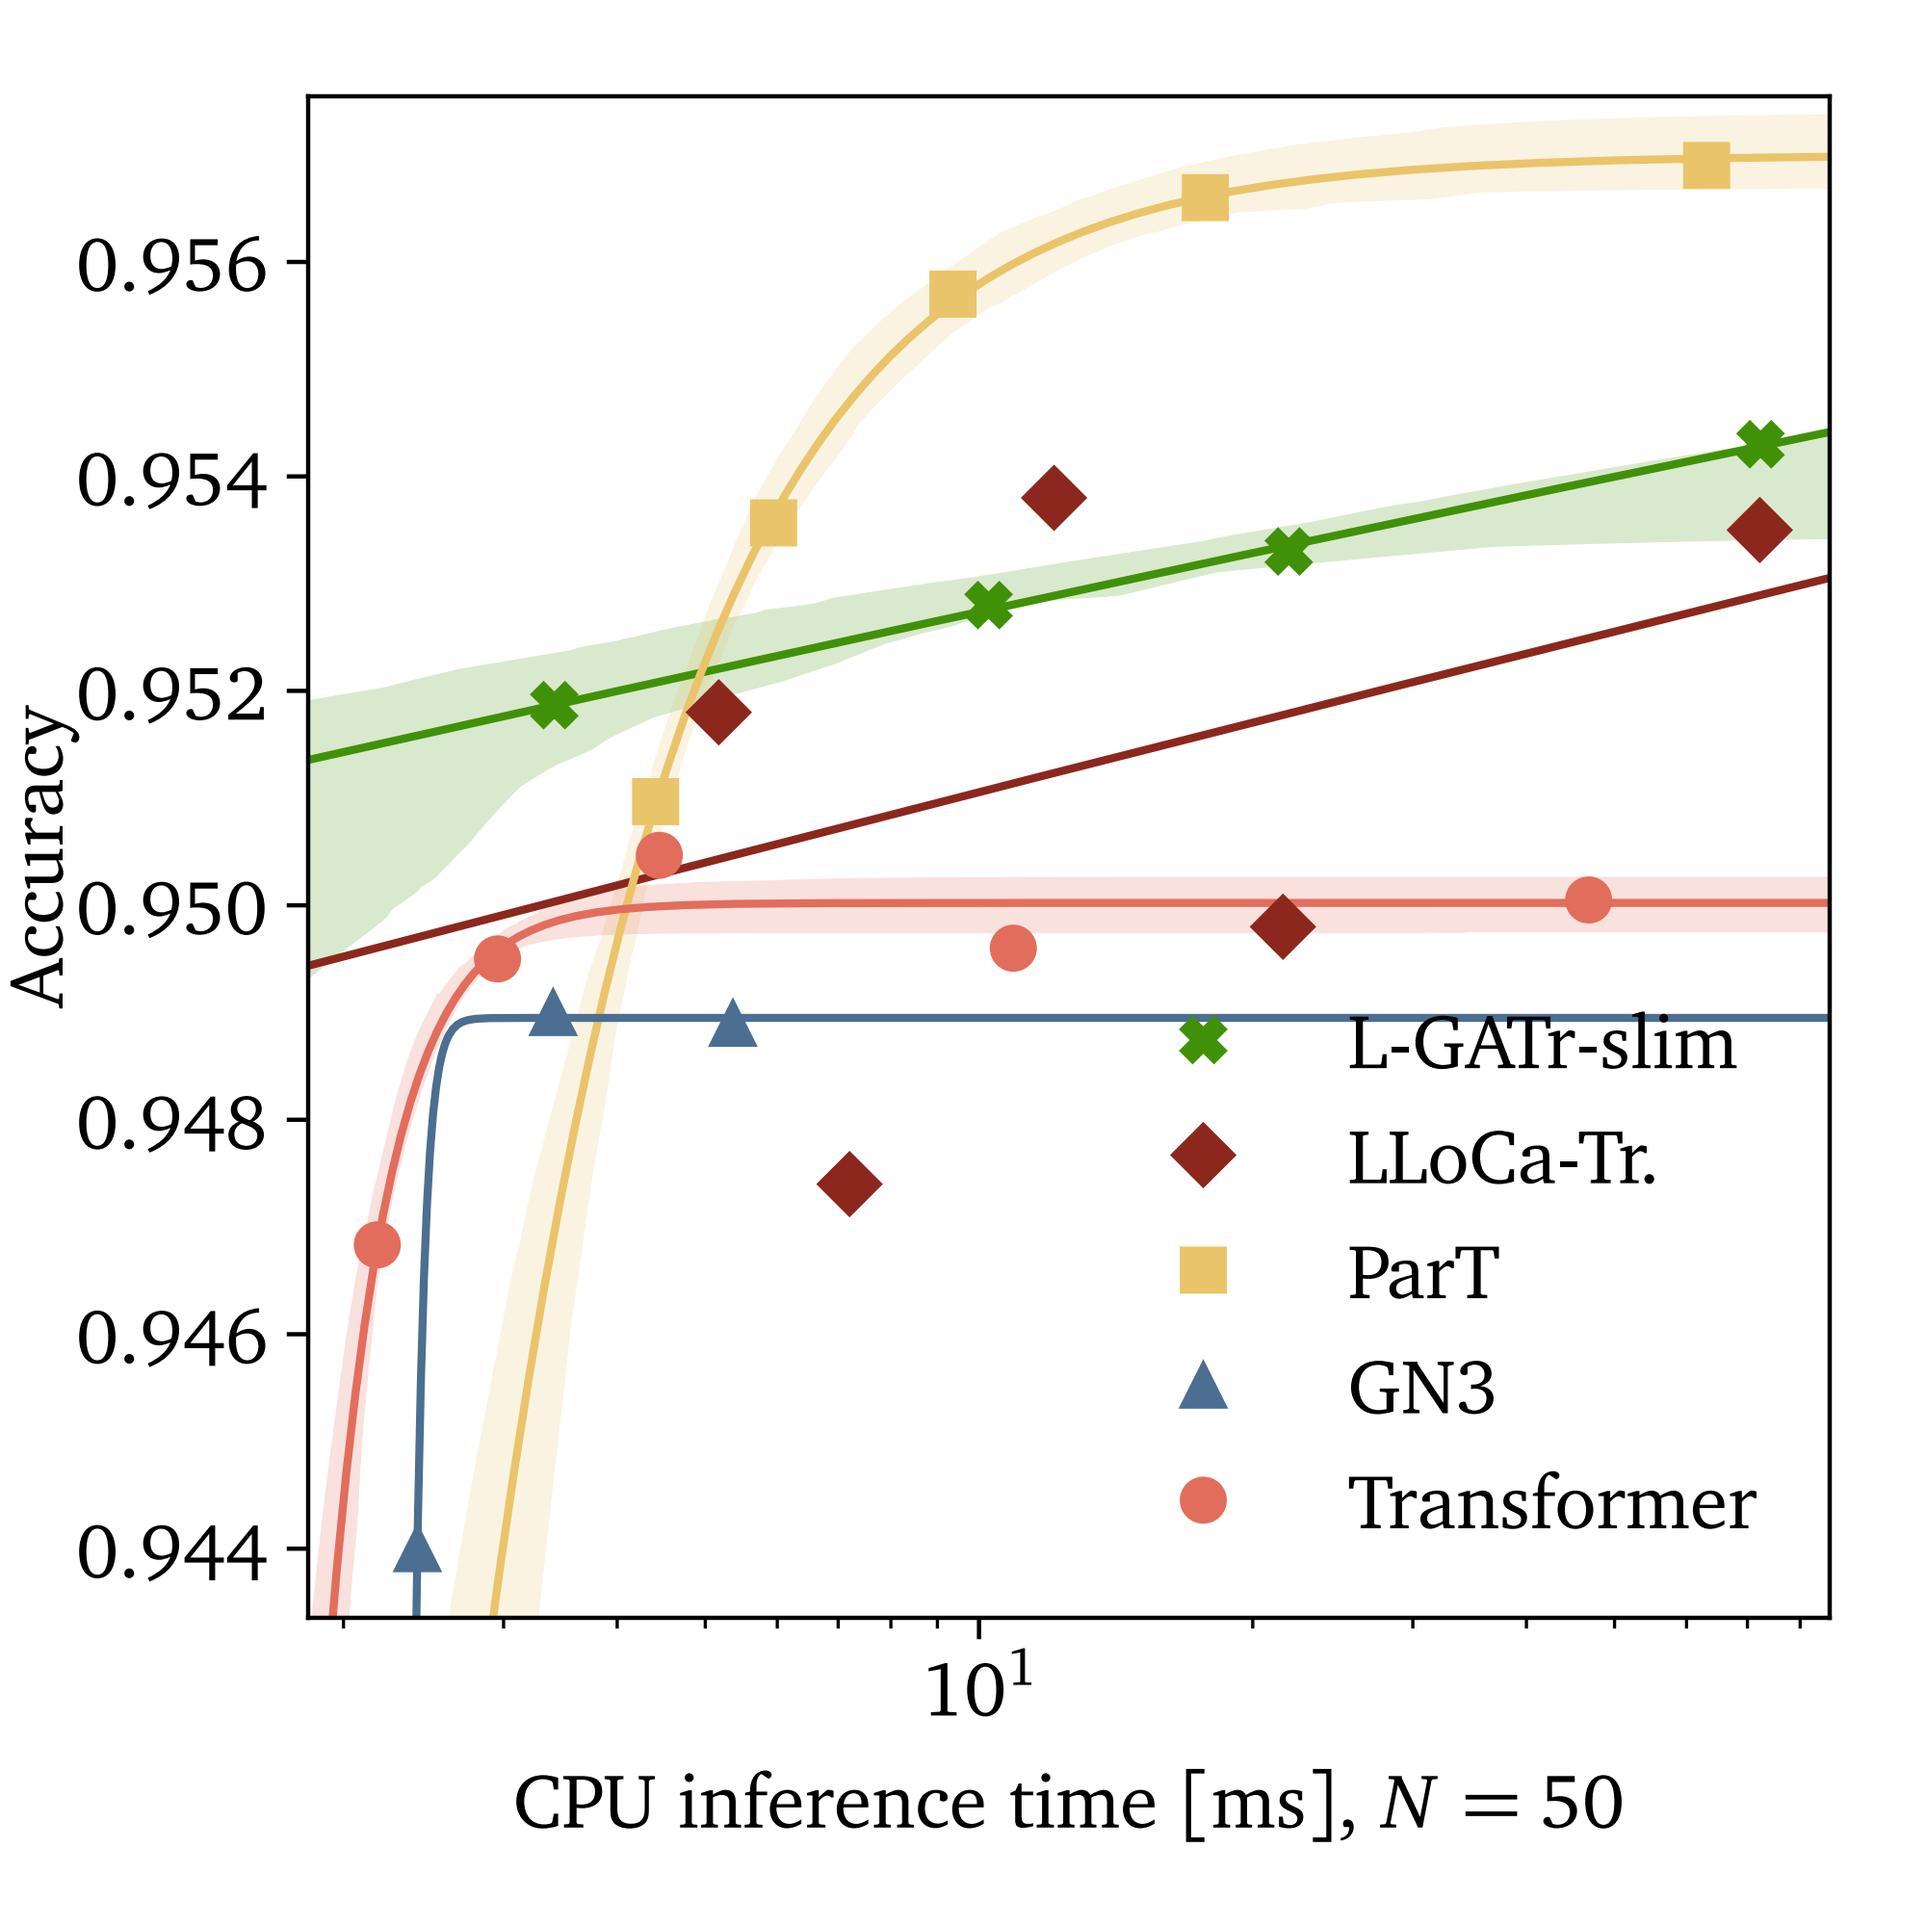
0.956
0.954
0.952
0.950
0.948
0.946
0.944
Accuracy
10^1
CPU inference time [ms], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
GN3
Transformer

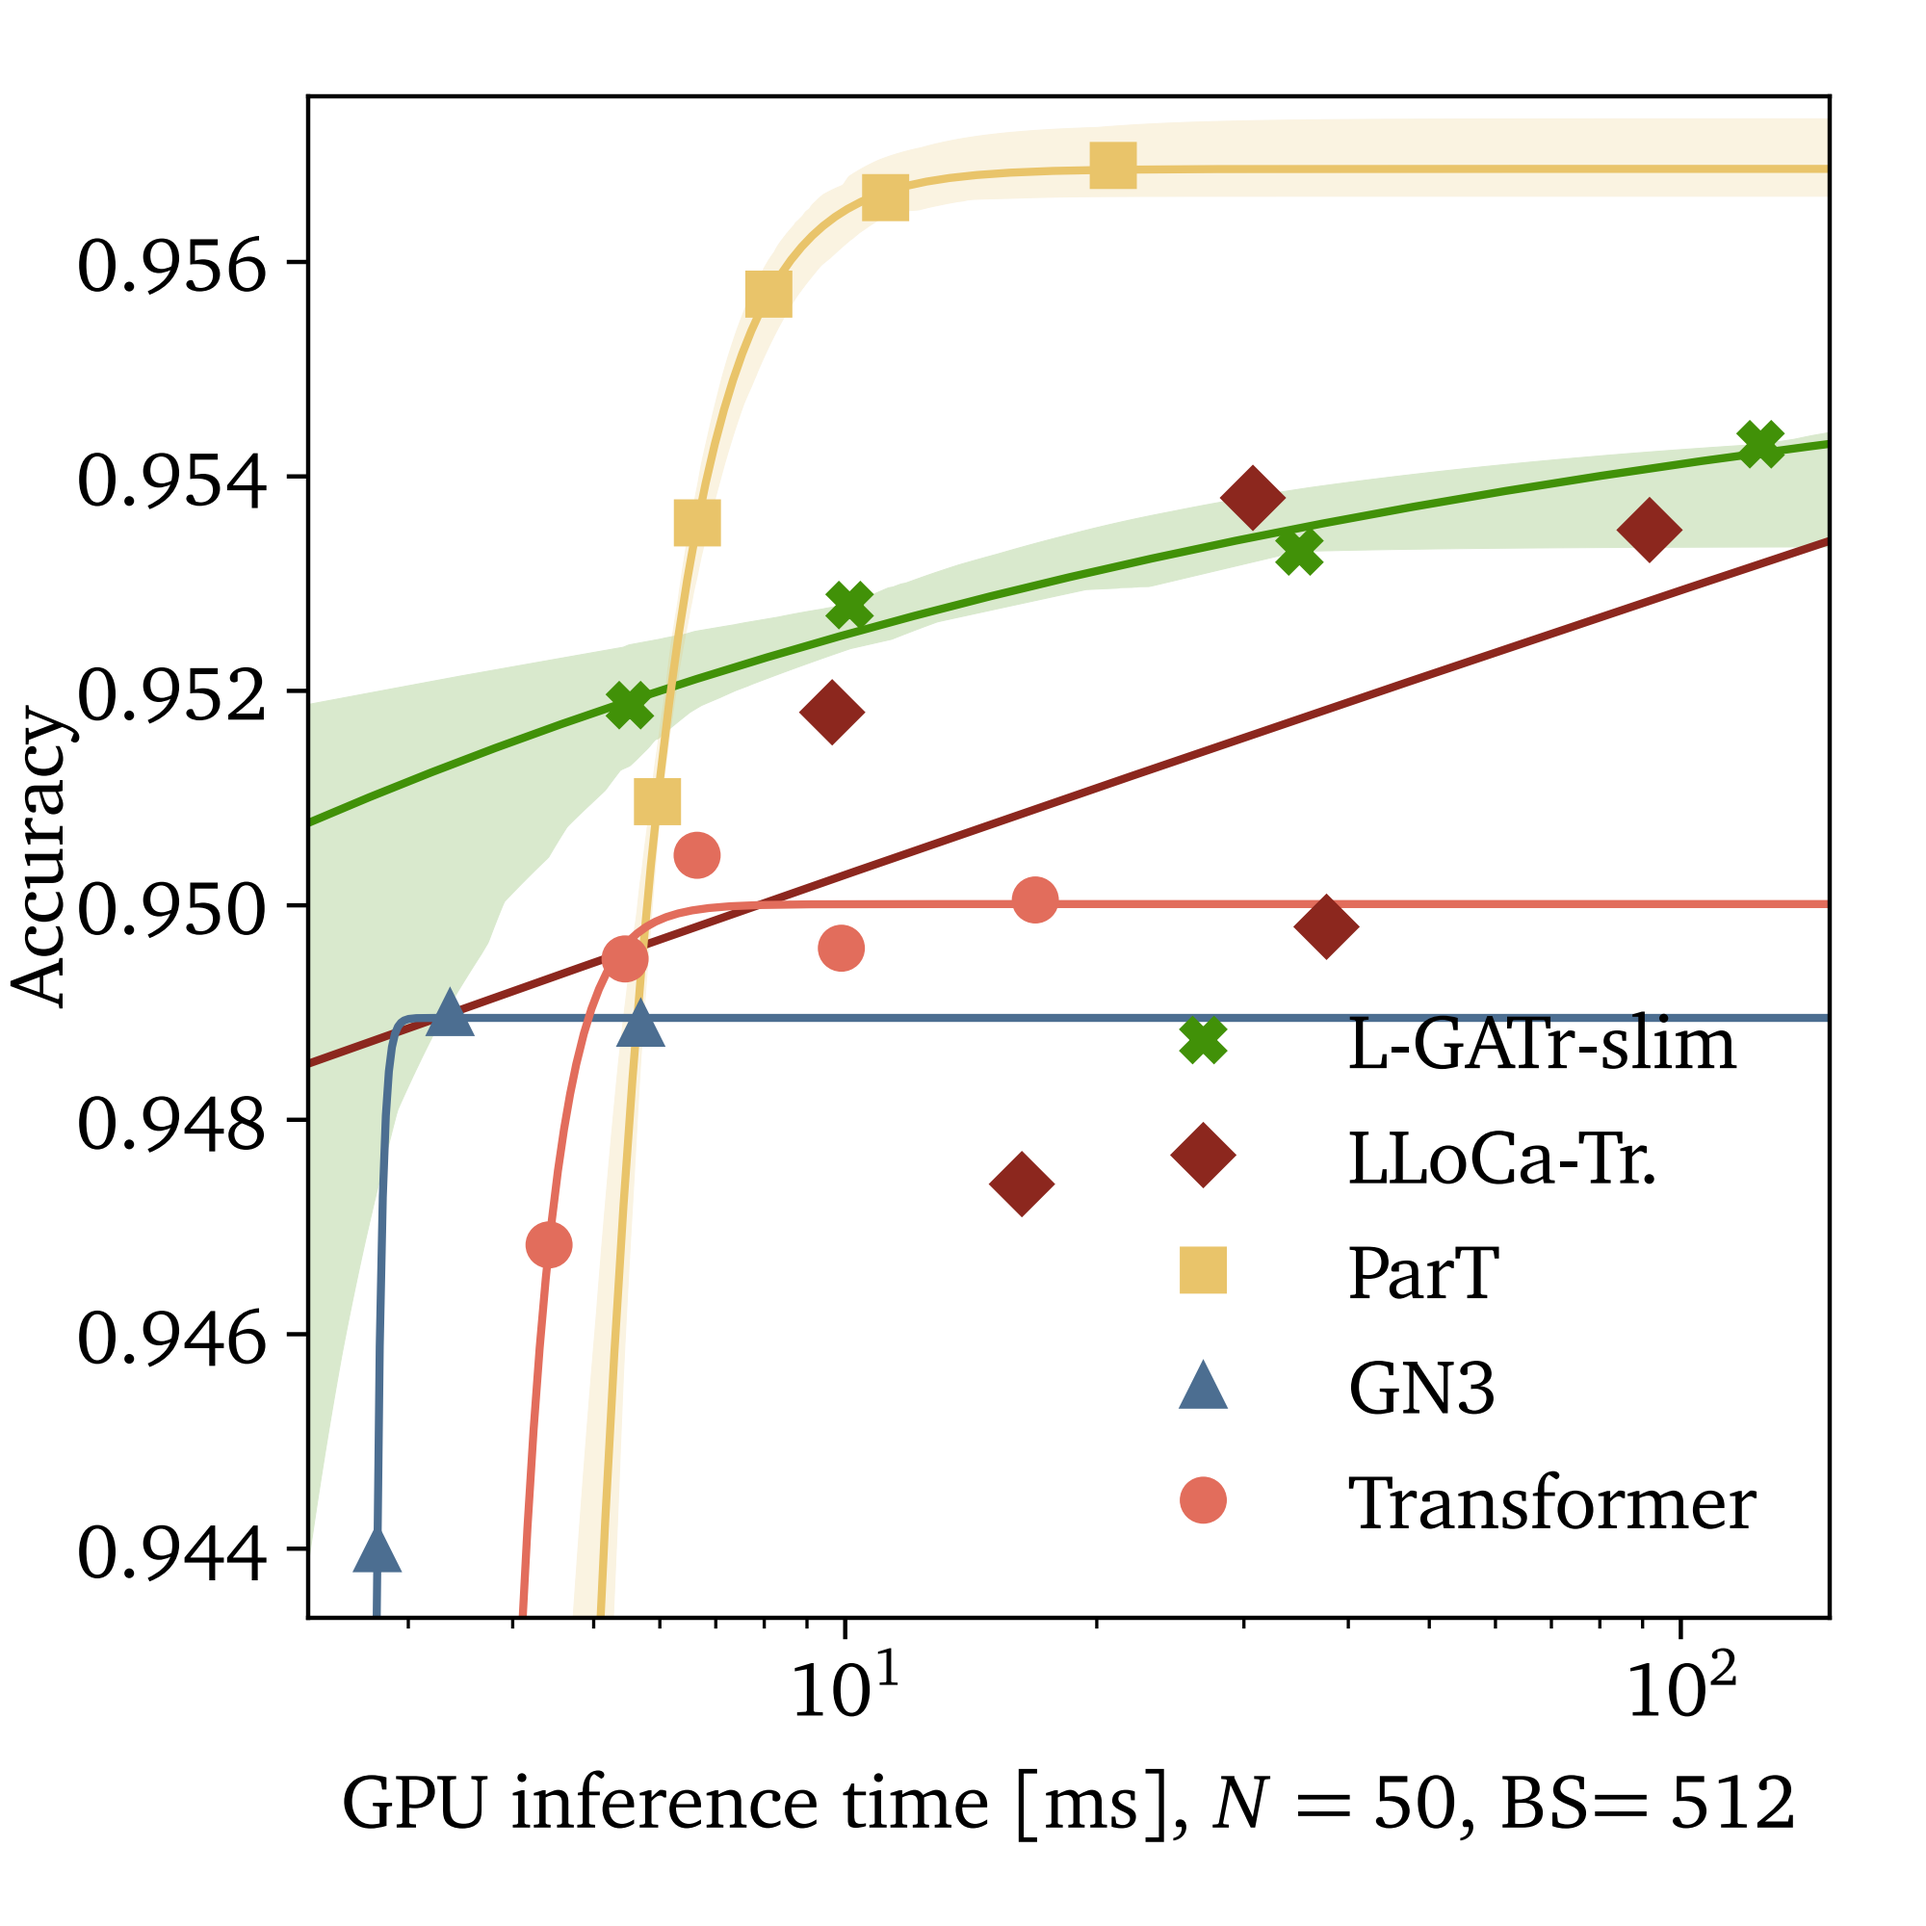
0.956
0.954
0.952
0.950
0.948
0.946
0.944
Accuracy
10^1
10^2
GPU inference time [ms], $N = 50$, BS= 512
L-GATr-slim
LLoCa-Tr.
ParT
GN3
Transformer

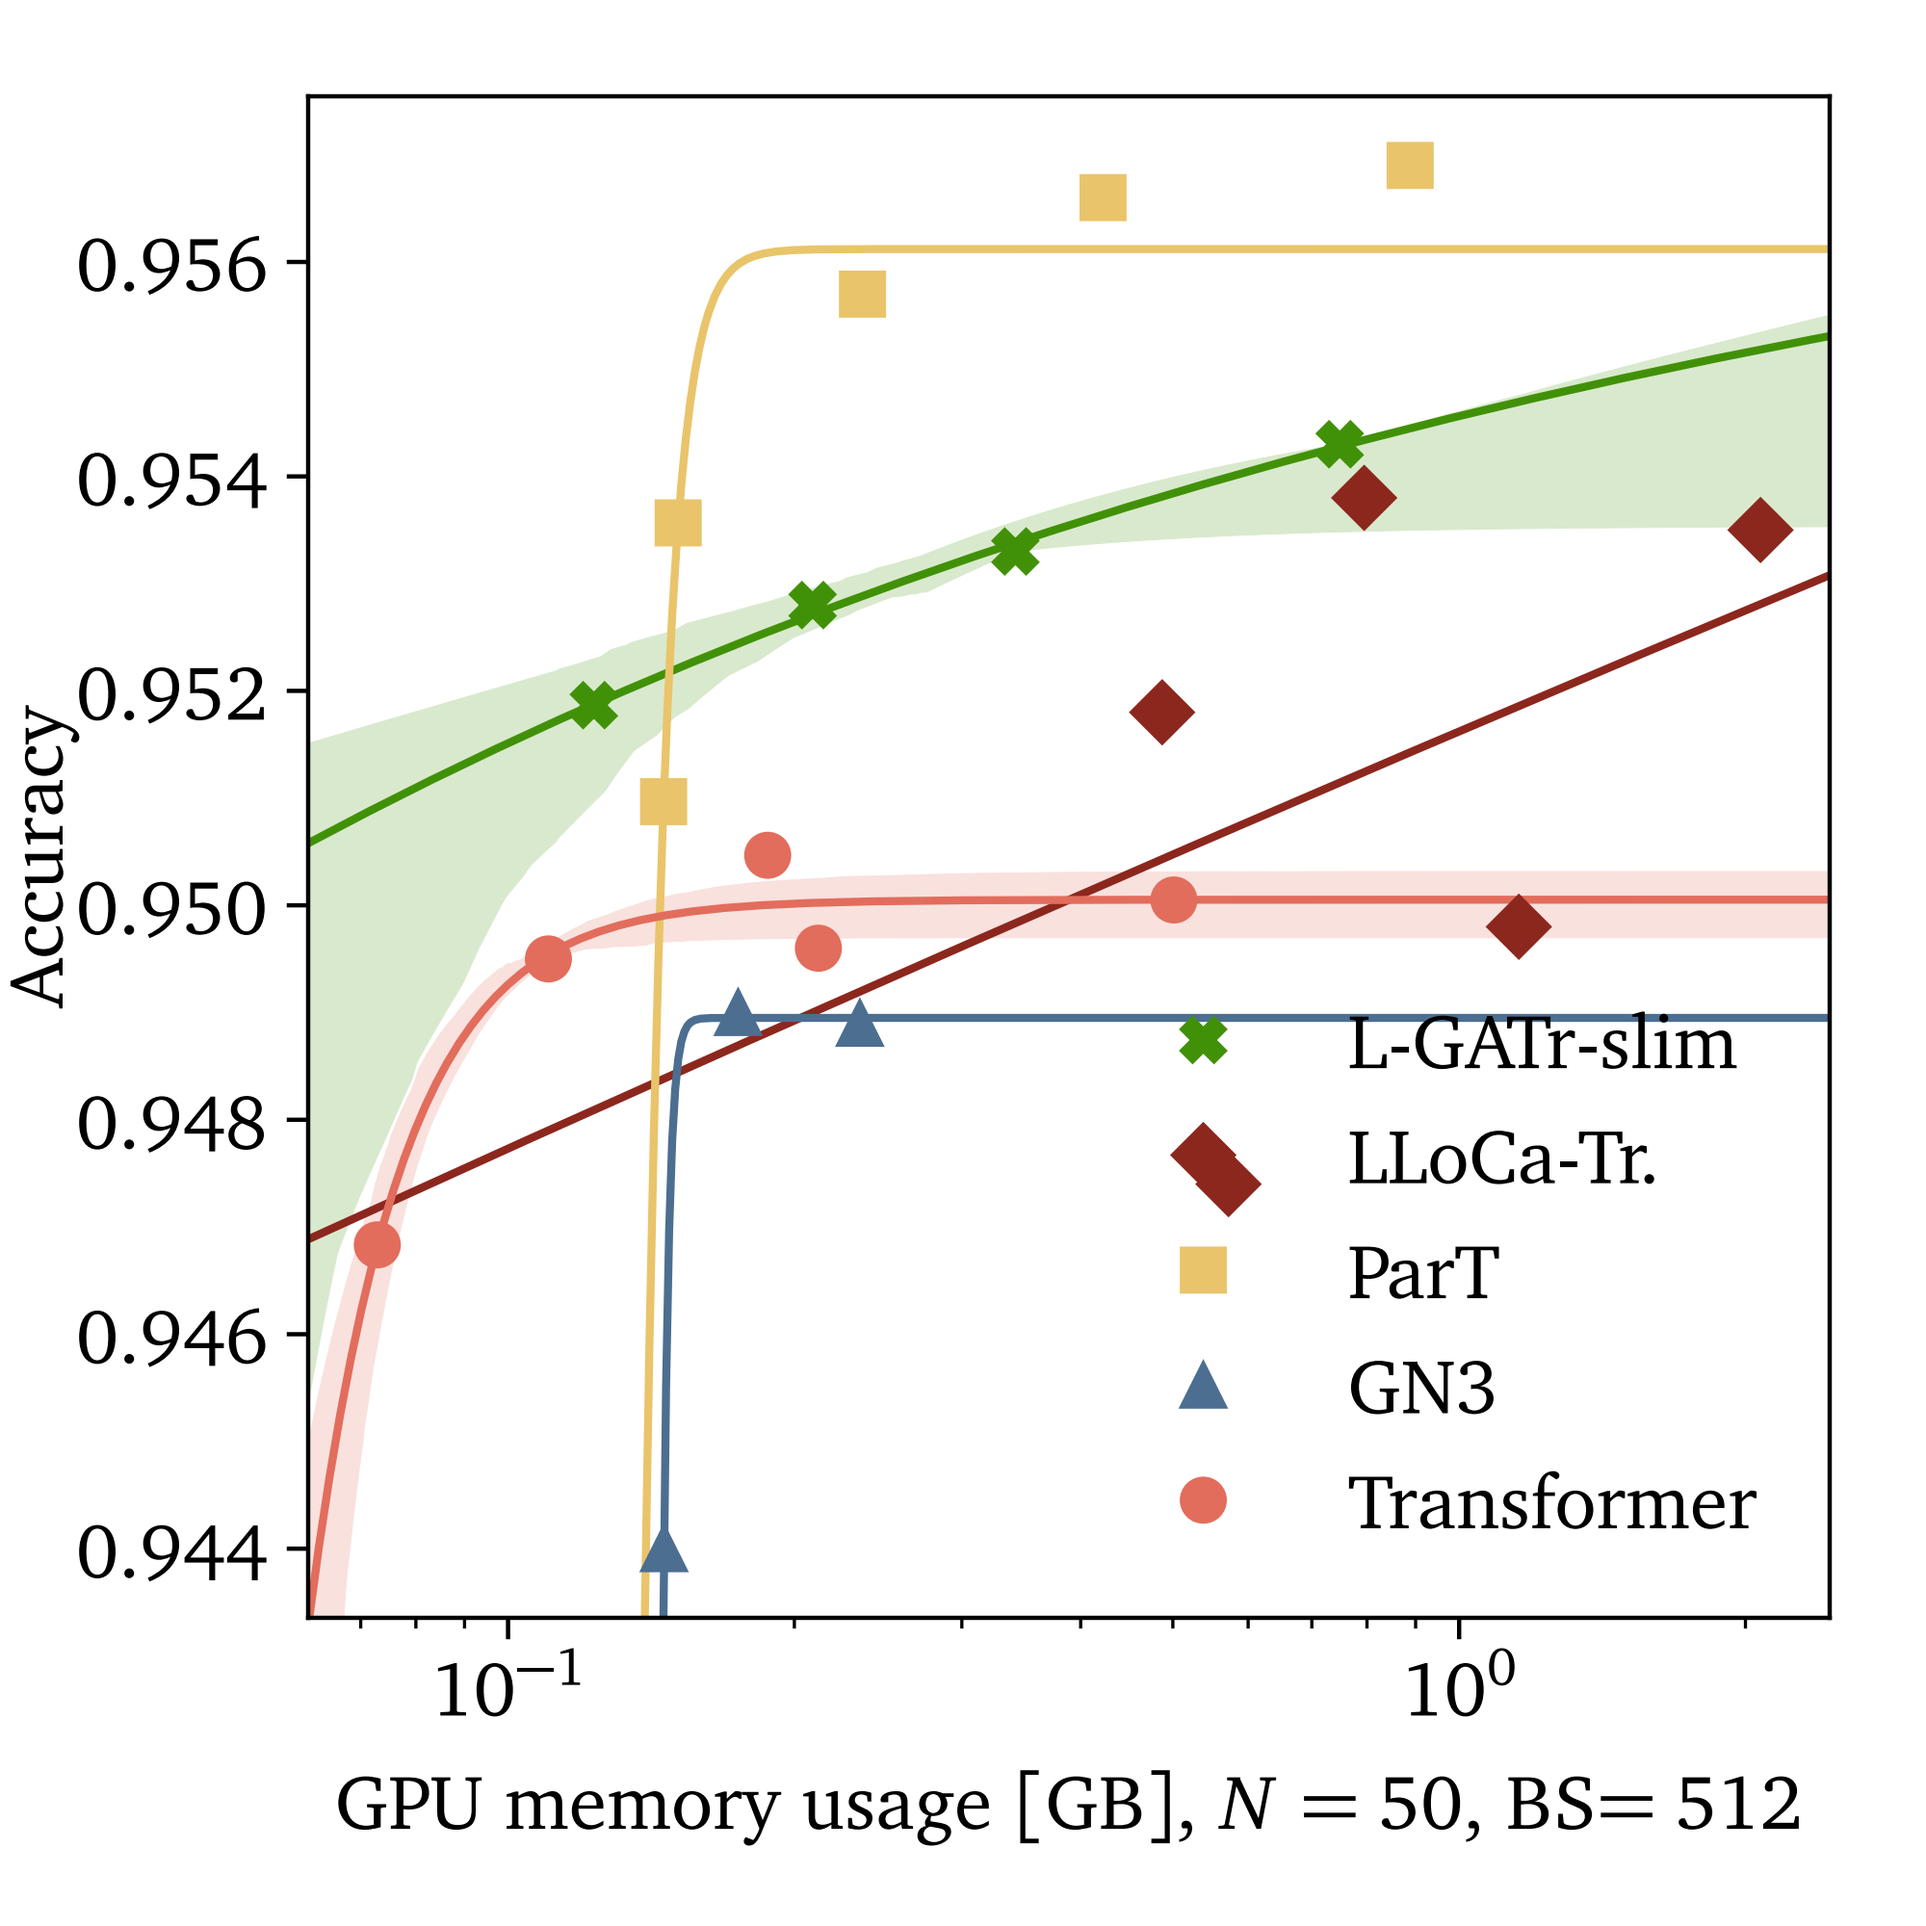
0.956
0.954
0.952
0.950
0.948
0.946
0.944
Accuracy
10^{-1}
10^{0}
GPU memory usage [GB], $N = 50$, BS$= 512$
L-GATr-slim
LLoCa-Tr.
ParT
GN3
Transformer

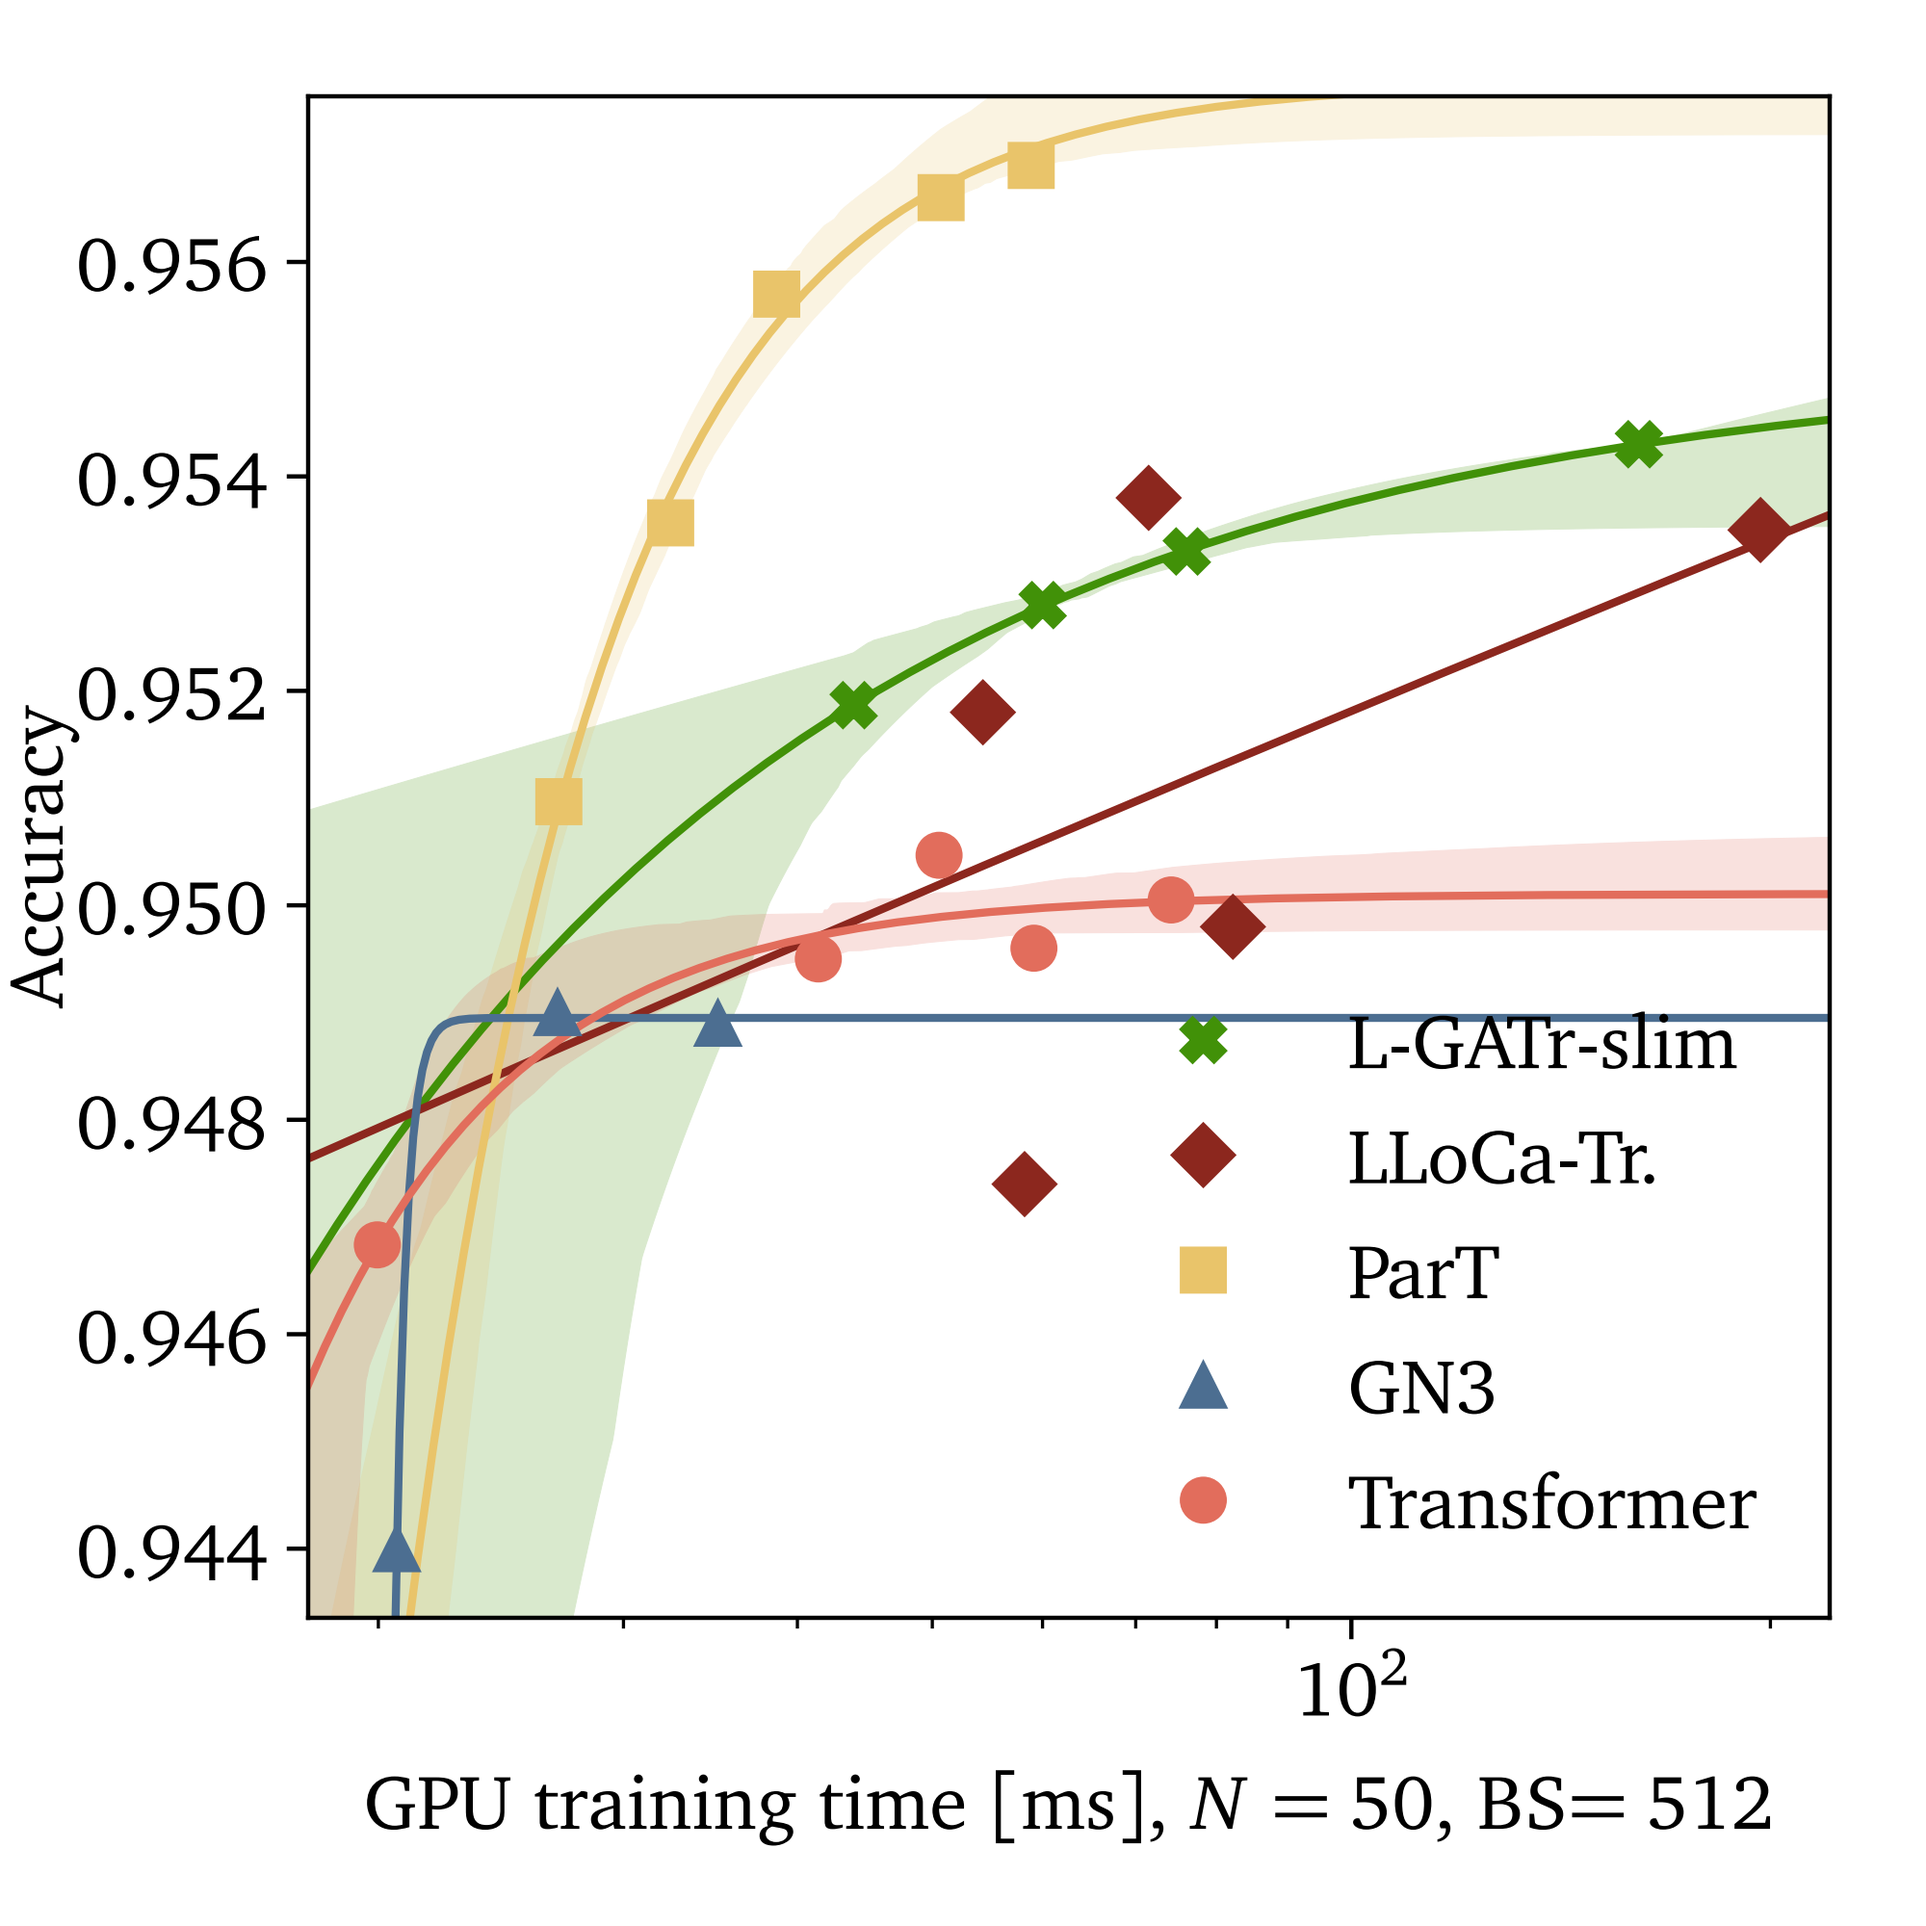
0.956
0.954
0.952
0.950
0.948
0.946
0.944
Accuracy
10^2
GPU training time [ms], $N = 50$, BS$= 512$
L-GATr-slim
LLoCa-Tr.
ParT
GN3
Transformer