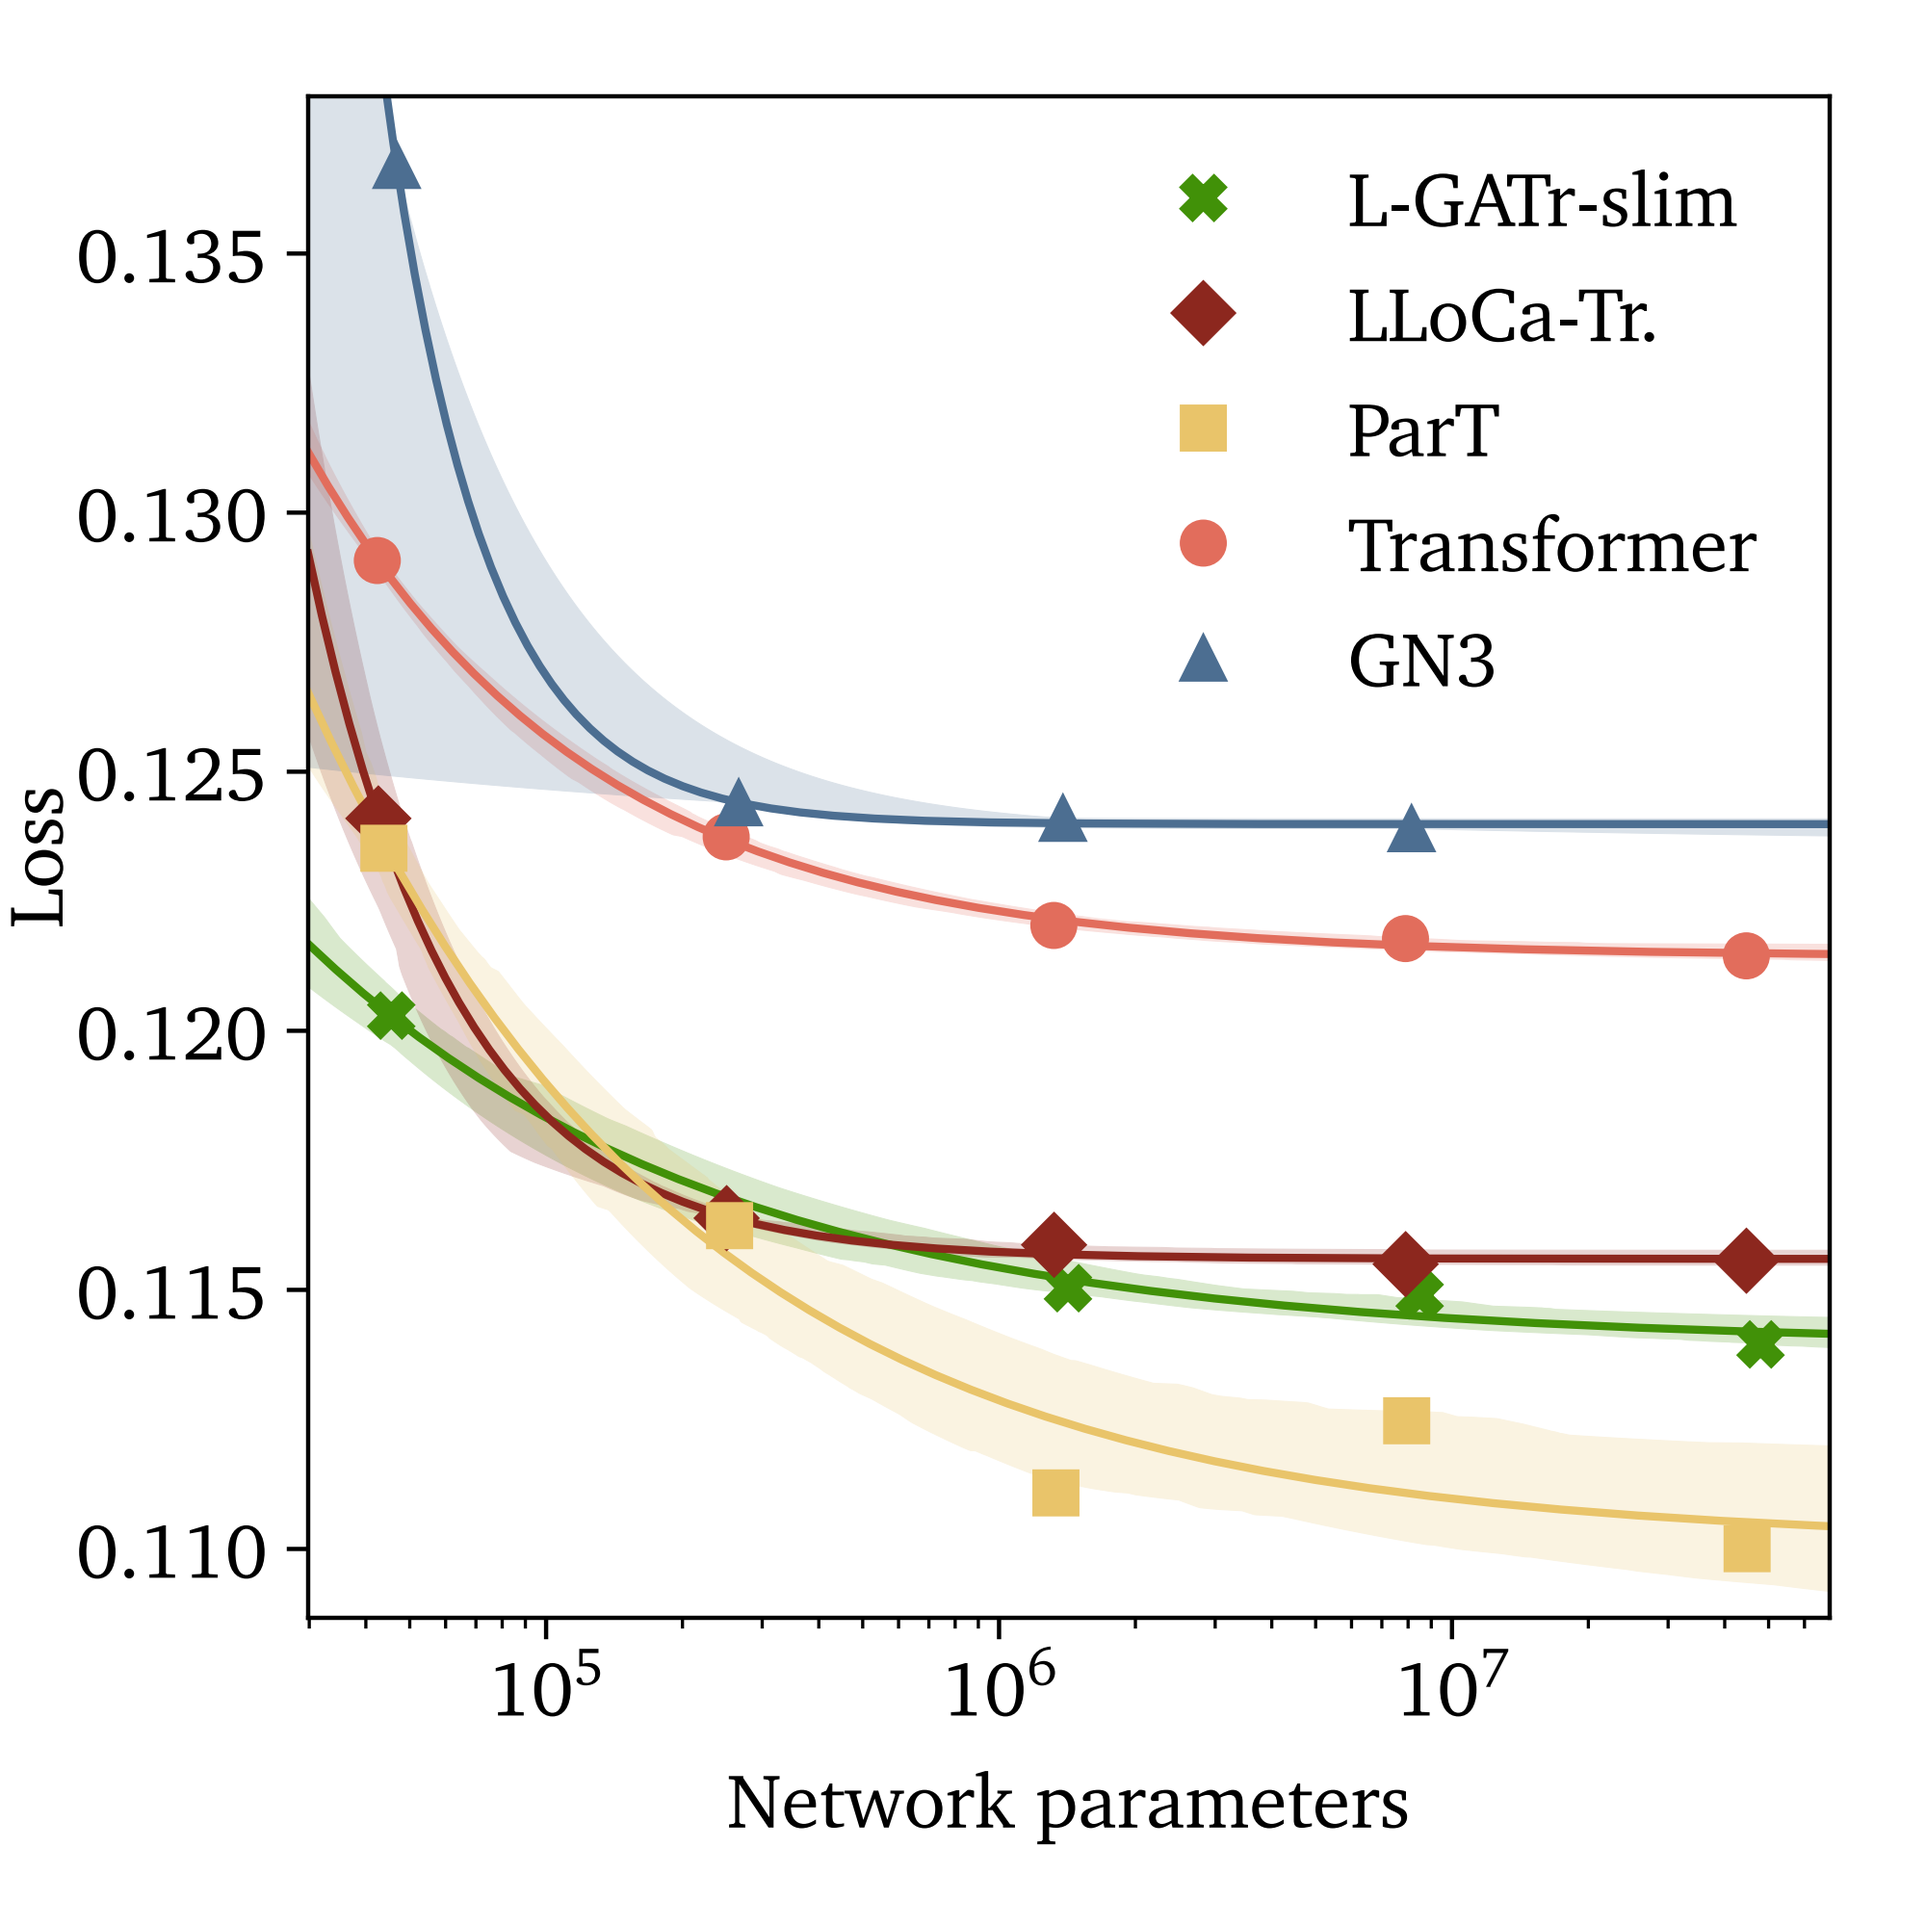
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GN3
Loss
0.135
0.130
0.125
0.120
0.115
0.110
$10^5$
$10^6$
$10^7$
Network parameters

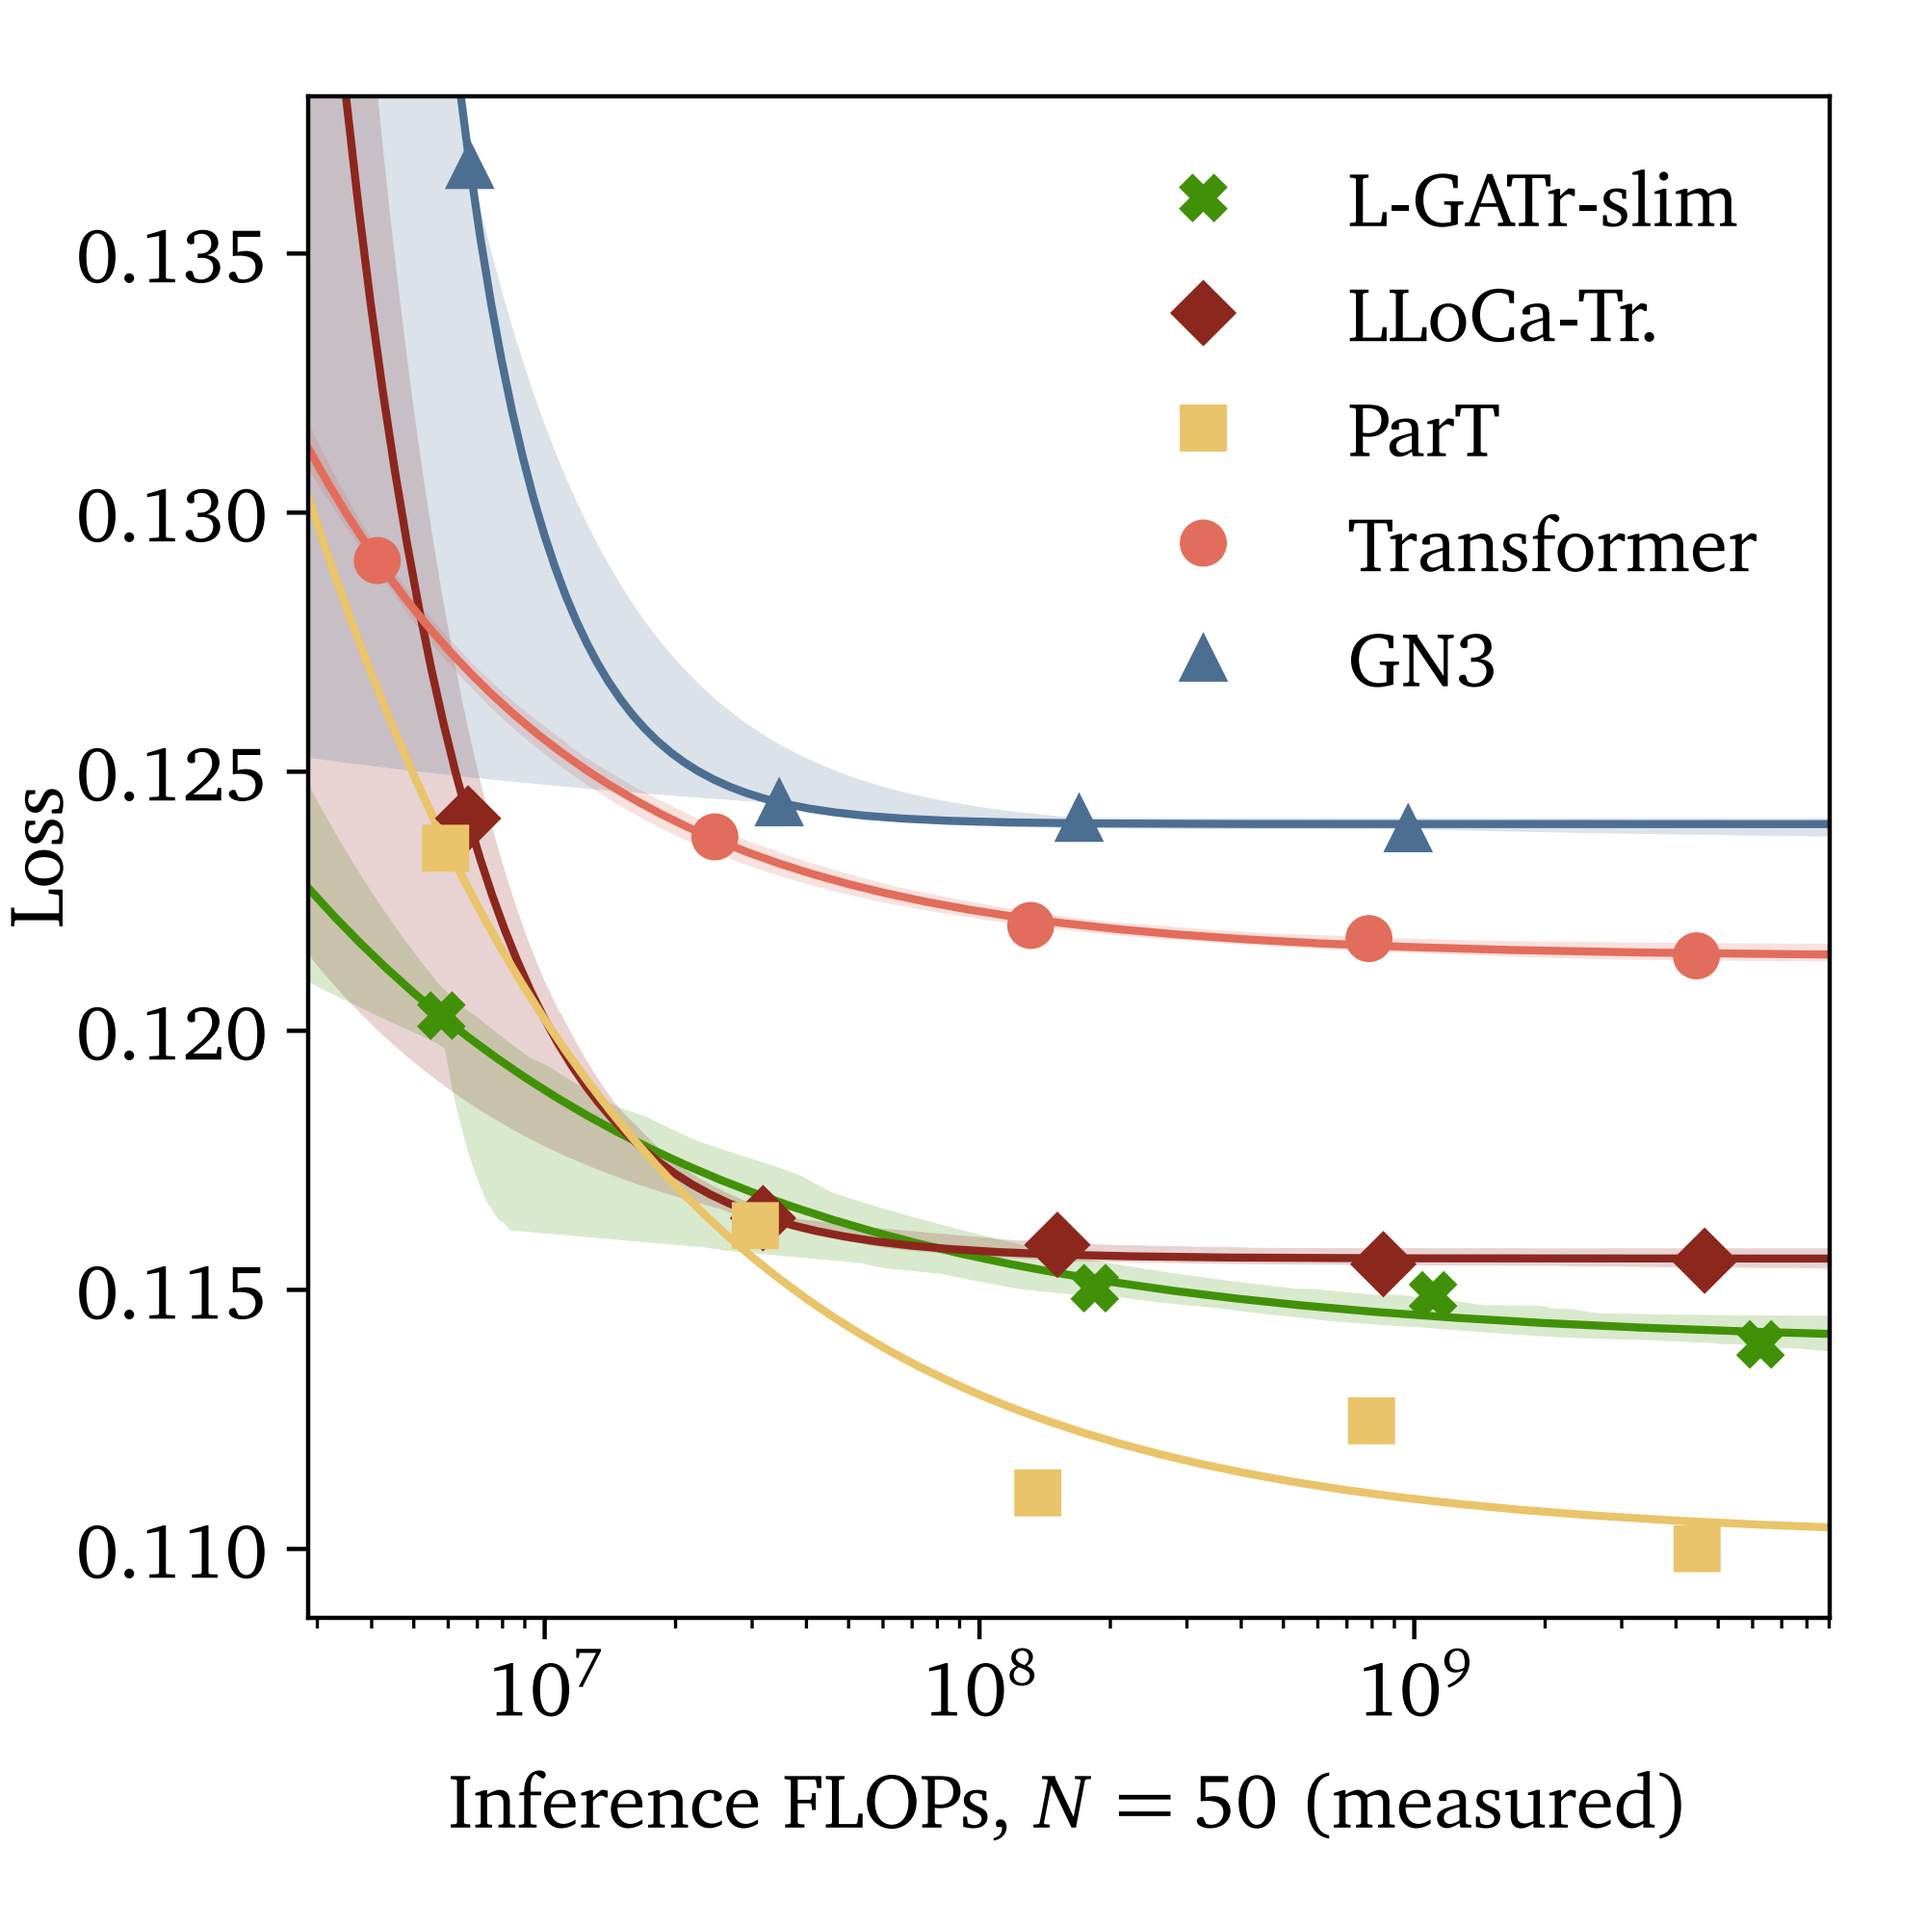
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GN3
Loss
0.135
0.130
0.125
0.120
0.115
0.110
$10^7$
$10^8$
$10^9$
Inference FLOPs, $N = 50$ (measured)

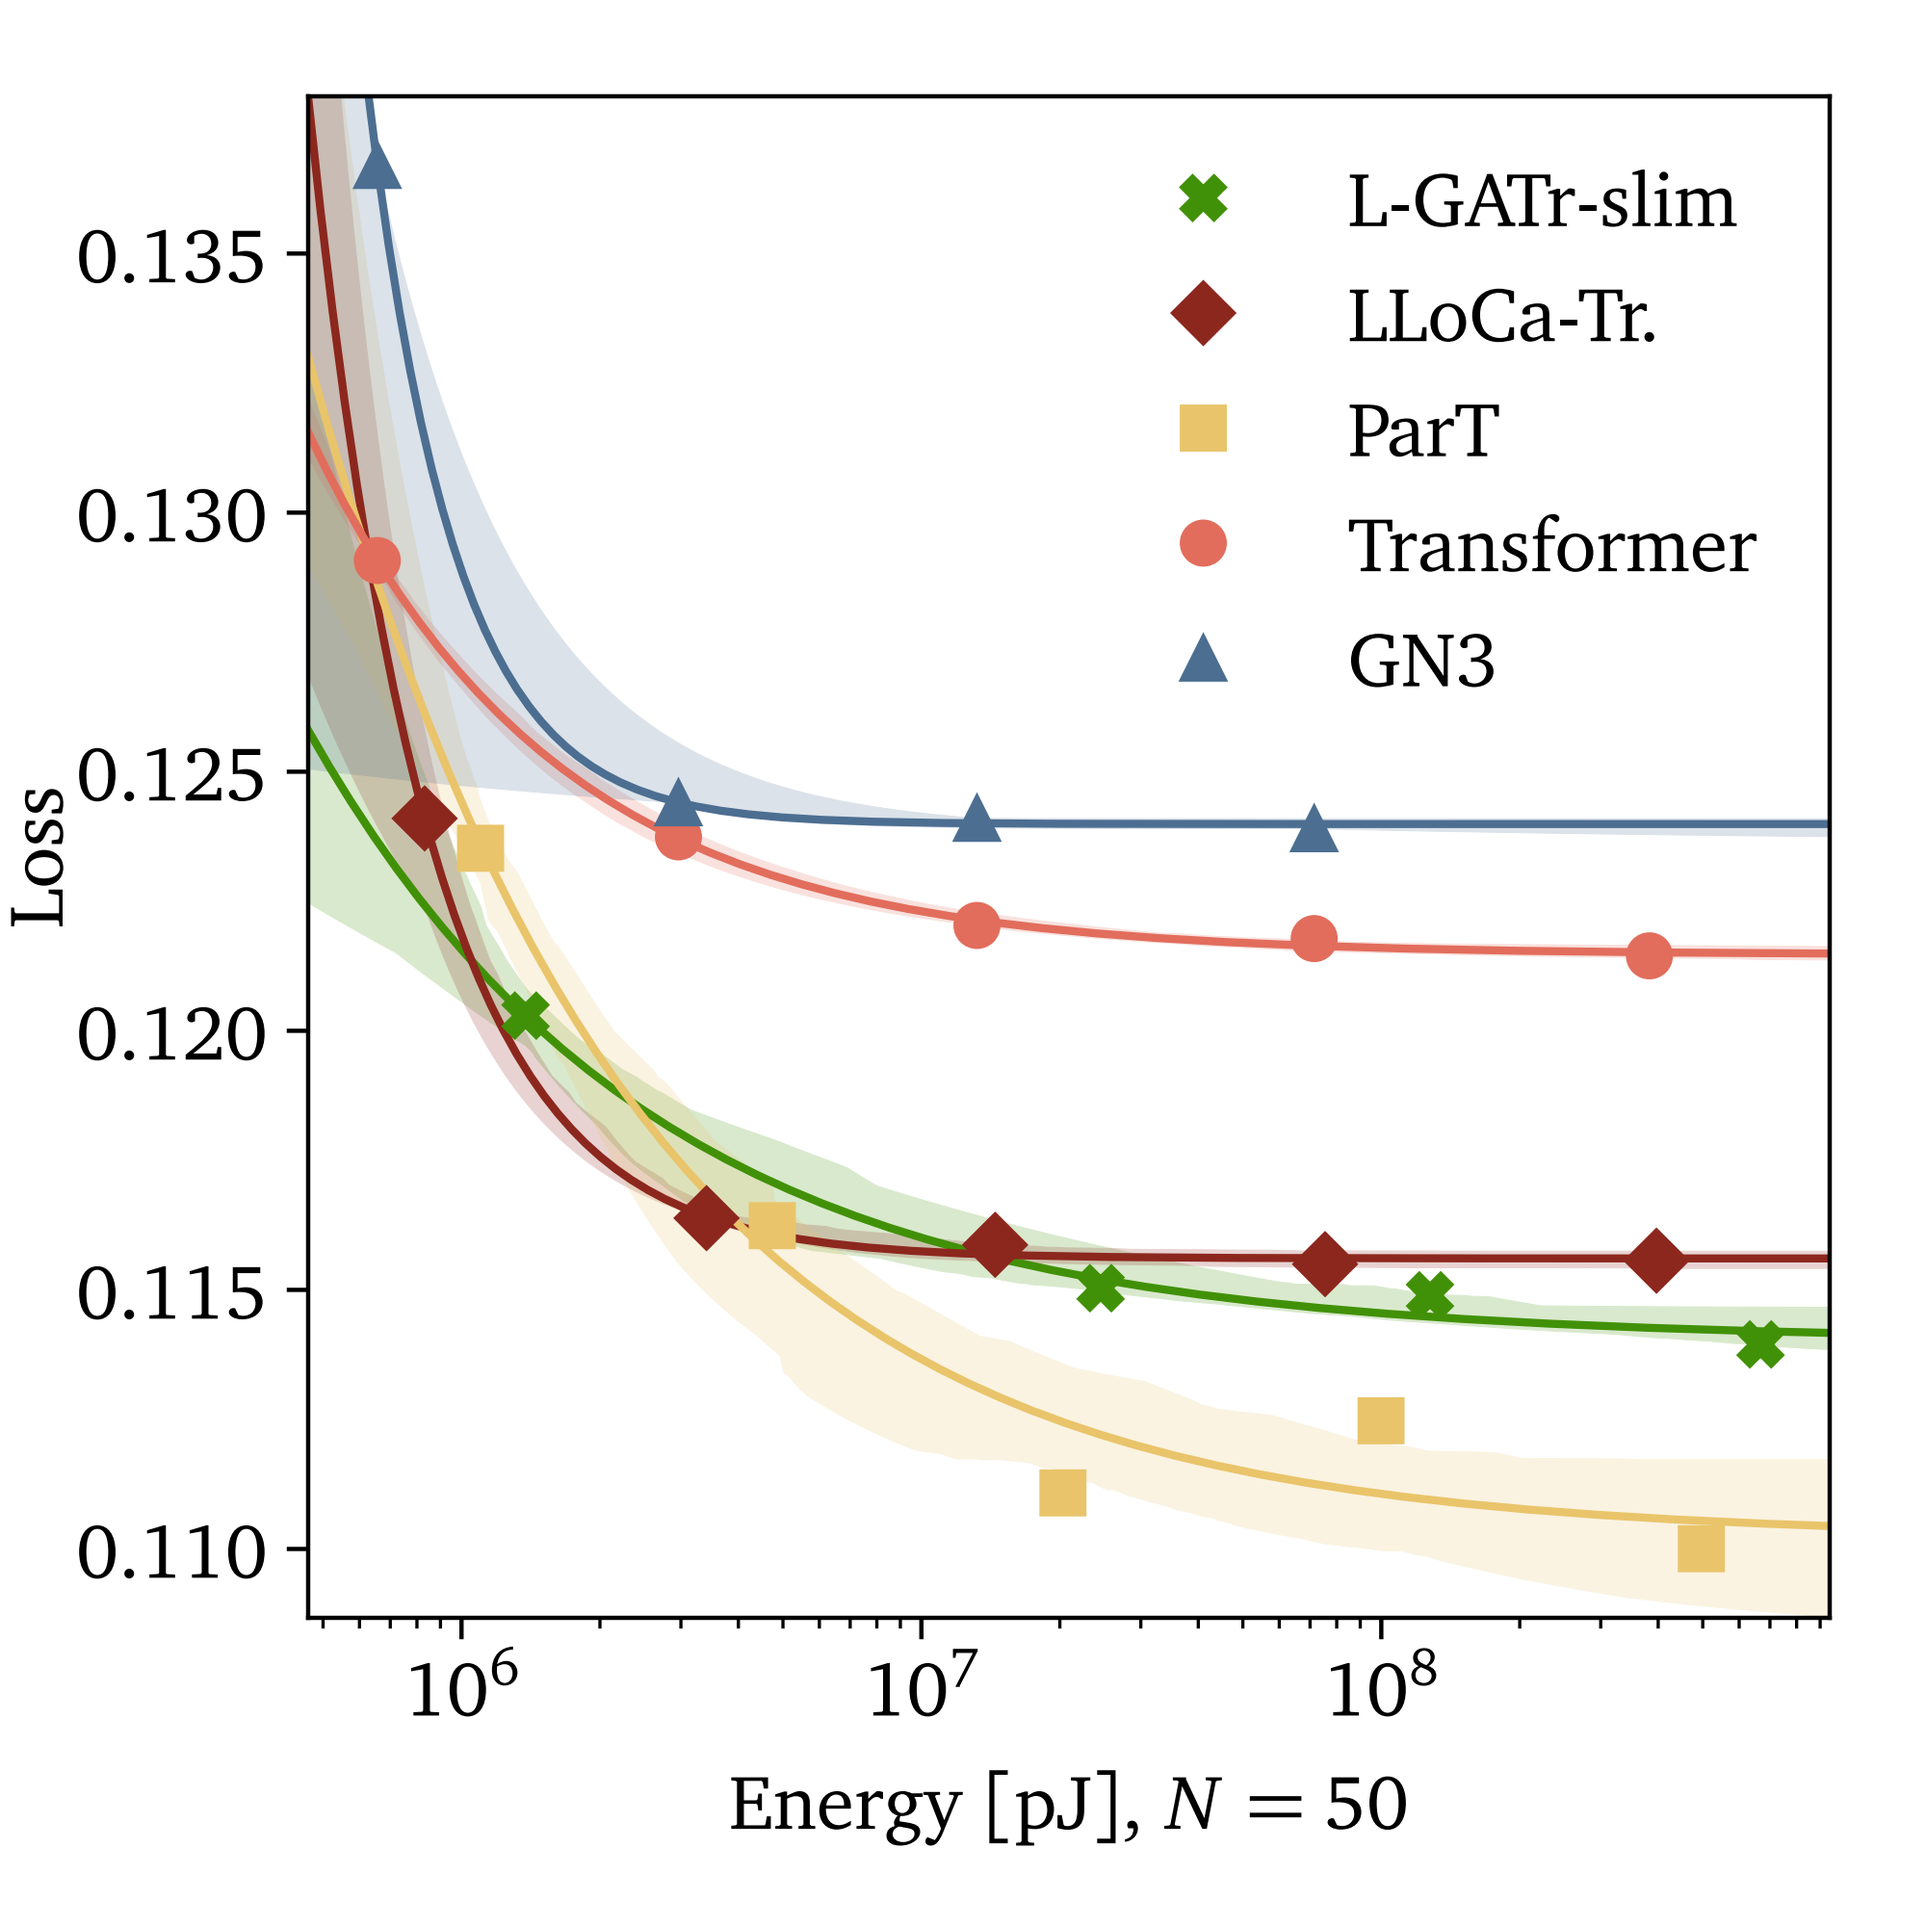
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GN3
Loss
0.135
0.130
0.125
0.120
0.115
0.110
$10^6$
$10^7$
$10^8$
Energy [pJ], $N = 50$

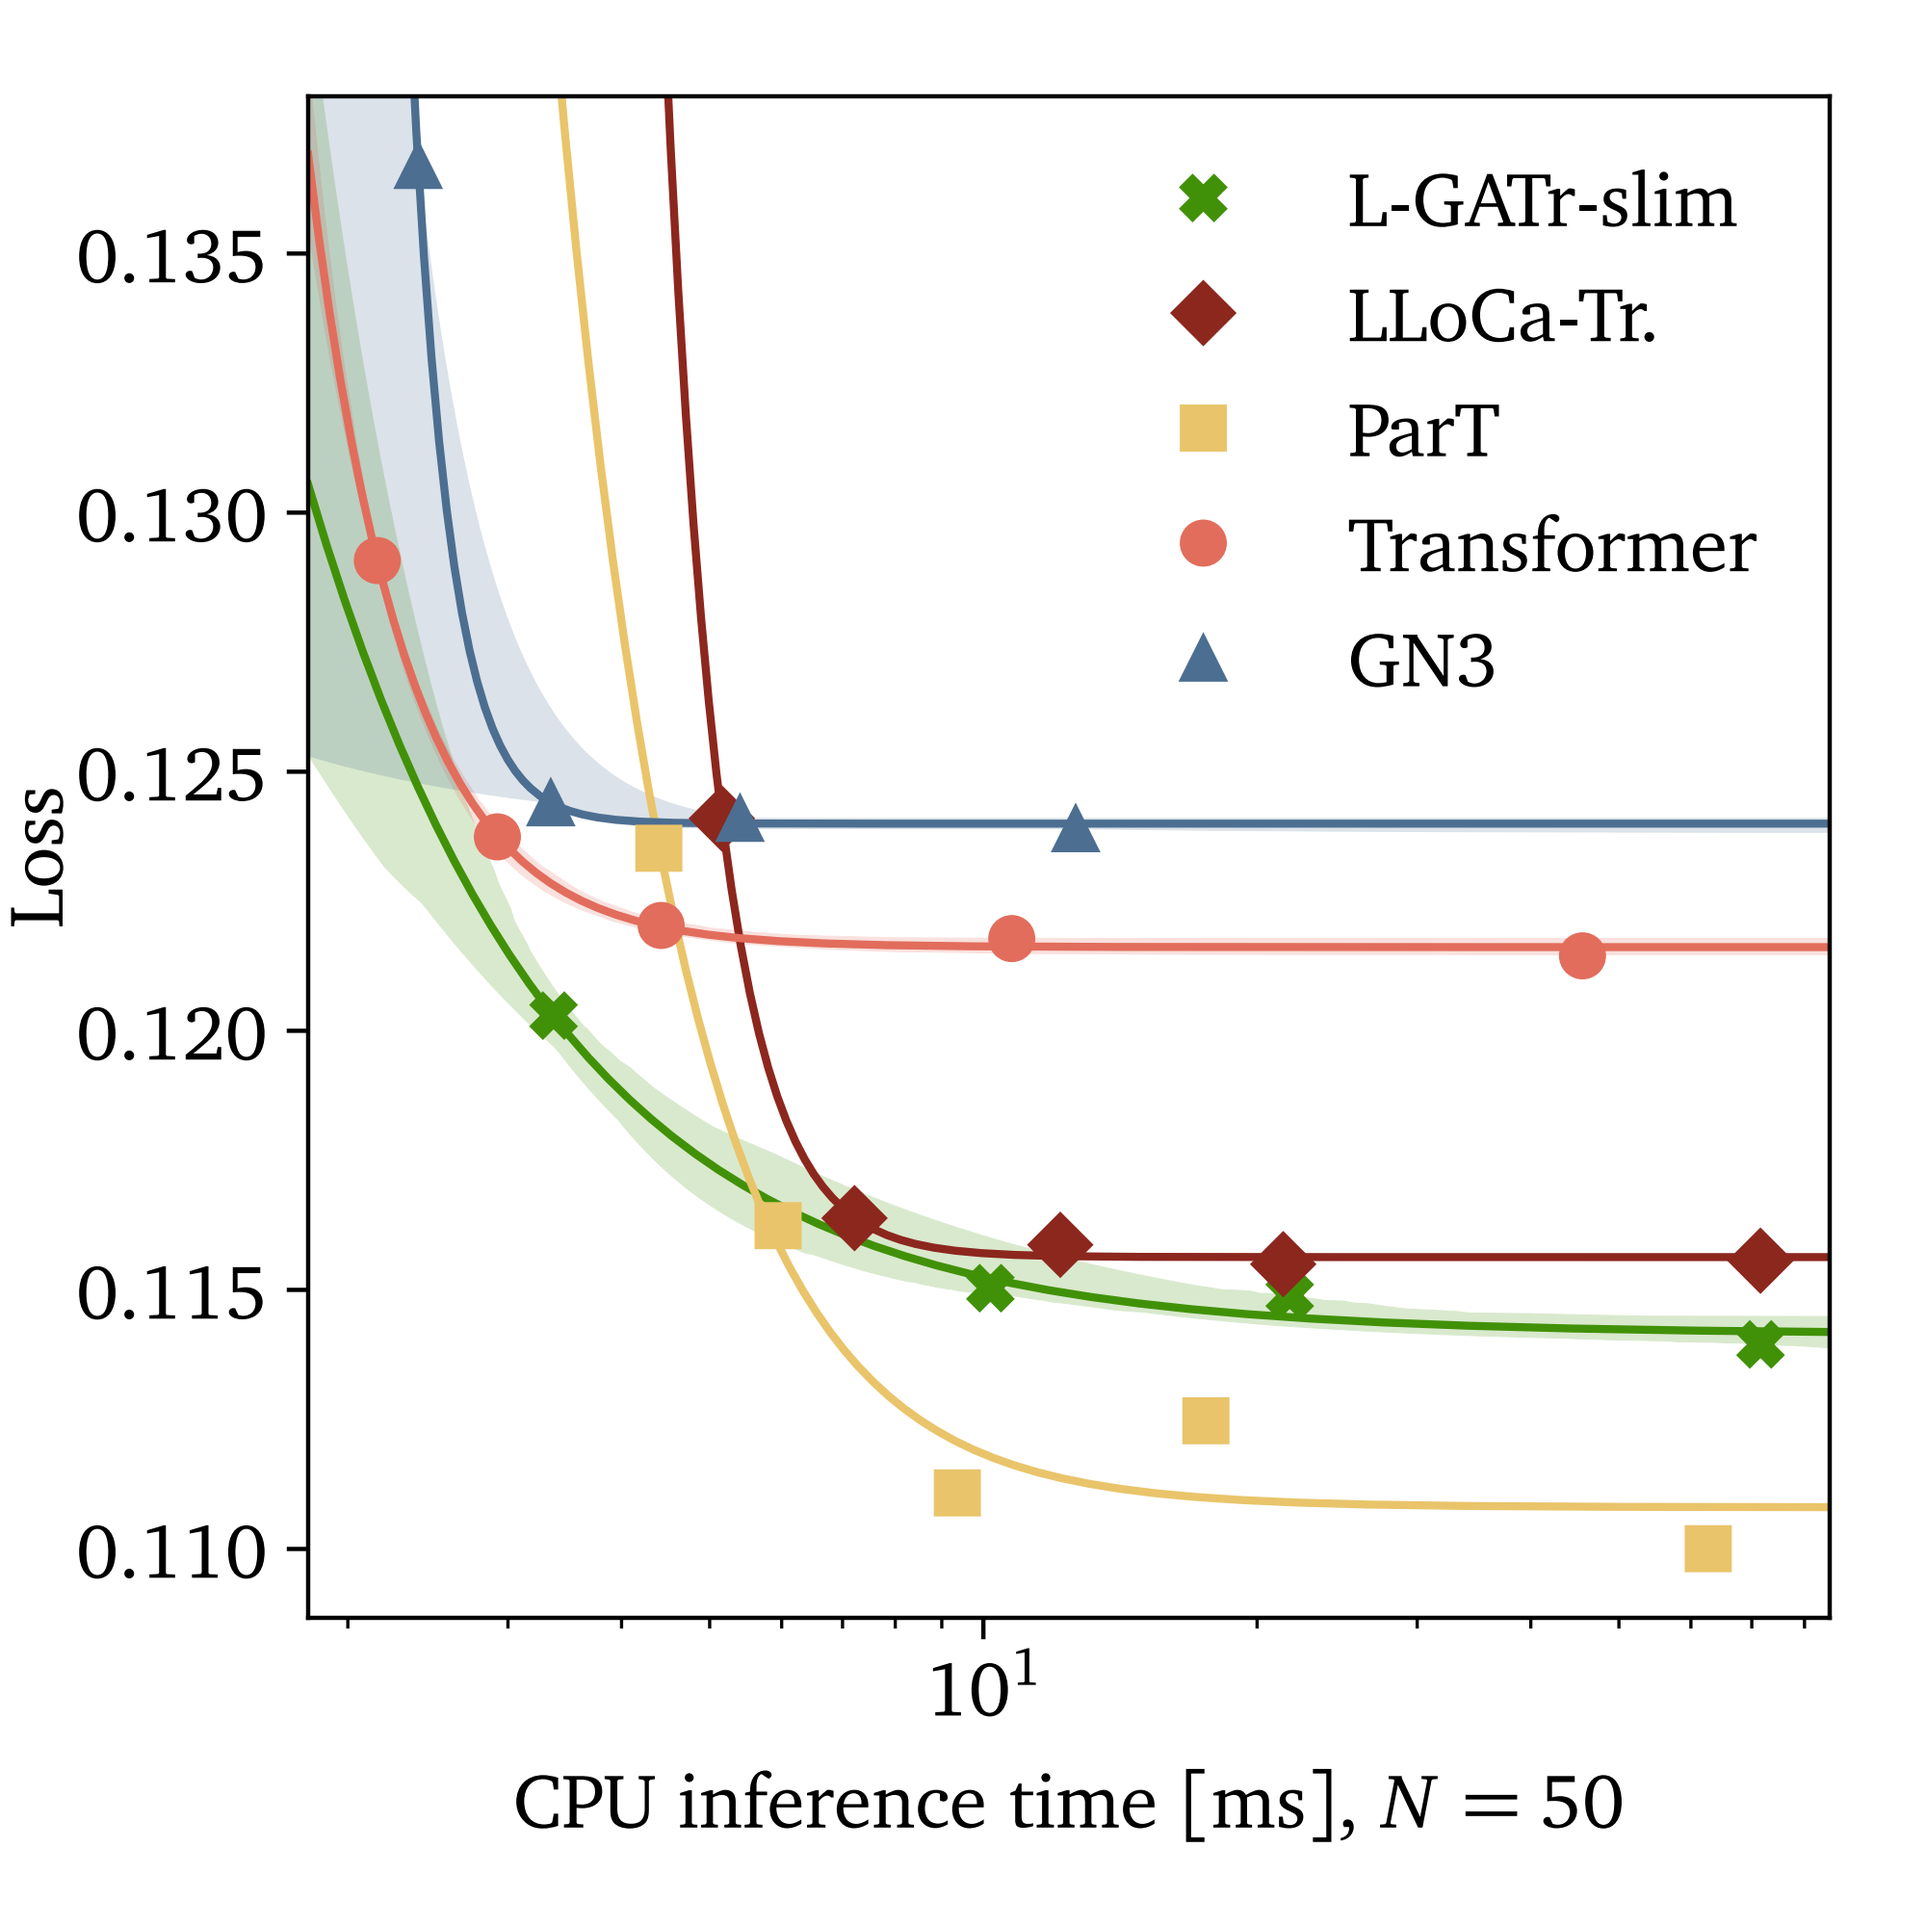
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GN3
Loss
0.135
0.130
0.125
0.120
0.115
0.110
$10^1$
CPU inference time [ms], $N = 50$

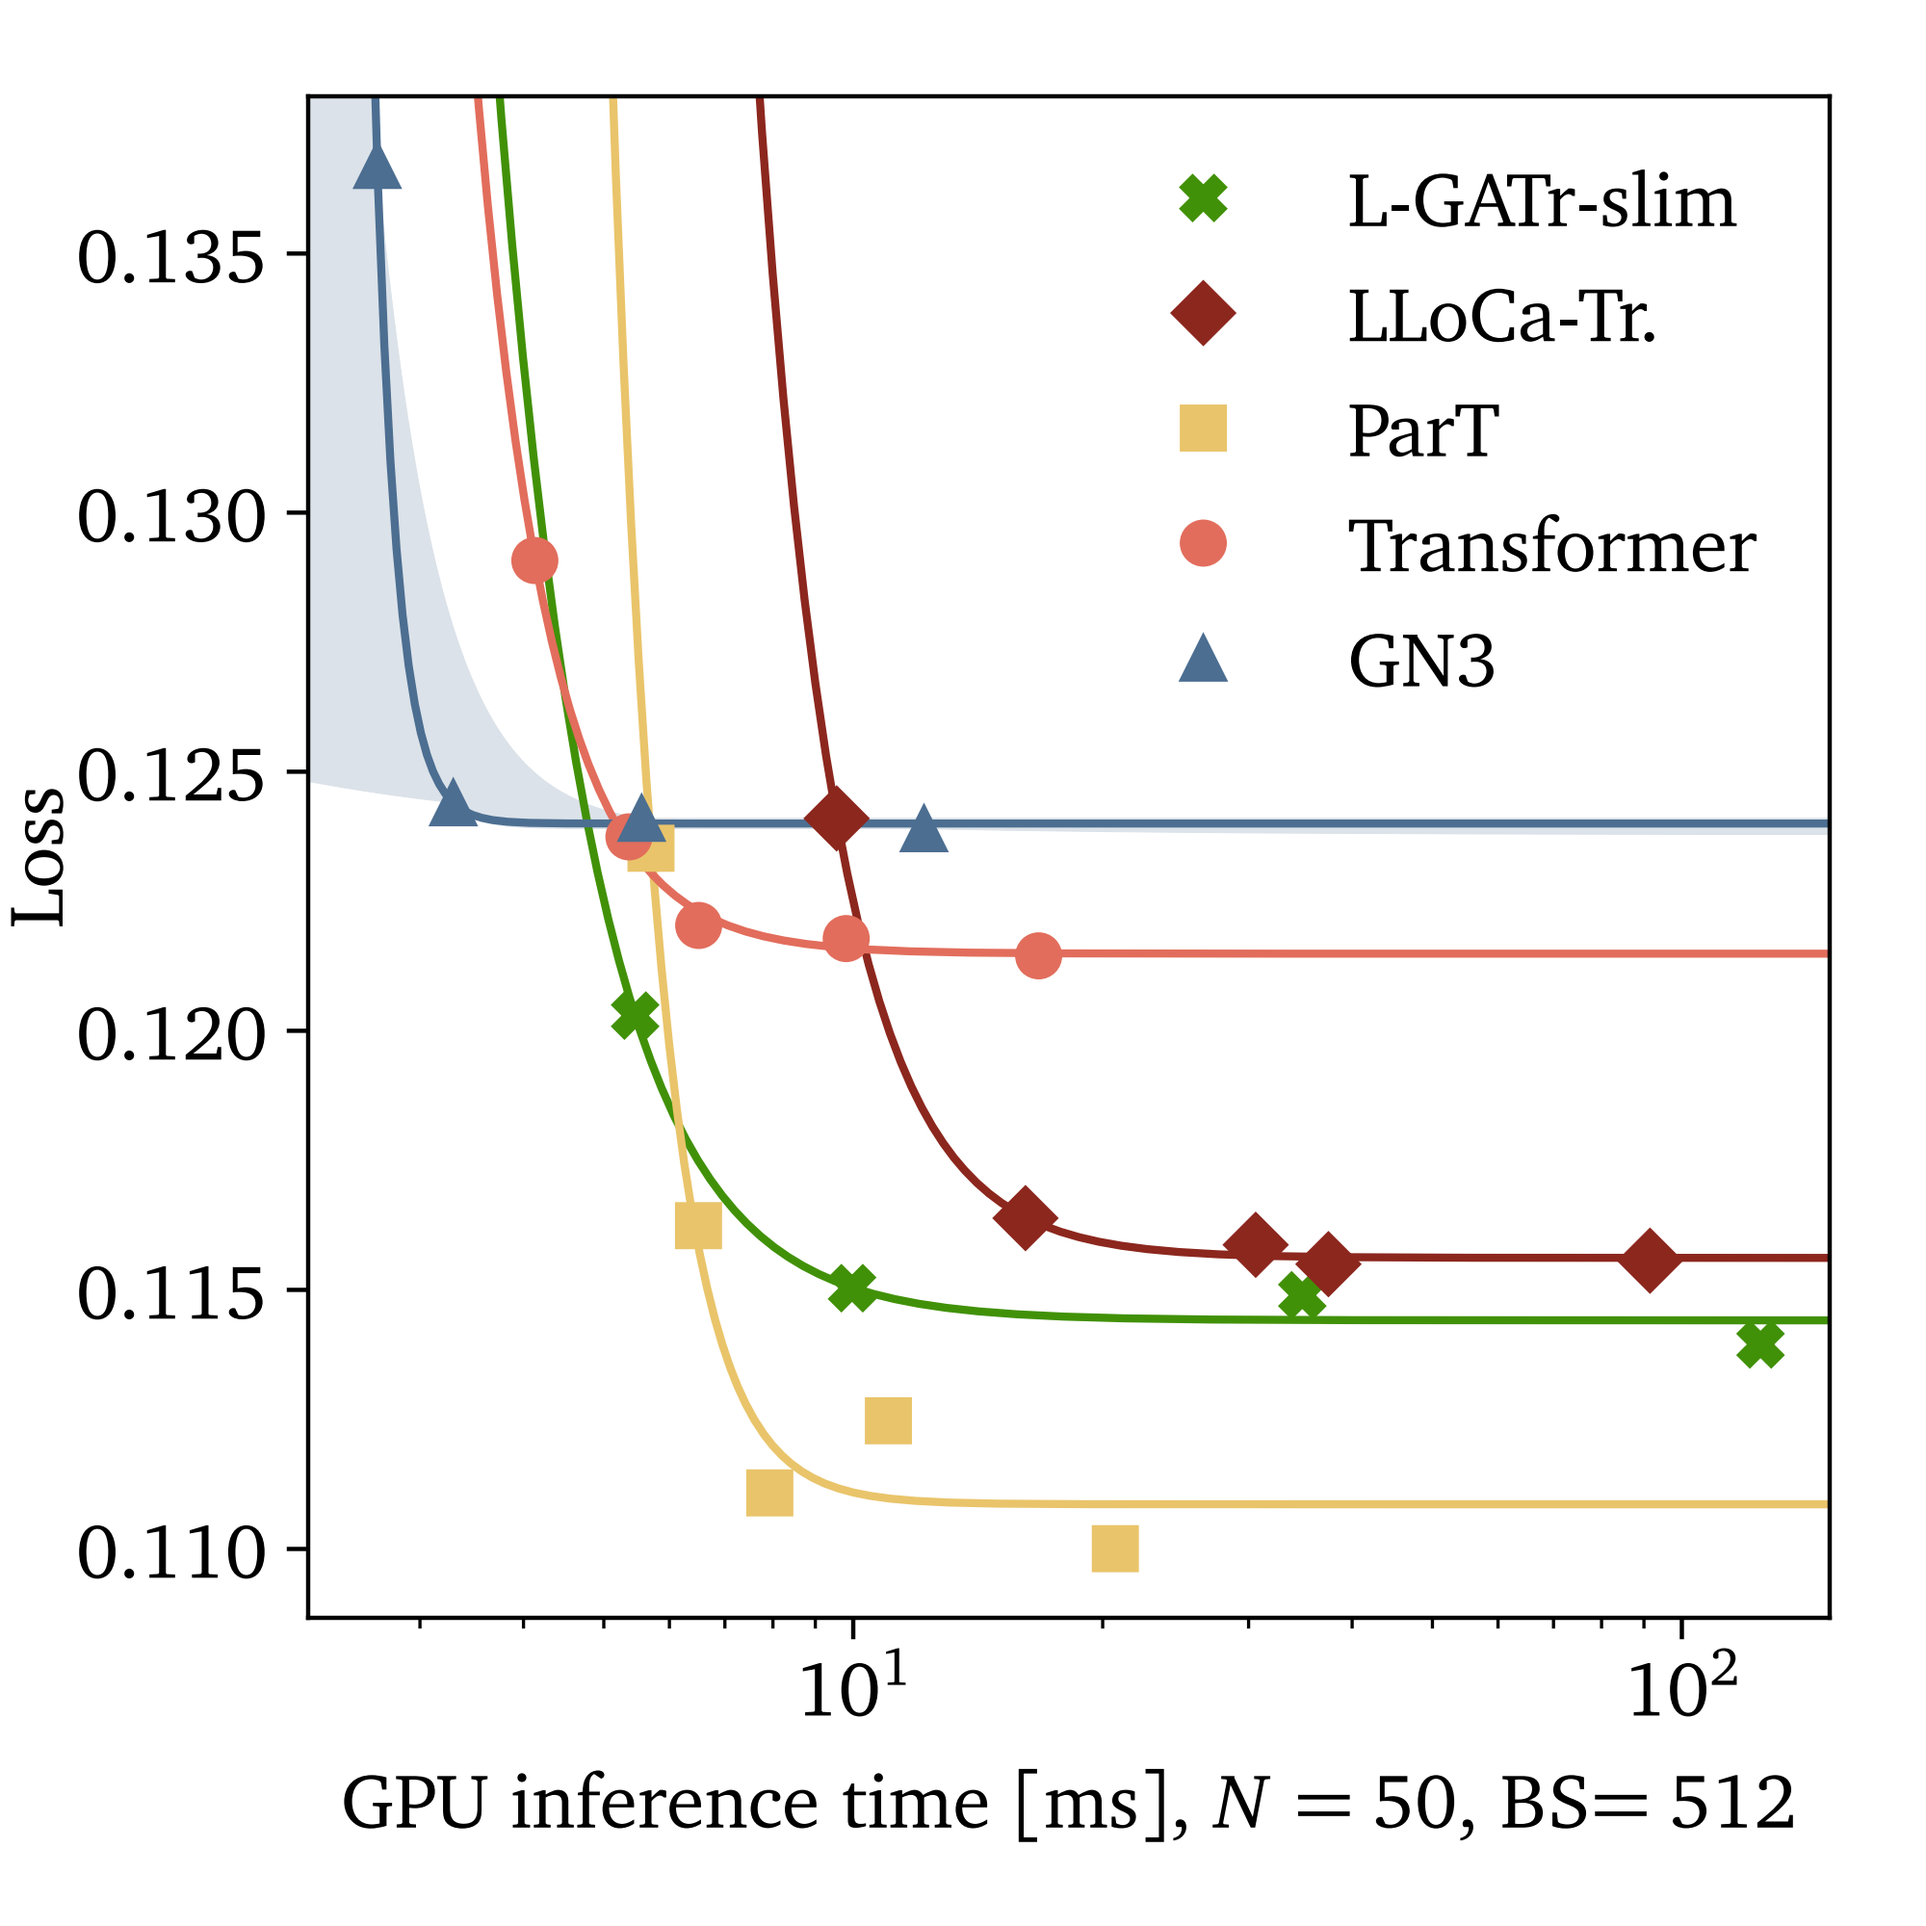
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GN3
Loss
0.135
0.130
0.125
0.120
0.115
0.110
$10^1$
$10^2$
GPU inference time [ms], $N = 50$, BS$= 512$

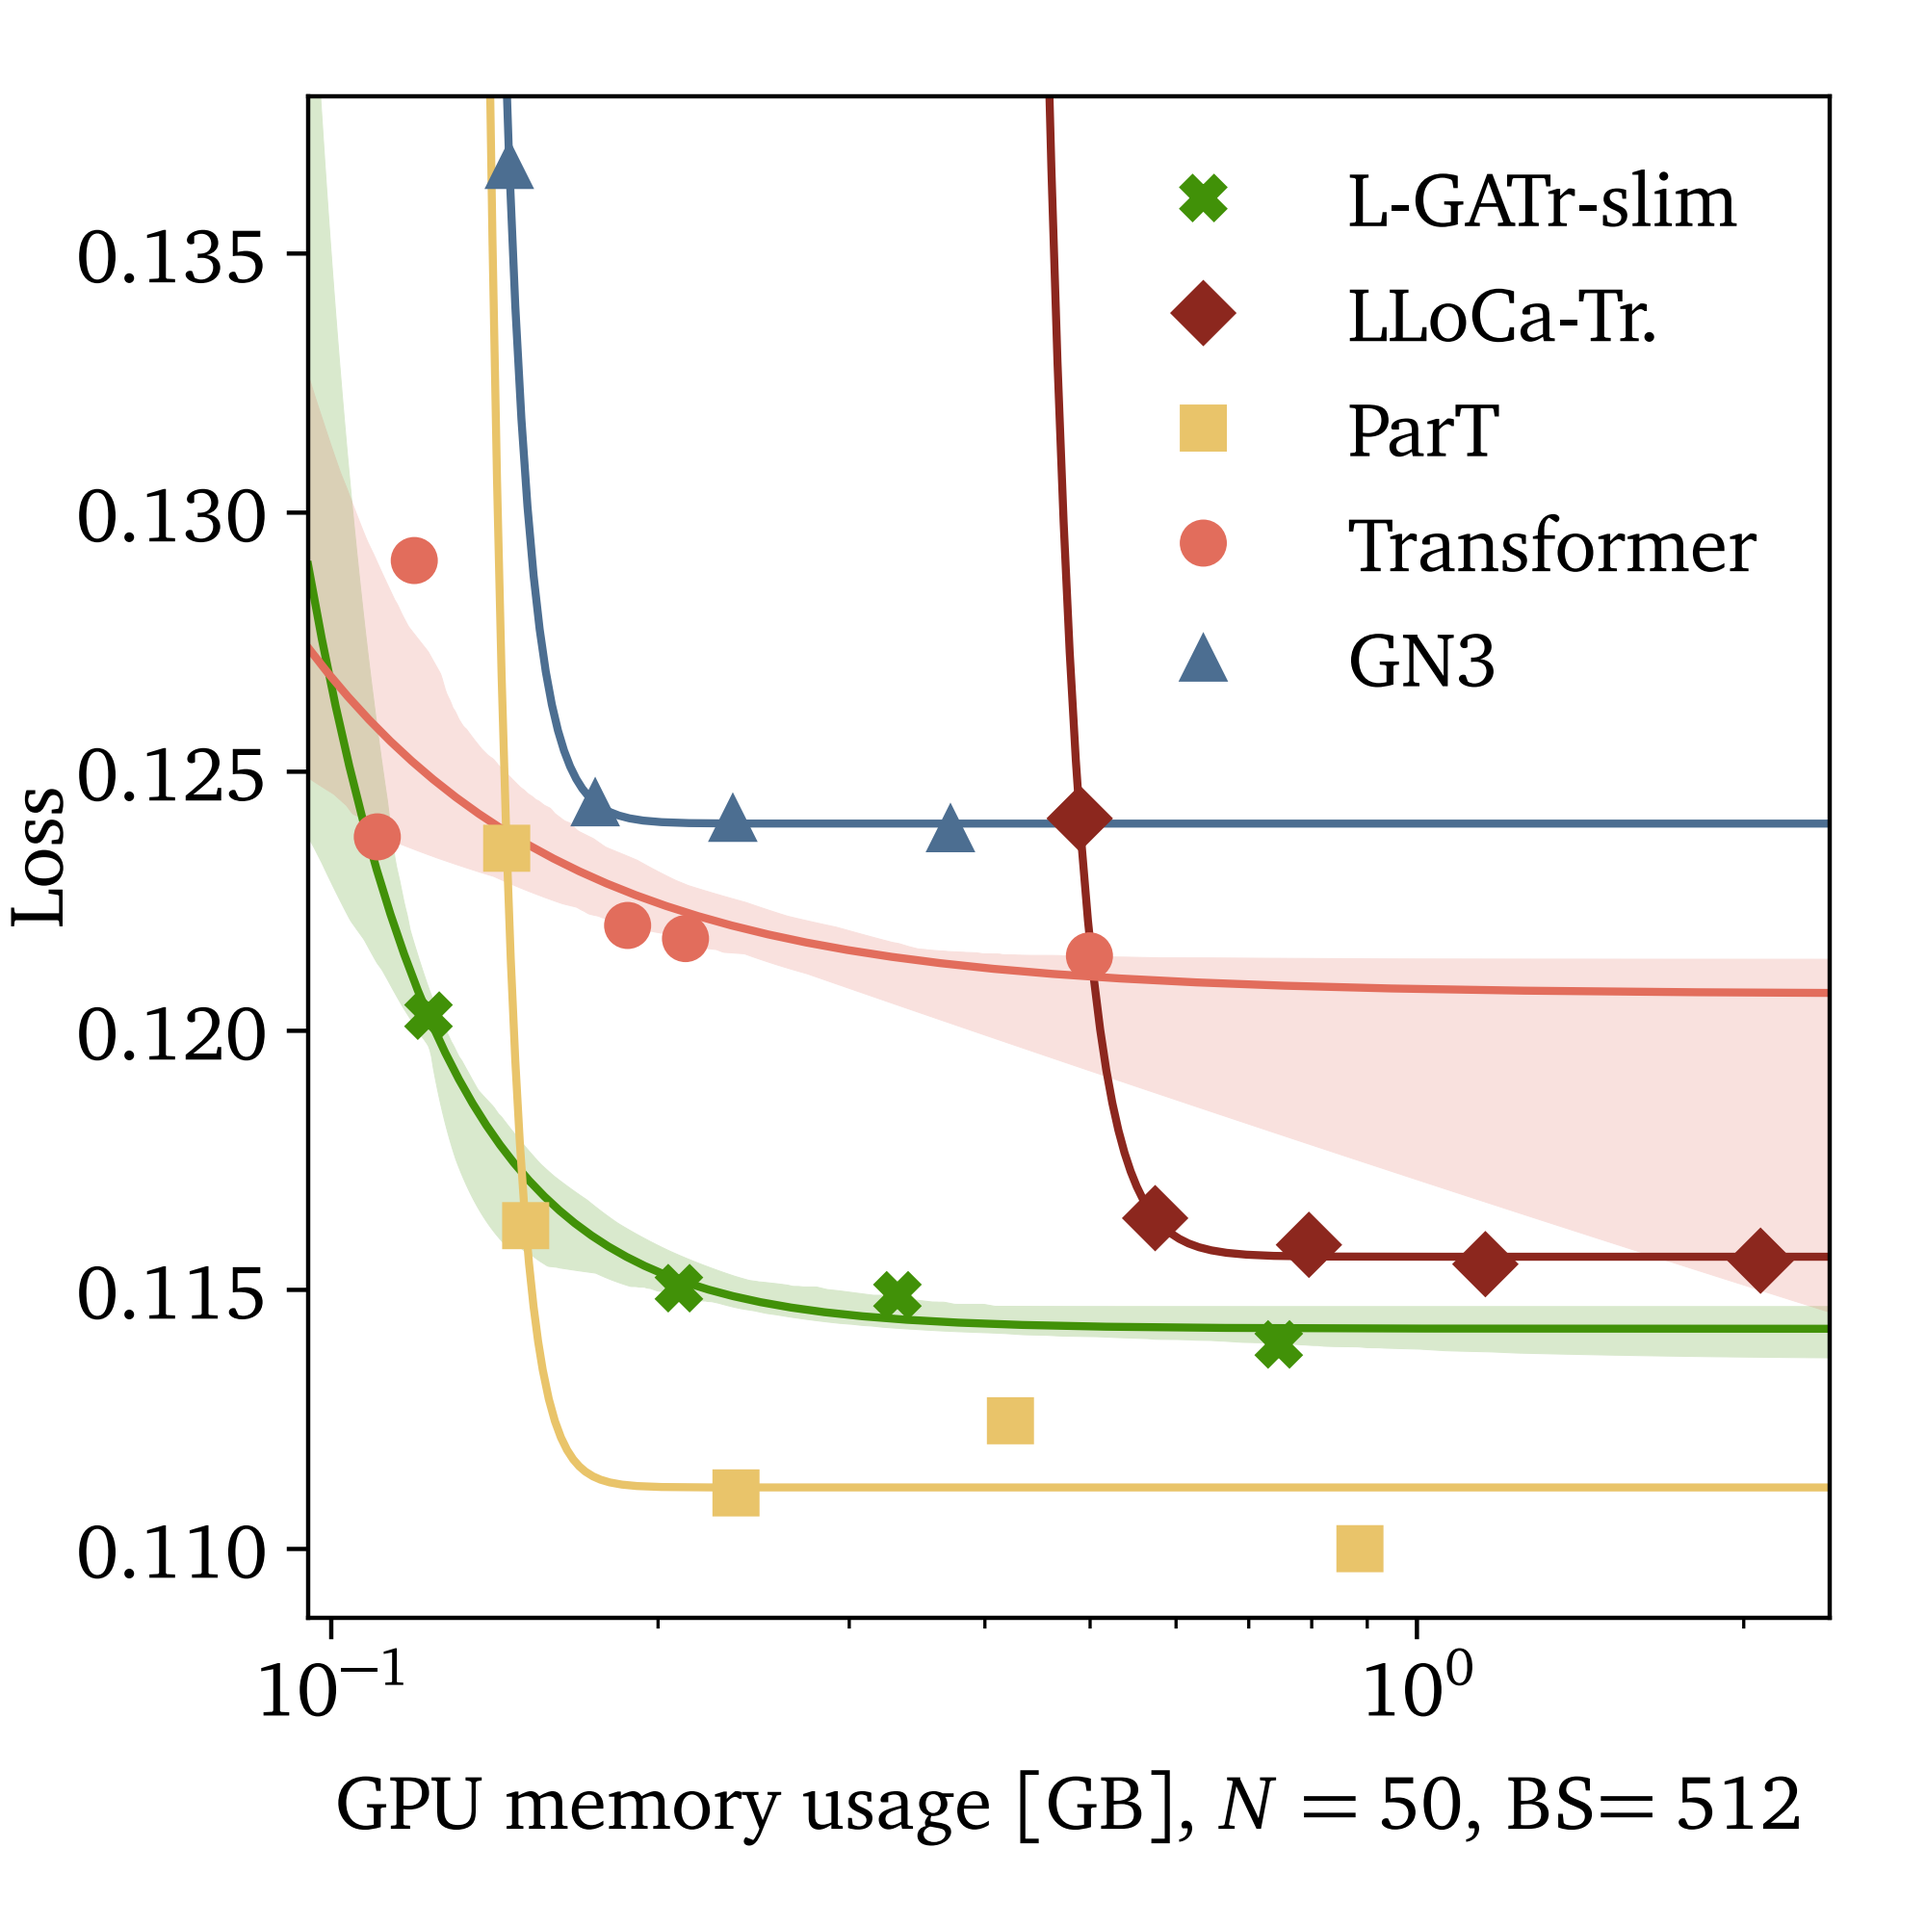
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GN3
Loss
0.135
0.130
0.125
0.120
0.115
0.110
$10^{-1}$
$10^{0}$
GPU memory usage [GB], $N = 50$, BS$= 512$

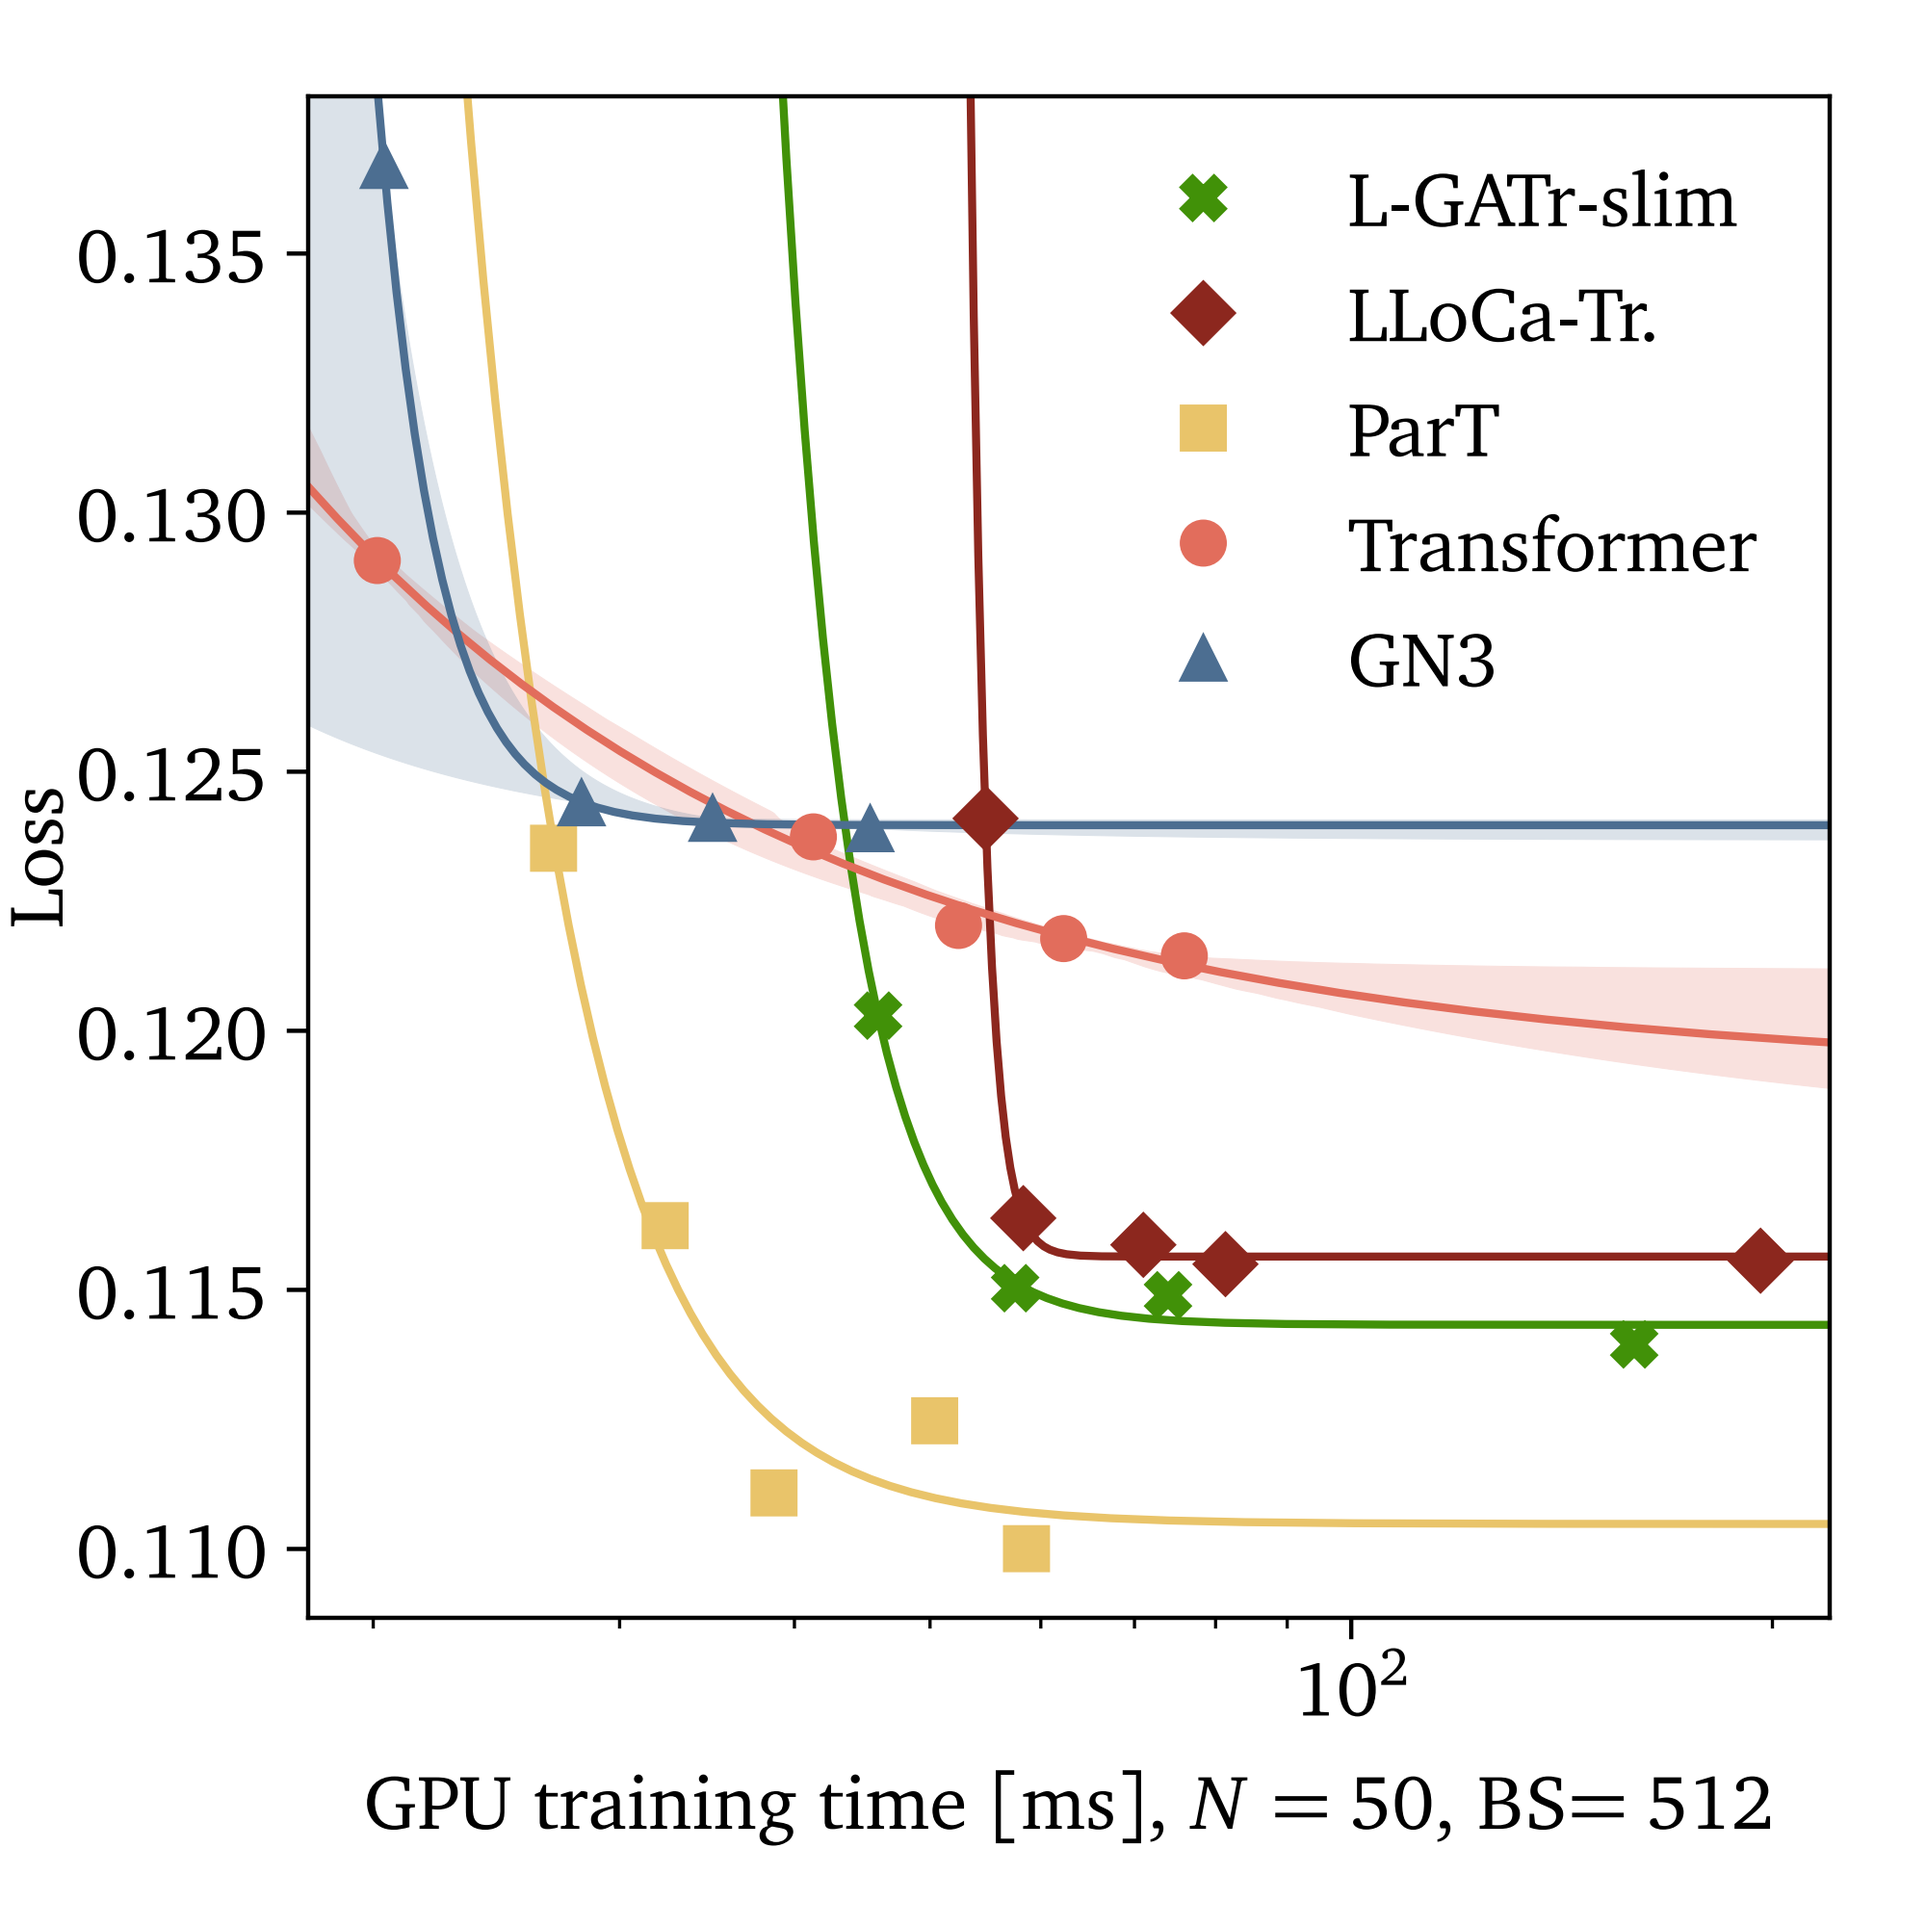
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GN3
Loss
0.135
0.130
0.125
0.120
0.115
0.110
$10^2$
GPU training time [ms], $N = 50$, BS$= 512$

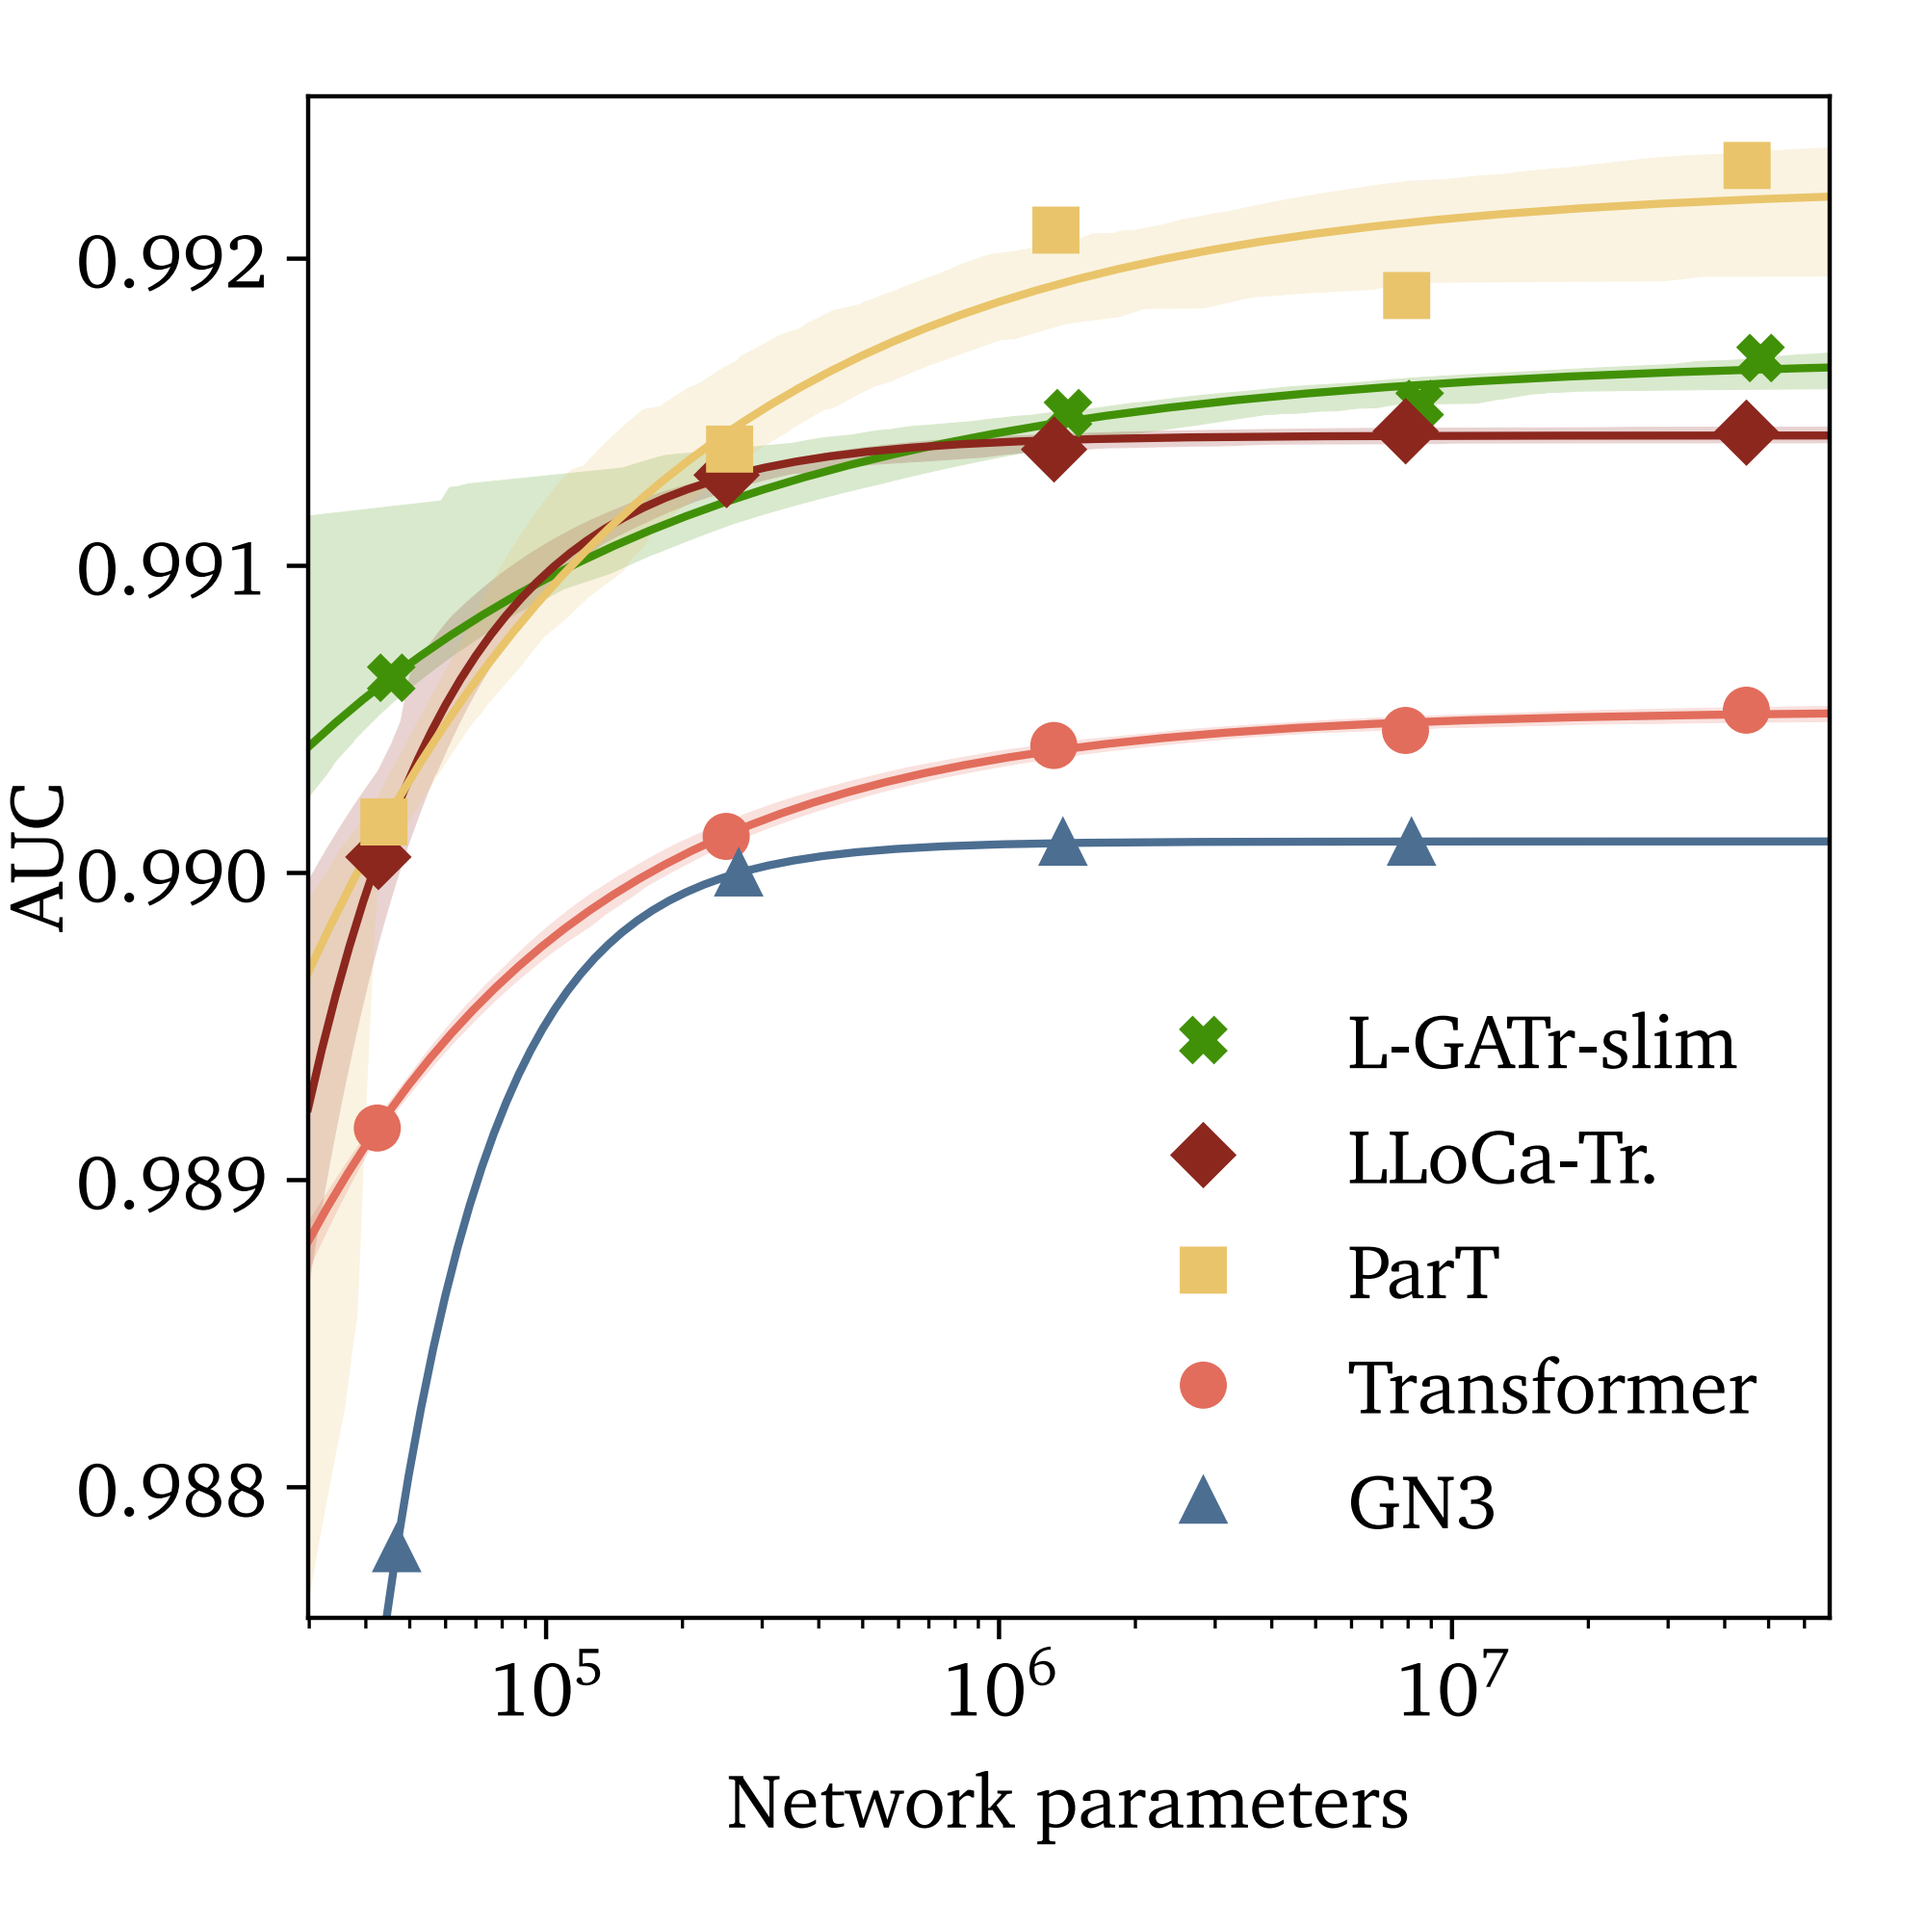
0.992
0.991
0.990
0.989
0.988
AUC
$10^5$
$10^6$
$10^7$
Network parameters
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GN3

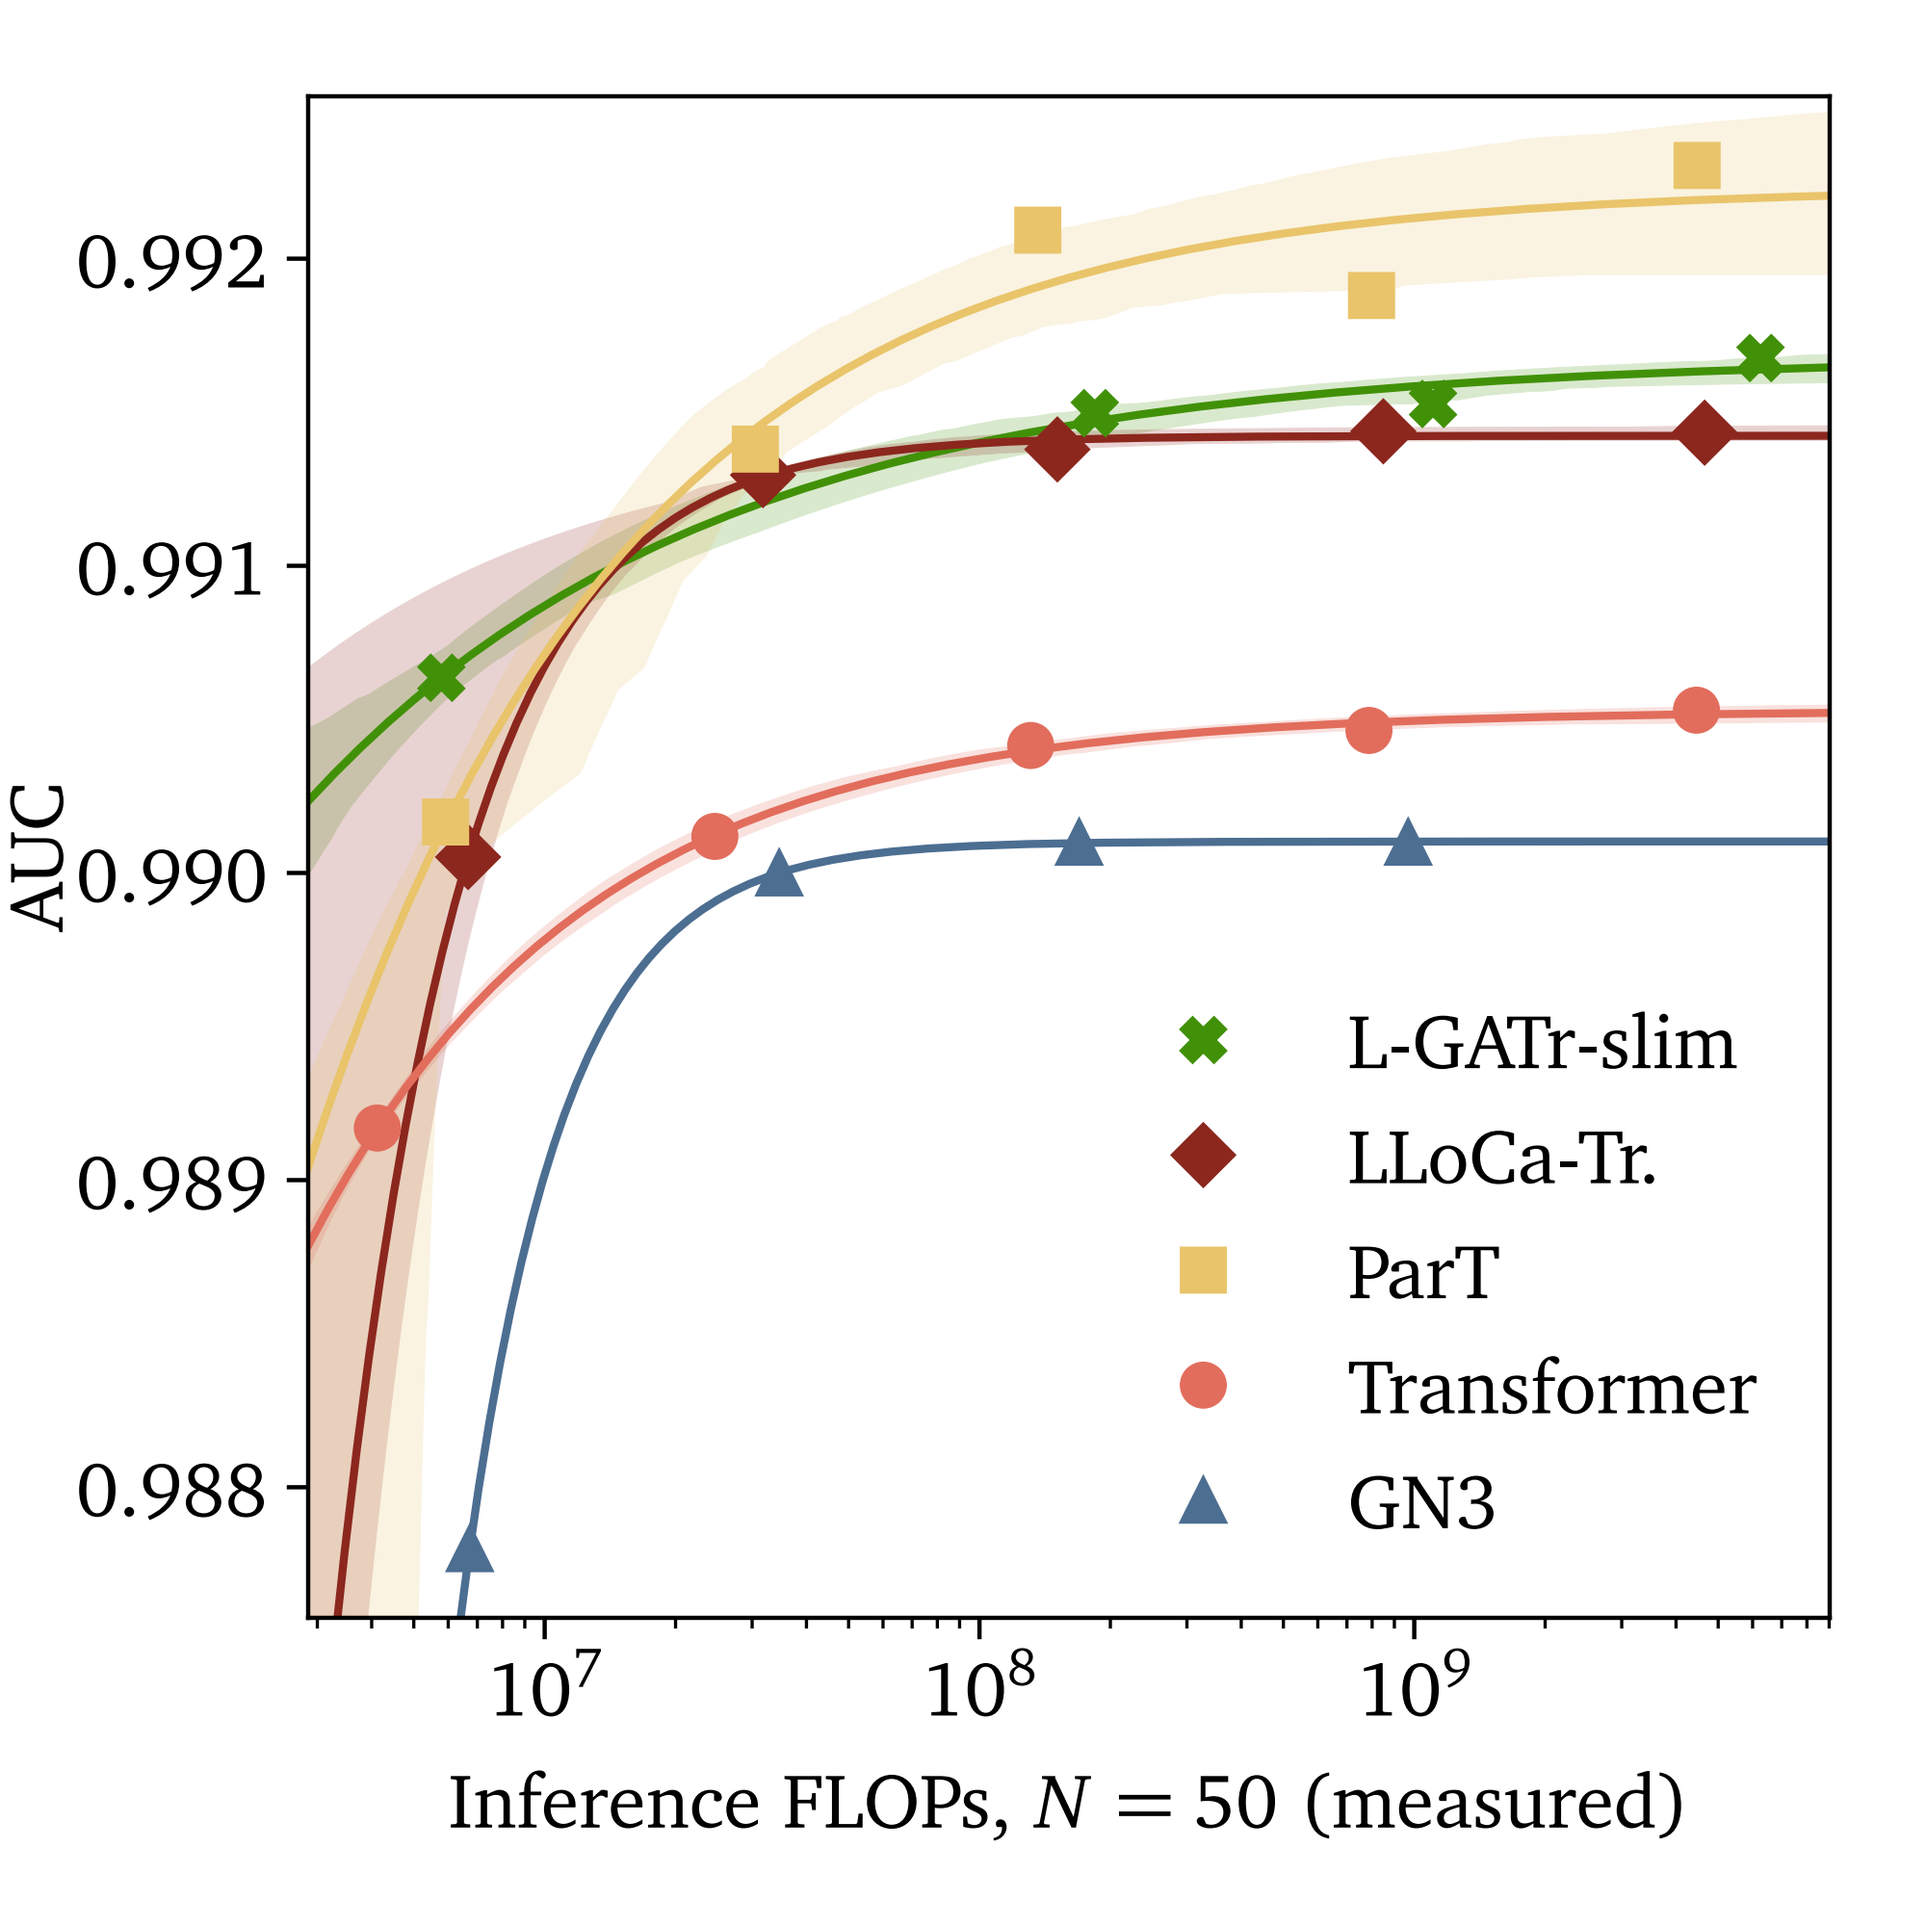
0.992
0.991
0.990
0.989
0.988
AUC
$10^7$
$10^8$
$10^9$
Inference FLOPs, $N = 50$ (measured)
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GN3

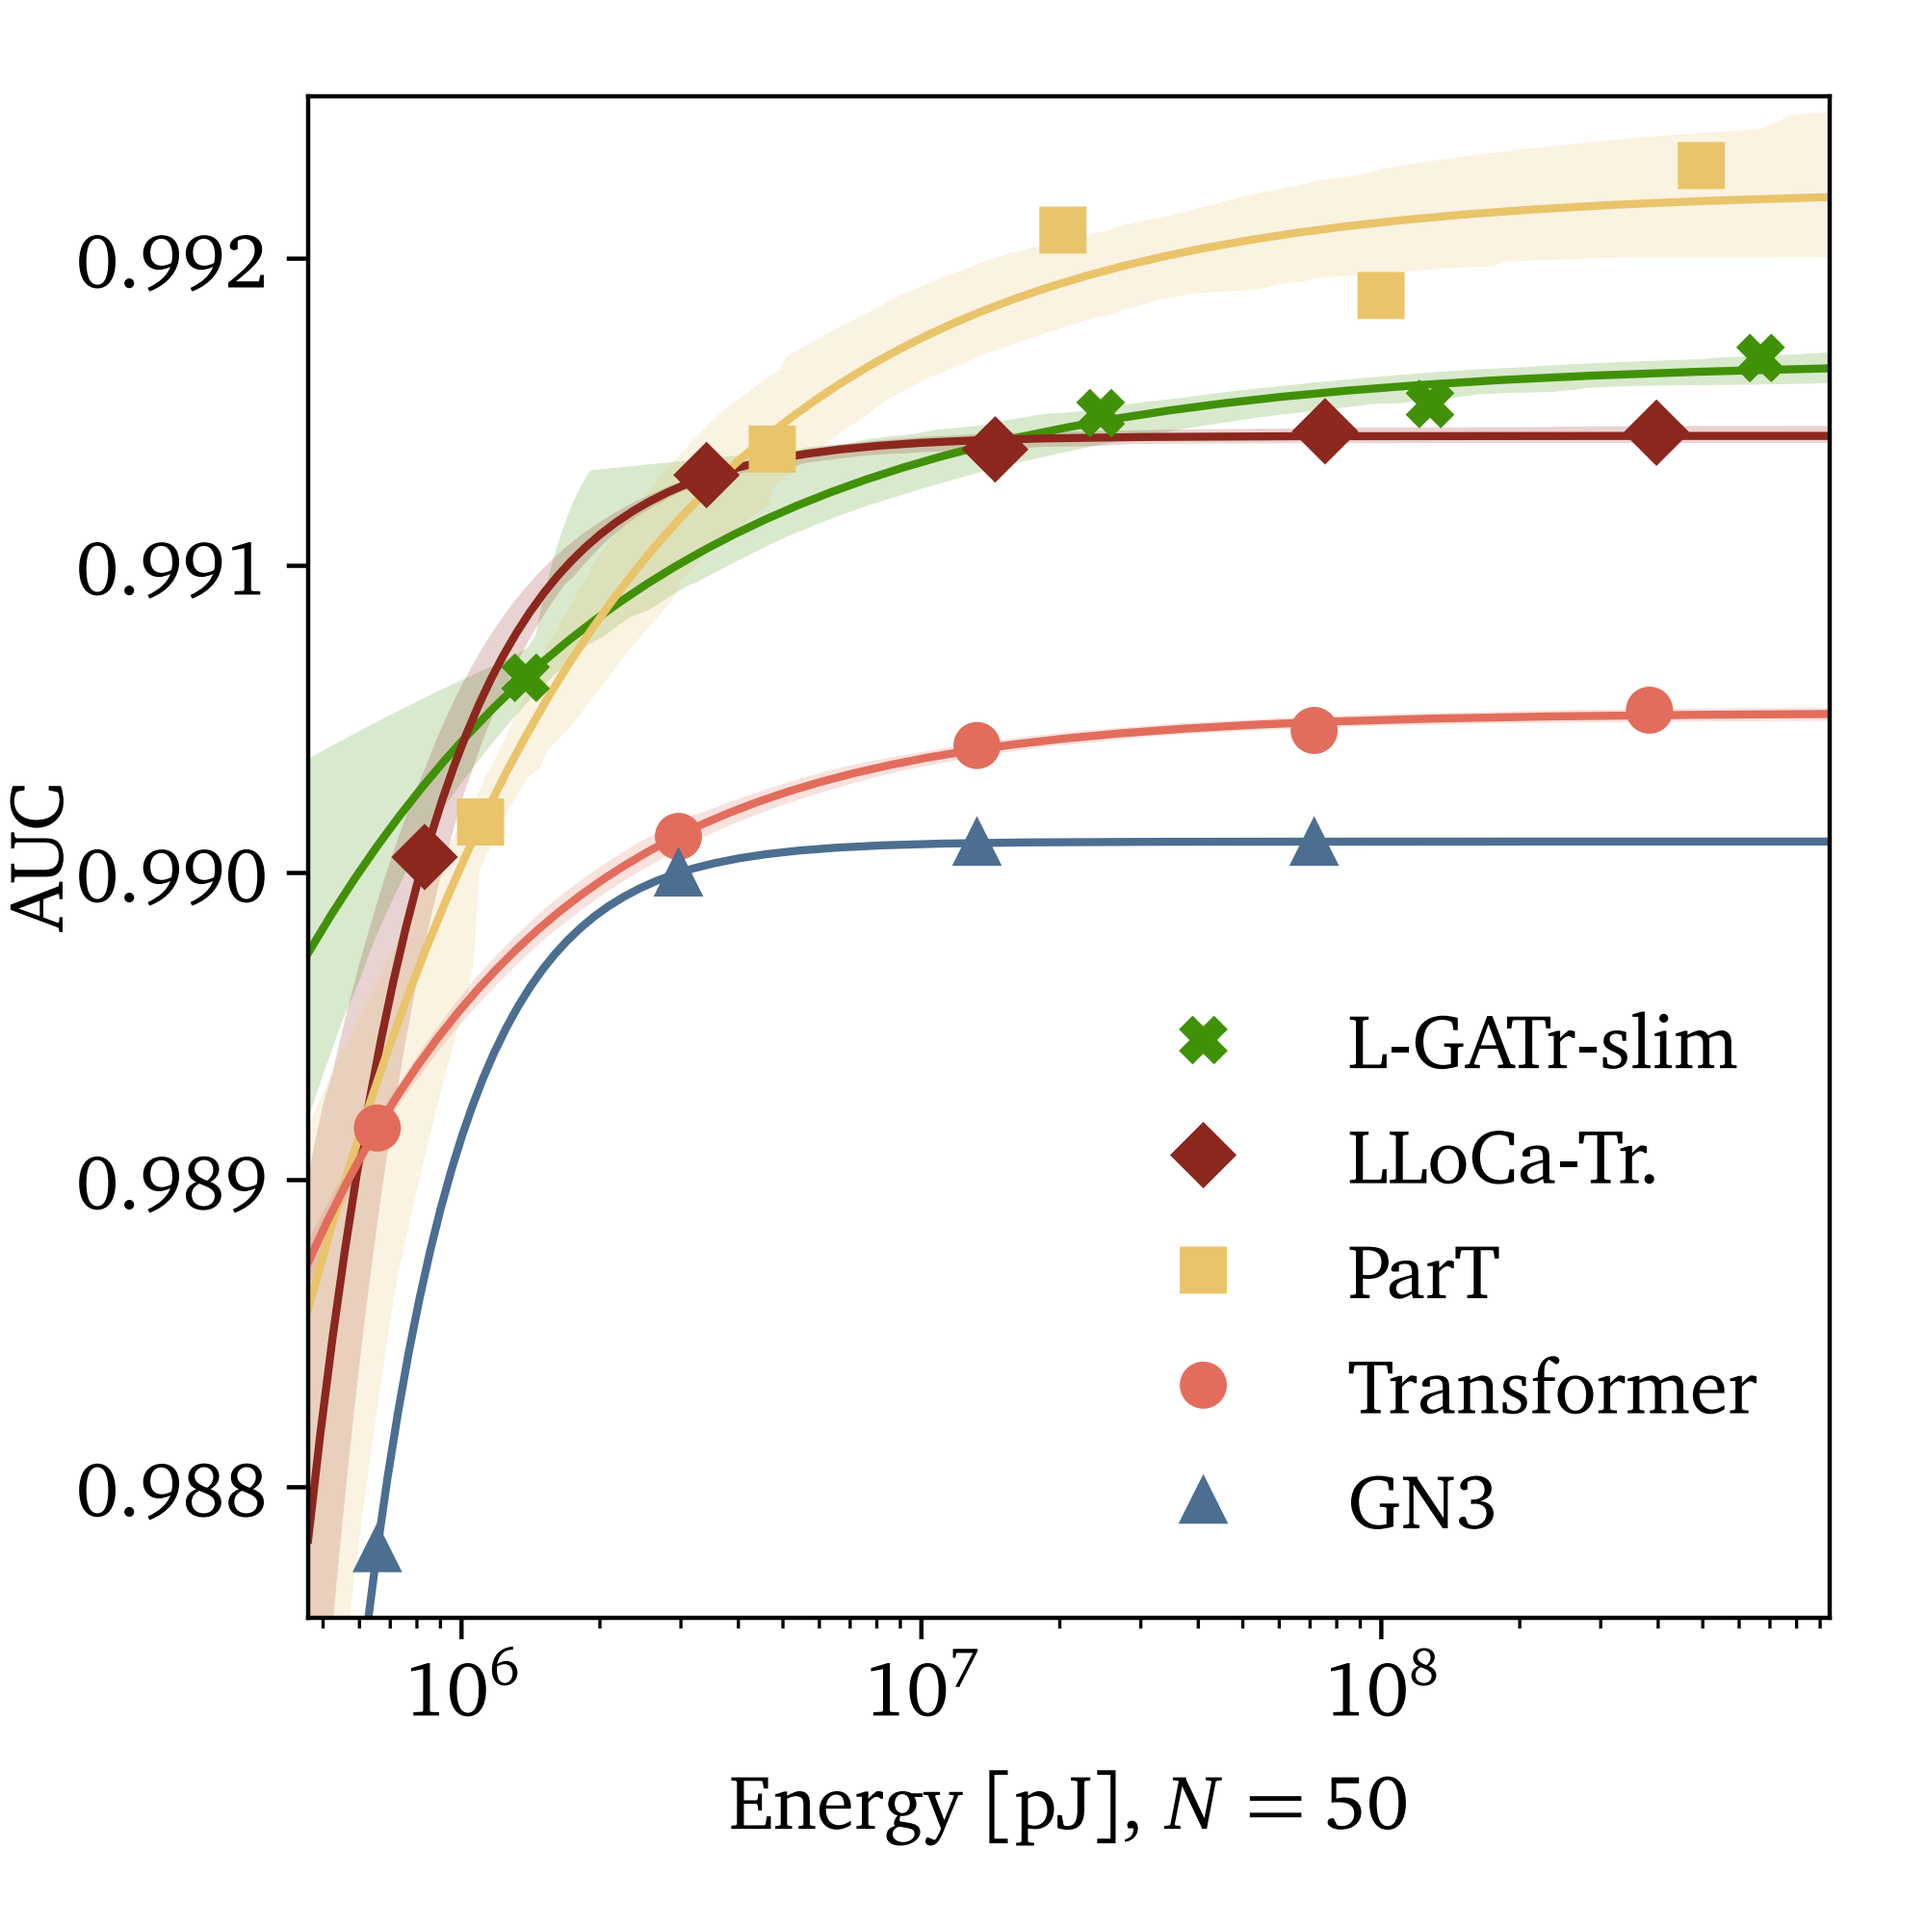
0.992
0.991
0.990
0.989
0.988
AUC
$10^6$
$10^7$
$10^8$
Energy [pJ], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GN3

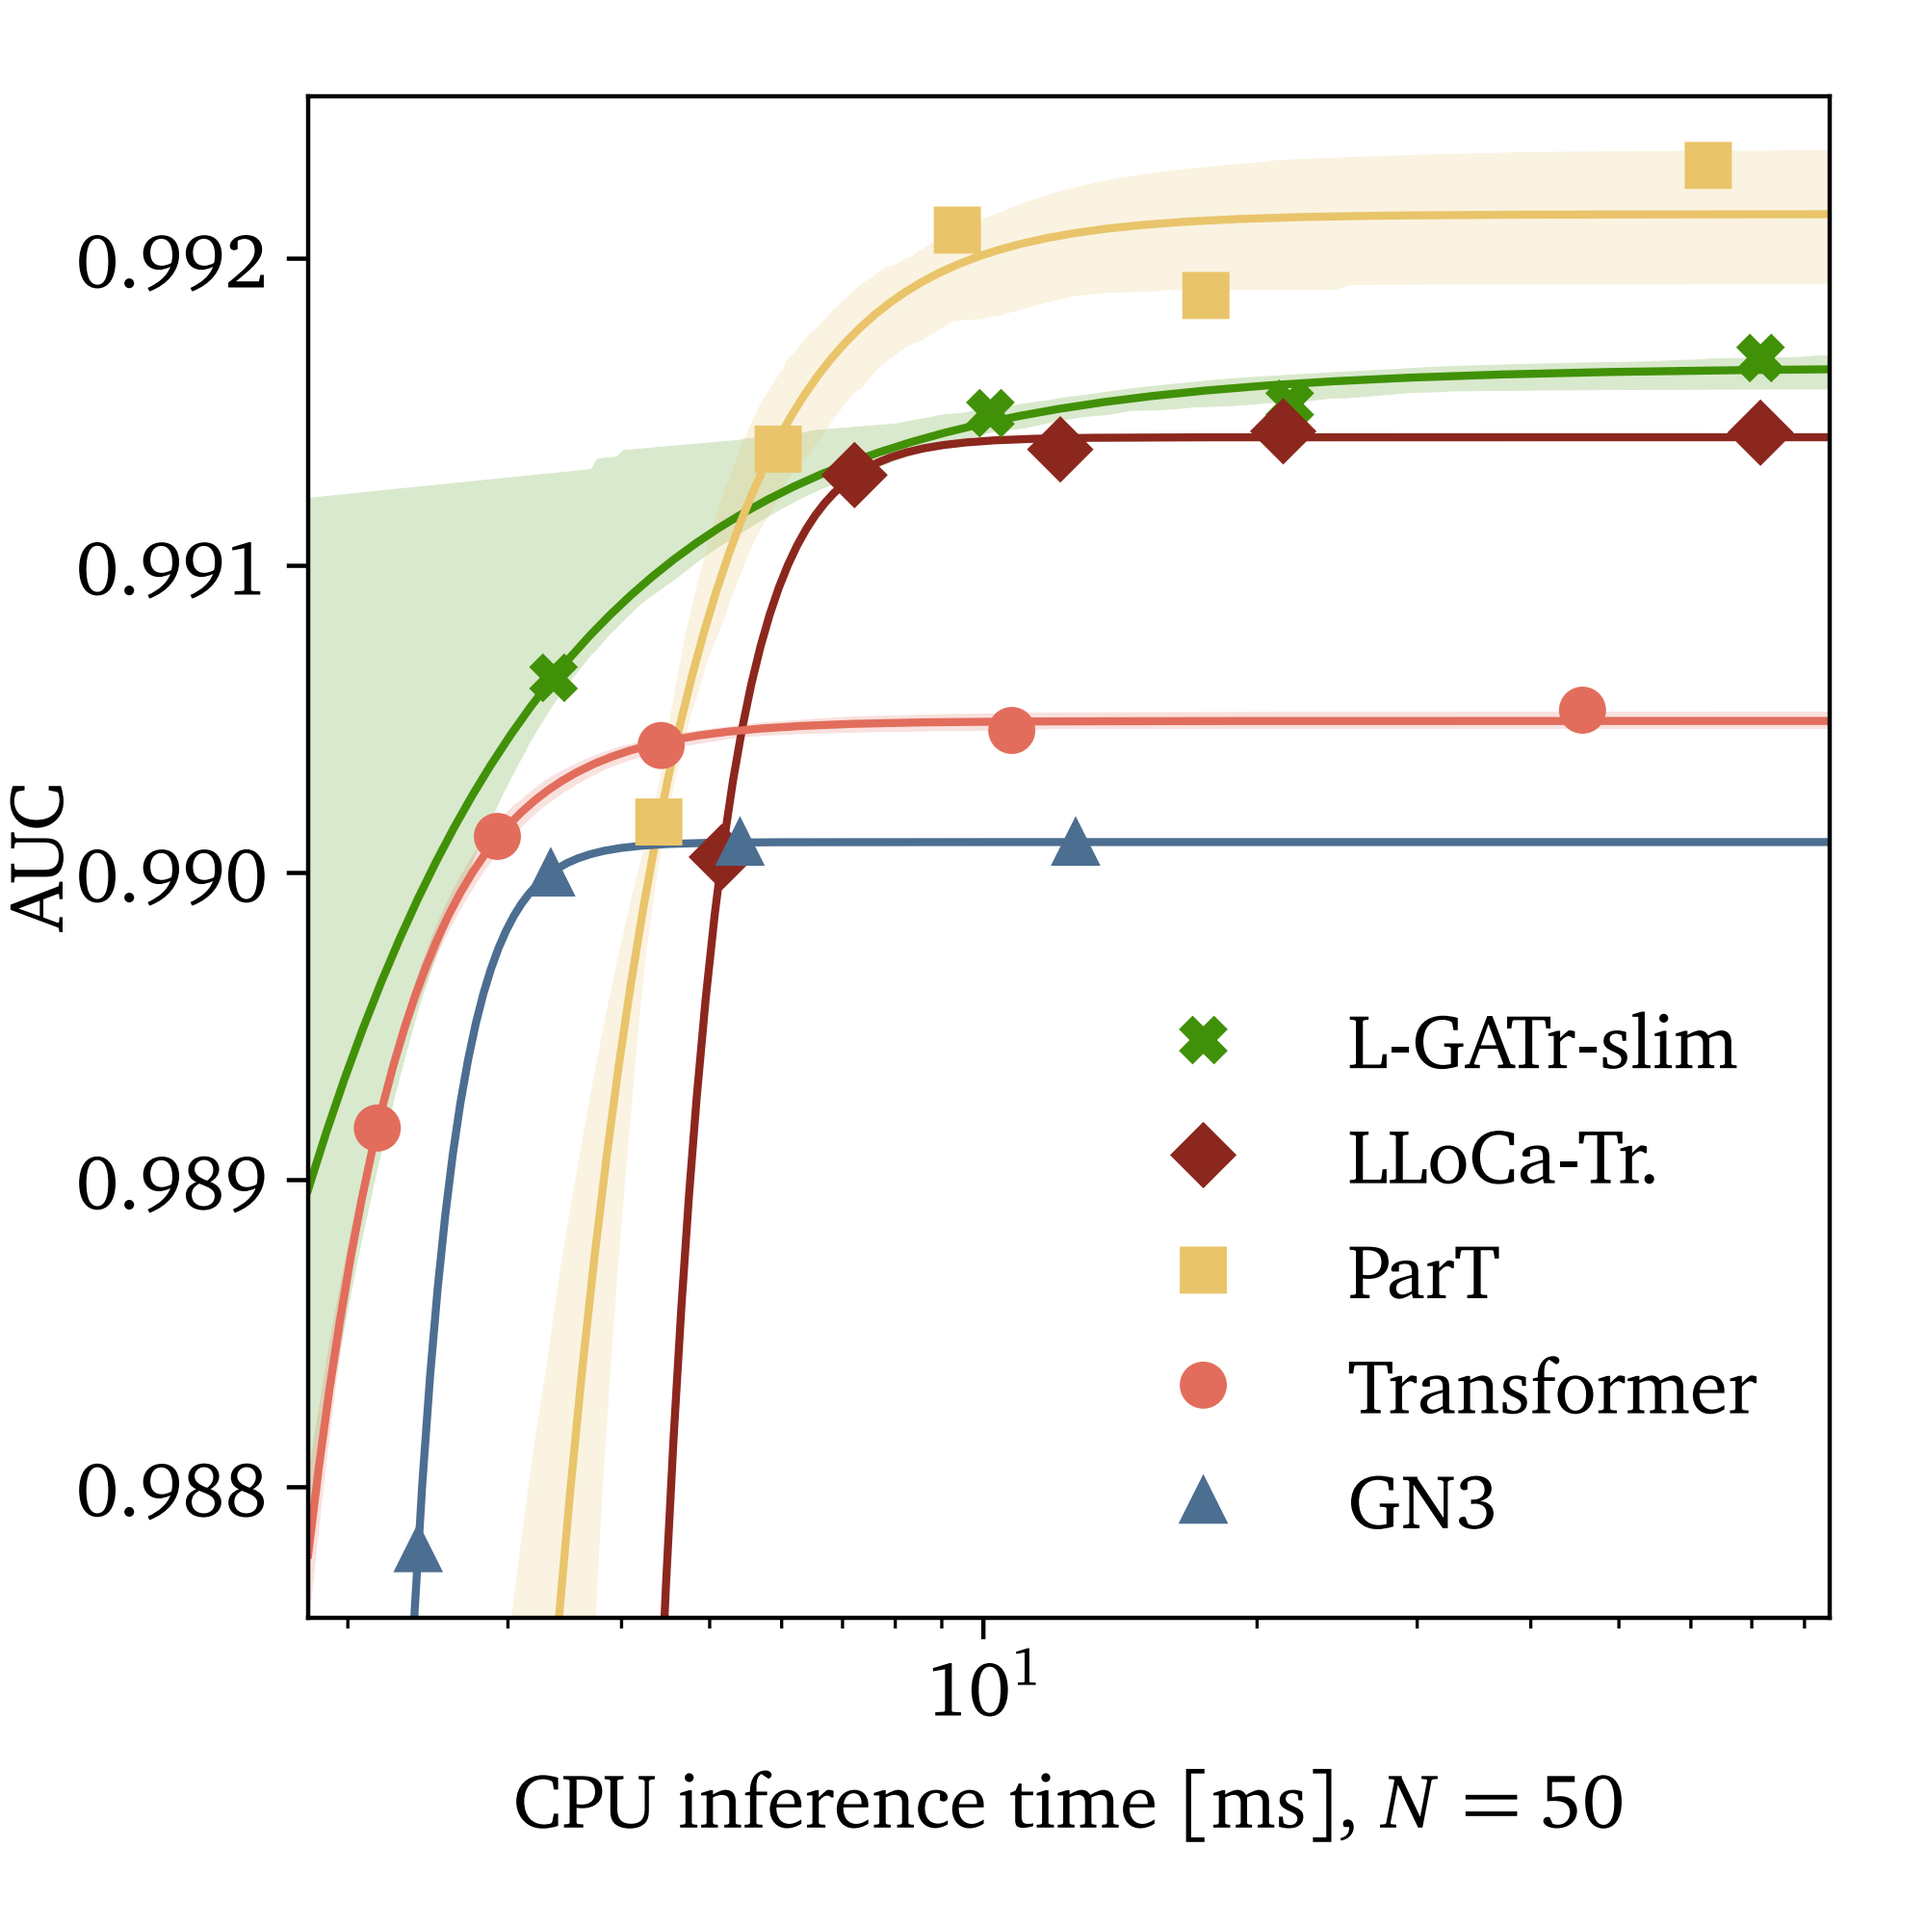
0.992
0.991
0.990
0.989
0.988
AUC
$10^1$
CPU inference time [ms], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GN3

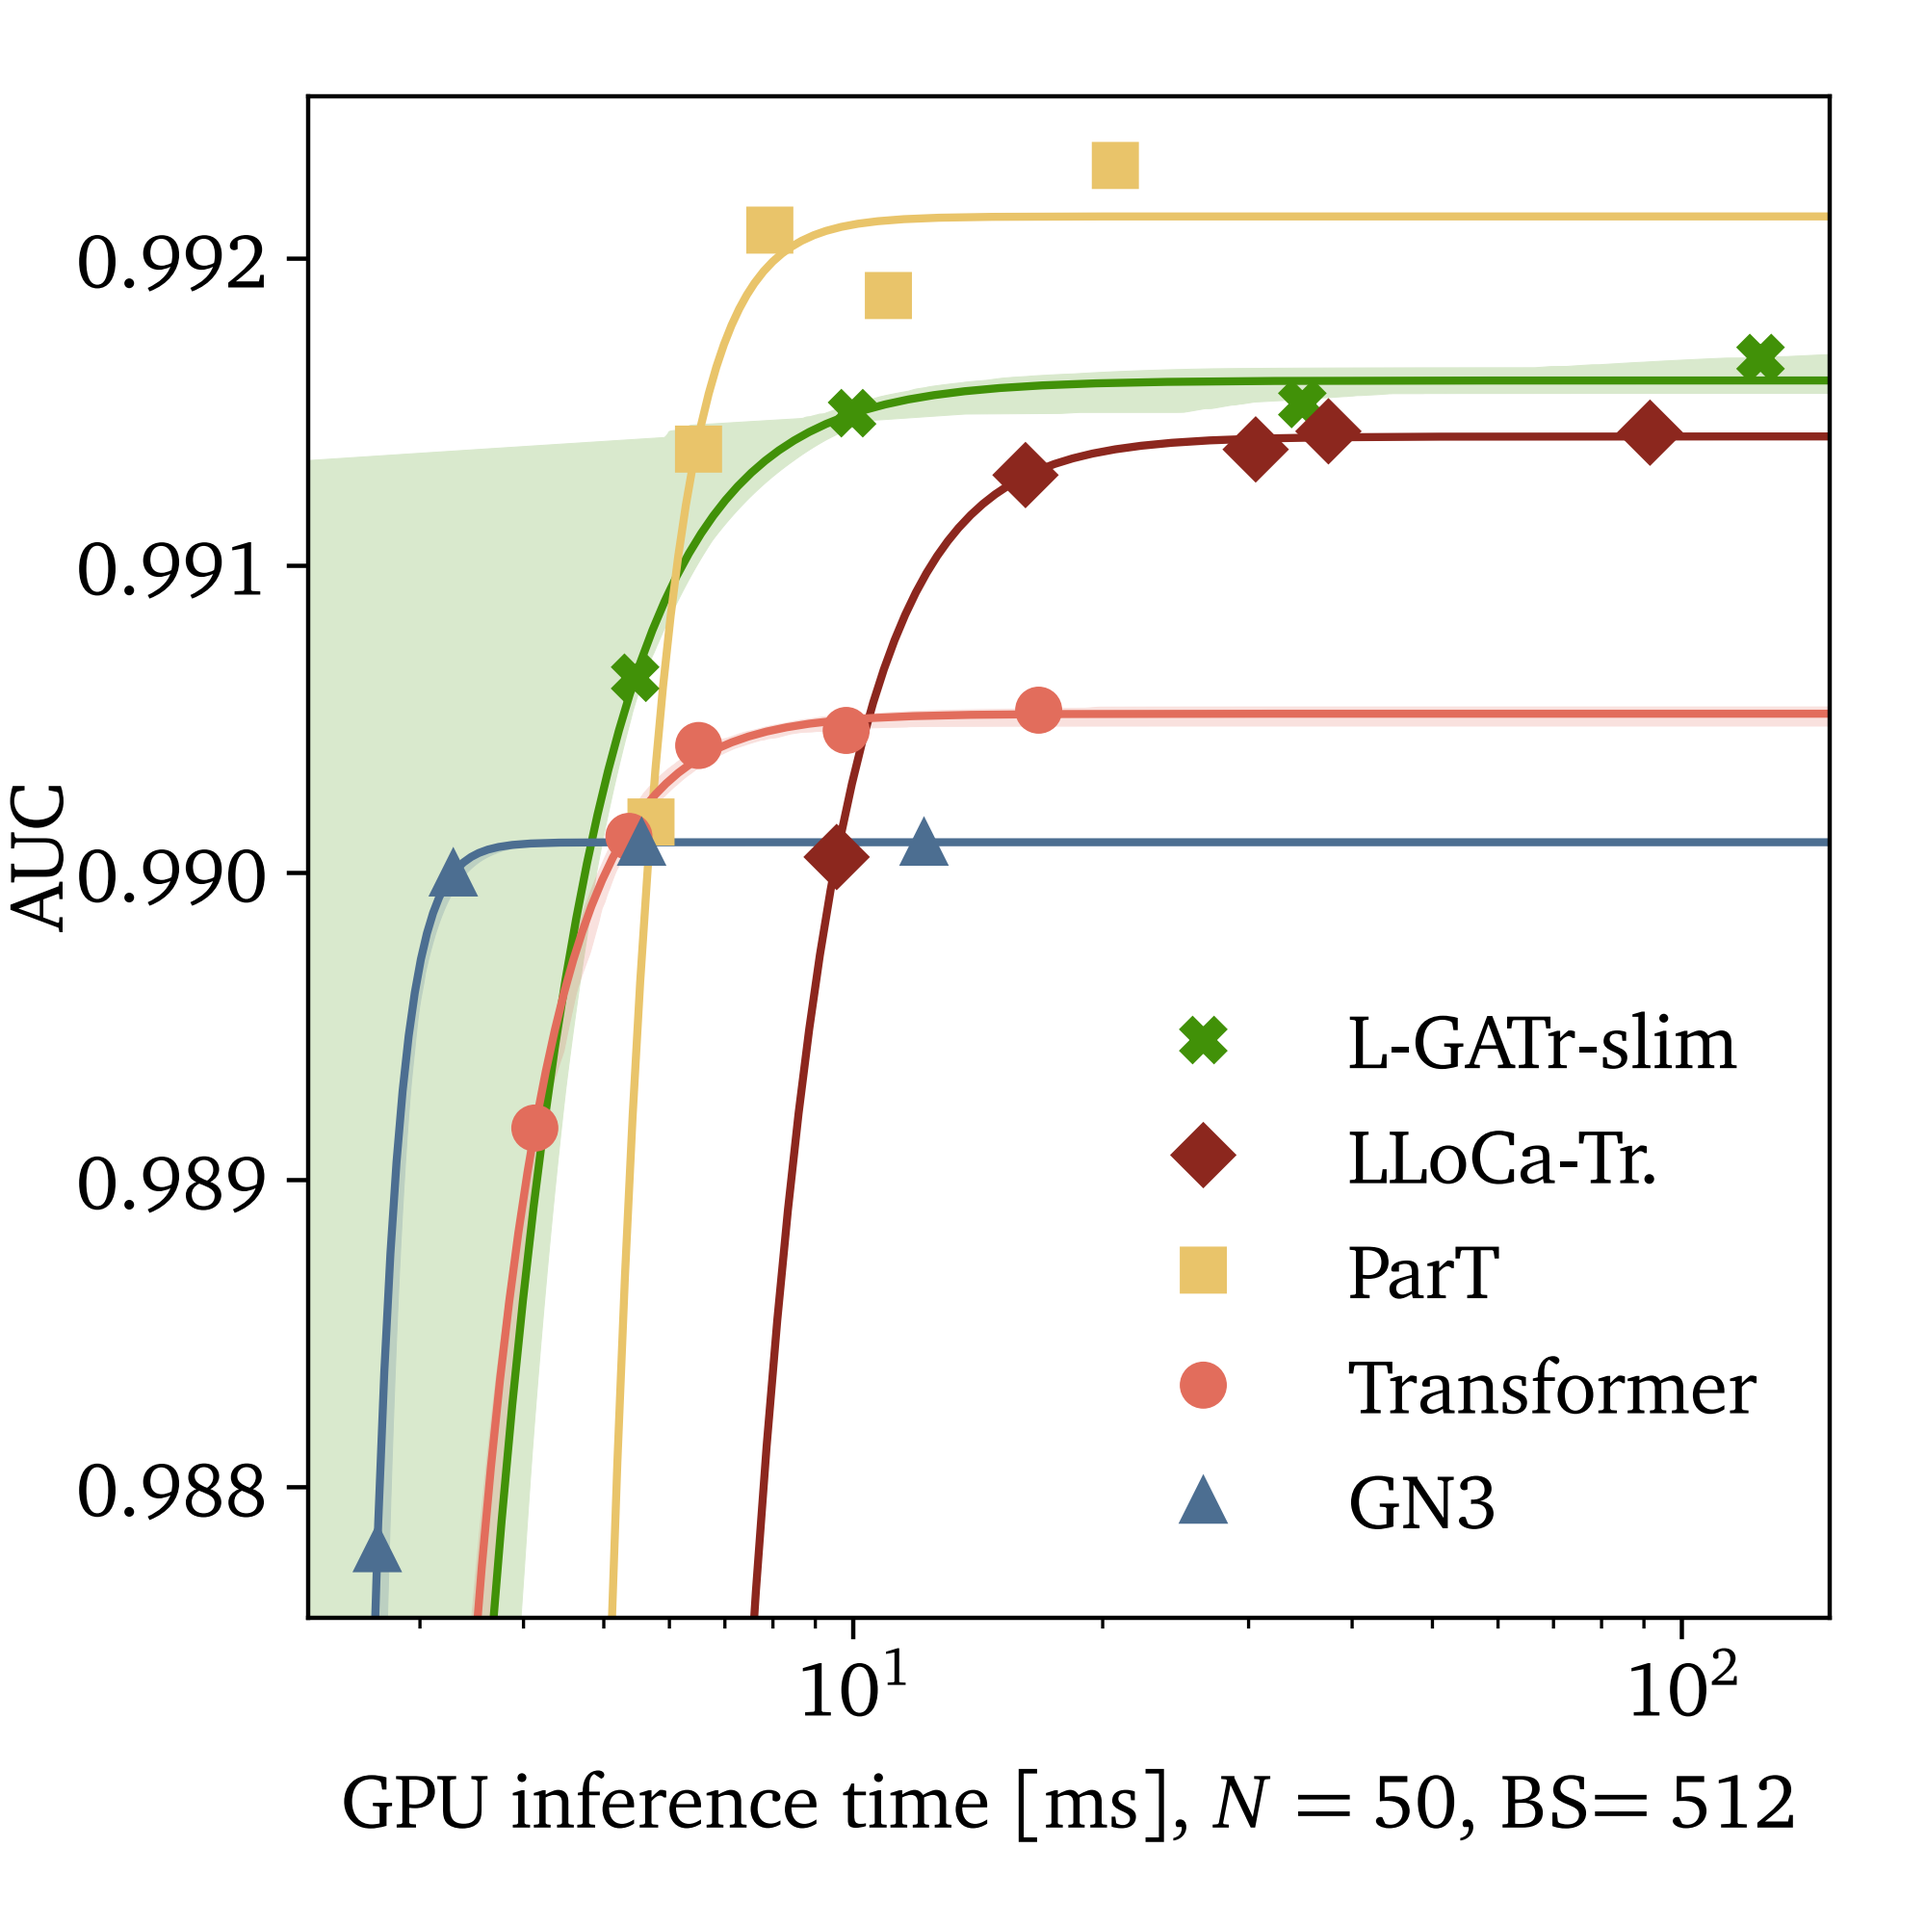
0.992
0.991
0.990
0.989
0.988
AUC
$10^1$
$10^2$
GPU inference time [ms], $N = 50$, BS$= 512$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GN3

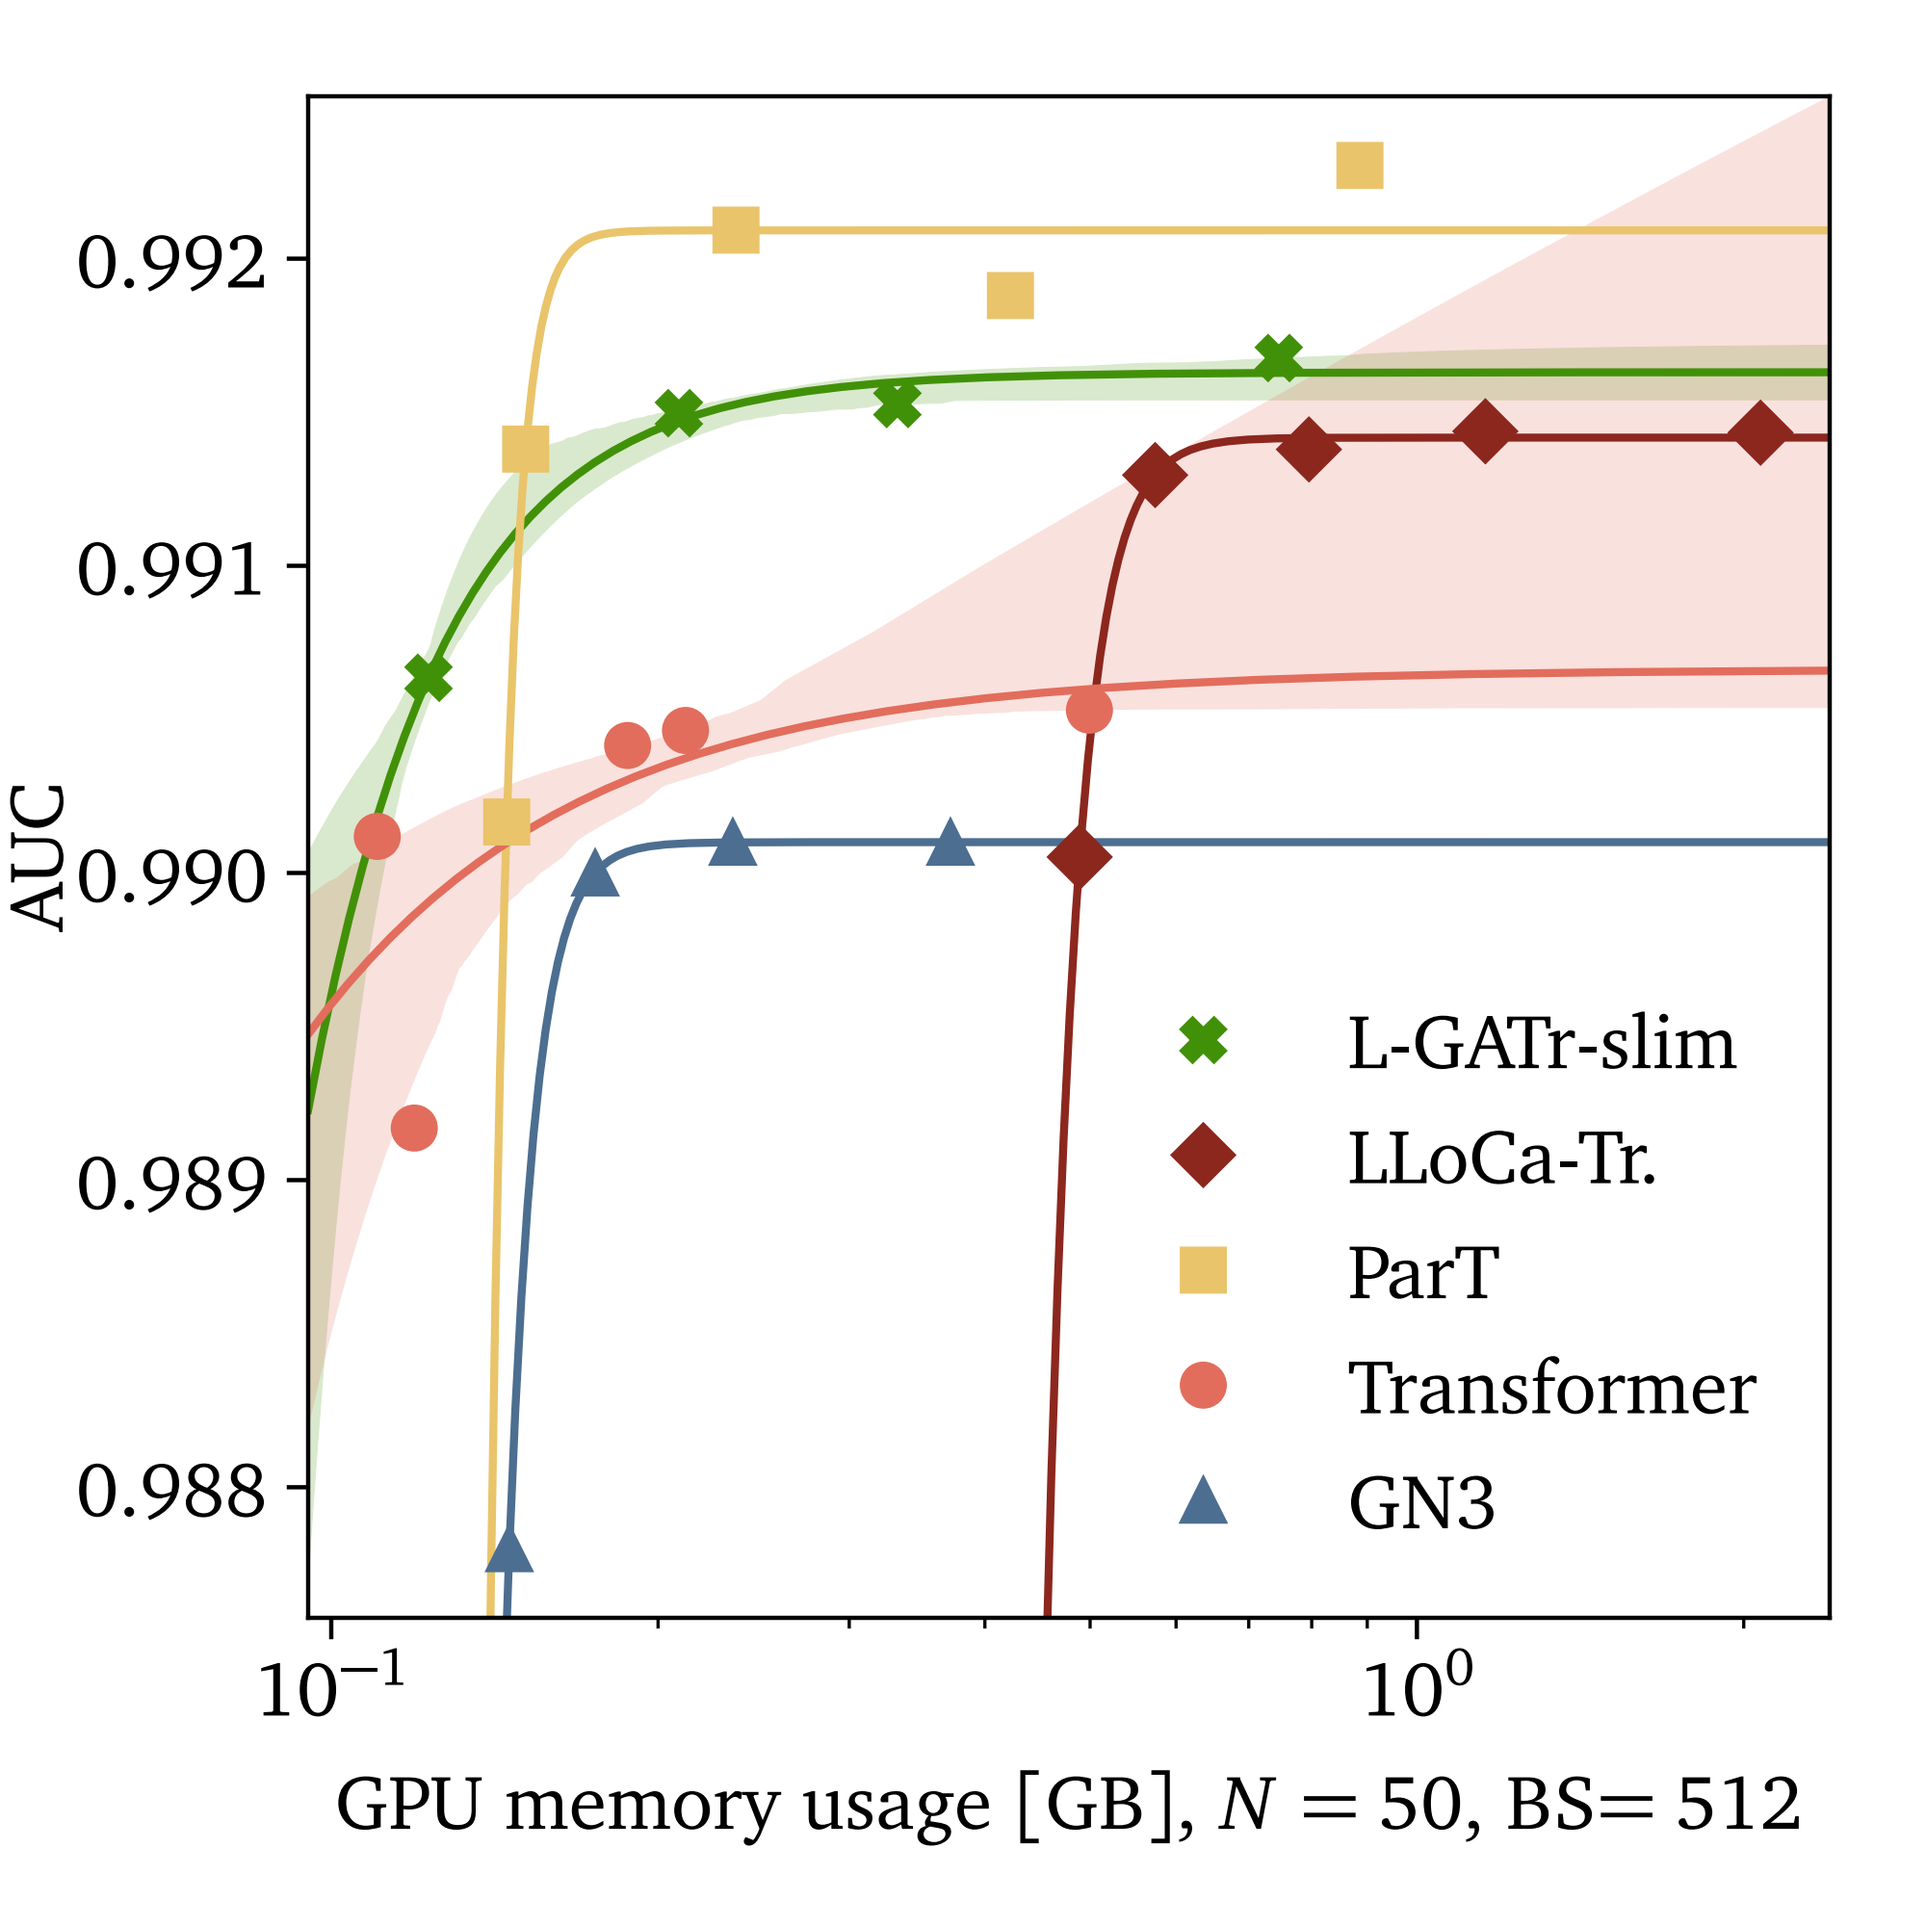
0.992
0.991
0.990
0.989
0.988
AUC
$10^{-1}$
$10^{0}$
GPU memory usage [GB], $N = 50$, BS$= 512$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GN3

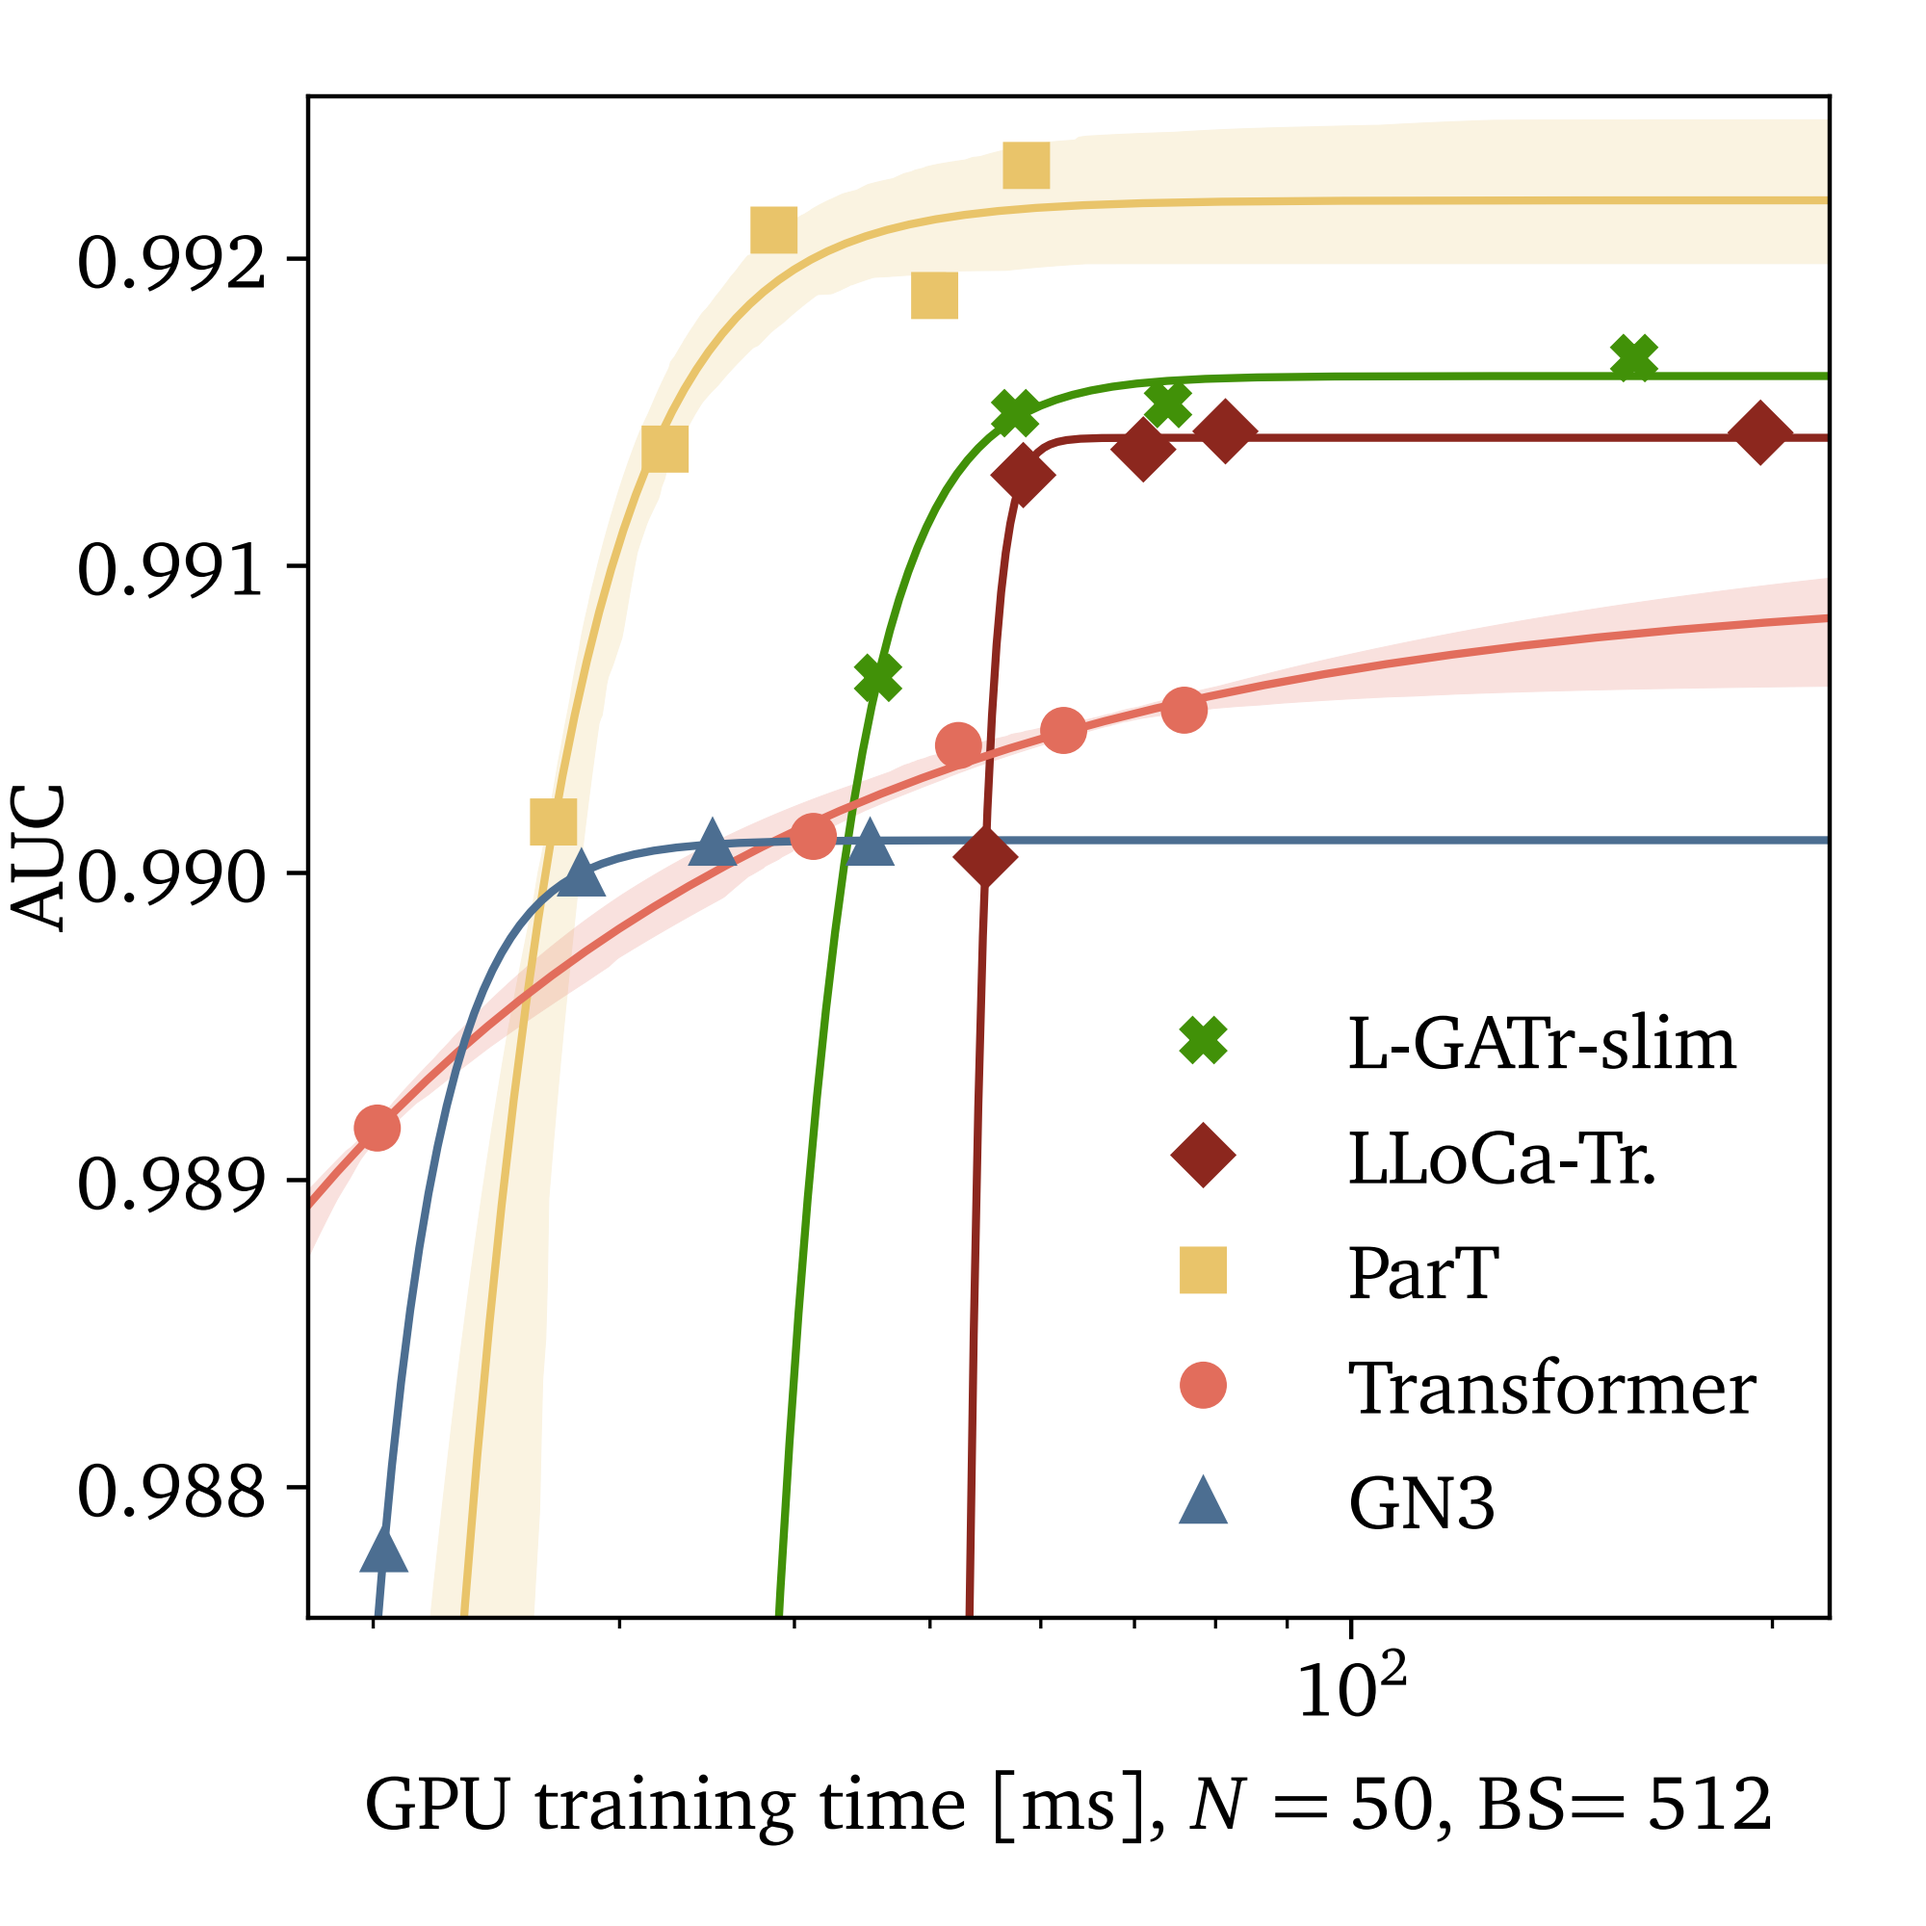
0.992
0.991
0.990
0.989
0.988
AUC
$10^2$
GPU training time [ms], $N = 50$, BS$= 512$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GN3

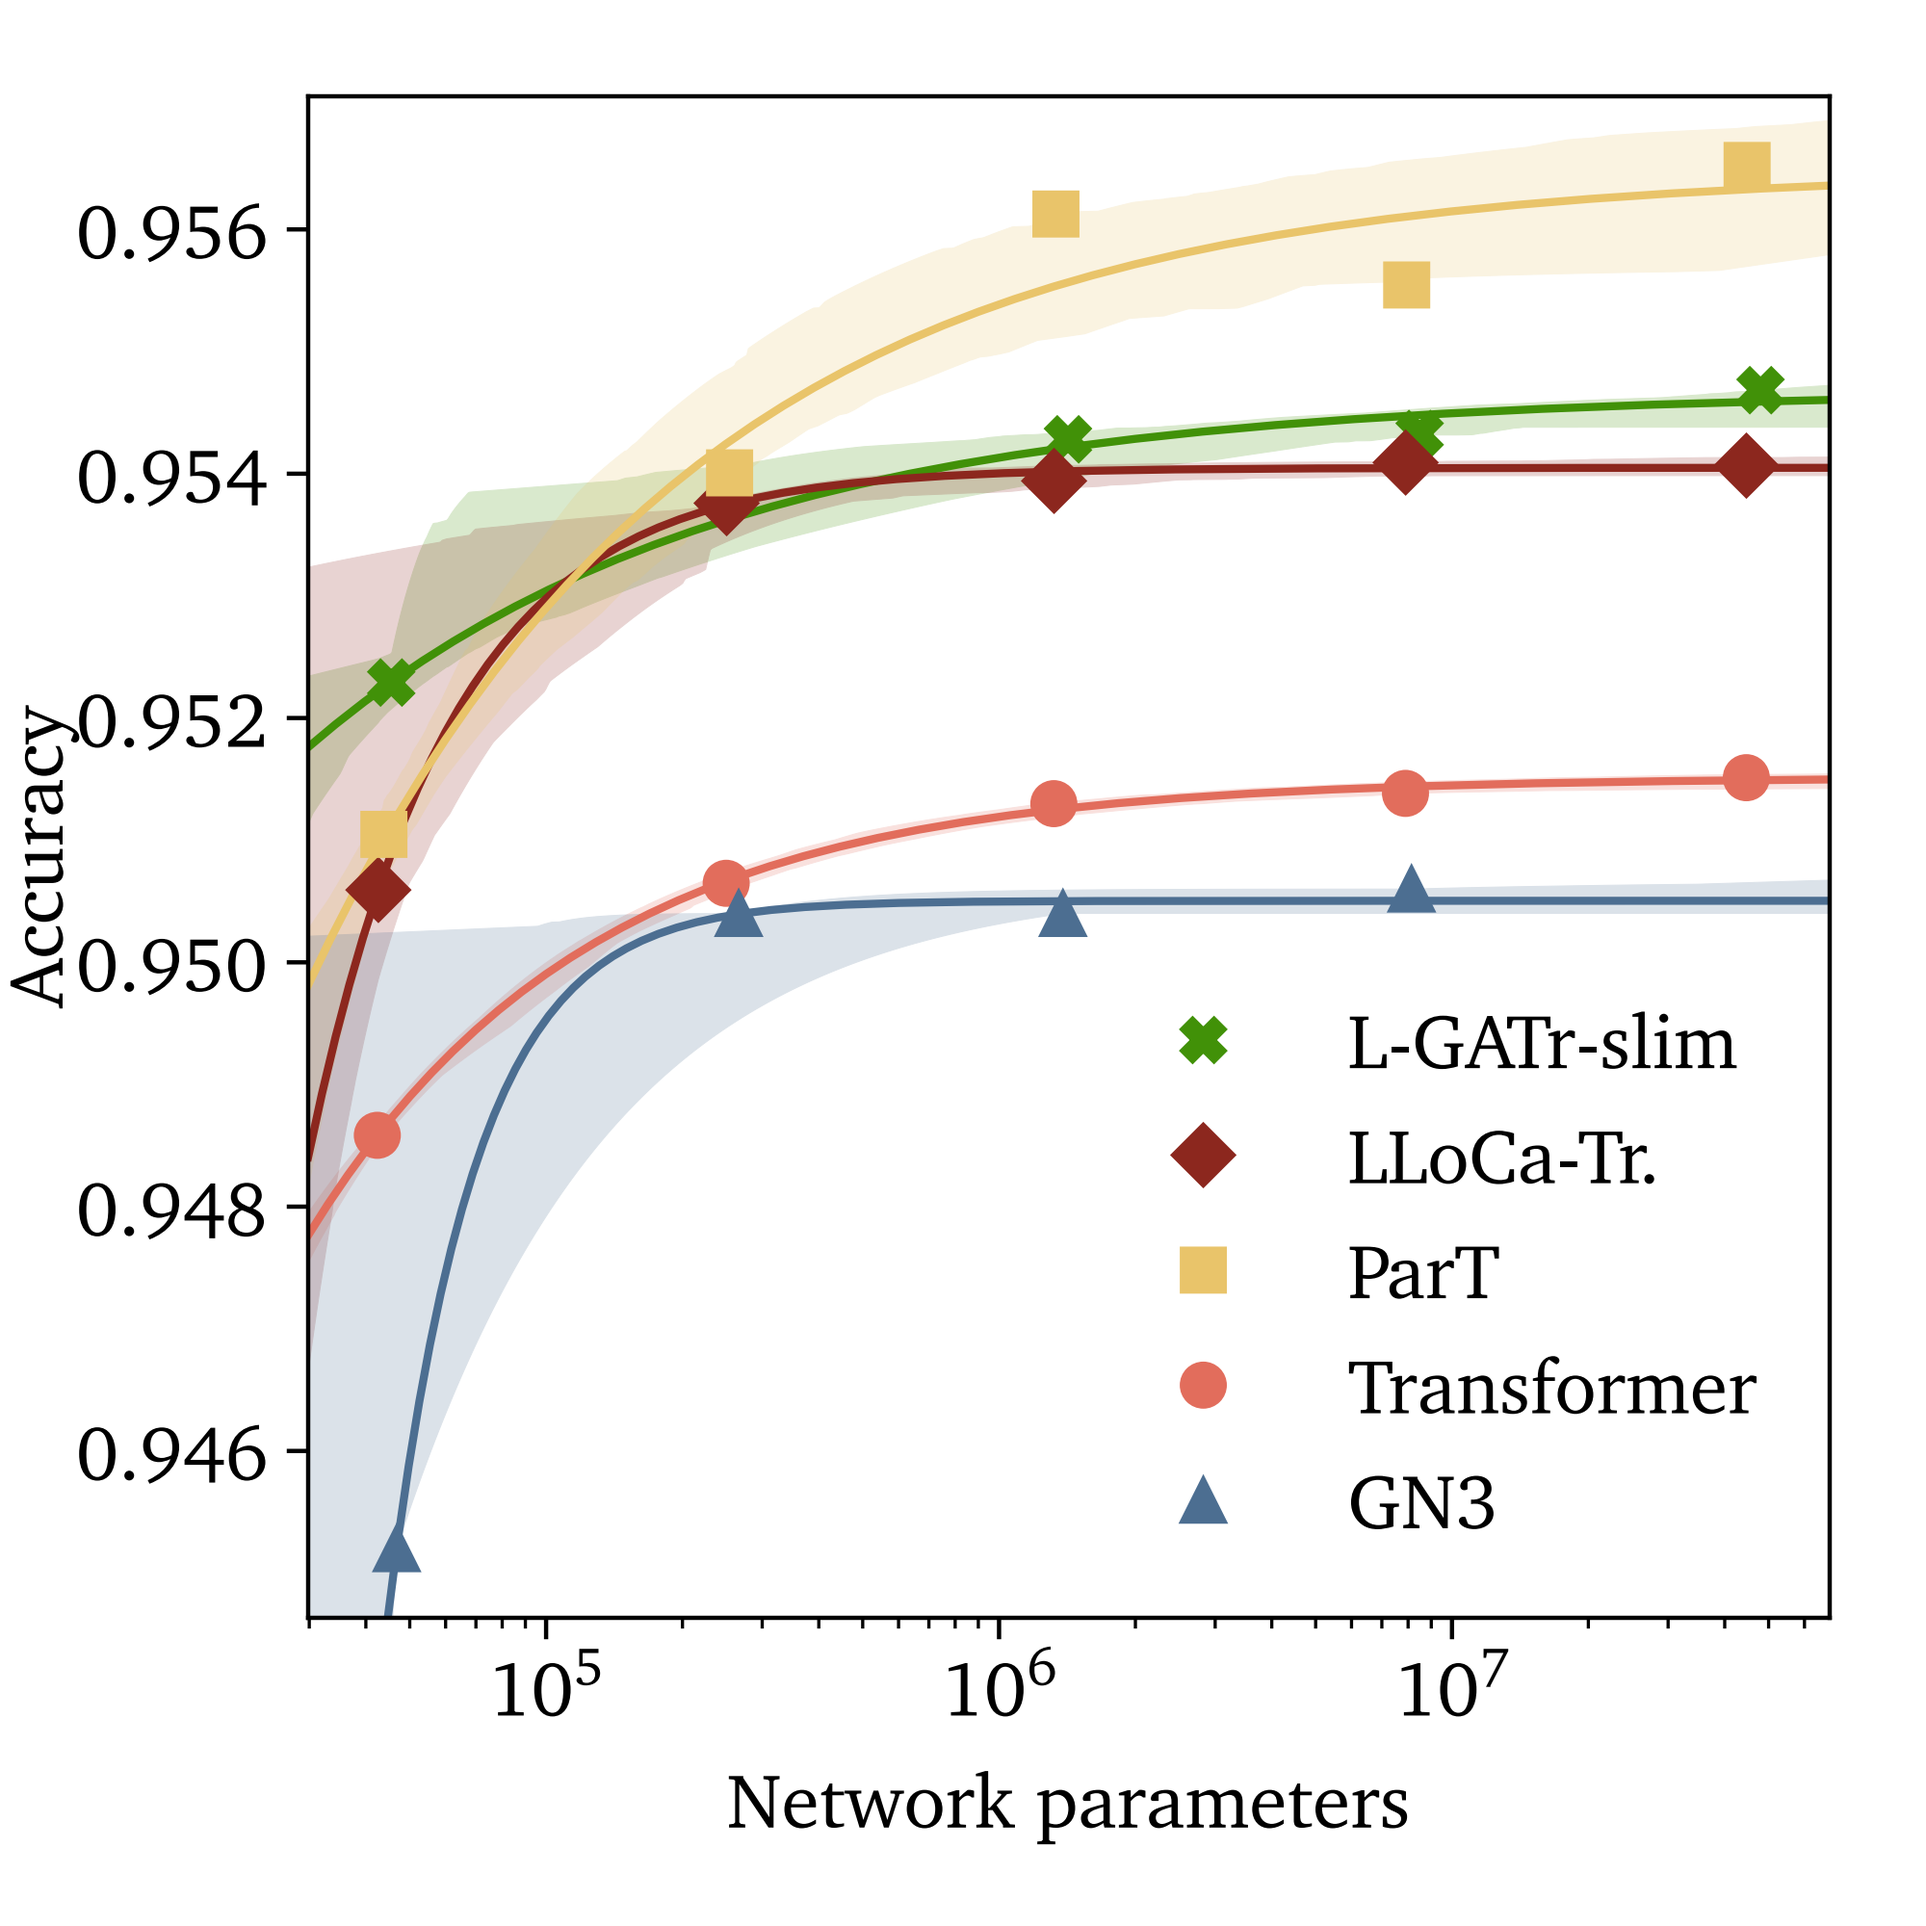
0.956
0.954
0.952
0.950
0.948
0.946
Accuracy
$10^5$
$10^6$
$10^7$
Network parameters
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GN3

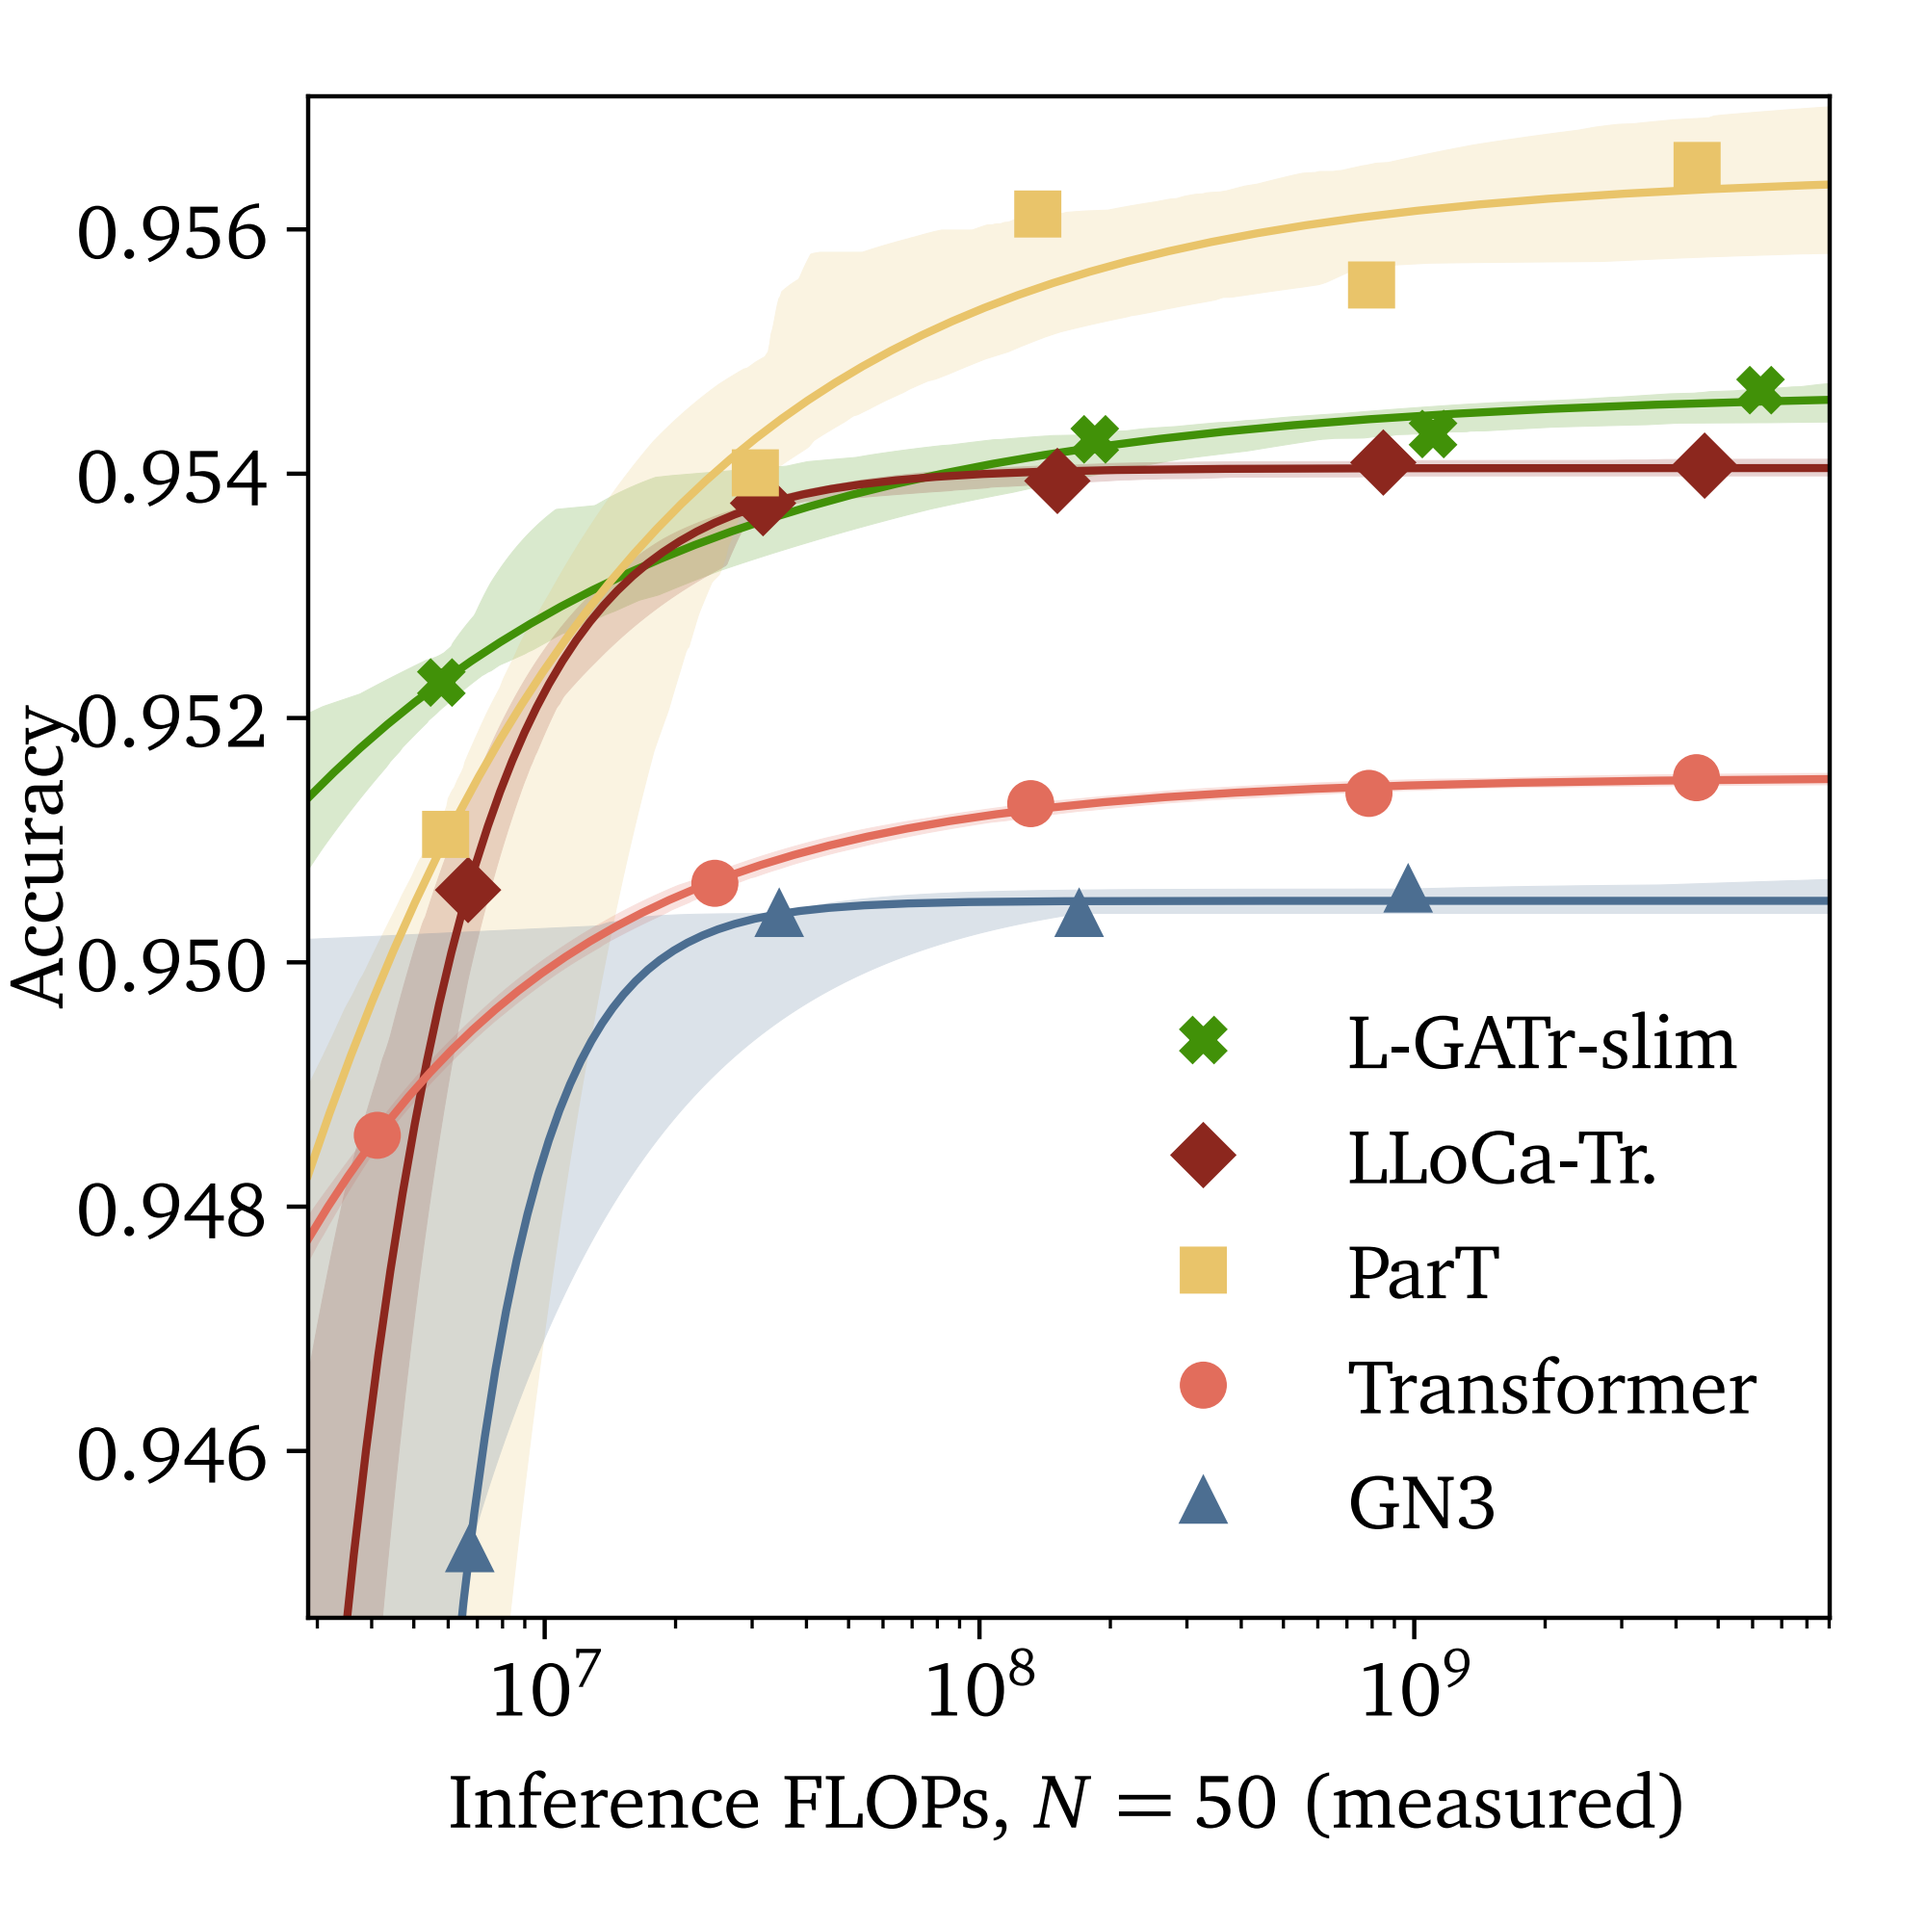
0.956
0.954
0.952
0.950
0.948
0.946
Accuracy
$10^7$
$10^8$
$10^9$
Inference FLOPs, $N = 50$ (measured)
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GN3

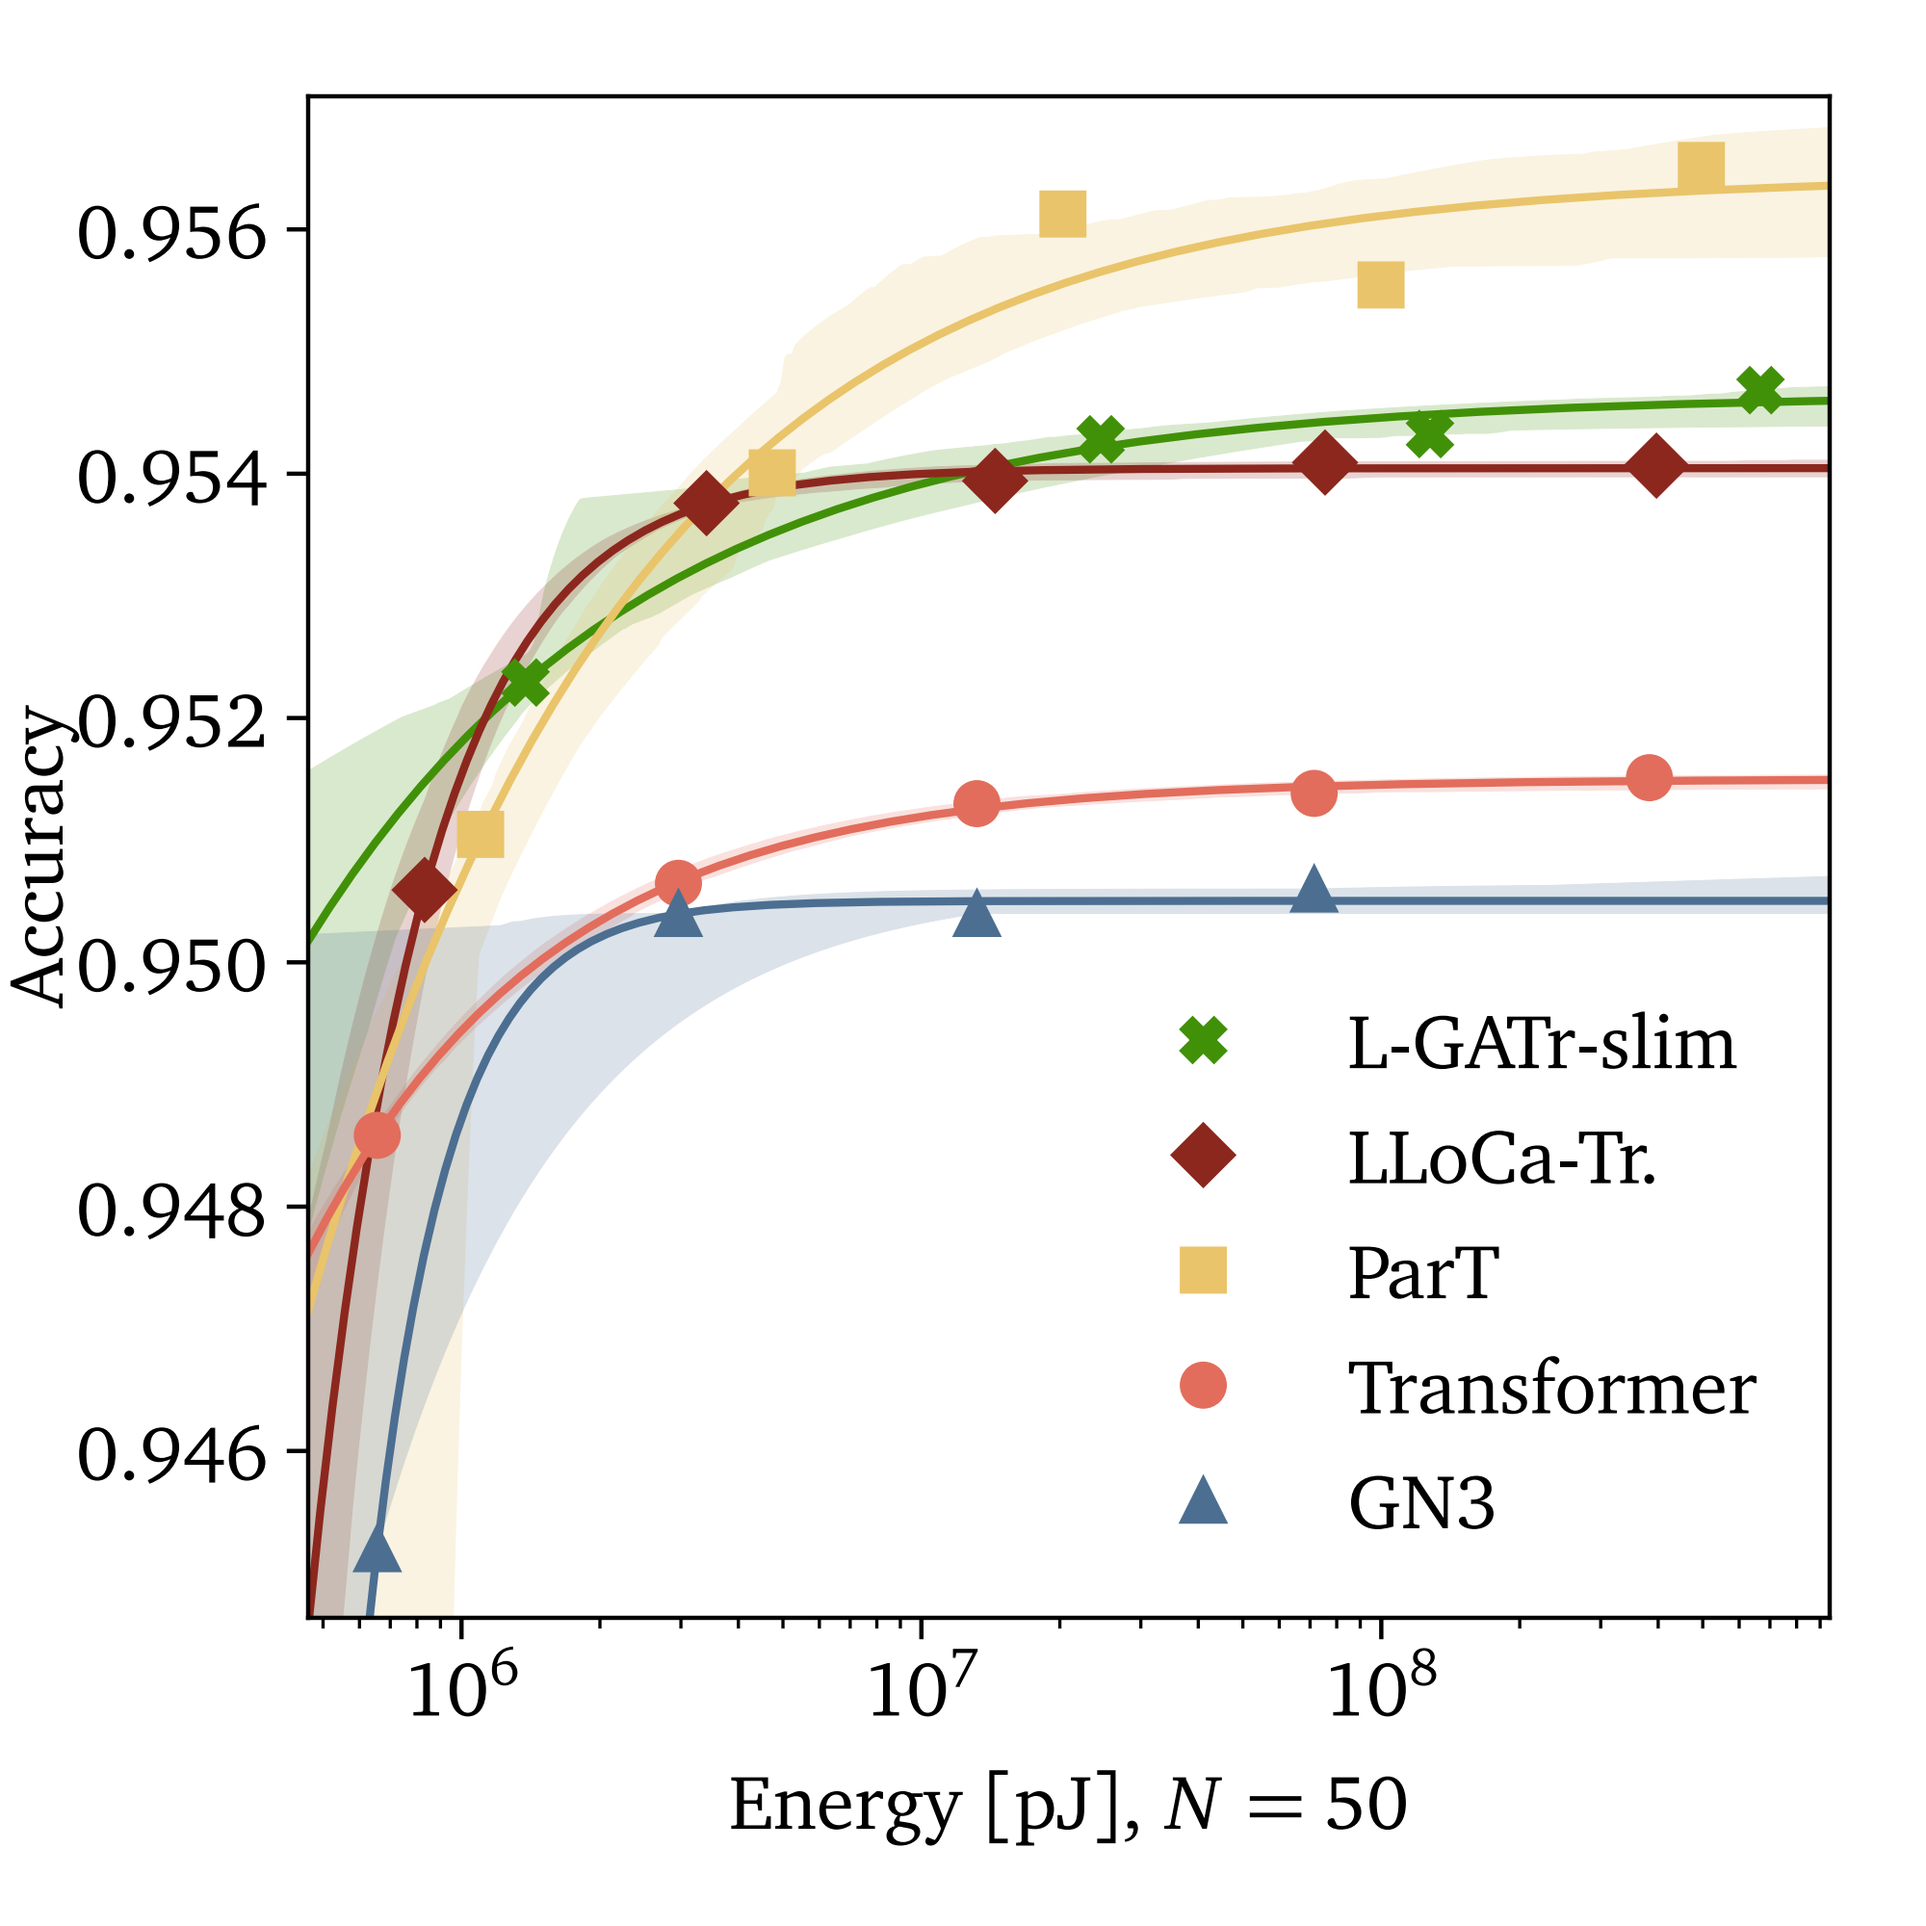
0.956
0.954
0.952
0.950
0.948
0.946
Accuracy
$10^6$
$10^7$
$10^8$
Energy [pJ], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GN3

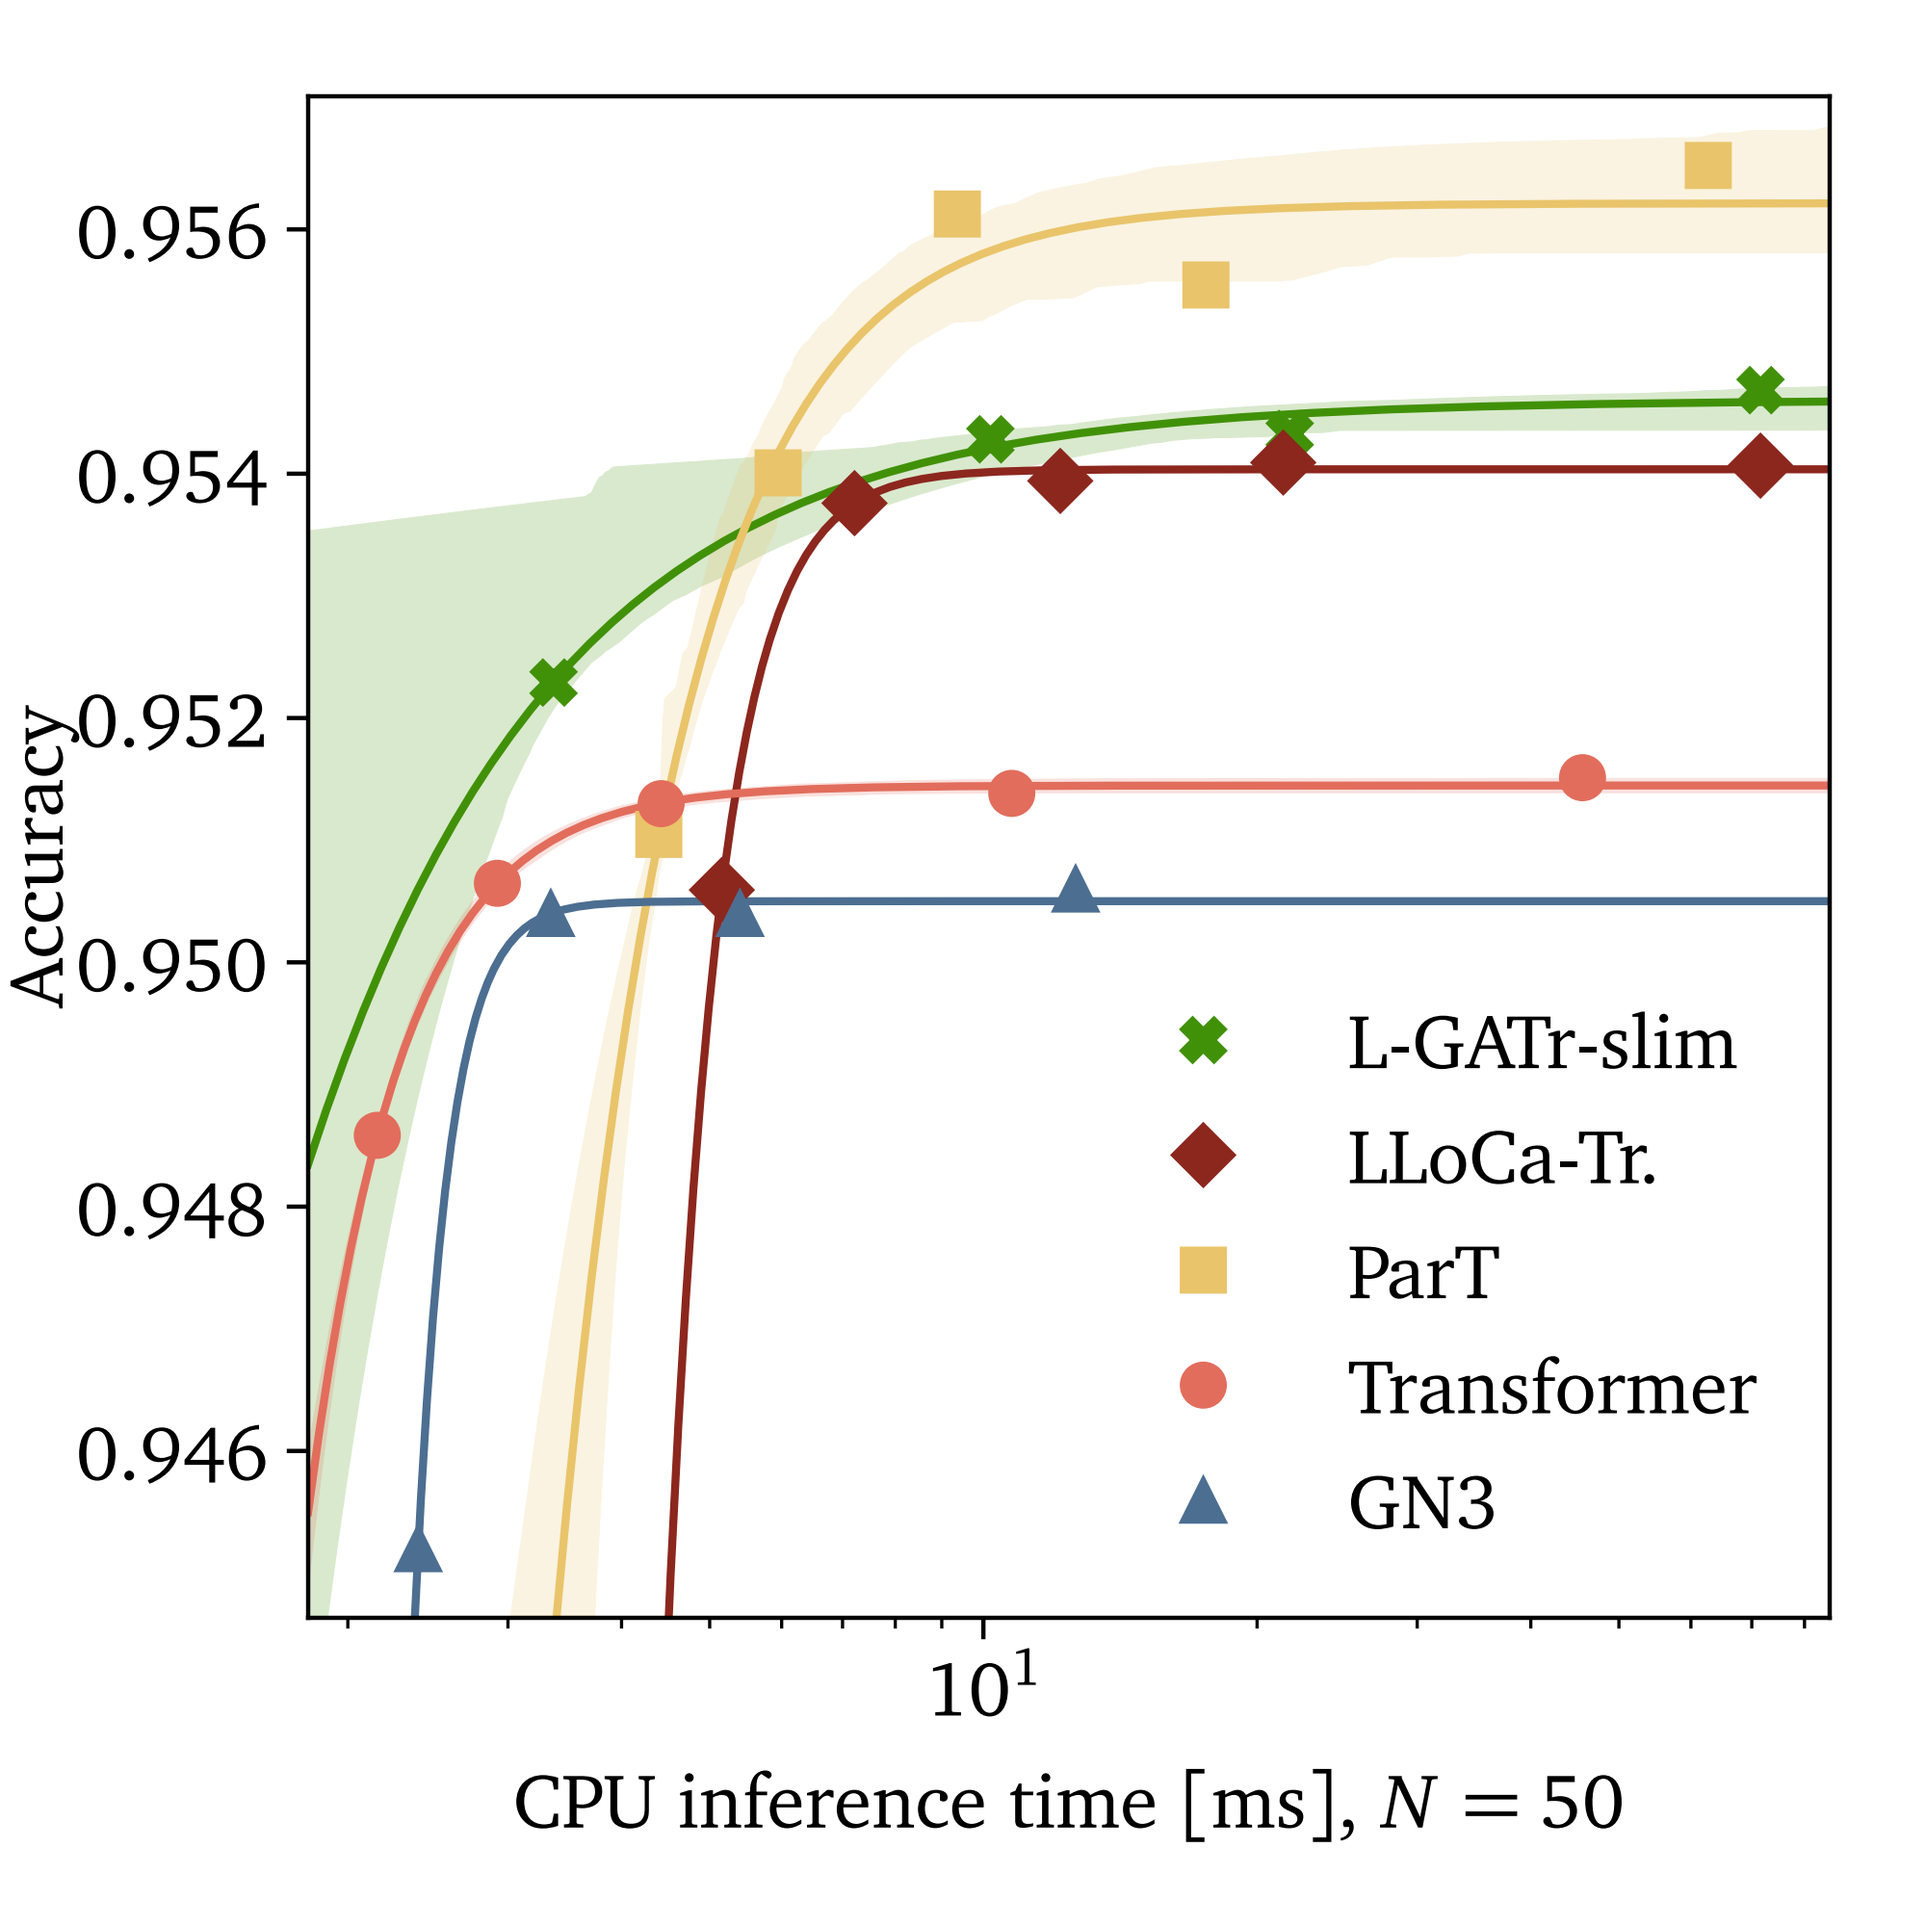
0.956
0.954
0.952
0.950
0.948
0.946
Accuracy
$10^1$
CPU inference time [ms], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GN3

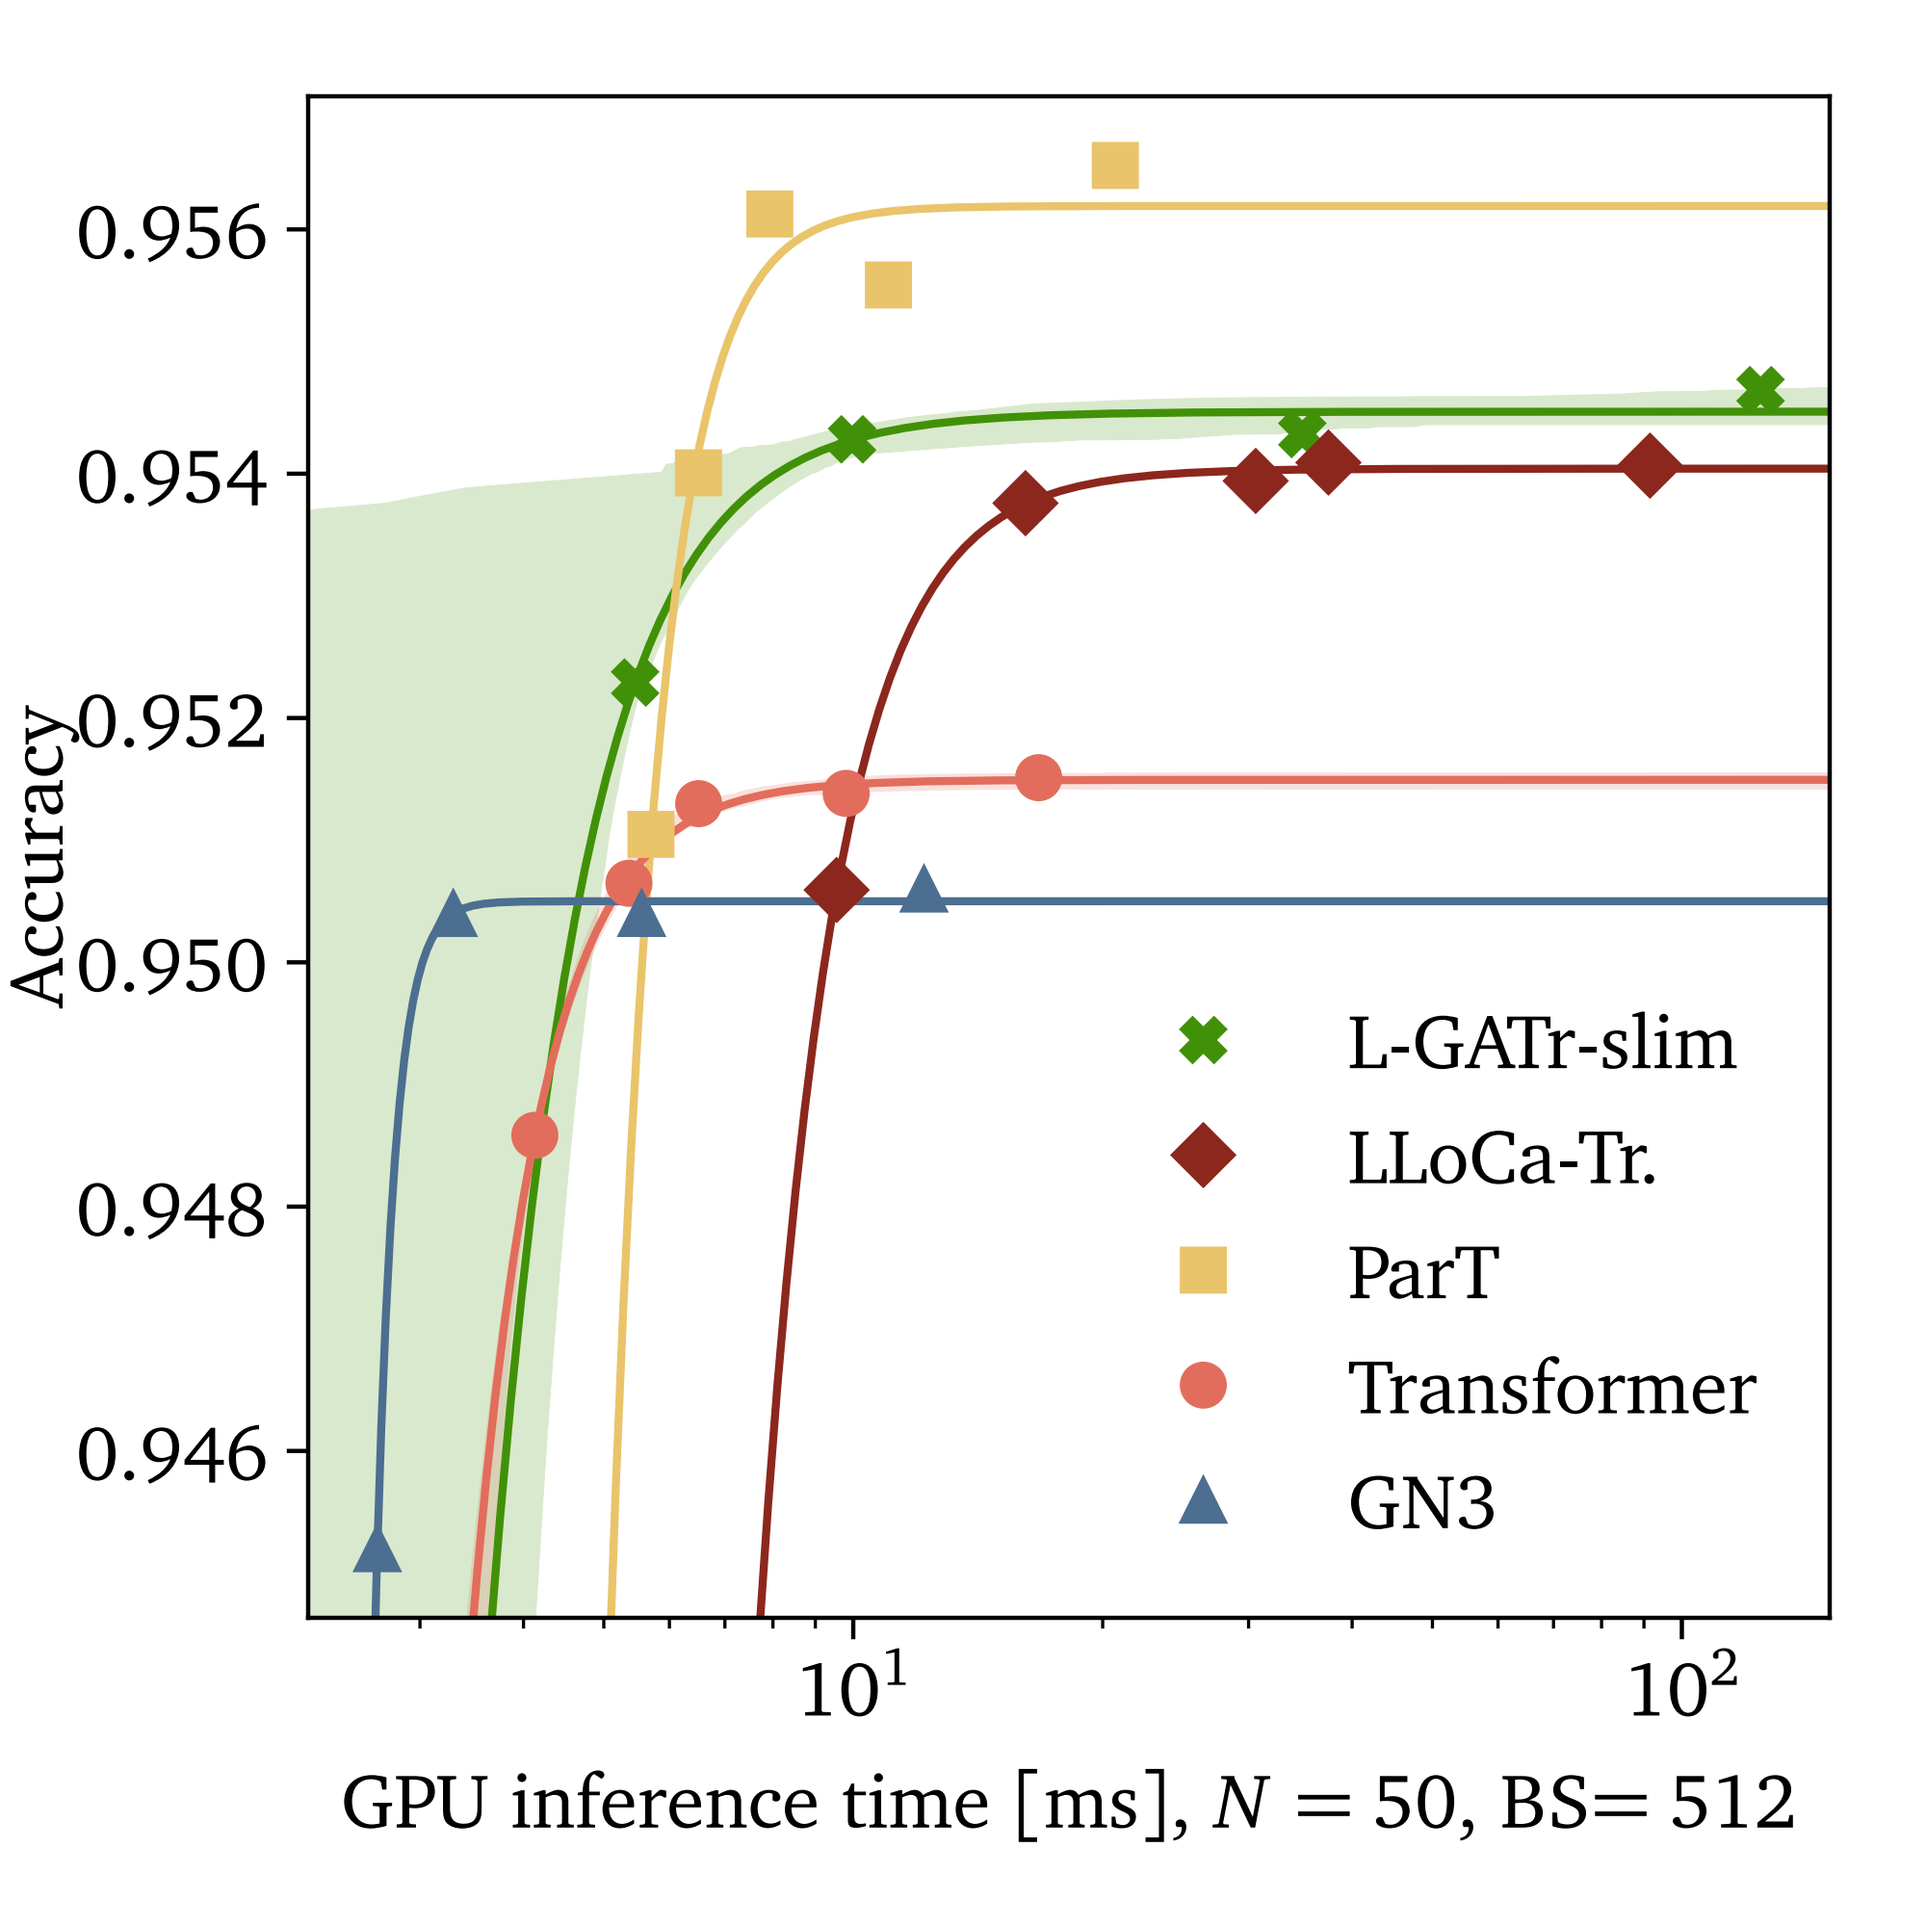
0.956
0.954
0.952
0.950
0.948
0.946
Accuracy
$10^1$
$10^2$
GPU inference time [ms], $N = 50$, BS= 512
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GN3

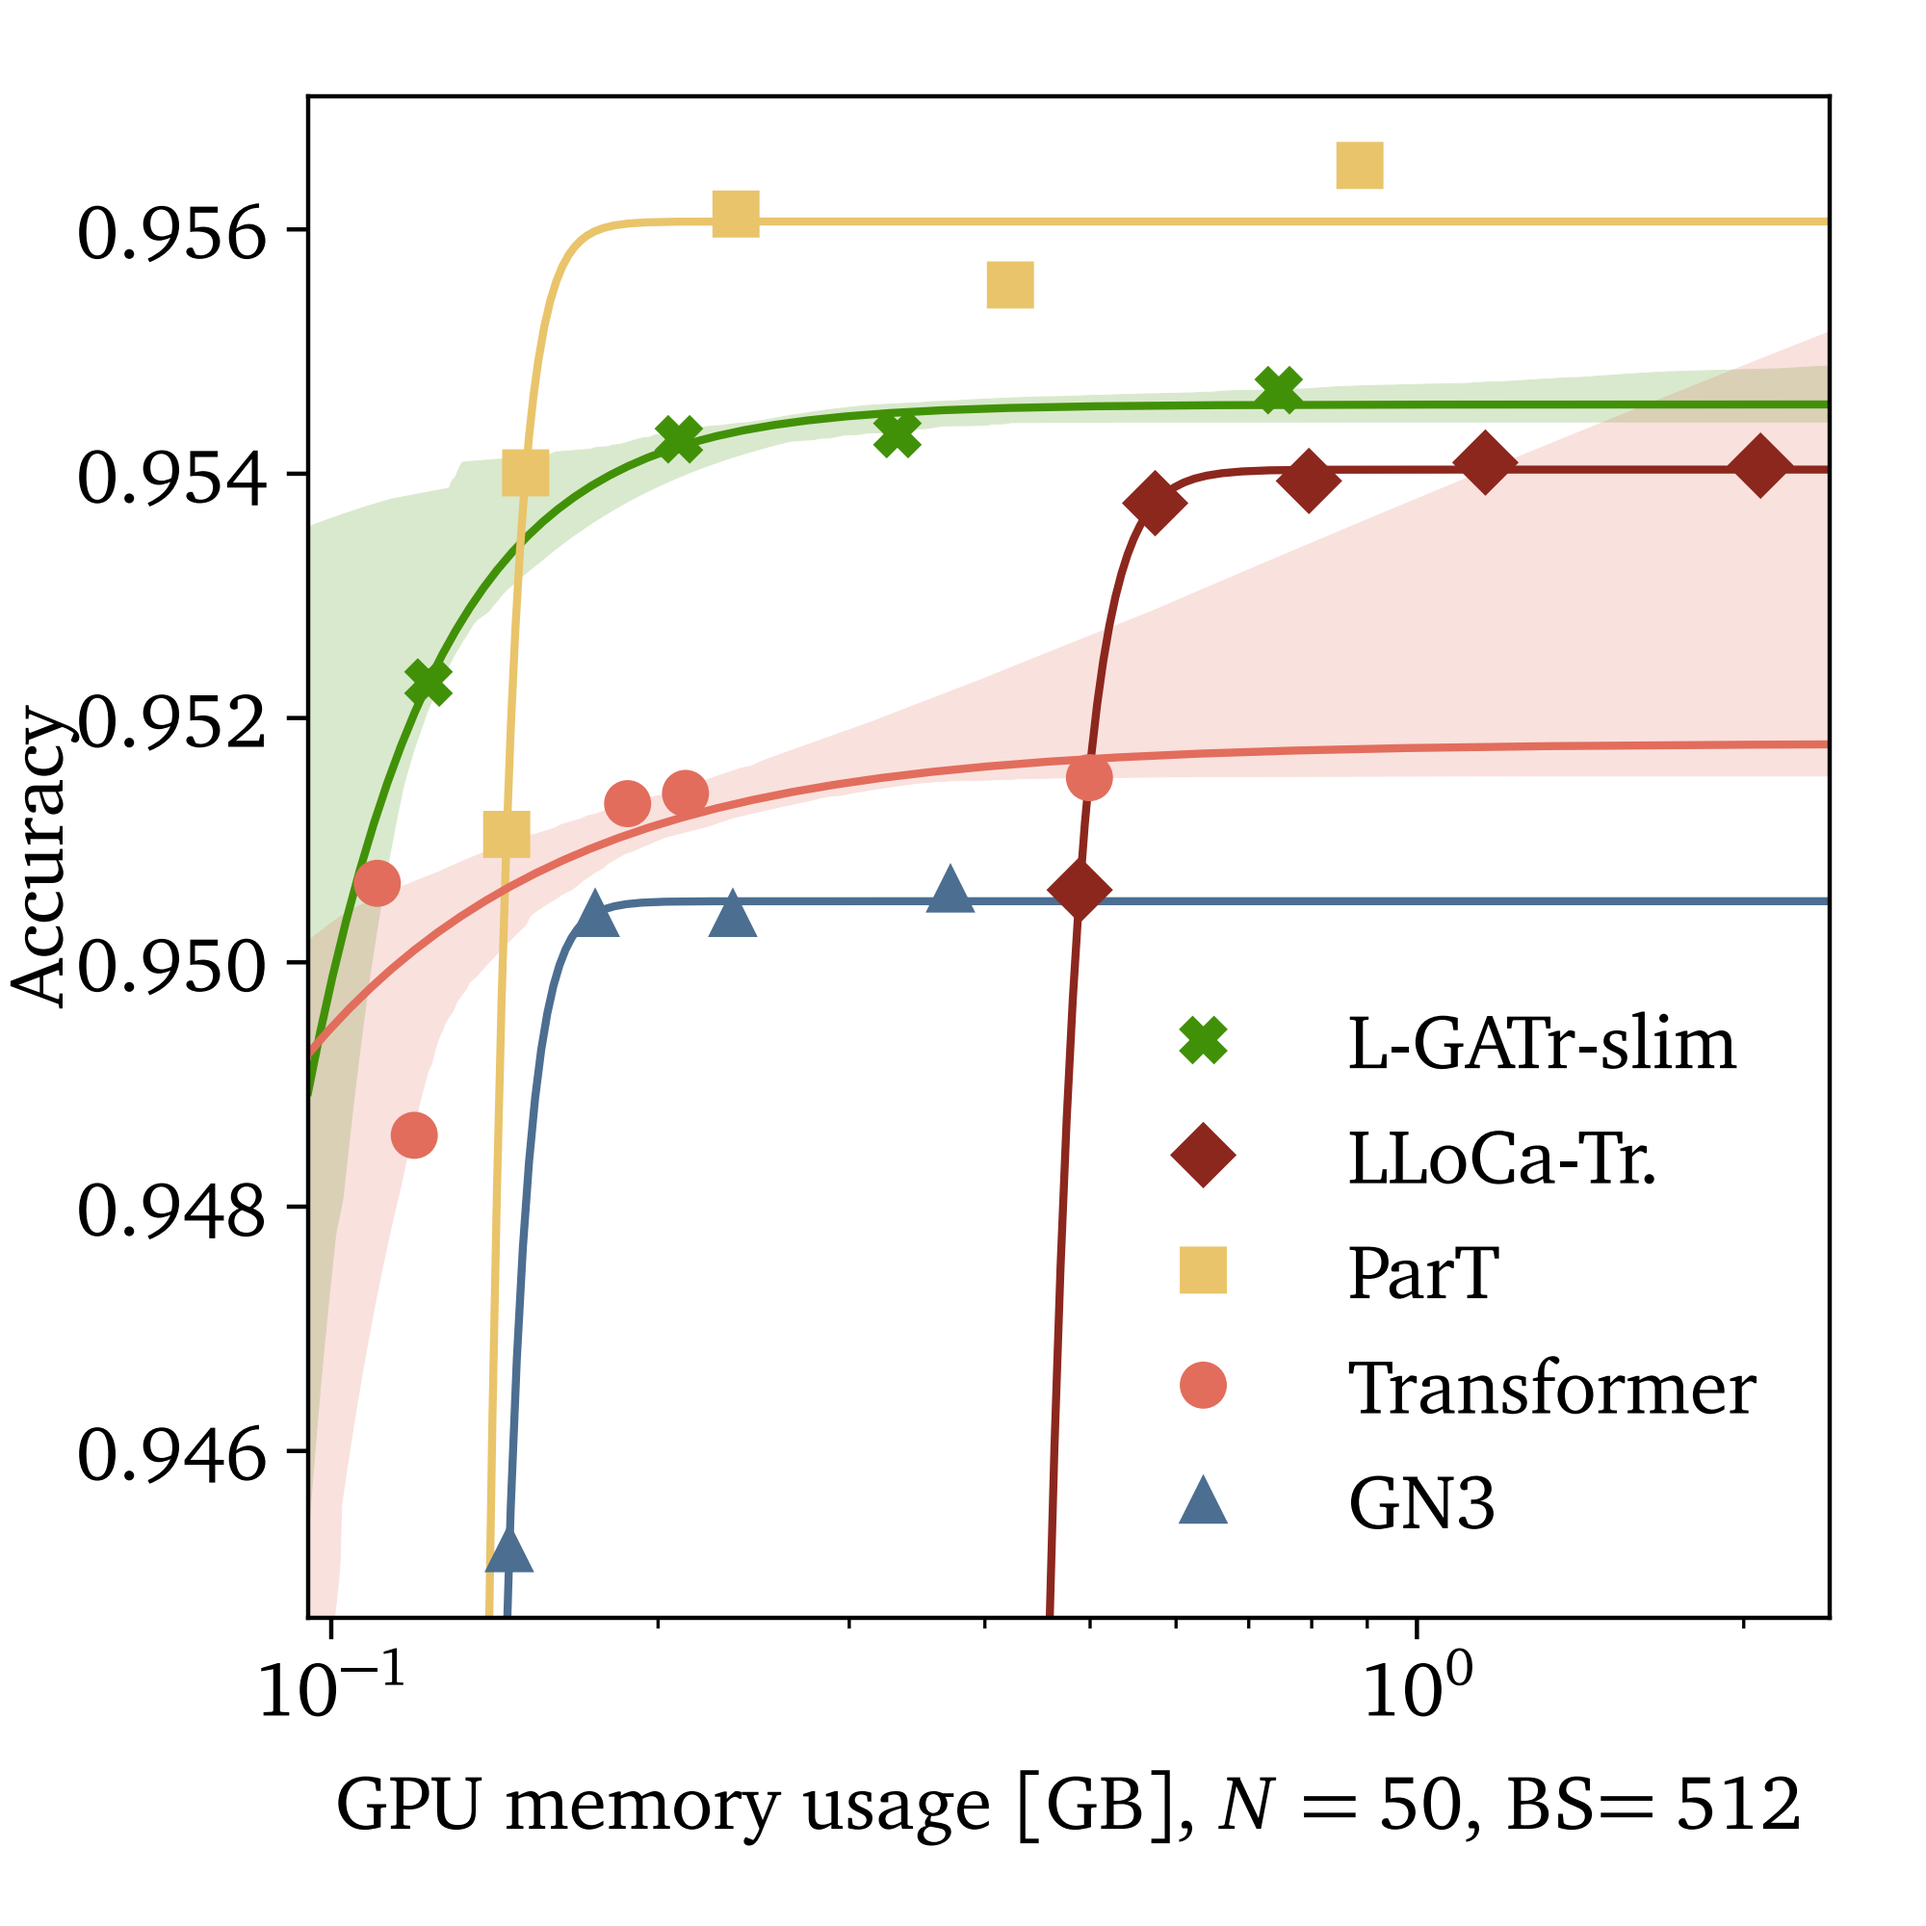
0.956
0.954
0.952
0.950
0.948
0.946
Accuracy
$10^{-1}$
$10^{0}$
GPU memory usage [GB], $N = 50$, BS$= 512$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GN3

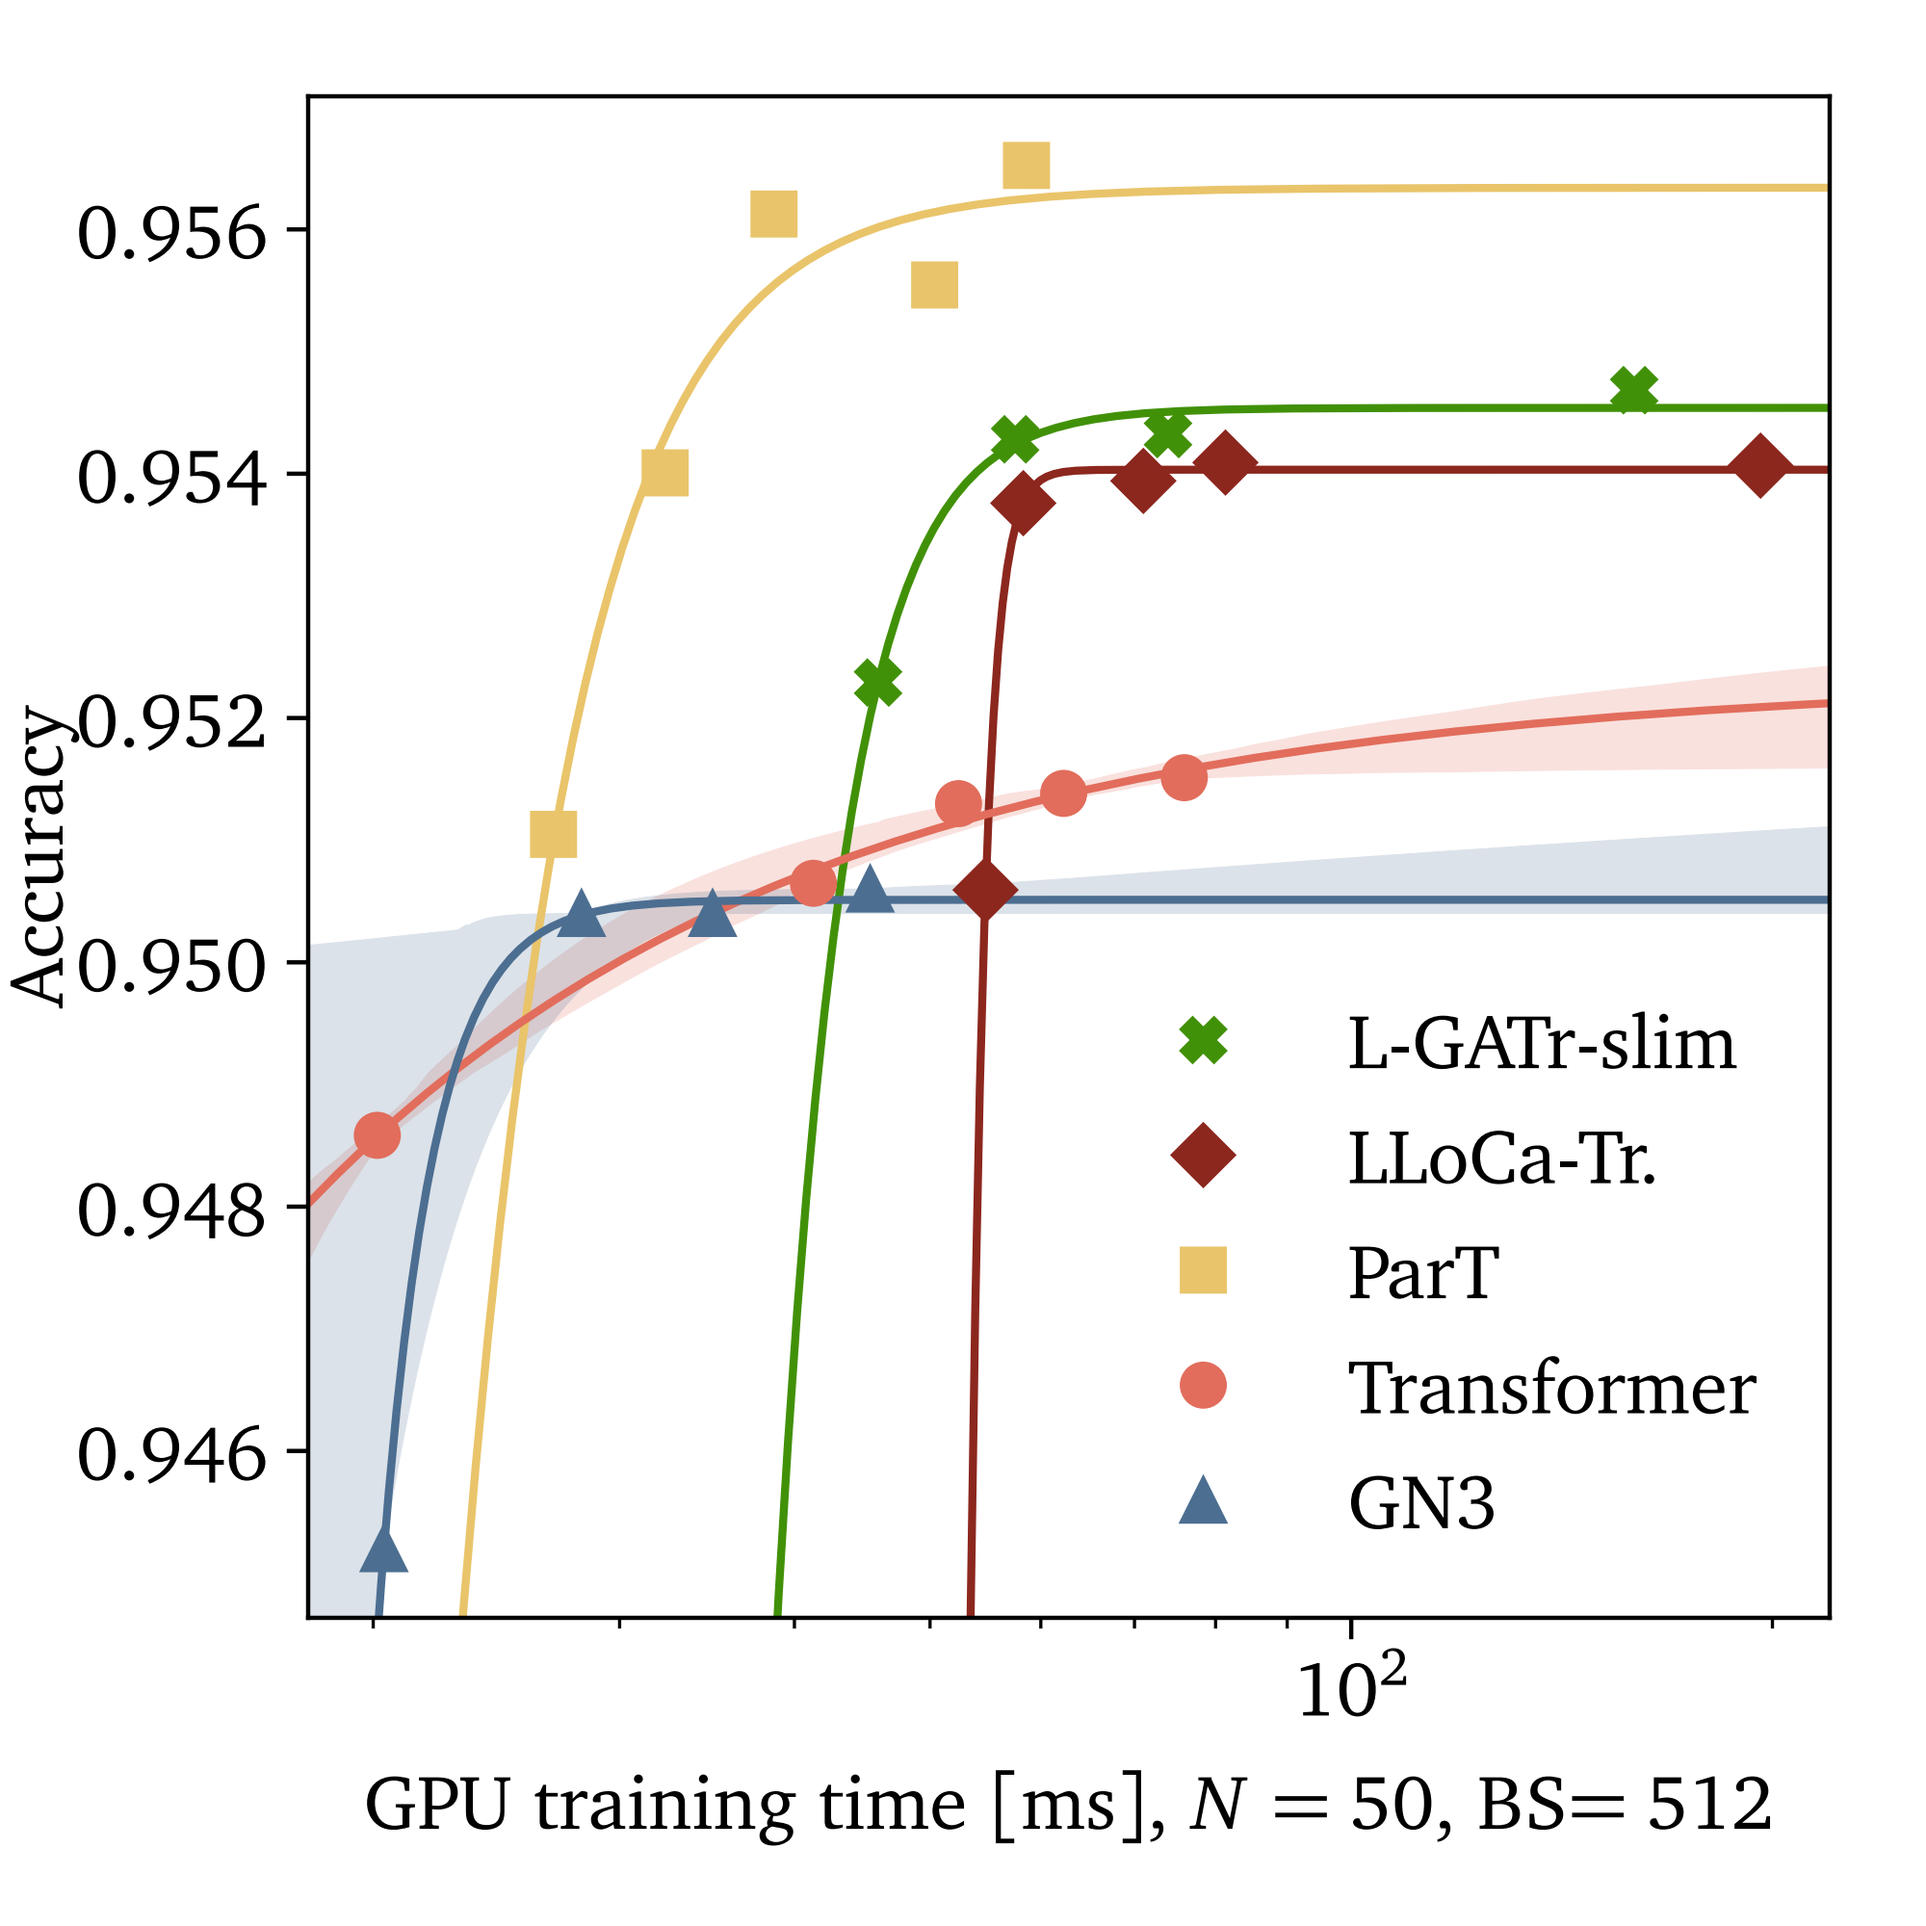
0.956
0.954
0.952
0.950
0.948
0.946
Accuracy
$10^2$
GPU training time [ms], $N = 50$, BS$= 512$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GN3