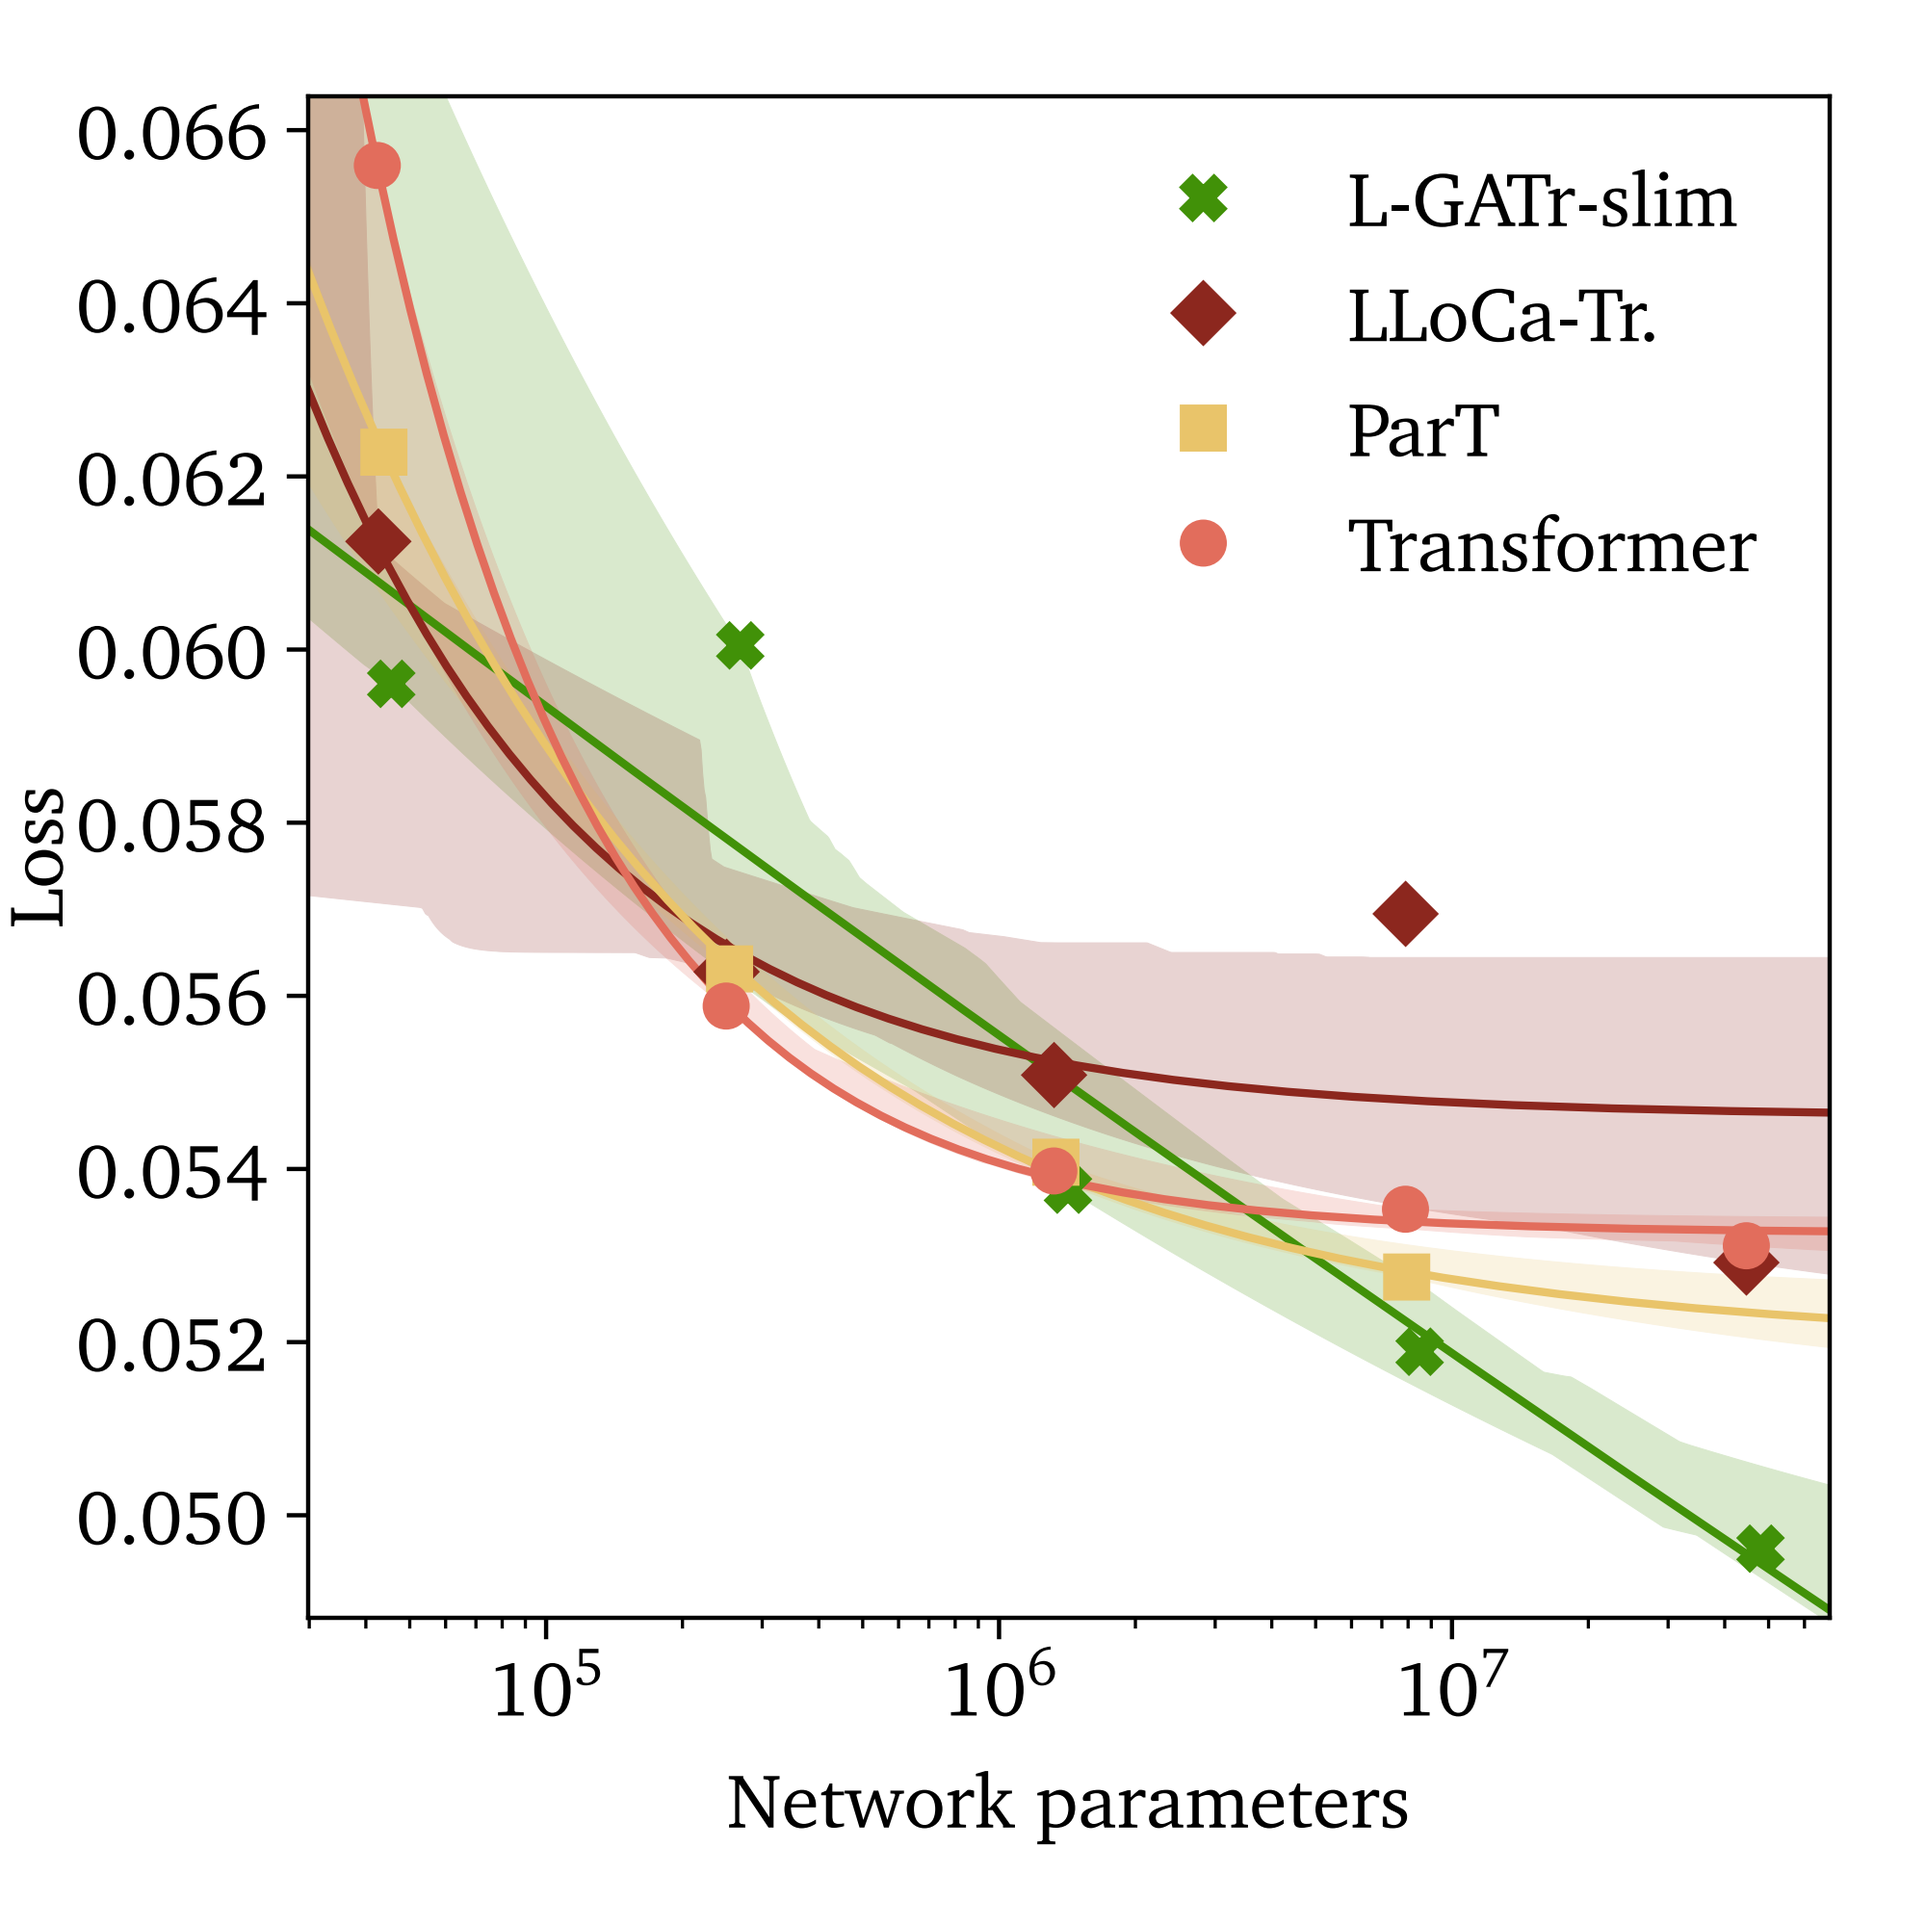
0.066
0.064
0.062
0.060
0.058
0.056
0.054
0.052
0.050
Loss
10^5
10^6
10^7
Network parameters
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

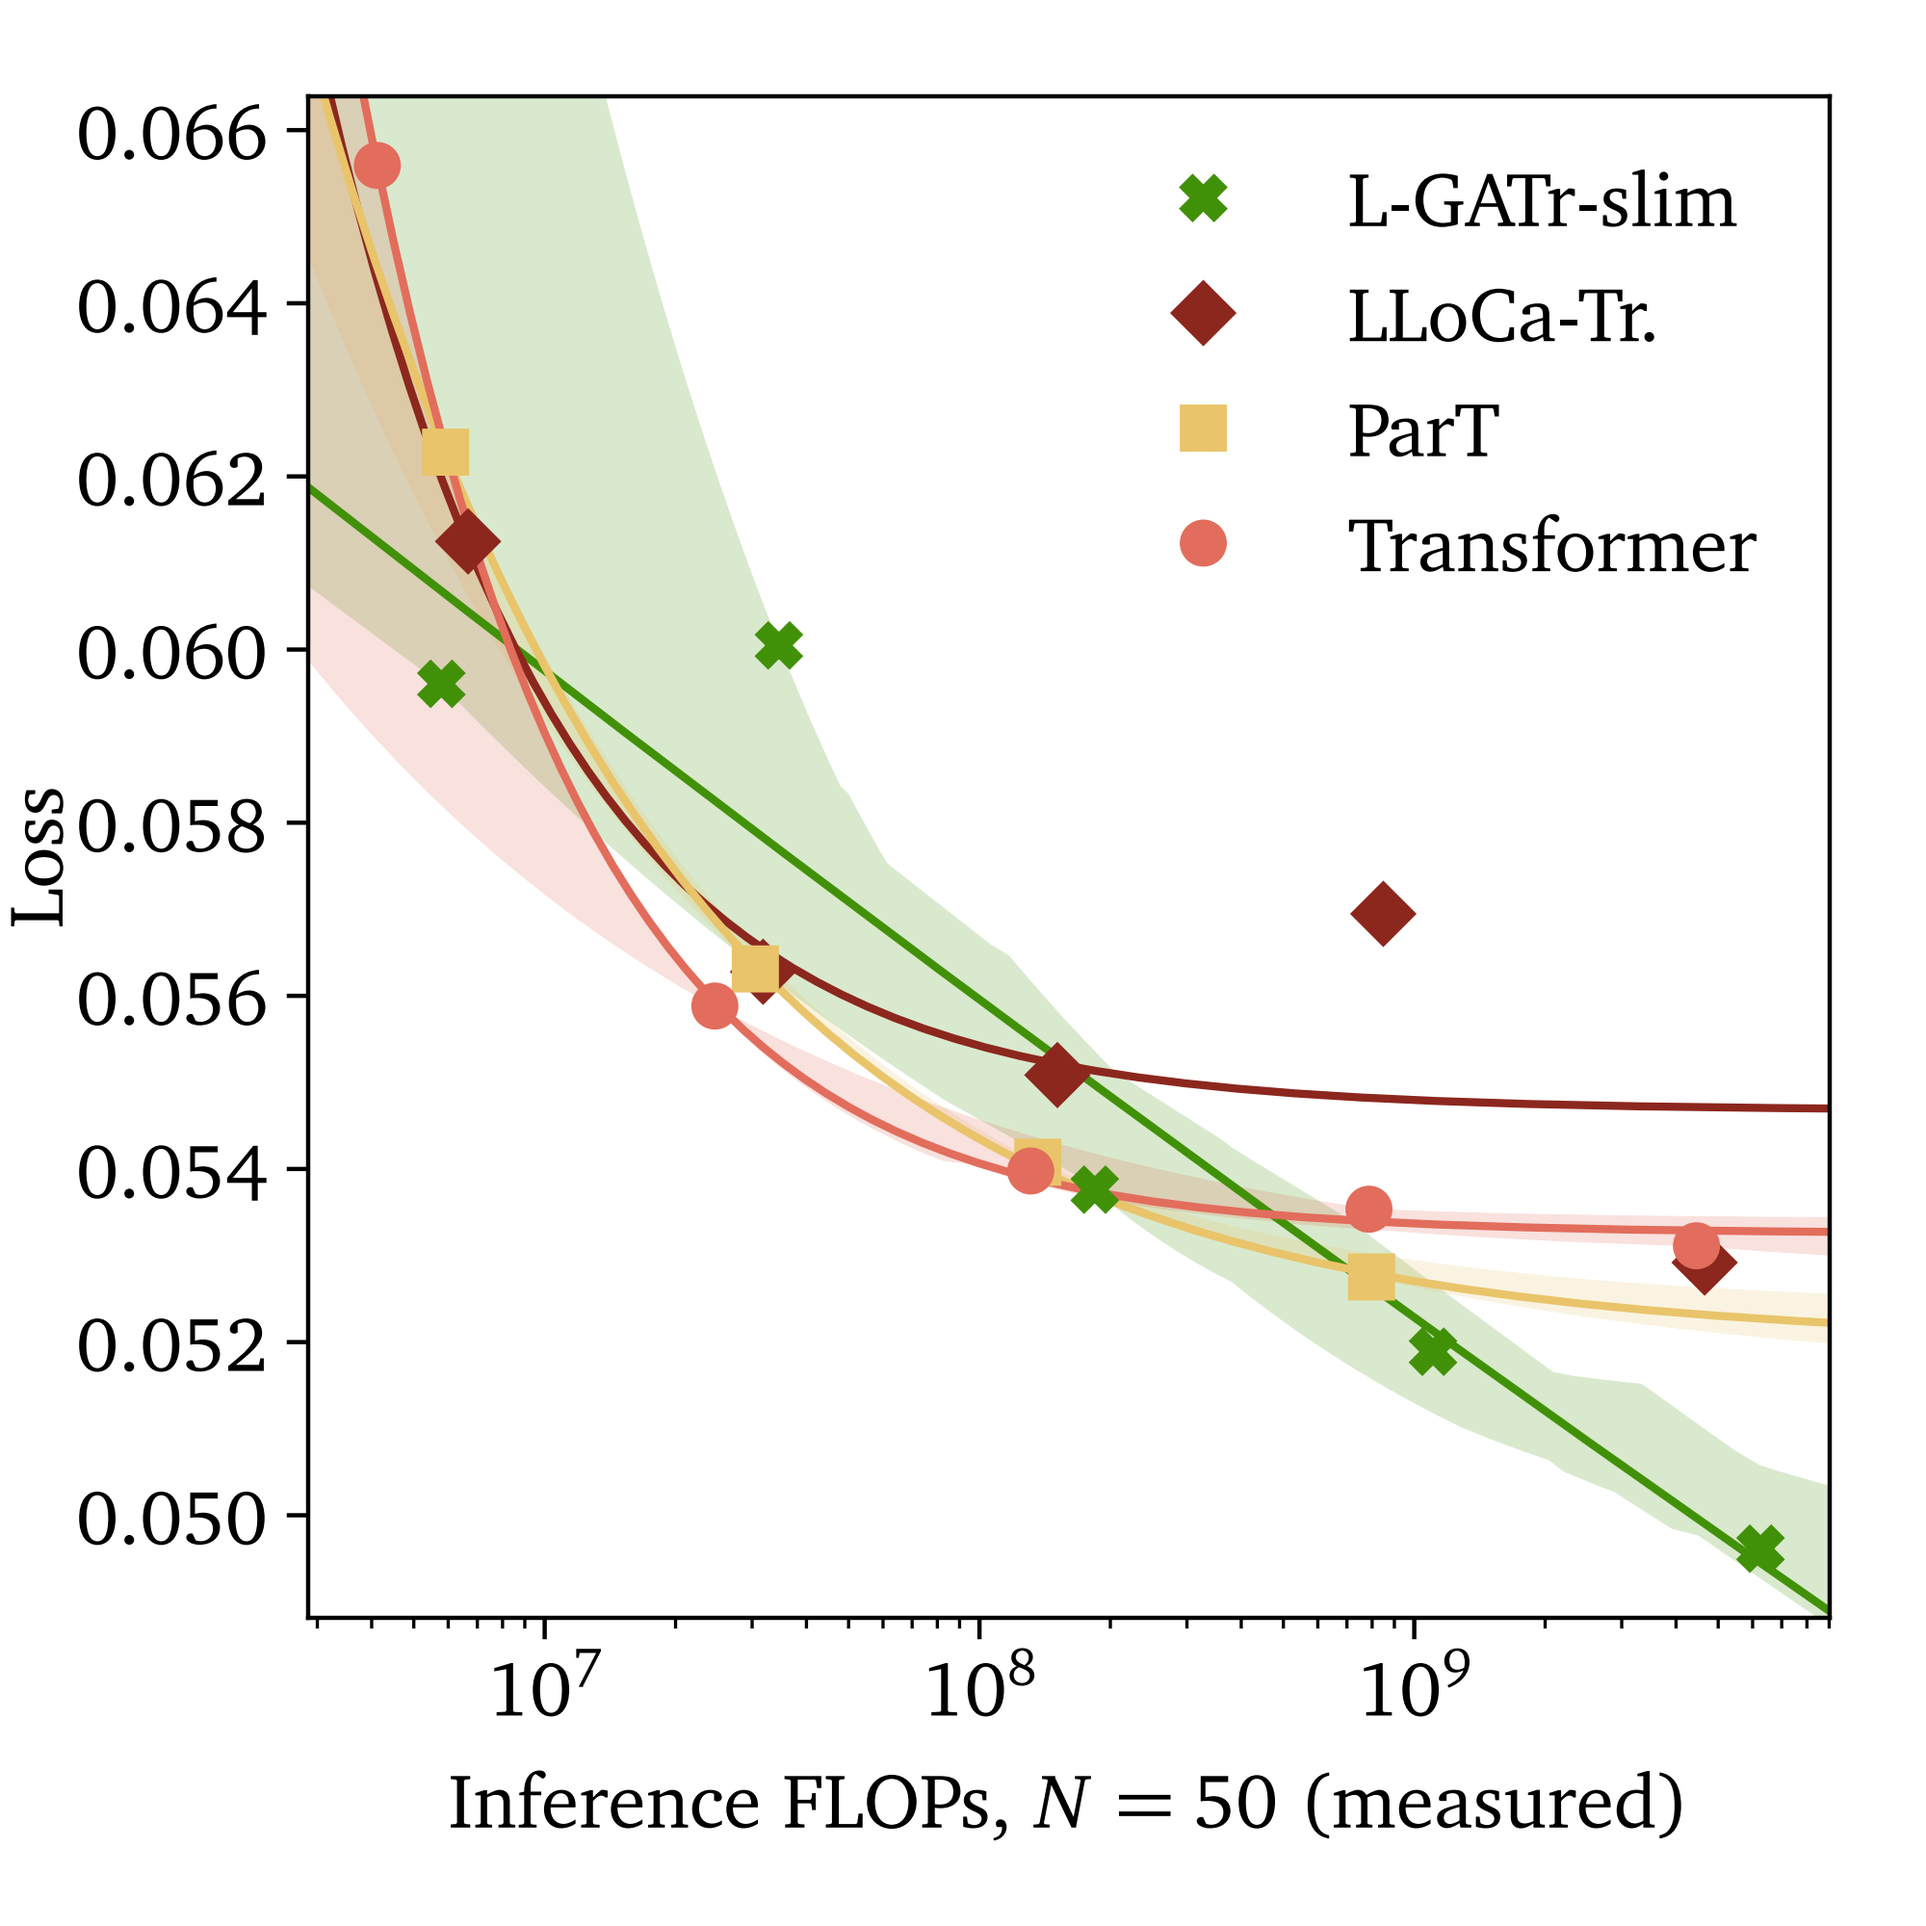
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.066
0.064
0.062
0.060
0.058
0.056
0.054
0.052
0.050
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (measured)

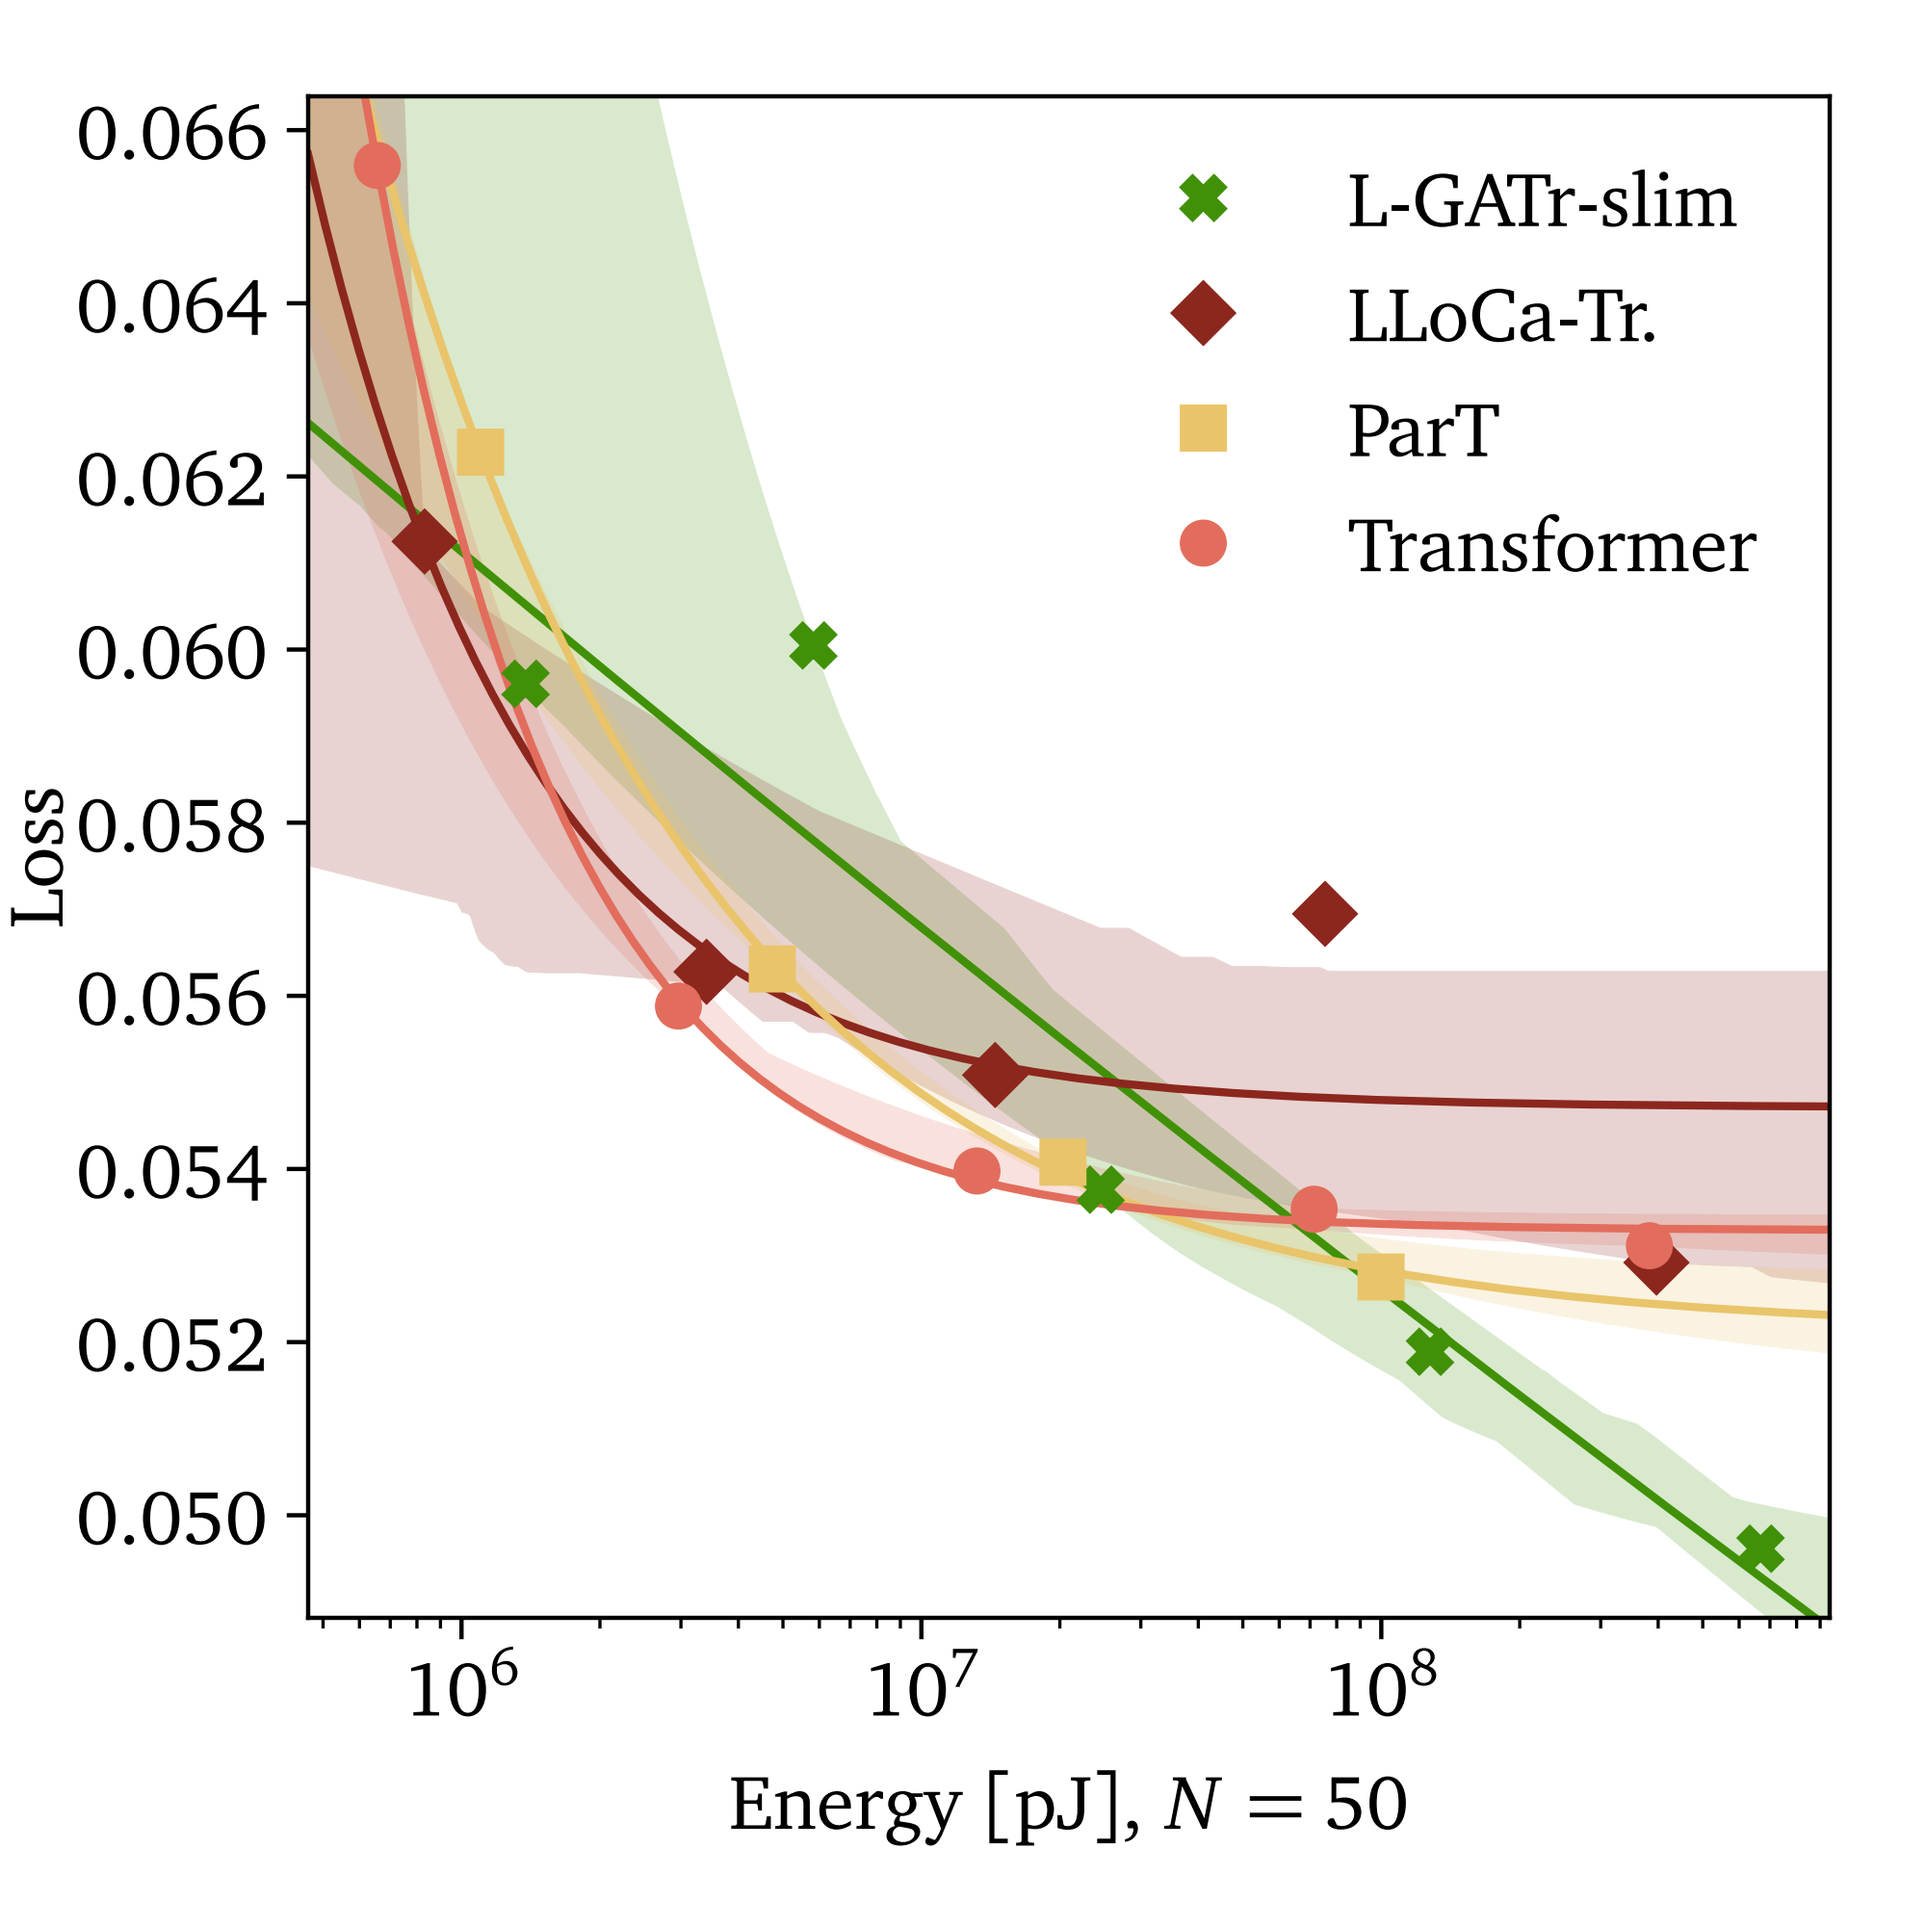
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.066
0.064
0.062
0.060
0.058
0.056
0.054
0.052
0.050
10^6
10^7
10^8
Energy [pJ], $N = 50$

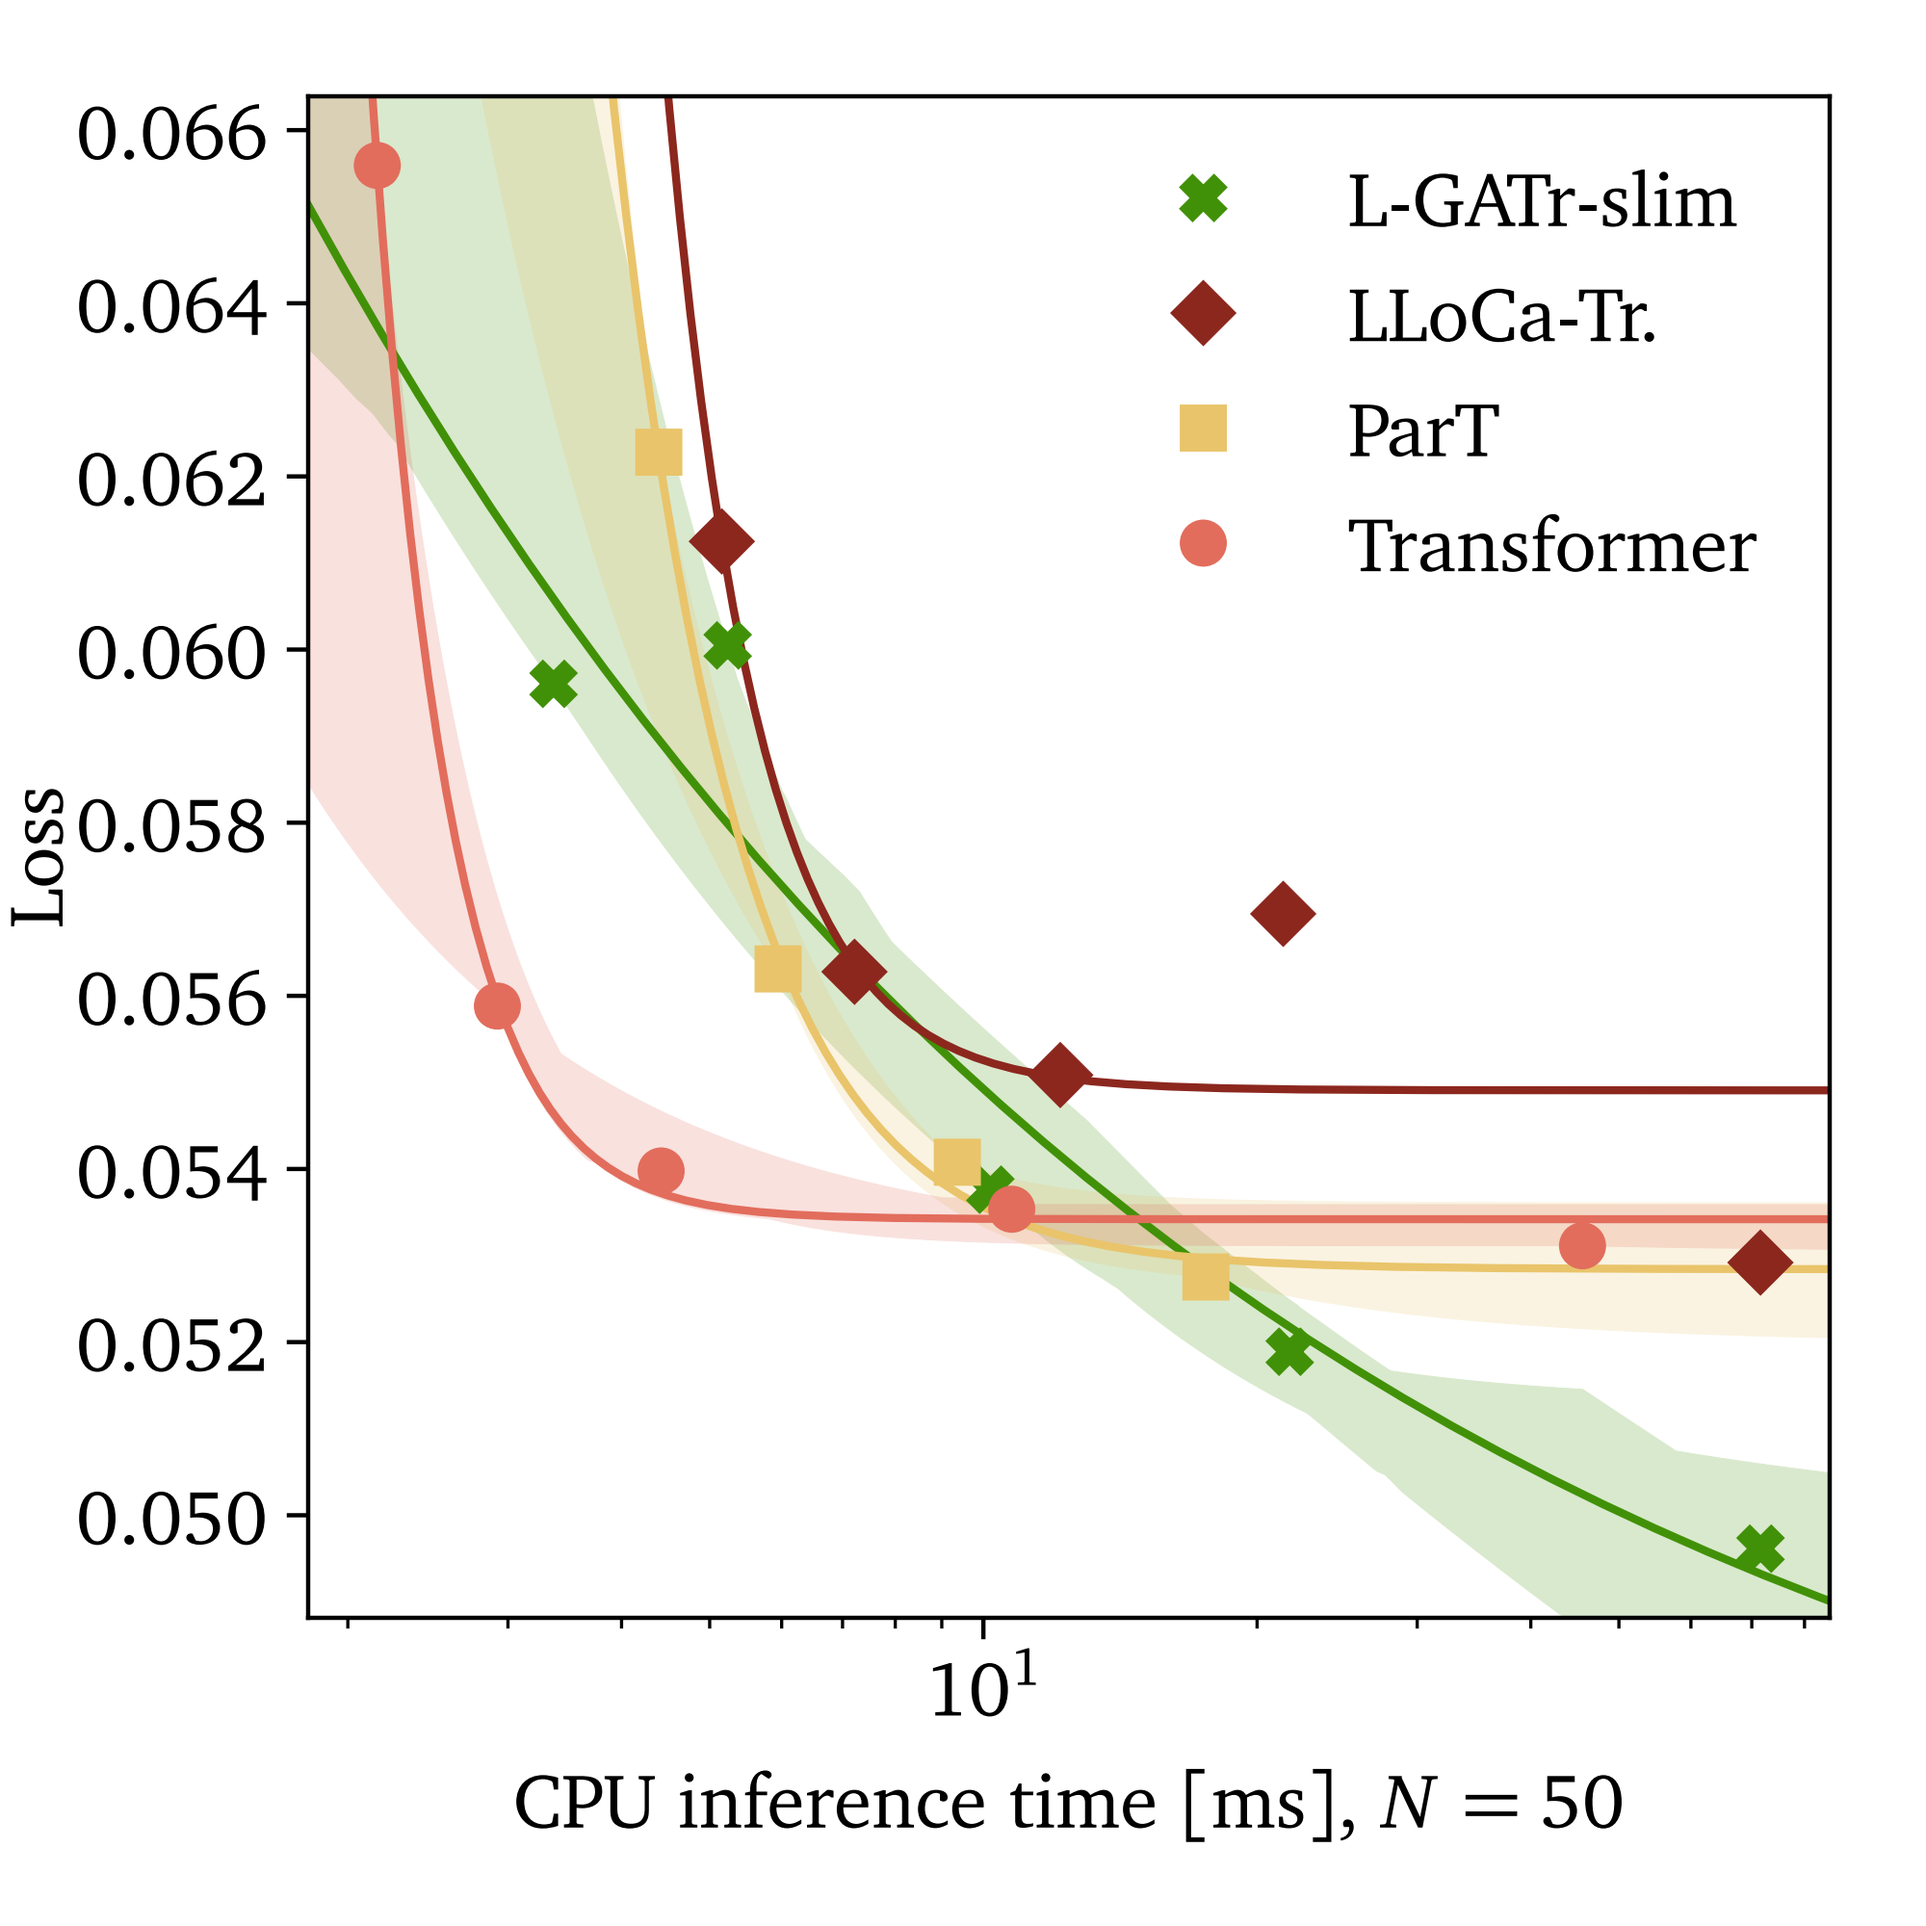
0.066
0.064
0.062
0.060
0.058
0.056
0.054
0.052
0.050
Loss
10^1
CPU inference time [ms], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

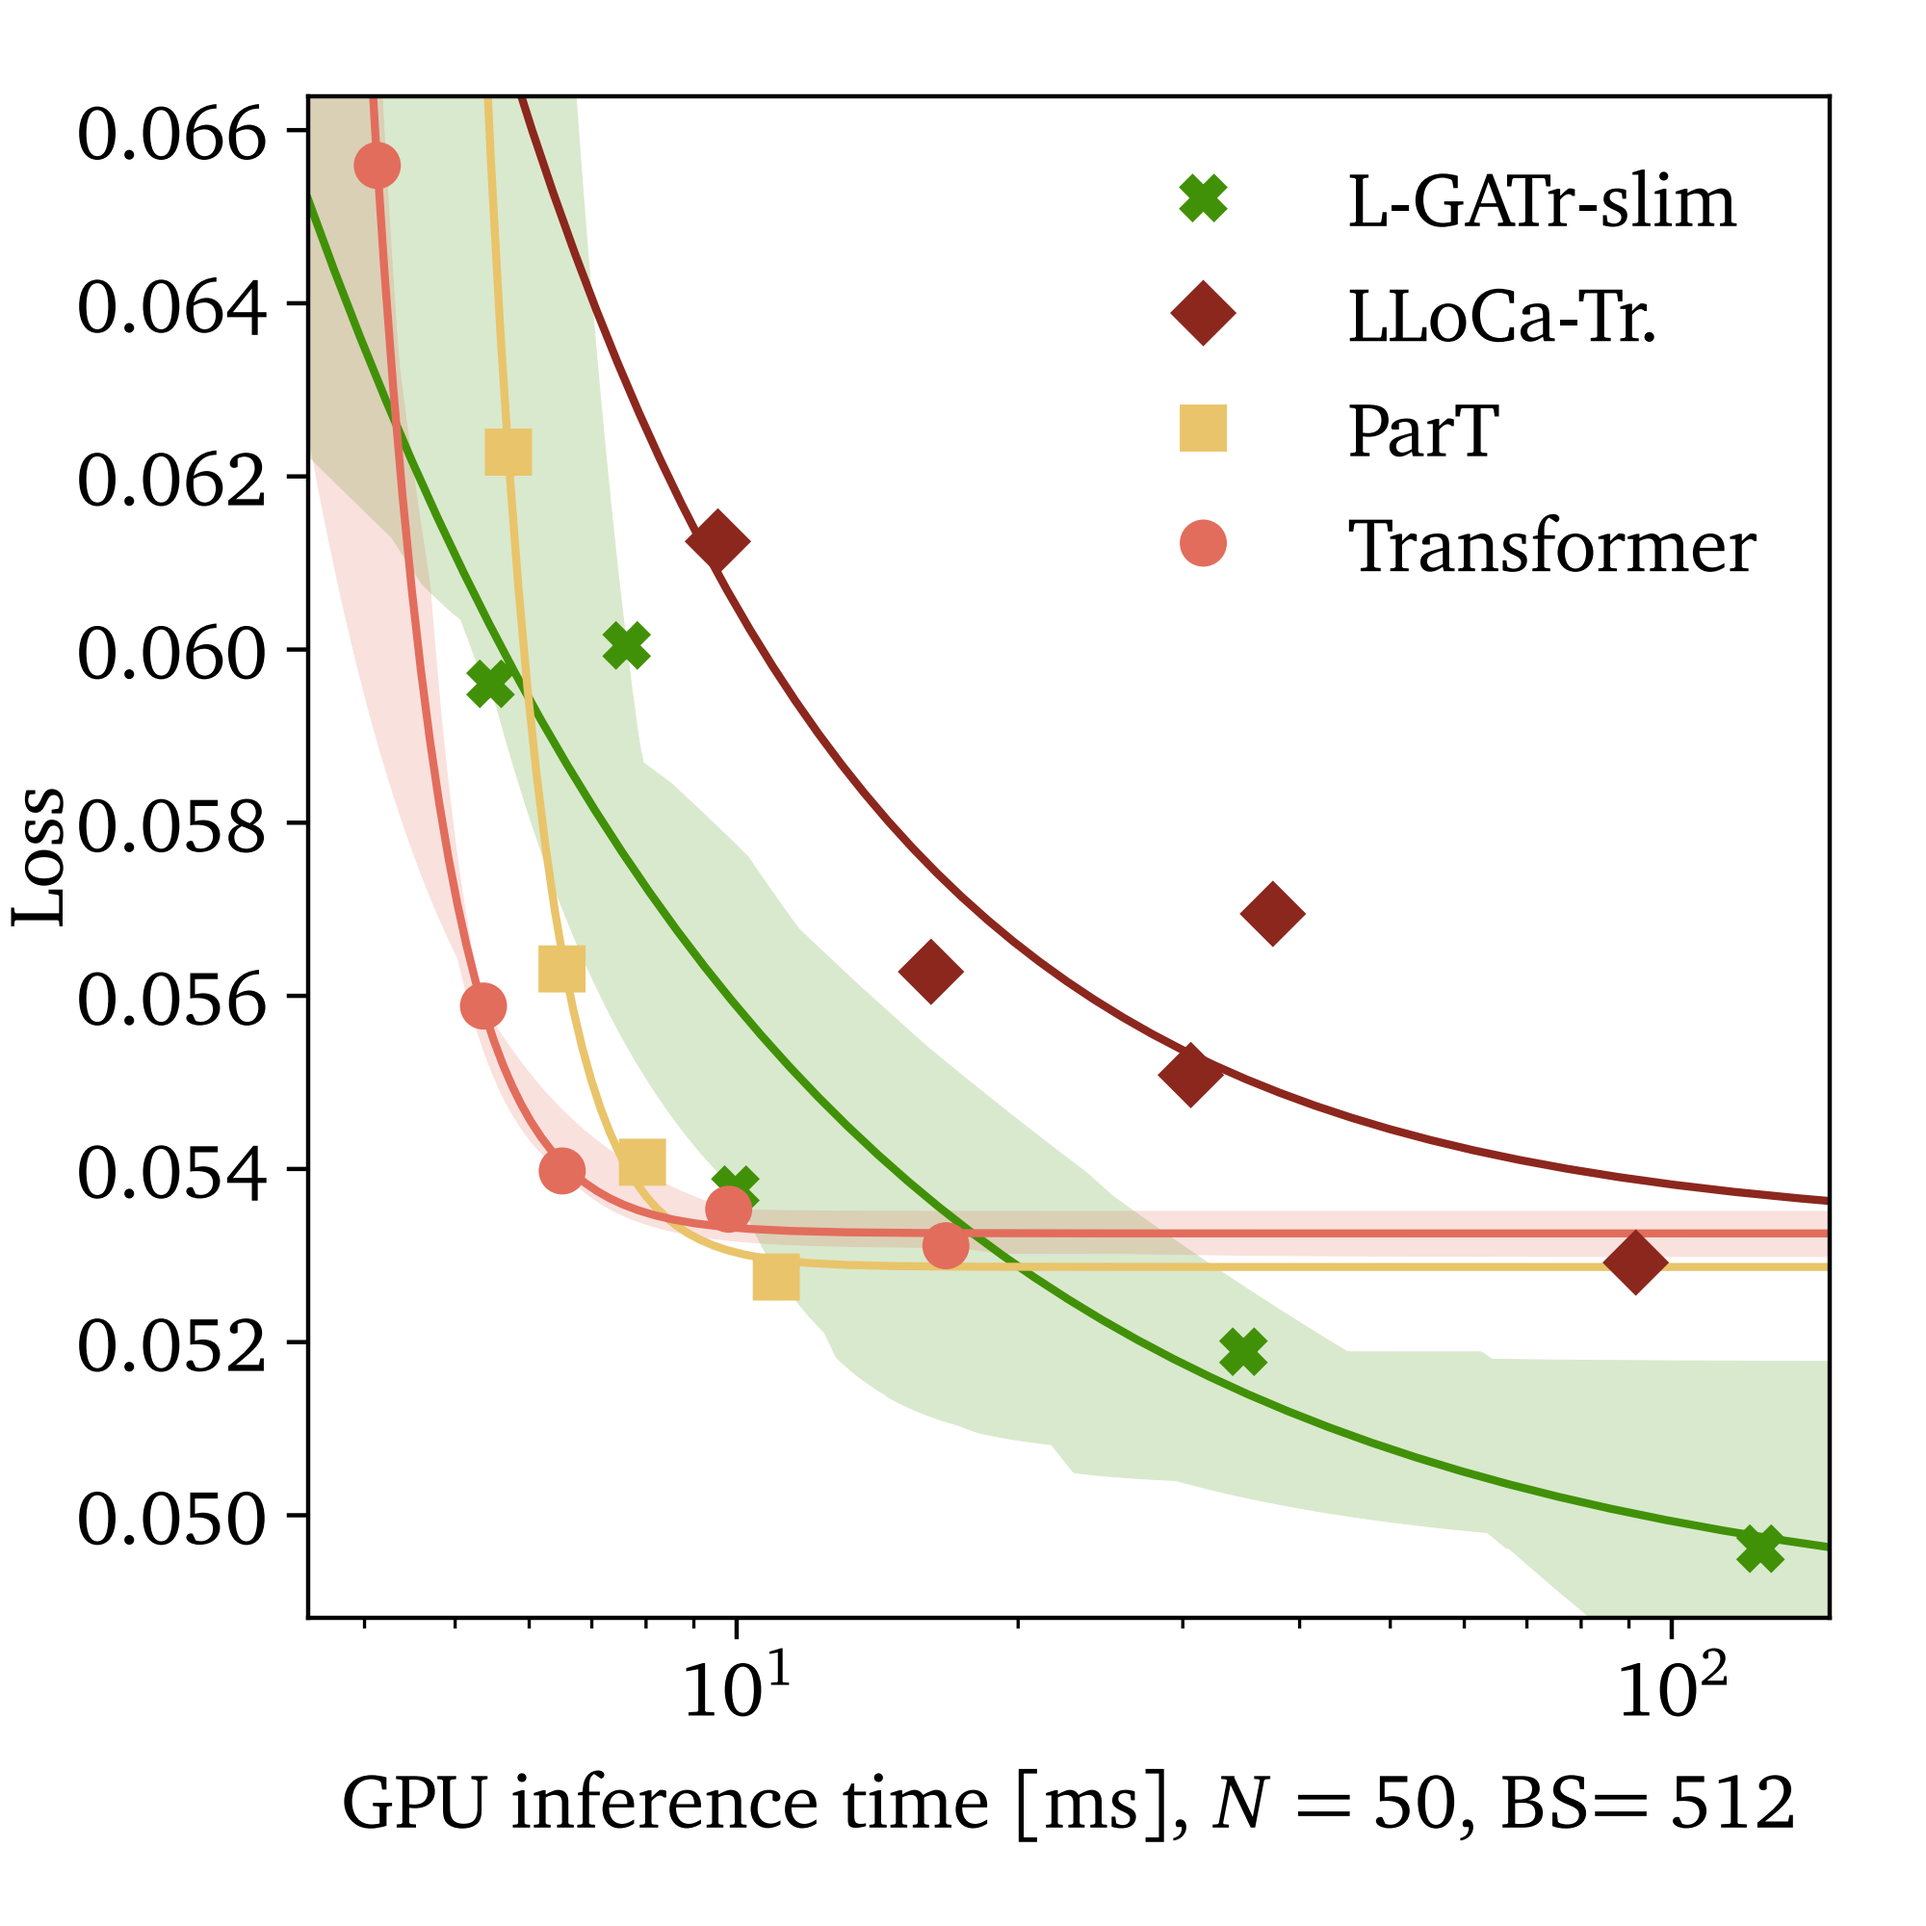
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.066
0.064
0.062
0.060
0.058
0.056
0.054
0.052
0.050
10^1
10^2
GPU inference time [ms], $N = 50$, BS$= 512$

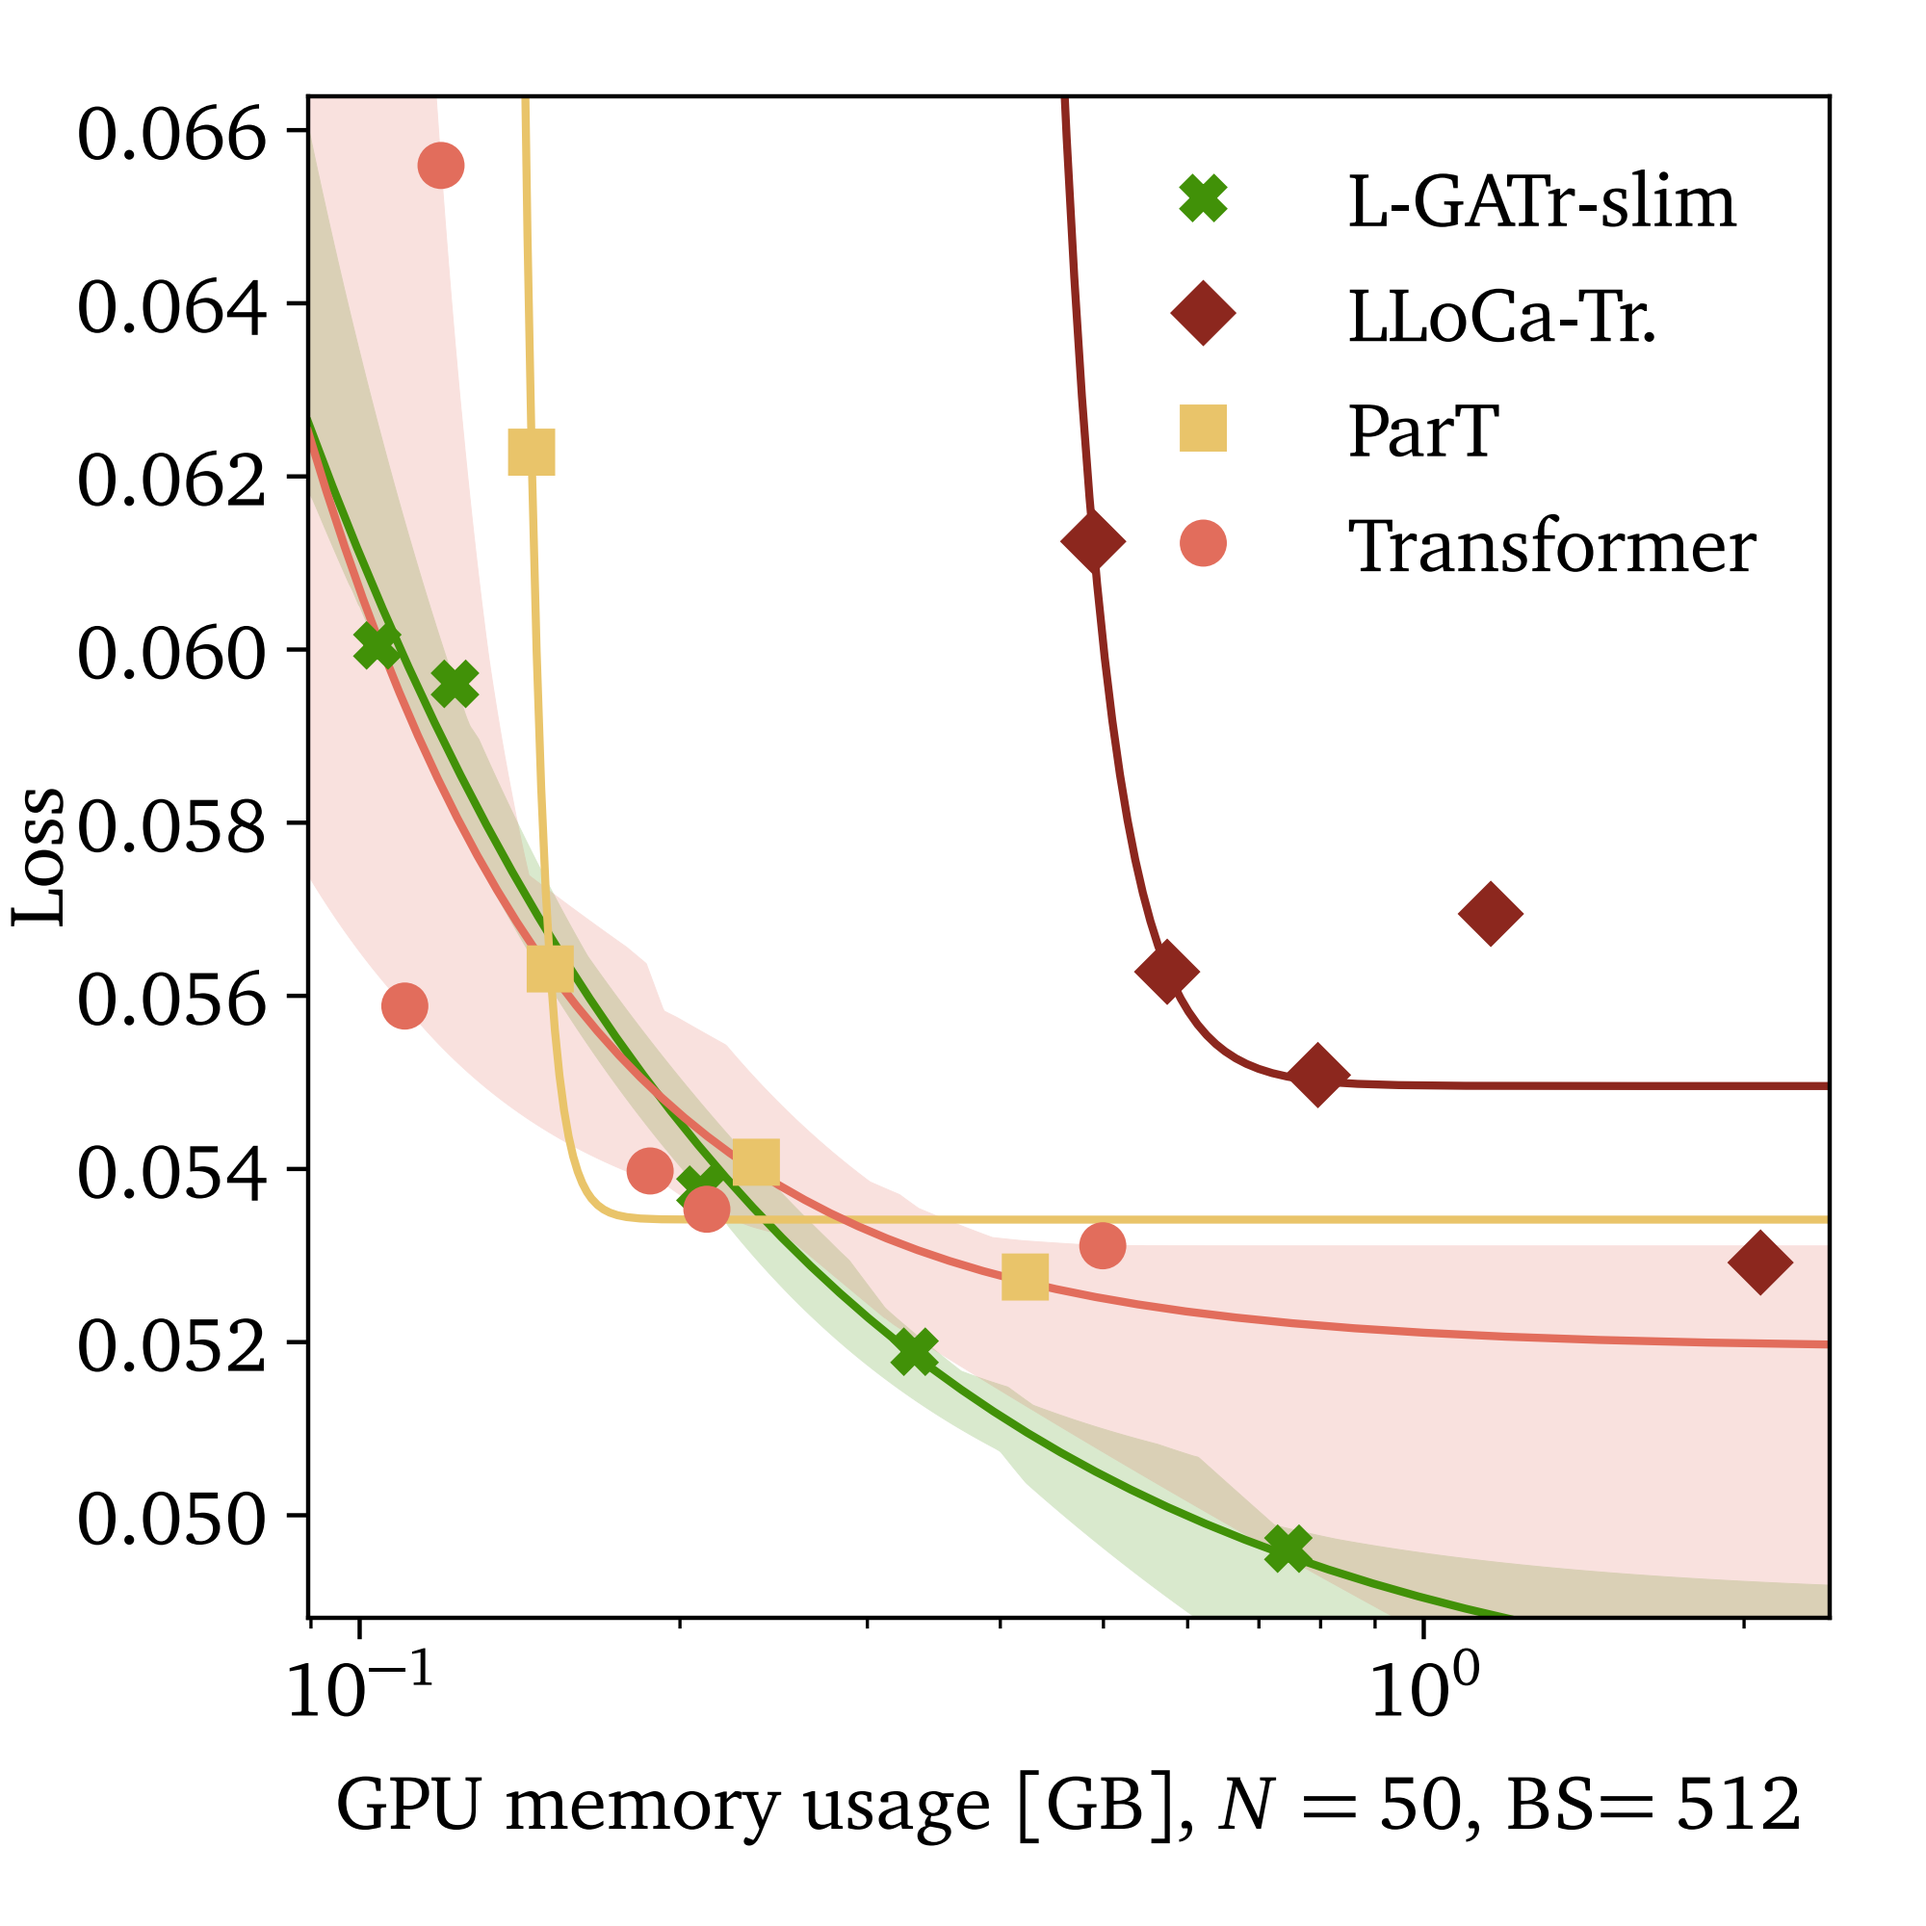
0.066
0.064
0.062
0.060
0.058
0.056
0.054
0.052
0.050
Loss
10^{-1}
10^{0}
GPU memory usage [GB], $N = 50$, BS$= 512$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

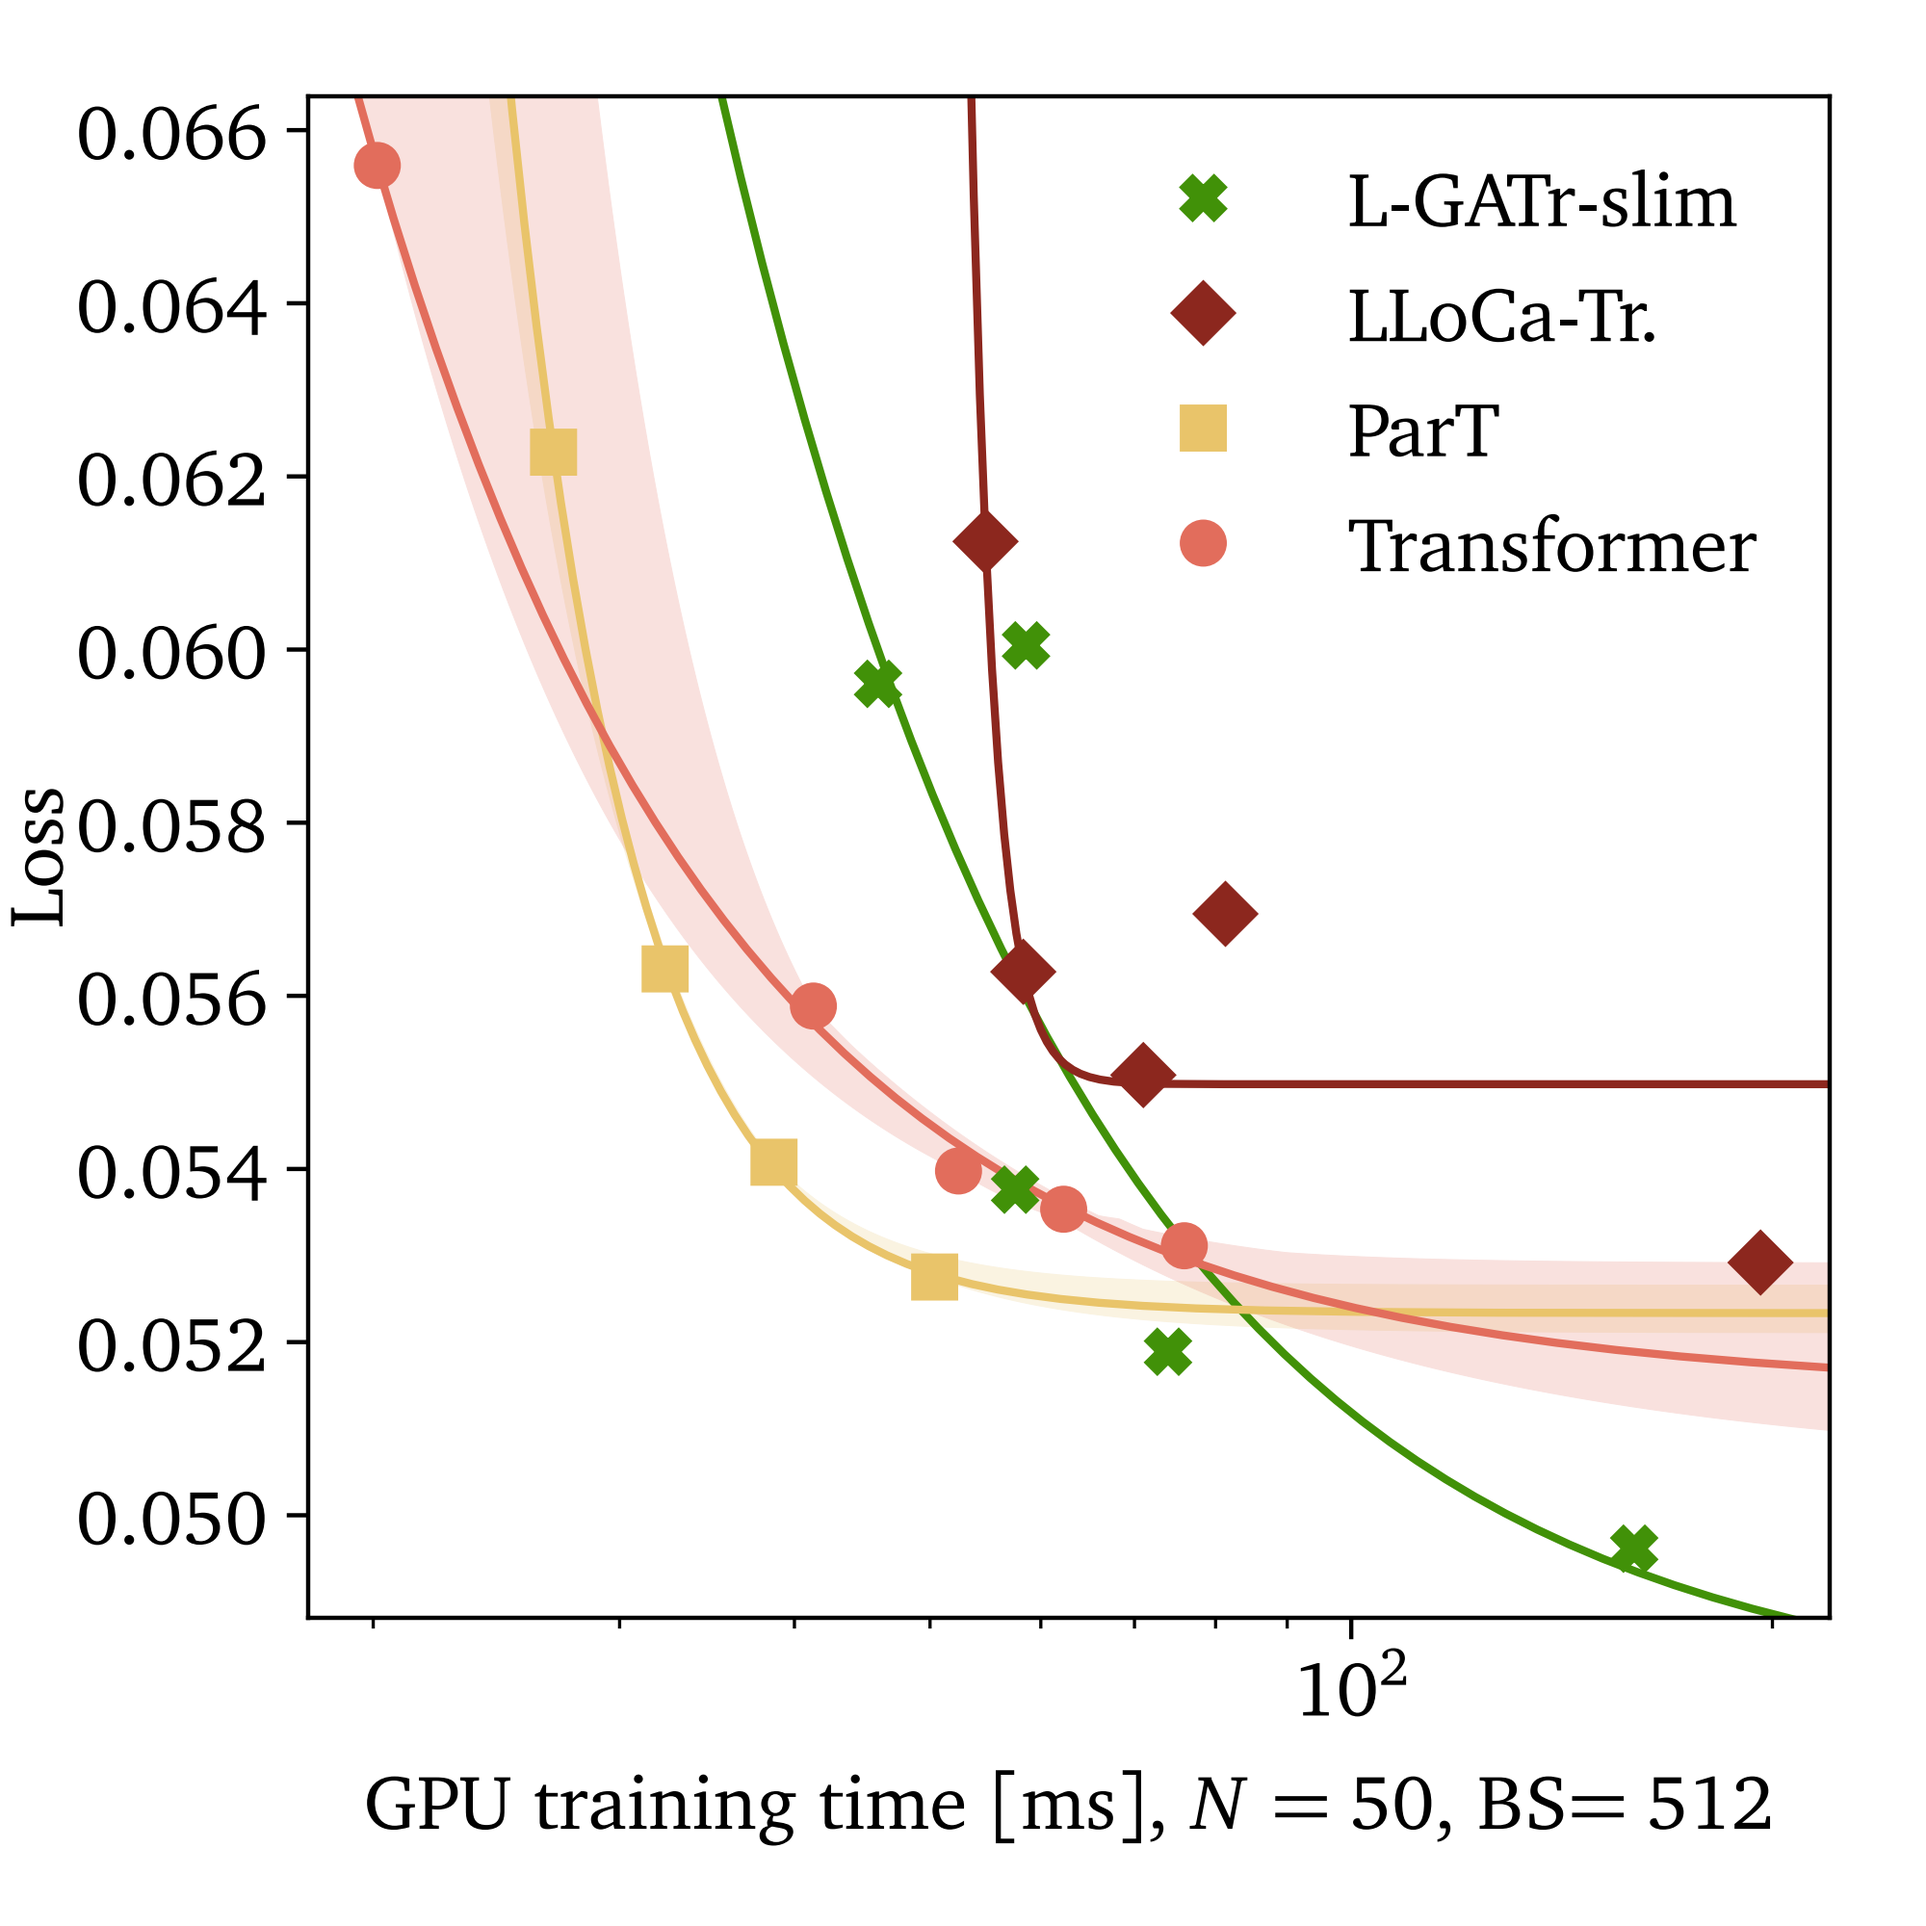
0.066
0.064
0.062
0.060
0.058
0.056
0.054
0.052
0.050
Loss
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
10^2
GPU training time [ms], $N = 50$, BS$= 512$

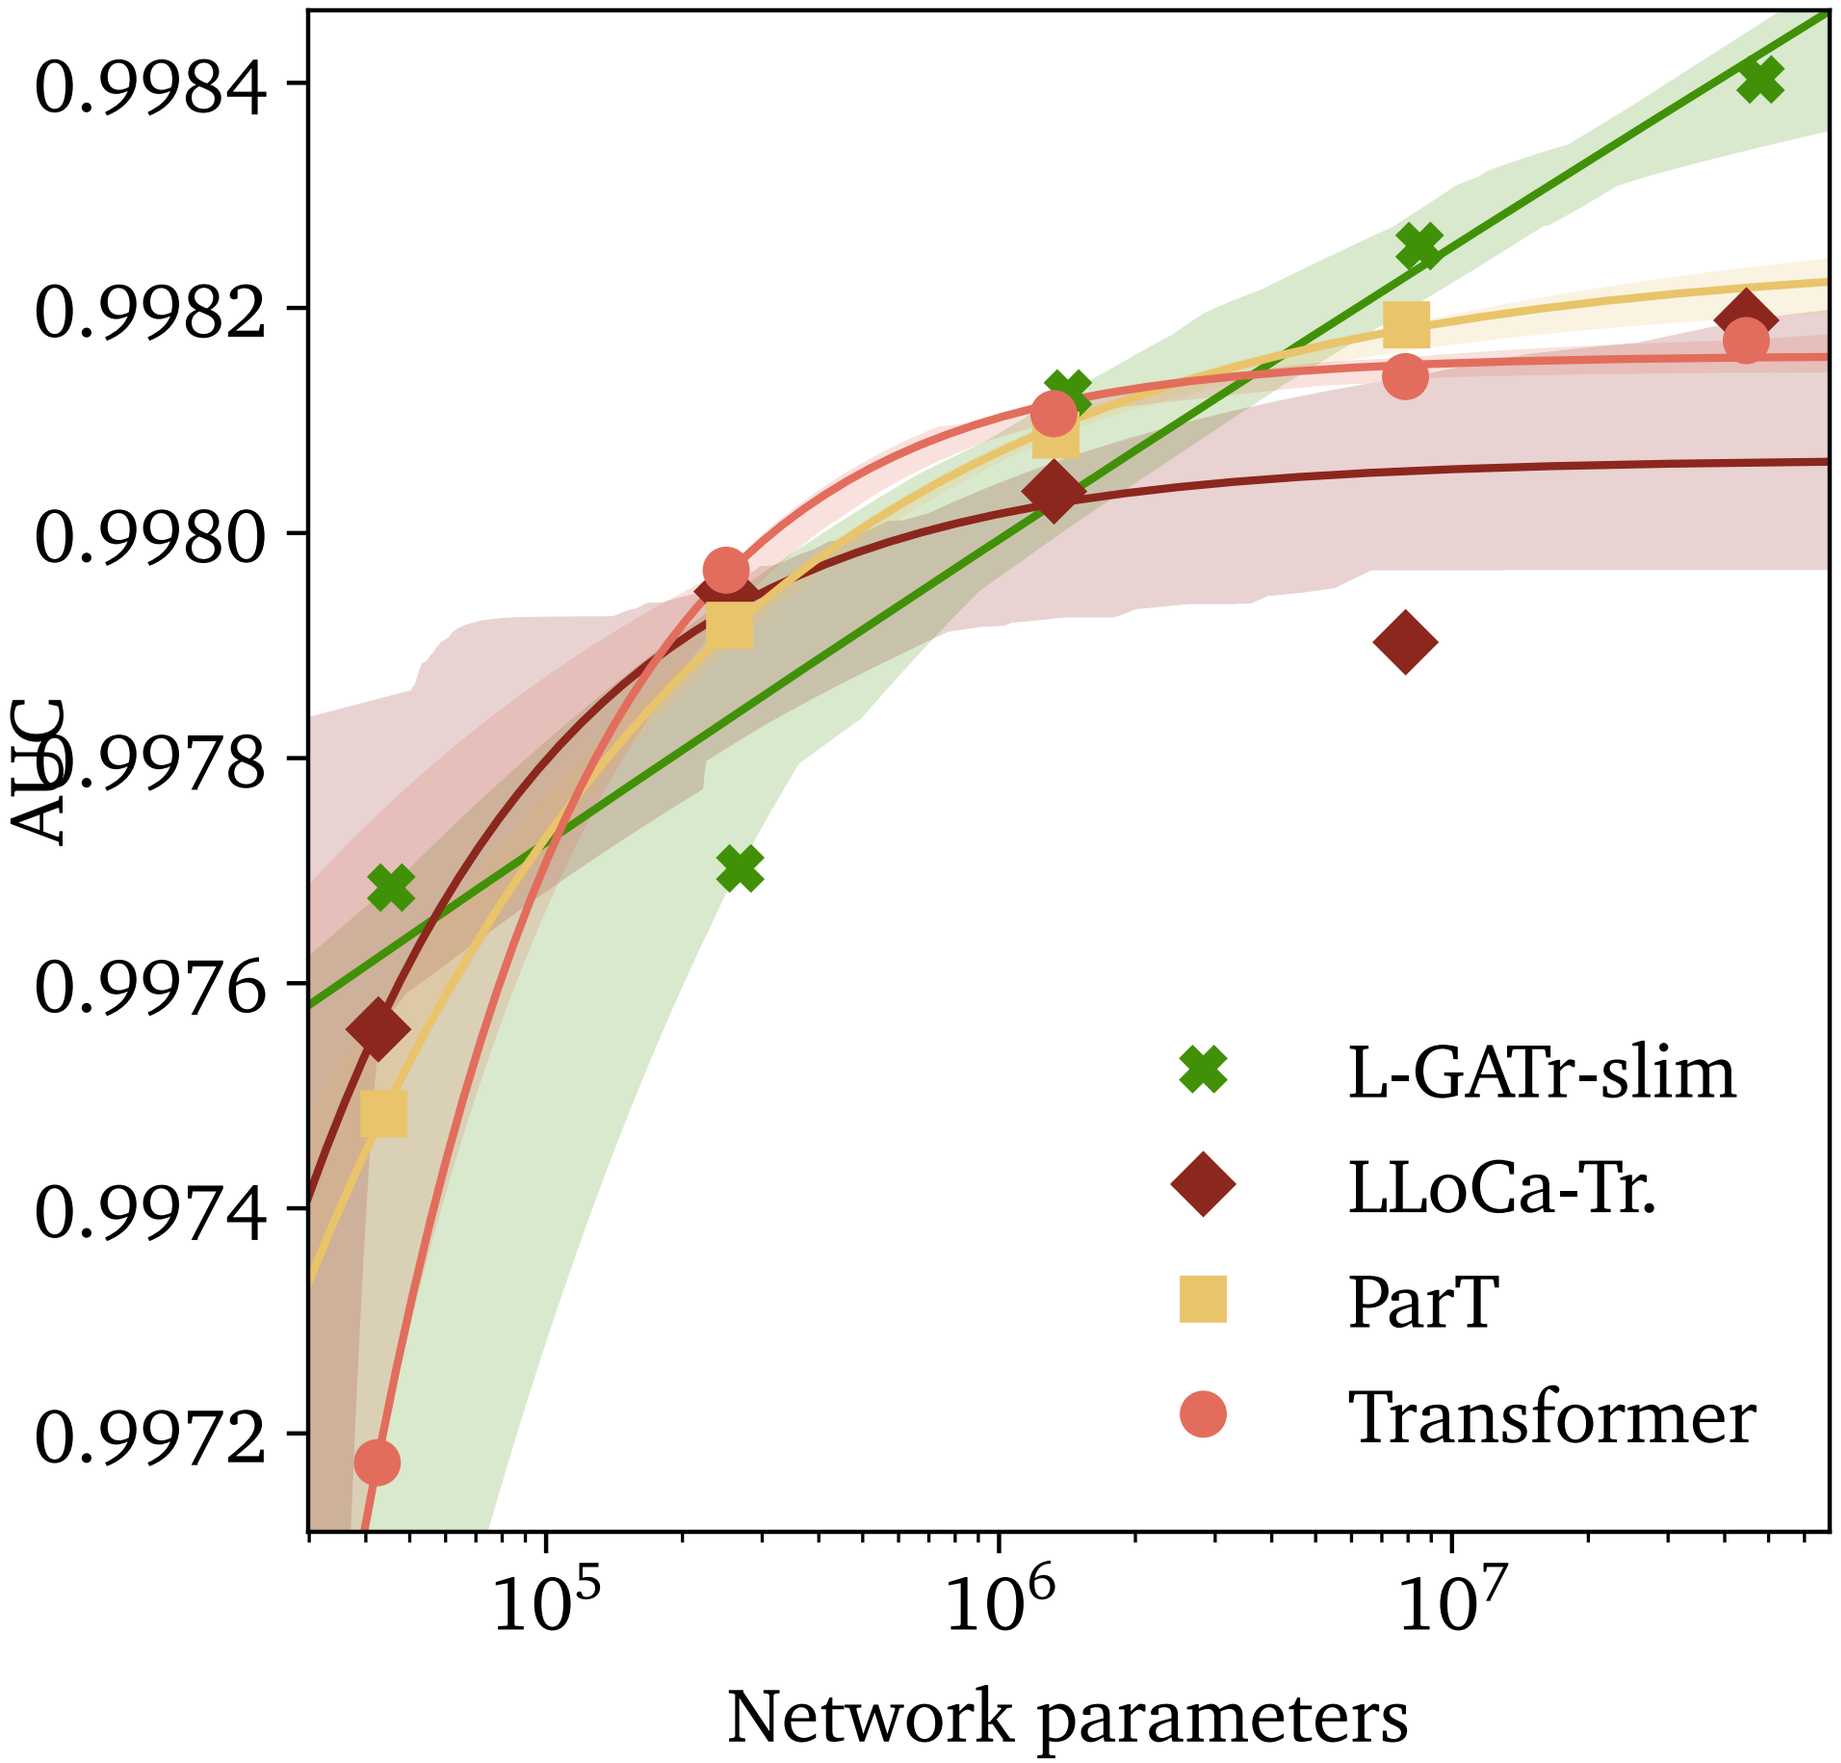
0.9984
0.9982
0.9980
0.9978
0.9976
0.9974
0.9972
AUC
10^5
10^6
10^7
Network parameters
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

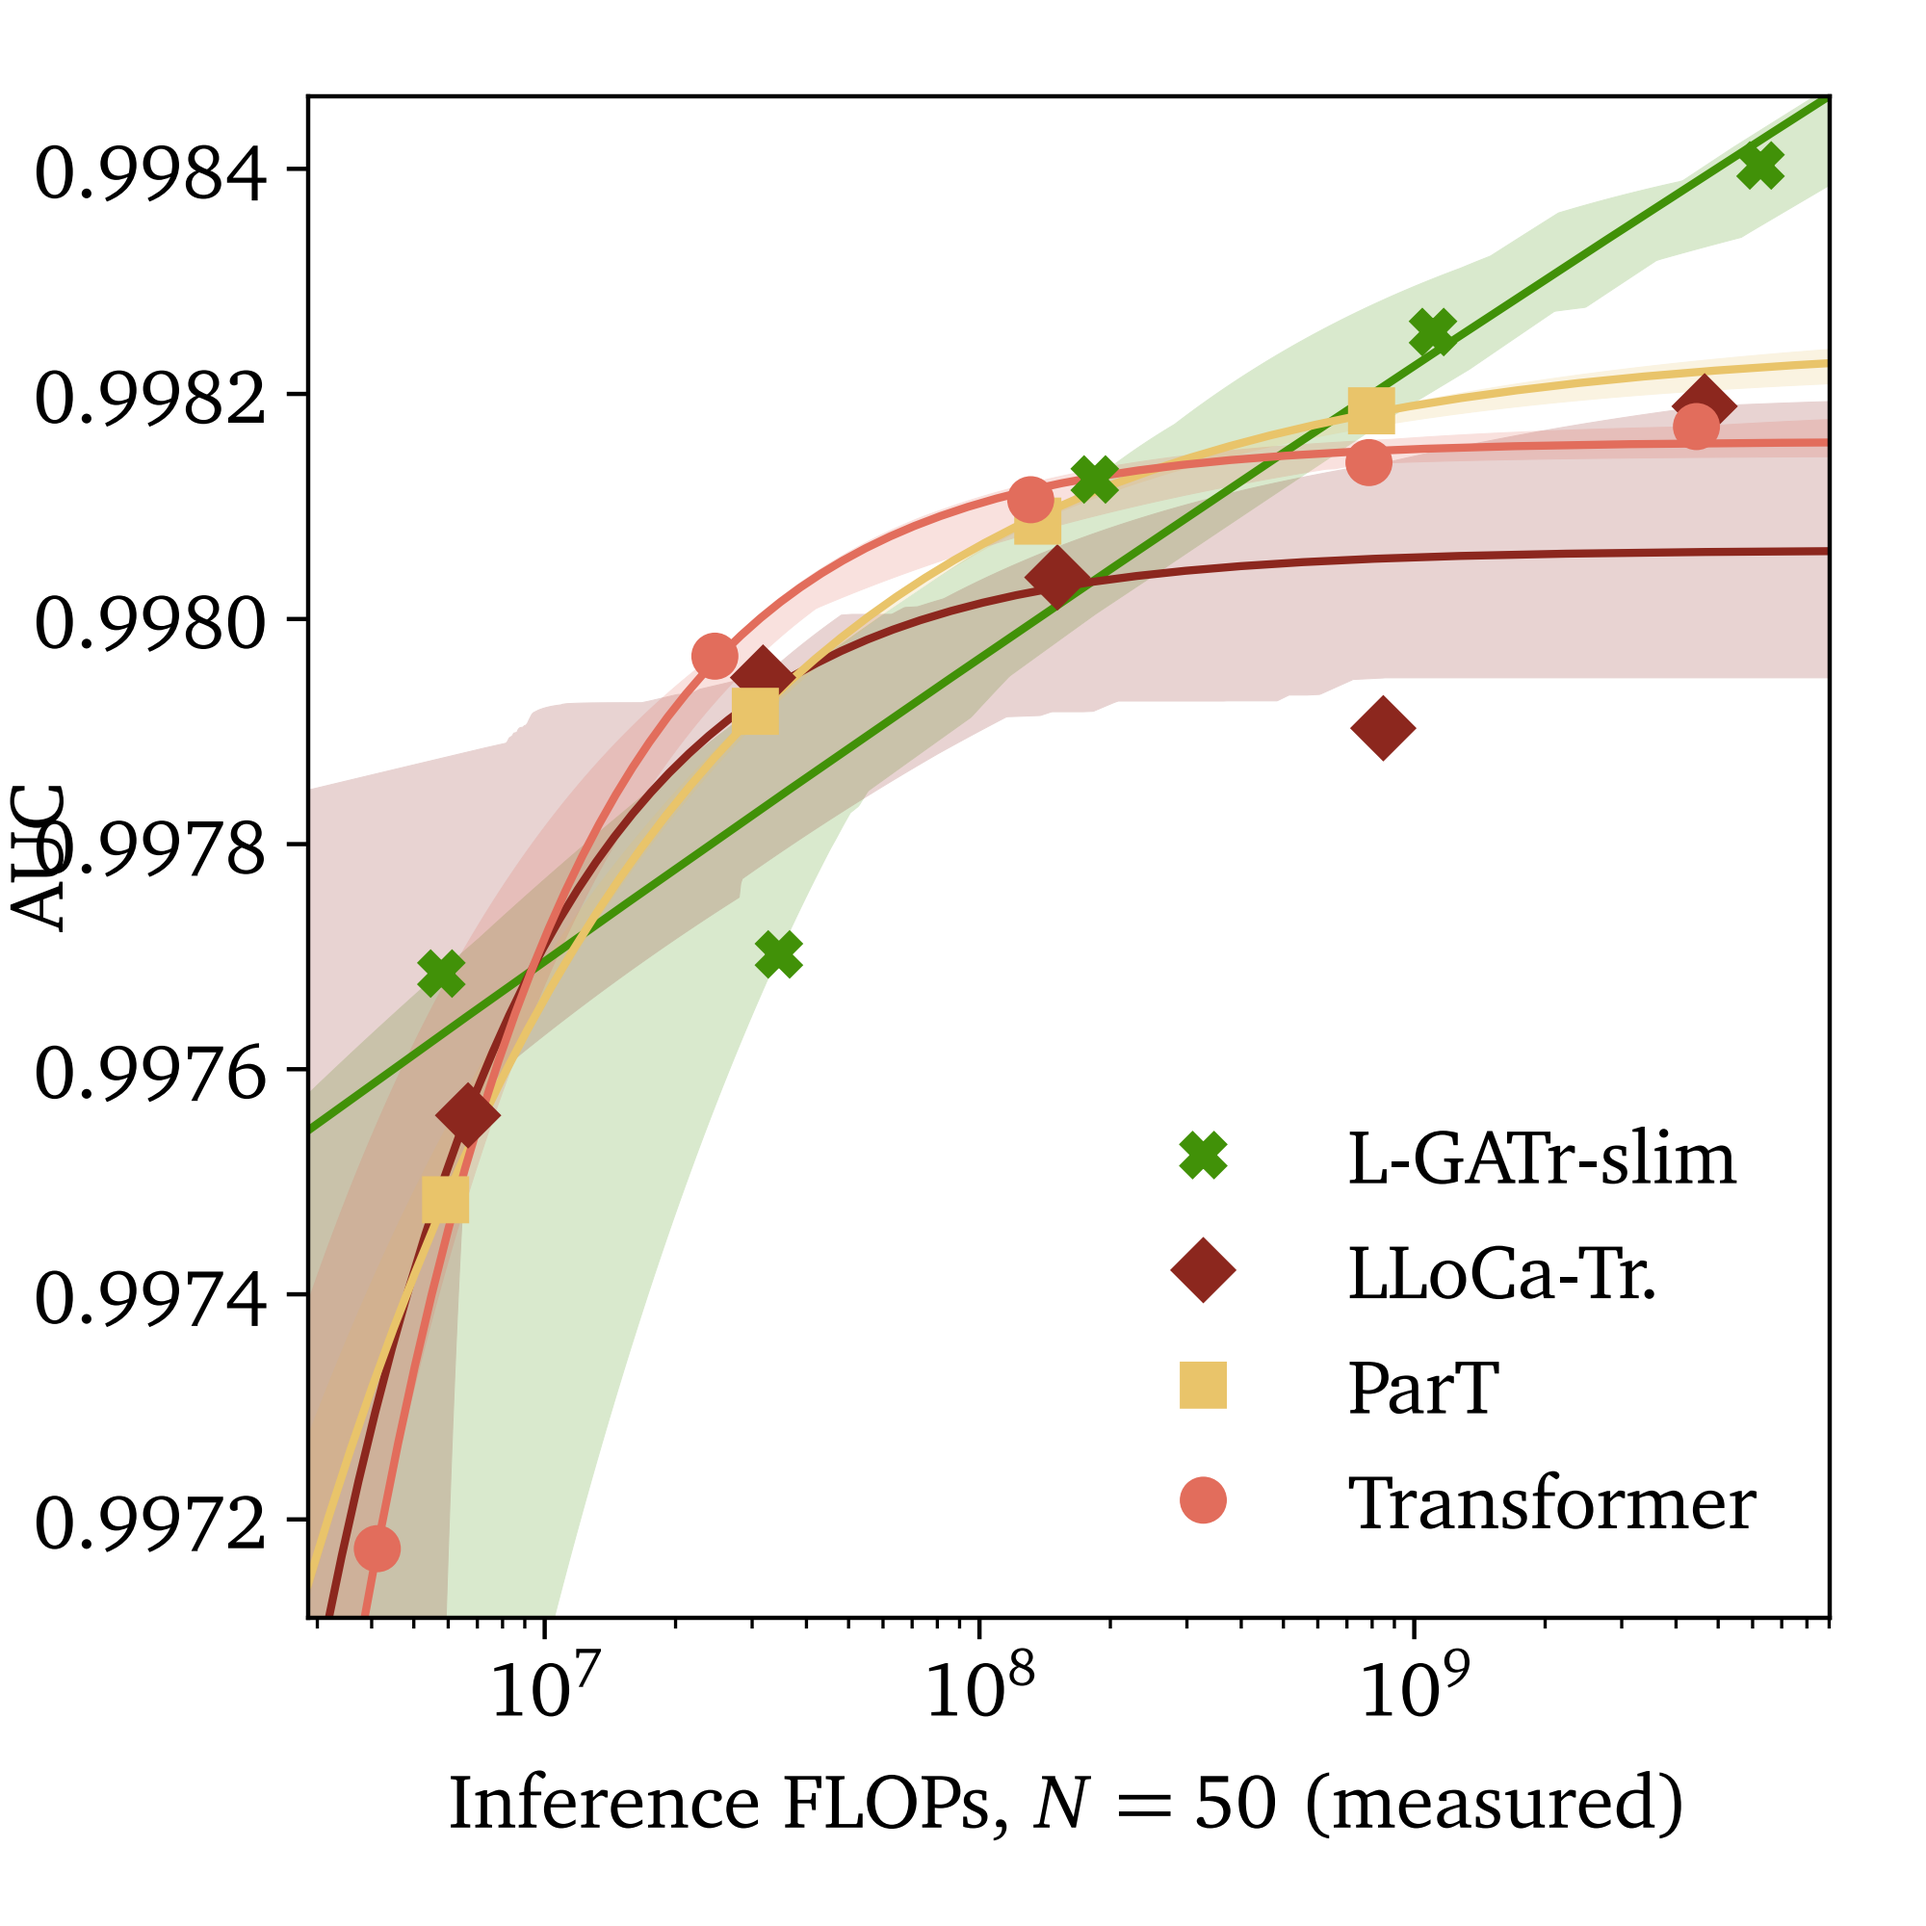
0.9984
0.9982
0.9980
0.9978
0.9976
0.9974
0.9972
AUC
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (measured)
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

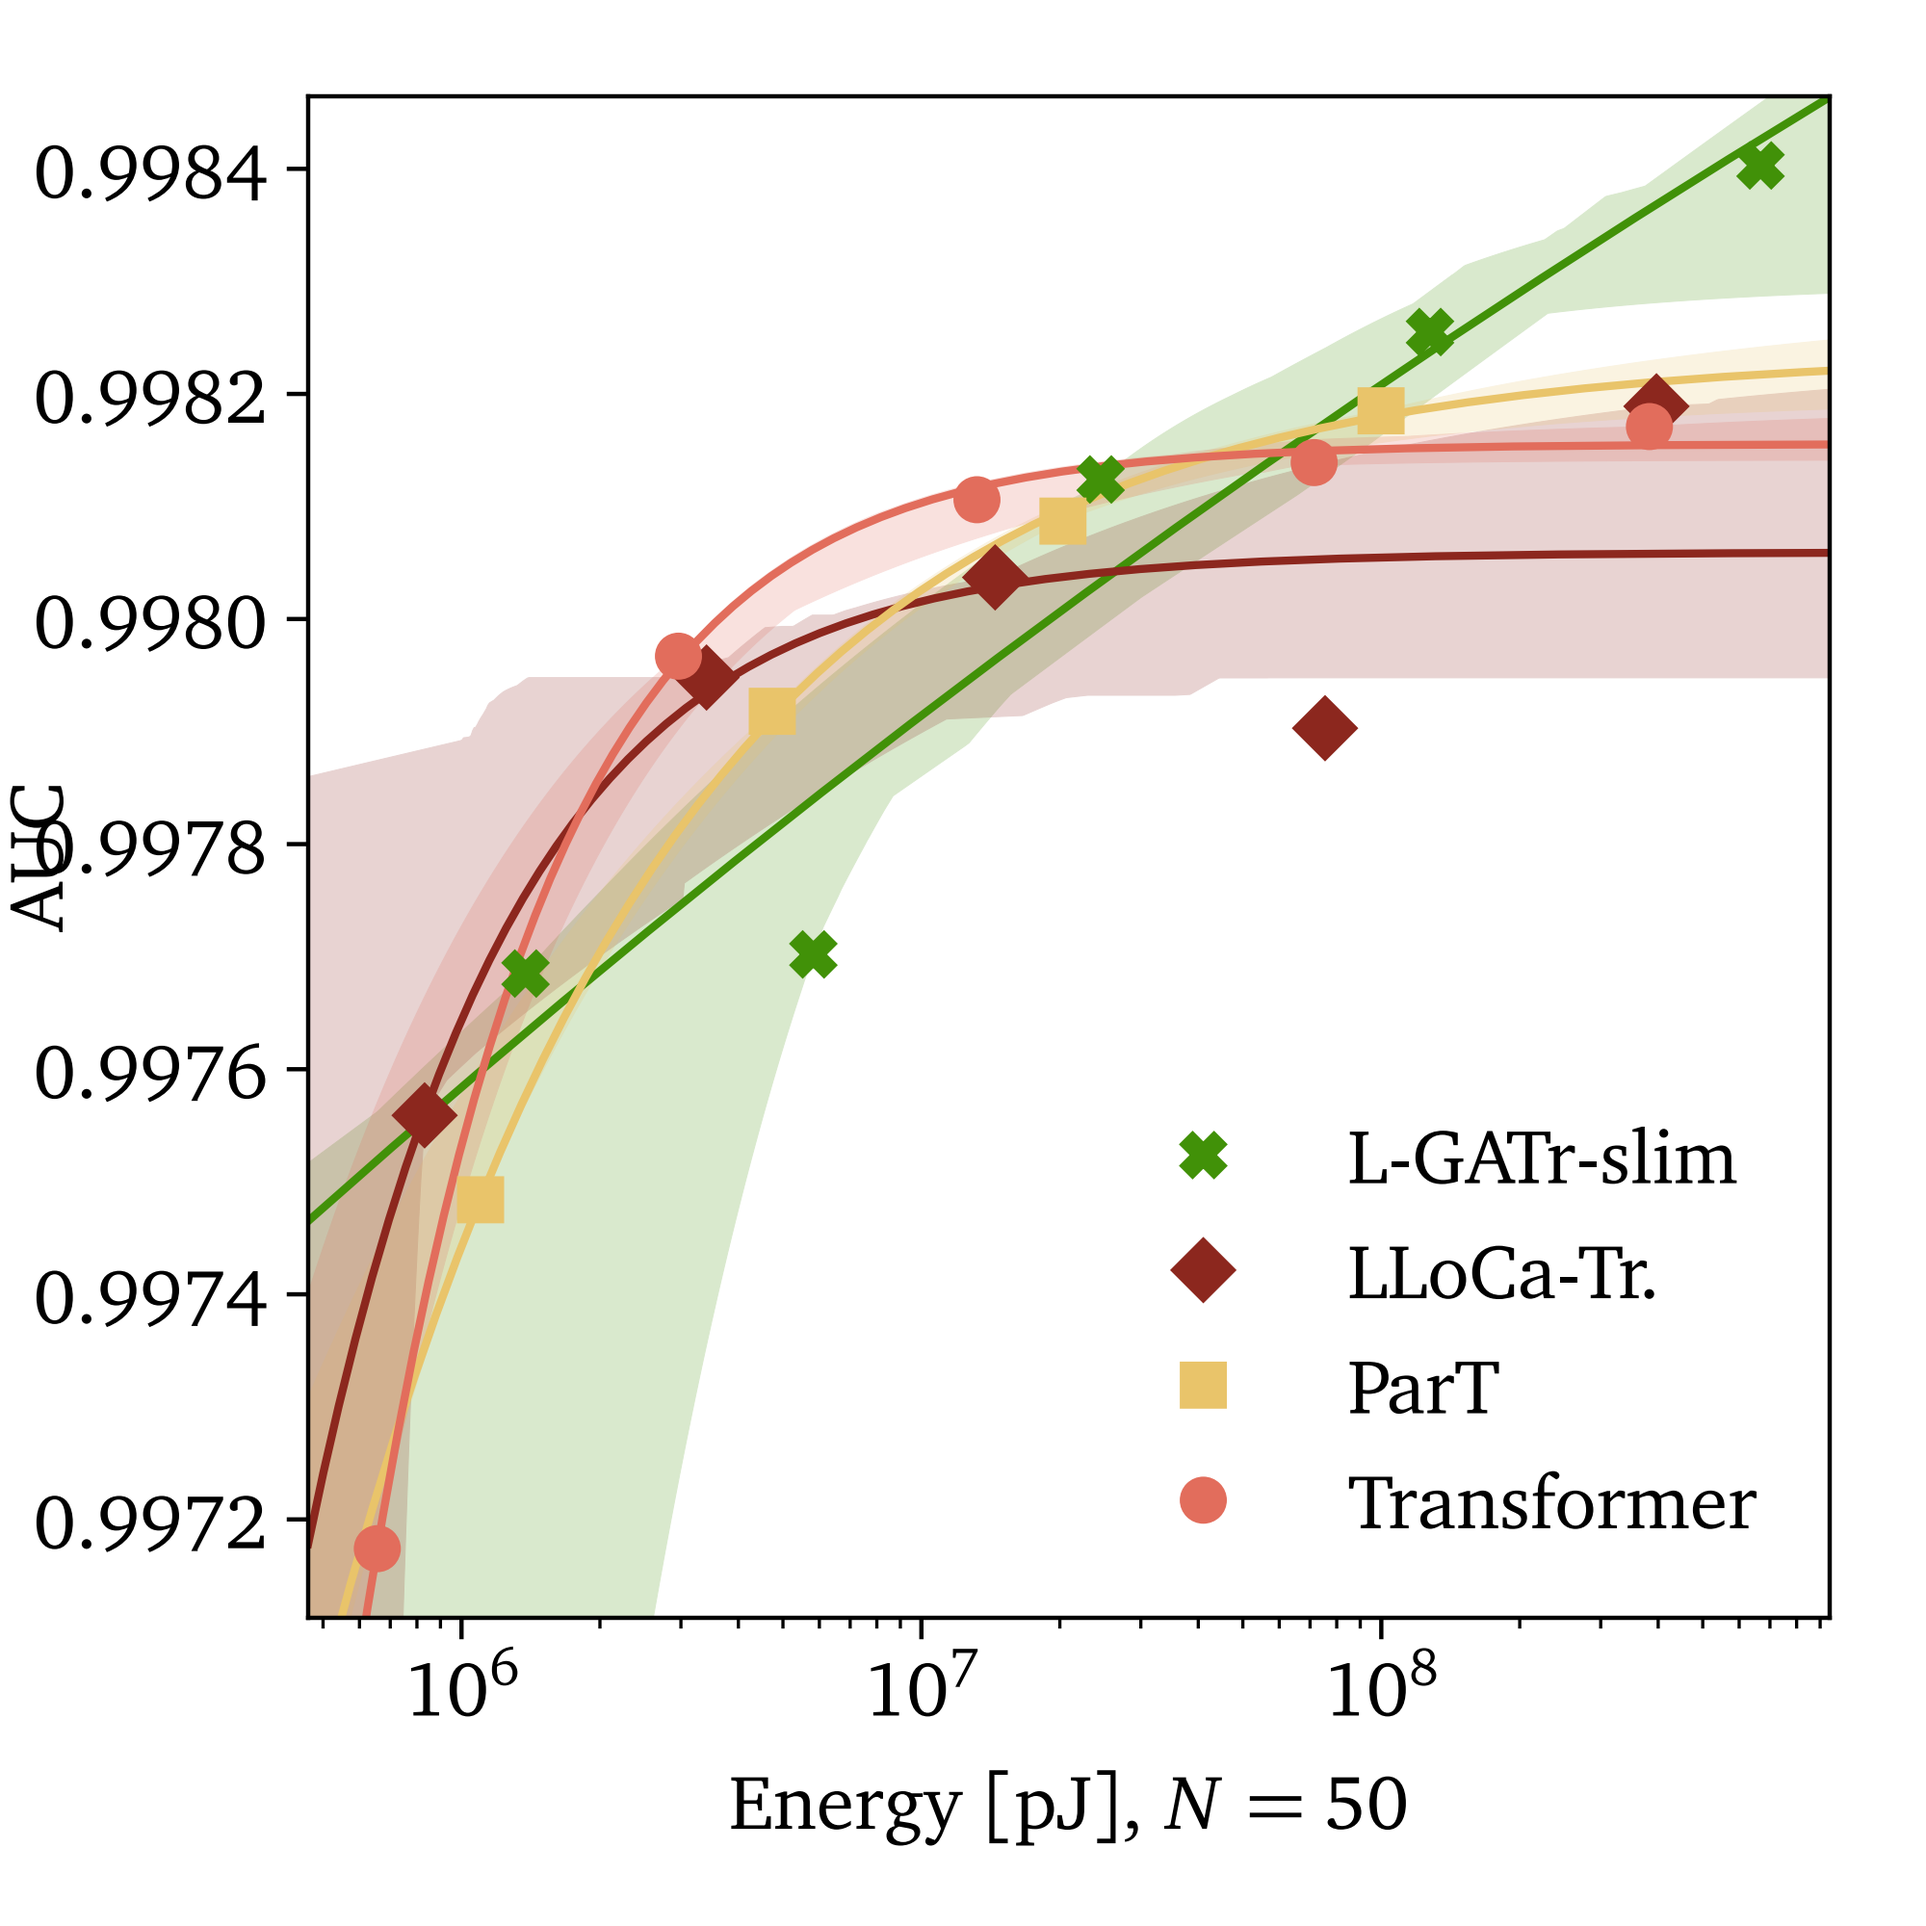
0.9984
0.9982
0.9980
0.9978
0.9976
0.9974
0.9972
AUC
10^6
10^7
10^8
Energy [pJ], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

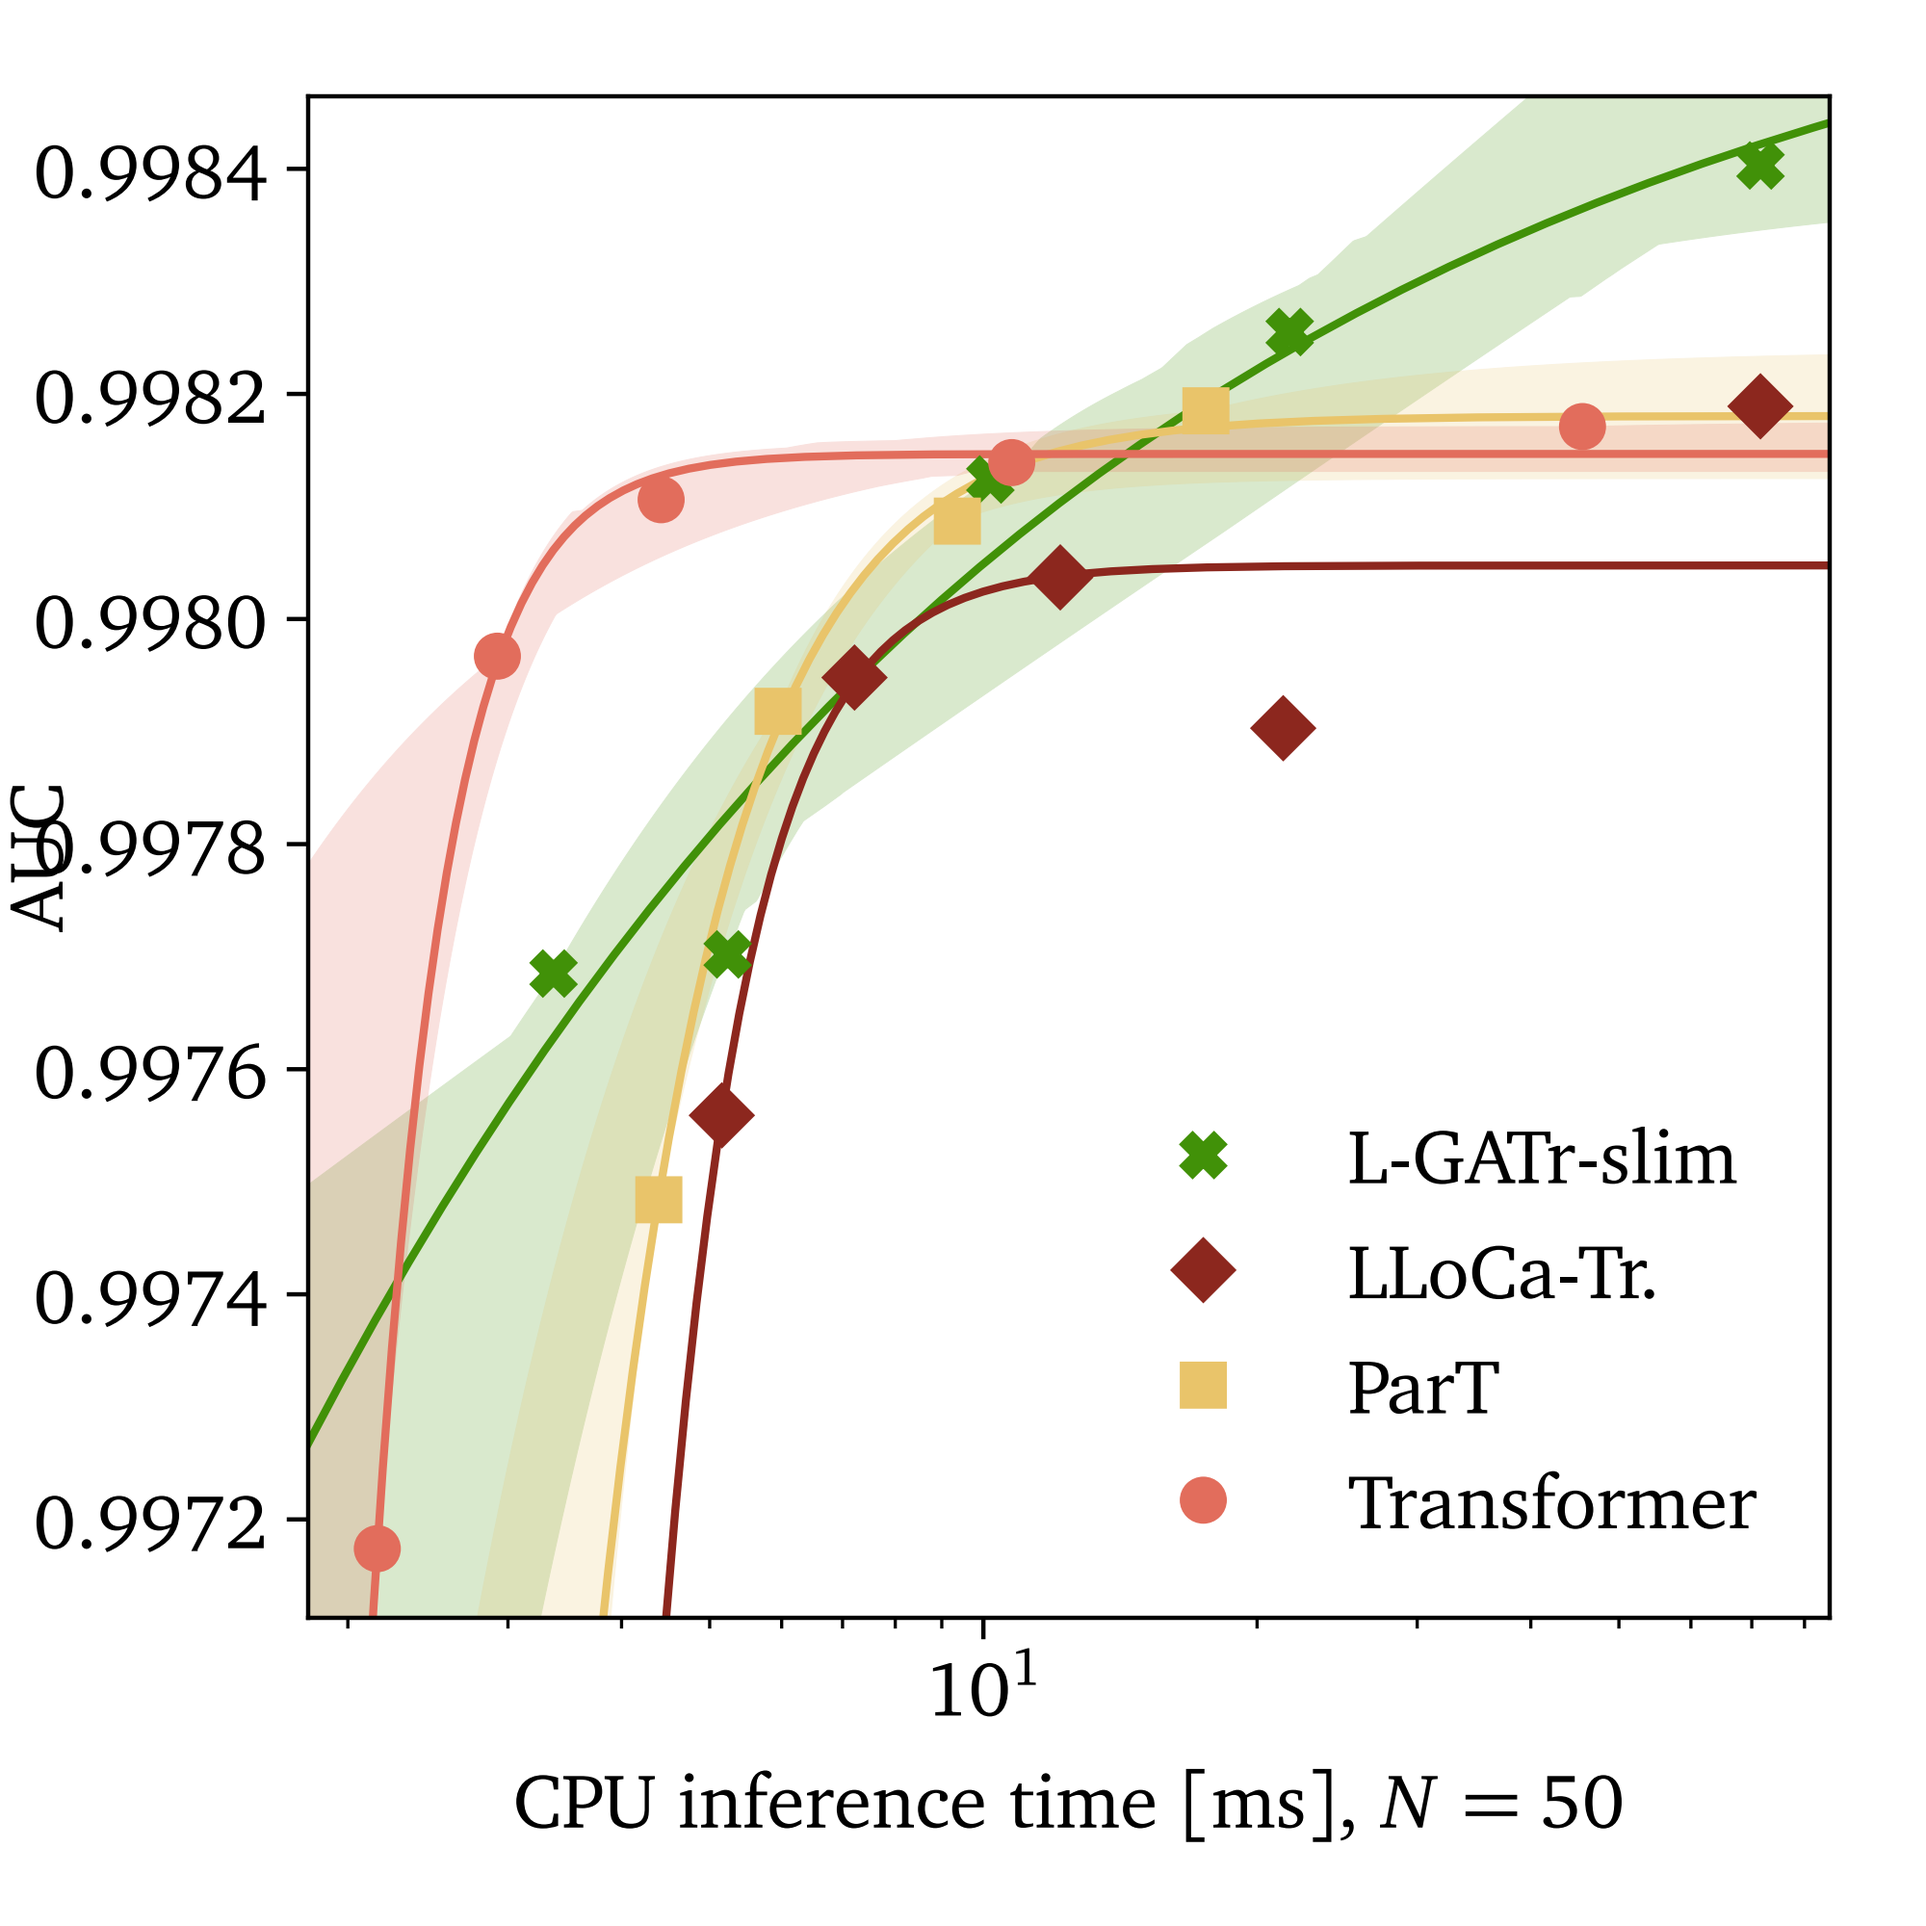
0.9984
0.9982
0.9980
0.9978
0.9976
0.9974
0.9972
AUC
10^1
CPU inference time [ms], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

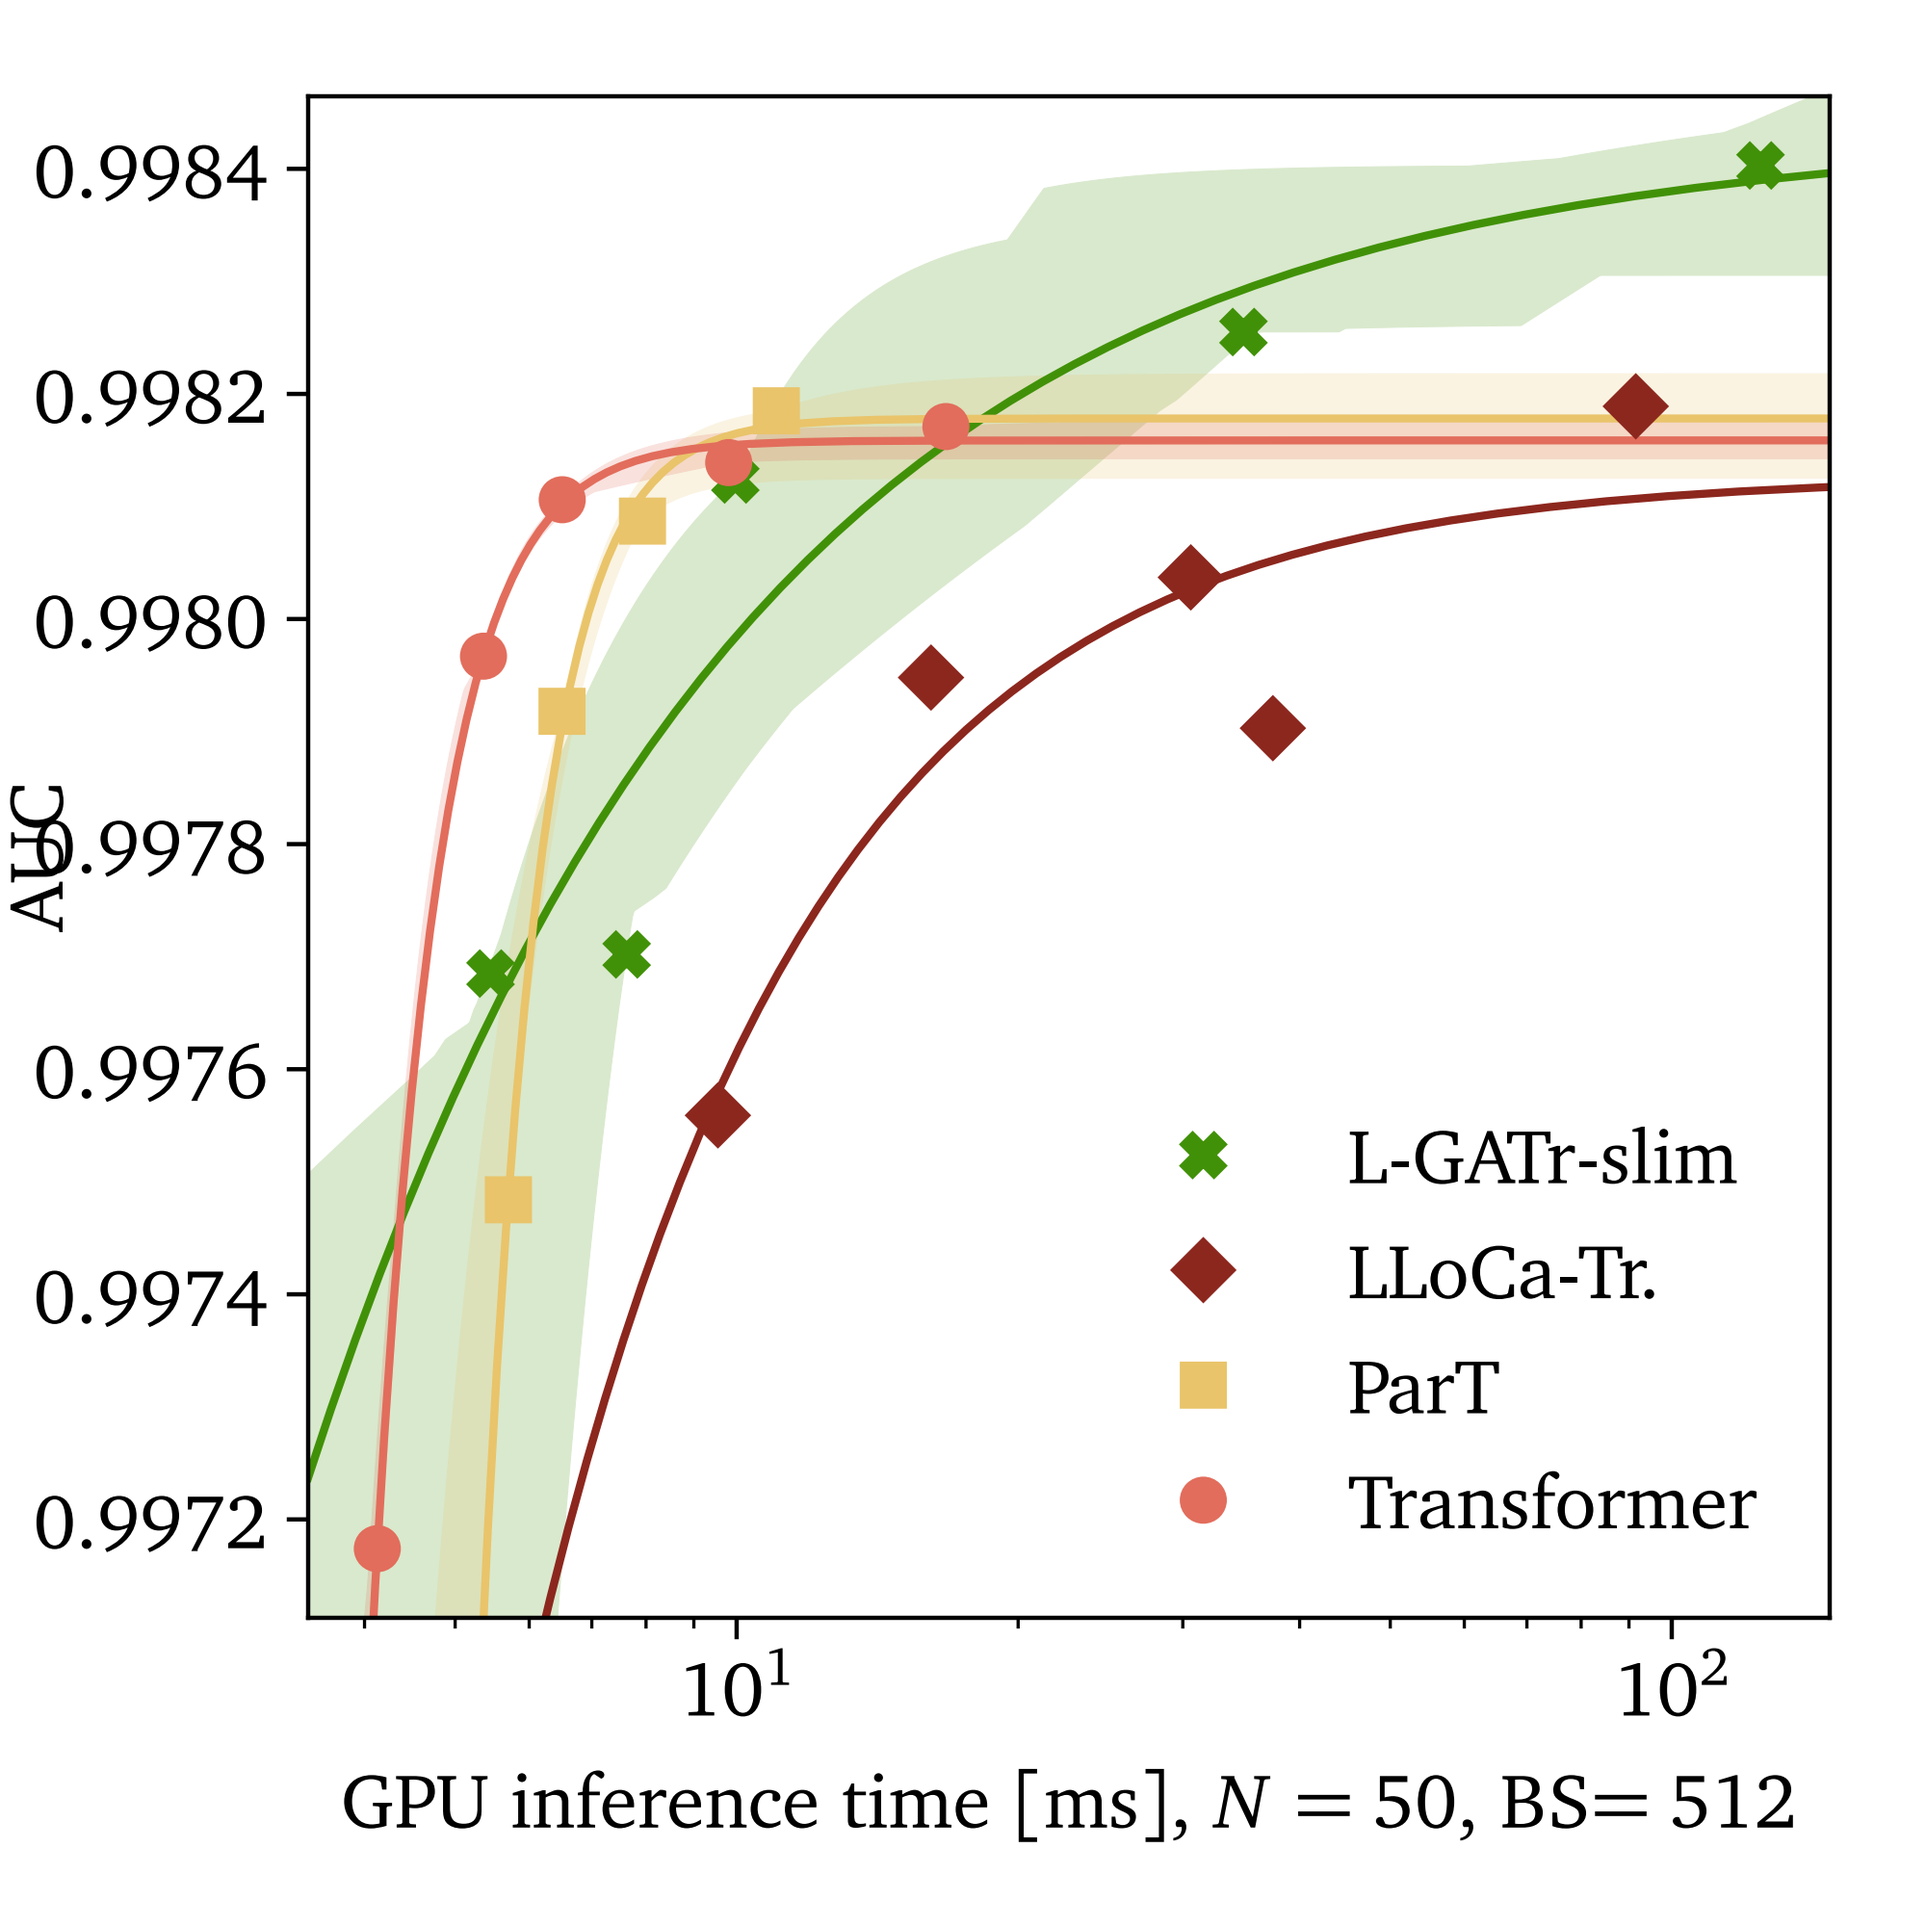
0.9984
0.9982
0.9980
0.9978
0.9976
0.9974
0.9972
AUC
10^1
10^2
GPU inference time [ms], $N = 50$, BS= 512
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

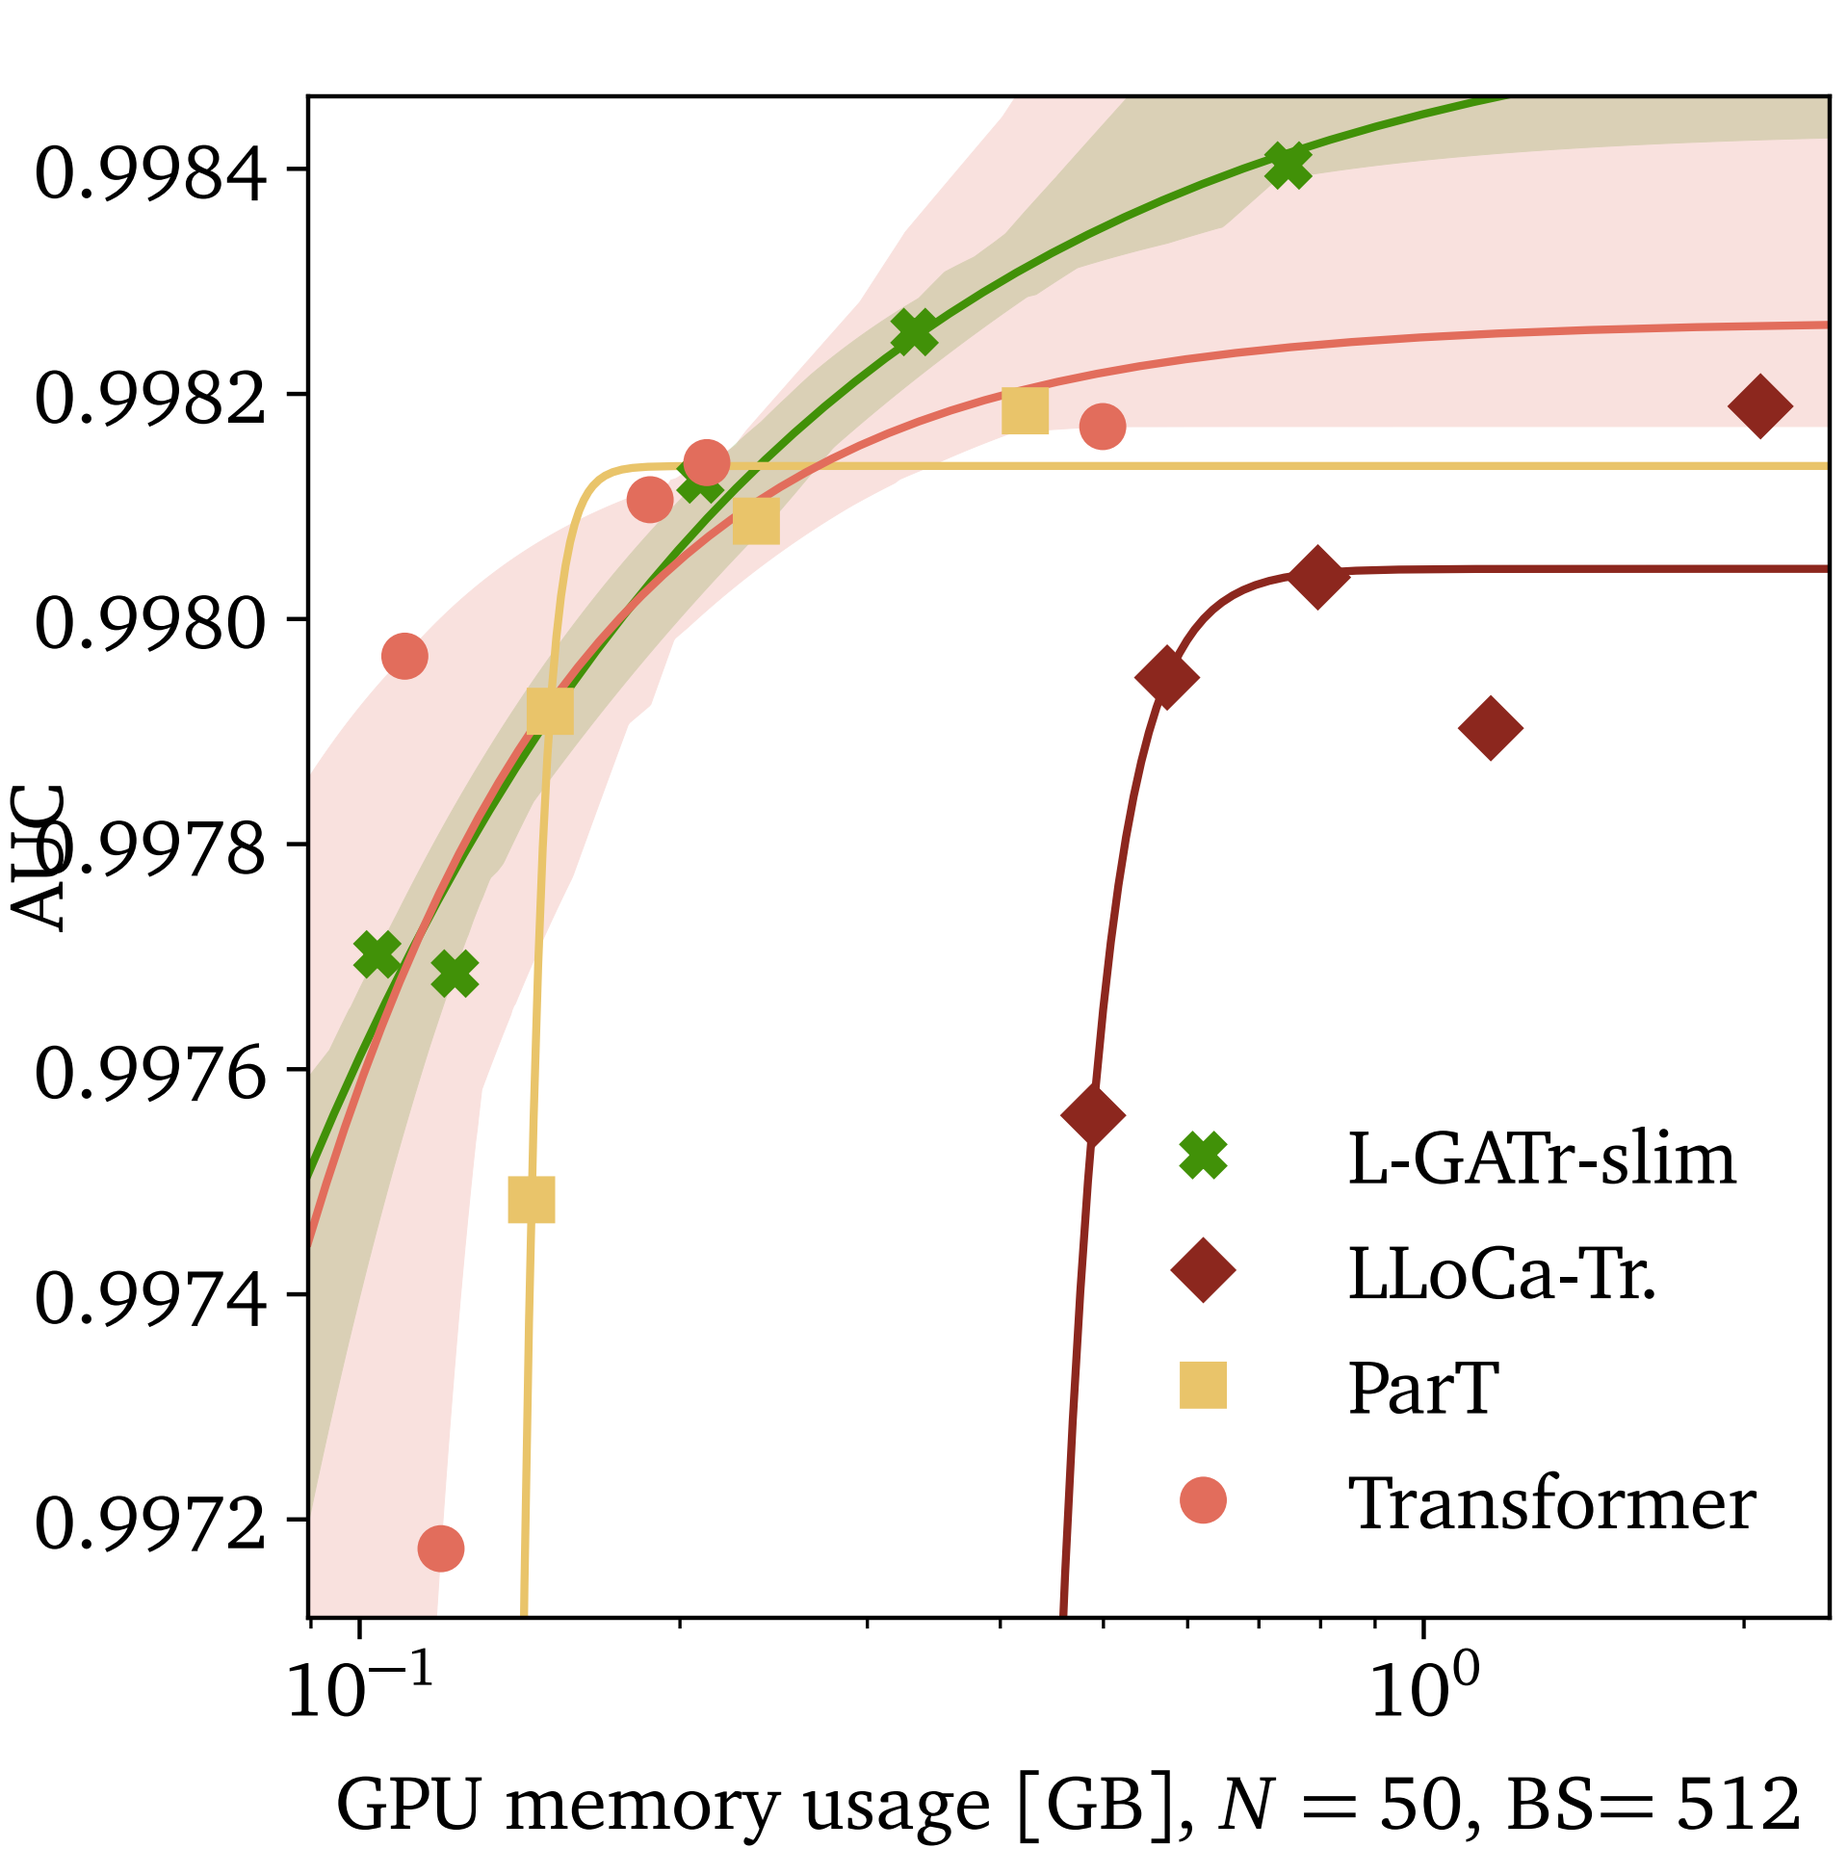
0.9984
0.9982
0.9980
0.9978
0.9976
0.9974
0.9972
AUC
10^{-1}
10^{0}
GPU memory usage [GB], $N = 50$, BS$= 512$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

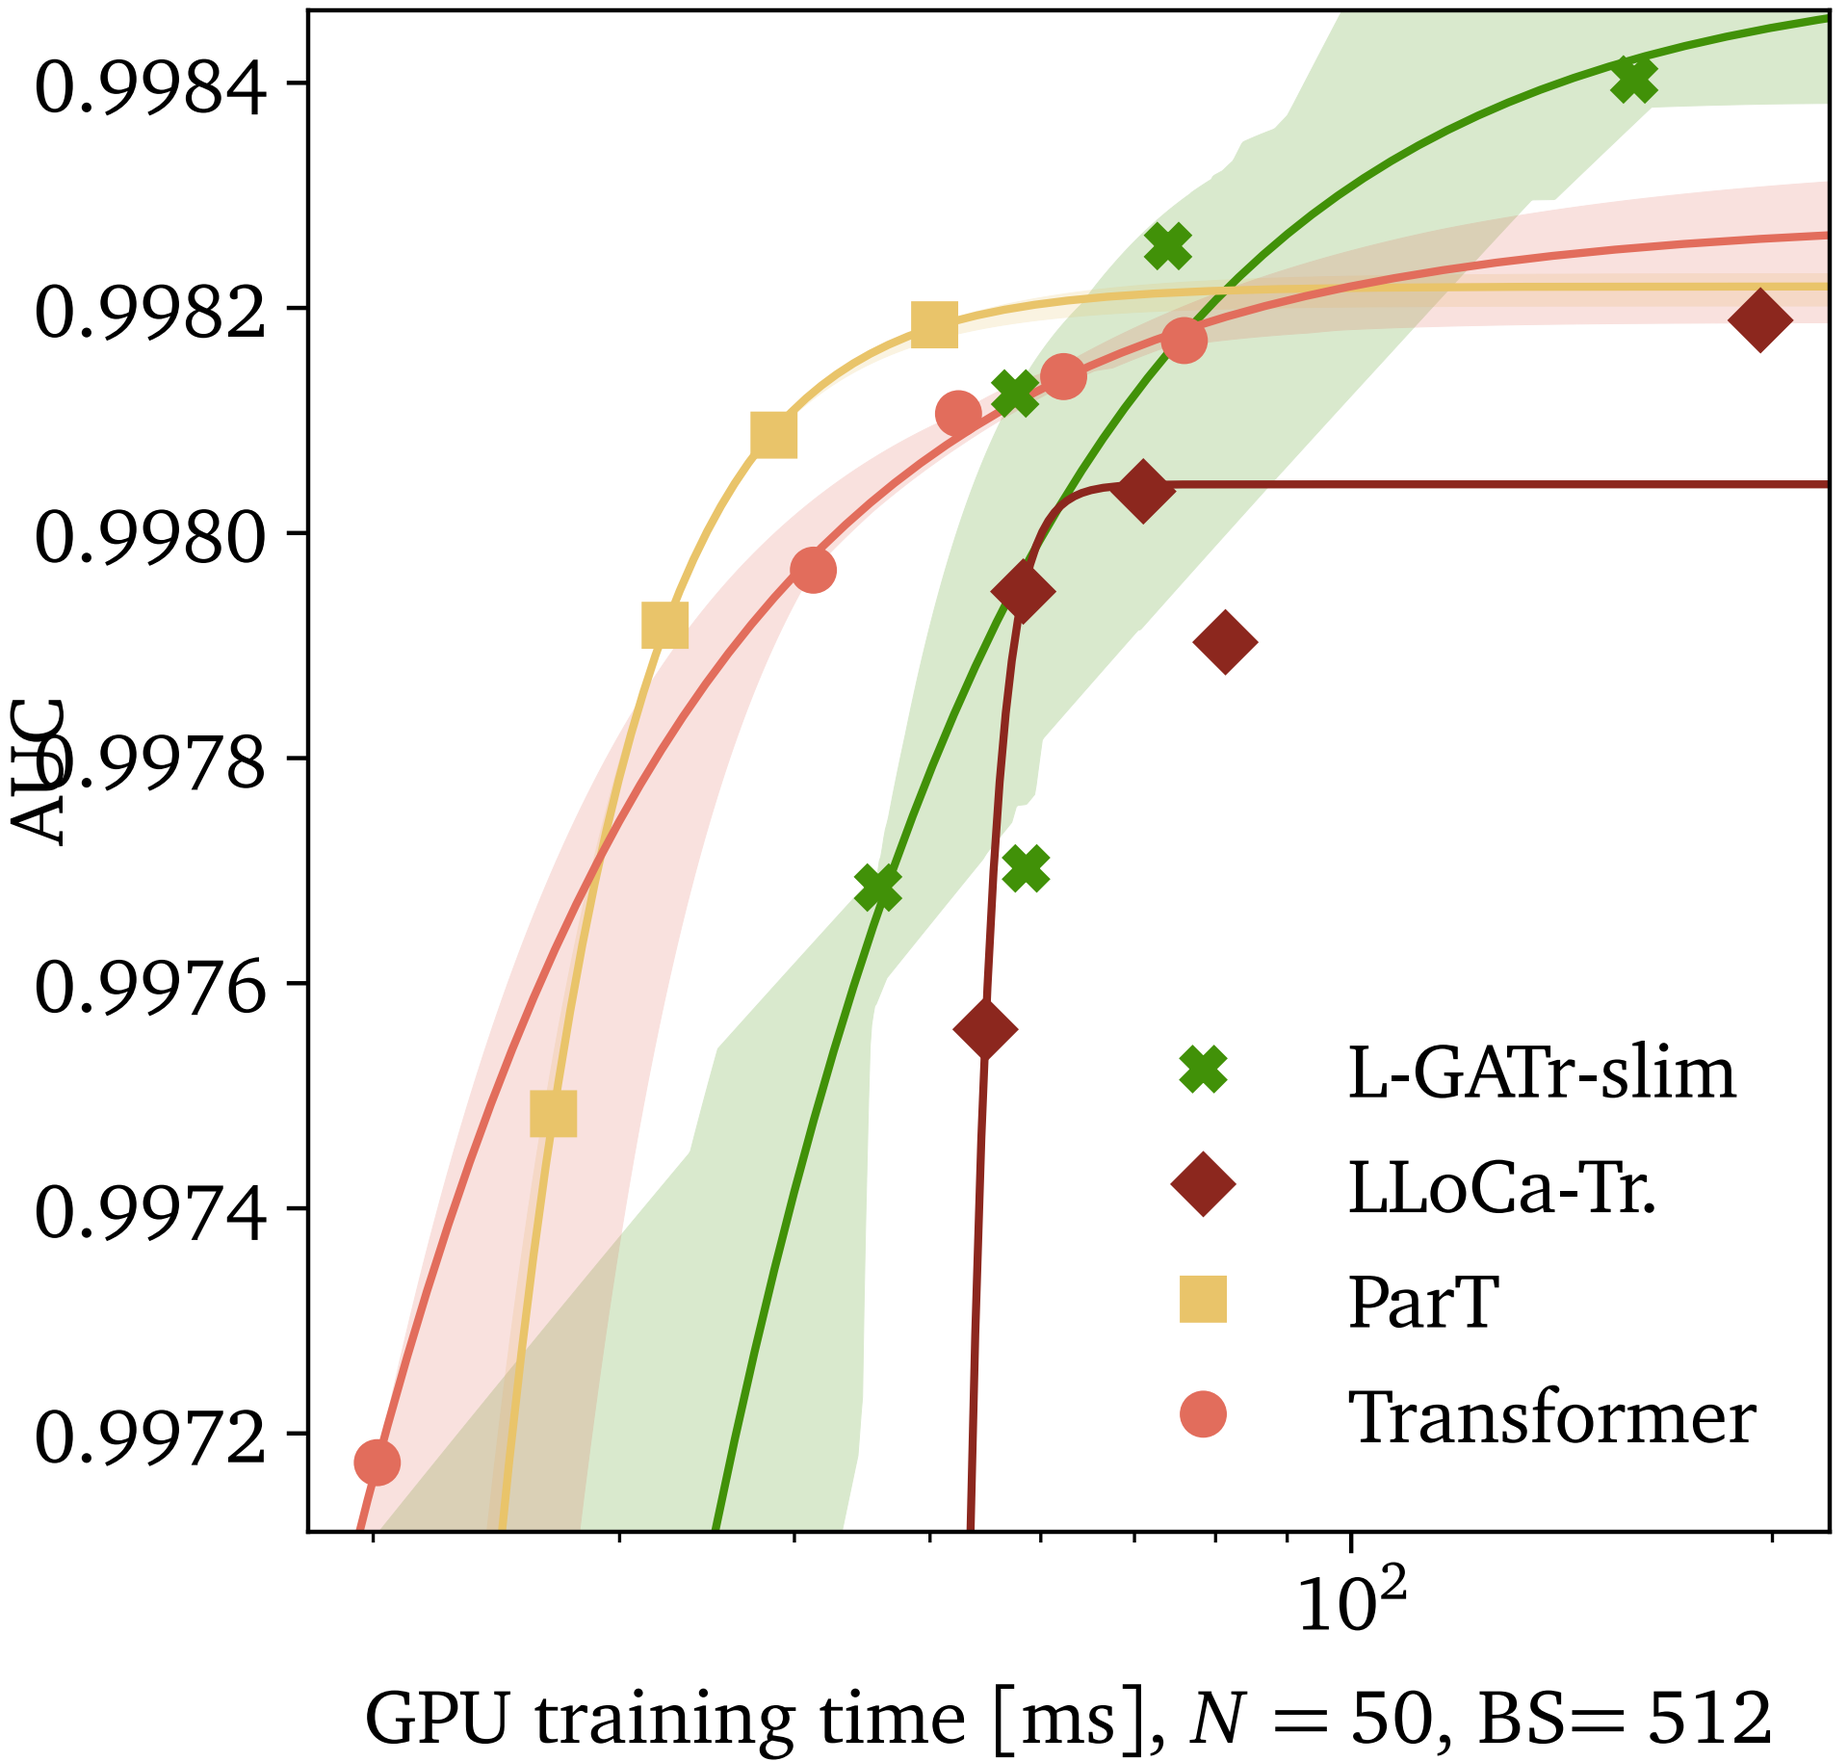
0.9984
0.9982
0.9980
0.9978
0.9976
0.9974
0.9972
AUC
10^2
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
GPU training time [ms], $N = 50$, BS$= 512$

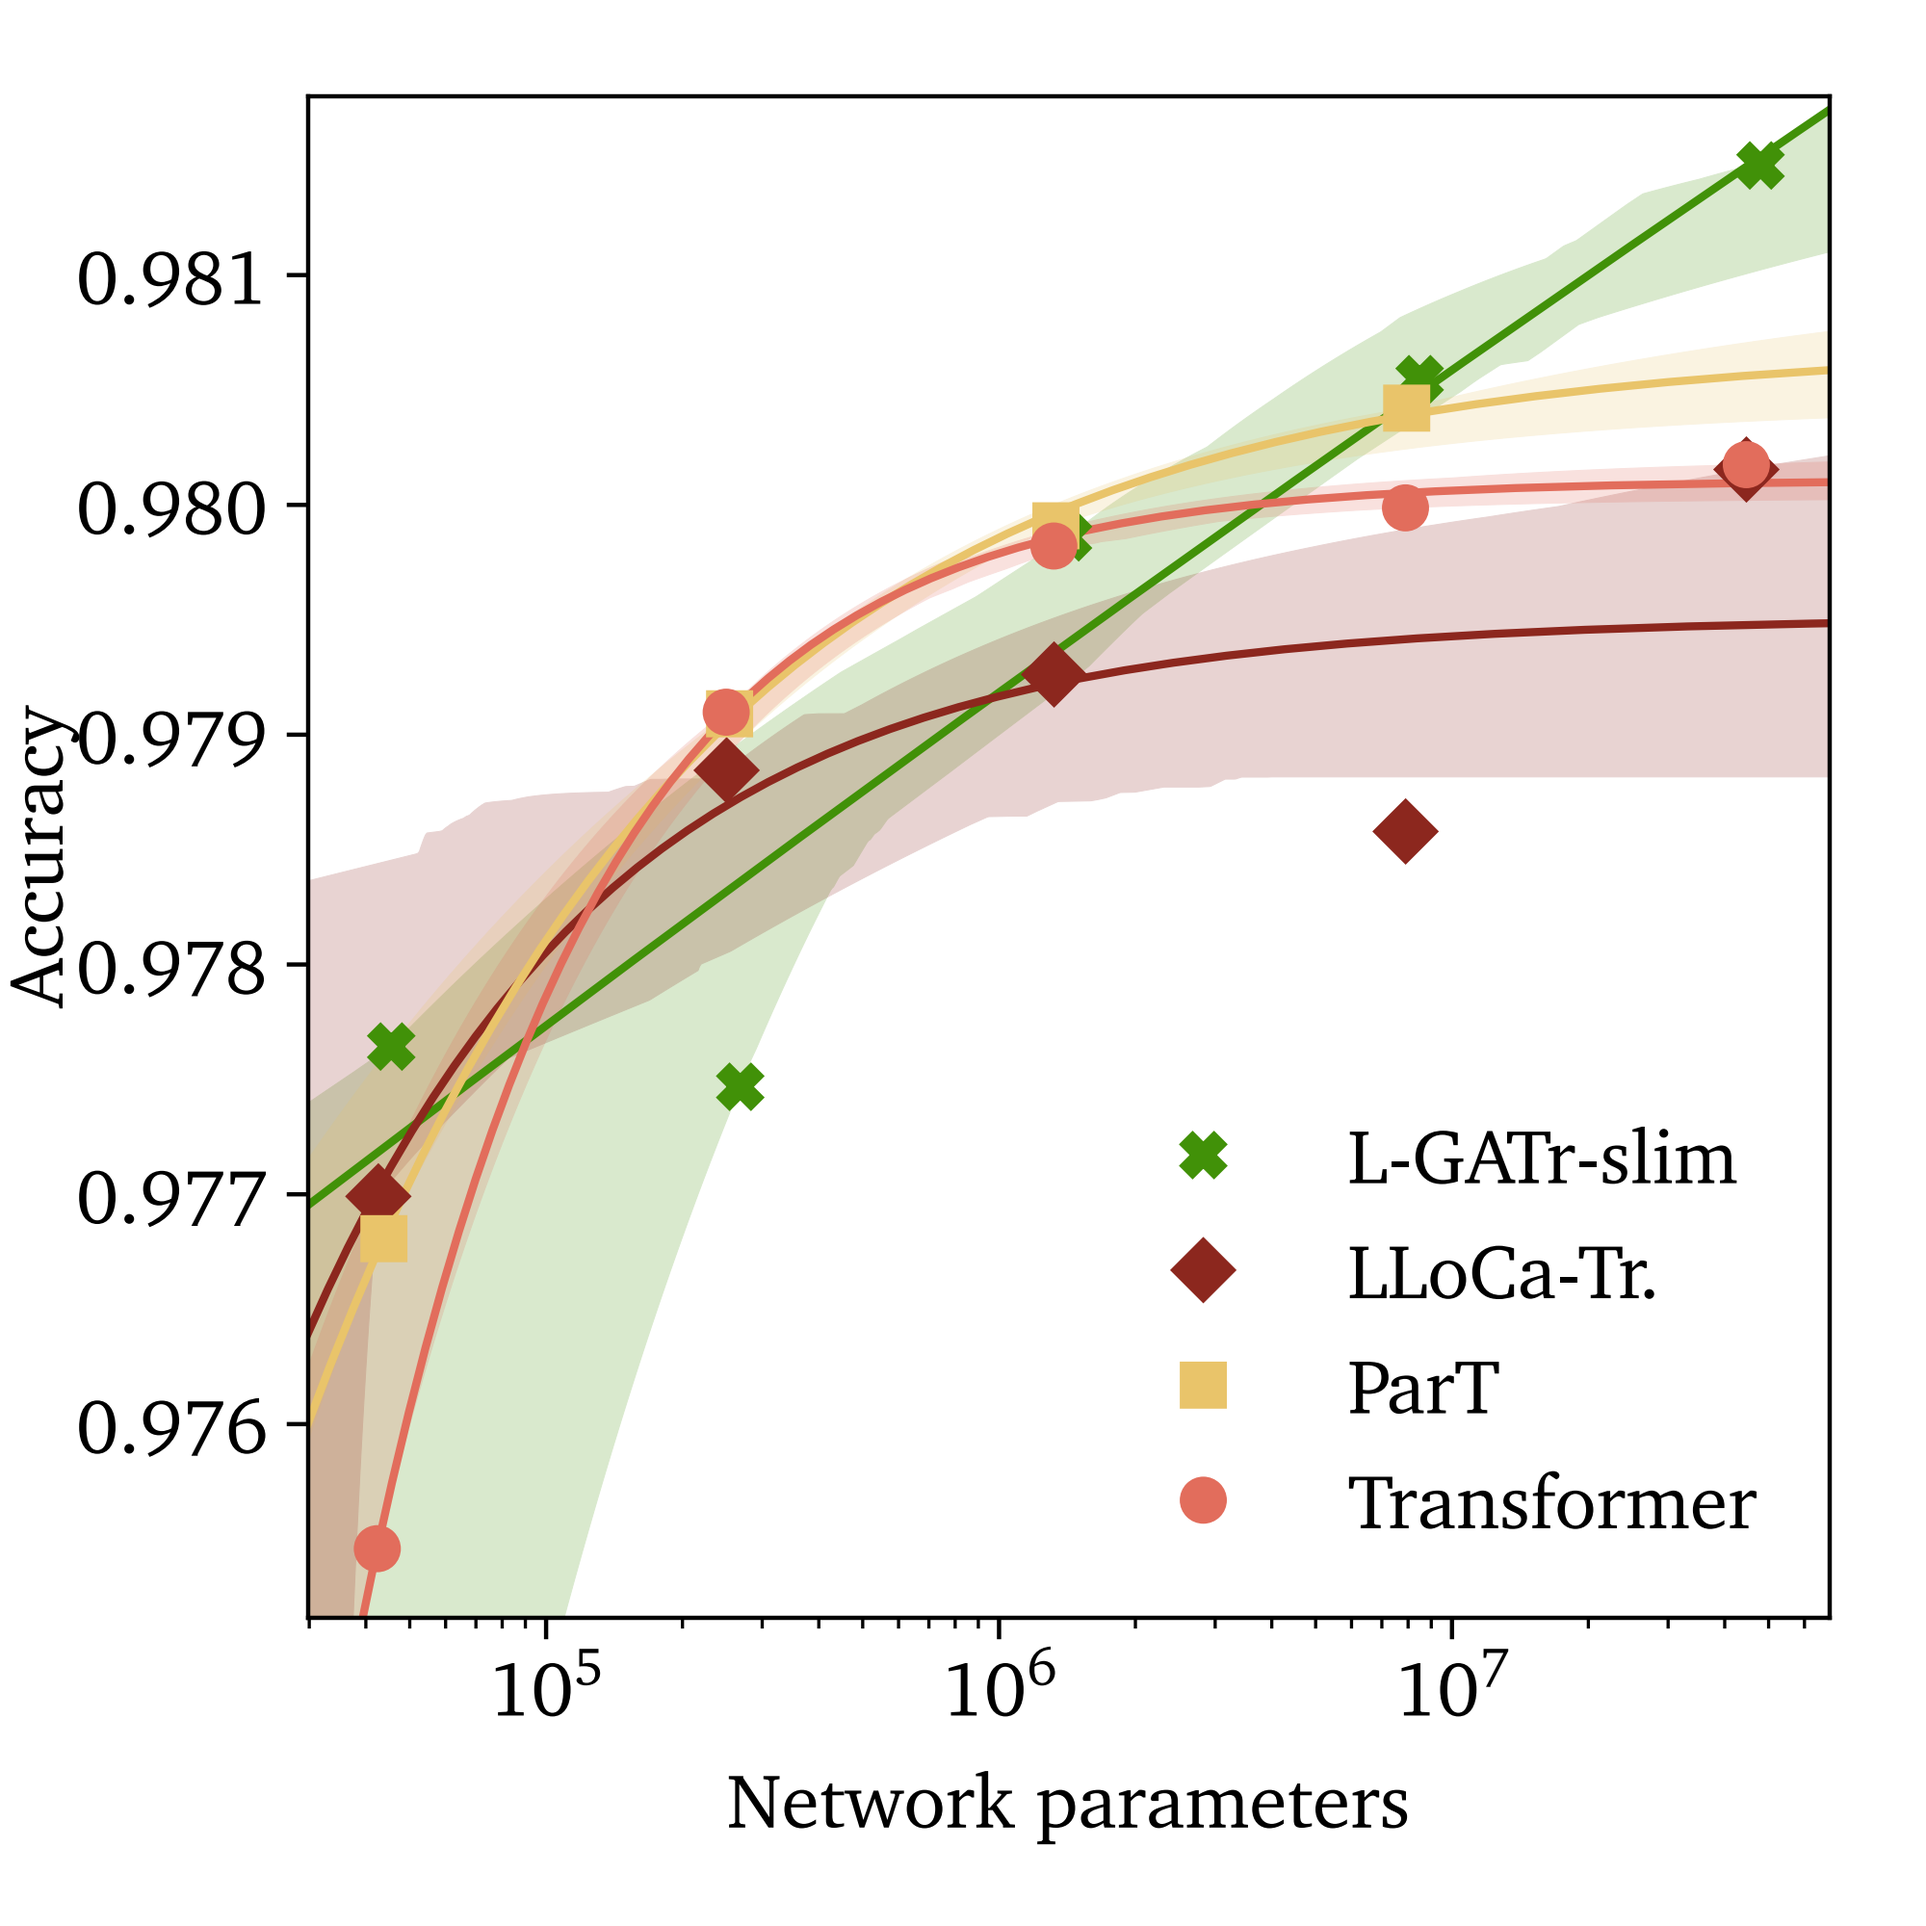
0.981
0.980
0.979
0.978
0.977
0.976
Accuracy
10^5
10^6
10^7
Network parameters
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

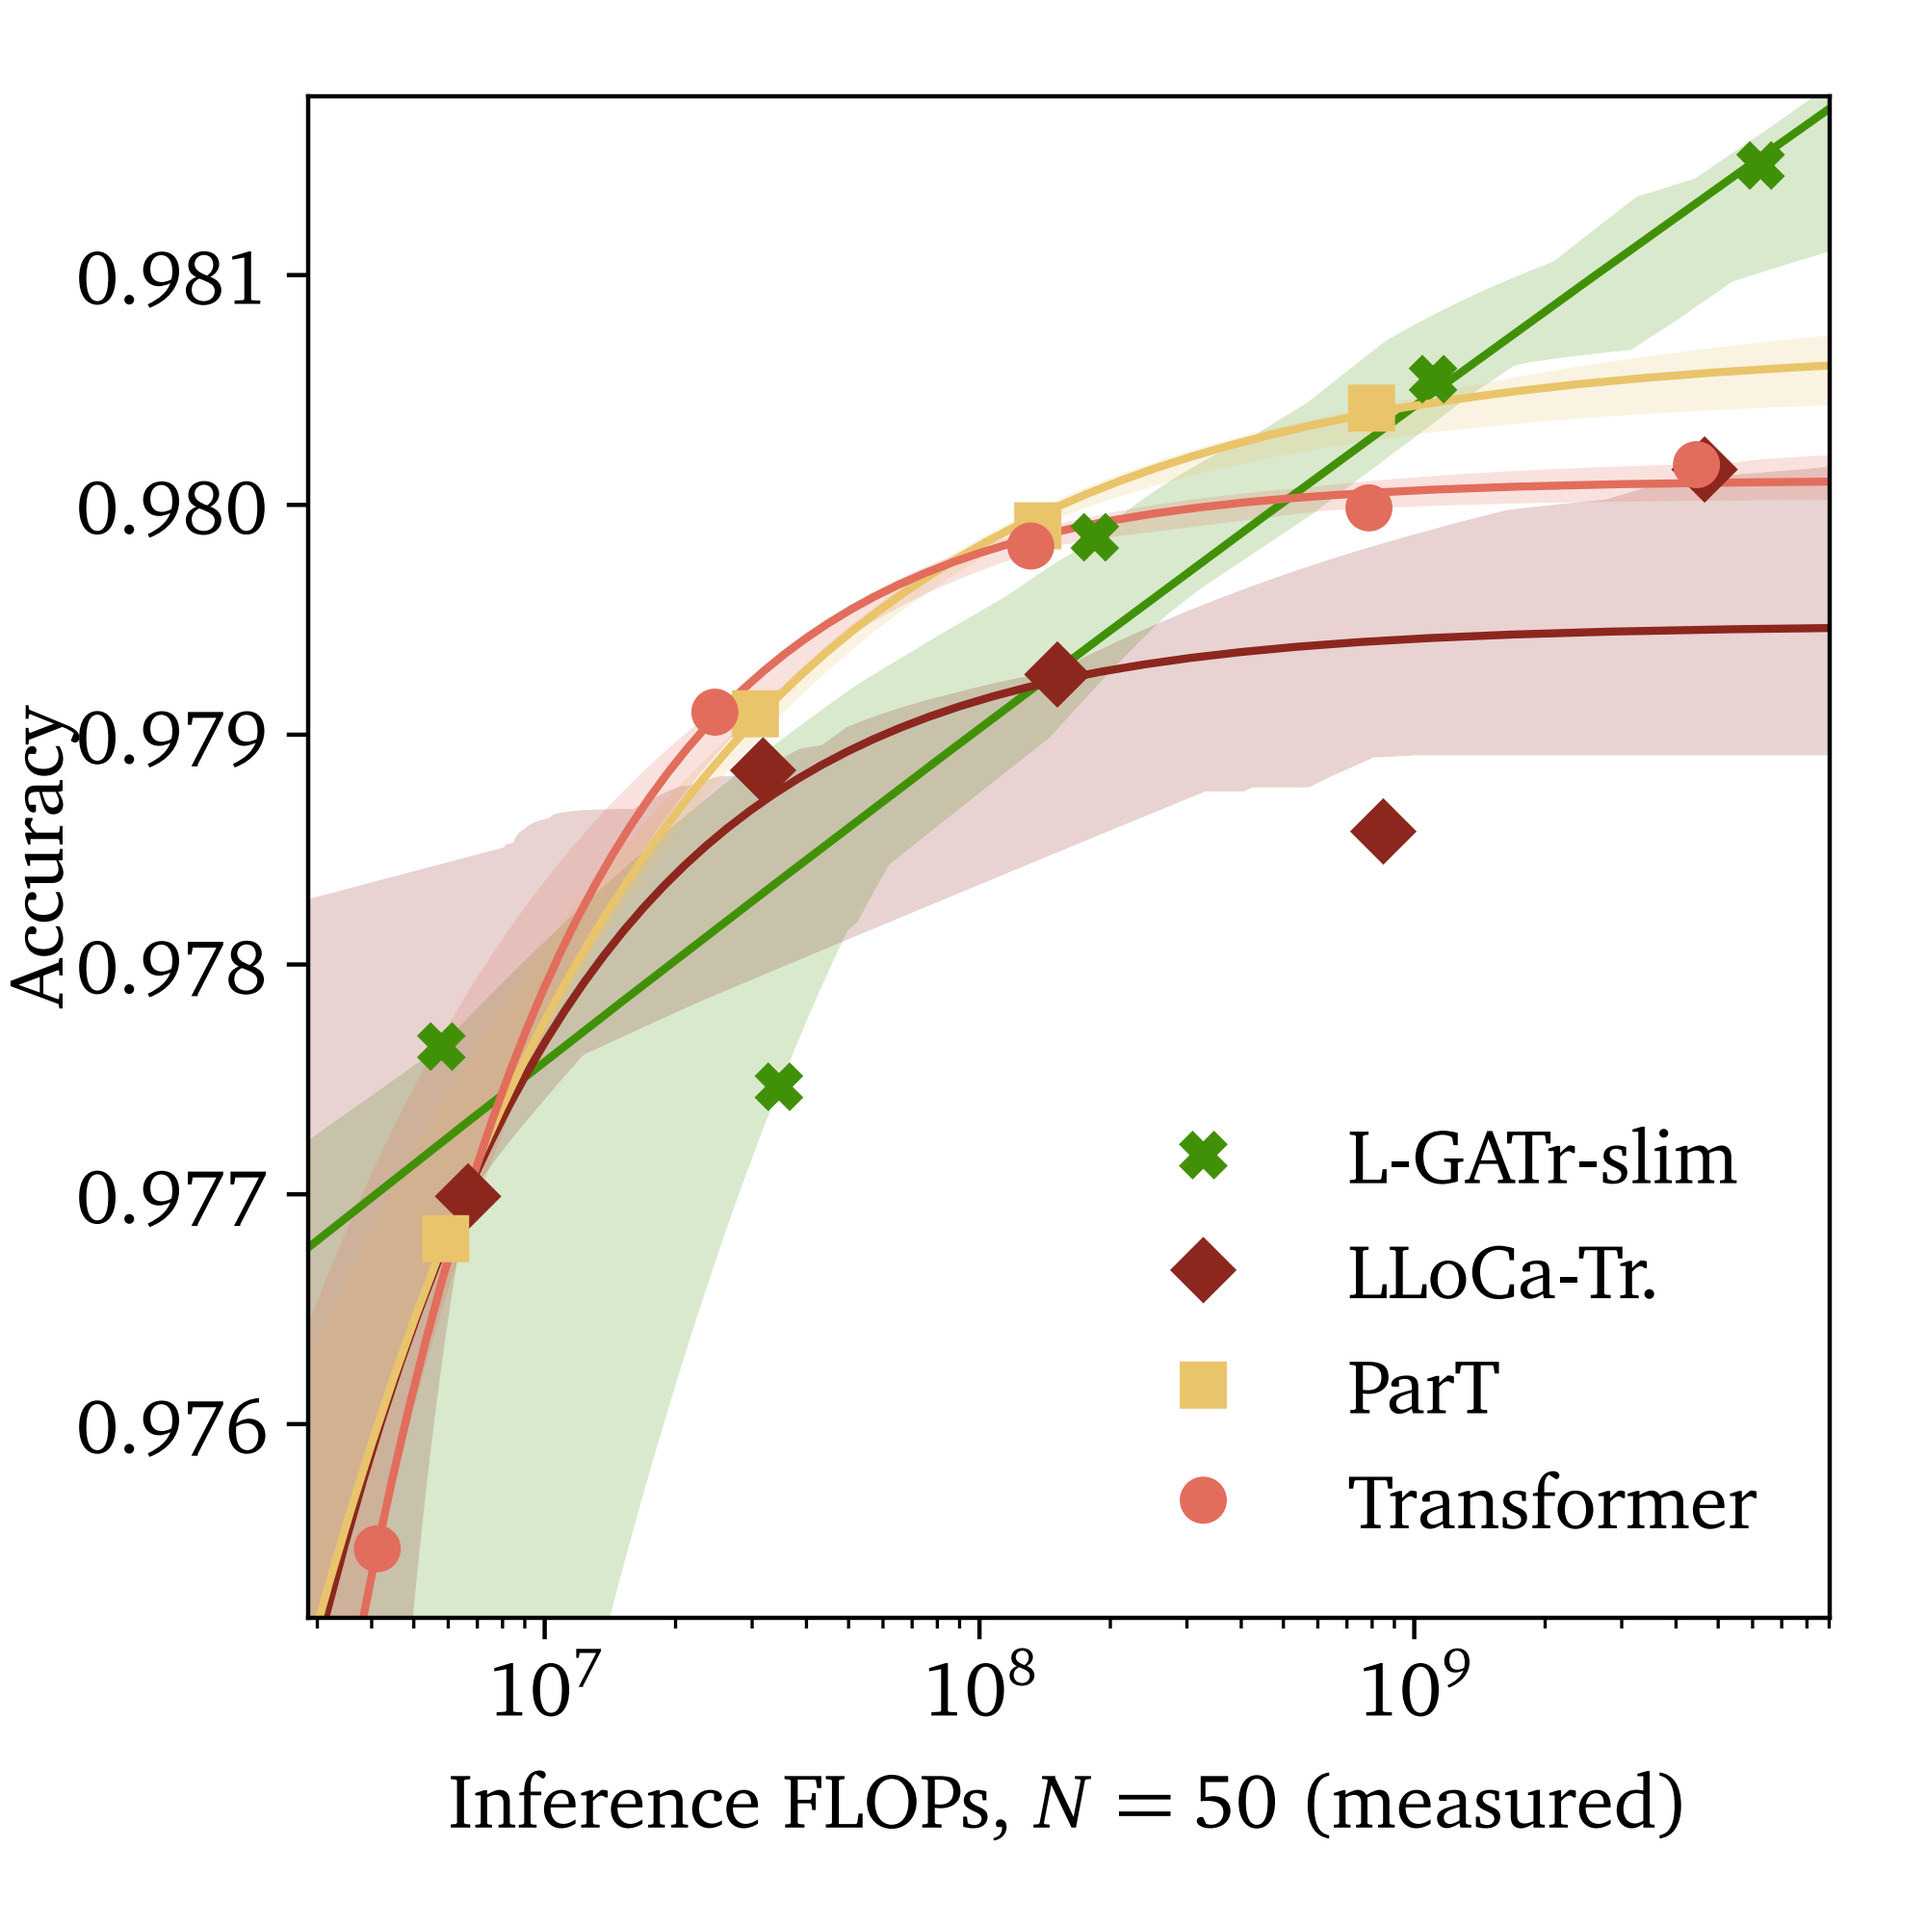
0.981
0.980
0.979
0.978
0.977
0.976
Accuracy
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (measured)
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

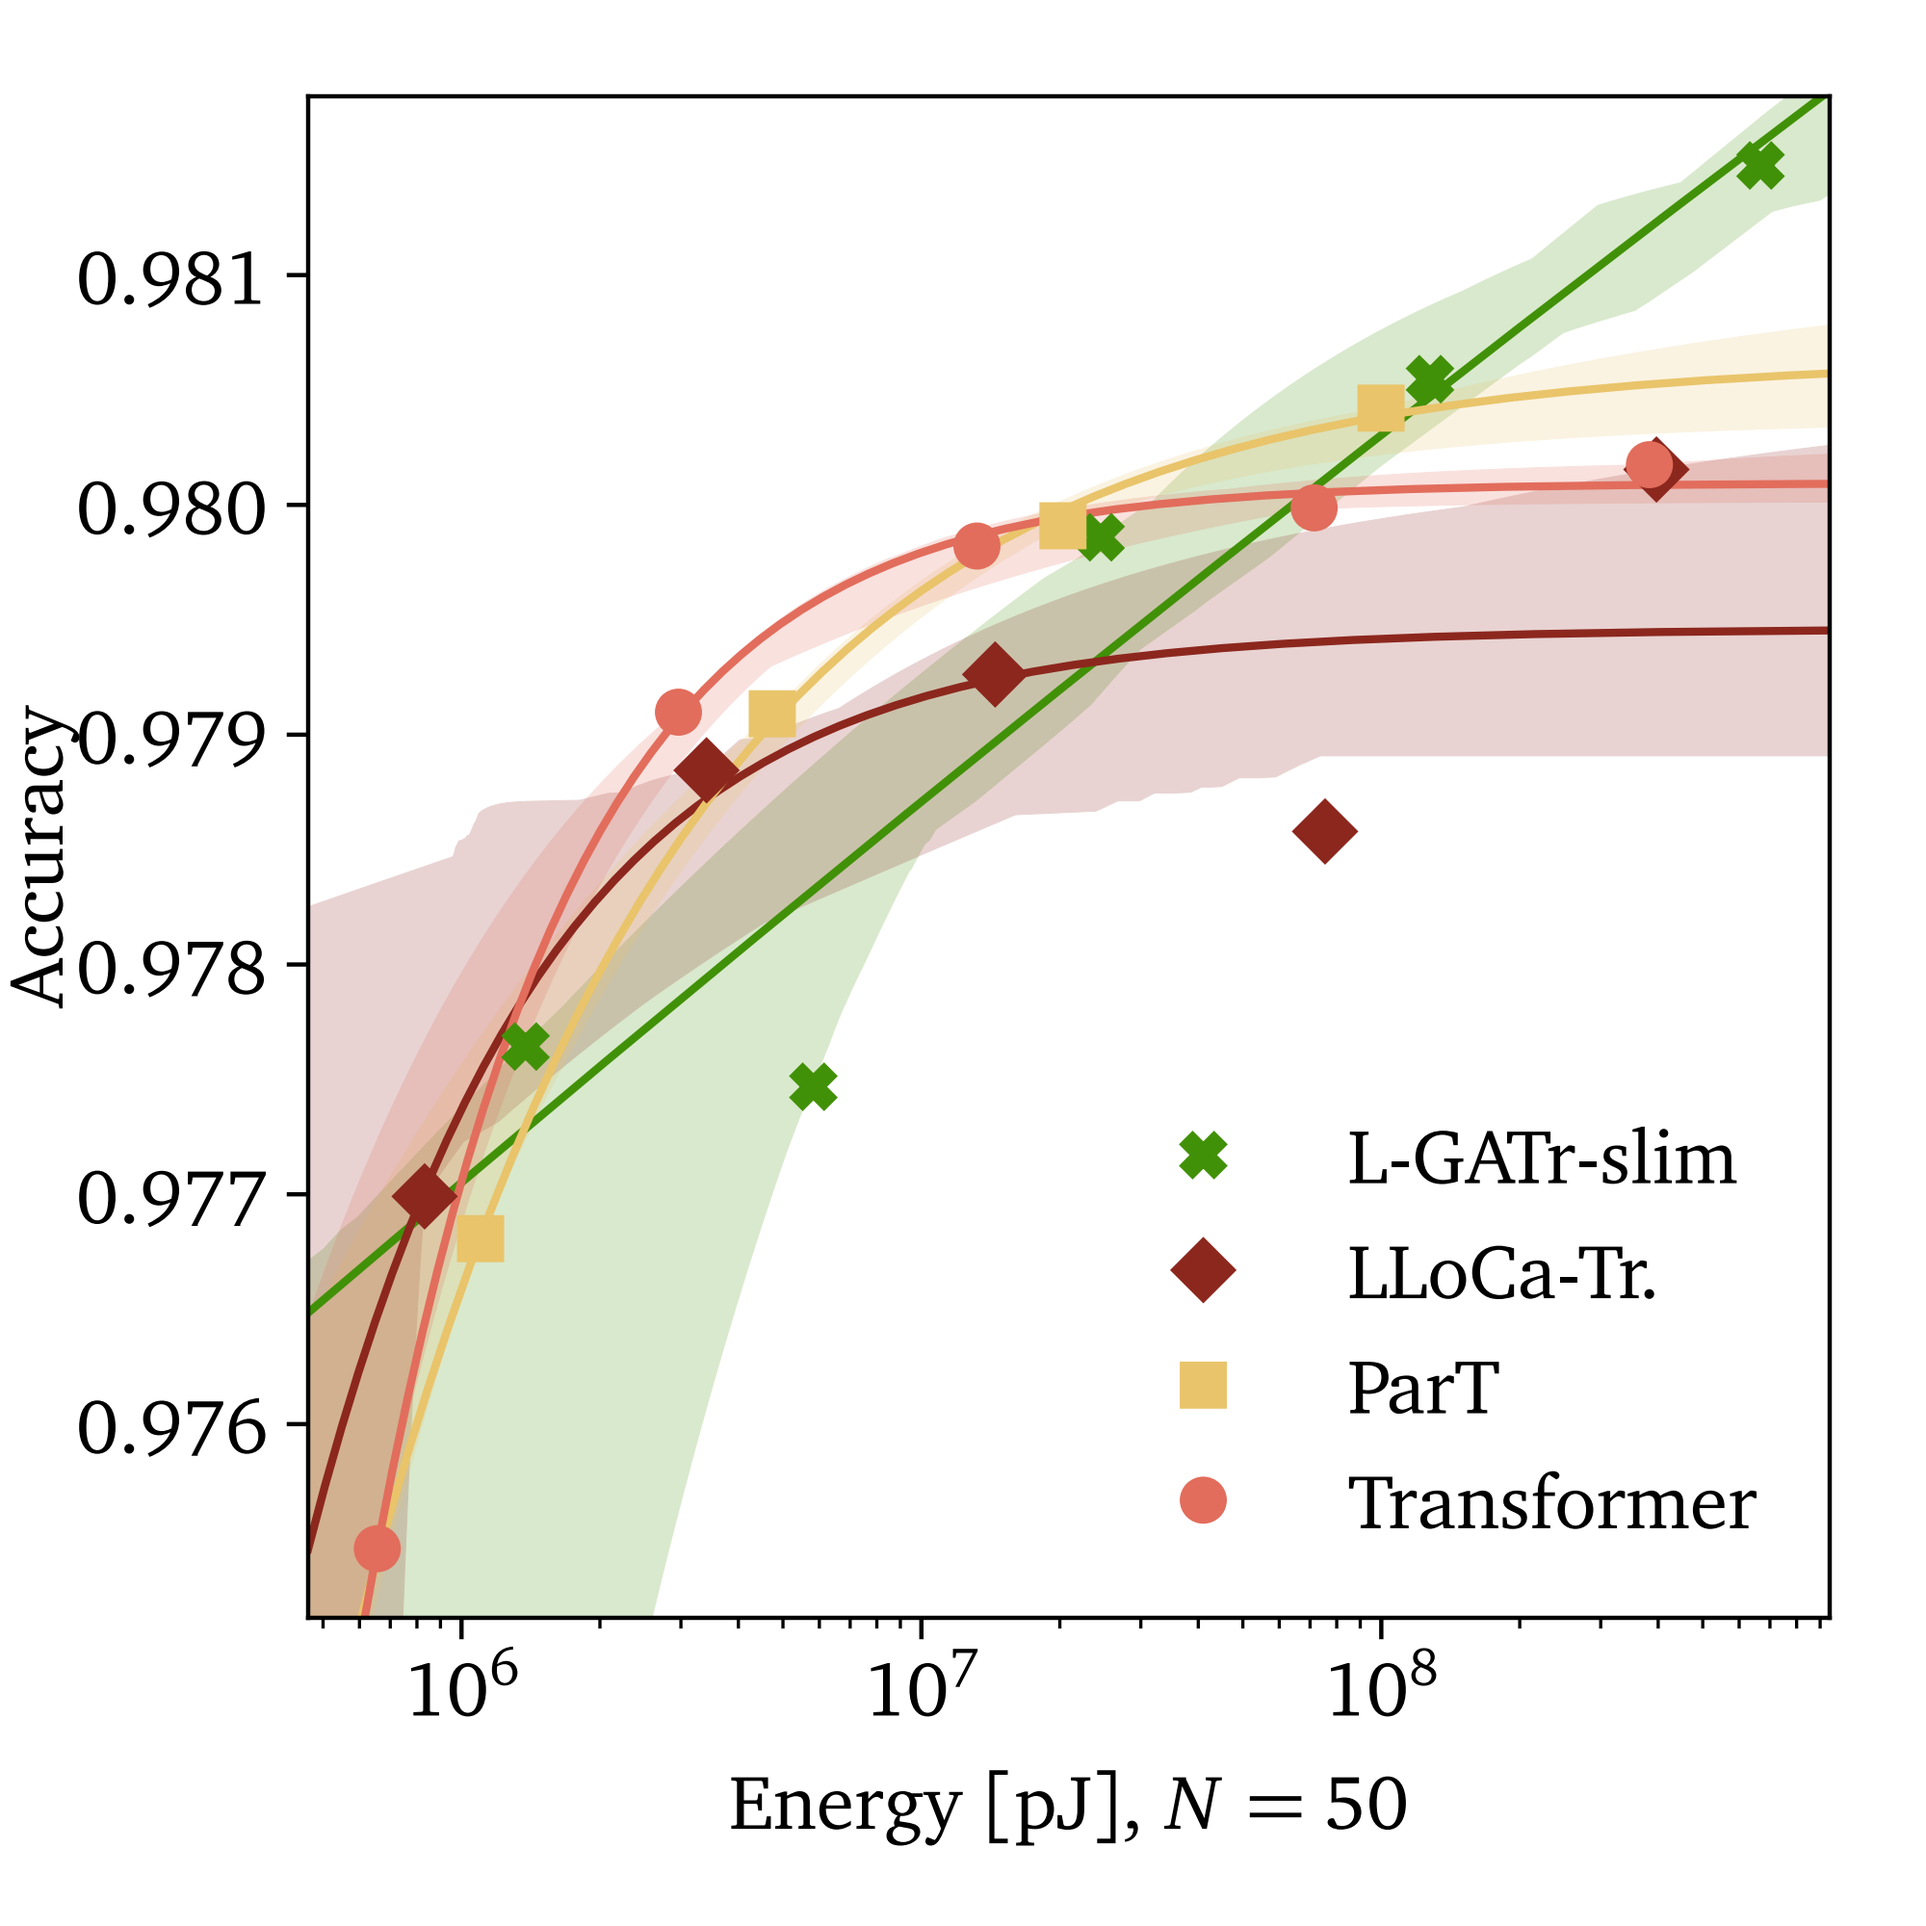
0.981
0.980
0.979
0.978
0.977
0.976
Accuracy
10^6
10^7
10^8
Energy [pJ], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

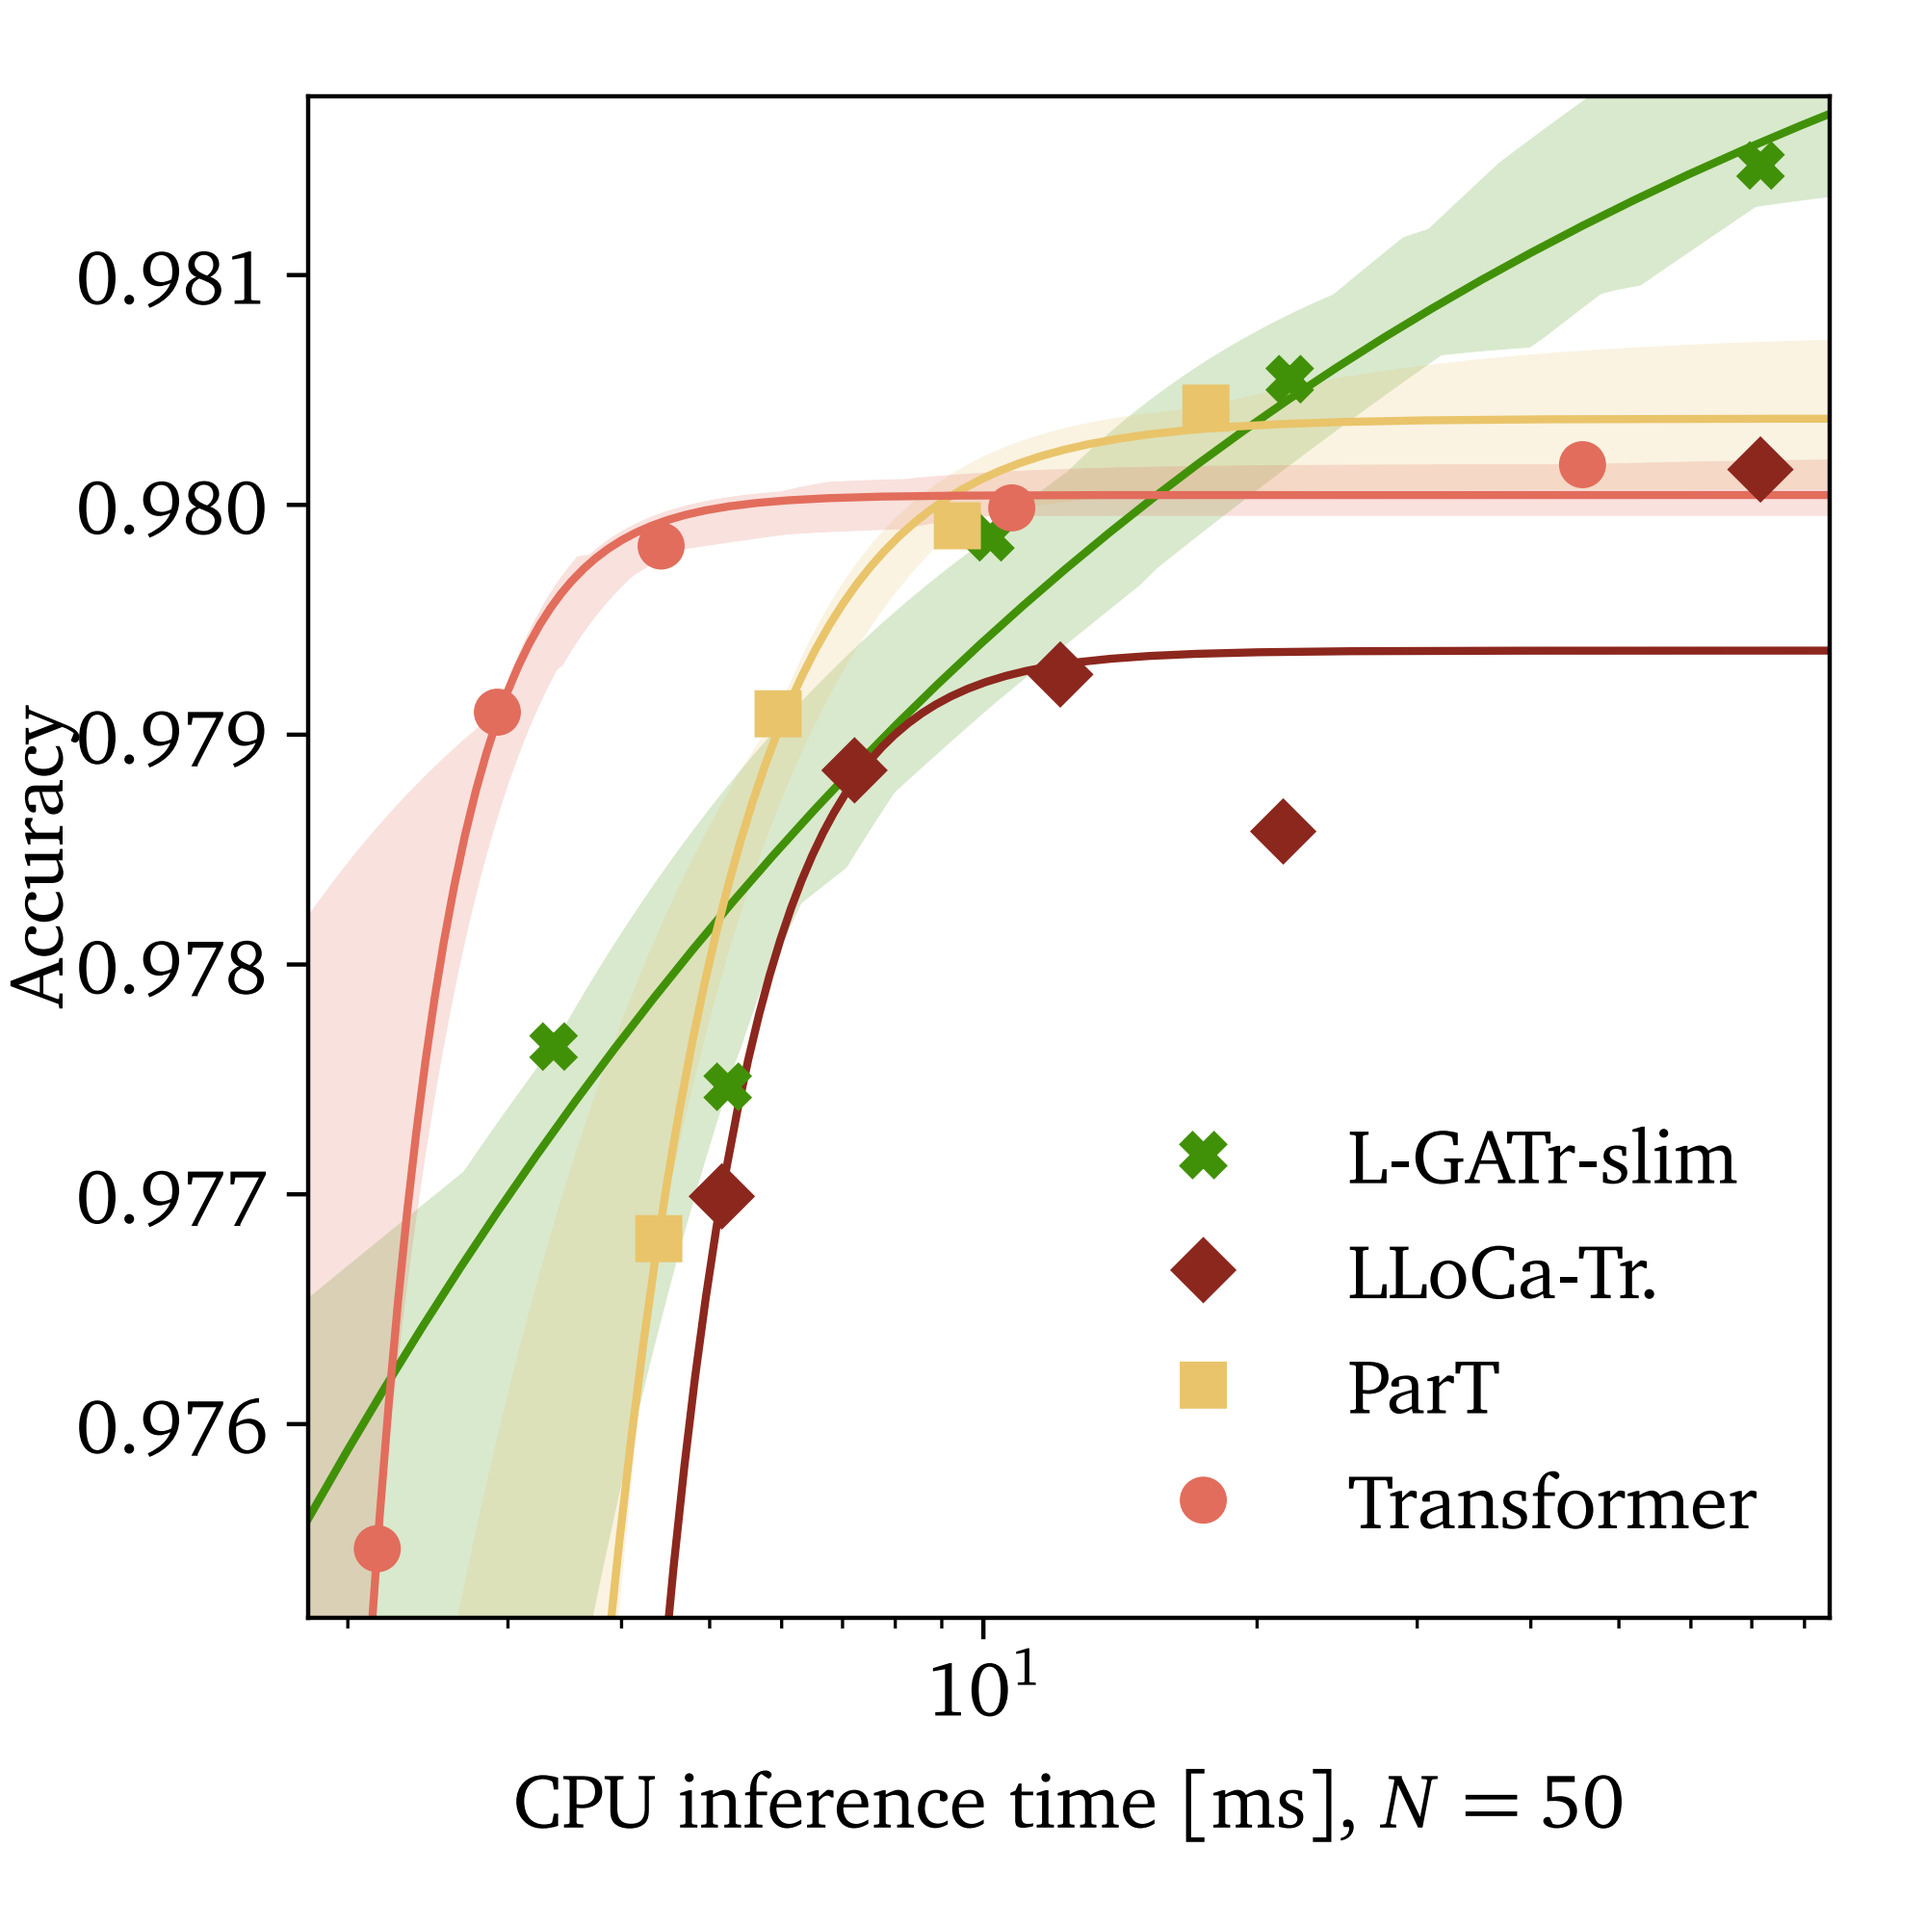
0.981
0.980
0.979
0.978
0.977
0.976
Accuracy
10^1
CPU inference time [ms], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

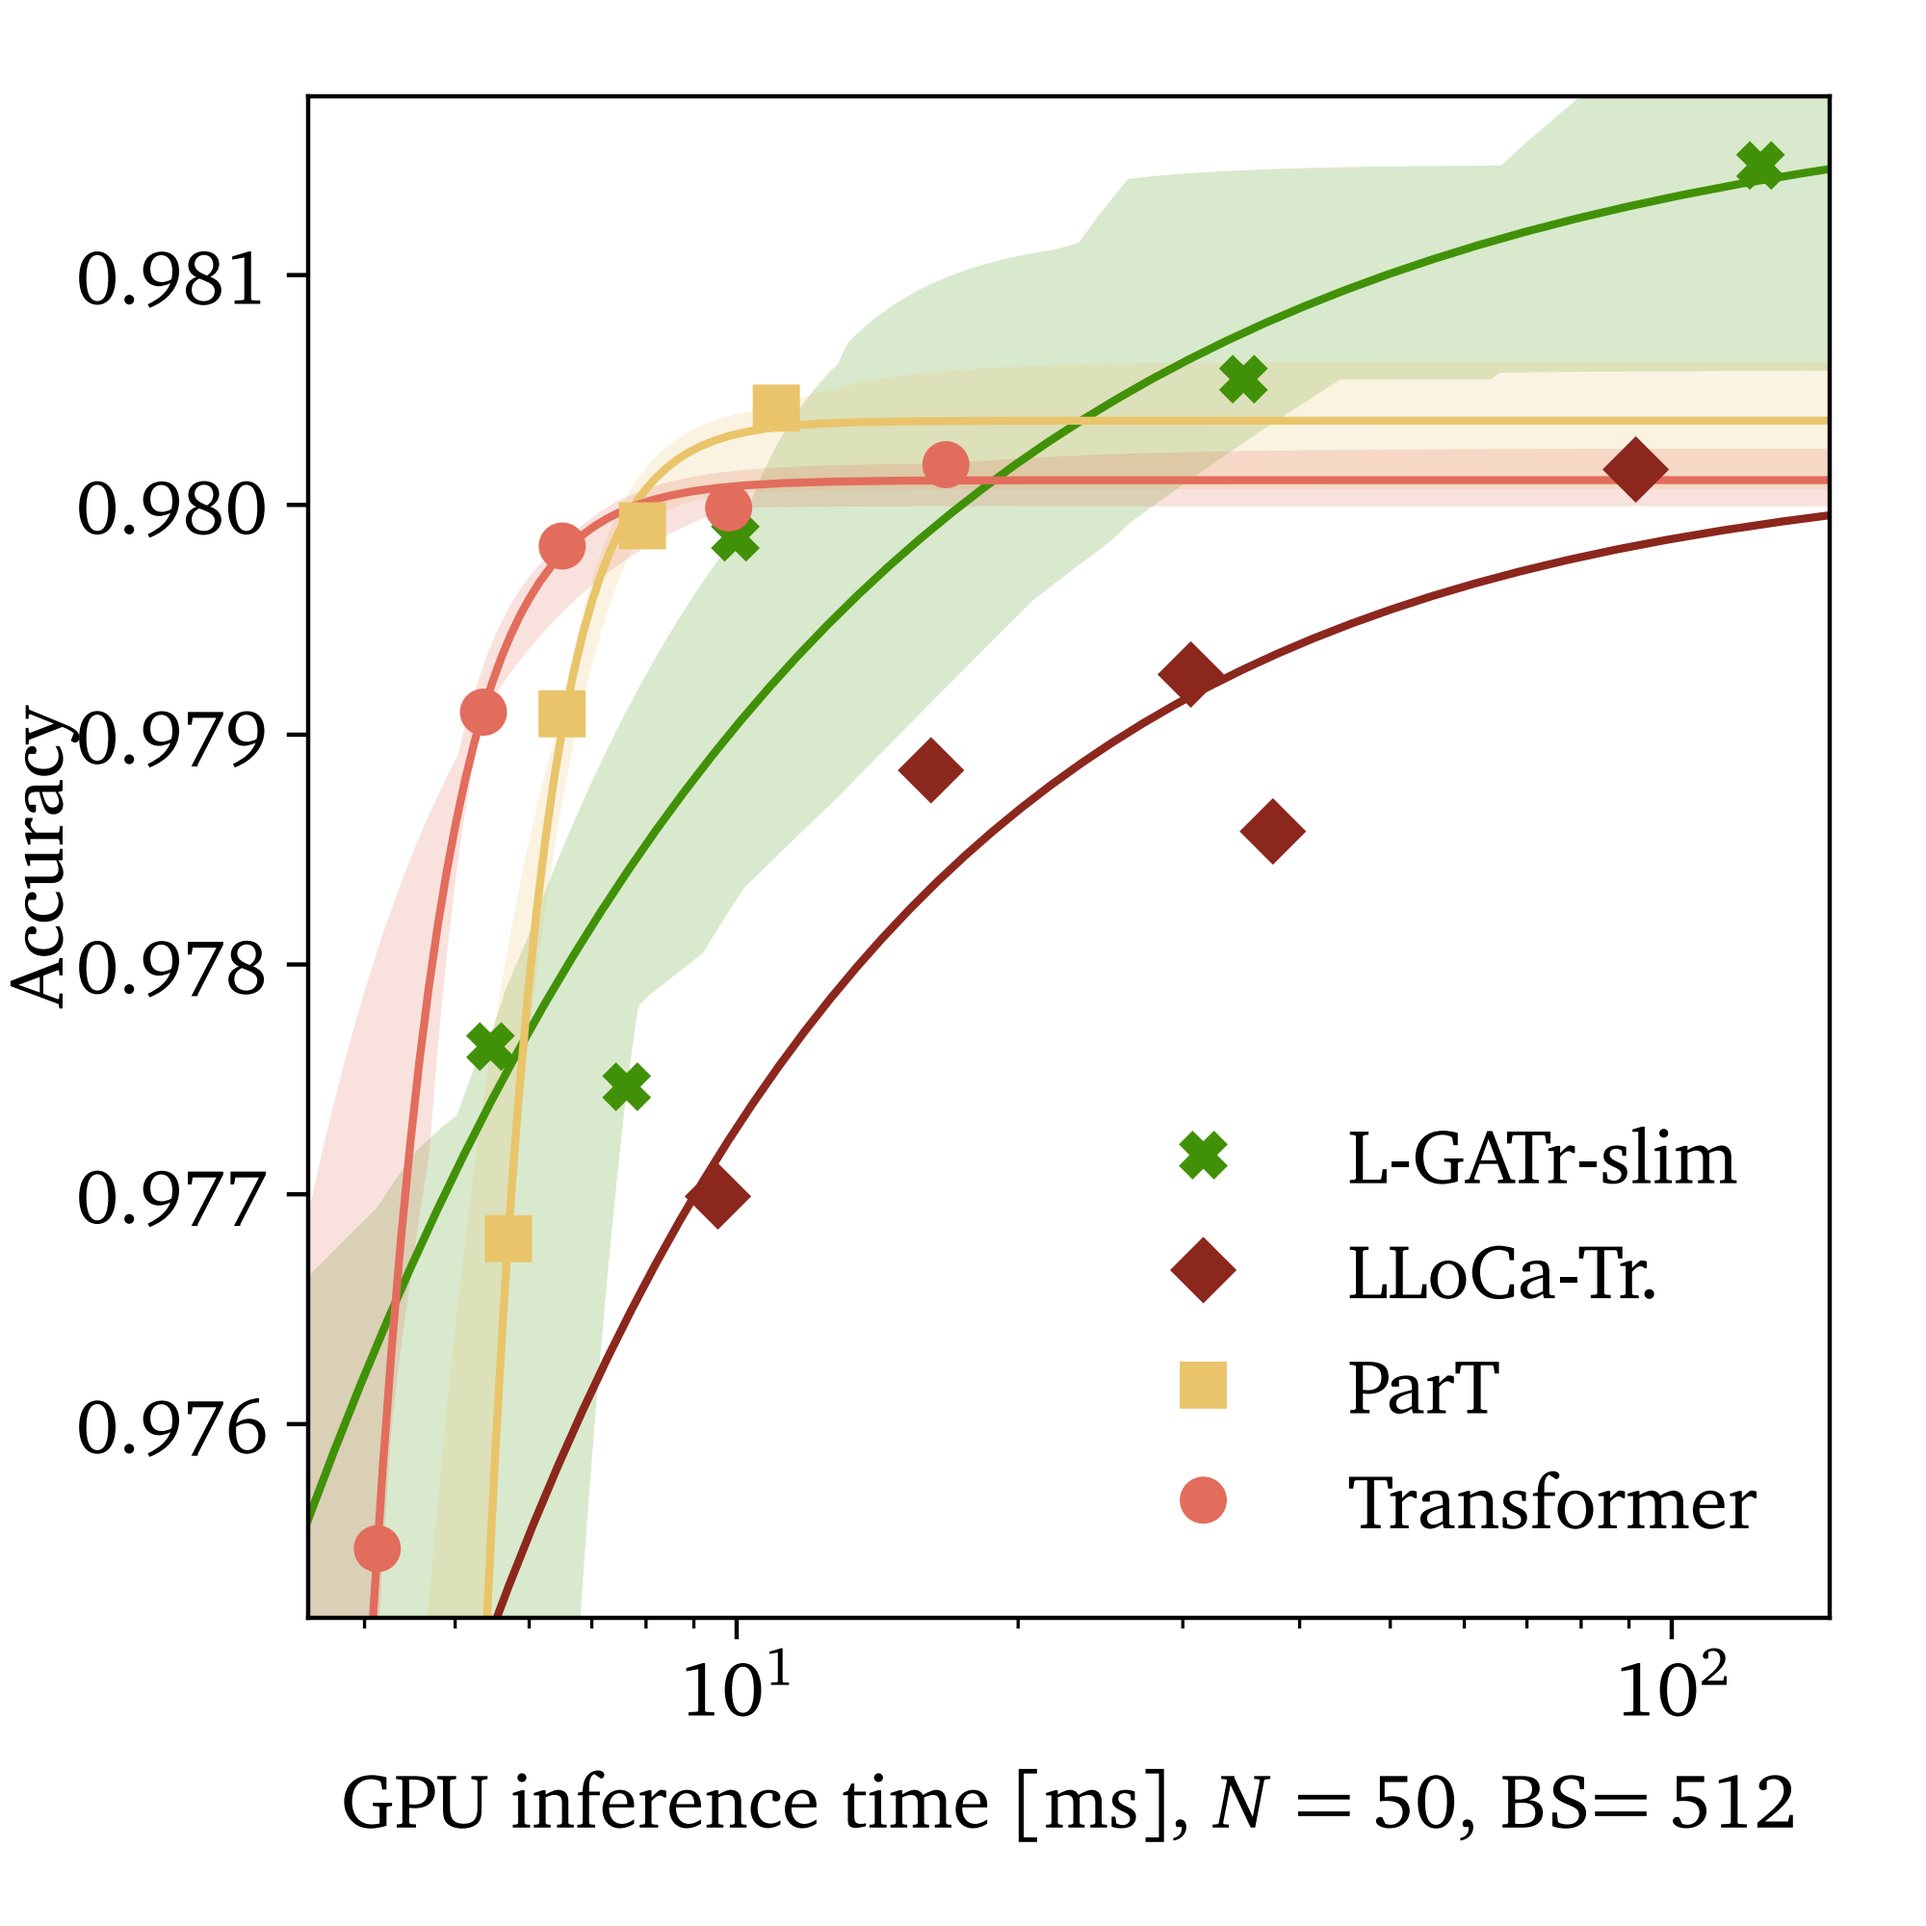
0.981
0.980
0.979
0.978
0.977
0.976
Accuracy
10^1
10^2
GPU inference time [ms], $N = 50$, BS= 512
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

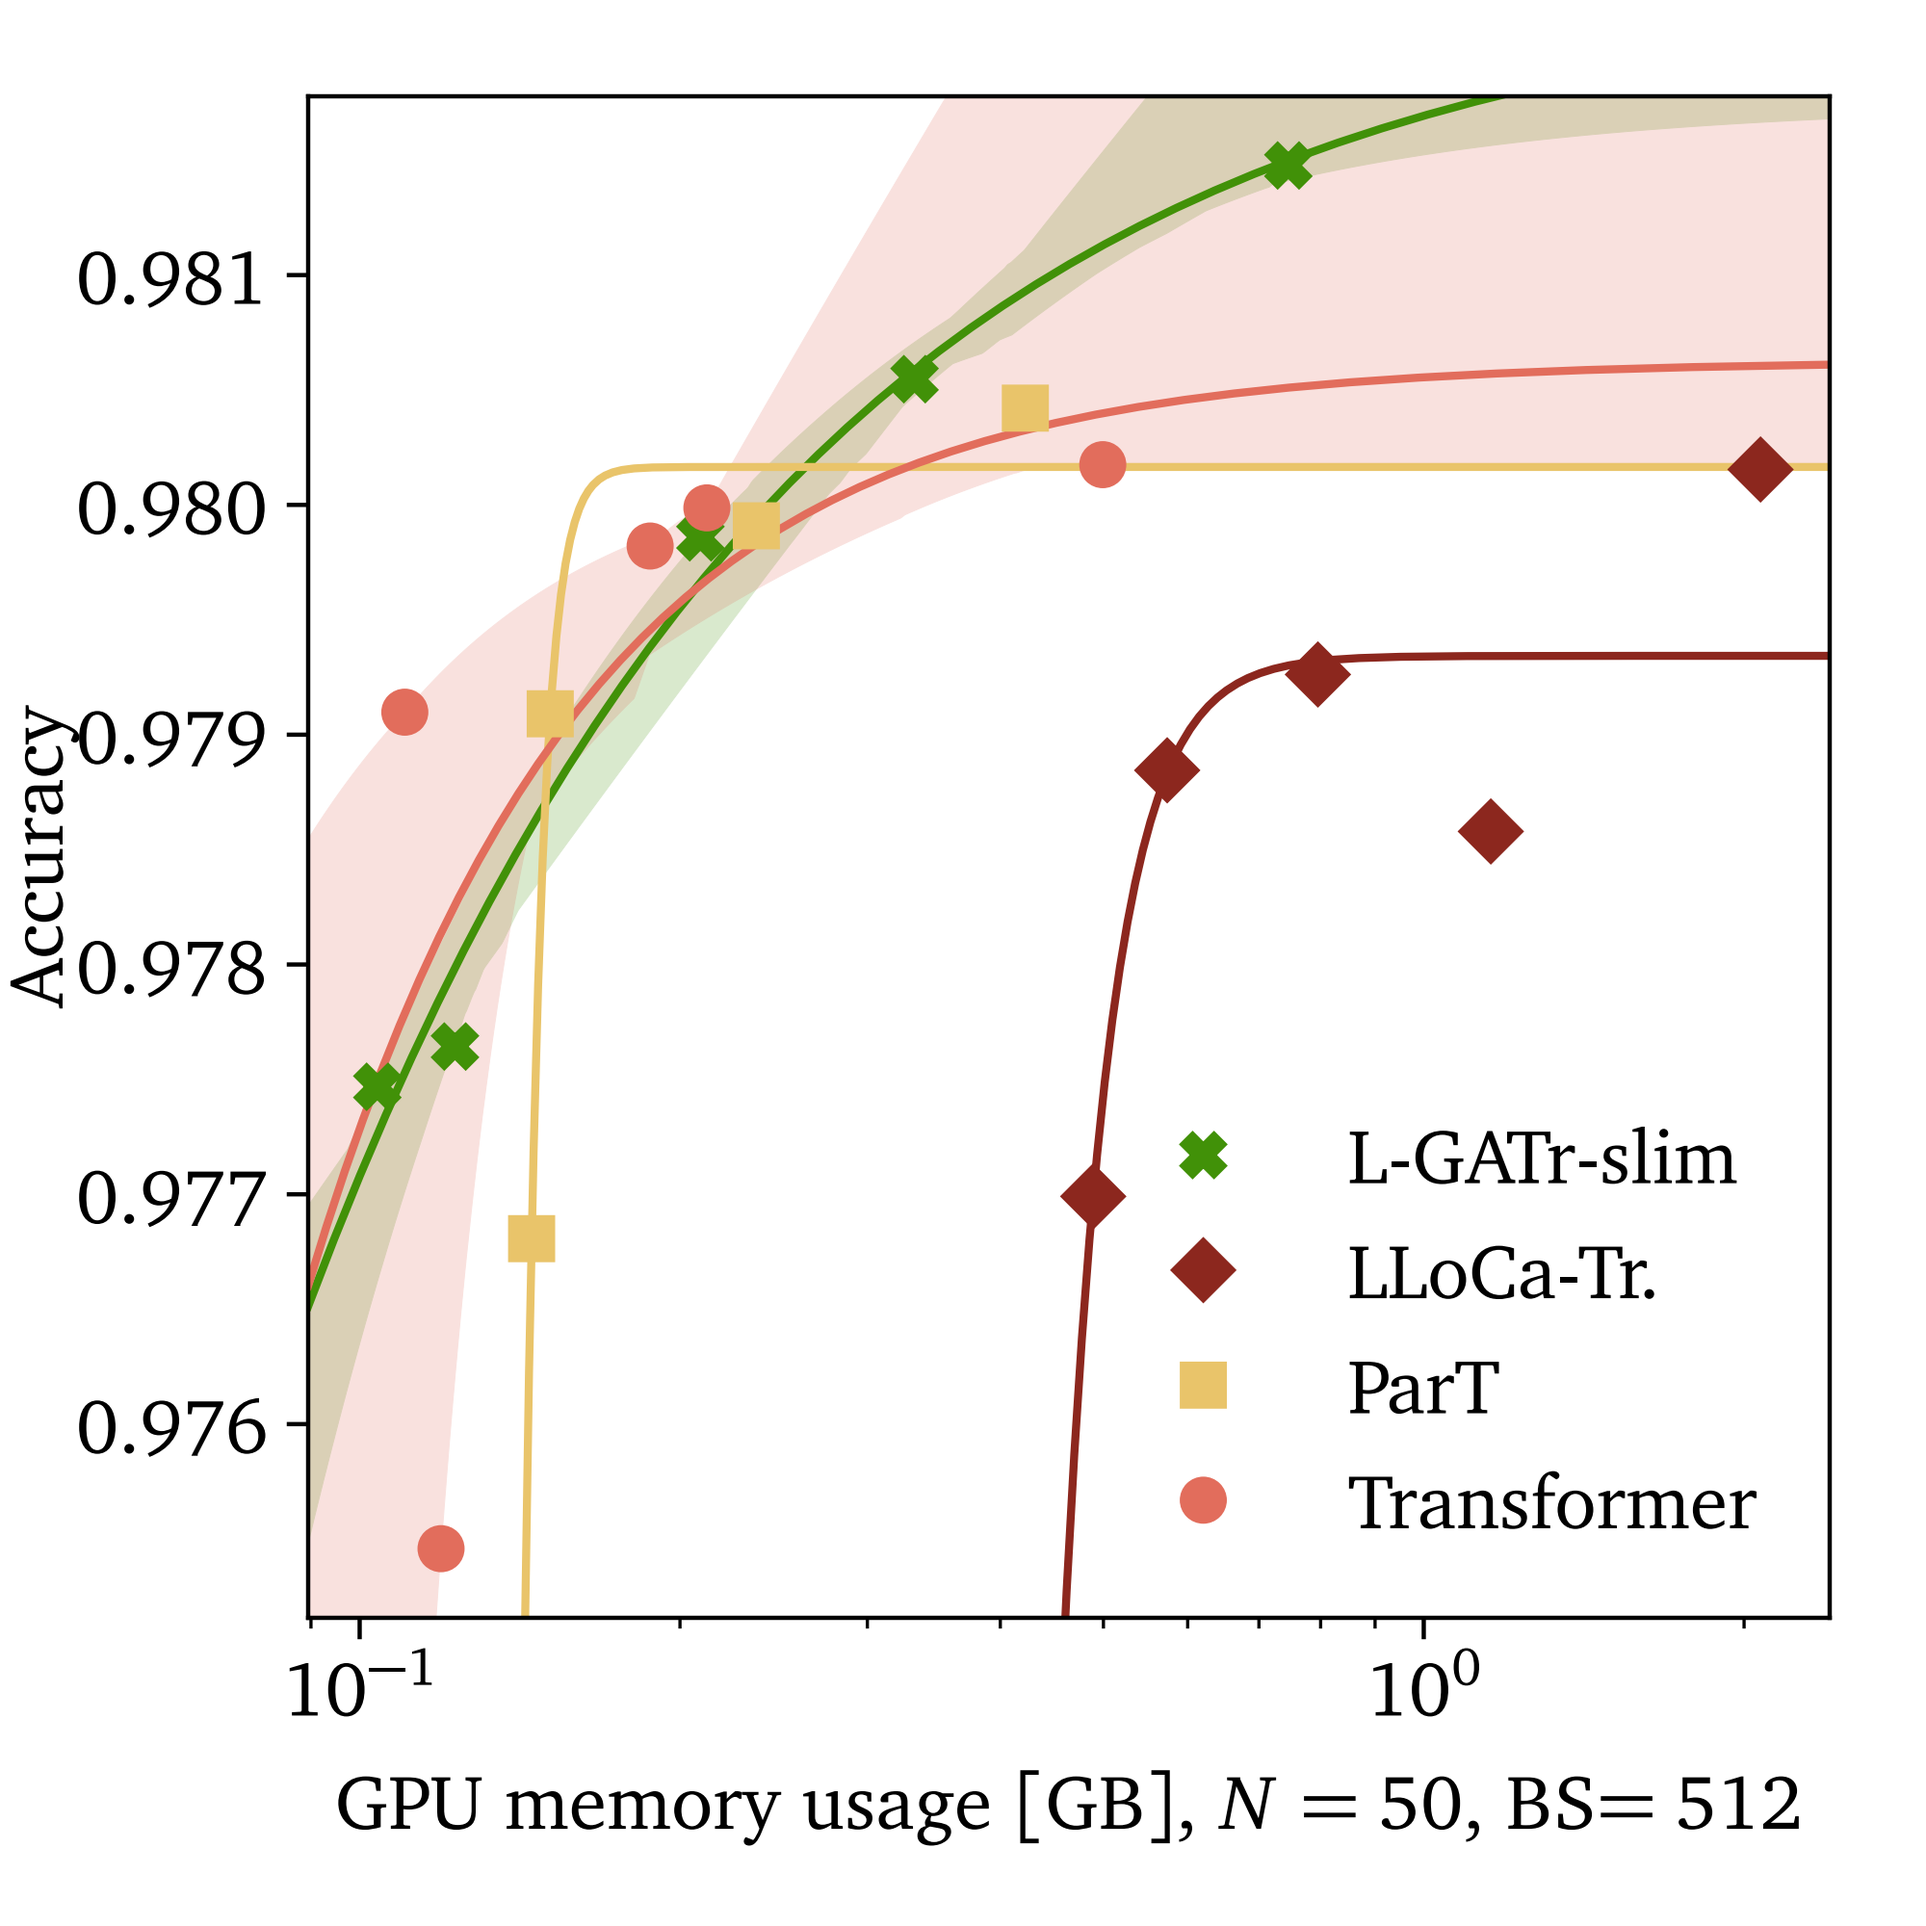
0.981
0.980
0.979
0.978
0.977
0.976
Accuracy
10^{-1}
10^{0}
GPU memory usage [GB], $N = 50$, BS$= 512$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

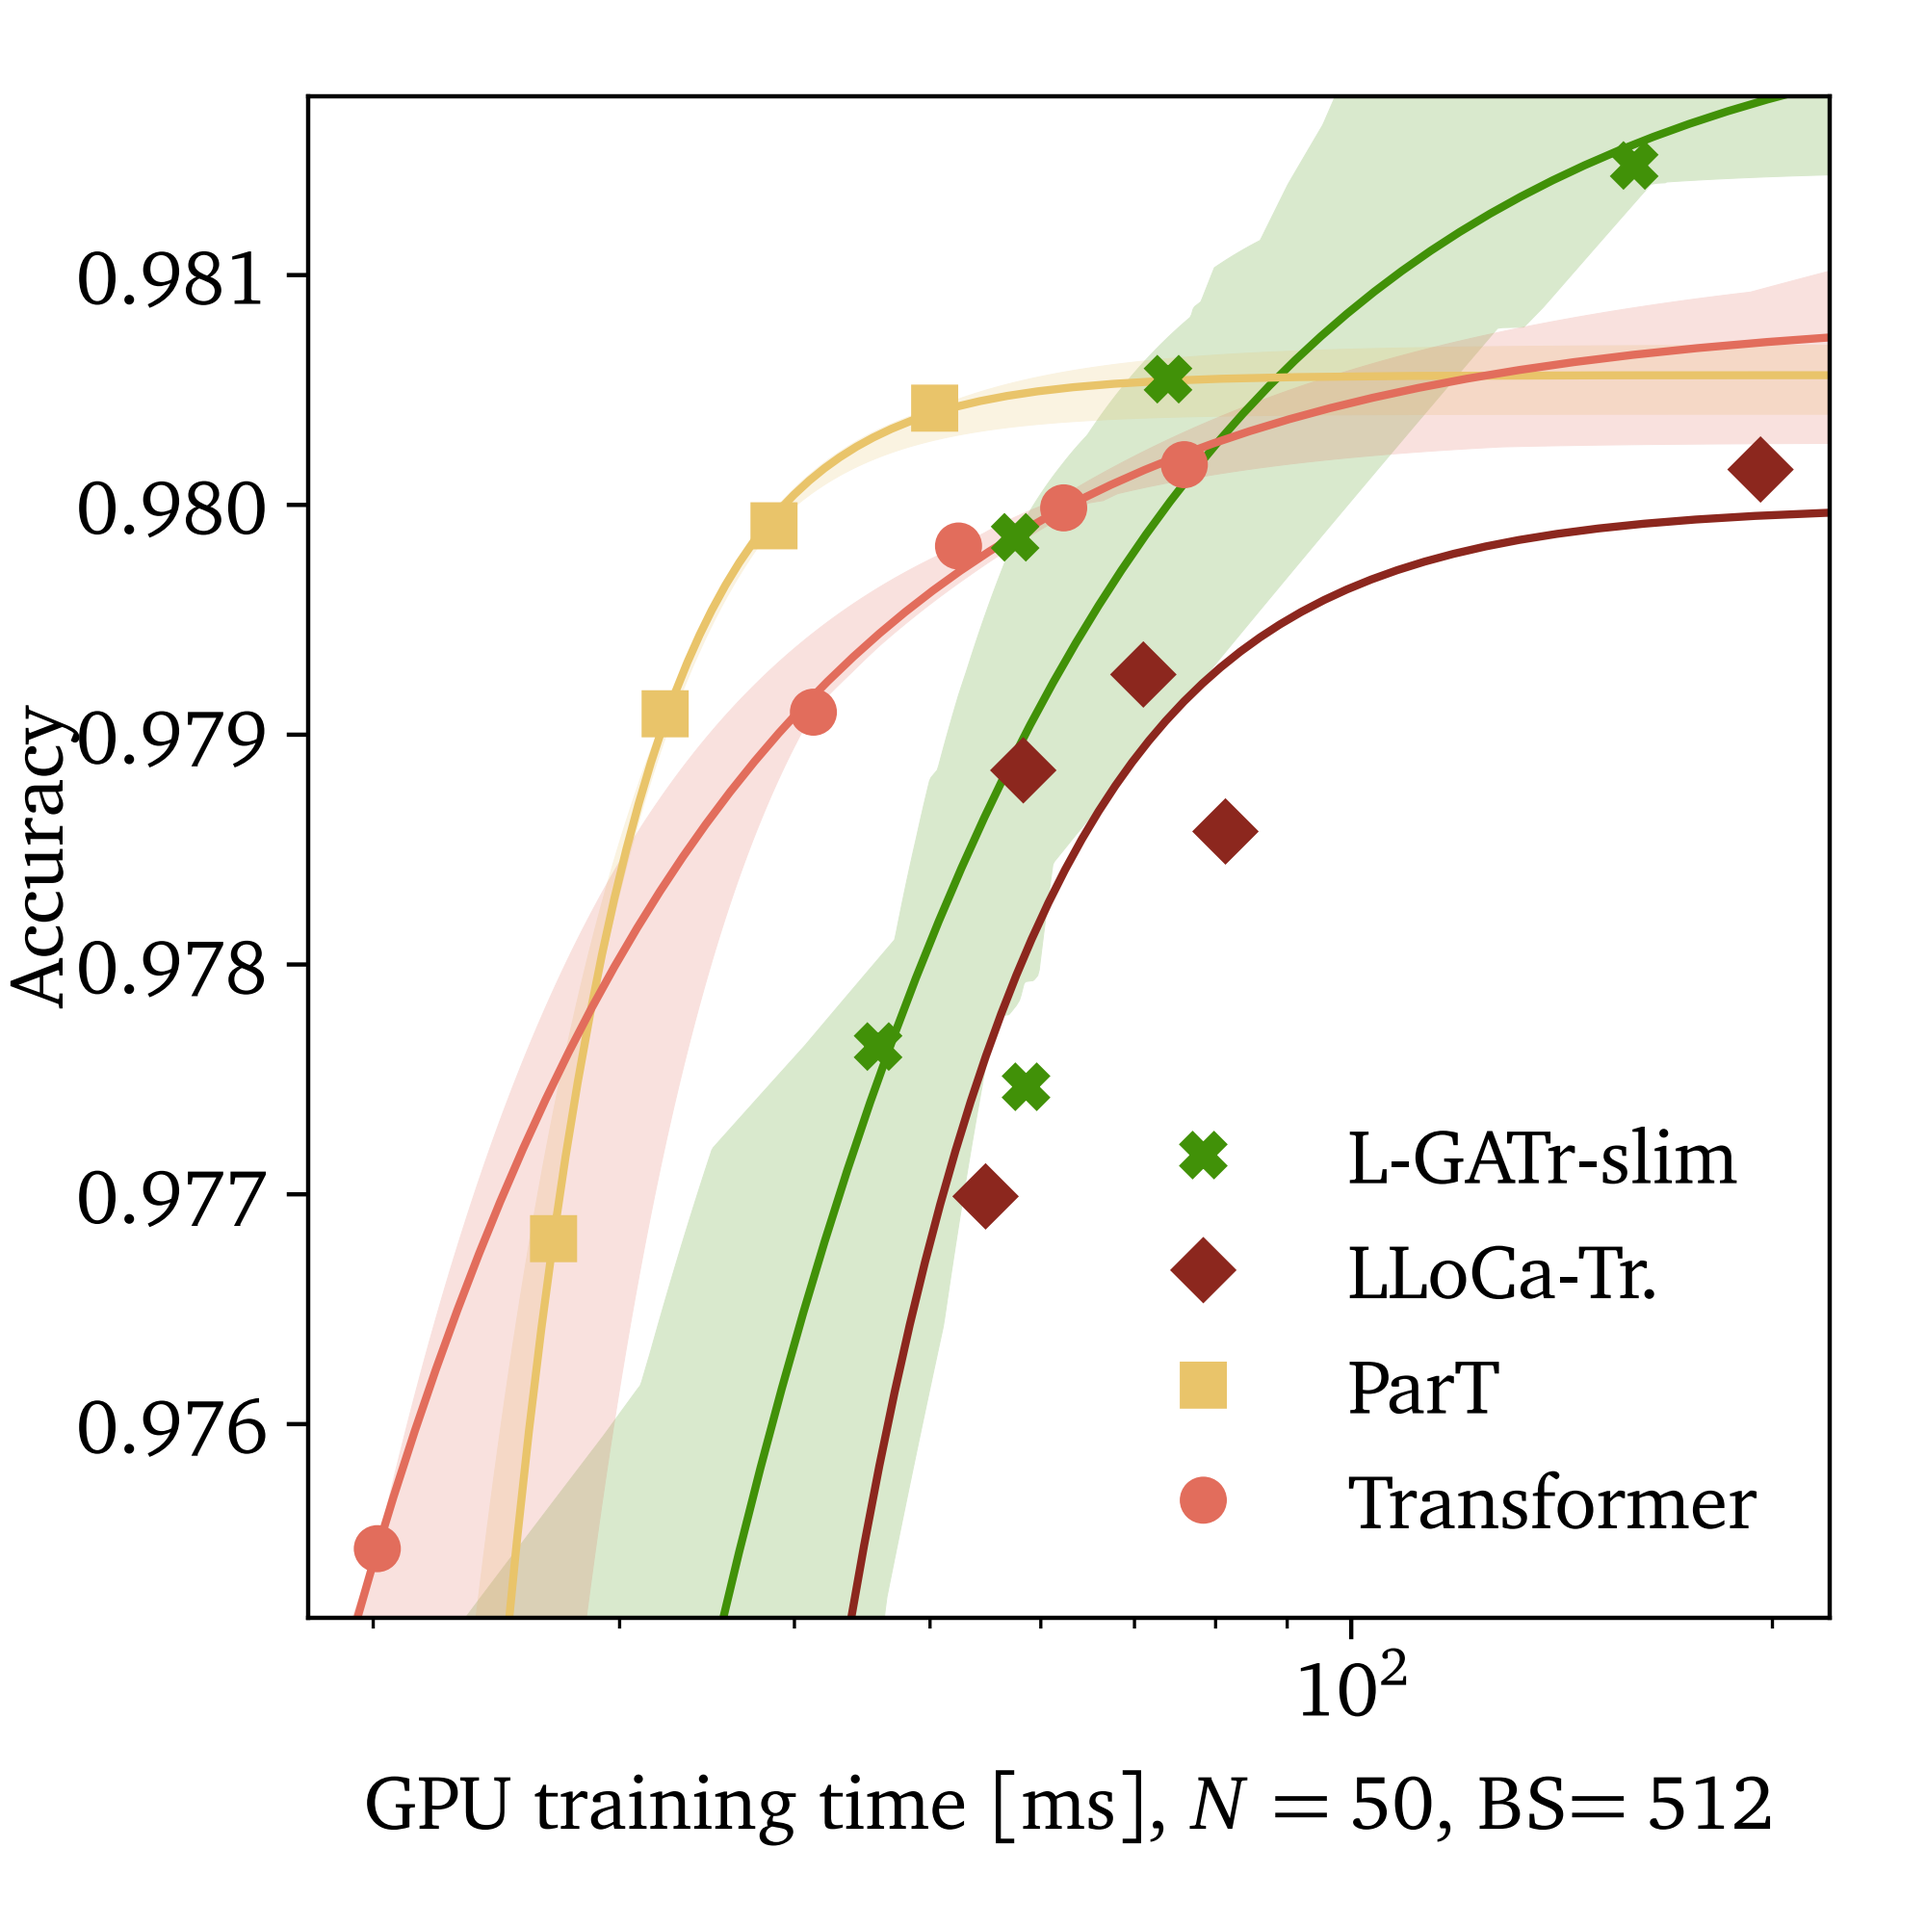
0.981
0.980
0.979
0.978
0.977
0.976
Accuracy
10^2
GPU training time [ms], $N = 50$, BS$= 512$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer