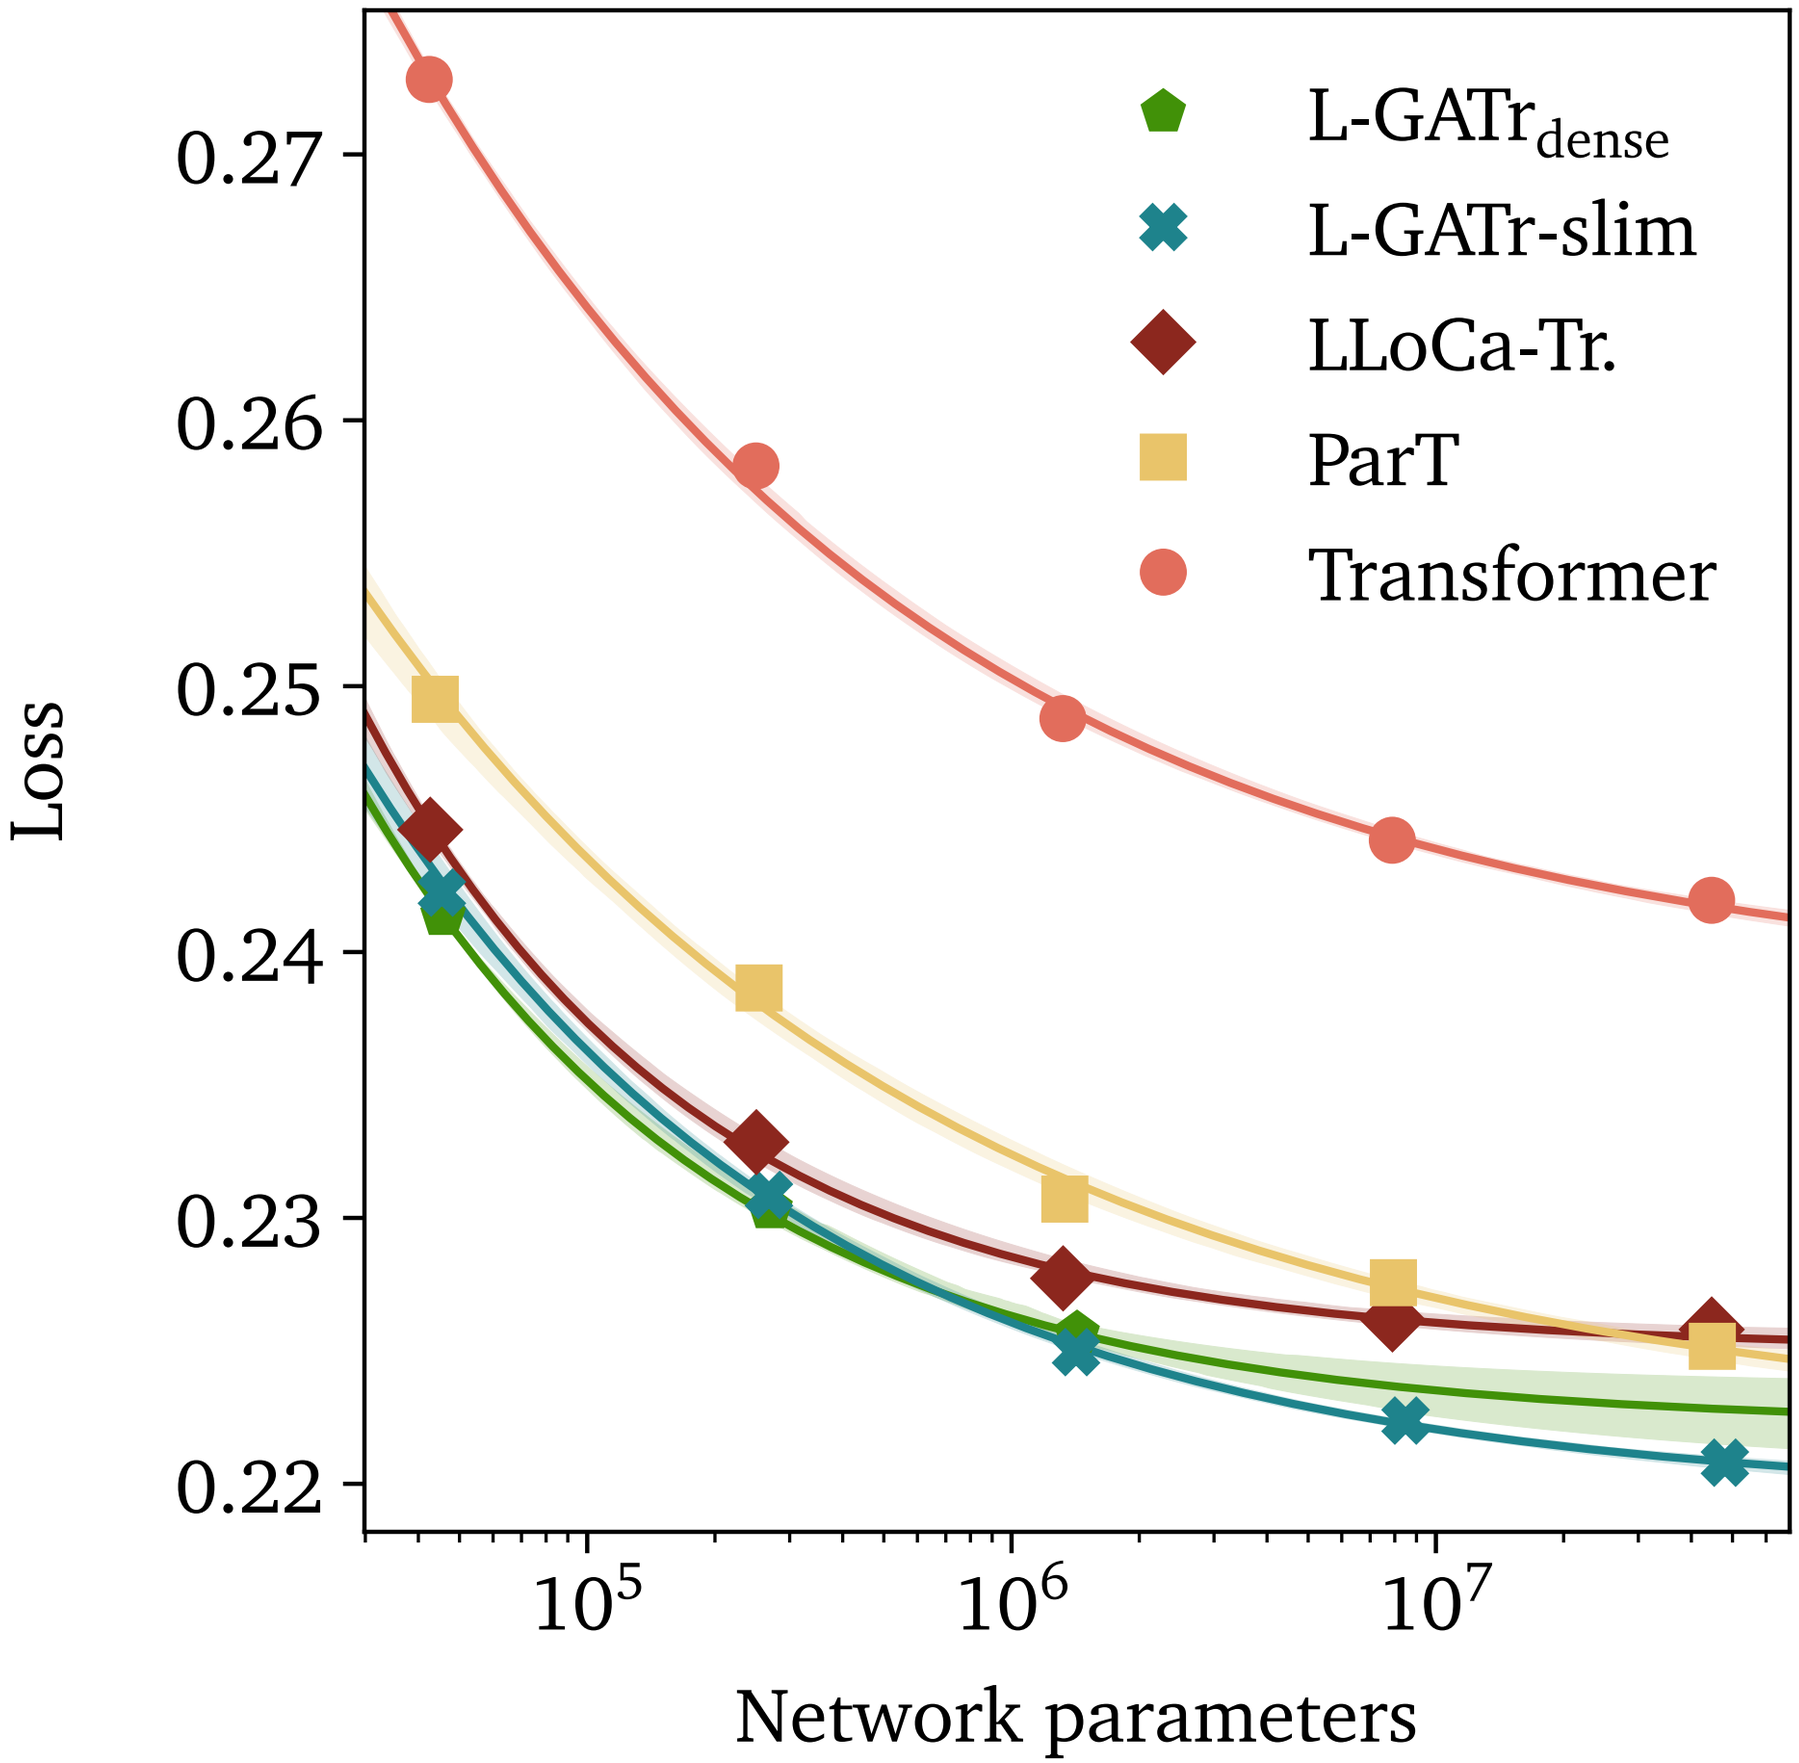
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.27
0.26
0.25
0.24
0.23
0.22
10^5
10^6
10^7
Network parameters

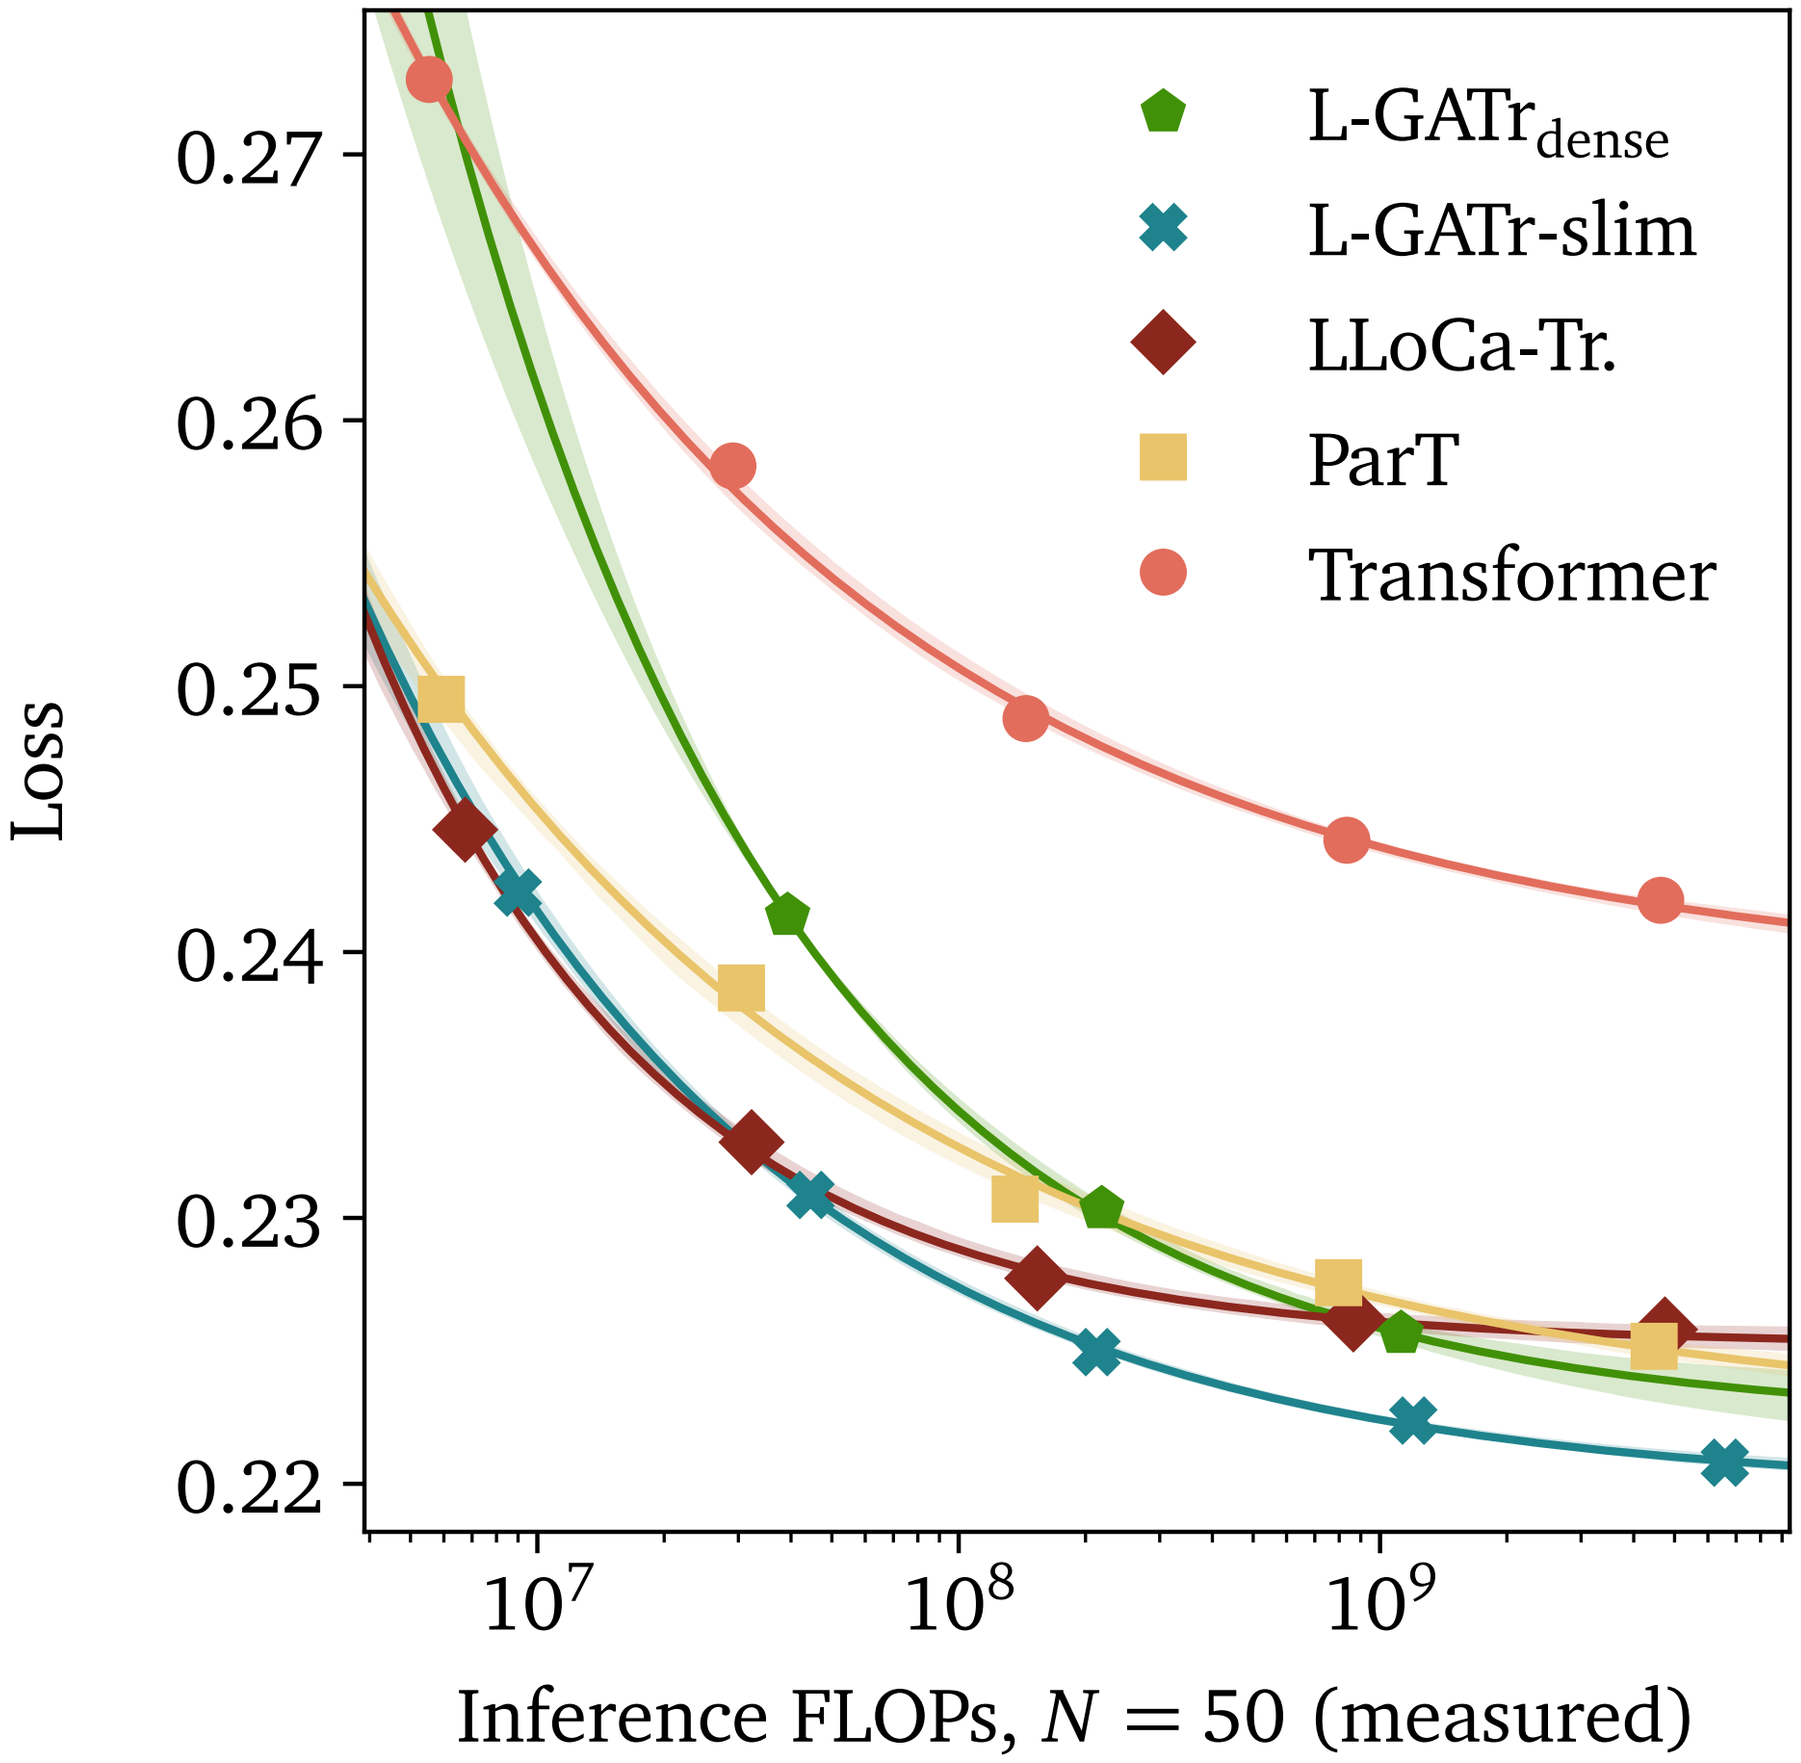
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.27
0.26
0.25
0.24
0.23
0.22
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (measured)

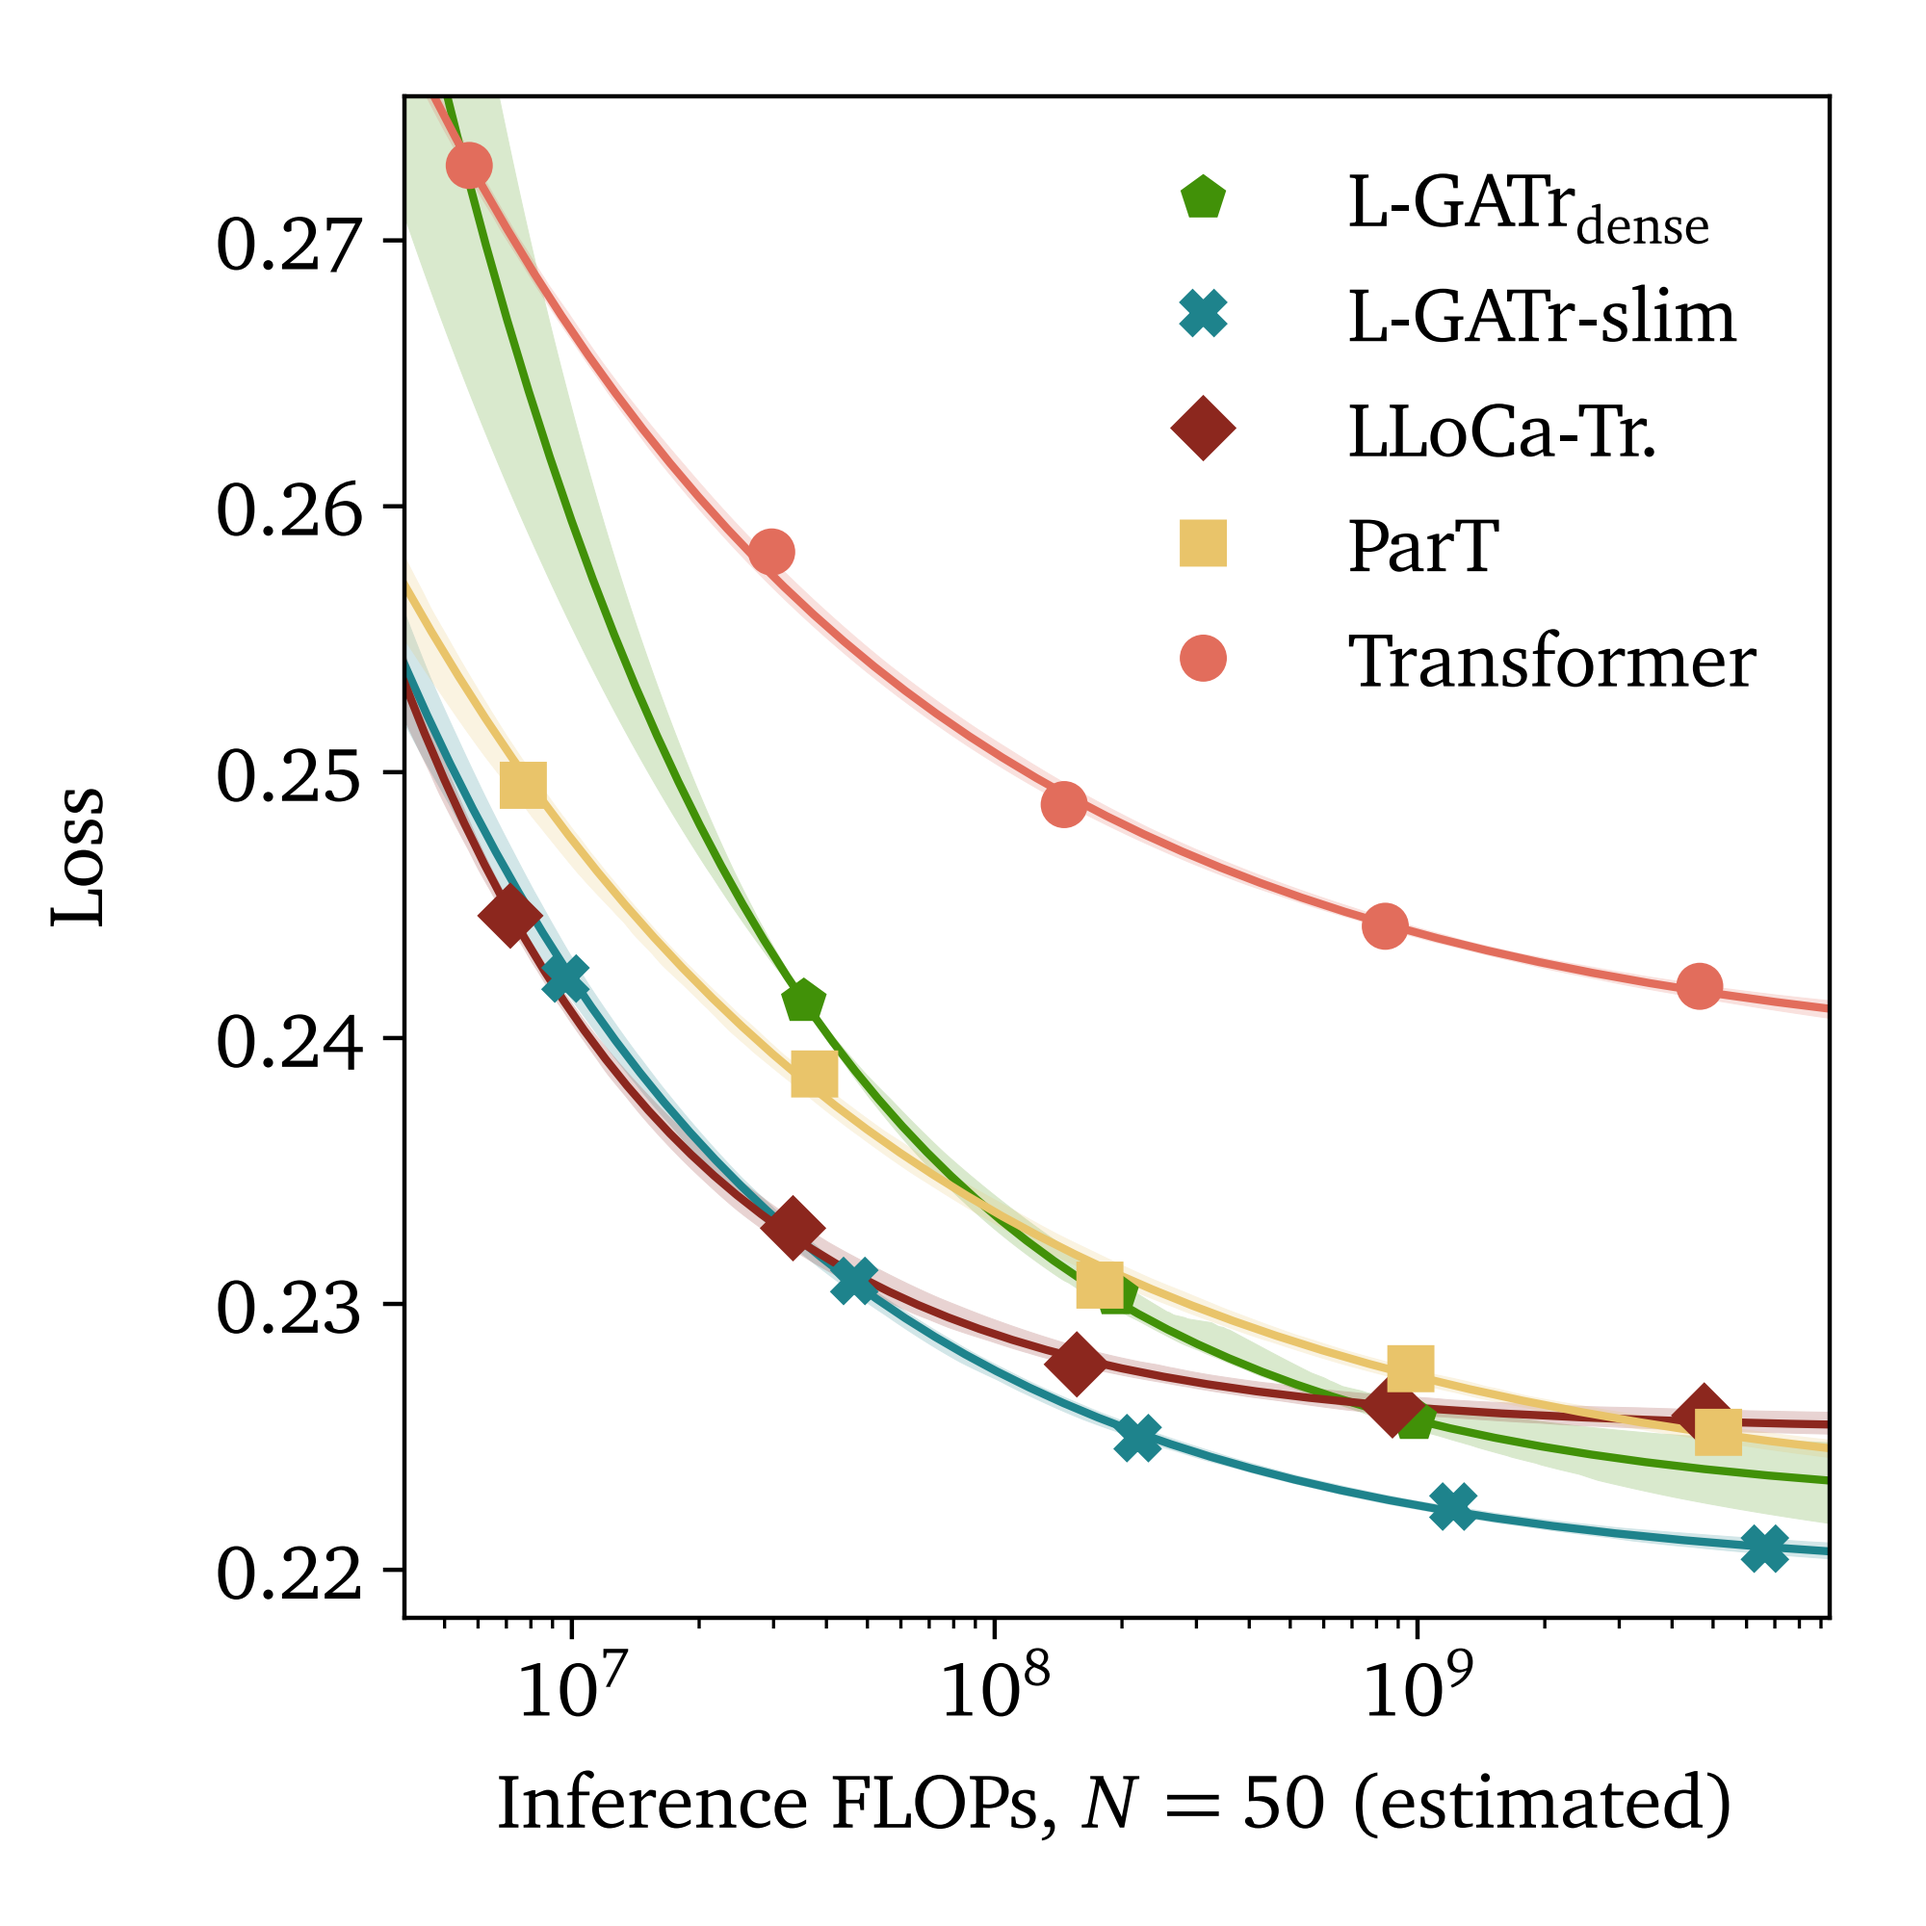
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.27
0.26
0.25
0.24
0.23
0.22
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (estimated)

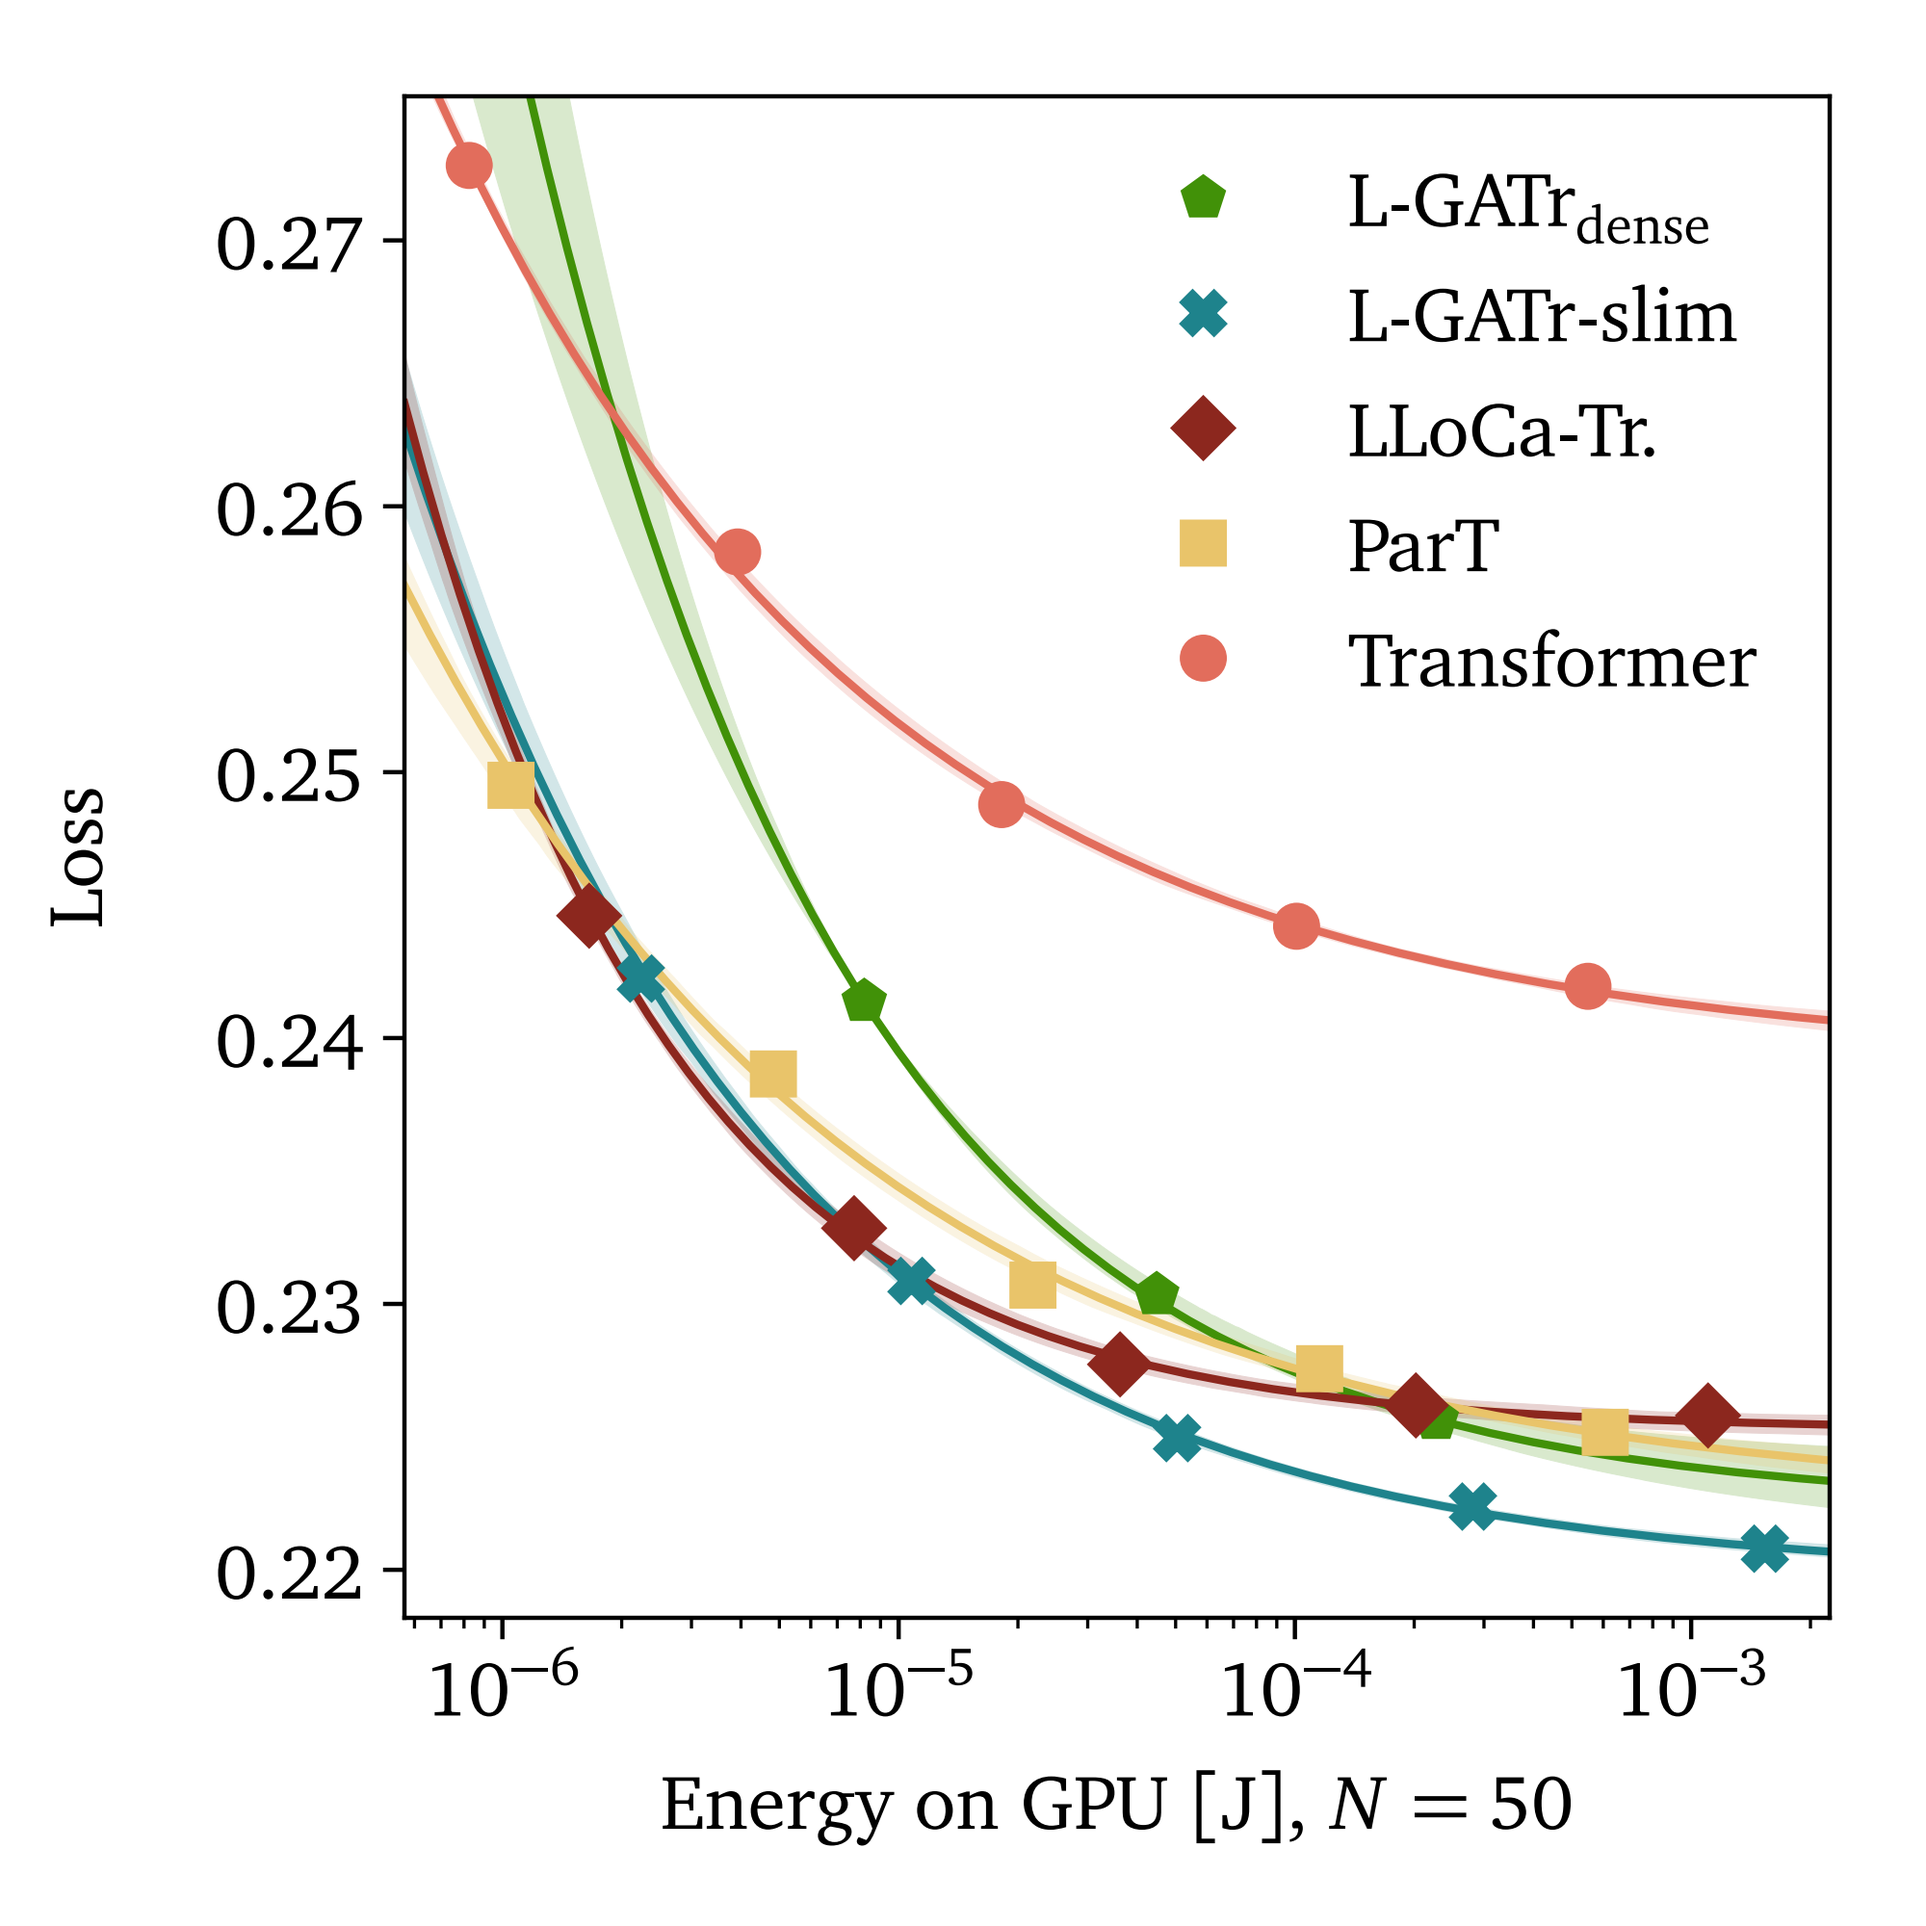
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.27
0.26
0.25
0.24
0.23
0.22
10^{-6}
10^{-5}
10^{-4}
10^{-3}
Energy on GPU [J], $N = 50$

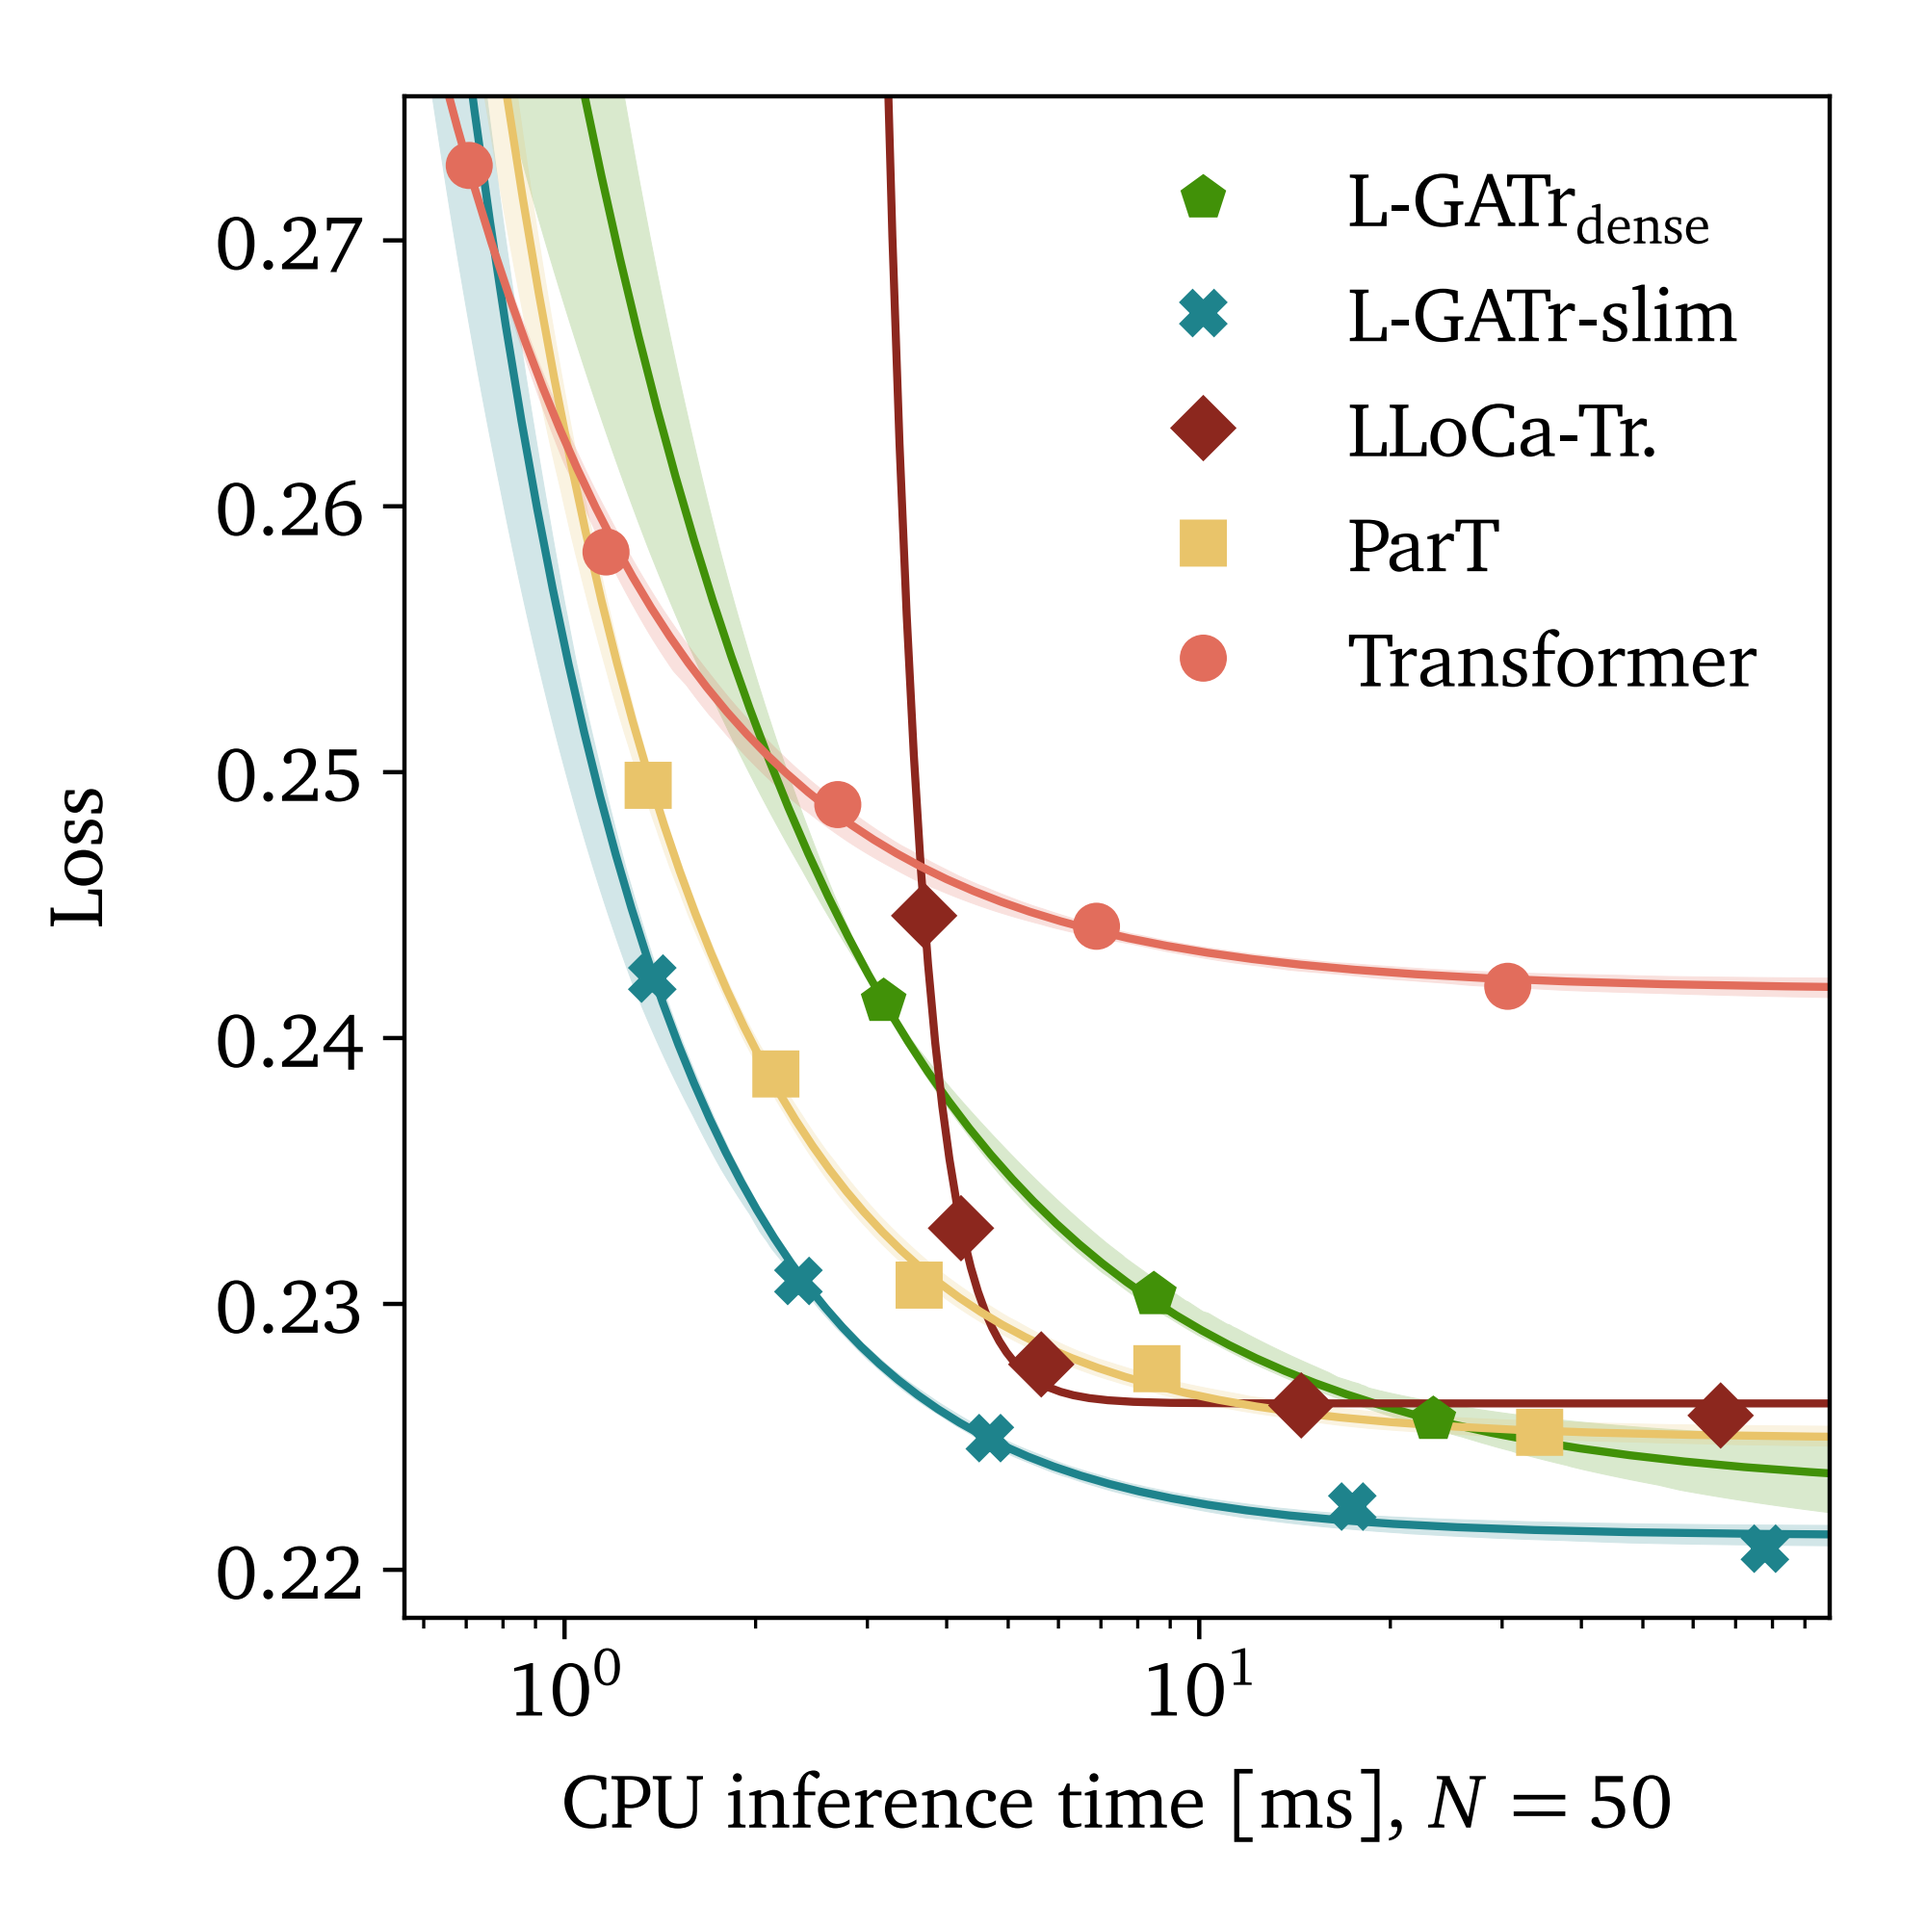
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.27
0.26
0.25
0.24
0.23
0.22
10^0
10^1
CPU inference time [ms], $N = 50$

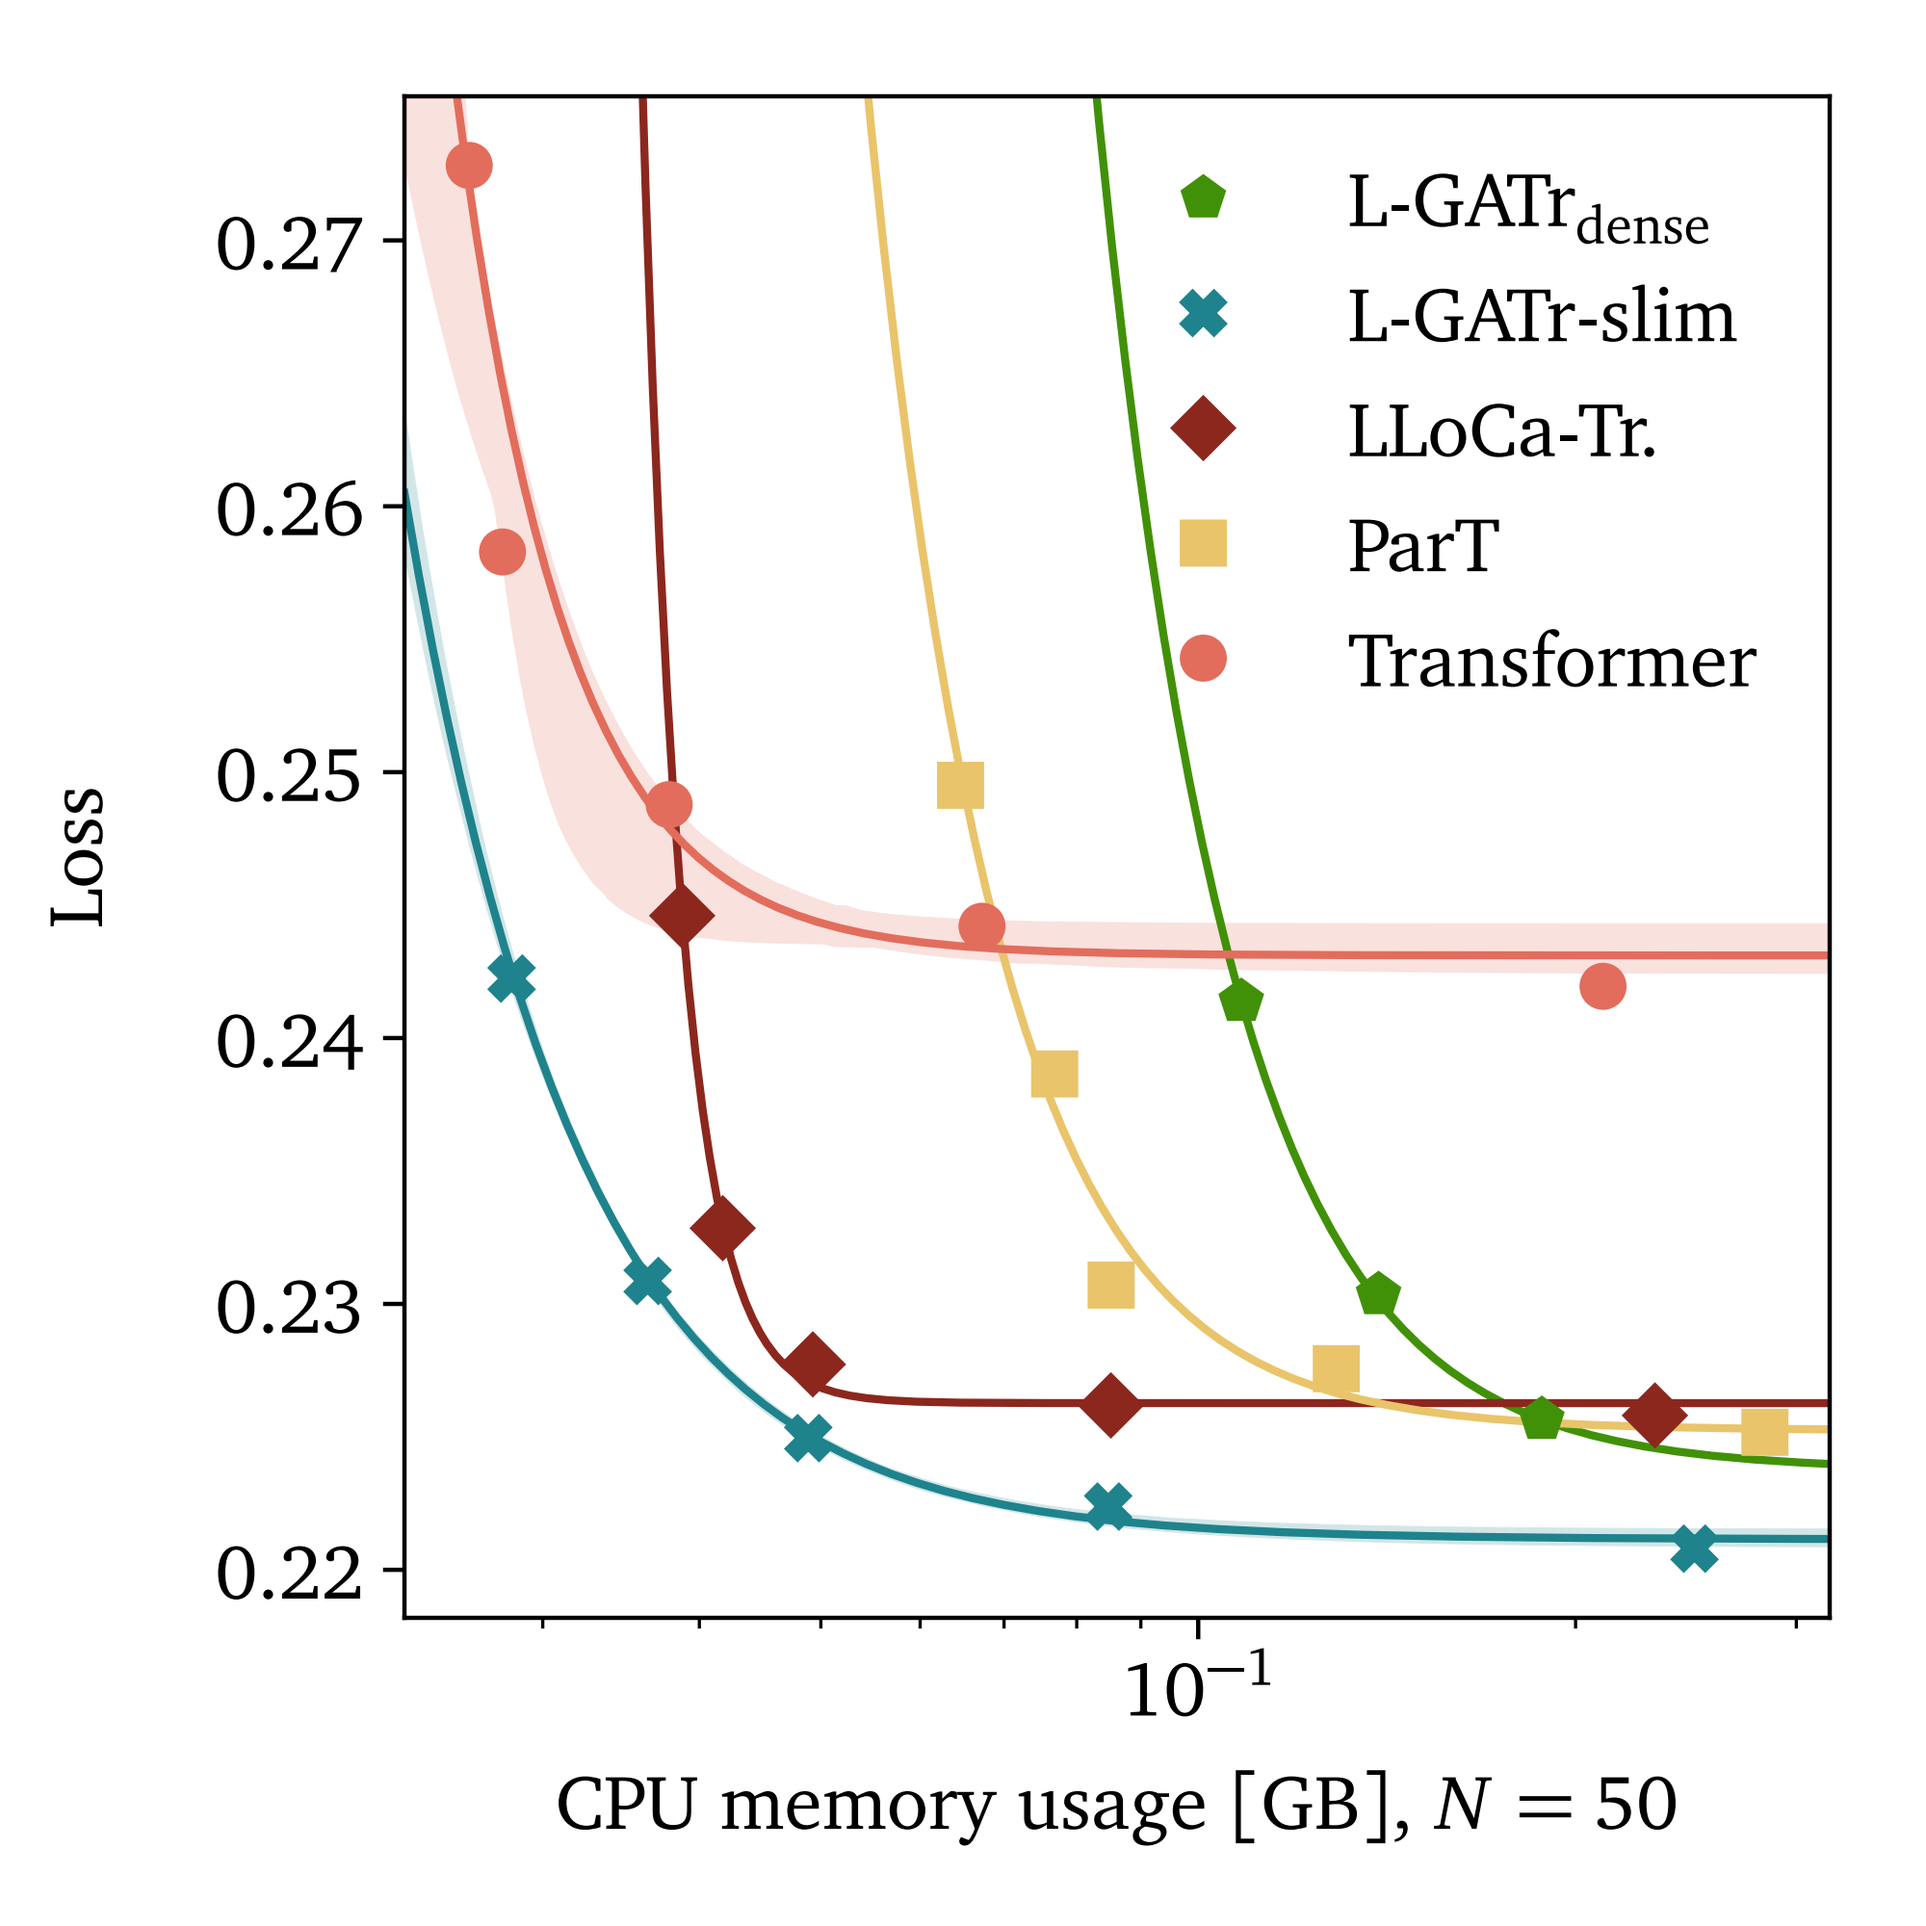
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.27
0.26
0.25
0.24
0.23
0.22
10^{-1}
CPU memory usage [GB], $N = 50$

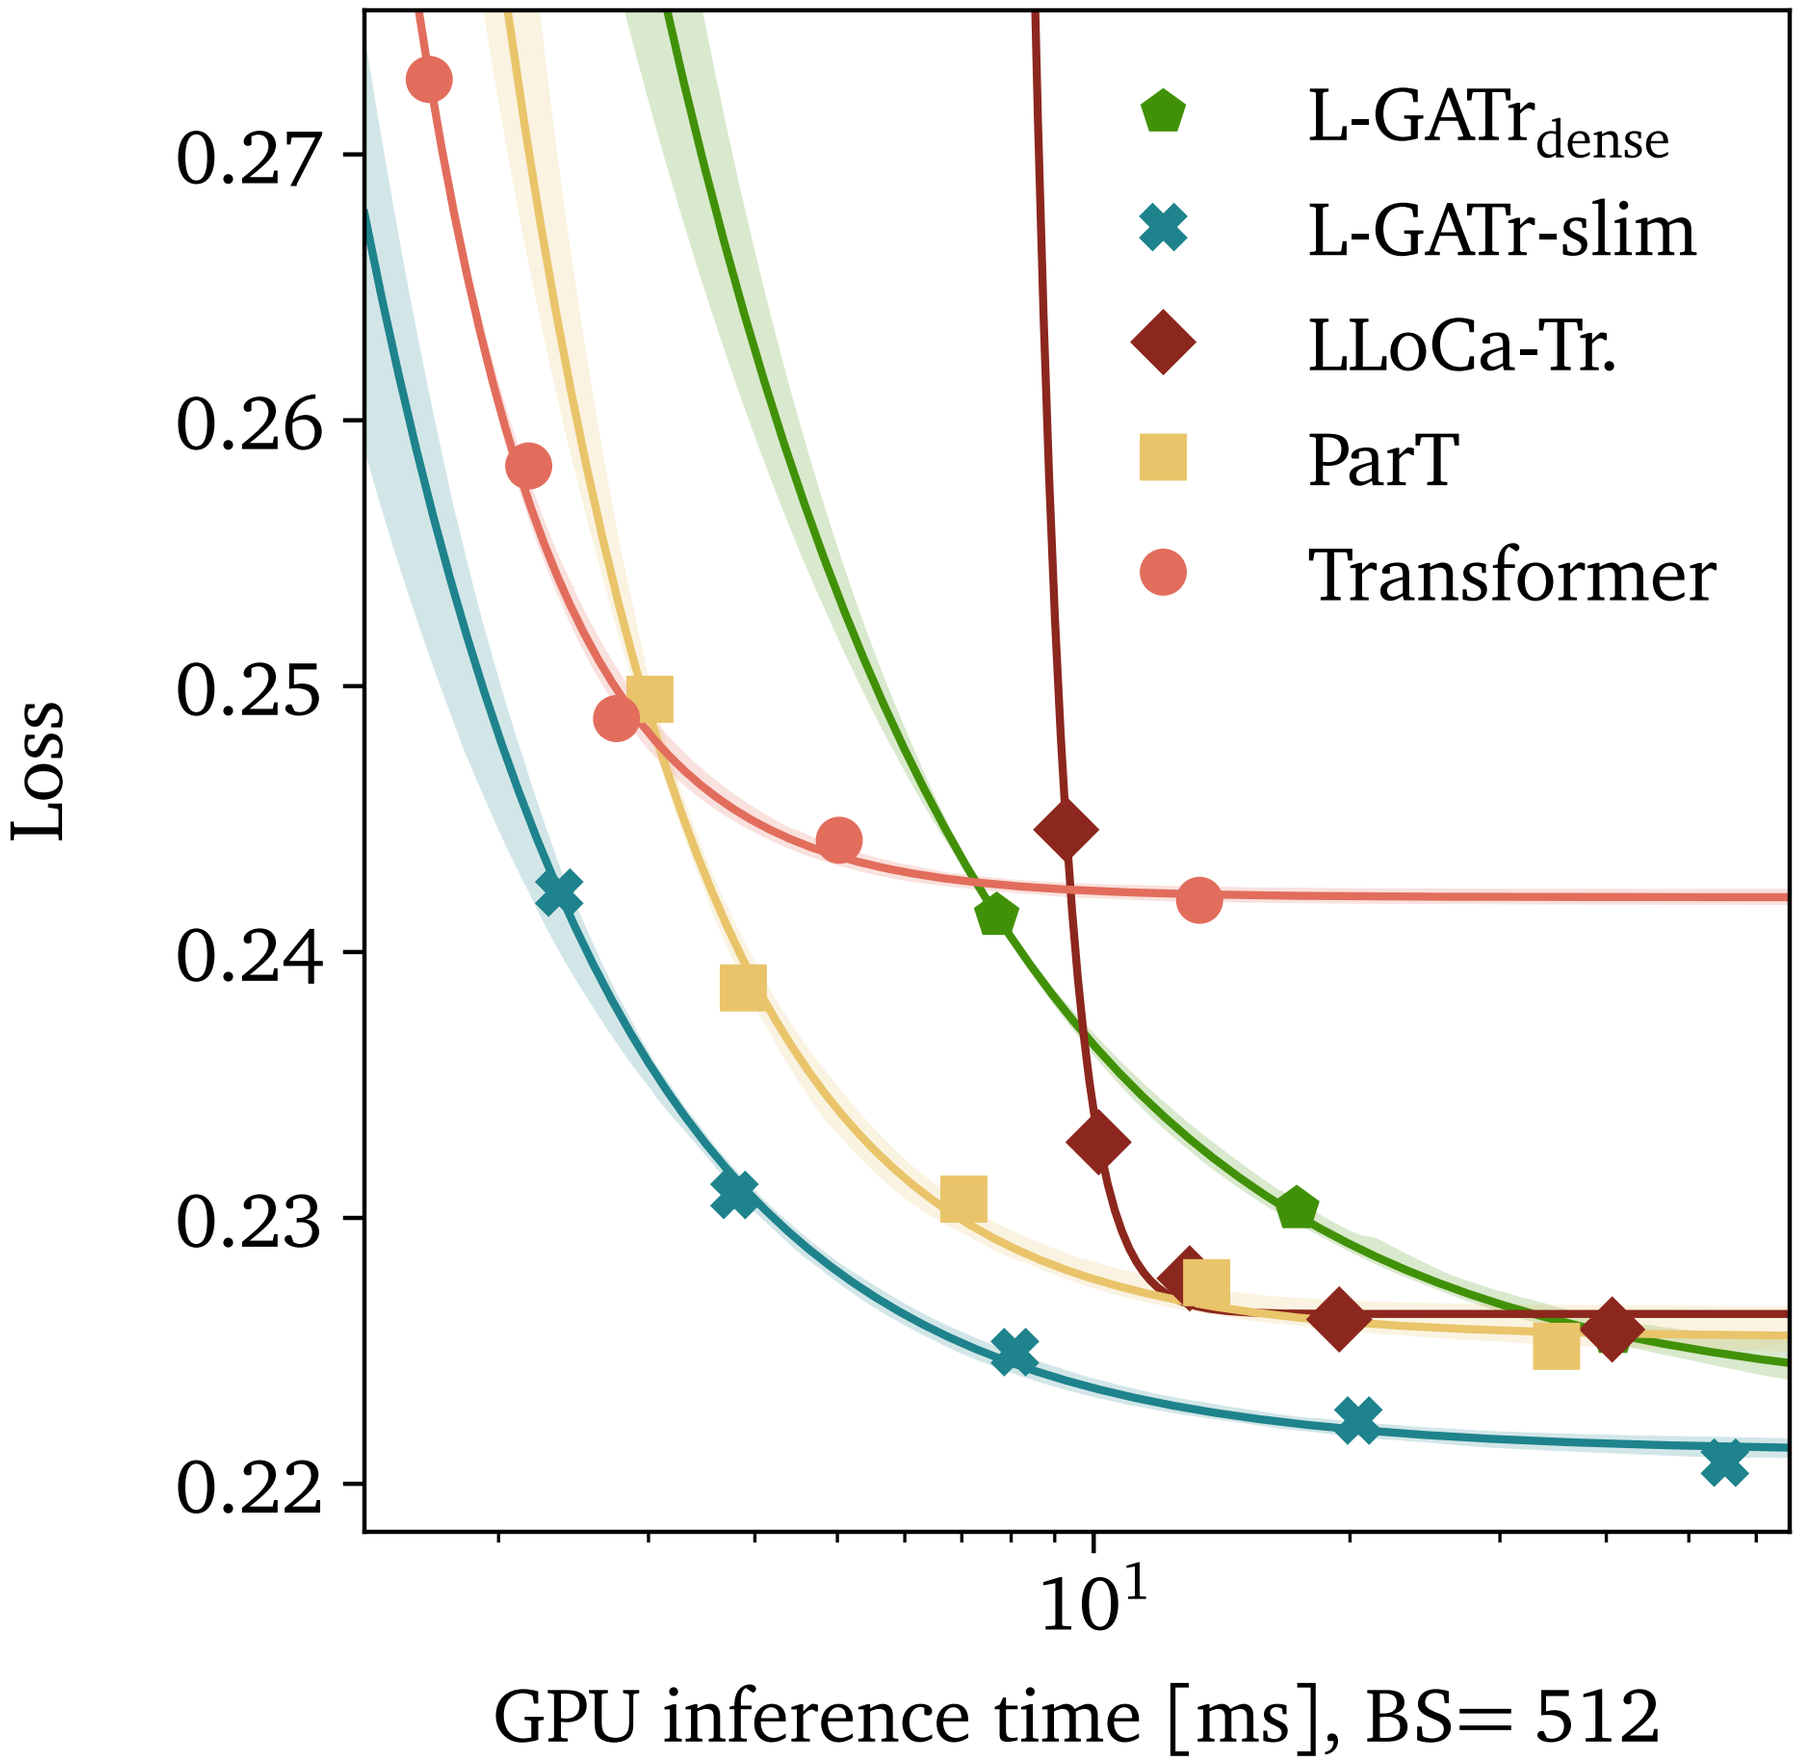
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.27
0.26
0.25
0.24
0.23
0.22
10^1
GPU inference time [ms], BS= 512

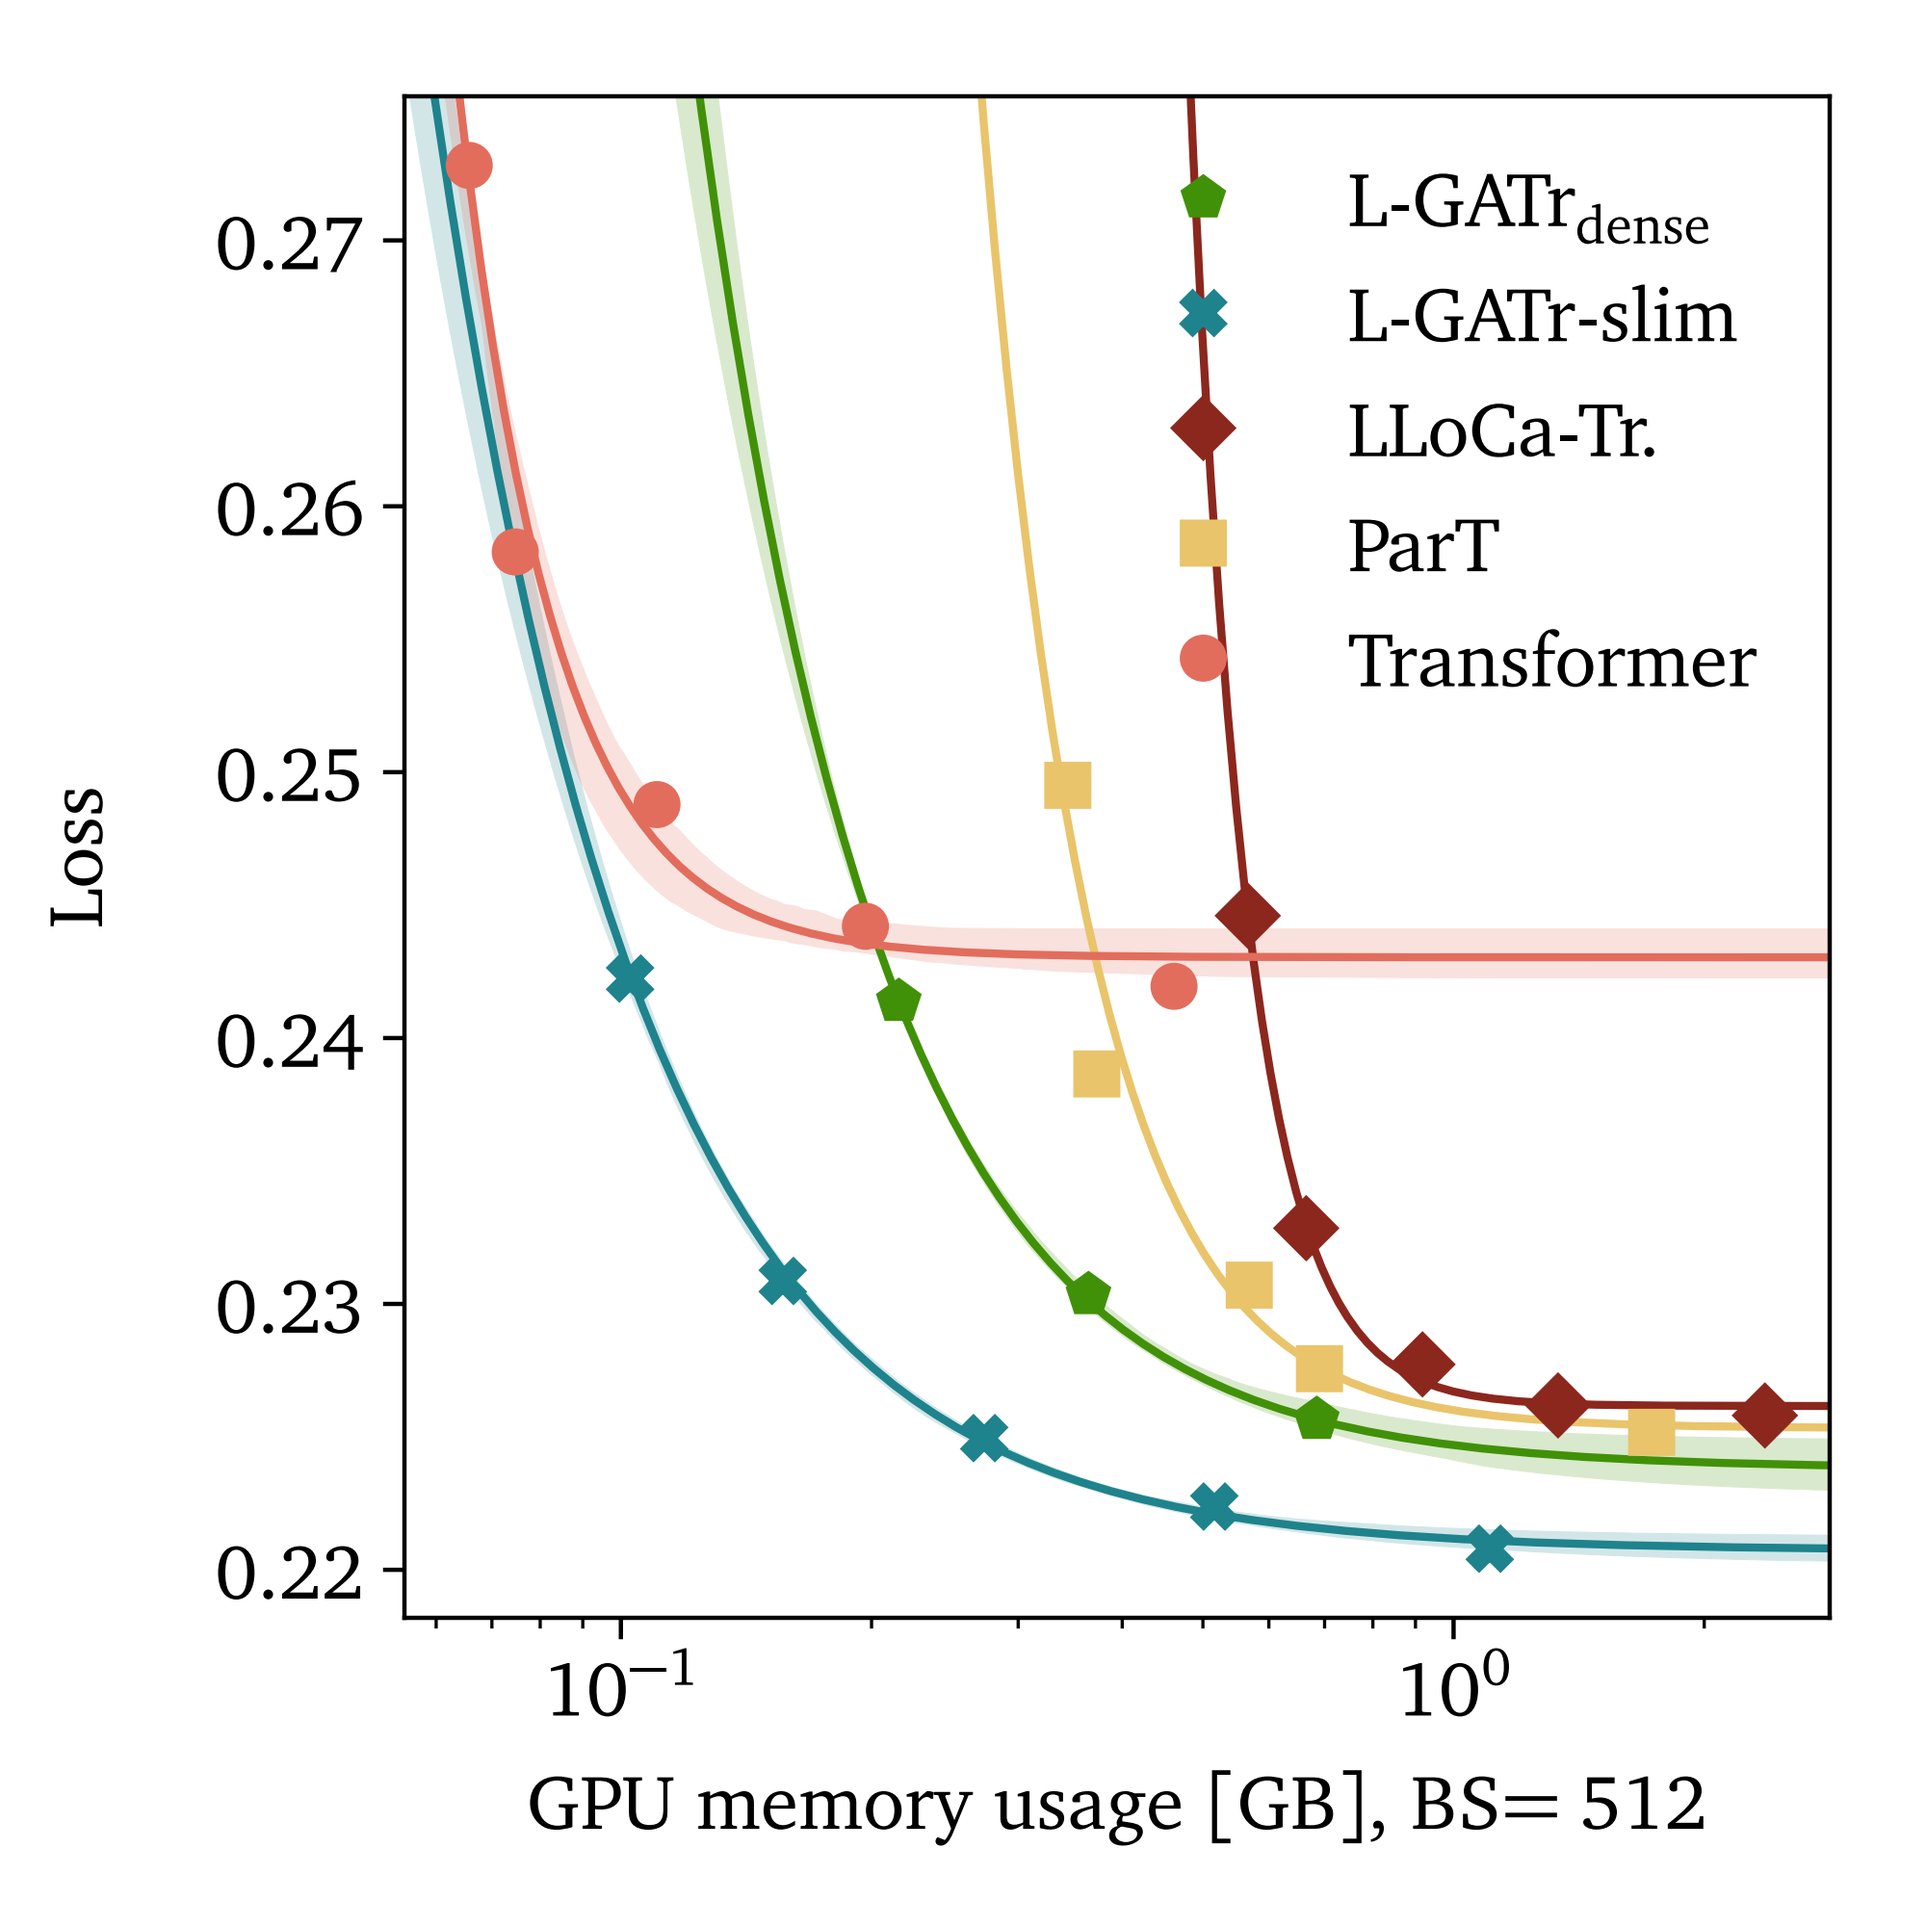
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.27
0.26
0.25
0.24
0.23
0.22
10^{-1}
10^{0}
GPU memory usage [GB], BS= 512

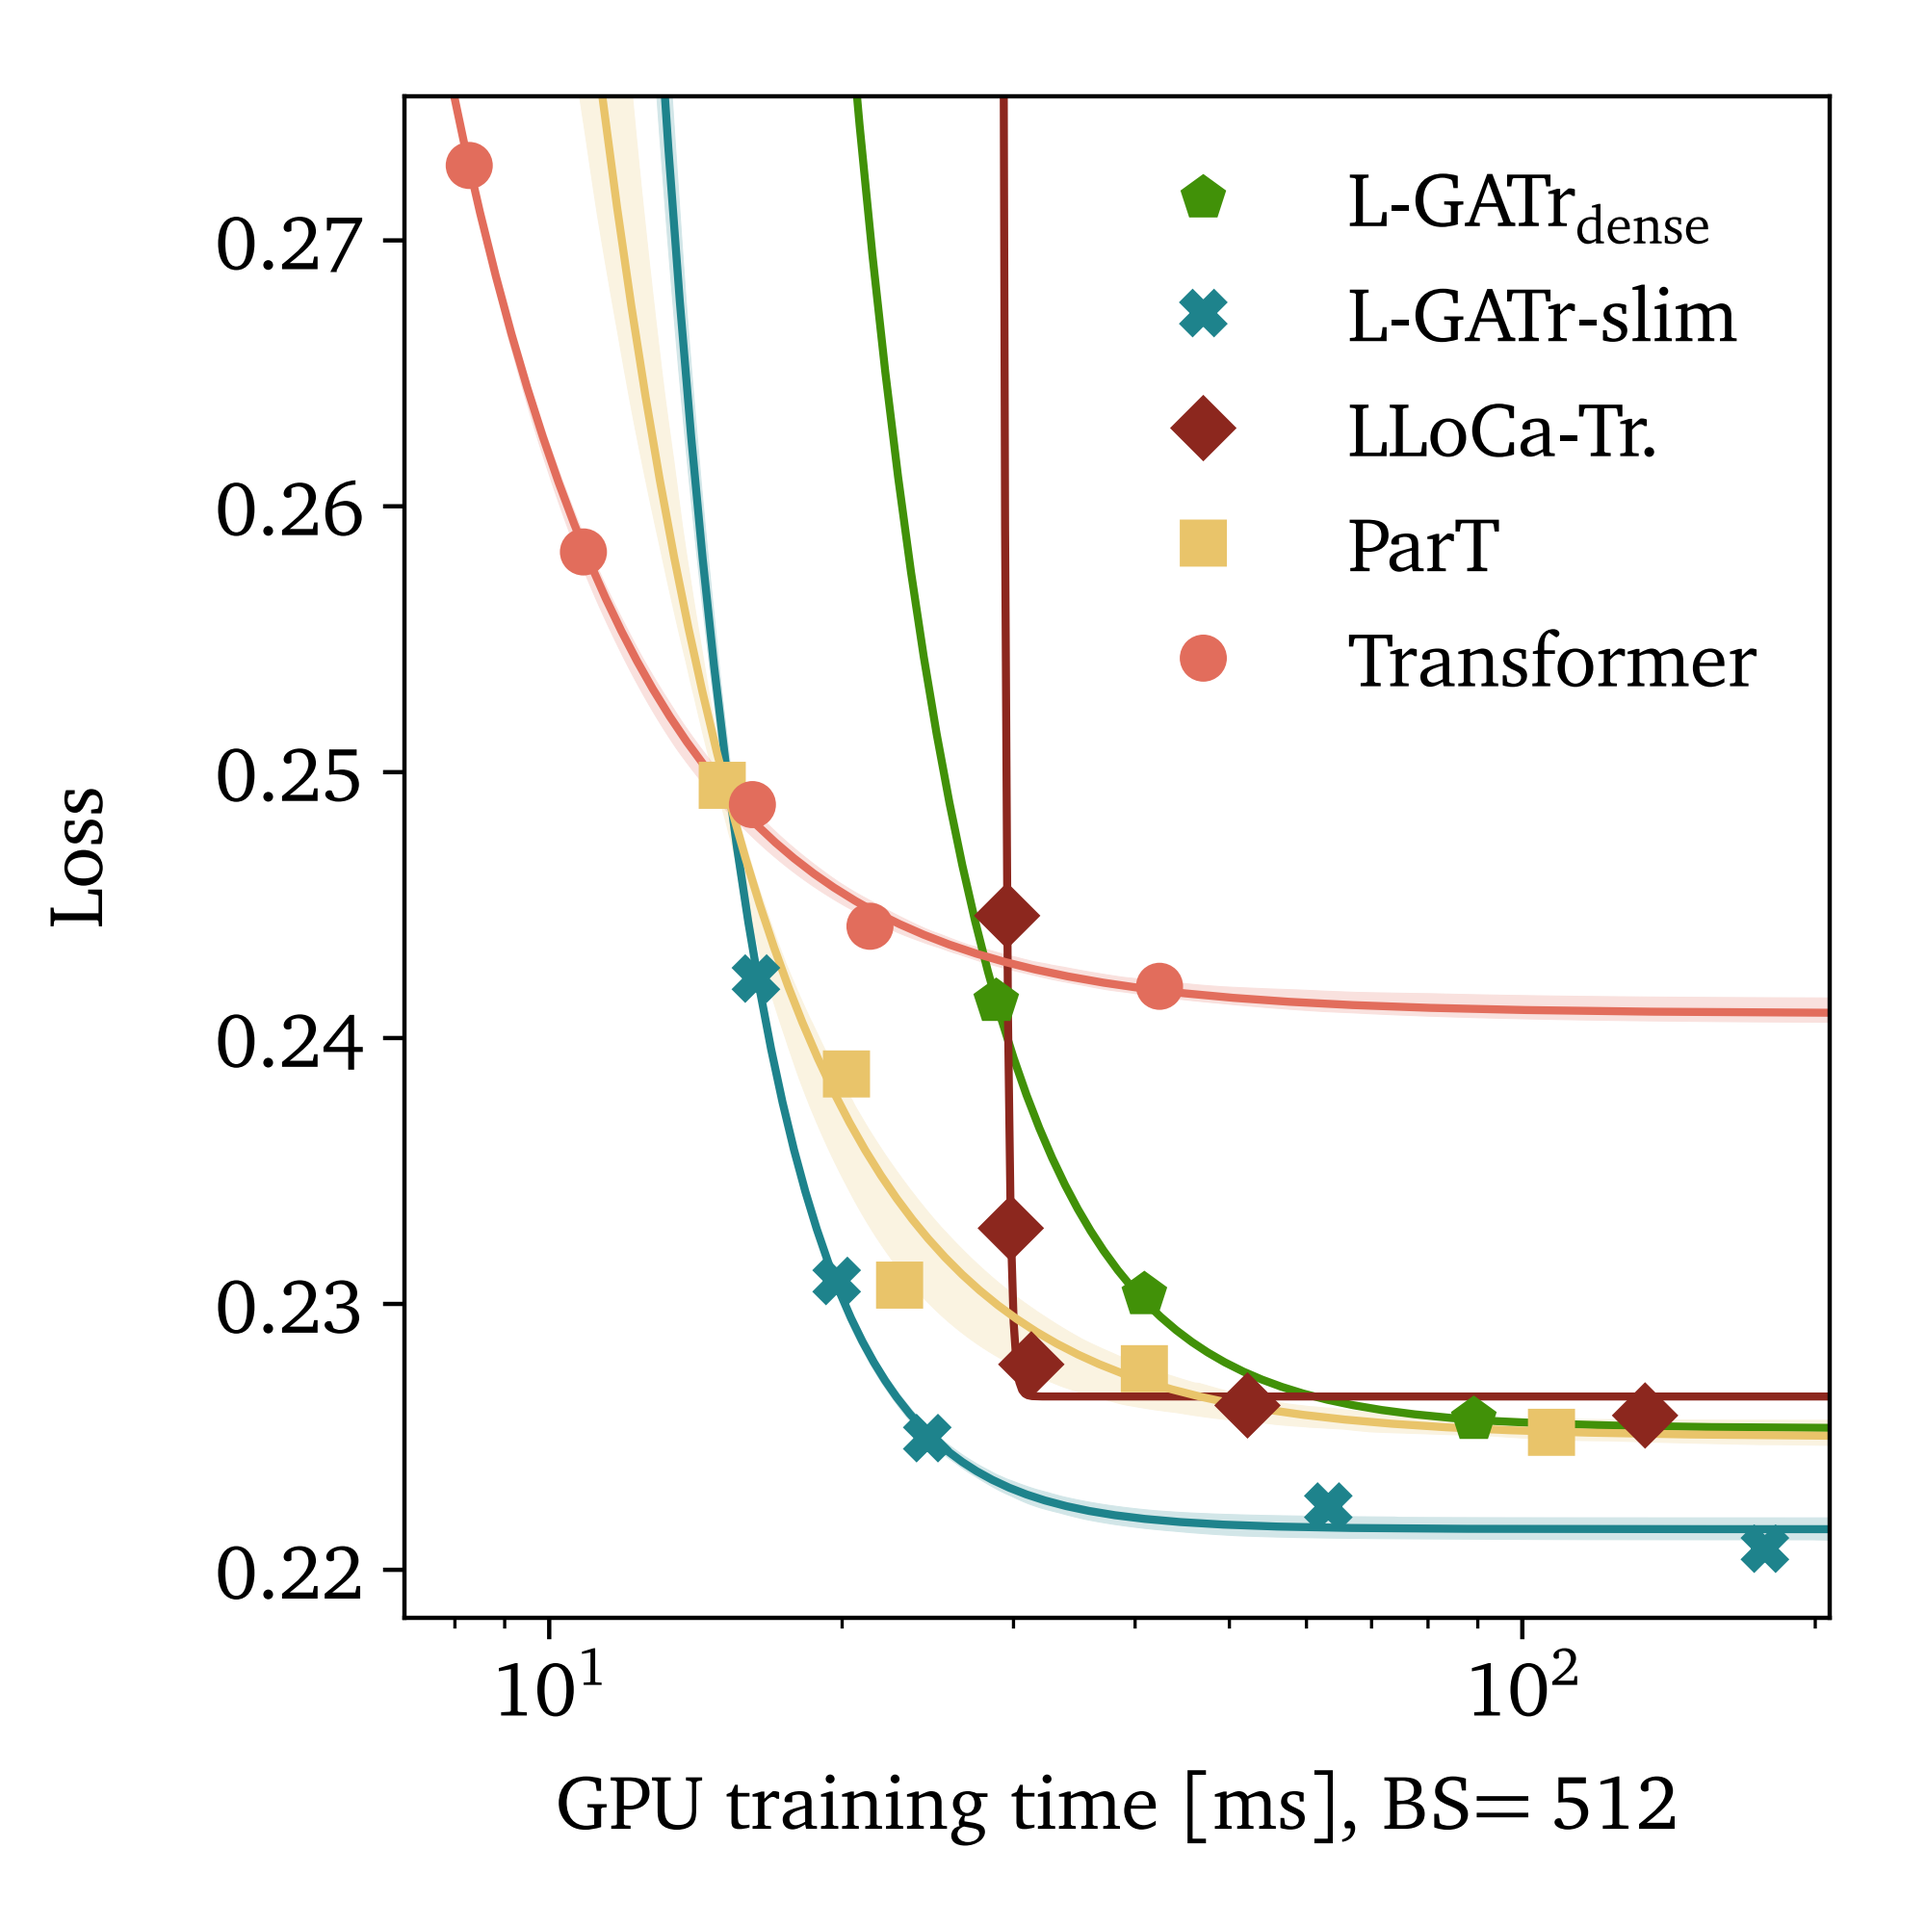
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.27
0.26
0.25
0.24
0.23
0.22
10^1
10^2
GPU training time [ms], BS= 512

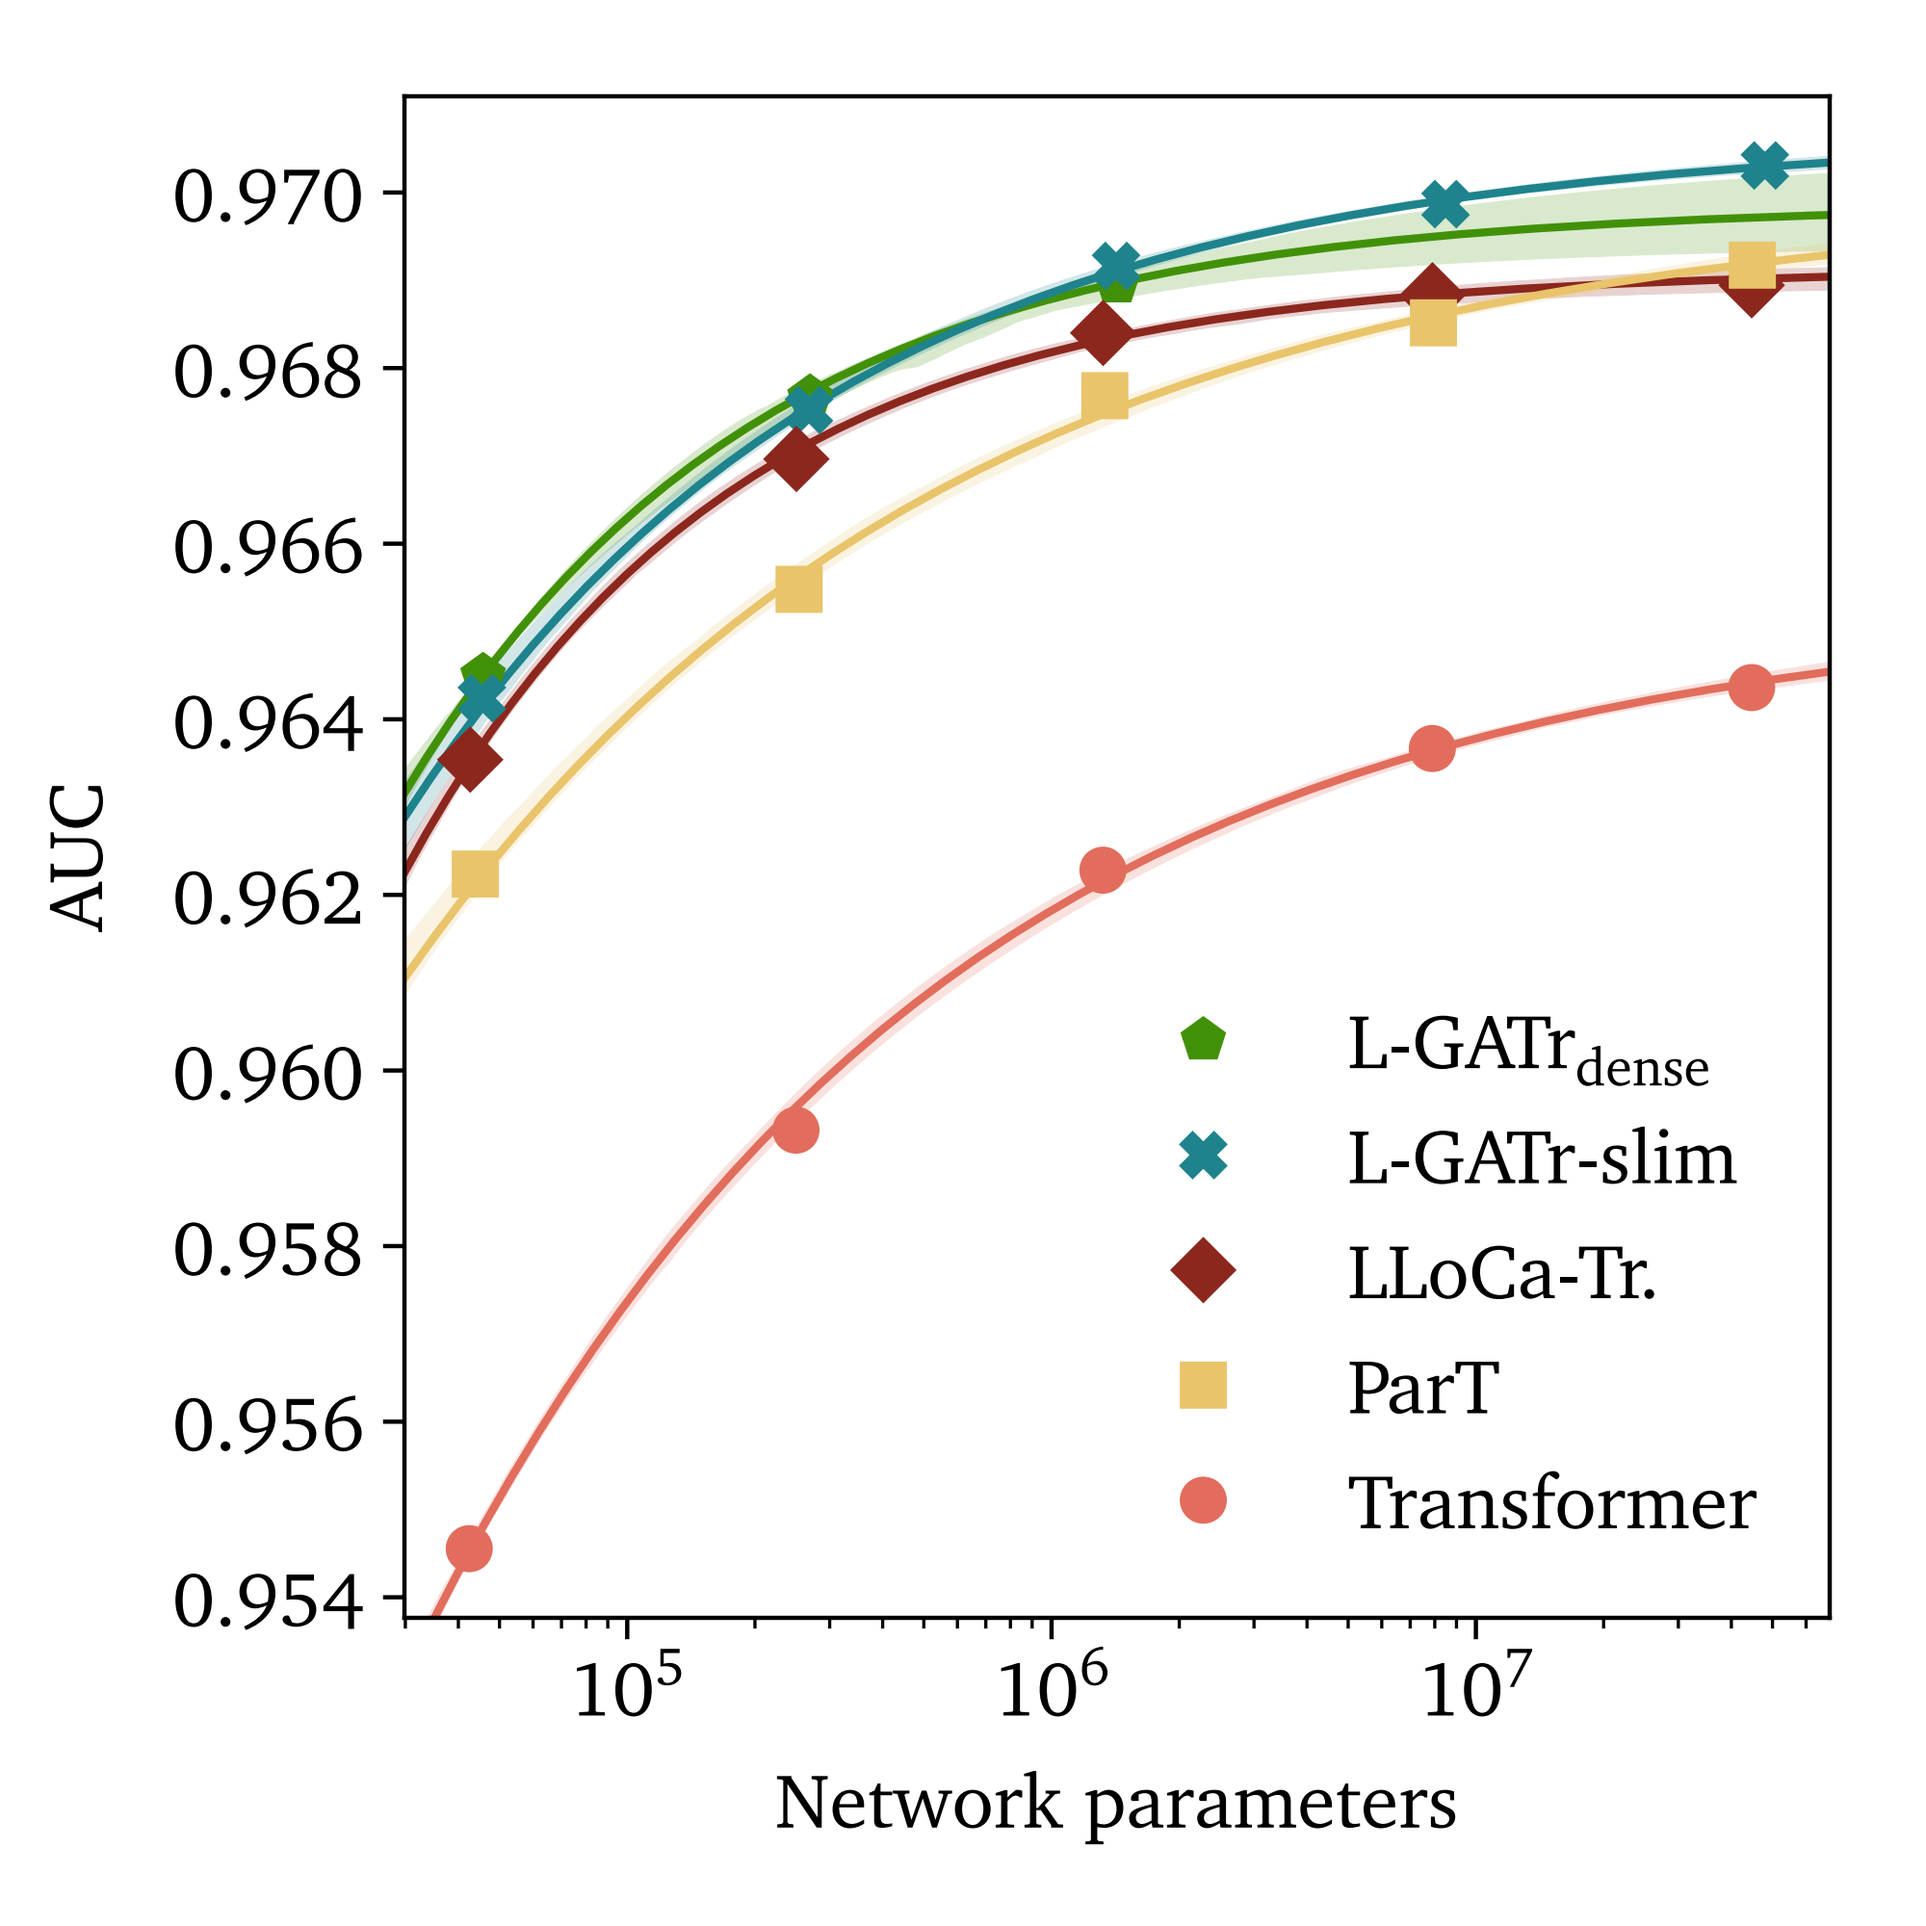
0.970
0.968
0.966
0.964
0.962
0.960
0.958
0.956
0.954
AUC
10^5
10^6
10^7
Network parameters
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

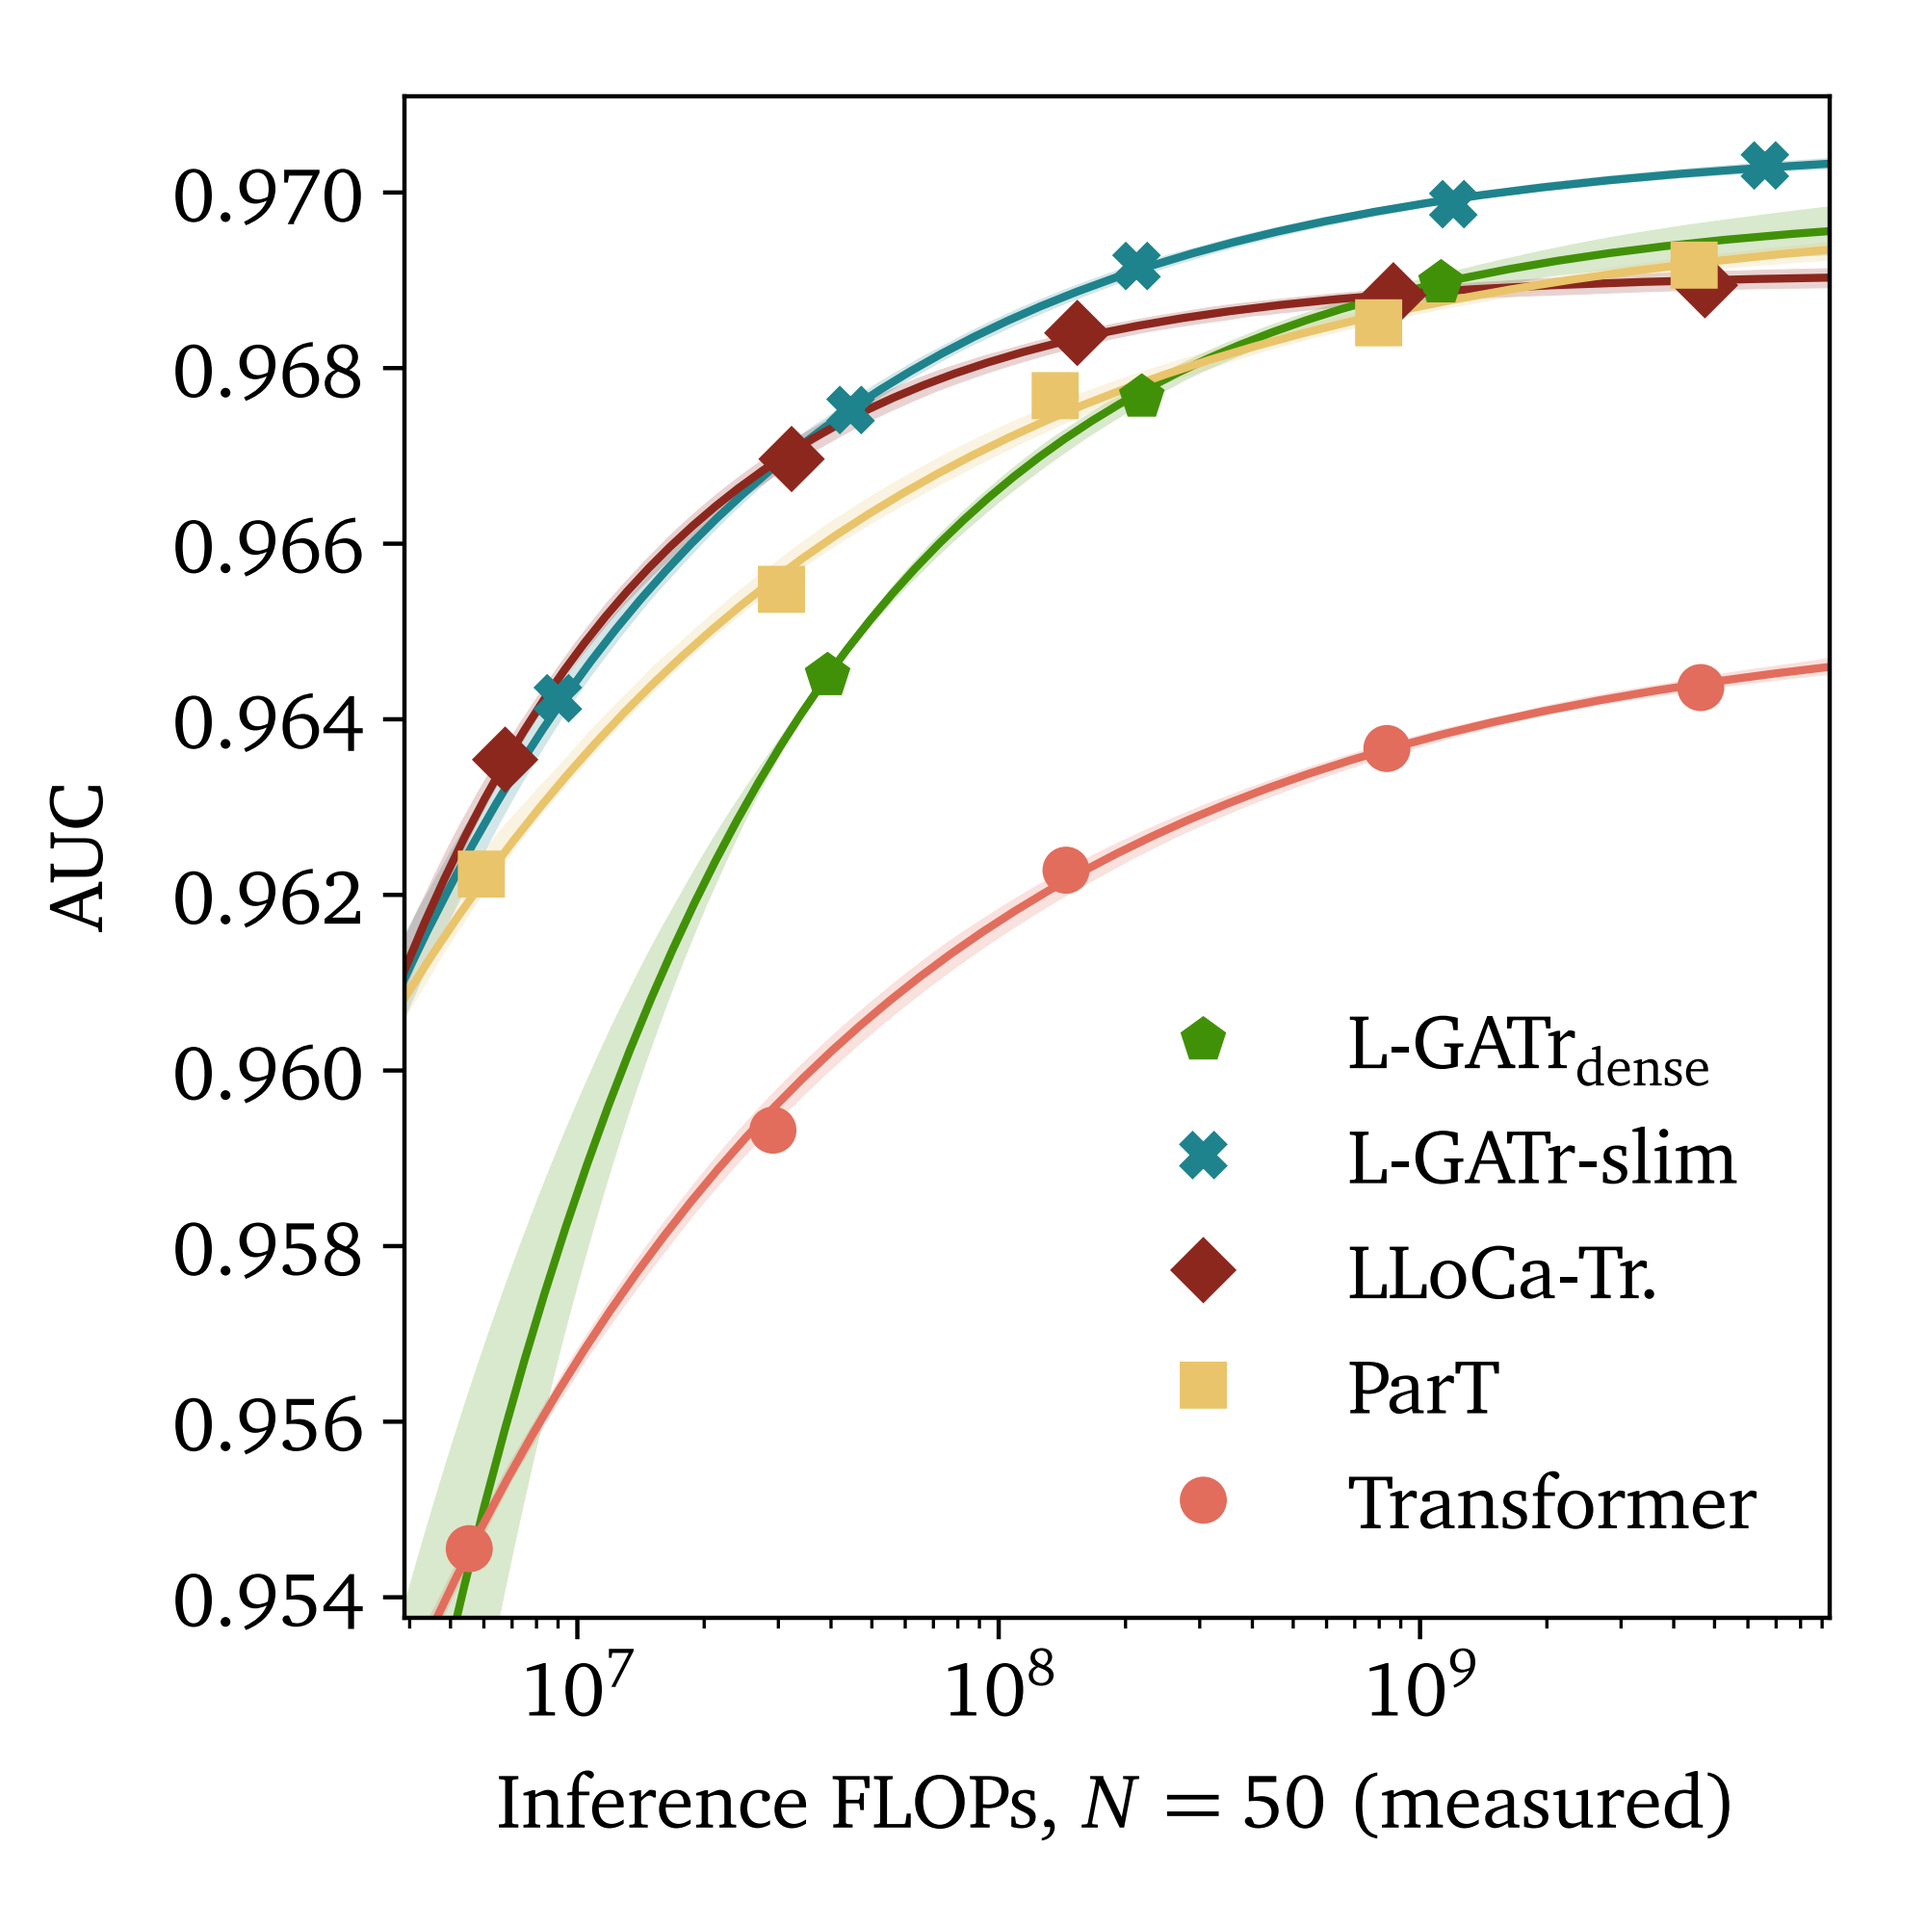
0.970
0.968
0.966
0.964
0.962
0.960
0.958
0.956
0.954
AUC
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (measured)
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

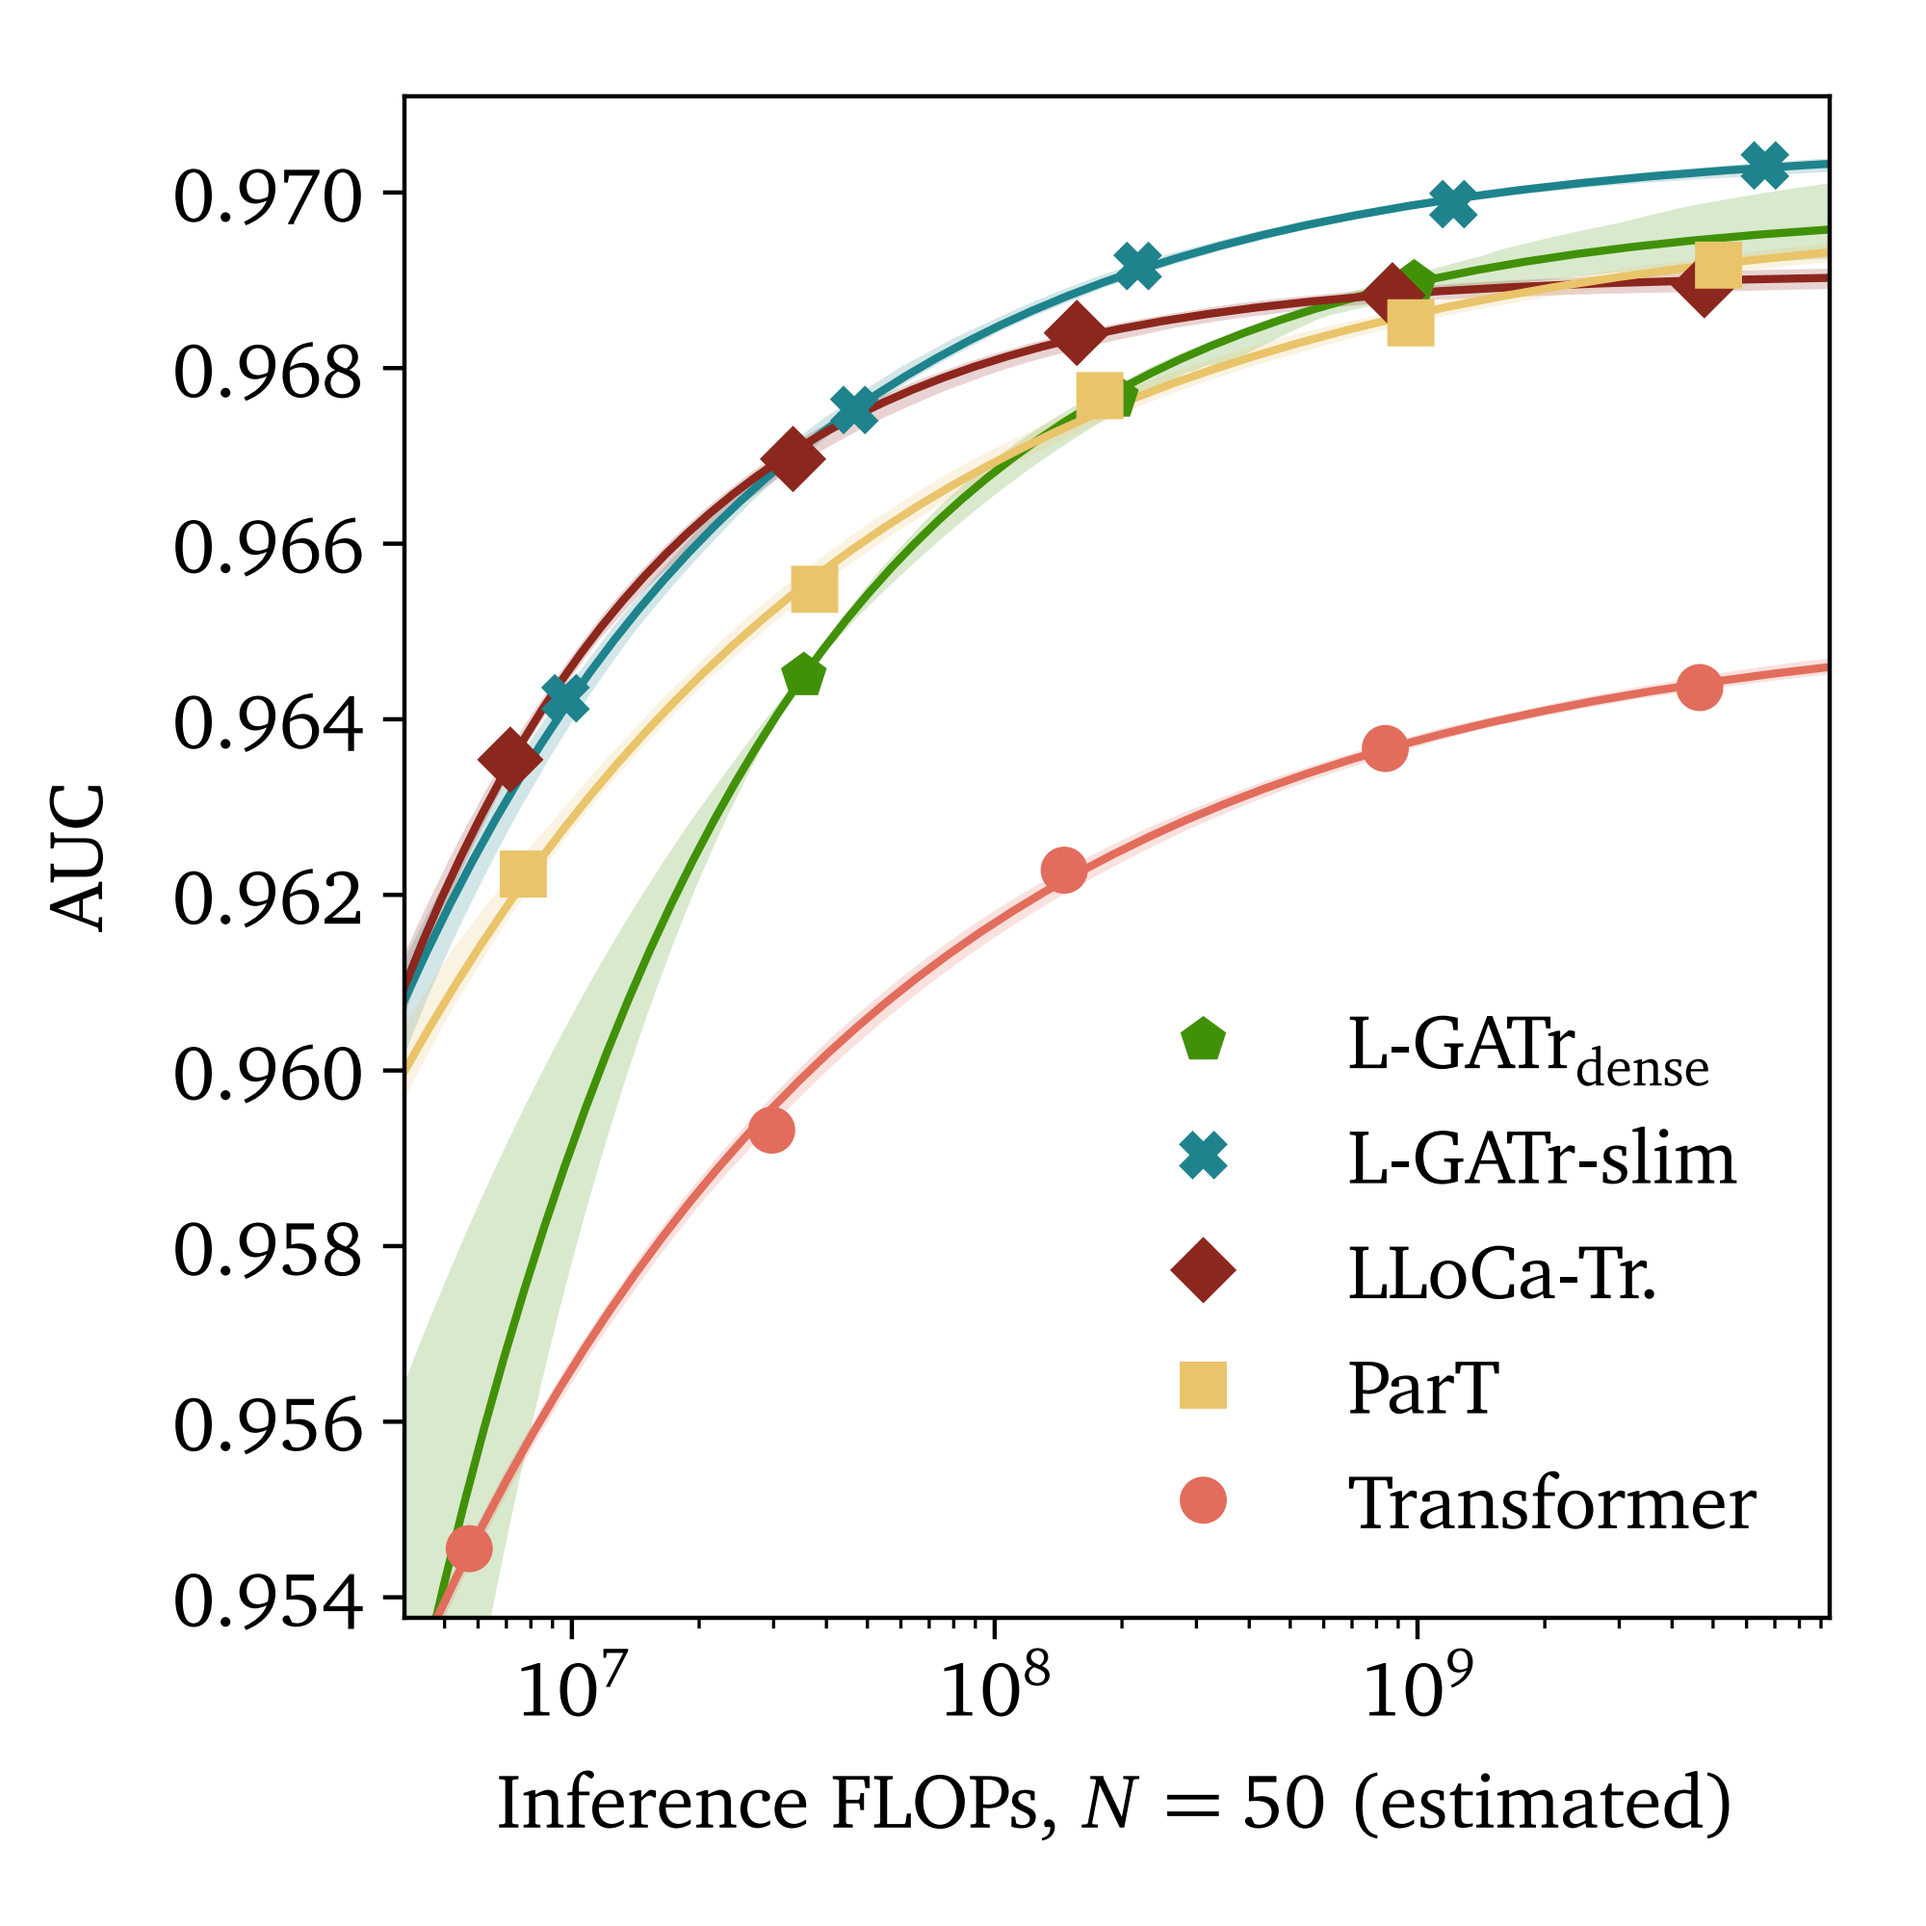
0.970
0.968
0.966
0.964
0.962
0.960
0.958
0.956
0.954
AUC
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (estimated)
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

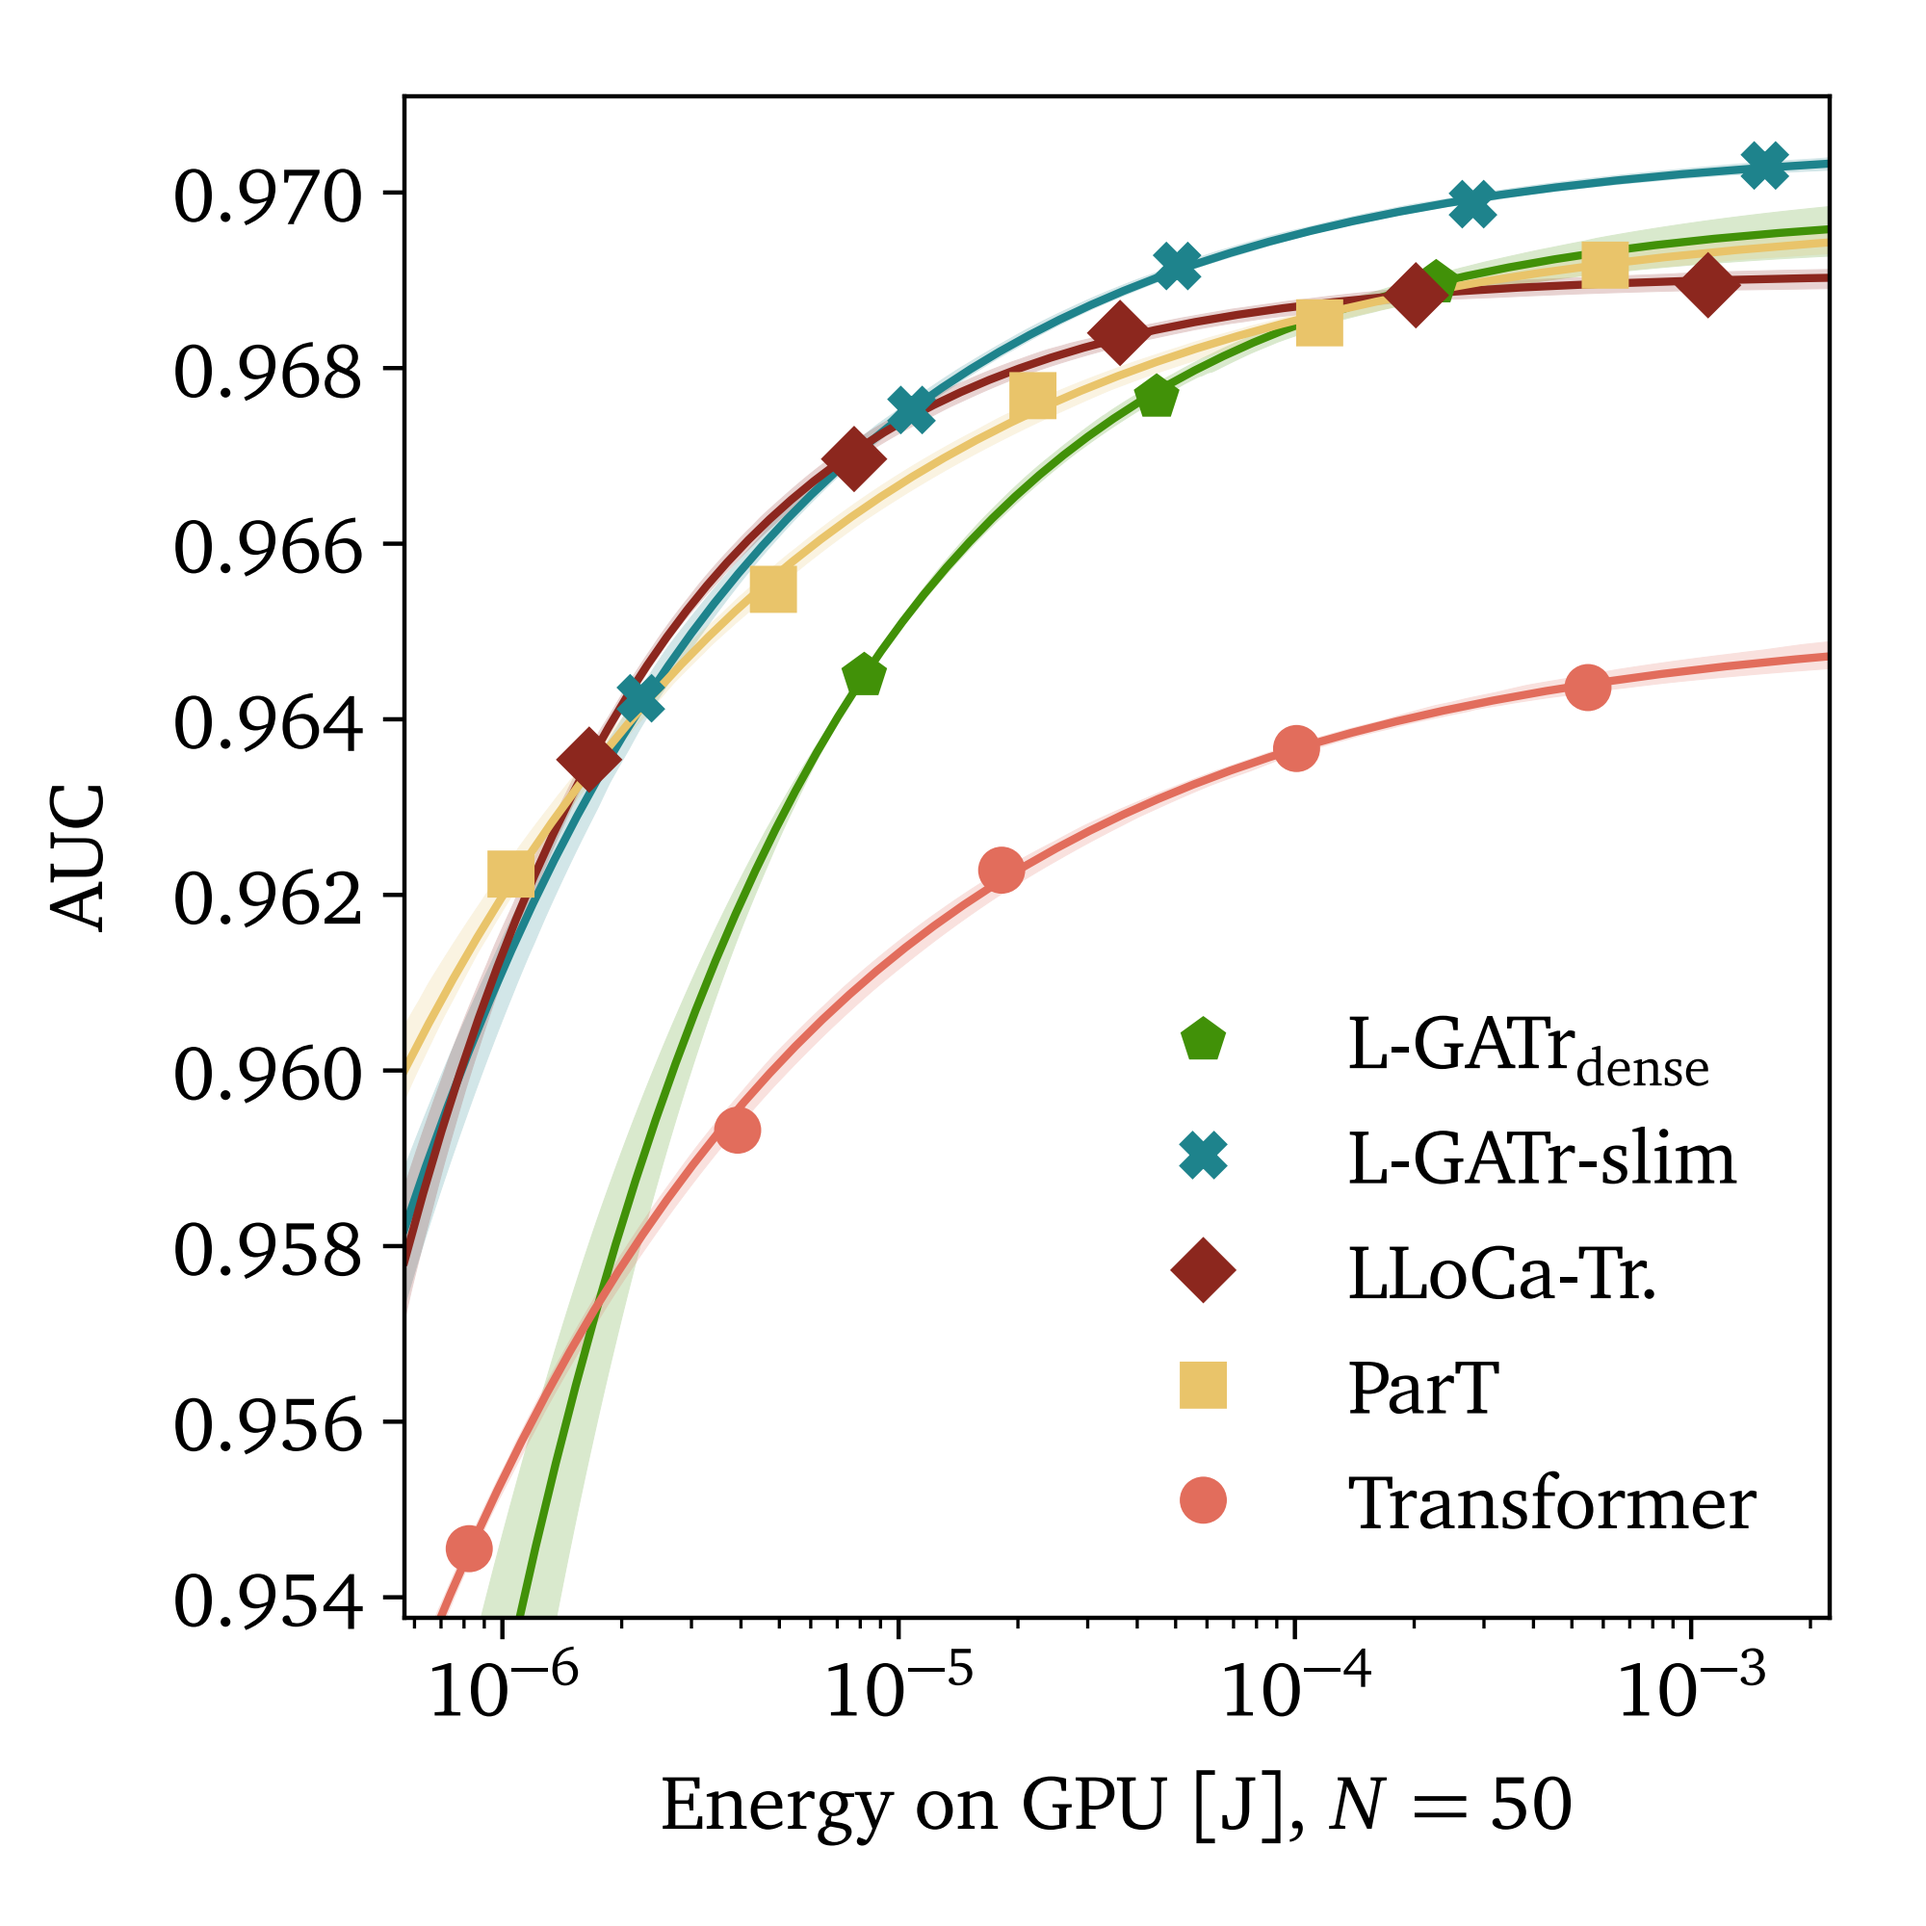
0.970
0.968
0.966
0.964
0.962
0.960
0.958
0.956
0.954
AUC
10^{-6}
10^{-5}
10^{-4}
10^{-3}
Energy on GPU [J], $N = 50$
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

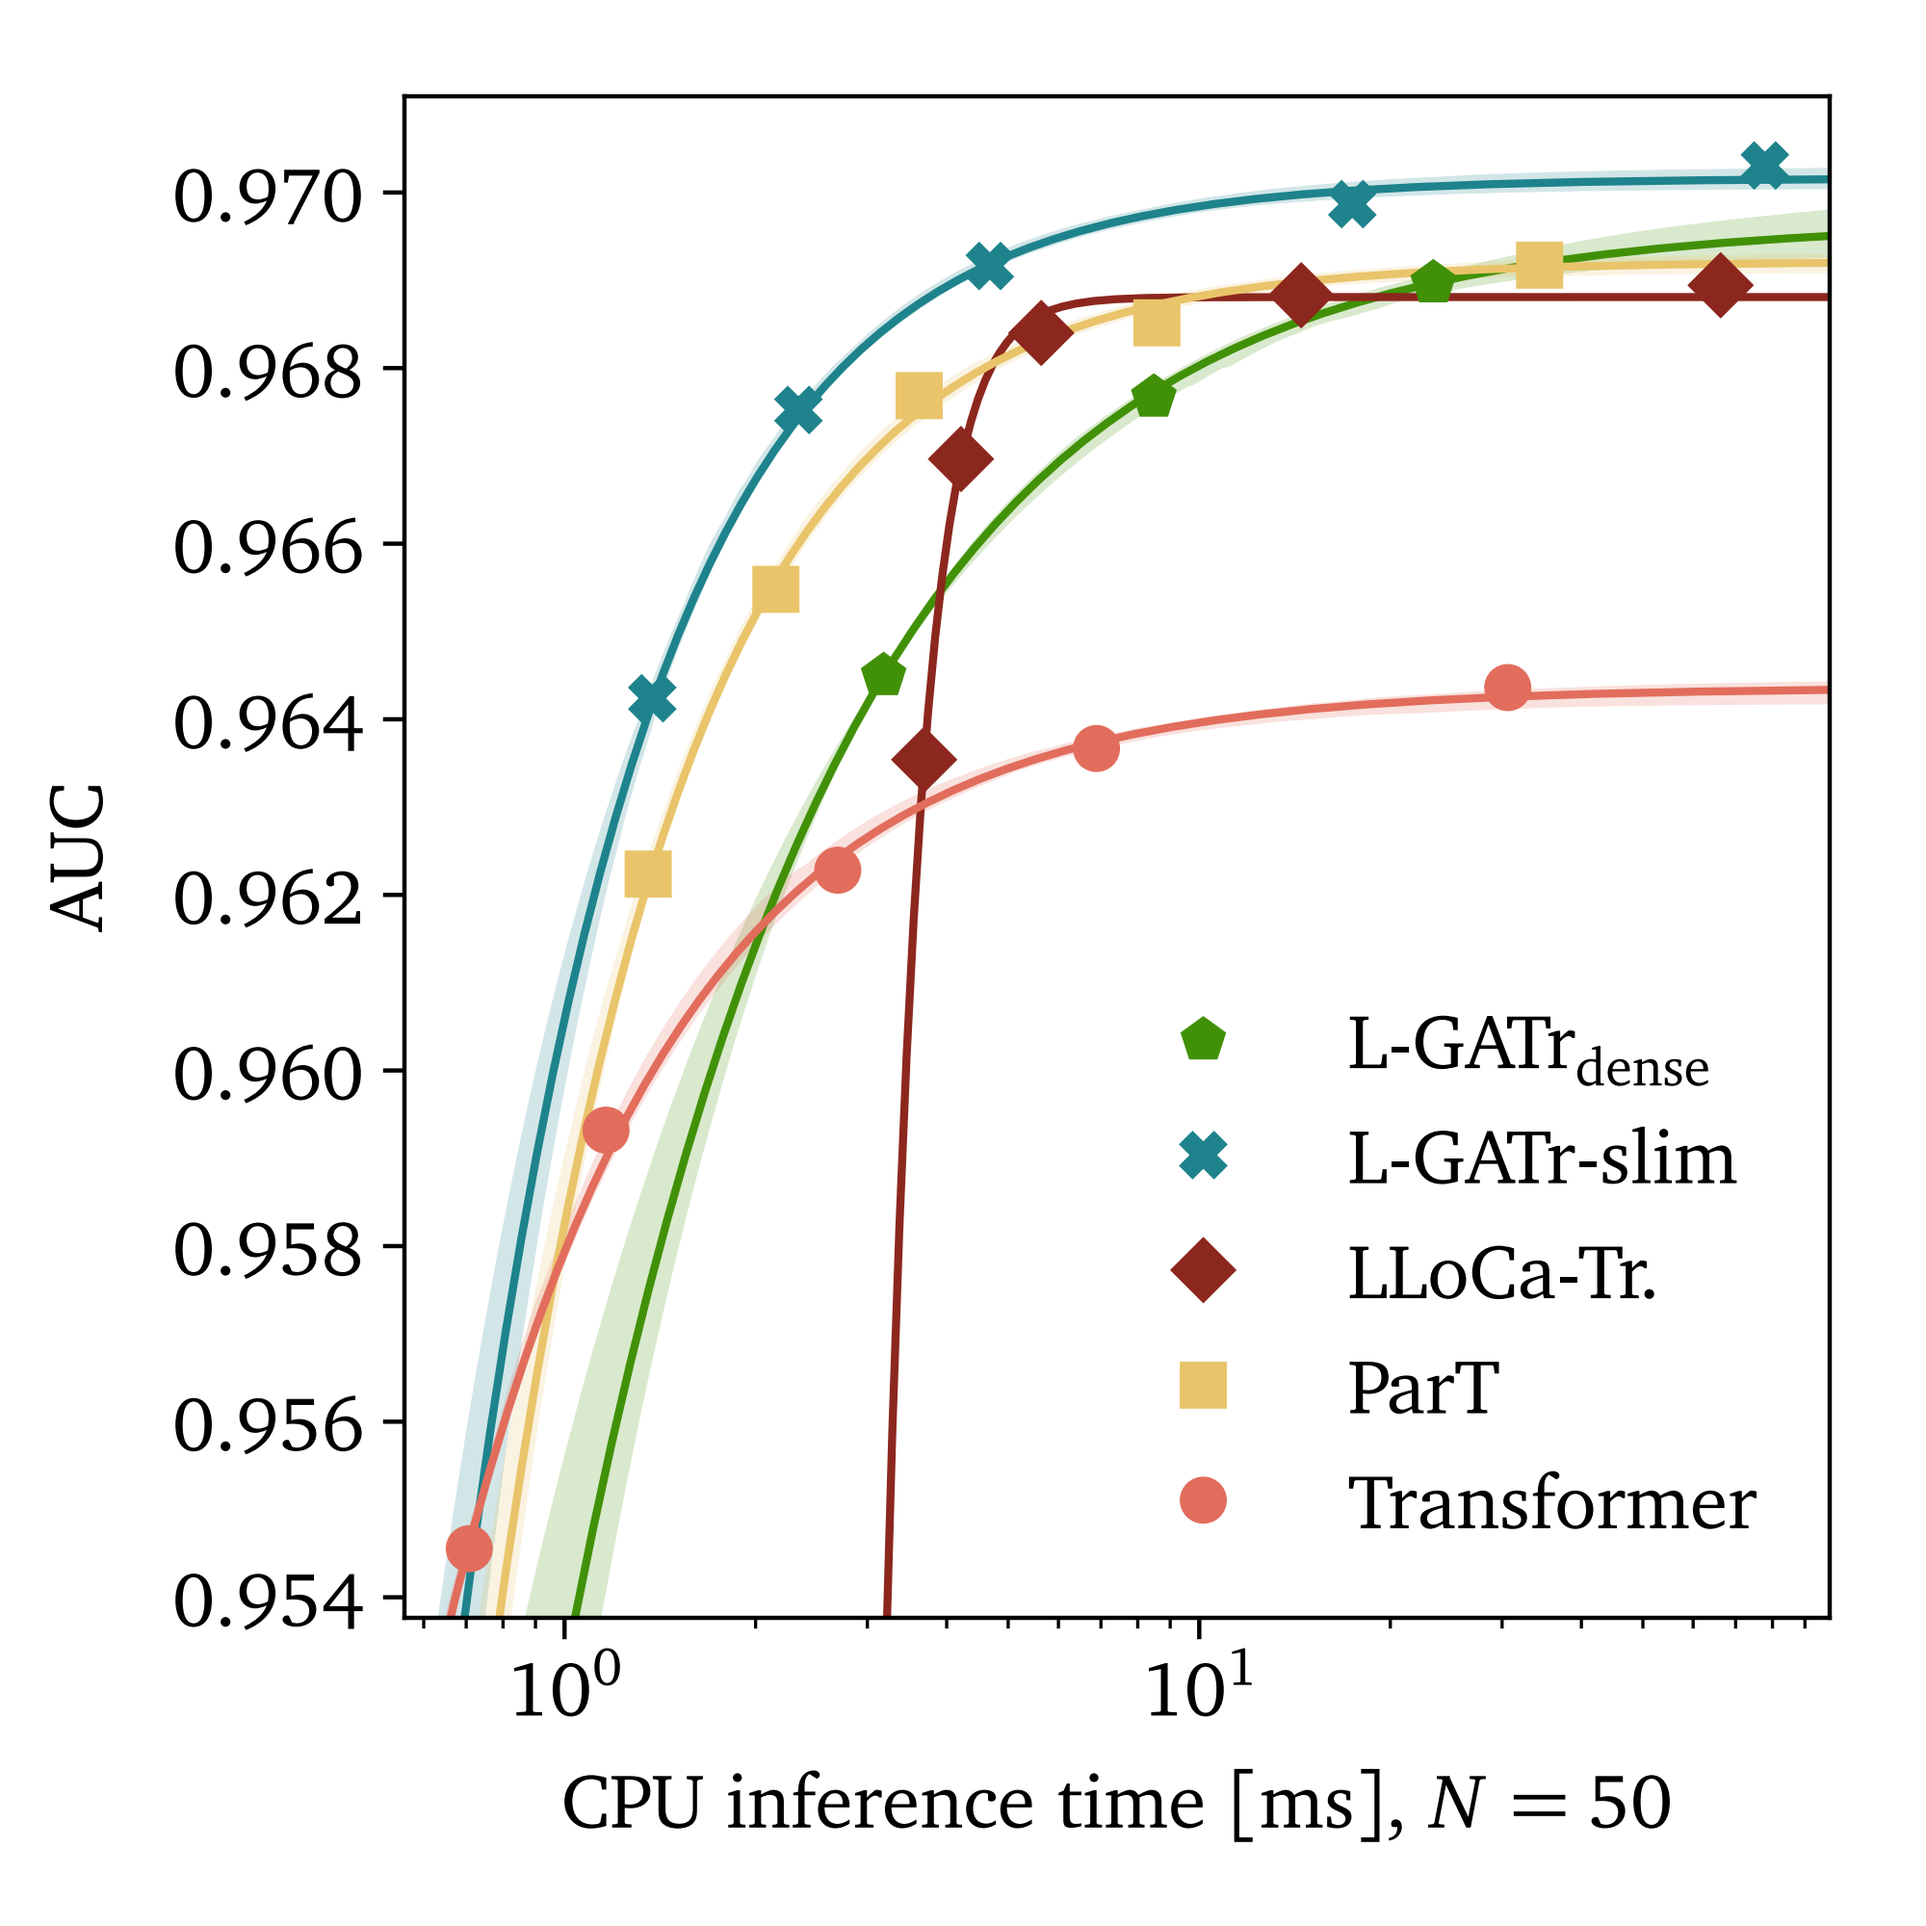
0.970
0.968
0.966
0.964
0.962
0.960
0.958
0.956
0.954
AUC
10^0
10^1
CPU inference time [ms], $N = 50$
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

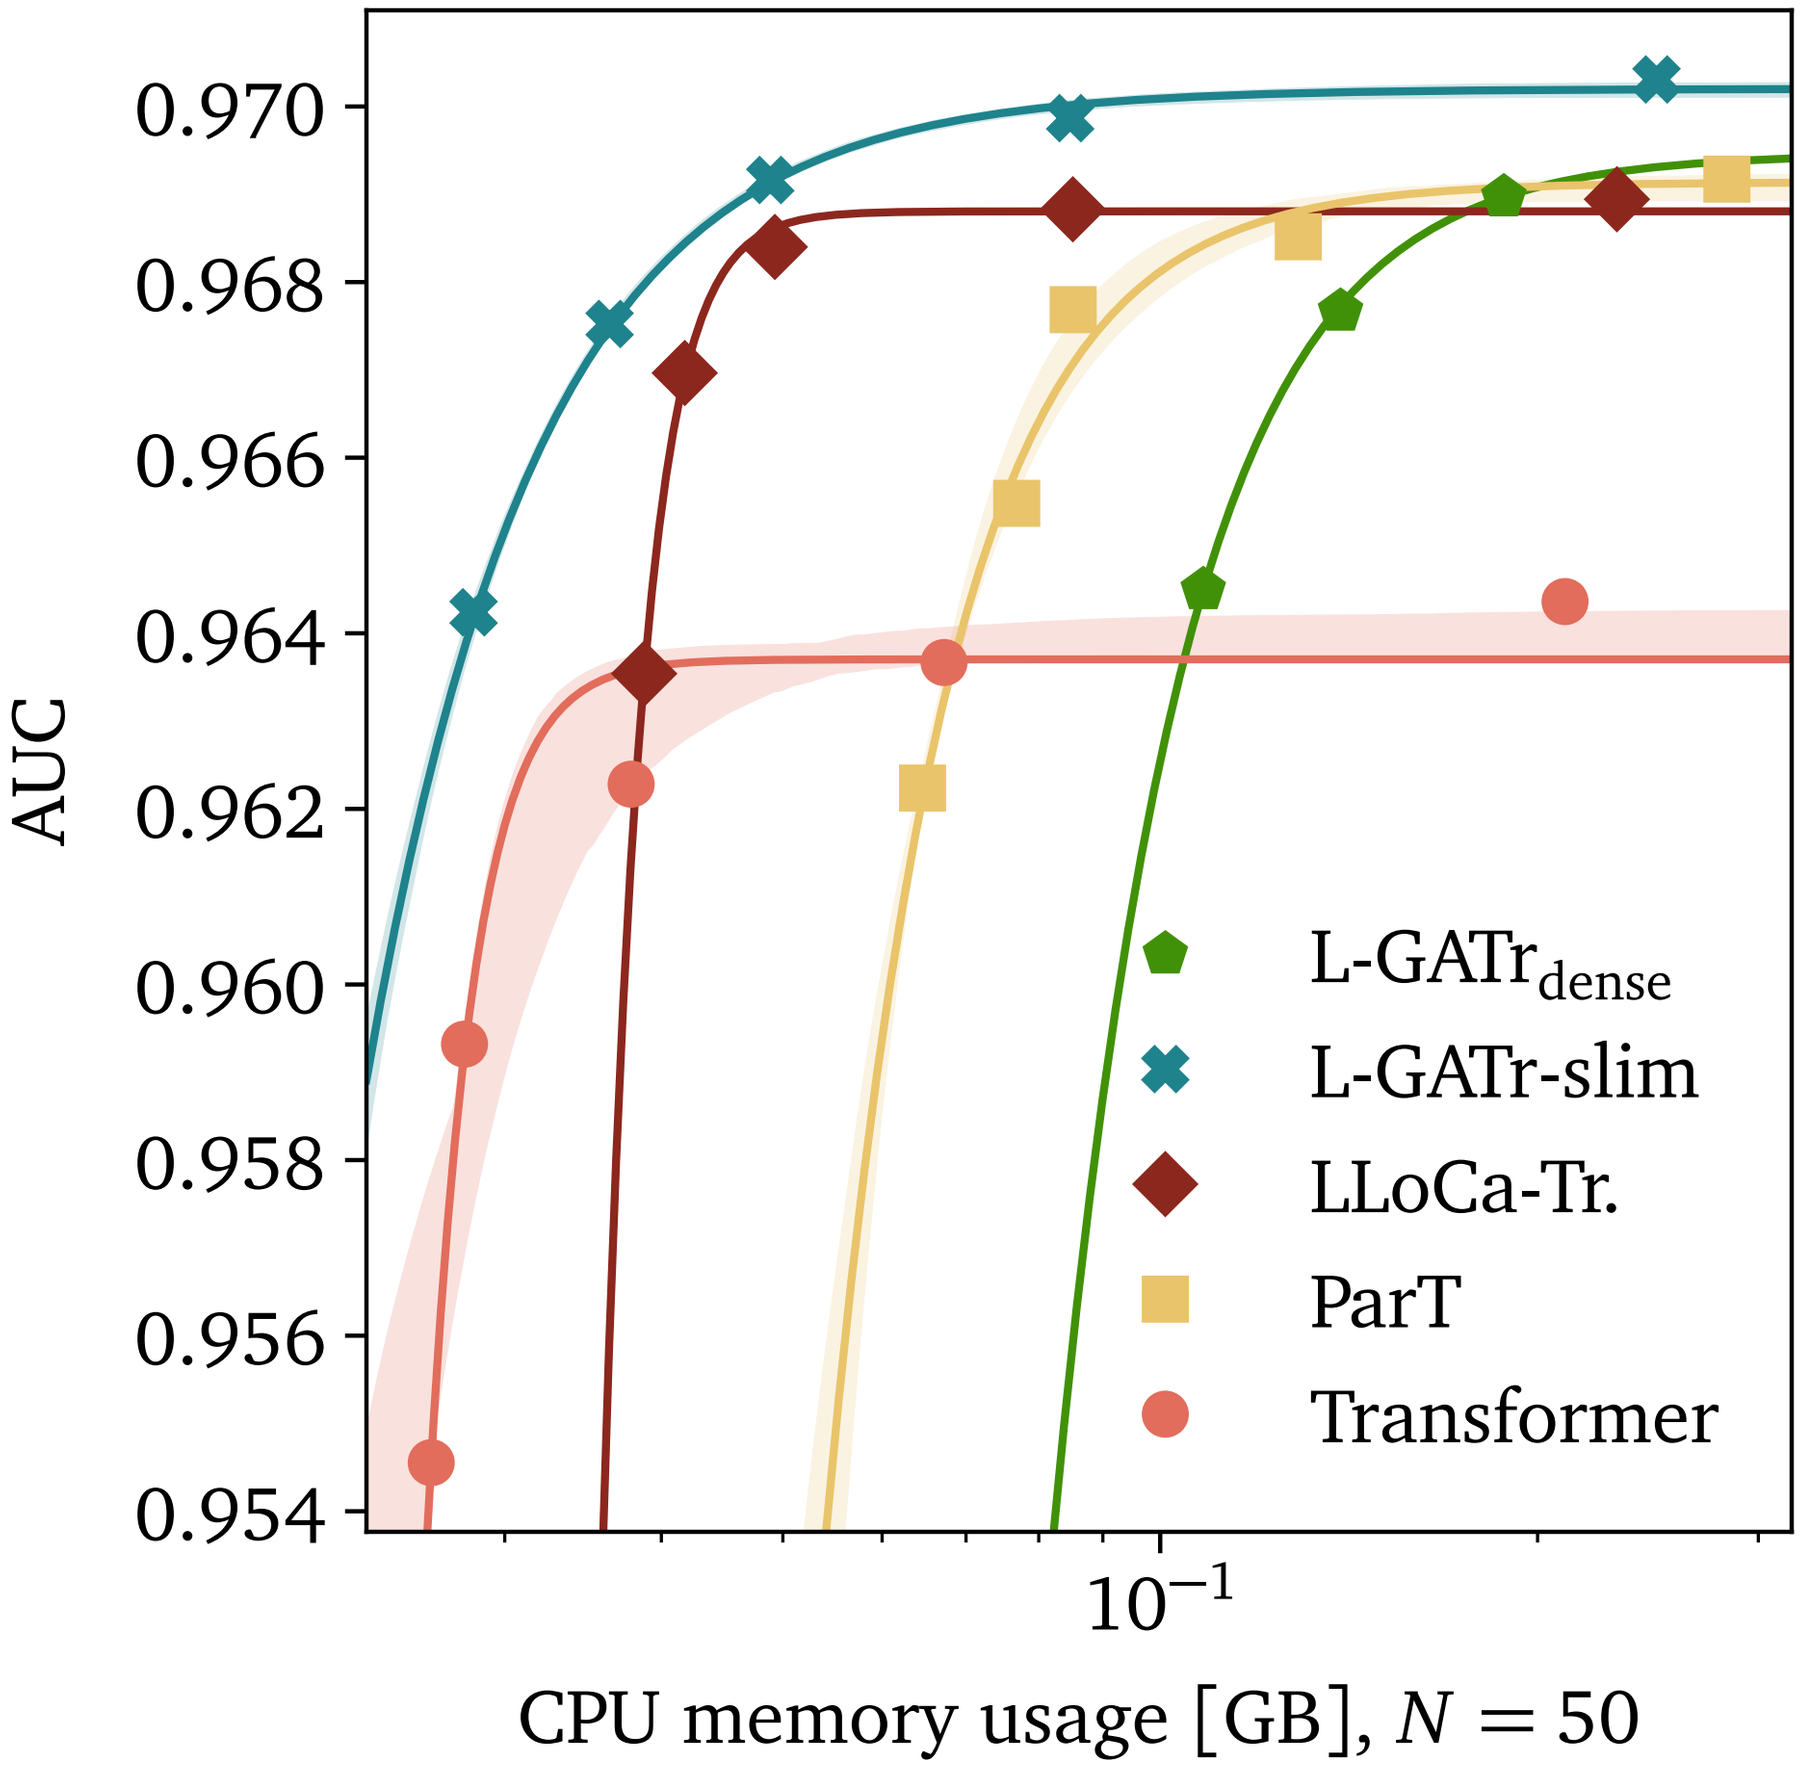
0.970
0.968
0.966
0.964
0.962
0.960
0.958
0.956
0.954
AUC
10^{-1}
CPU memory usage [GB], $N = 50$
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

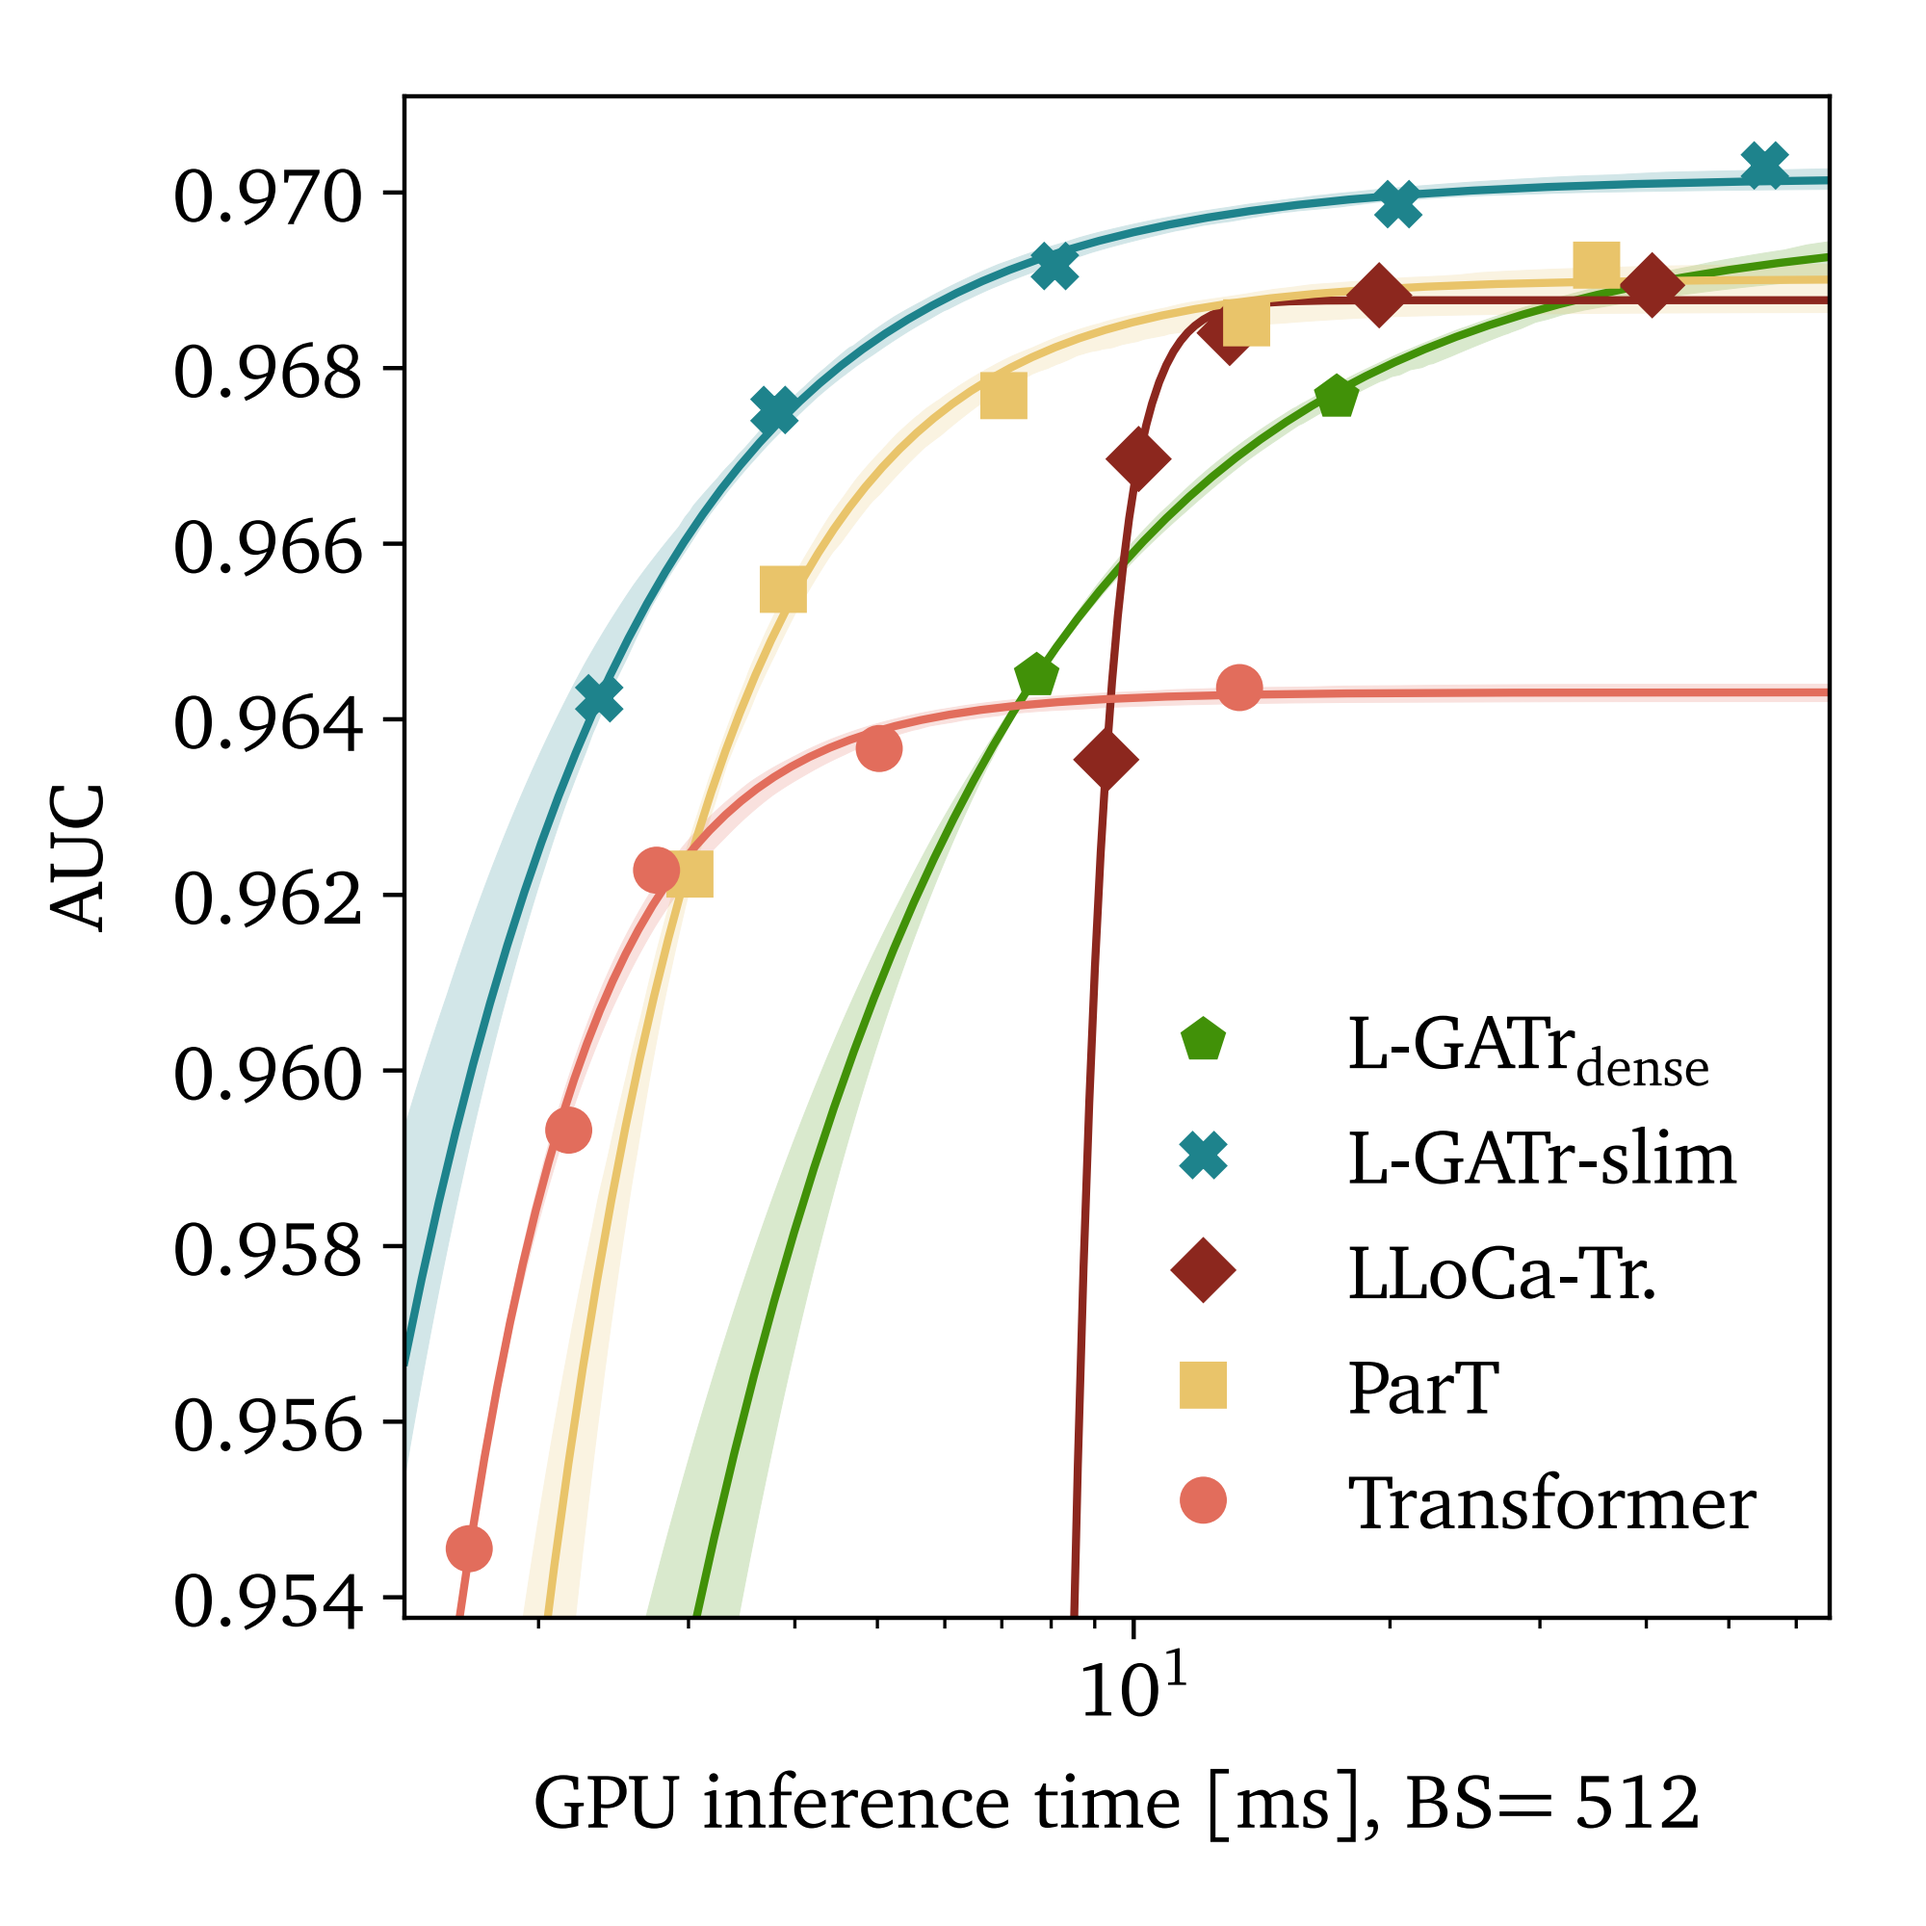
0.970
0.968
0.966
0.964
0.962
0.960
0.958
0.956
0.954
AUC
10^1
GPU inference time [ms], BS= 512
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

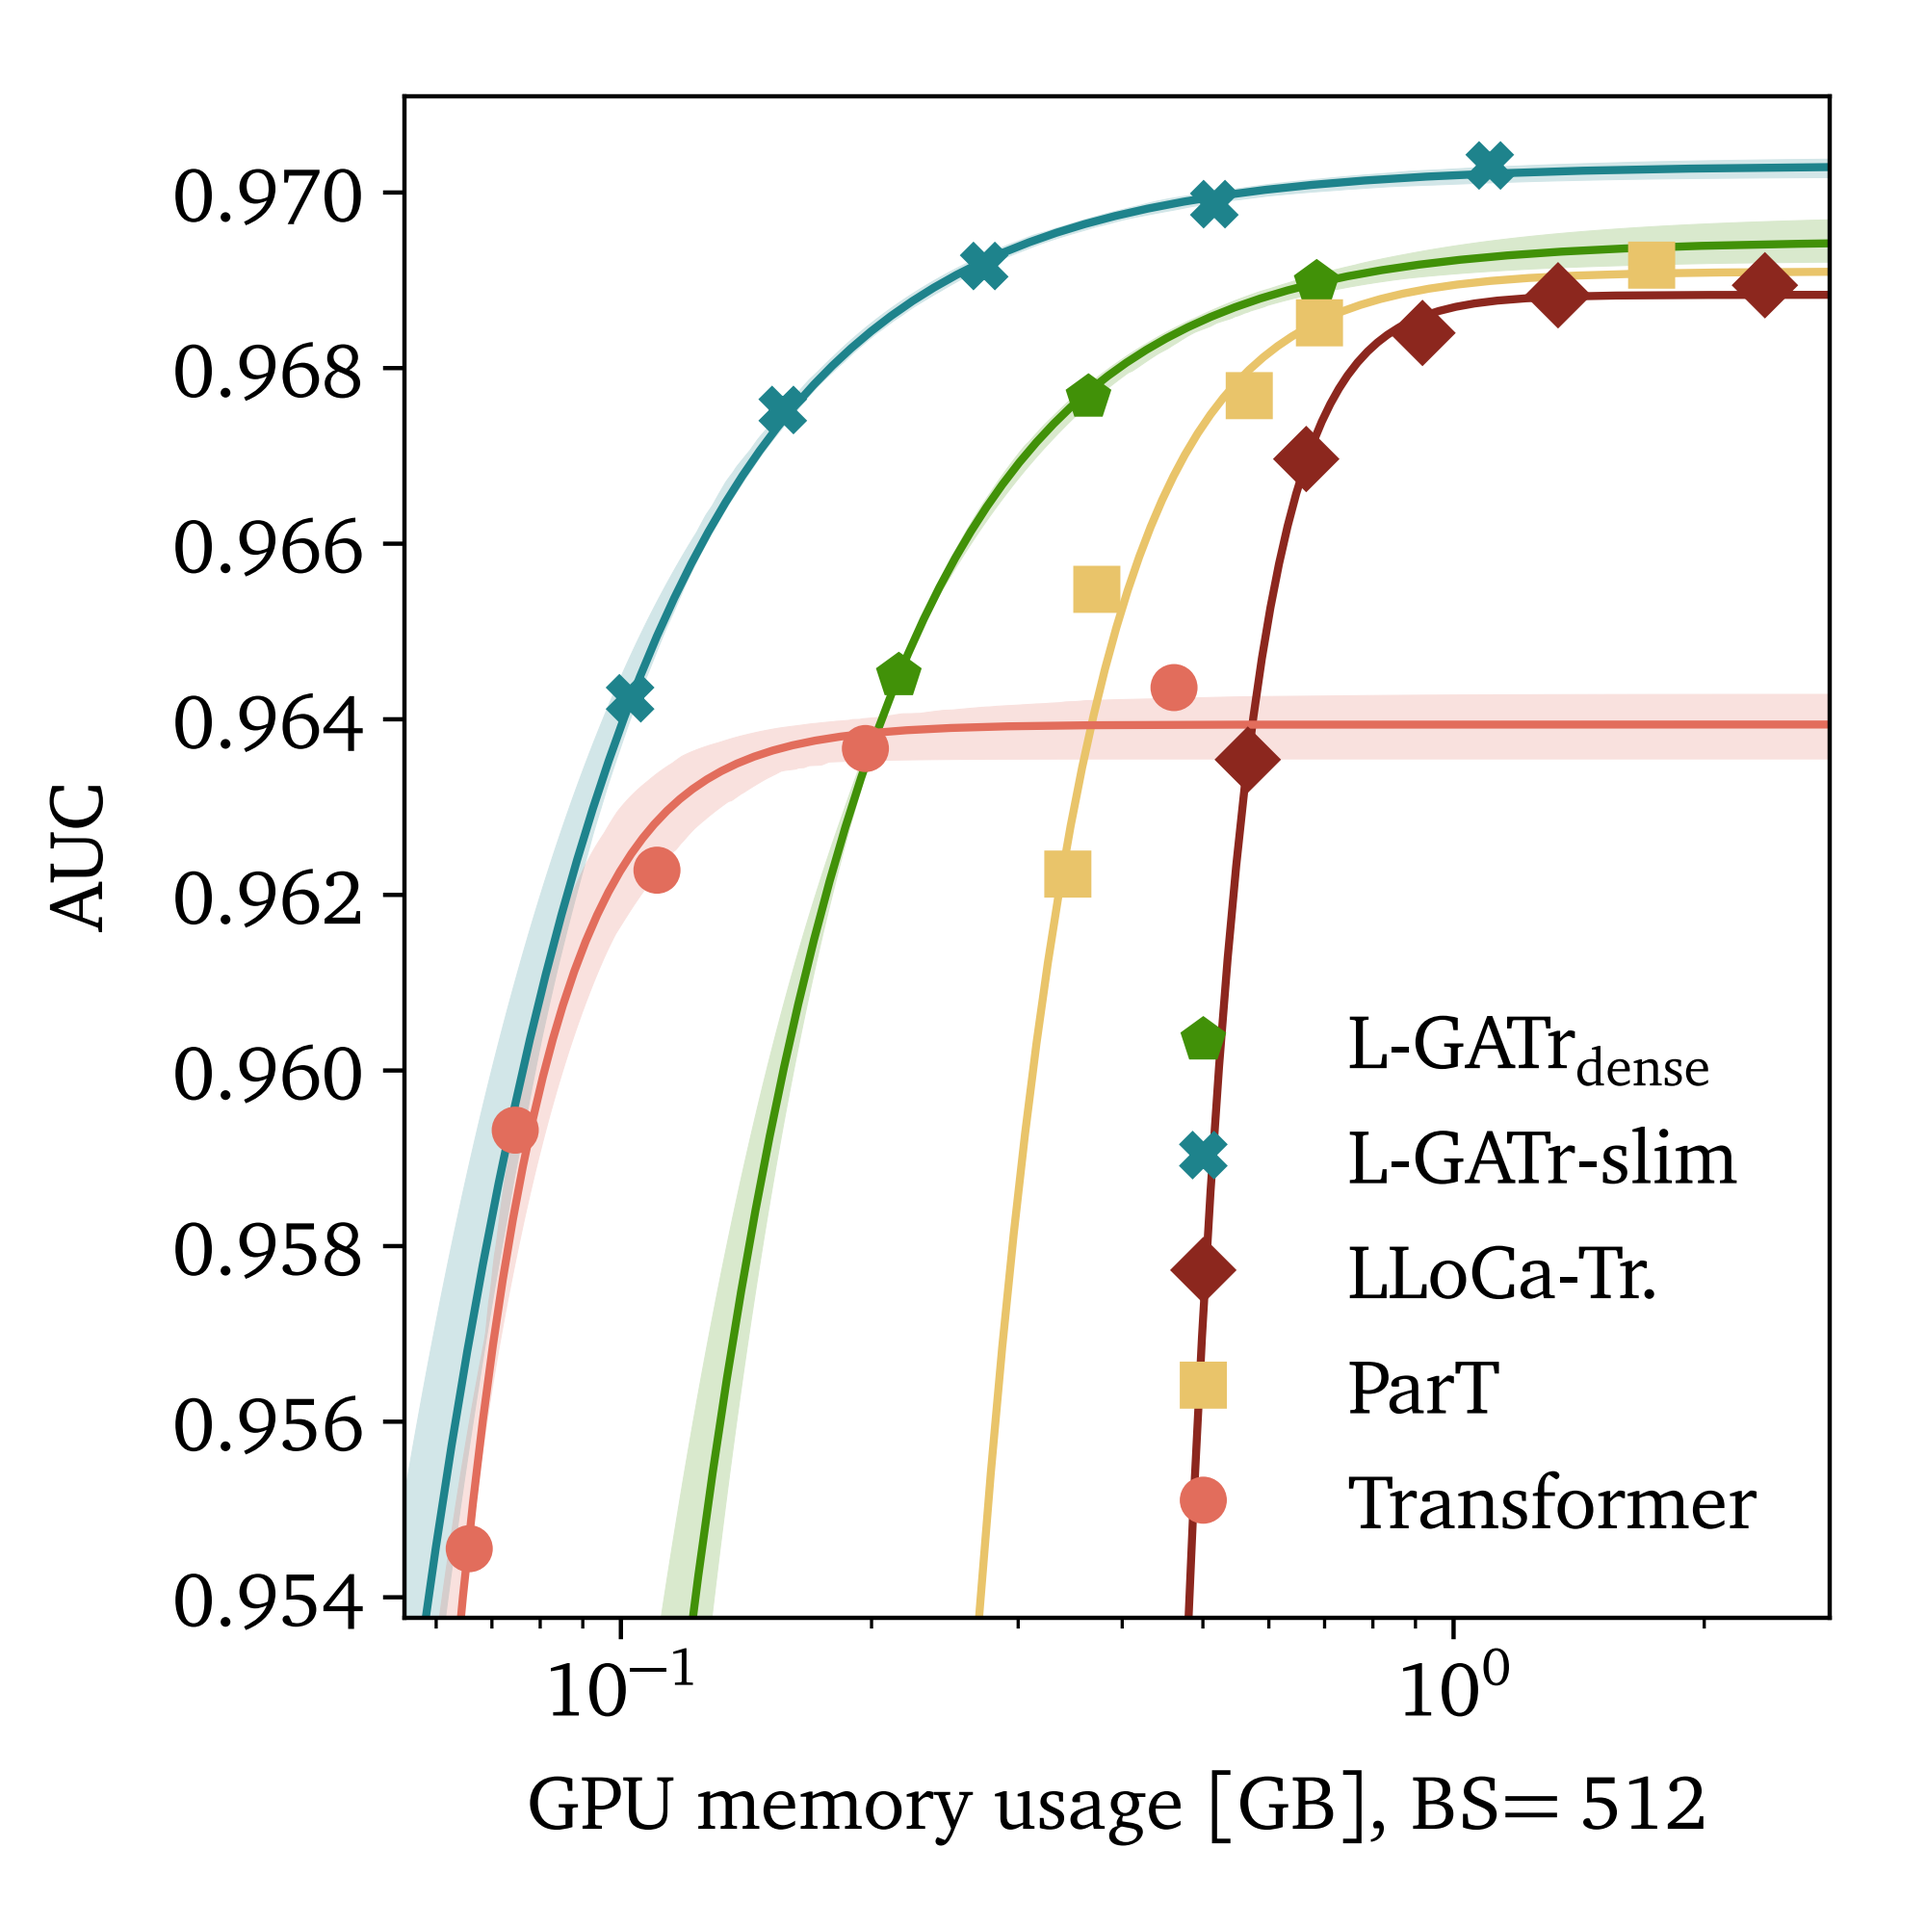
0.970
0.968
0.966
0.964
0.962
0.960
0.958
0.956
0.954
AUC
10^{-1}
10^{0}
GPU memory usage [GB], BS= 512
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

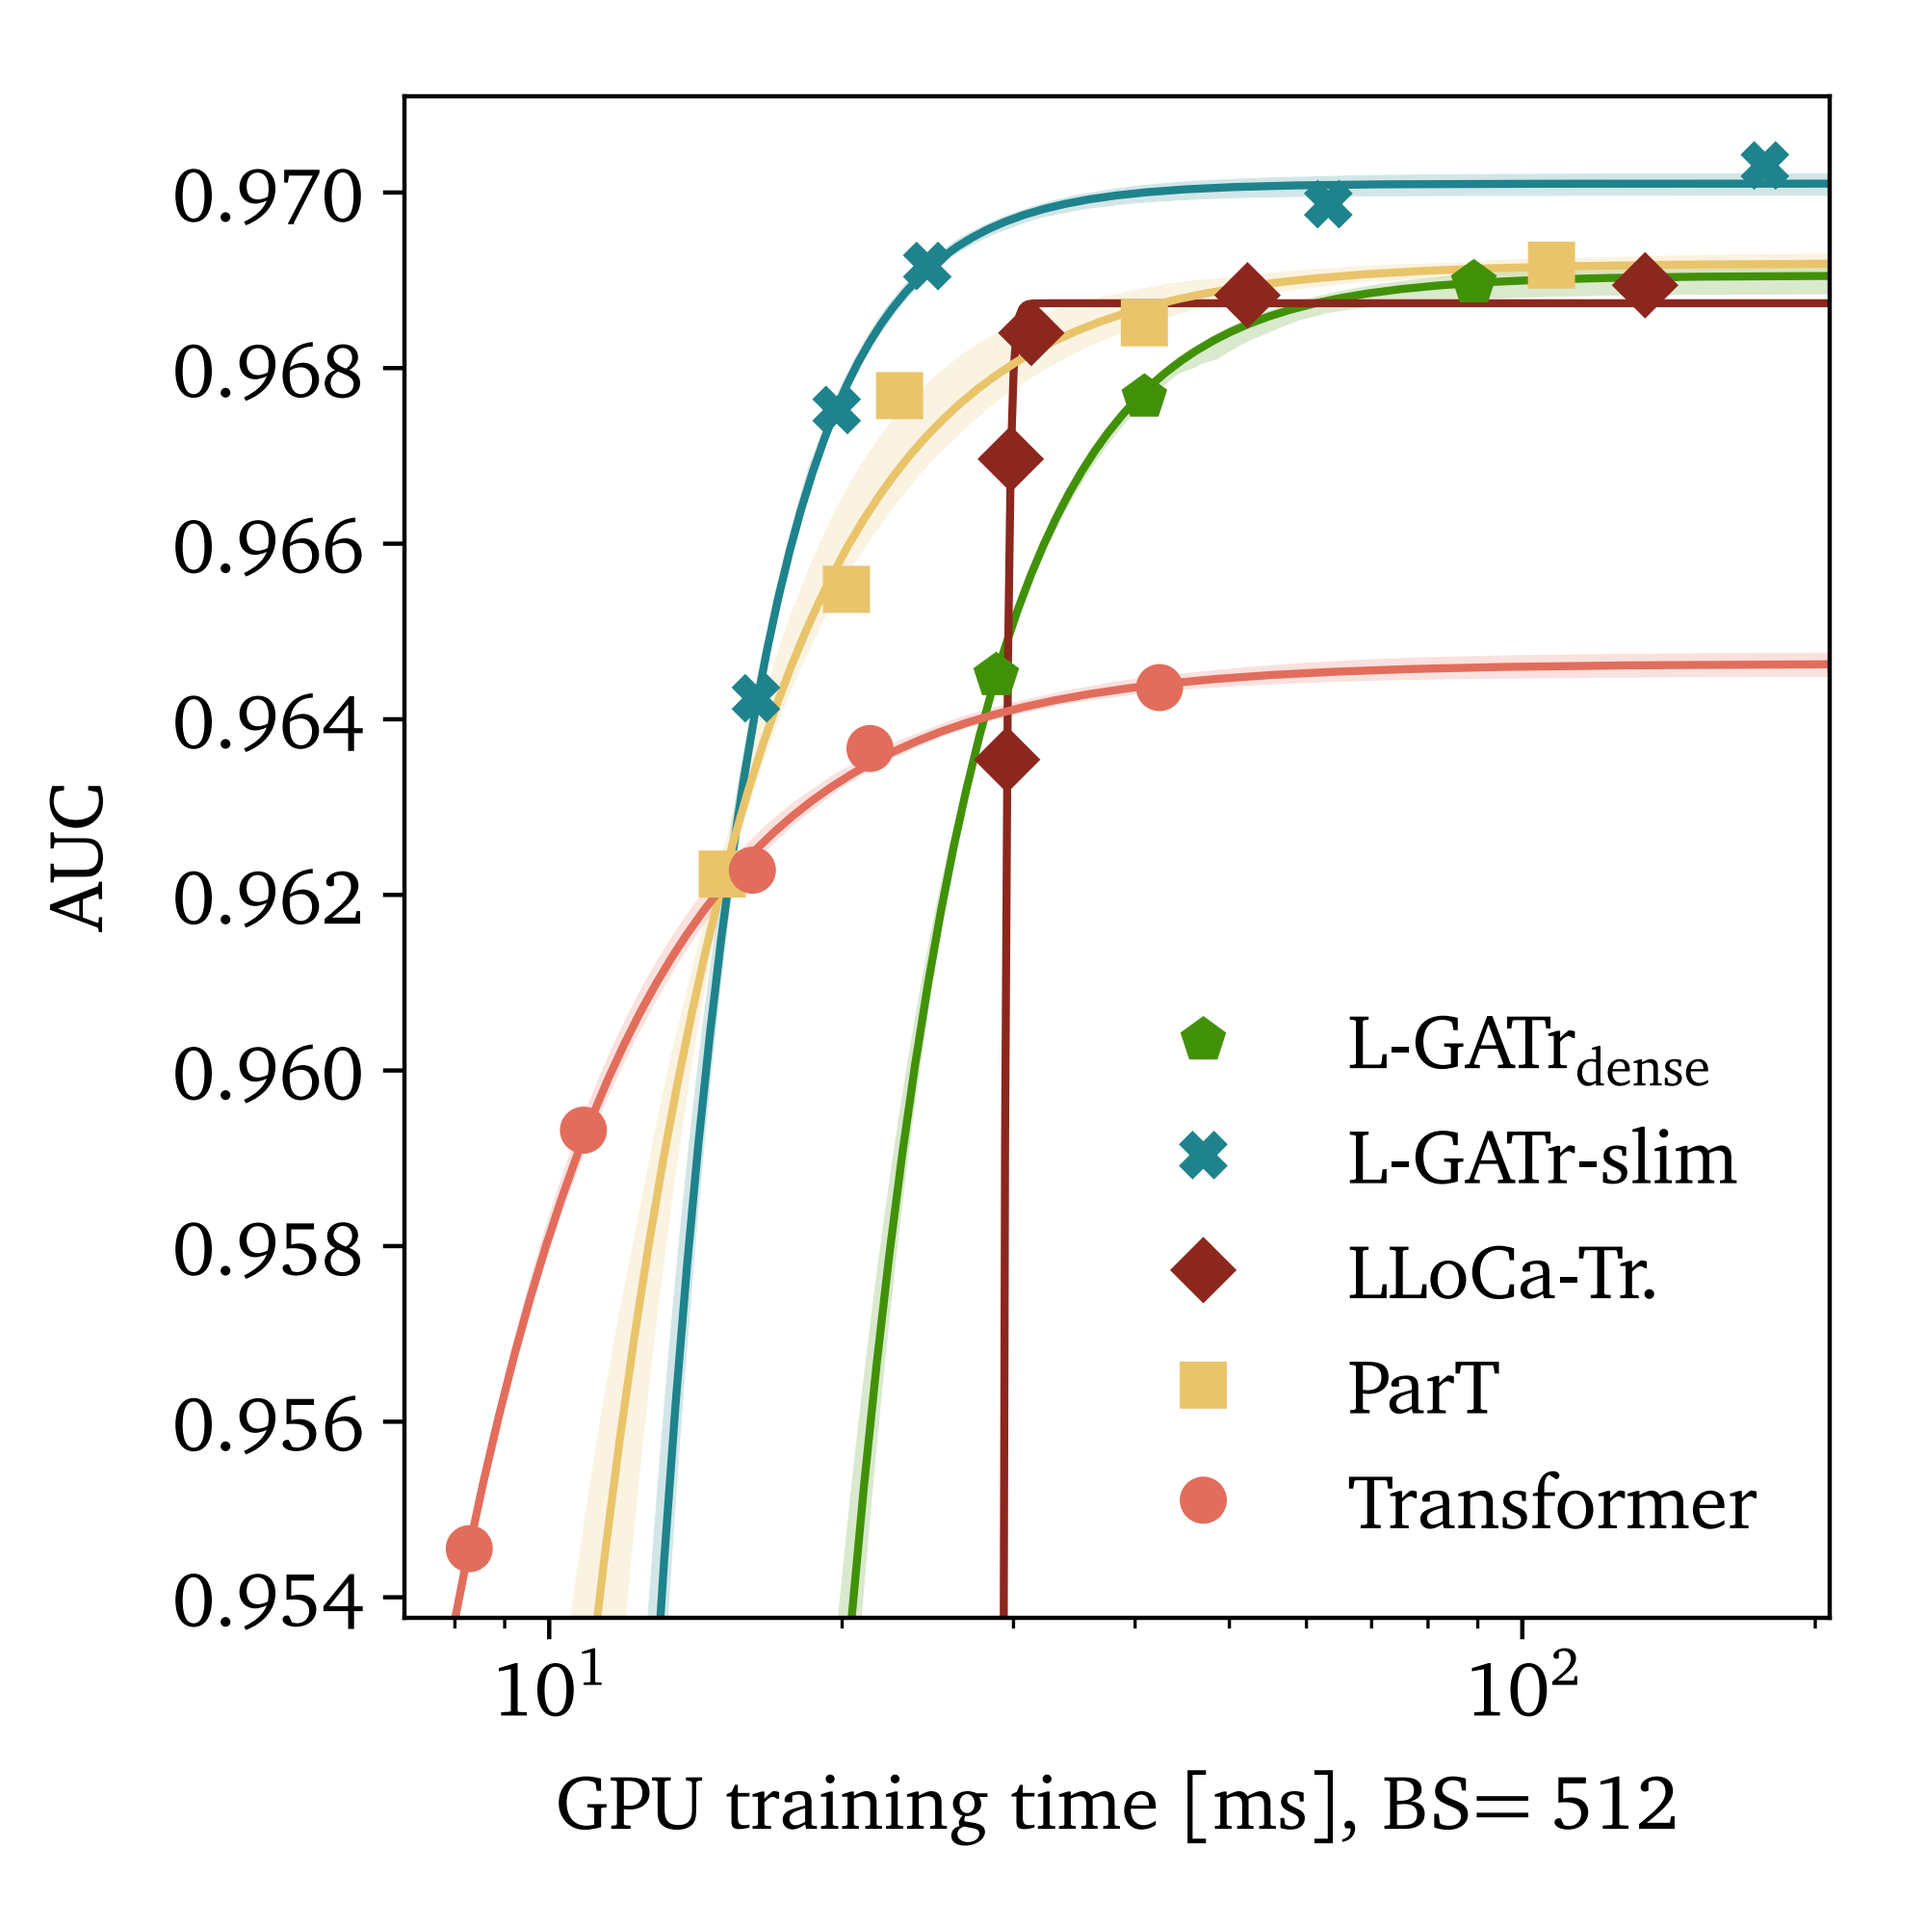
0.970
0.968
0.966
0.964
0.962
0.960
0.958
0.956
0.954
AUC
10^1
10^2
GPU training time [ms], BS= 512
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer