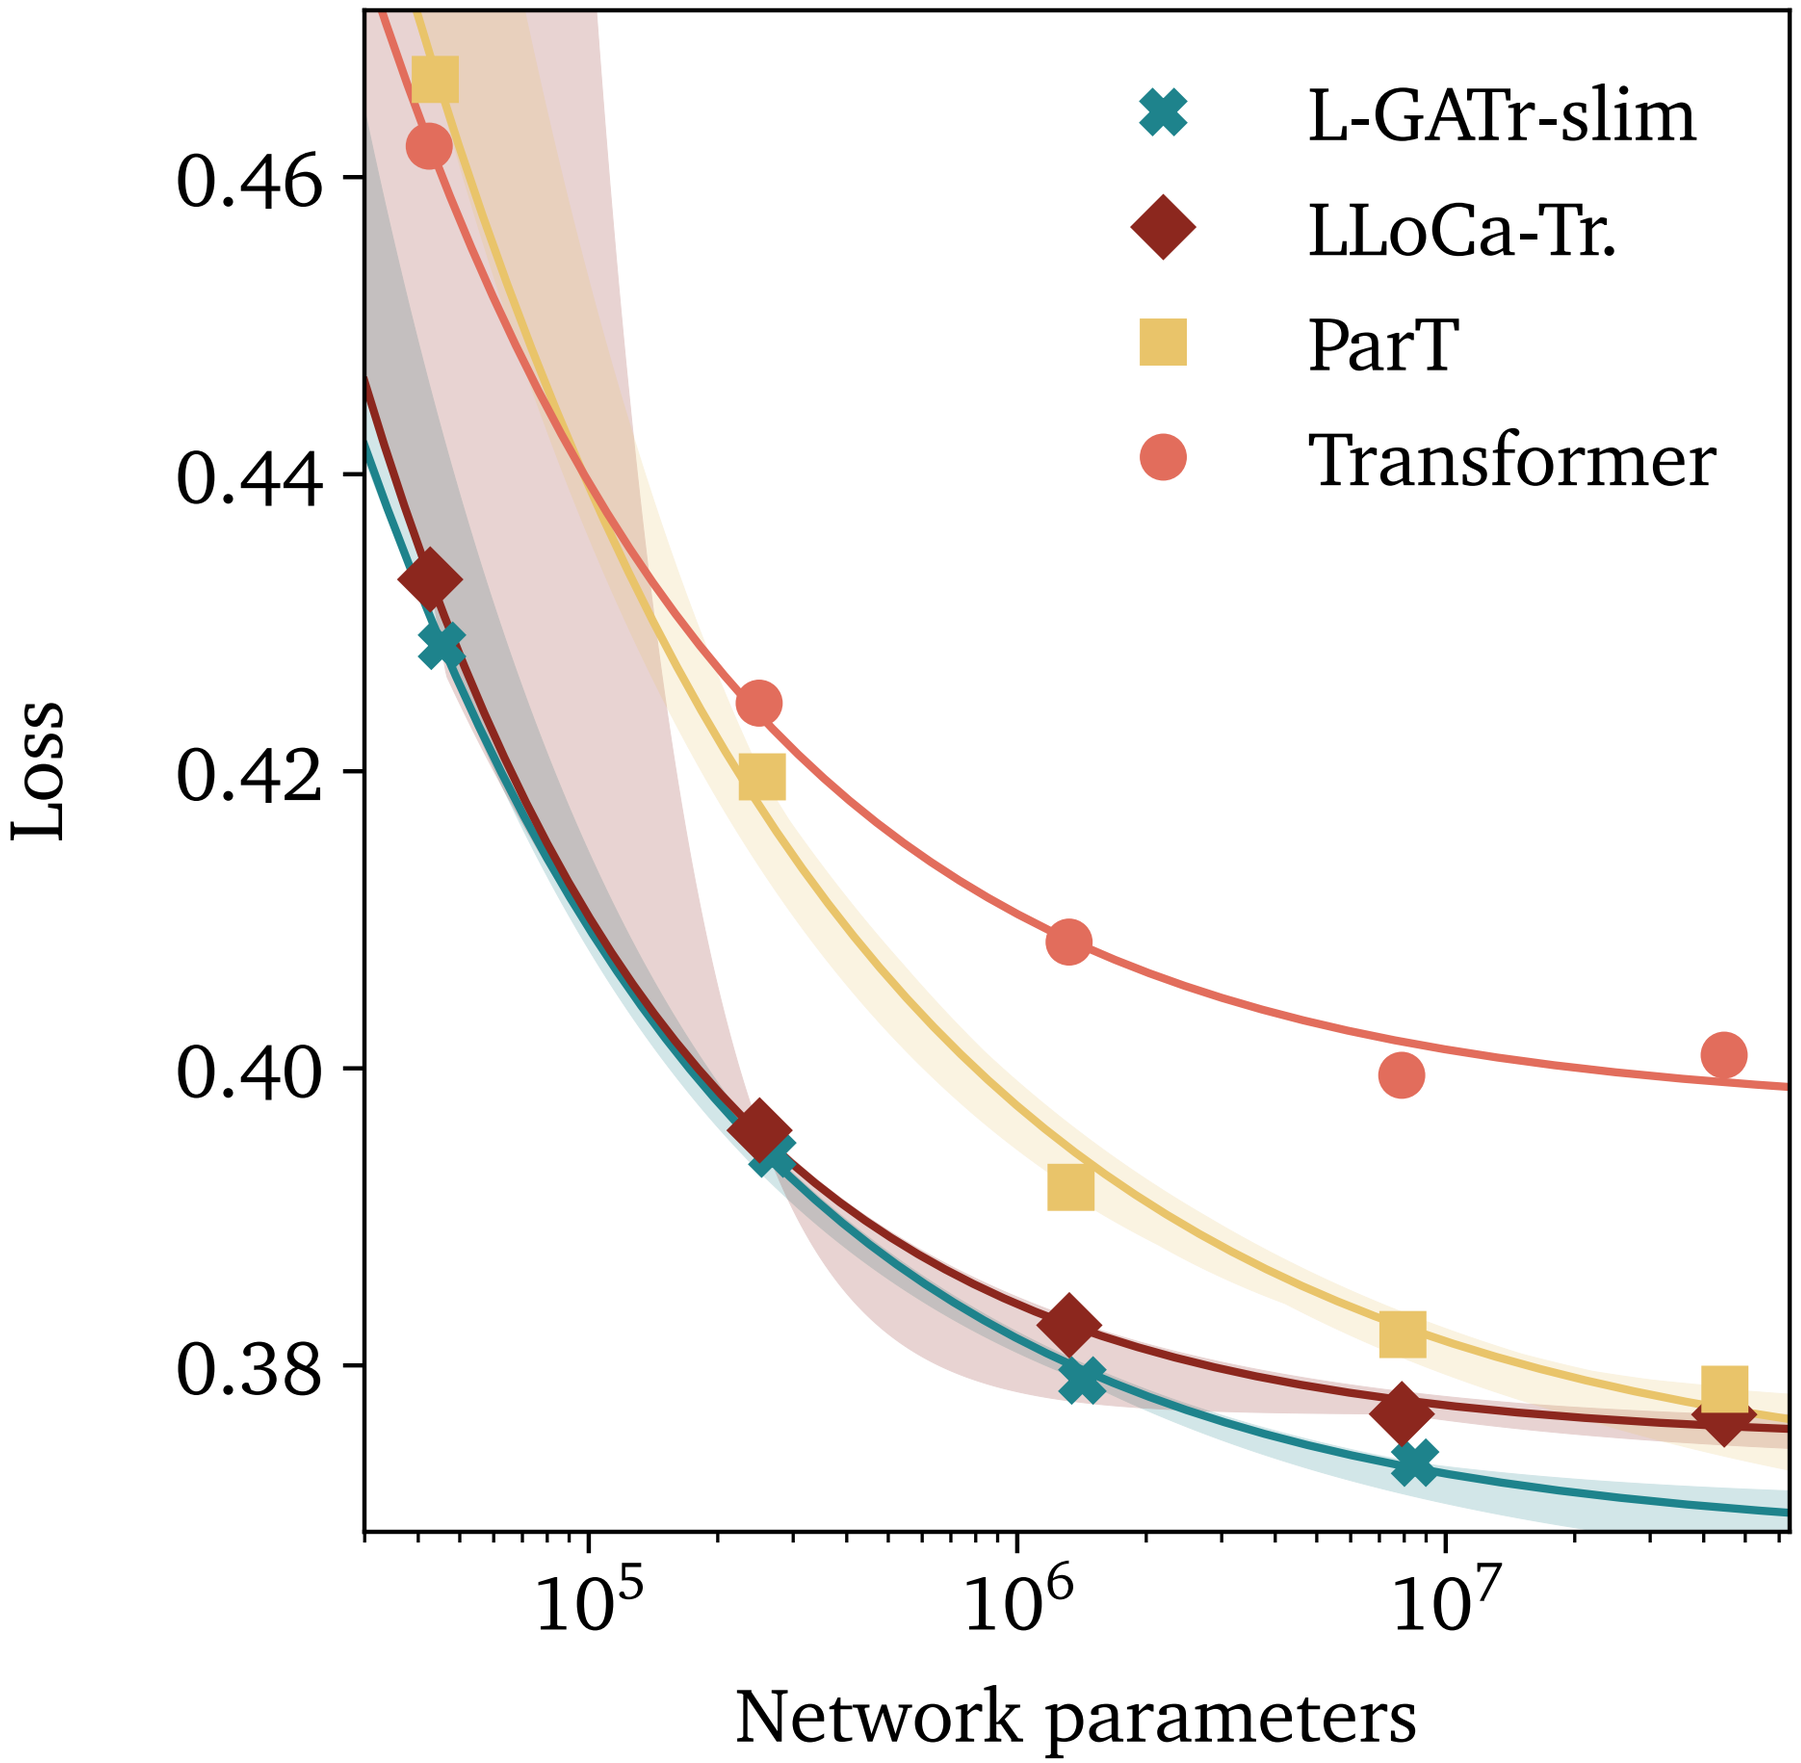
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.46
0.44
0.42
0.40
0.38
10^5
10^6
10^7
Network parameters

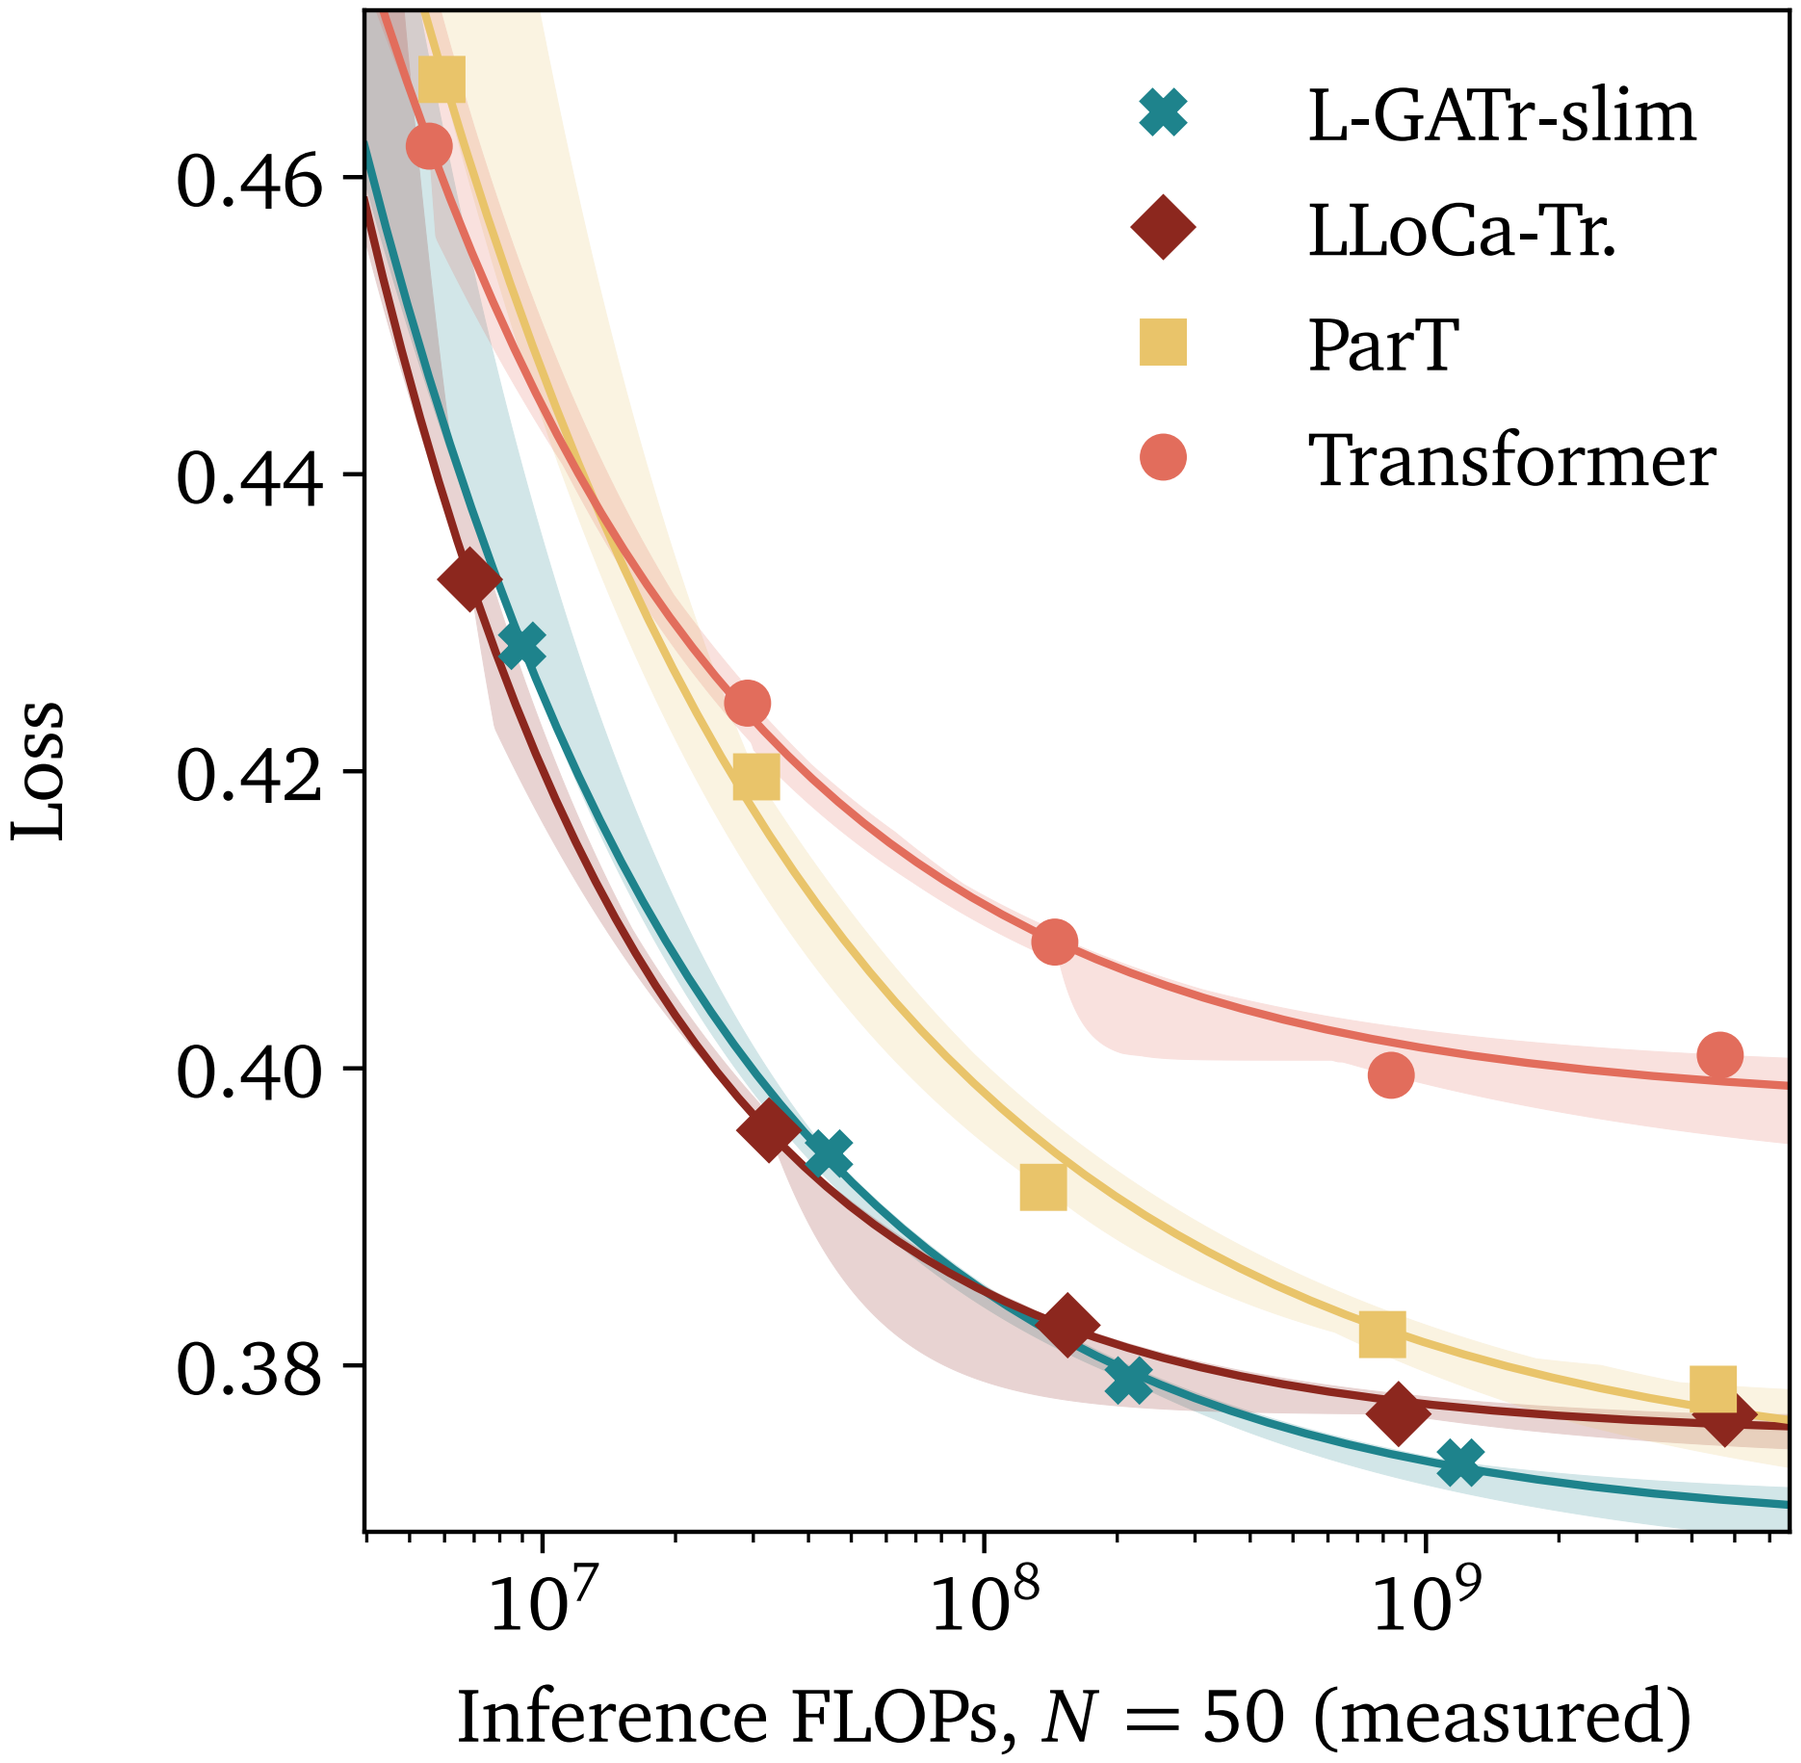
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.46
0.44
0.42
0.40
0.38
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (measured)

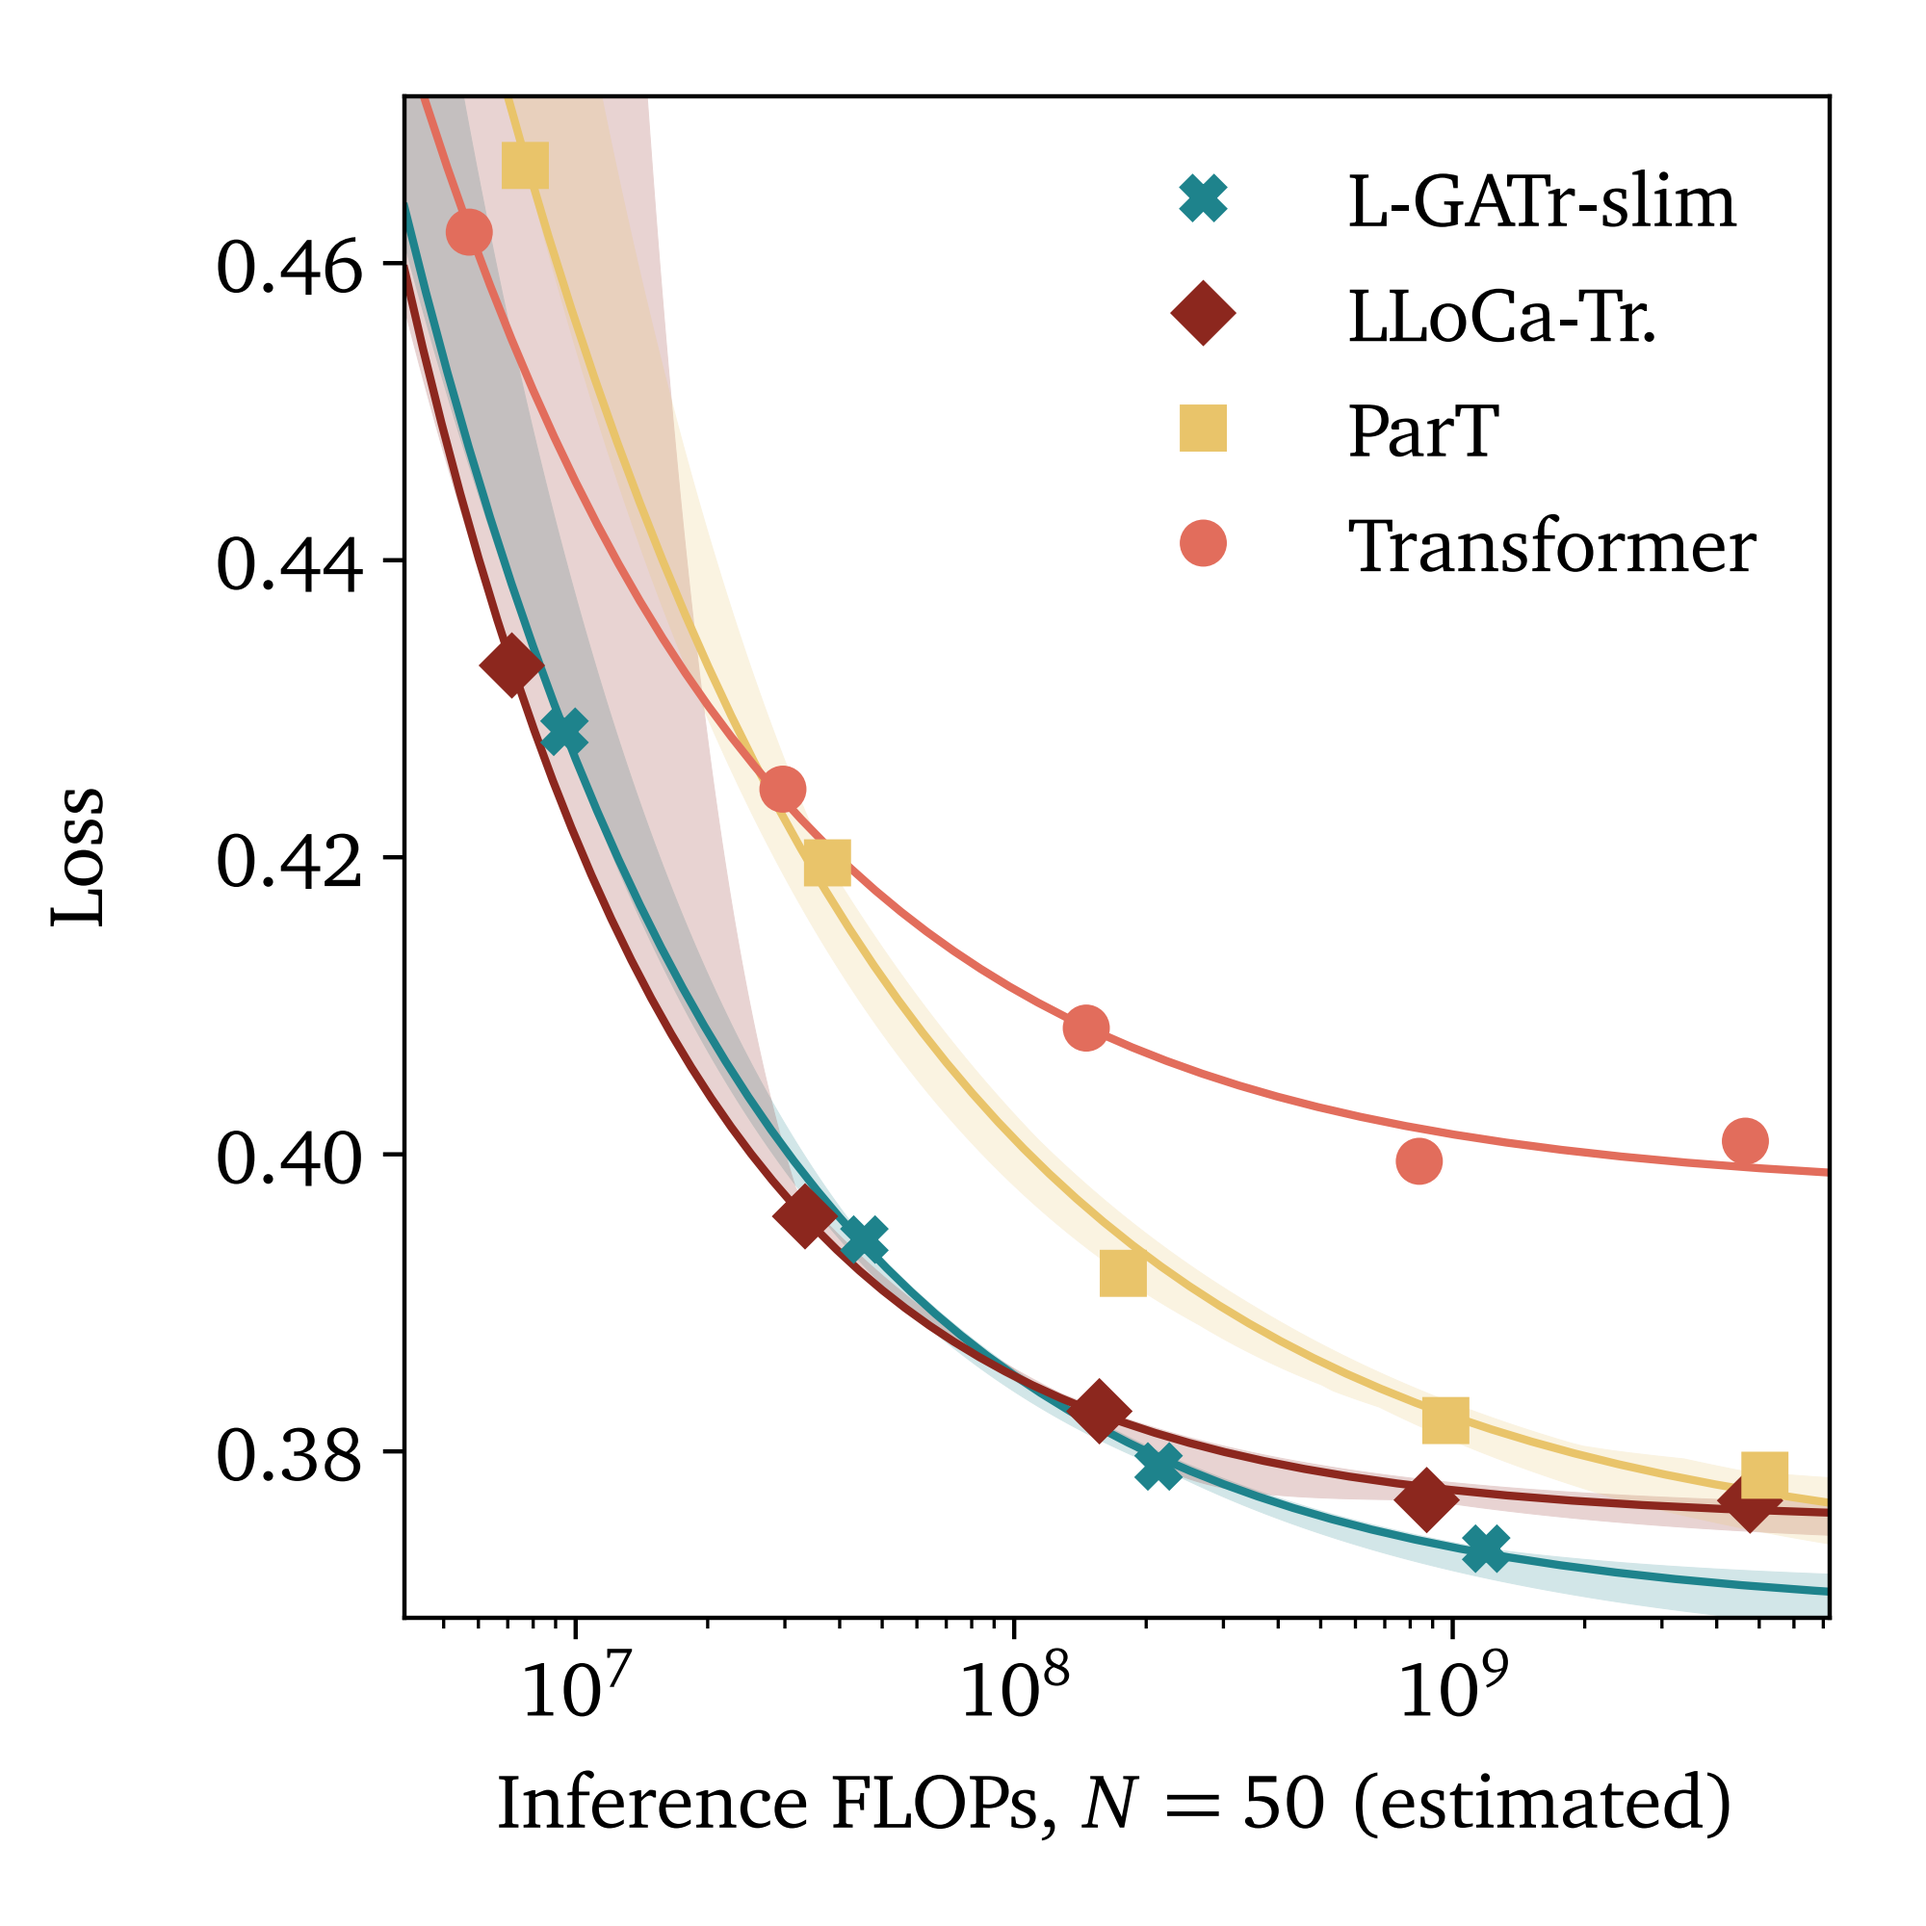
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.46
0.44
0.42
0.40
0.38
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (estimated)

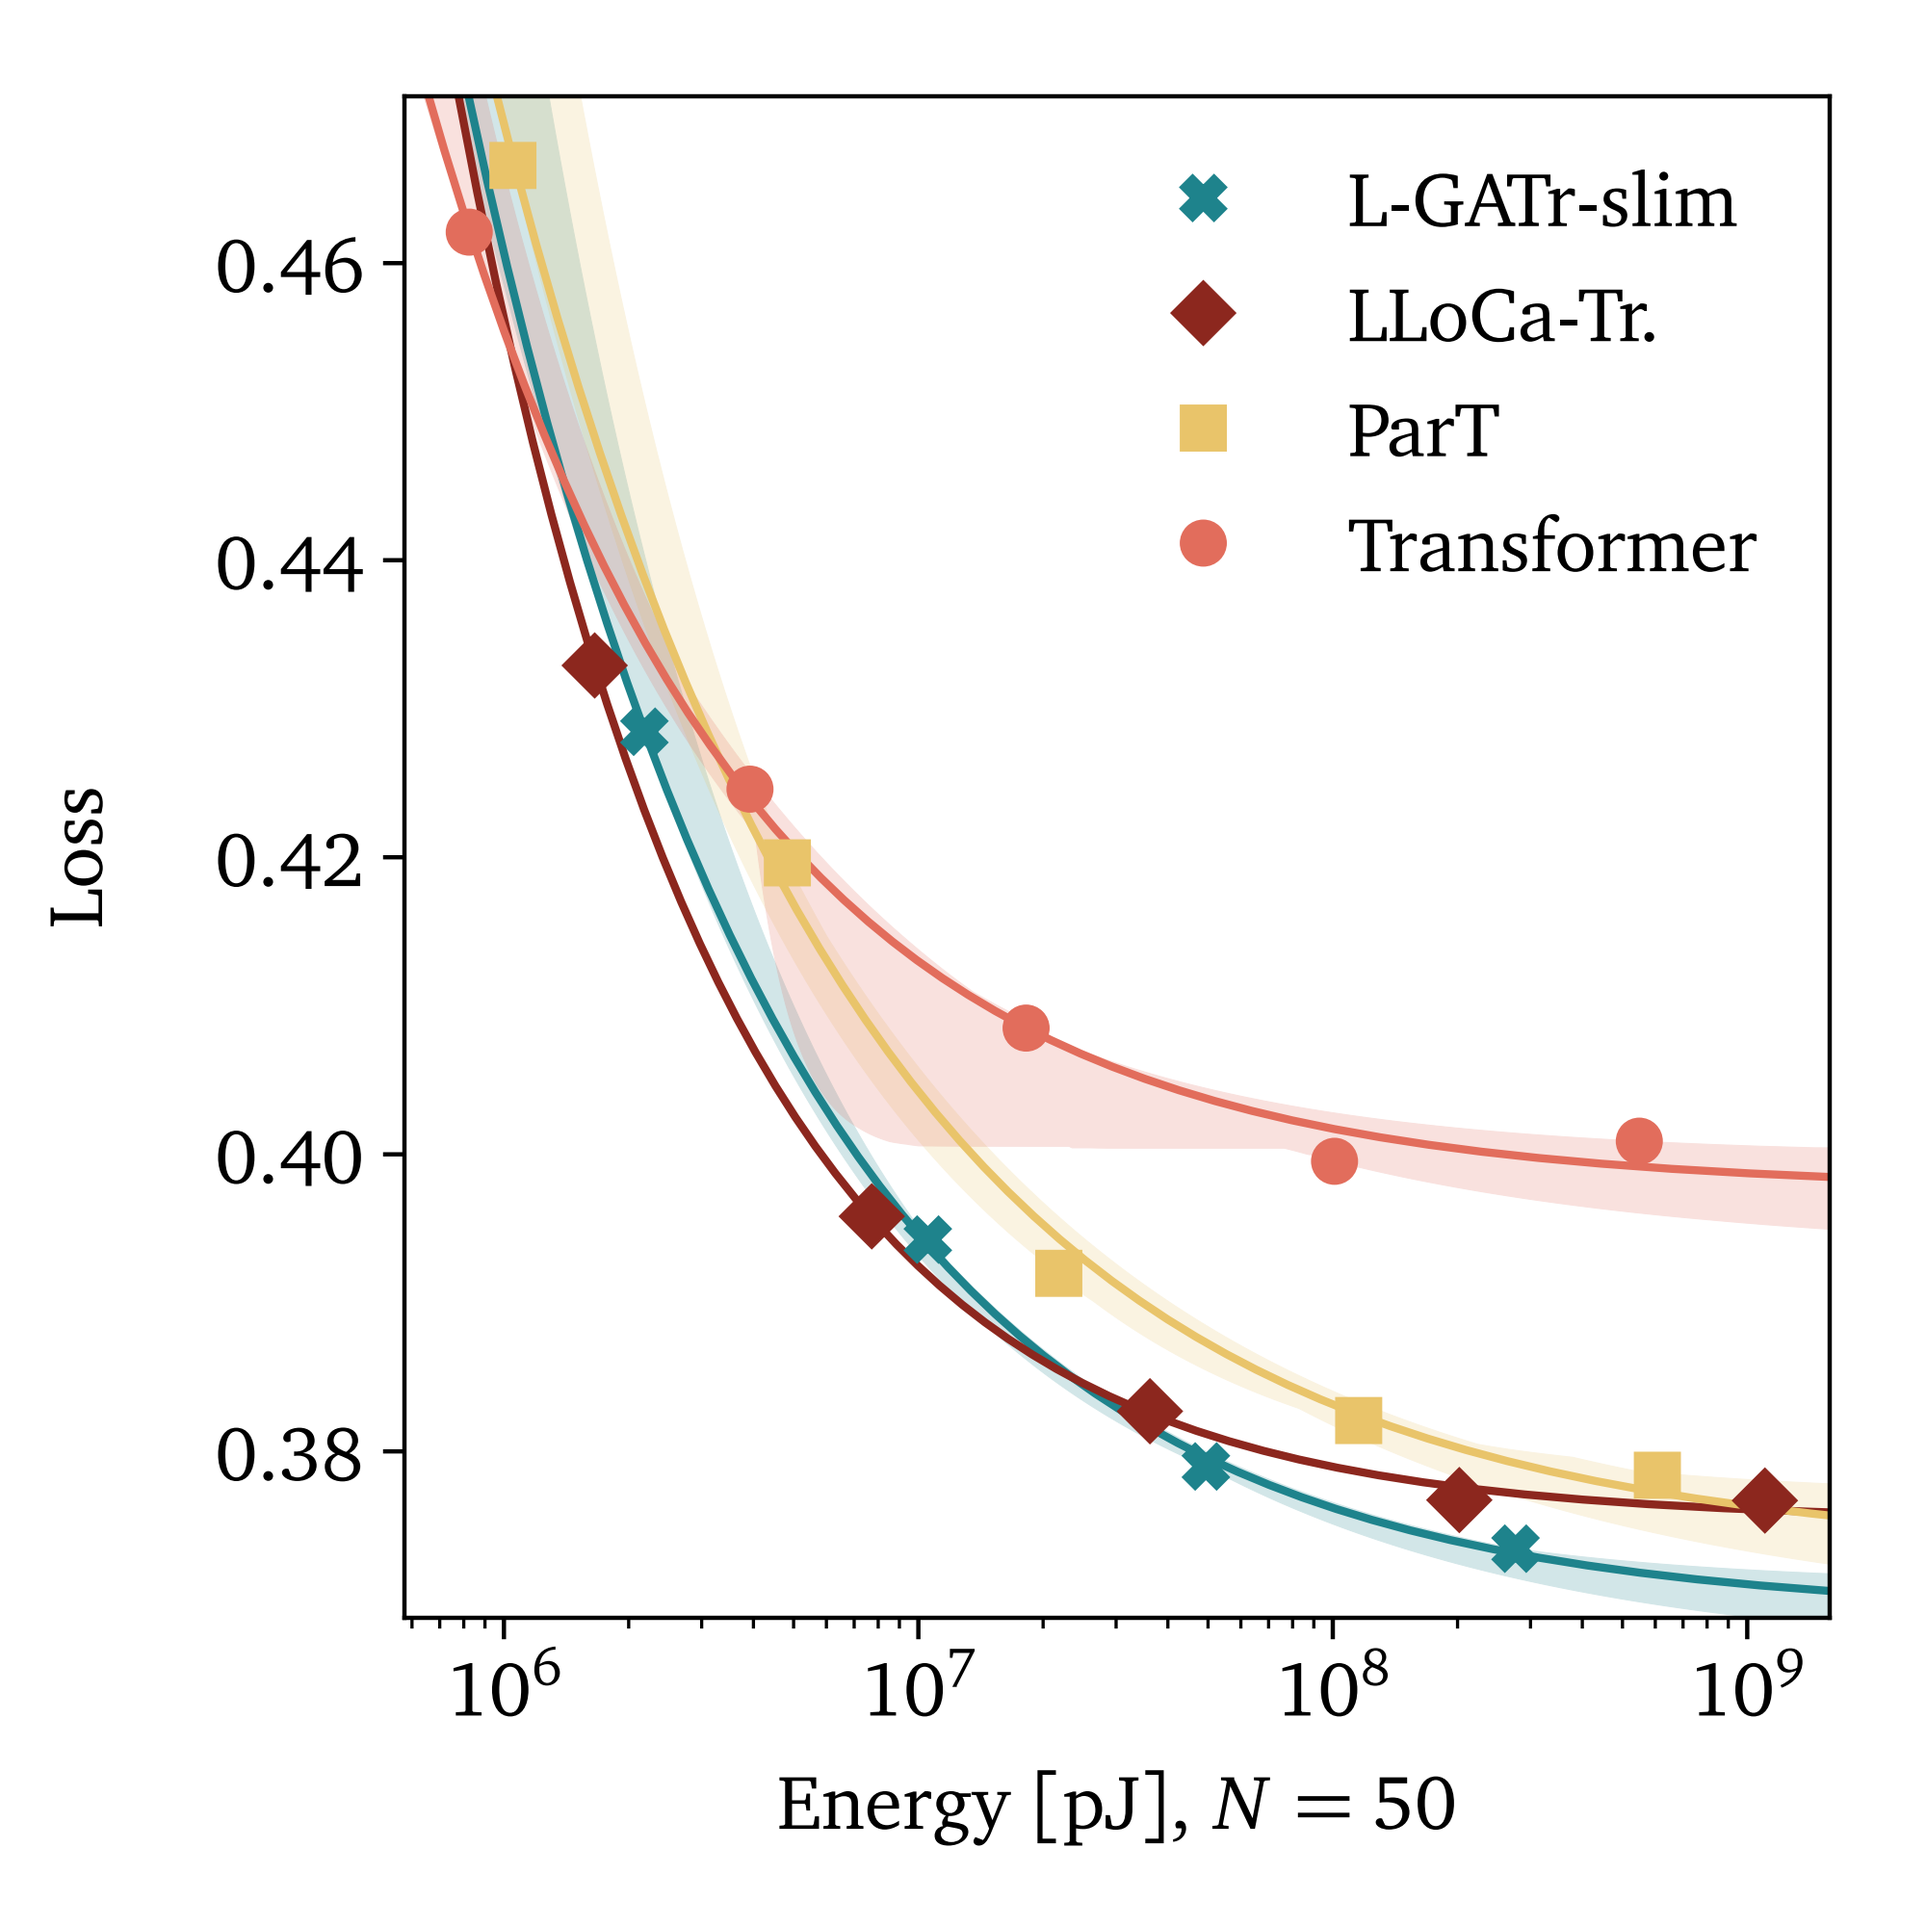
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.46
0.44
0.42
0.40
0.38
10^6
10^7
10^8
10^9
Energy [pJ], $N = 50$

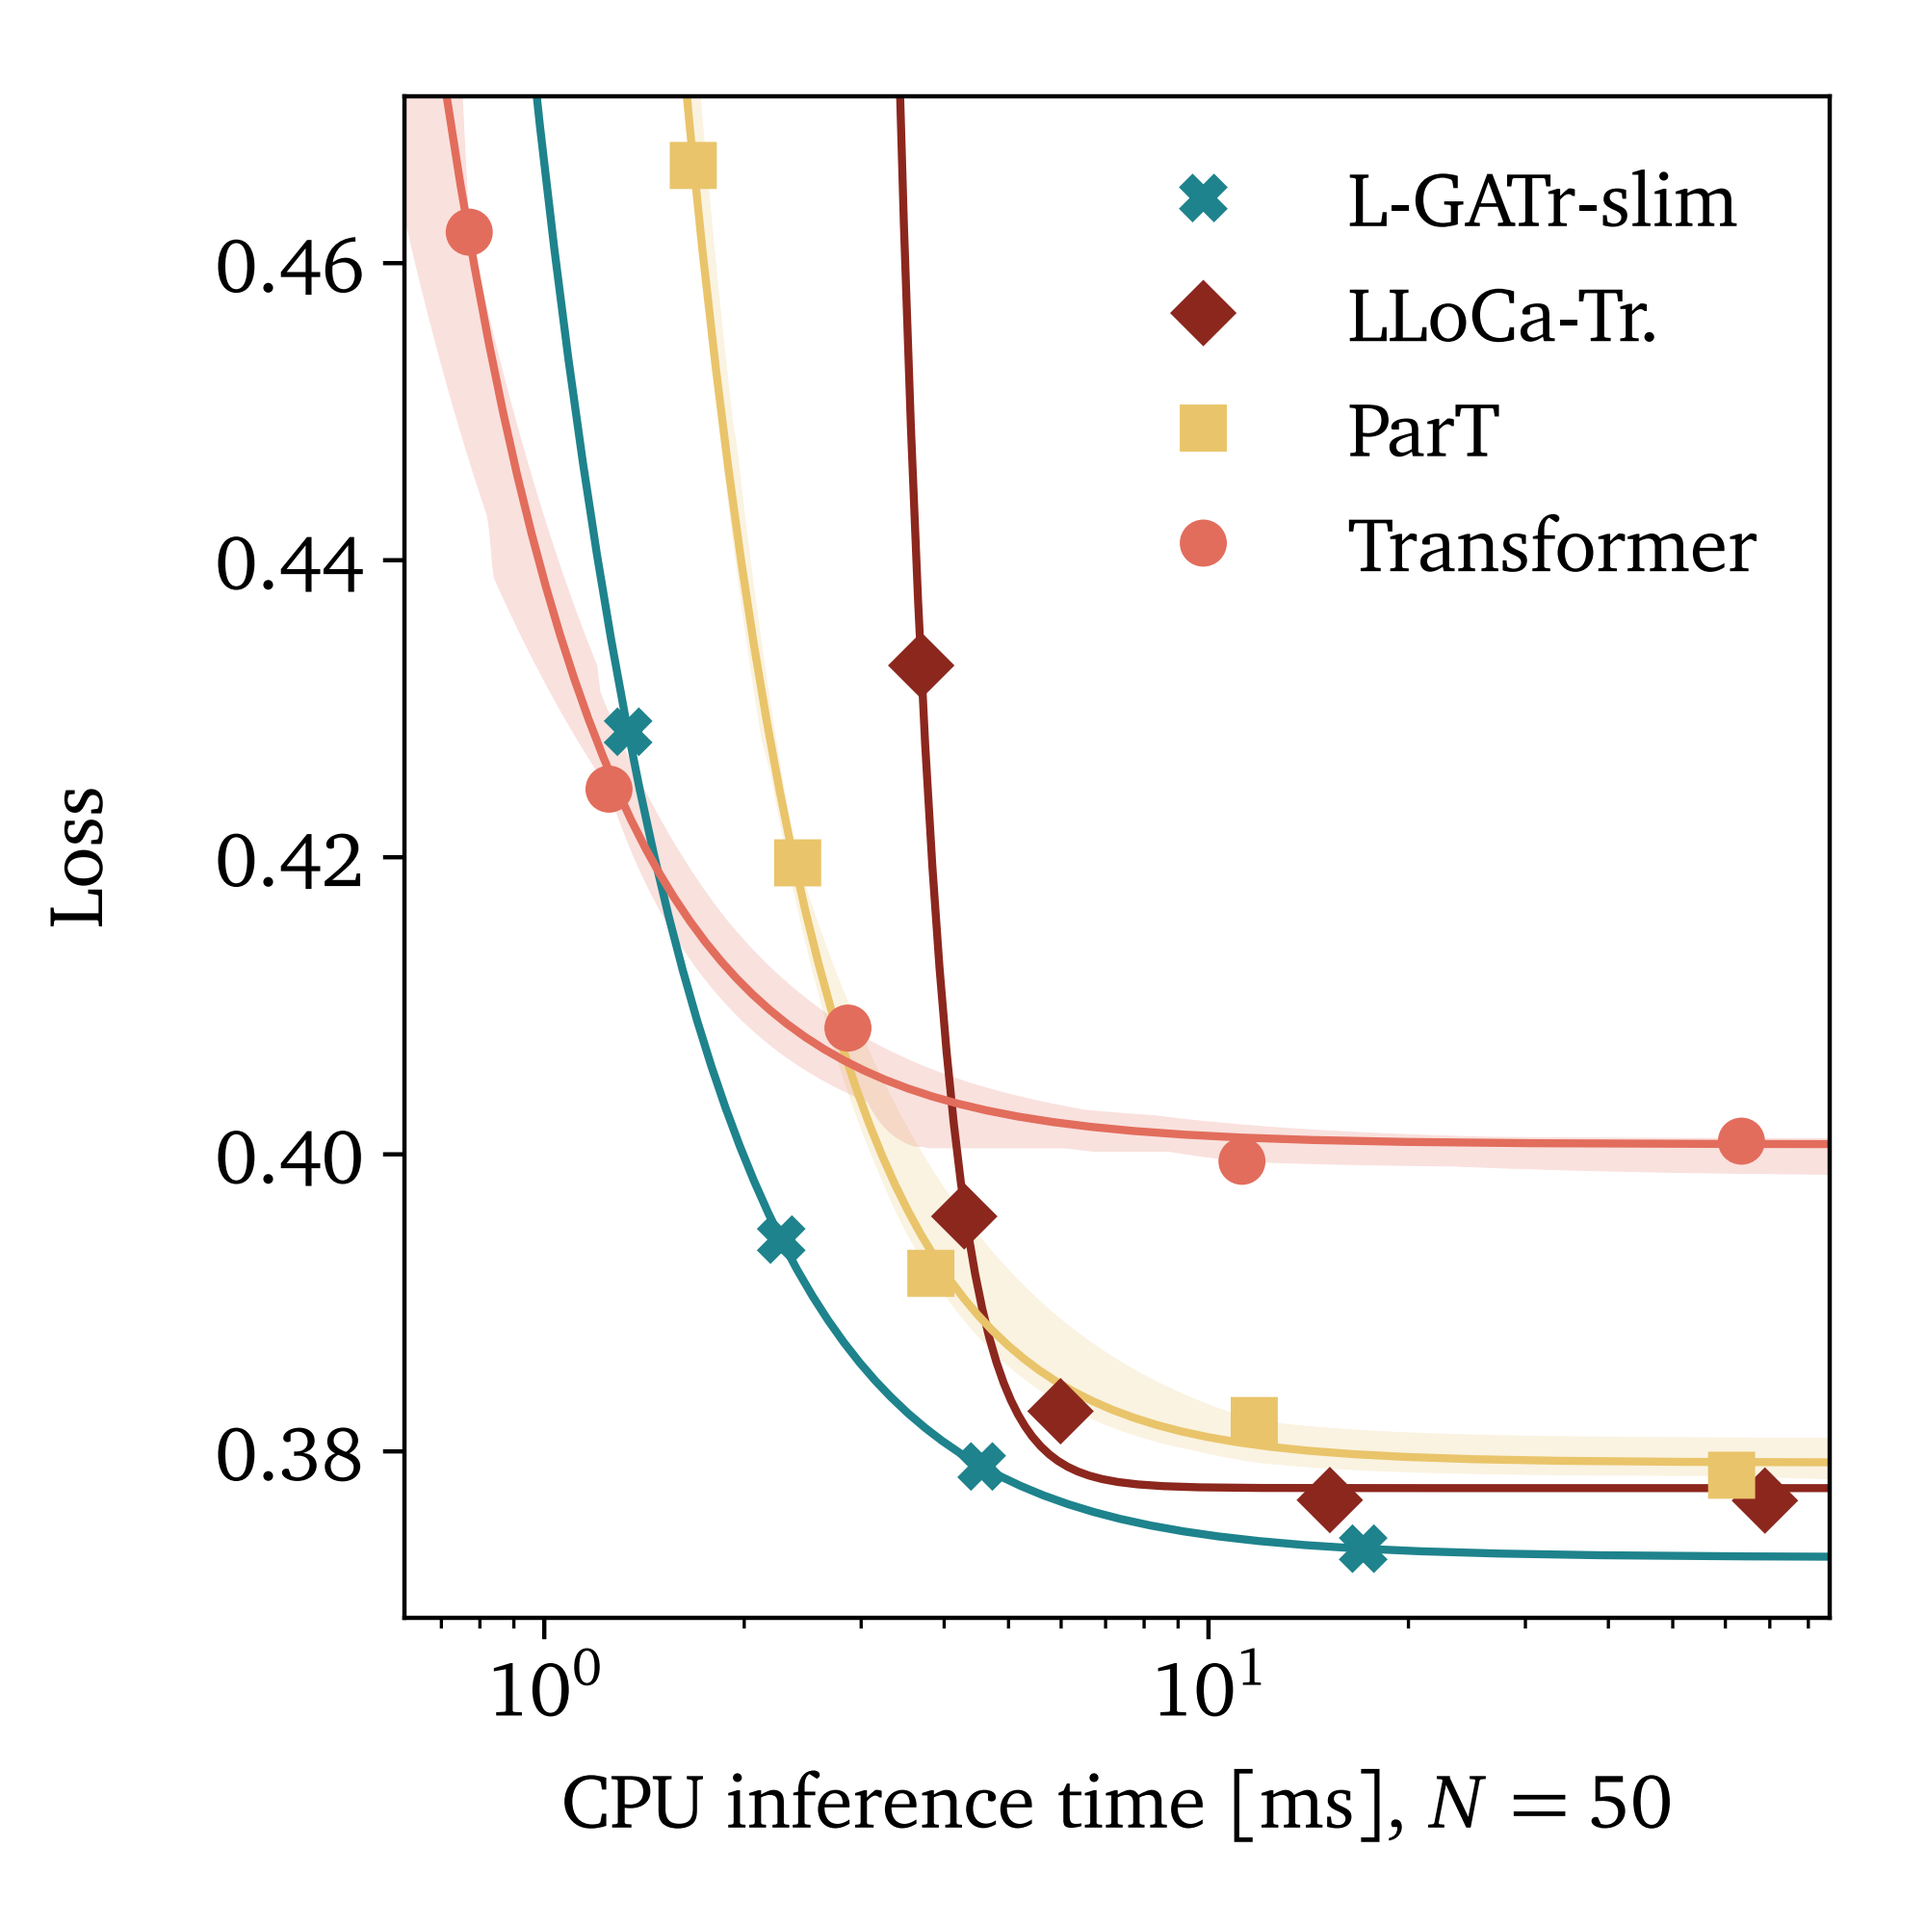
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.46
0.44
0.42
0.40
0.38
10^0
10^1
CPU inference time [ms], $N = 50$

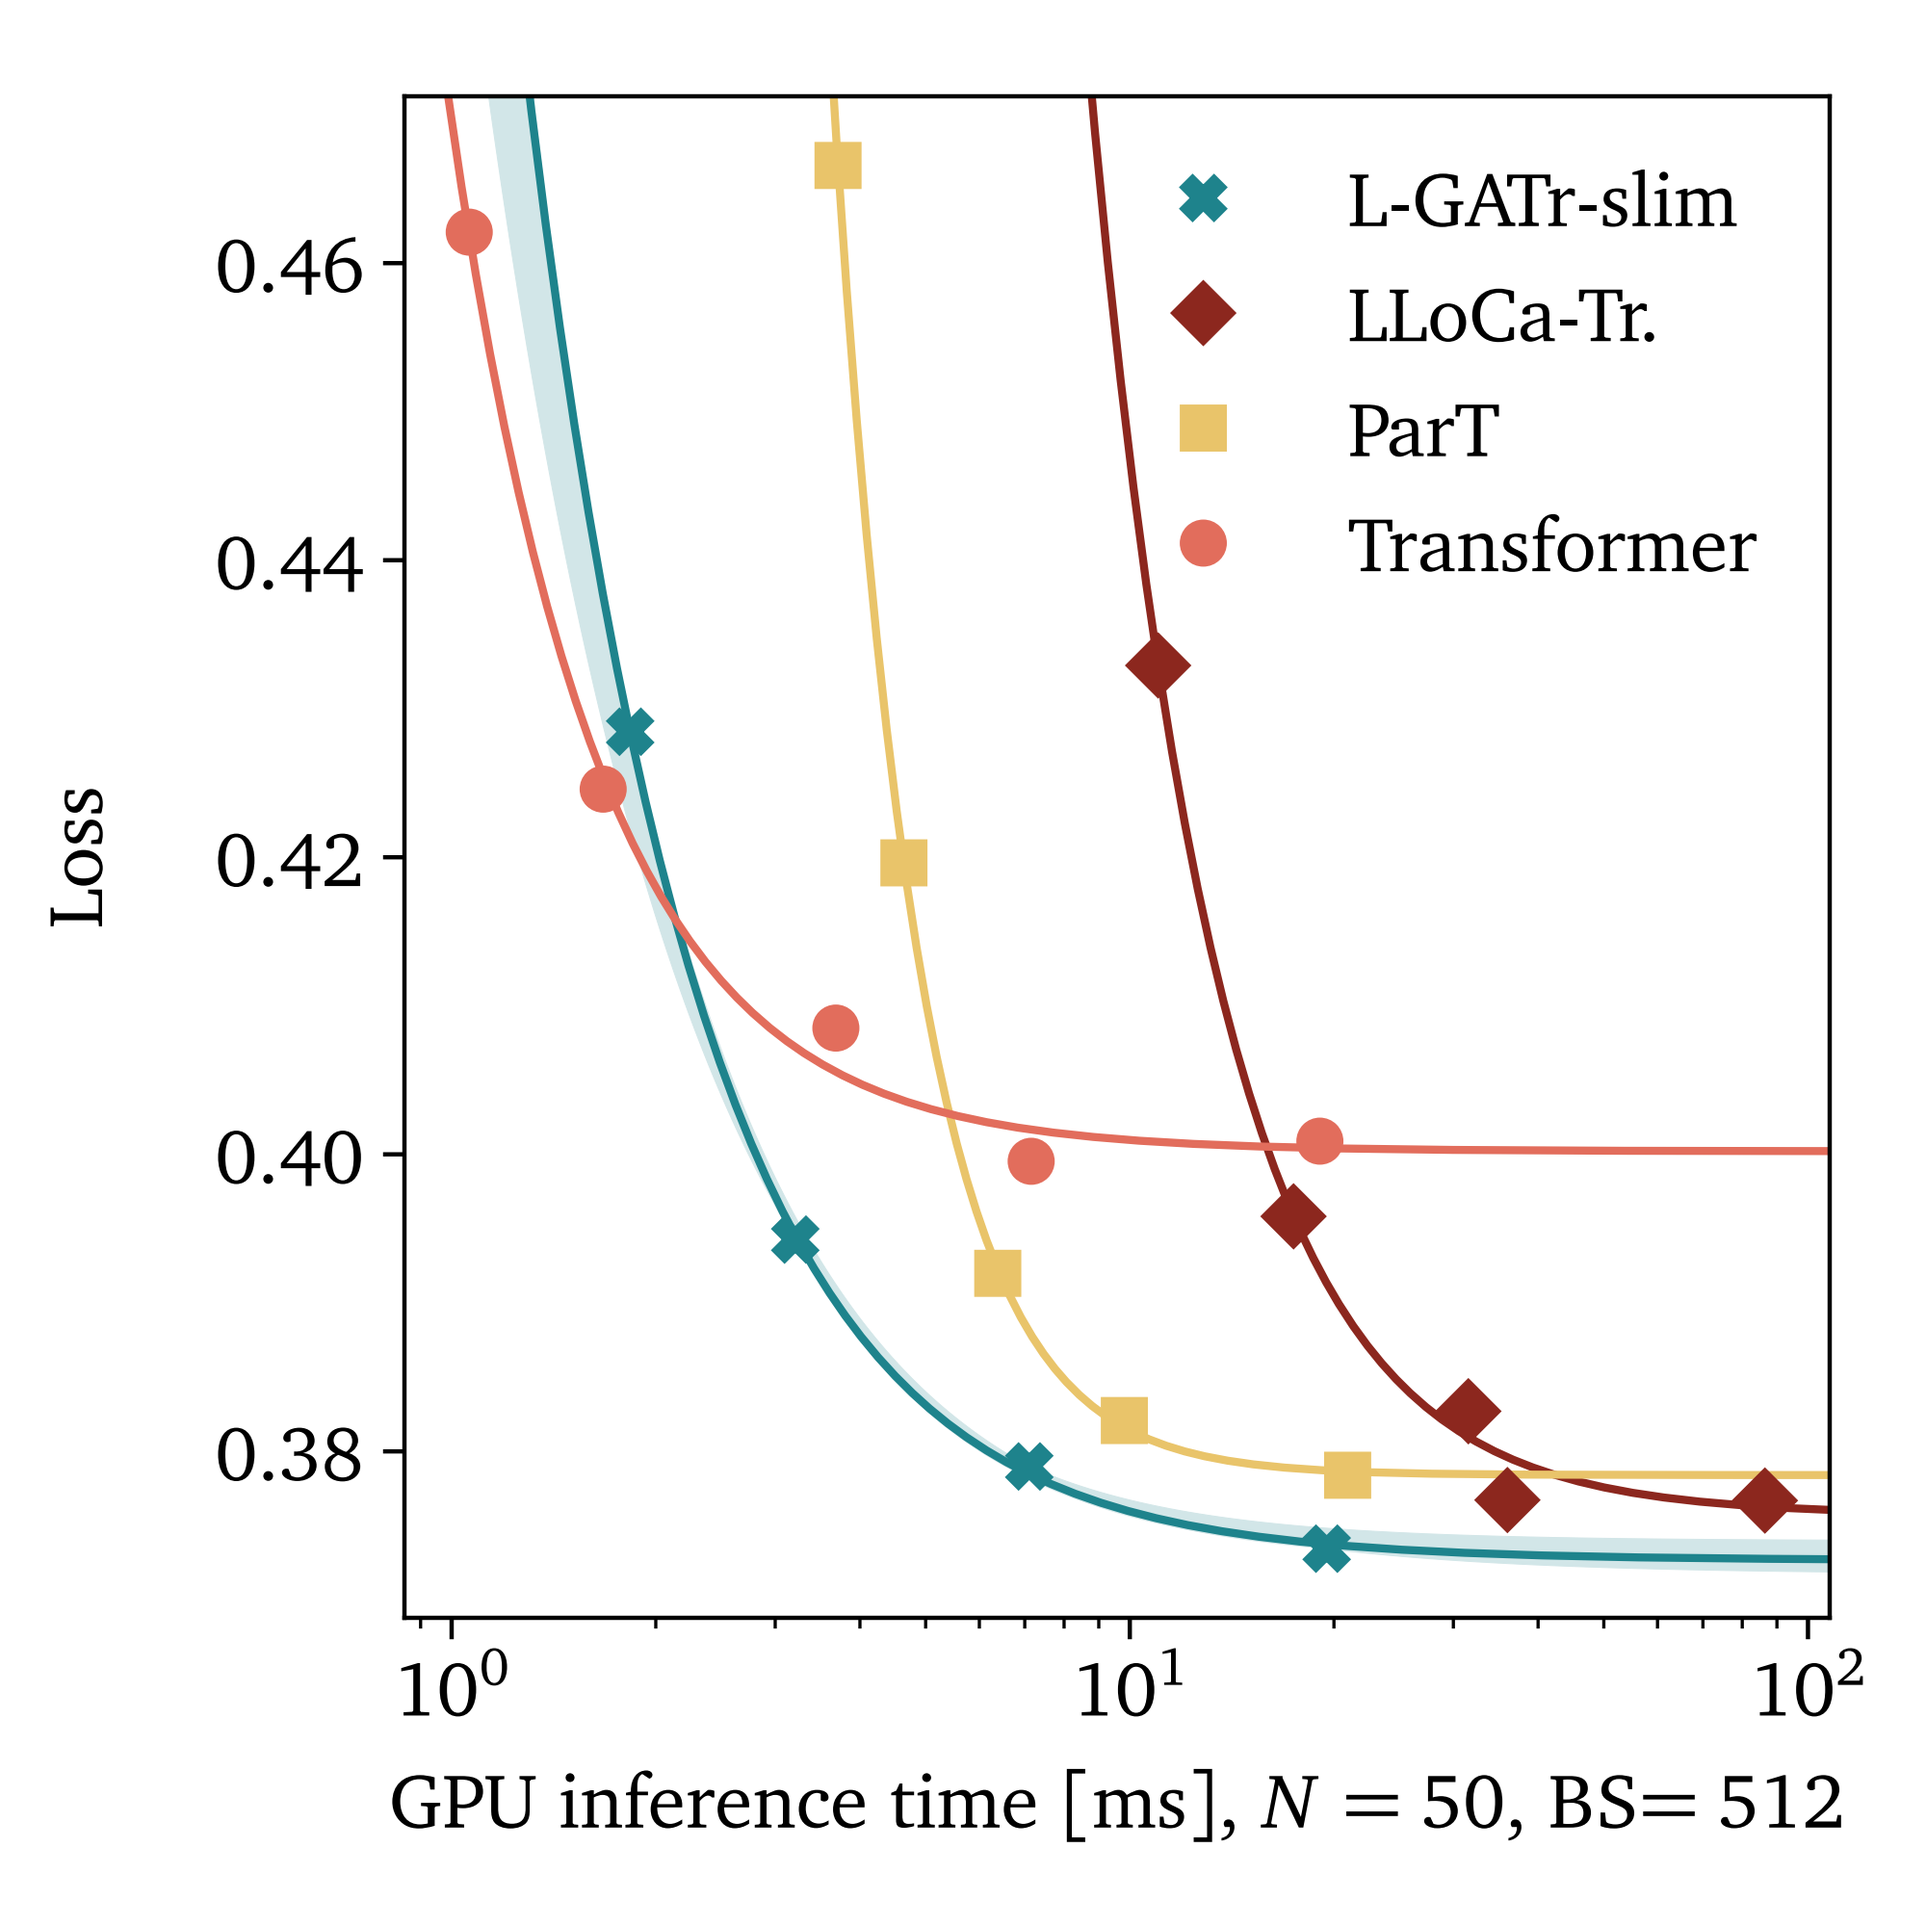
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.46
0.44
0.42
0.40
0.38
10^0
10^1
10^2
GPU inference time [ms], $N = 50$, BS$= 512$

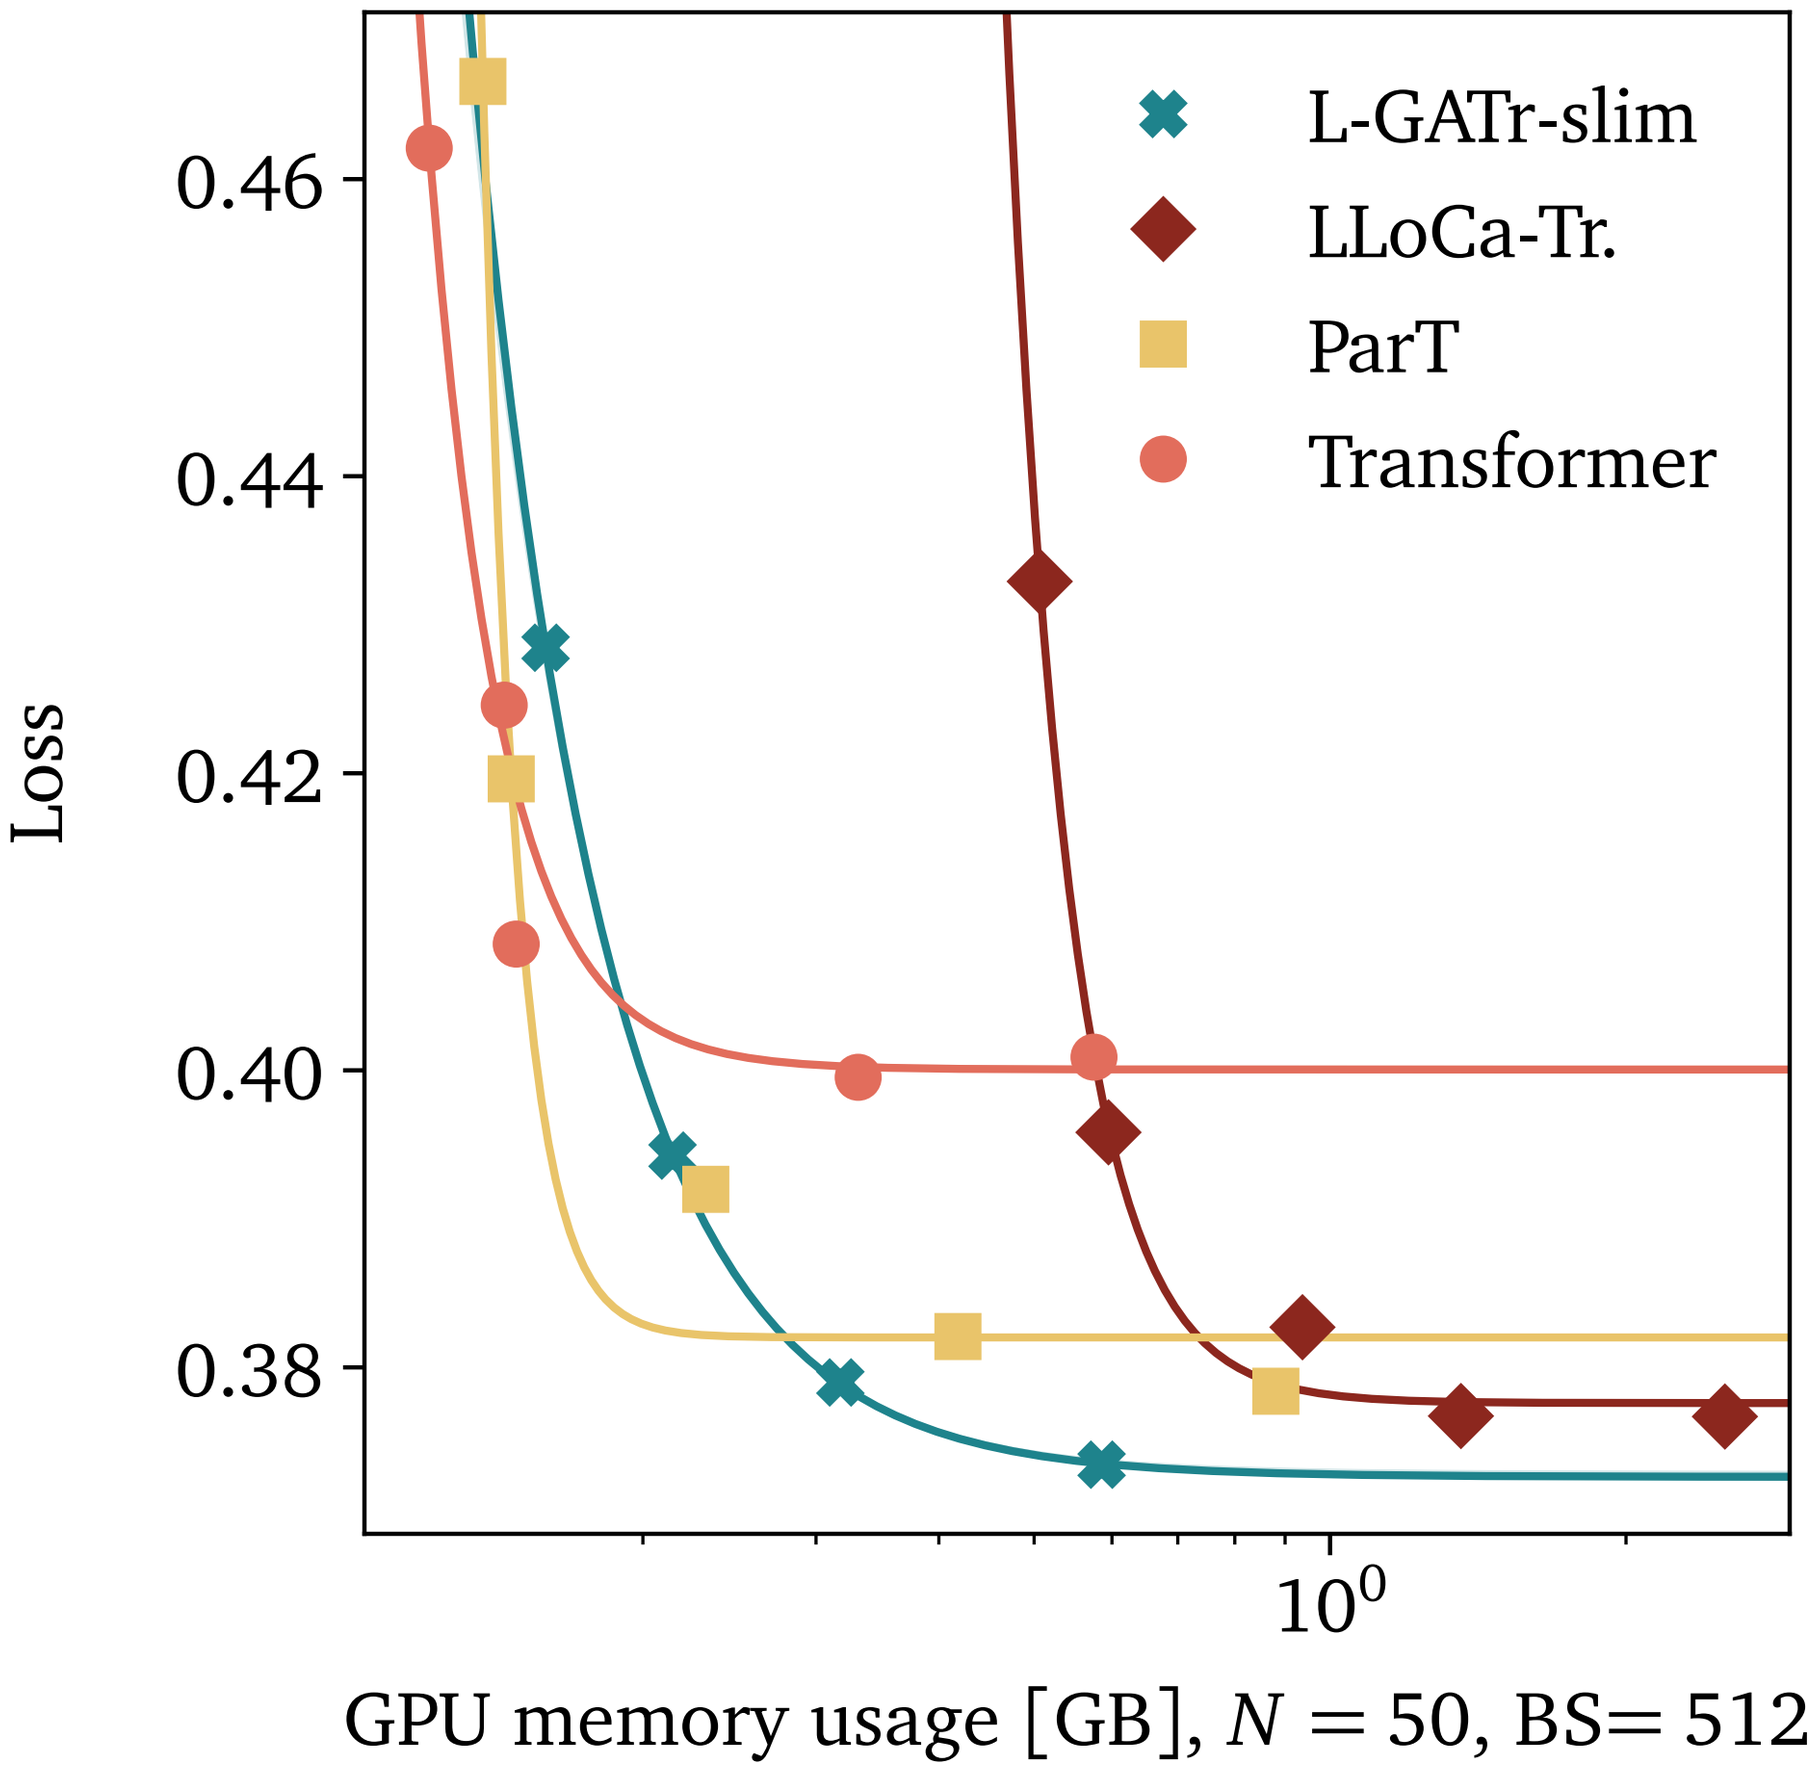
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.46
0.44
0.42
0.40
0.38
10^0
GPU memory usage [GB], $N = 50$, BS$= 512$

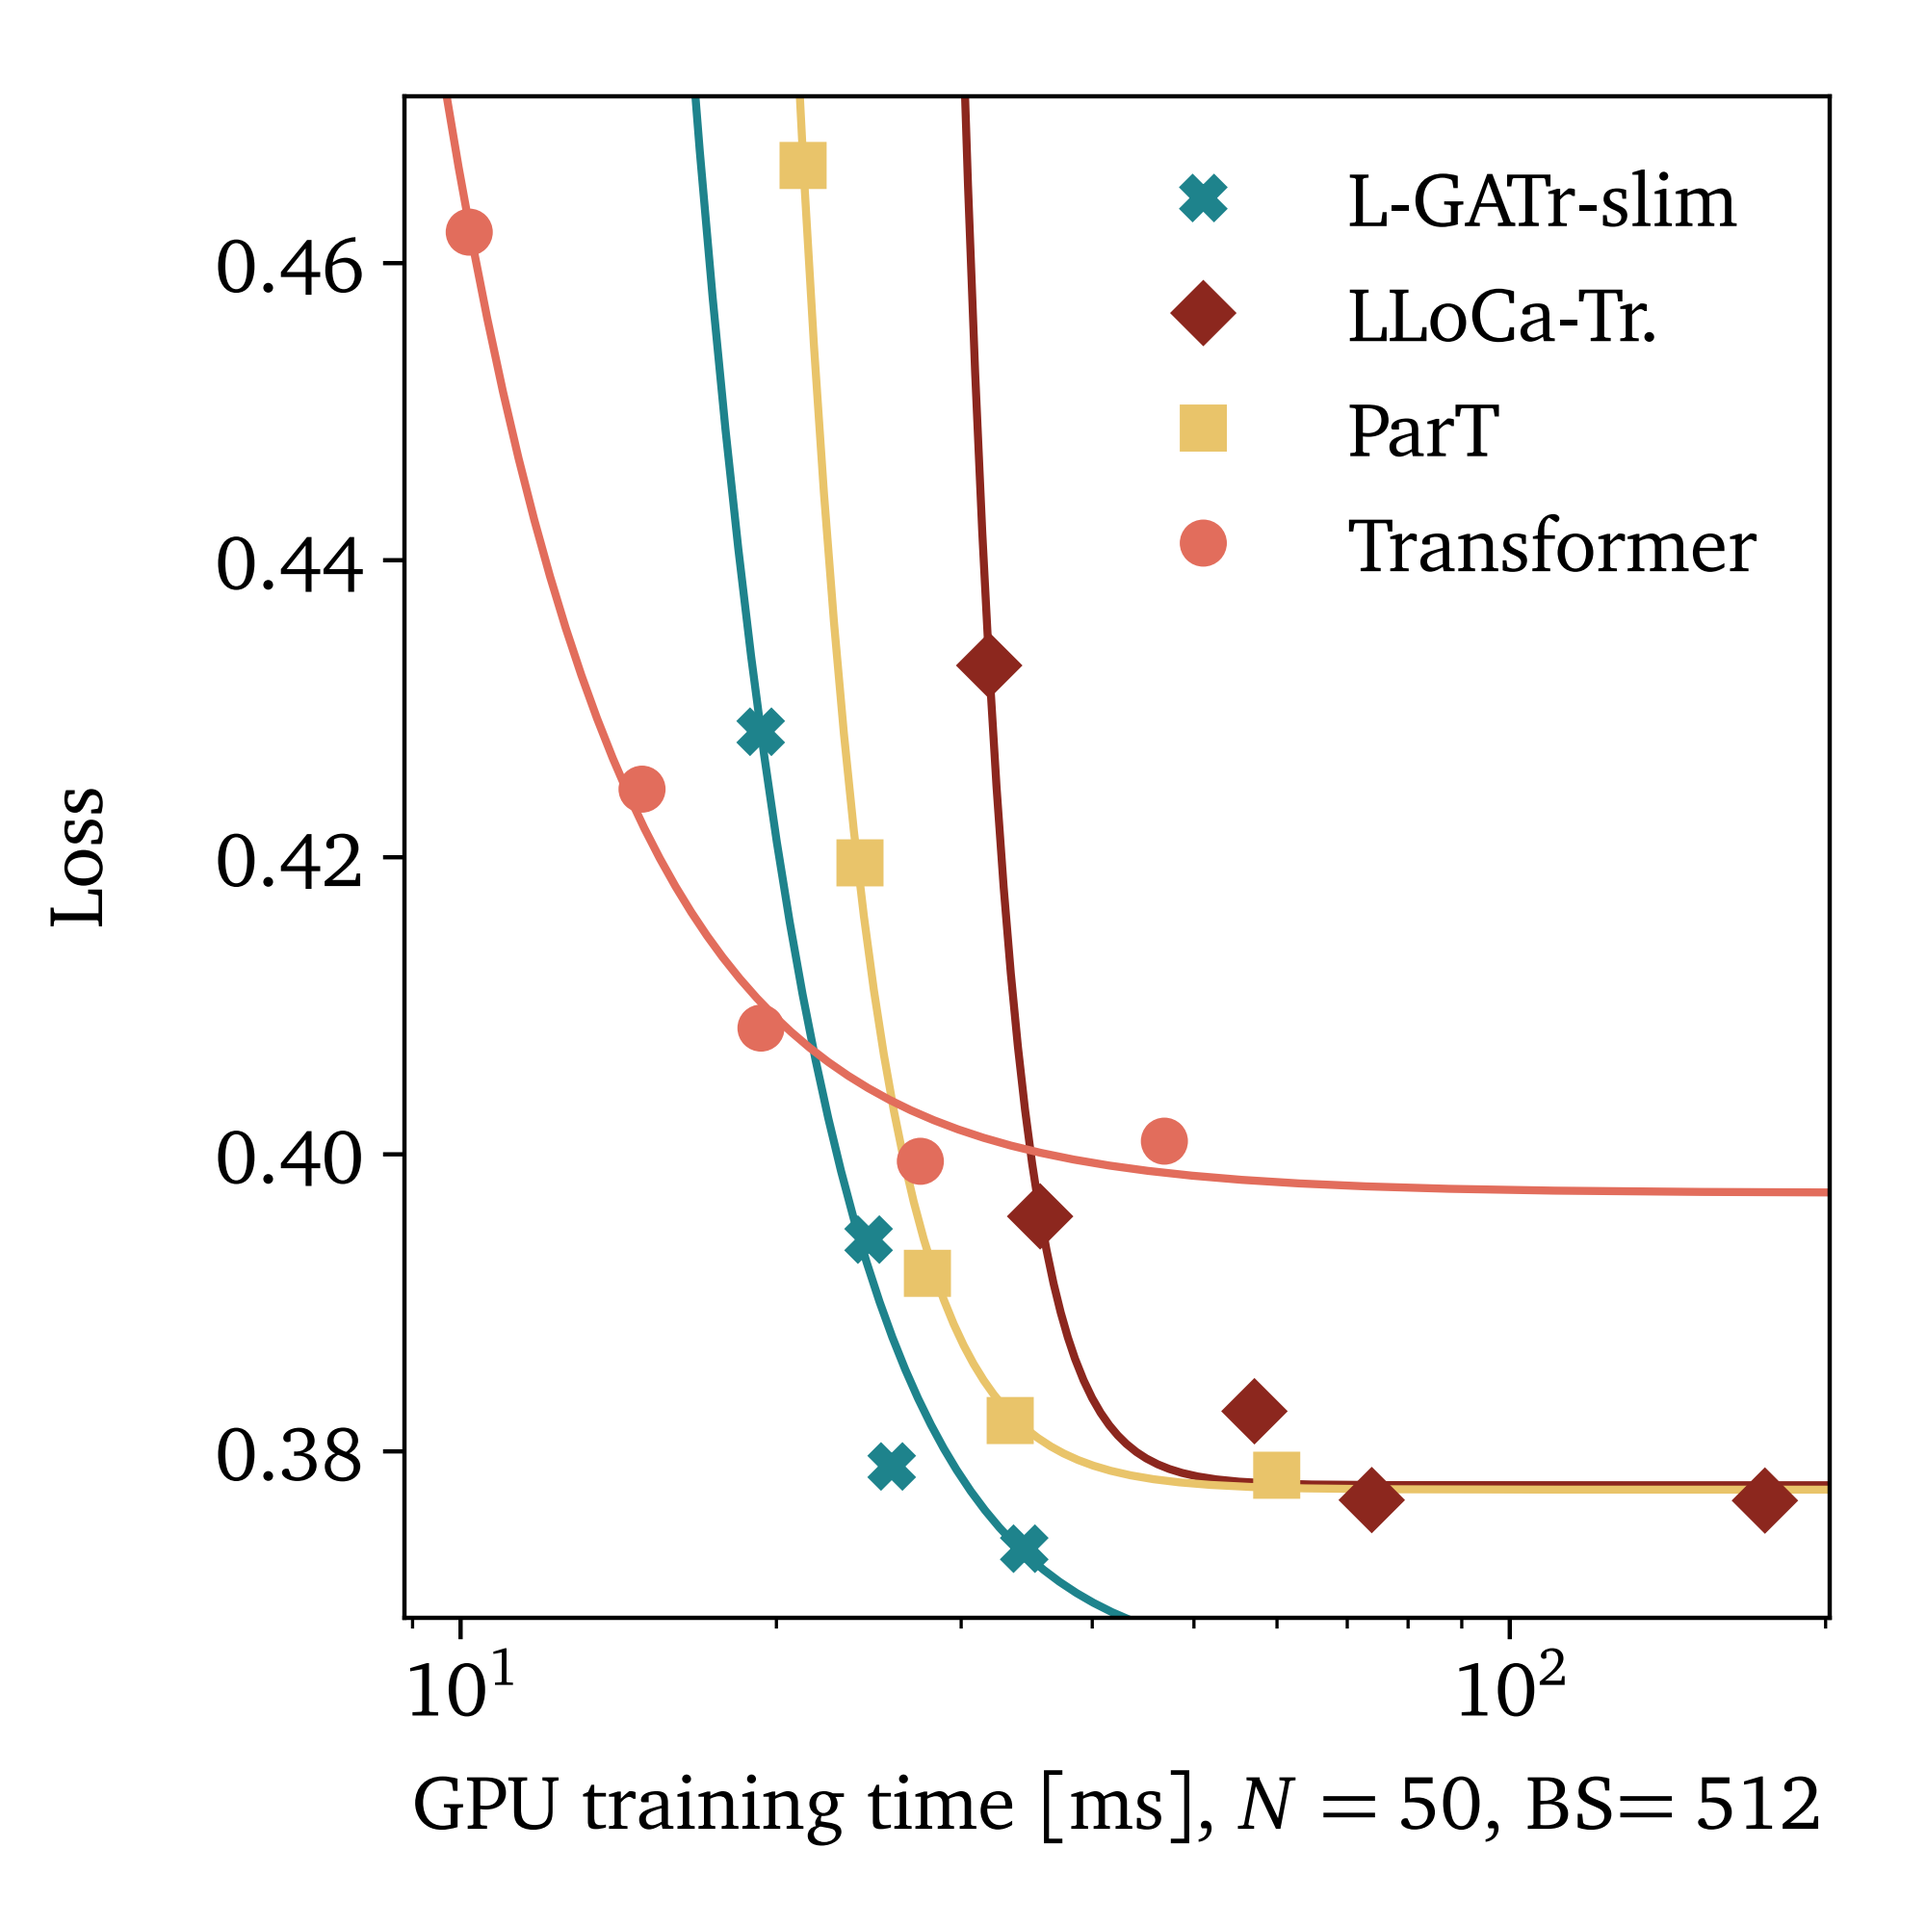
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.46
0.44
0.42
0.40
0.38
10^1
10^2
GPU training time [ms], $N = 50$, BS$= 512$

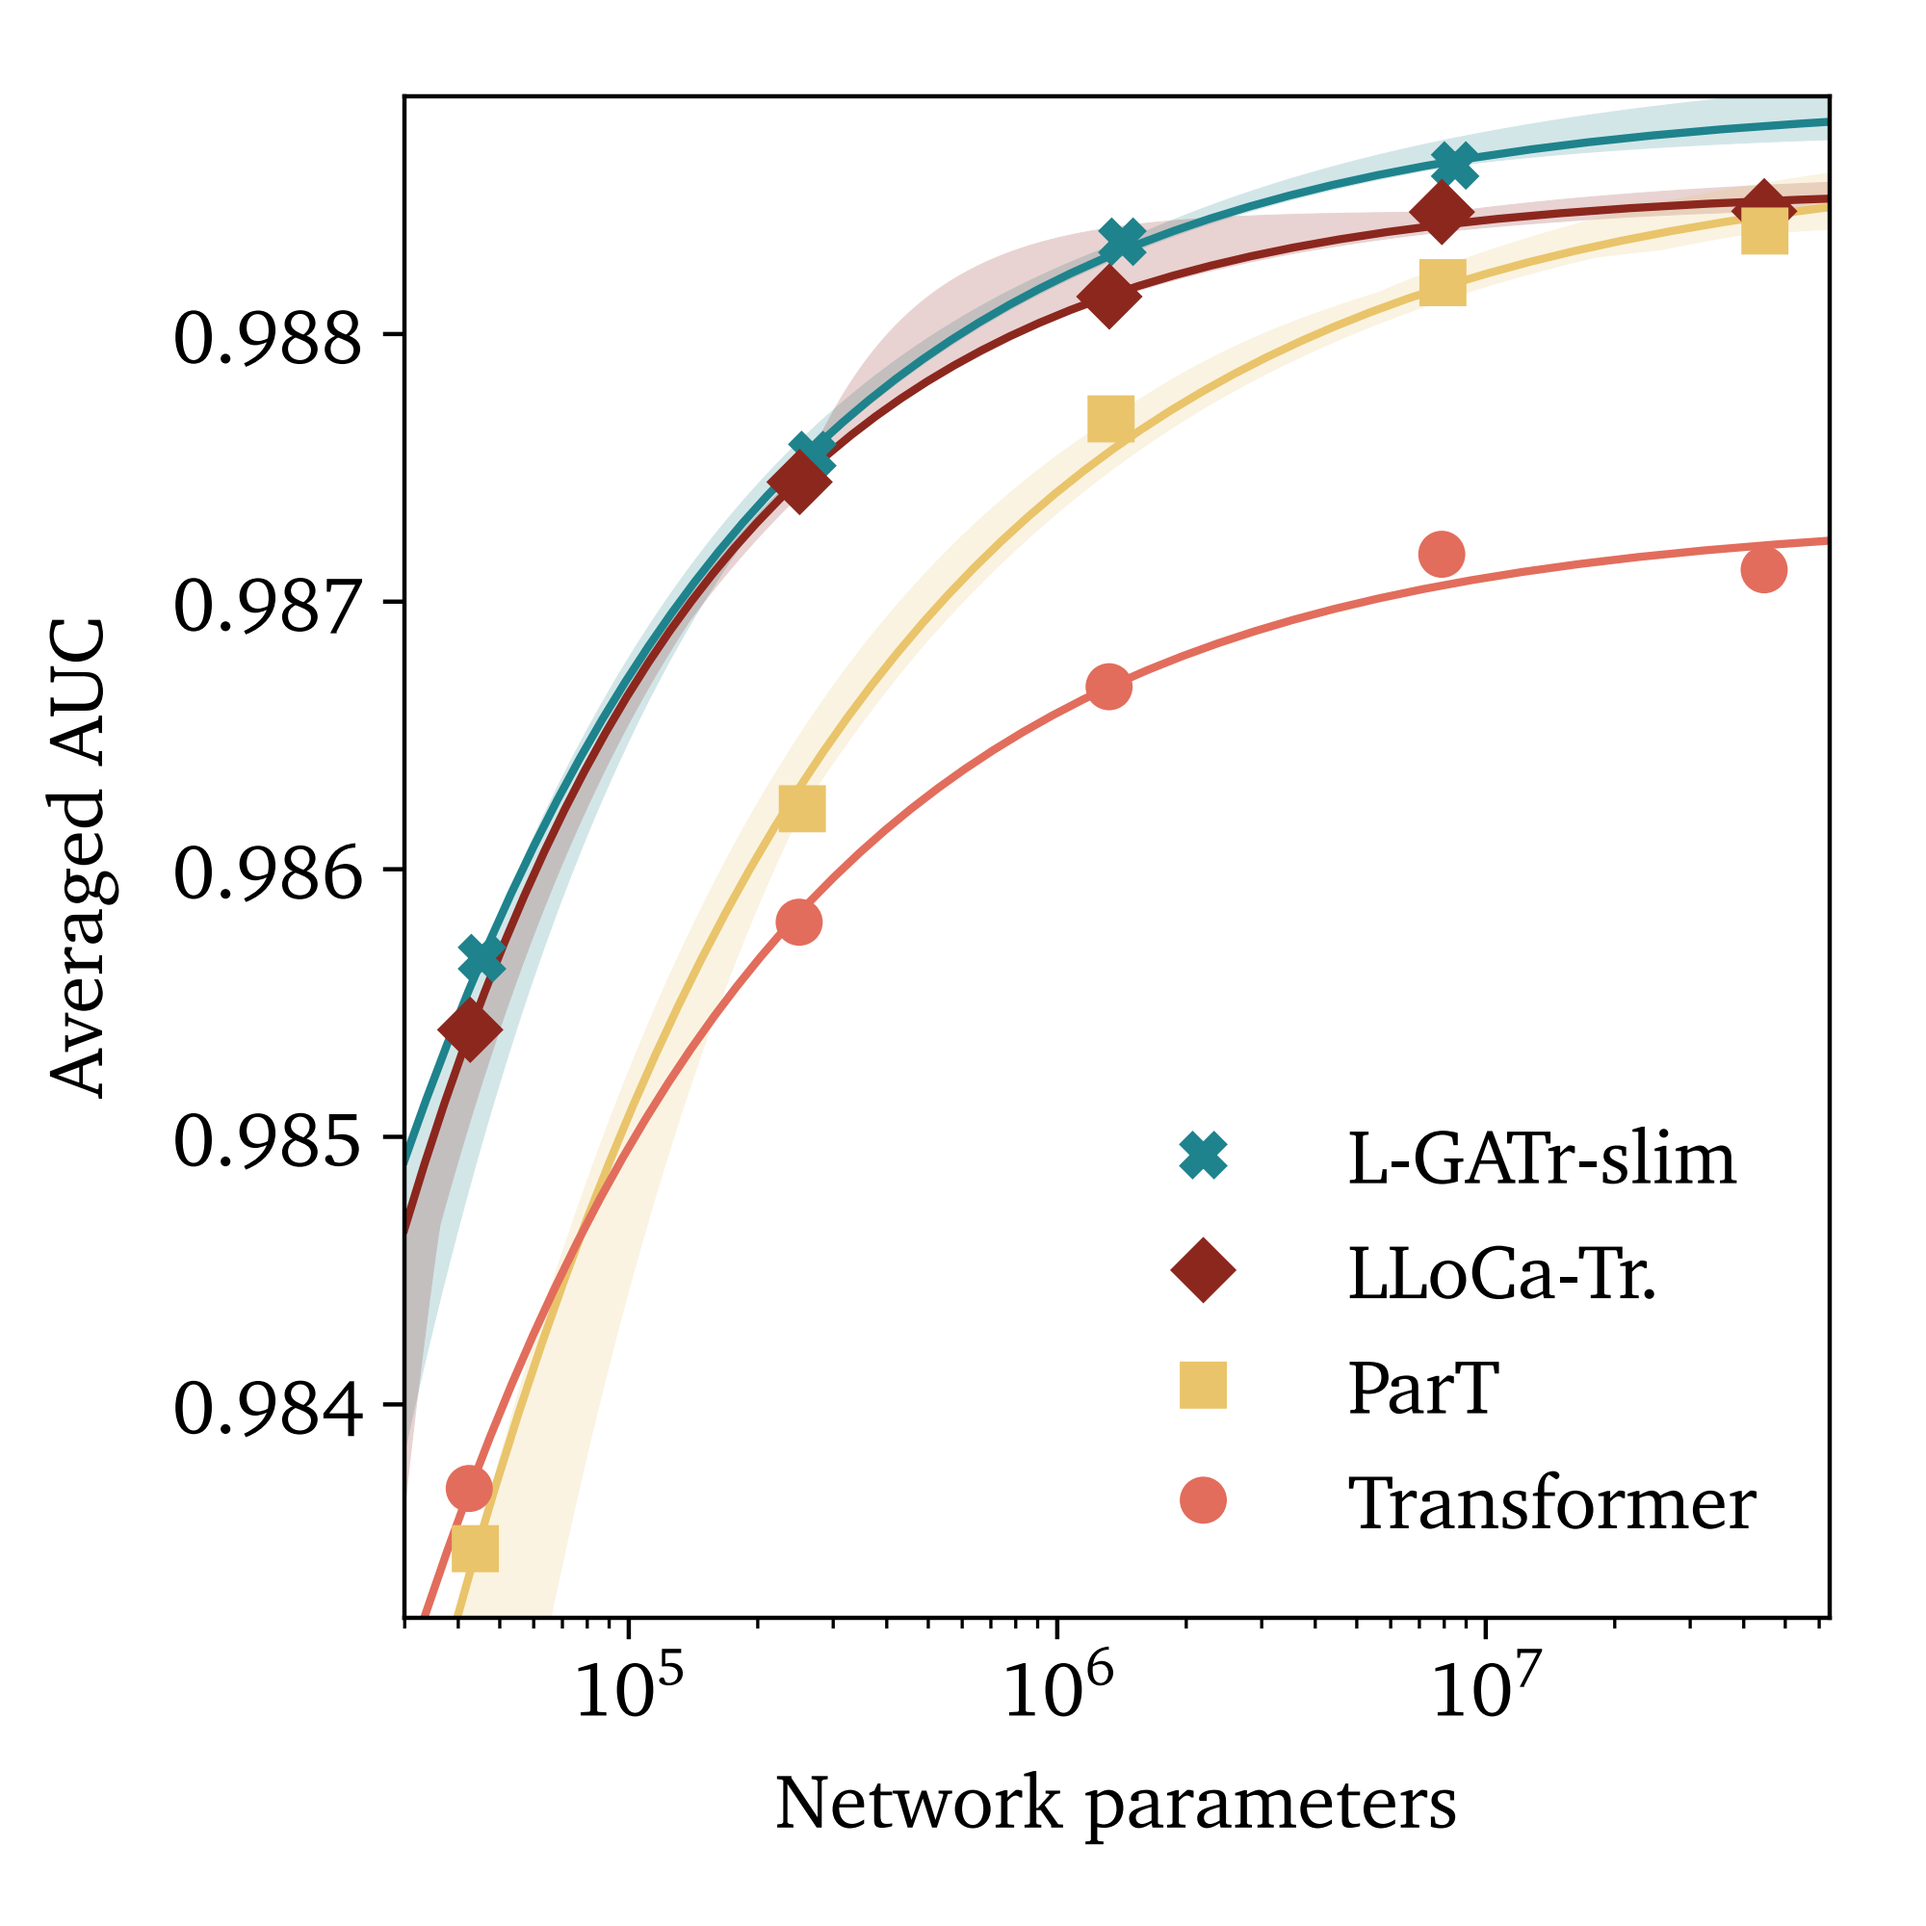
Averaged AUC
0.988
0.987
0.986
0.985
0.984
10^5
10^6
10^7
Network parameters
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

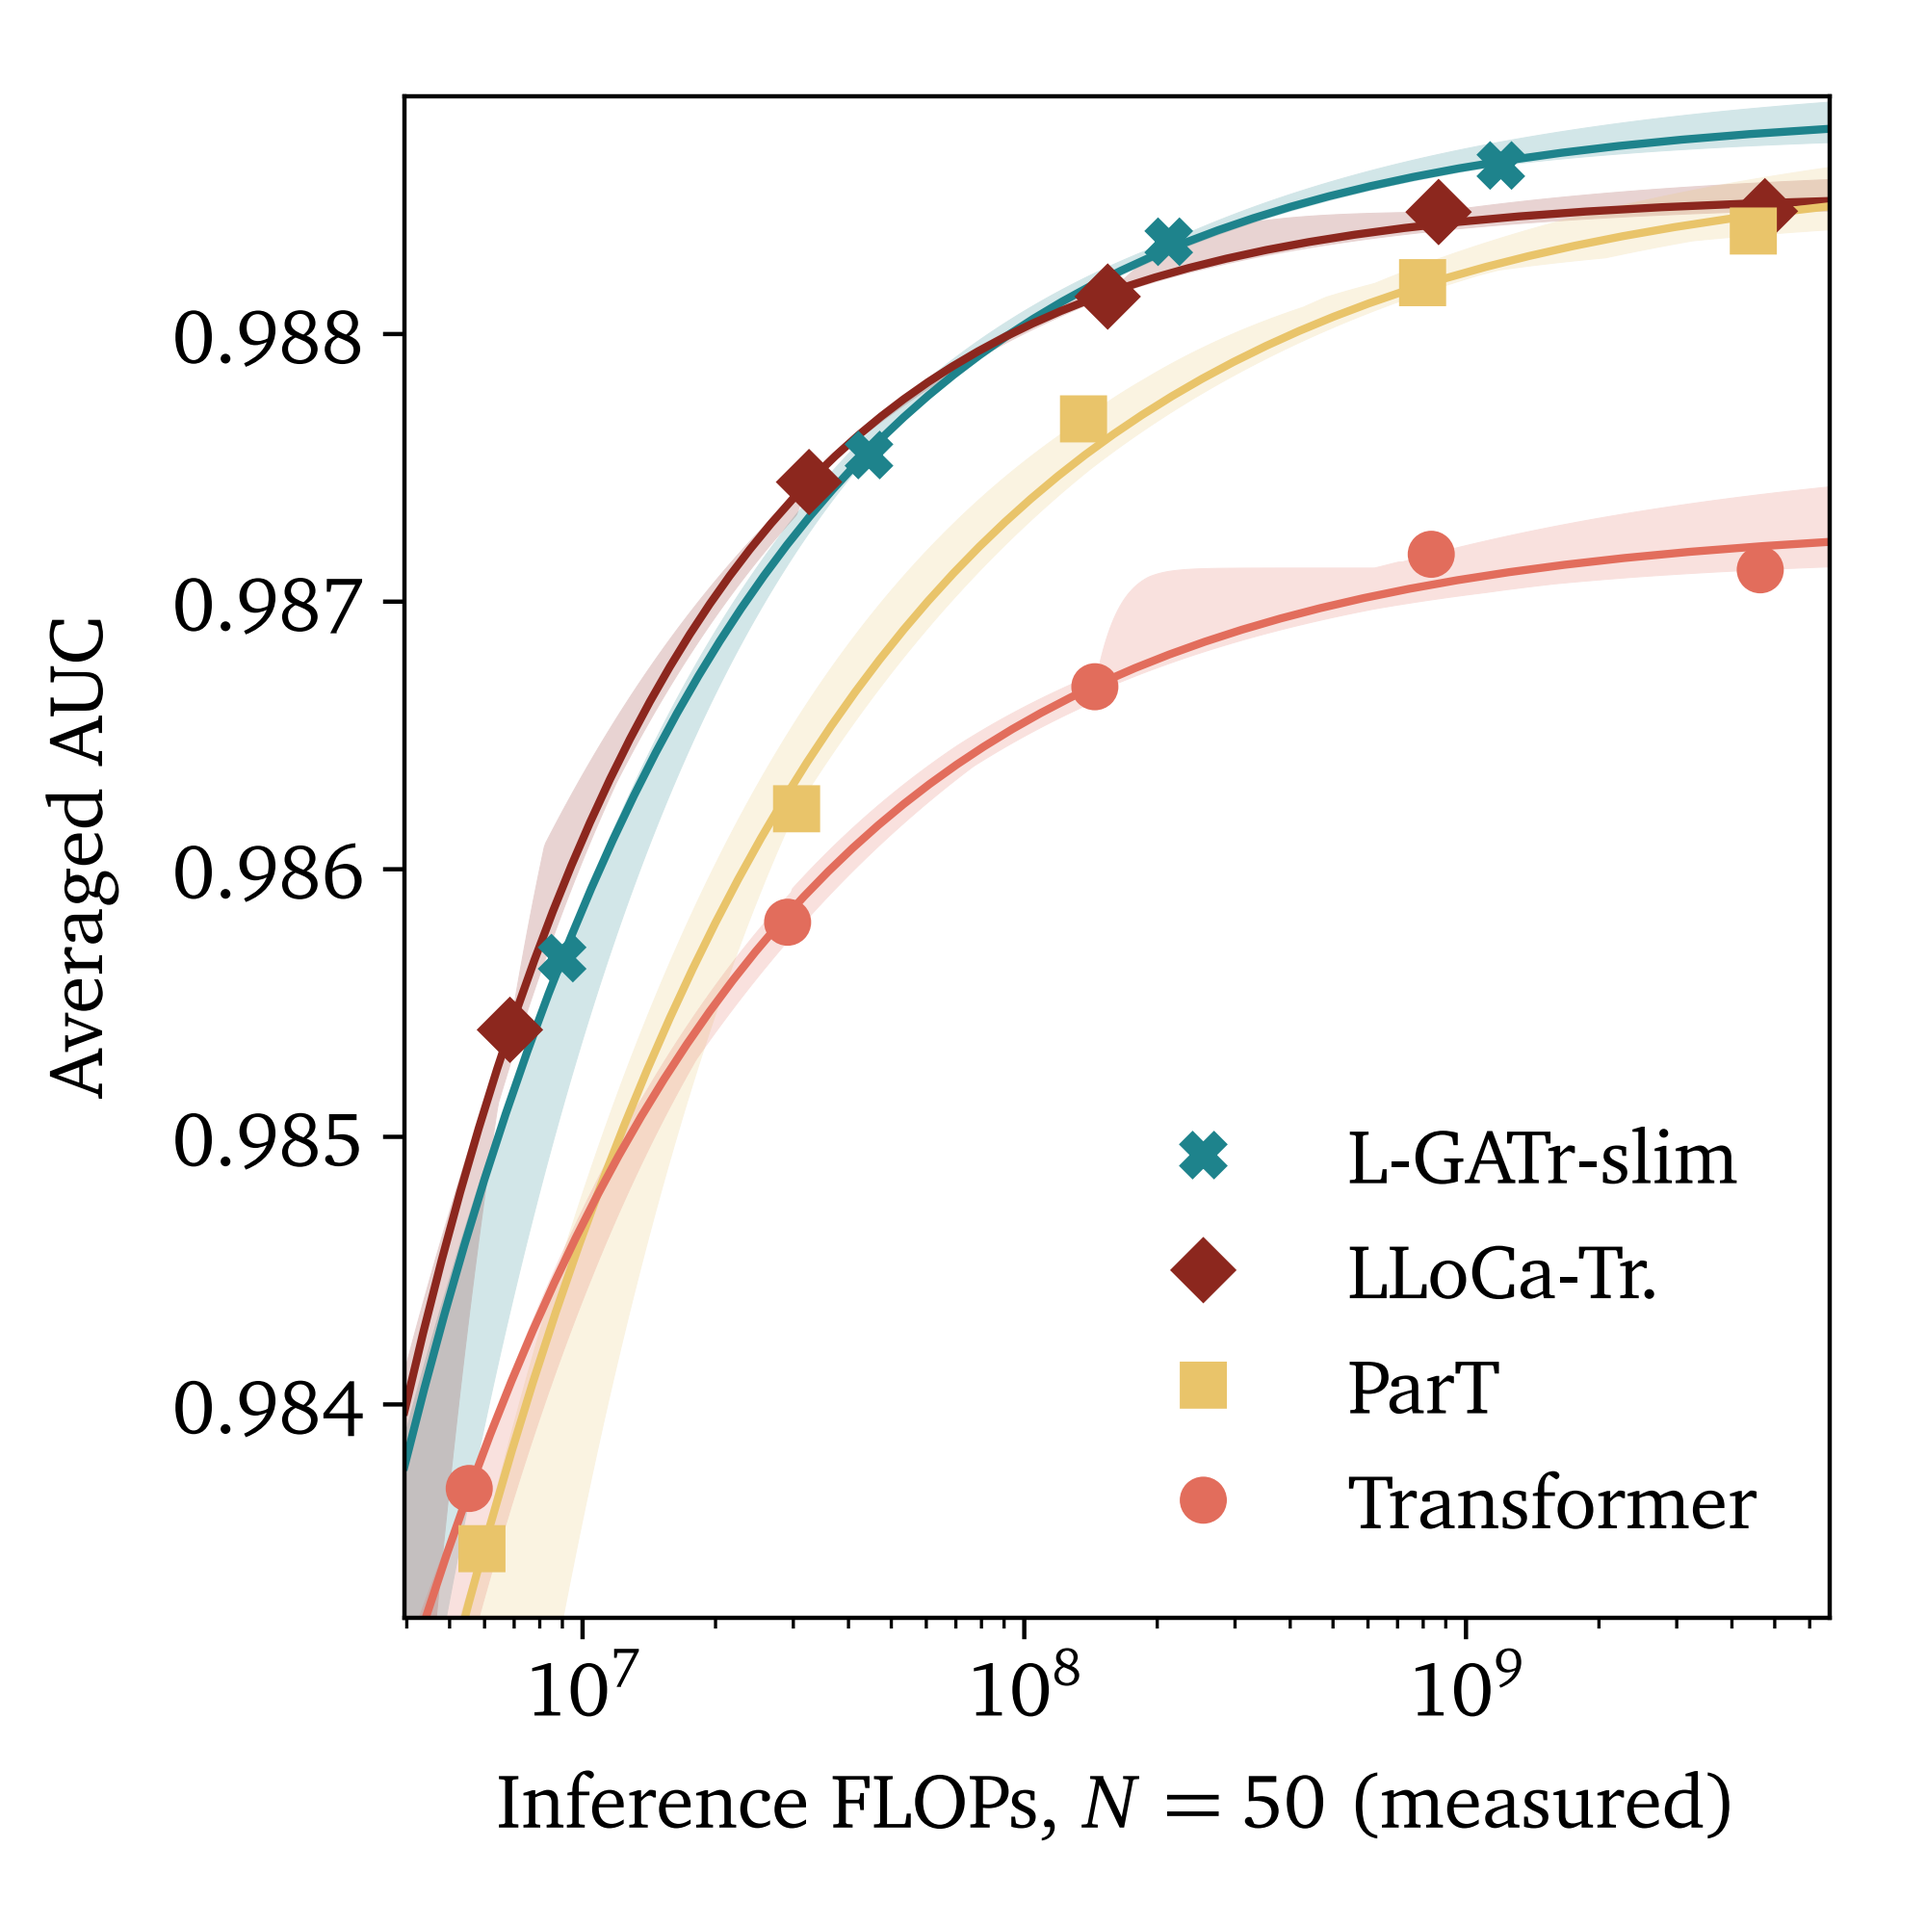
0.988
0.987
0.986
0.985
0.984
Averaged AUC
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (measured)
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

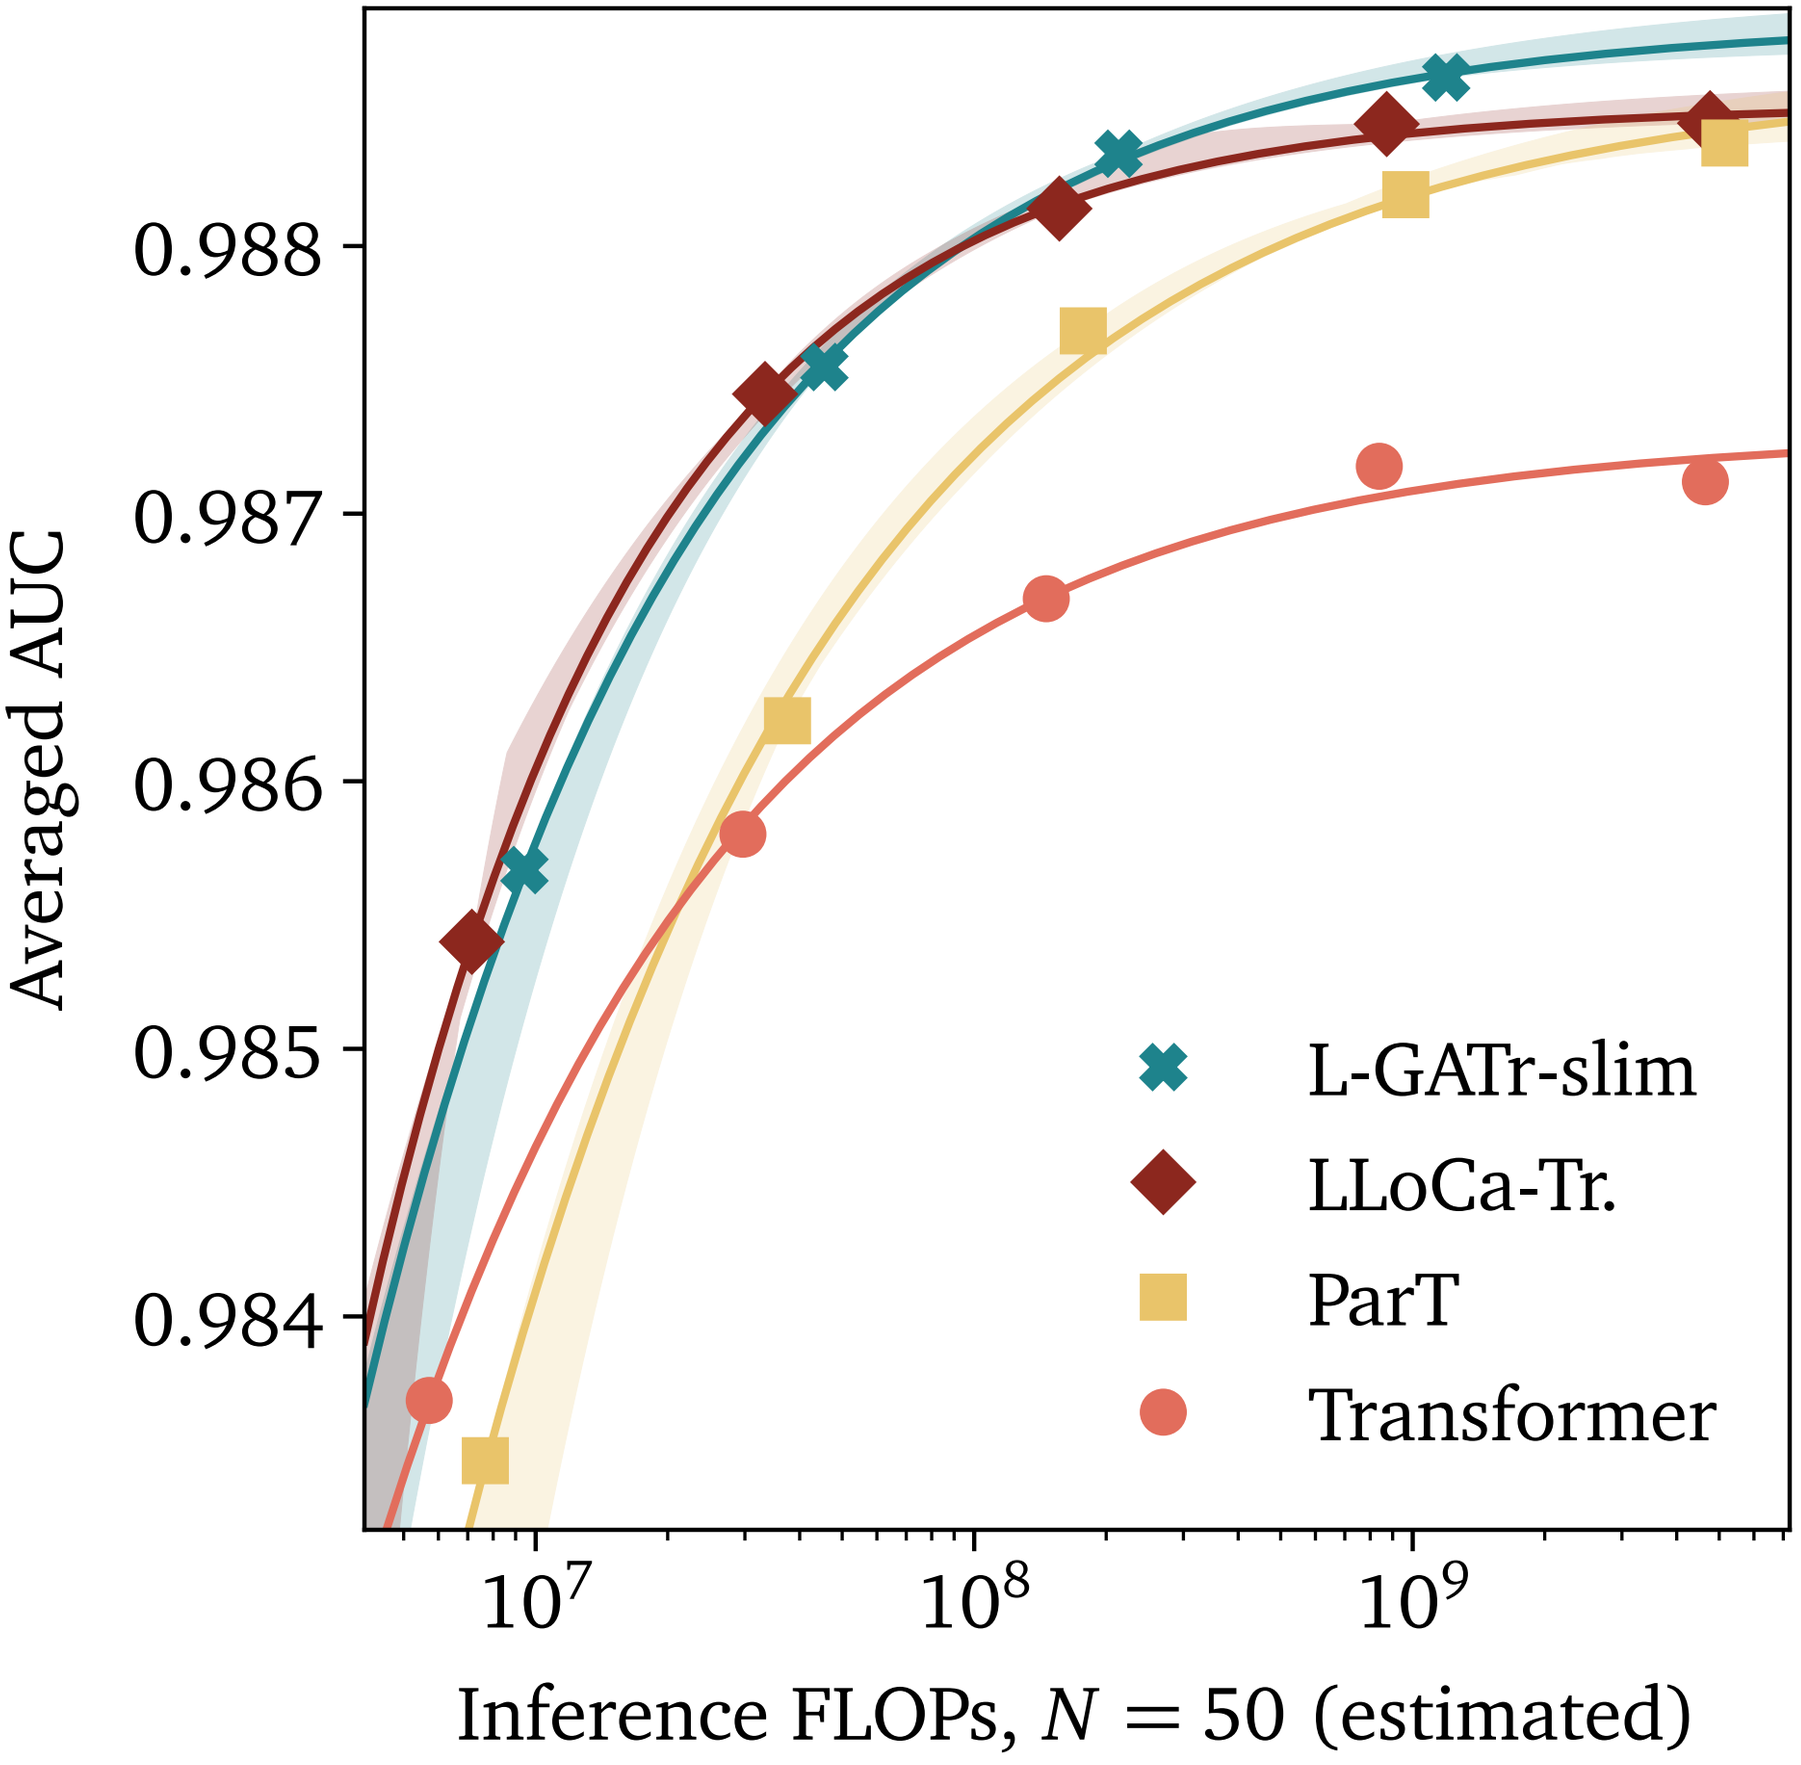
0.988
0.987
0.986
0.985
0.984
Averaged AUC
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (estimated)
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

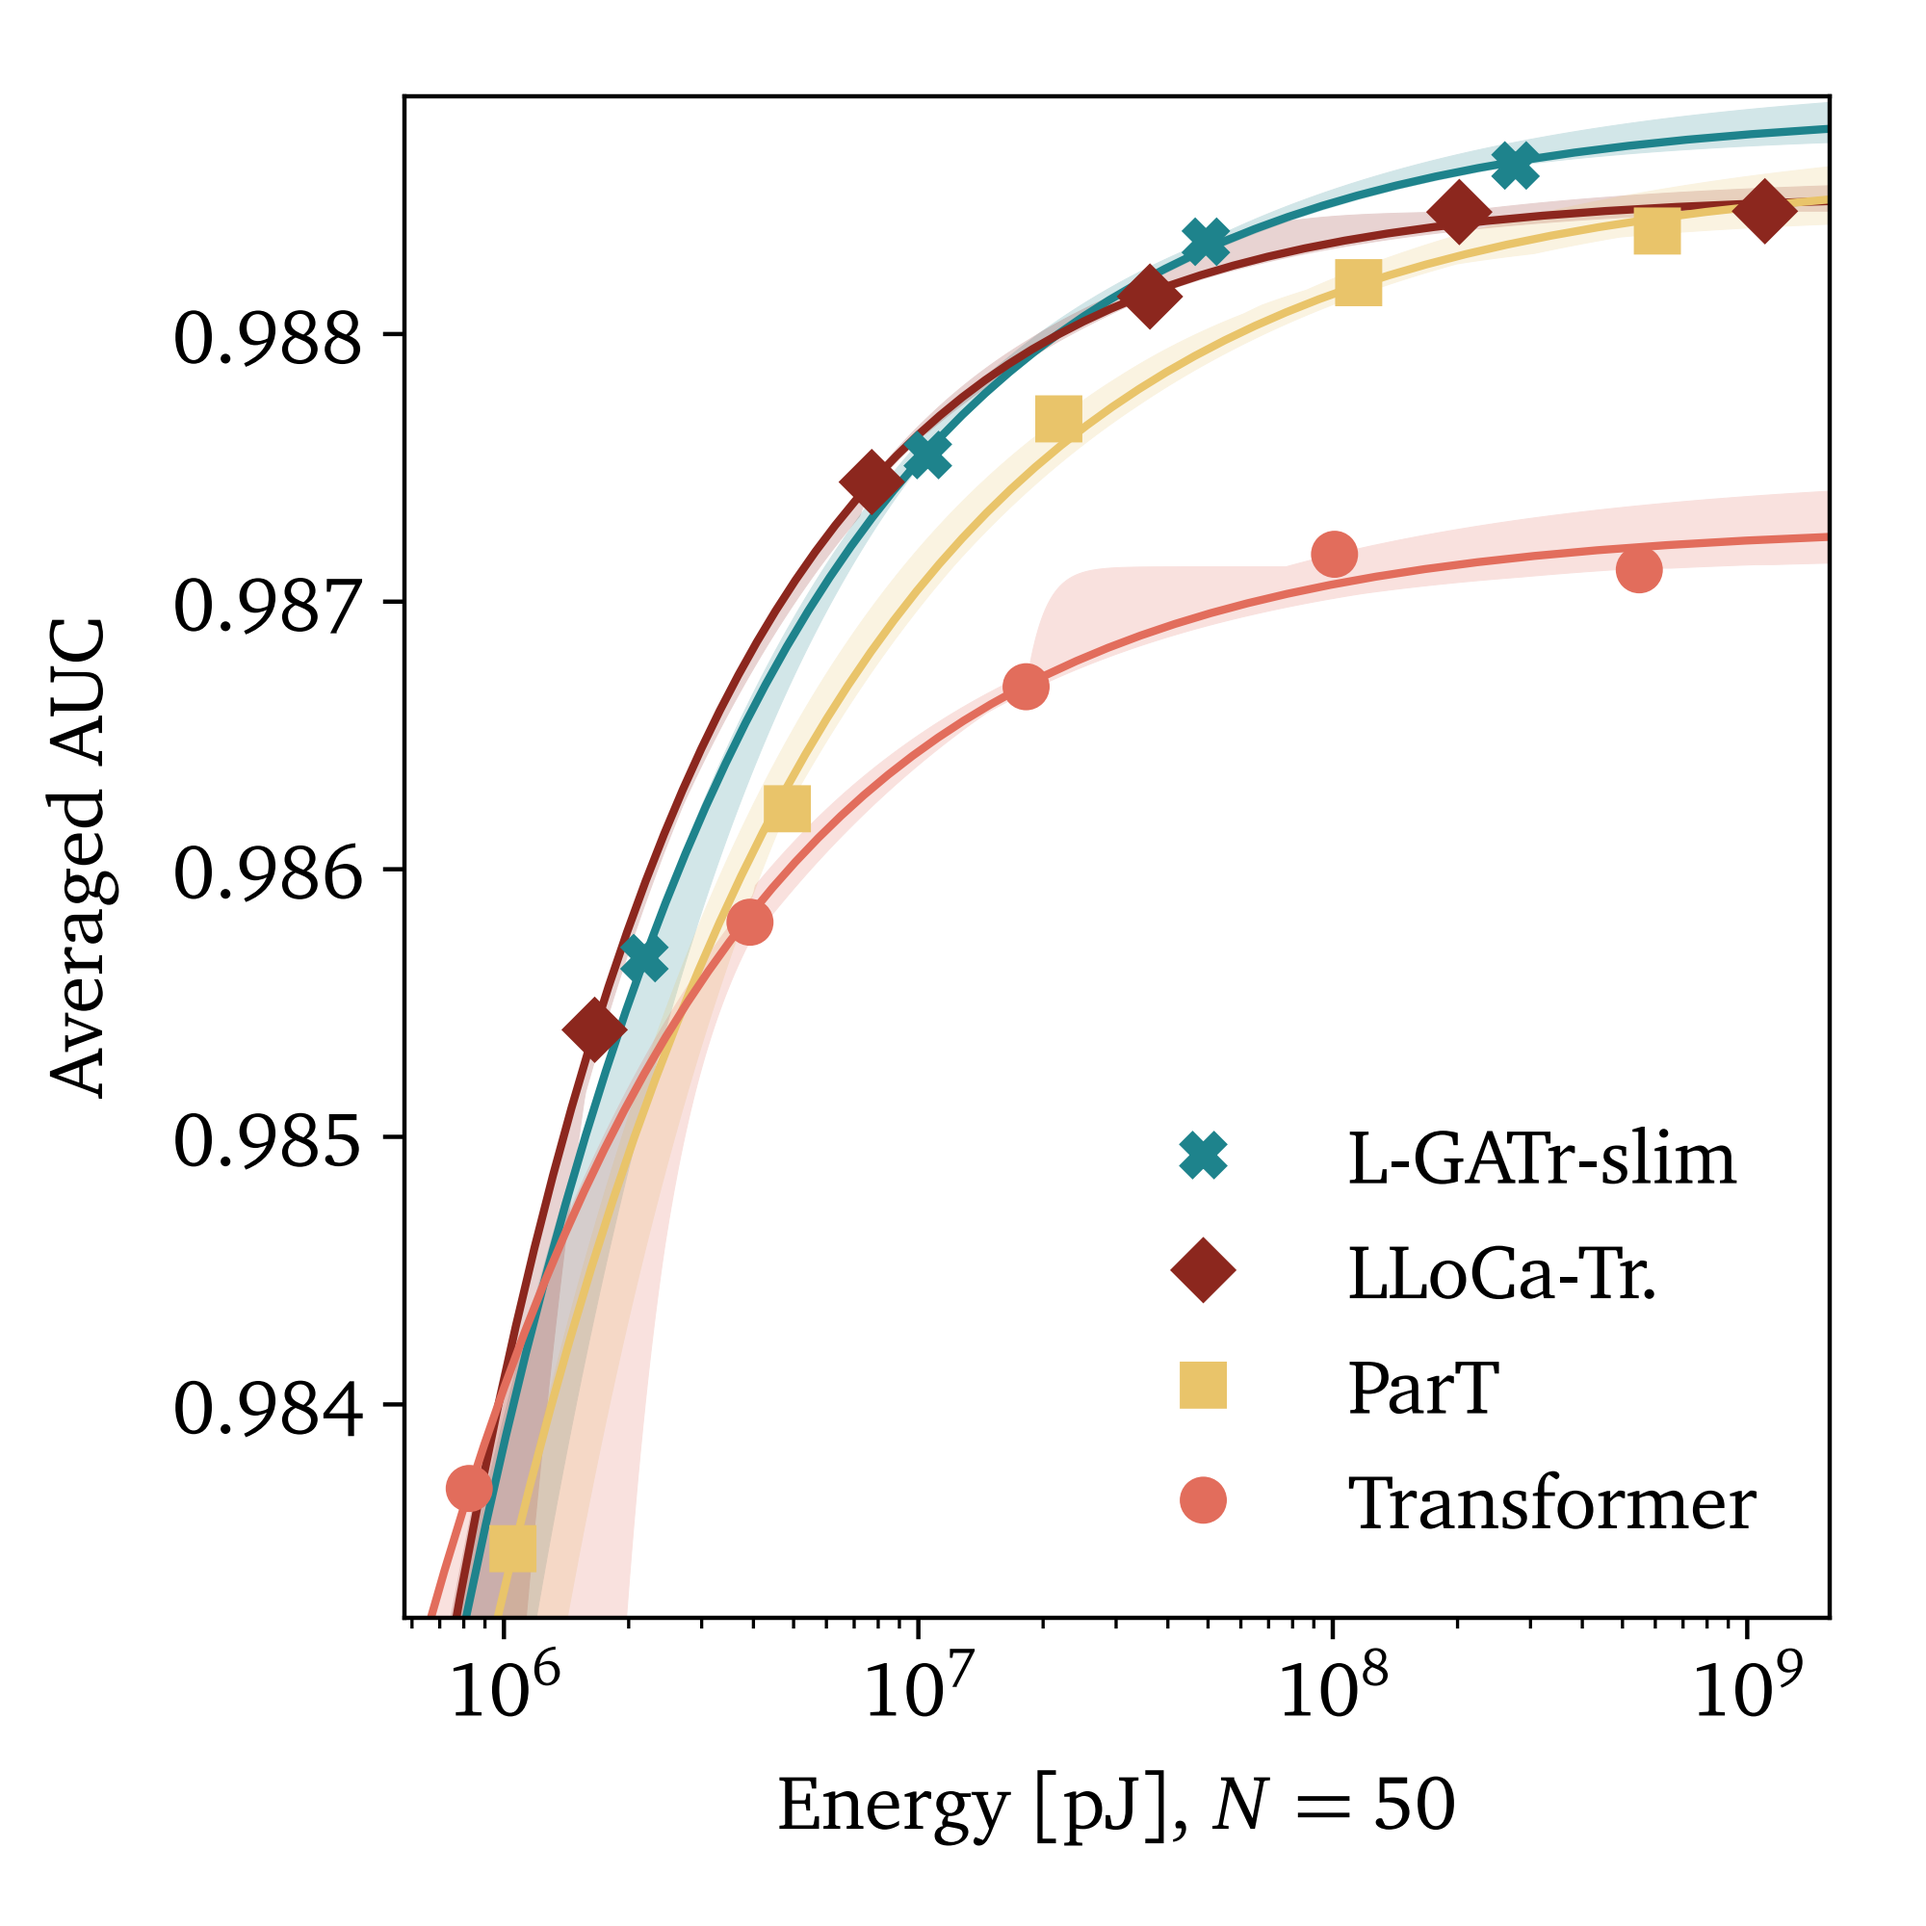
0.988
0.987
0.986
0.985
0.984
Averaged AUC
10^6
10^7
10^8
10^9
Energy [pJ], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

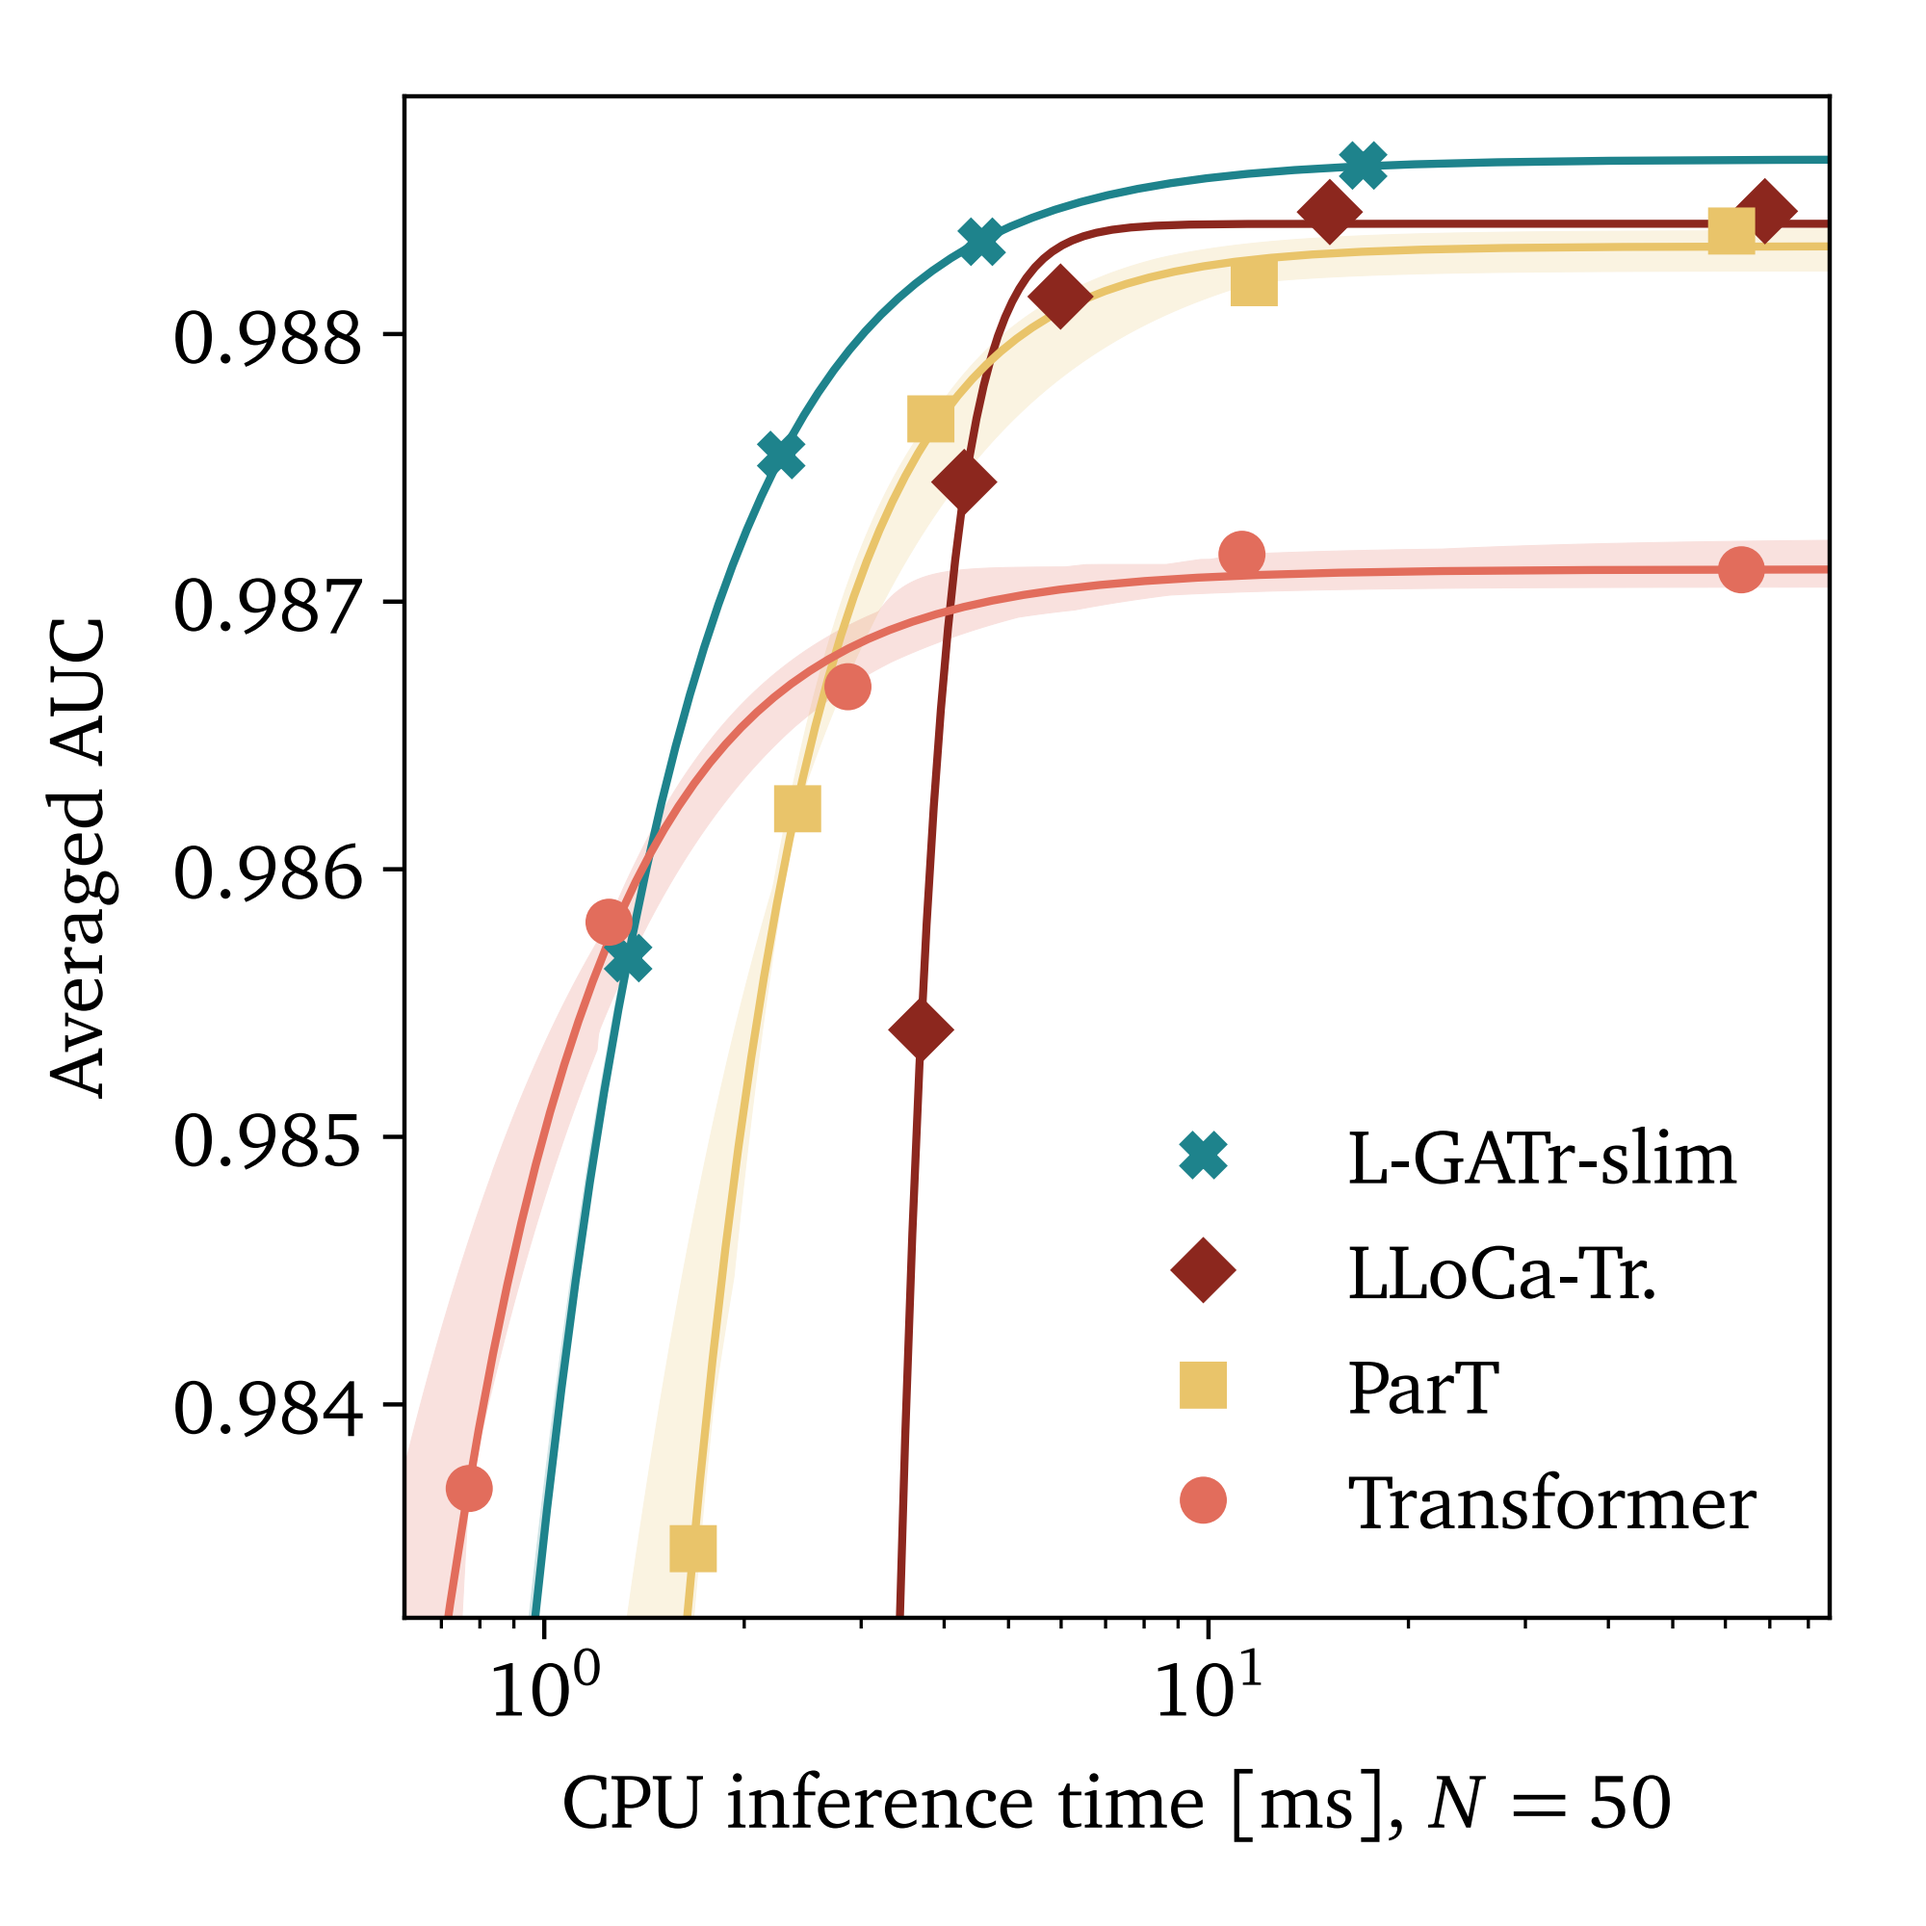
0.988
0.987
0.986
0.985
0.984
Averaged AUC
10^0
10^1
CPU inference time [ms], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

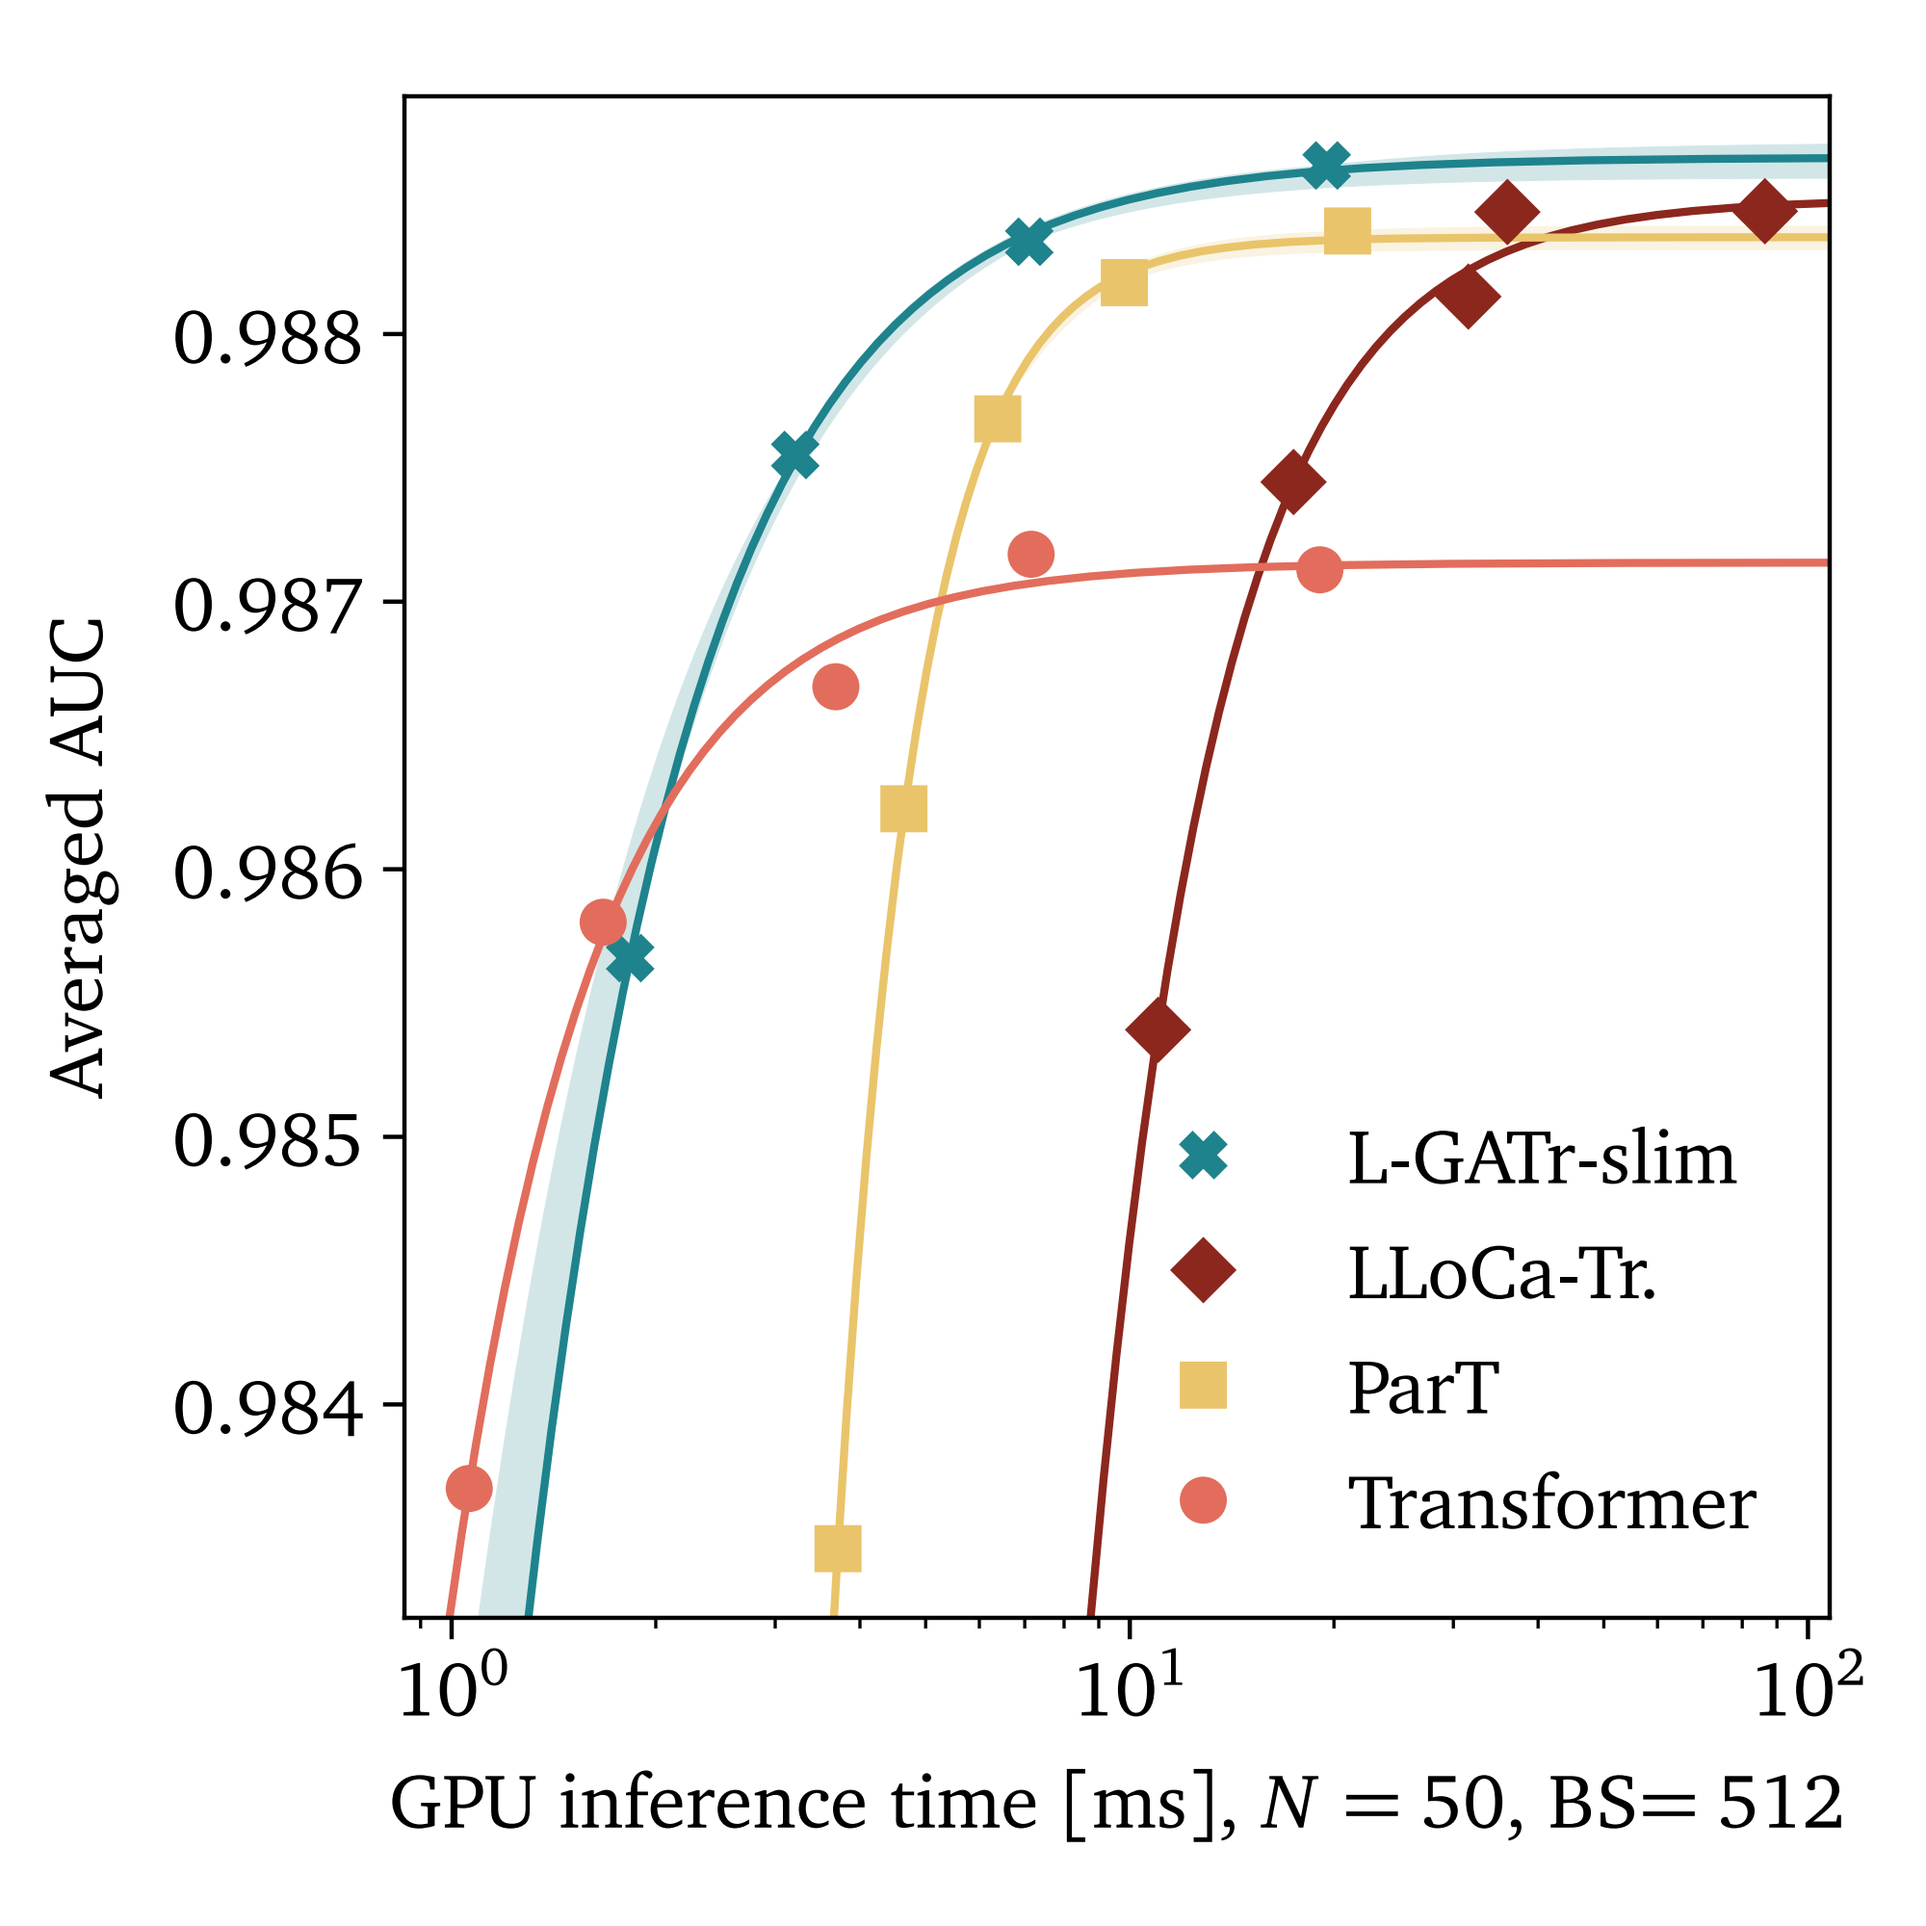
0.988
0.987
0.986
0.985
0.984
Averaged AUC
10^0
10^1
10^2
GPU inference time [ms], $N = 50$, BS= 512
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

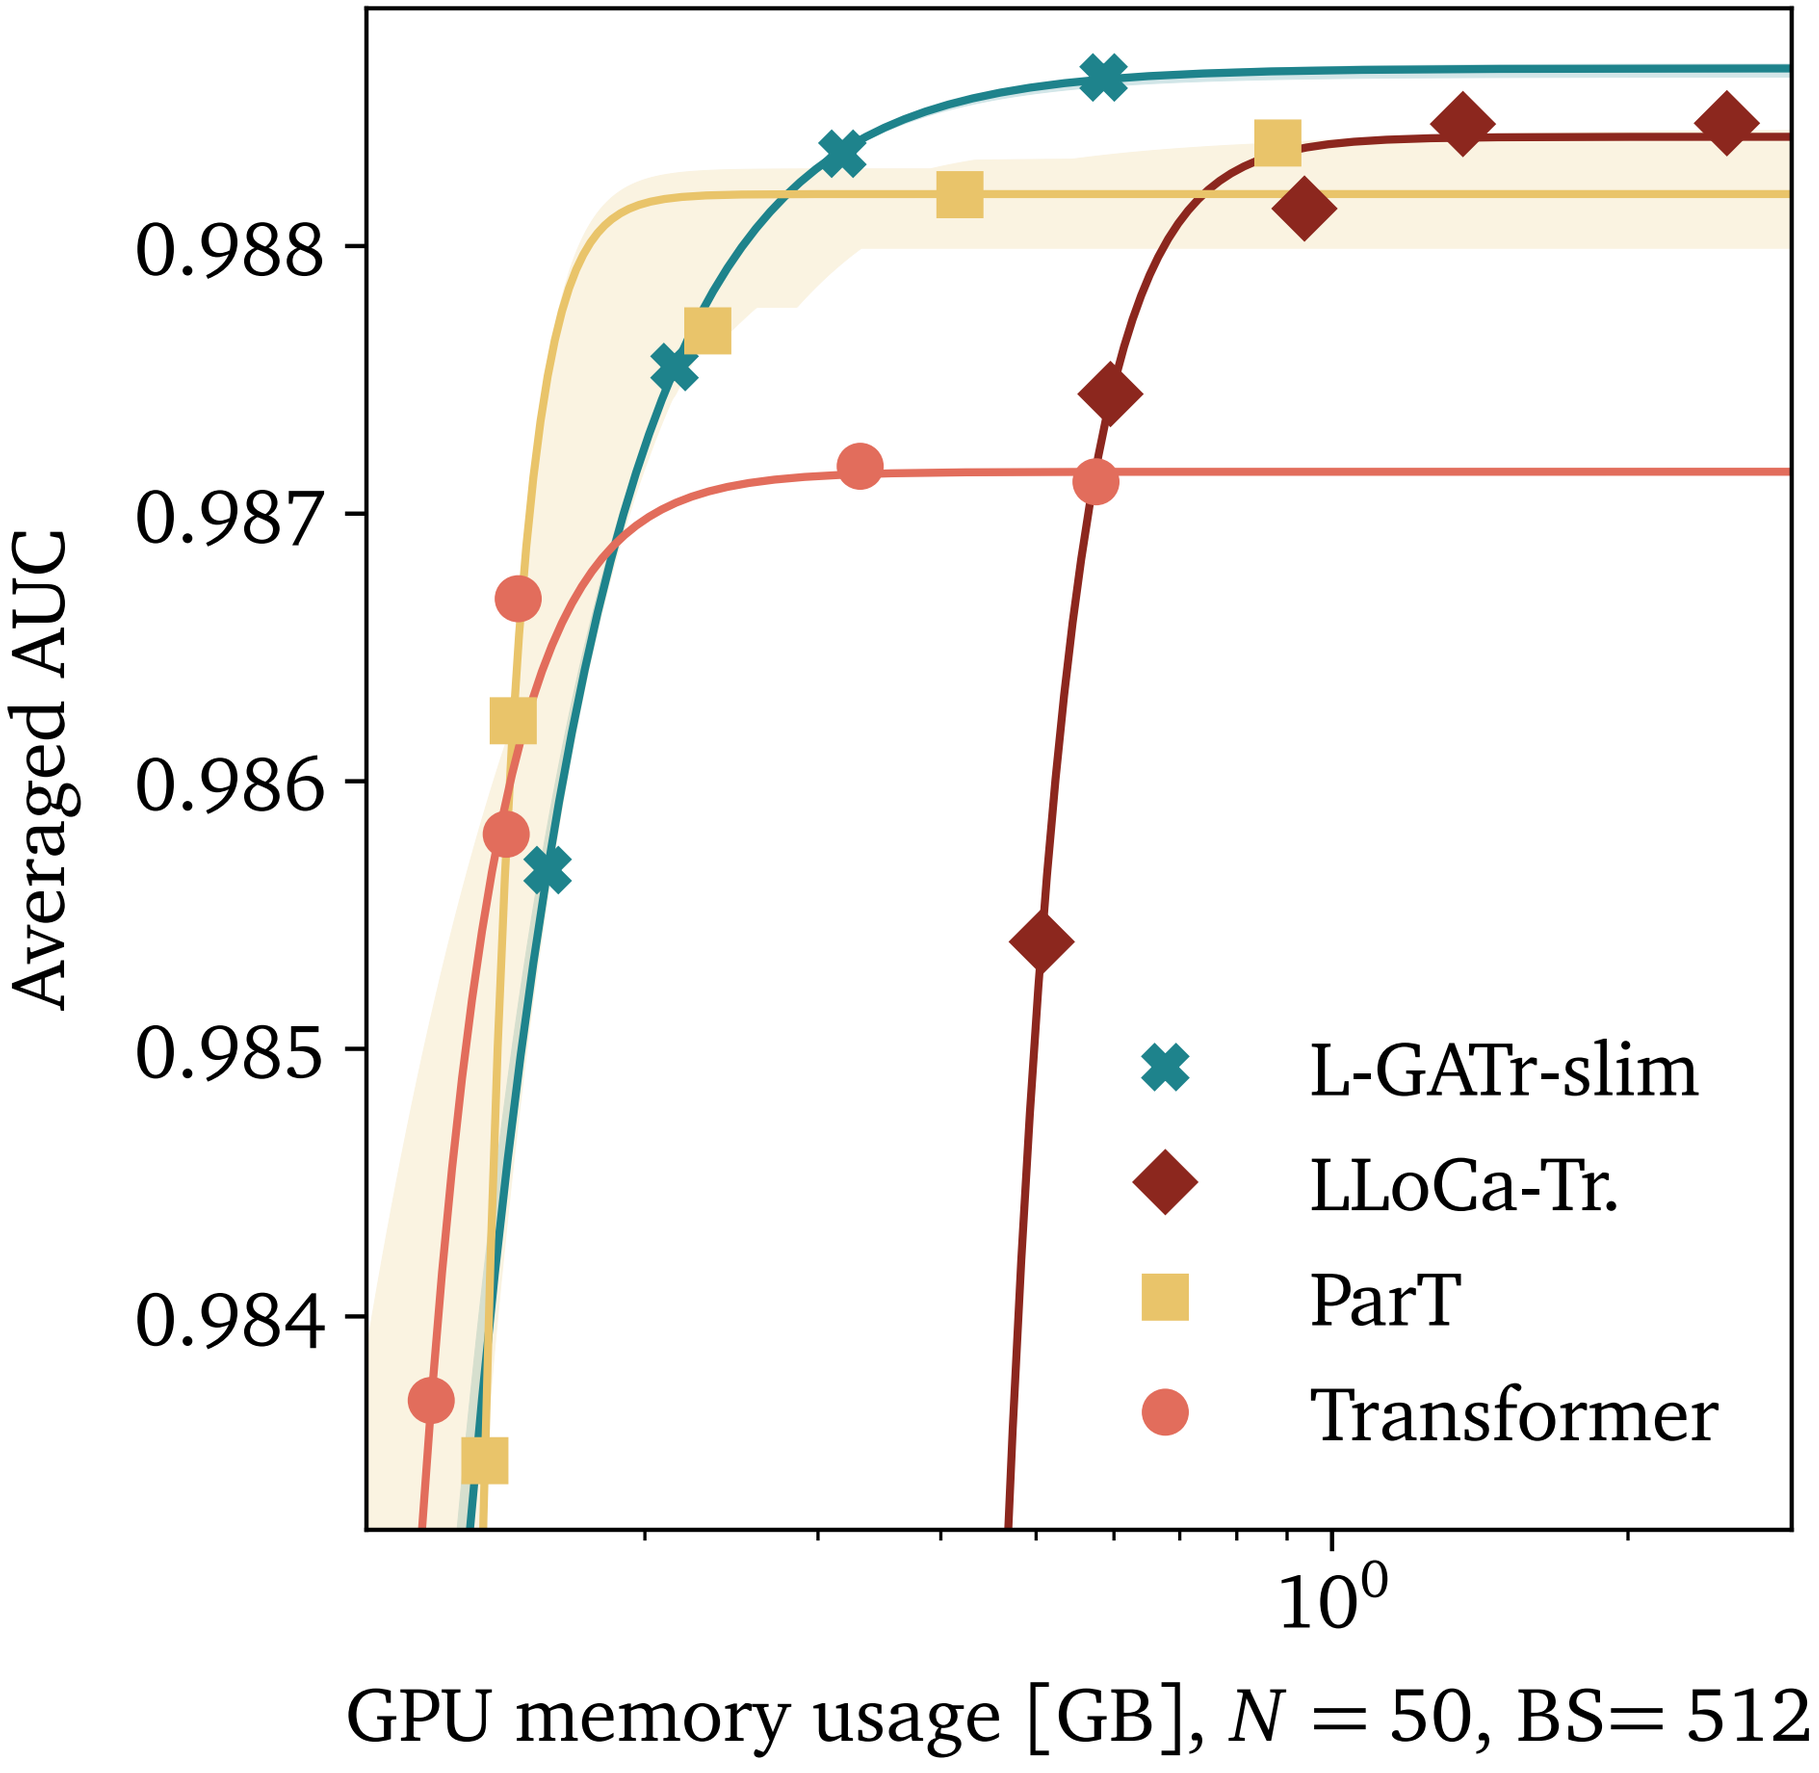
0.988
0.987
0.986
0.985
0.984
Averaged AUC
10^0
GPU memory usage [GB], $N = 50$, BS= 512
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

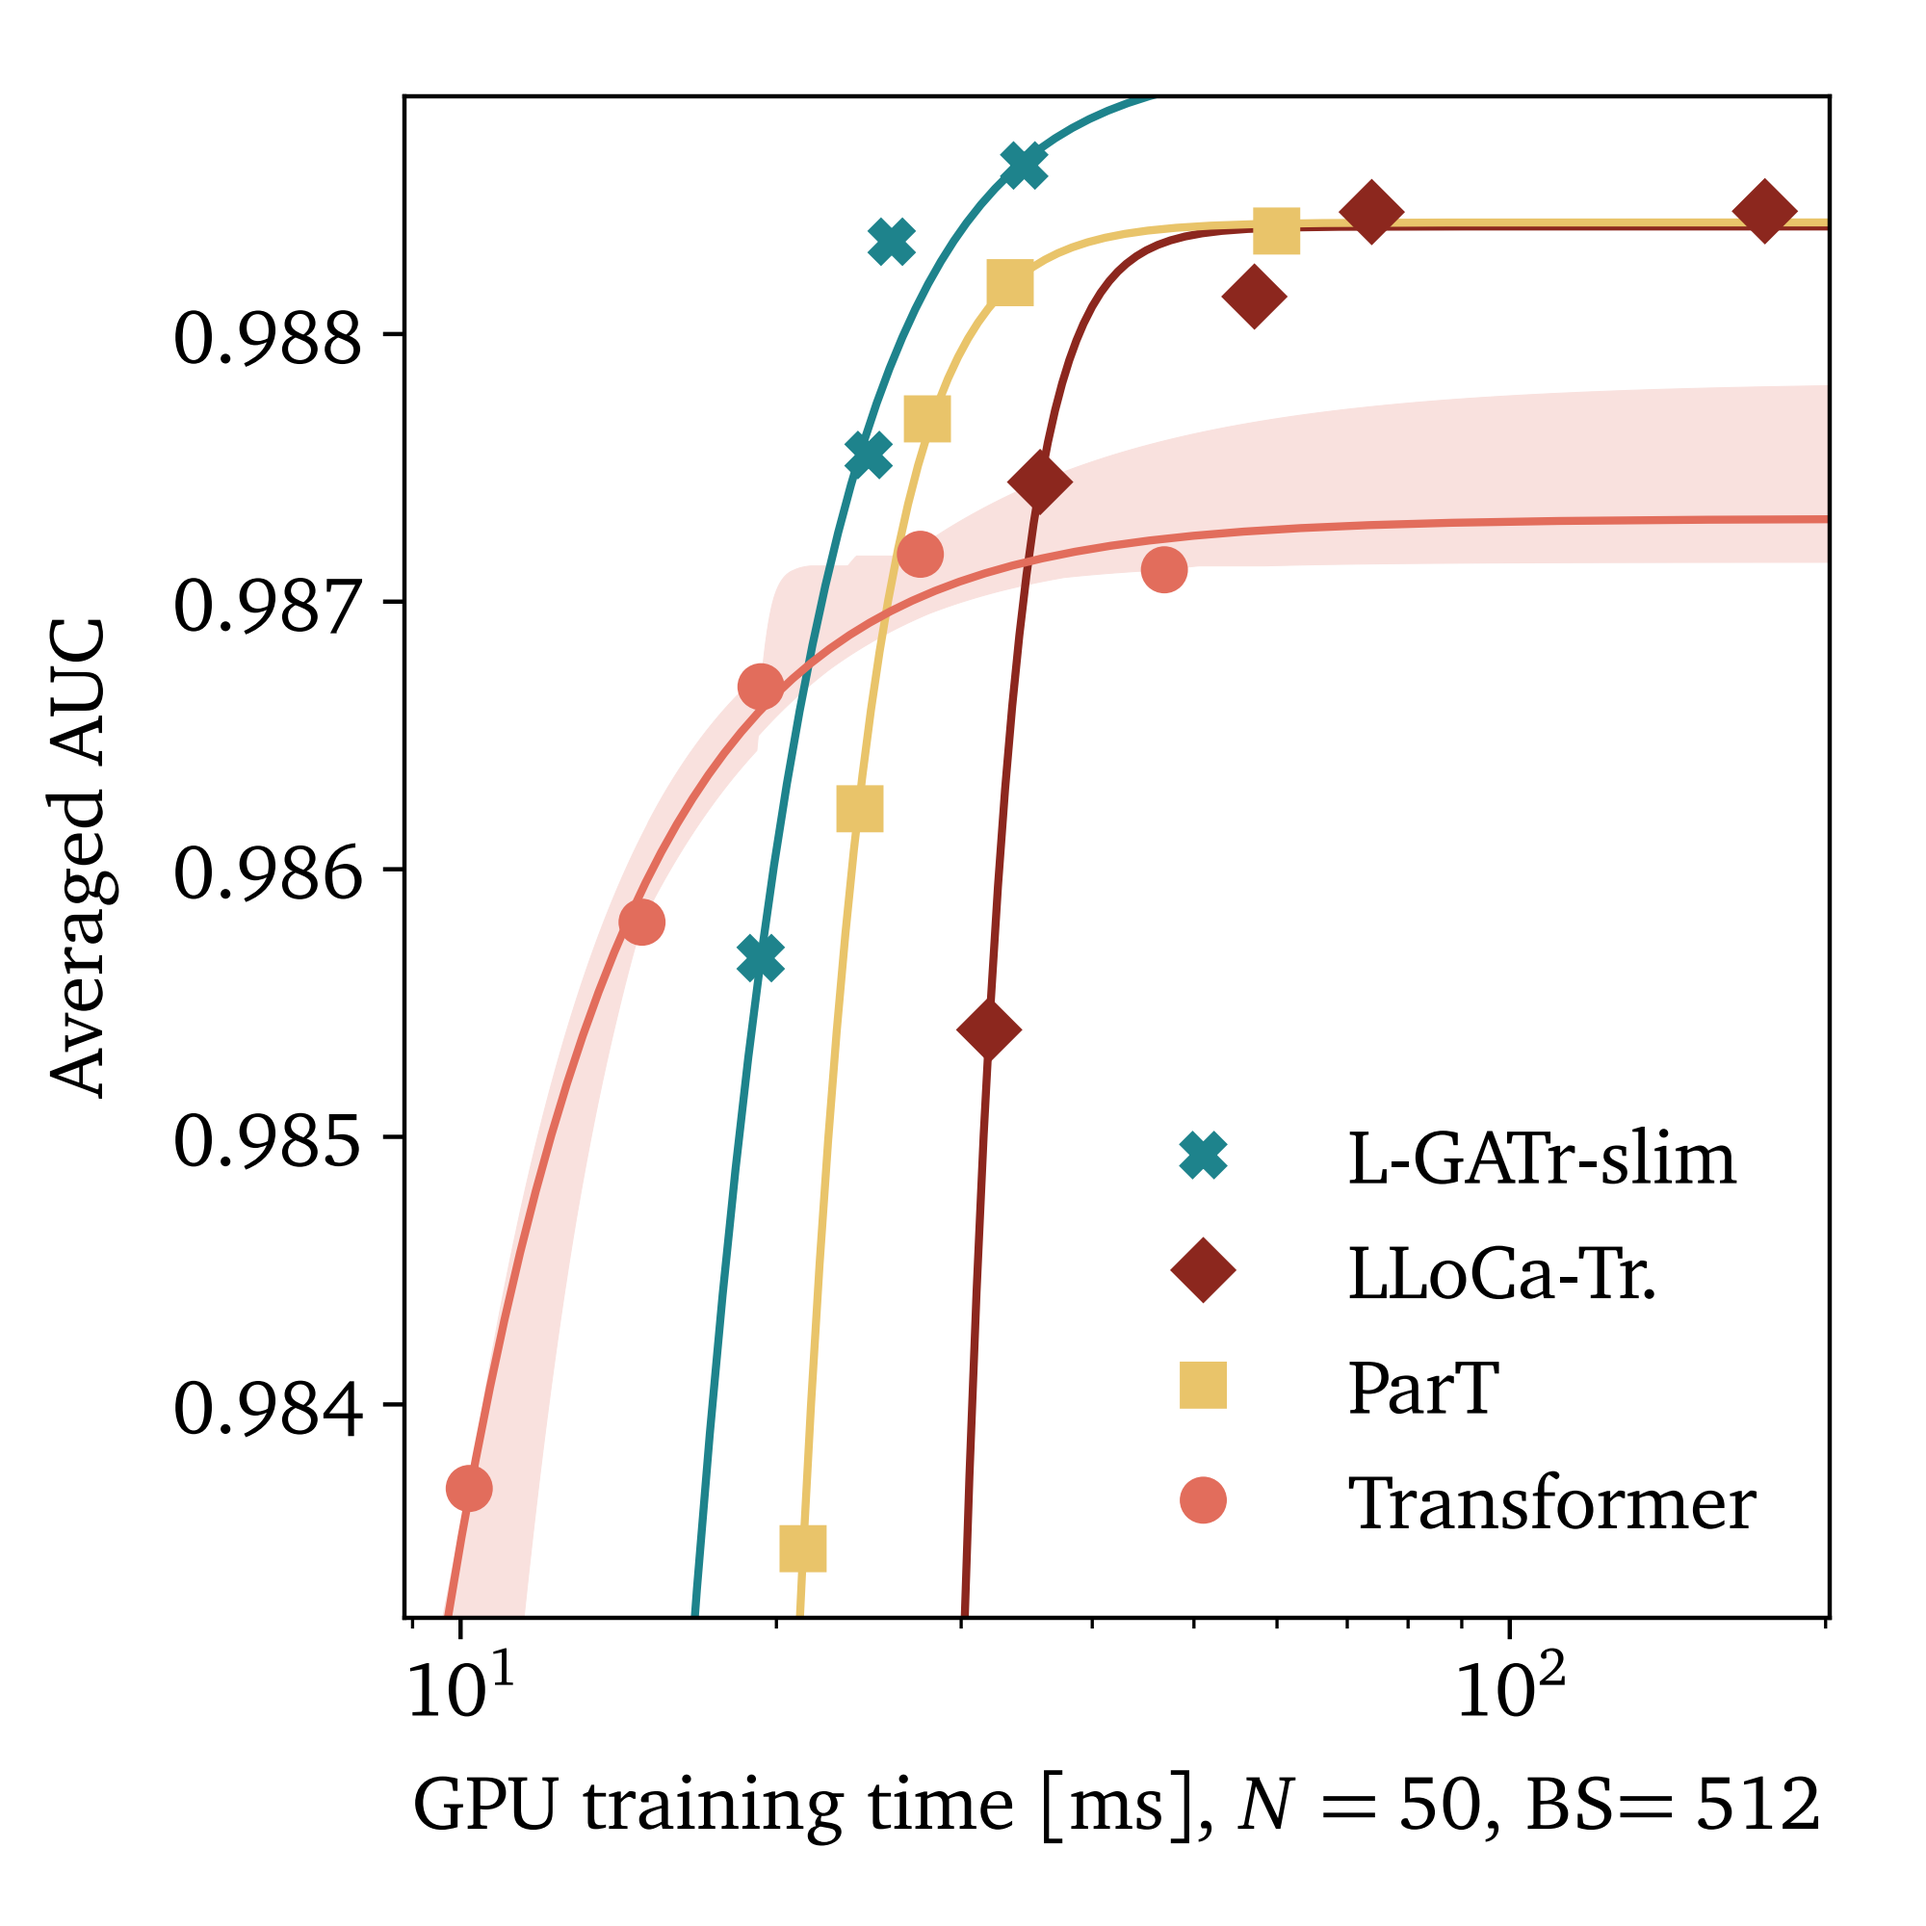
Averaged AUC
0.988
0.987
0.986
0.985
0.984
10^1
10^2
GPU training time [ms], $N = 50$, BS= 512
L-GATr-slim
LLoCa-Tr.
ParT
Transformer