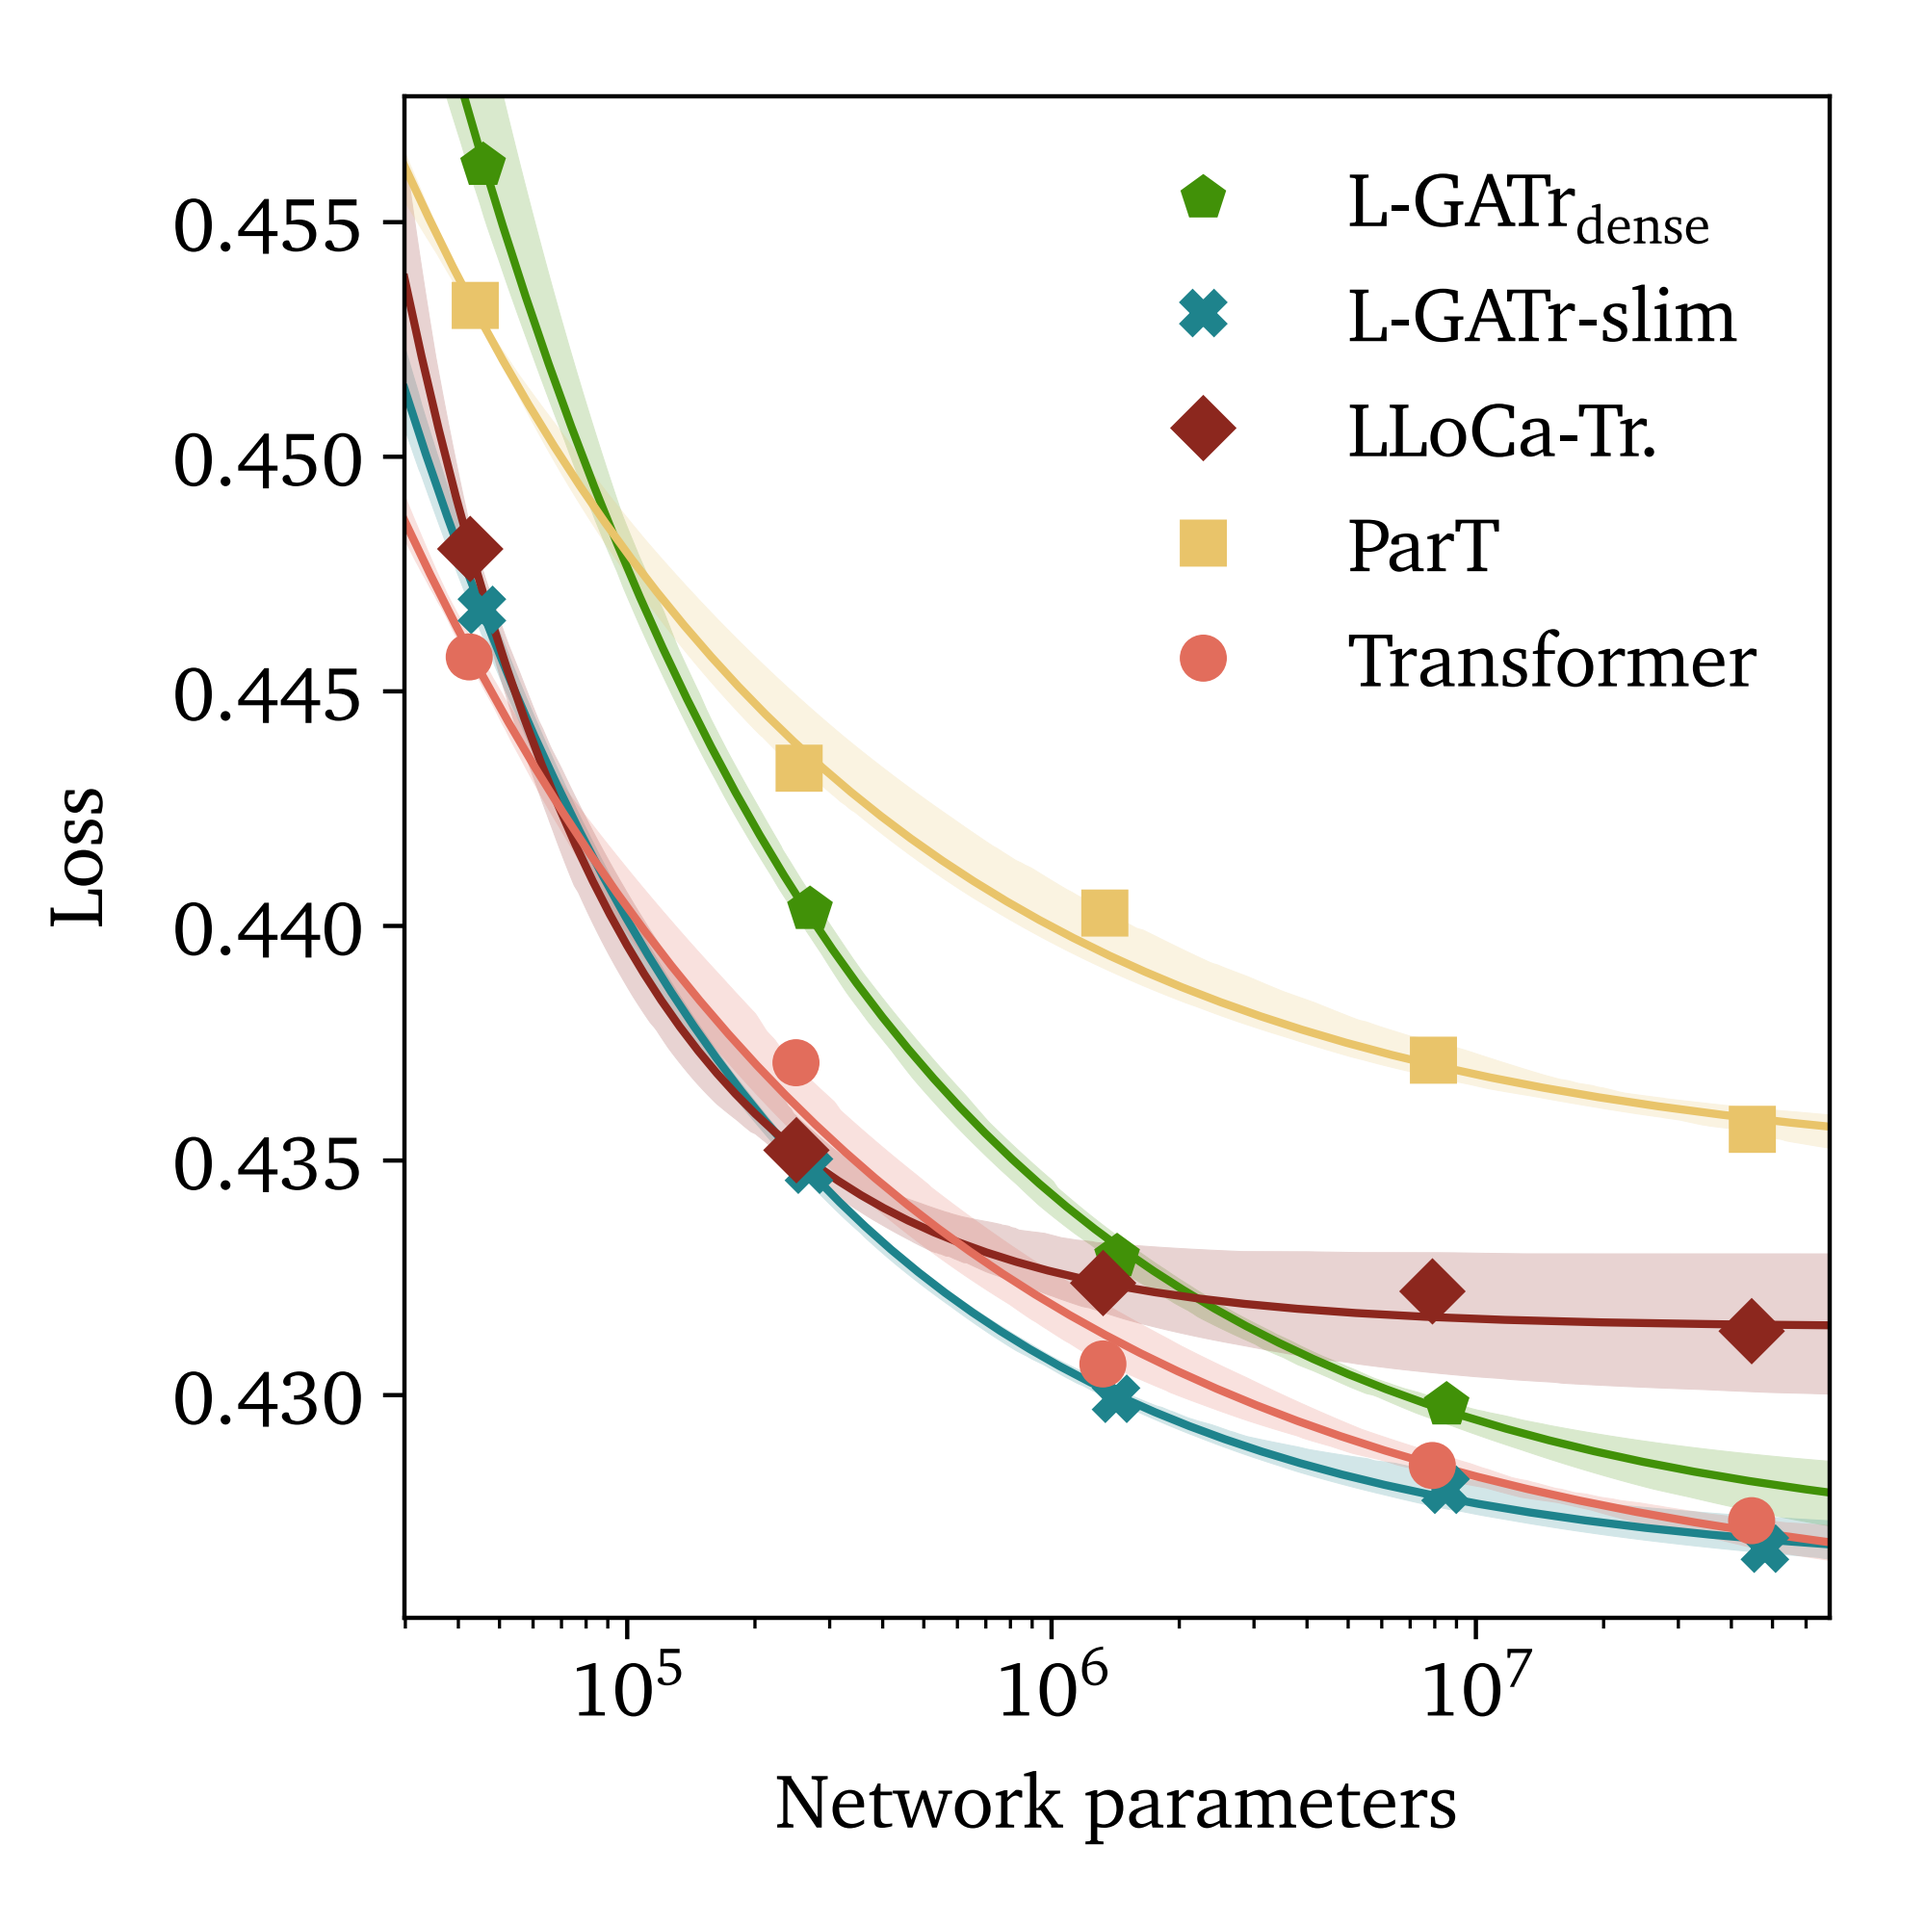
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.455
0.450
0.445
0.440
0.435
0.430
10^5
10^6
10^7
Network parameters

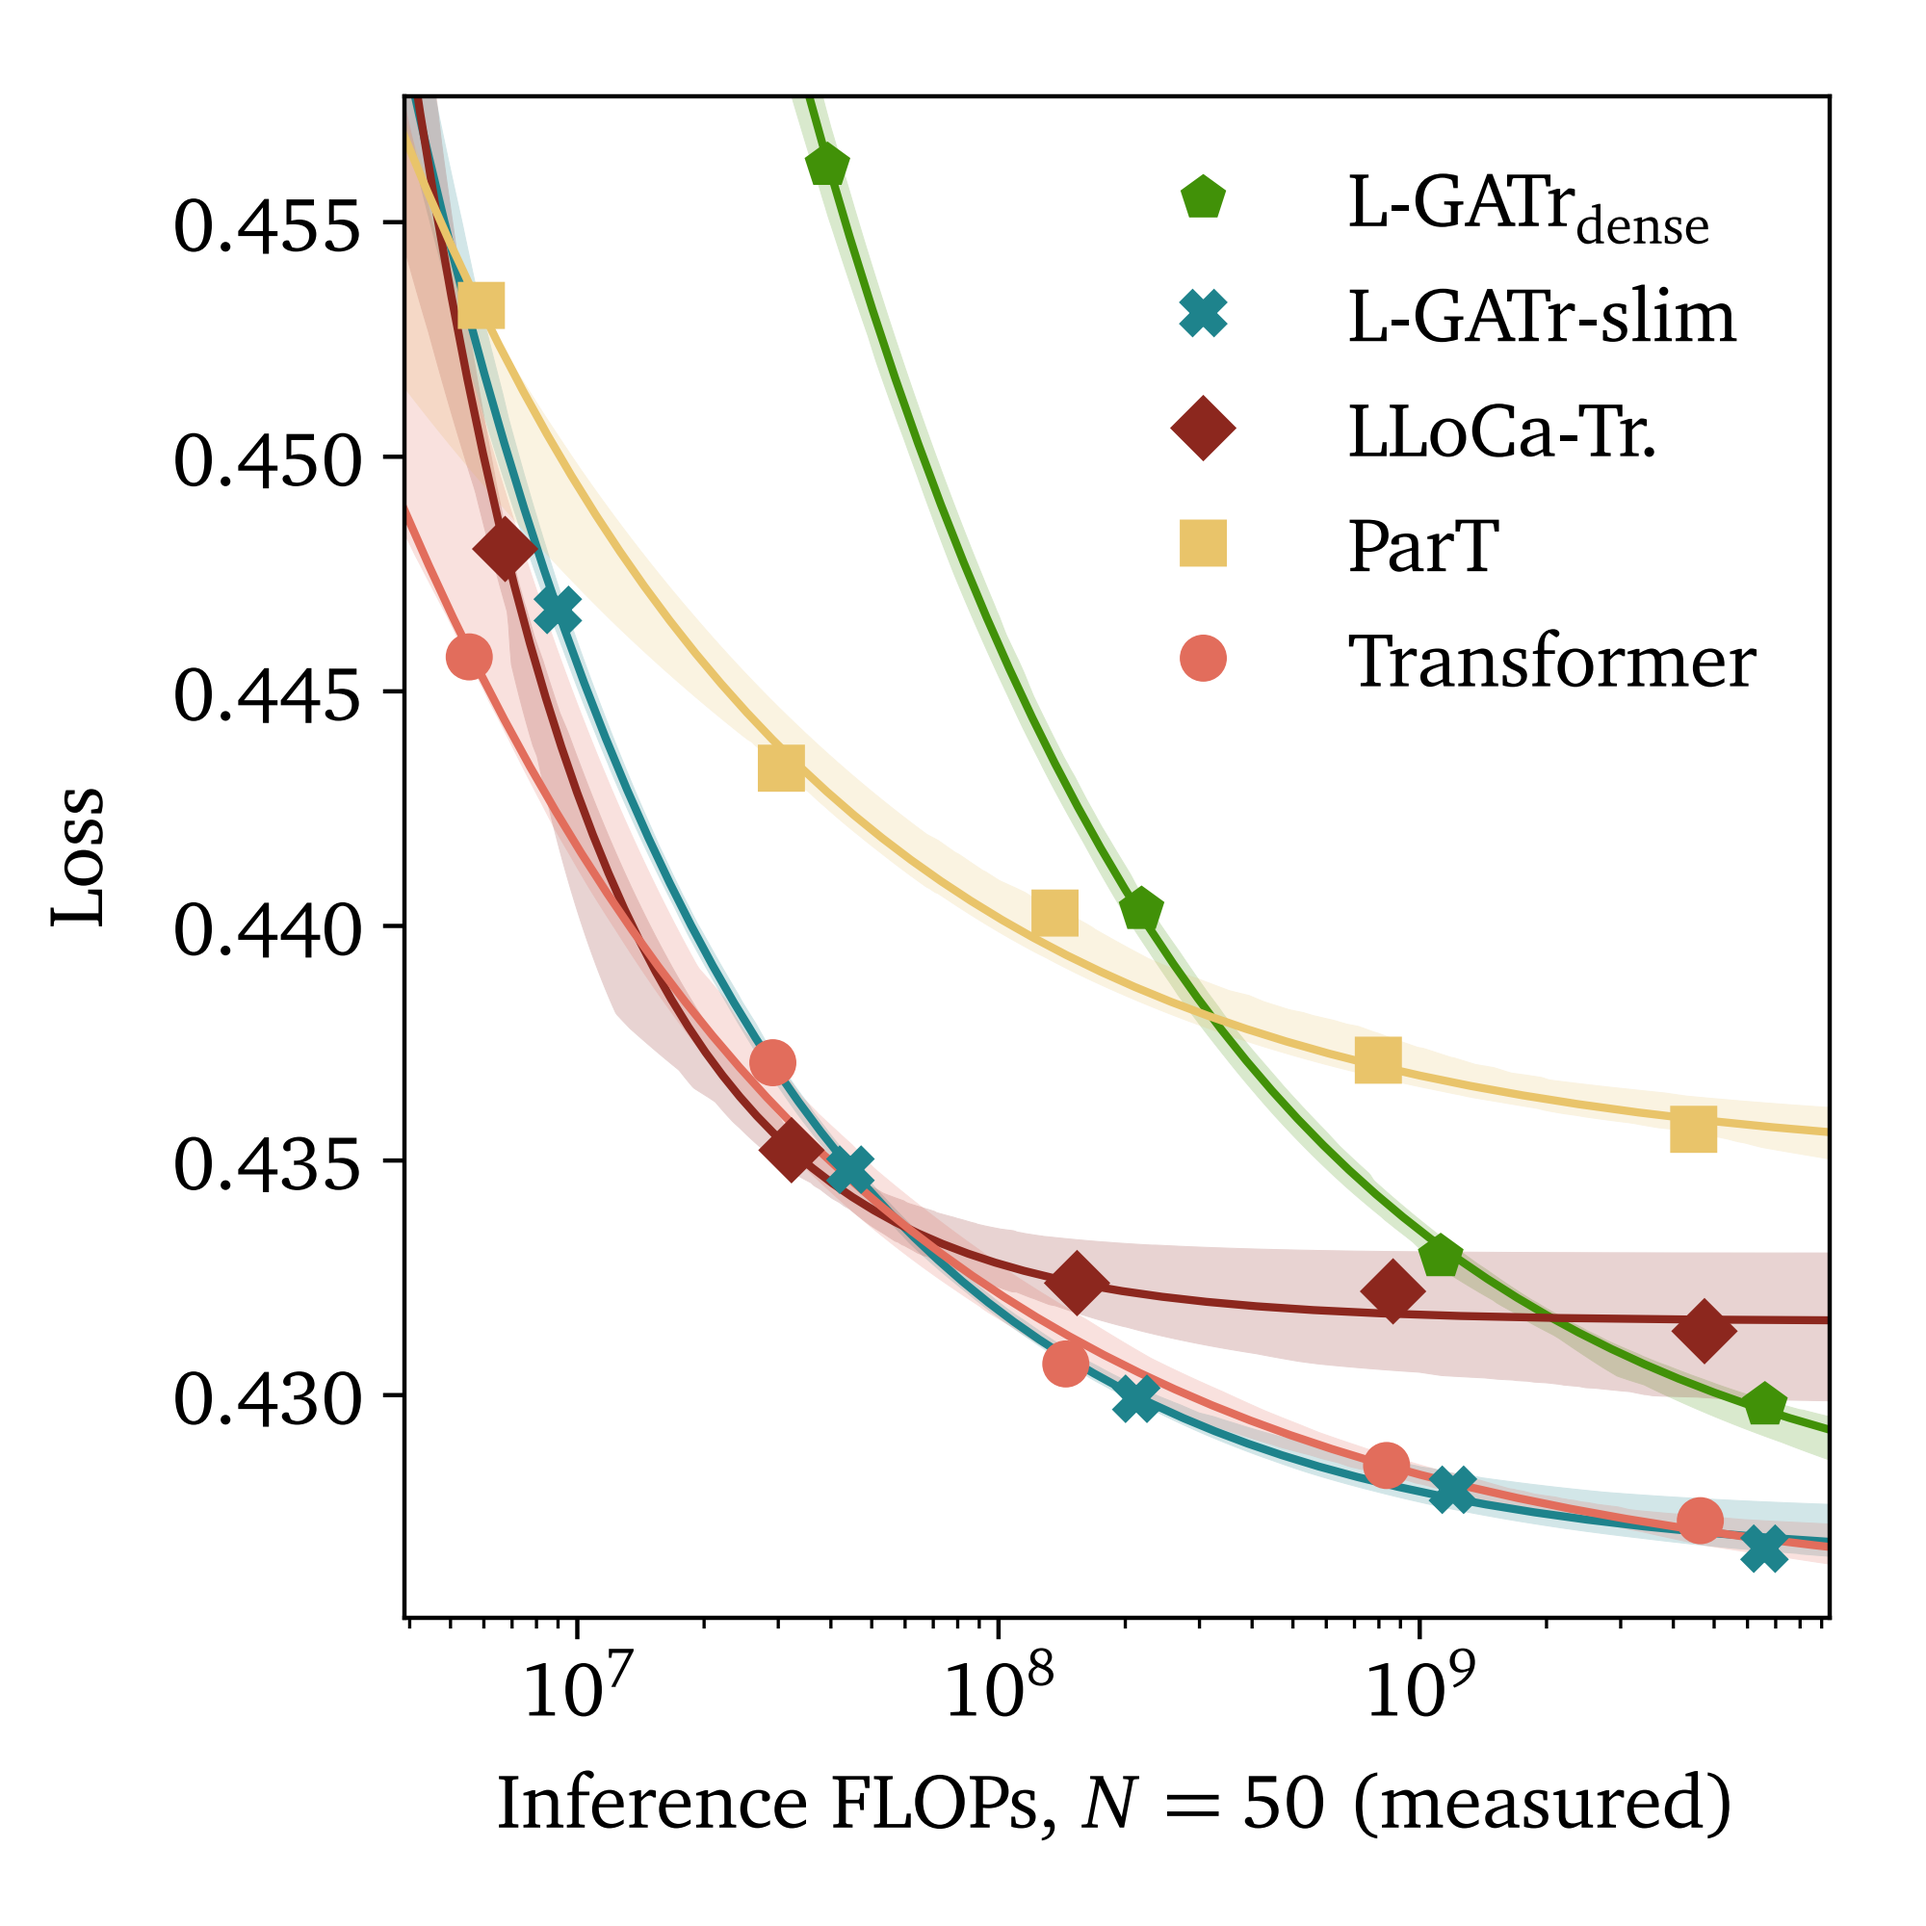
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.455
0.450
0.445
0.440
0.435
0.430
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (measured)

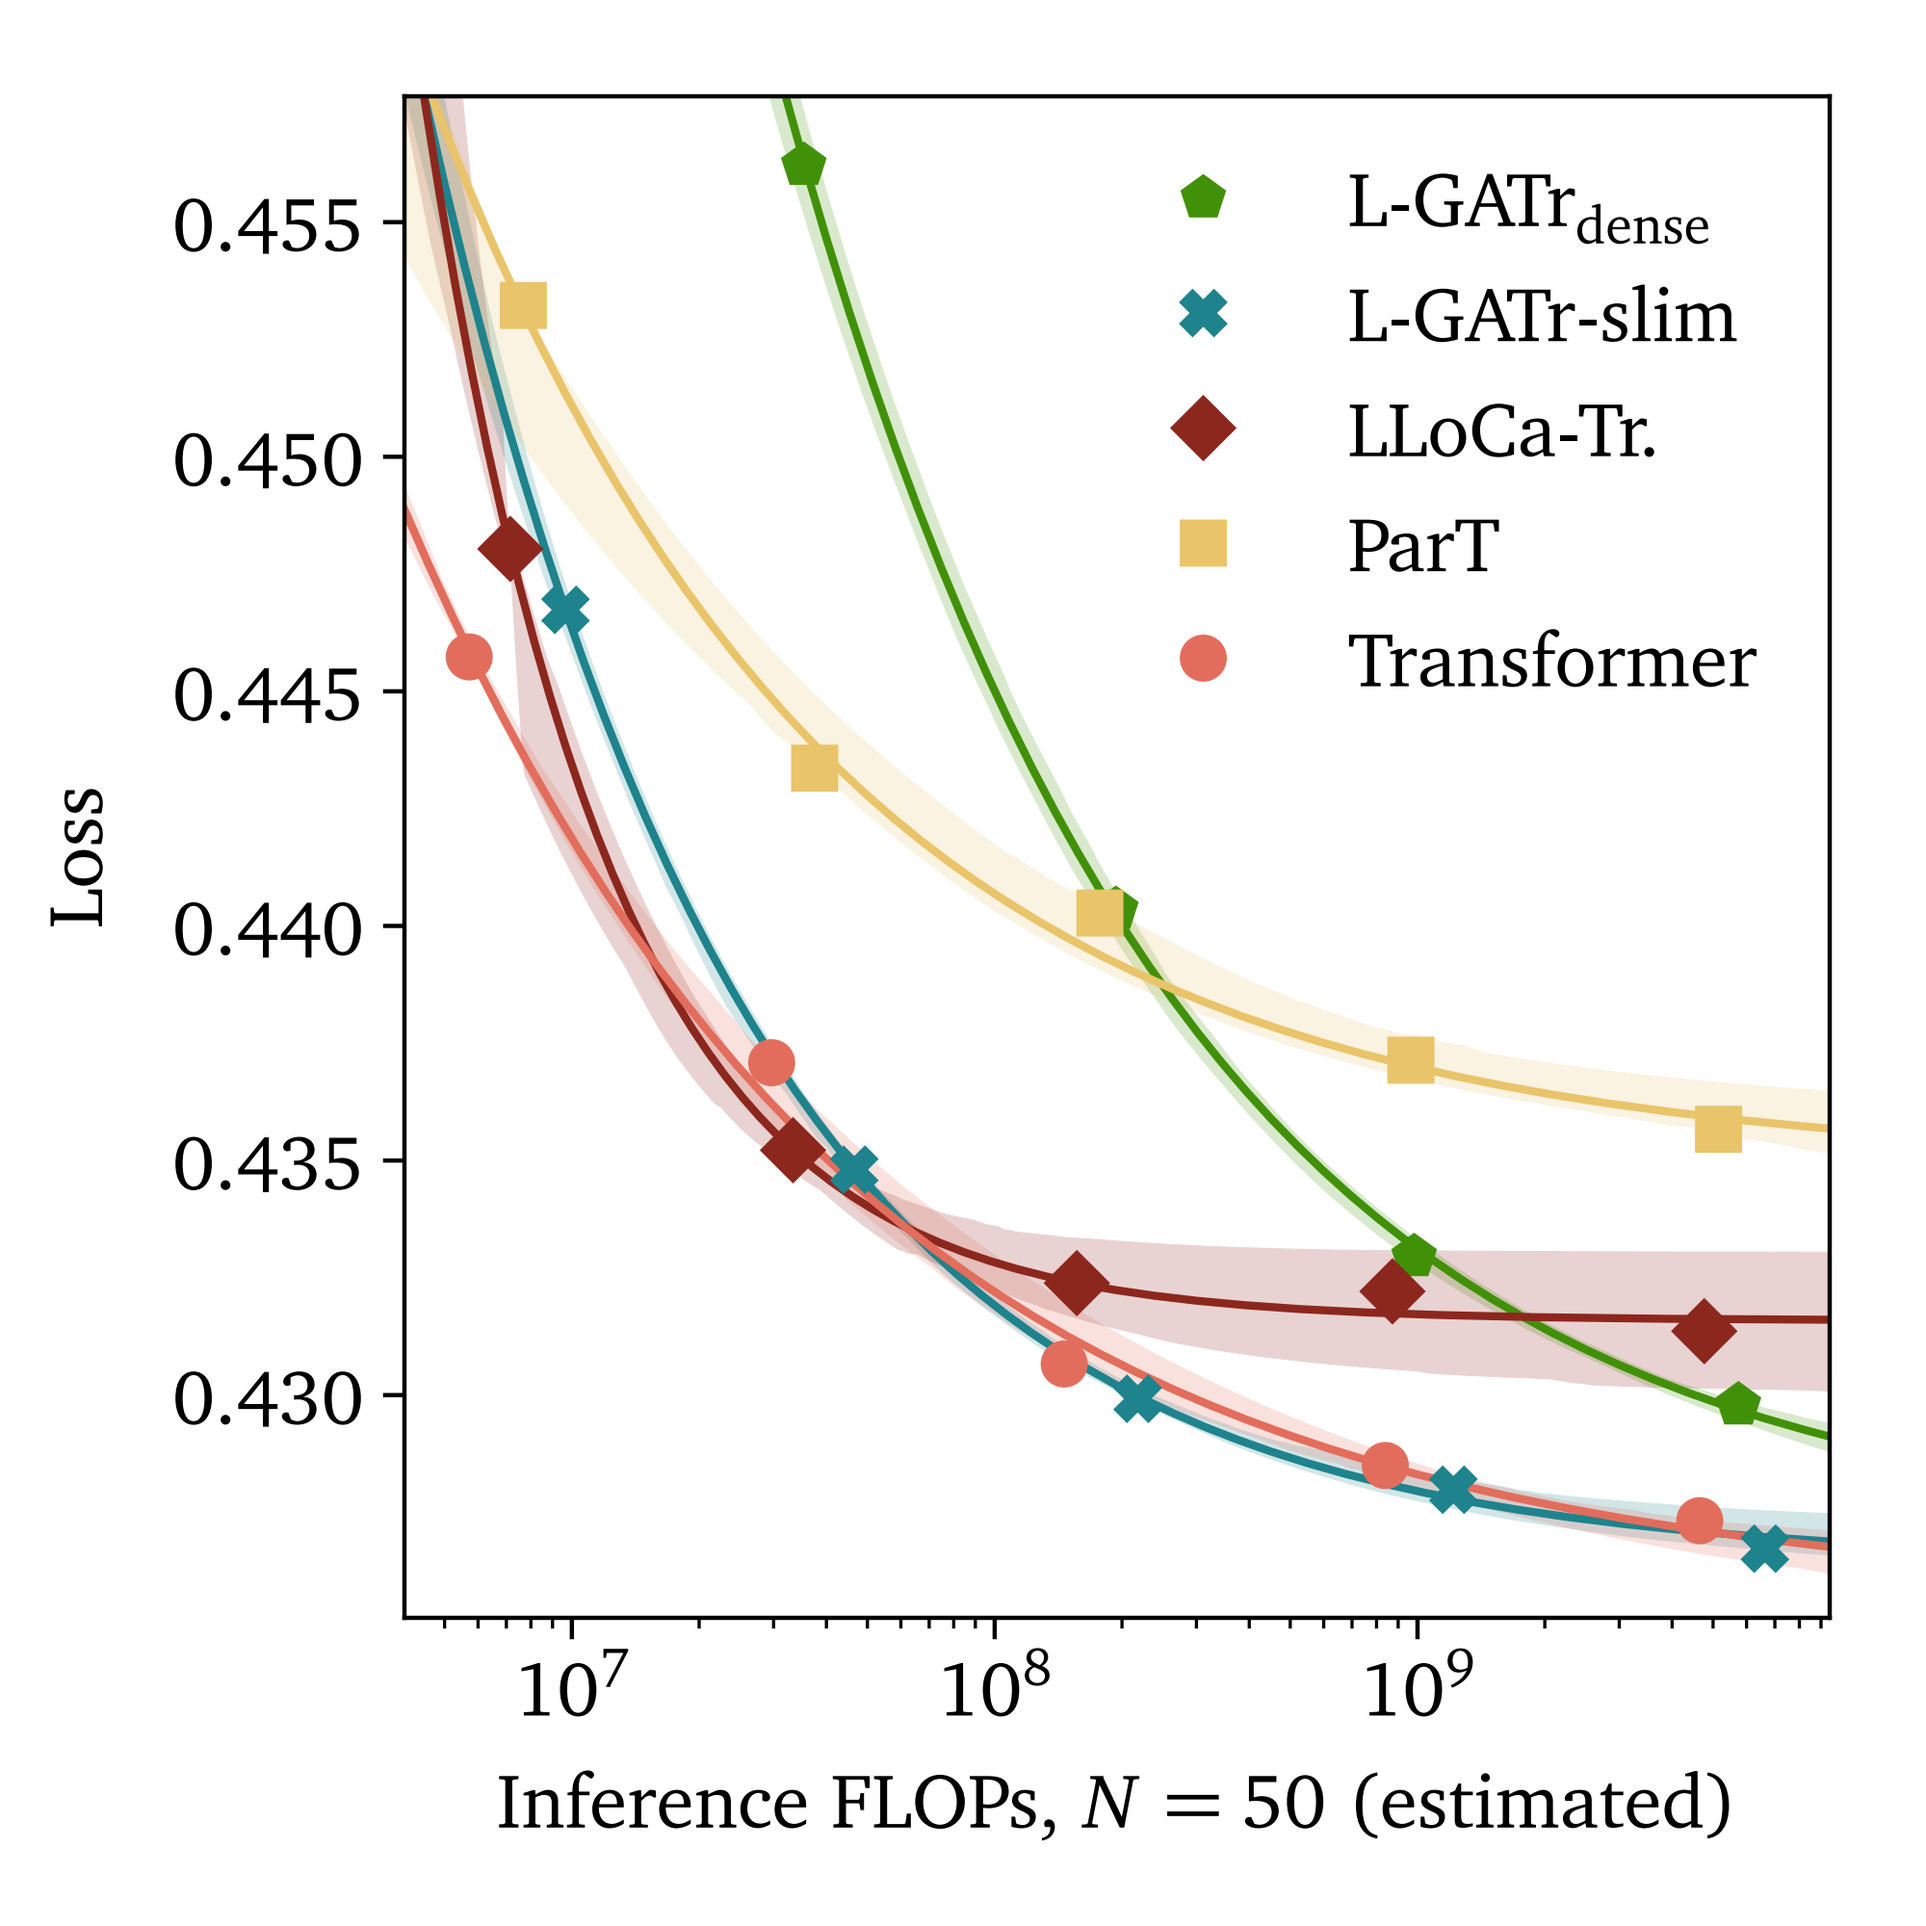
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.455
0.450
0.445
0.440
0.435
0.430
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (estimated)

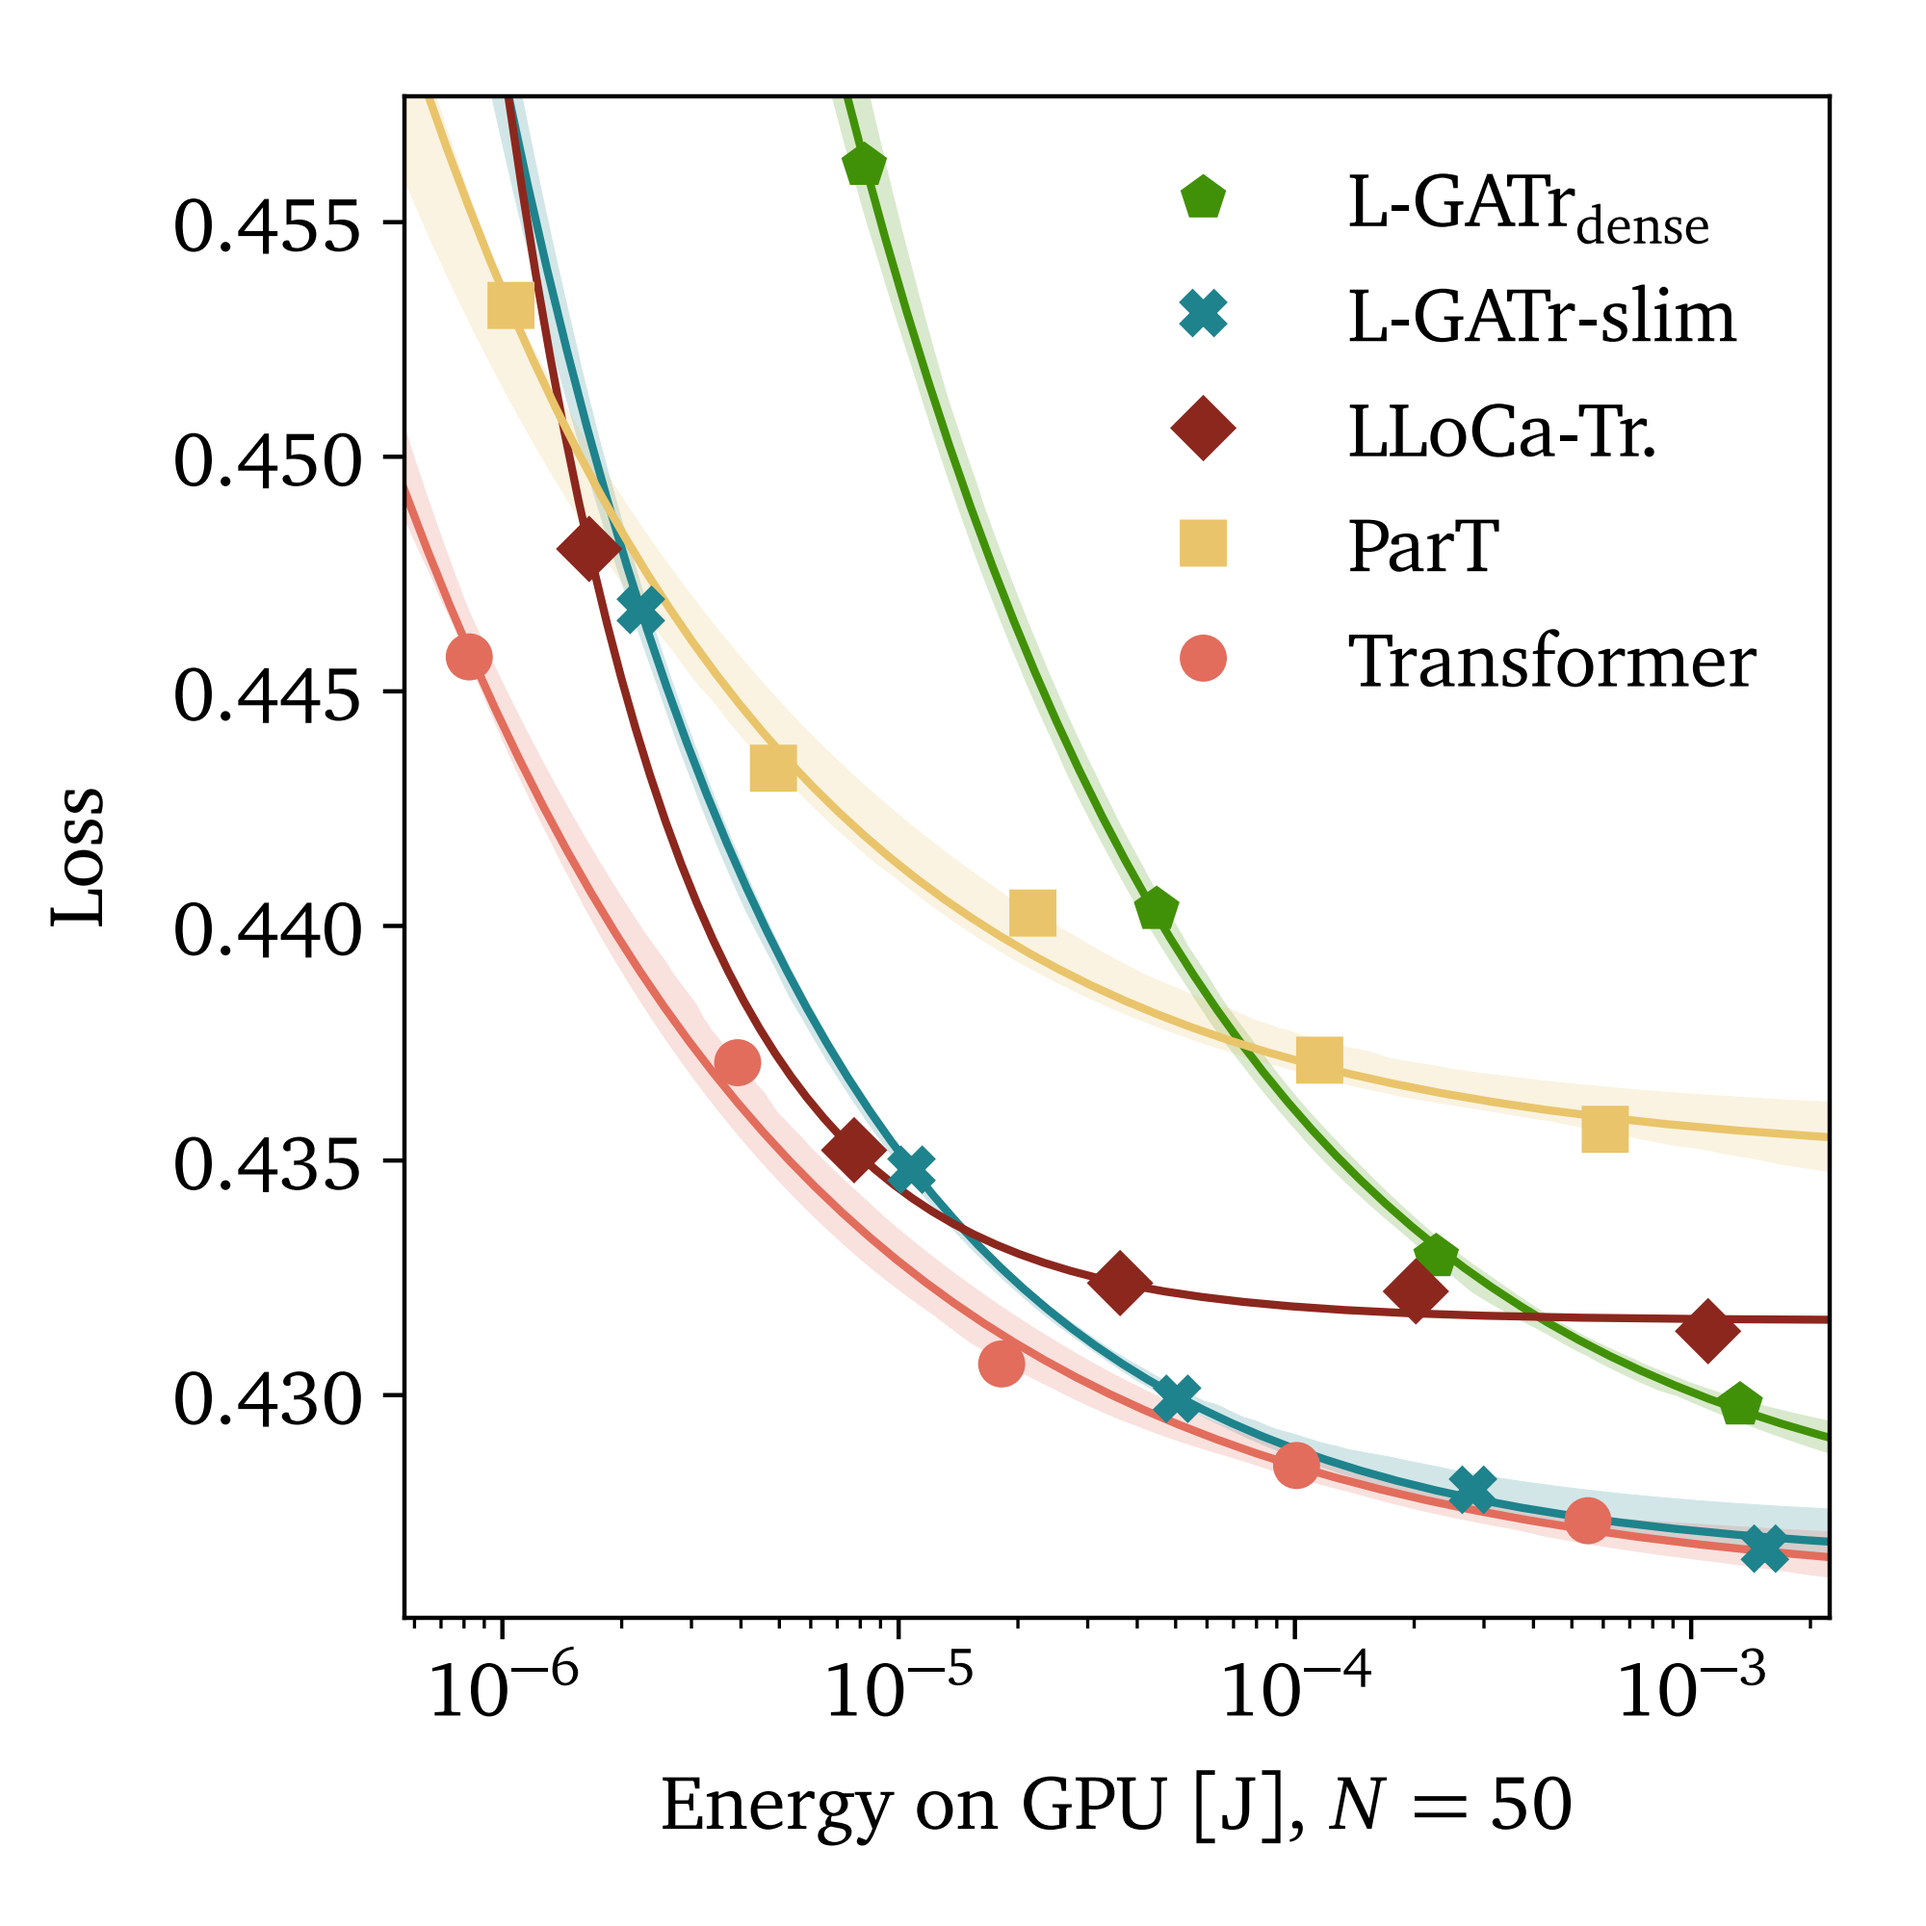
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.455
0.450
0.445
0.440
0.435
0.430
10^{-6}
10^{-5}
10^{-4}
10^{-3}
Energy on GPU [J], $N = 50$

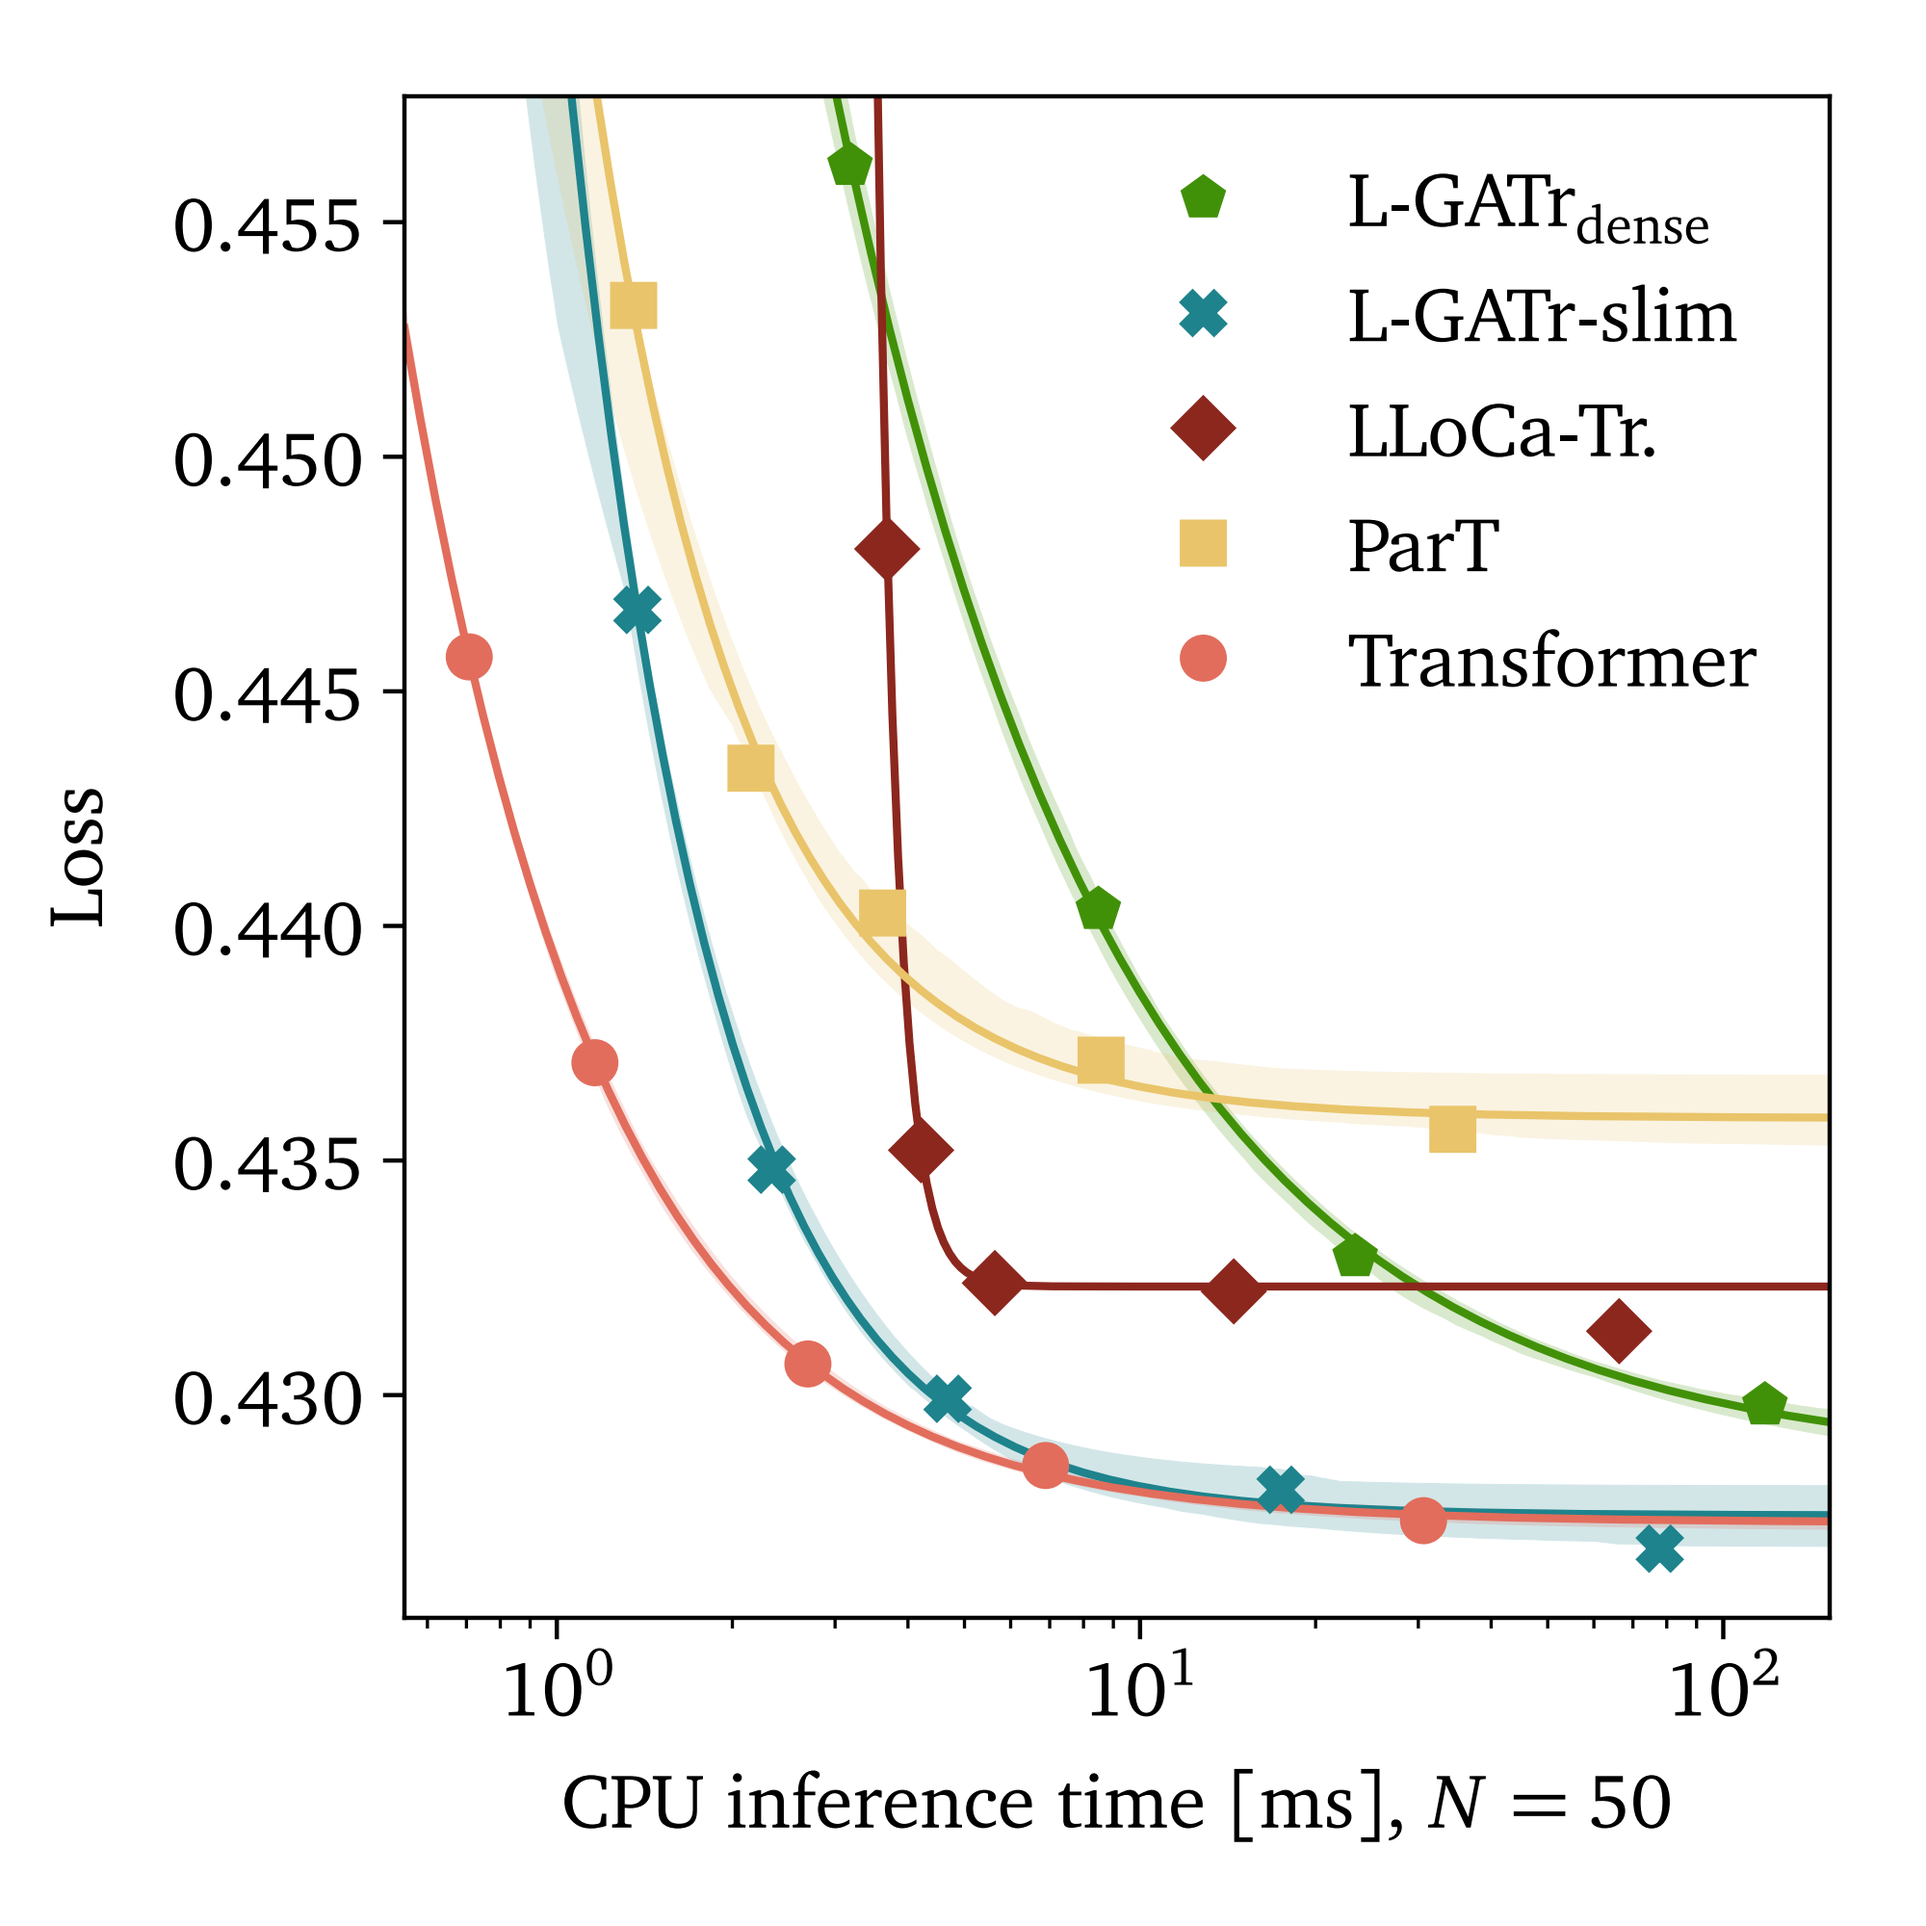
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.455
0.450
0.445
0.440
0.435
0.430
10^0
10^1
10^2
CPU inference time [ms], $N = 50$

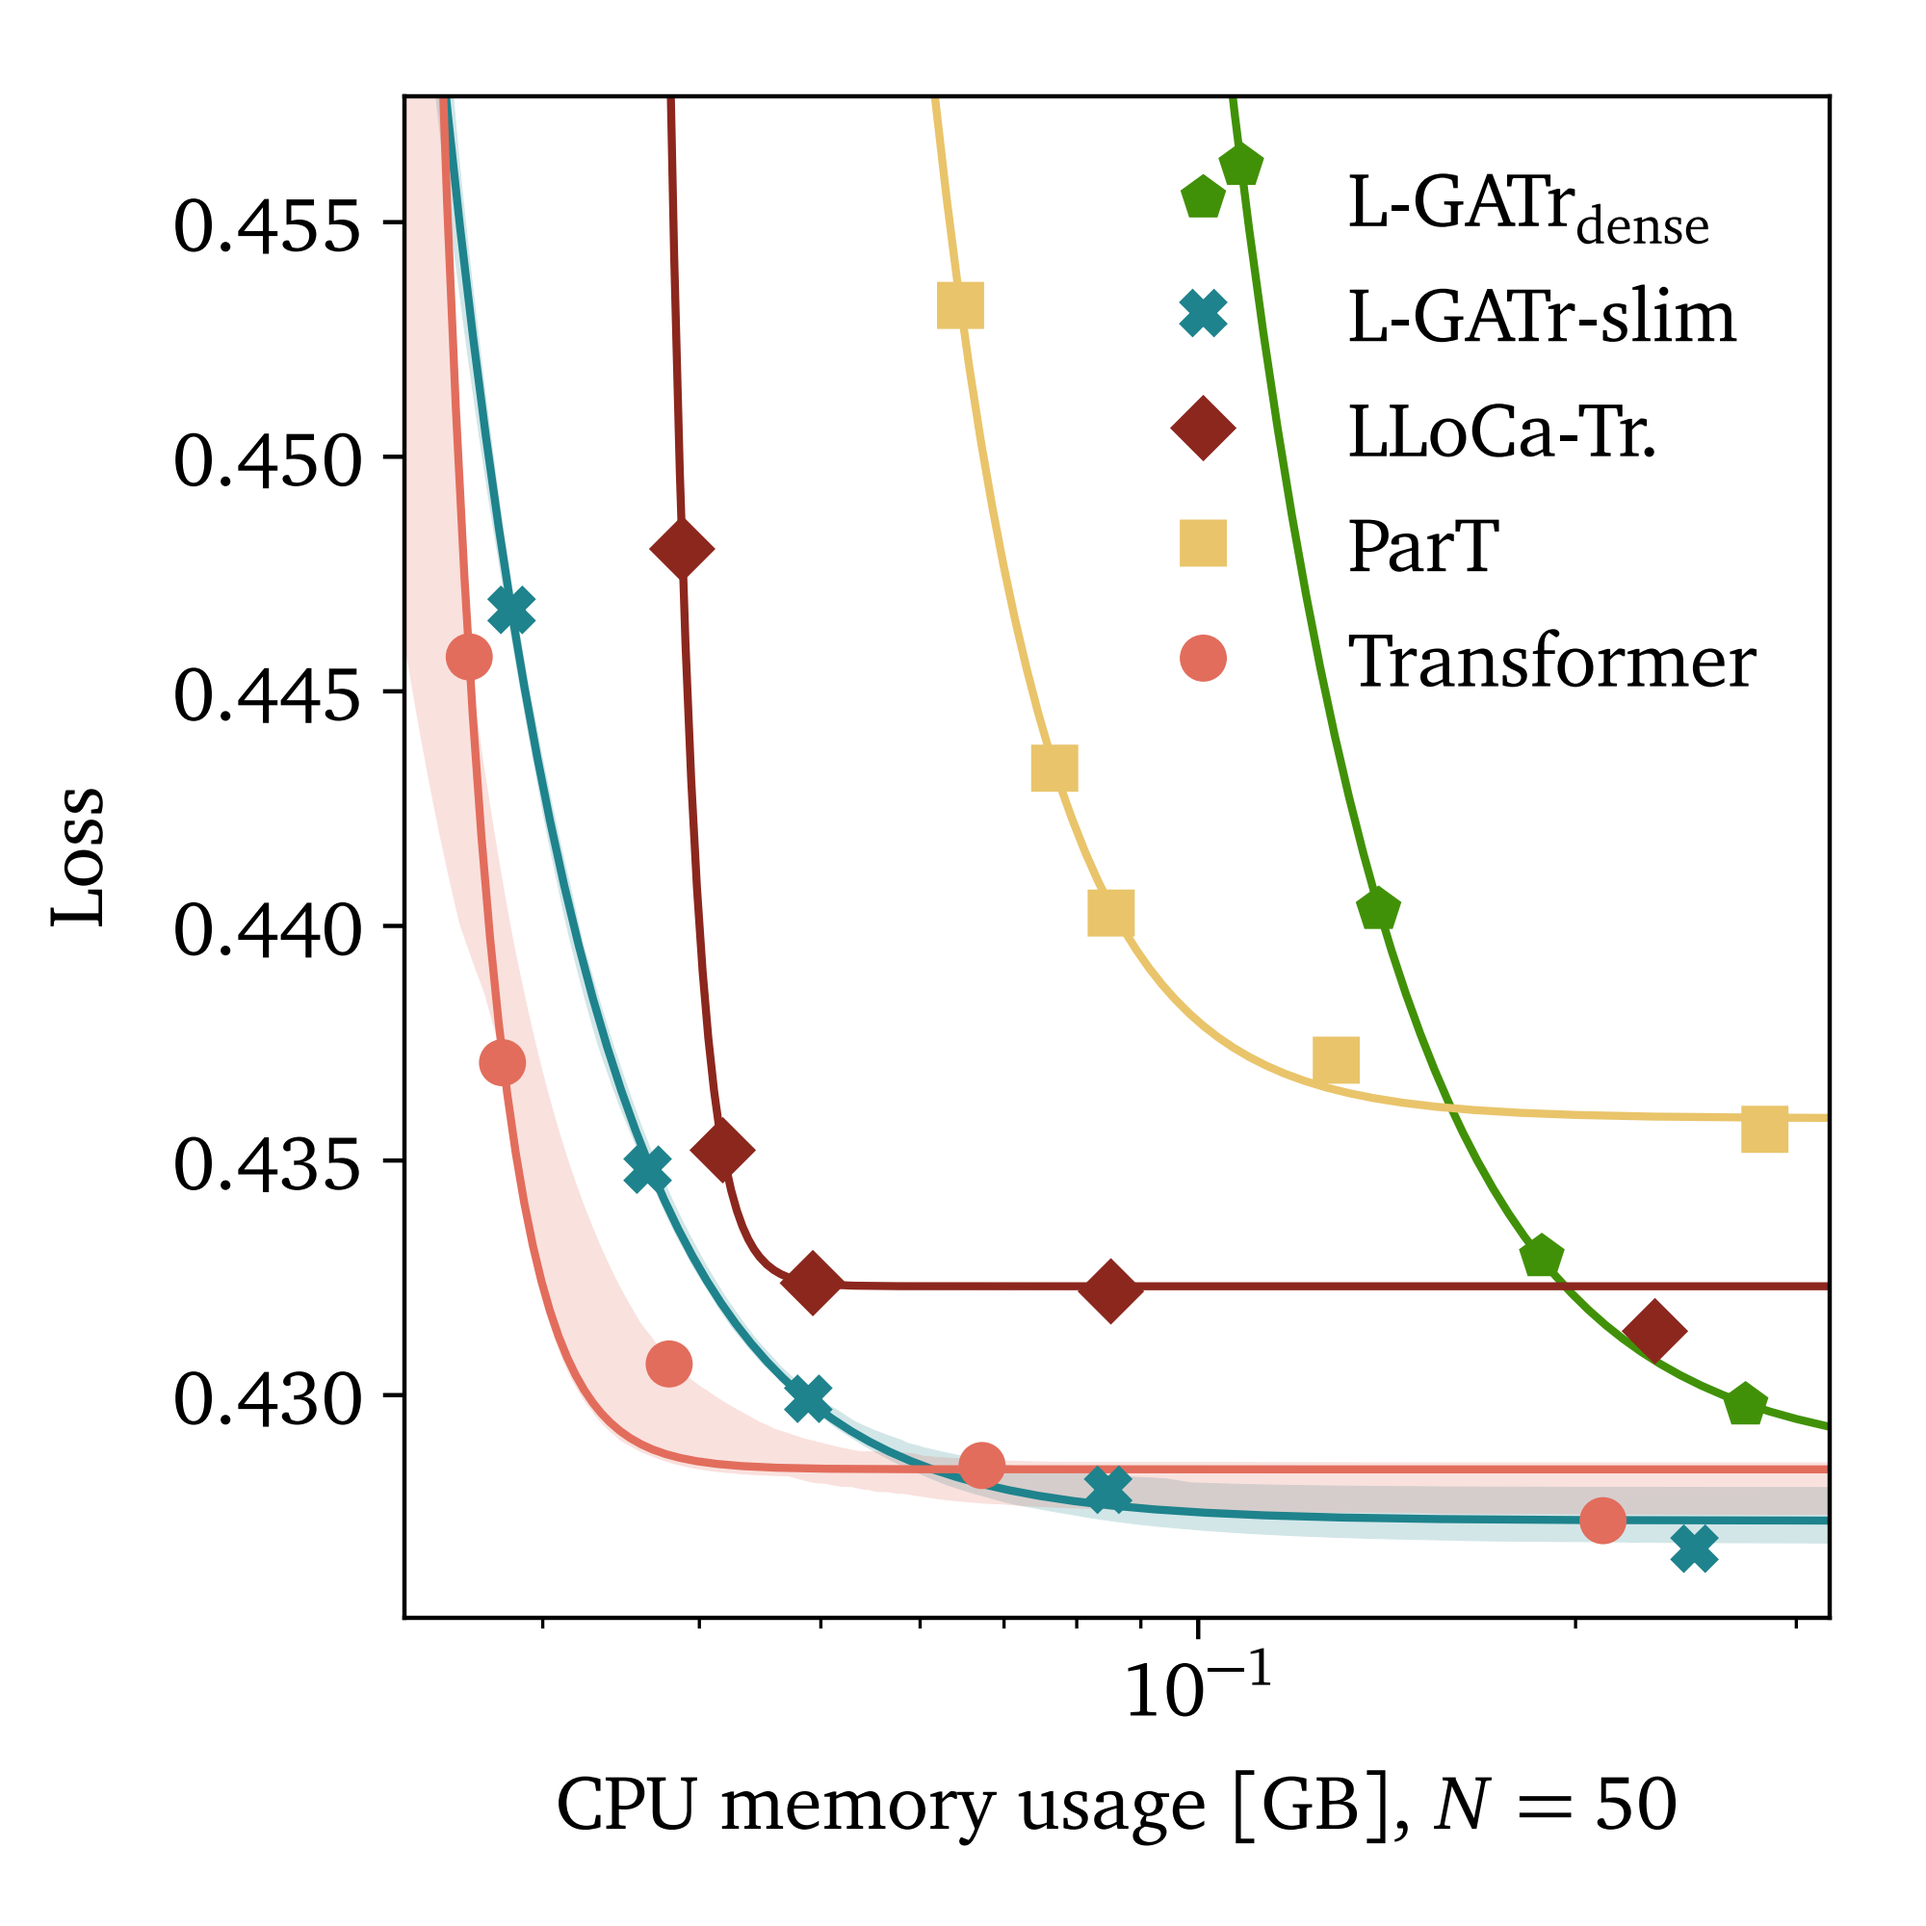
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.455
0.450
0.445
0.440
0.435
0.430
10^{-1}
CPU memory usage [GB], $N = 50$

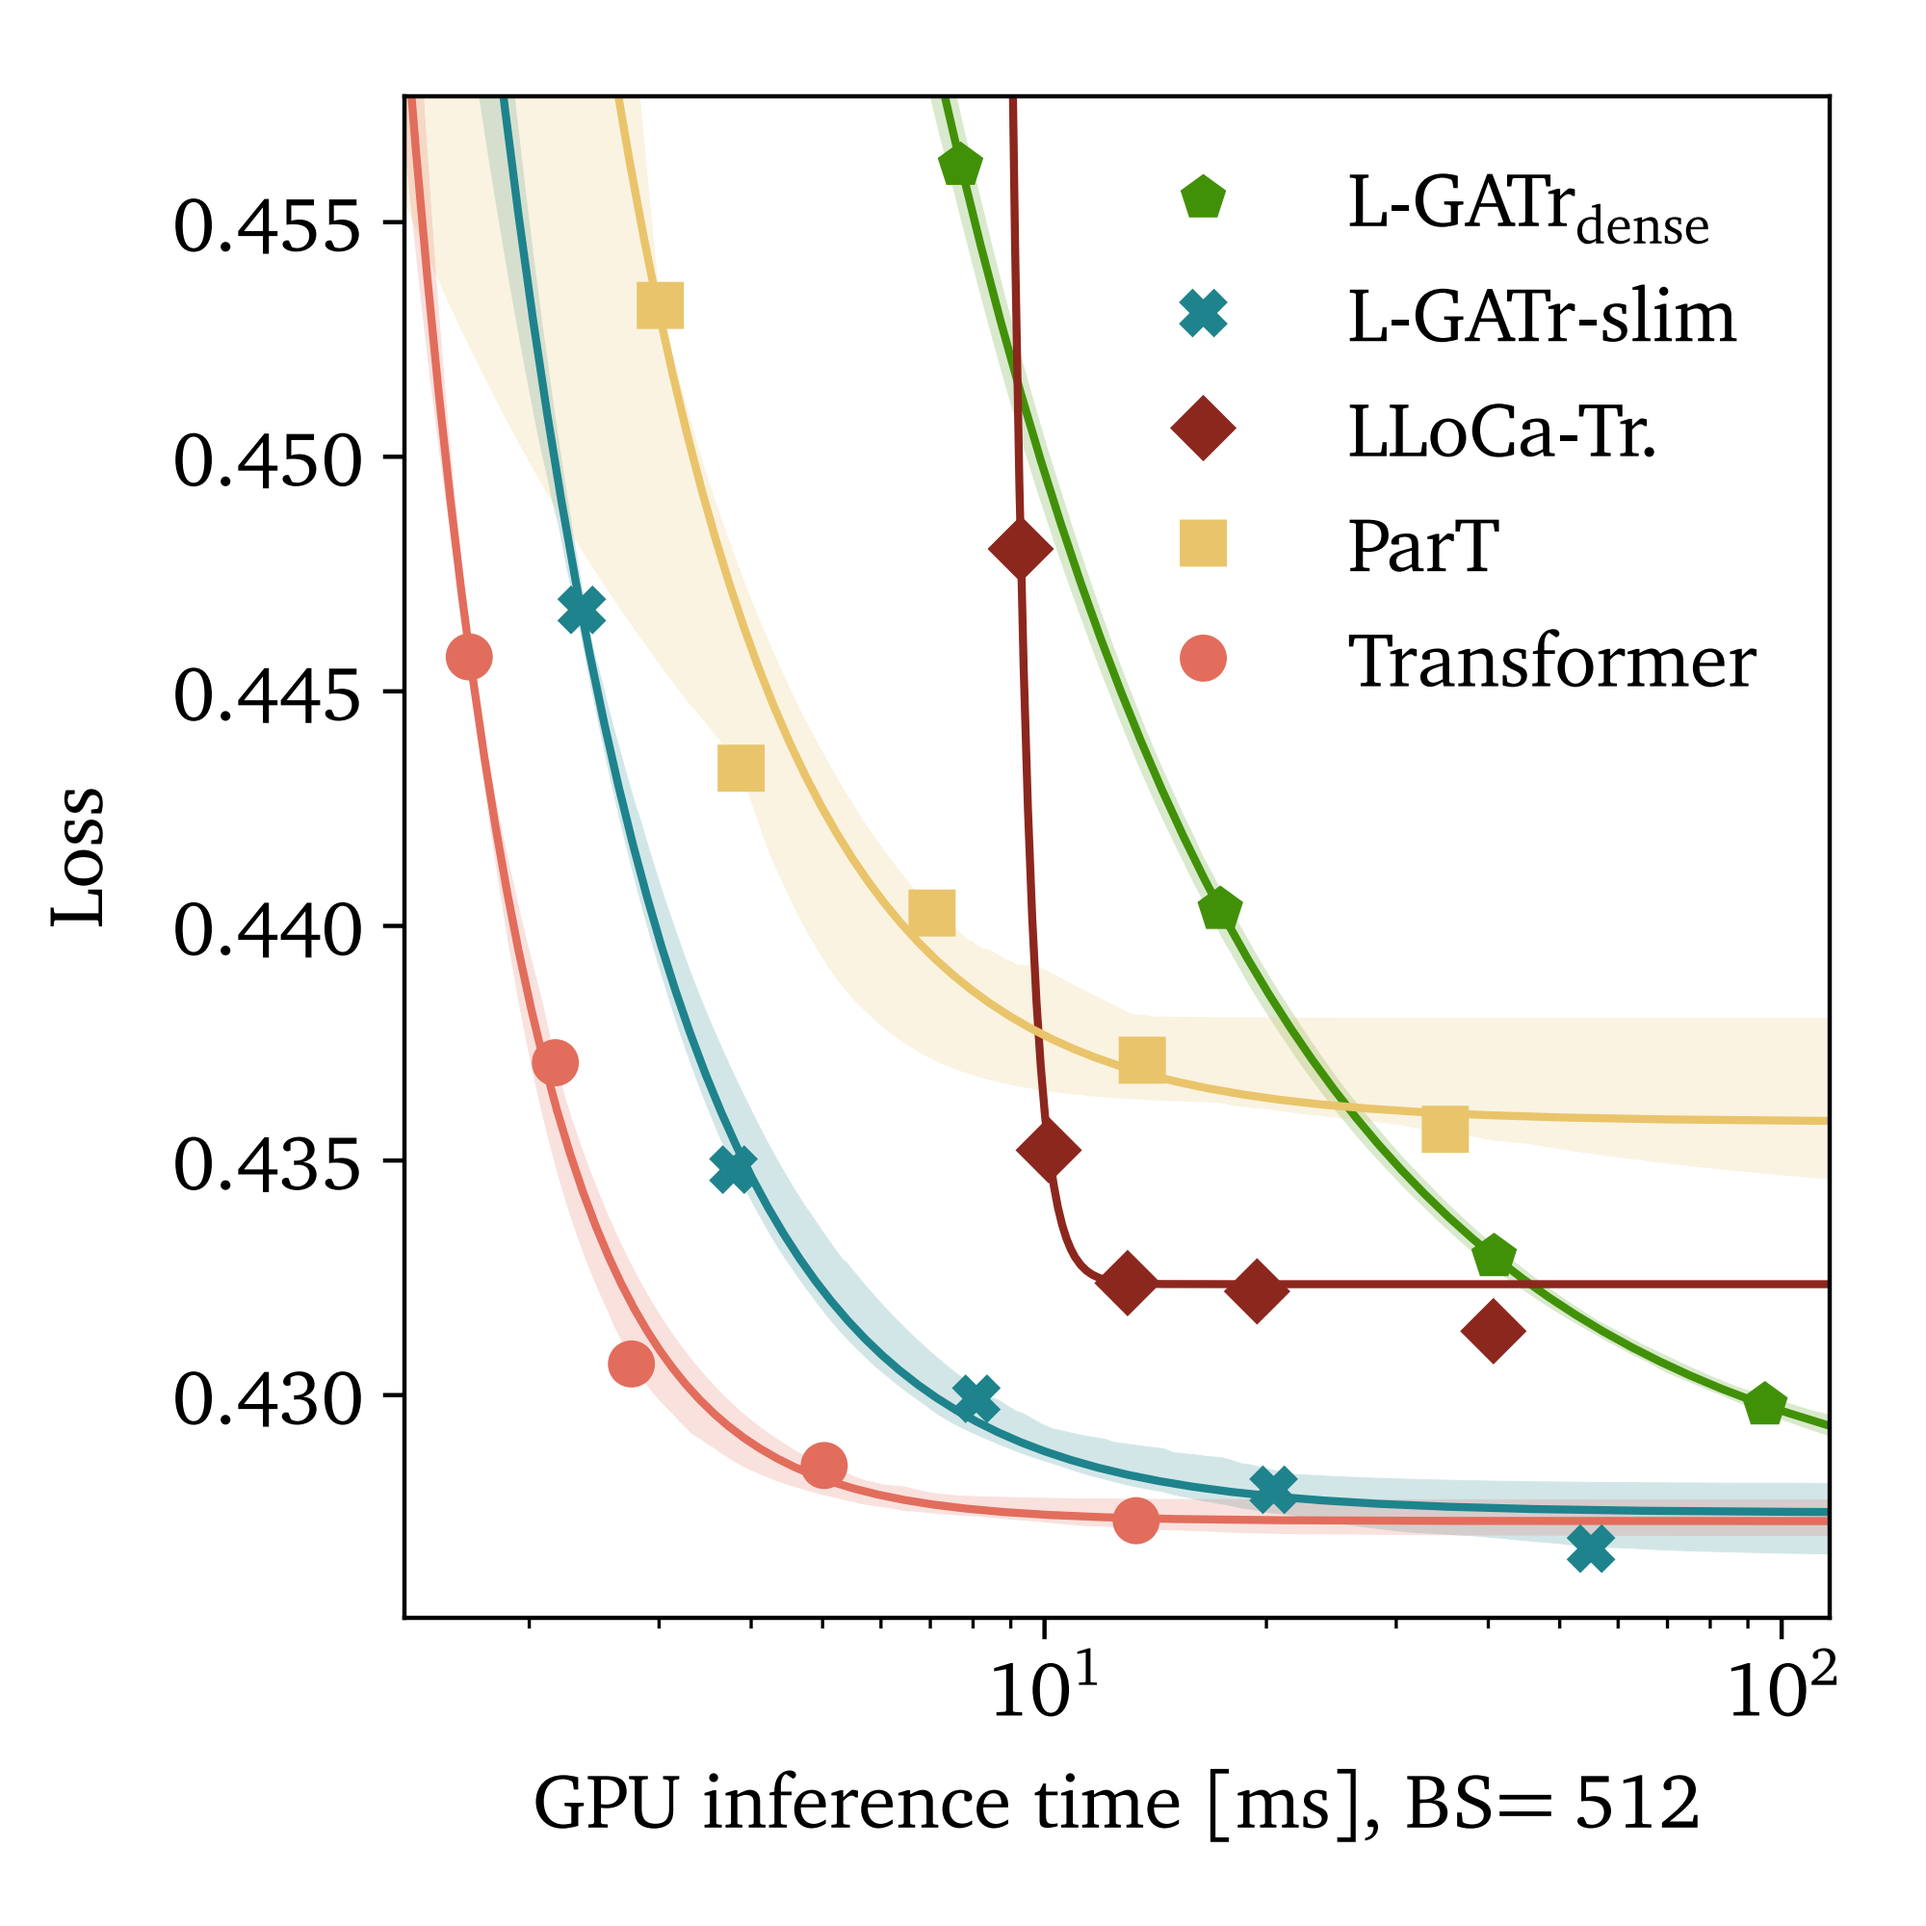
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.455
0.450
0.445
0.440
0.435
0.430
10^1
10^2
GPU inference time [ms], BS= 512

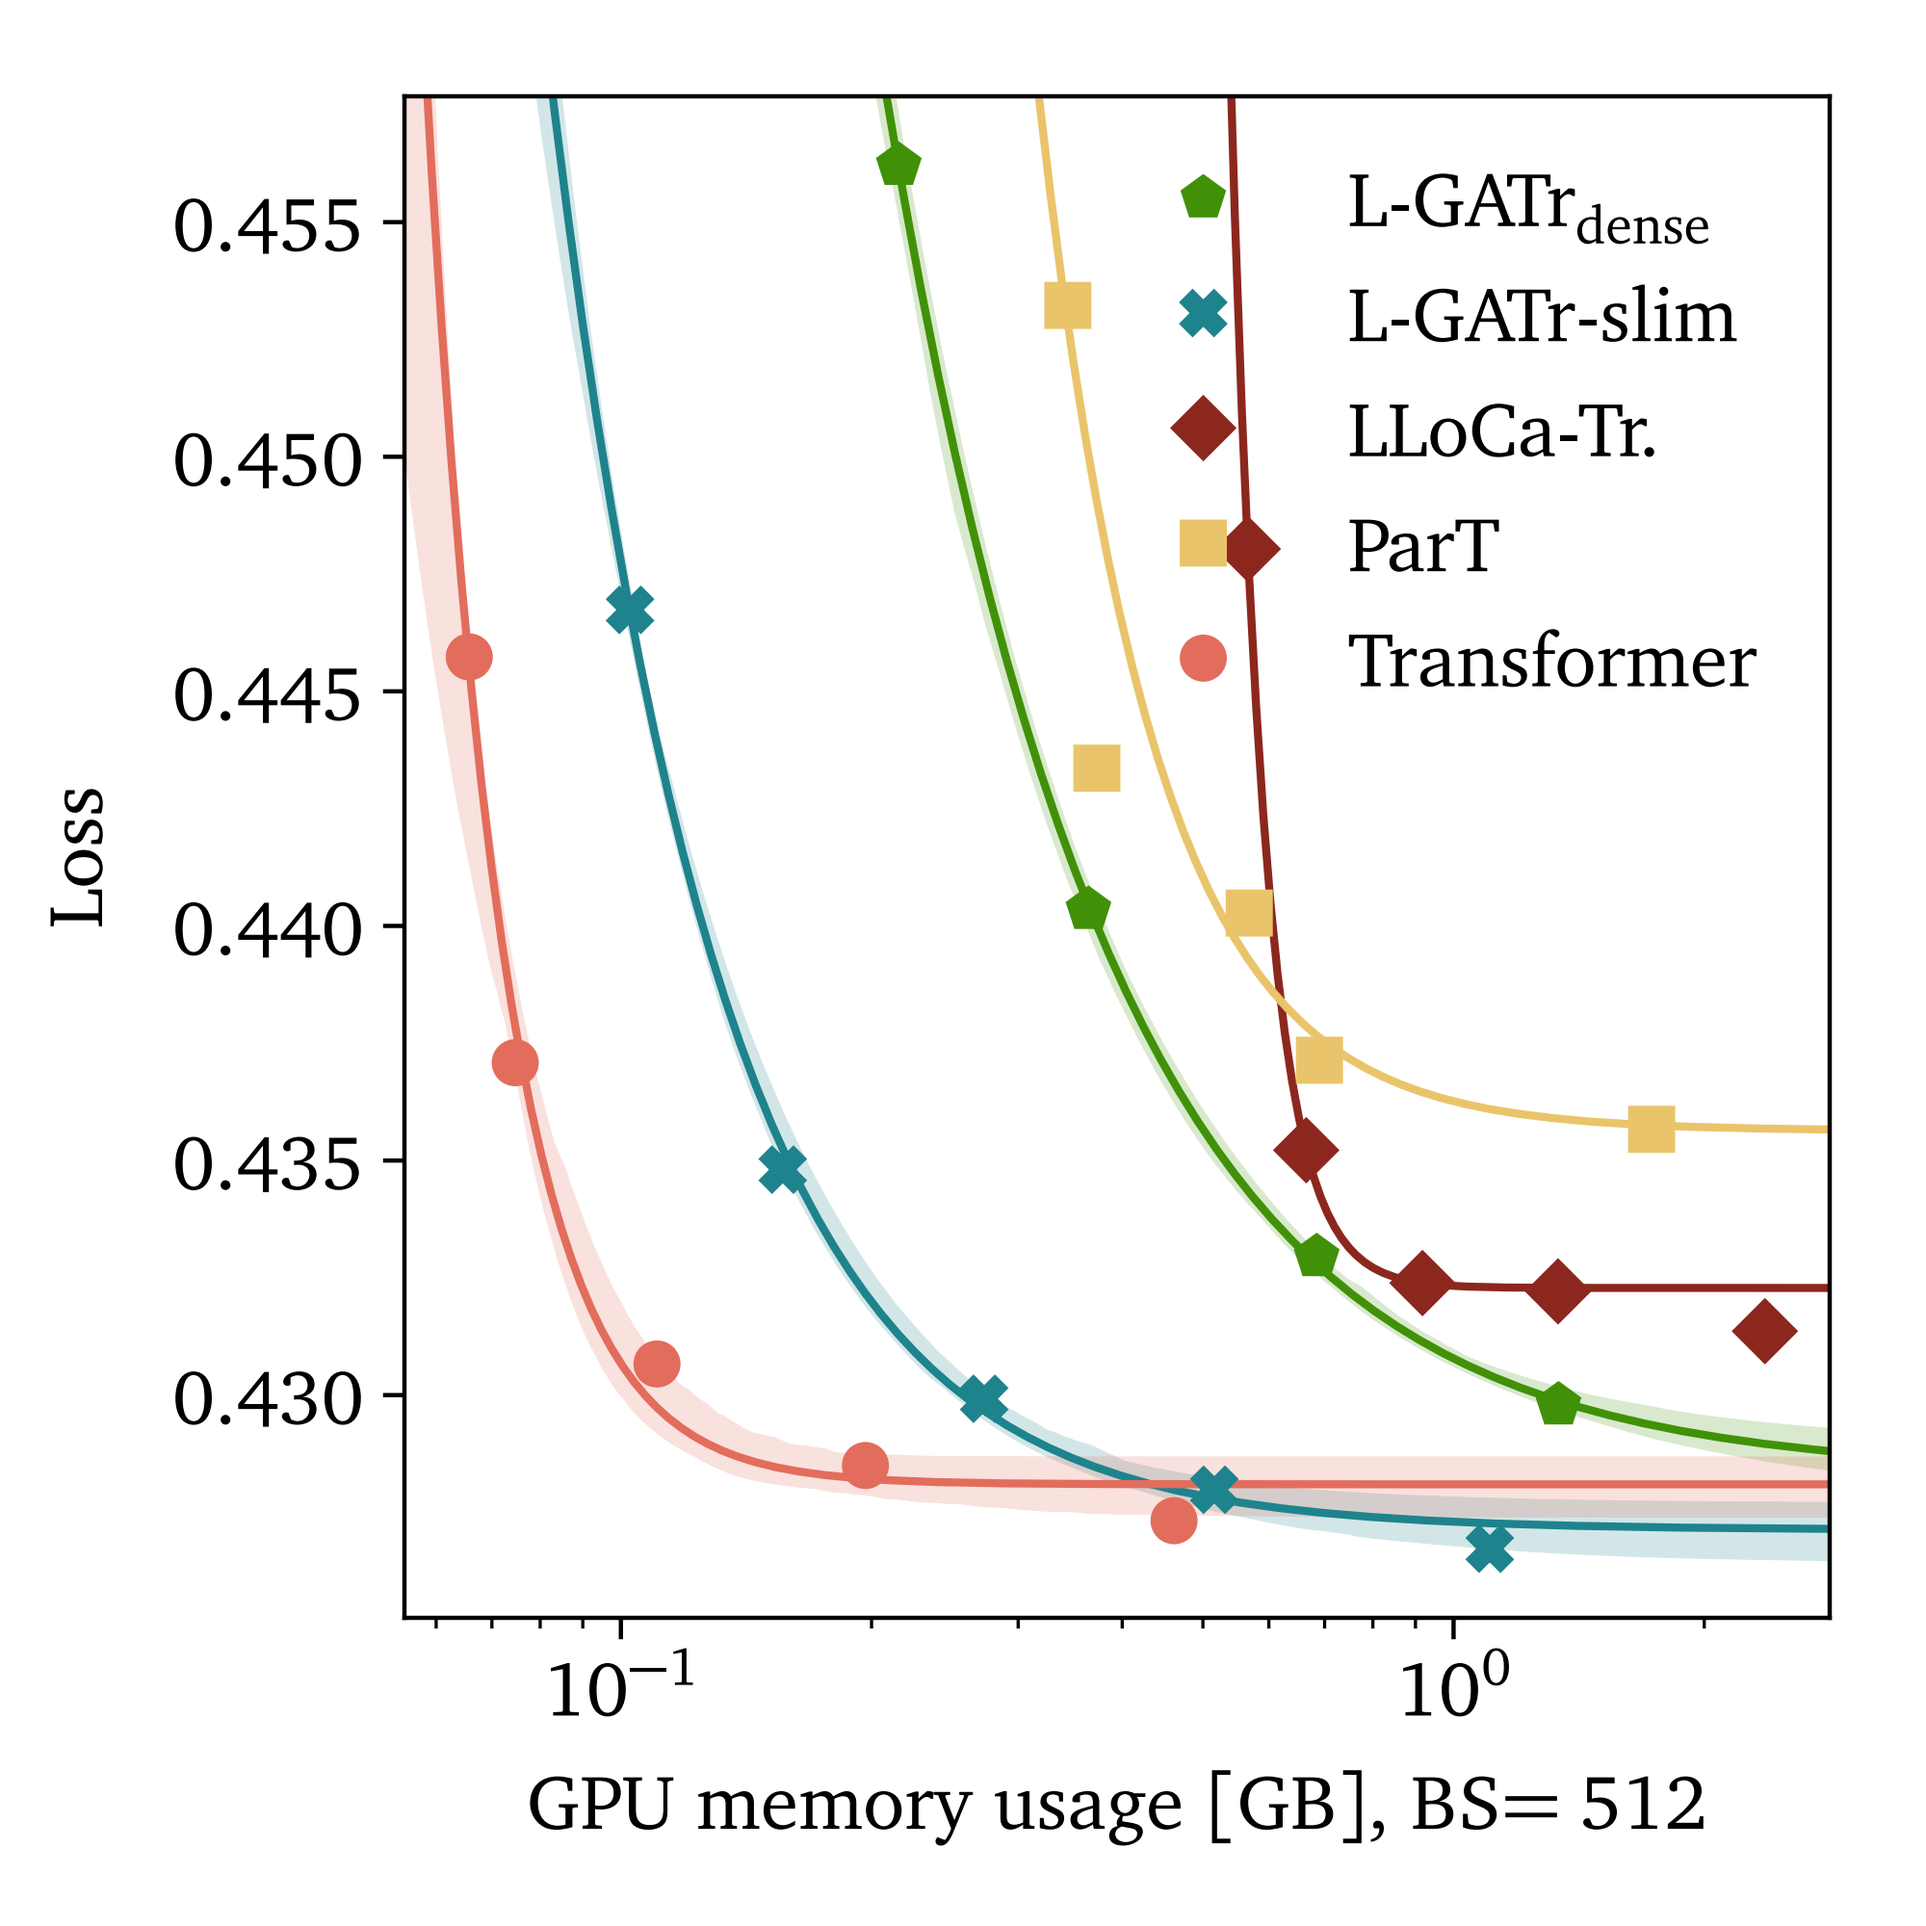
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.455
0.450
0.445
0.440
0.435
0.430
10^{-1}
10^{0}
GPU memory usage [GB], BS= 512

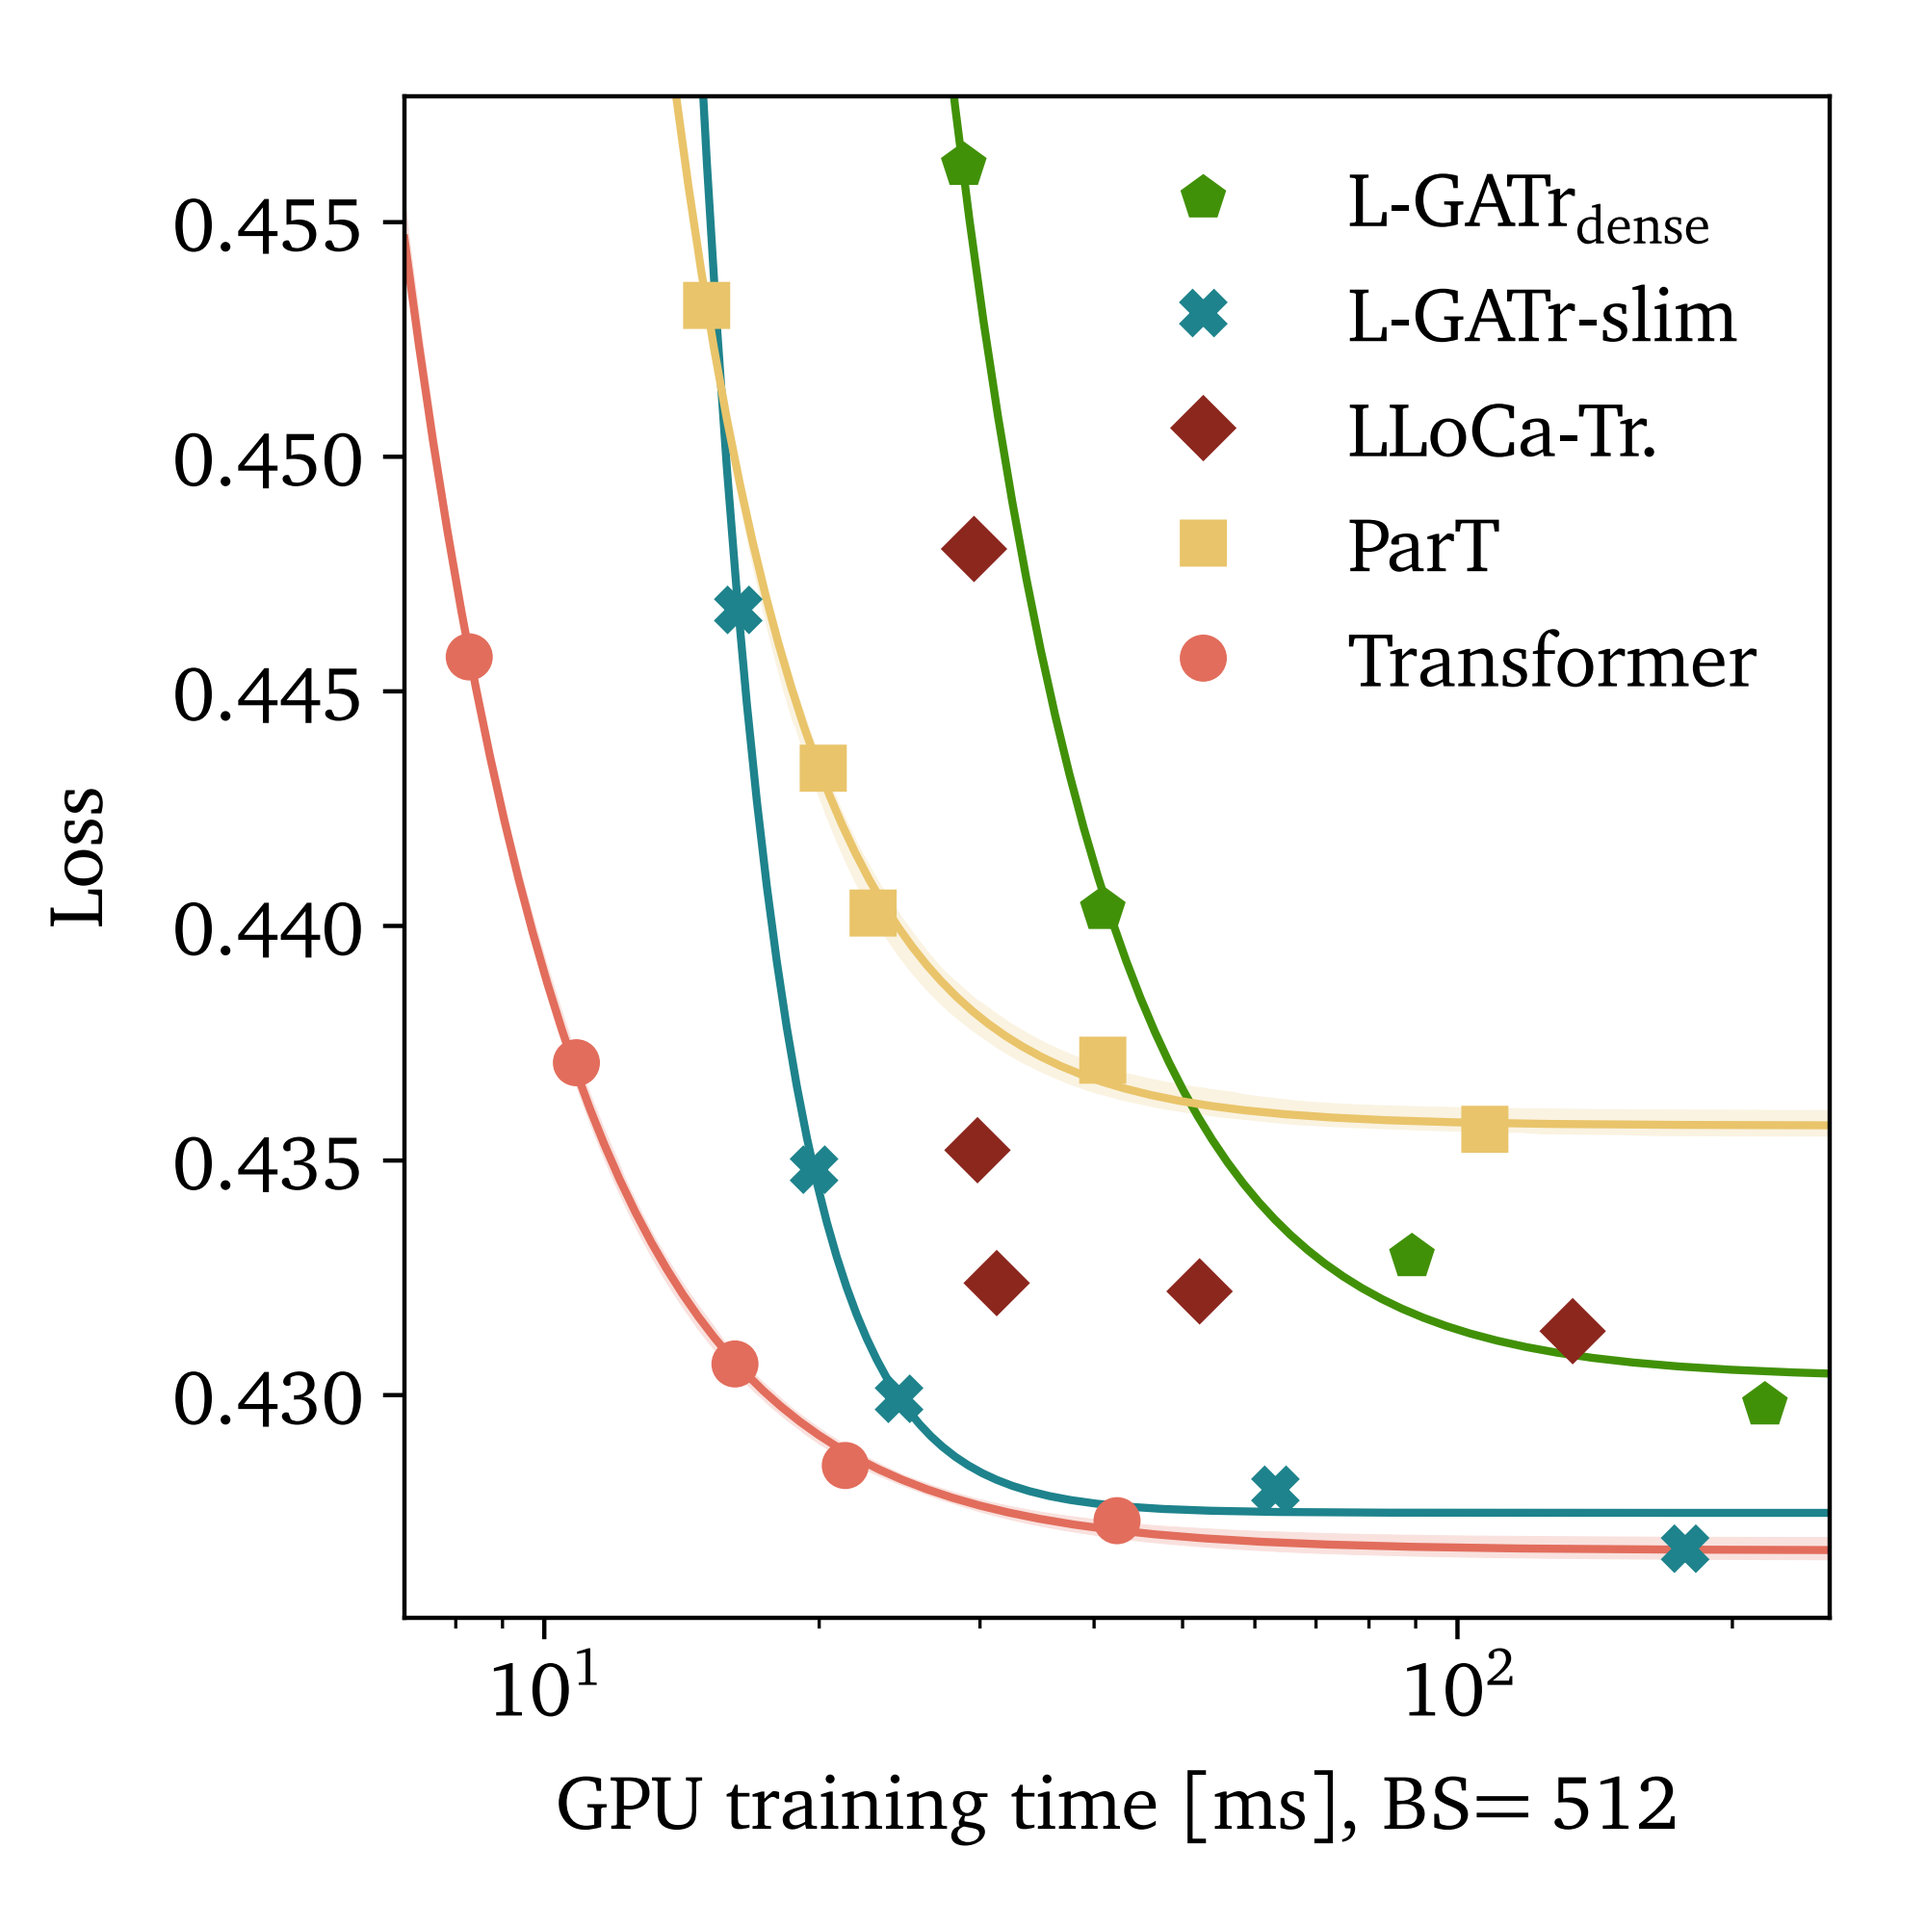
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Loss
0.455
0.450
0.445
0.440
0.435
0.430
10^1
10^2
GPU training time [ms], BS= 512

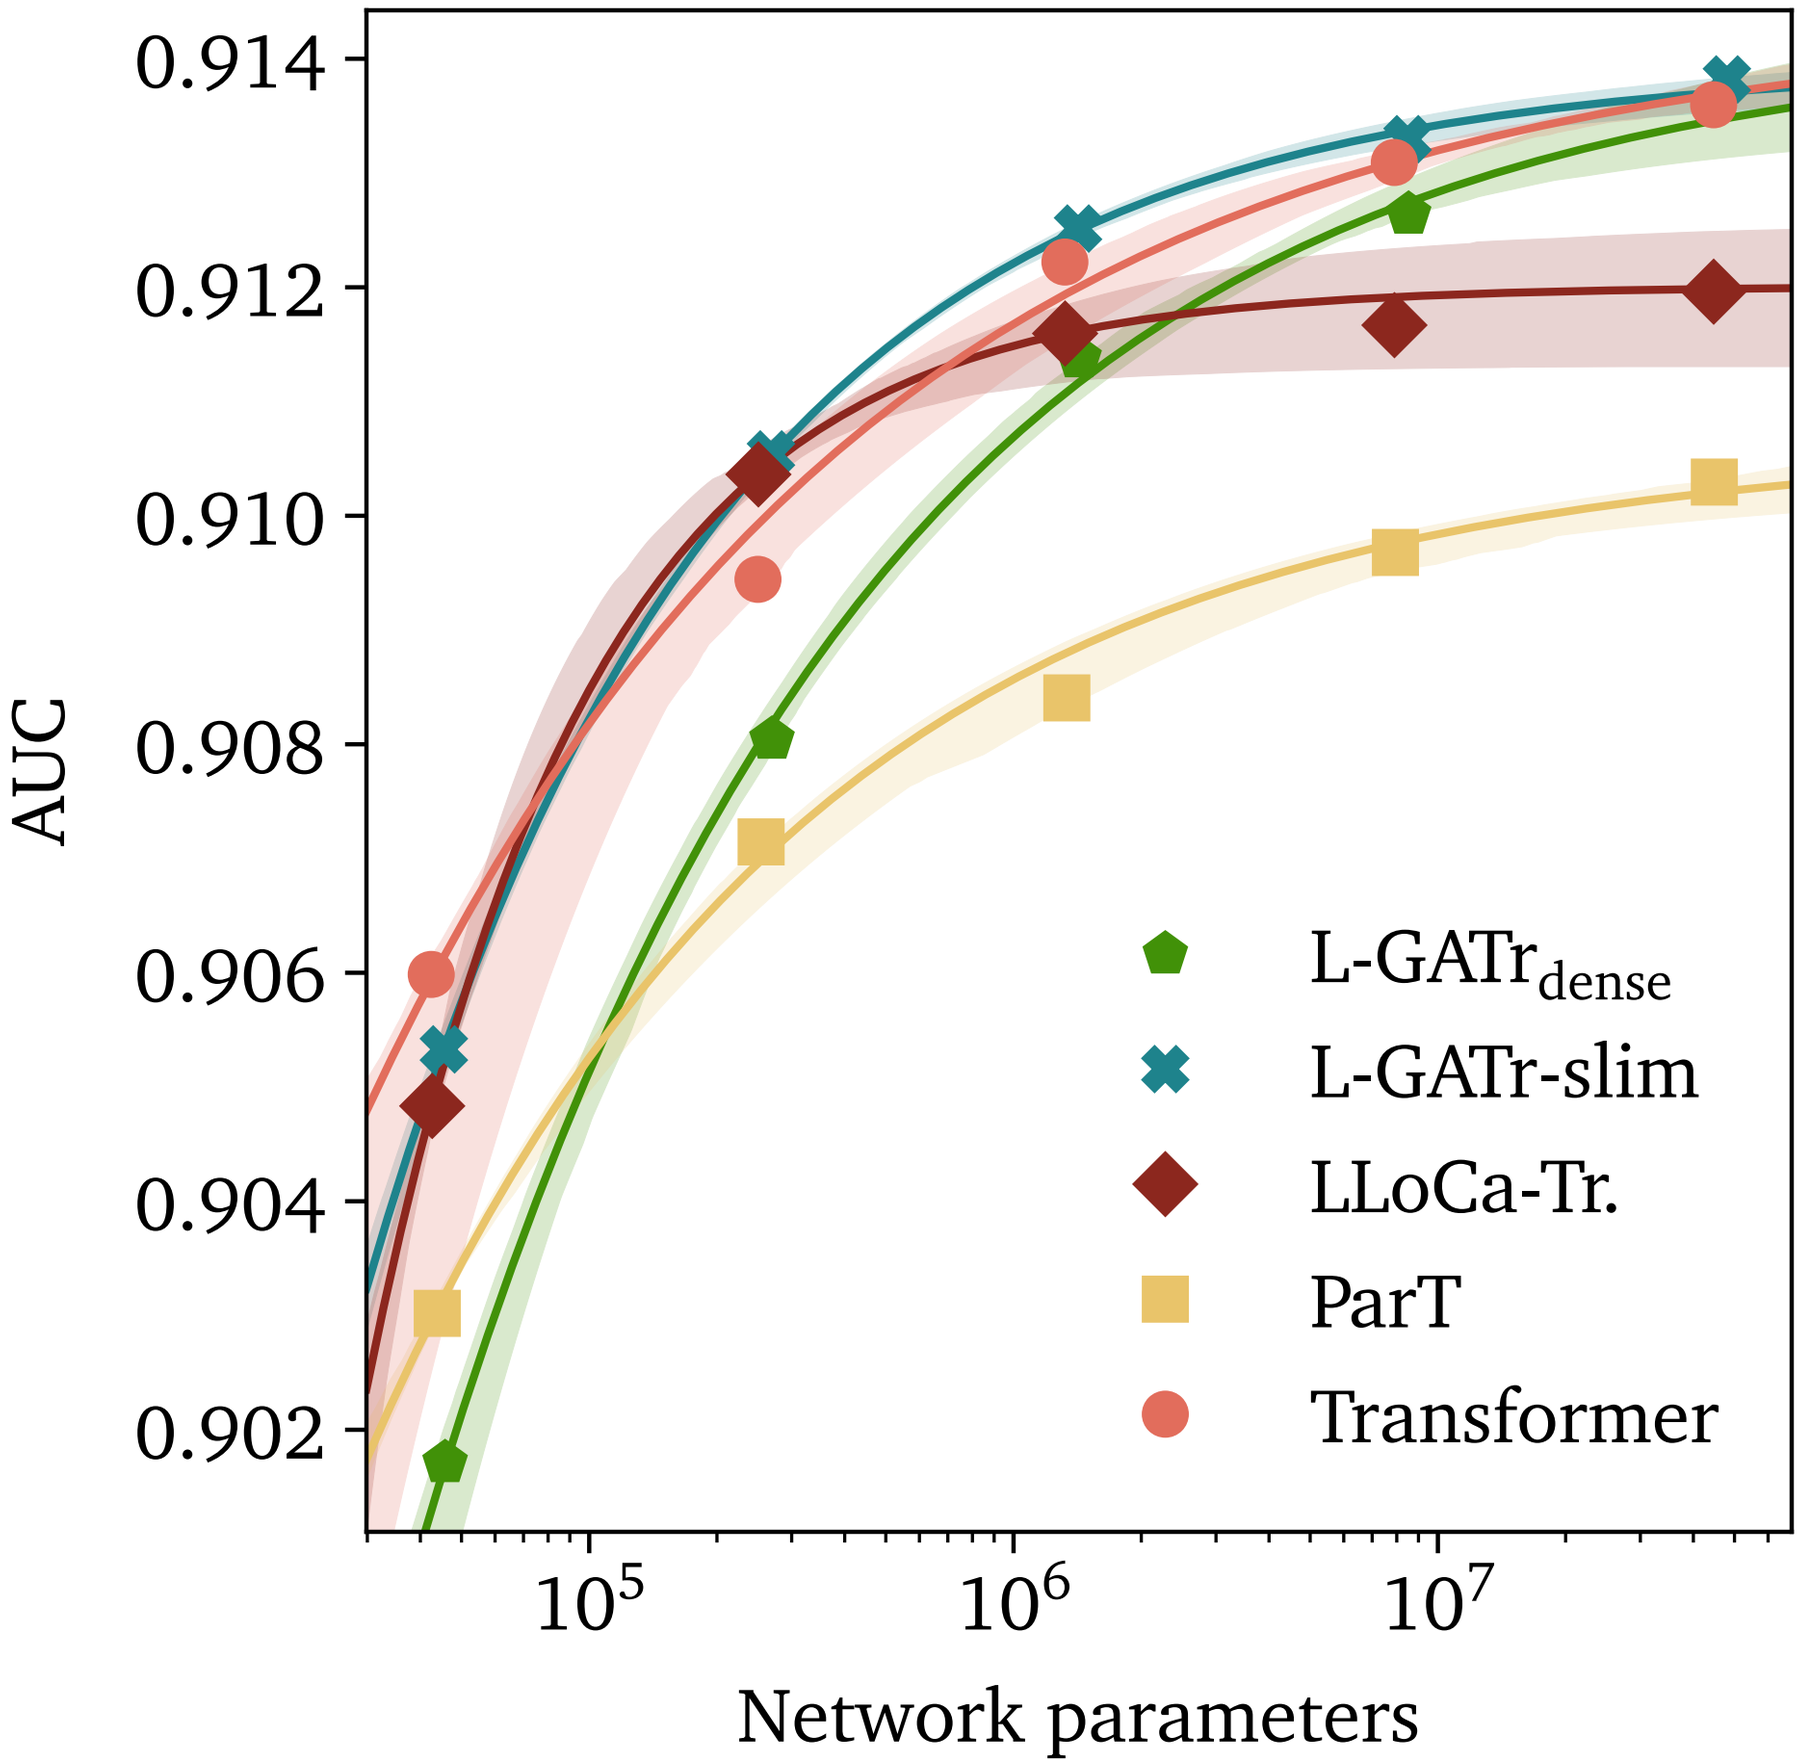
0.914
0.912
0.910
0.908
0.906
0.904
0.902
AUC
10^5
10^6
10^7
Network parameters
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

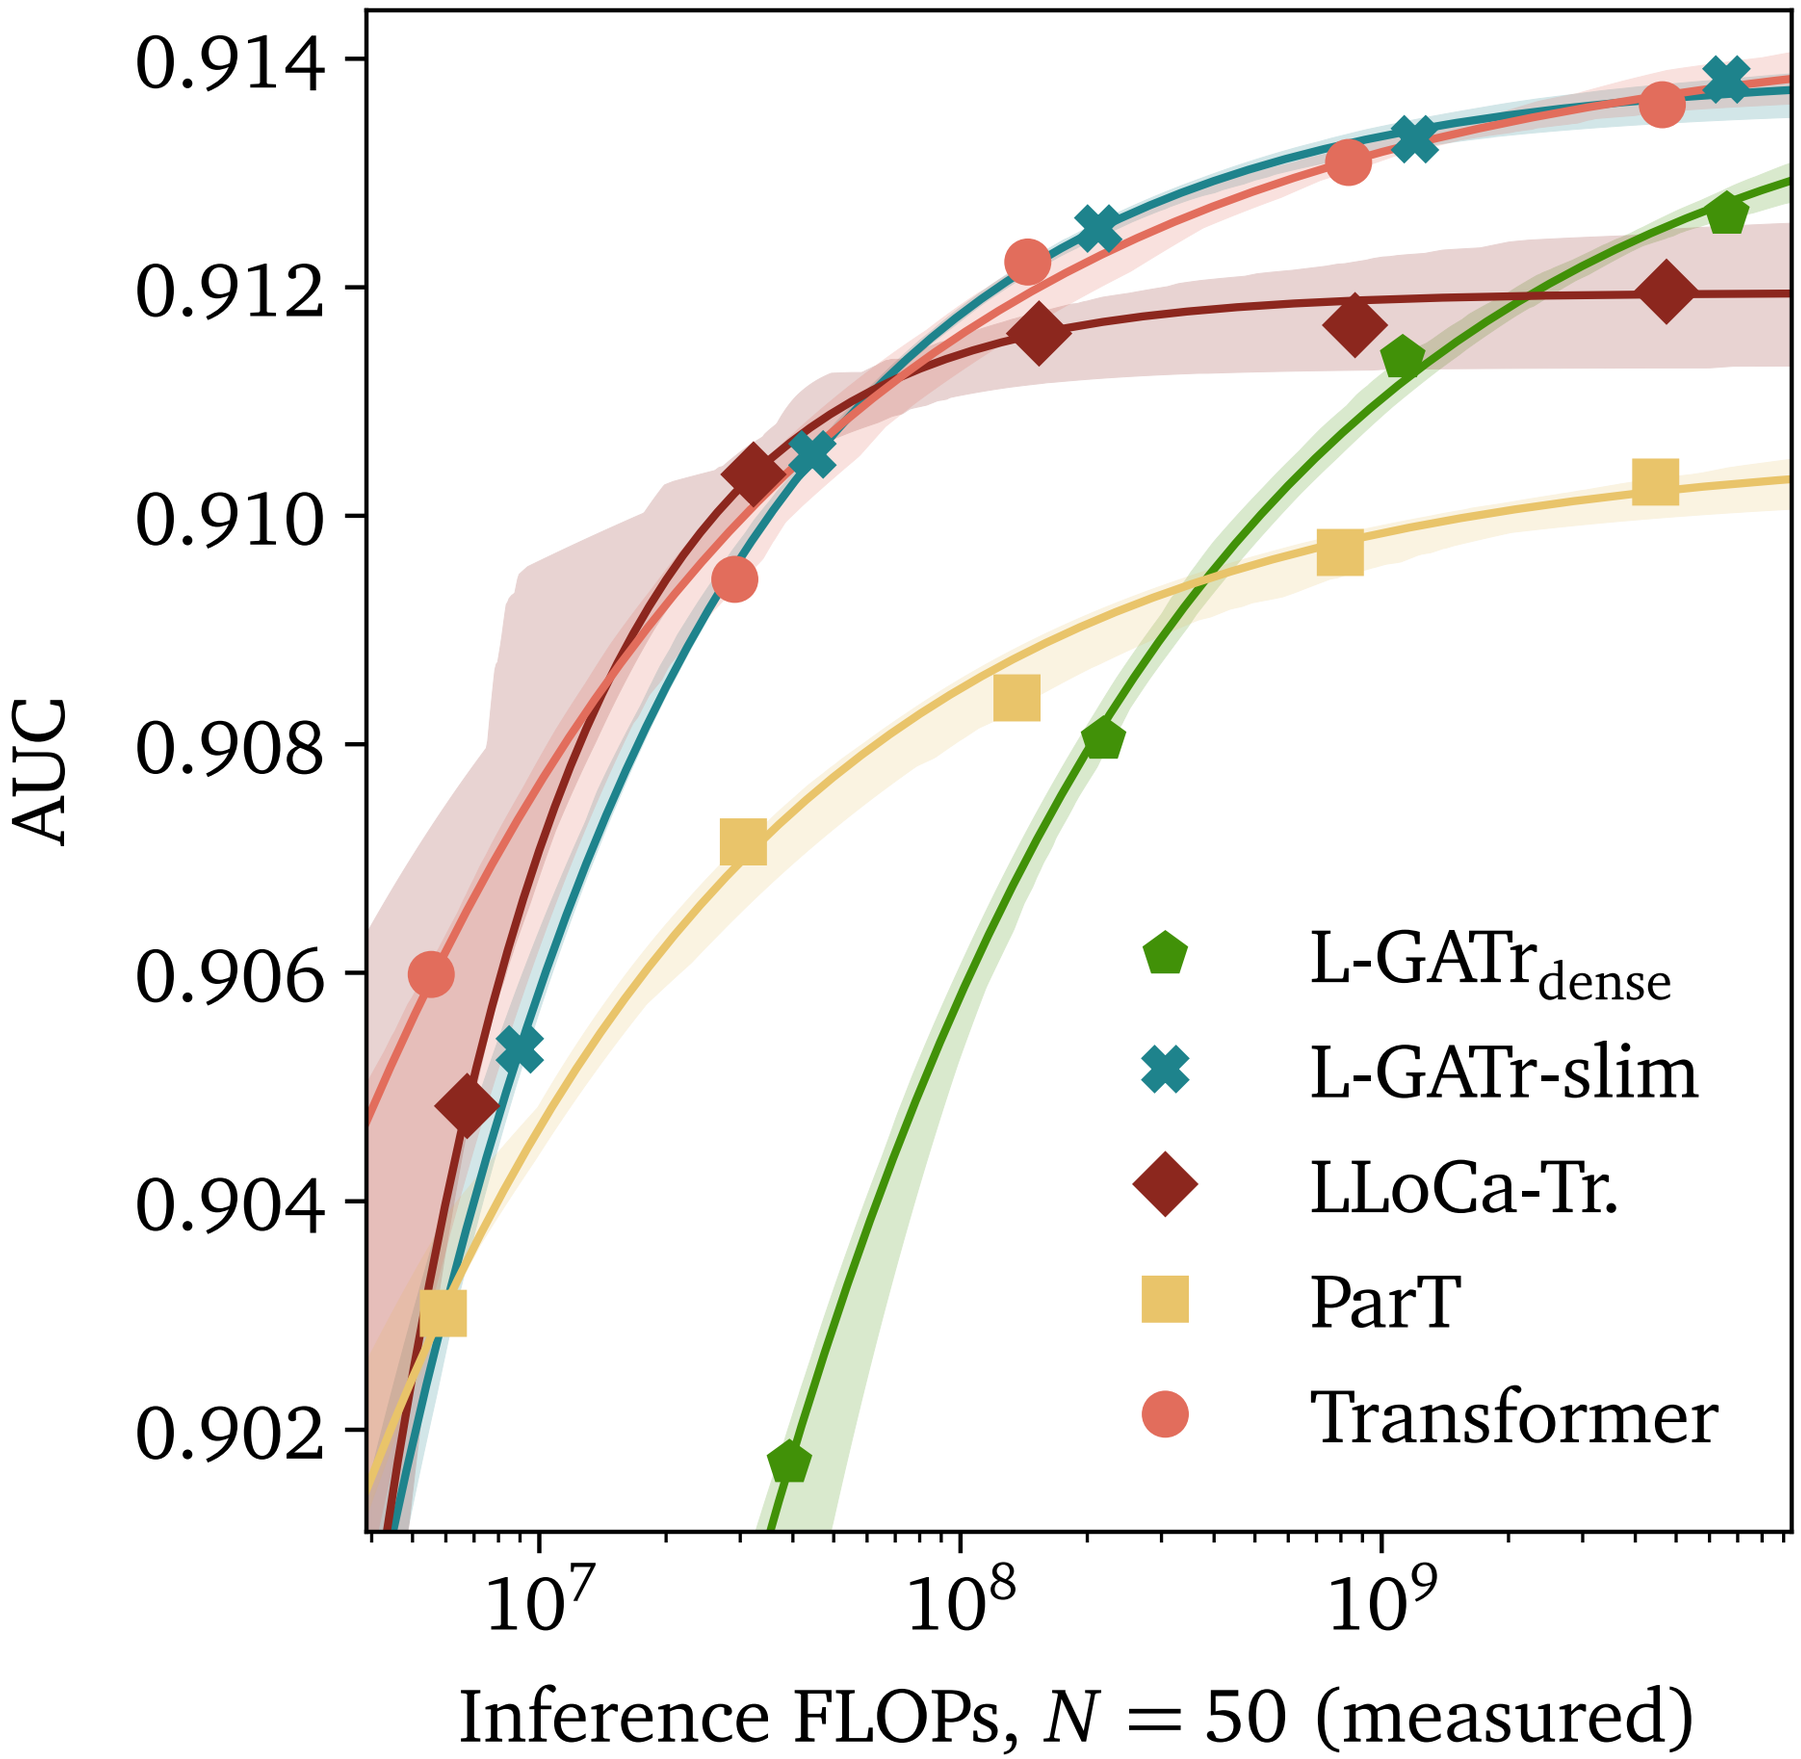
0.914
0.912
0.910
0.908
0.906
0.904
0.902
AUC
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (measured)
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

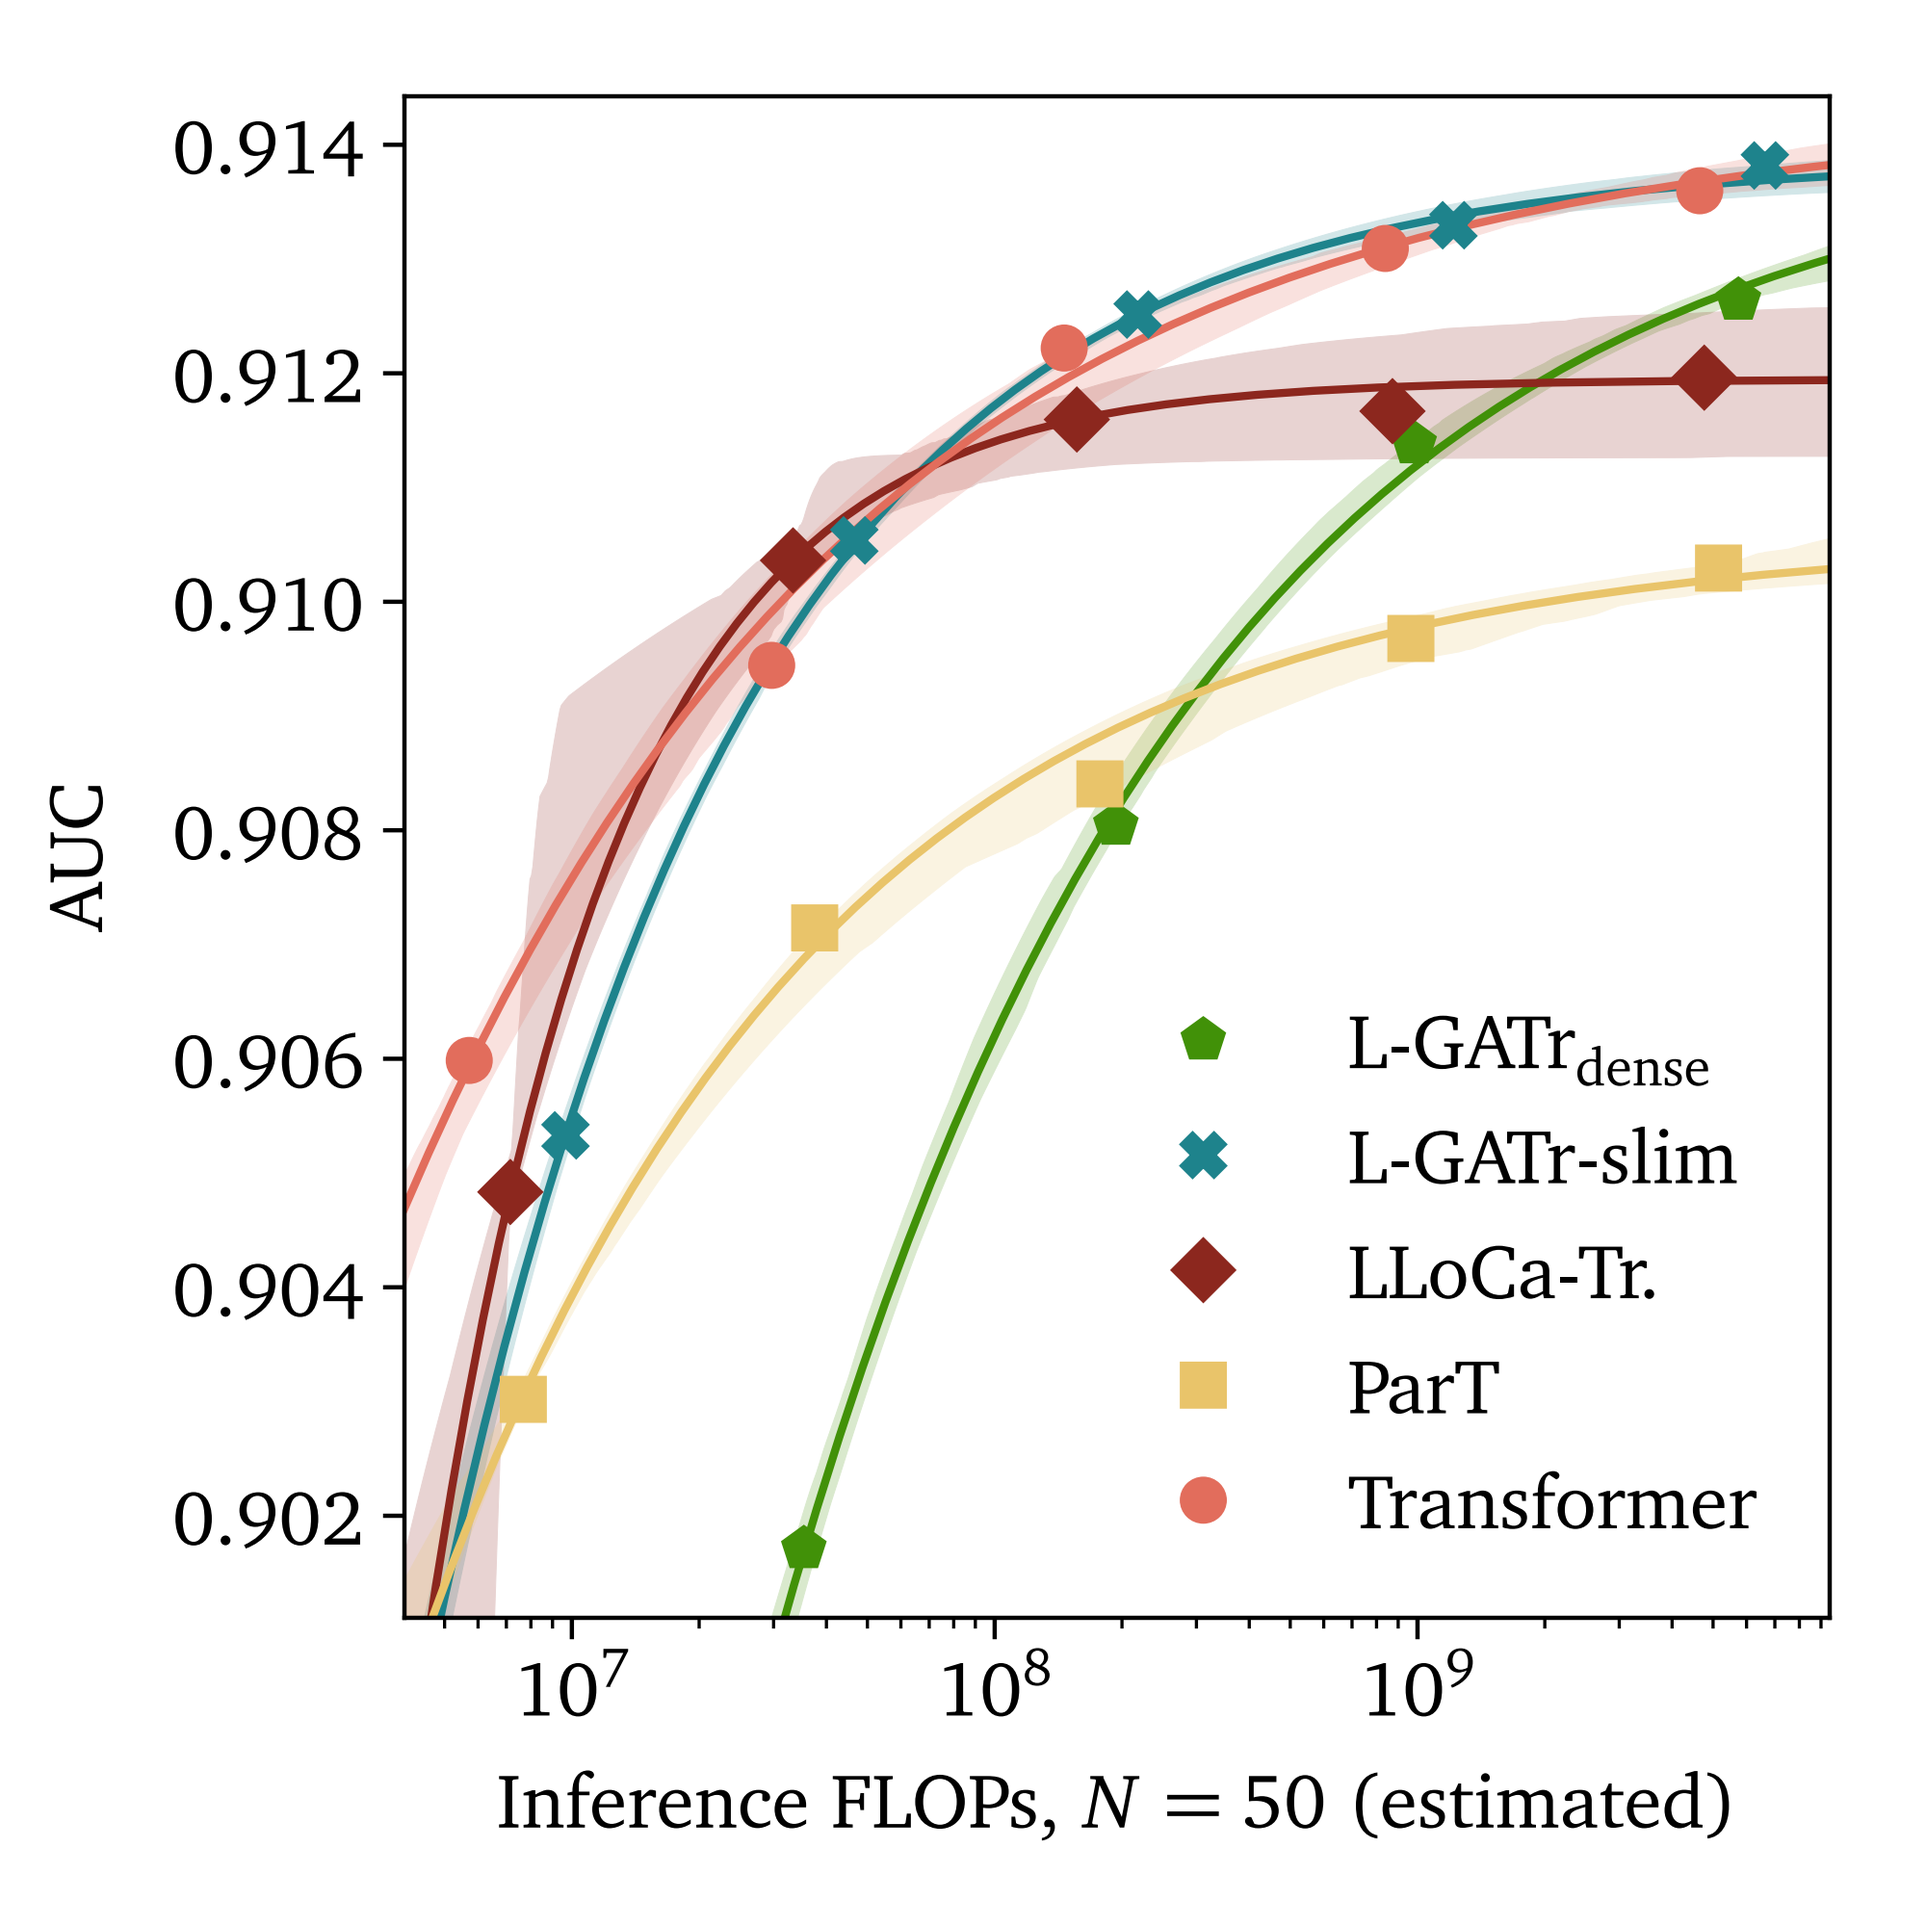
0.914
0.912
0.910
0.908
0.906
0.904
0.902
AUC
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (estimated)
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

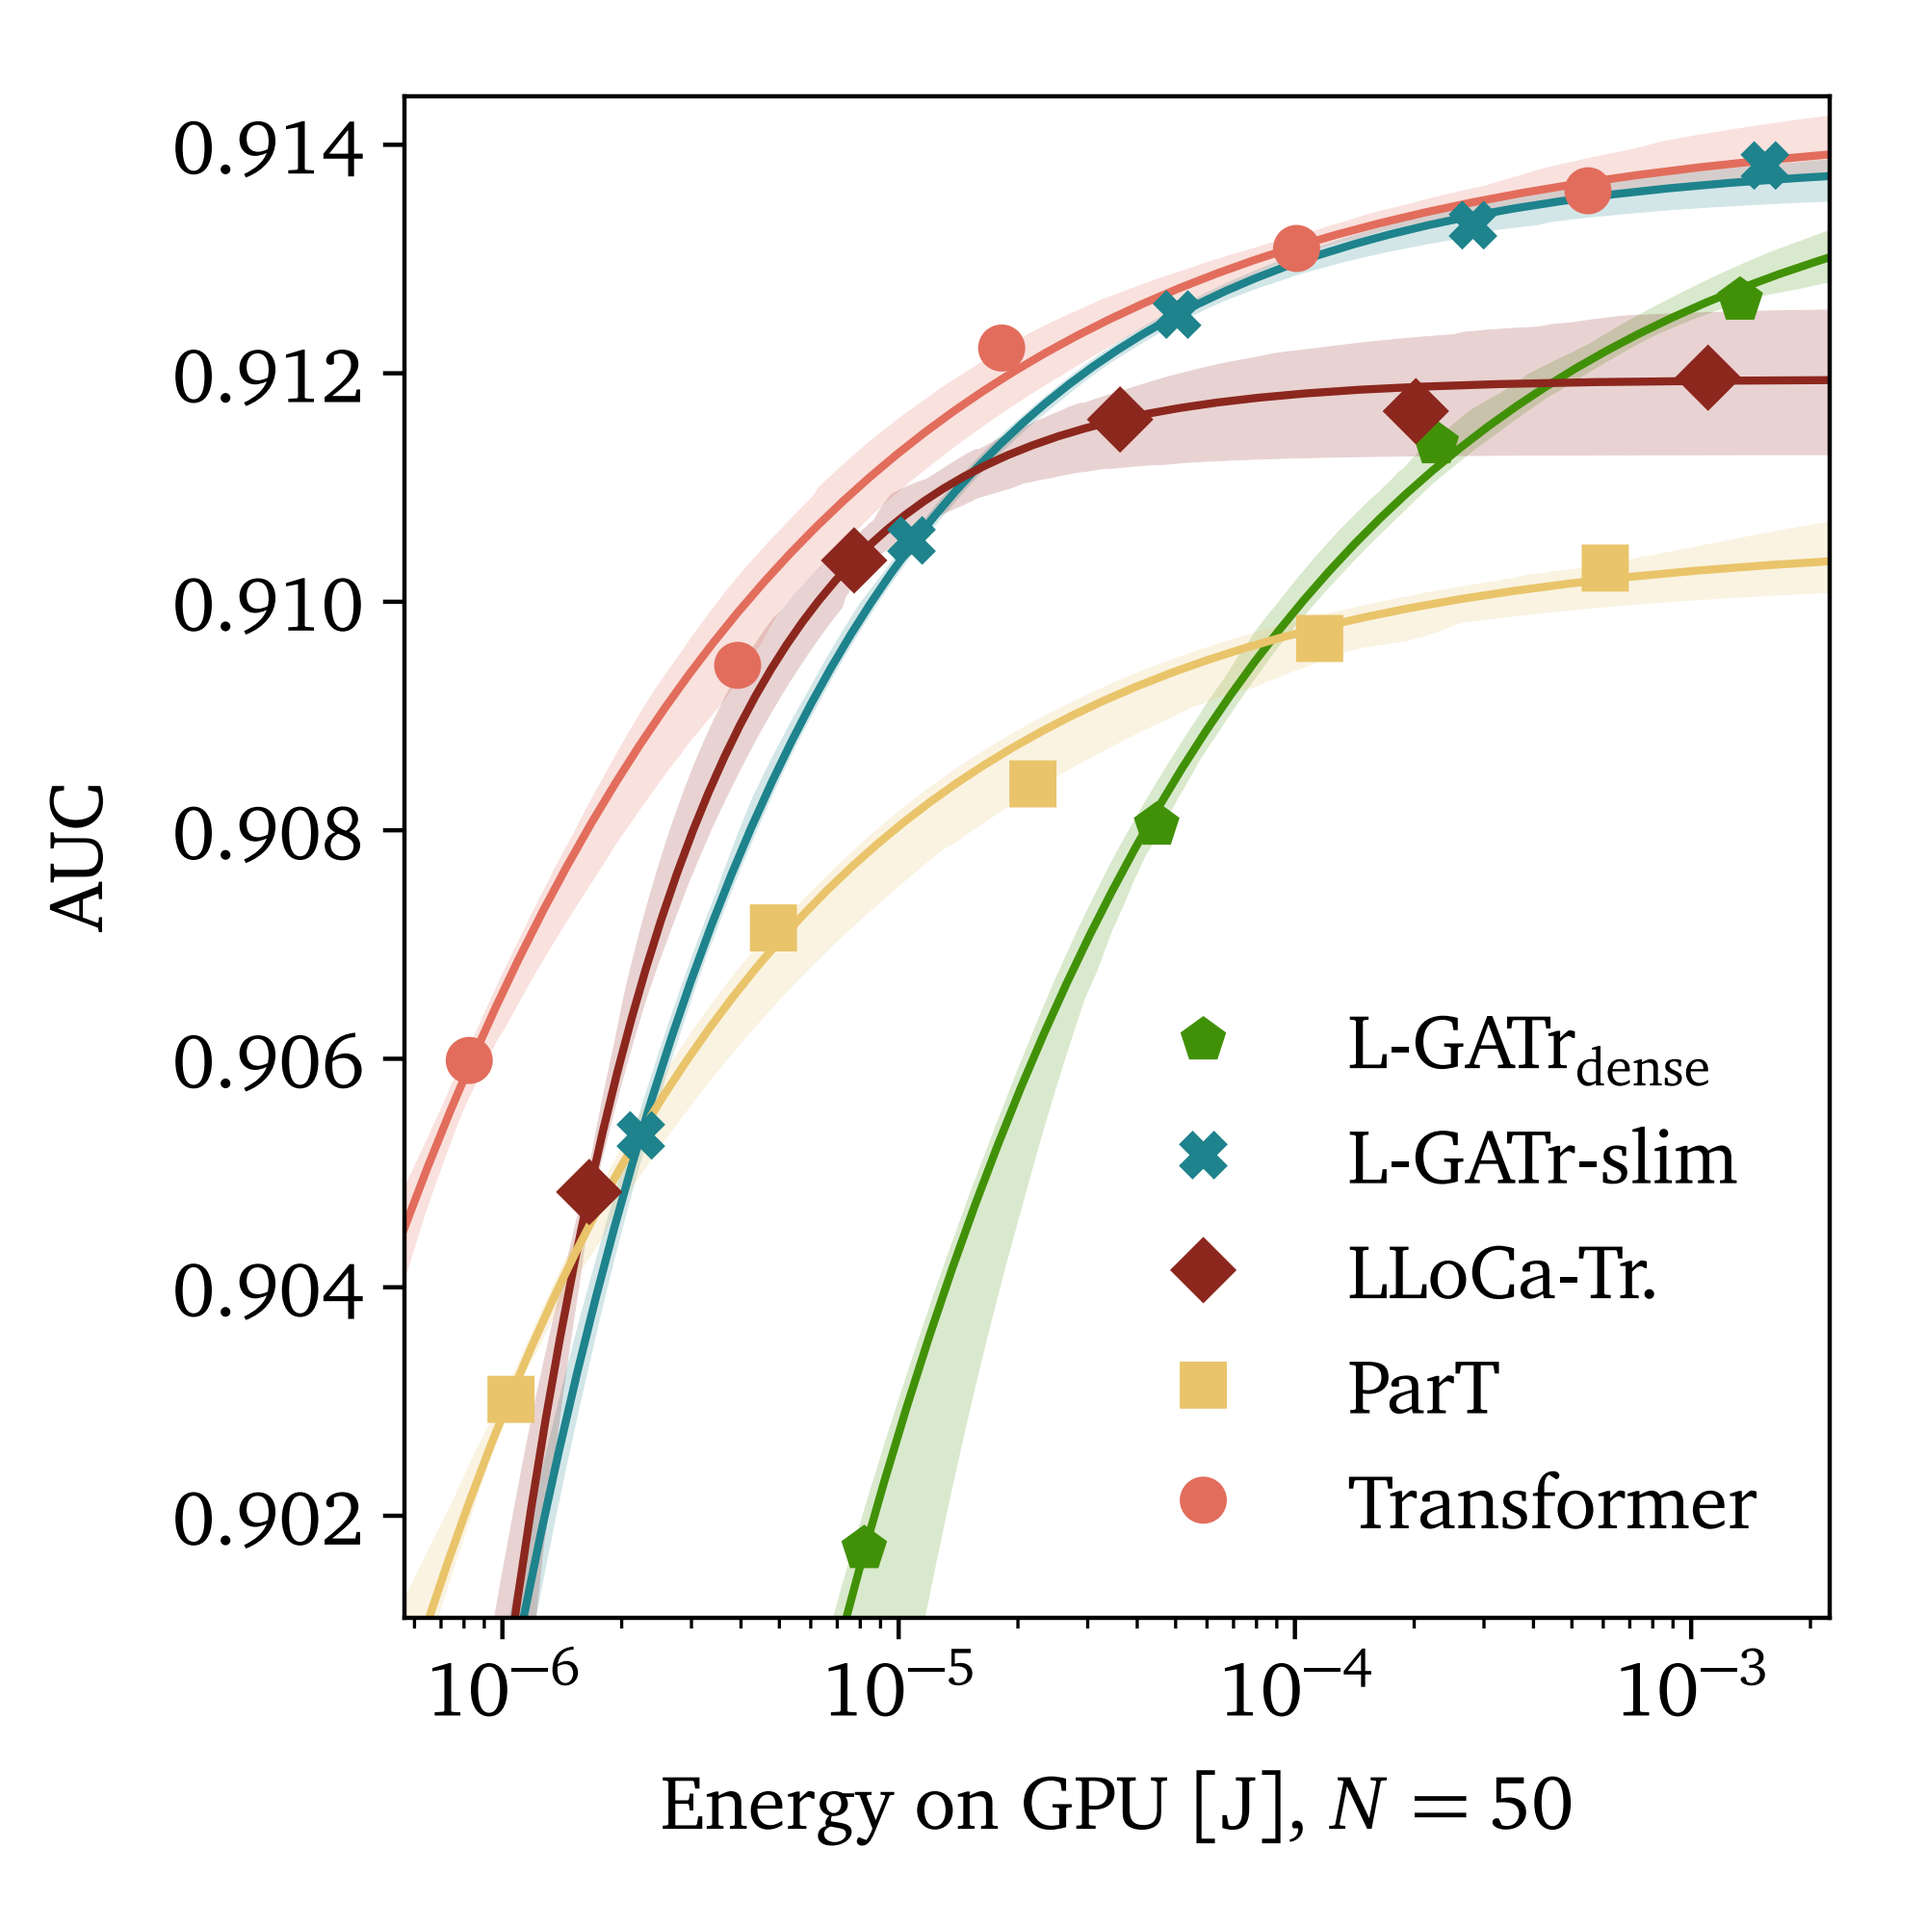
0.914
0.912
0.910
0.908
0.906
0.904
0.902
AUC
10^{-6}
10^{-5}
10^{-4}
10^{-3}
Energy on GPU [J], $N = 50$
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

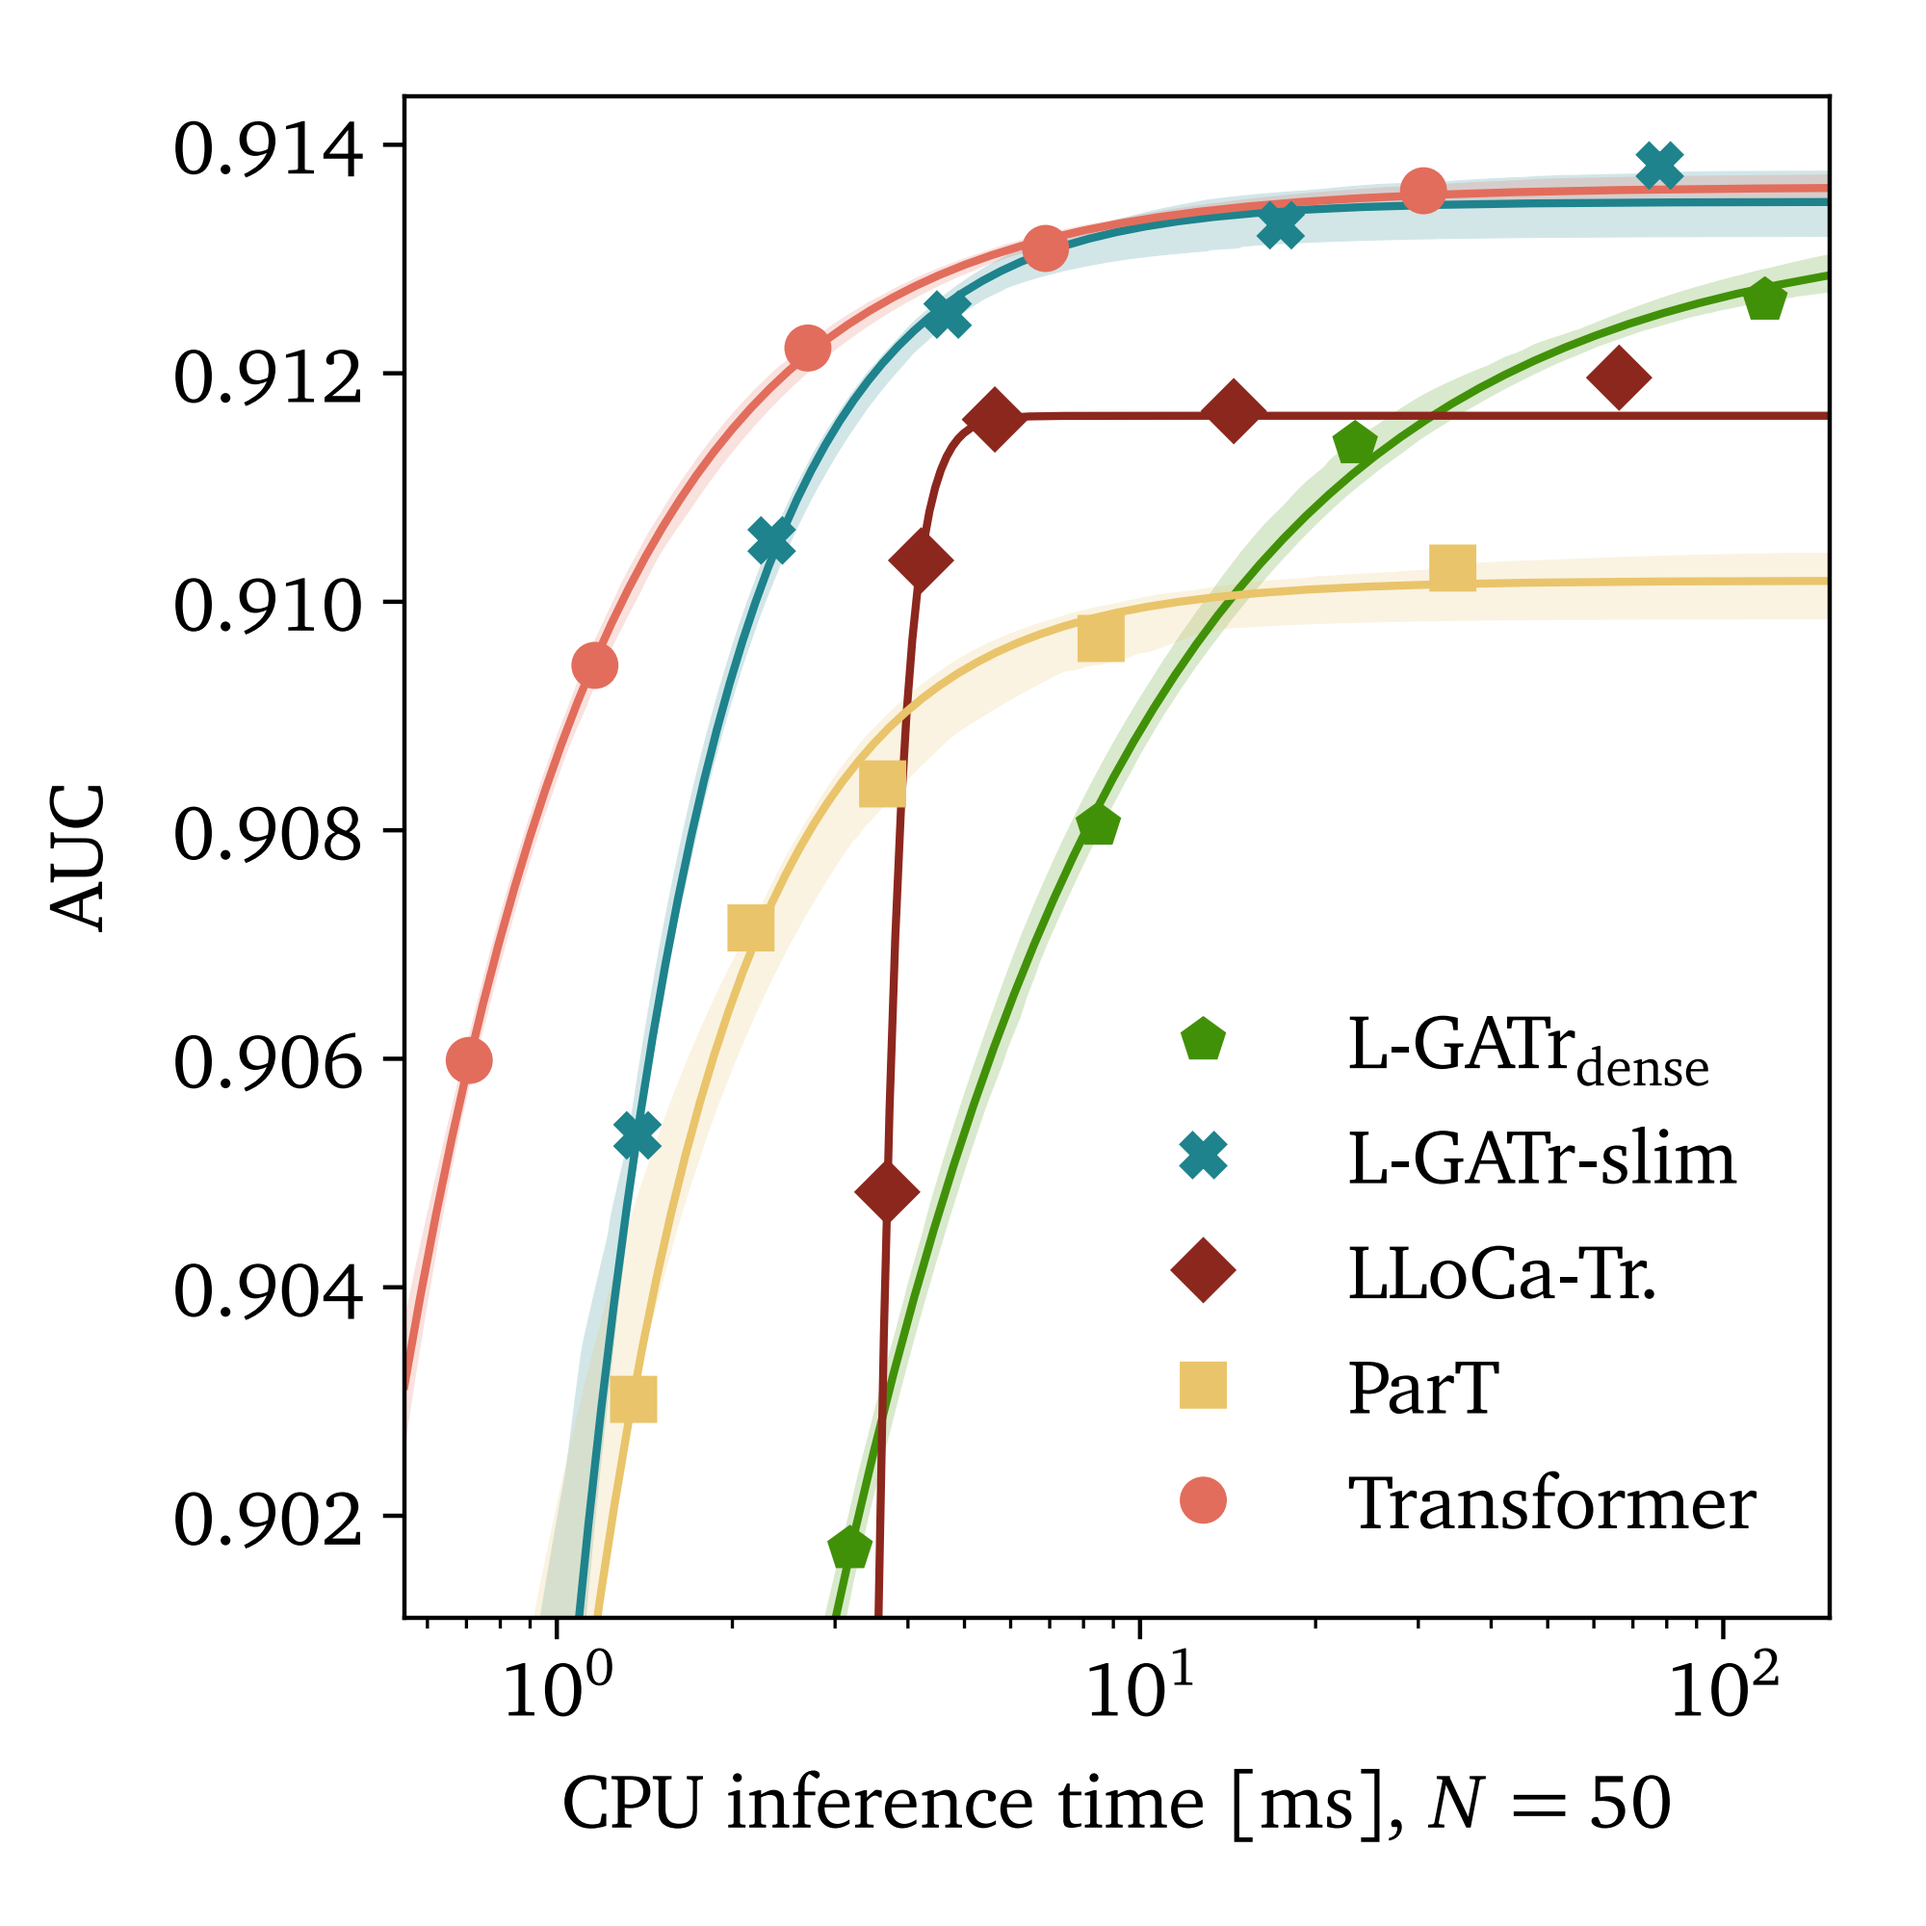
0.914
0.912
0.910
0.908
0.906
0.904
0.902
AUC
10^0
10^1
10^2
CPU inference time [ms], $N = 50$
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

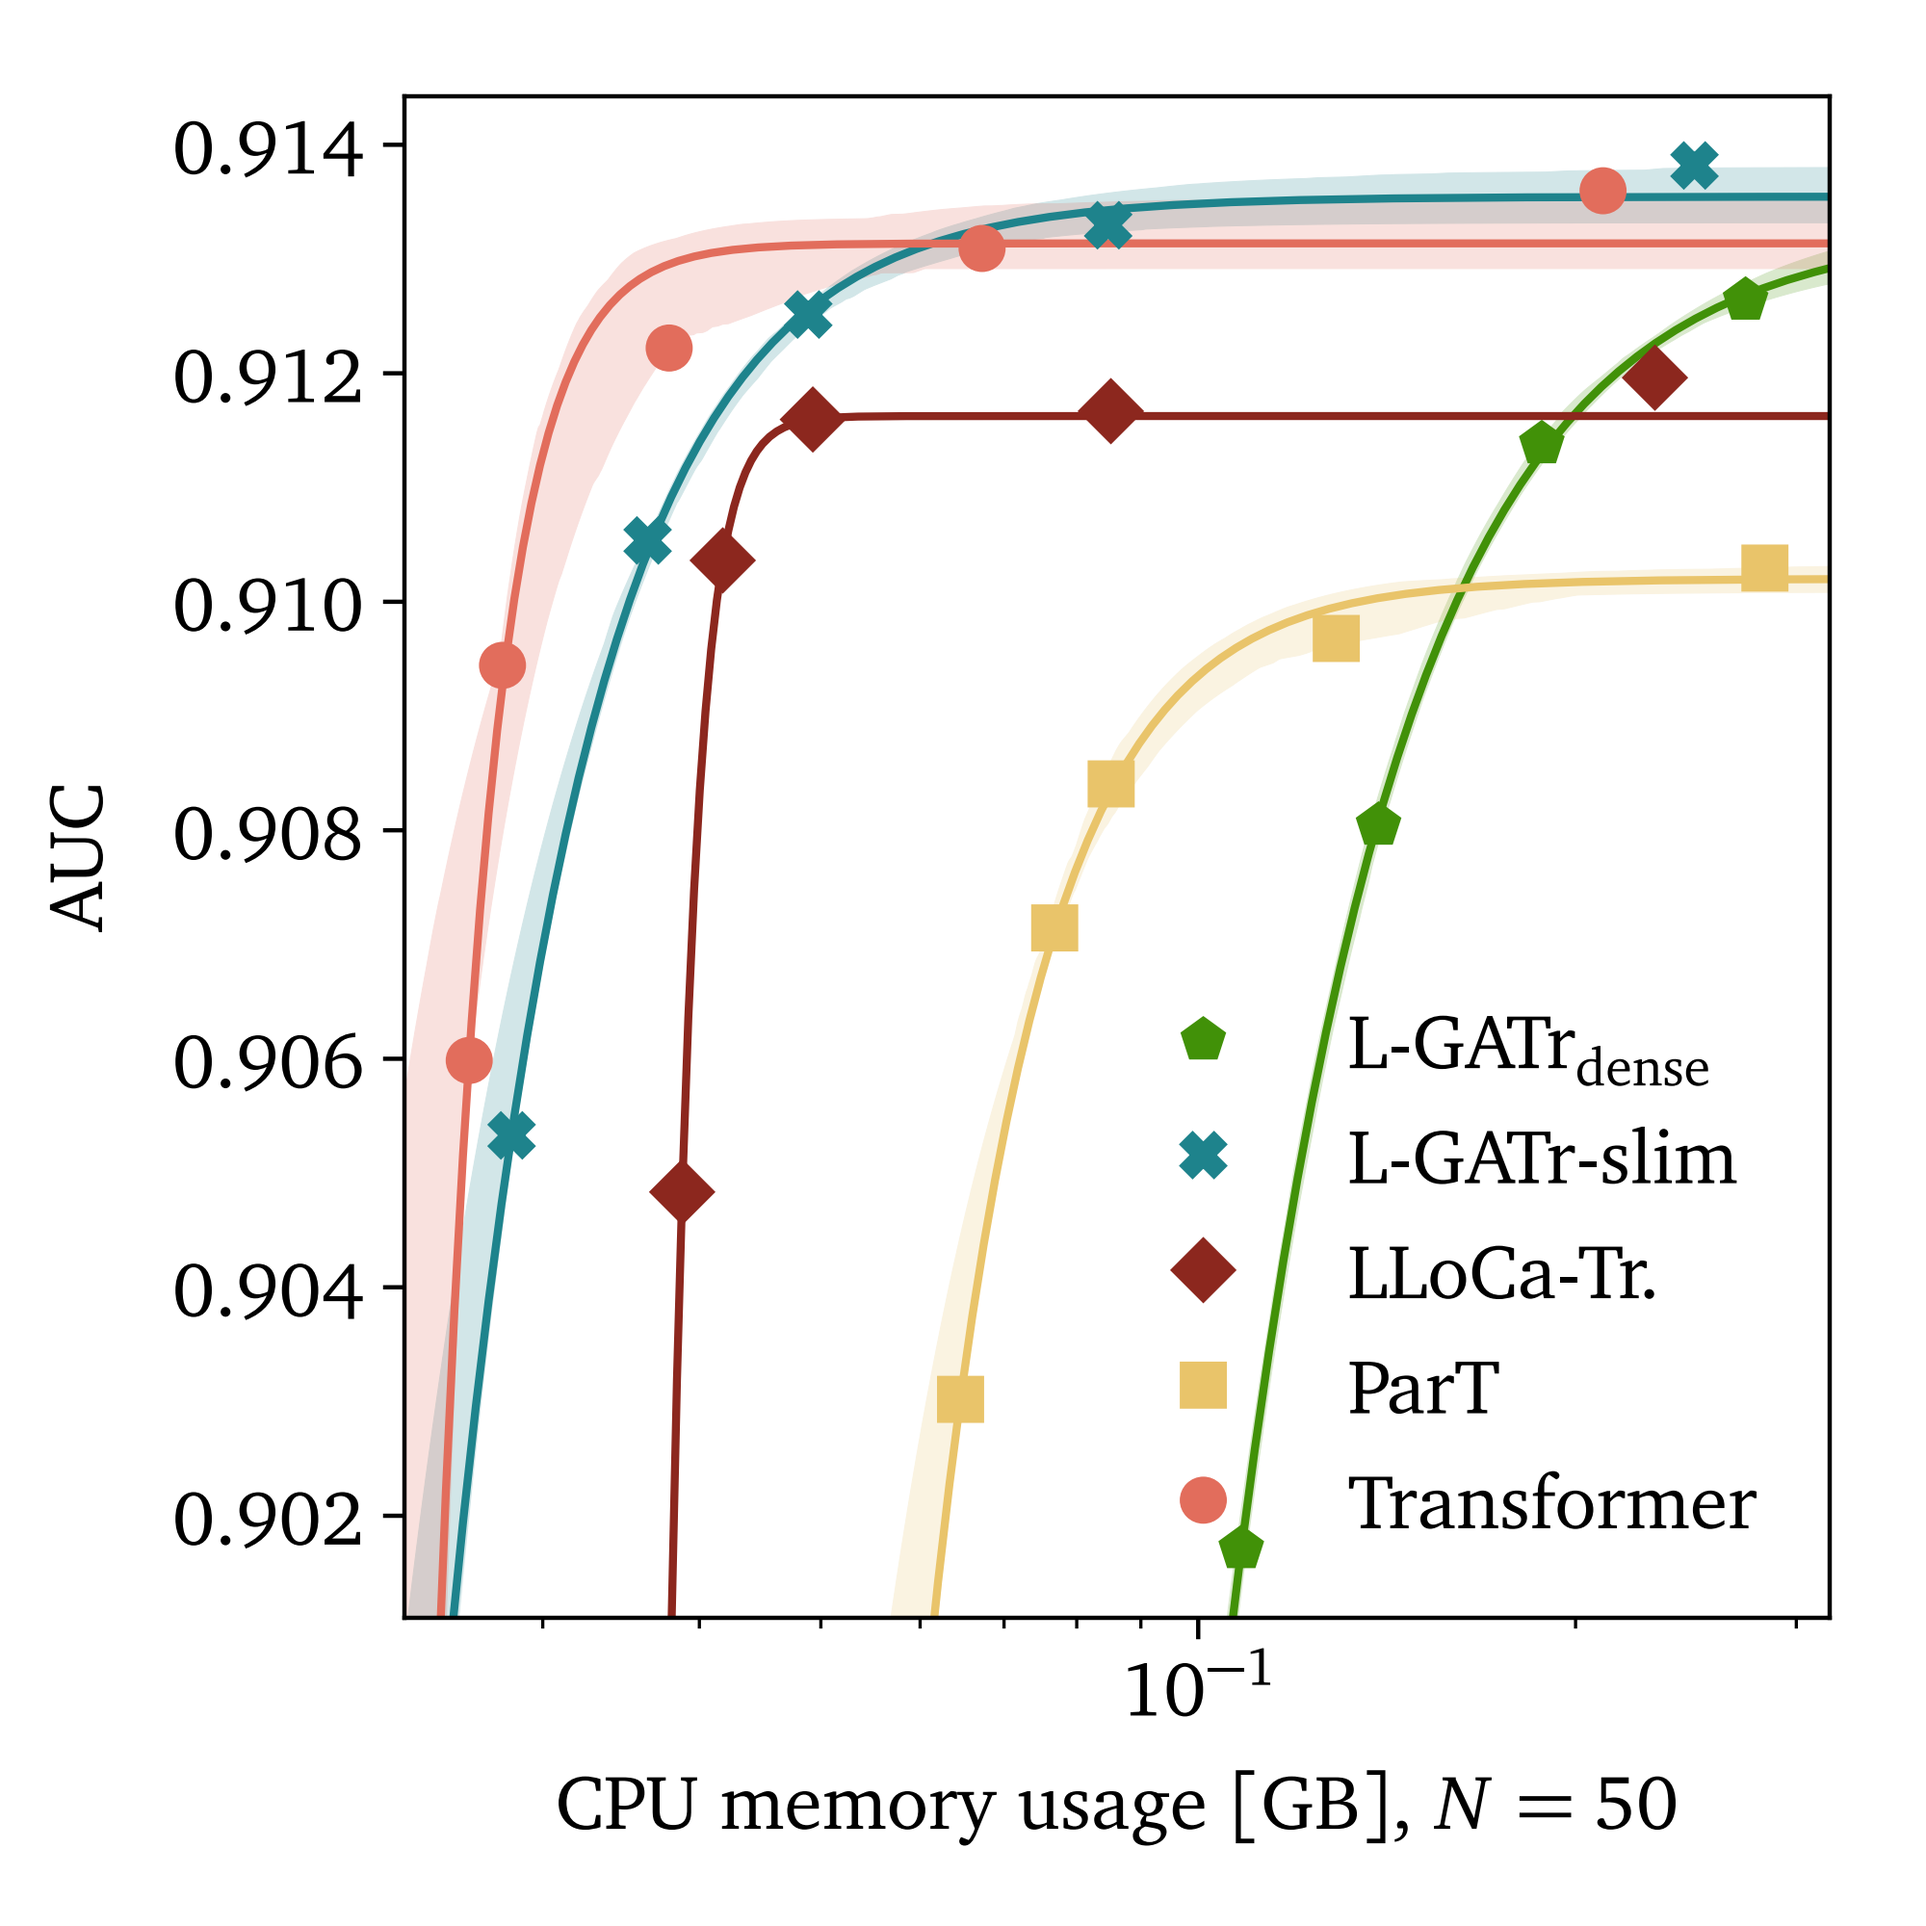
0.914
0.912
0.910
0.908
0.906
0.904
0.902
AUC
10^{-1}
CPU memory usage [GB], $N = 50$
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

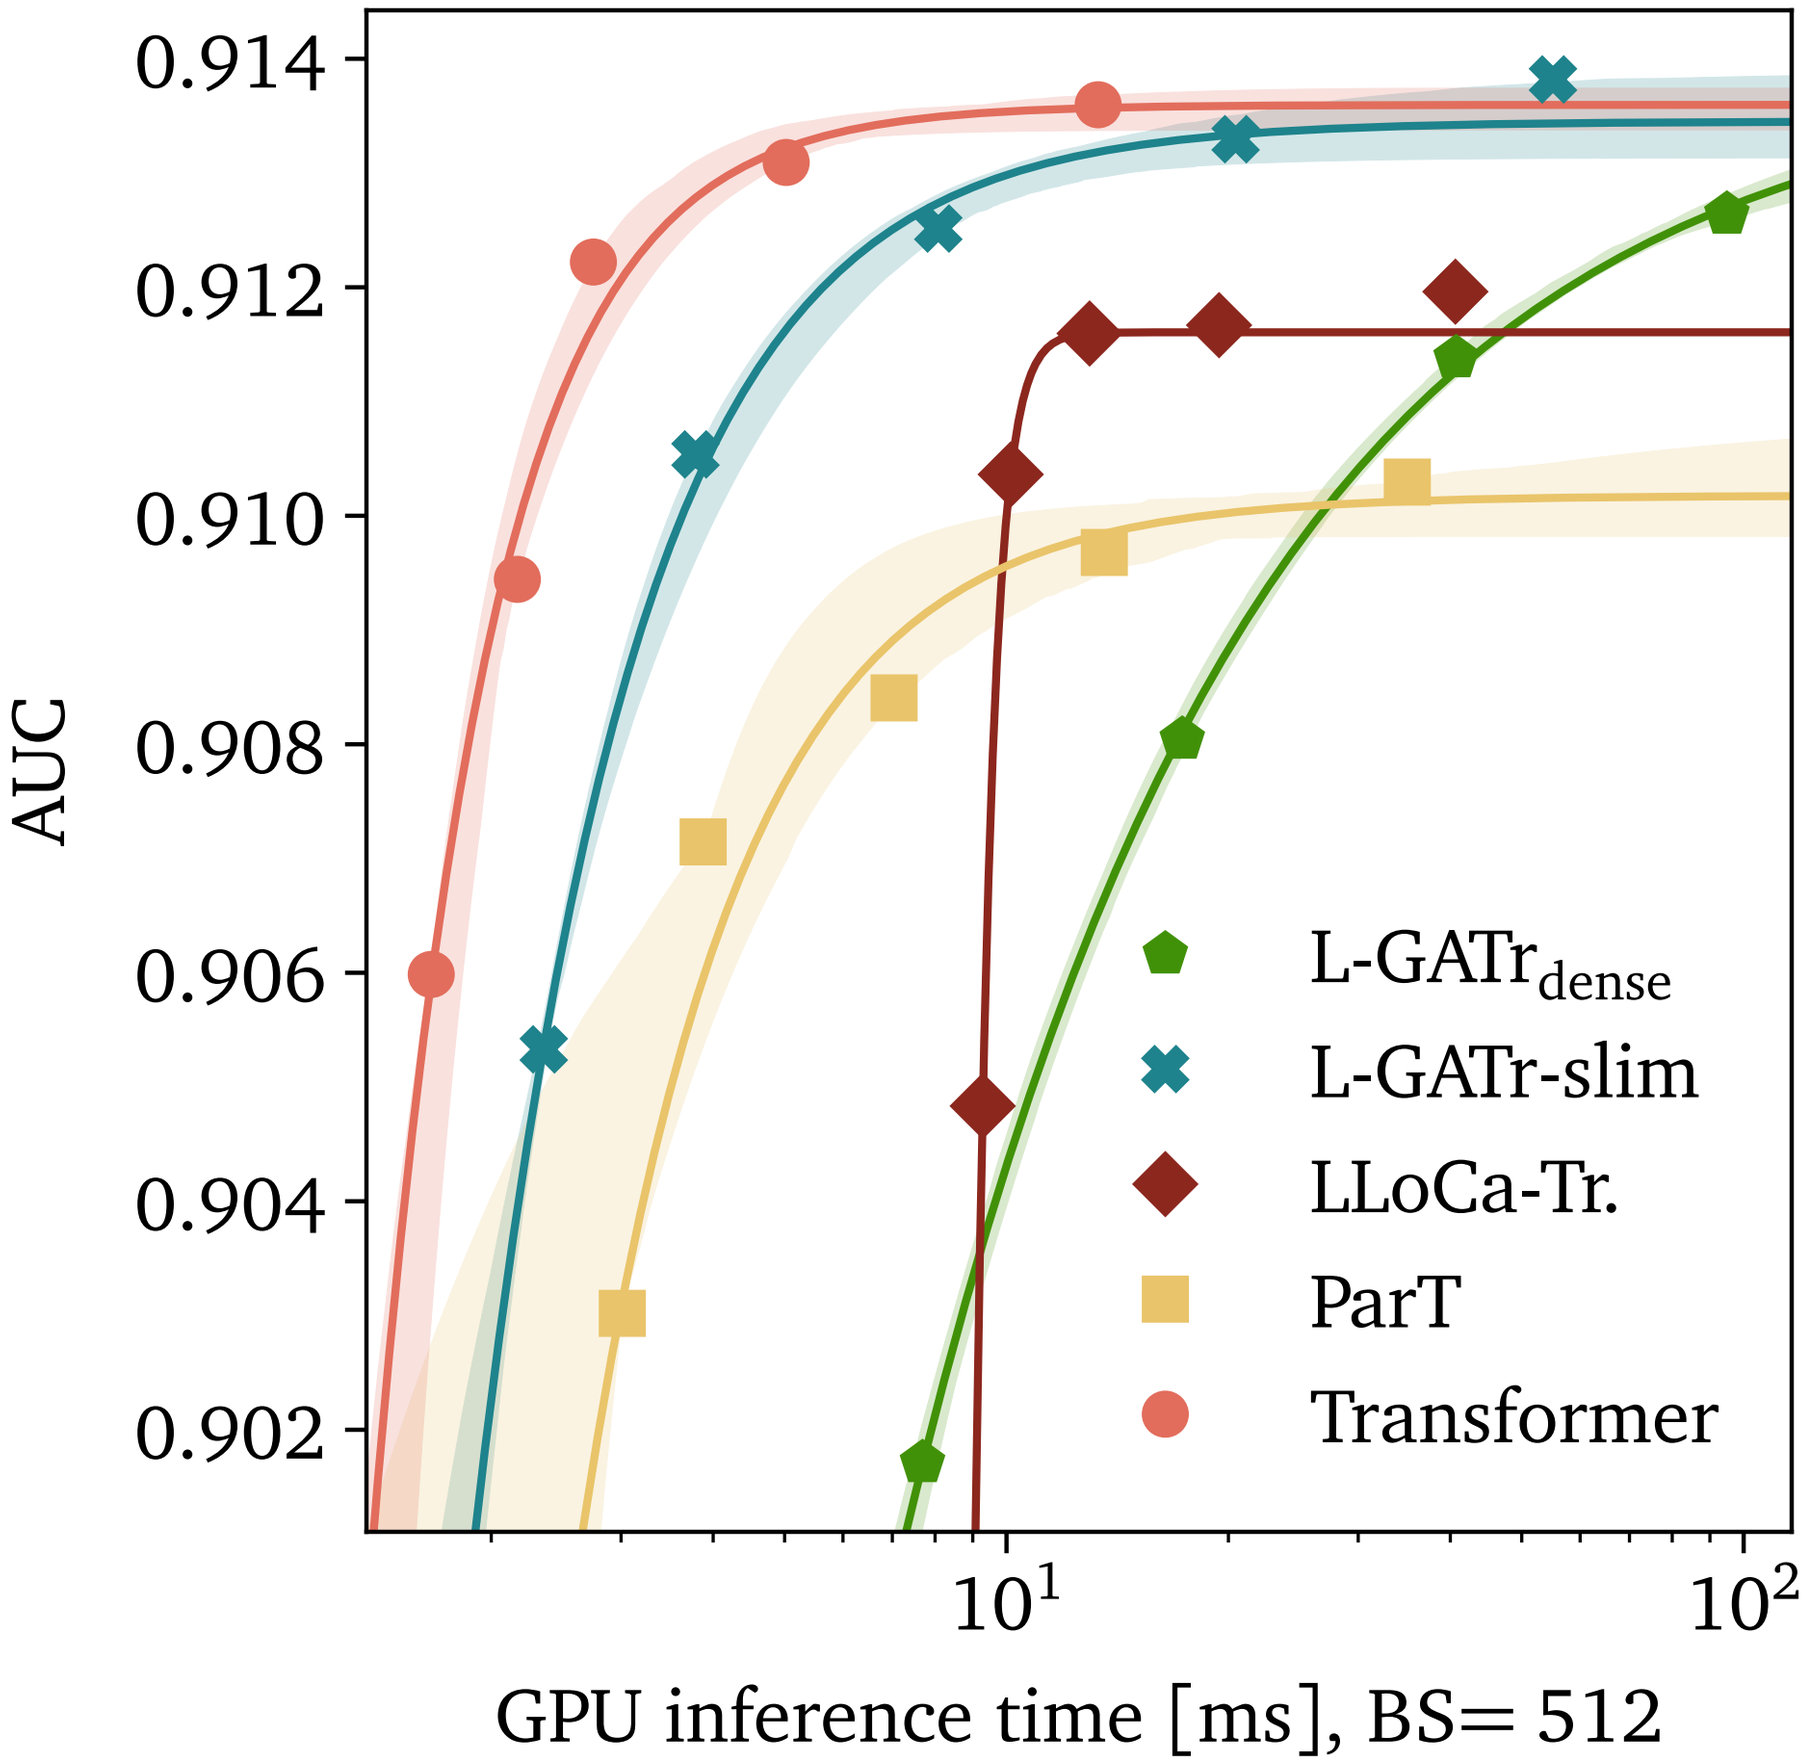
0.914
0.912
0.910
0.908
0.906
0.904
0.902
AUC
10^1
10^2
GPU inference time [ms], BS= 512
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

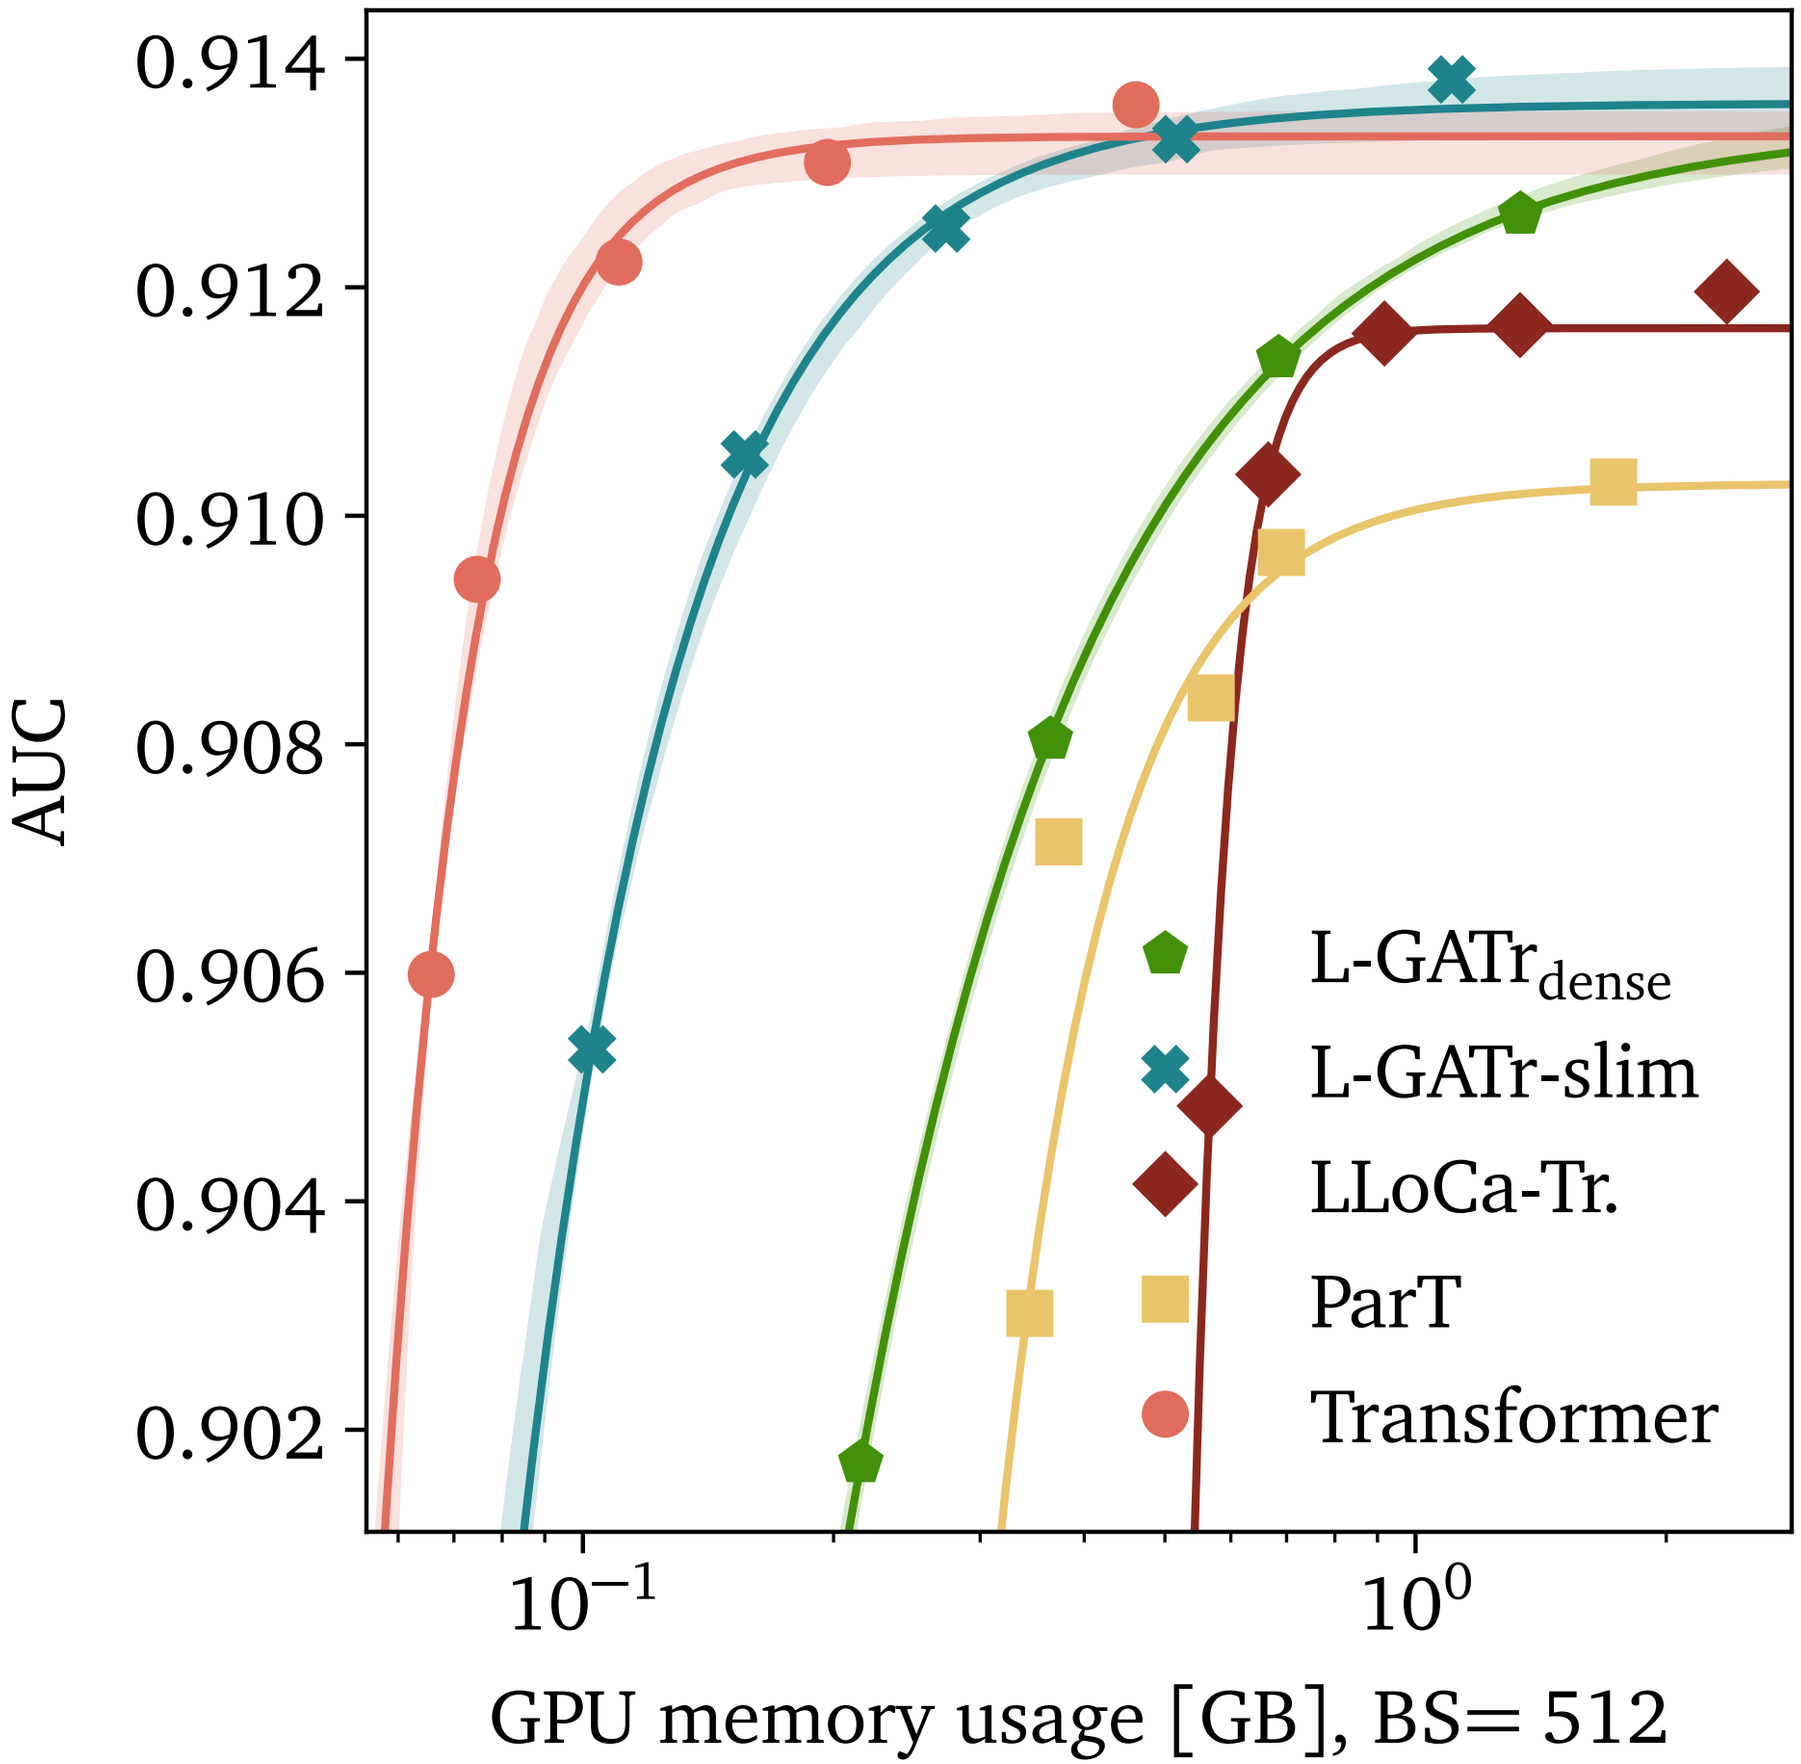
0.914
0.912
0.910
0.908
0.906
0.904
0.902
AUC
10^{-1}
10^{0}
GPU memory usage [GB], BS= 512
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

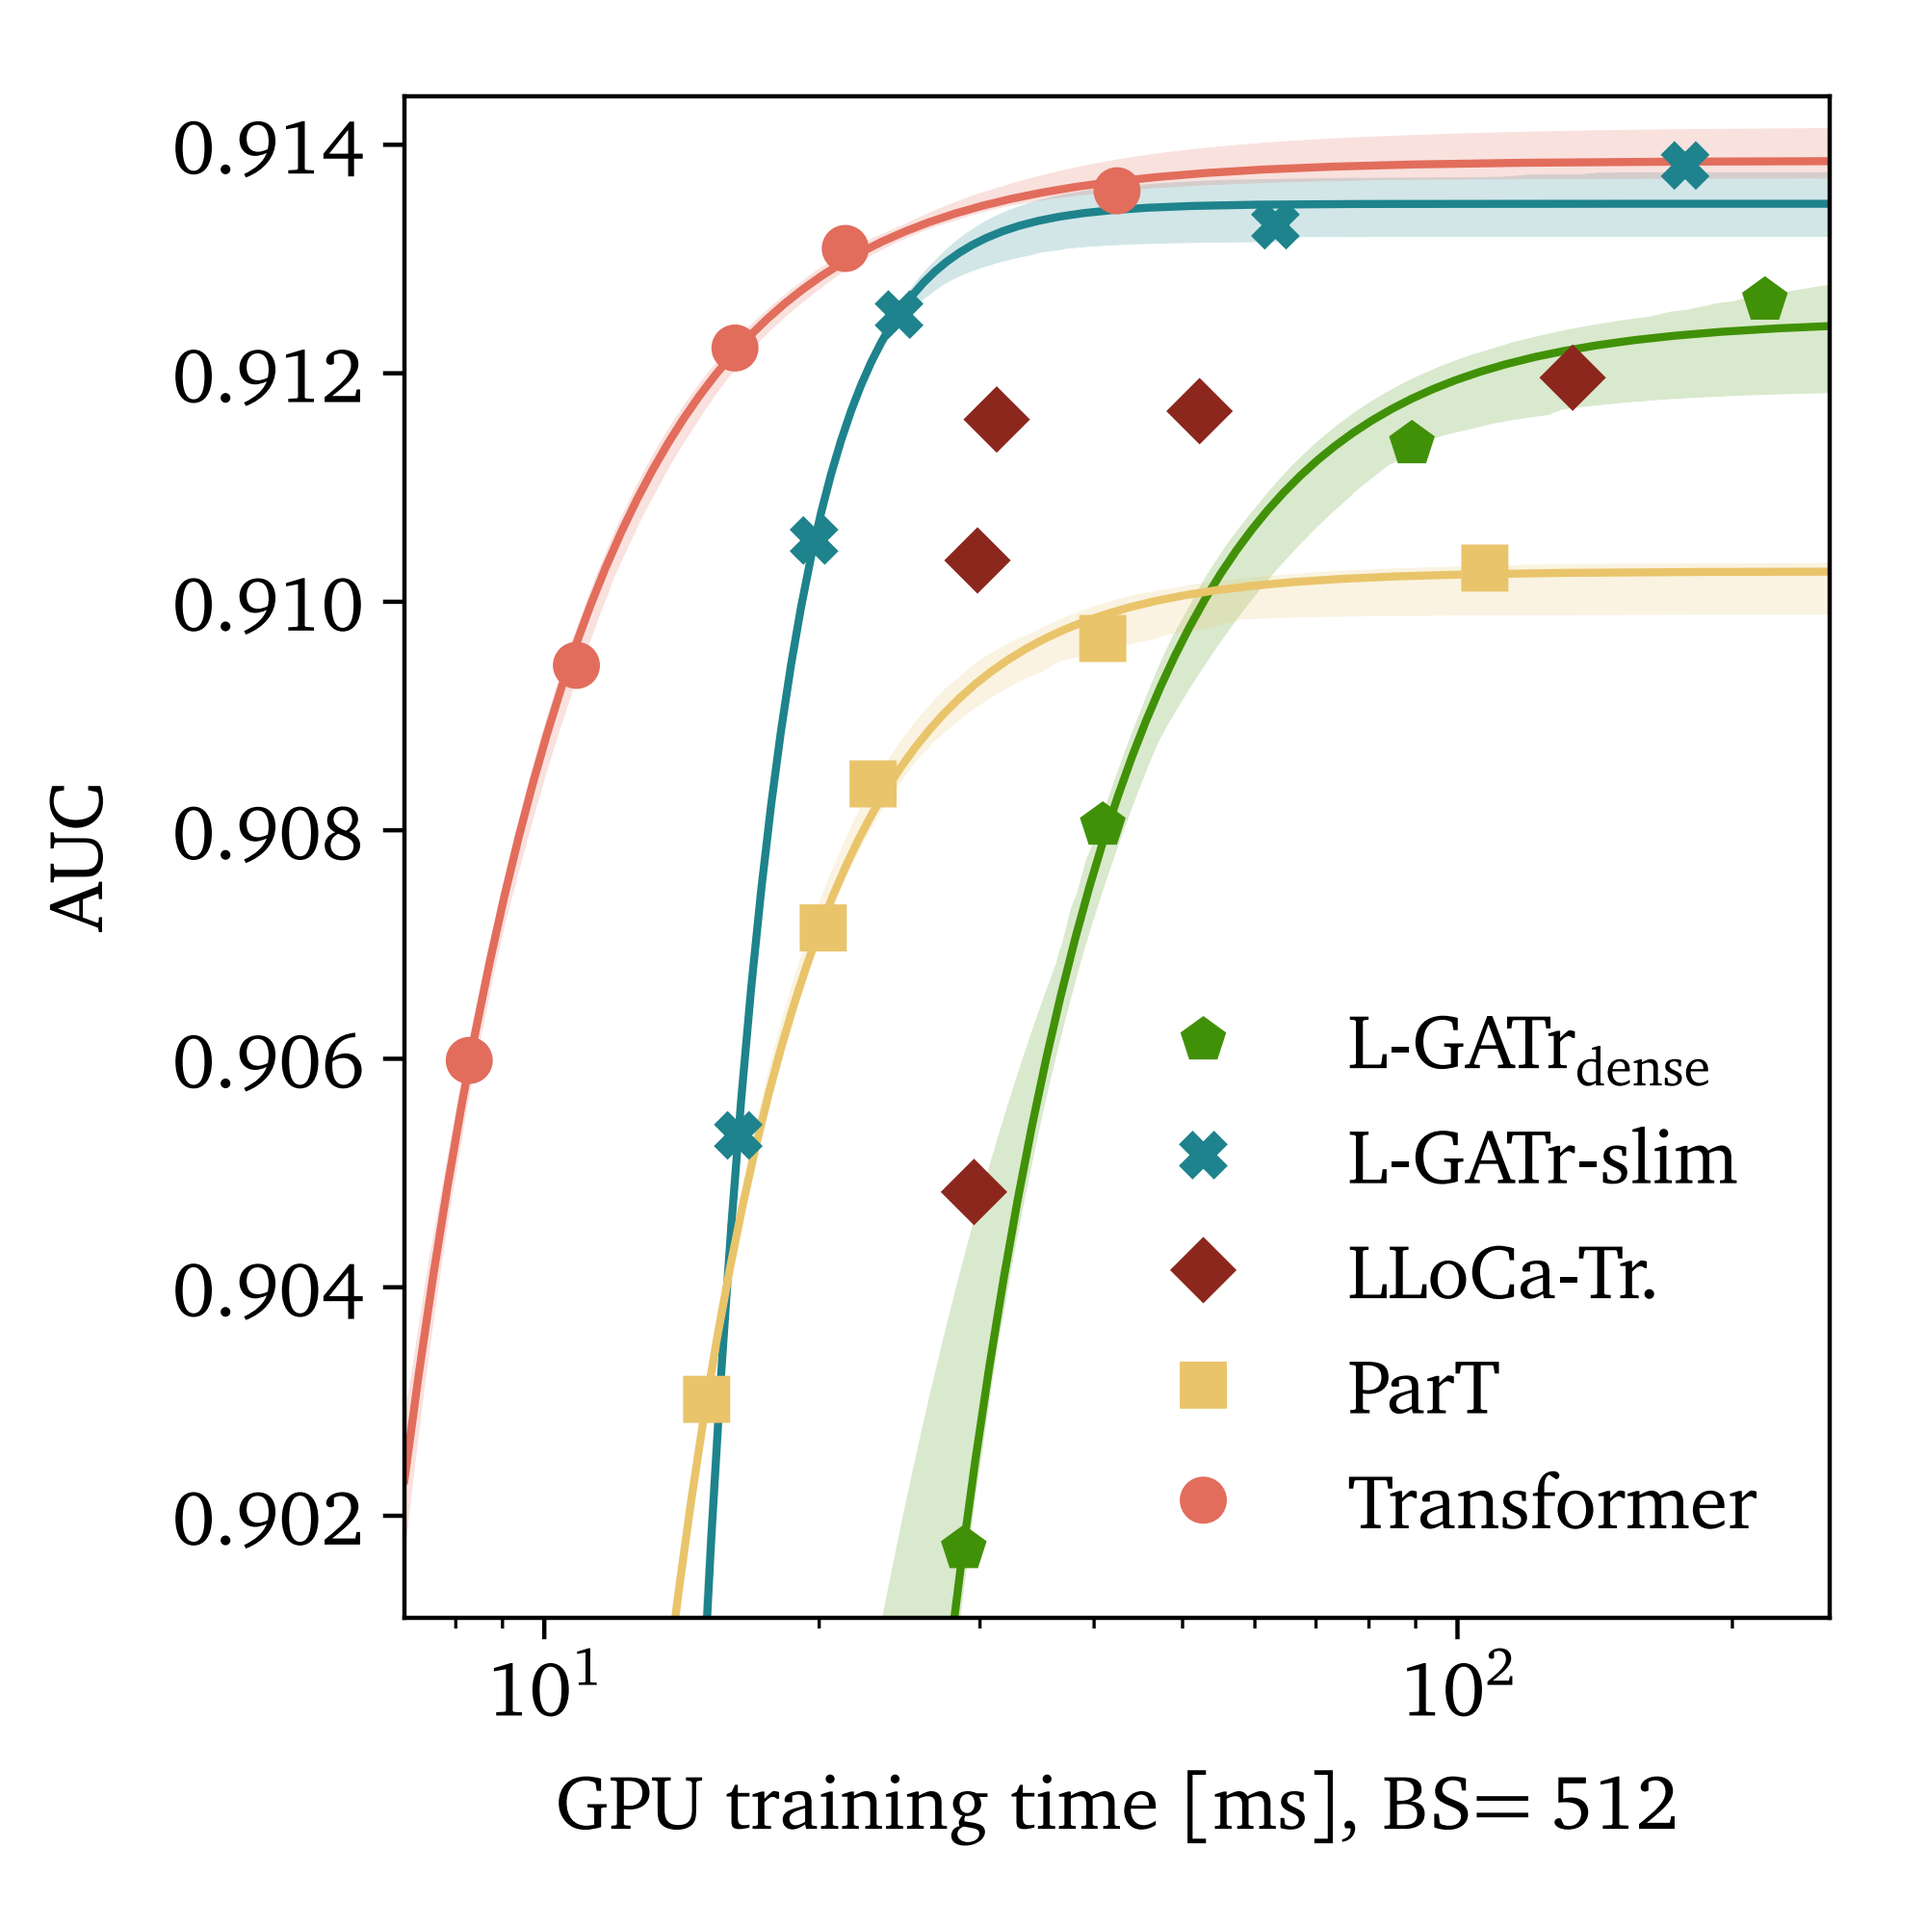
0.914
0.912
0.910
0.908
0.906
0.904
0.902
AUC
10^1
10^2
GPU training time [ms], BS= 512
L-GATr$_{\text{dense}}$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer