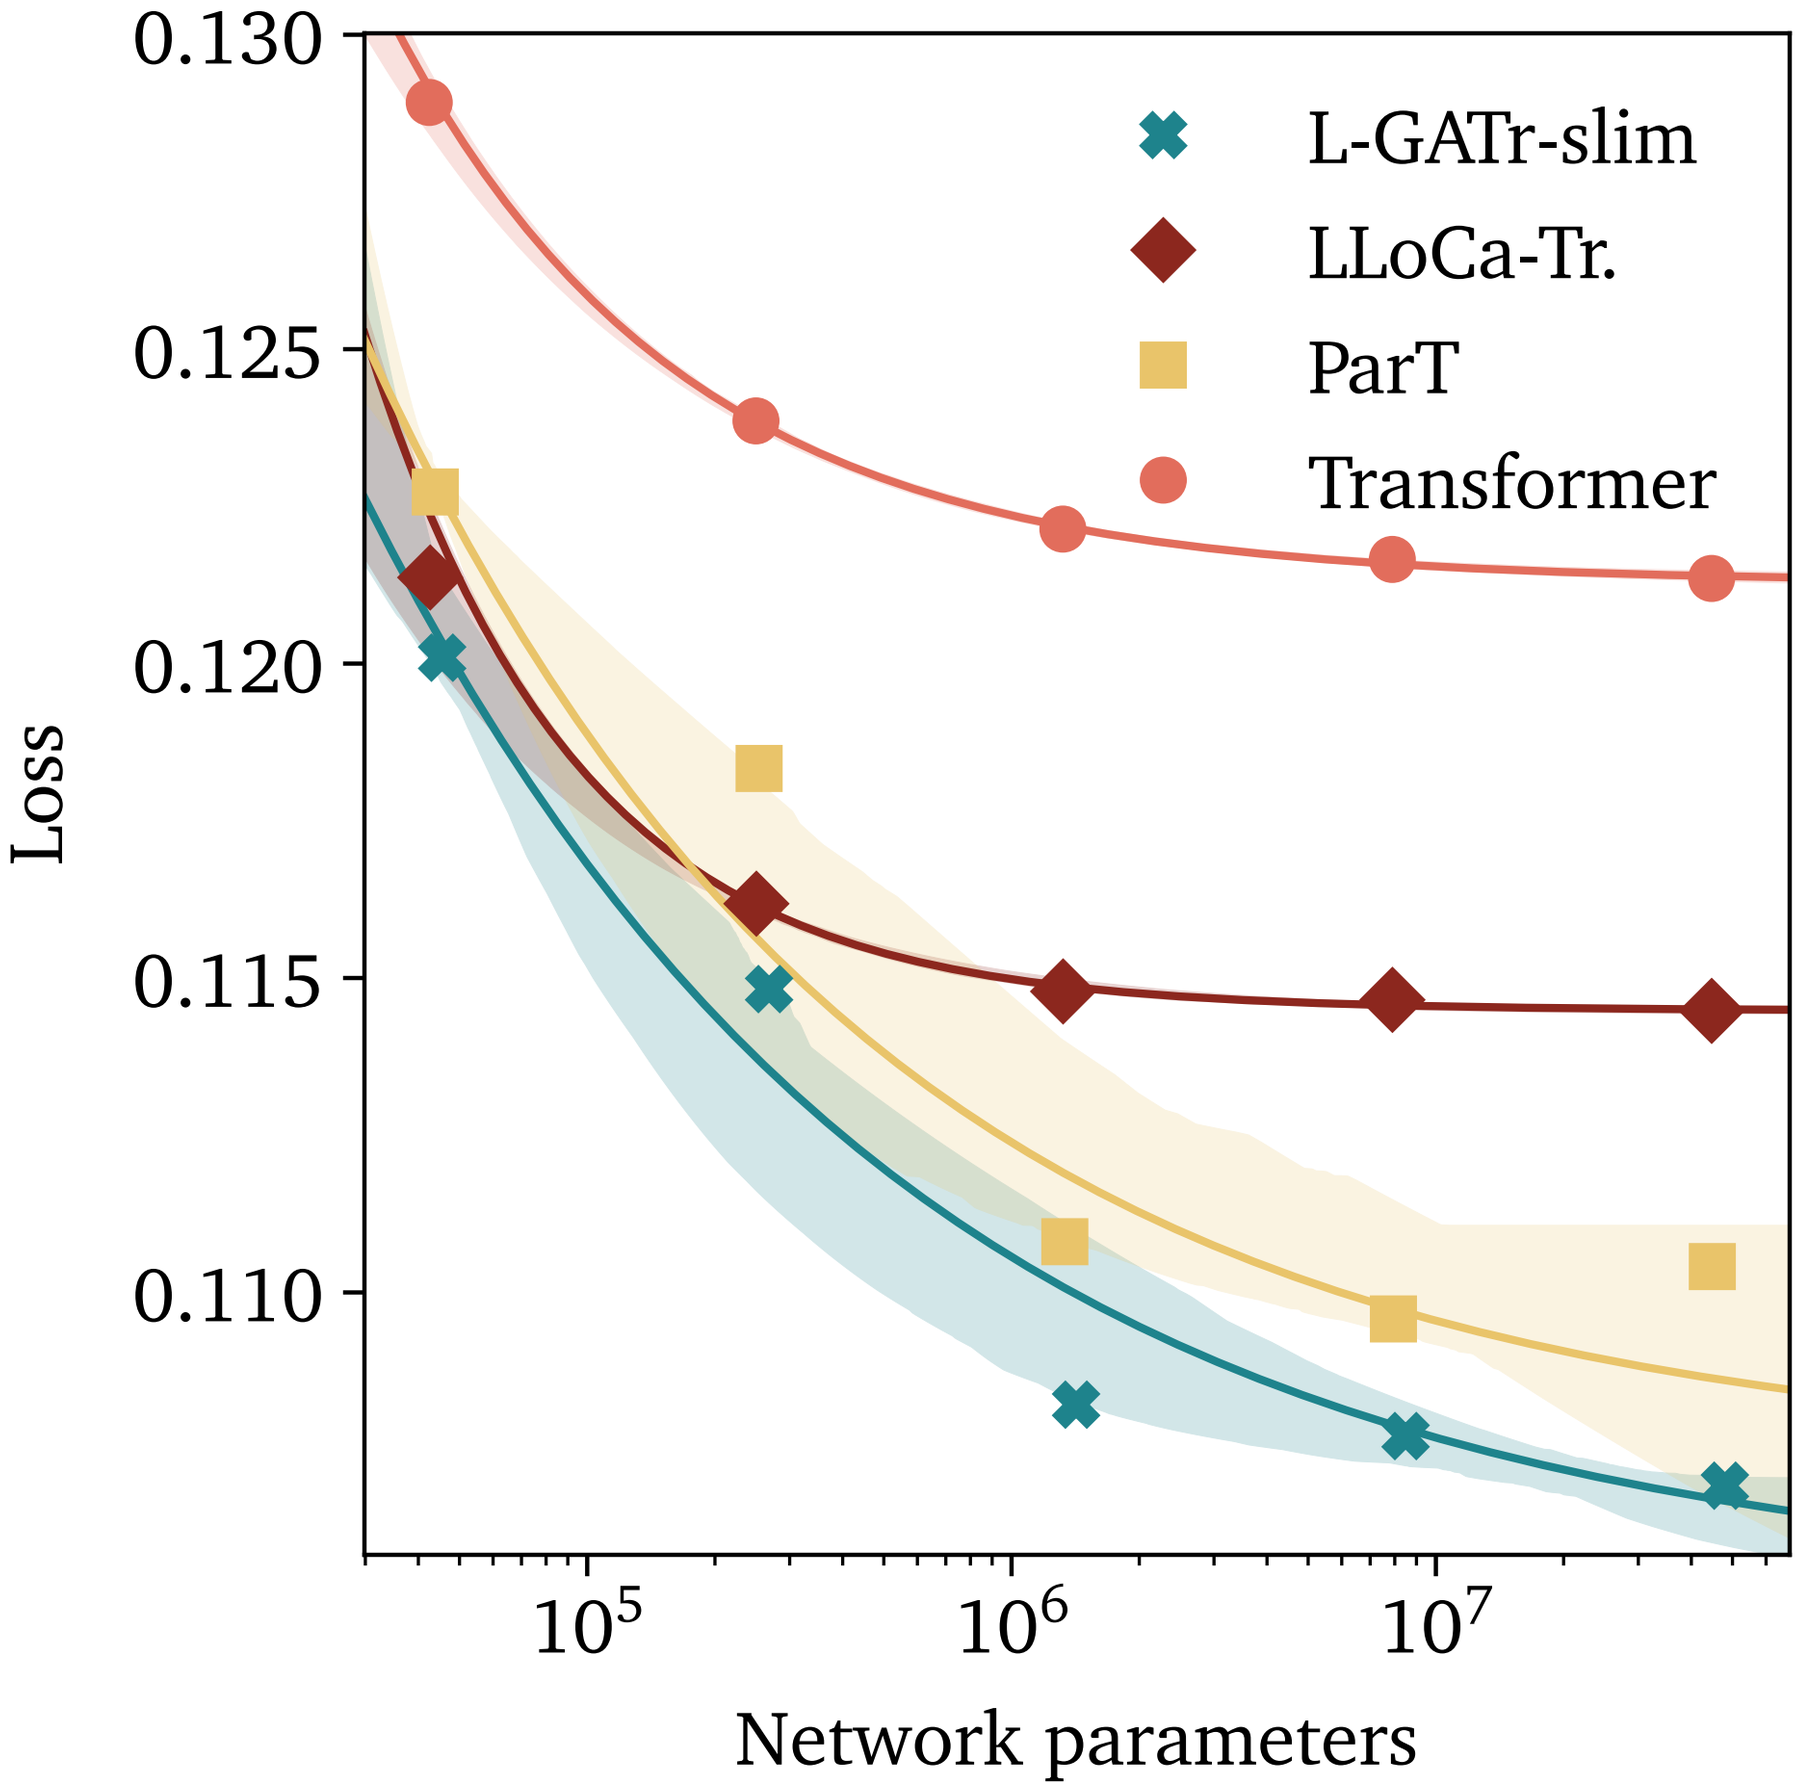
0.130
0.125
0.120
0.115
0.110
Loss
10^5
10^6
10^7
Network parameters
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

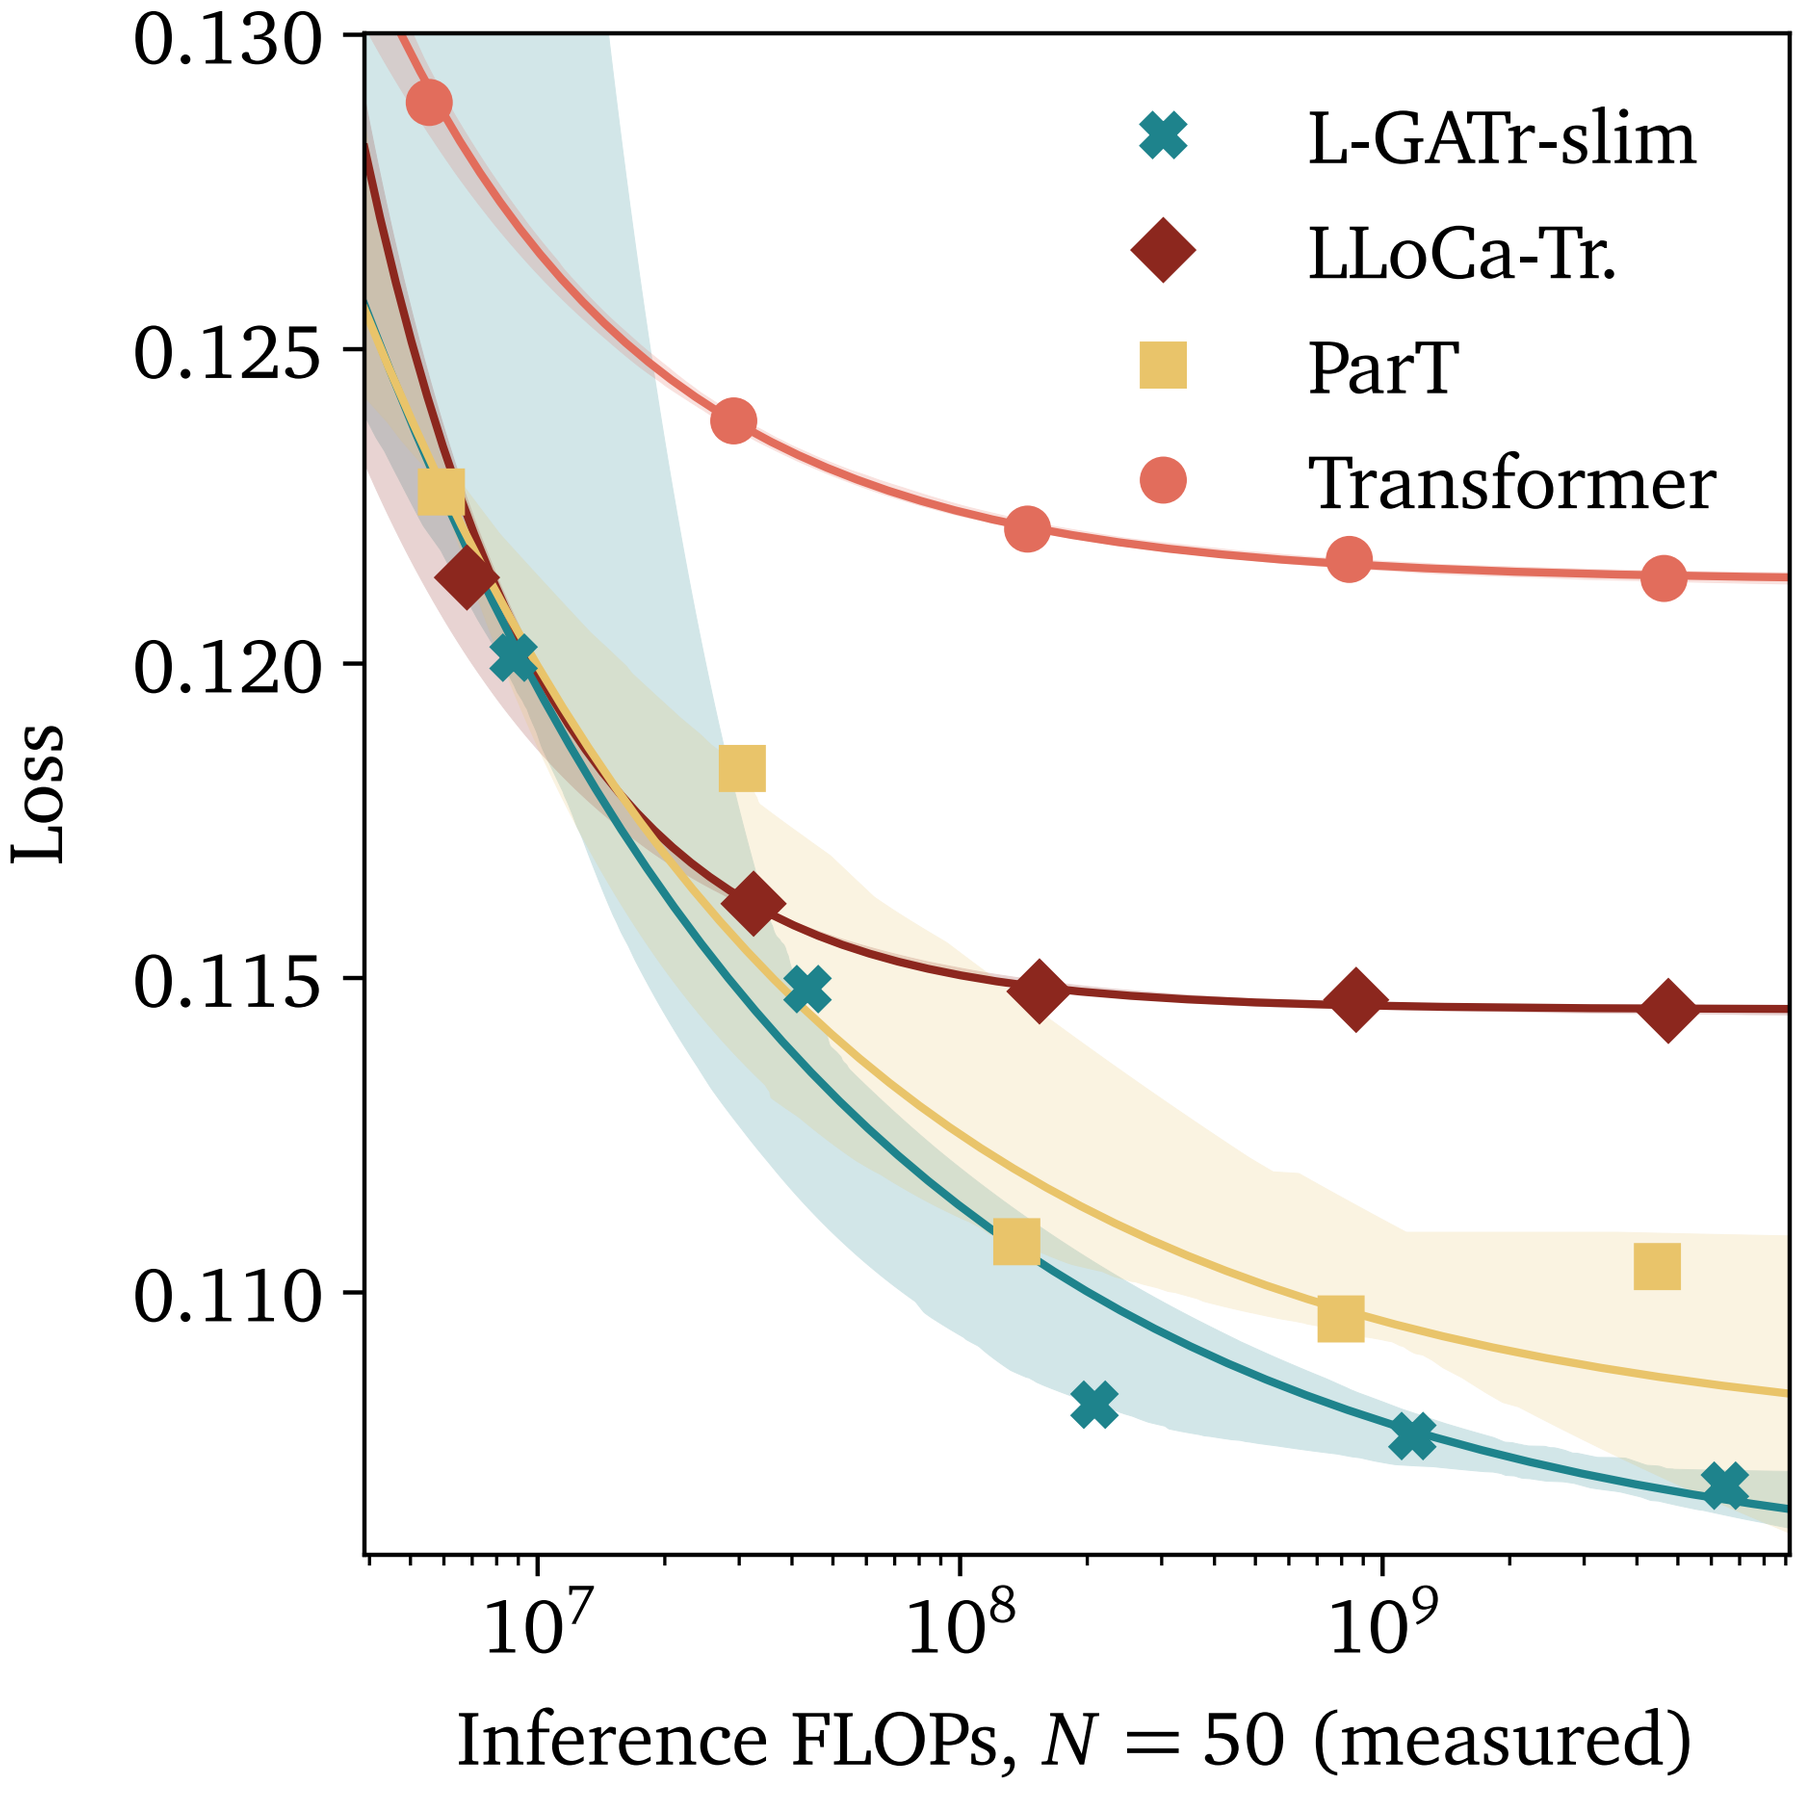
0.130
0.125
0.120
0.115
0.110
Loss
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (measured)

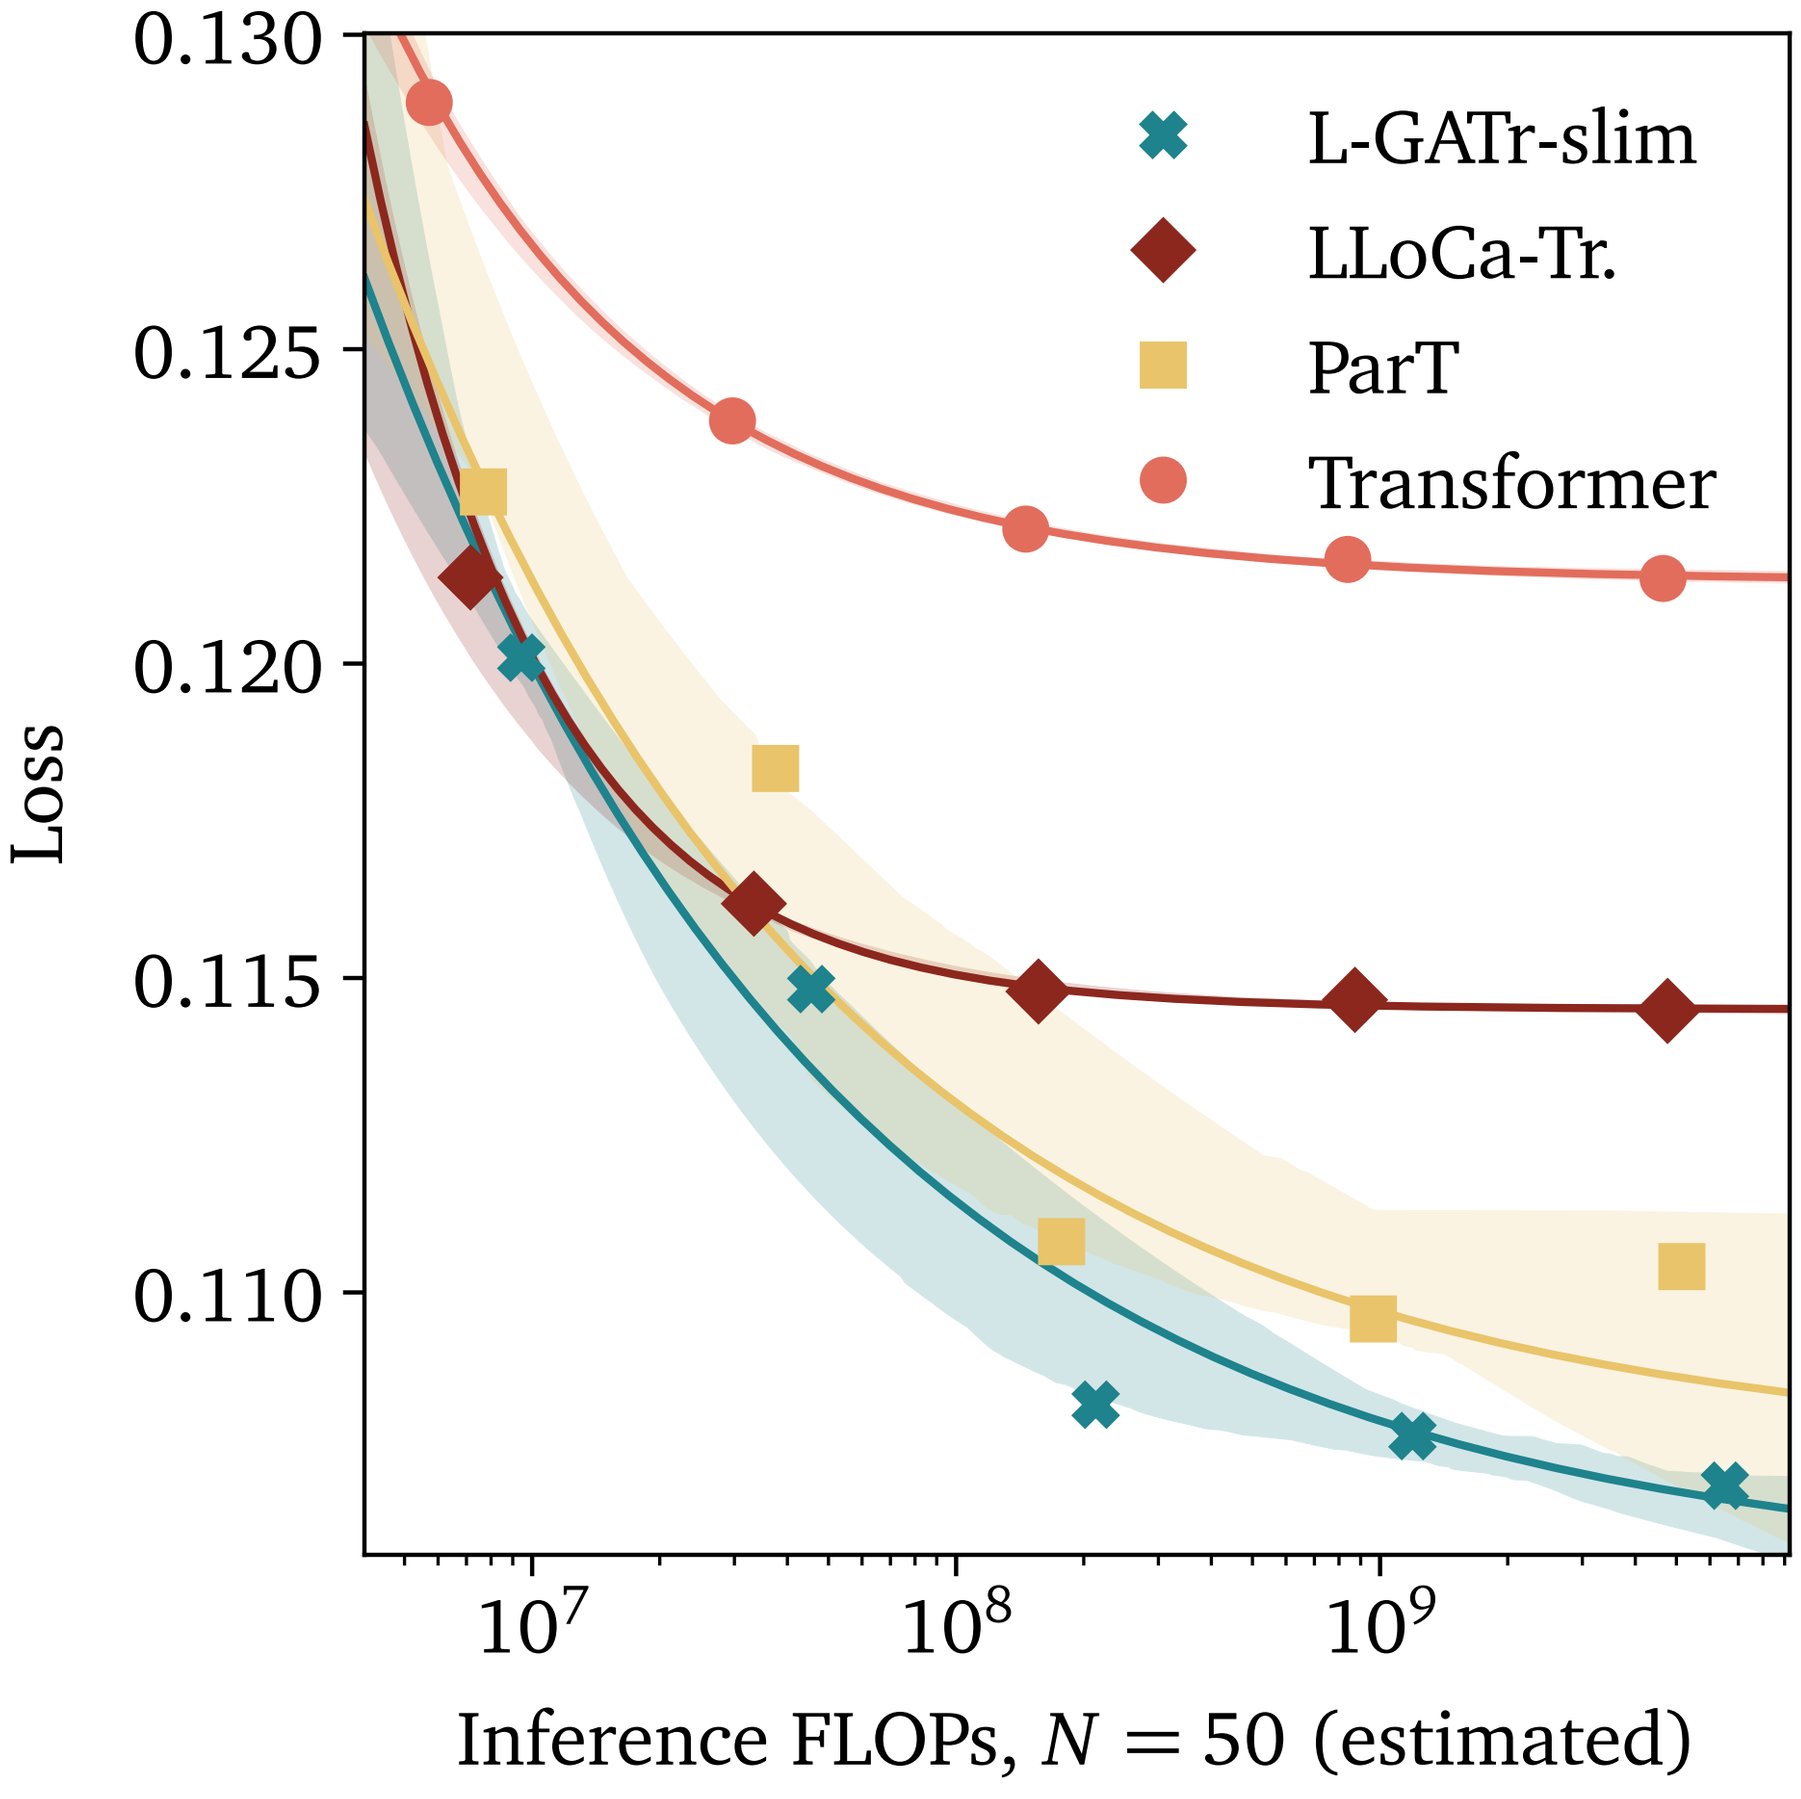
0.130
0.125
0.120
0.115
0.110
Loss
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (estimated)

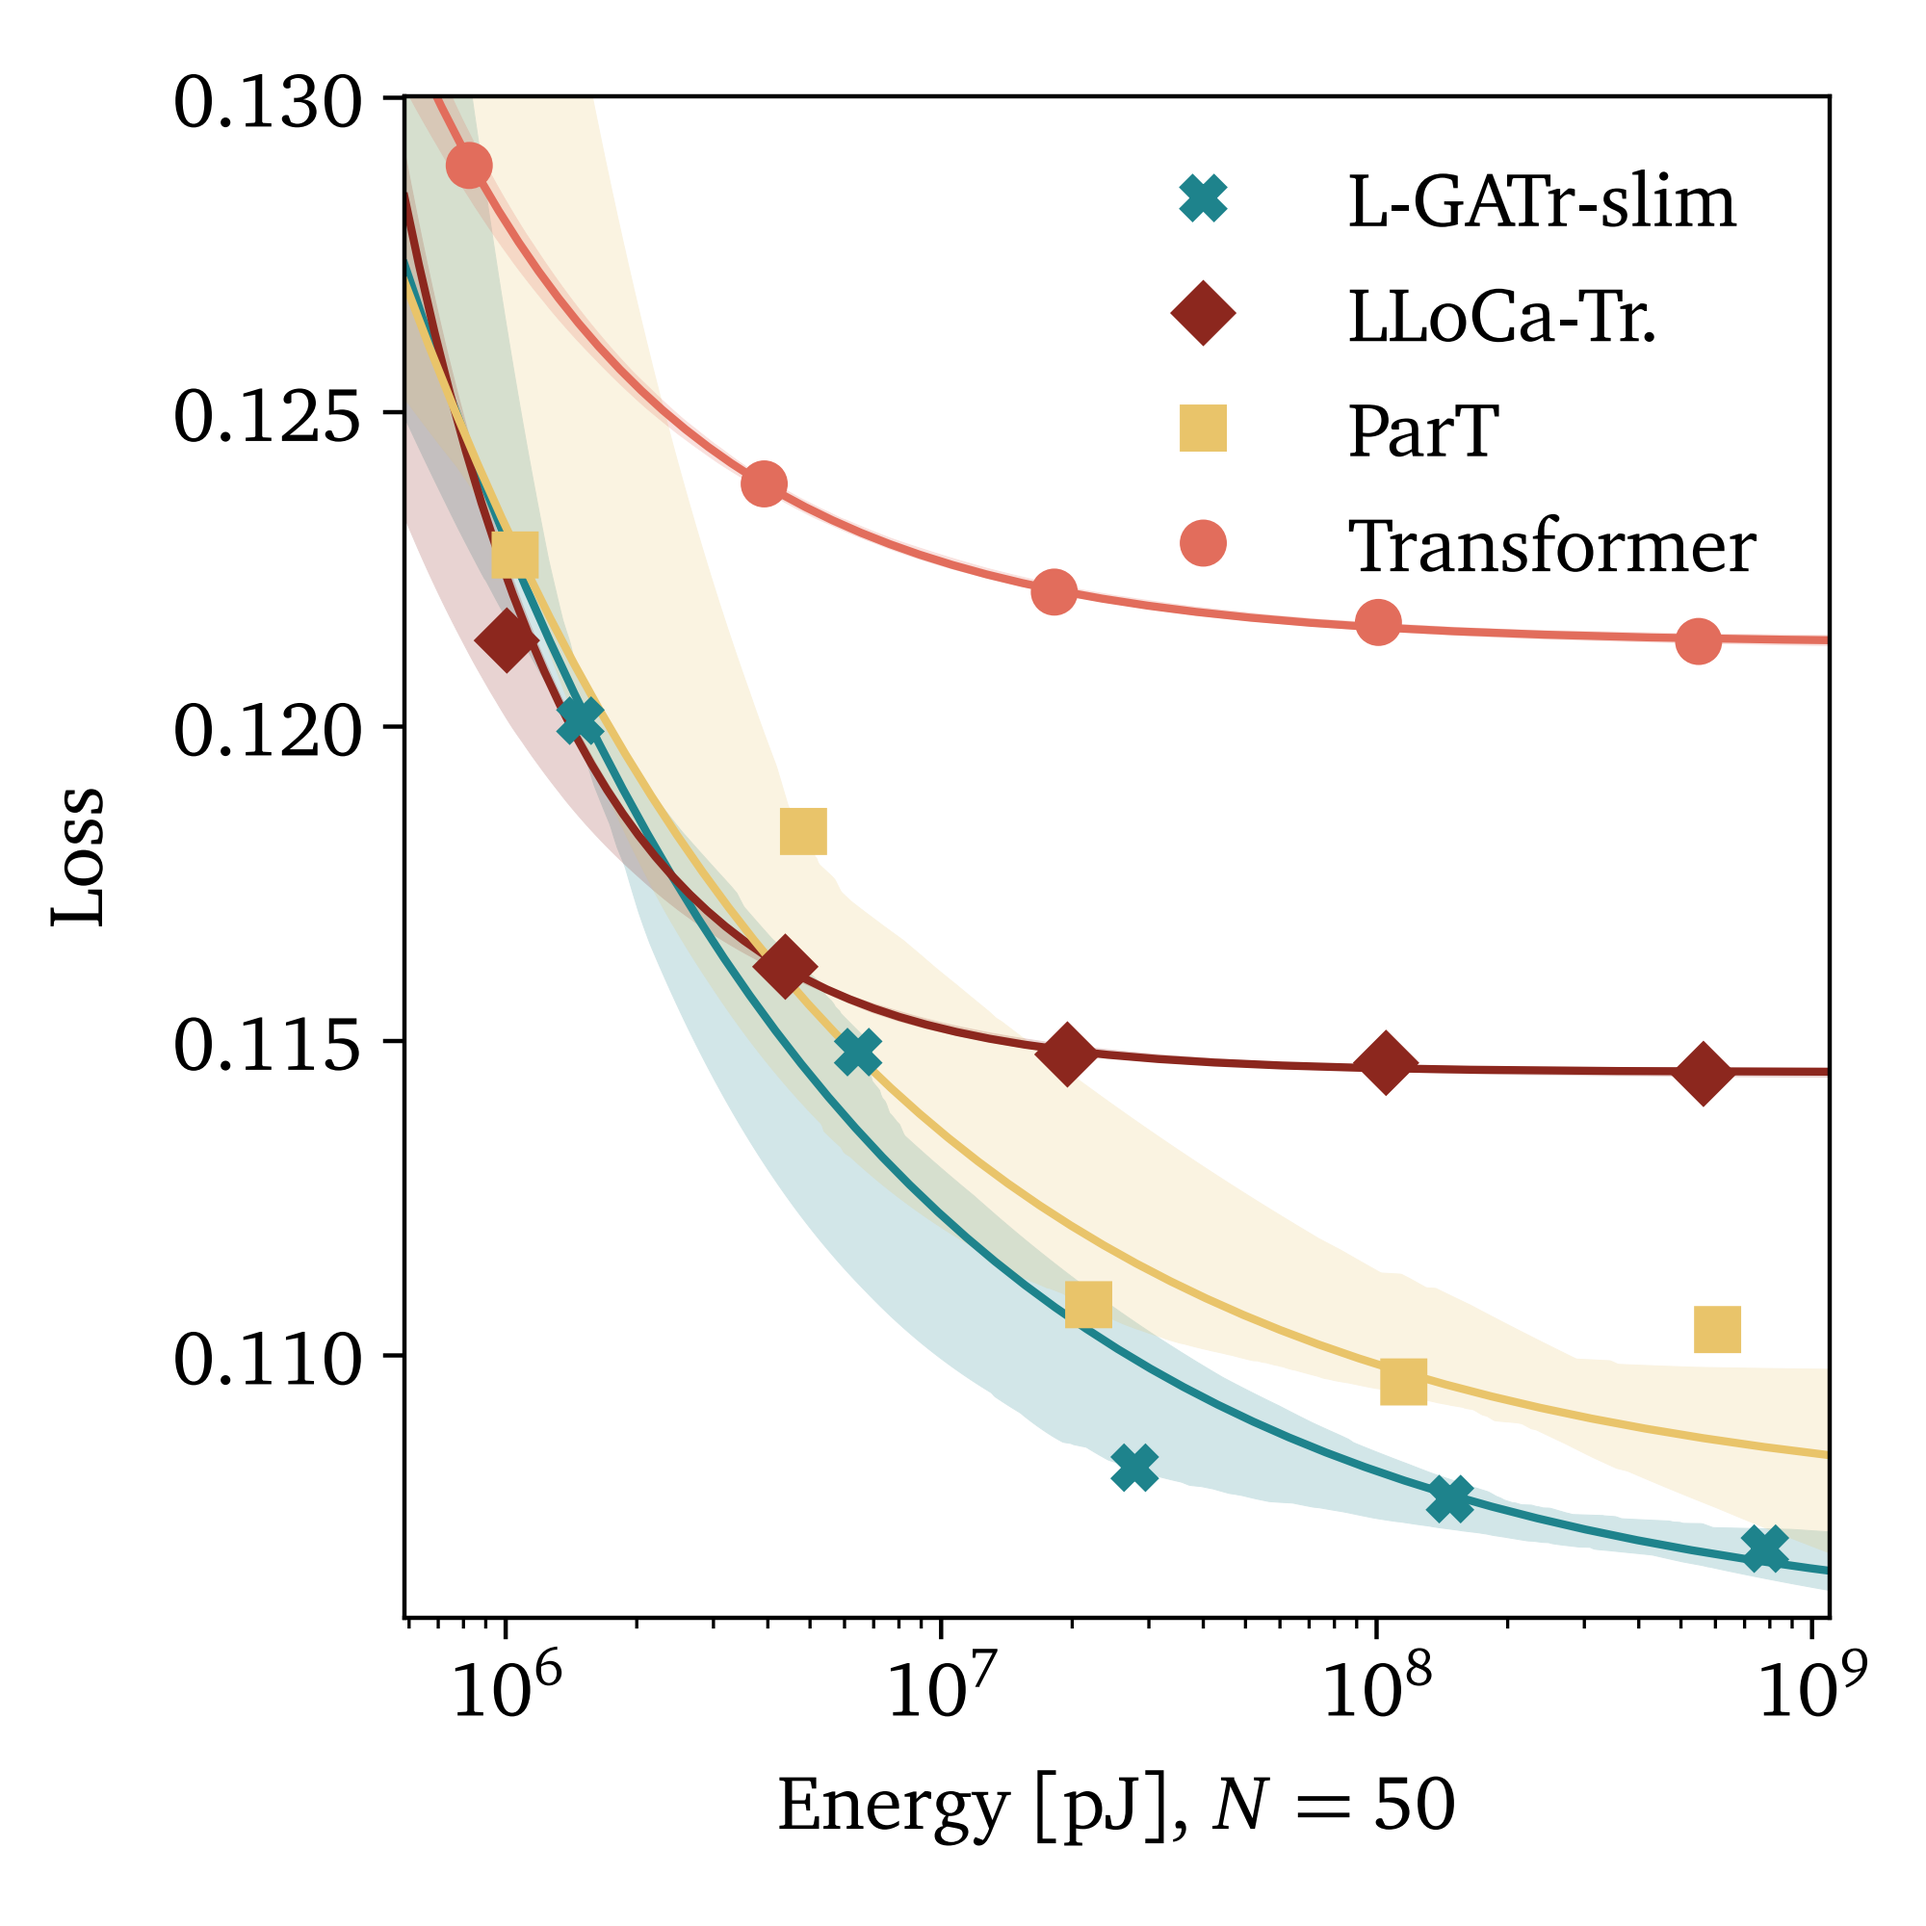
0.130
0.125
0.120
0.115
0.110
Loss
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
10^6
10^7
10^8
10^9
Energy [pJ], $N = 50$

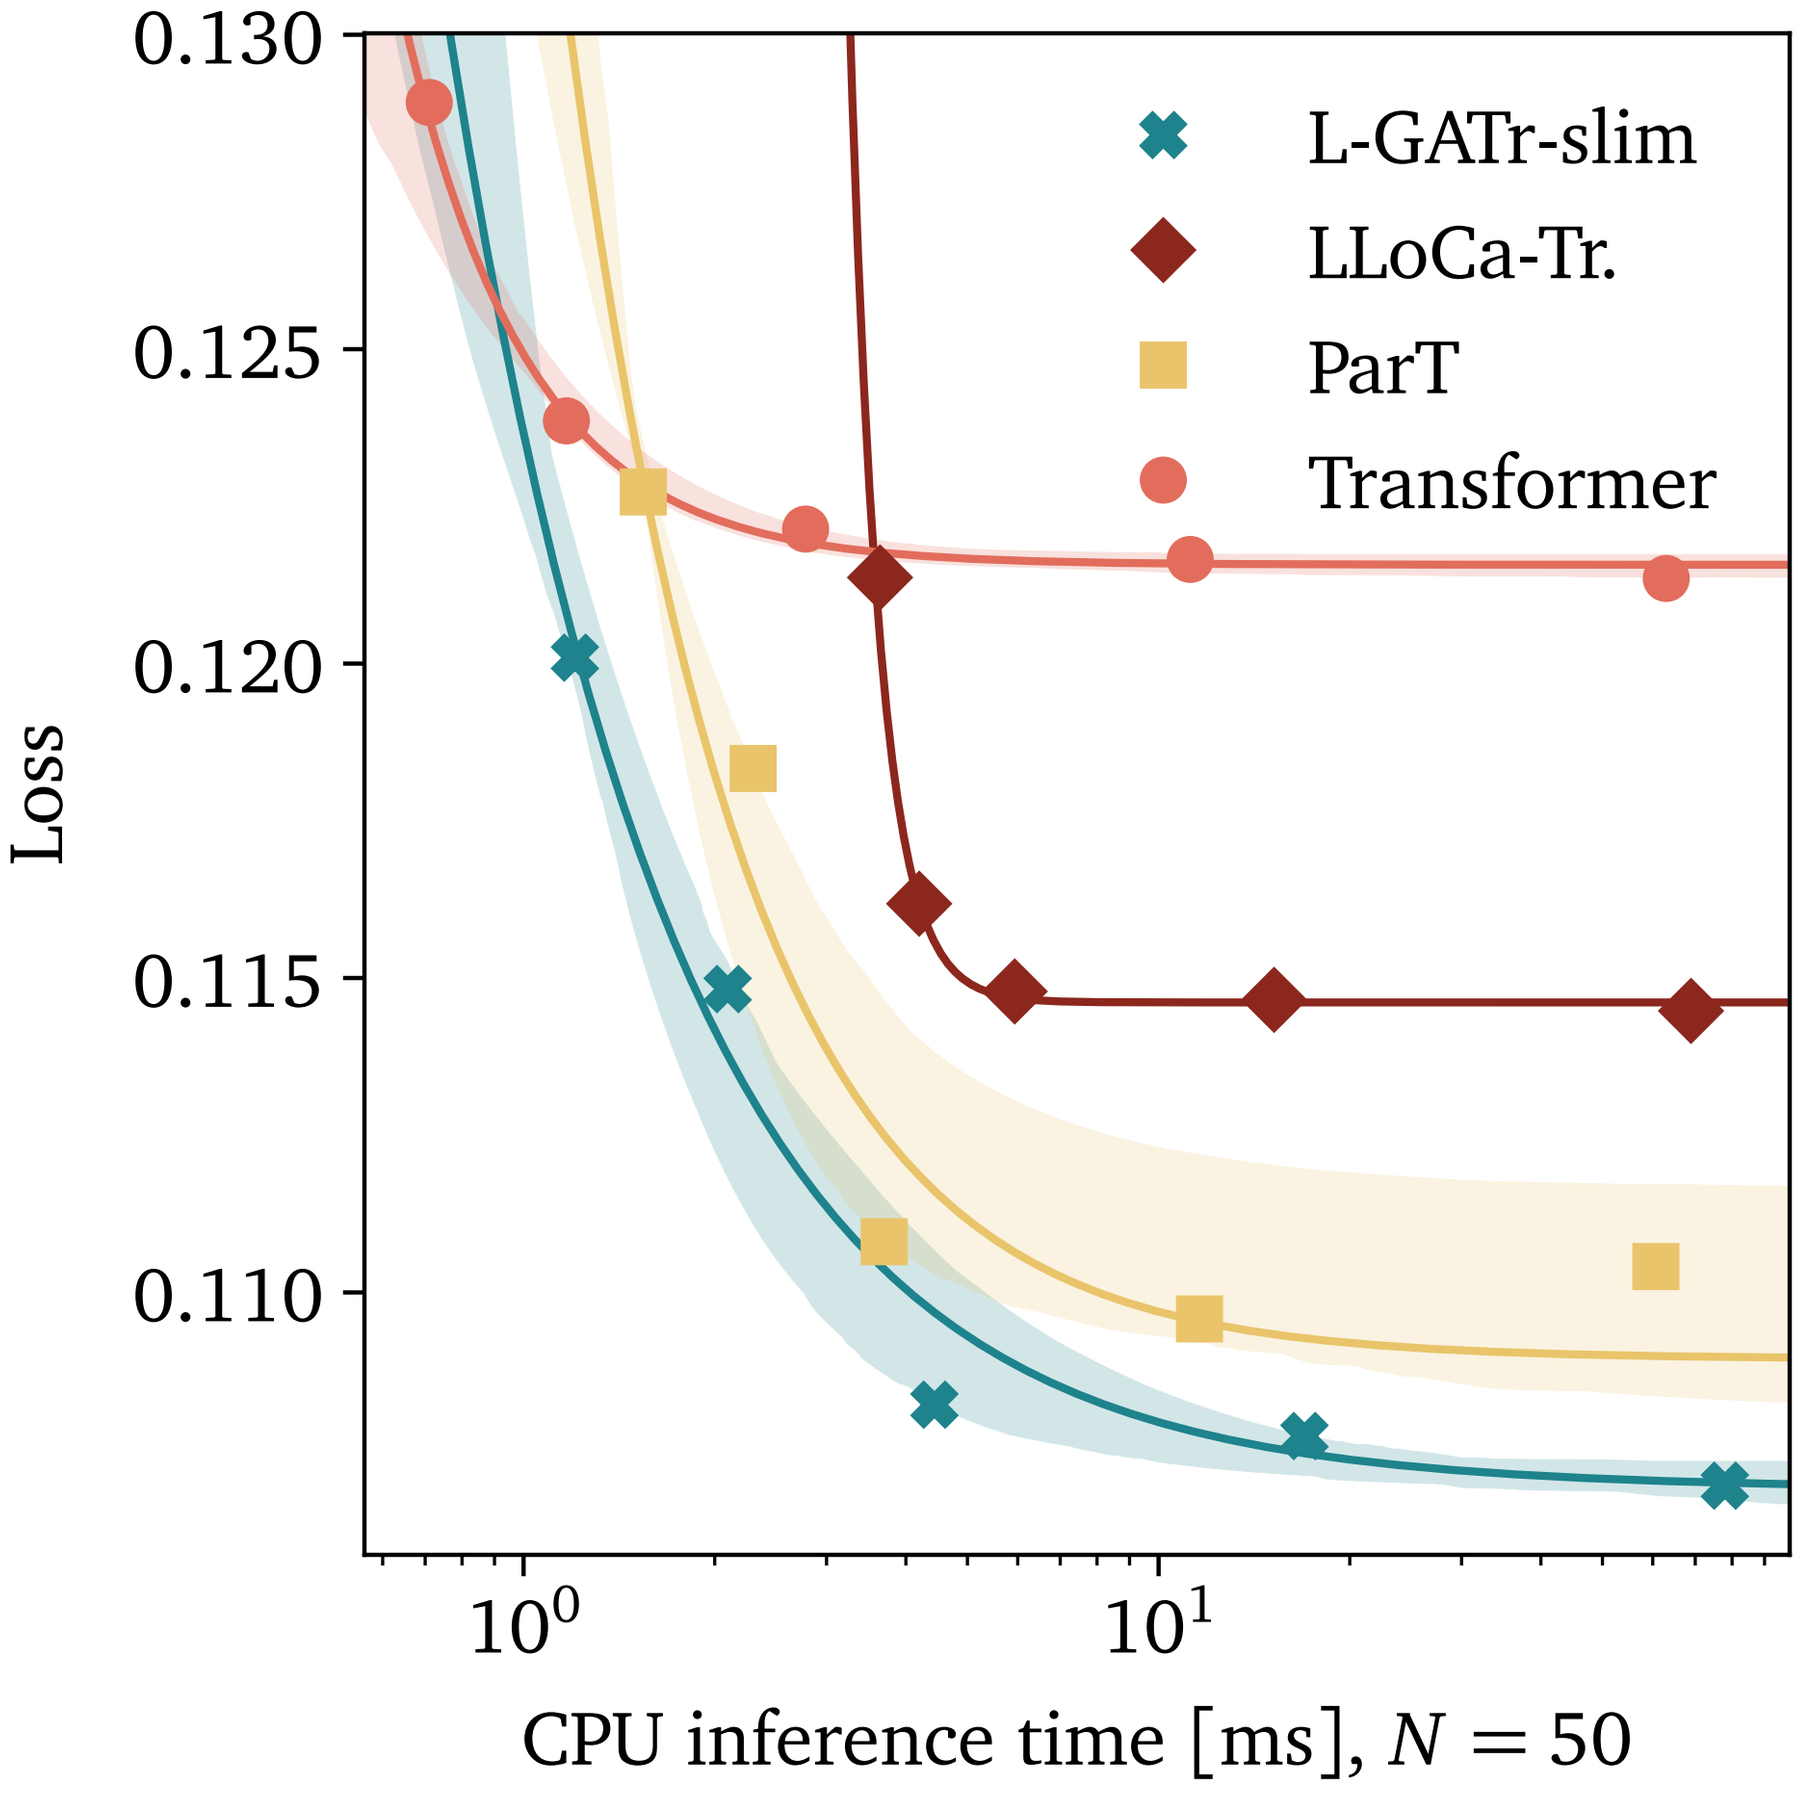
0.130
0.125
0.120
0.115
0.110
Loss
10^0
10^1
CPU inference time [ms], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

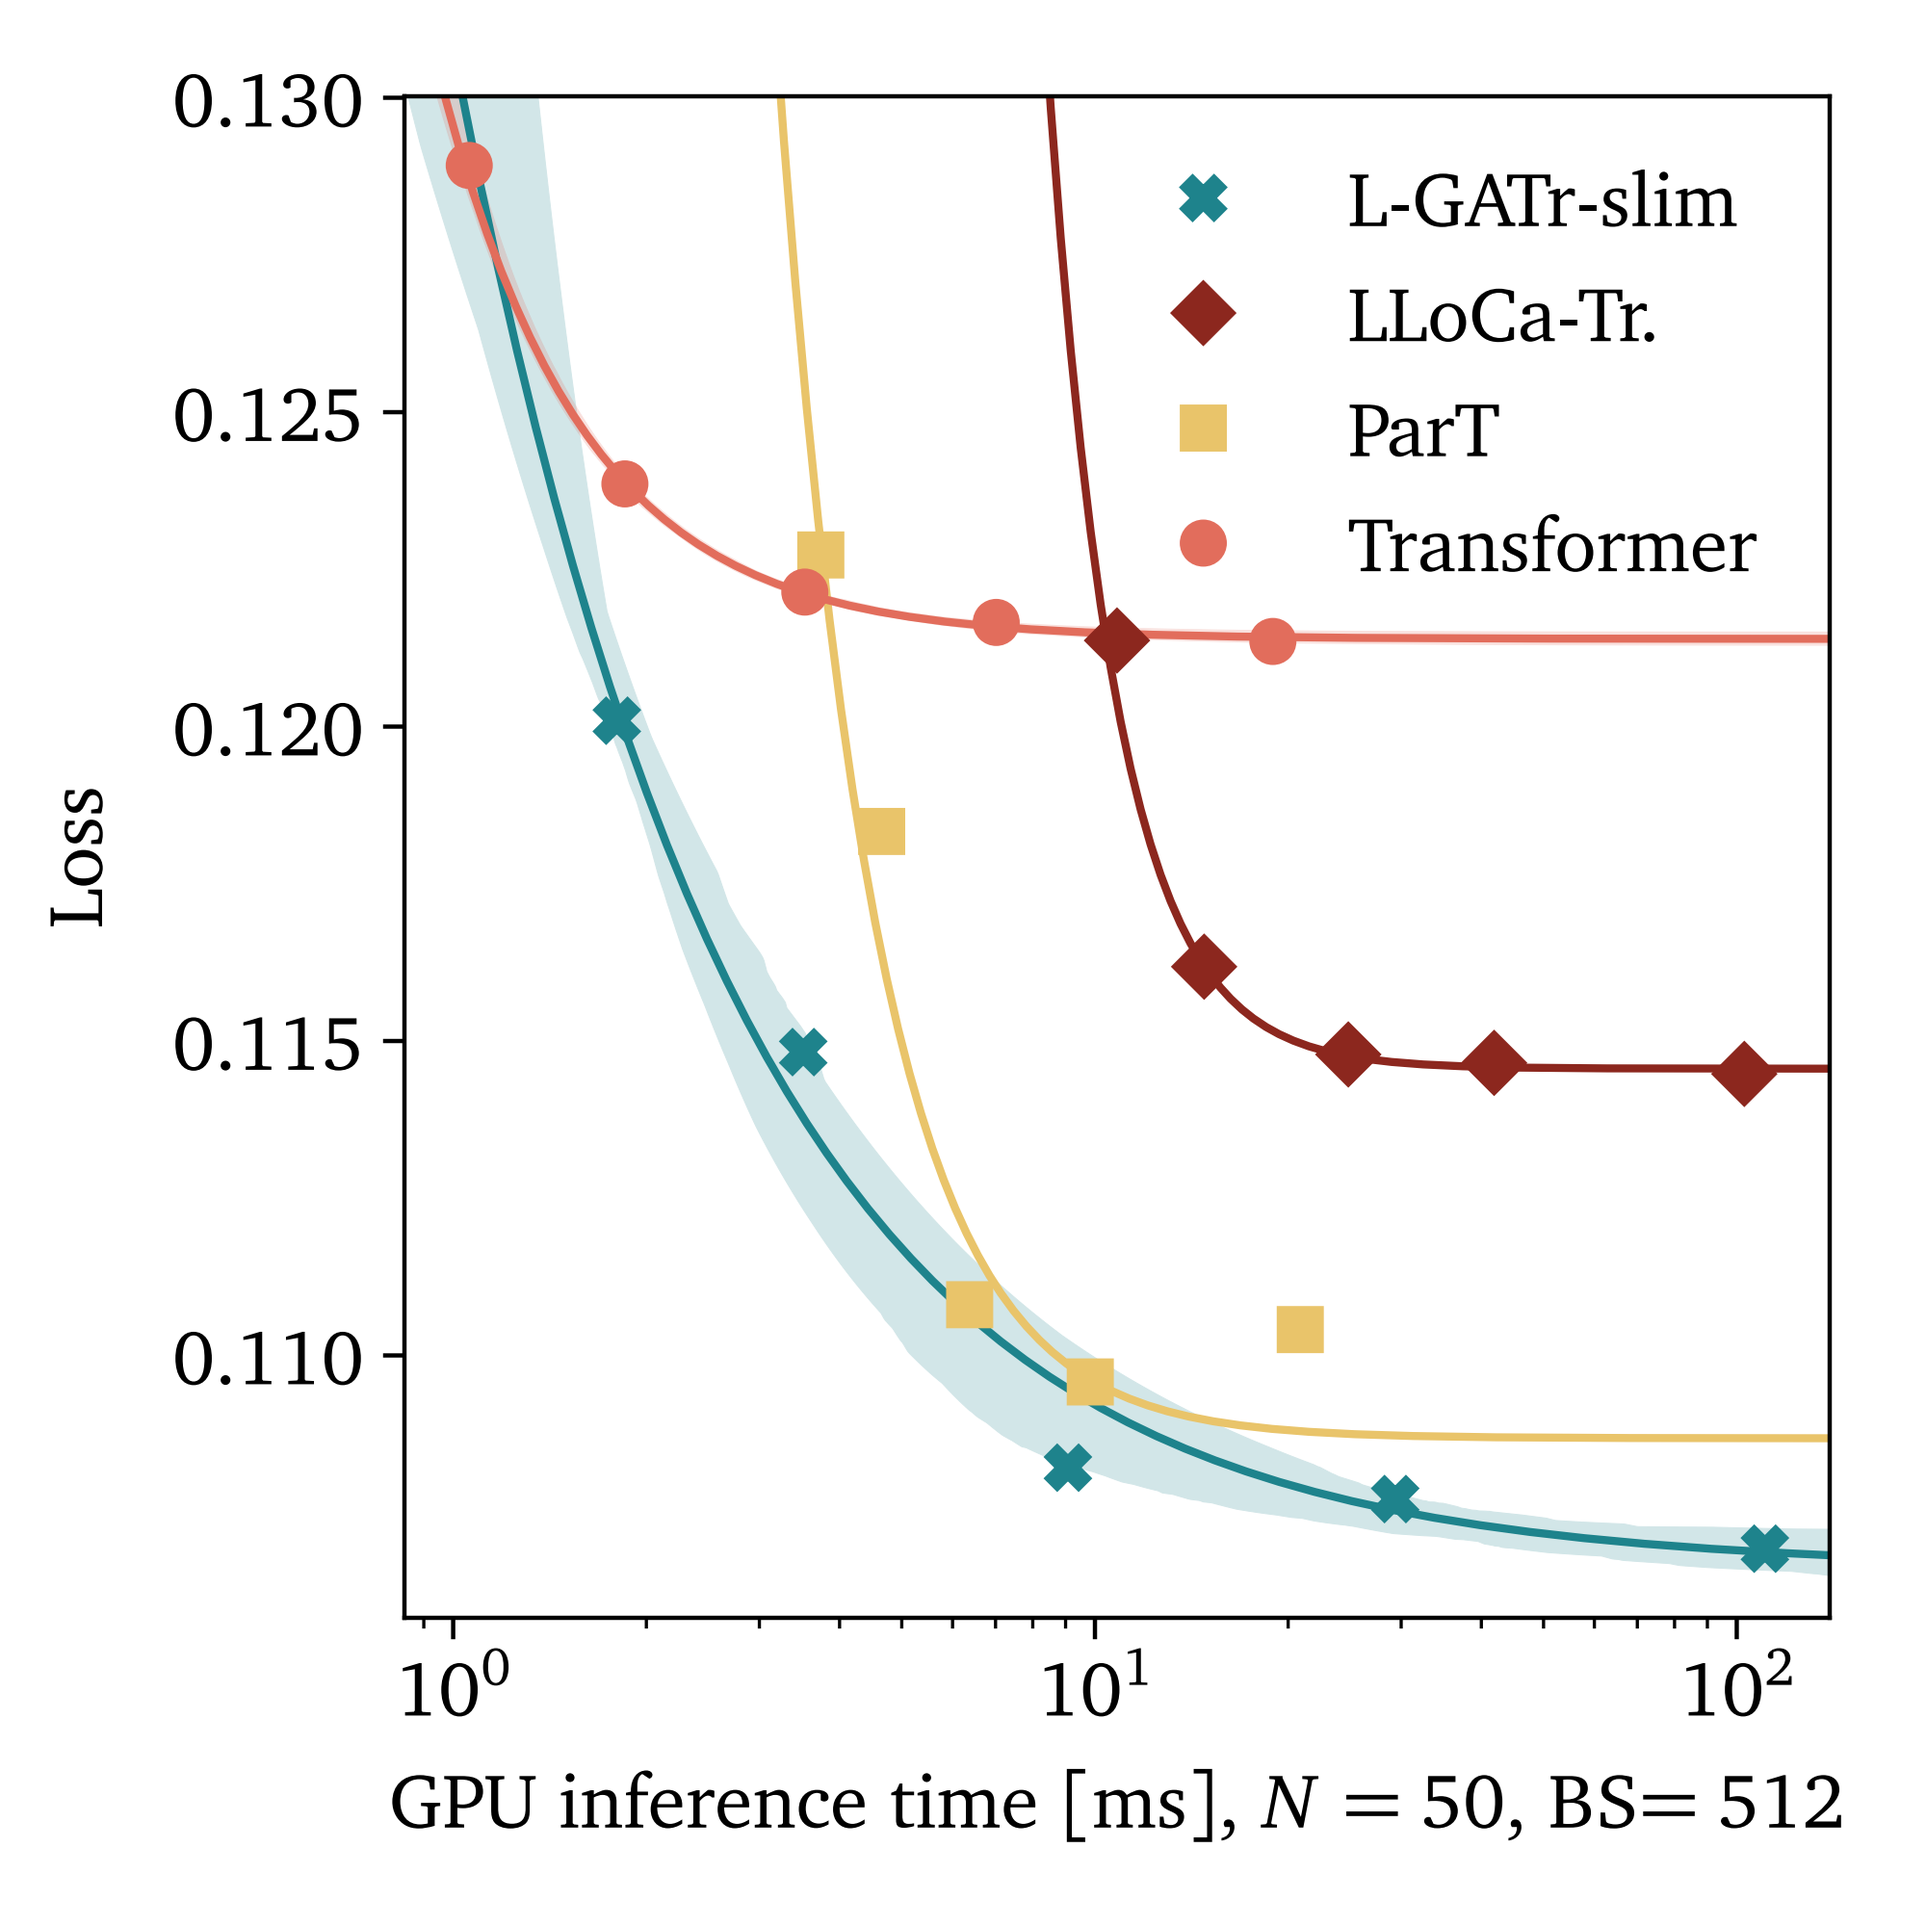
0.130
0.125
0.120
0.115
0.110
Loss
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
10^0
10^1
10^2
GPU inference time [ms], $N = 50$, BS$= 512$

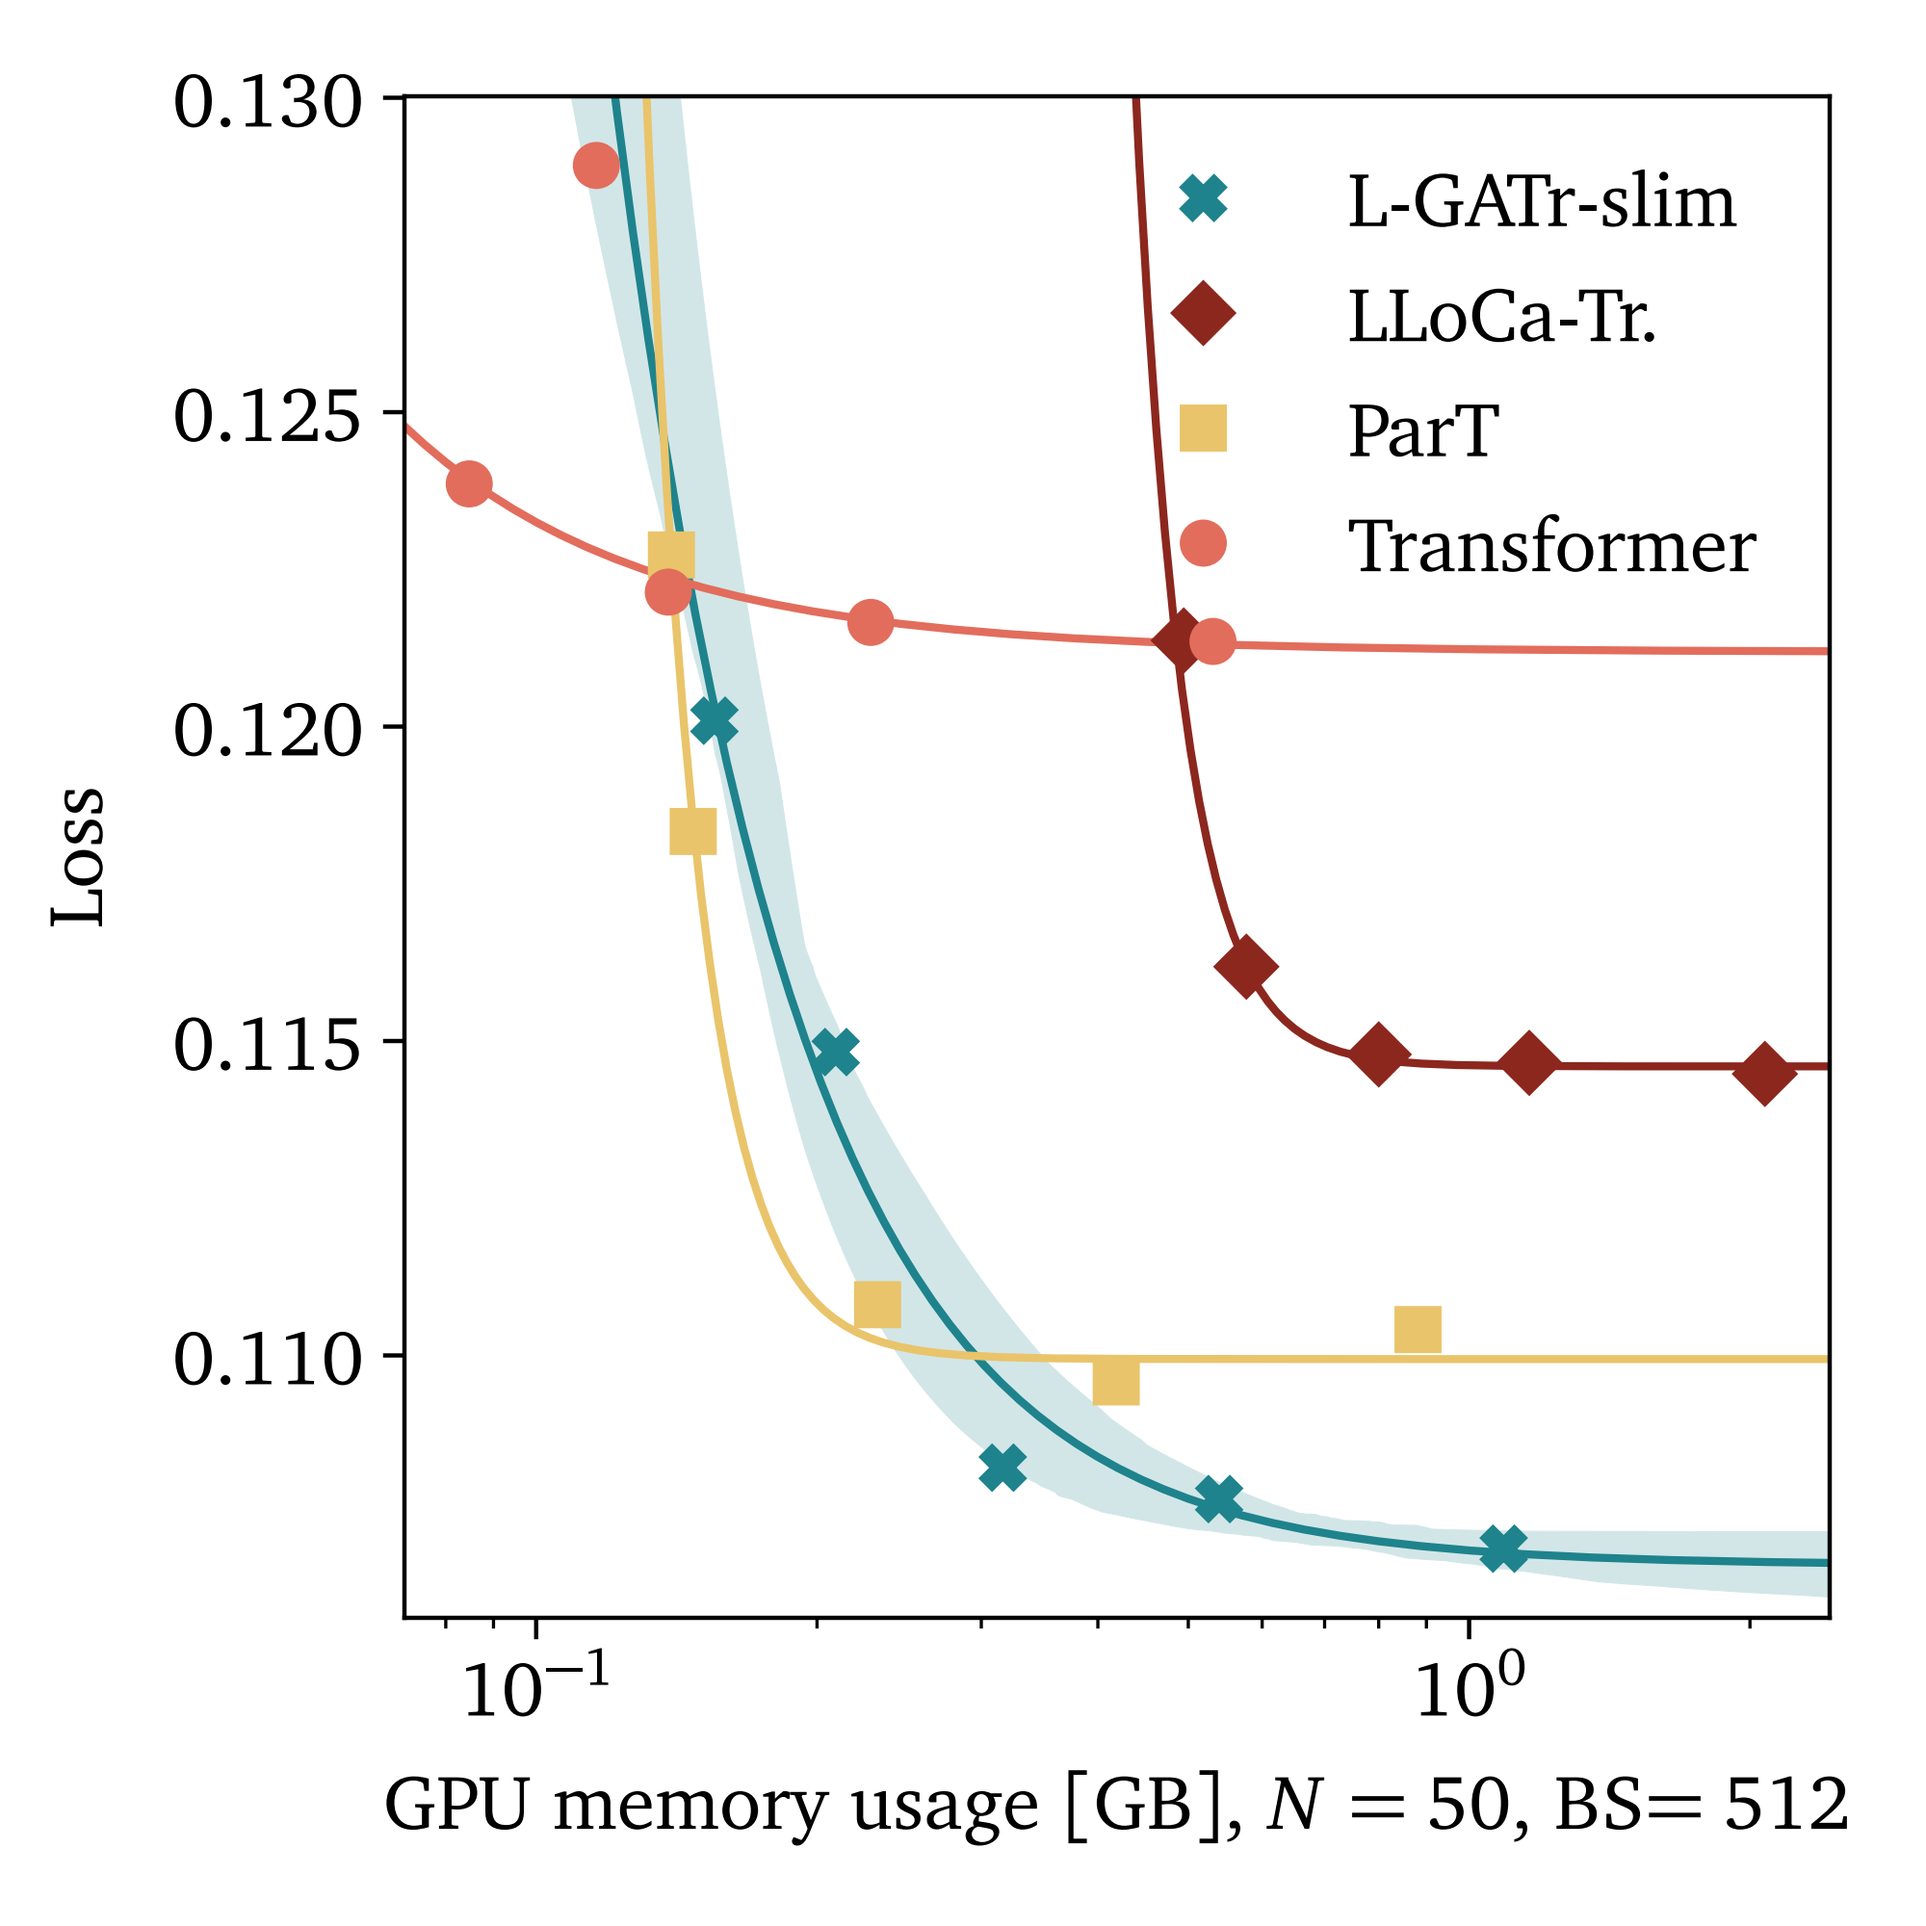
0.130
0.125
0.120
0.115
0.110
Loss
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
10^{-1}
10^{0}
GPU memory usage [GB], $N = 50$, BS$= 512$

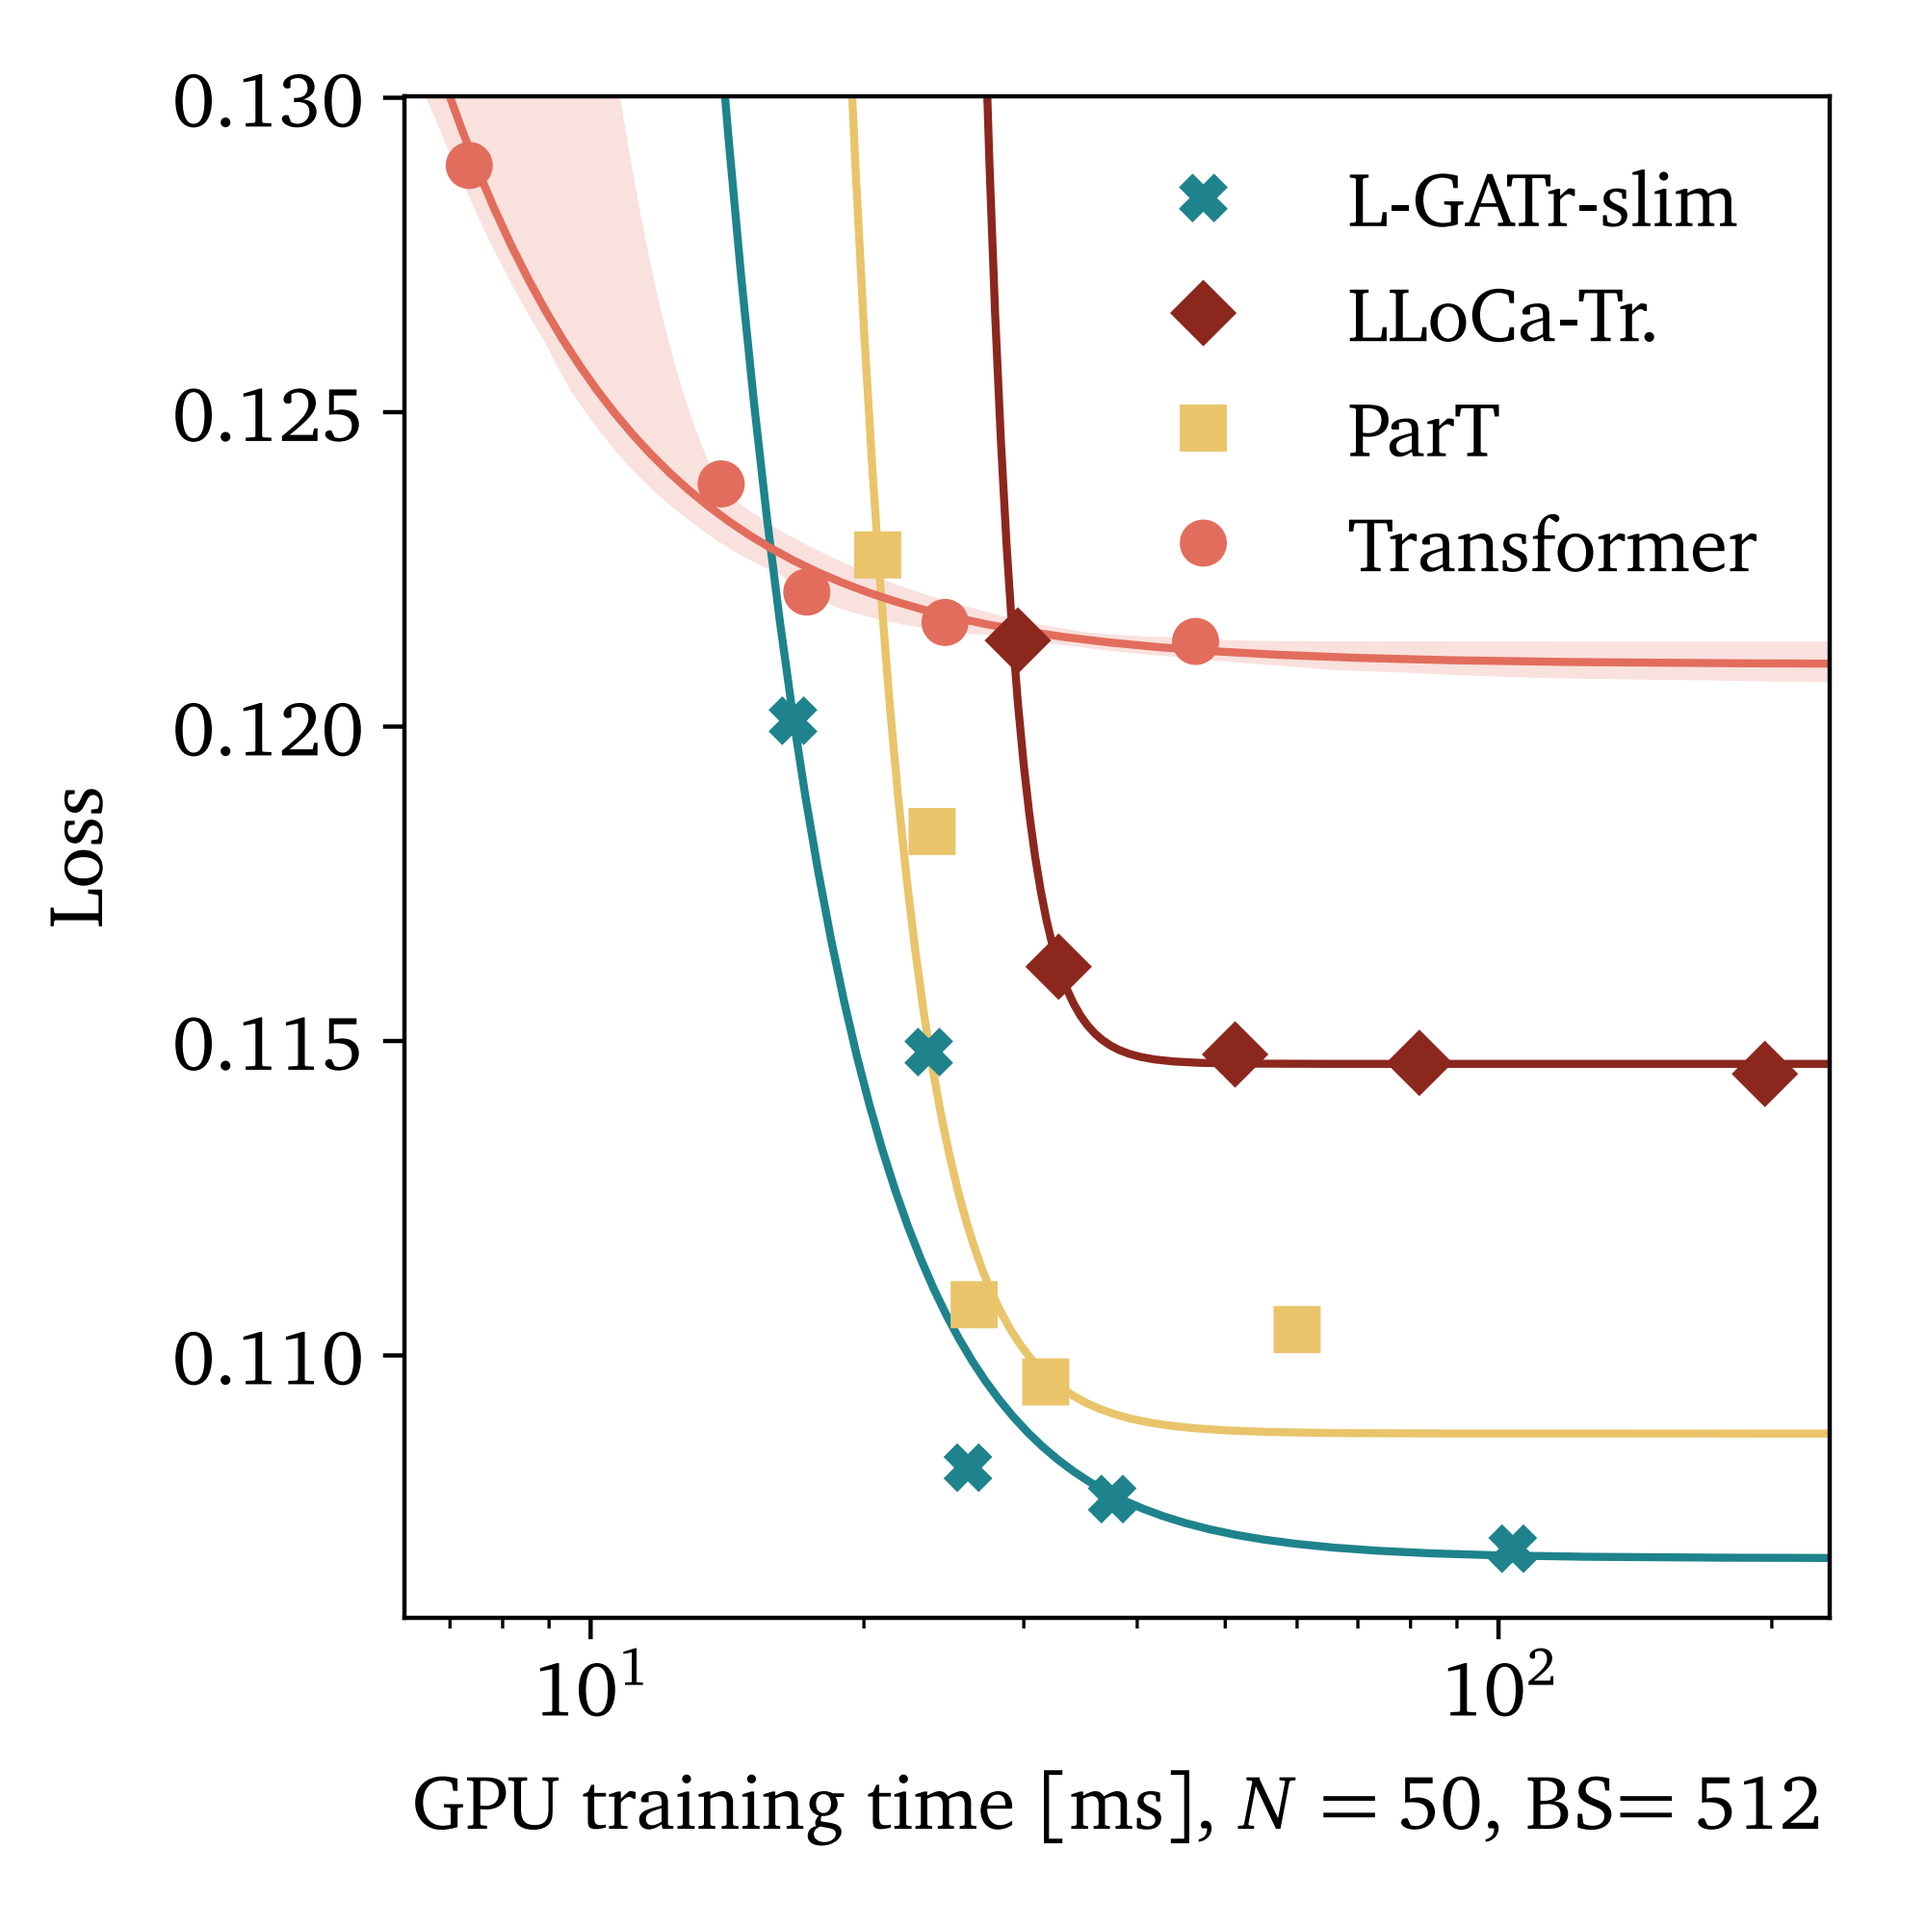
0.130
0.125
0.120
0.115
0.110
Loss
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
10^1
10^2
GPU training time [ms], $N = 50$, BS$= 512$

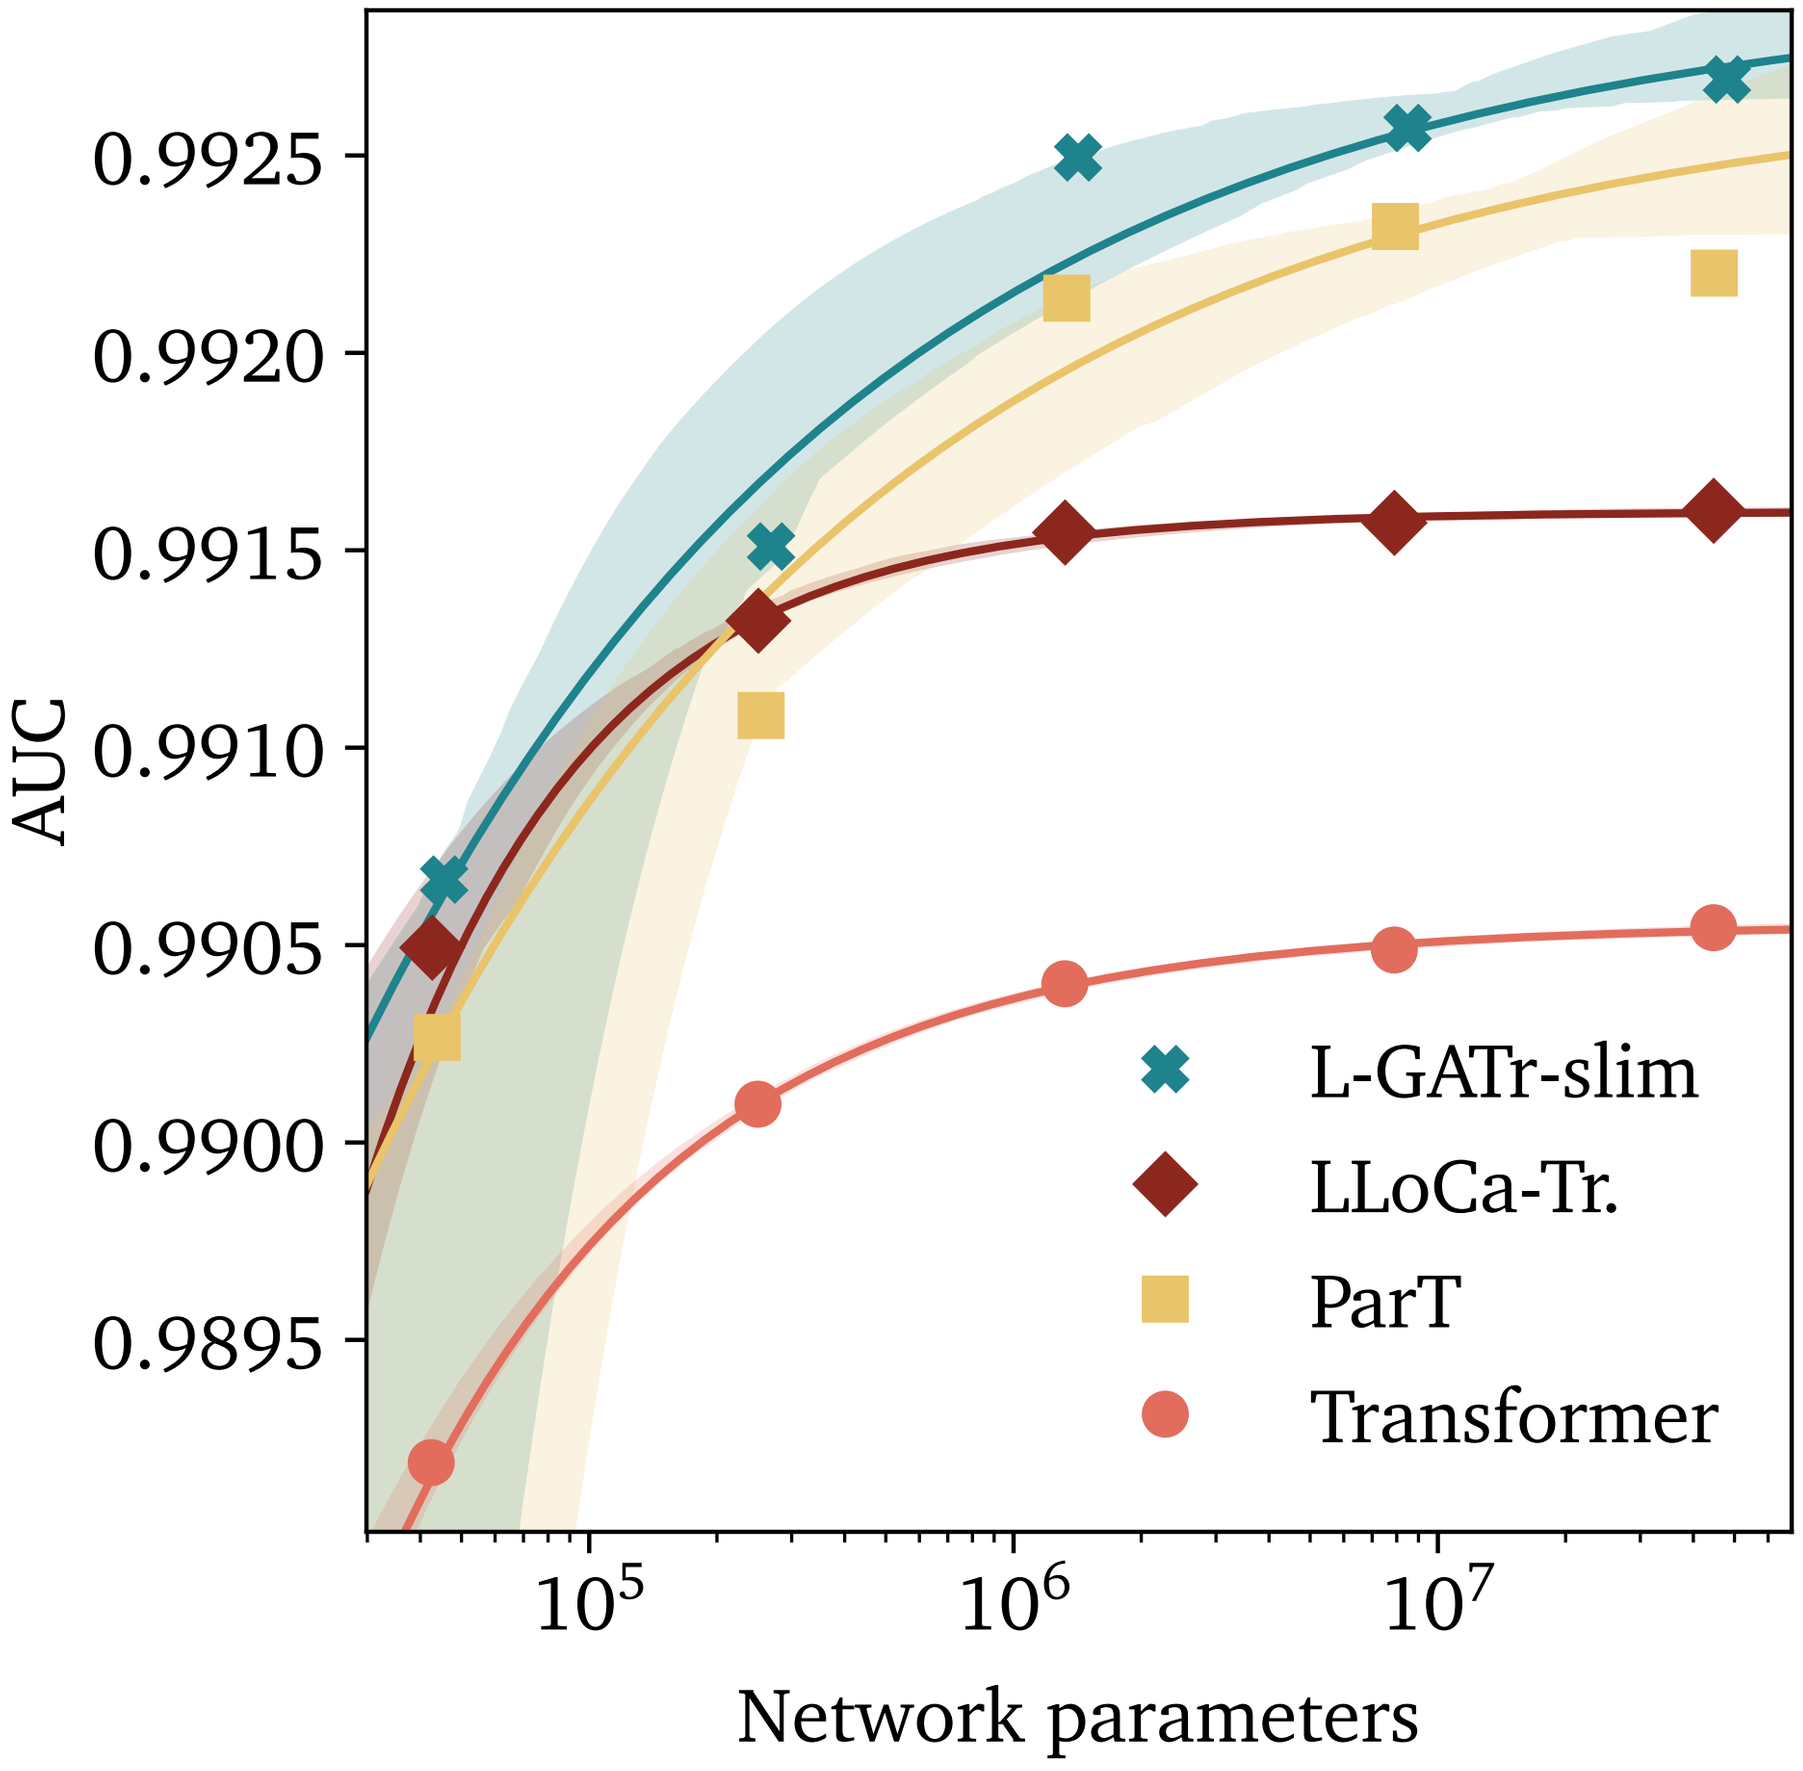
0.9925
0.9920
0.9915
0.9910
0.9905
0.9900
0.9895
AUC
10^5
10^6
10^7
Network parameters
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

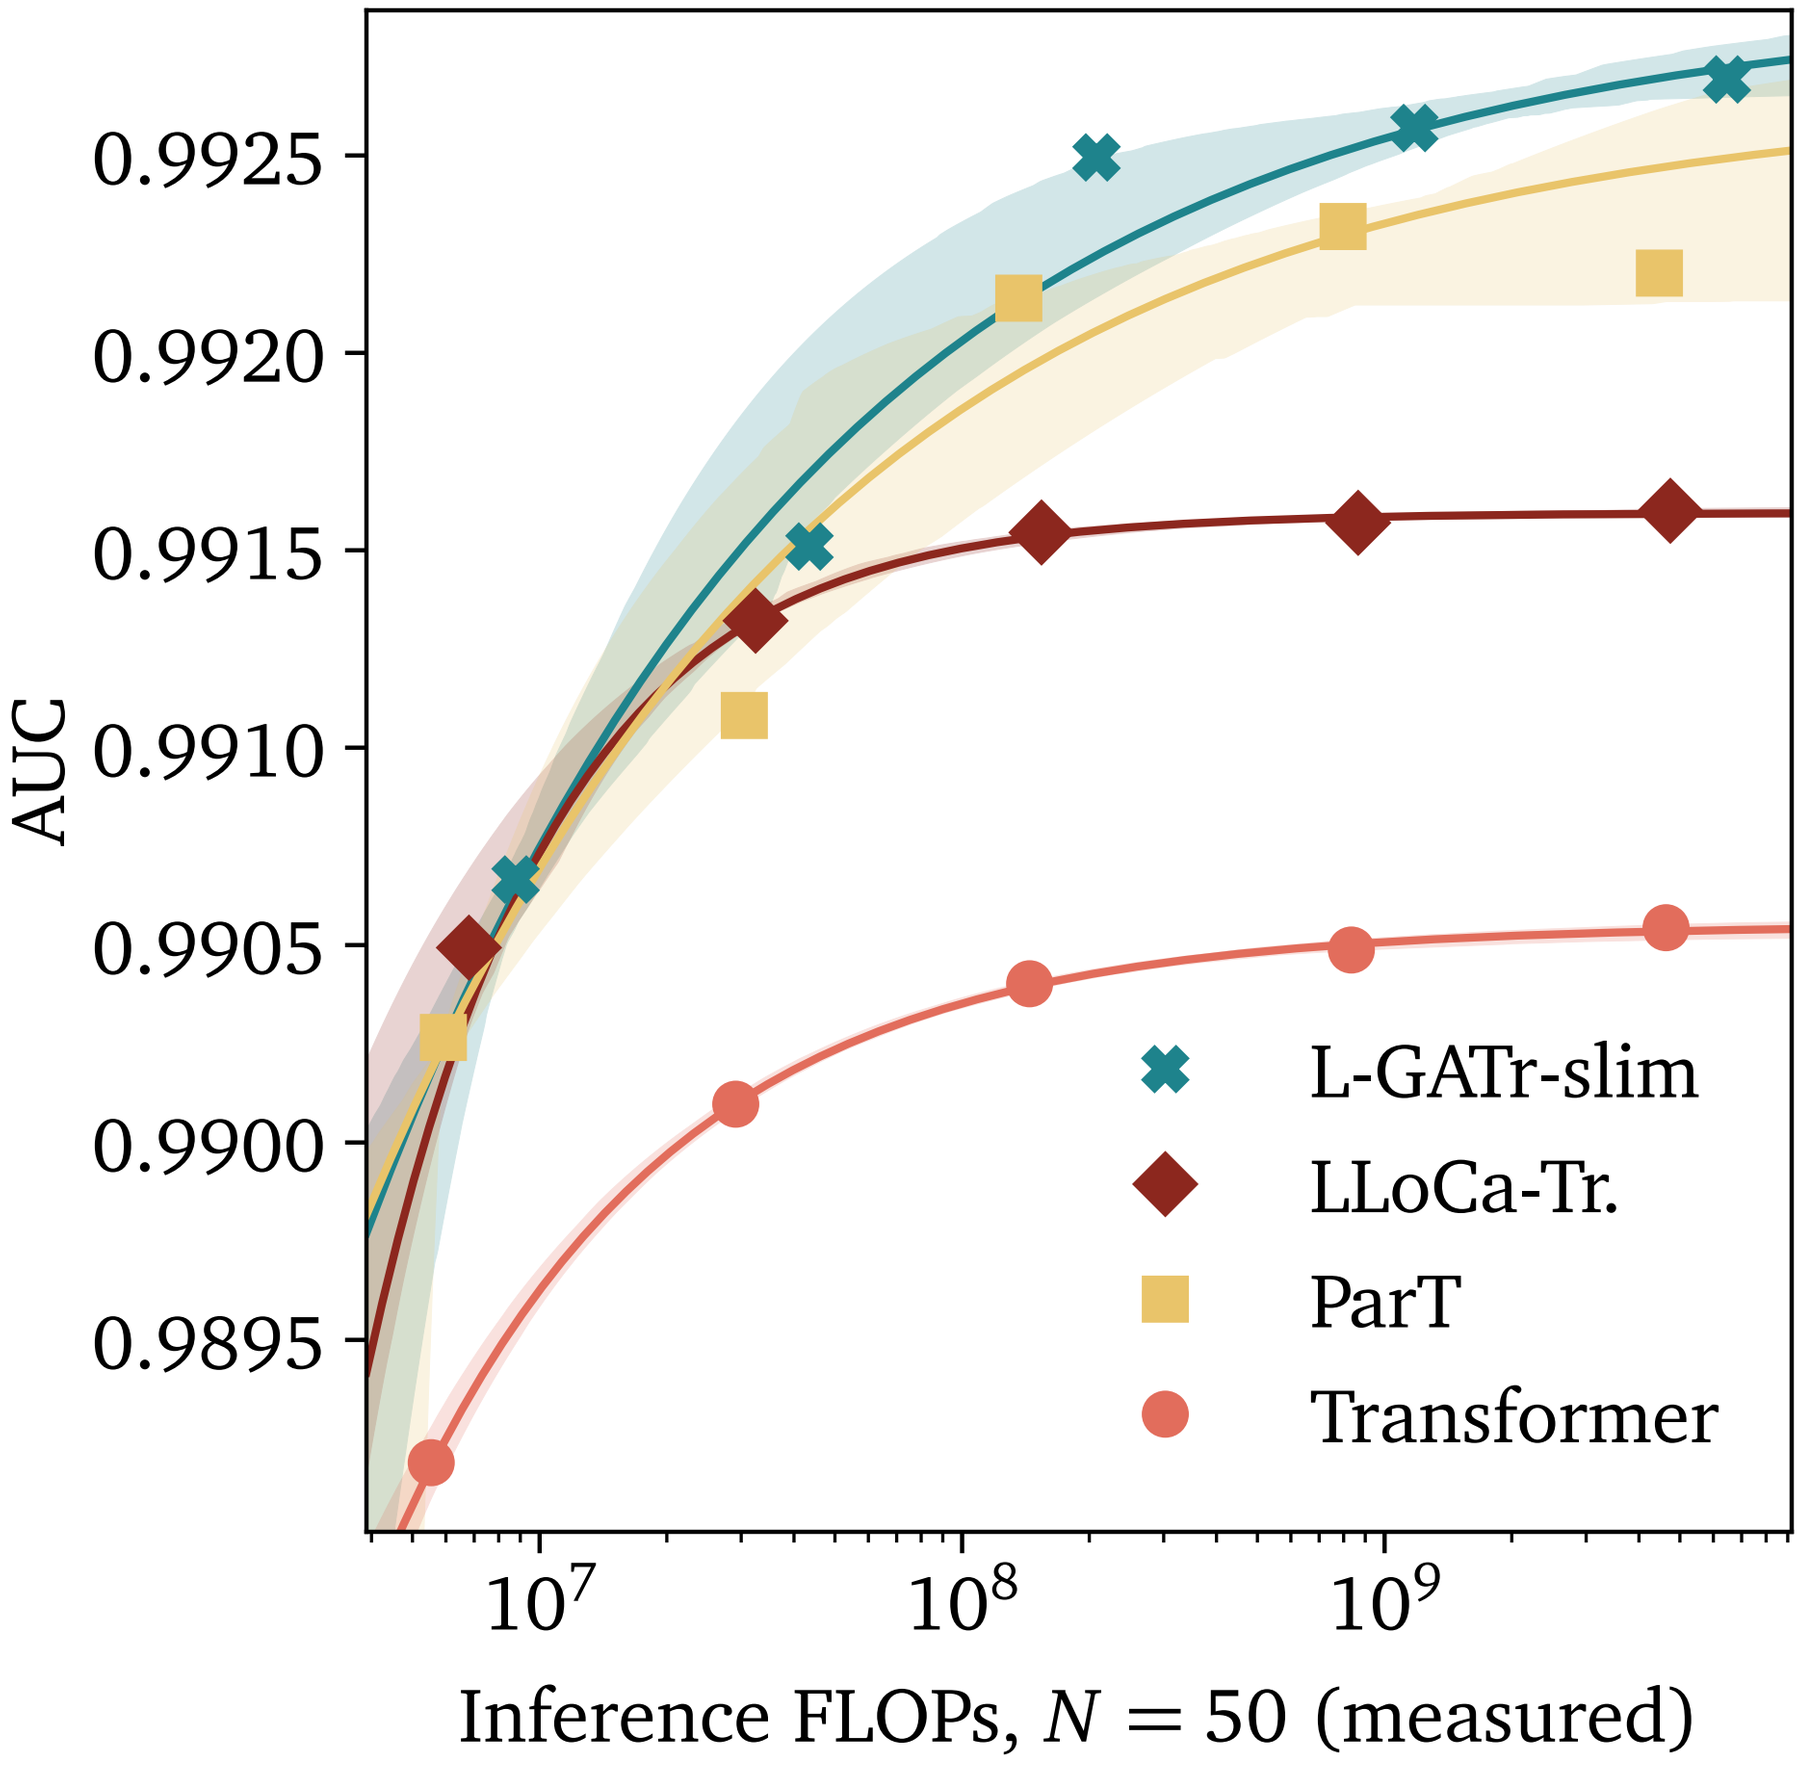
0.9925
0.9920
0.9915
0.9910
0.9905
0.9900
0.9895
AUC
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (measured)
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

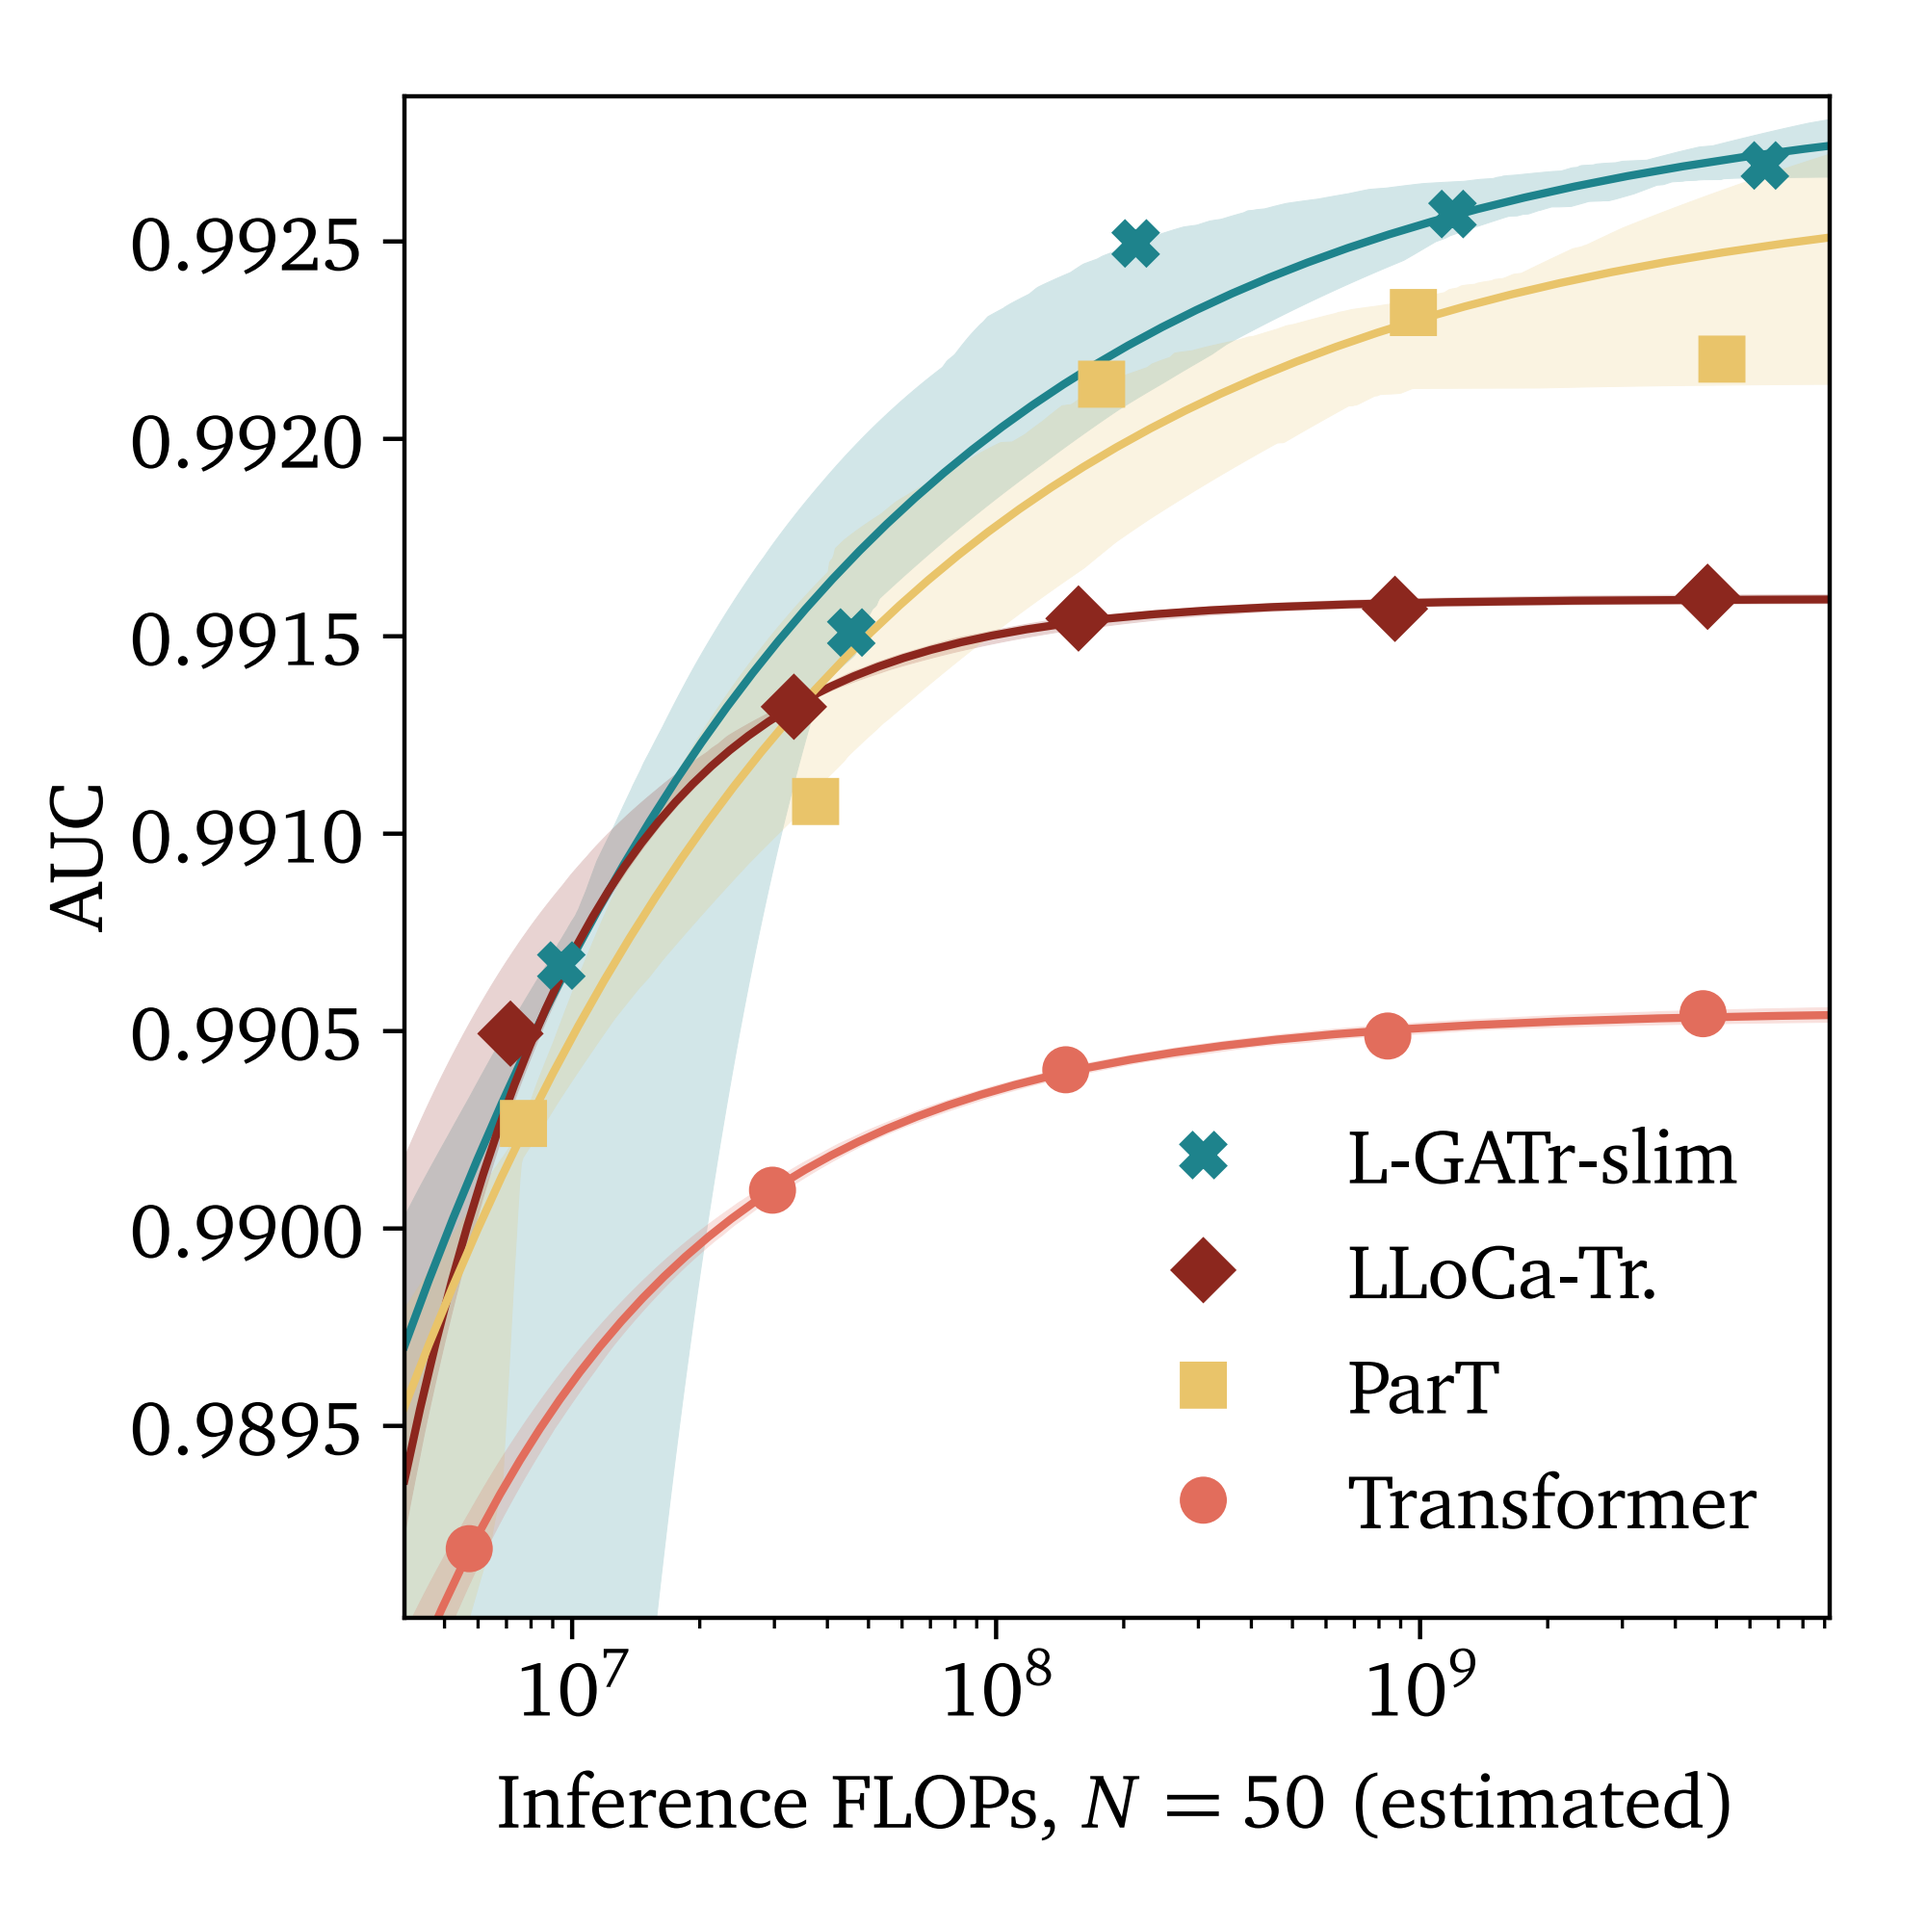
0.9925
0.9920
0.9915
0.9910
0.9905
0.9900
0.9895
AUC
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (estimated)
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

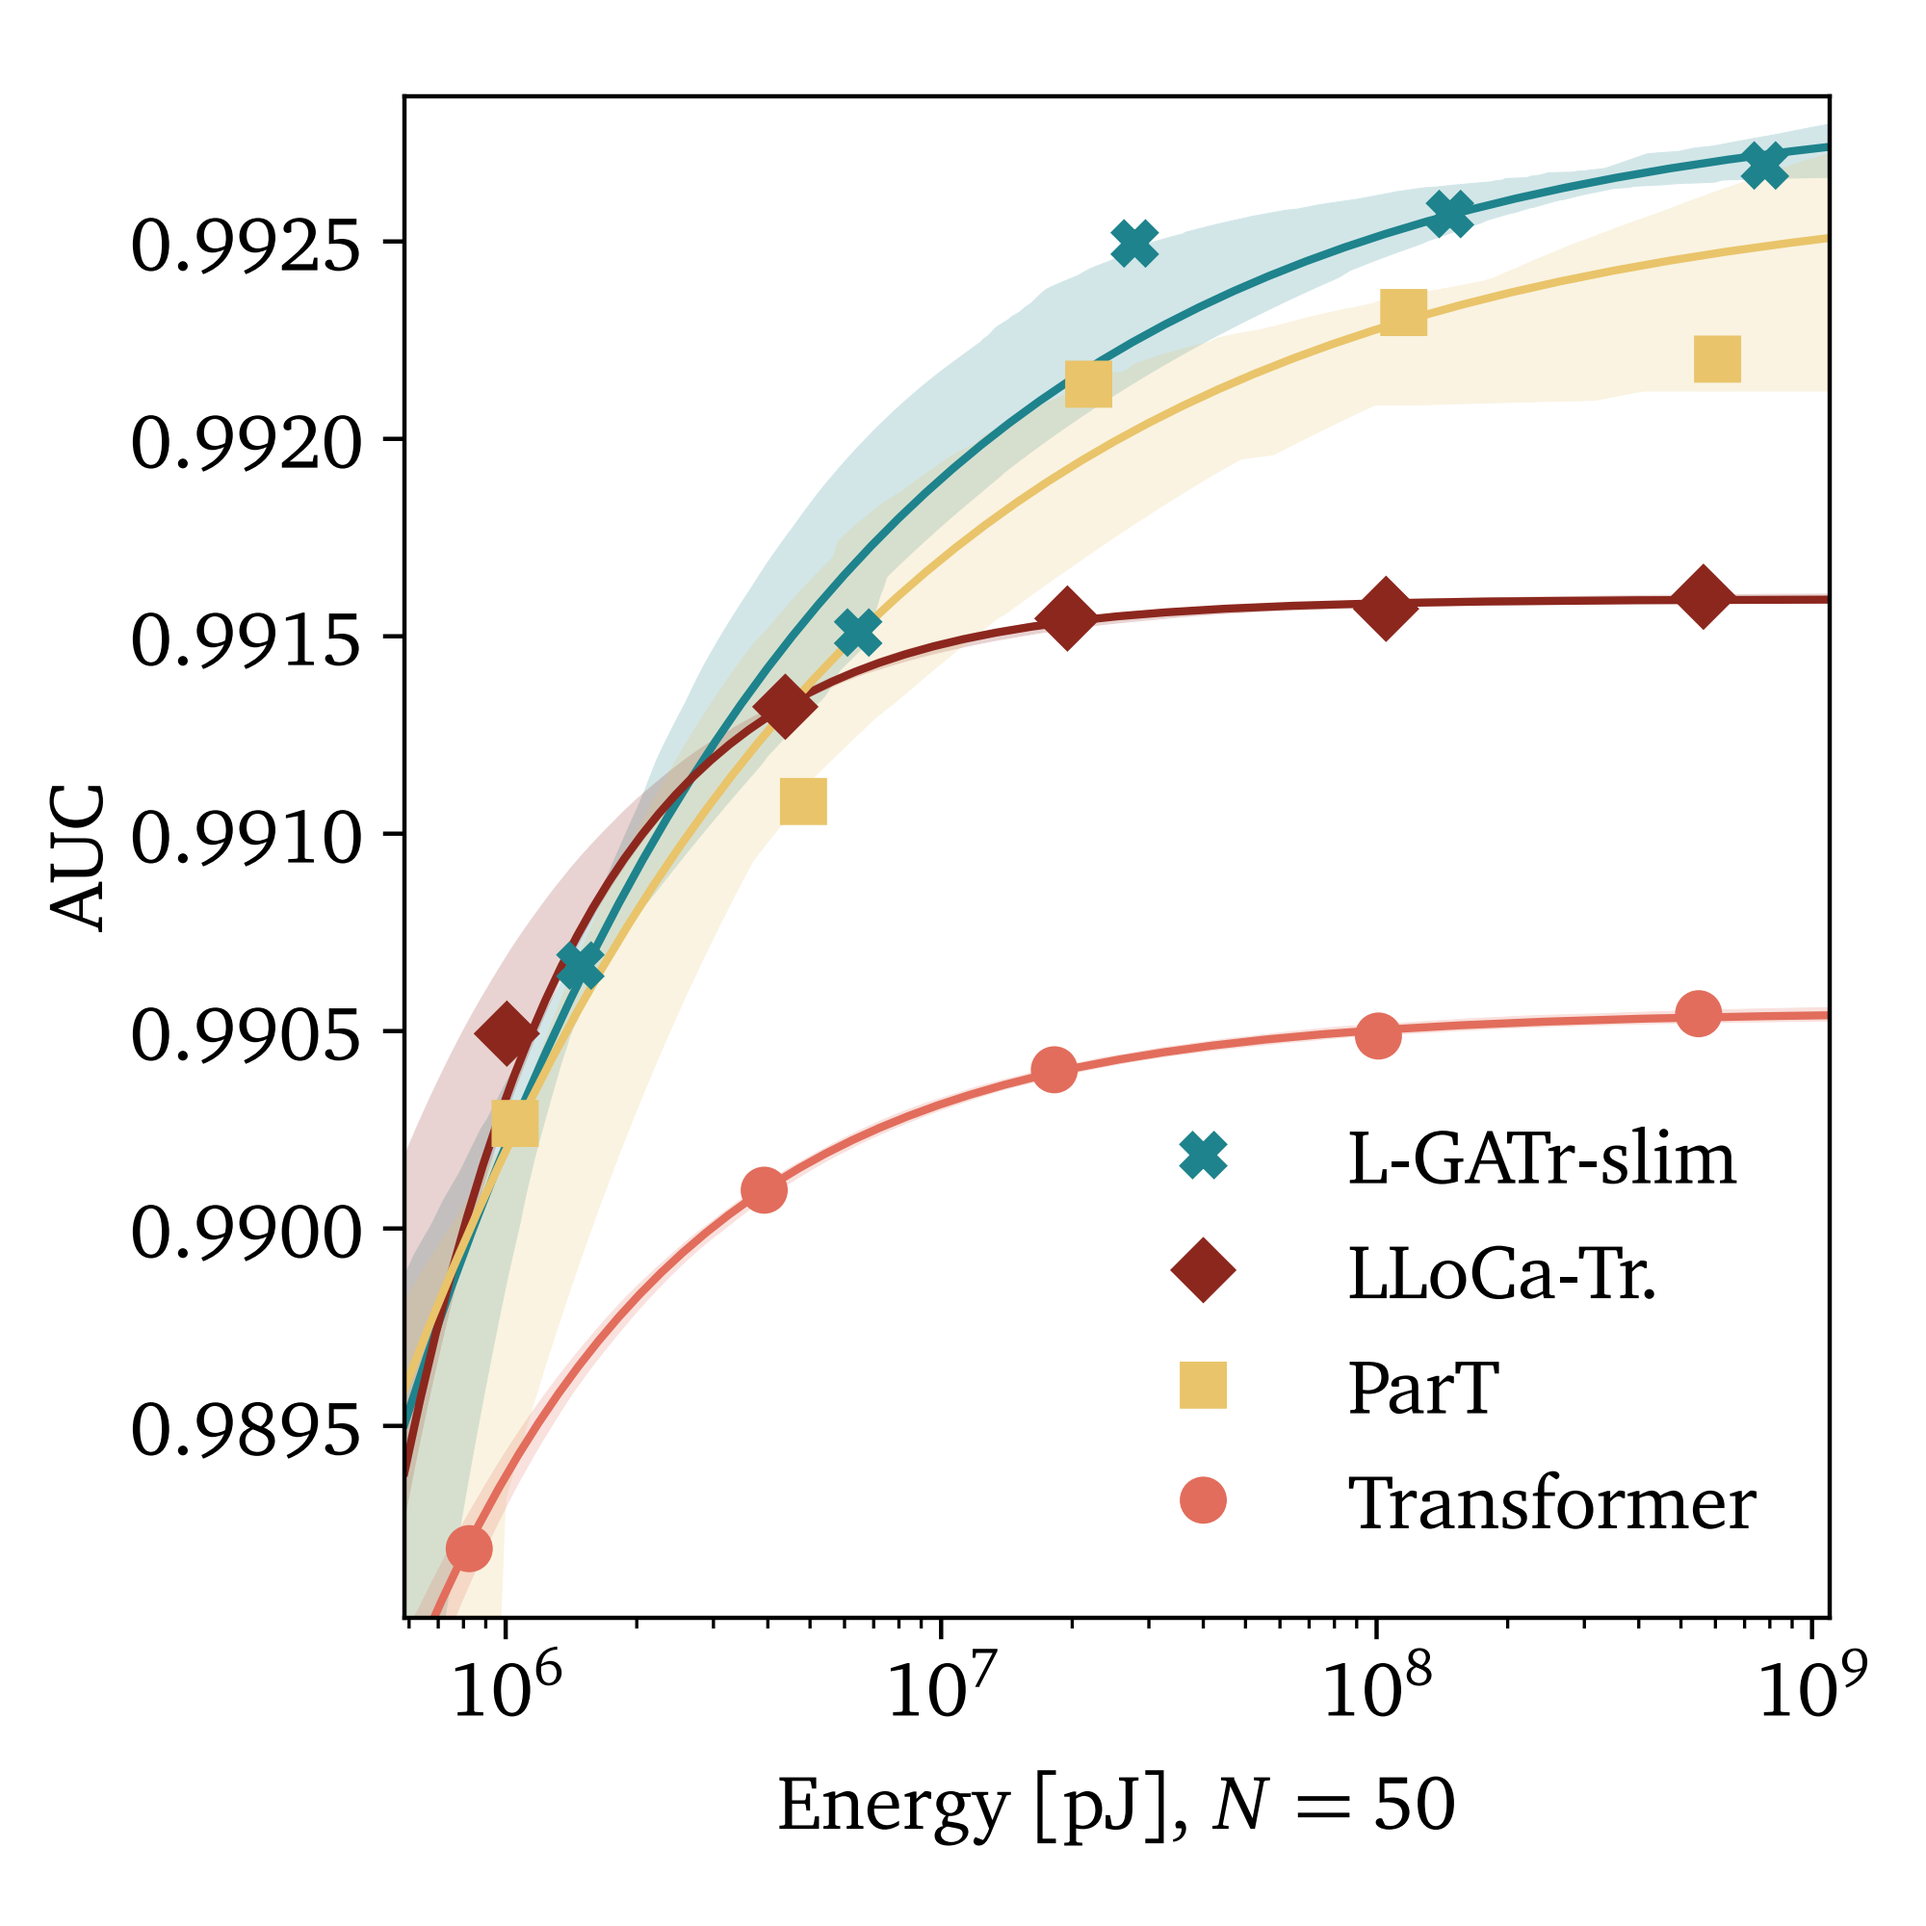
0.9925
0.9920
0.9915
0.9910
0.9905
0.9900
0.9895
AUC
10^6
10^7
10^8
10^9
Energy [pJ], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

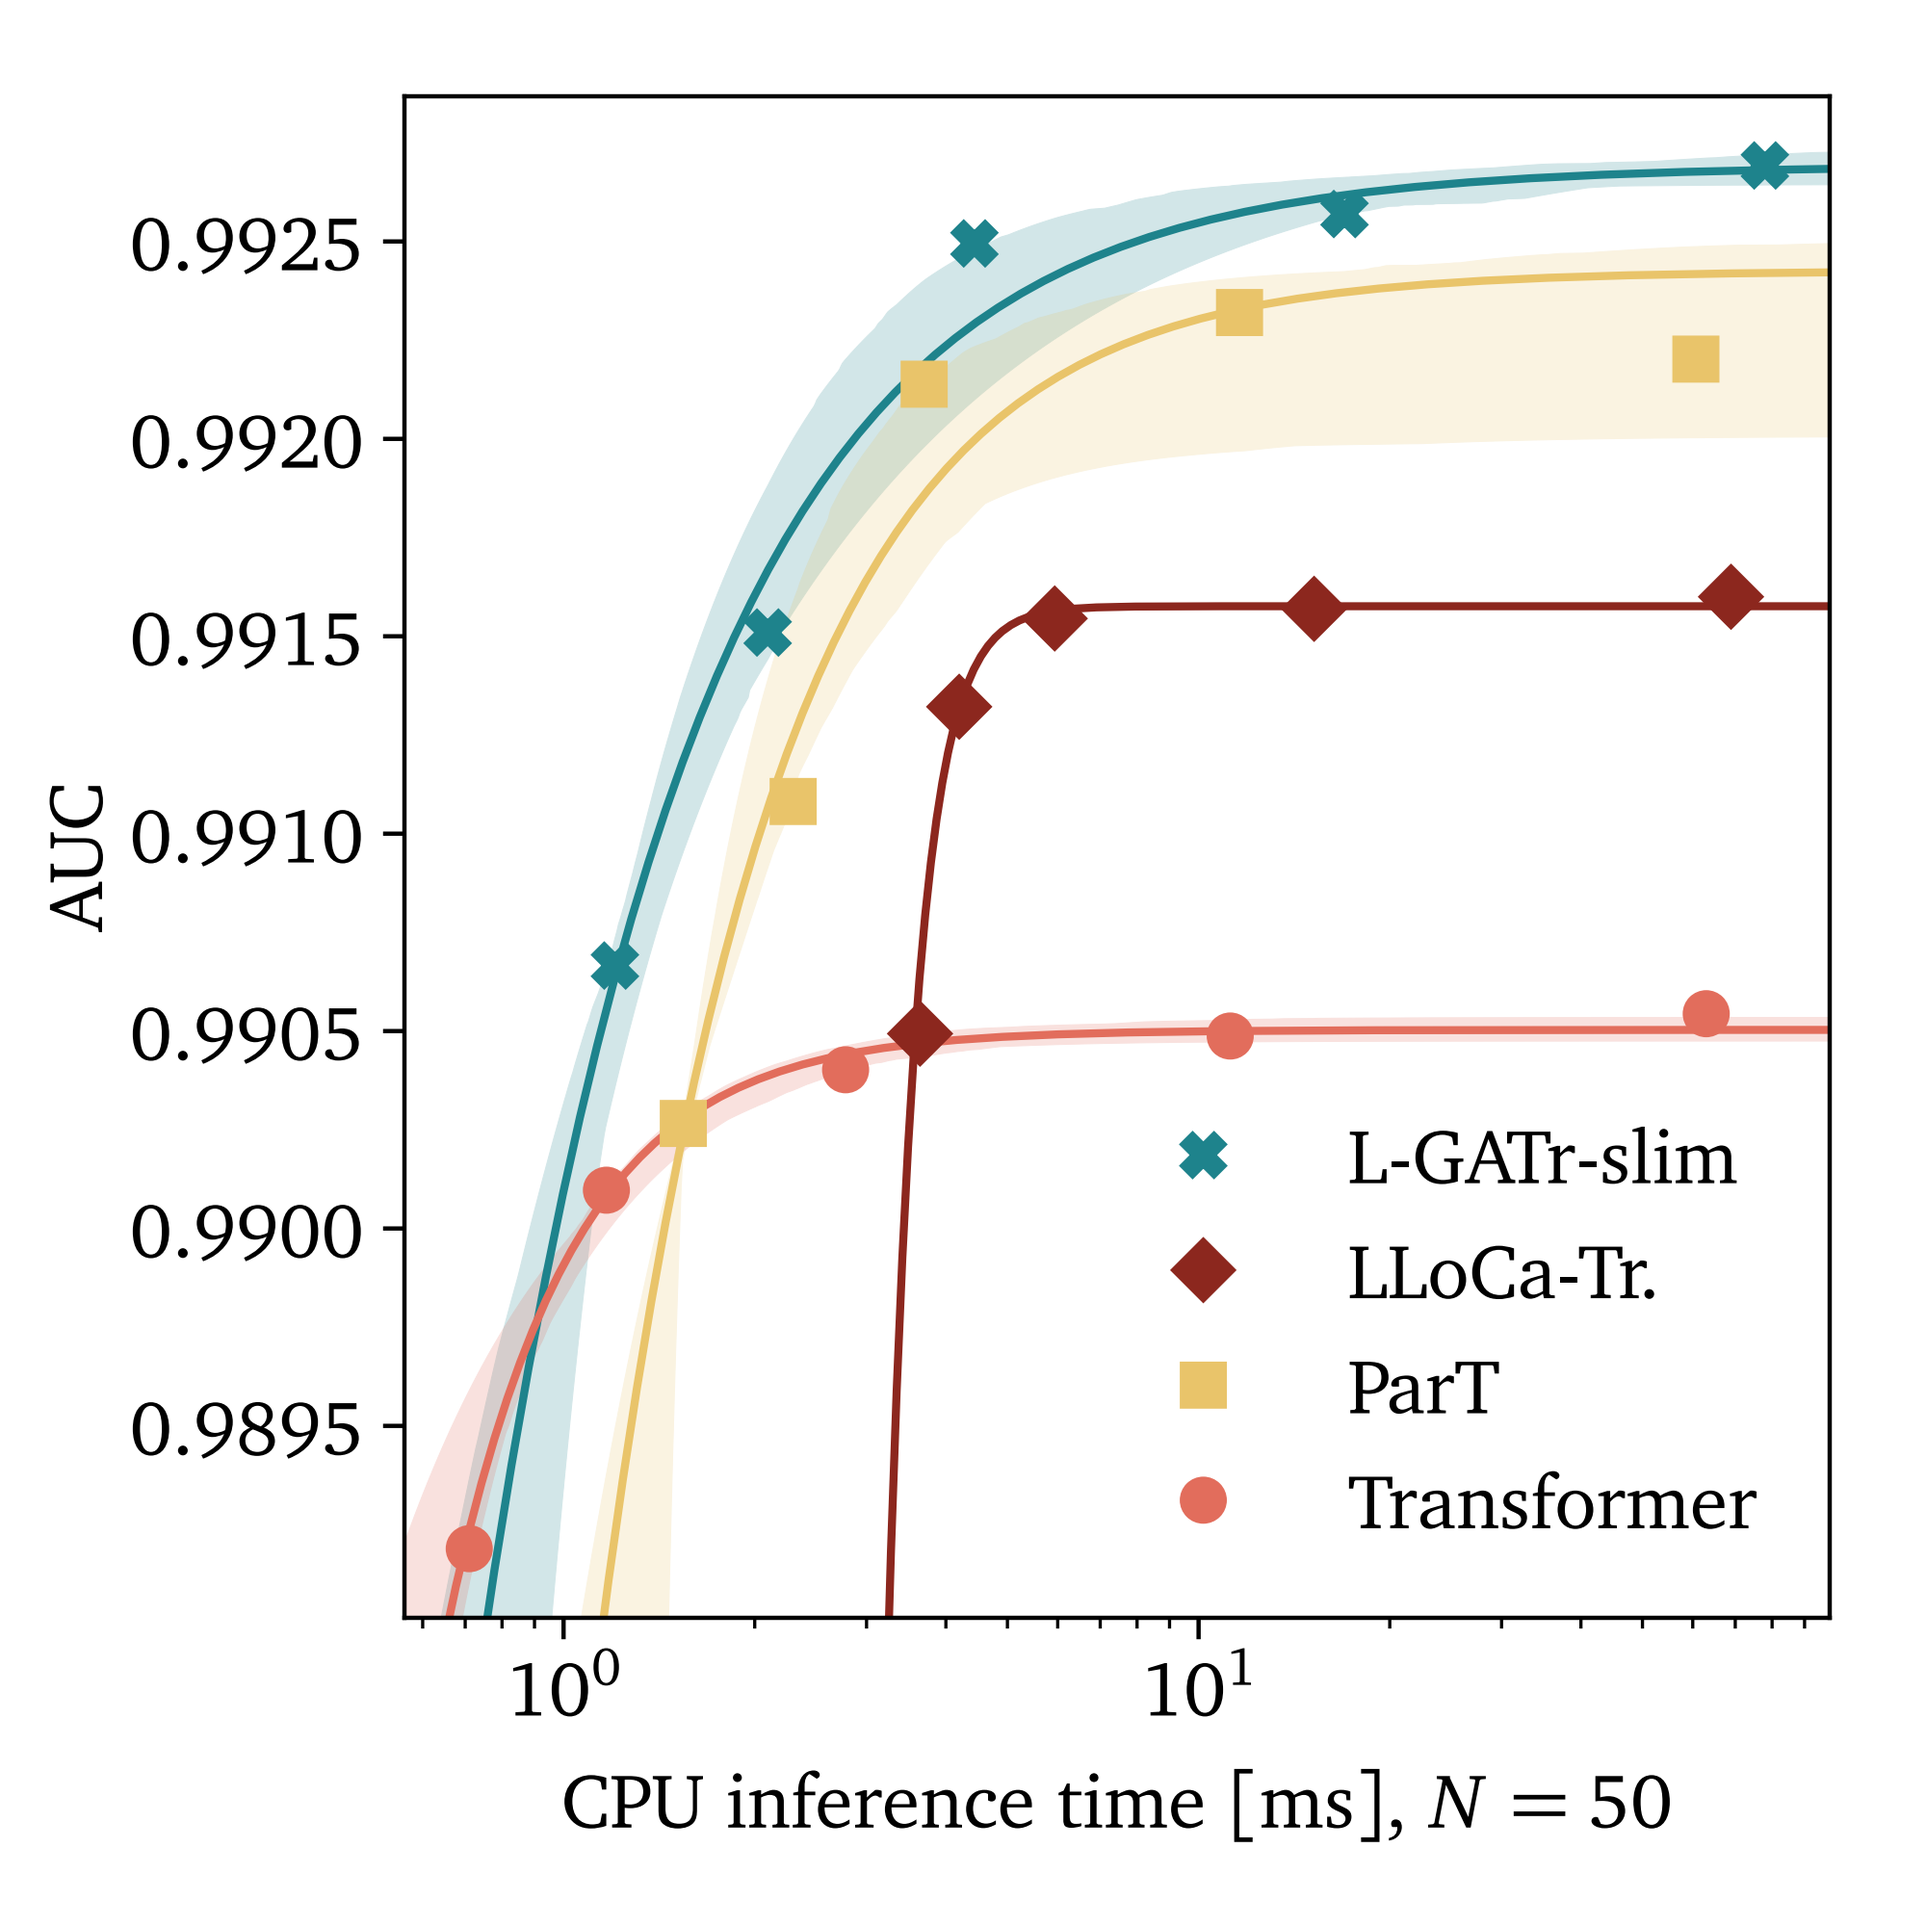
0.9925
0.9920
0.9915
0.9910
0.9905
0.9900
0.9895
AUC
10^0
10^1
CPU inference time [ms], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

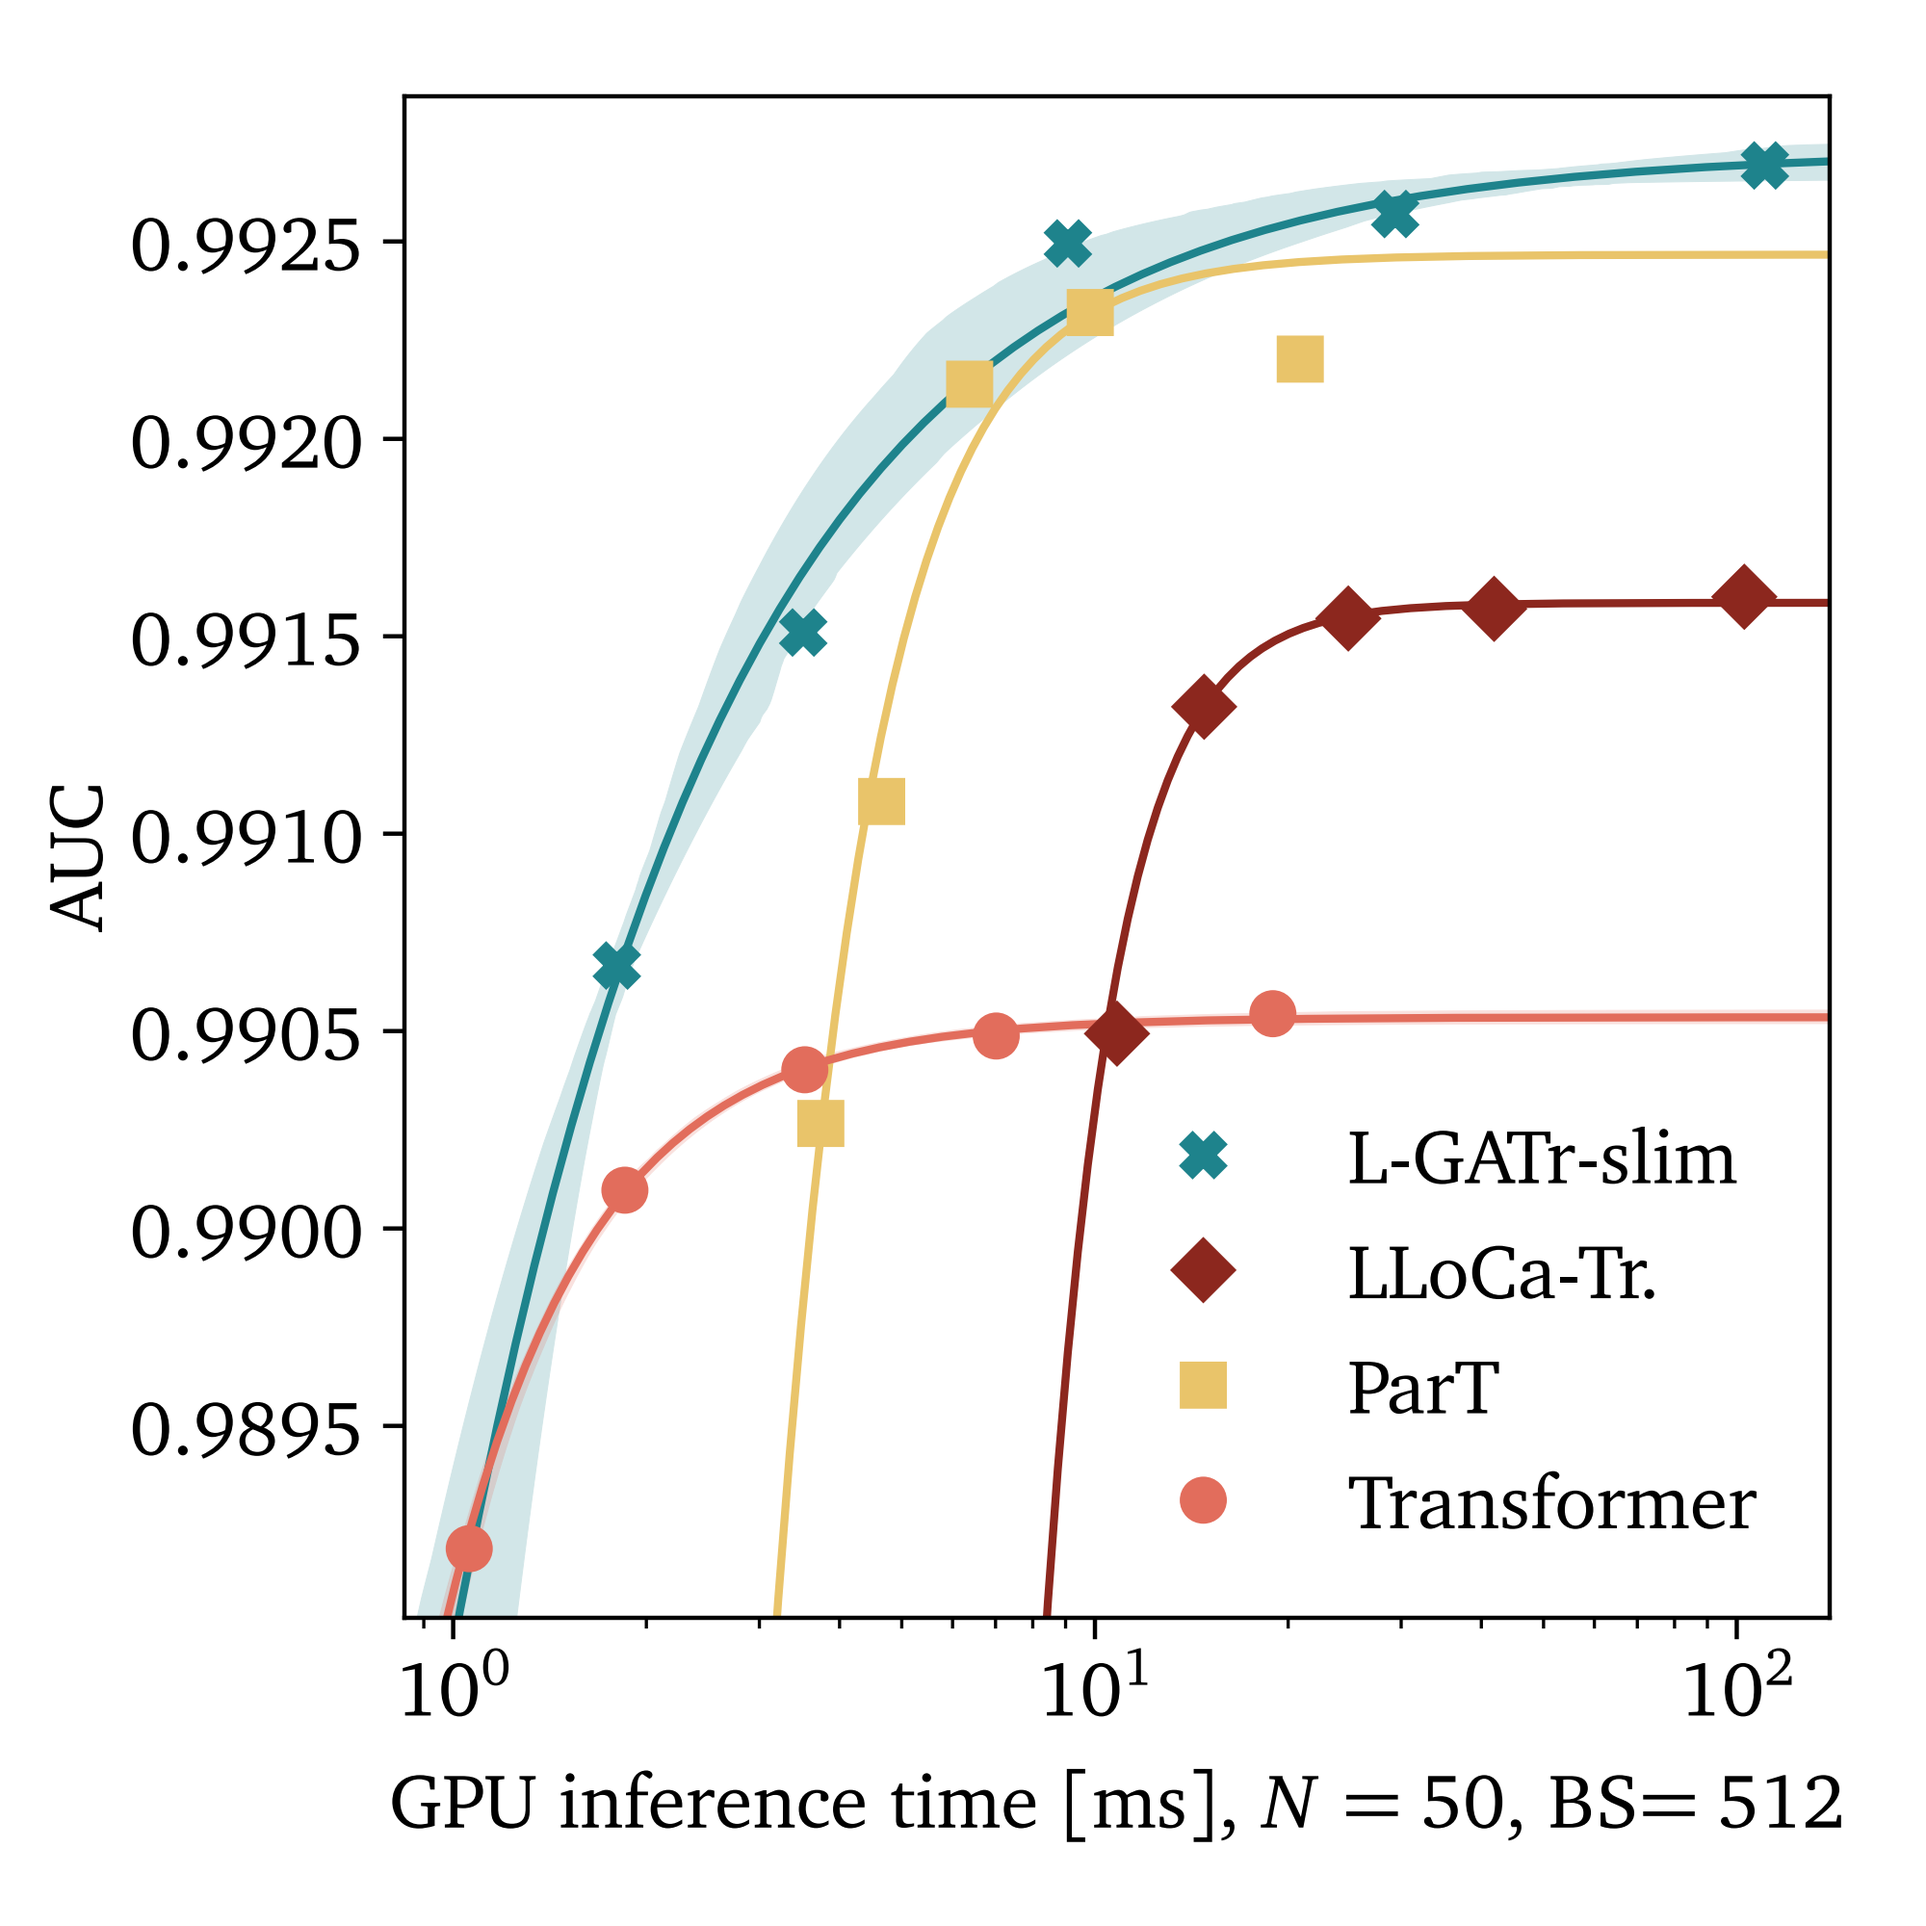
0.9925
0.9920
0.9915
0.9910
0.9905
0.9900
0.9895
AUC
10^0
10^1
10^2
GPU inference time [ms], $N = 50$, BS= 512
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

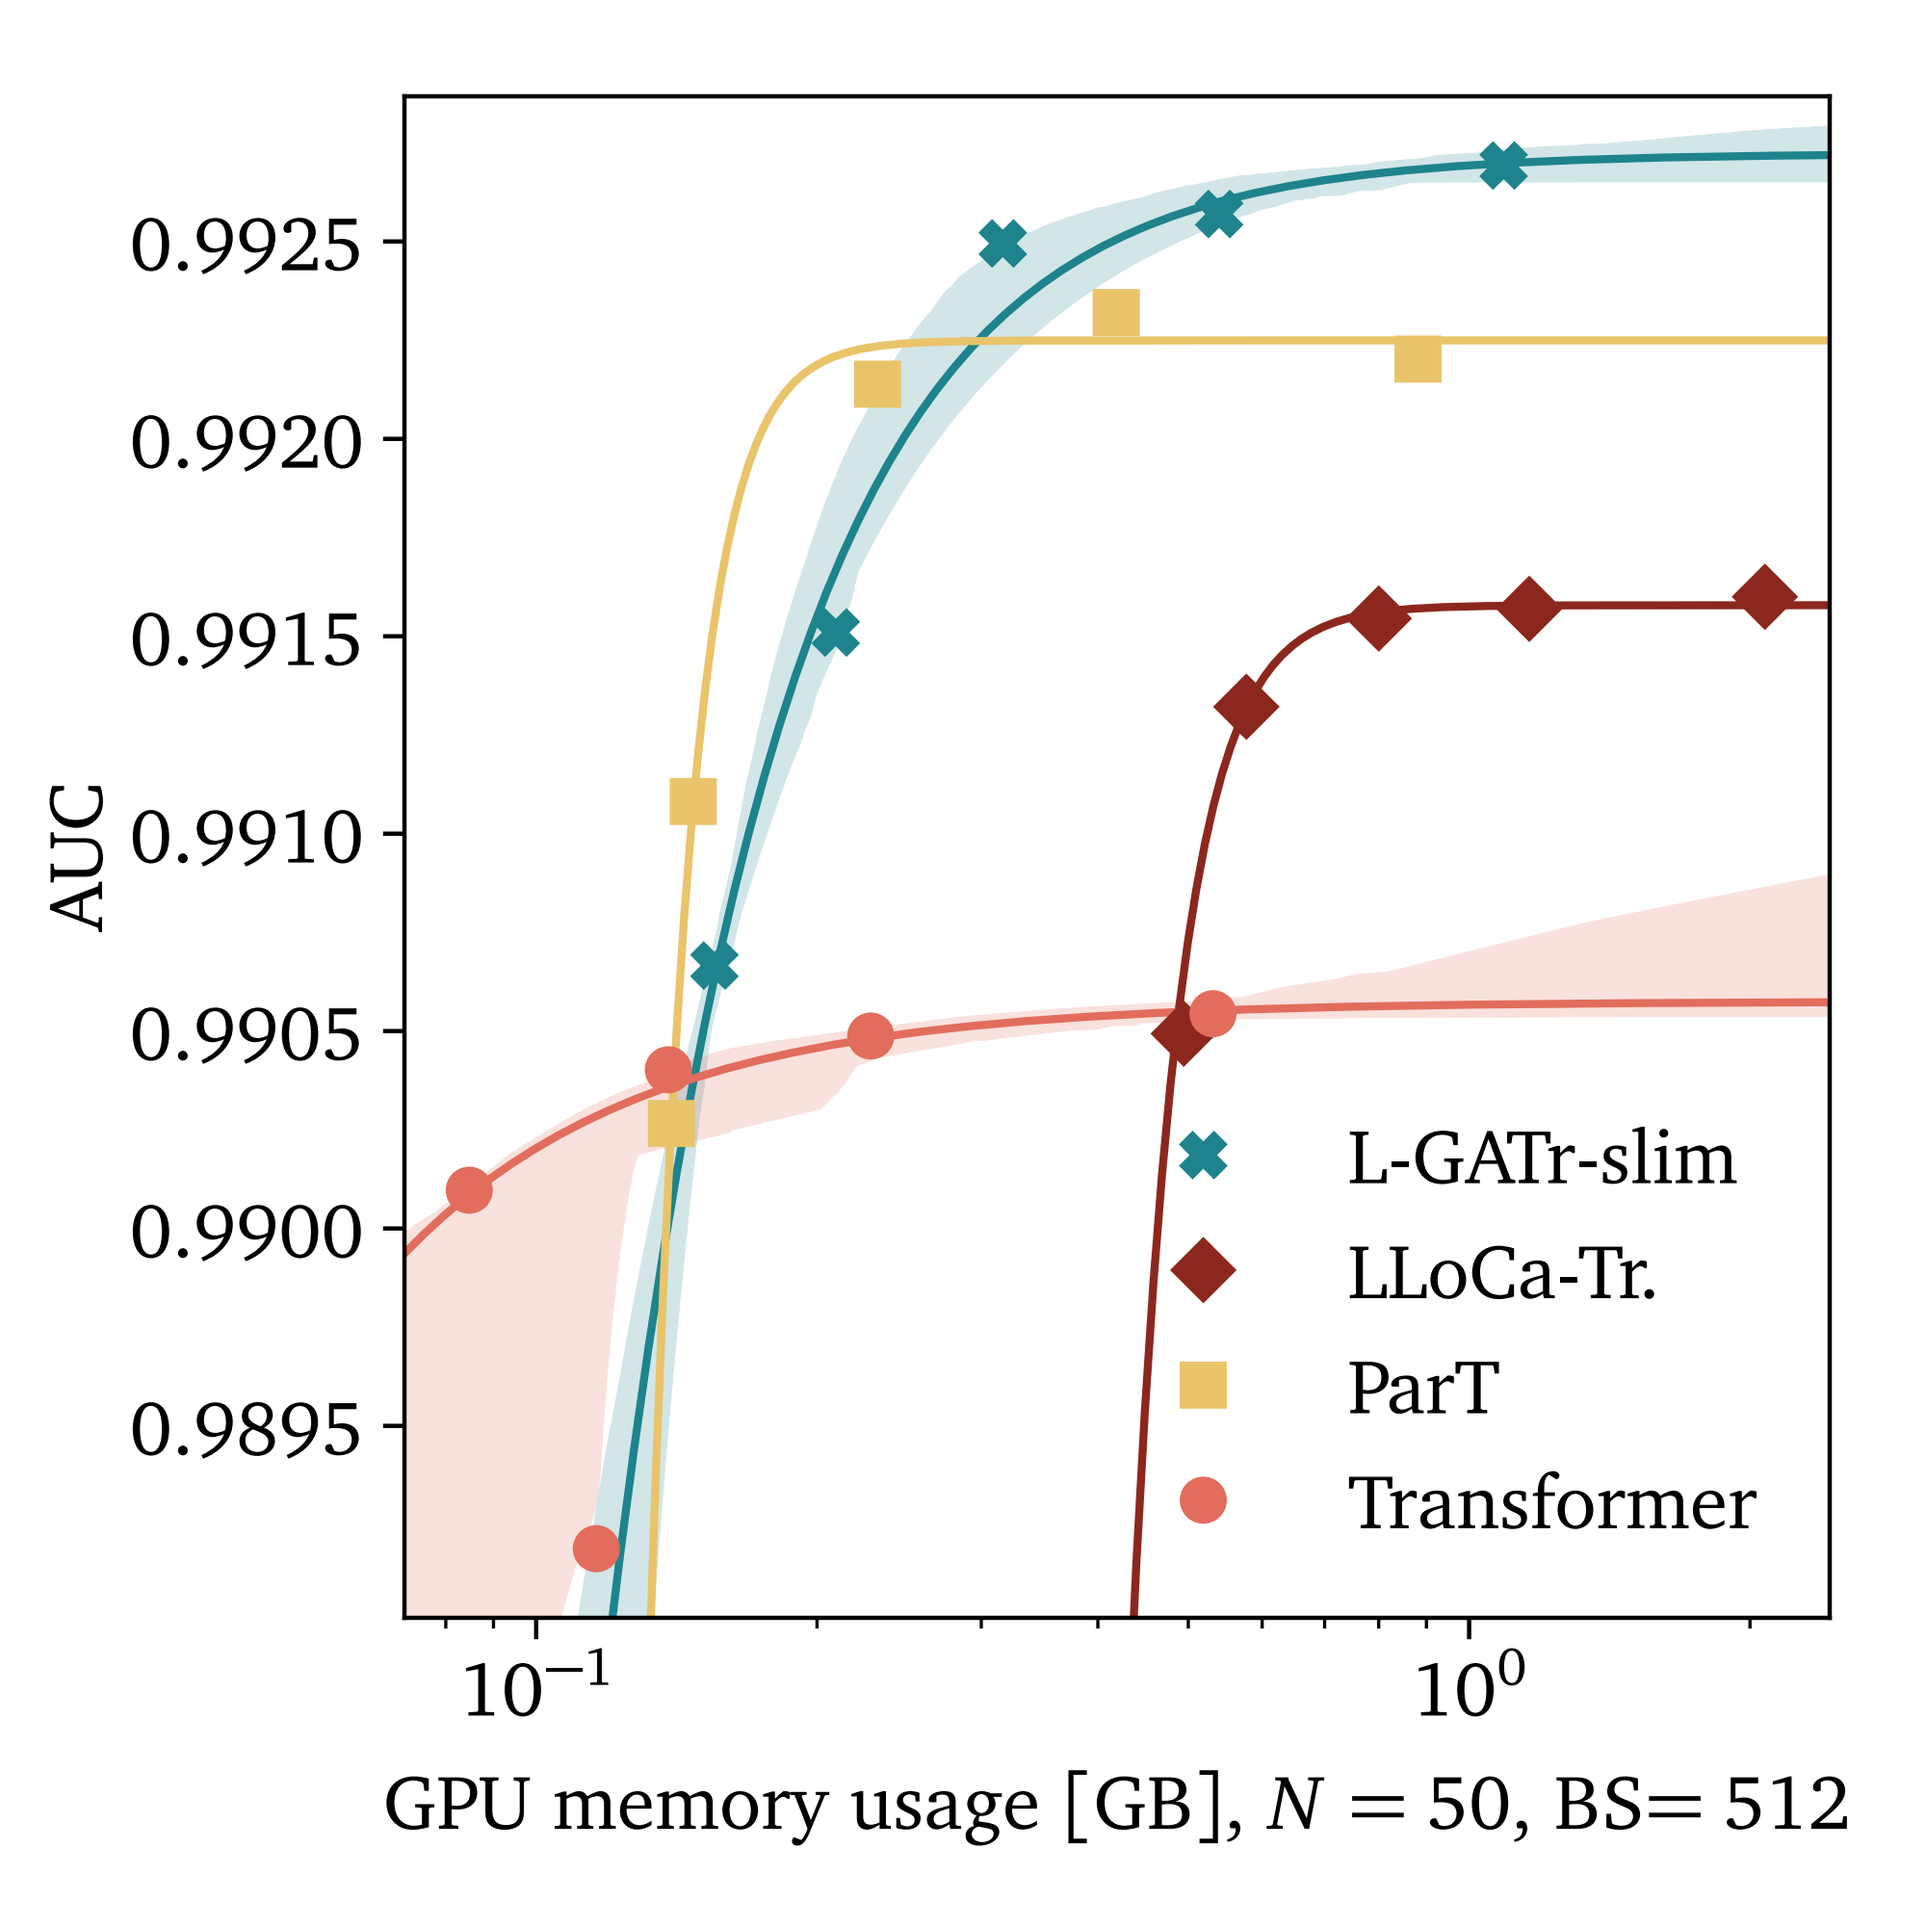
0.9925
0.9920
0.9915
0.9910
0.9905
0.9900
0.9895
AUC
10^{-1}
10^{0}
GPU memory usage [GB], $N = 50$, BS= 512
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

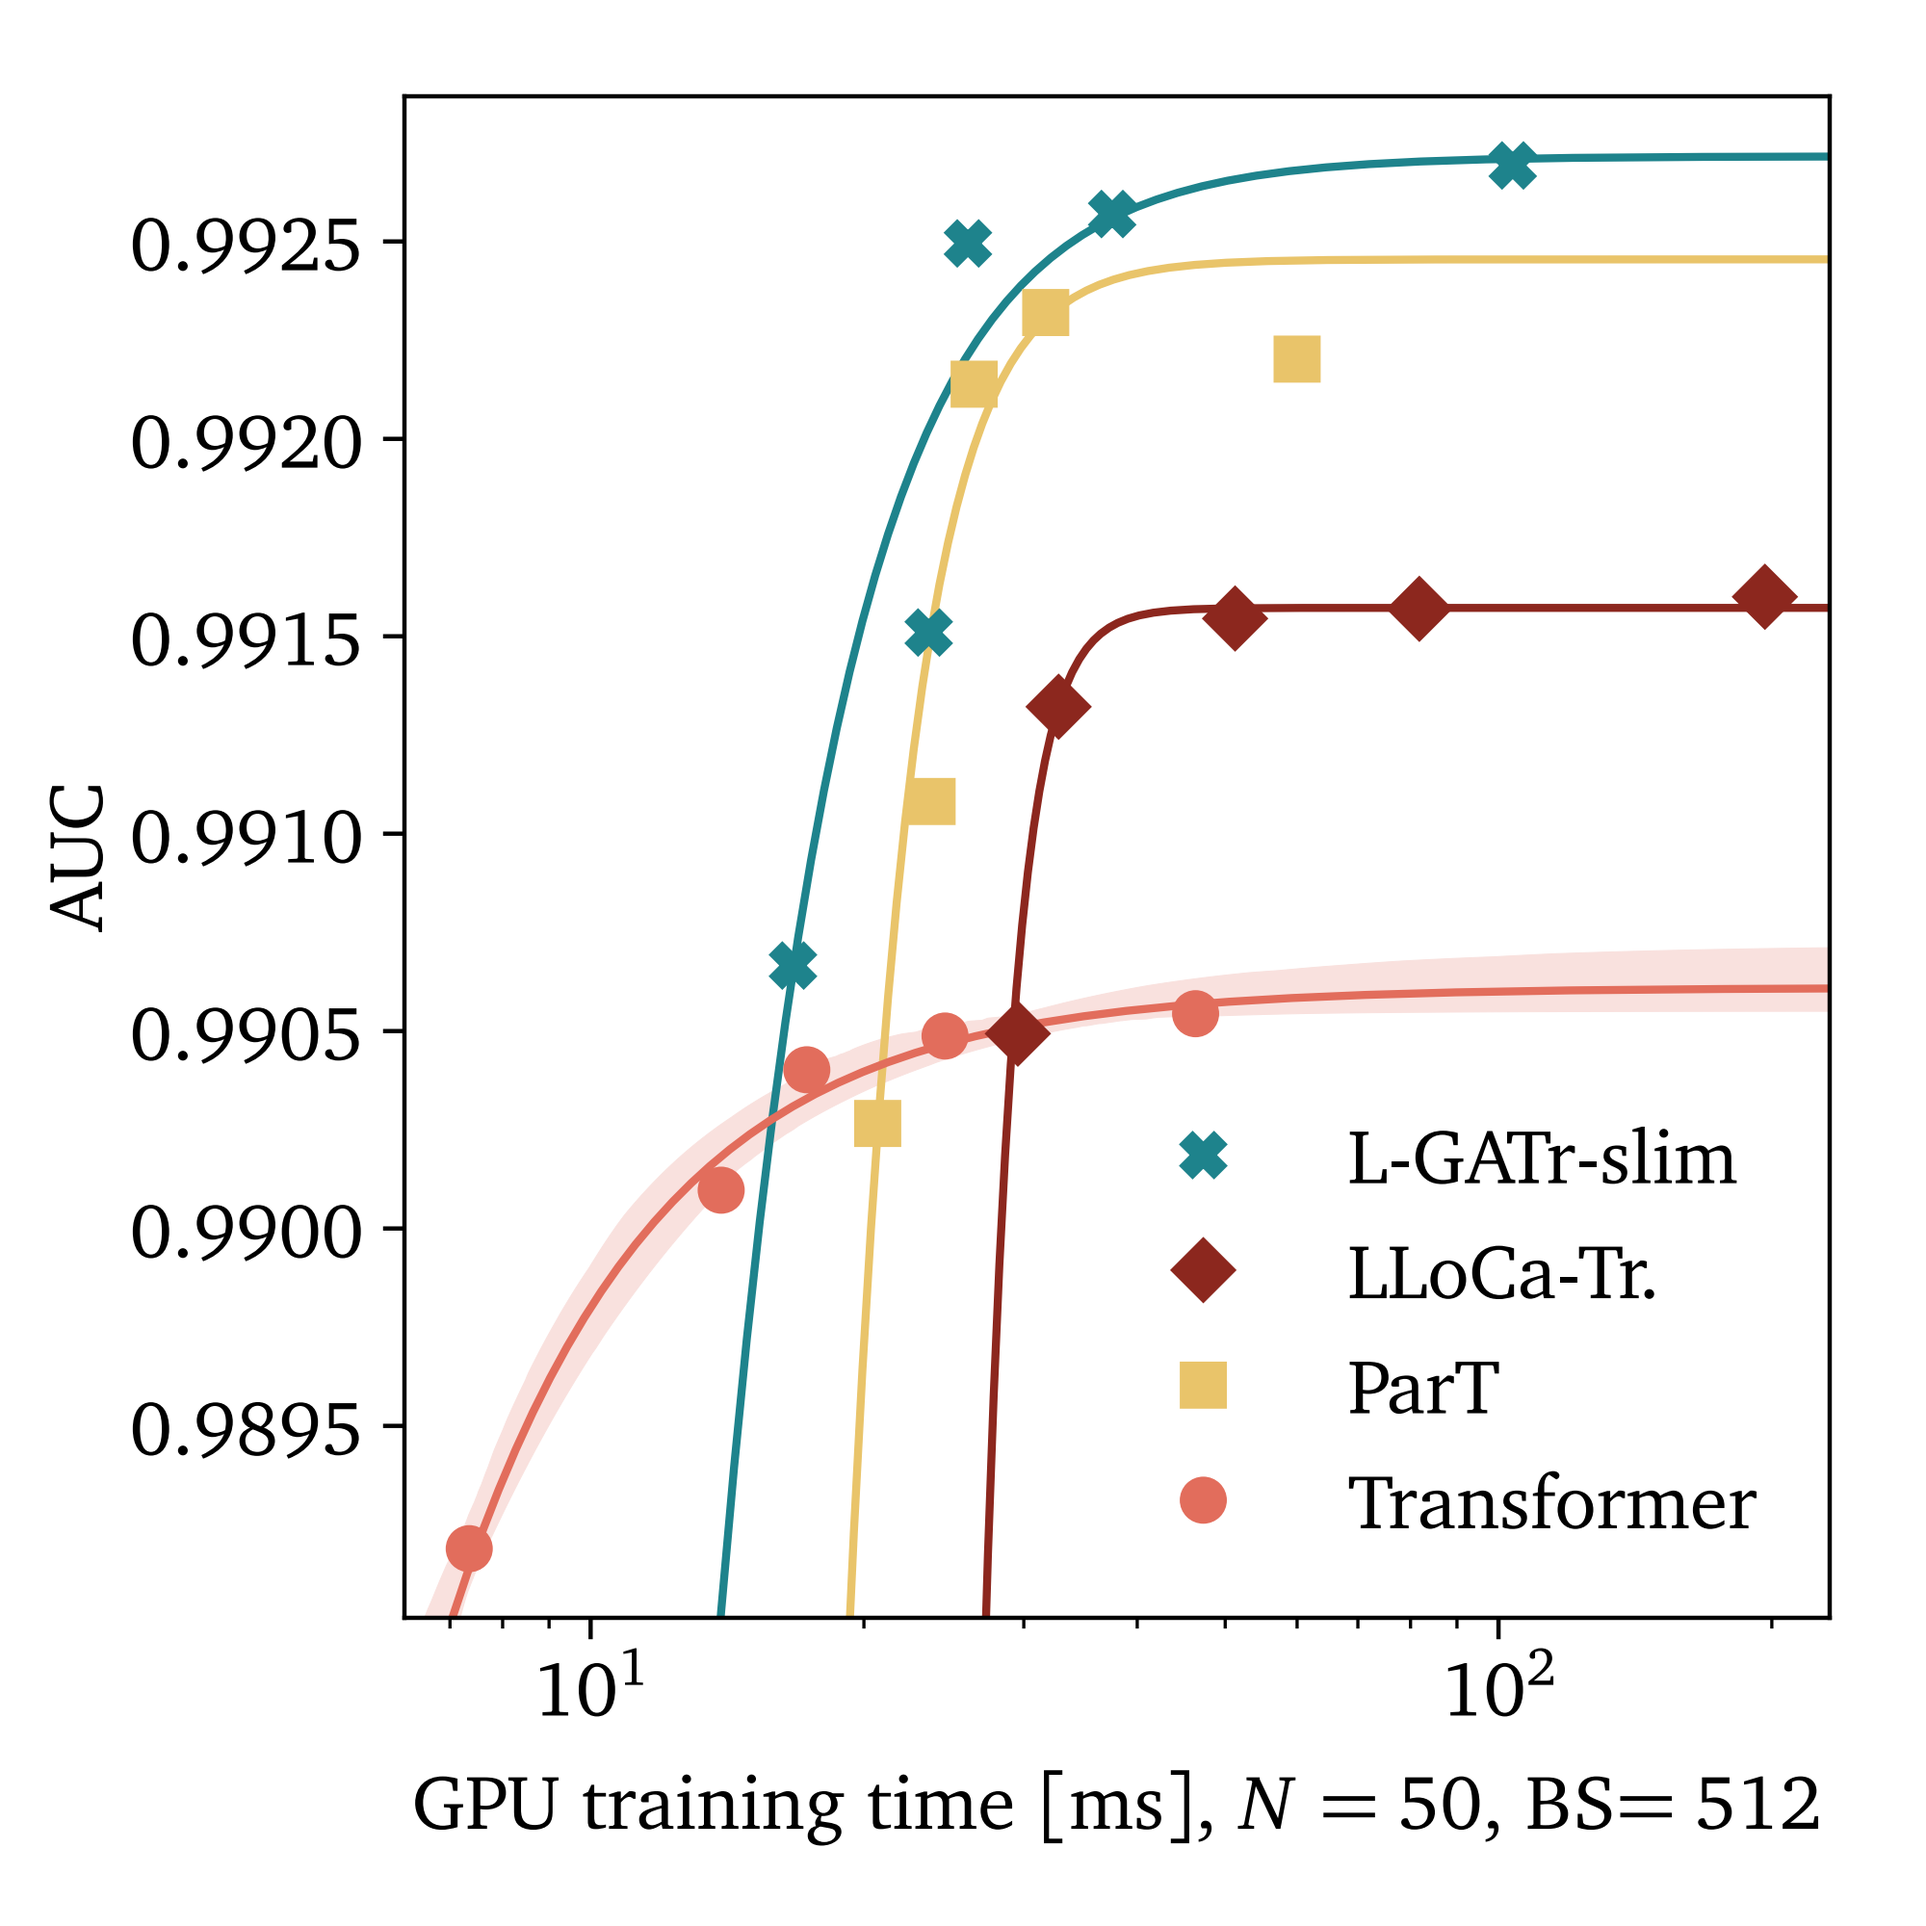
0.9925
0.9920
0.9915
0.9910
0.9905
0.9900
0.9895
AUC
10^1
10^2
GPU training time [ms], $N = 50$, BS$= 512$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer

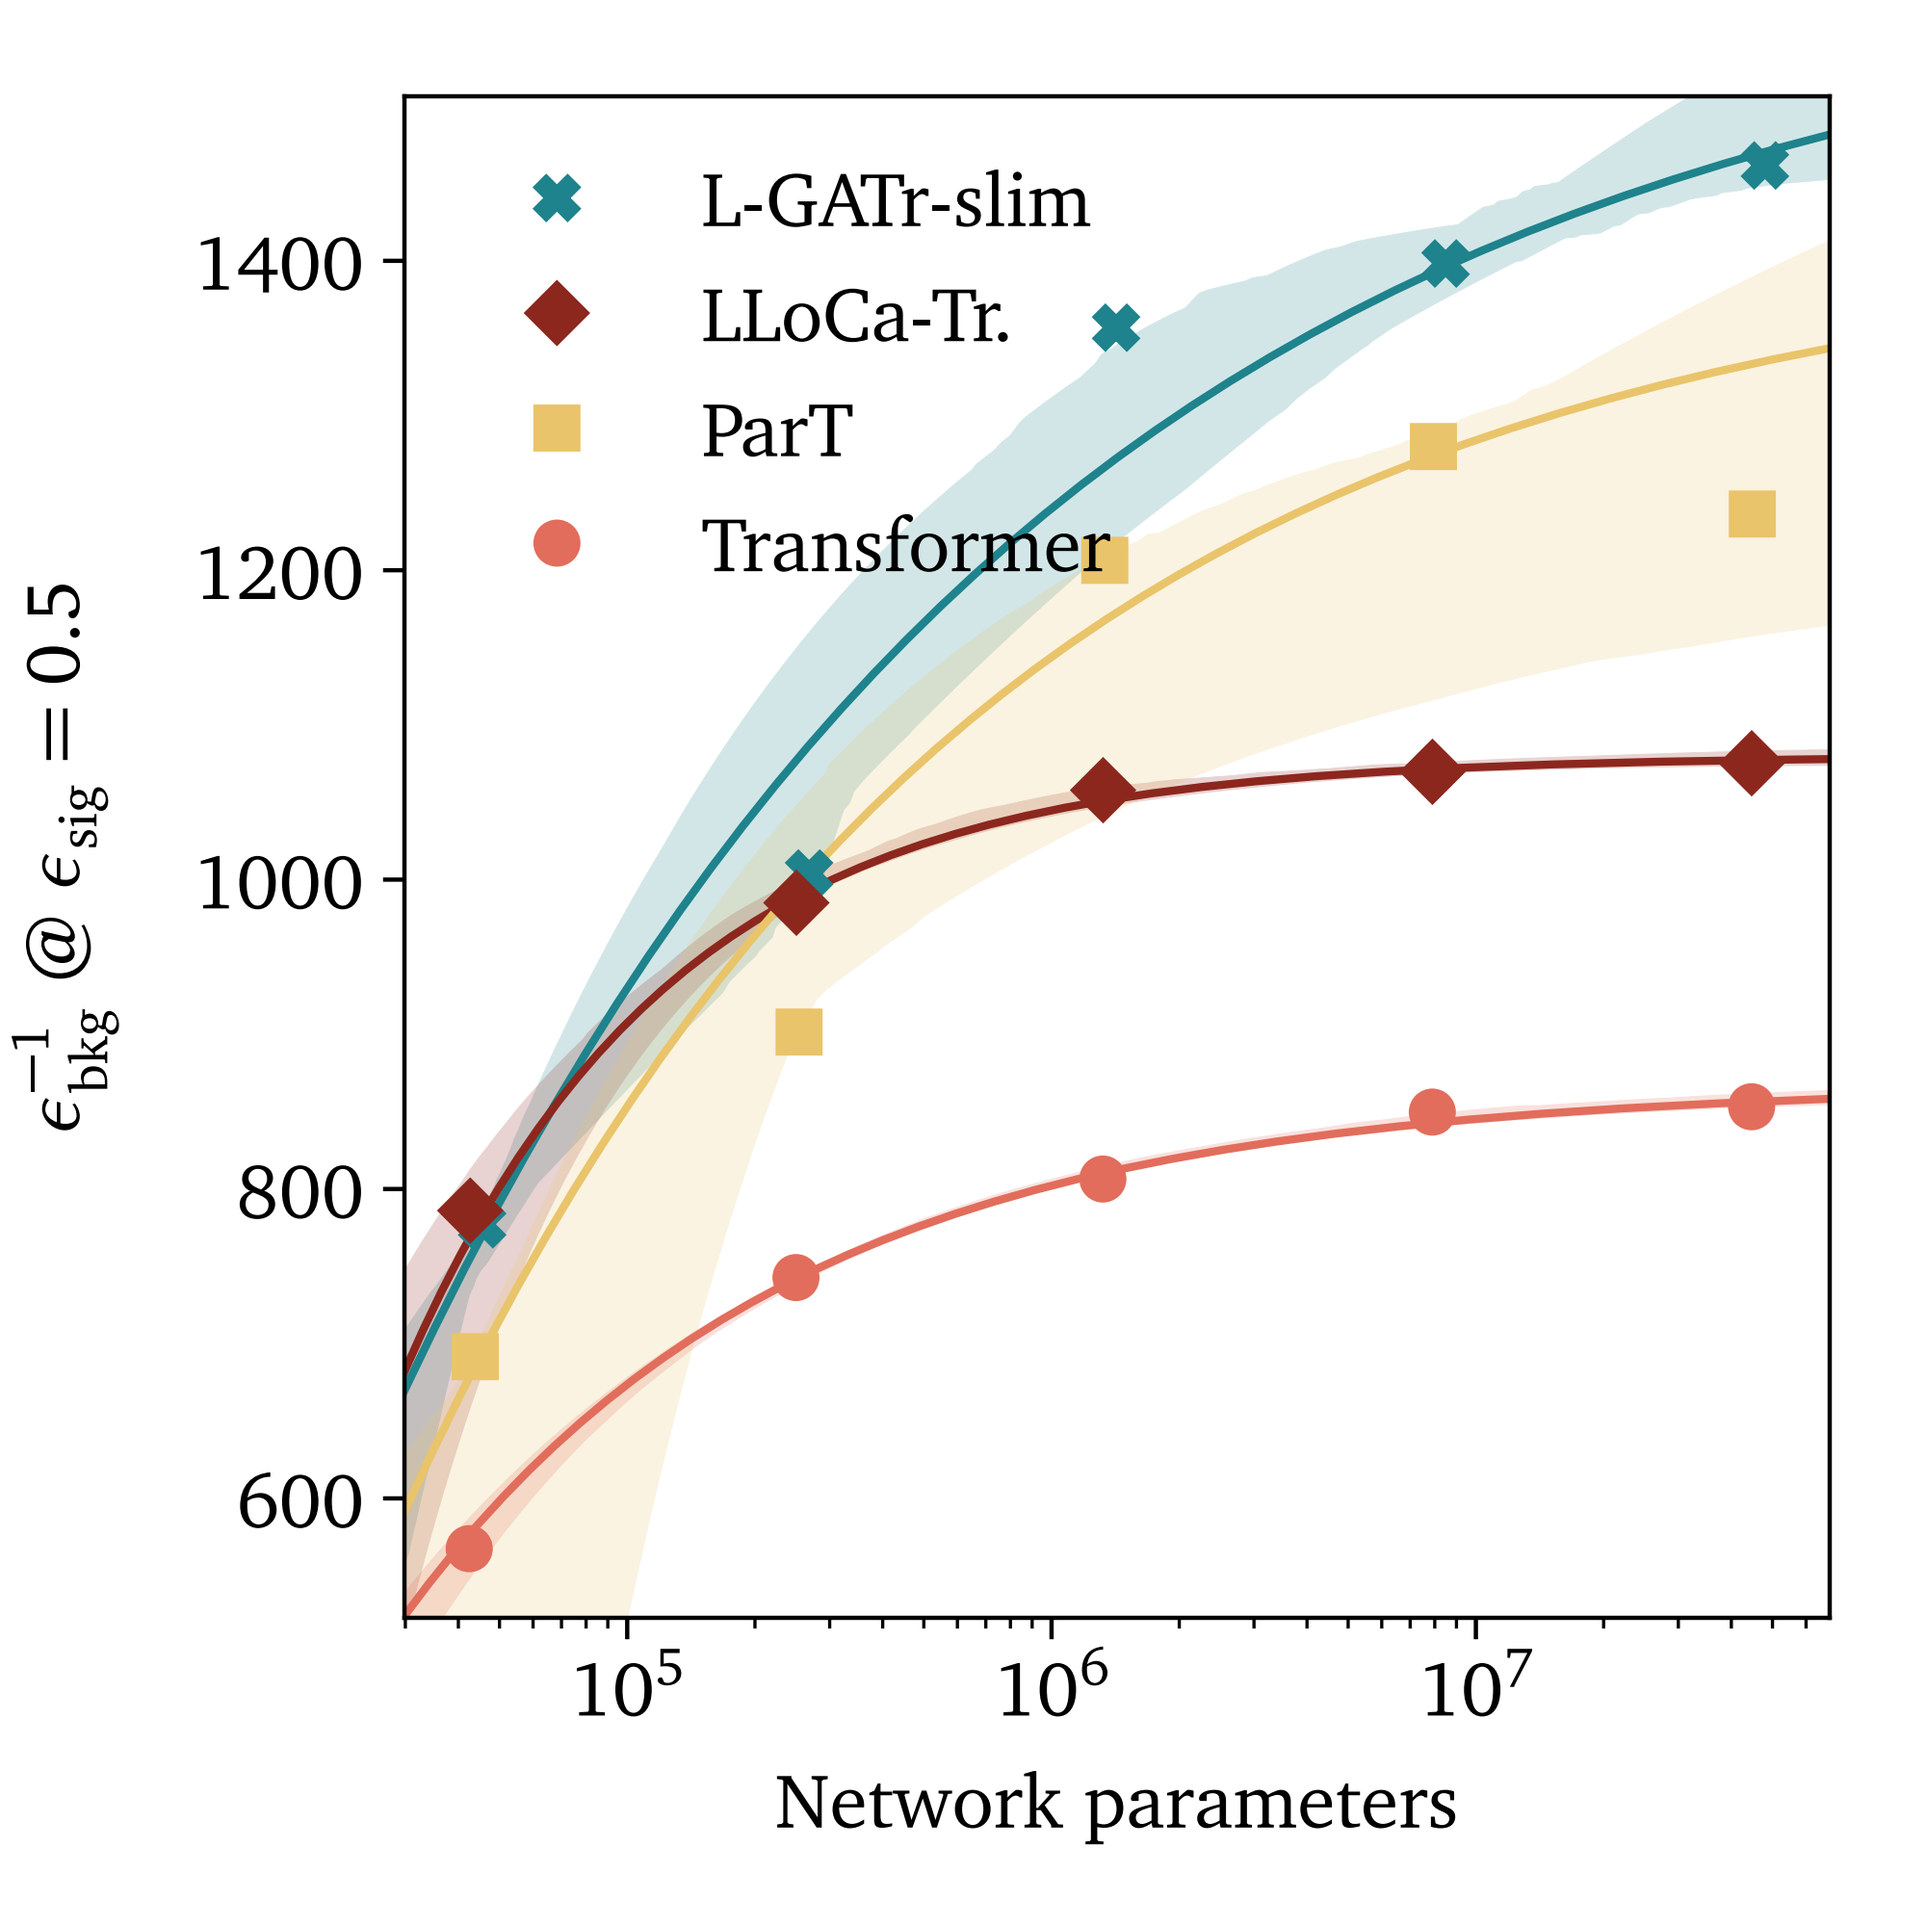
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
1400
1200
1000
800
600
$\epsilon_{\text{bkg}}^{-1}$ @ $\epsilon_{\text{sig}} = 0.5$
10^5
10^6
10^7
Network parameters

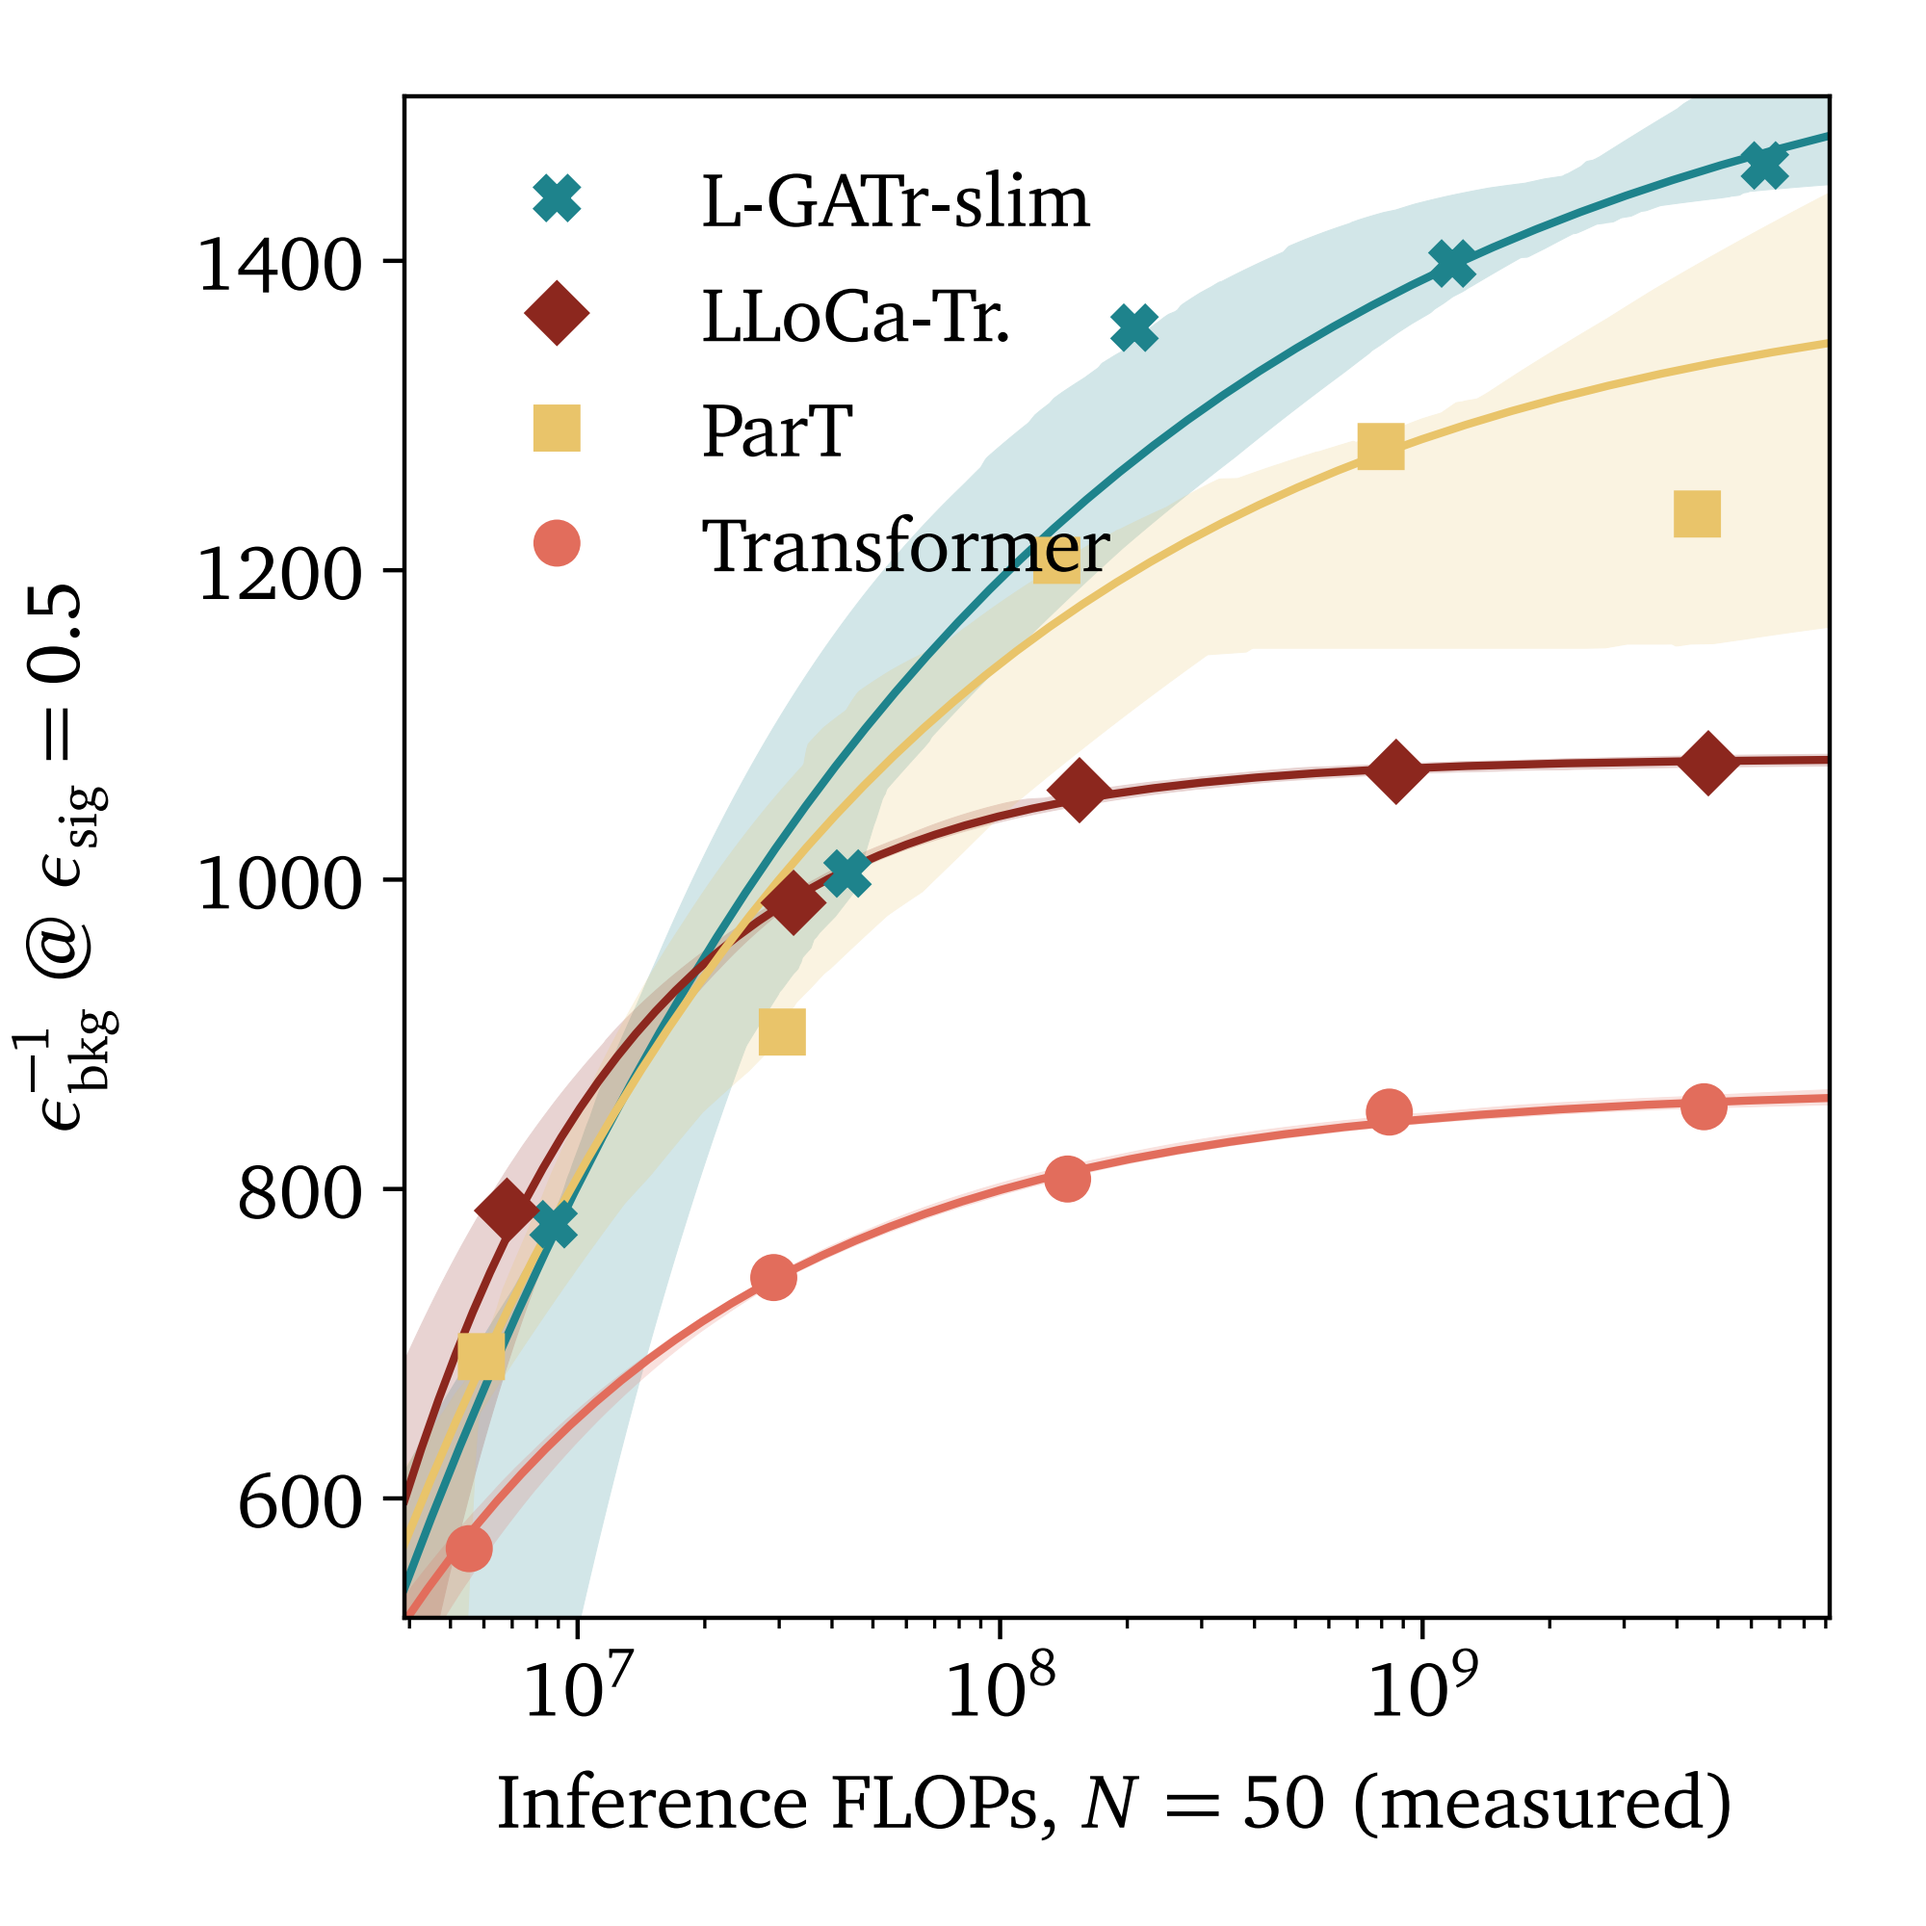
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
1400
1200
1000
800
600
$\epsilon_{\text{bkg}}^{-1}$ @ $\epsilon_{\text{sig}} = 0.5$
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (measured)

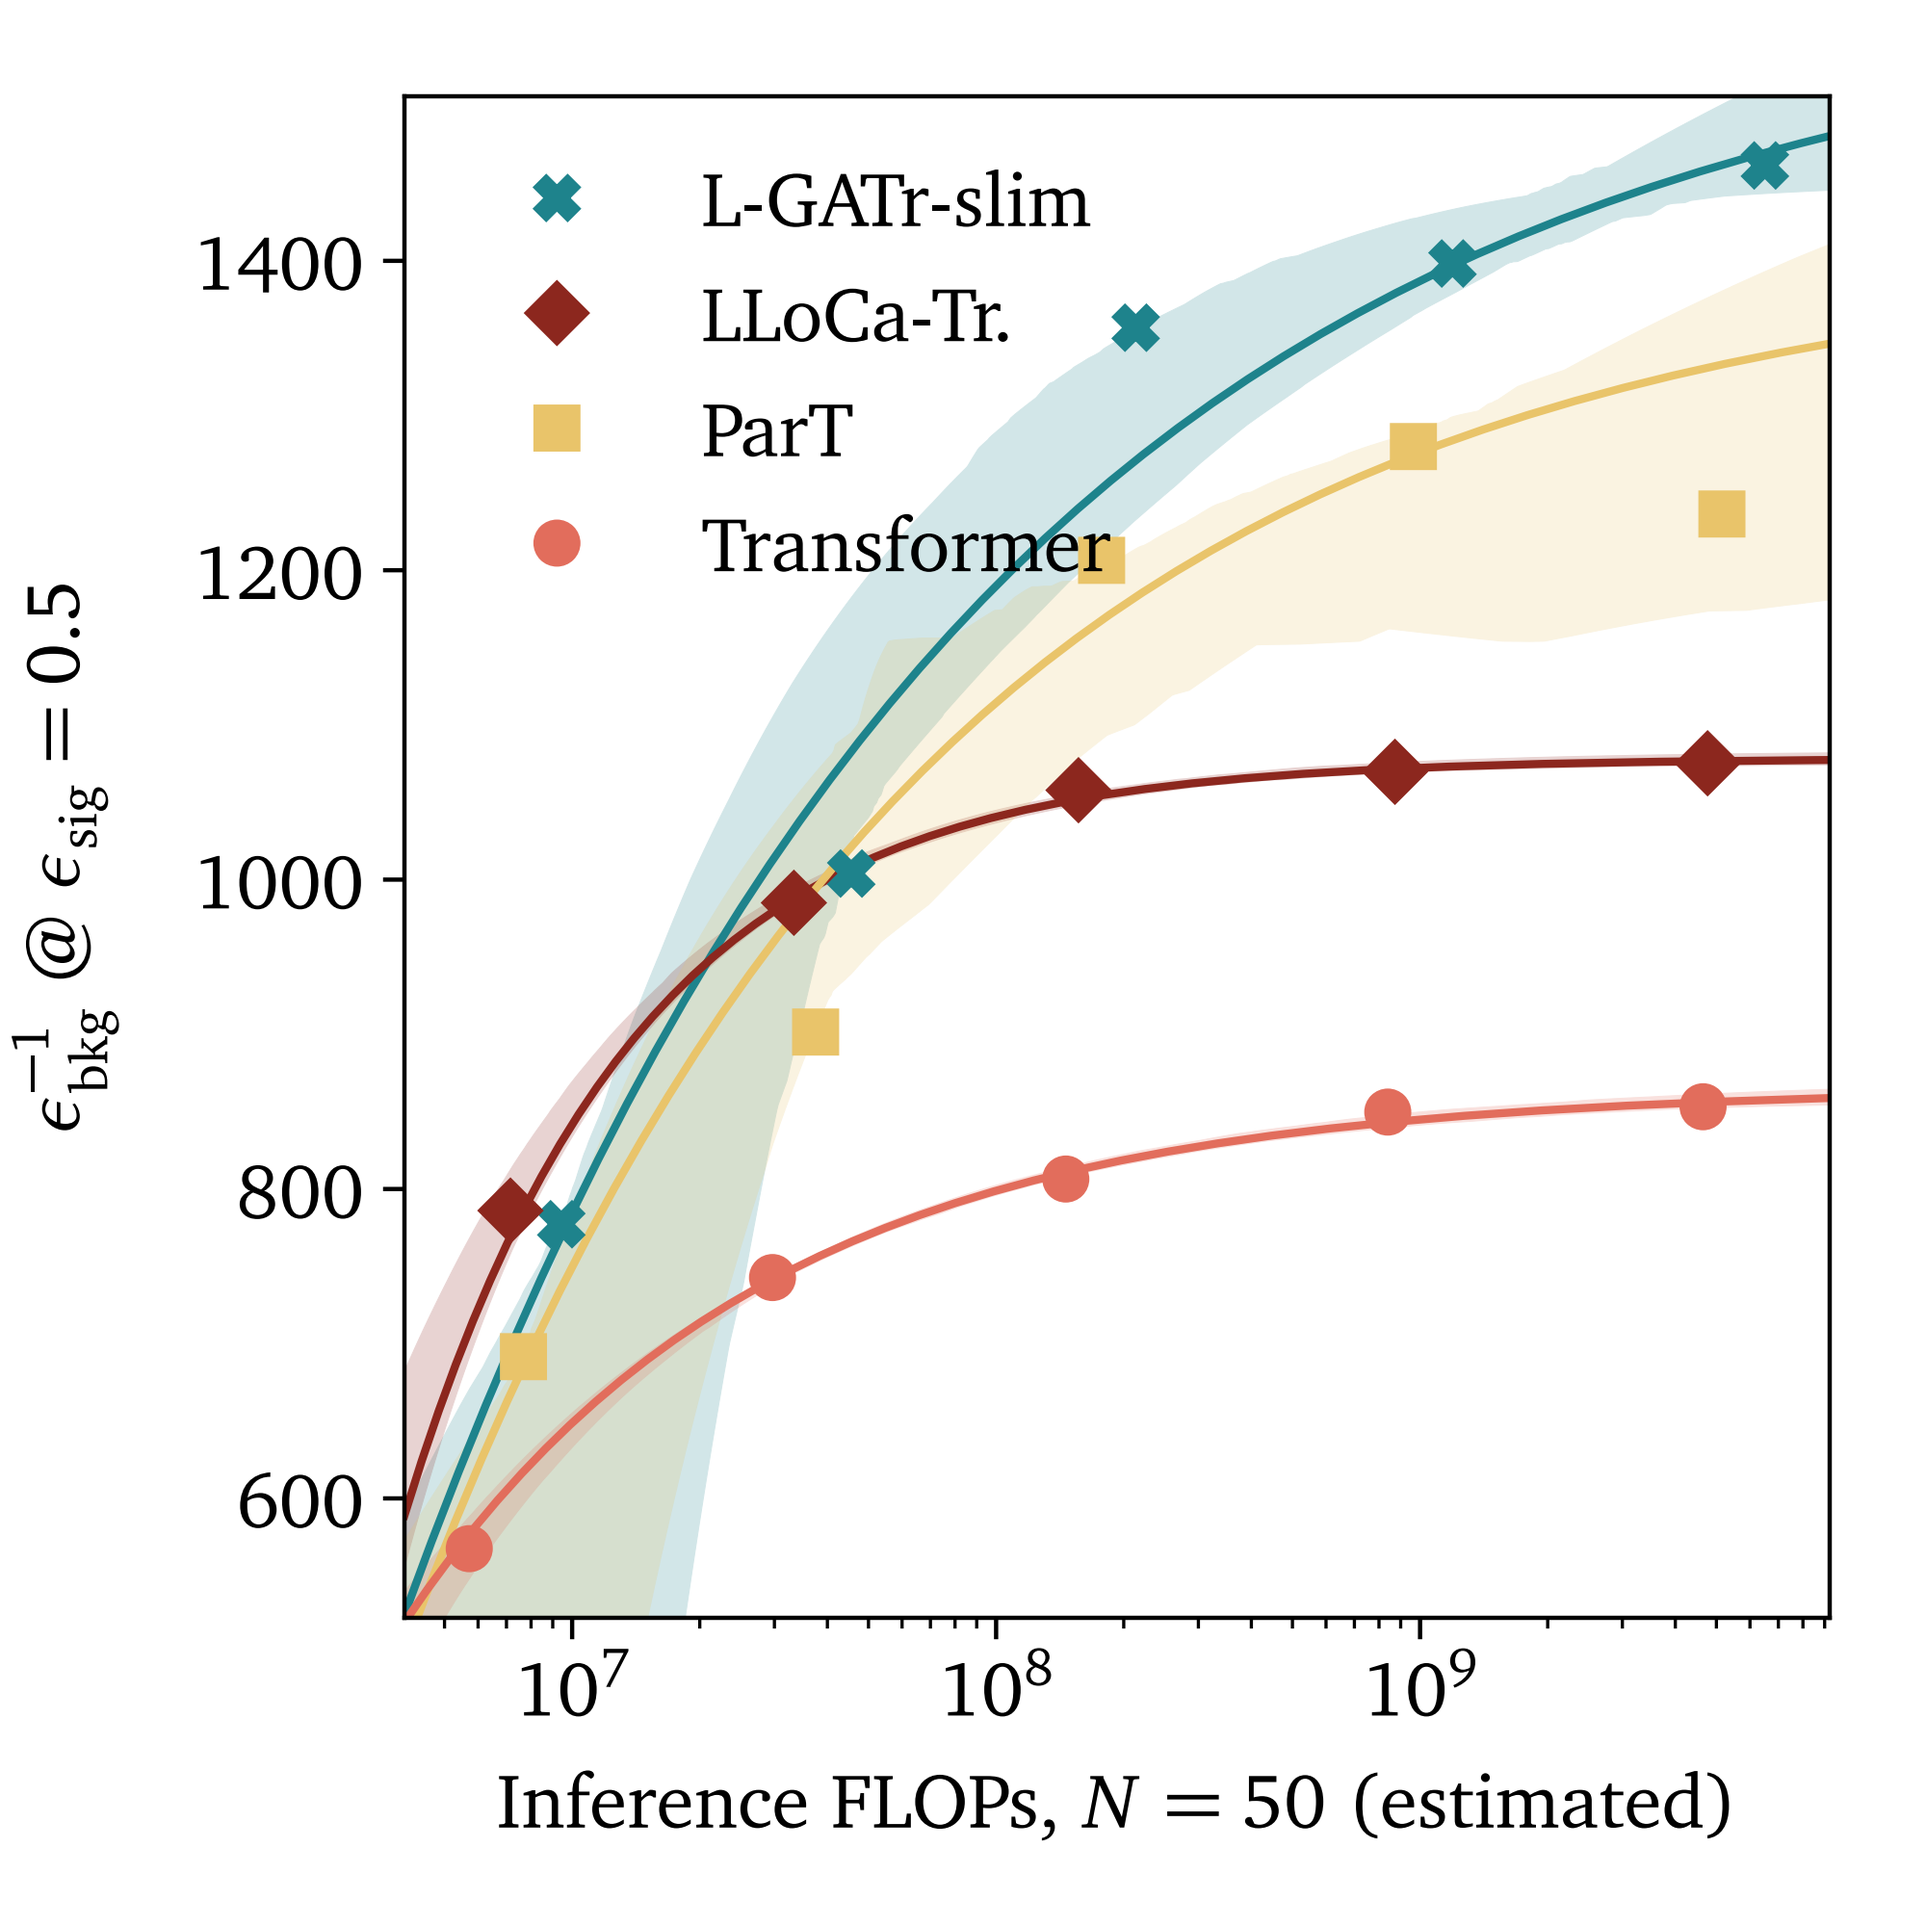
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
1400
1200
1000
800
600
$\epsilon_{\text{bkg}}^{-1}$ @ $\epsilon_{\text{sig}} = 0.5$
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (estimated)

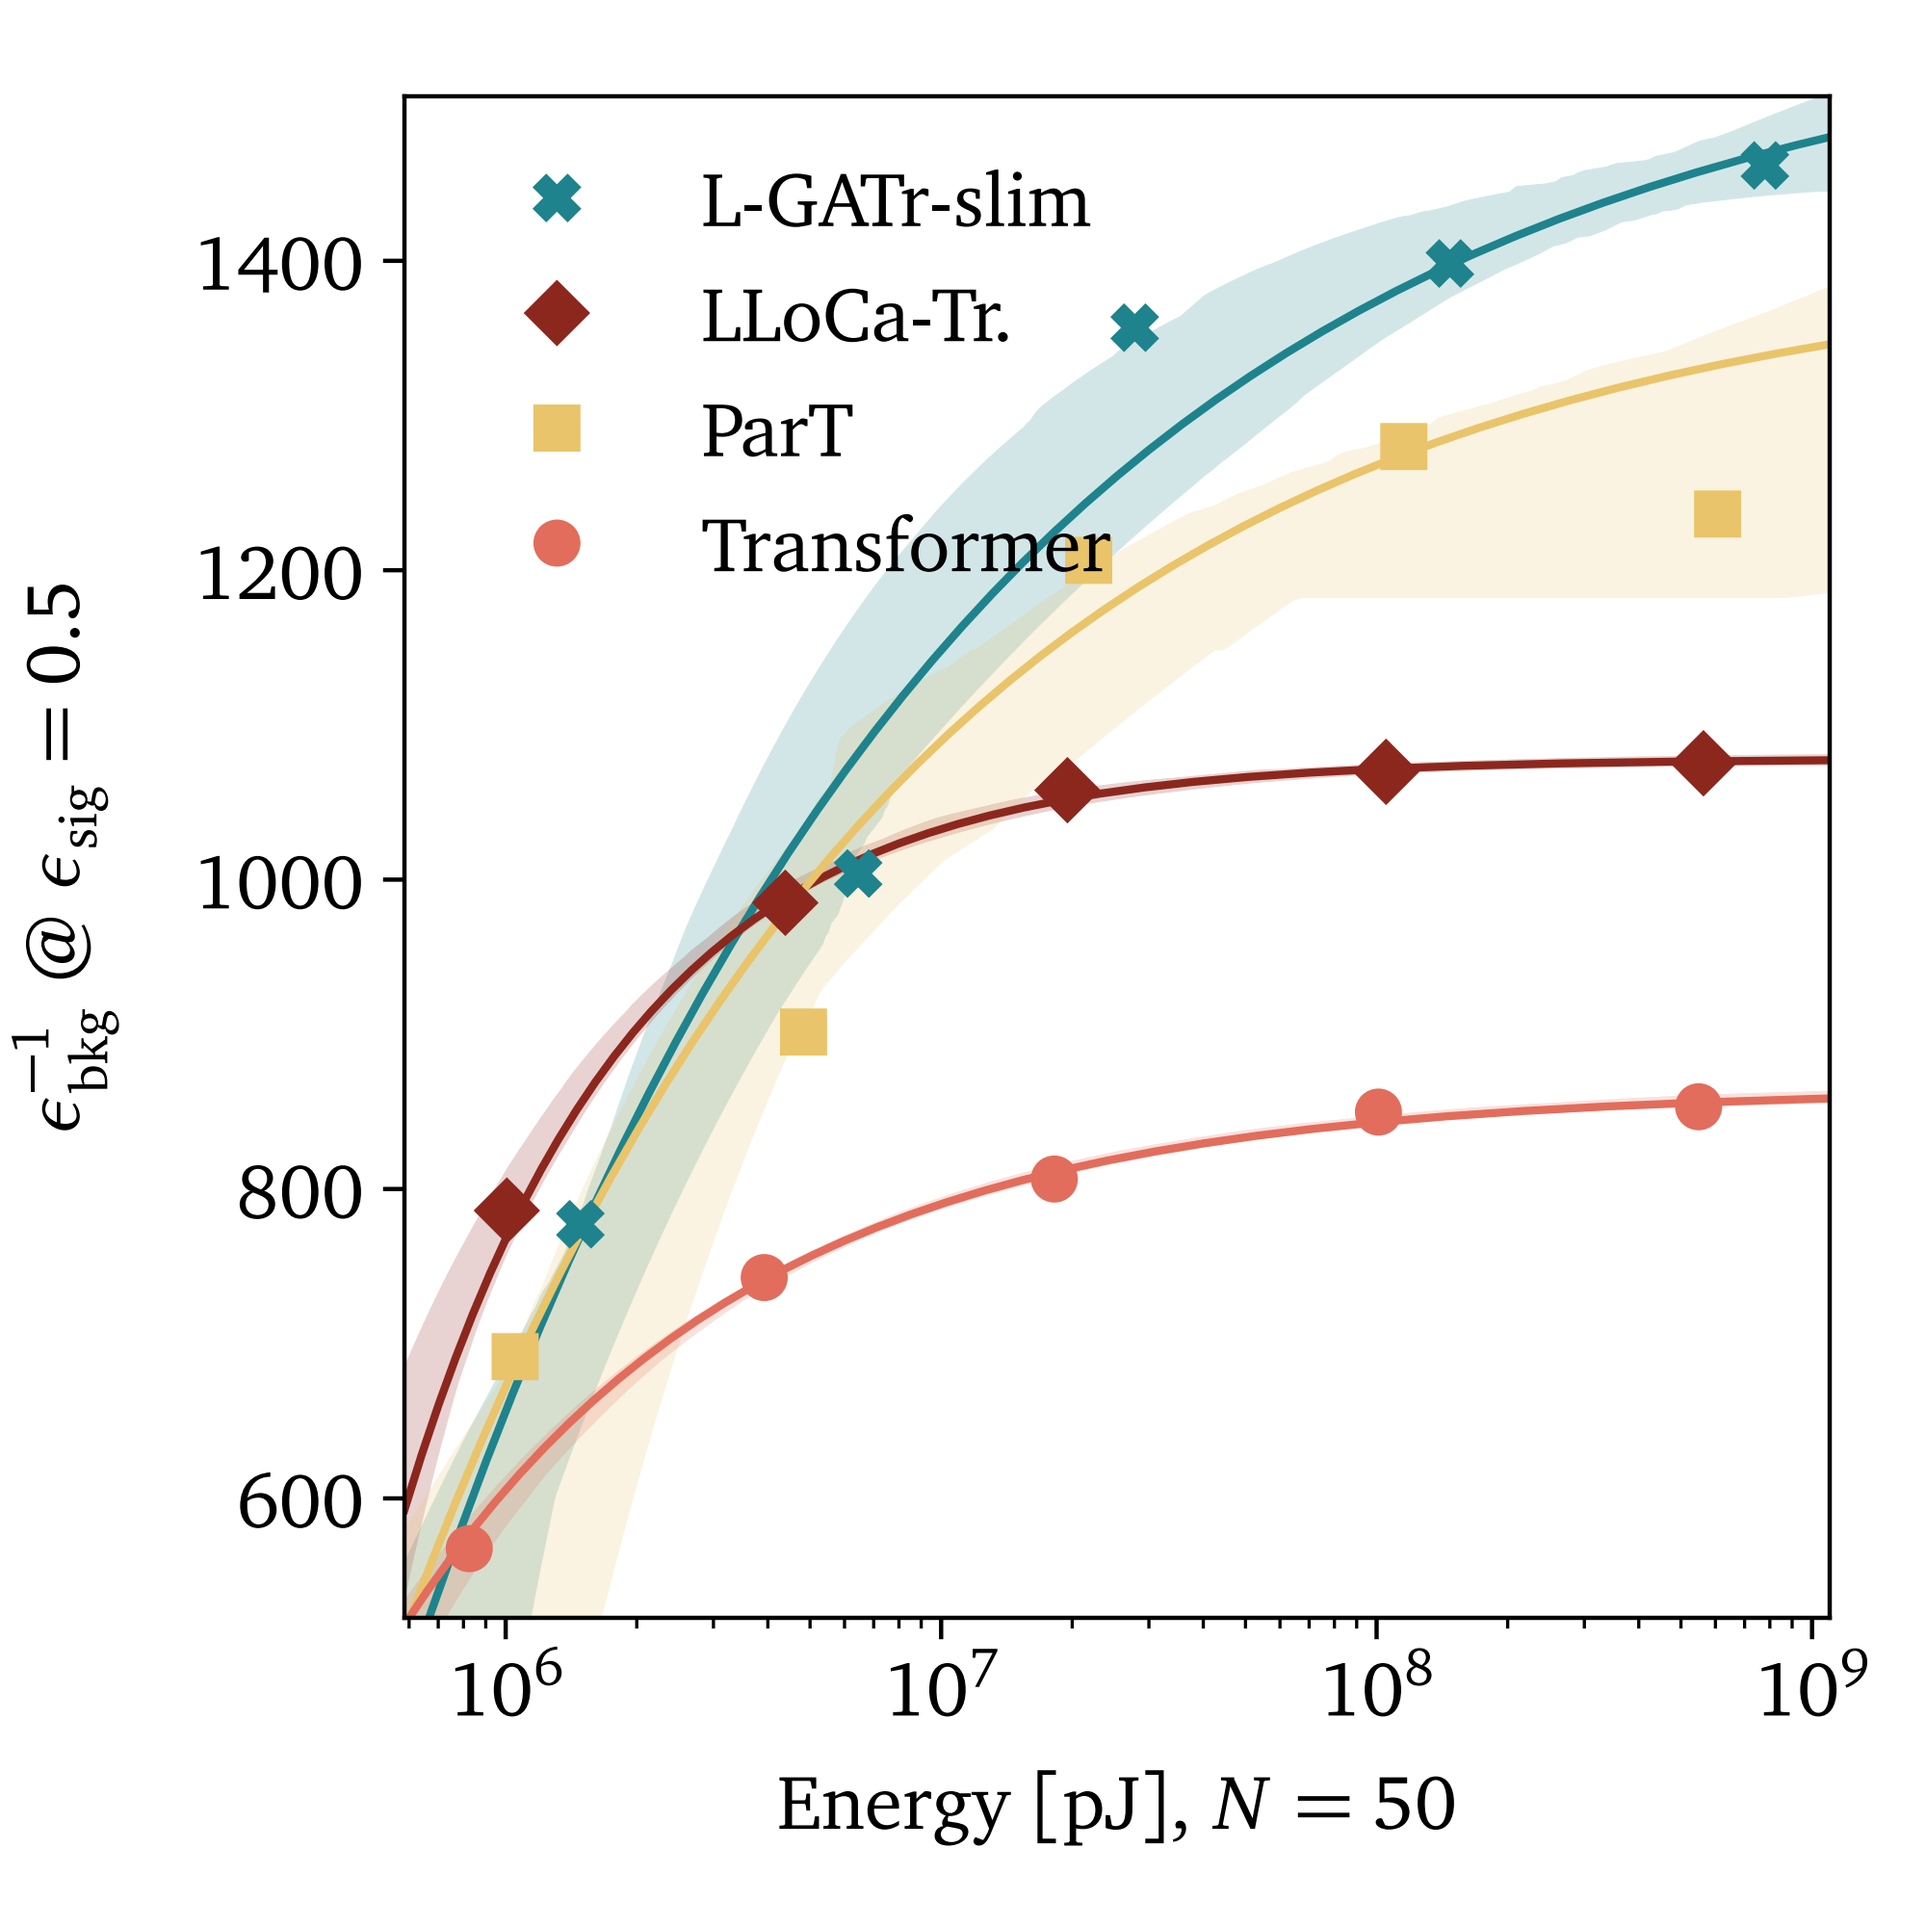
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
1400
1200
1000
800
600
$\epsilon_{\text{bkg}}^{-1}$ @ $\epsilon_{\text{sig}} = 0.5$
10^6
10^7
10^8
10^9
Energy [pJ], $N = 50$

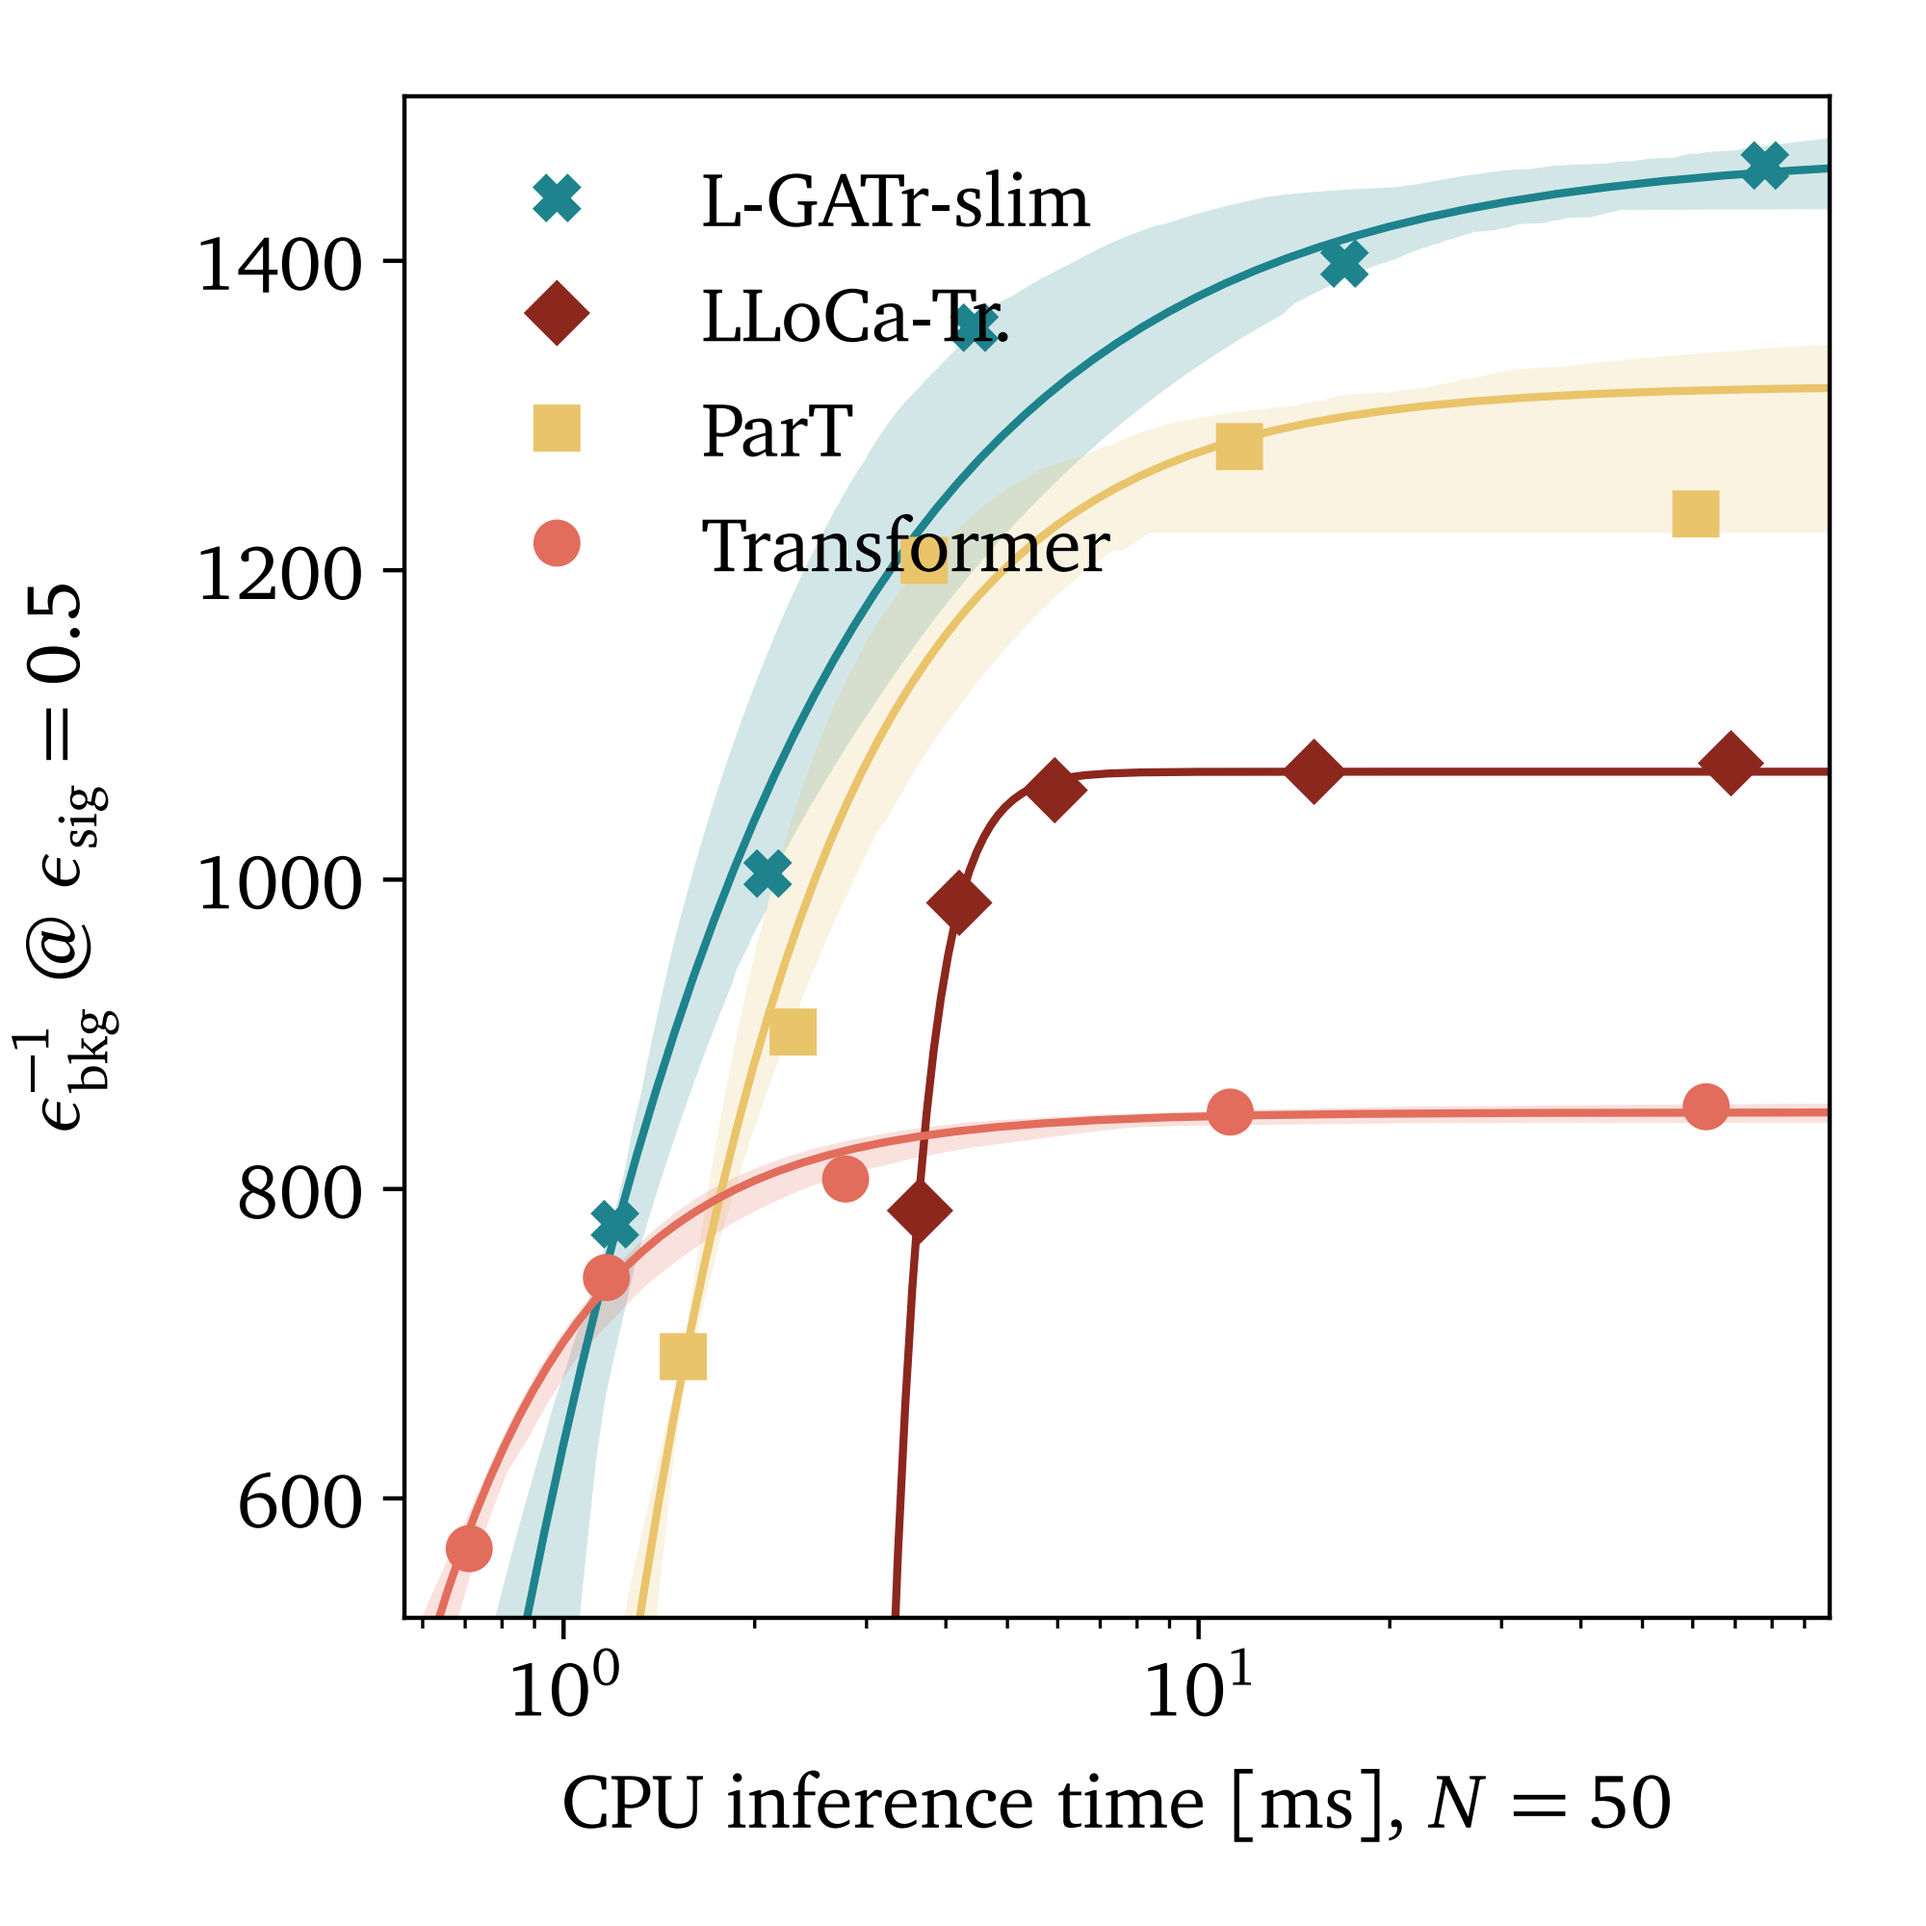
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
1400
1200
1000
800
600
$\epsilon_{\text{bkg}}^{-1}$ @ $\epsilon_{\text{sig}} = 0.5$
10^0
10^1
CPU inference time [ms], $N = 50$

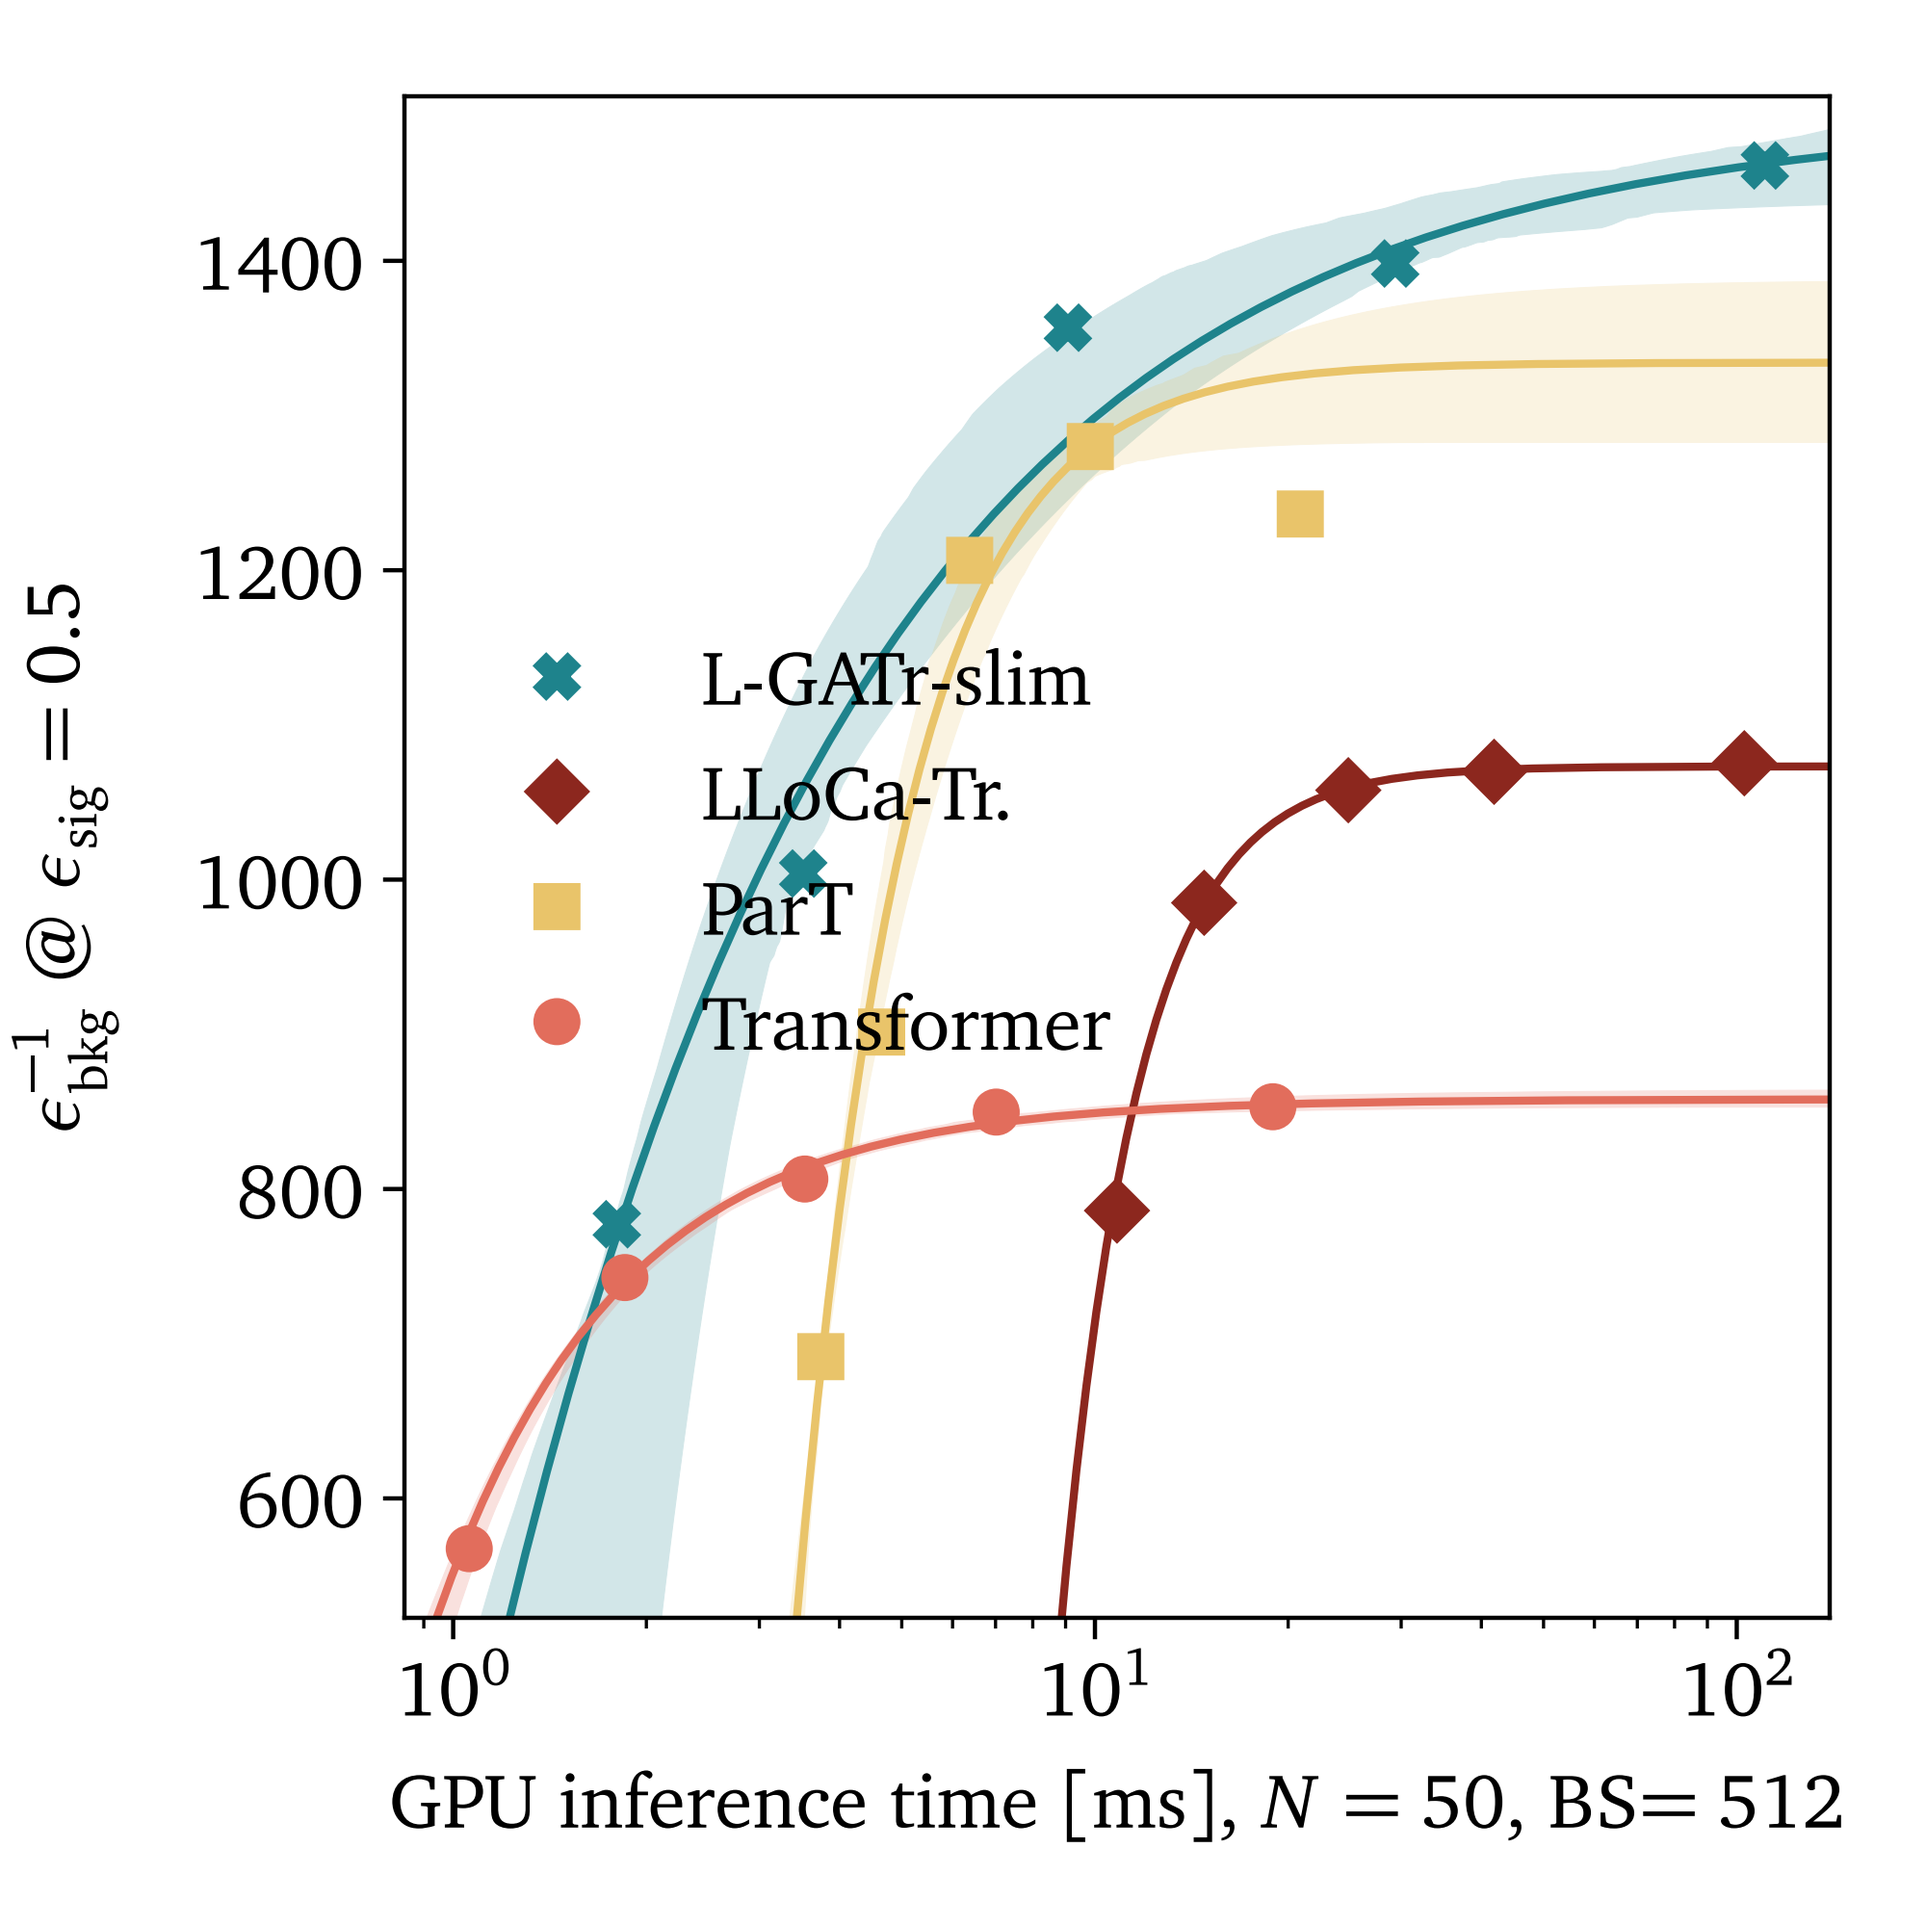
1400
1200
1000
800
600
$\epsilon_{\text{bkg}}^{-1}$ @ $\epsilon_{\text{sig}} = 0.5$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
10^0
10^1
10^2
GPU inference time [ms], $N = 50$, BS= 512

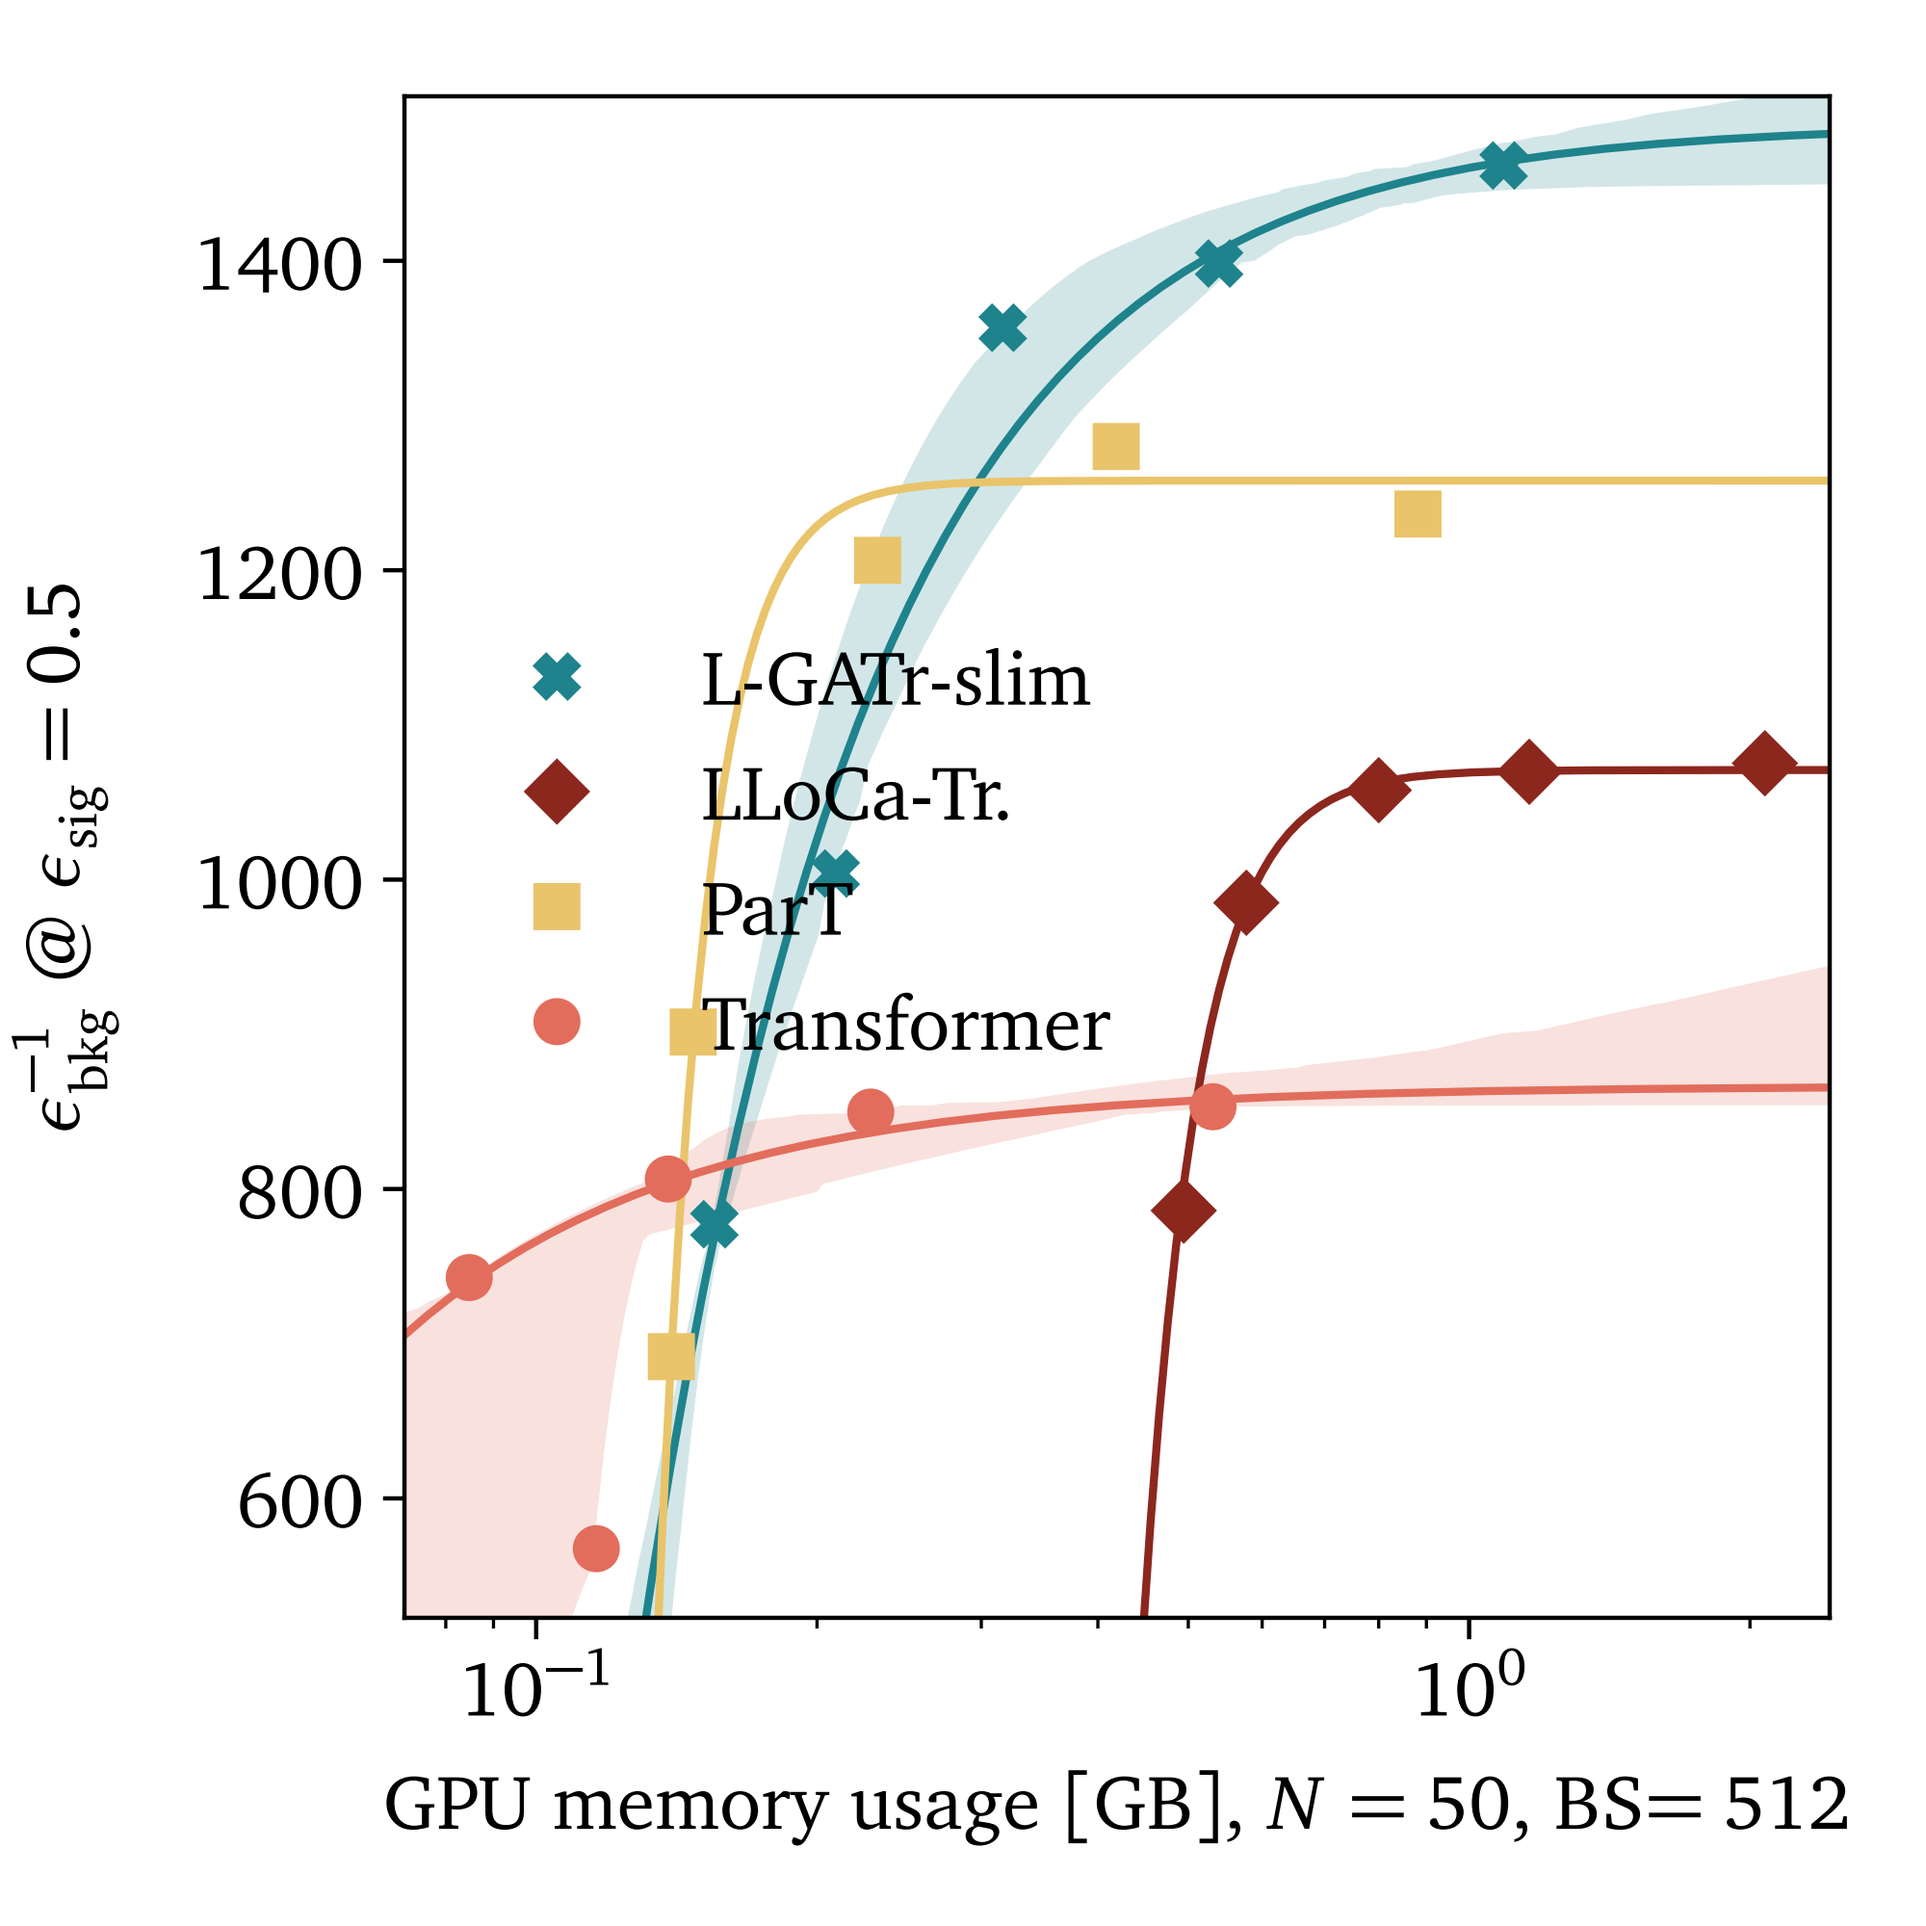
1400
1200
1000
800
600
$\epsilon_{\text{bkg}}^{-1}$ @ $\epsilon_{\text{sig}} = 0.5$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
10^{-1}
10^{0}
GPU memory usage [GB], $N = 50$, BS= 512

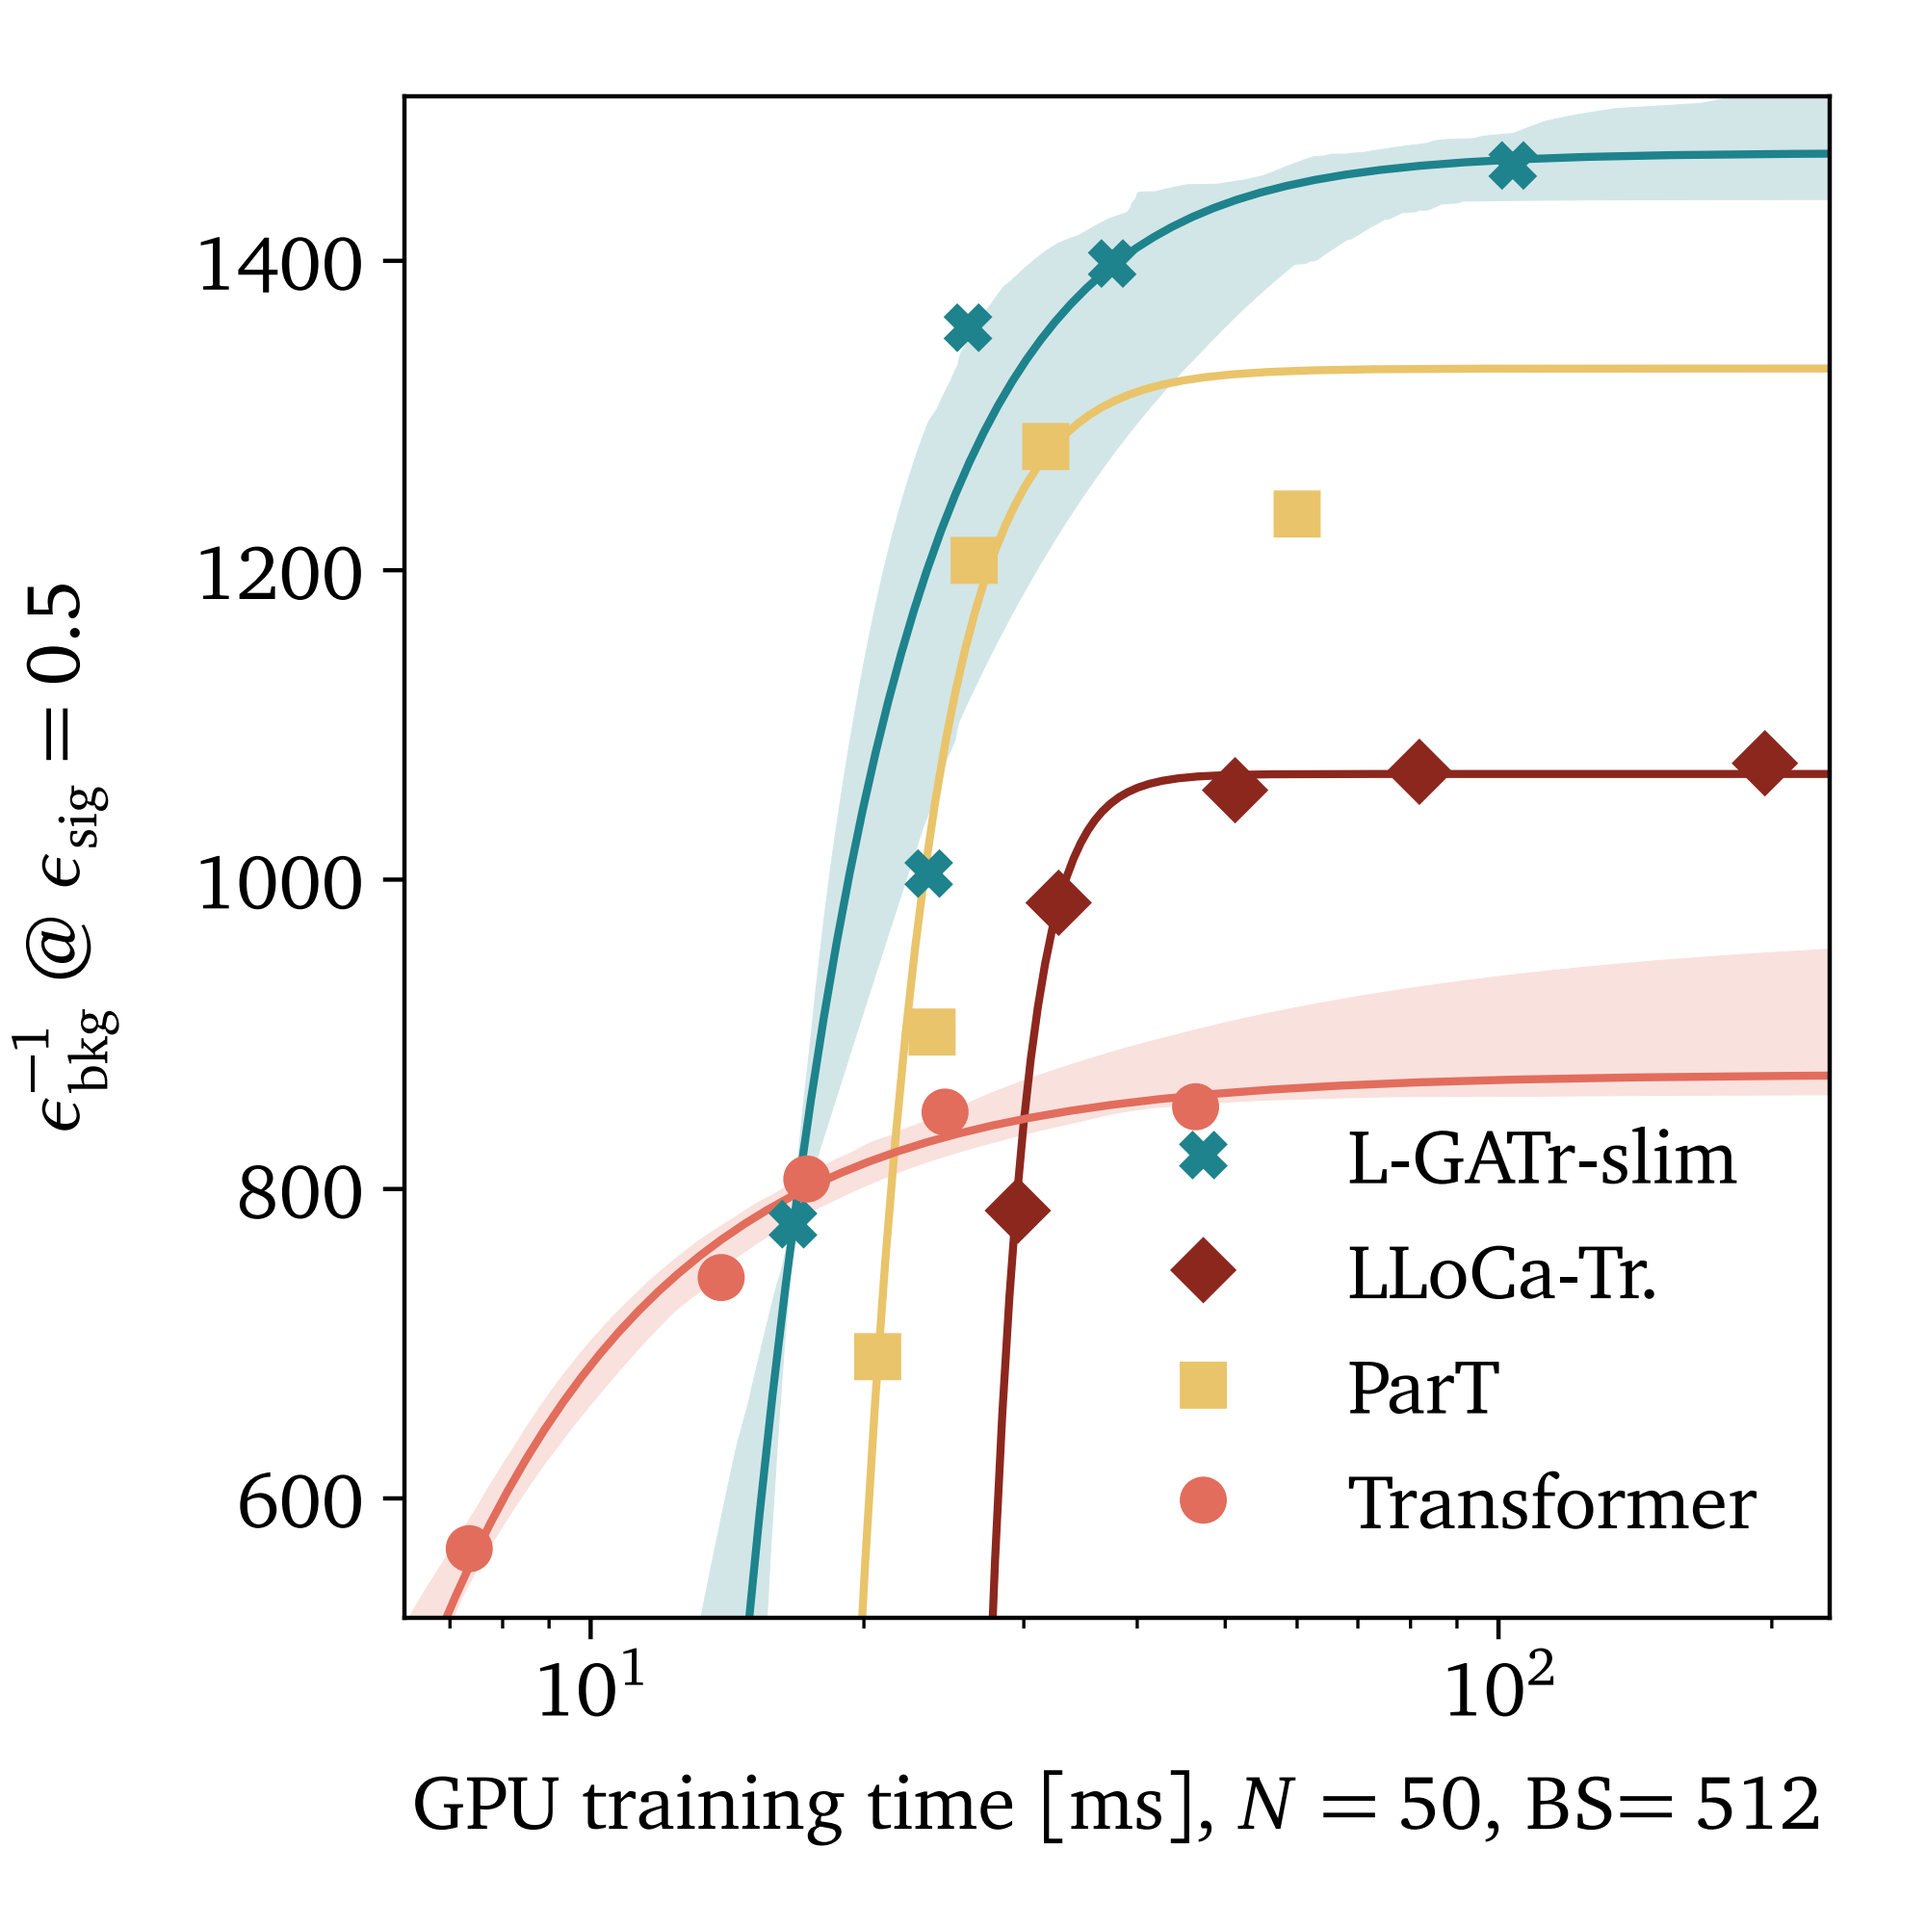
1400
1200
1000
800
600
$\epsilon_{\text{bkg}}^{-1}$ @ $\epsilon_{\text{sig}} = 0.5$
10^1
10^2
GPU training time [ms], $N = 50$, BS= 512
L-GATr-slim
LLoCa-Tr.
ParT
Transformer