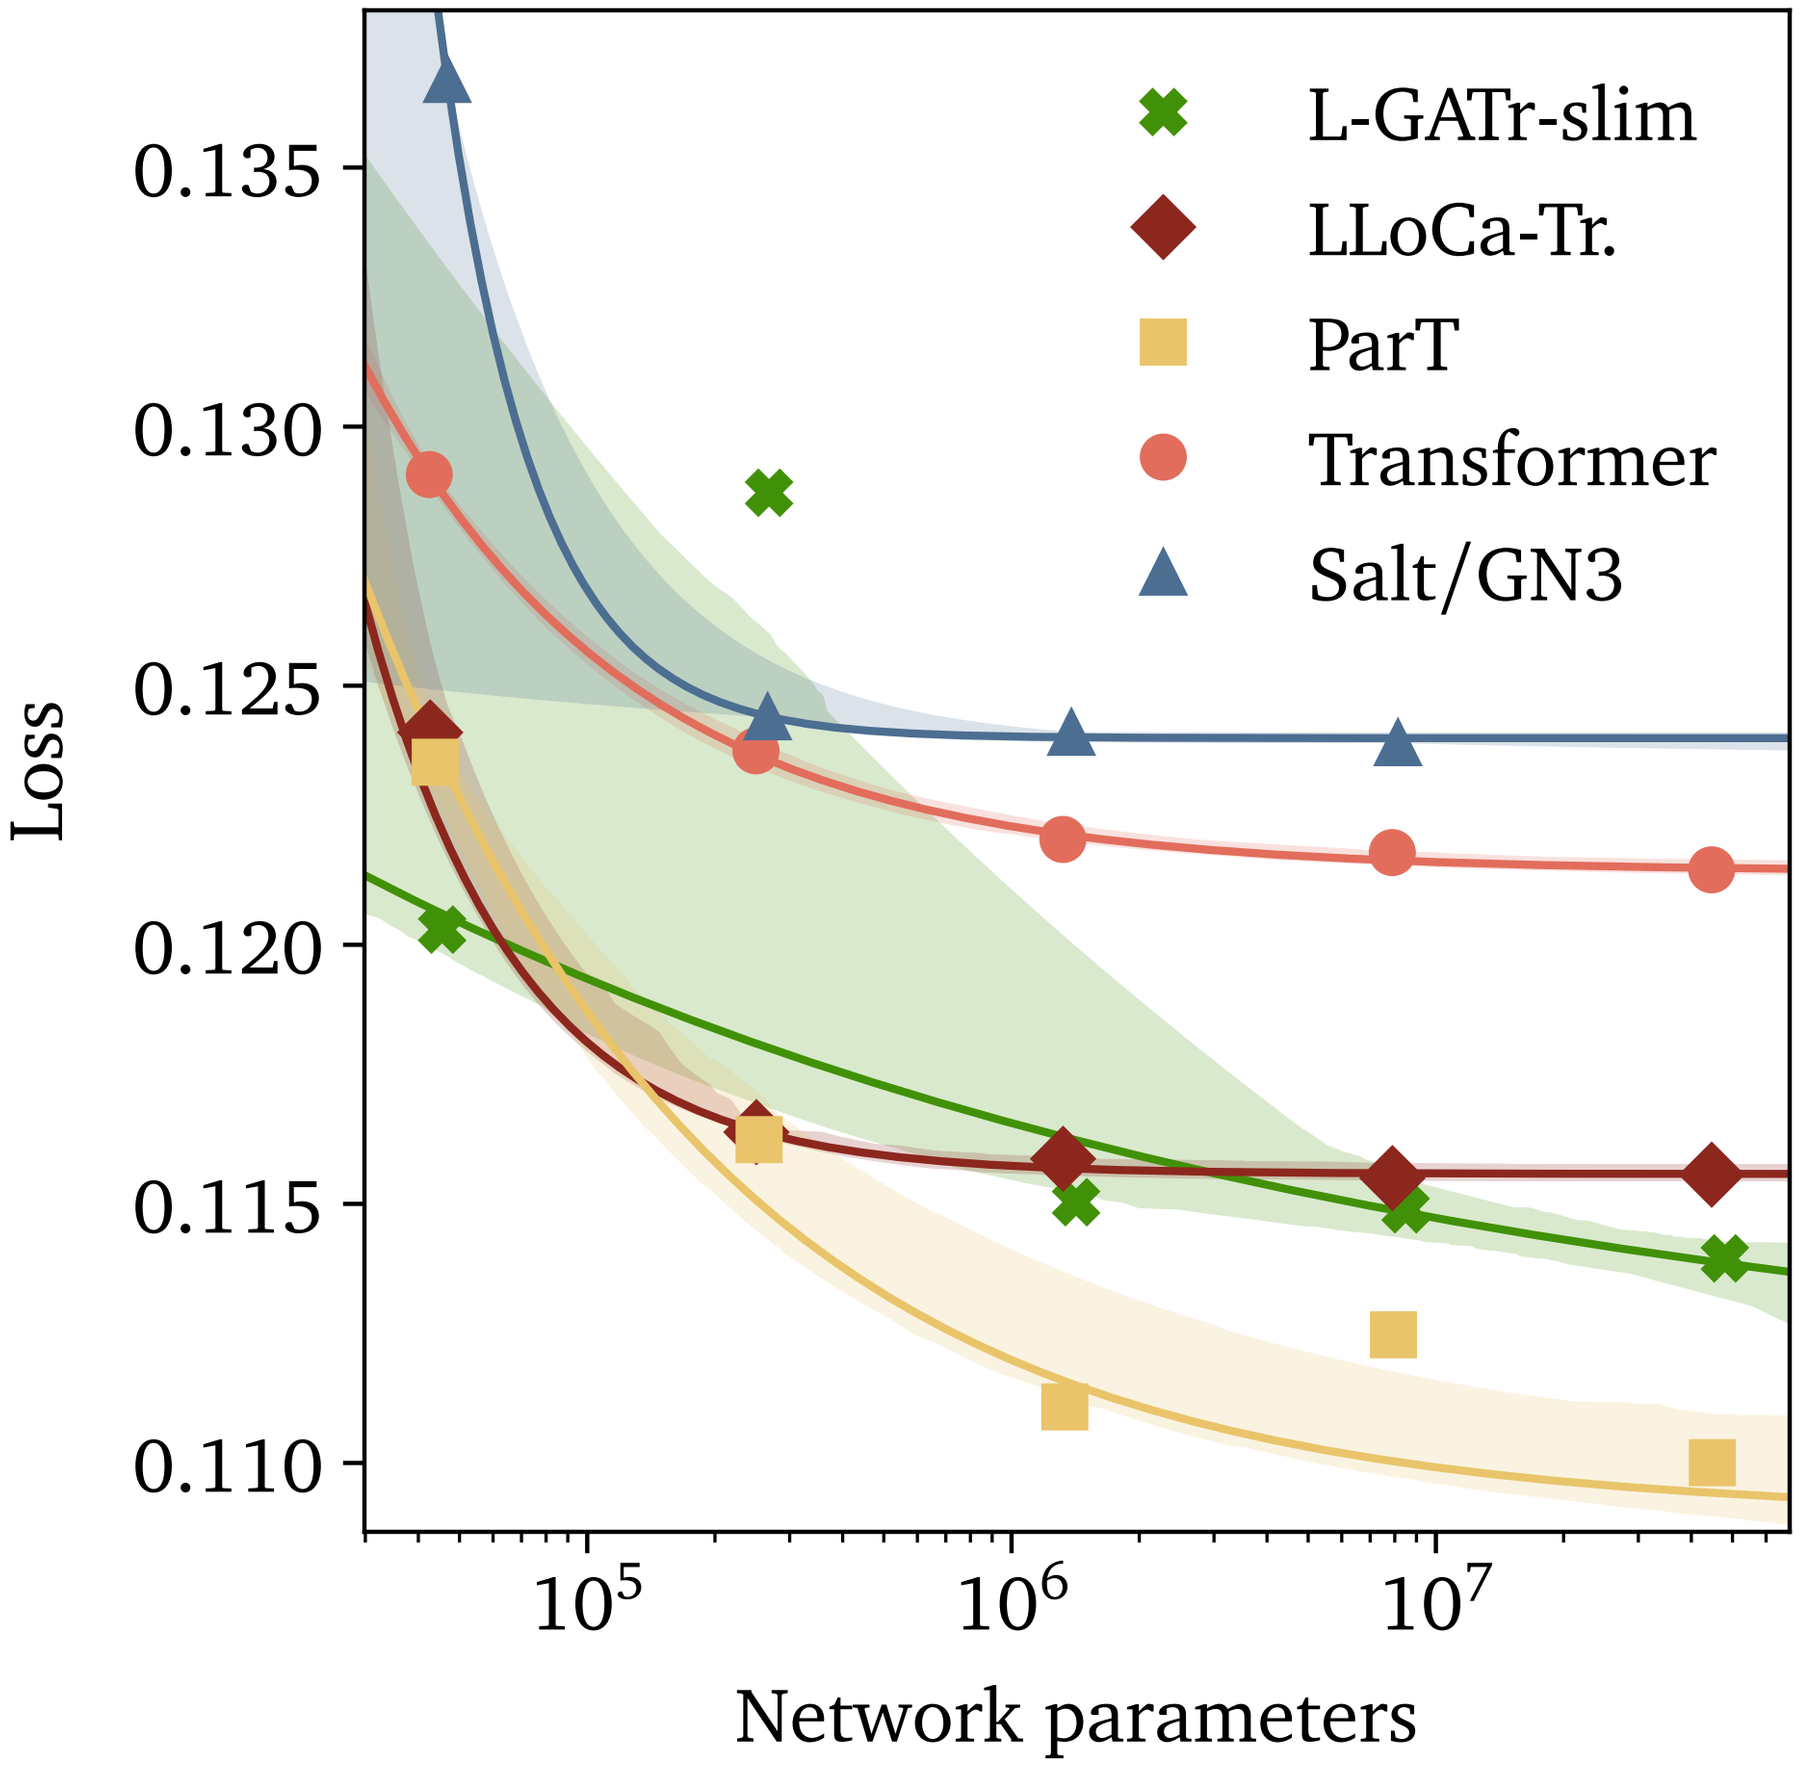
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3
Loss
0.135
0.130
0.125
0.120
0.115
0.110
10^5
10^6
10^7
Network parameters

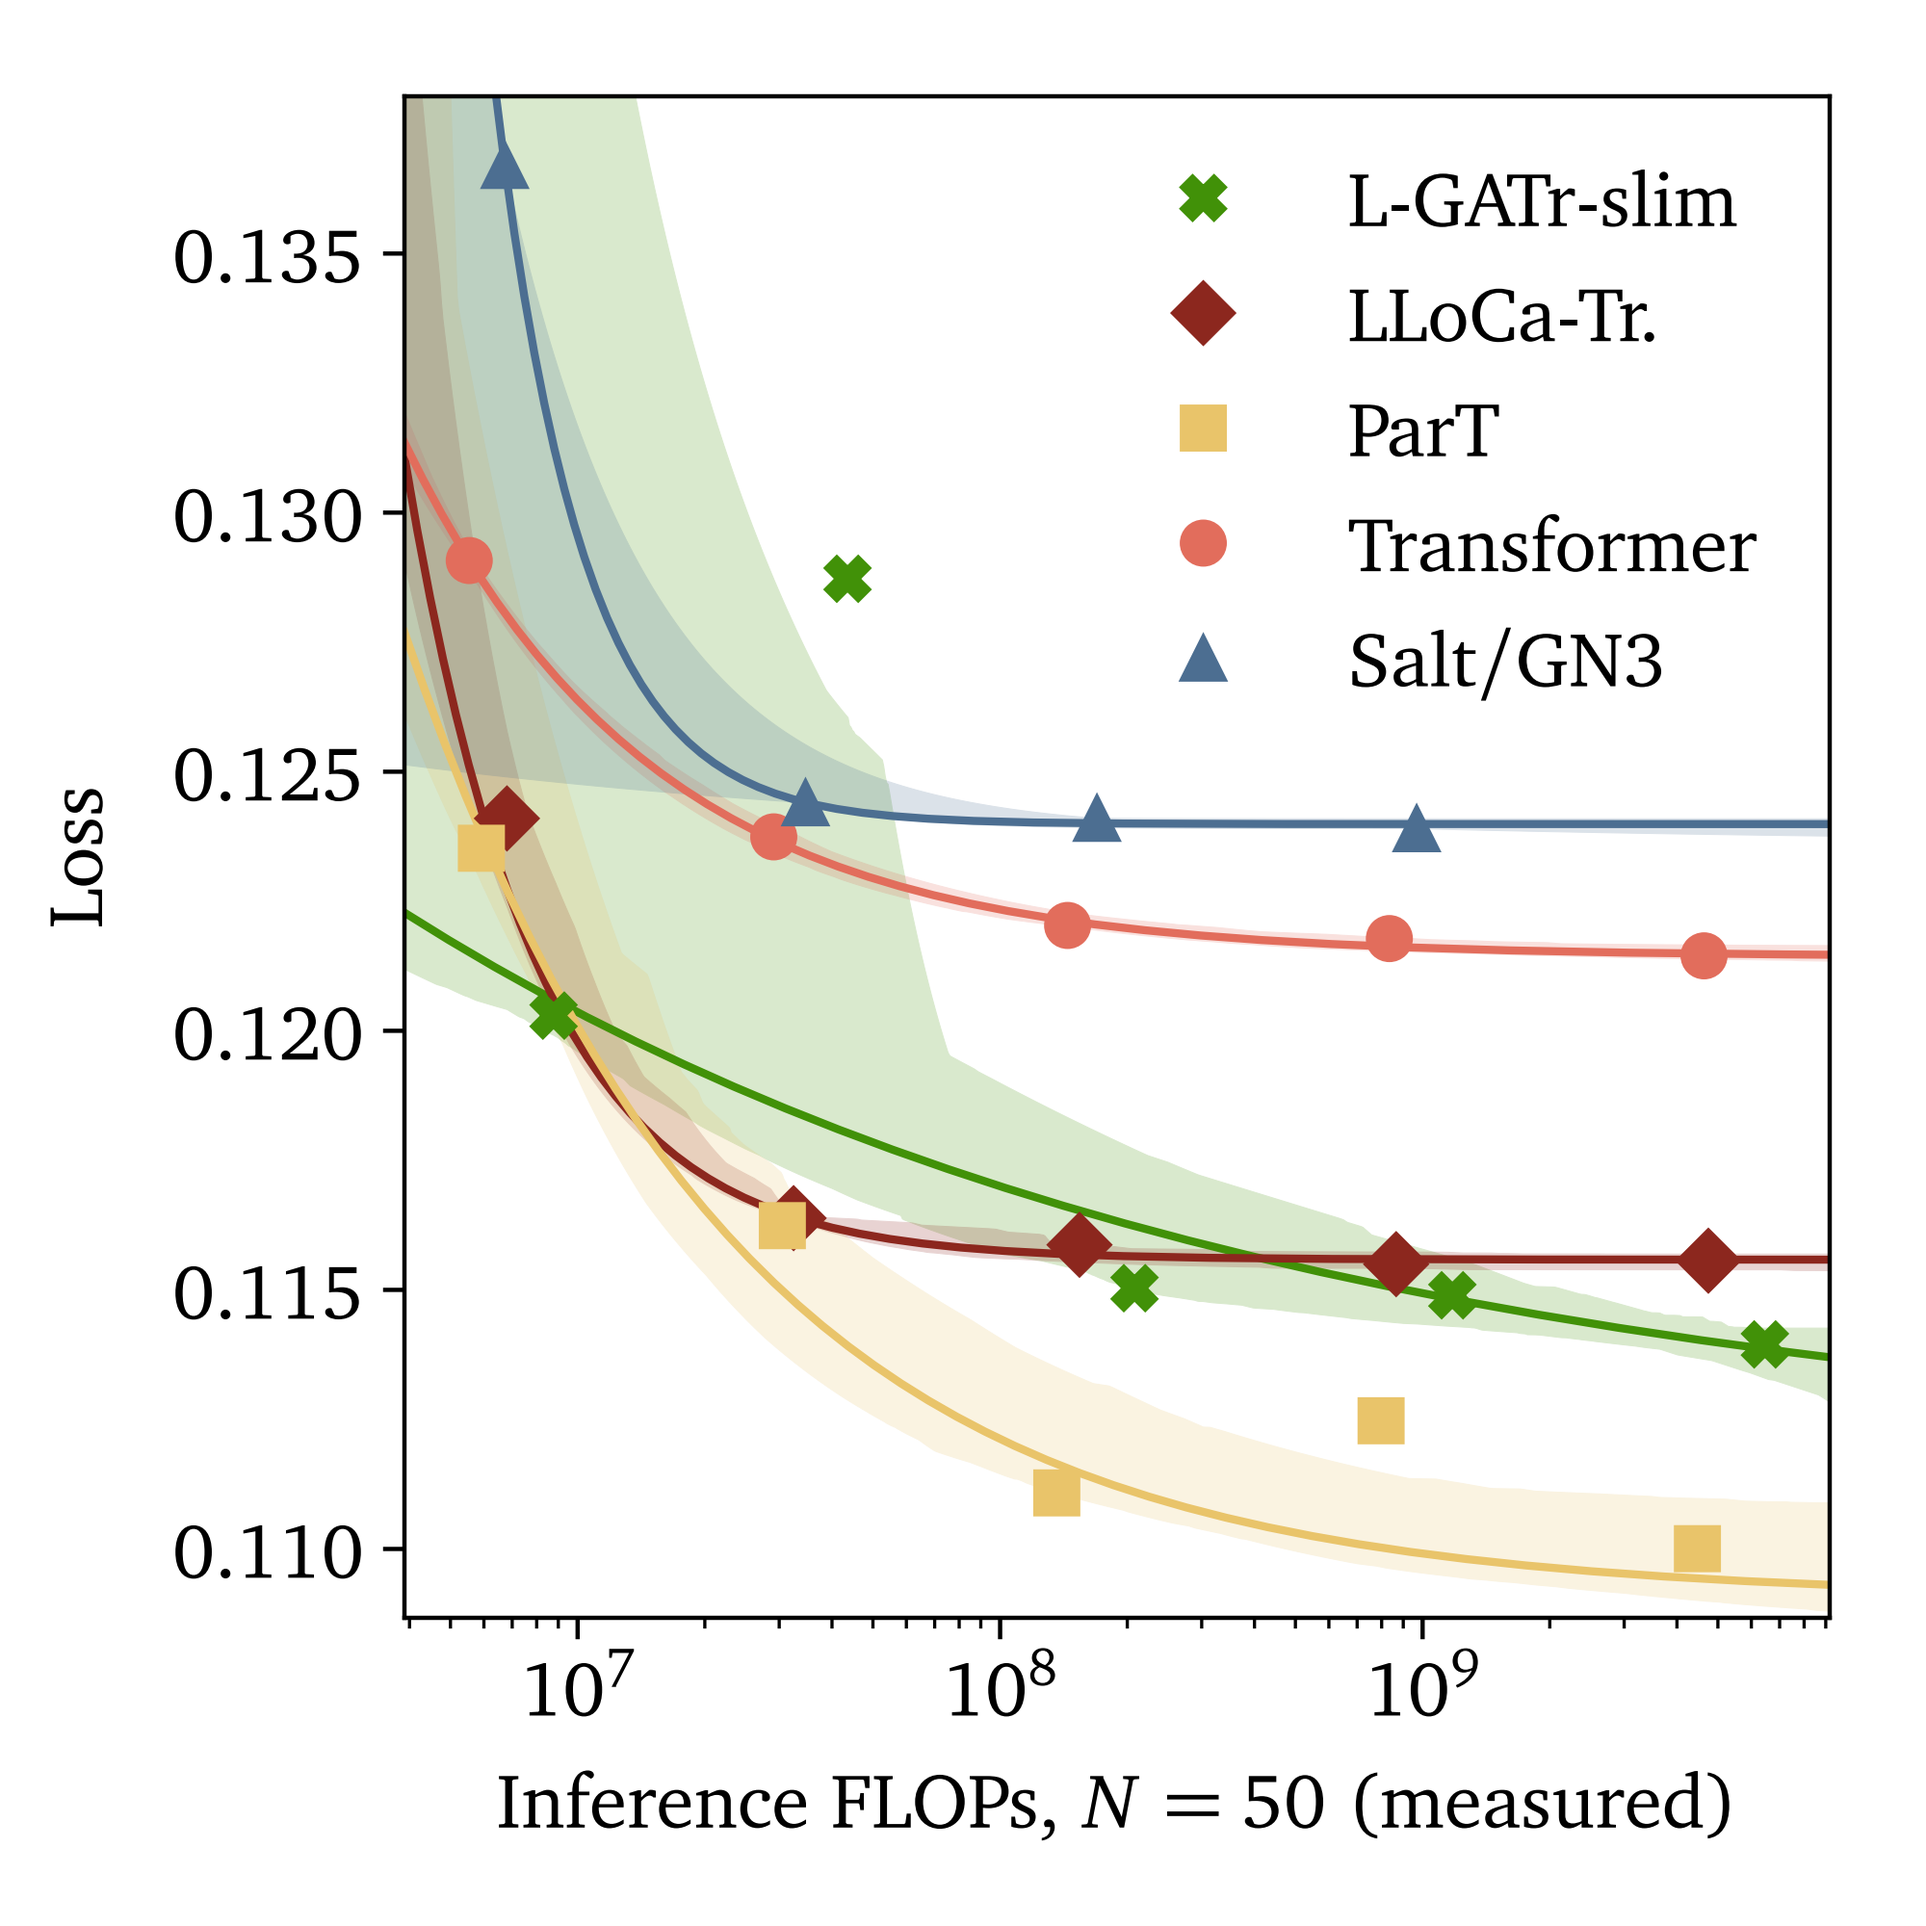
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3
Loss
0.135
0.130
0.125
0.120
0.115
0.110
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (measured)

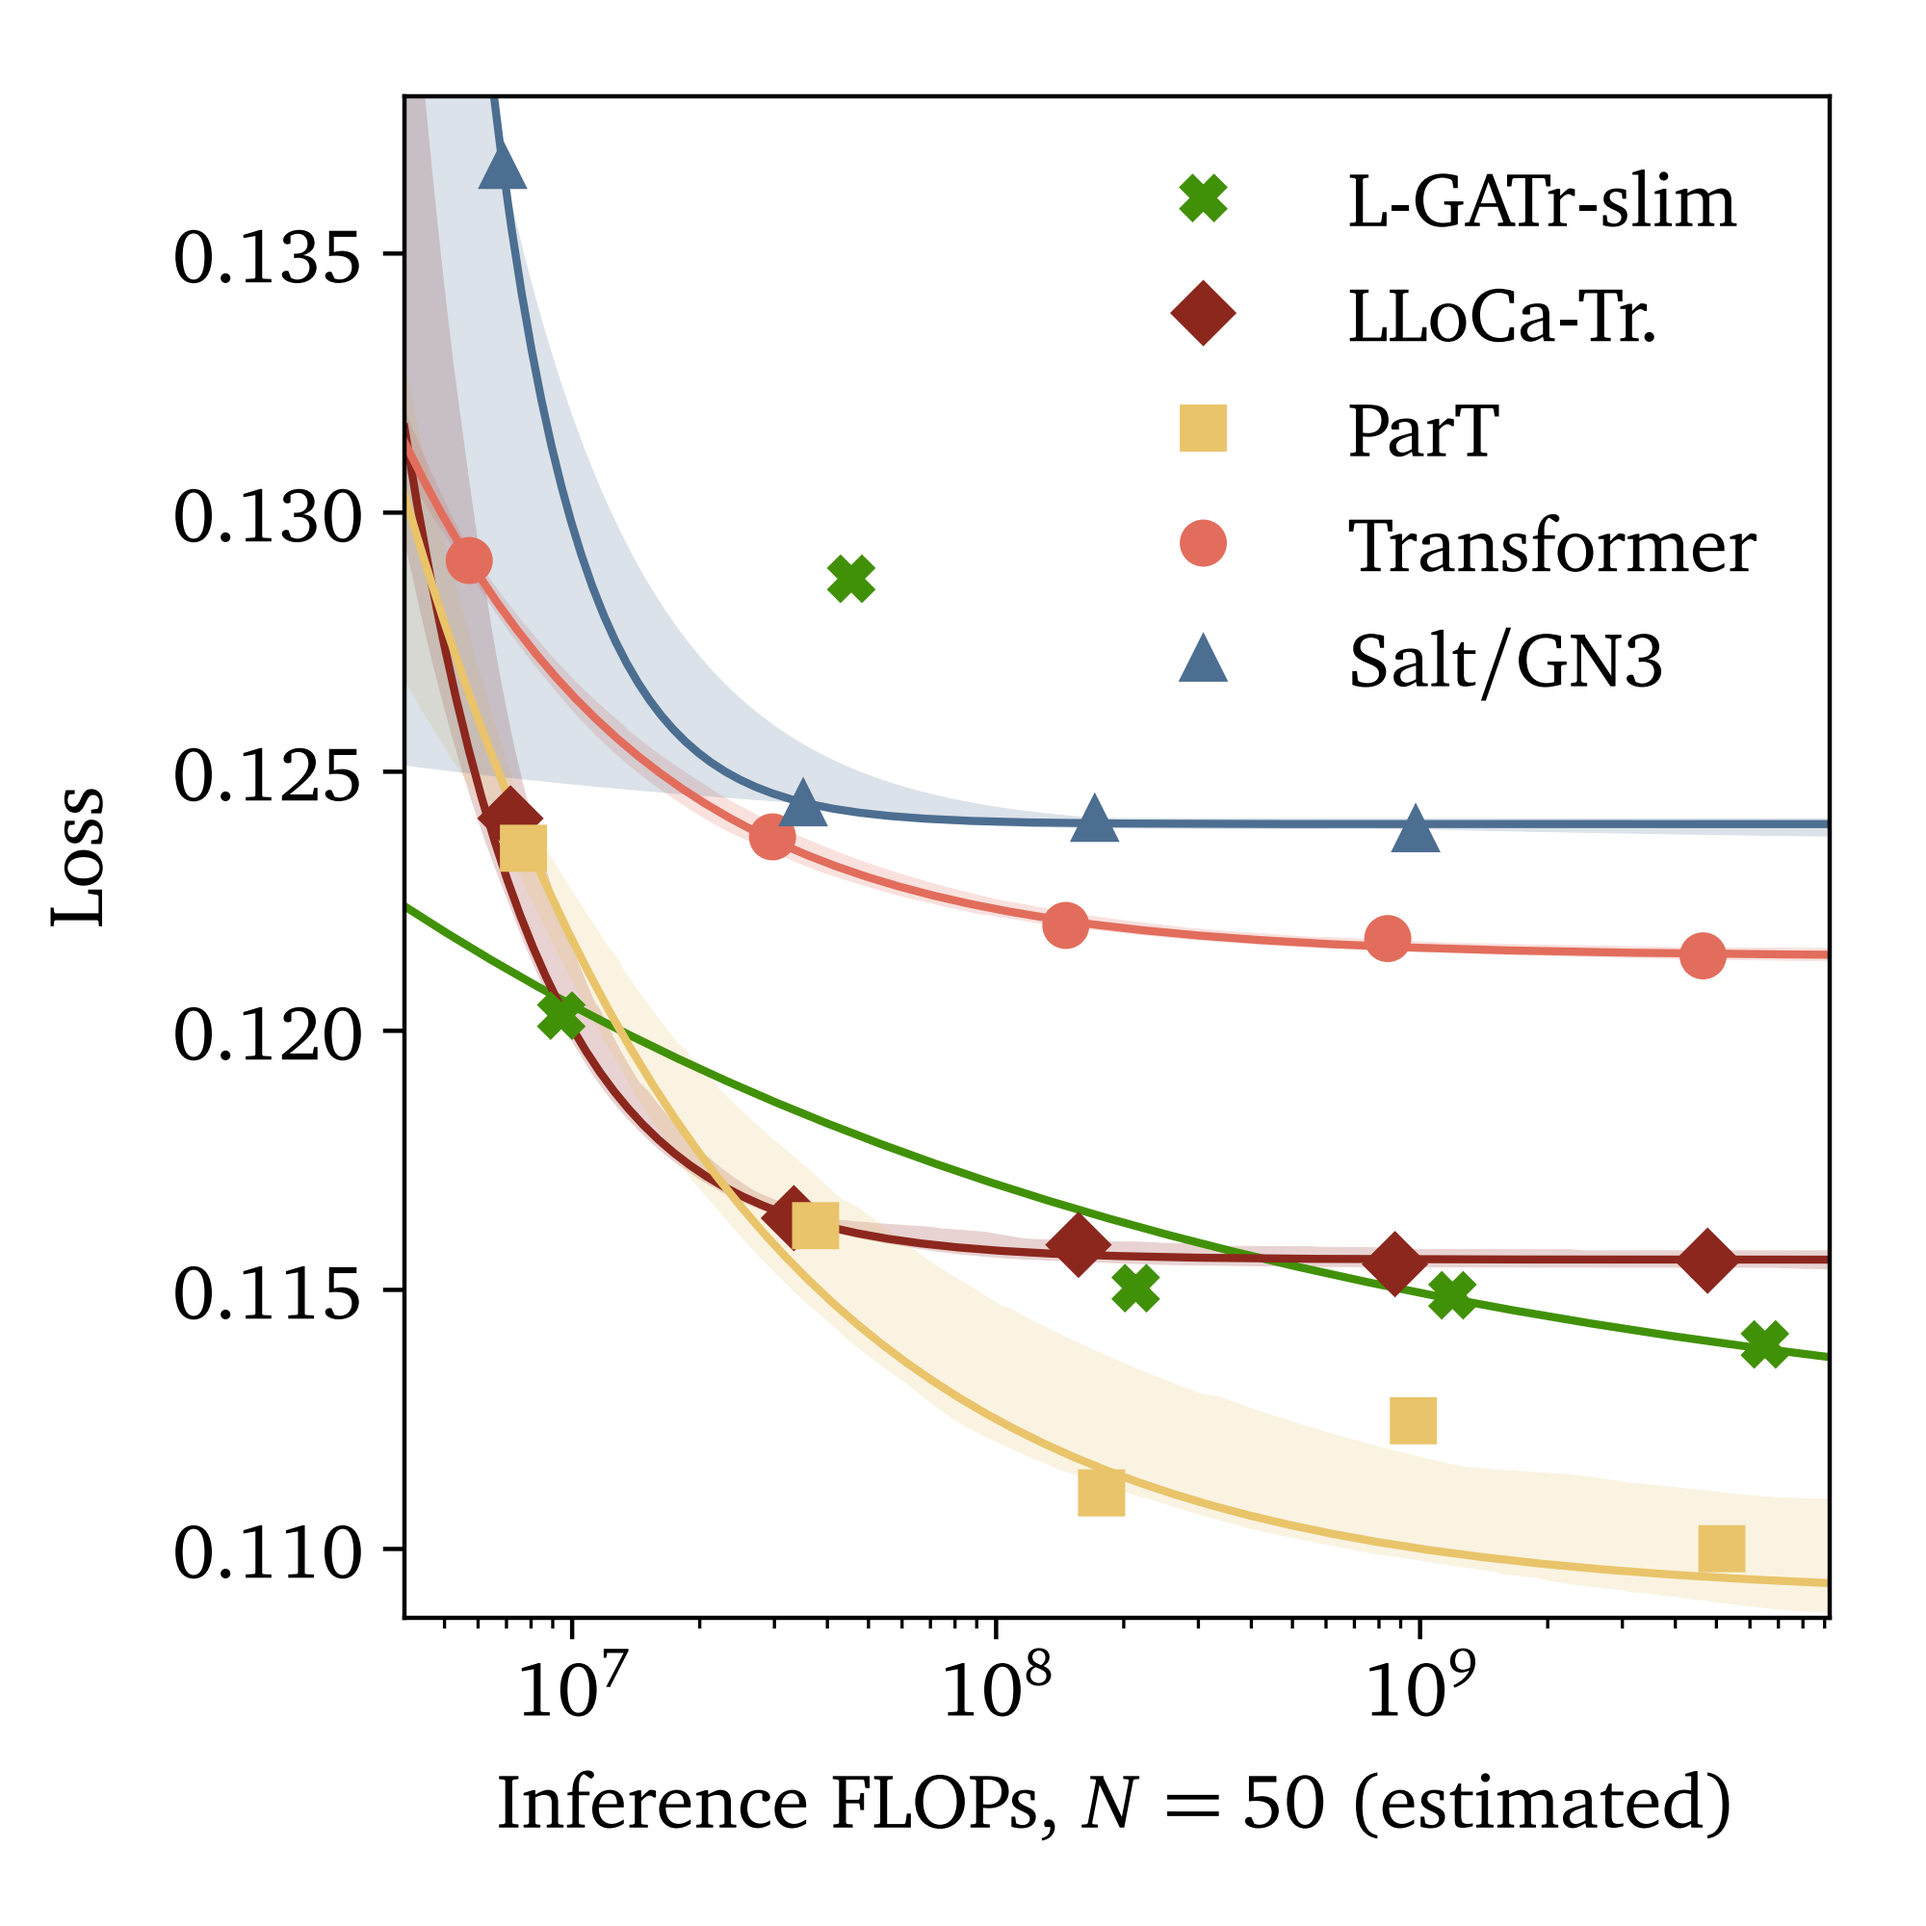
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3
Loss
0.135
0.130
0.125
0.120
0.115
0.110
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (estimated)

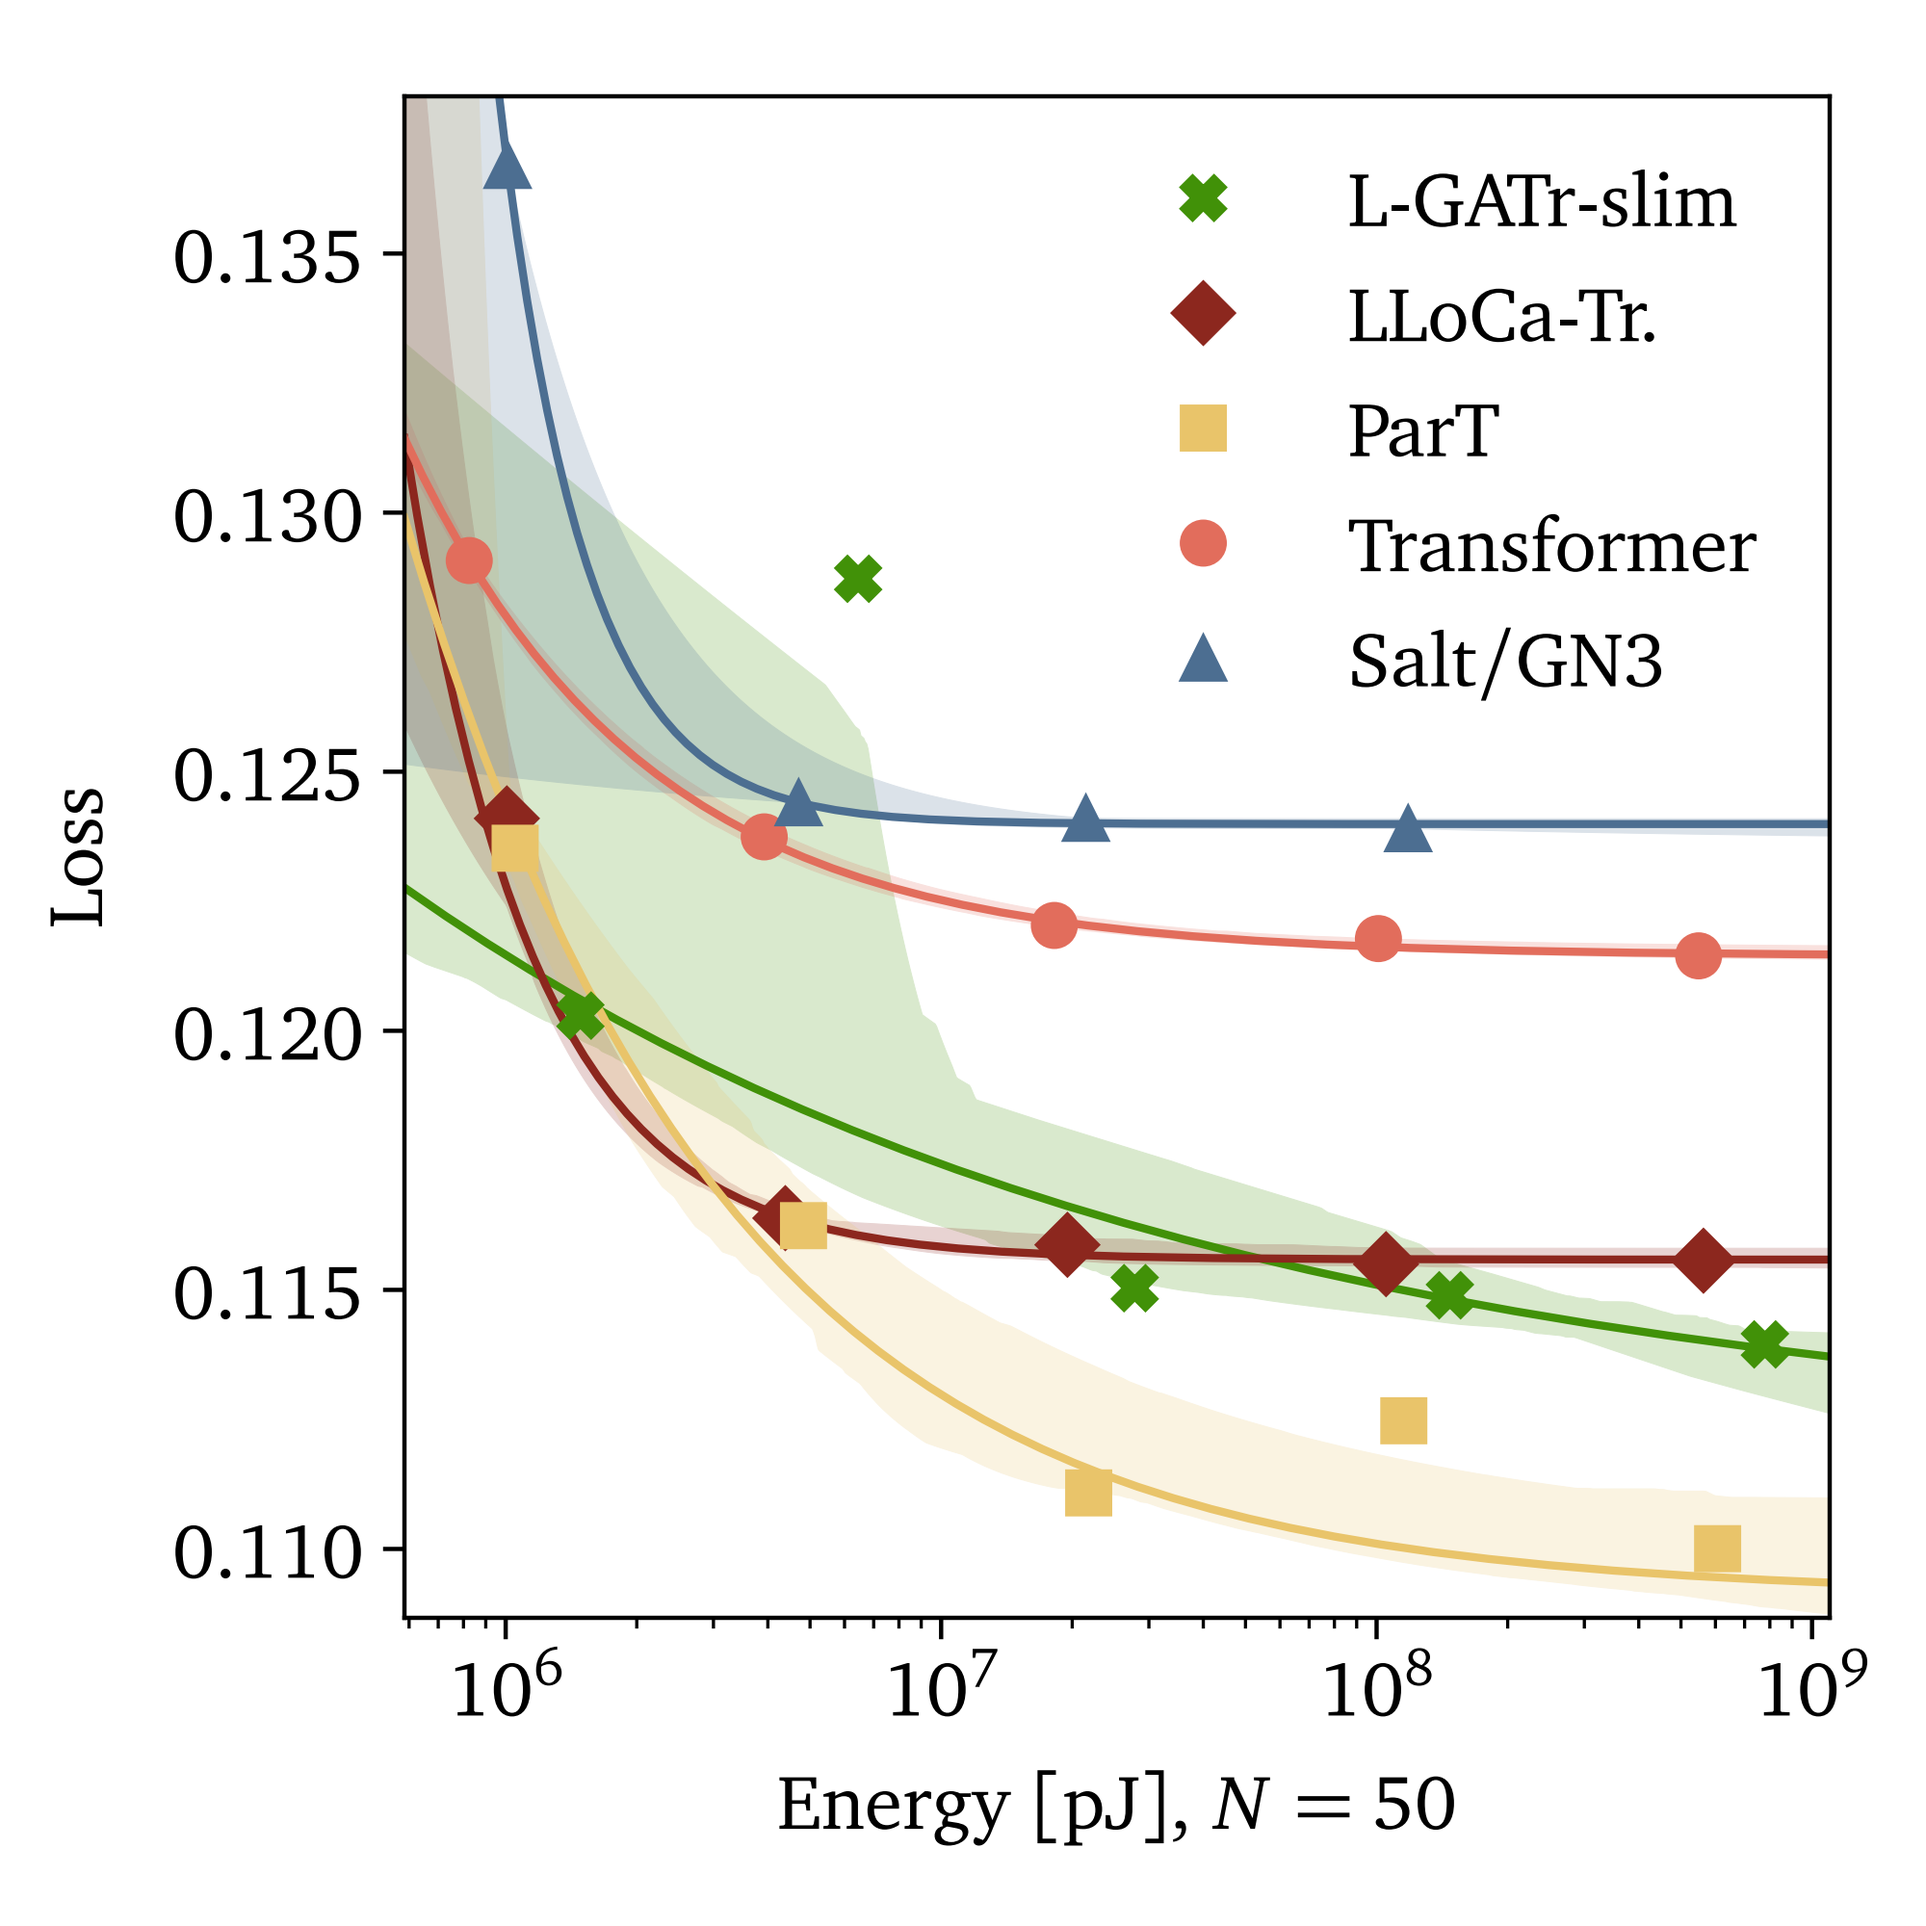
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3
Loss
0.135
0.130
0.125
0.120
0.115
0.110
10^6
10^7
10^8
10^9
Energy [pJ], $N = 50$

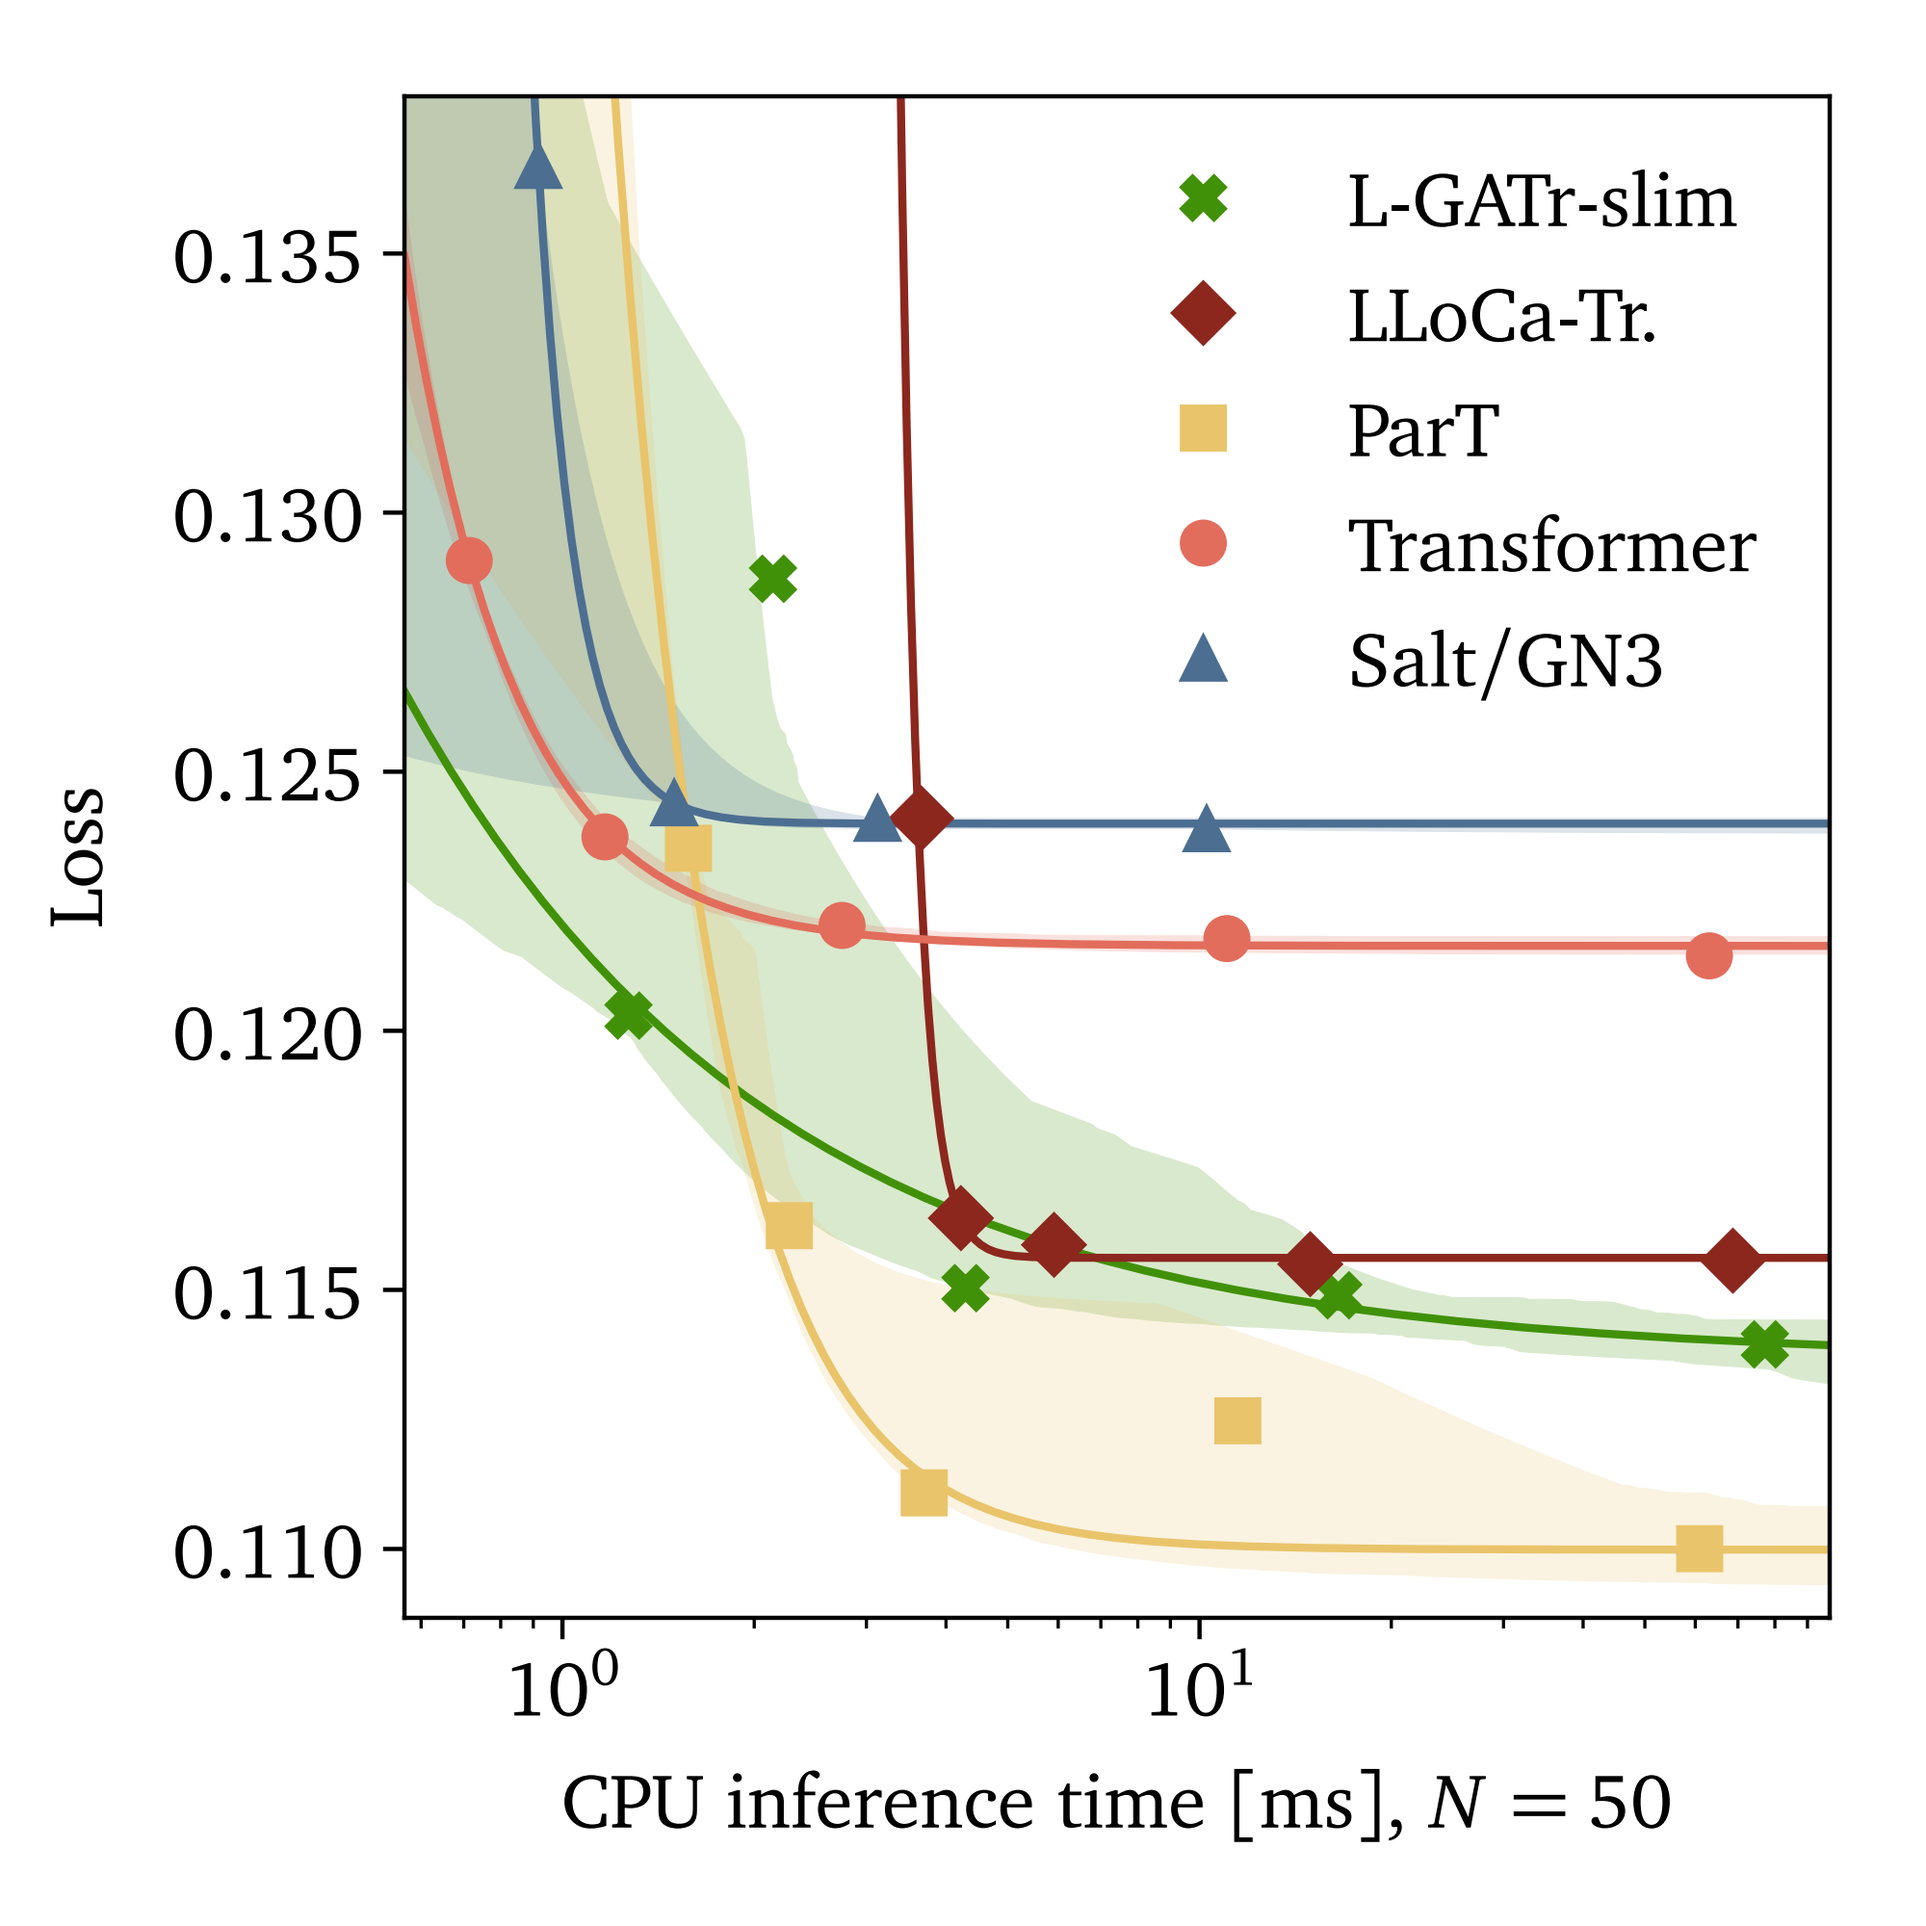
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3
Loss
0.135
0.130
0.125
0.120
0.115
0.110
10^0
10^1
CPU inference time [ms], $N = 50$

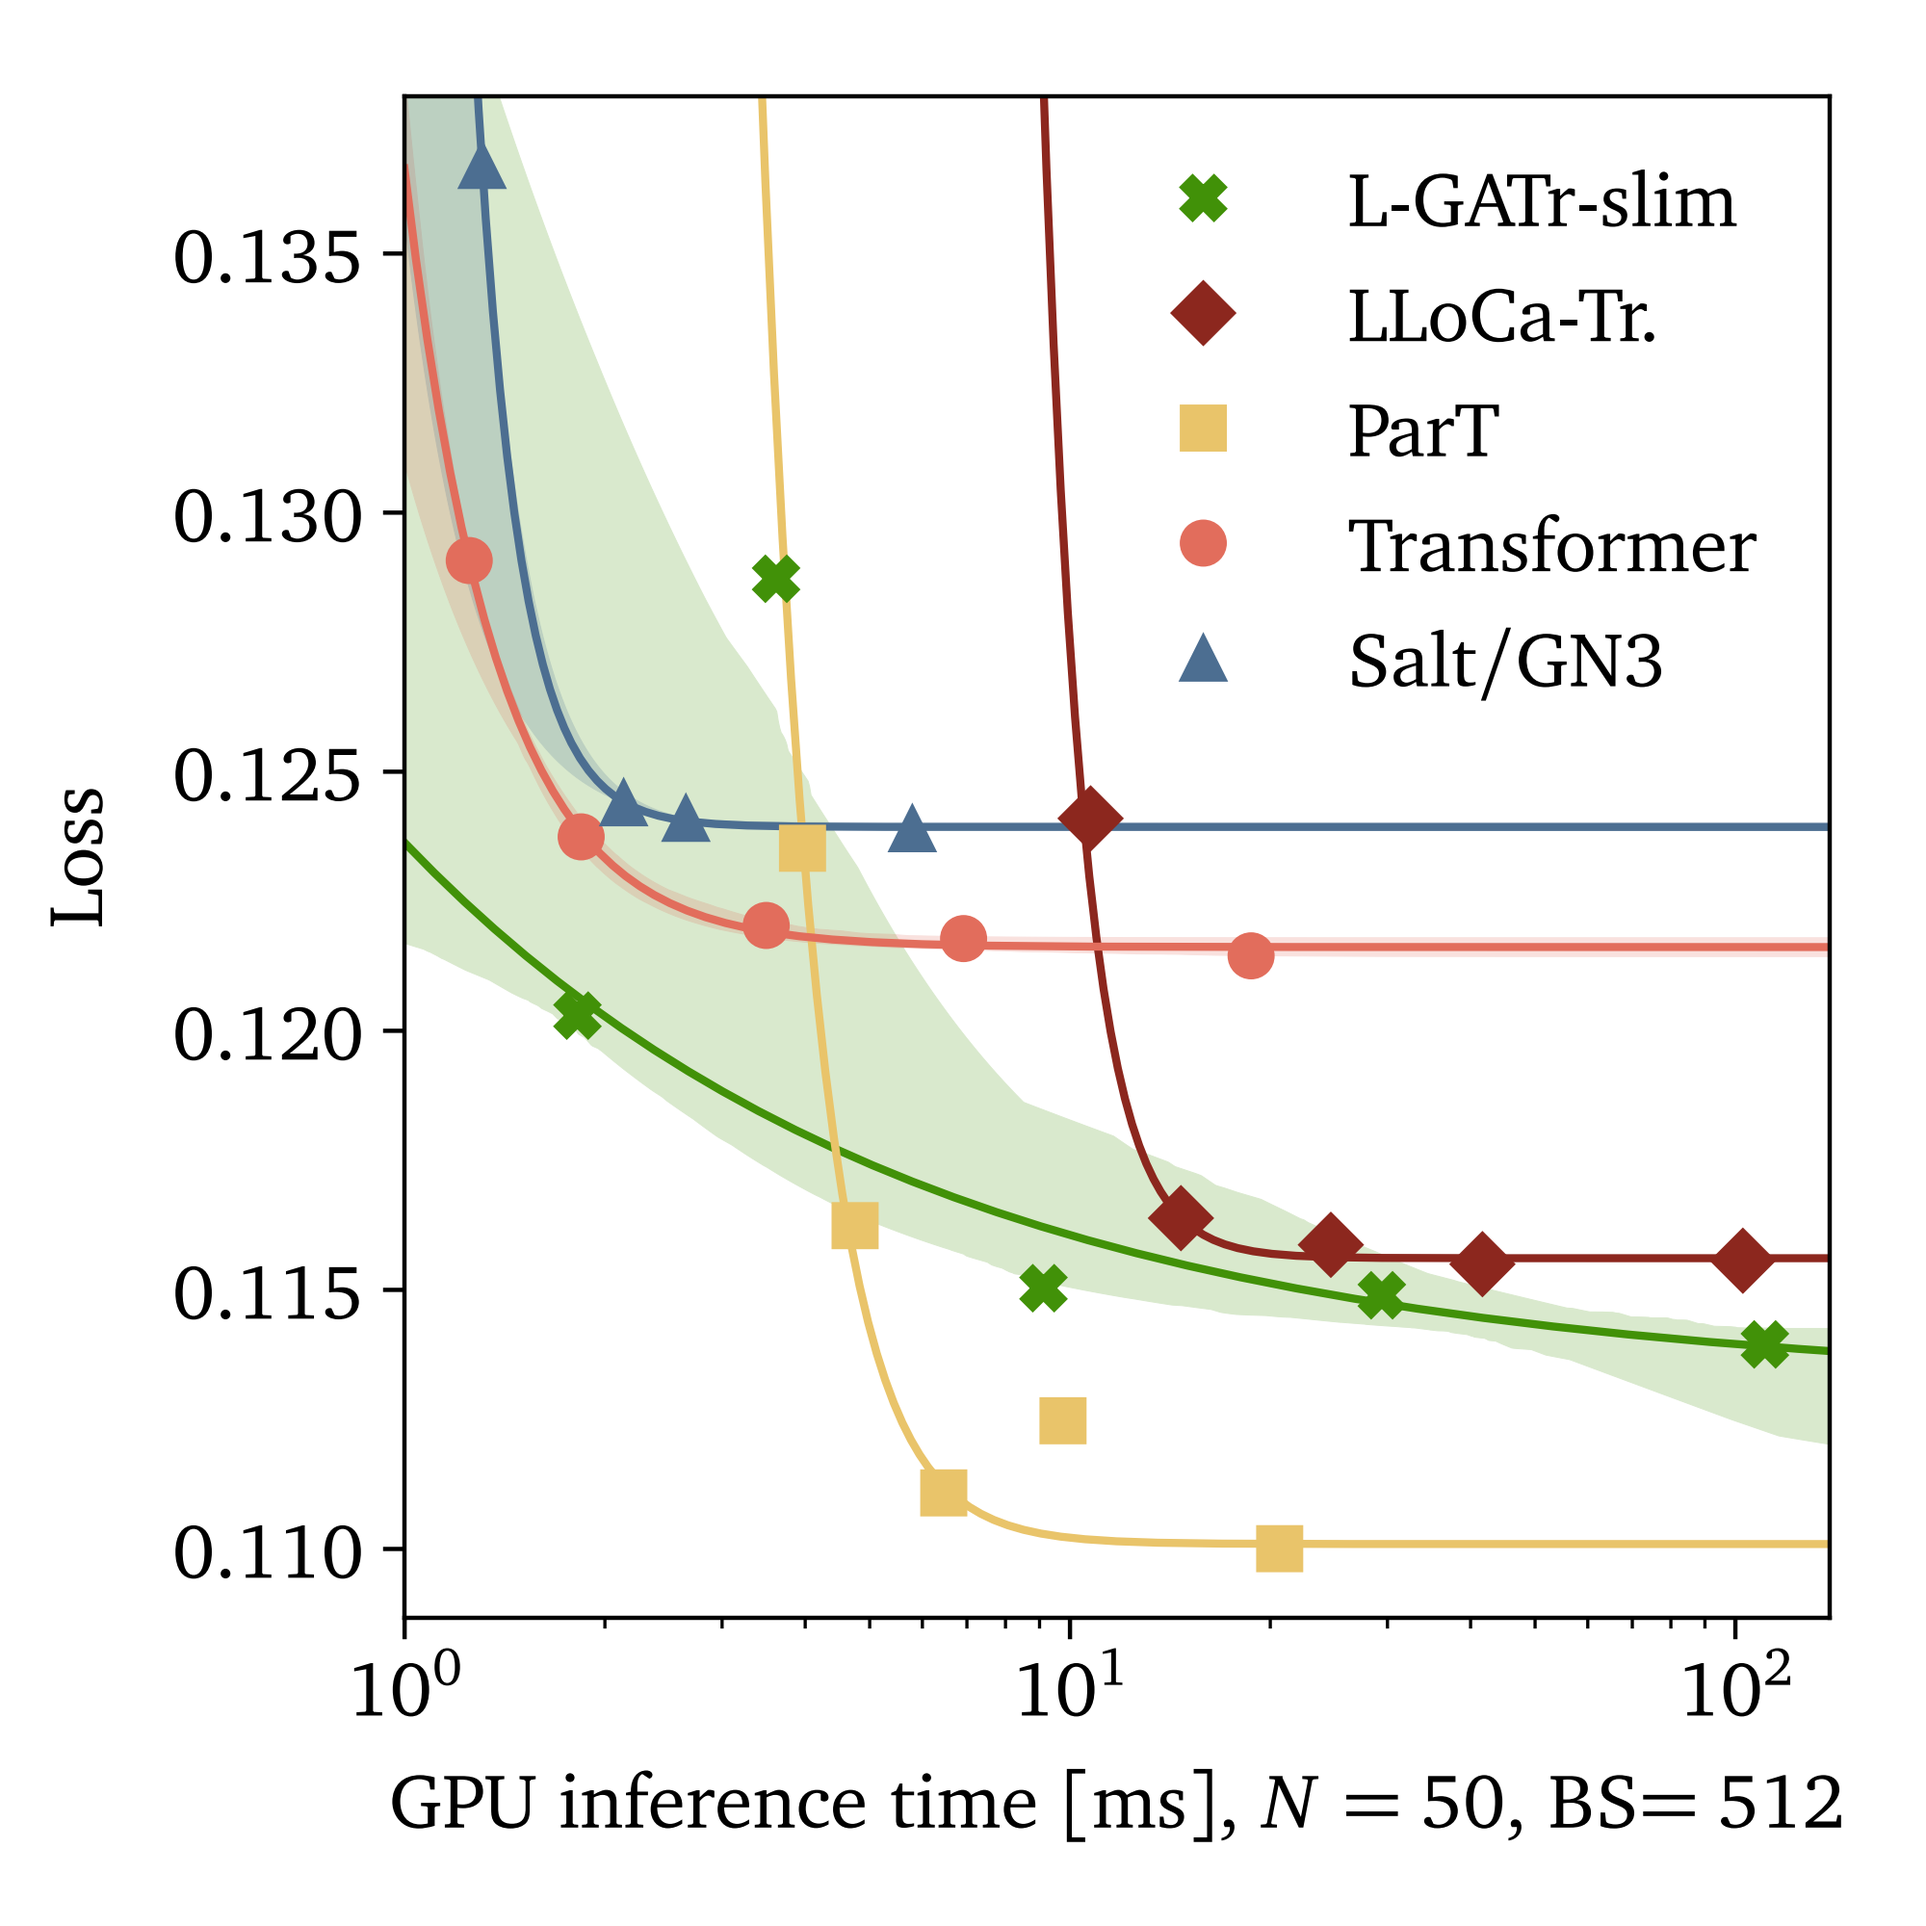
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3
Loss
0.135
0.130
0.125
0.120
0.115
0.110
10^0
10^1
10^2
GPU inference time [ms], $N = 50$, BS$= 512$

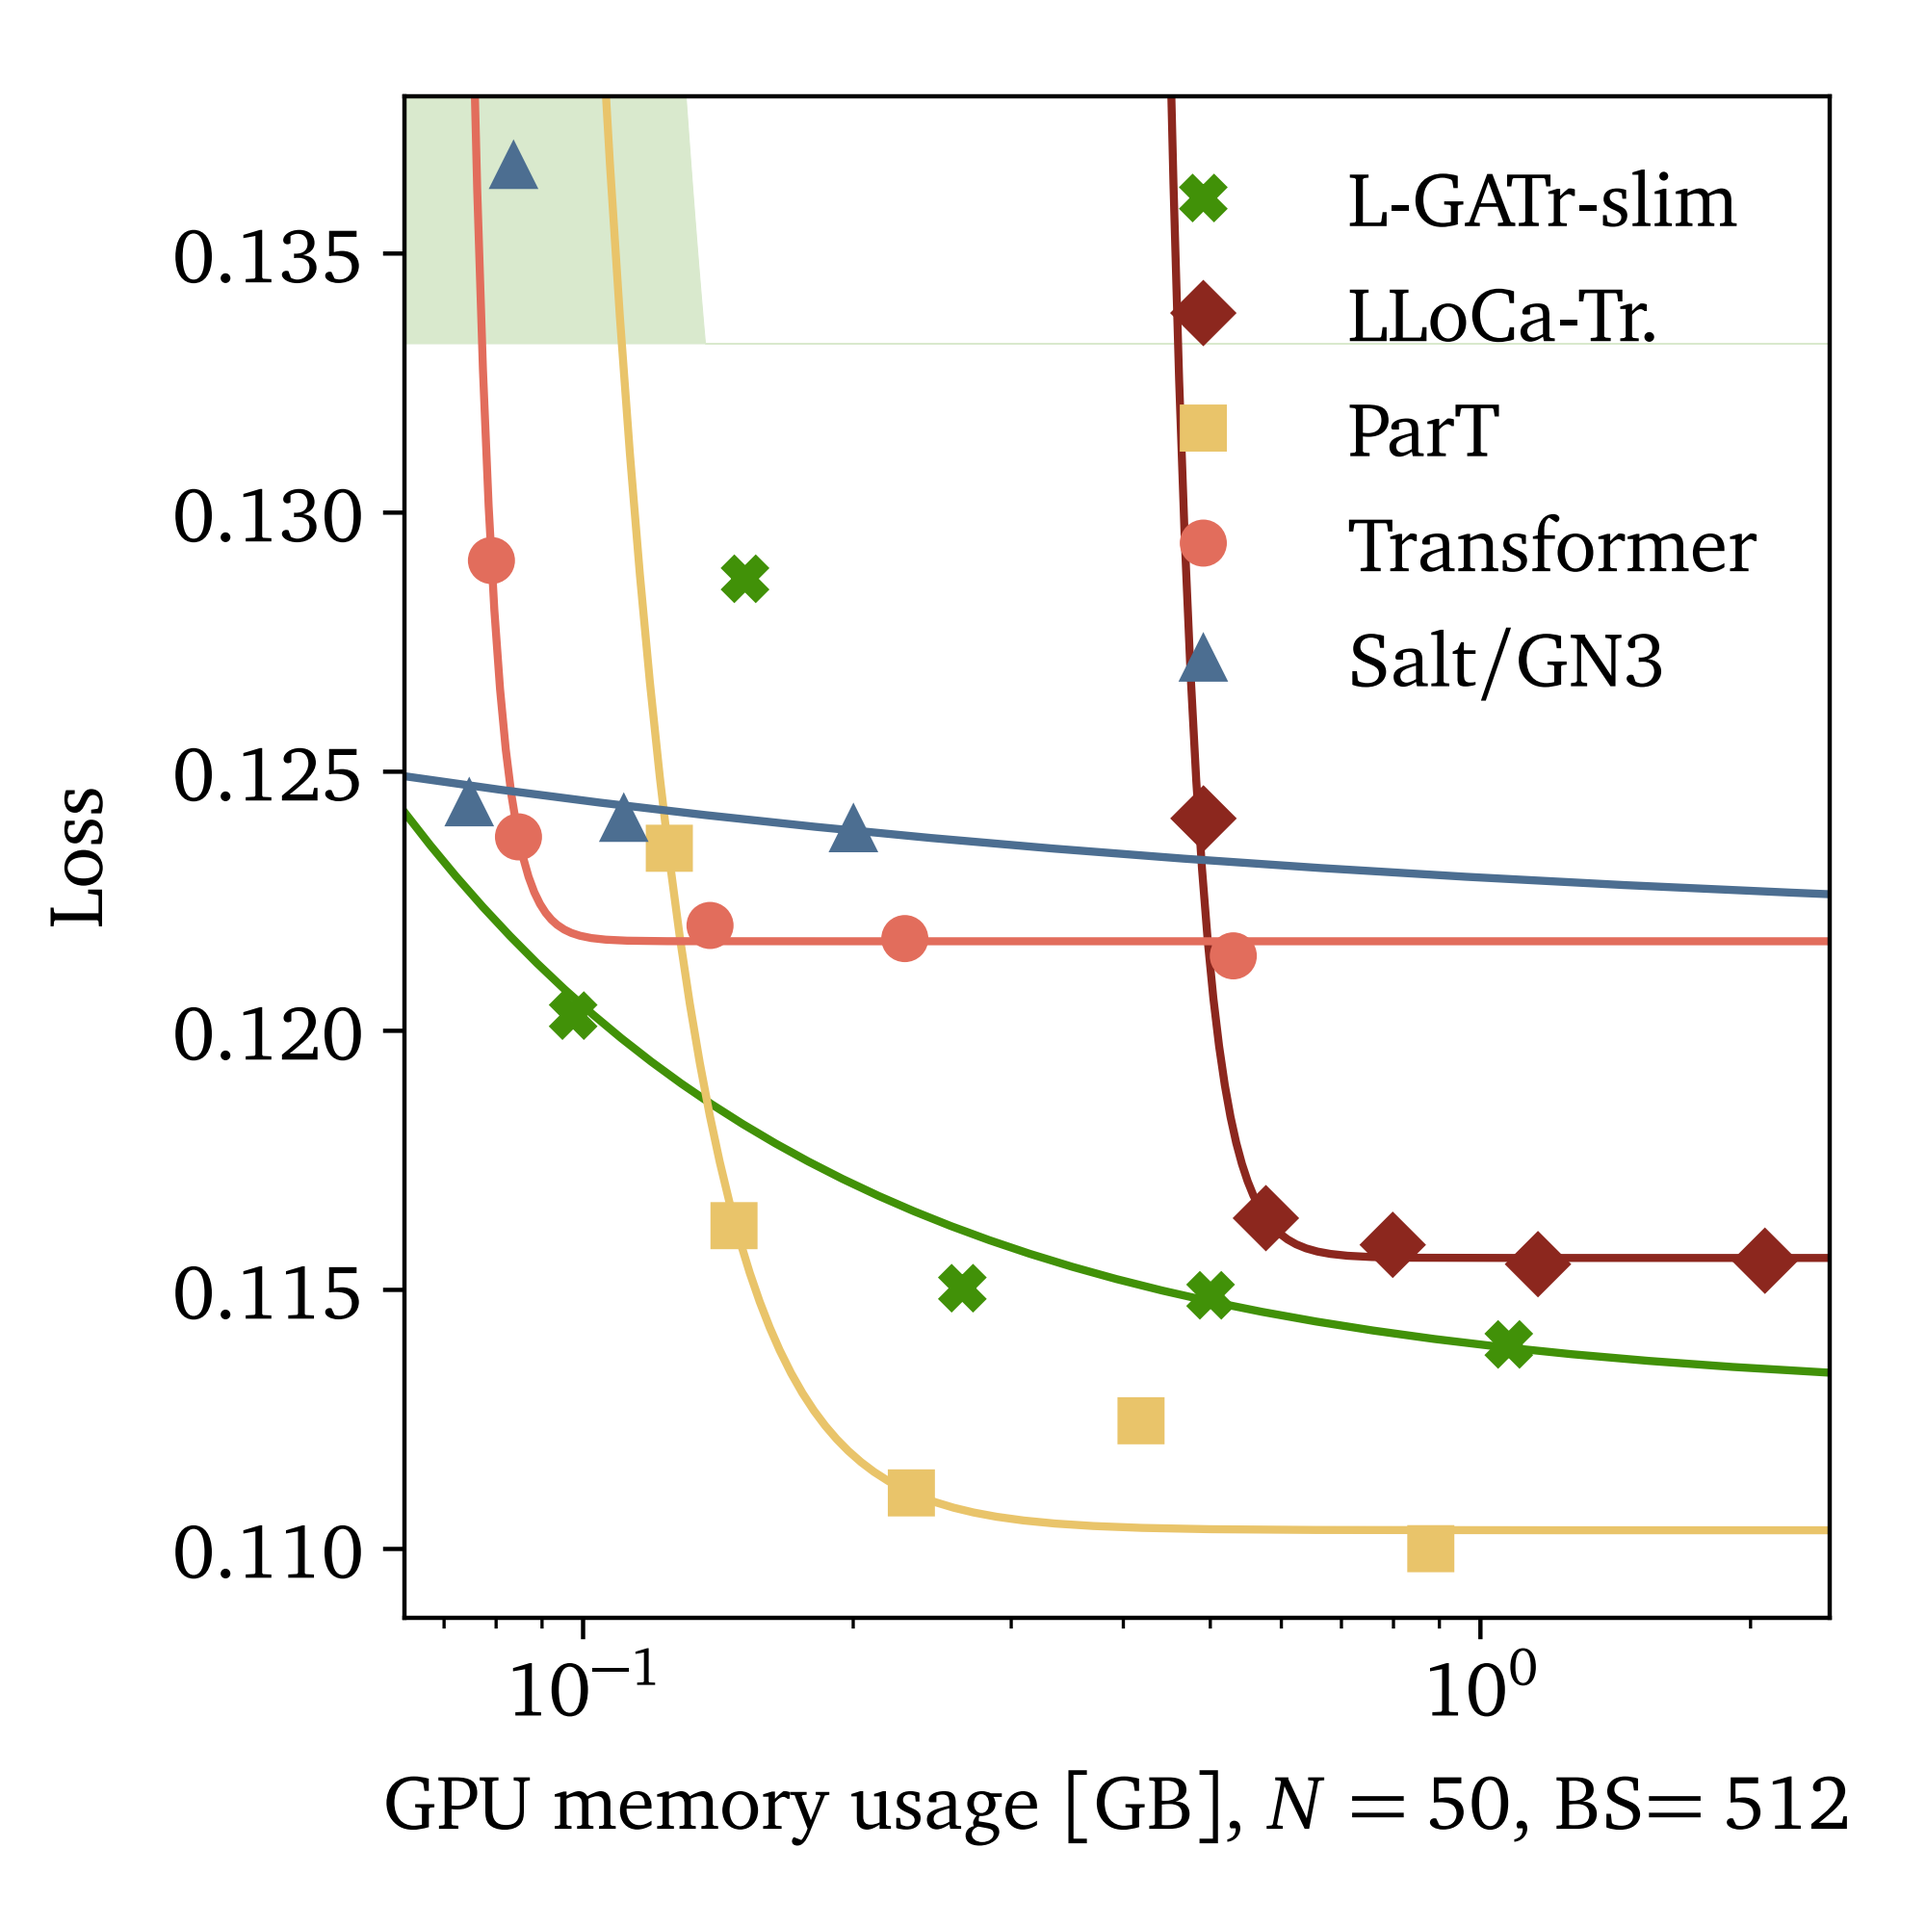
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3
Loss
0.135
0.130
0.125
0.120
0.115
0.110
10^{-1}
10^{0}
GPU memory usage [GB], $N = 50$, BS= 512

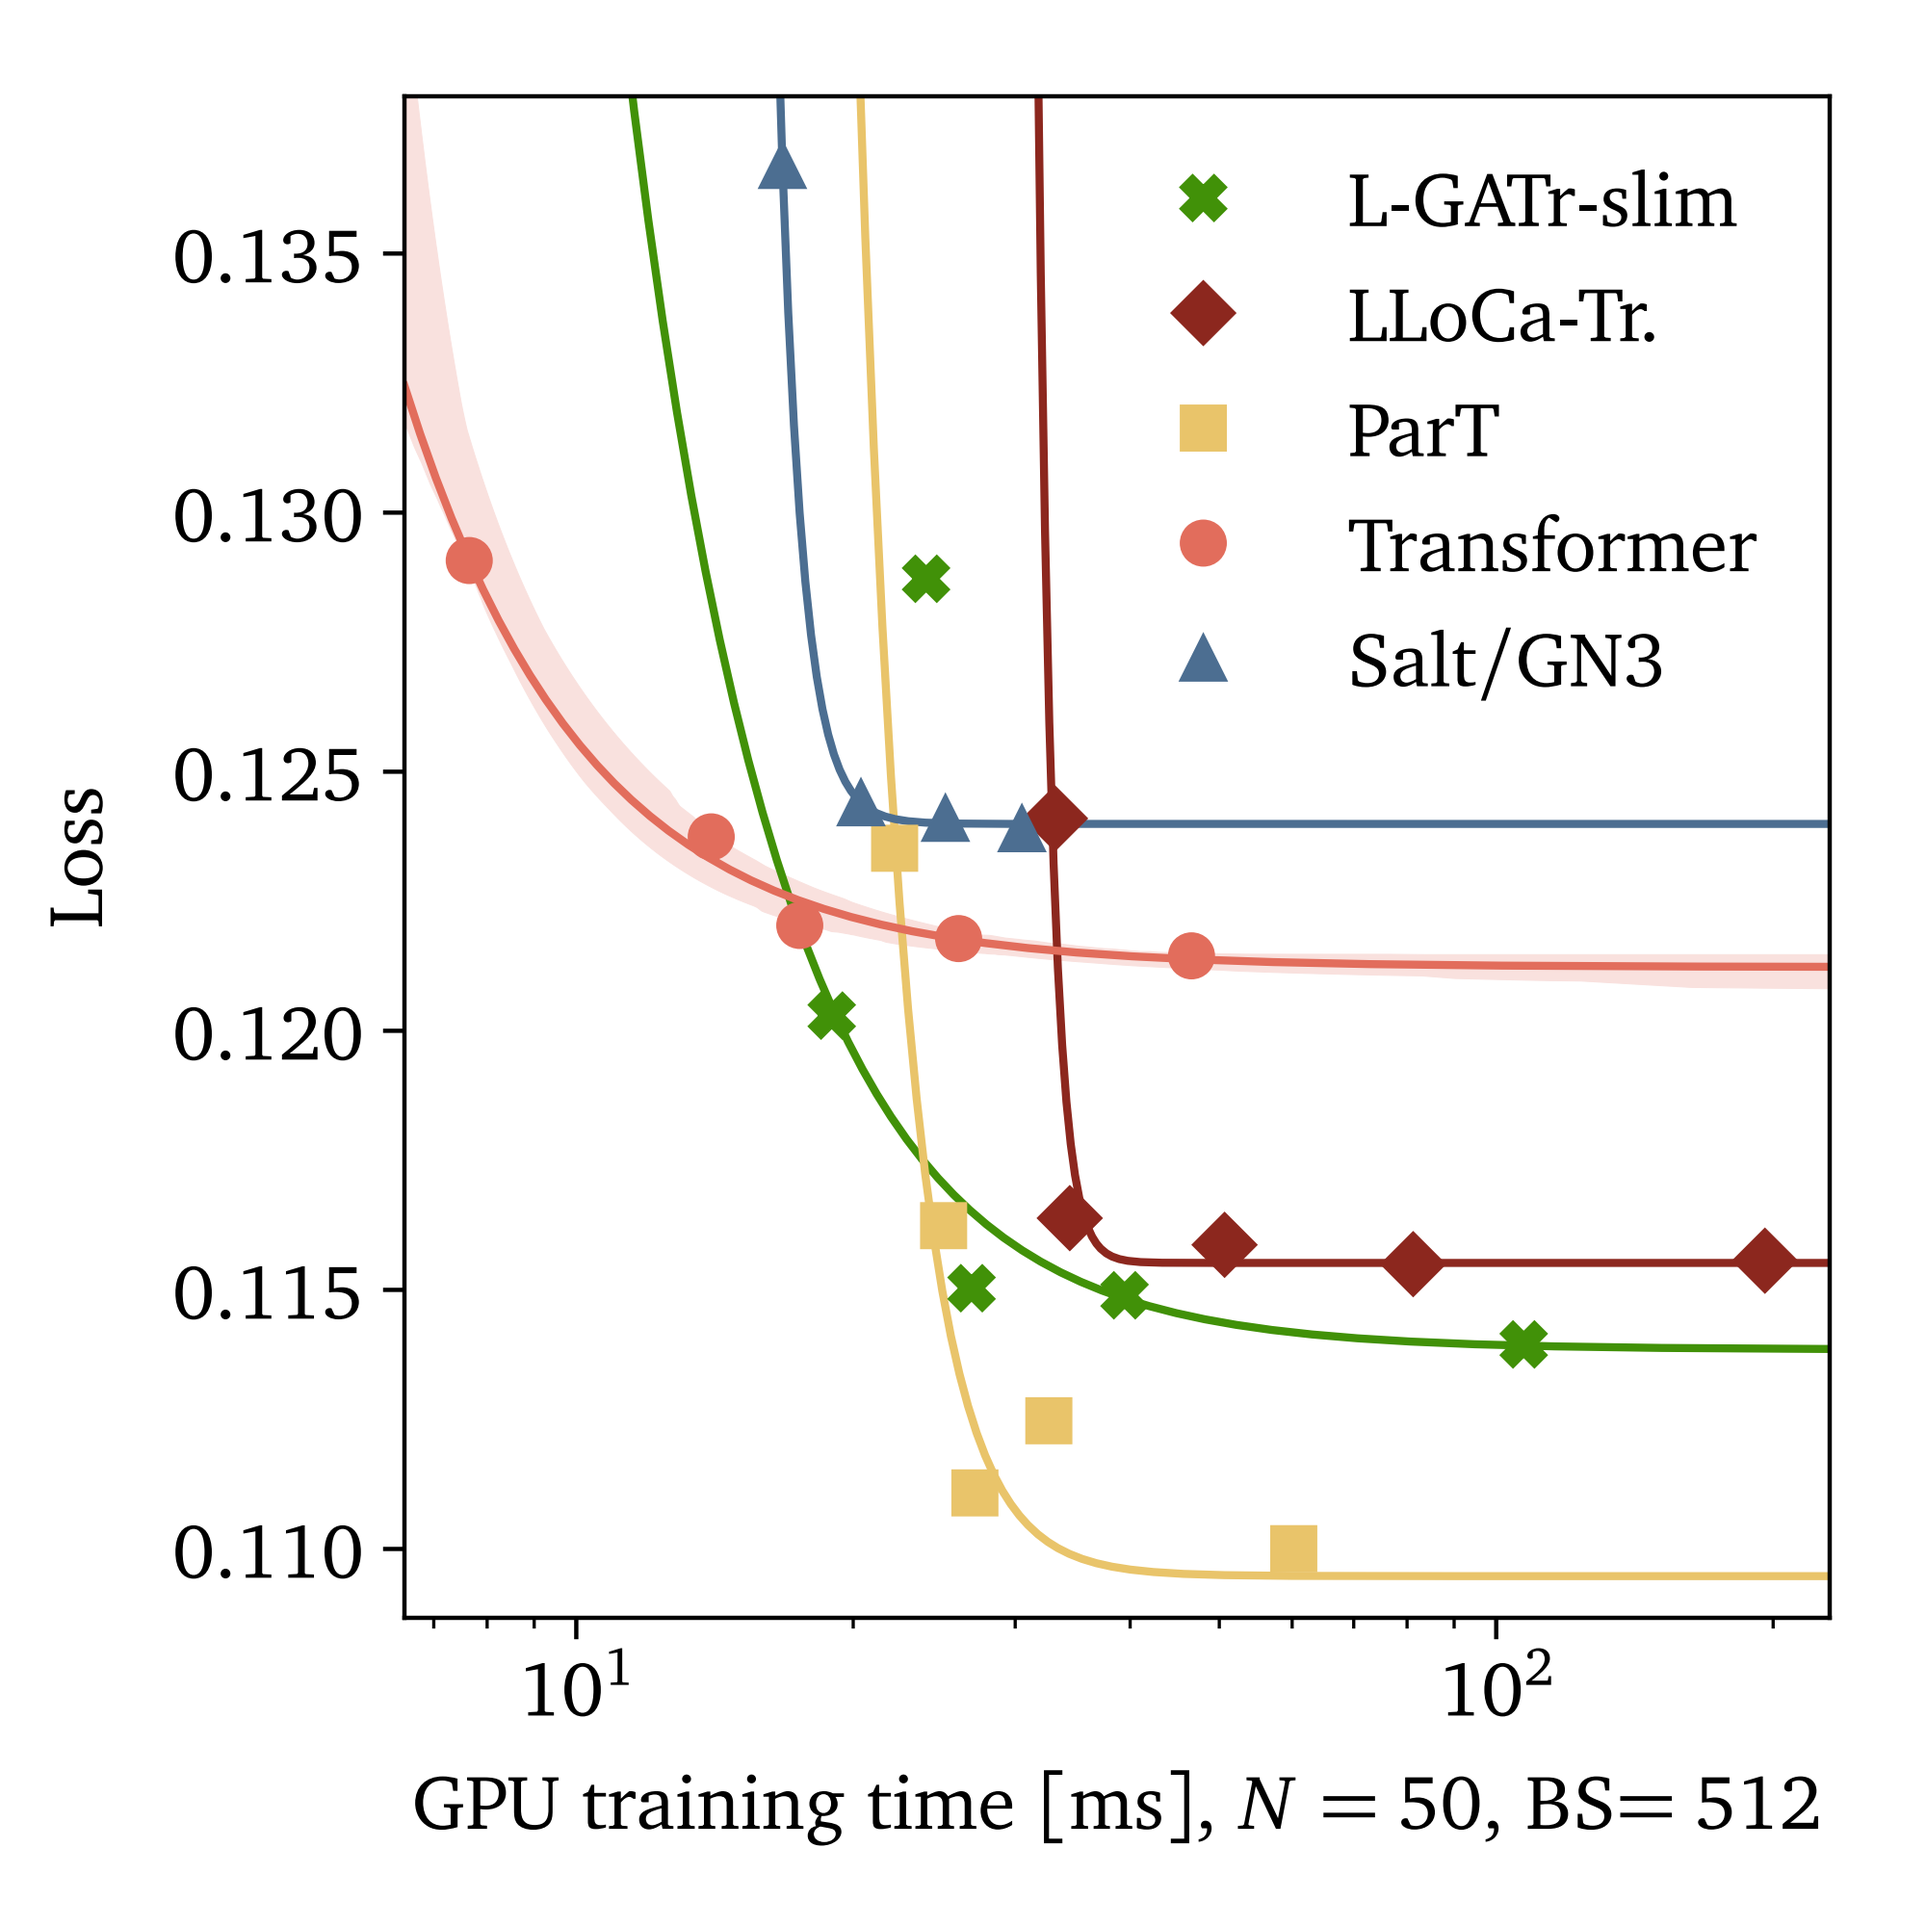
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3
Loss
0.135
0.130
0.125
0.120
0.115
0.110
10^1
10^2
GPU training time [ms], $N = 50$, BS$= 512$

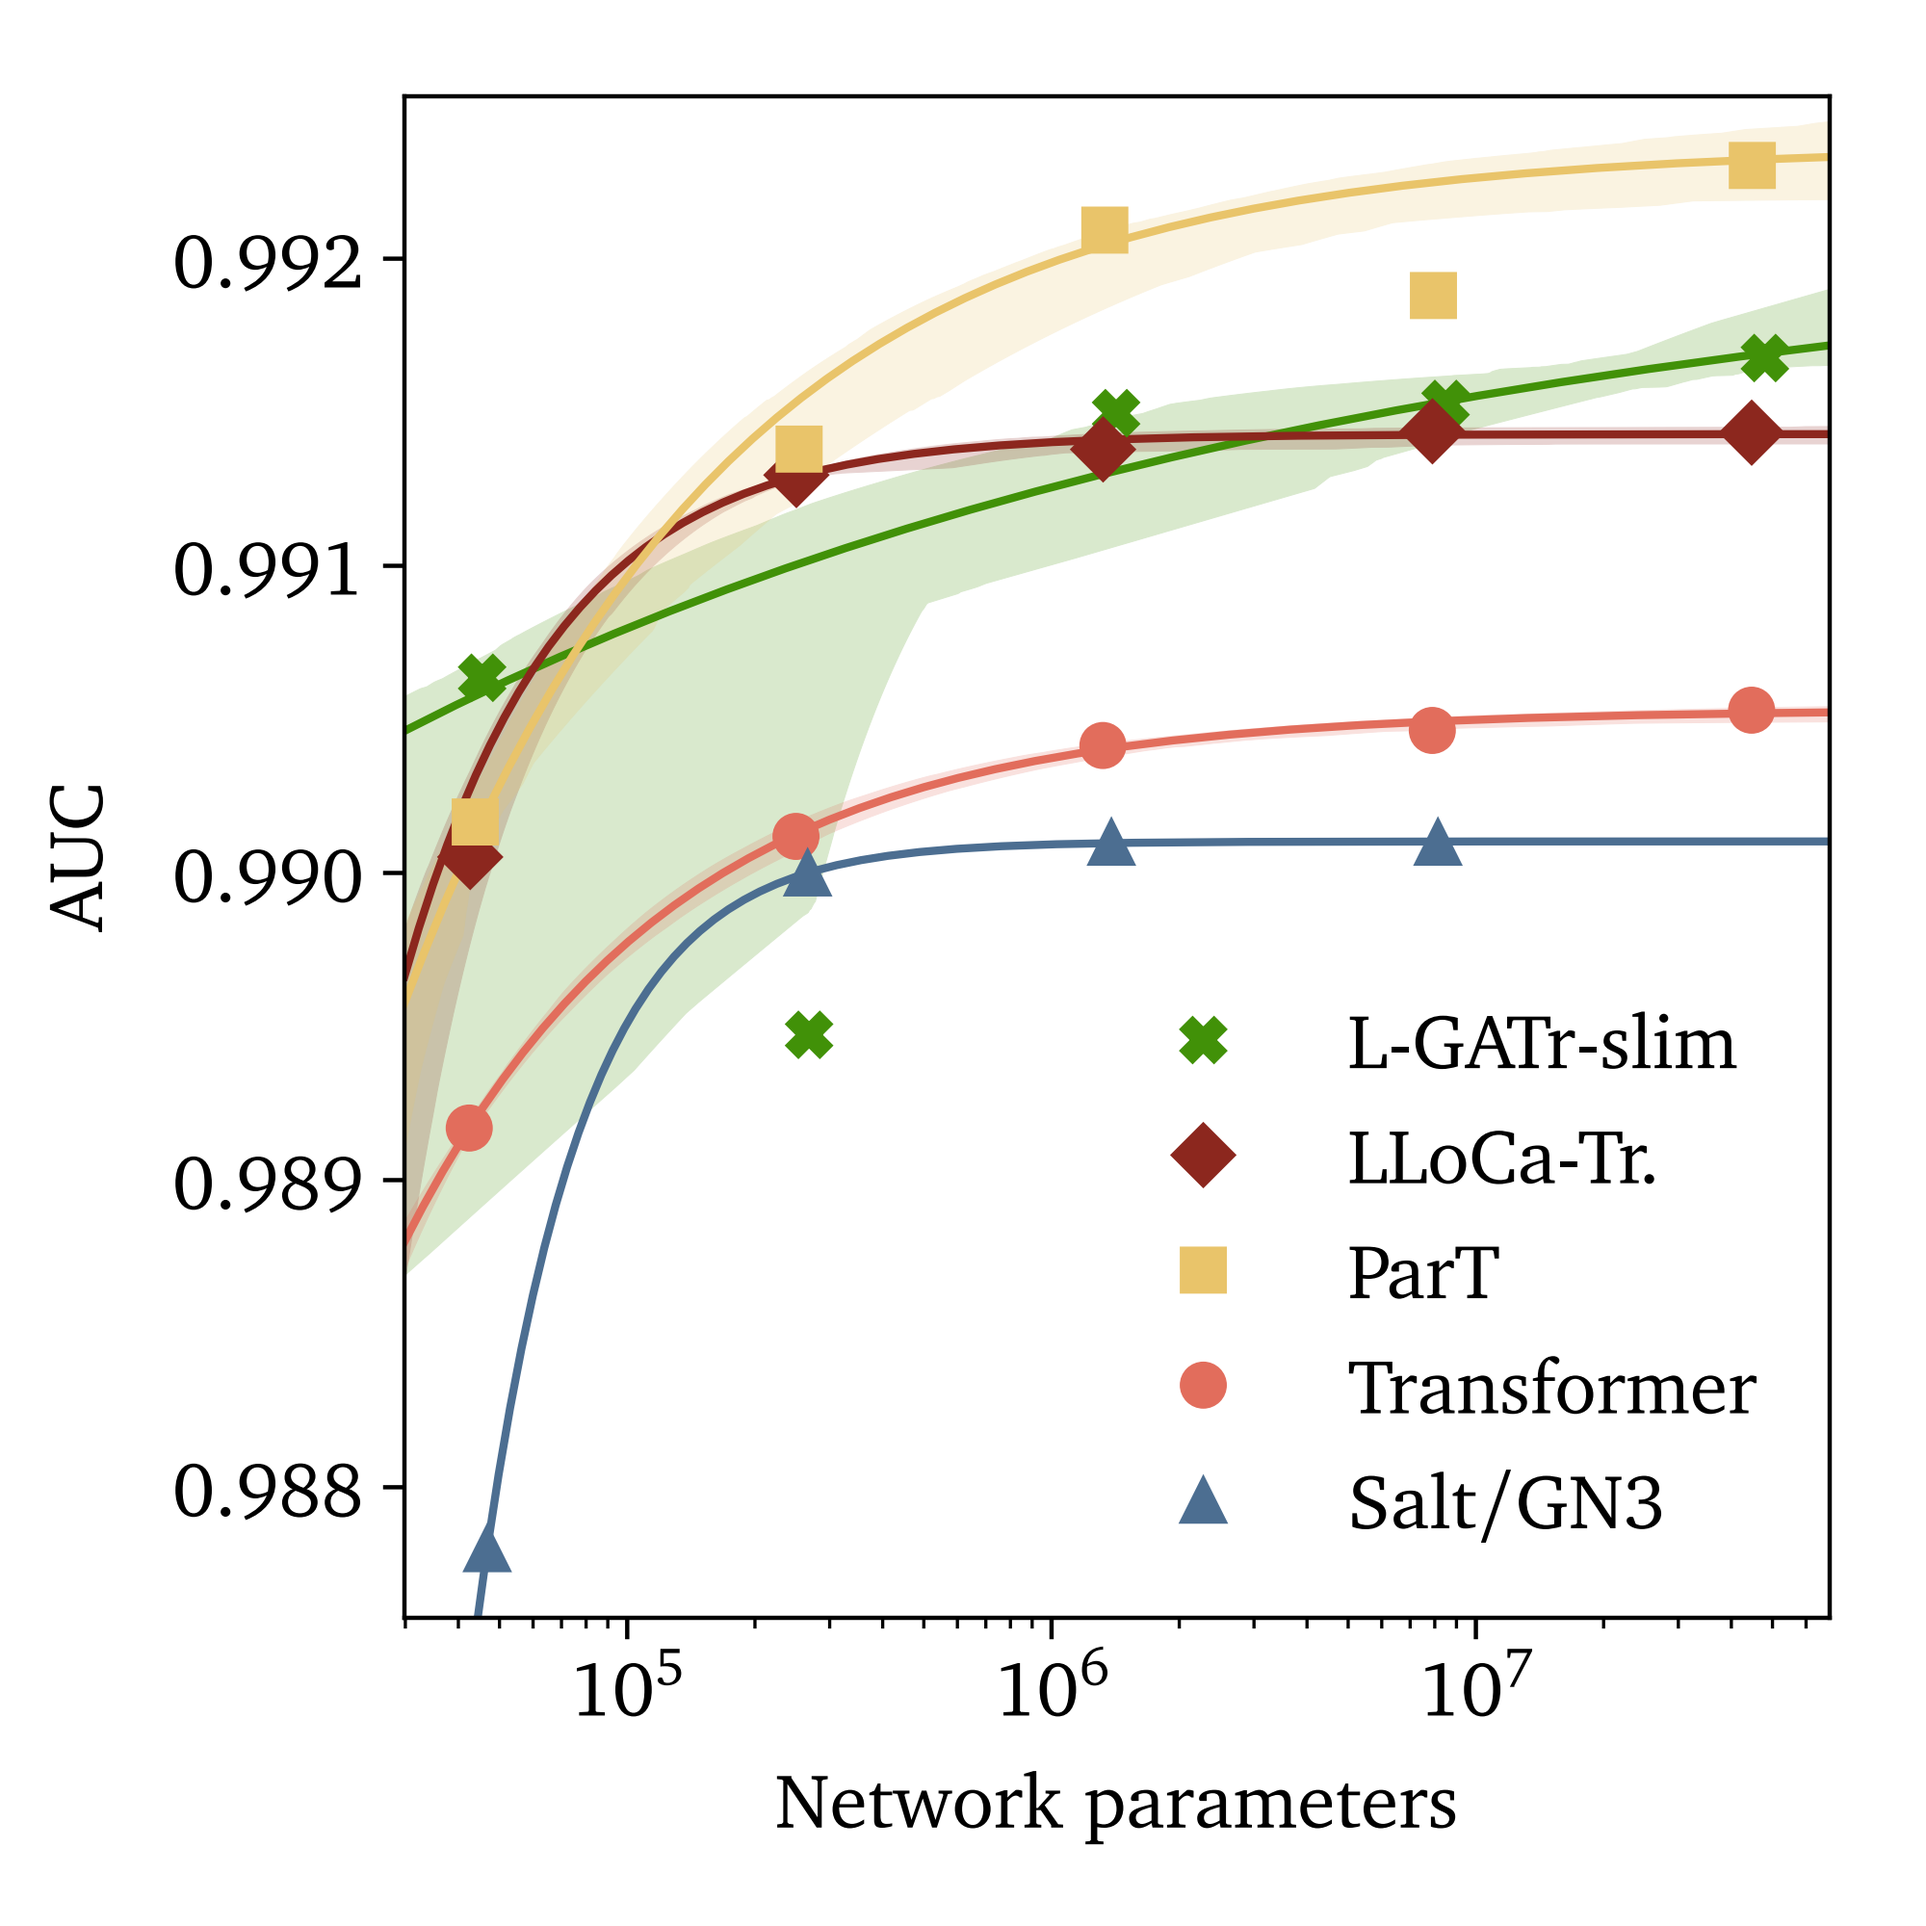
AUC
0.992
0.991
0.990
0.989
0.988
10^5
10^6
10^7
Network parameters
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3

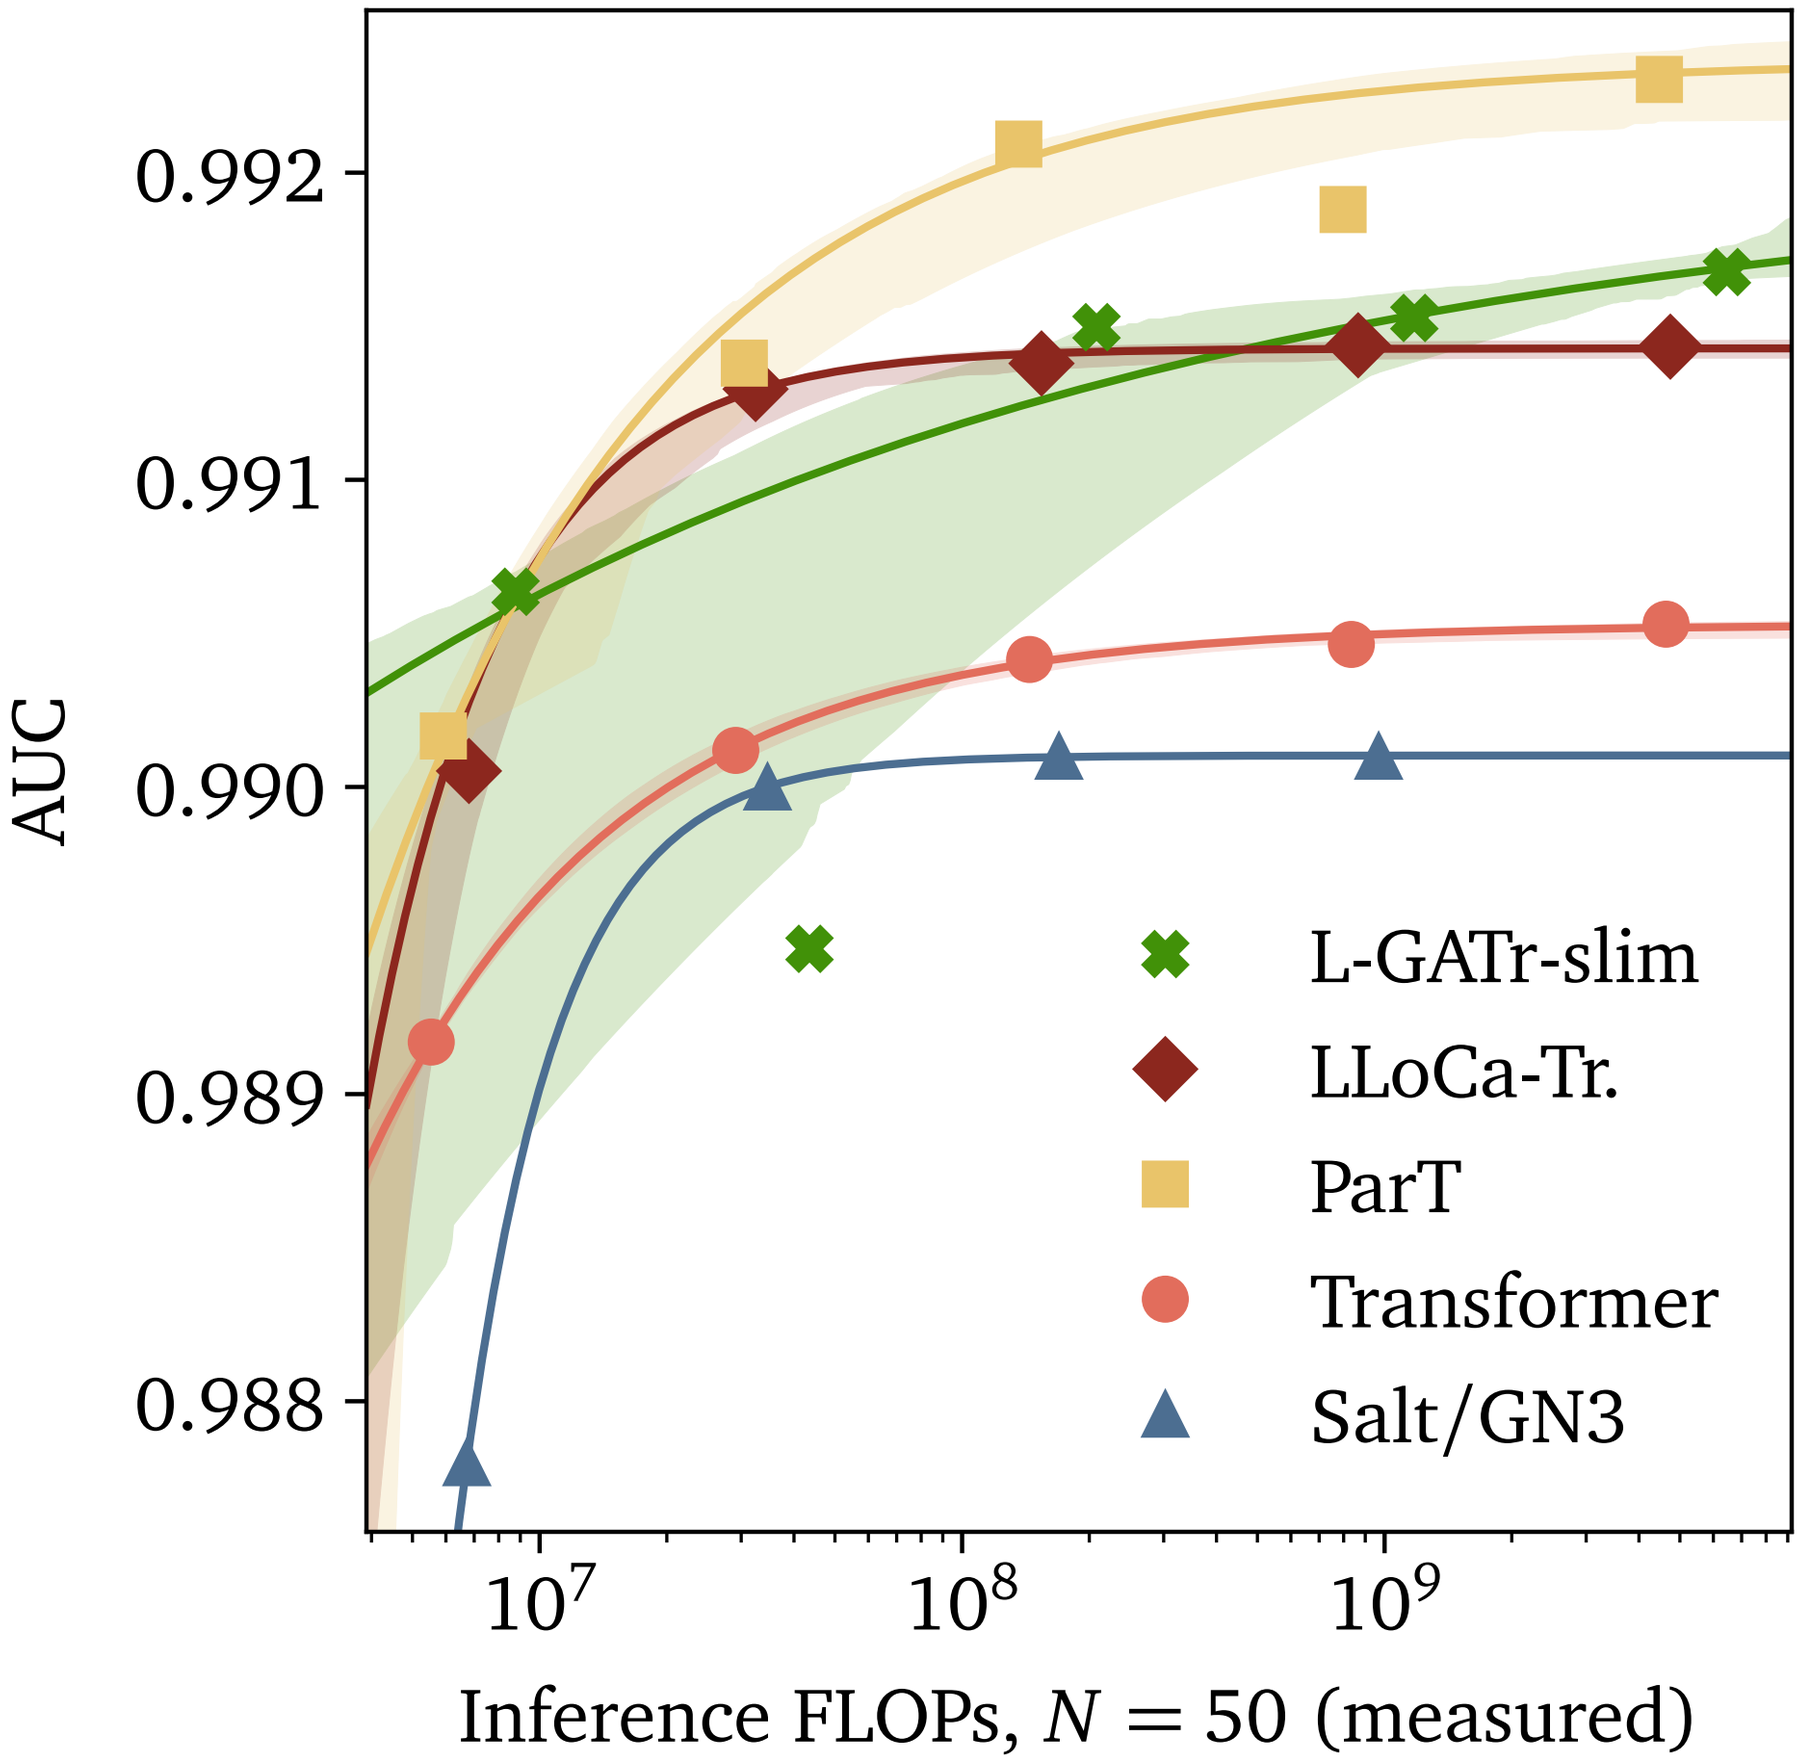
0.992
0.991
0.990
0.989
0.988
AUC
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (measured)
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3

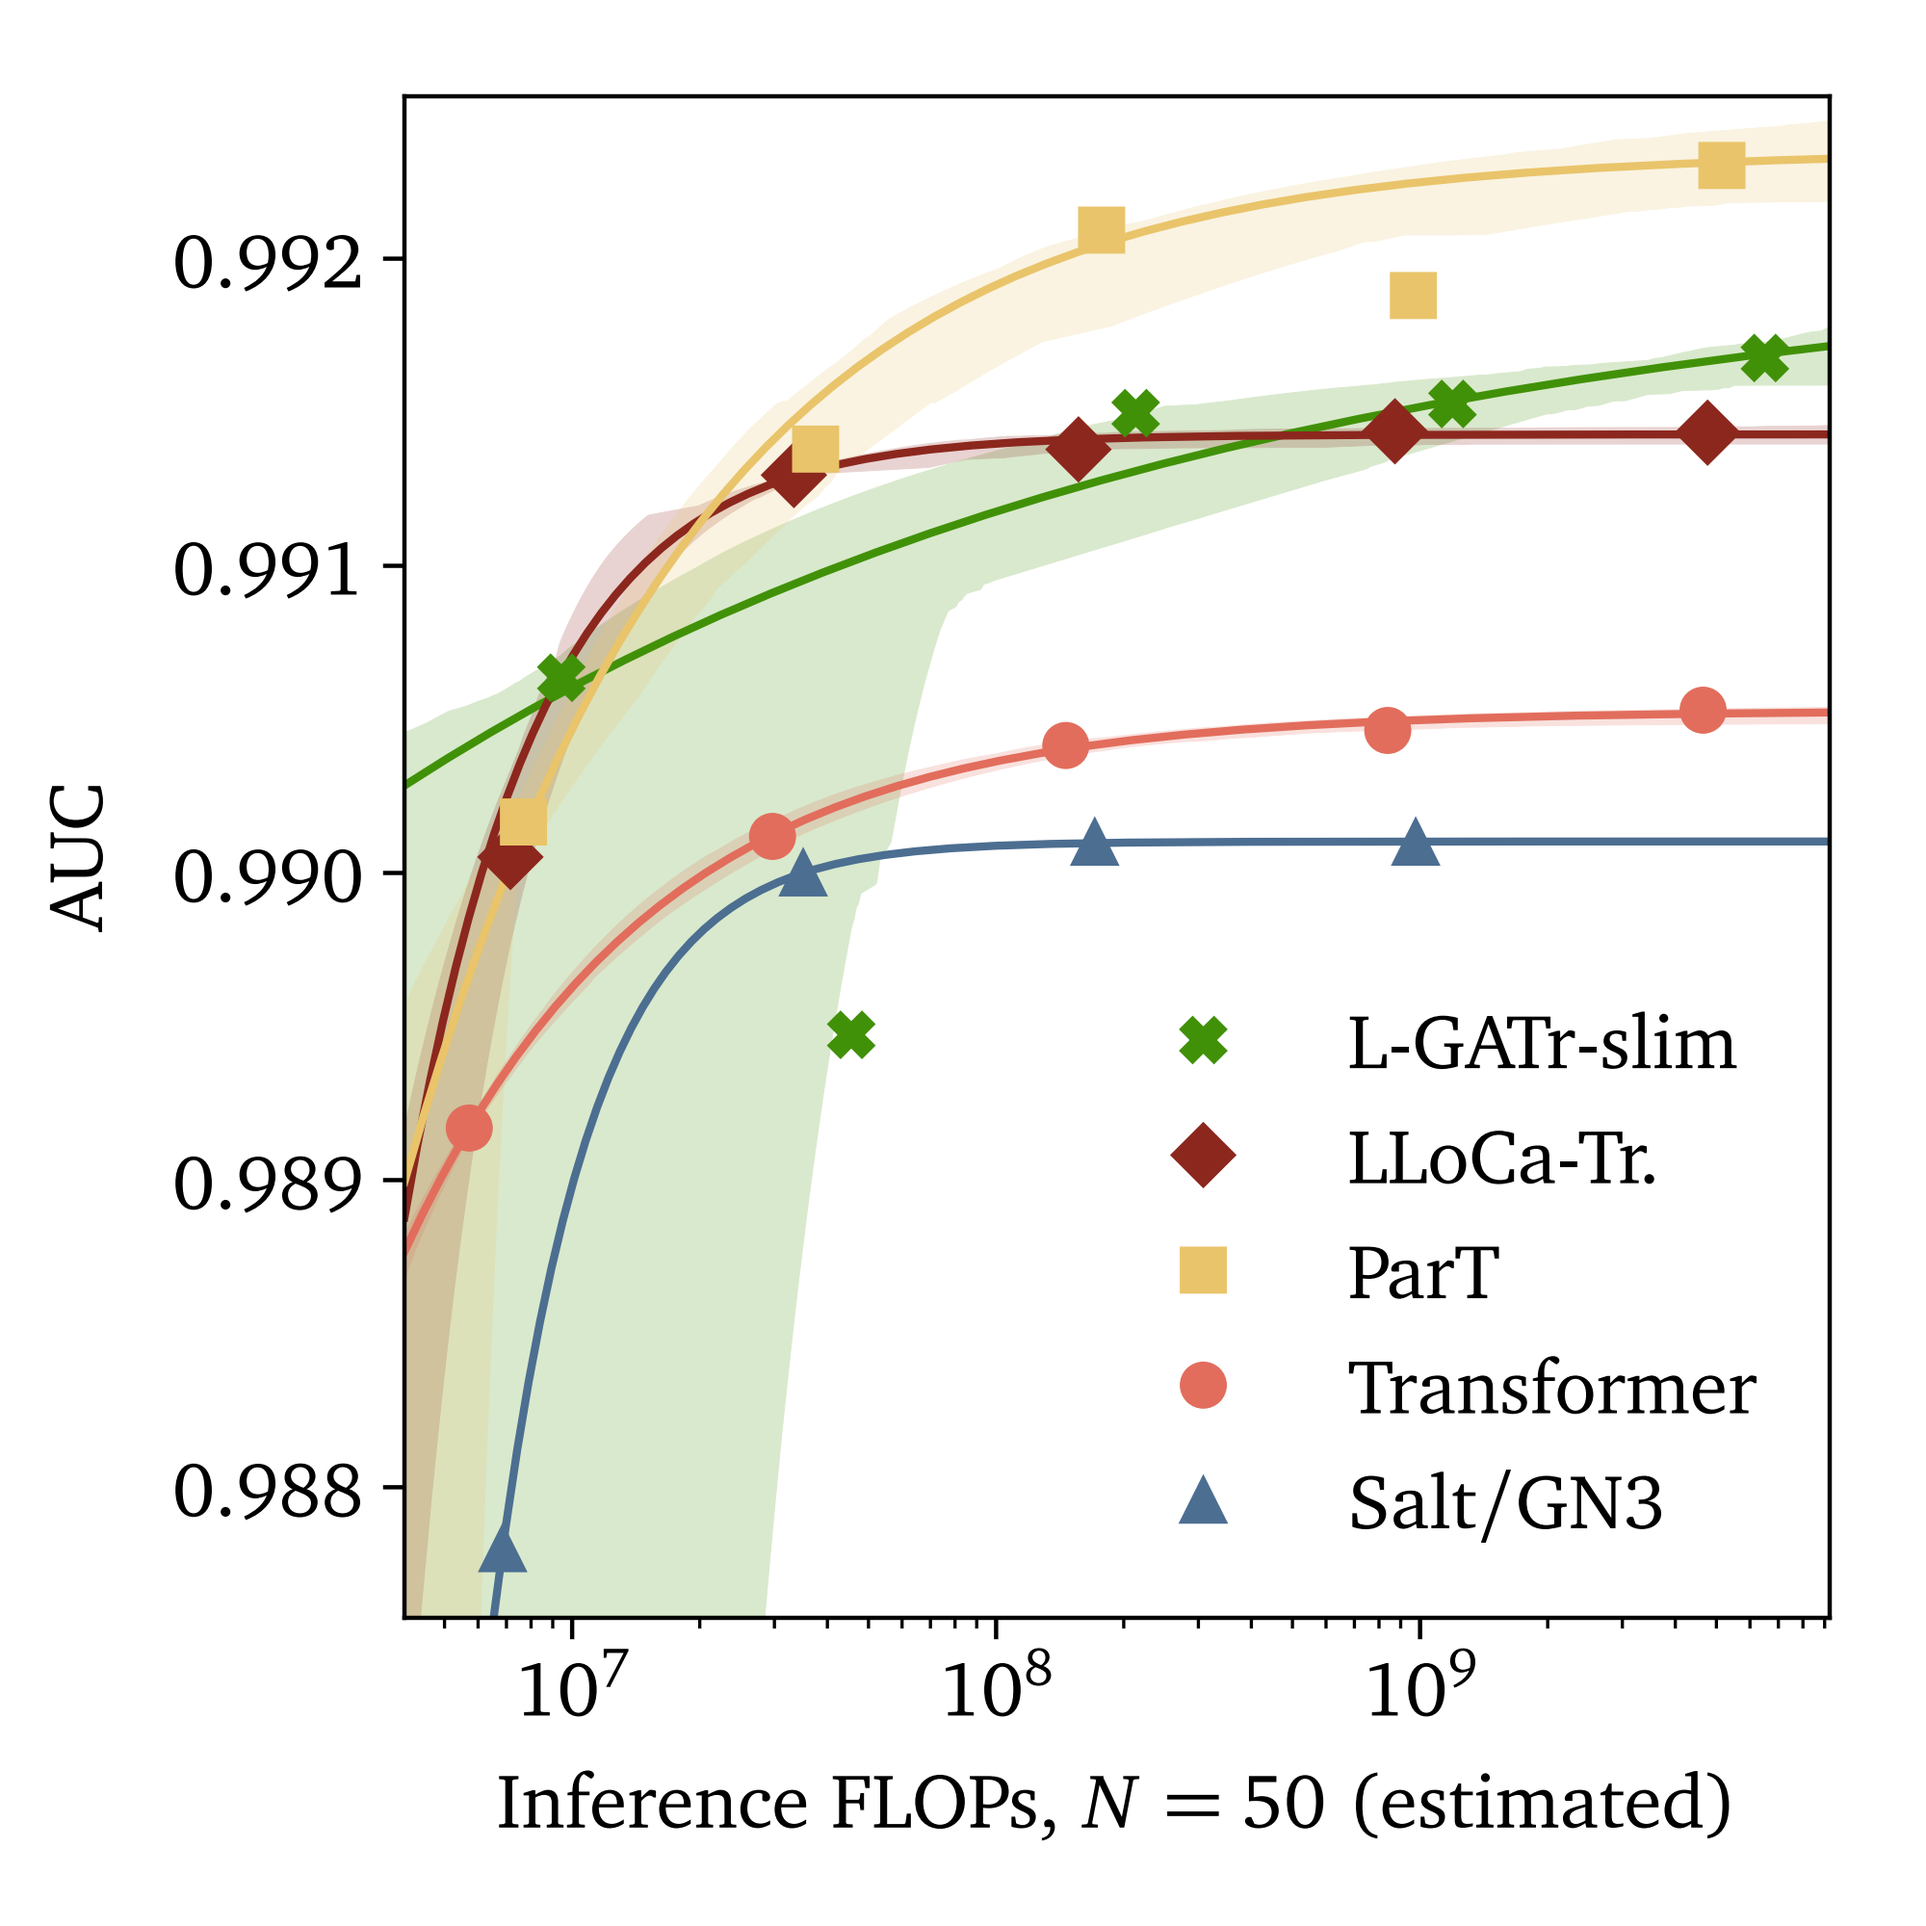
AUC
0.992
0.991
0.990
0.989
0.988
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (estimated)
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3

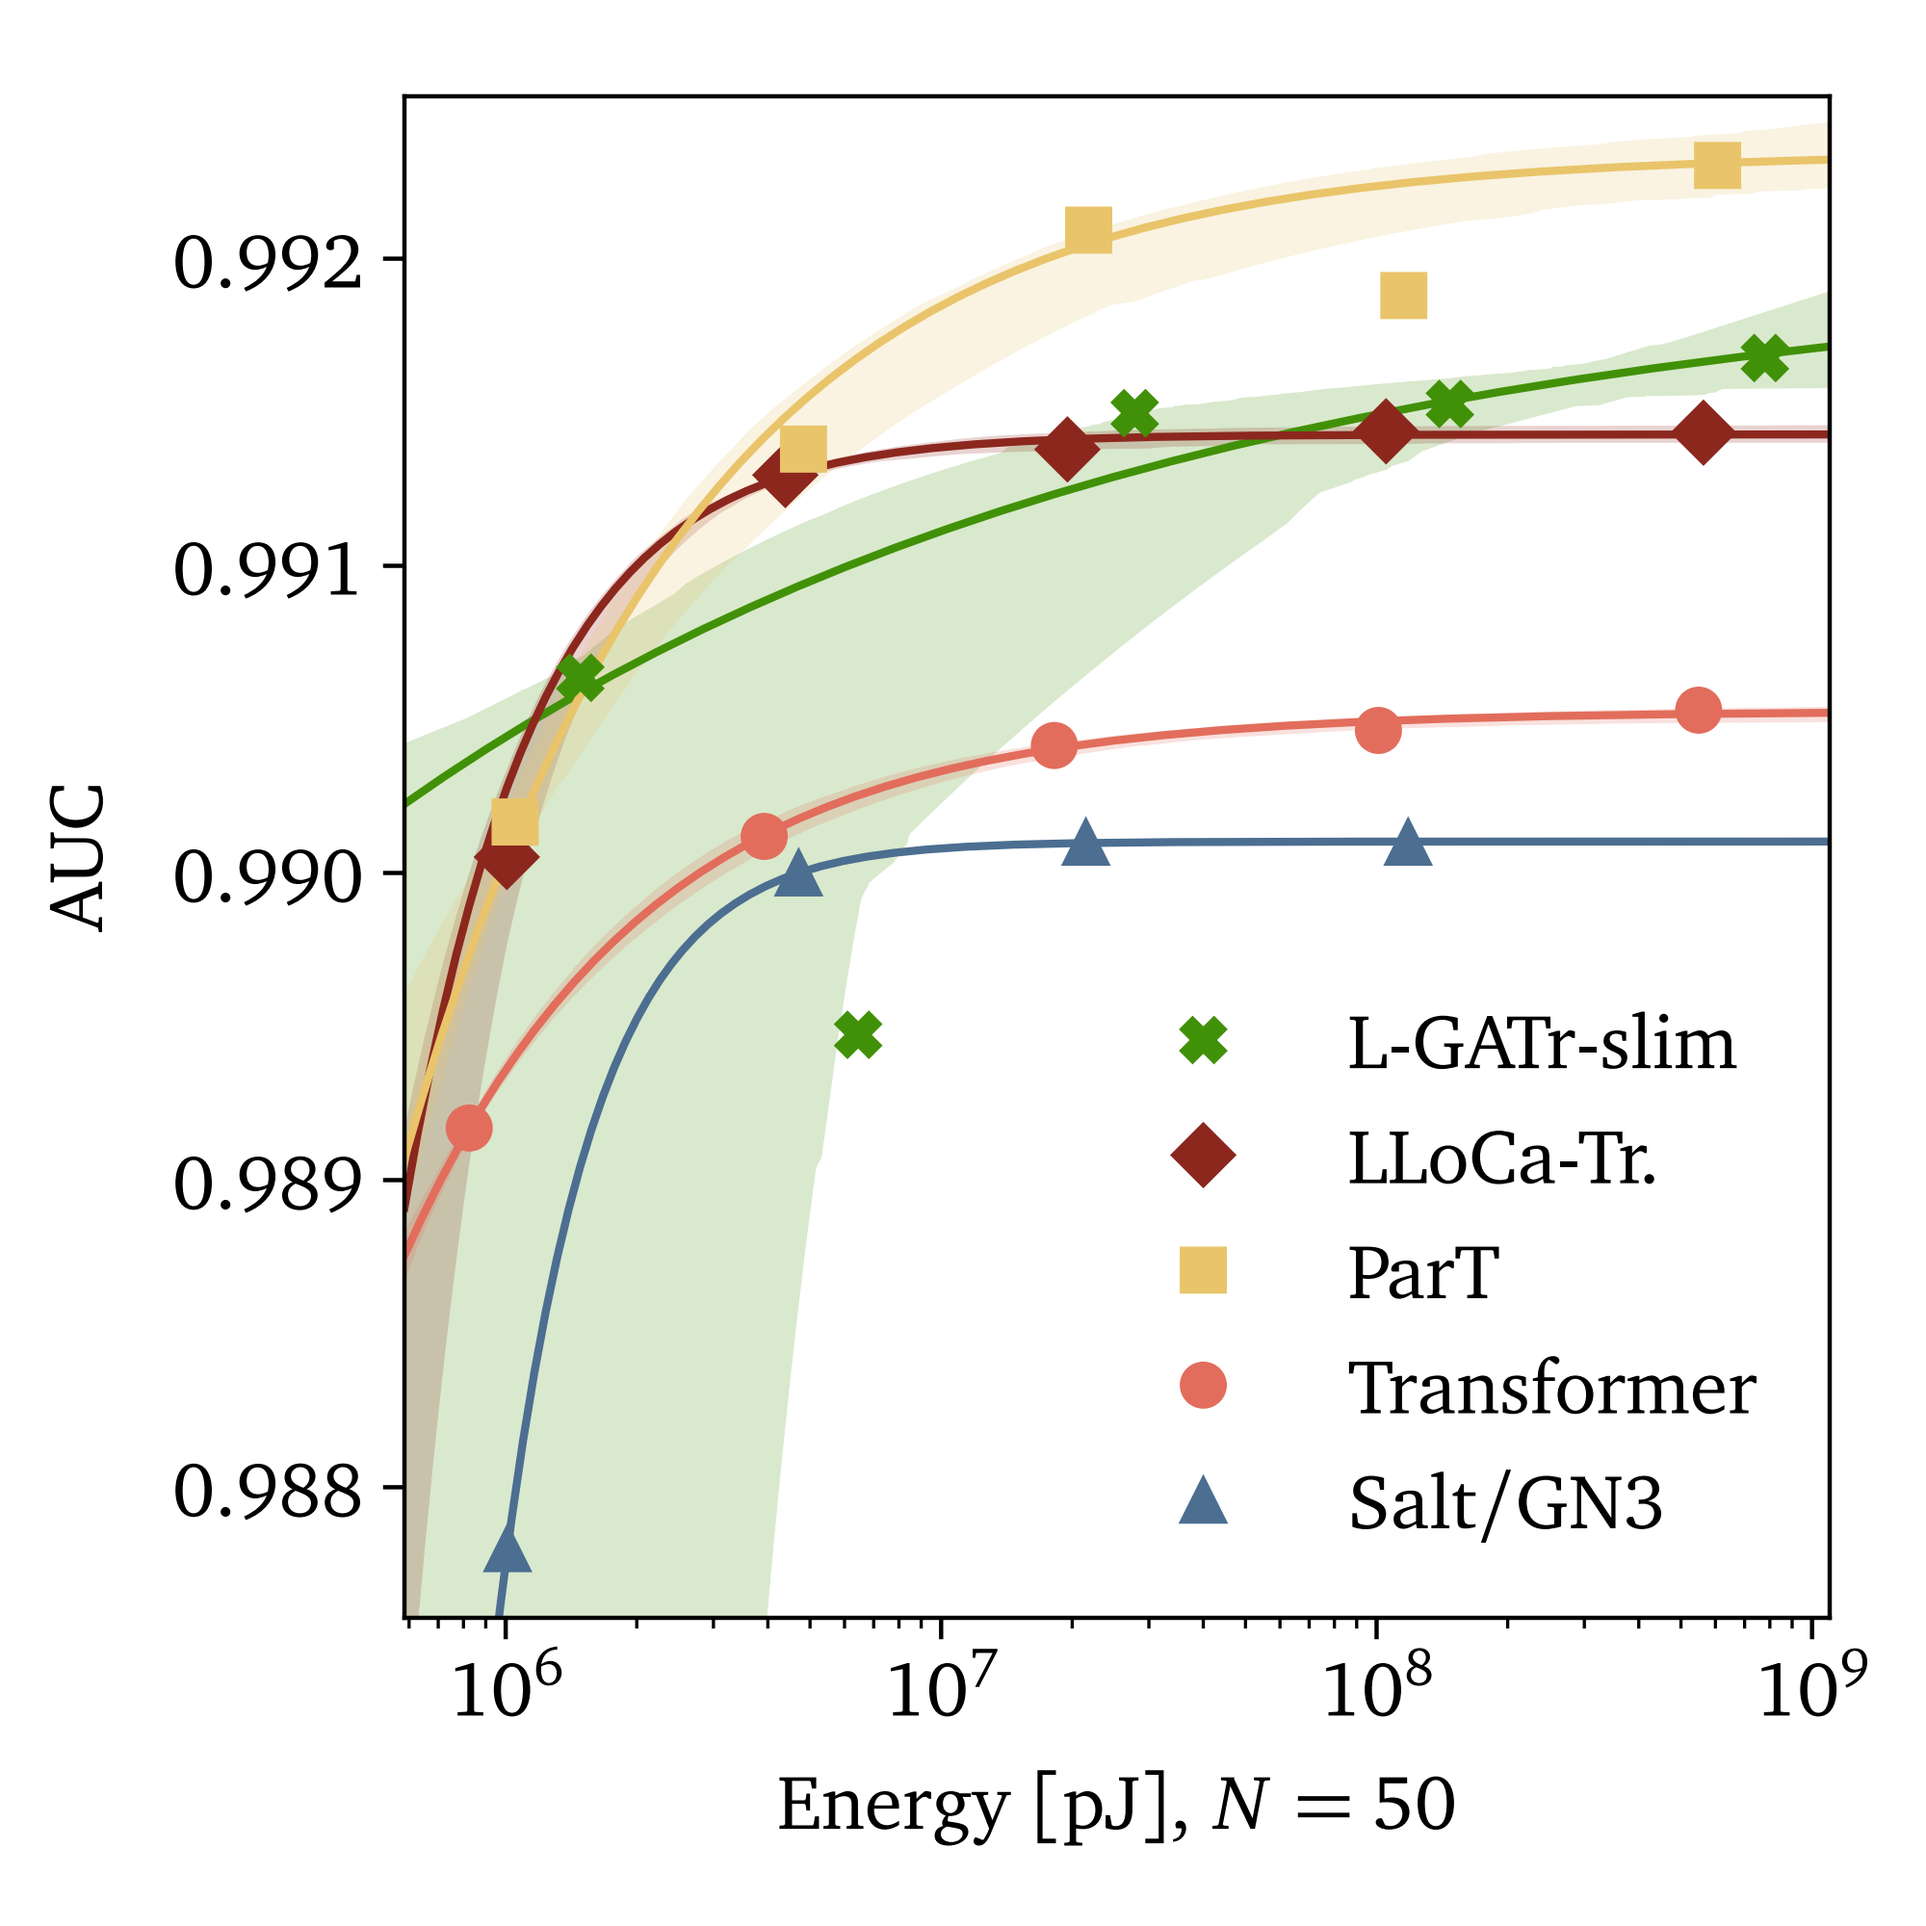
0.992
0.991
0.990
0.989
0.988
AUC
10^6
10^7
10^8
10^9
Energy [pJ], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3

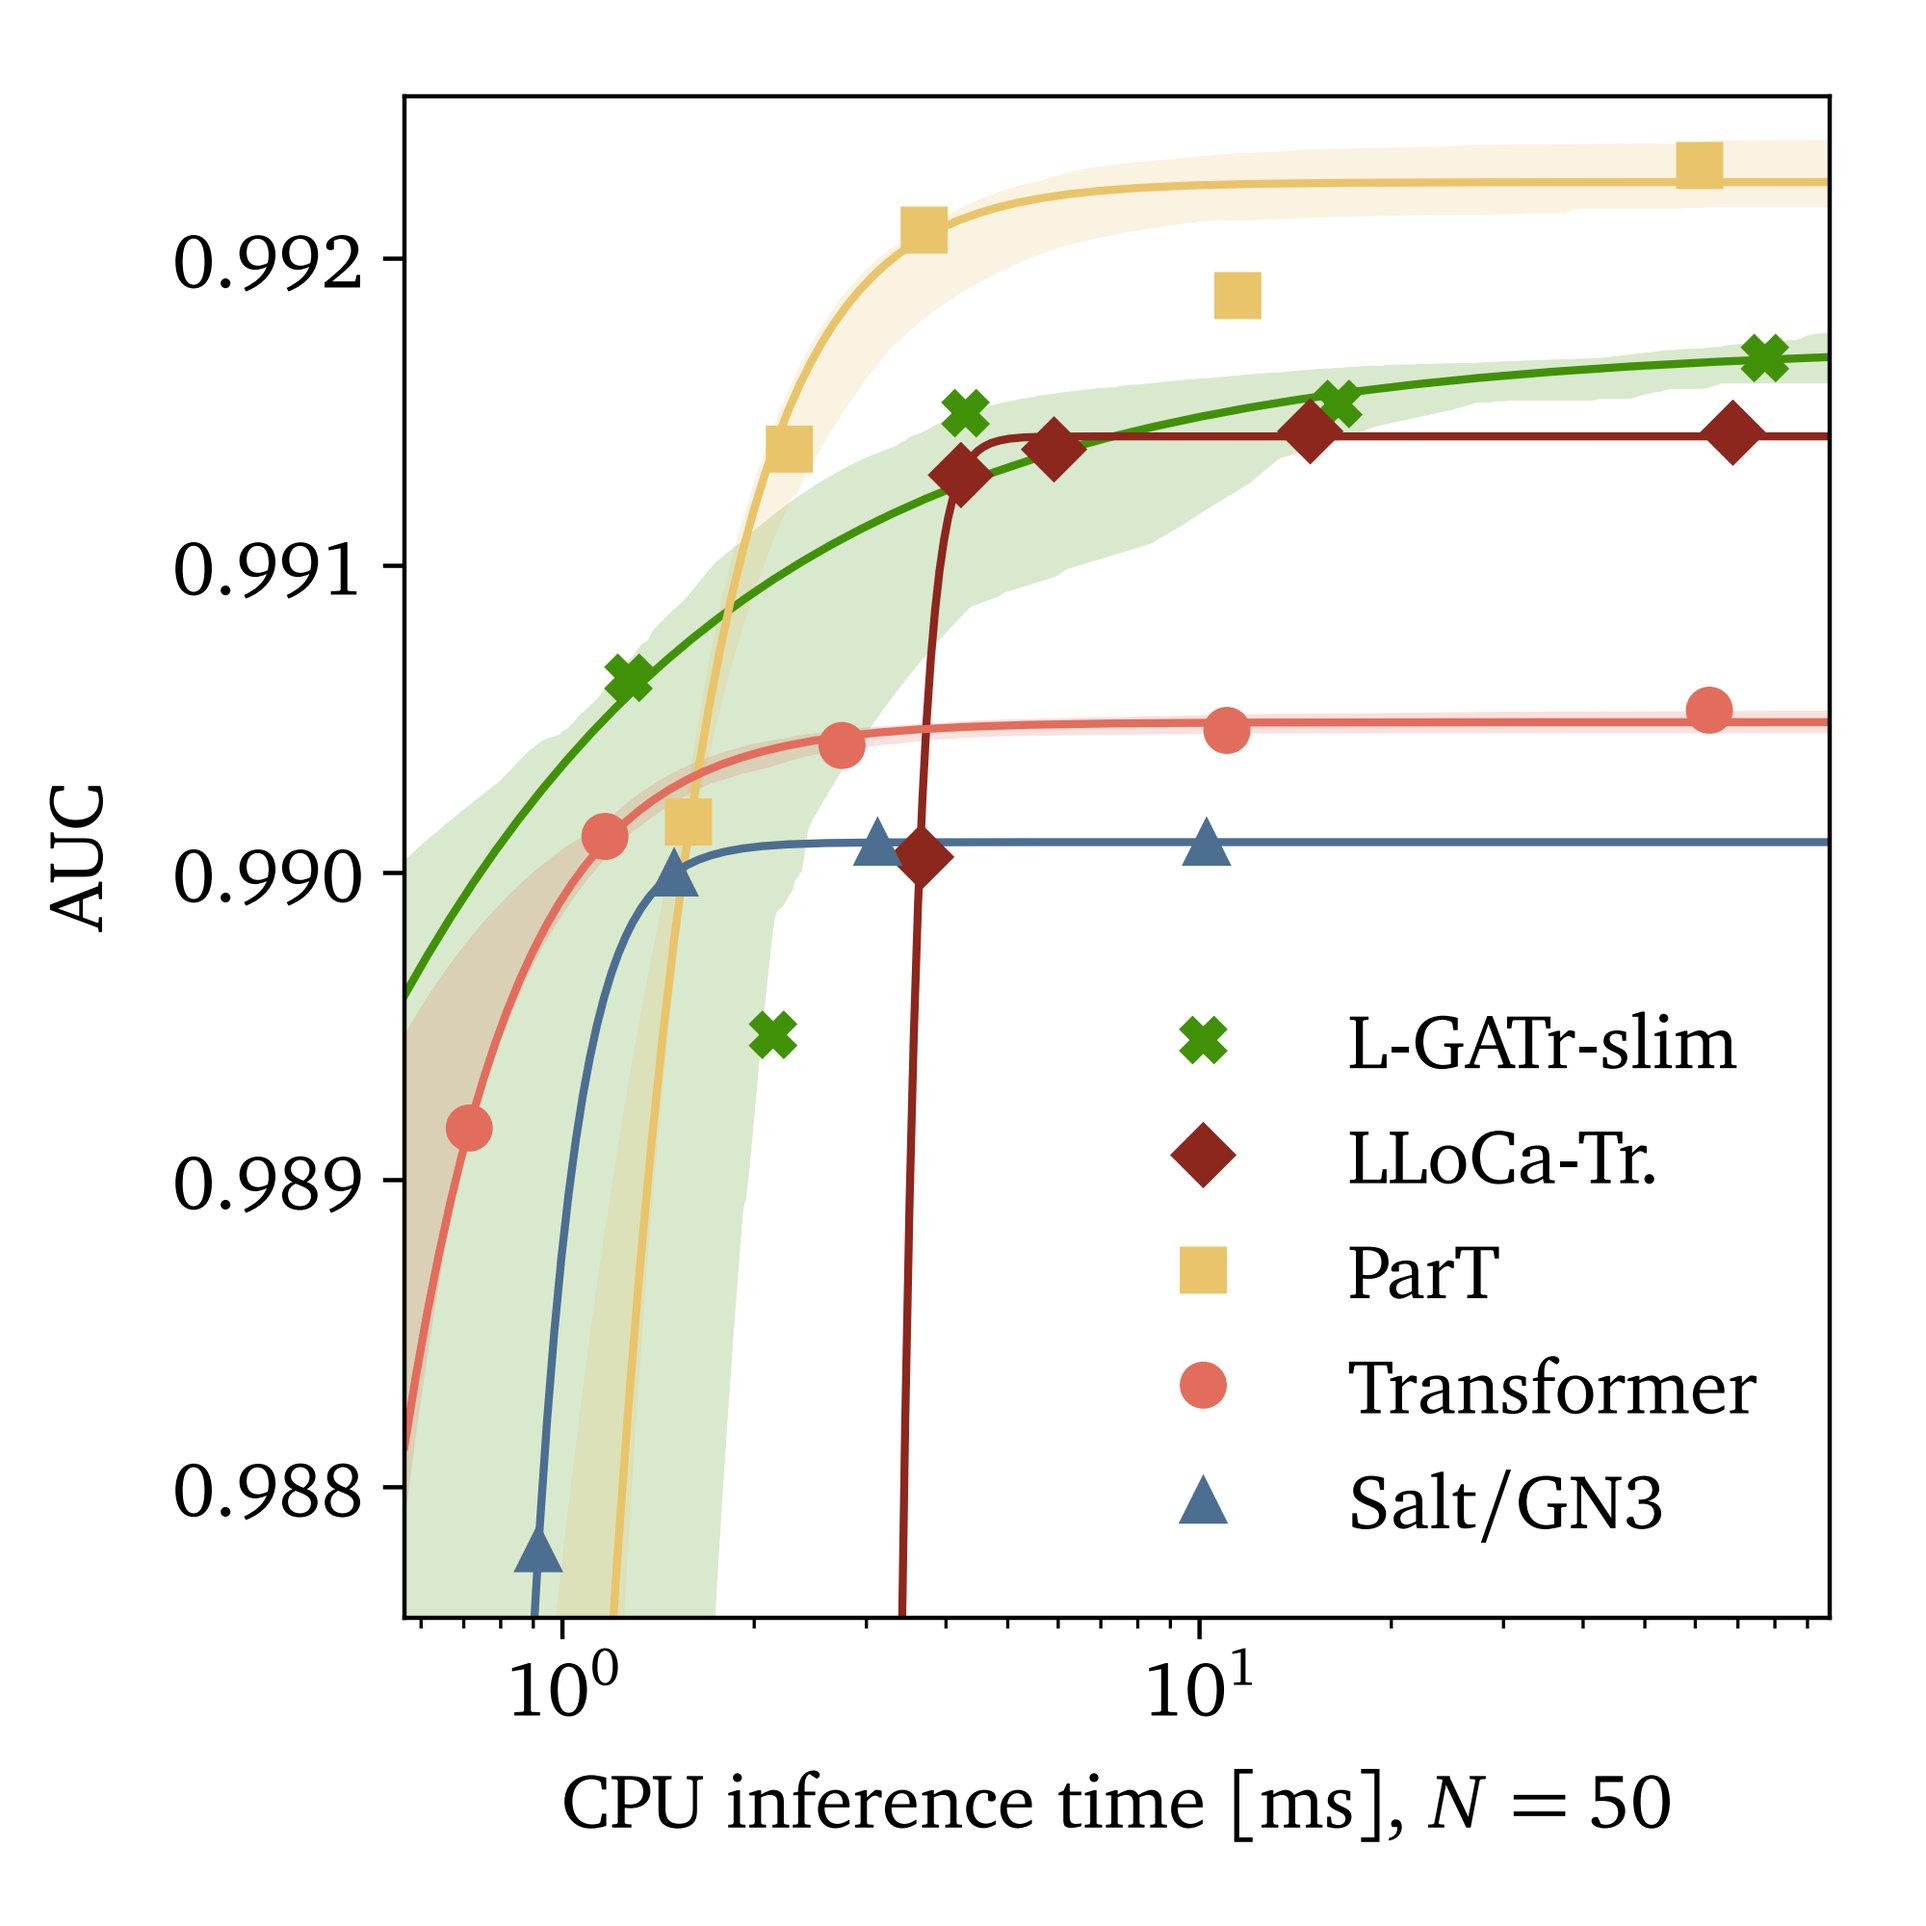
0.992
0.991
0.990
0.989
0.988
AUC
10^0
10^1
CPU inference time [ms], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3

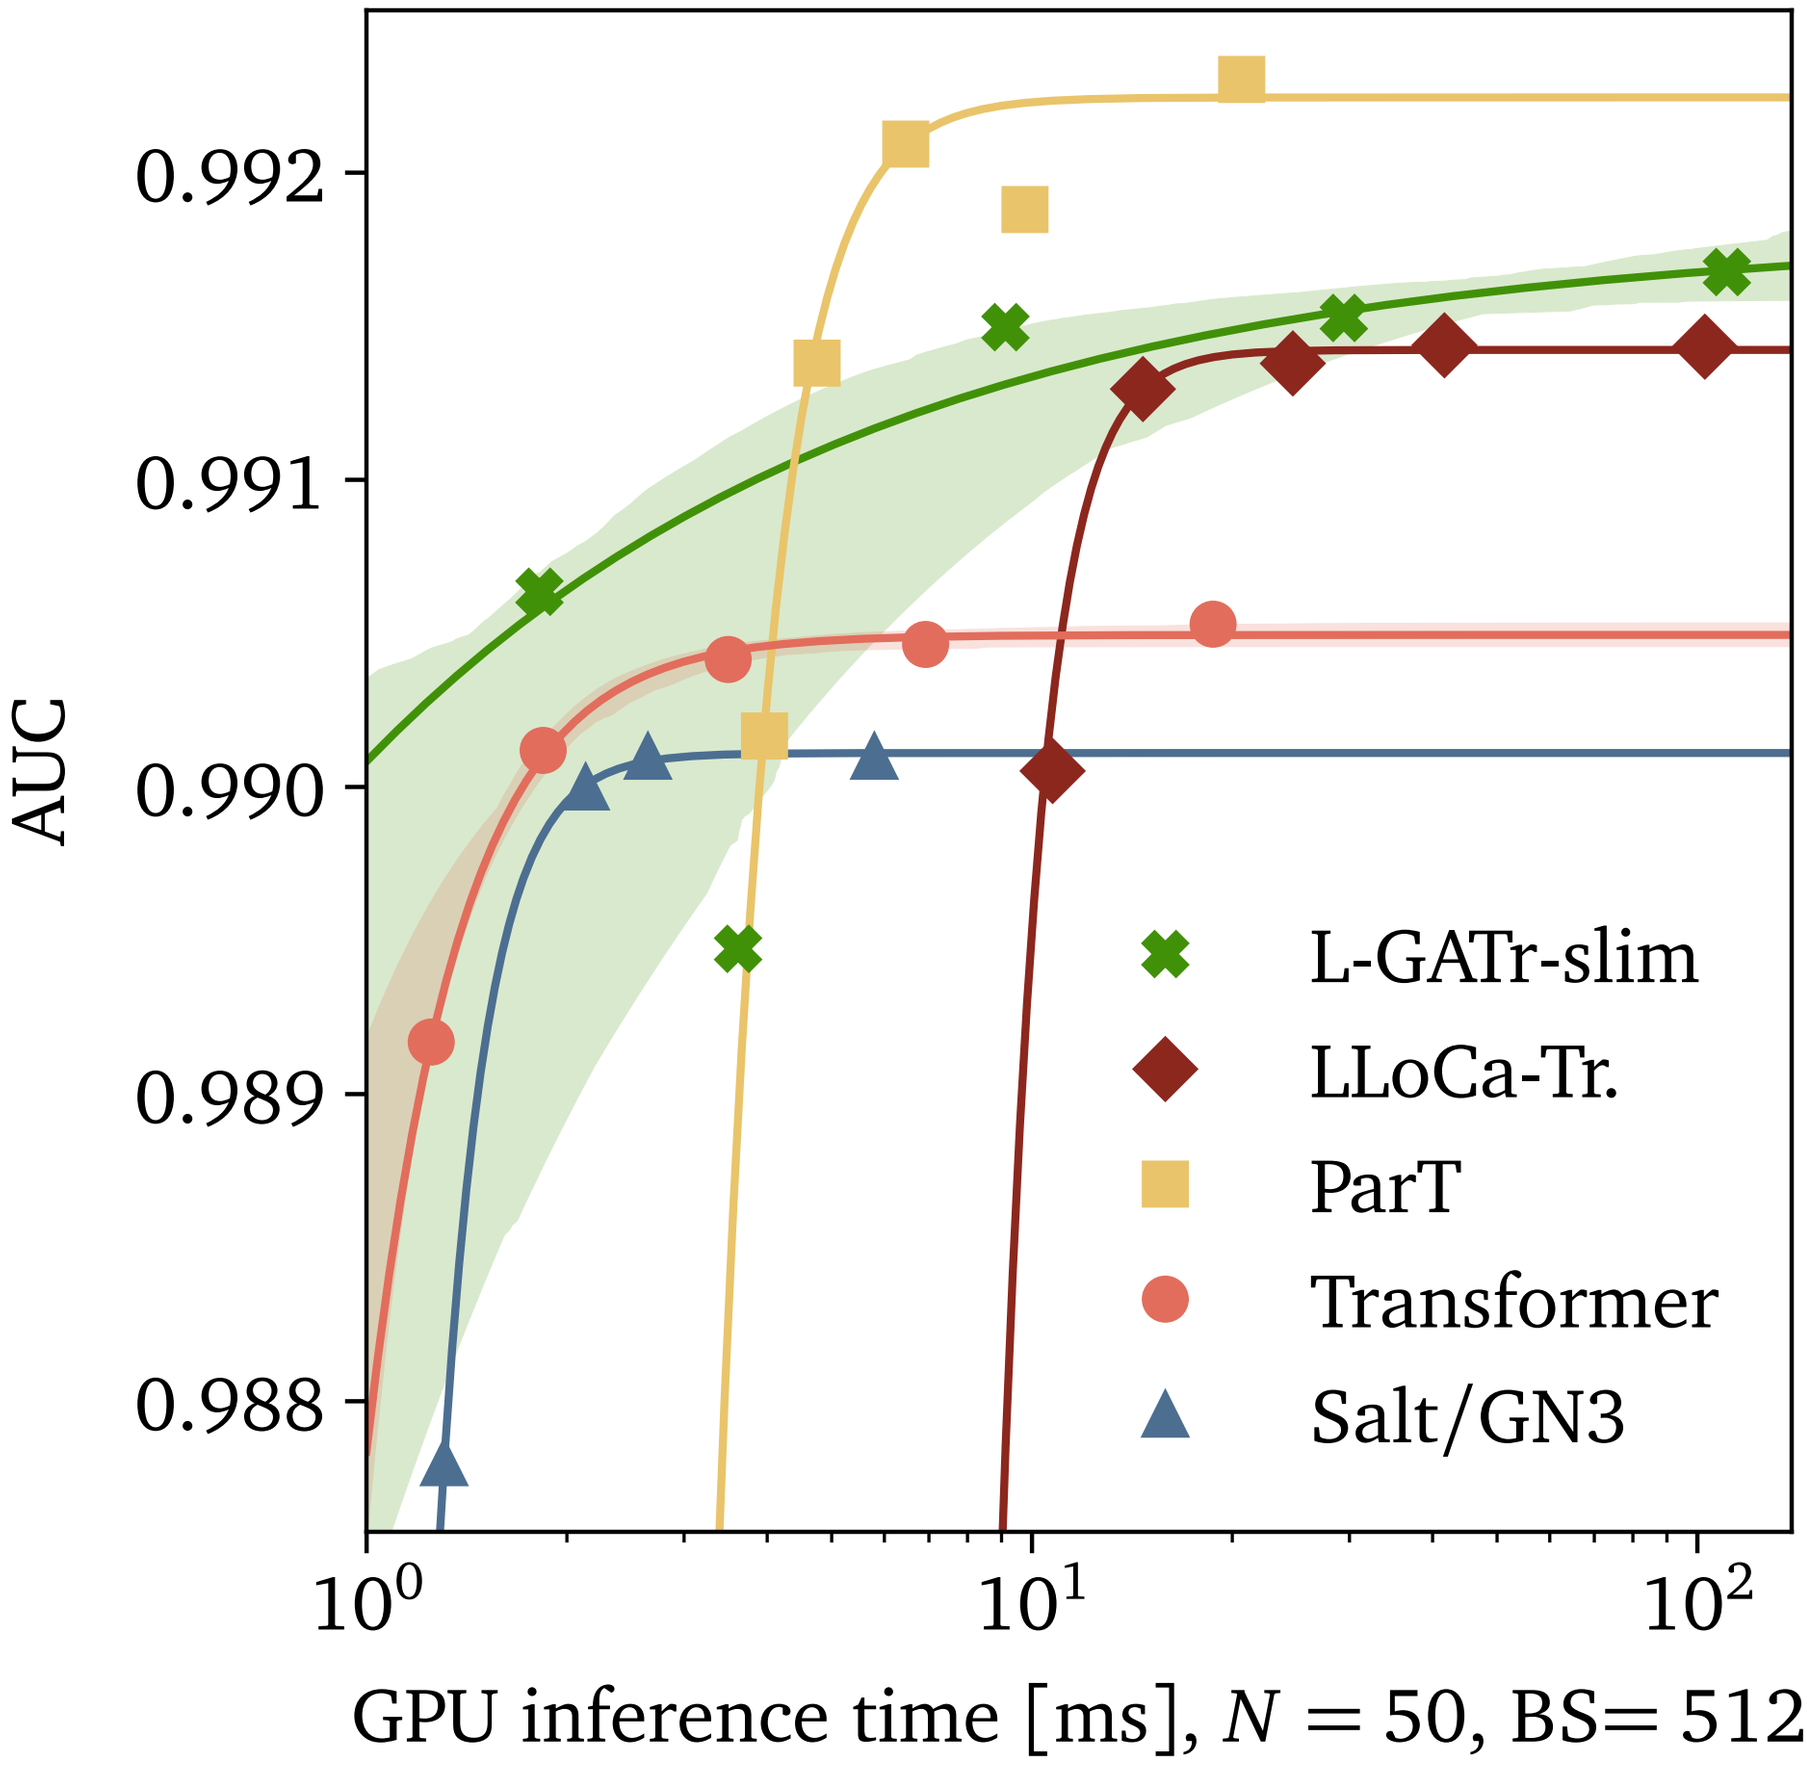
0.992
0.991
0.990
0.989
0.988
AUC
10^0
10^1
10^2
GPU inference time [ms], $N = 50$, BS= 512
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3

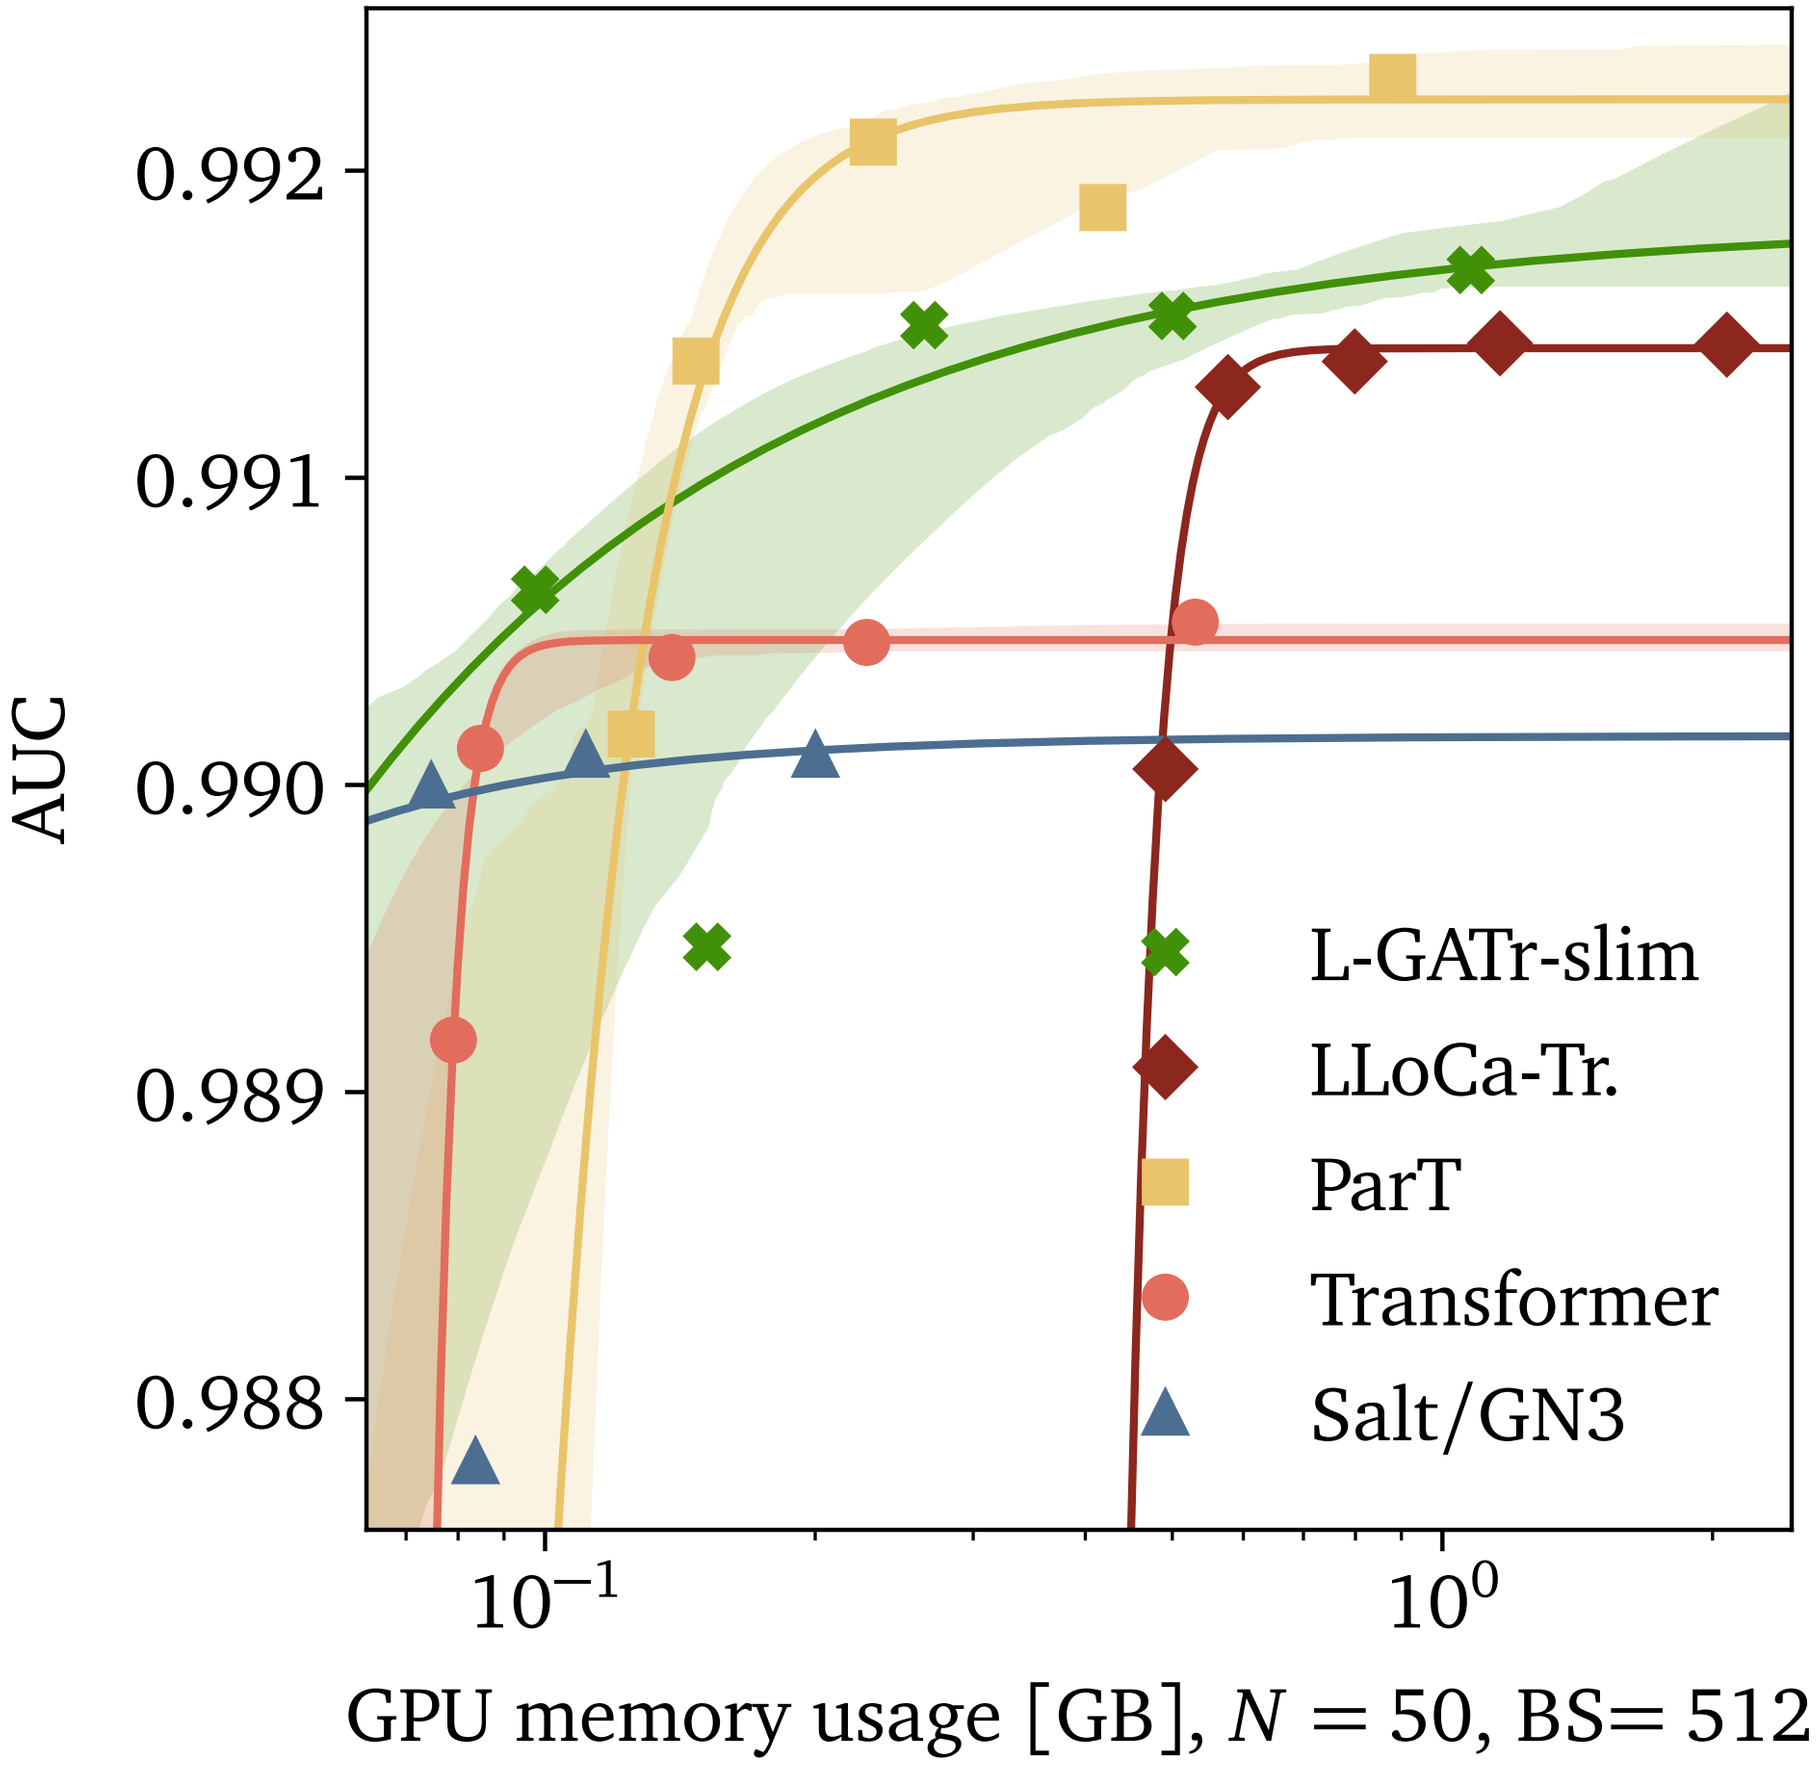
0.992
0.991
0.990
0.989
0.988
AUC
10^{-1}
10^{0}
GPU memory usage [GB], $N = 50$, BS= 512
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3

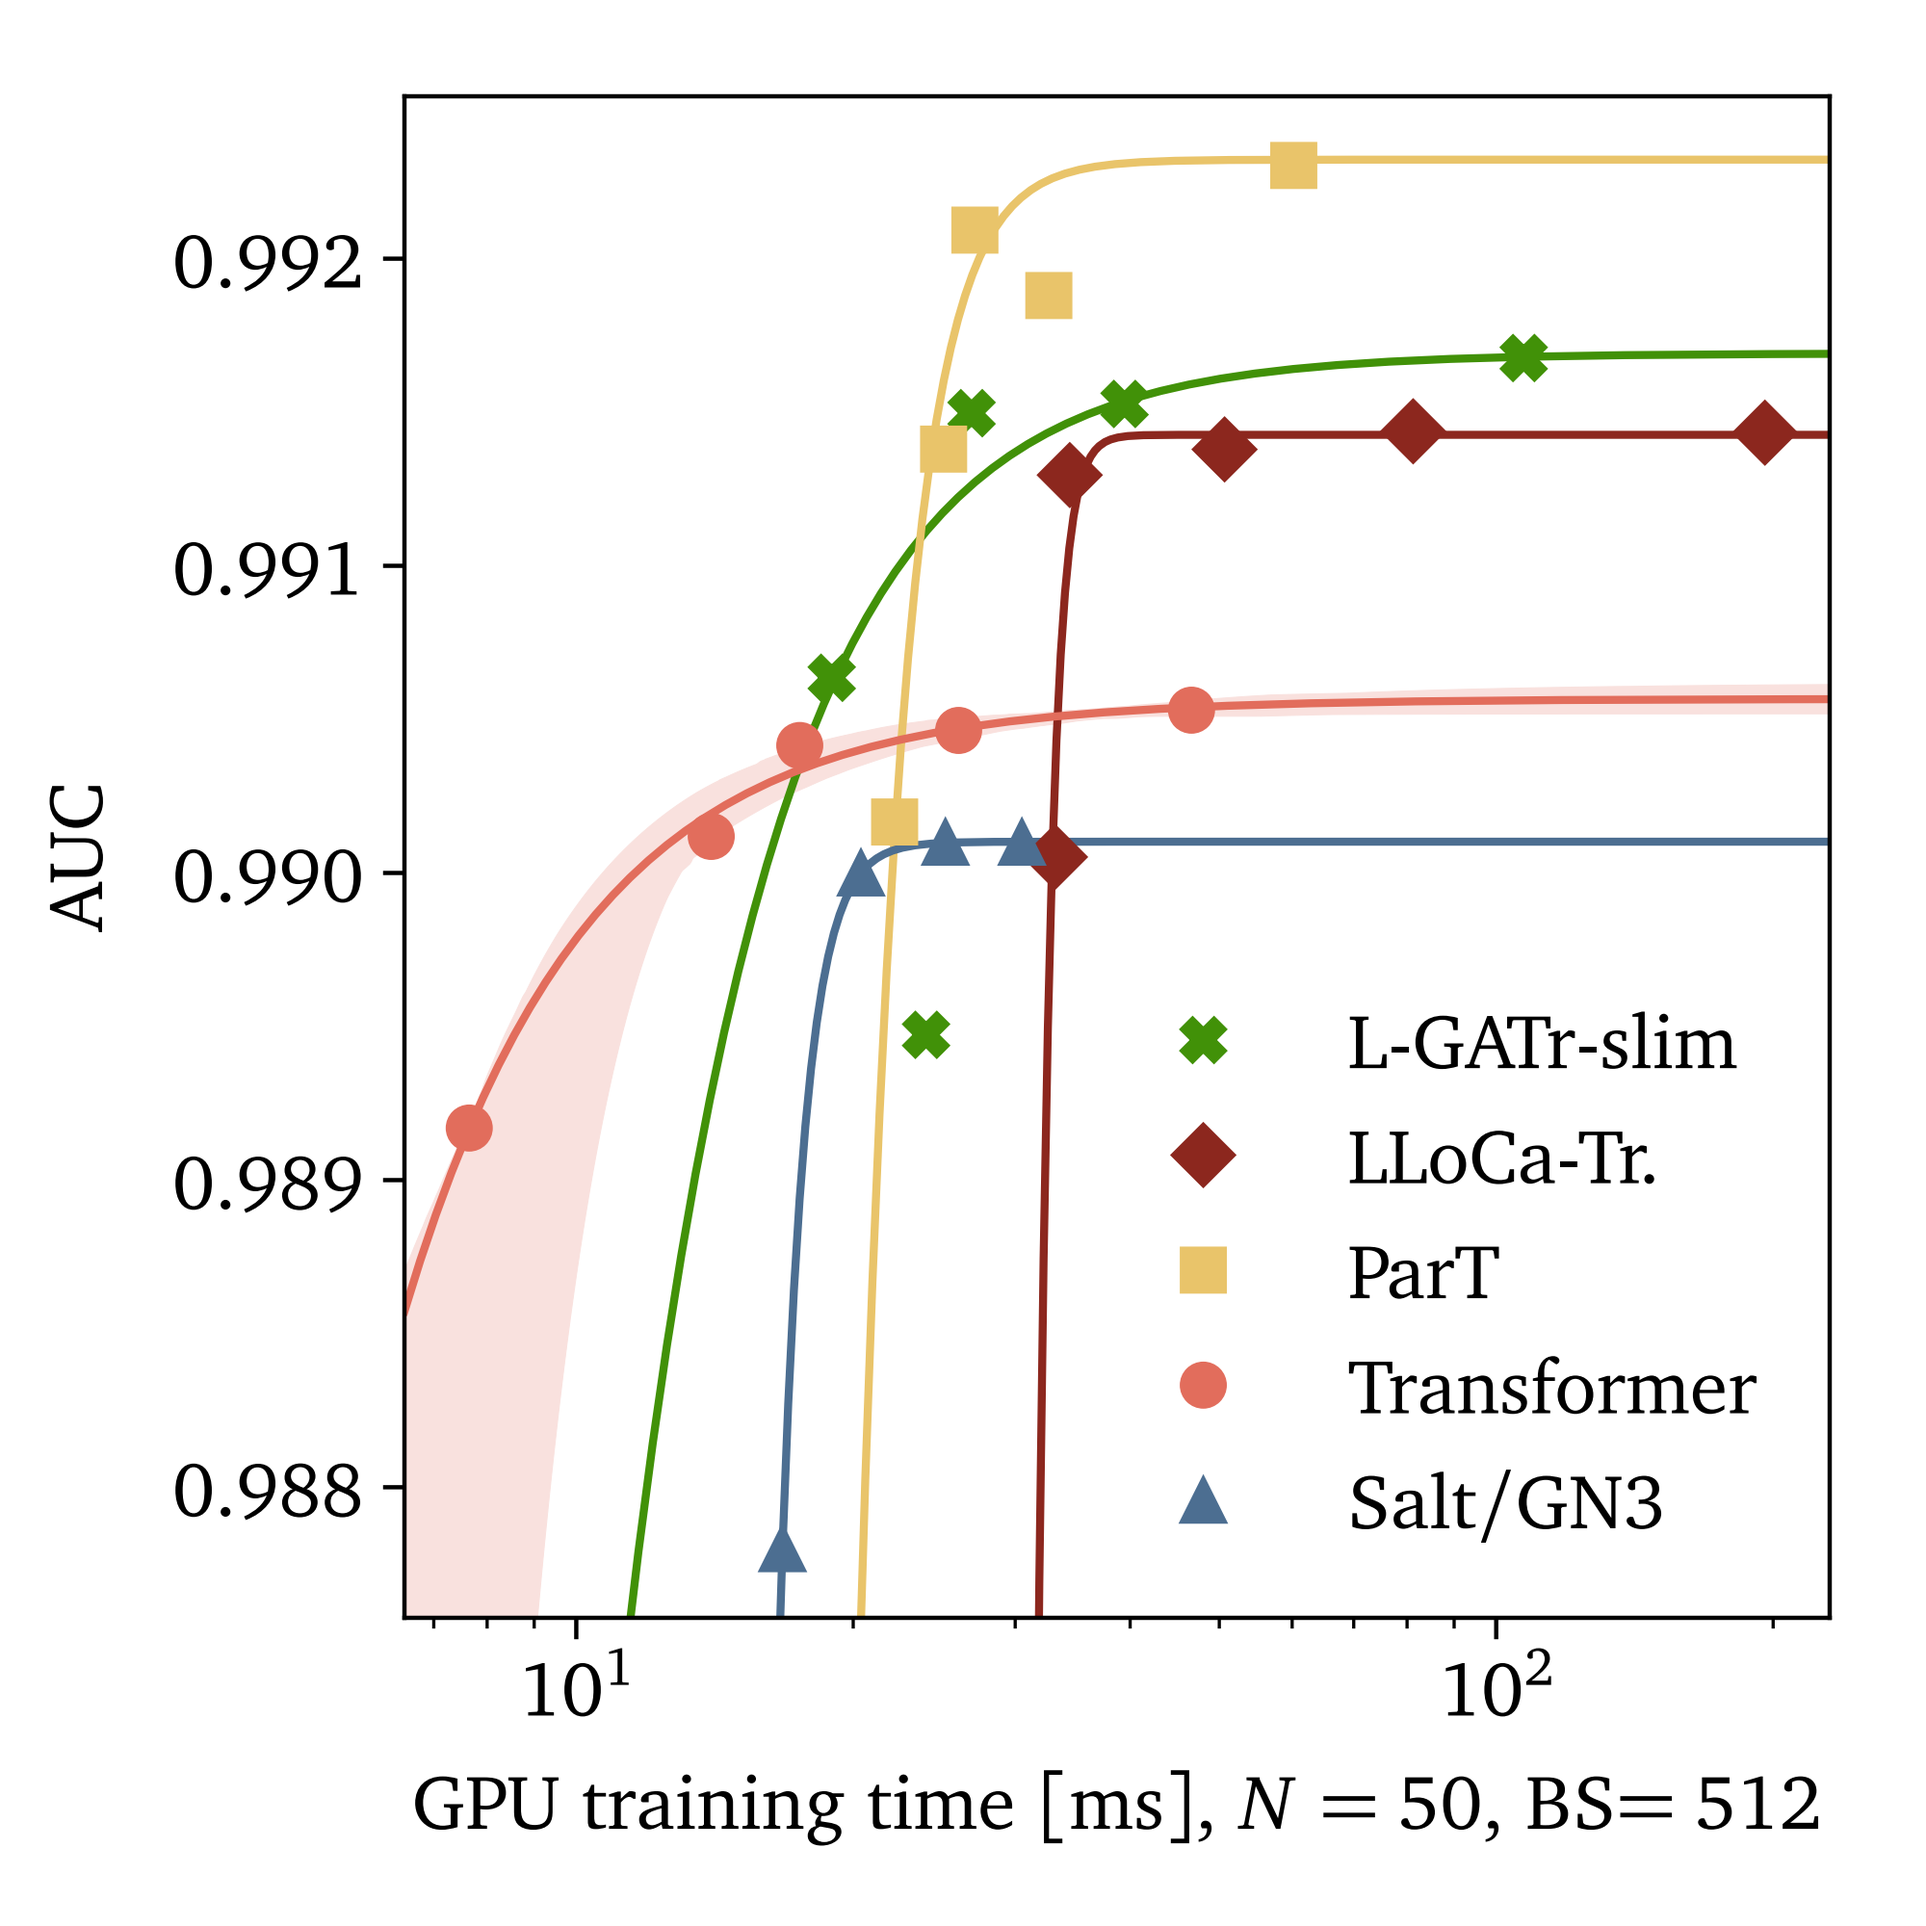
0.992
0.991
0.990
0.989
0.988
AUC
10^1
10^2
GPU training time [ms], $N = 50$, BS$= 512$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3

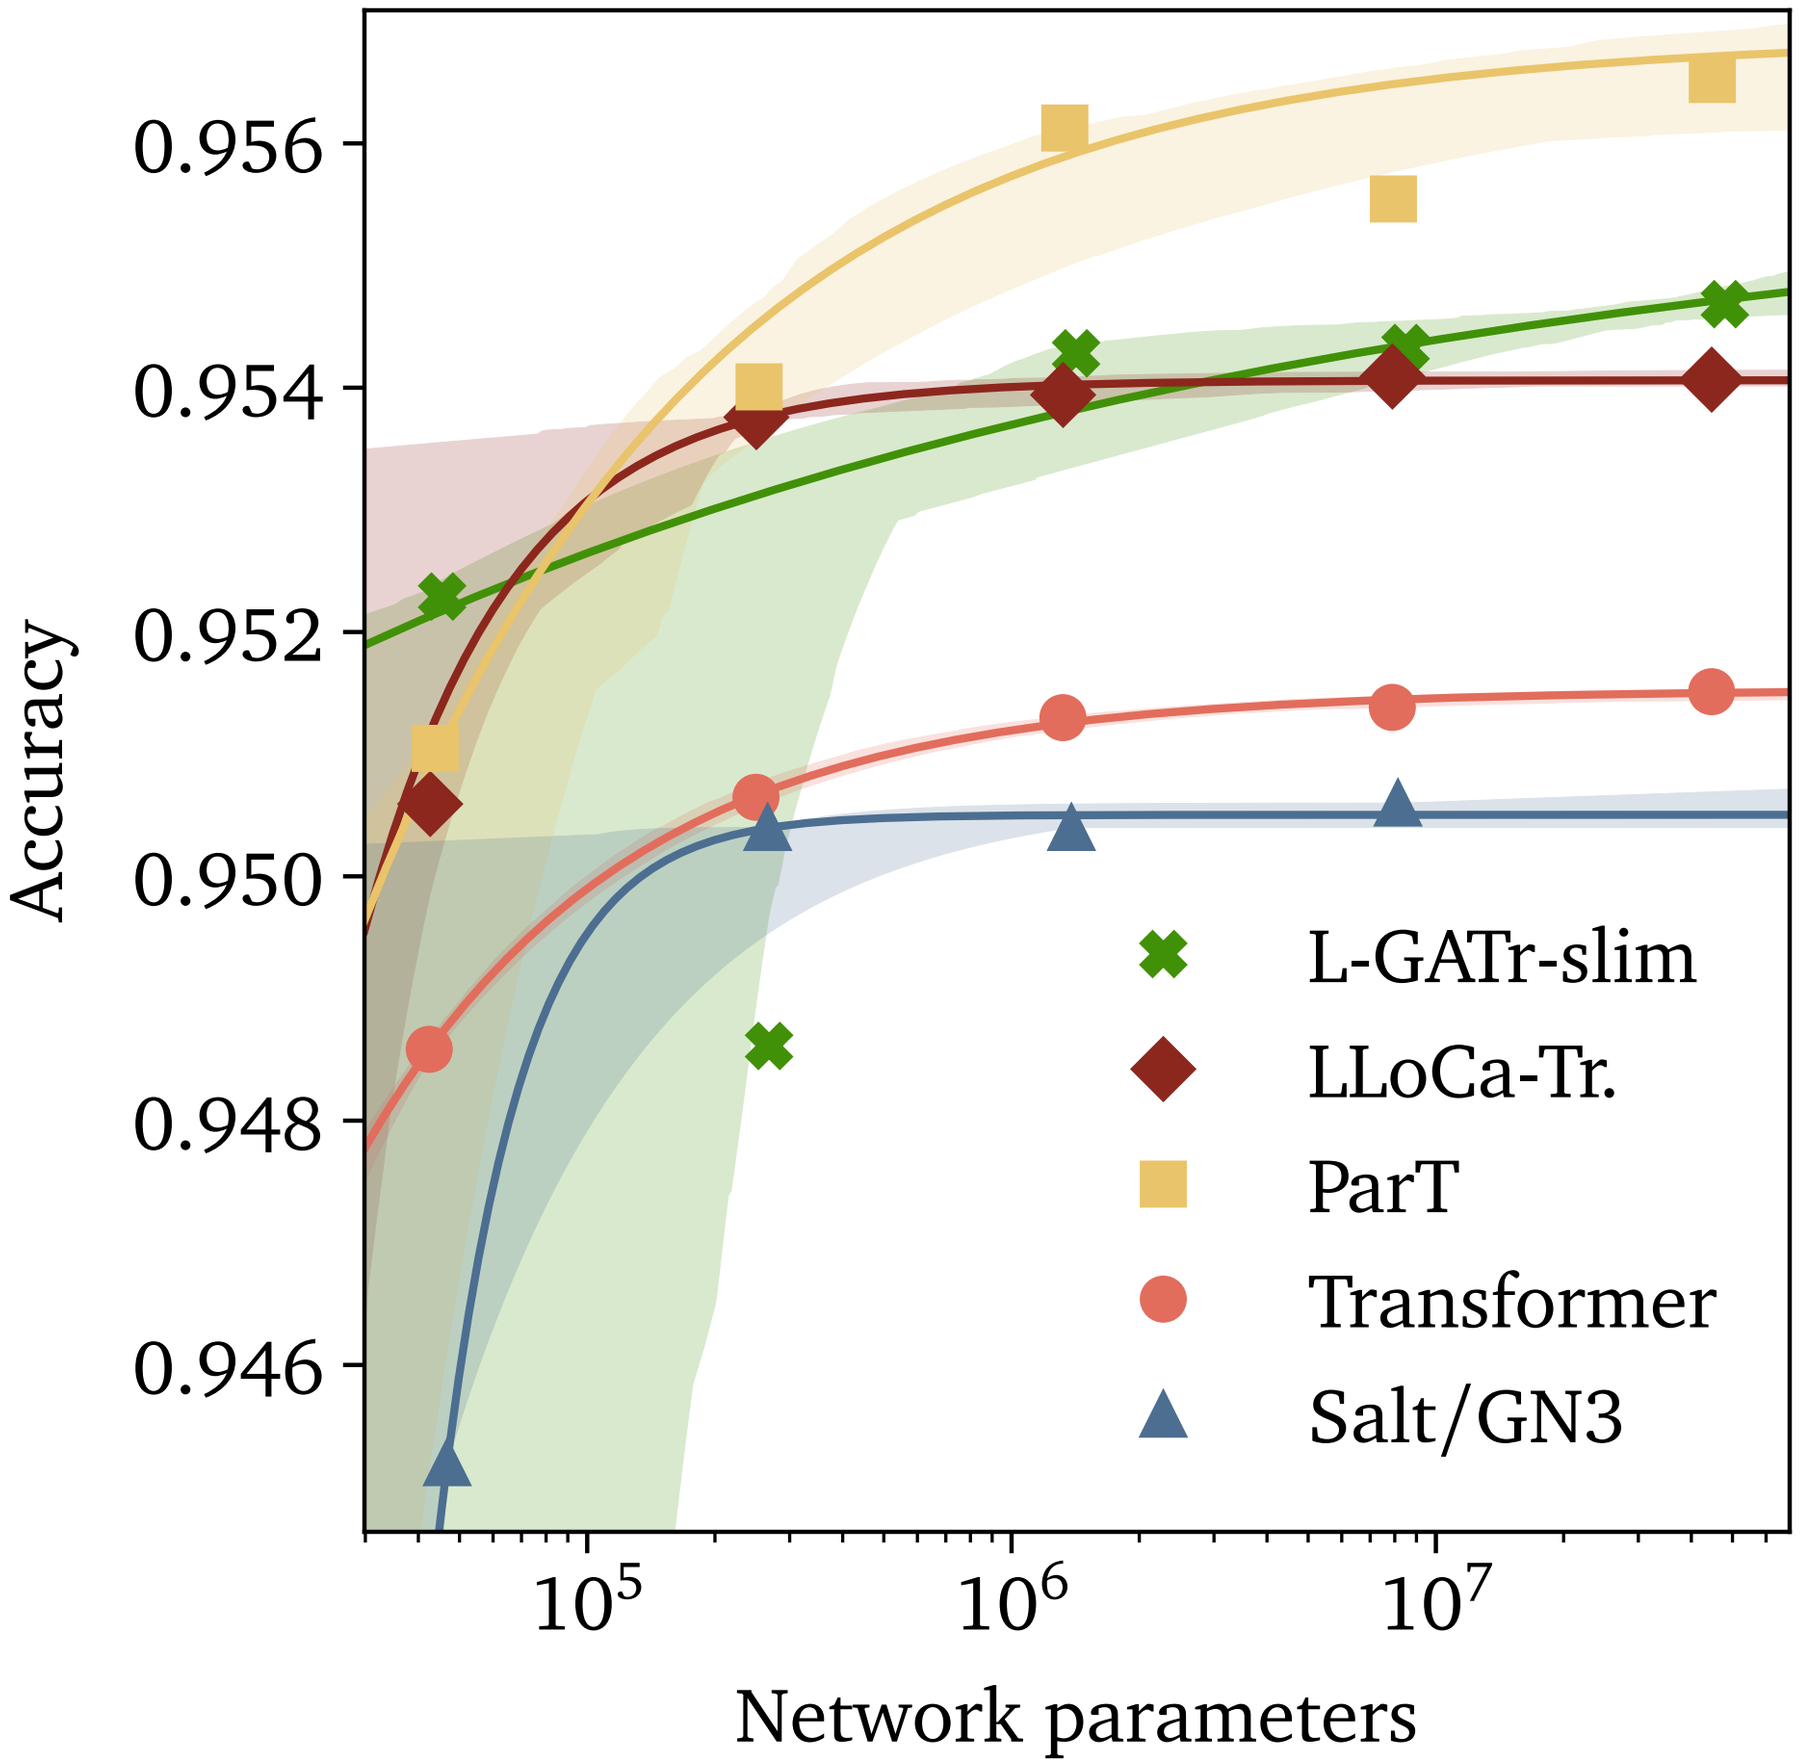
0.956
0.954
0.952
0.950
0.948
0.946
Accuracy
10^5
10^6
10^7
Network parameters
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3

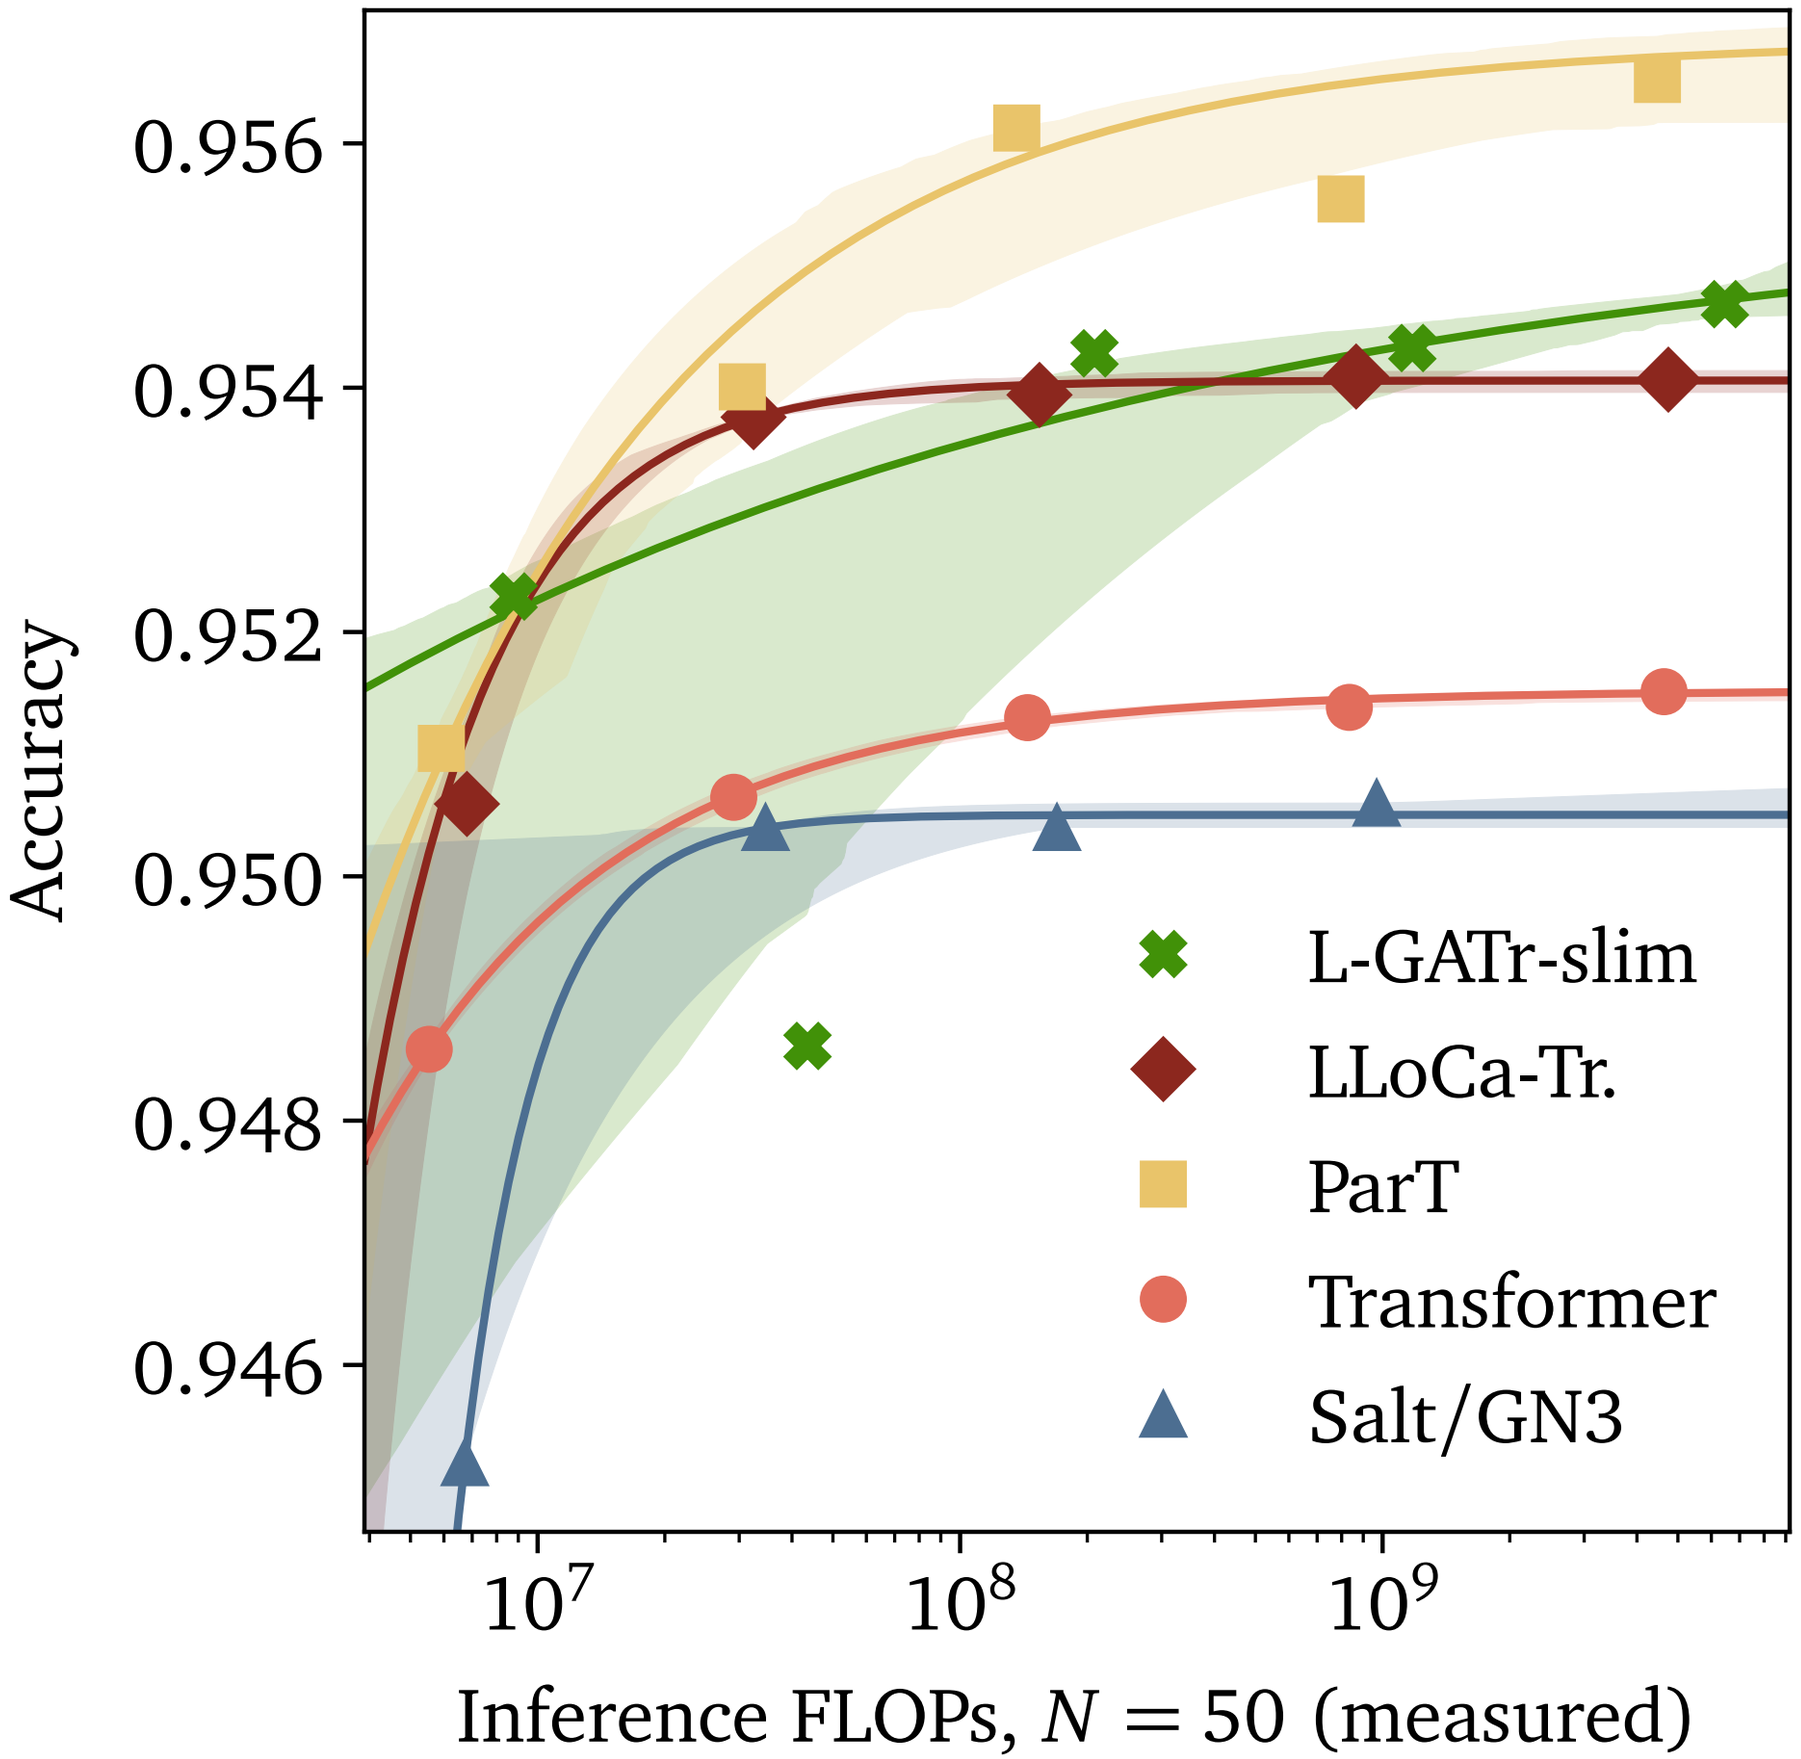
0.956
0.954
0.952
0.950
0.948
0.946
Accuracy
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (measured)
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3

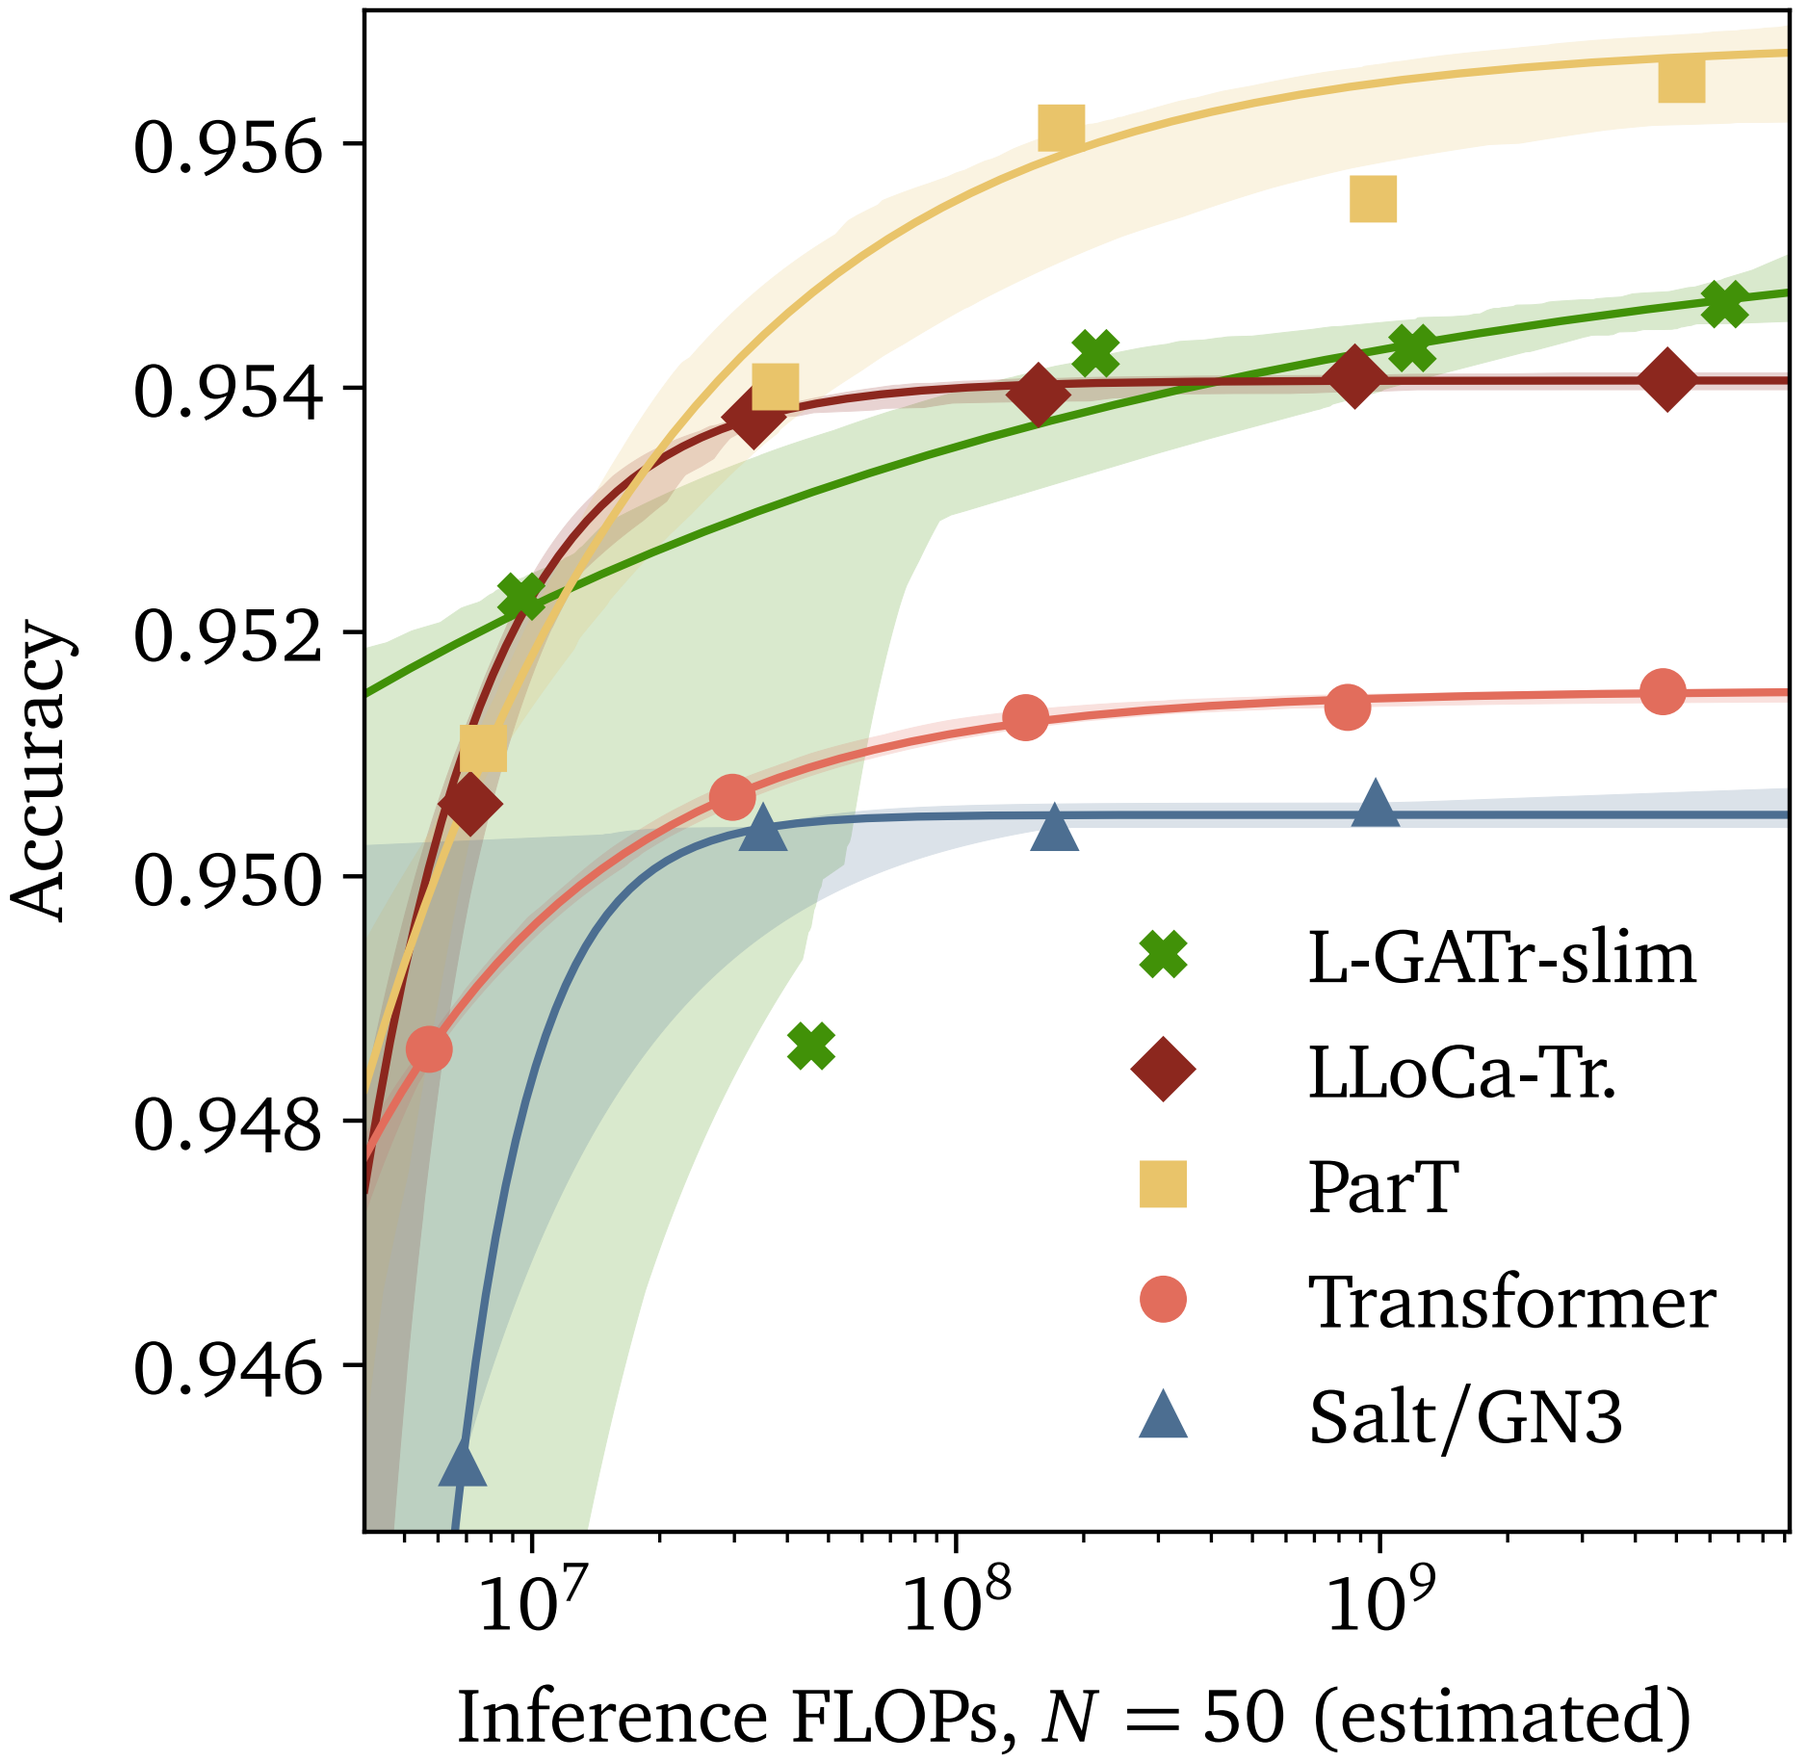
0.956
0.954
0.952
0.950
0.948
0.946
Accuracy
10^7
10^8
10^9
Inference FLOPs, $N = 50$ (estimated)
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3

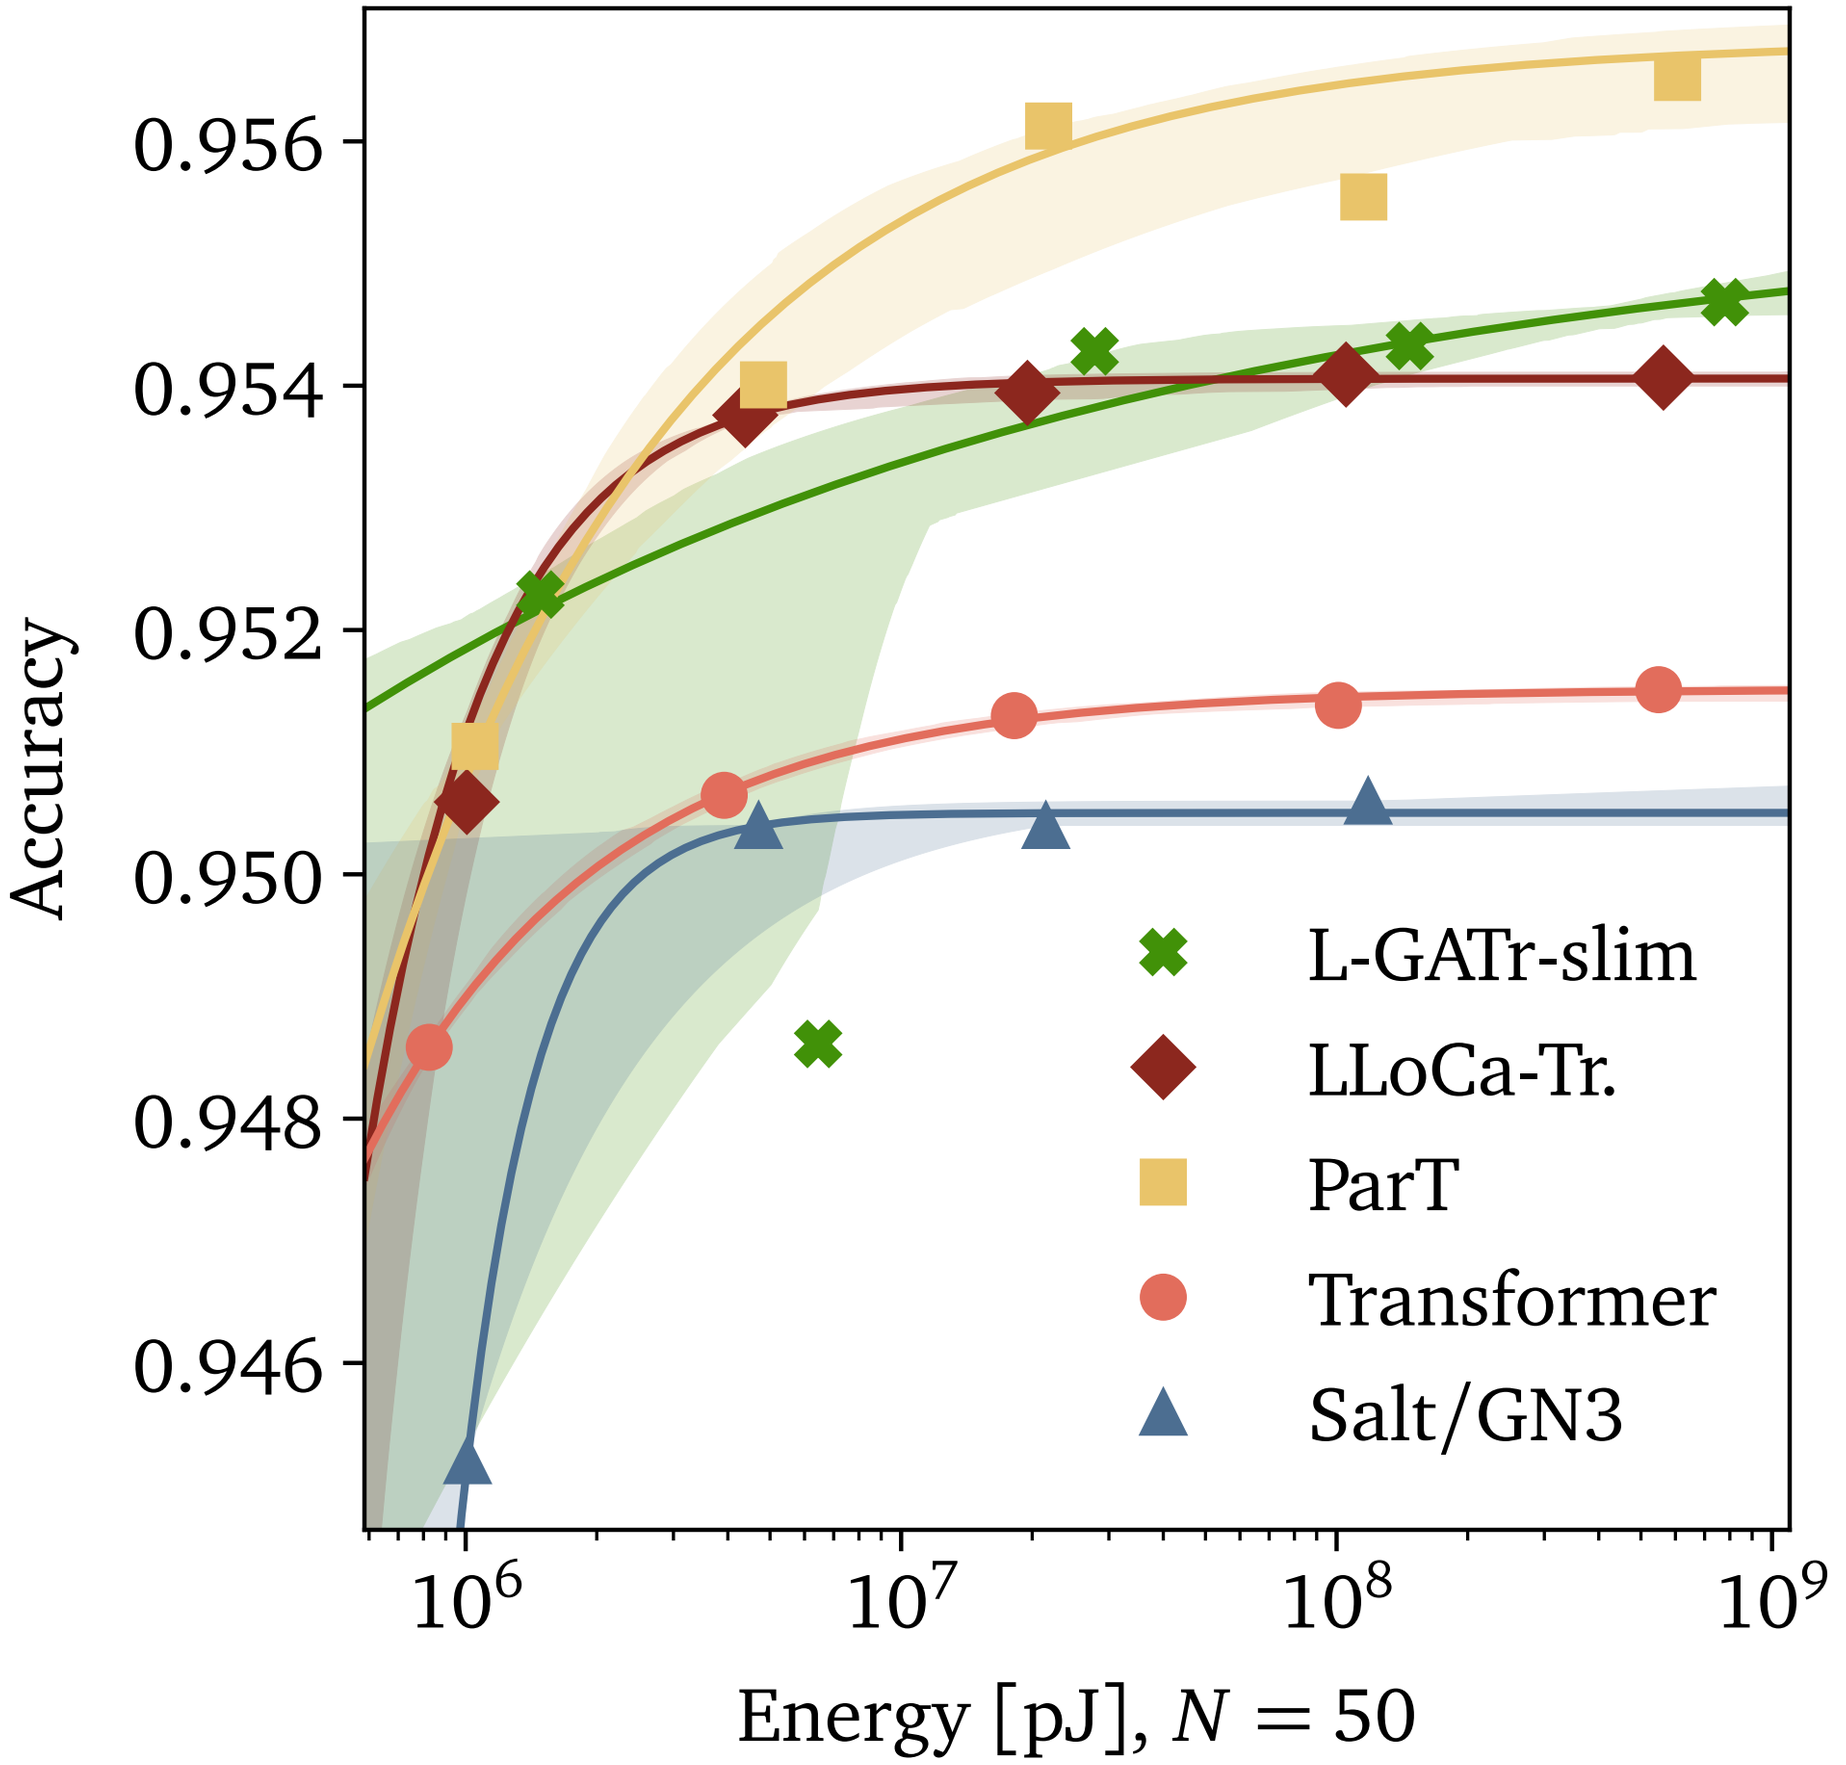
0.956
0.954
0.952
0.950
0.948
0.946
Accuracy
10^6
10^7
10^8
10^9
Energy [pJ], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3

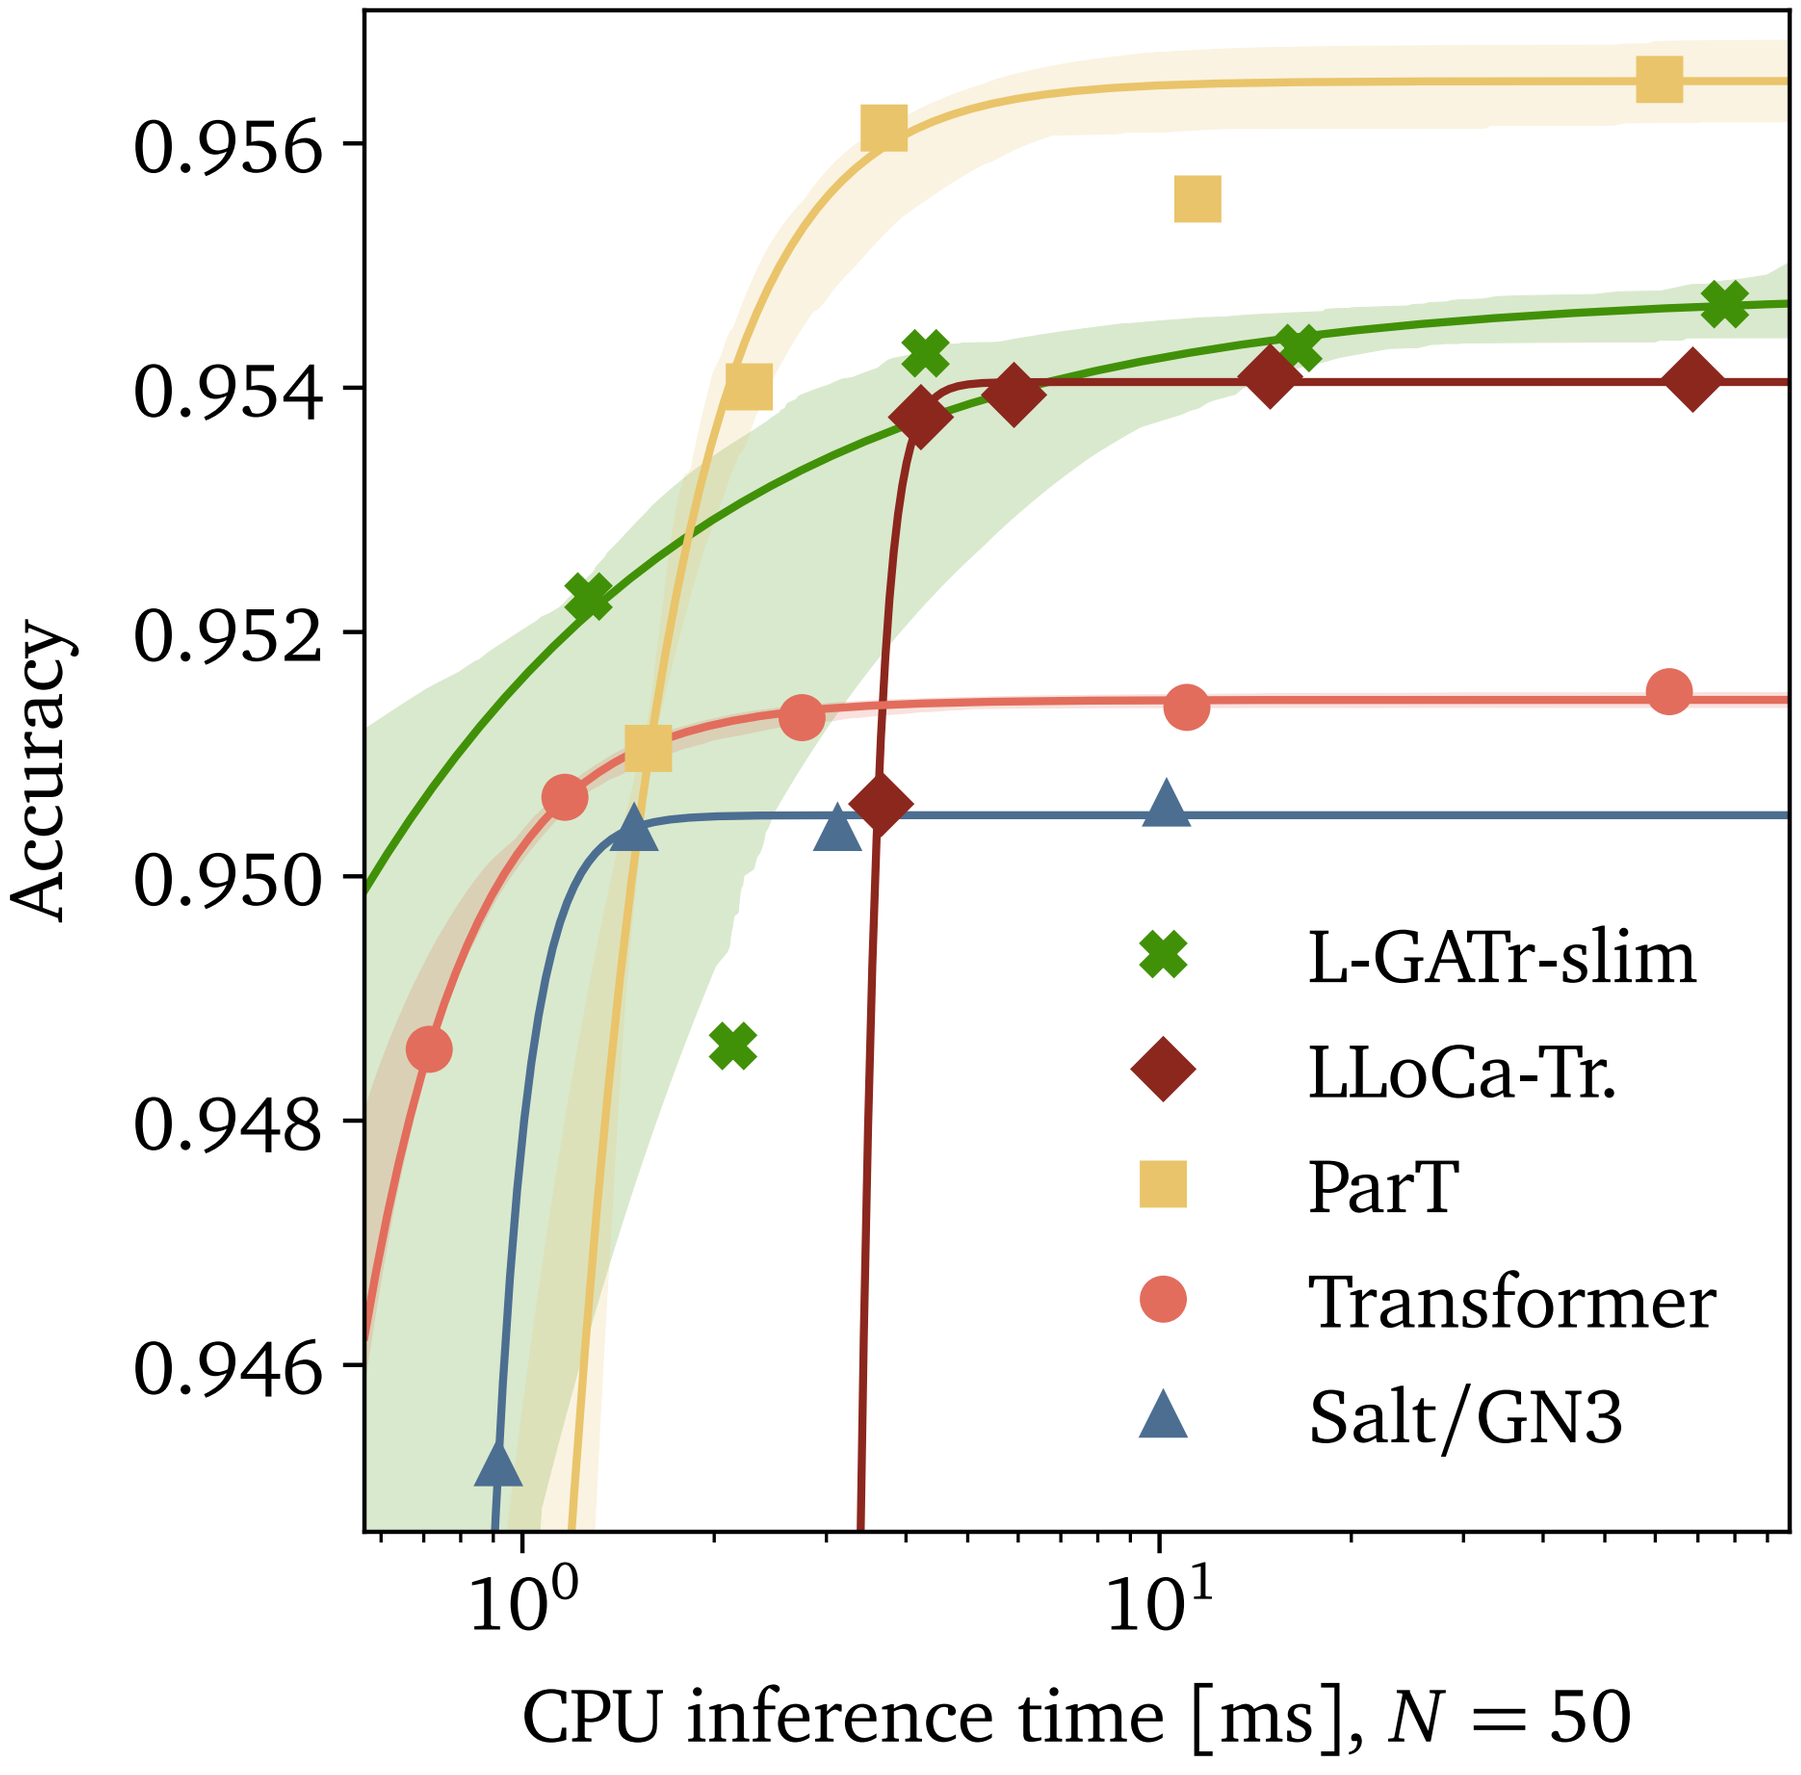
0.956
0.954
0.952
0.950
0.948
0.946
Accuracy
10^0
10^1
CPU inference time [ms], $N = 50$
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3

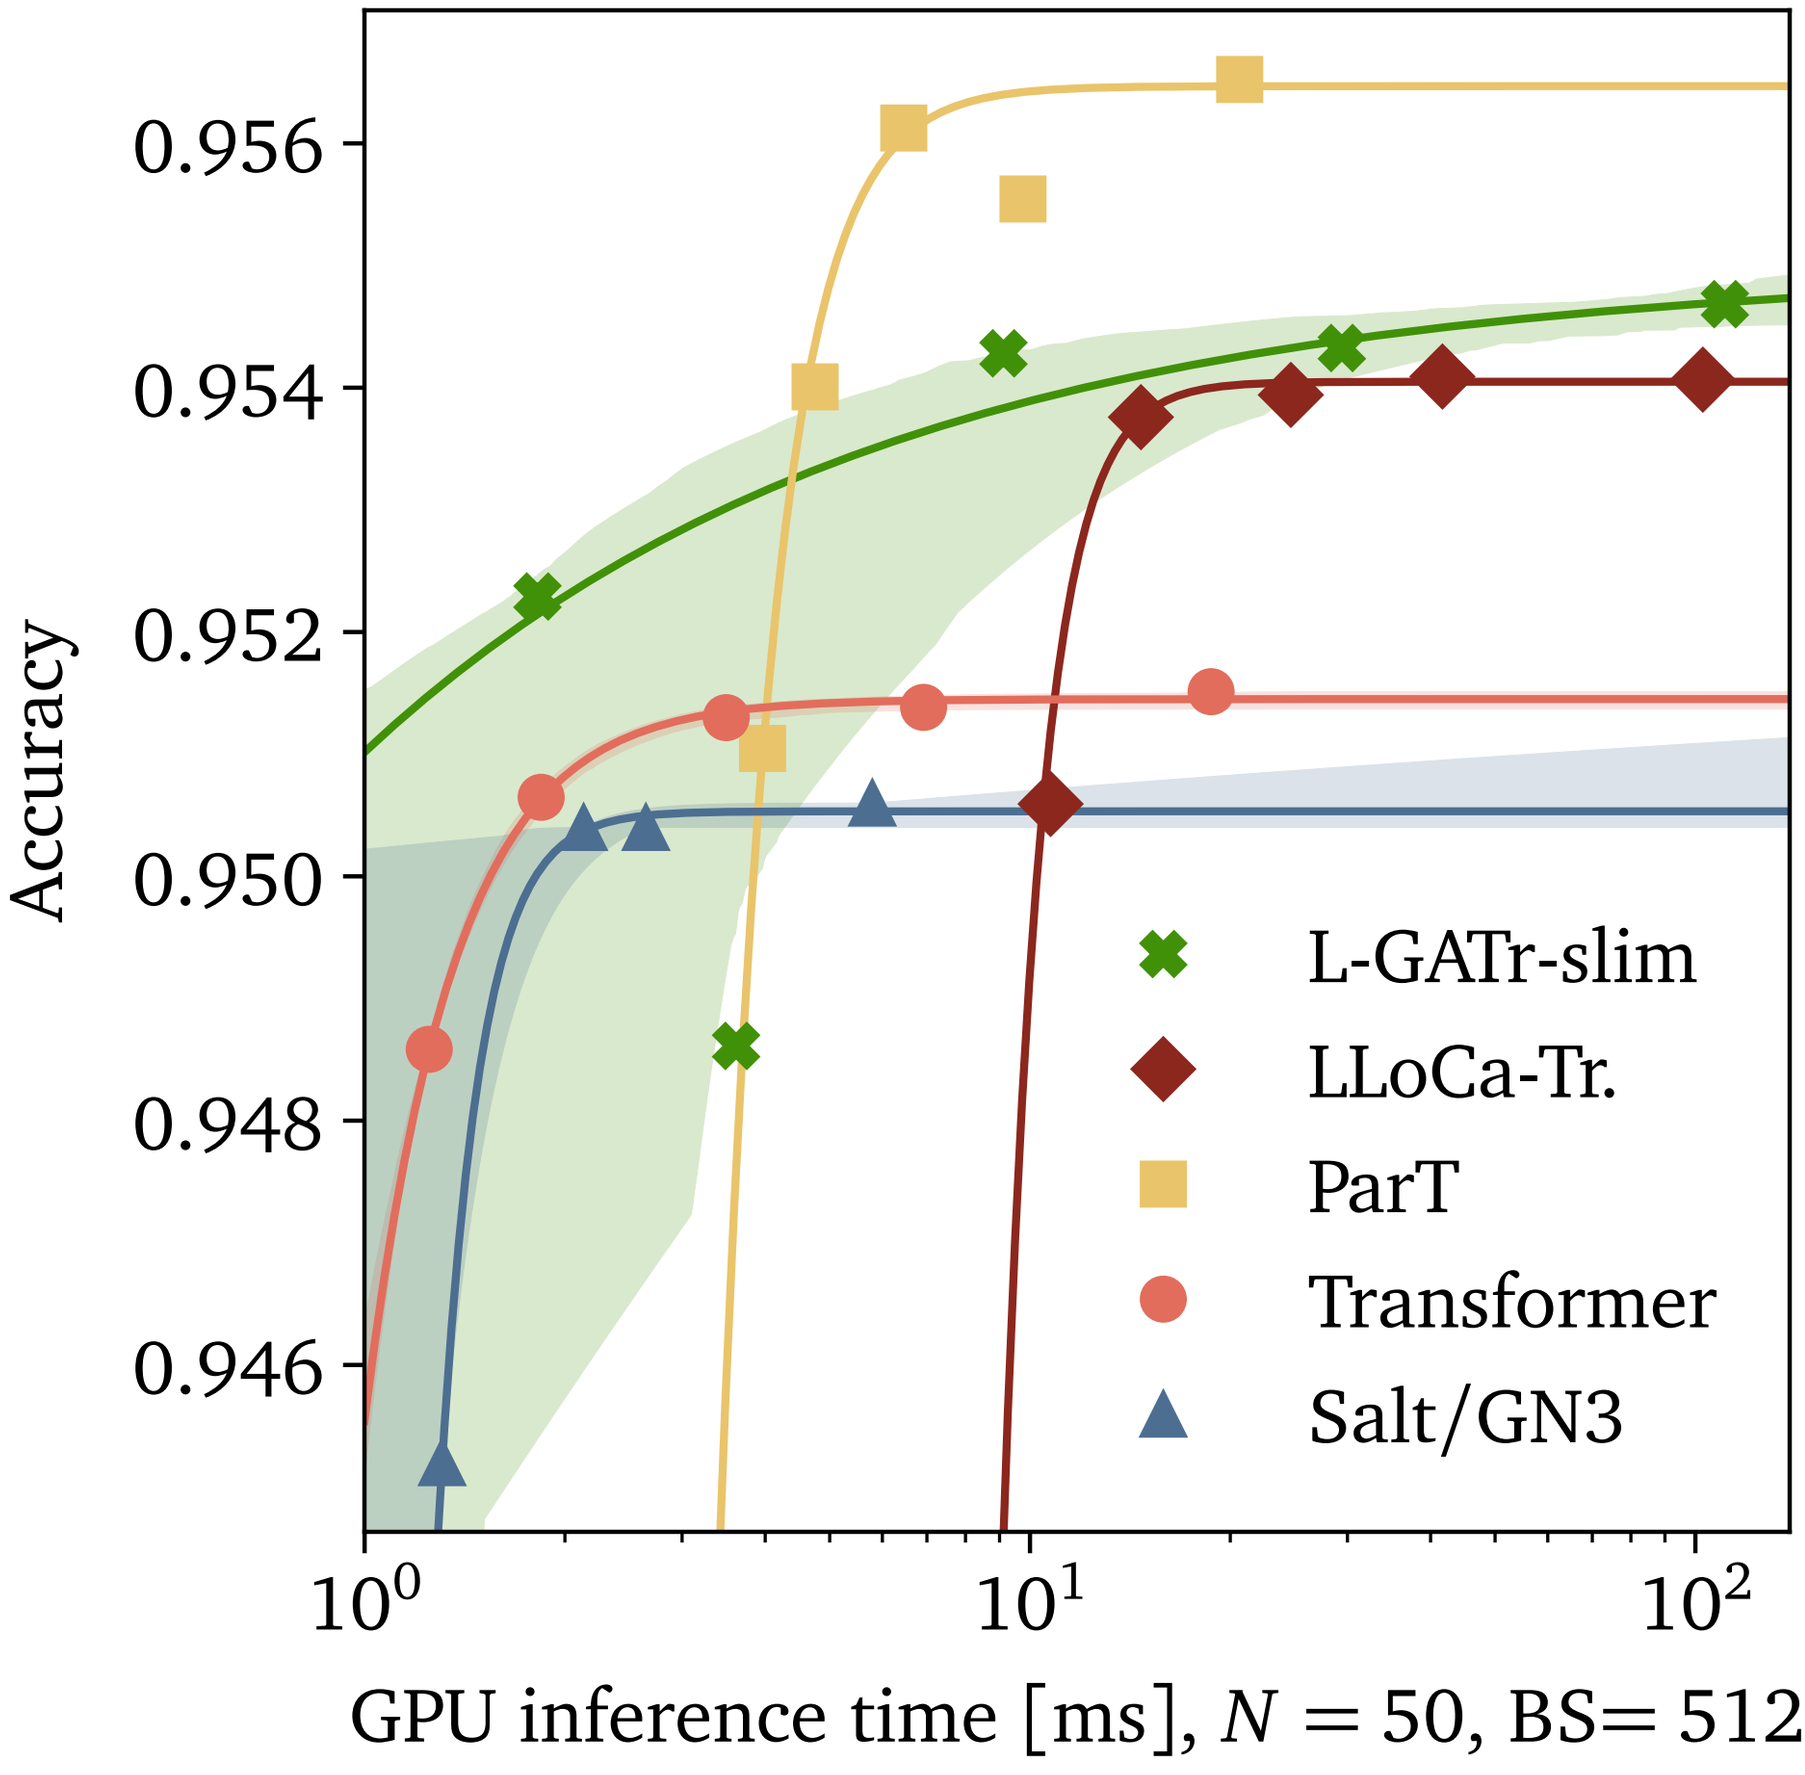
0.956
0.954
0.952
0.950
0.948
0.946
Accuracy
10^0
10^1
10^2
GPU inference time [ms], $N = 50$, BS= 512
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3

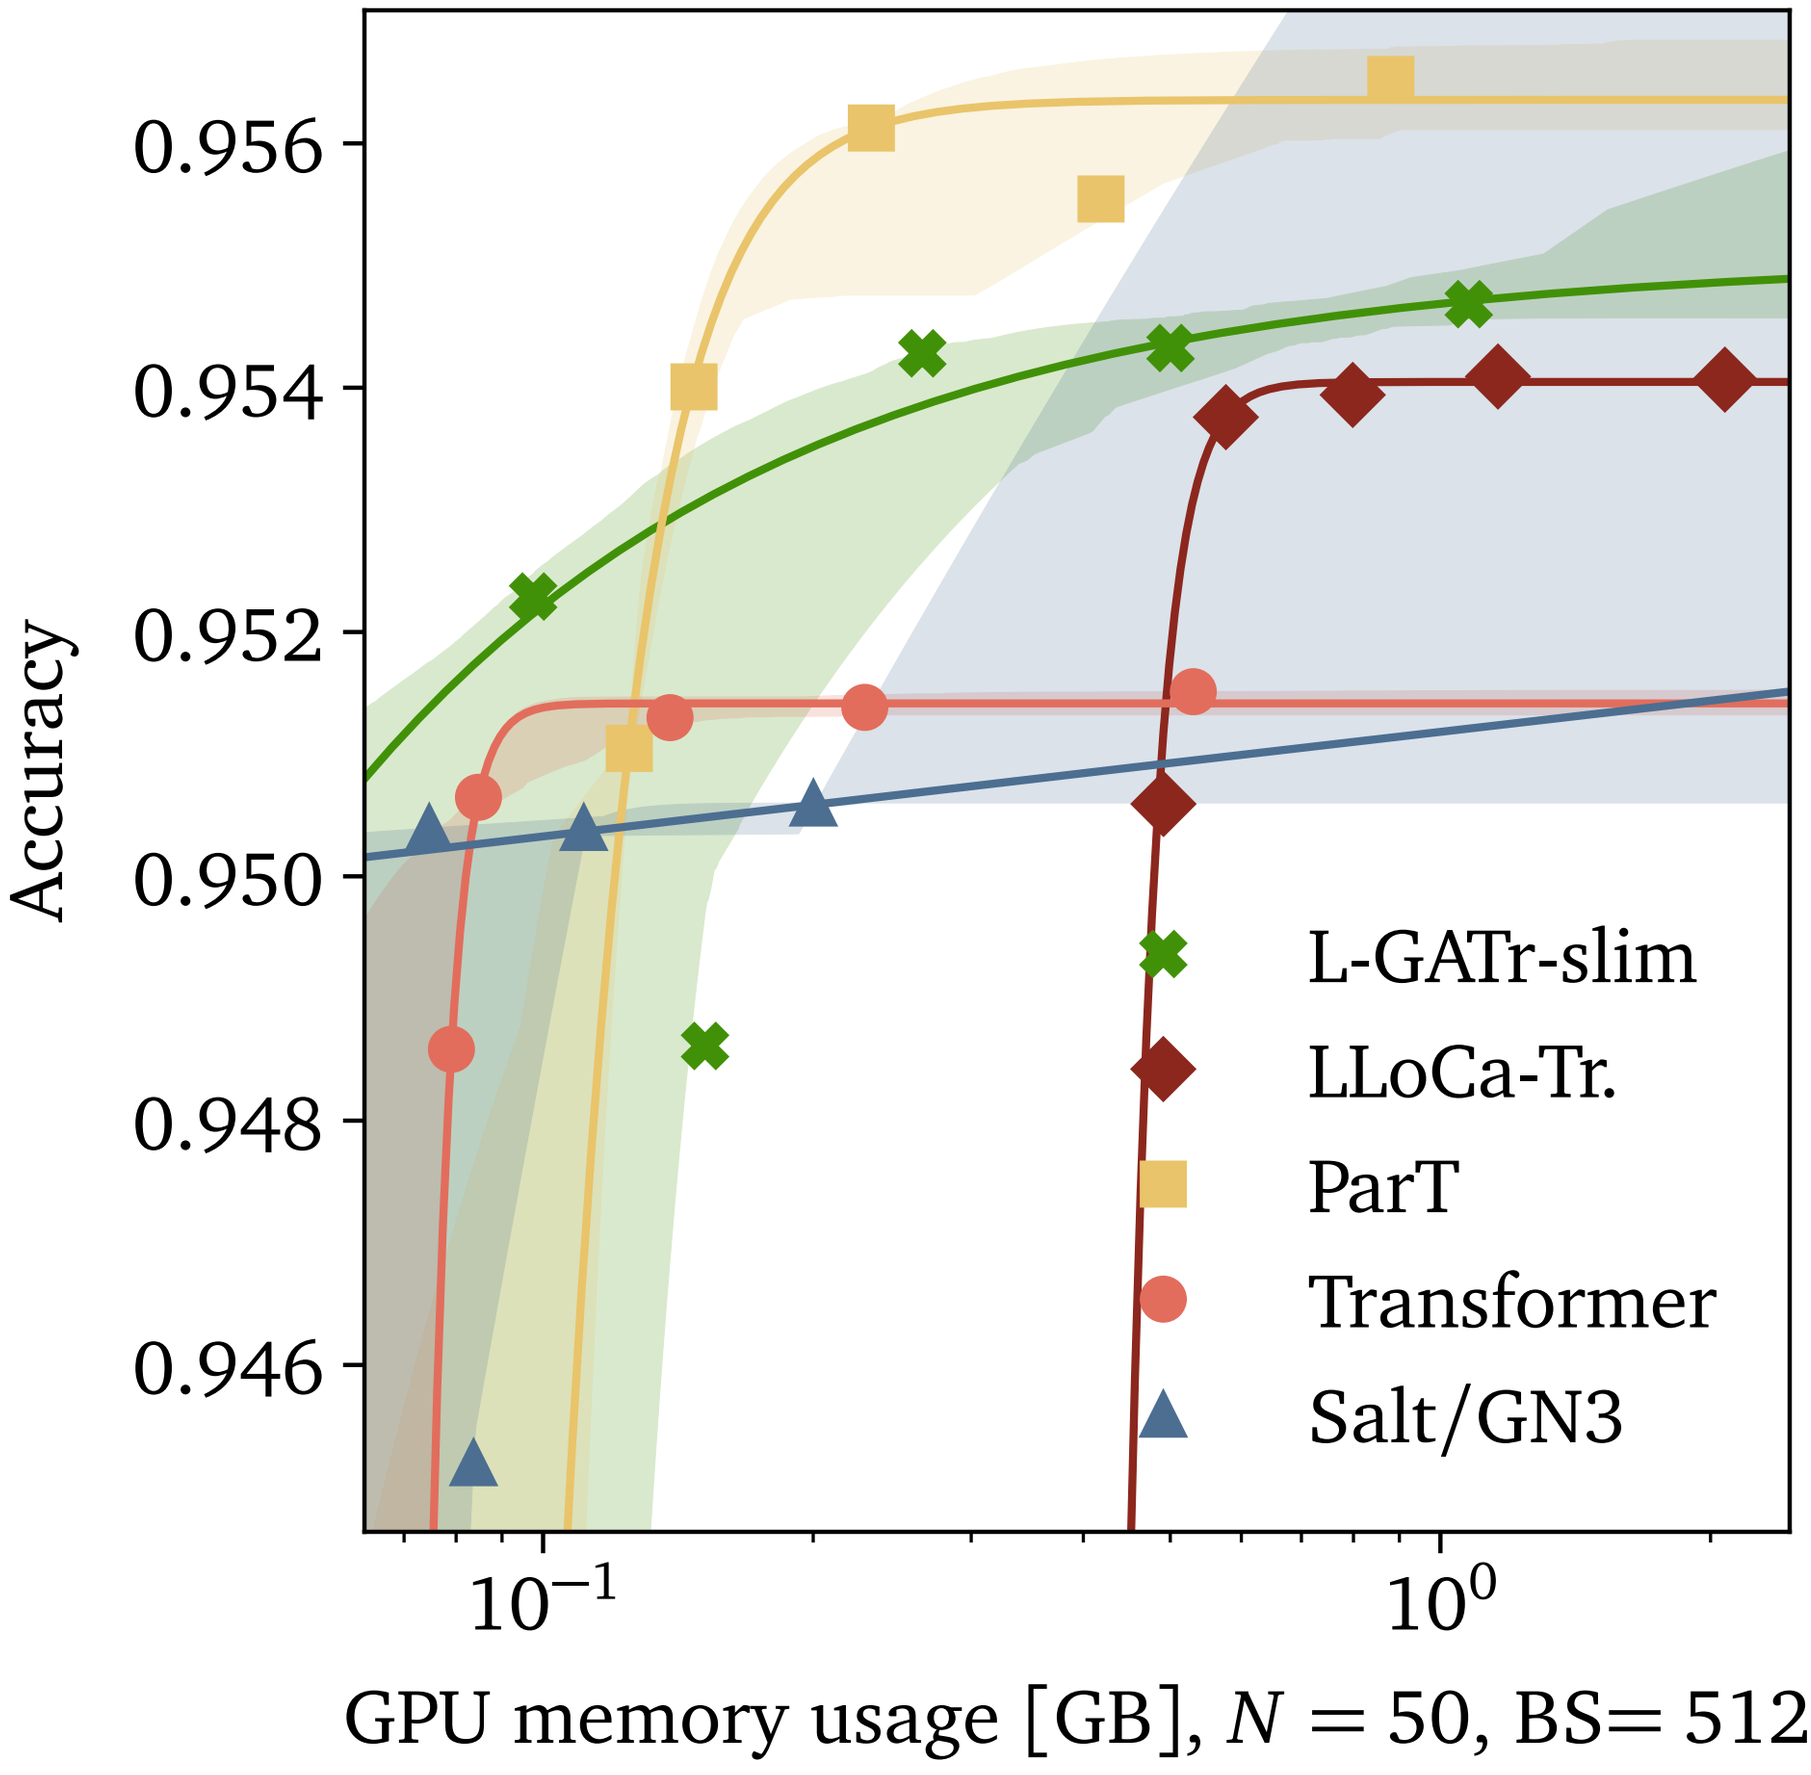
0.956
0.954
0.952
0.950
0.948
0.946
Accuracy
10^{-1}
10^{0}
GPU memory usage [GB], $N = 50$, BS= 512
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3

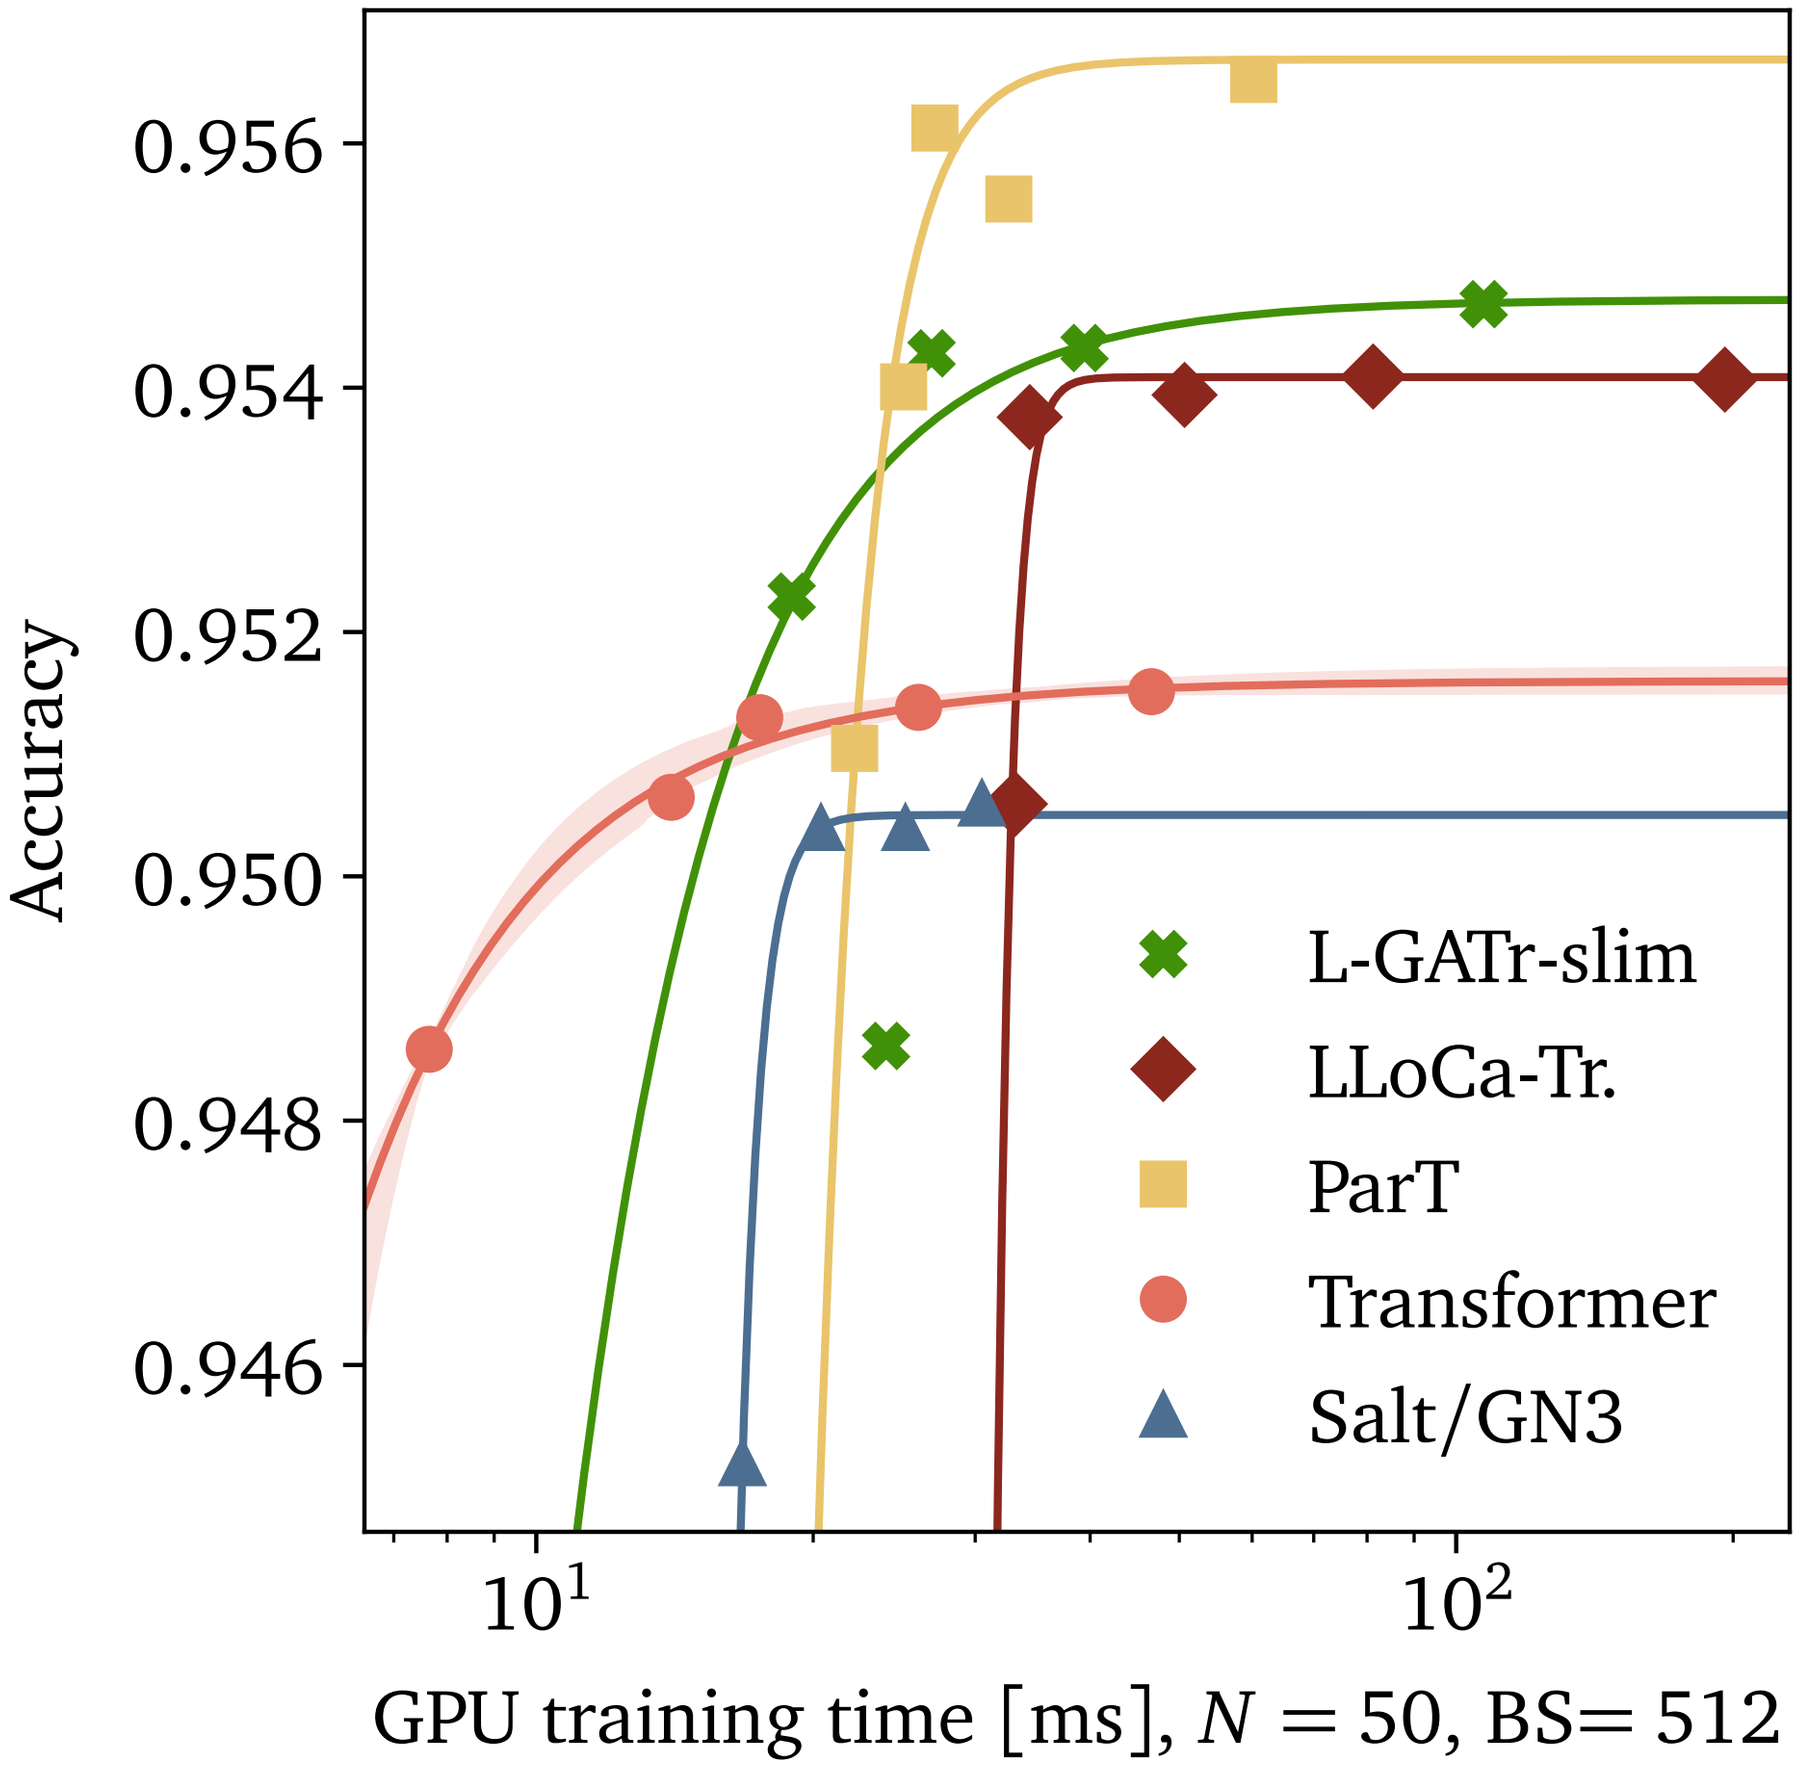
0.956
0.954
0.952
0.950
0.948
0.946
Accuracy
10^1
10^2
GPU training time [ms], $N = 50$, BS= 512
L-GATr-slim
LLoCa-Tr.
ParT
Transformer
Salt/GN3